

Is there a desk here?

Ten agents, four estates, five months of work. The short answer: yes — but not the desk you think you built, and what stands between you and an automated one is plumbing, not research.

GAZBOT V7 · MNQ + MGC · PAPER · DUQ191770 · 45 trading days live · 2026-09-11

THE VERDICT

You are on to something. It is not the thing you have been tending.

Five months produced **one detector that works unattended** (`drift.py`), **one filter that makes it testable** (`WATCH_RT`), and **a measurement estate better than most funded managers run**. None of those three is connected to the thing that trades.

Every component that is connected — the polled market-order exits, the manual entry button, the six-gate tournament — is where the money went. That is a wiring diagram, not a verdict on the idea.

The honest number: on the only uncontaminated sample — single-lot trades in the current architecture — the systematic desk is +\$126.50 over 432 trades, $t = 0.09$. Not profitable. Not losing. Sitting on its own cost line, which is where real systems start rather than where they end.

01 How to read the numbers

The system and its documentation were largely written by the same AI running this audit. Agents were therefore briefed to treat every document — `CLAUDE.md`, `STATE.md`, the memory files — as a claim to verify, never a fact. Where a document and the data disagreed, the data won and the disagreement became a finding.

Every figure carries its provenance:

- **RE-DERIVED** Re-computed by the lead from raw rows, or read directly from the code line. All load-bearing claims got this treatment.
- **AGENT** An agent computed it from primary data and showed its method; not independently reproduced.
- **QUALIFIED** True only under a stated assumption, named inline.

Three agent claims did not survive checking and are corrected in §07 rather than quietly dropped — including a false alarm about a live naked position that does not exist.

The central finding

Your exits beat the machine. Your entries are worse than a coin flip. And the skill in your hands is not what anyone assumed it was.

WHAT THE OPERATOR'S HAND CONTRIBUTES — DECOMPOSED

COMPONENT	FINDING	EFFECT	ENCODABLE?
Exit timing	Real skill; two independent samples agree	+\$65.70/lot	Yes — it is a stop
Trade selection	41% went his way first; −\$13,235 held to clock	−\$68.94/pos	Already encoded, and better
Sizing	40 of 41 positions exactly 4 lots	none	Nothing to encode
Regime presence	17 of 41 entries outside validated US hours	negative	Yes — session filter

The exit edge is $n=80$, $t=2.28$, 95% CI [$+\$10$, $+\$123$] **AGENT**. Real, but two months short of the sample it needs.

THE MECHANISM NOBODY HAD NAMED

The edge is **entirely loss-cutting**. On the 21 positions the machine would have lost, the hand saved $+\$10,718$. On the 22 it would have won, the hand gave up $-\$2,663$ by exiting early.

You are not top-ticking winners. **You are personally acting as the stop-loss on a desk that, by your own instruction, has no stop.** That is the most encodable object in trading — and it is why "claiming cannot be backtested" was true but aimed one level too low. The 119 calibrations searched for a rule to replace your hands on your trades. The finding is that your trades are the problem.

Confirming from the other direction — rider positions, flagged rows excluded **RE-DERIVED**:

1st trade of the day	$n=23$	$+\$2,505.01$	78.3% win
2nd and beyond	$n=30$	$-\$5,069.94$	50.0% win
sequence 2	$n=10$	$-\$2,261.94$	
sequence 3	$n=7$	$-\$2,851.50$	
sequence 4	$n=6$	$-\$302.00$	

A one-trade-per-day rule needs no judgement, no watching and no edge discovery. It is three lines of code, and on this record it flips the rider from negative to positive.

The detector, meanwhile, **has not opened a position since 2026-08-26** AGENT . For eleven sessions every entry was a button press — trading over the top of your own strategy, while the strategy was doing better.

03 Two thirds of the loss never happened

The IBKR paper engine fabricates fills: every lot beyond the first in a multi-lot order executes at exactly 0.1% of price, adverse. An agent solved the generating rule — worst = `tick_round(first × (1 ± 0.001))`, 44 of 46 exact, and **91 of 91 deviations adverse, not one favourable**.

LIVE DESK P&L, DECONTAMINATED – 879 TRADES, 45 DAYS

BASIS	P&L	WHAT IT MEANS
As booked	−\$5,477	The headline you have lived with
Fabricated lots reversed	−\$1,919	62 lots — 2.8% of all lots ever filled
Clean subset only	−\$2,030	812 of 879 trades
At 1–3 lots, repriced	+\$6,200	853 trades
At 4 lots	−\$7,742	25 trades; five carry −\$7,394
Single-lot, current architecture	+\$126.50	t = 0.09. The cleanest read available

AND LIVE FILLS WOULD BE BETTER, NOT WORSE

An agent walked your own captured L2 ladder — 18,692 snapshots. A real 4-lot MNQ market order costs about \$0.43. The paper engine charged \$58–\$177 for the same order: **130x to 400x too much**.

Execution cost was never the problem. What remains is directional and concentrated in twenty-five four-lot trades — a sizing problem, which is fixable.

04 What it earns unattended

This is the figure the decision rests on, since you will not be watching. There are only six zero-intervention rider days, so that sample is noise and is not quoted. The defensible figures:

UNATTENDED CONFIGURATIONS, RANKED BY EVIDENTIAL STRENGTH

CONFIGURATION	RESULT	T	VERDICT
Tournament automatic exits n = 742 trades	-\$9.07/trade	-3.70	Established negative. 11 of 13 gate families lose
Drift desk, unfiltered n = 42 sessions	+\$77/session	0.45	Would need ~6.1 years to prove
Drift + rt < 0.50 n = 15 sessions	+\$452/session	3.11	The only testable lead on the desk

That last row is WATCH_RT, at day_rider.py:256 RE-DERIVED:

```
WATCH_RT = 0.50

def watch_verdict(rt: float) -> str:
    """One line for the operator at confirmation: is today worth
    watching by hand?

    Reporting only. Callers must NOT branch trading behaviour on this."""
```

It is referenced **nowhere else in the codebase**. The one filter that survived the full anti-overfit battery over 231 sessions prints a sentence at you and does nothing. It needs n=12 for power and has 15. Every other lead on this desk needs years.

QUALIFIED

QUALIFIED Those 42 sessions probably overlap the study the filter came from, so this is a replication, not out-of-sample evidence. The confirmed band sits at rt 0.450–0.485 — "barely confirmed", counter-intuitive enough to smell of overfitting. Pre-register it — threshold, size, stop rule and stop date fixed in advance — and it becomes a zero-cost experiment with a defined pass/fail.

Why the exits keep failing

The trail was switched off on 2026-09-11 on the reasoning that a polled level is not a stop. That reasoning was right and **aimed at the wrong component**. Reading the ladder, `day_rider.py`: 1509 **RE-DERIVED**:

```
for _k, _tpt in enumerate(_ladder):
    if _k in done or lots_open <= 0 or ahead < _tpt: continue
    _tr = ib.placeOrder(contract, MarketOrder("SELL" if d > 0 else "BUY", 1))
```

The take-profit ladder is **not a resting limit**. It is a once-a-minute poll that fires a market order — and so was the trail. One was removed and the other left, so the desk now has less automatic protection than before. The proof is in the ledger **RE-DERIVED**:

```
id 868    LONG 1    29320.25 -> 29240.75    -$160.50    TARGET_100
```

The +50pt rung sat at 29370.25, and the tape printed 29396.00 forty seconds earlier. That take-profit filled **129.5 points below its own level** and booked a loss.

Faster polling does not fix it — 60s to 5s measures as non-monotonic noise. What dominates is sampling phase: the same rule on the same tape swings \$651 depending on which second of the minute the poll lands.

THE FIX IS ALREADY IN YOUR CODEBASE

LimitOrder at :1007. StopOrder at :1802, gated off by `PLACE_VENUE_STOP = False`. The tournament already rests stops at the broker — 312 stp-* executions in the fills table.

Four OCA-grouped limits at 50/100/200/300pt plus one protective stop, **resting at IBKR**. Polling disappears, sampling phase disappears, and the fabricated multi-lot exit fill disappears. One change closes three findings.

What is genuinely good

Not consolation. Several of these were called better than funded practice, independently and unprompted, by agents that did not know each other's findings.

THE MEASUREMENT ESTATE

Four independent auditors reconstructed five weeks of trading from venue fills, detected a broker-side artifact, and replayed the strategy causally — on someone else's codebase, in an afternoon. Almost no one-person desk can be audited that way. Ledger arithmetic is perfect: 879 of 879 rows internally consistent, zero exceptions.

THE TICK REPRICER

Re-run from scratch on a random 250-trade sample: **249 of 250 exact to the cent**. Defensive against the IBKR sentinel, defers rather than freezes a partial window, replays each chandelier variant on its own trail.

INTELLECTUAL HONESTY, STRUCTURALLY ENFORCED

\$1,551 of fictitious profit kept in the table labelled EXCLUDE: rather than deleted. The rows the desk discards are net positive — it is biased against itself. Controls protected by name with written reasons. A pre-registration file states "an unfavourable result is the most valuable outcome this can produce". One agent: "I have seen funds that do not do this."

KILLING YOUR OWN HEADLINES

The tunnel service's founding 1.75× claim was withdrawn by the people who shipped it. An agent re-derived it independently at 1.02× **and found the arithmetic that manufactured the error**: ATR at a break is only 0.607× an active minute's ATR, so the denominator alone produces $1/0.607 = 1.65\times$. The withdrawal was correct.

THE DATA ESTATE – YOUR HYPOTHESIS WAS RIGHT, FOR A REASON YOU HAD NOT STATED

AGENT Priced against vendor list rates, **the raw data is worth about \$250**: 40 days of MBP-10 depth for MNQ + MGC costs \$2–\$55, ten years of 1-minute bars across NQ/MNQ/GC/MGC is a \$199 subscription, and real-time L2 to keep capturing is \$15/month. The bytes are a commodity.

The pipeline that makes those bytes trustworthy would cost \$150k–\$400k to rebuild. **That ratio is roughly 1,000 : 1, and it is the evidence for your belief that the R&D facility is the real asset**. The durably valuable core is narrower than the whole repository, though — roughly the 6,000-line capture-and-archive layer, not the 107k lines of which much is single-use study code.

One hard constraint worth knowing before you ever plan to sell or license it: CME's Historical Distribution Licence is \$35,220/year per exchange — \$70,440 for CME + COMEX — and its agreement bans redistribution "in any format, to any other party" including derived data products. You are fully within your rights capturing at \$12.10/month under a display-device licence, but **this estate can never be sold or licensed onward**. Its worth is captive to this operation.

FALSE ALARM – THERE IS NO NAKED SHORT

One agent reported an open 38-lot naked position sitting for five days. It reconstructed that from the fills table — precisely what the desk's cardinal rule forbids. The venue reads $\text{venue} +0 = \text{tournament} +0 + \text{rider} +0$, verified twice by the independent read-only reconciler. The -38 is the missing 09-07 cover. **Bookkeeping, not risk.**

INFLATED HEADLINE

An agent put the first-trade-of-day bucket at $+\$3,974.51$. Direct query gives $+\$2,505.01$ — its filter admitted phantom rows. The effect replicates strongly; the number did not.

WEAKENED ON CHALLENGE

The V5 archive was reported as independently reproducing the negative result at $-\$2.53/\text{trade}$, $t=-2.26$, $n=3,008$. On challenge the agent corrected itself: V5 booked zero commission, so that applies an assumed fee. Gross it is $-\$0.93/\text{trade}$, $t=-0.83$. **Two negative samples — not two significant ones.** The direction replicates; the significance rests on a cost constant.

1 · THE BOX IS OPEN TO THE INTERNET

RE-DERIVED Root SSH with password authentication, no firewall, and the dashboard listening on all interfaces returning HTTP 200 from the public IP.

```
permitrootlogin yes      passwordauthentication yes
ufw: inactive            iptables -P INPUT ACCEPT
LISTEN 0.0.0.0:8087      <- the dashboard
```

That app has no authentication — `web.py:1152` says so itself — and it serves `/api/control/dayrider-buy`, which writes the file the inotify unit turns into a broker order in **1.35 seconds**. nginx has no `auth_basic`, so it is reachable via the domain too. Every safety control on this desk stops the desk misbehaving; none asks who sent the request.

2 · The fail-safe is inverted. **AGENT** `gate-reactivate` arms every gate at 22:00Z regardless of whether the router is alive. On 2026-09-10 the router went blind at 16:30Z with all six gates off; at 22:00Z five were armed and could not be benched for **10h50m**. The stated failure mode is "benched is the cheap error". The actual one is "armed and unbenchable, overnight".

3 · The shadow book has a 55-second look-ahead. **AGENT** Entry price is stamped at bar close while the decision is made up to 55s later. It flatters every delayed-confirm arm and defames every tick-triggered one. Corrected, `abs_veto_55s` falls from +\$9.60 to +\$1.17/trade and **nothing survives multiple-comparison correction**. The same defect has been wrongly condemning `capitulation_long` at −\$15.63/trade when the honest figure is −\$0.57.

4 · The census measures 3-hour runs and calls them 15-minute runs. **RE-DERIVED** `run_census.py:36` sets `W = 180` as a bar count commented "15-min window in 5s bars". With `--take` the selector picks 1-minute bars, so `W` becomes 180 minutes — while the participation window stays hardcoded at `start + W * 5`, i.e. 15 minutes. A gate that traded minute 40 of a run scores as having sat out. **That is the accounting behind the greenfield programme's founding "we sat out \$8,025"**.

5 · The books do not tie. **RE-DERIVED** The fills table implies a net position of −38 against a venue of 0, because zero fills were recorded on 09-07 when the runaway short was covered. 89% of booked trades have no `exec_id` linkage. The largest loss in the desk's history — roughly −\$8,400 — appears in no table at all.

6 · The captured order book contains a twenty-hour lie. **RE-DERIVED** In `depth.db`, the MNQ best ask sat frozen at 30151.75 across **200,708 consecutive snapshots spanning 19.6 hours**. Any study reading the book through that window gets garbage, silently. An agent root-caused it to a weekend-restart failure — **so it will recur**. Three more data defects sit alongside it: a driftlab study anchored to the wrong DST offset on 47.5% of its sessions that goes live on 2026-11-01, an 89.4% orphan-fill rate specific to the current gate slate, and no `live↔shadow` pairing at all — which is why the 55-second look-ahead above went eight weeks undetected.

7 • The router costs more than it can be shown to add. AGENT Of 4,609 gate fires it benched, **80% would have failed the gate's own ATR floor anyway** — a slower, costlier copy of a one-line threshold. Measured benefit \$35–350/month with a confidence interval including zero, against \$352/month of inference. Its reasoning quality is genuinely high (98.4% numerical fidelity, 0.13% rulebook violations); its substrate is a consumer OAuth subscription with a combined failure MTBF of 4.8 days and a repair path requiring a human at a keyboard.

Everything the audit produced, in the order it should be done. Nothing here is a research programme — **eighteen of the twenty-three items are plumbing.**

#	ACTION	EFFORT
TIER 0 – TODAY, REGARDLESS OF THE GO-LIVE DECISION		
A1	Bind the dashboard to 127.0.0.1 Closes the public unauthenticated trade endpoint. nginx keeps serving dashboard.gazbot.dev — zero impact on your access.	5 min
A2	Install Tailscale on the VPS and your devices Prerequisite for A3 — makes locking down SSH safe rather than a lockout risk, and fixes the dead-desktop access problem at the same time.	15 min
A3	Key-only SSH, disable root password login, enable ufw Do only after A2 is confirmed working from a second device.	20 min
A4	Rotate every credential exposed by the June breach Logged as "in progress" three months ago and never confirmed. The box holds your keys.	30 min
A5	Put authentication in front of /v7/ Basic auth or Cloudflare Access. Mandatory before any funded account.	20 min
TIER 1 – BLOCKING FOR ANY LIVE MONEY		
B1	Move every exit to a resting broker order Four OCA'd limits at 50/100/200/300pt plus one protective stop. Both order classes already exist in the code. Closes the polling defect, the sampling-phase lottery and the fabricated multi-lot exit fill at once.	half day
B2	Re-arm a protective stop A 4-lot naked position with a measured 444-point adverse excursion is the only thing here capable of ending the account.	included in B1
B3	Cap the rider at 1-2 lots +\$6,200 across 853 trades at 1-3 lots; −\$7,742 across 25 at 4 lots. Size is the lever.	1 line
B4	One trade per day — ban same-day re-entry 1st of day +\$2,505 (n=23, 78% win); 2nd+ −\$5,070 (n=30, 50%).	3 lines
B5	Fix the inverted fail-safe in reactivate_gates.py Must not arm gates when the router has been silent > 30 min.	4 lines
B6	Clear the failed eod-flatten unit and page on its failure The component STATE.md calls the real floor has been red since Wednesday with nothing saying so.	15 min
TIER 2 – THE INSTRUMENTS ARE MIS-AIMED UNTIL THESE LAND		

#	ACTION	EFFORT
C1	Fix <code>run_census.py</code> W/STEP units, then re-run the census Derive the window from the chosen bar's seconds; fix the participation window and the printed headline. Re-points the whole greenfield effort at a real target.	4 lines + re-run
C2	Record the true entry wall-clock in the shadow book Add <code>entry_ms</code> ; price from it, not from bar start. Then re-derive every shadow headline — patching is not enough.	1 day
C3	Bound <code>drift.read()</code> 's window AND <code>bar_ts <= ?</code> . Unbounded, it silently look-aheads in any retrospective study and returns a confident wrong direction.	2 lines
C4	Give <code>backfill_to_lake.py</code> a roll policy The dedup partitions on its own sort key, so 12 of 291 days carry the expiring contract with 118–289pt level errors.	1 hour
C5	Fix the fee model and re-price \$1.50 per contract, not per trade. Add ~\$1.00/RT of spread to any mid-priced backtest — MNQ's median spread is 0.50pt, not 0.25.	1 hour
C6	Populate <code>fills.commission</code> and reconcile a month against a broker statement 1,802 of 1,802 rows are zero. Every cost number is an assumption never checked.	half day
C7	Fix <code>exec_id</code> linkage and assert <code>sum(fills) == venue</code> 89% of booked trades cannot be traced to a fill.	half day
C8	Route the three unfiltered P&L consumers through the <code>data_quality</code> filter <code>hour_watch.py:86</code> , <code>router_watch.py:226</code> , <code>telegram_bot.py:163</code> .	30 min
C9	Detect and flag frozen quotes in the capture layer 200,708 consecutive L2 snapshots carried a stale best ask for 19.6 hours and nothing noticed. A staleness assertion at write time turns a silent corruption into an alarm — and the weekend-restart cause means it recurs until fixed.	2 hours
C10	Fix the driftlab DST anchoring before 2026-11-01 47.5% of sessions are anchored to the wrong offset, and the clocks change on that date. This one has a deadline.	1 hour
C11	Pair every live trade with its shadow twin, and run the gap test on a schedule No pairing exists today, which is precisely why the 55-second look-ahead survived eight weeks. One scheduled query would have caught it in July.	half day

#	ACTION	EFFORT
TIER 3 – THE EXPERIMENT THAT ACTUALLY ANSWERS THE QUESTION		
D1	Wire up what already works drift.py for entry, WATCH_RT < 0.50 as the filter, one trade per day, 13:30-15:00Z only. Every component exists and is tested; none is connected.	1 day
D2	Pre-register, then run 12 confirmed sessions in paper Threshold, size, stop and stop-date fixed in advance. Judge on the artifact, not the exit code. ~8 weeks to the first honest answer this desk will have produced.	8 weeks

TIER 4 – THE ROUTER		
E1	Ship a deterministic ER/ATR/hysteresis filter as the live switch driver Testable offline, free, cannot expire, provable fail-safe.	1 day
E2	Demote the LLM to shadow and run the randomised A/B It keeps deciding and being scored; nothing is applied. The only way to settle whether it adds anything. Cost falls ~99%.	half day
E3	Move off the OAuth subscription to an API key with prompt caching Fixes 94% of observed downtime and ~60% of cost — if the LLM ever returns to the live path.	1 hour

ON THE \$30,000

Not yet — and not for the reason you would expect. The strategy is not disproven; it is untested, because the instruments measuring it were mismeasuring and the components that work were never switched on.

What disqualifies today is the risk architecture: a 4-lot naked position with no stop, a documented 444-point adverse excursion, and a weekend runaway that on a \$30k account would have been a margin call rather than a drawdown.

Do Tier 0 and Tier 1, wire up D1, run the twelve sessions. Then you are deciding on evidence instead of on mood — a far better position than five months of work usually buys.

Method. Ten agents briefed independently, each instructed to treat project documentation as claims requiring verification. Databases opened read-only; no orders placed, no services restarted, no files modified during the audit. Load-bearing claims were re-derived by the lead before inclusion; three did not survive and are corrected in §07.

Sample sizes throughout are small — 45 trading days, one instrument, one volatility regime, on a paper engine that demonstrably fabricates fills. Every positive result in this report has a confidence interval that touches or nearly touches zero. That is the honest state of the evidence, and it is why the recommendation is a twelve-session test rather than a deployment.