

What actually showed promise

GAZBOT V7 · Saturday 2026-09-12 · one day's work: 50 research agents, 36 hypotheses, 10.5 years of minute tape, 26 years of daily across 14 markets. Costs charged throughout: **MNQ 1.25 points all-in per round trip** (\$1.50 commission + the half-point spread crossed once = \$2.50). **MGC 0.45 points = \$4.50.**

The honest headline: nothing found today is ready for money, and I am not going to dress that up. But that is a different sentence from "nothing is worth doing", and I spent most of today confusing the two. Below is everything with a mechanism and a live number, split by instrument, ranked by how sure I am.

Two items are not edges at all — they are money you are currently leaving on the table, and they are the most valuable things on this page. **One signal cleared a confidence interval today**, and it is on gold. Everything else is shadow material: worth forward evidence, not worth capital.

1. Ship — these cannot lose money

Not bets. Strictly better versions of things the desk already does.

#	WHAT	THE NUMBER	WHY IT IS CERTAIN
1	Stop crossing the spread. Rest at the touch instead of sending a marketable limit. SHIP	+\$0.478 / leg ≈\$4,000 / yr	35 of 35 sessions positive , worst +\$0.388. CI [+ \$0.405, +\$0.554]. It is a cheaper version of an order you already send — 38% of MNQ's entire friction. Posting INSIDE the spread was the actual hypothesis and it failed (−\$0.0097/leg); only the touch works.
2	Replace ATR14 with a HAR volatility forecast everywhere it feeds a stop width or an entry floor. SHIP	R² 0.589 vs −0.057	The best-replicated finding of the day — four independent agents hit it from different angles. Beats a random walk in 6 of 6 held-out years. Every stop on this desk is currently sized off a number with negative predictive power.

Caveat on #1: three risks are unpriced by any historical tape — our own order is not in the replayed book, the miss-fallback, and adverse selection. It should go to shadow before it goes to the live order path, and passive entry was killed here once before.

2. Gold (MGC) — where the promise actually is

#	WHAT	GROSS EDGE	95% CI	READ
3	Short-side thrust fade (abs_veto_short) on gold SHADOW	+1.137 pt +\$6.87 net/trade	[+0.105, +2.180]	The only cell all day whose interval clears zero , n=1,468 over 13 months, against a same-entry sign-flip control. And it is the same family as abs_veto_55s, the best arm in the live MNQ book at +\$4,745. The short thrust fade is

the one thing that **repeats across instruments**.

4	The regime-join at a 120-minute hold, on gold SHADOW	+0.15 to +0.28 pt	spans zero	The only net-positive cells the regime study produced on any instrument (+0.16/+0.17). Gold's edge/cost ratio is 16-62% against MNQ's 12-27%, purely because a gold round trip costs 2.8× less in points. A signal too small to pay MNQ's fee can clear gold's — that is the whole argument for gold.
5	The 14:00Z fixed-time short + its 13:35 / 15:00 controls SHADOW	+\$41.57/trade in sample	—	Beat all six purpose-built gold gates. Its mechanism is already refuted (the DST test says it is not the London fix) and both neighbours fail out of sample, so it is a spike. Only forward evidence settles a spike; the two controls are the point.

3. Nasdaq (MNQ) — thinner, and one thing that is definitely true

#	WHAT	NUMBER	READ
6	Joining a 60-minute break is adverse	-0.86 pt/day CI [-1.25, -0.47]	n≈195,000 over 10.5 years — a real, significant, negative result. Do not build momentum breakout entries. Correction from yesterday: I said the fade must therefore be the other side. It is not — the mirror control earns more out of sample, so the fade's profit was 2.5:1 target/stop geometry, not direction. The fade is less bad, not an edge.
7	Short-side thrust fade on MNQ SHADOW	+\$6.79/trade CI spans zero	Same family as #3, positive point estimate on n=1,141, not significant. Already live in the MNQ shadow book as abs_veto_55s at +\$4,745.
8	The regime-join direction call SHADOW	+0.15 to +0.34 pt every t ≈ 1	Positive in 12 of 12 cells across two instruments and eleven years against a sign-flip control — and a quarter of what MNQ's fee costs. Real and too small here, which is exactly why it is being trialled on gold.
9	Quarterly expiry week UNEXPLAINED	+8.79 pt/session to a holder	n=203 sessions. Holding earns +8.79 while cutting loses 2.41 — a 16× swing on 7.6% of sessions, and nobody has explained it. Not a strategy yet. Worth a dig.
10	Volatility affordability UNTESTED	+7.38 pt/trade to an oracle	Described by the agent that killed it as "real and huge". It died because a day-ahead forecast captures 0.0% of it — the effect lives on the range surprise. But the desk watches the range develop all day. The nowcast version has never been tested and this is the largest single number in the study.

4. The structural finding, and why it governs everything above

On 239 sessions the smallest per-trade edge that reaches statistical significance is **2.84 points at a 5-minute hold, rising to 23.57 points at a full-session hold** — against a gross edge ceiling in this instrument of **1.05 to 1.50 points**. Holding longer escapes the fixed fee and escapes statistical power at the same rate, so they cancel exactly.

On one year of one instrument there is no holding period at which this desk can both profit and prove it. That is not a reason to stop; it is the reason the answer keeps coming back "not proven" and why the bar for shadow must be lower than the bar for money.

HOLD	MOVE SD	FEE AS % OF MOVE	EDGE NEEDED TO PROVE IT
5 minutes	0.18%	2.3%	0.001%
1 session	1.62%	0.3%	0.064%
1 week	3.61%	0.1%	0.323%
1 month	7.41%	0.1%	1.352%
1 quarter	12.83%	0.0%	4.056%

Breadth is the one lever that beats this, and it was confirmed at theory today: 14 markets at an average strategy correlation of +0.099 give **6.1 effective bets** and a realised multiplier of **2.27×** against a predicted 2.47×. The mechanism works. It is multiplying a per-market Sharpe of 0.13 into 0.39, and vol-targeted buy-and-hold on the same markets does 0.54 — so breadth is real, measurable, and on this evidence still too small. It also requires multi-week overnight holds.

5. What died, so you can see the work

FAMILY	TESTED	CAUSE OF DEATH, IN ONE LINE
Reversion & fading	6	Mirror controls matched or beat them; the participation gate was a 2015–18 effect that decayed to zero by 2022.
Daily / multi-day momentum	6	The contract roll. An unadjusted front month prints the whole quarterly cost-of-carry as one bar — excising roll weeks flipped the best breakout from +17.9 to −54.9 pt/yr.
Session, clock & calendar	6	A session-boundary bug (a CME session opens at 18:00 ET, so "the 09:30 open" was the prior evening's price) manufactured an entire published finding.
Volatility	6	Every route converts a good forecast into nothing: sizing ΔSharpe +0.0029, gating worse in 13 of 15 cells.
Microstructure	6	All negative except the touch entry (#1). The paper engine's fill bias does not change any gate ranking — Spearman 0.9909 against 5.2M real book snapshots.
Portfolio	6	

The desk's nine "validated" abs_veto arms are **1.29 effective bets** at $p + 0.75$ — and held-out, that one bet is $t = -0.01$.

Plus, outside the hunt: the cut-and-flip (the cut is $+0.1$ pt/session CI $[-4.5, +4.5]$, the flip -0.2), the ATR-scaled scalping grid (every target/stop cell nets exactly minus the cost; a \$5 take-profit needs a 75% hit rate and the tape gives 39.7%), and trend across 14 markets losing to buy-and-hold 0.39 to 0.54.

6. Five things I got wrong today, corrected

WHAT I TOLD YOU	WHAT IT ACTUALLY IS	WHY I WAS WRONG
A random London long earns $+11.25$ pt	$+1.80$ pt, $t = 1.33$	The control drew its entries only from the 52 sessions its own signal fired on. Those drifted $+8.27/\text{hr}$; the other 188 drifted -0.62 .
Holding naked bleeds -26 to -38 pt/session	$+0.9$ pt over 10.5 years	That was the 2025–26 window, not the strategy. The window was the finding.
14-market trend beats buy-and-hold 0.84 to 0.65	0.39 to 0.54 — it loses	Fourteen markets keep fourteen holiday calendars; a 60-day rolling window needs 60 consecutive values, so NQ computed on 56% of rows and Russell on none.
The break is adverse, so the fade is the edge	The fade is less bad, not an edge	Its mirror control earns more out of sample — the profit was 2.5:1 geometry, not direction.
Gold's cheaper cost means the MNQ gates belong there	3 of the 4 are negative on gold	I inferred it from a cost ratio measured on a different signal and never tested these gates. Tested now: only the short thrust fade survives.

7. What is running right now

WHAT	WHERE	STATUS
Gold shadow book	mgc_slate()	3 arms $\rightarrow$ 10 . Δ 3 of the 4 gate ports measure negative on gold — recommend dropping all but mgc_absveto_short. Your call.
The tape reader	gazbot7-tape-reader, every 2 min	Reads the whole day and judges the open position — HOLD / CLAIM / CUT / ADD. Read-only , no order path, asserted by test. Grading itself from Monday.
Regime-join trial	prereg_regime_join.json	Registered, frozen, bar stated up front. Not yet wired to a timer — needs your go.
Judged-cut trial	prereg_judged_cut.json	Your cut-and-flip, with the cut decided by judgement instead of a fixed distance. The fixed 40pt version is the control and is known to be zero.
Book-thinning trial	prereg_book_thinning.json	Running since 09-11. 60 sessions.

8. Tomorrow

- **Decide on the three negative gold gates** — drop or leave. I shipped them without asking; that was my error and I have not touched them since.
 - **Say go on the two ship-now items** (touch entry to shadow, HAR swap on the non-trading path). Together they are worth about \$4,000 a year and neither is a bet.
 - **The one genuinely unexplored dig:** volatility affordability as a nowcast rather than a forecast. Largest number in the study, and the reason it died does not apply to the intraday version.
 - **The overnight question is yours and it is the big one.** Breadth is the only lever that beats the power arithmetic, and it needs multi-week holds. If "never hold overnight" stands, that lever is closed — and this report is the end of the search rather than the middle of it.
-

Every number here is reproducible from `reports/edge_hunt_2026-09-12/` (36 hypothesis directories, each with its scripts and VERDICT.md), `reports/regime_2026-09-12/`, and the committed scripts `bt_regime_transition.py`, `bt_cut_and_flip.py`, `bt_cut_flip_long.py`, `bt_breadth_daily.py`, `bt_gates_on_gold.py`, `null_expectancy_rescore.py`. Full working: `/v7/static/edge_hunt_2026-09-12.html`