

The Shadow Desk — What's Working, What's Bleeding, What To Do Monday

Week of Mon 6 -> Fri 10 Jul 2026 · every one of the 35 shadow sims on the honest tick-repriced fills, adversarially verified

§1 The shadow desk this week — read this first

You told me to point the whole report at the **32 shadow sims**, in plain English, and to cover the things the first cut missed — where the week's big runs were and whether we showed up, the footprints, the exit change we made on Wednesday, and how we'd filter the live gates so they stop chasing knives. This is that report. Here's the shape of it before you dive in.

THE ONE THING TO TAKE AWAY

The shadow desk splits cleanly down the middle, and the line is **momentum vs fading**. **Eight sims made money on honest fills (+\$2,440); the rest bled**. Every winner rides a move; every big loser fades one, sits in dead tape, or rests a limit order and gets picked off. And the sharpest version of that — **momentum done with confirmation makes +\$2,344 while "ride any strength" loses -\$1,017** — is the one finding an independent skeptic re-ran and it **held**. So "how do we make money Monday" has a plain answer: **feed the confirmed thrust, starve the fade**.

-\$8,599

HONEST NET, ALL 35 SIMS

+\$7,700 on perfect fills — the fills are the story

+\$2,344

CONFIRMED MOMENTUM

vs -\$1,017 for "ride any strength"

+\$992

THE STAR: MNQ THRUST-RIDER

beat its own perfect-fill ceiling

-\$4,249

THE WORST: PASSIVE_CHOP

no green pocket anywhere — keep it dead

35

SIMS DISSECTED

+ 7 contracts of big-run footprints

§1.1 Where the money was — and did we show up?

This time I used the **full run census**, not a top-8 (you were right to call that out). Across the seven contracts there were **540 big runs this week, and we sat out 413 of them**. But here's the honest

surprise when you look at all of them: **the runs themselves aren't the leak**. The with-the-move momentum sims that did show up already print **+\$2,441 on honest fills** — while the fade book gave back **-\$11,040 fighting the very same moves**. The money isn't in the runs we missed; it's in the fades we should never have taken. And most runs have **no readable tell before they fire** — the tape and even the L2 book stay flat or point the wrong way until the break — so the play is to **join a move once it's clearly moving, not predict it**. The one clean, repeatable exception across contracts is the **US cash-open window (13:00-15:00 UTC)**, where half the week's biggest bursts live.

§1.2 So how do we grab them from Monday?

THE CAPTURE PLAN — THE POINT OF THE WHOLE REPORT

Don't build a new edge. **Point the one proven entry we already have — confirmed-momentum + chandelier**

(`tw\_mnq\_thrust\_cont`, +\$992) — at the FLOW-LED continuation runs we keep sitting out on MNQ and MES (price already leaning with a positive VWAP slope, ATR expanding — that's the one cluster with a clean pre-run signature). In the same breath, **stop the reversion sims fading into the vacuum runs** — that's where the **-\$11,040** went. For the US-open bursts on M2K/MYM/MBT, shadow a simple opening-range break (never pre-position — the open whips both ways). And the no-tape contracts (MCL, and MBT's book) are a **data-capture job, not tradeable yet**. All shadow-first; arm nothing with size until it proves a second week.

ONE RULE THAT MUST NOT GET LOST — MOMENTUM IS TWO-SIDED

Every momentum gate trades **both ways**: it **longs** a rising thrust and **shorts** a falling one. **Never turn a momentum gate long-only**. This week the momentum **short** side (+\$1,435) actually **beat** its long side (+\$1,006) — on MES and MGC the short side

was the only green side — so a long-only version would have binned the bigger half of the edge. The thing we kill is **FADING** (selling rips / buying dips against the move), never **shorting**; the "shorts bled" this week were the fade book's shorts fighting up-moves, a completely different behaviour. (Prerequisite for ever taking momentum live: re-enable shorts on the desk.)

§1.3 The exit change, before and after

On Wednesday afternoon we threw out the nine-mechanism exit stack for two simple ones — reversion gets a 1.5-ATR stop and a 2R target, momentum gets a 1-ATR stop and a single trailing chandelier. The honest read after the switch: **momentum keeps every dollar on that one chandelier — it matches the old eight-mechanism spine for free** (62% of trades close at about +\$10). Reversion still bleeds, and the logged worry is confirmed — the **2R fade target only fires on about one trade in ten**; a nearer target trims the bleed but never fixes it, because **the fade's problem is the entries, not the exit**. Simpler exits are as good or better; the work is on the entry side.

§1.4 The live gates, and the knife-filter

Both live gates (``vwap_dip``, ``vwap_pullback``) are mean-reversion, and their whole shadow brood bled −\$5,157 this week. Here's the counter-intuitive bit worth your attention: **the "only fade in quiet tape" filter is backwards** — the book bled worst in the quietest tape. The fix isn't a calm-tape band, it's a **direction/trend veto** — stop fading into a live move. On the knife-filter you asked about: I owe you the honest result — **aggressor-tape confirmation on entry is null ($p=0.40$), and the flow-flip exit that looked great at one-minute dies at tick resolution**. The obvious bolt-on doesn't pay. The one use still open is **L2 absorption to gate the knife-**

catching fades — but that needs another week or two of depth capture before it's even testable.

\$1.5 What survived the skeptic

Every bankable claim got re-run by an independent skeptic. Most clever per-sim tweaks **fell** — on one week a filter is usually just "trade less" in disguise. What **held** was the big structural stuff: the **confirmed-vs-loose momentum split** (the skeptic reproduced +\$1,828), and that the worst faders genuinely bleed and should be stood down, not tuned. Read the amber rows below as leads for the shadow book, not instructions.

CLAIM	AGENT SAID	RE-RUN	VERDICT	WHY
The OFF switch: when should each sim NOT fire?	£3500	£6647	HELD	REAL. On honest net real_pnl for Mon6-Fri10 Jul (n=3797, desk - \$8599): rvol<0.8 = 58.9% of entries (claim says two-thirds; actually ~three-fifths) and - \$8414 = 97.9% of the net loss; muzzling it desk-wide lifts the desk to -\$18...
tw_deadtape_standdown: it stood down in the quiet and got ru	£529	£528	HELD	The arithmetic reproduces to the dollar (veto MBT/MNQ/MCL = 108 trades, -\$669.9/~£528 in-sample swing; calm residual 152 trades, -\$292 still red), so the claim as written holds, but it is a post-hoc 'drop the three worst-losi...
pullback_loose: the longs are fine, the shorts are the wound	£402	£402	HELD	The +£402/wk headline reproduces exactly and the load-bearing move — stand down on shorts — is directionally REAL (pullback_loose shorts bleed every single day this week, - \$509/84, and -\$803/144 all-data: a fader getting run over)...
	£150	£1828	HELD	HELD — reproduces to the dollar on honest tick-repriced fills, but

Momentum is confirmed-vs-loose, one contract at a time				it's a thin-n, MNQ-concentrated one-week read, not yet a bankable forward edge. REPRODUCTION (data/shadow_real.db::shadow_real_pn real_pnl, JOIN alphabot.db...
tw_ride_strength: it rides strength fine going up, and gets	£541	£541	FELL	The ~£541 reproduces exactly (killing the -\$705.69 short book at ~0.767 GBP/USD), but it is loss-avoidance, not an edge: long-only still nets -\$40 gross (worse after fees — the sim is cost-blind), the surviving long side is proppe...
passive_mgc_hours: the shorts are fading rips that run it ov	£433	£434	FELL	The ~£433 headline roughly reproduces — MGC's short leg really lost about \$555.72 (~£434, already net of round-trip fees) this one week — but it does NOT hold as an edge: it's a thin-n (31 filled shorts), single trend-we...
How We Grab The Runs From Monday	£400	£975	FELL	CURVE-FIT / MISLABELED PARTITION — does not hold as stated. The figures reproduce exactly, but + \$2,441 is just the 9 momentum strategies with pnl>=0, and - \$11,040 is the arithmetic complement (week_total -8,599 minus the winner...
dip_loose: three markets, one edge, and it's all in MNQ	£335	£0	FELL	The numbers reproduce to the dollar on honest real_pnl (-\$184.5 to + \$242.6, ~£335 in-sample swing) and it's cost-aware not cost-blind, but it's essentially a "trade-less" move whose entire £335 is avoided MES/ M...
MCL "Short With Strength" — It Shorts Into A Tape	£321	£4	FELL	The direction reproduces — trending MCL shorts bled the whole -\$461 while chop/shock netted +\$100 gross — but the claimed "+\$78 green book" does NOT hold: it's a one-week, in-sample, single-threshold cut on the stra...

pullback_quiet: it only trades when nobody's home, and	£307	£0	FELL	Not real: the actual entry_rvol maxes at 0.199 so a rvol>=0.8 floor vetoes ALL 106 trades (not 75 of 92) — the claim is a mislabeled, curve-fit "trade-less" cut (effective threshold ~0.15, chosen in-sample on this sin...
MGC coil-break: it pays, but only when the break has already	£190	£0	FELL	It does not hold: the whole 134-trade sim genuinely nets +\$136.92 (~£106) this week, but the claimed +£190/wk gain is not a real edge — it is unreproducible (shadow_real.db stores no VWAP or any entry feature, and gate_feature_log...
The Breakout Sim Pays — But Two Symbols Are Eating a Third o	£173	£0	FELL	Not real: the true drop-MCL+MYM delta is +\$260 (USD, net of fees) not £173, and it is a one-week curve-fit whose entire +\$154 "kept book" rests on two MNQ trades -- strip them and it is -\$146, four of the five "kept...
tw_thin_drift_standdown: it keeps shorting into the drift an	£170	£168	FELL	The arithmetic reproduces exactly and cost-honestly (vetoing below-VWAP shorts removes 13 trades worth -\$215.56, lifting the week from -\$210.98 to +\$4.58, a ~£168 swing that matches the claimed ~£170), but it does NOT hold as a fo...
Strength-Ceiling Pullback: the shorts are the disease, thin	£164	£164	FELL	The arithmetic reproduces exactly and net-of-fees (22 shorts = - \$207.79 real / -\$46.52 ceiling ≈ £164), but the refinement does NOT hold as a forward edge: it is a thin-n, 100%-in-sample "trade less" call on a single wee...
tw_mnq_pb_expand_veto — the quiet bars are killing it	£163	£-22	FELL	The arithmetic reproduces exactly (veto 9 trades 0-for-9 for -\$208.58, keep 4 for -\$28.48), but it is a curve-fit, thin-n, trade-less artifact — the 0.042 threshold is fit into the single in-sample gap, the surviving 4-trade bucke...

**The biggest runs on MNQ
this week — did we show
up?**

£150

£-278

FELL

REPRODUCES but does not survive
as a bankable edge. Every
descriptive fact checks out on honest
cost-inclusive real_pnl: 76 runs at
1.5xATR, 17 caught for +\$993.5 real,
fought ~+\$18, and thrust_cont +
\$992 / momentum +\$495 / orb +\$...

§1.6 So what do we do Monday?

THE PLAN, IN ORDER

1. Lean the book toward **confirmed momentum on the equity/gold micros, led by MNQ** — and only when there's a real push and real amplitude already in the tape (the one pocket that was green all five days). **2.** Stand down the confirmed bleeders — passive_chop first (it has no favourable pocket anywhere), then the fade family; and **never let the desk fade a live move**. **3.** Bin the "ride any strength" family (tw_ride_strength, tw_mcl_short_with_strength) — it's the -\$1,017 drag. **4.** Everything else stays shadow-only for another week; arm nothing with size until it survives a second, independent week. Feed the confirmed thrust, starve the fade — slowly, on what proves itself twice.

§1.7 How to read the rest

First, **where the money was** — the full run census on each contract (every big run, not a top-8), whether we showed up, and the footprint before each. Then **the through-lines**, opening with "**how we grab the runs from Monday**" (the capture plan across all contracts), then momentum confirmed-vs-loose, the exit before/after, the live gates and the knife-filter, **the L2 order book — our first honest look** (the data you caught me leaving out), the OFF switch, greenfield, the sizing ladder, and "**your leads this week**" — every mid-week musing walked to its verdict. Then **the winners**, one plain

dossier each (lean into these), and **the bleeders** (fix or retire these). Each sim dossier carries a coloured verdict at the top — green works, amber means one filter might fix it, red means retire. The green verdicts, the confirmed-momentum finding, the OFF-switch, and the capture plan are what to act on.

Contents

Every entry is a live link — tap it (page number is the PDF page).

§1 — The shadow desk this week — read this first	1
§1.1 — Where the money was — and did we show up?	2
§1.2 — So how do we grab them from Monday?	3
§1.3 — The exit change, before and after	4
§1.4 — The live gates, and the knife-filter	4
§1.5 — What survived the skeptic	5
§1.6 — So what do we do Monday?	8
§1.7 — How to read the rest	8
§2 — The biggest runs on MNQ this week — did we show up?	18
§2.1 — Every run this week, in order	19
§2.2 — The four clusters, and how blunt the participation is	25
§2.3 — What the ten minutes before each cluster looked like	26
§2.4 — The L2 book (MNQ, Thu 07-09 14:17 → Fri 07-10 21:00 UTC)	27
§2.5 — The honest money left on the table	28
§2.6 — How we grab more of them from Monday	28
§3 — The biggest runs on MES this week — did we show up, and what did they look like?	
§3.1 — The week in one picture — 63 runs, biggest at the extremes	31
§3.2 — 1 · The full census — every run, not a top-8	32
§3.3 — 2 · Participation — we showed up for 10, sat out 53	37
§3.4 — 3 · The footprint — what the 10 minutes before each cluster looked like	39
§3.5 — 4 · The L2 order book — is the vacuum visible before it opens?	40
§3.6 — 5 · The honest money left on the table	42
§3.7 — 6 · How we grab them from Monday	42
§3.8 — SHADOW-FIRST FLOW-LED & the momentum burst — widen what already works	42
§3.9 — DON'T CHASE VACUUM — unforecastable from what we capture	43

§3.10 — REACT, DON'T PREDICT OPEN / NEWS — the spike isn't tradeable, the second leg might

§4 — The biggest runs on MGC this week — did we show up?	45
§4.1 — 1 · The full census — every run, not a top-8	45
§4.2 — 2 · Did we show up? Blunt version	51
§4.3 — 3 · The footprint, clustered — what did the 10 minutes before look like?	51
§4.4 — 4 · The L2 order book — could we see the vacuum coming?	53
§4.5 — 5 · The honest money left on the table	54
§4.6 — 6 · How we grab them from Monday	55
§5 — The biggest runs on M2K this week — did we show up?	57
§5.1 — 1 — The full census: all 64 runs	57
§5.2 — 2 — Did we show up? Bluntly, less than half	63
§5.3 — 3 — The footprint before a run: it erupts from quiet	64
§5.4 — 4 — The L2 book: nothing usable for M2K this week	65
§5.5 — 5 — The honest money left on the table	66
§5.6 — 6 — How we grab them from Monday	66
§5.7 — The OPEN-hour break shadow-first	66
§5.8 — The quiet overnight cluster can't forecast — don't try	67
§6 — The biggest runs on MYM — did we show up?	69
§6.1 — 1 The full census — every run $\geq 1.5 \times \text{ATR}$ this week	70
§6.2 — 2 Did we show up? — participation and the honest money	75
§6.3 — 3 The footprint, clustered — what the 10 minutes before looked like	77
§6.4 — 4 The L2 book — MYM has none, so I used MNQ as a proxy	79
§6.5 — 5 The honest money left on the table	80
§6.6 — 6 How we grab them from Monday	81
§7 — The biggest runs on MCL this week — did we show up?	83
§7.1 — 1 · The full census — all 78 runs	83
§7.2 — 2 · Did we show up? Bluntly, mostly no	89
§7.3 — 3 · The footprint before each cluster	90
§7.4 — 4 · The L2 book — nothing to show, and that's the point	91
§7.5 — 5 · The honest money left on the table	92
§7.6 — 6 · How we grab them from Monday	92
§8 — The biggest runs on MBT this week — did we show up?	95
§8.1 — 1 · The full census — every run, not a top-eight	95
§8.2 — 2 · Participation — we caught 19, and lost the money anyway	104
§8.3 — 3 · The footprint — when MBT runs, and why we're in the wrong session	105

§8.4 — 4 · The book — MBT has none, and the tape doesn't lead	106
§8.5 — 5 · The honest money left on the table	107
§8.6 — 6 · How we grab them from Monday	108
§8.7 — Lead — the US-open continuation break SHADOW-FIRST	108
§8.8 — Watch, don't build — the other three clusters	109
§9 — How We Grab The Runs From Monday	111
§9.1 — 1 The desk-wide scoreboard — where the money ran this week	112
§9.2 — 2 The honest money — the direction was right, the fades killed it	114
§9.3 — 3 The Monday capture play for each cause cluster	115
§9.4 — 4 Forecastable vs not — be blunt	116
§9.5 — 5 The single highest-value build	117
§9.6 — 6 Monday action list	118
§10 — Momentum is confirmed-vs-loose, one contract at a time	121
§10.1 — The per-contract matrix — where momentum actually pays	122
§10.2 — The mechanism split — what "confirmed" means here	124
§10.3 — Honest health warnings	125
§11 — The exit swap: two pure archetypes, before and after	127
§11.1 — The board, before and after	127
§11.2 — Inside the momentum book: the give-back works, the entries don't	128
§11.3 — Is the 2R fade target too far? Yes — but that's a tune, not the fix	130
§11.4 — The plain takeaway	131
§12 — The two live gates, their shadow children, and the knife-filter	133
§12.1 — What the two gates actually are	133
§12.2 — How they behaved this week, on money that's real	134
§12.3 — Which shadow sims are children of which gate	135
§12.4 — The condition filter — and the counter-intuitive part	136
§12.5 — The knife-filter — the obvious bolt-on that doesn't pay	137
§13 — The order book — our first honest look	139
§13.1 — What we actually captured	139
§13.2 — Does the book point the way? No — if anything it points the wrong way	140
§13.3 — The knife test: does fading into a wall lose more?	141
§13.4 — What the book did around the big moves	142
§13.5 — Honest takeaway	143
§14 — The OFF switch: when should each sim NOT fire?	145

§14.1 — The cheap desk-wide gate: is the tape tradeable at all?	146
§14.2 — Per-sim pockets: where each big one actually lives	148
§14.3 — momentum shadow — a real amplitude plateau FLOOR IT	148
§14.4 — passive chop — no pocket exists KILL	149
§14.5 — pullback loose — same thin-tape disease FLOOR IT	150
§14.6 — Plateau or knife-edge? — the one-week honesty check	150
§15 — Greenfield: if we can't catch the big misses, what would?	152
§15.1 — The misses, laid out honestly	152
§15.2 — The first test kills the obvious idea	153
§15.3 — So I flipped it — and built the fade instead	154
§15.4 — The out-of-sample test buries it	155
§15.5 — And the order book had nothing to add	156
§15.6 — The plain takeaway	156
§16 — The Sizing Ladder — Has Any Sim Earned More Size?	157
§16.1 — The green sims built real money this week	157
§16.2 — Now score them against the ladder — and the ladder says "not yet"	158
§16.3 — Two robustness caveats before anyone gets excited	160
§16.4 — The Departure Lounge	
§17 — Your leads this week — what we found	165
§17.1 — The ledger	165
§18 — The MNQ thrust-rider: it caught the runs and the stops were free	170
§18.1 — Why it works: big winners on the ratchet, free losers on the stop	170
§18.2 — What separates its winners from its losers: bar size	171
§18.3 — The honest size of the win	172
§19 — Momentum Shadow: it made money by letting winners run and getting stopped fast	
§19.1 — Did it make money — yes, and here's exactly how	173
§19.2 — Why it wins vs loses — the bar has to be moving	174
§19.3 — The one thing for Monday: put a floor under the bar	175
§20 — The Breakout Sim Pays — But Two Symbols Are Eating a Third of It	177
§20.1 — What it's actually doing: let the runner run, cut the rest	178
§20.2 — Where it wins vs where it bleeds	178
§20.3 — The equity curve — climbed to +\$300, gave a third back	179
§20.4 — The one thing for Monday	179
§21 — M2K loose-momentum: a winner that keeps chasing spikes it shouldn't	181

§21.1 — Did it make money — yes, and it's a pure momentum shape	181
§21.2 — Why it wins and why it bleeds — it's late to the hot ones	182
§21.3 — The one thing for Monday — stand down when rvol is hot	182
§21.4 — The honest caveat	183
§22 — MGC coil-break: it pays, but only when the break has already left home	184
§22.1 — What it's actually doing	185
§22.2 — The one thing for Monday: only fire when the break has left VWAP	186
§23 — MGC Momentum-Thrust: it made money, but the whole week is six trades	188
§23.1 — 1. The money — it's real, and it's almost all short	188
§23.2 — 2. Why it won — Goldilocks amplitude	189
§23.3 — 3. The one thing for Monday — veto the extremes	190
§24 — MES thrust-continuation: a real little winner, and one fat-bar leak	192
§24.1 — What it's actually doing	192
§24.2 — The one thing for Monday: stand down on fat bars	193
§25 — MYM open-thrust: green, but one trade carried the whole week	195
§25.1 — 1) Did it make money? Yes — \$28 honest, and it beat its ceiling	195
§25.2 — 2) Why — it works when the tape is trending, it bleeds when it chops	196
§25.3 — 3) The one thing for Monday — stand down in chop, only fire when trending	197
§26 — passive chop: it fades the chop, and the chop keeps going	199
§26.1 — What it's actually doing	199
§26.2 — Is there a gate that saves it? I looked. There isn't.	200
§27 — tw deadtape standdown: it stood down in the quiet and got run over in the fast	204
§27.1 — Why it bleeds: it leans on fast, event-driven names and the stop eats it	204
§27.2 — The one thing for Monday: stand down on the fast names	205
§28 — tw clock standdown: a slow, all-week bleed with no rescue gate	207
§28.1 — 1) Did it make money? No — it bled all week	207
§28.2 — 2) Why it loses: small wins, bigger losses, and it stops out constantly	208
§28.3 — 3) The one thing for Monday: cut its three worst symbols — but it's still red, so retire it	
§29 — tw ride strength: it rides strength fine going up, and gets run over going short	211
§29.1 — What it's actually doing	211
§29.2 — The honest-money curve: the short book is the whole drawdown	212
§29.3 — The one thing for Monday	213
§30 — passive mgc hours: the shorts are fading rips that run it over	214
§30.1 — 1. Did it make money — no, and it's a one-sided wound	215

§30.2 — 2. Why — it fades a rip above VWAP and the rip keeps going	215
§30.3 — 3. The one thing for Monday — stand the shorts down	216
§31 — pullback loose: the longs are fine, the shorts are the wound	218
§31.1 — Did it make money? No — and it's all shorts	219
§31.2 — Why — it fades rips that just run it over	219
§31.3 — The one thing for Monday	220
§32 — pullback quiet: it only trades when nobody's home, and that's the whole problem	222
§32.1 — The bleed never let up	223
§32.2 — Why it loses: it fires into dead tape and the stop pays for it	223
§32.3 — The one thing for Monday: don't fire unless someone's actually in the tape	224
§33 — MCL "Short With Strength" — It Shorts Into A Tape That's Still Trending, And The T	227
§33.1 — What it's actually doing — shorting strength and getting steamrolled	227
§33.2 — The one thing for Monday — stand down when the tape is trending	228
§34 — The MBT Euphoria-Fade: Shorting Into A Freight Train	230
§34.1 — 1. The money — I re-ran it myself	230
§34.2 — 2. Why — it fades a rip that just runs it over	231
§34.3 — 3. The one thing for Monday — there isn't a gate, retire it	232
§35 — Strength-Ceiling Pullback: the shorts are the disease, thin fills are the fever	234
§35.1 — Why it bled: two different diseases wearing one coat	235
§35.2 — The one thing for Monday: stand down the shorts	236
§36 — tw mnq pb expand veto — the quiet bars are killing it	238
§36.1 — Why it bled — it keeps firing when nobody's home	239
§36.2 — The condition split — bar range is the only thing that sorts them	240
§36.3 — The one thing for Monday	240
§37 — tw thin drift standdown: it keeps shorting into the drift and gets run over	242
§37.1 — Where the money went	243
§37.2 — Why it wins vs why it loses	243
§37.3 — The one thing for Monday	244
§38 — The MGC pullback that went 0-for-7	246
§38.1 — 1) The money — it just bled, straight down	246
§38.2 — 2) Why — bad entries, and two "winners" that were pure fill fantasy	247
§38.3 — 3) The one thing for Monday — there isn't one. Retire it.	248
§39 — dip loose: three markets, one edge, and it's all in MNQ	250
§39.1 — 1) The money — it bled, but only in two rooms	251

§39.2 — 2) Why — gold's fills eat the edge, MES has no edge to eat	251
§39.3 — Honest cumulative P&L by symbol — MNQ climbs, gold sinks	252
§39.4 — 3) The one thing for Monday — trade MNQ only	253
§40 — MCL down-flush sim: it fires into dead tape and pays for it	254
§40.1 — What it's actually doing	254
§40.2 — The one thing for Monday	255
§41 — tw xconfirm veto — a fader that keeps getting run over	257
§41.1 — What it's actually doing	257
§41.2 — Every trade this week, worst to best (fat bars ringed in gold)	258
§41.3 — The one thing for Monday: stand down when the bar is fat	259
§42 — tw mes block us open — a fade that gets run over, for a fill it can't get	260
§42.1 — What it's actually doing	260
§42.2 — The one thing for Monday — and the honest answer is: retire it	261
§43 — MGC Midband Pullback: It Buys Into Selling and Gets Run Over	263
§43.1 — Why it bled: the winners trail out, the losers get stopped	263
§43.2 — The one thing for Monday: don't fade into the order flow	264
§44 — tw mes atr ceiling 003 — the ceiling's set too high, so it keeps buying the fat bars t	266
§44.1 — Where the money actually went	266
§44.2 — Why — thin tape it ratchets out green, fat tape it gets run over	267
§44.3 — The one thing for Monday — bring the ceiling down to atr pct 0.020	268
§45 — passive_chop_calm: it fades the quiet, but the quiet is overnight and the fills are ugl	269
§45.1 — What it's actually doing	269
§45.2 — When it fires vs when it pays	270
§45.3 — The one thing for Monday	271
§46 — The MNQ Greenband-Block Short: Four Fired, Four Bled	272
§46.1 — Did it make money? No. And the ceiling proves it wasn't just bad fills.	272
§46.2 — Why it bled: it's shorting into chop, below VWAP, when nobody's trading	273
§46.3 — The one thing for Monday: don't fire it in chop	274
§47 — The MYM drive-chaser that only pays above the line	275
§47.1 — Did it make money? No — and the shape tells you why	275
§47.2 — Why it bleeds: small winners, full-size stops	276
§47.3 — The answer: it chases drives while the tape is sagging	276
§47.4 — The one thing for Monday	277
§48 — tw m2k open momentum — the open thrust that never thrusts	279

§48.1 — 1. Did it make money? No — it bled every single time.	279
§48.2 — 2. Why — it fires into fat, fast bars and gets run over	280
§48.3 — 3. The one thing for Monday — retire it (or at best, veto fat bars)	281
§49 — The Open-Surge Sim Shorts the Rip It's Meant to Buy	283
§49.1 — Why it bled: it fired five shorts and got run over	283
§49.2 — The condition split — long vs short	284
§49.3 — The one thing for Monday	285
§50 — The MBT pullback that loses even when it's right	286
§50.1 — Did it make money? No — it went 0-for-8	286
§50.2 — Why: the entries don't follow through, and the fills finish them off	287
§50.3 — Trade by trade: the fantasy vs what you actually got	287
§50.4 — Both lines are underwater — the gap is the fills	288
§50.5 — The one thing for Monday: retire it	288
§51 — tw m2k vst — a good little dip-buyer that keeps taking one trade too many	290
§51.1 — Did it make money? No — the stops give back everything the ratchet earns	290
§51.2 — Why — it buys the dip when it wins, and chases the rip when it loses	291
§51.3 — The one thing for Monday — stand down when the trade fires on the wrong side of VWAP	
§52 — The Small Fry — too thin to judge, worth a glance	295
§52.1 — M2K absorption-only WATCH	295
§52.2 — MBT spike-fade IGNORE	295

CHAPTER

Where The Money Was — The Week's Big Runs, And Did We Show Up

⚠ A SKEPTIC RE-RAN THE MONEY CLAIM — IT FELL (READ AS A LEAD, NOT A RULE)

REPRODUCES but does not survive as a bankable edge. Every descriptive fact checks out on honest cost-inclusive real_pnl: 76 runs at 1.5xATR, 17 caught for +\$993.5 real, fought ~+\$18, and thrust_cont +\$992 / momentum +\$495 / orb +\$318 / ride_strength SHORT -\$527 all confirmed — so it is NOT curve-fit and NOT cost-blind. BUT it fails as bankable money on three counts. (1) Cherry-picked winning subset: the +\$993 is gross winnings of run-aligned trades pulled out of a book that LOST money this week — honest MNQ real_pnl was -\$353 (report window, 458 trades) to -\$458 (full 7d, 467 trades); you cannot bank the +993 without the fades that gave it all back. (2) Thin-n: top 2 runs = 54% and top 4 = 81% of the \$993; the 'FLOW-LED already prints' claim rests on 4 caught runs = +\$438 of which 2 trades carry +\$432. (3) The actionable FLOW-LED continuation is an unproven shadow-first hypothesis — the report's own honest-null shows the pre-run tape does not forecast direction (8/19), leaving only a fragile +0.19-ATR VWAP slope on n=19. Real insight that survives: the run-catchers are genuinely +EV and the fades bleed. Real number: the honest MNQ book this week was a LOSS of about GBP 278 (-\$353.5 report window; -GBP 361 on the full 7 days), not the ~GBP 150/wk implied.


§2 The biggest runs on MNQ this week — did we show up?

MNQ printed **76 runs of 1.5×ATR or bigger** in the last seven days (a 15-minute move of at least 73.9 points, against a typical 15-min range of 49.2). That is not a top-eight story — it is the whole tape. We showed up for **17** of them, we were positioned against **10**, and we sat out the other **49**. When we did align

with a run we banked **+\$993** of honest money; the shapes that print it are already in the book. The bleed is not the runs — it is the fading we do in the quiet between them.

§2.1 Every run this week, in order

Each bar is one run, up or down by size, along the week left-to-right. **Green** = a sim fired with the run (caught), **red** = a sim fired against it (fought), navy = we sat it out. The tall red bar mid-chart is the 07-08 08:15 flush — a -335-point drop we were long into.



THE FULL CENSUS

All 76 runs. **Move** is the directional swing in points; **\$** is that move on one MNQ lot (\$2/pt) — the hindsight ceiling if you nailed the whole thing. **real \$** is the honest money the aligned (or opposing) sims actually made near that run. **Flow** is net aggressor volume during the run; **amp** is pre-run amplitude. Green rows we caught, red rows we fought.

TIME (UTC)	DIR	MOVE	\$ 1LOT	US	REAL \$	FLOW	AMP	CLUSTER
07-05 23:47	DN	-84	168	sat out		+230	0.10	VACUUM

07-06 00:16	DN	-110	219	caught	+1	+281	0.19	VACUUM
07-06 00:34	UP	+77	154	caught	+6	-75	0.28	VACUUM
07-06 01:13	DN	-90	180	FOUGHT	+3	+24	0.14	VACUUM
07-06 01:40	DN	-77	154	sat out		+15	0.09	VACUUM
07-06 03:40	DN	-89	178	caught	+2	-173	0.08	FLOW- LED
07-06 04:03	UP	+77	154	sat out		-7	0.11	VACUUM
07-06 06:23	UP	+74	148	caught	+0	-128	0.07	VACUUM
07-06 08:00	DN	-80	159	caught	+4	-505	0.07	FLOW- LED
07-06 13:00	DN	-94	189	sat out		+23	0.05	OPEN/ NEWS
07-06 13:36	UP	+185	370	caught	-52	-958	0.51	OPEN/ NEWS
07-06 13:59	UP	+90	180	FOUGHT	+3	-482	0.25	OPEN/ NEWS
07-06 15:03	DN	-96	191	sat out		-105	0.16	FLOW-LED
07-06 15:18	UP	+113	226	caught	+12	+20	0.17	UNCLASS
07-06 17:16	DN	-87	174	caught	+48	+180	0.08	VACUUM
07-06 18:01	UP	+92	184	caught	+54	-34	0.09	VACUUM

07-06 19:46	DN	-87	174	FOUGHT	-67	+54	0.10	VACUUM
07-07 00:25	DN	-96	192	FOUGHT	-38	+16	0.12	VACUUM
07-07 01:12	DN	-74	148	sat out		-24	0.09	UNCLASS
07-07 05:28	UP	+94	188	sat out		+72	0.17	FLOW-LED
07-07 11:55	UP	+88	177	sat out		+53	0.13	FLOW-LED
07-07 13:24	DN	-104	207	FOUGHT	+6	-6	0.11	OPEN/ NEWS
07-07 13:43	DN	-132	265	caught	-77	+64	0.35	OPEN/ NEWS
07-07 14:10	DN	-152	303	sat out		-765	0.21	OPEN/ NEWS
07-07 14:44	UP	+86	173	sat out		+461	0.24	OPEN/ NEWS
07-07 15:40	UP	+113	226	sat out		-1221	0.22	VACUUM
07-07 16:08	UP	+138	276	sat out		+646	0.16	FLOW-LED
07-07 18:28	DN	-75	150	sat out		-71	0.11	FLOW-LED
07-07 18:46	DN	-106	212	caught	+42	+106	0.14	VACUUM
07-07 19:04	DN	-82	165	sat out		+95	0.17	VACUUM
07-07 19:19	UP	+83	166	sat out		-458	0.24	VACUUM

07-07 19:53	UP	+91	182	sat out		-302	0.19	VACUUM
07-07 20:46	DN	-86	172	sat out		-74	0.04	FLOW-LED
07-08 00:01	UP	+143	286	sat out		-353	0.19	VACUUM
07-08 00:35	UP	+112	224	sat out		+152	0.14	FLOW-LED
07-08 01:23	DN	-81	162	caught	+71	+12	0.10	VACUUM
07-08 07:20	UP	+88	176	caught	+123	-97	0.12	VACUUM
07-08 08:15	DN	-334	669	FOUGHT	+8	+151	0.13	VACUUM
07-08 08:40	UP	+124	247	sat out		-295	0.38	VACUUM
07-08 08:55	DN	-143	286	sat out		-192	0.20	FLOW-LED
07-08 10:13	UP	+74	149	sat out		+105	0.10	FLOW-LED
07-08 11:29	UP	+114	228	sat out		-107	0.07	VACUUM
07-08 11:53	UP	+89	178	sat out		+133	0.18	FLOW-LED
07-08 13:25	UP	+209	418	sat out		-313	0.20	OPEN/ NEWS
07-08 13:59	DN	-112	223	sat out		-175	0.35	OPEN/ NEWS
07-08 14:35	DN	-142	284	sat out		-274	0.25	OPEN/ NEWS

07-08 15:11	DN	-161	322	sat out		+387	0.13	VACUUM
07-08 15:26	UP	+94	188	sat out		-400	0.34	VACUUM
07-08 15:53	UP	+77	154	sat out		+212	0.17	FLOW-LED
07-08 16:08	UP	+151	302	sat out		-177	0.18	VACUUM
07-08 16:37	UP	+114	228	sat out		-358	0.23	VACUUM
07-08 17:15	UP	+98	196	caught	+80	-20	0.11	VACUUM
07-08 17:34	DN	-92	184	sat out		+14	0.08	VACUUM
07-08 17:57	UP	+79	158	sat out		+68	0.23	FLOW-LED
07-08 19:42	UP	+92	184	sat out		-132	0.11	VACUUM
07-08 19:57	DN	-77	154	FOUGHT	-1	+450	0.21	VACUUM
07-08 23:30	UP	+80	161	caught	+145	+156	0.10	FLOW-LED
07-09 00:49	DN	-85	170	sat out		-296	0.17	FLOW-LED
07-09 01:47	DN	-94	188	sat out		+76	0.09	VACUUM
07-09 02:06	UP	+103	206	caught	+287	+167	0.15	FLOW-LED
07-09 02:31	DN	-80	160	sat out		-132	0.08	FLOW-LED

07-09 03:18	UP	+95	190	FOUGHT	-36	-510	0.10	VACUUM
07-09 05:43	UP	+76	152	FOUGHT	-26	-92	0.07	VACUUM
07-09 10:39	DN	-85	170	sat out		-17	0.08	UNCLASS
07-09 13:31	UP	+184	369	sat out		-356	0.25	OPEN/ NEWS
07-09 13:52	DN	-283	566	FOUGHT	+163	+6802	0.32	OPEN/ NEWS
07-09 14:21	UP	+164	328	caught	+249	-615	0.46	OPEN/ NEWS
07-09 14:53	UP	+112	223	sat out		-524	0.29	OPEN/ NEWS
07-09 23:56	DN	-87	174	sat out		+10	0.09	VACUUM
07-10 13:21	UP	+106	212	sat out		+104	0.08	OPEN/ NEWS
07-10 13:40	UP	+108	216	sat out		+412	0.38	OPEN/ NEWS
07-10 14:18	DN	-106	212	sat out		-1073	0.16	OPEN/ NEWS
07-10 14:49	DN	-99	198	sat out		+994	0.15	OPEN/ NEWS
07-10 15:05	UP	+136	272	sat out		-359	0.18	VACUUM
07-10 15:27	DN	-139	278	sat out		+670	0.22	VACUUM
07-10 15:45	UP	+112	224	sat out		+171	0.27	FLOW-LED

Verified against the 5-second bars myself: the 07-08 08:15 down-run shows open 29409 → low 29051, a −356-point real drop (census said −334) — confirmed, and it was the biggest of the week. The 07-06 13:36 up-run swung 29786 → 29984, a clean +198-point excursion (census +185) — confirmed. One caveat: for the 07-09 13:52 news candle the 5s bars are sparse (21 bars where 181 belong, the known 5s flush-loss), so I trust the scorecard's 250ms read there, not the bars.

§2.2 The four clusters, and how blunt the participation is

CLUSTER	RUNS	CAUGHT	FOUGHT	SAT OUT	CEILING \$	WHAT IT IS
VACUUM	36	9	7	20	7,633	snapped back against the prior tape
OPEN/NEWS	18	3	3	12	4,936	13:00–15:00 UTC data/open expansion
FLOW-LED	19	4	0	15	3,561	price already leaning, then continued
UNCLASS	3	1	0	2	544	quiet, no clean cause

The blunt version: on the 36 VACUUM snapbacks we fought seven of them — we were fading into the move as it detonated. On the 19 FLOW-LED continuations we fought **zero** and simply watched 15 of them go by. That asymmetry is the whole chapter: we get run over by reversals and we no-show the continuations.

§2.3 What the ten minutes before each cluster looked like

Averaging the signal funnel over the 600 seconds before every run, by cluster:

CLUSTER	RVOL	ATR%	VWAP DIST%	VWAP SLOPE (ATR)	READ
VACUUM	0.58	0.064	−0.05	−0.53	price drifting the wrong way first — a snapback
FLOW-LED	0.61	0.061	+0.00	+0.19	price already leaning the run's way
OPEN/NEWS	0.79	0.098	−0.01	−0.00	biggest amplitude, no directional lean

Two honest things fall out. First, **rvol is below 1.0 in every cluster** (~0.6) — these runs ignite out of quiet, so relative volume is not the trigger you'd hope for. Second, the only cluster with a clean, same-direction pre-run signature is **FLOW-LED**: a positive VWAP slope (+0.19 ATR) with price sitting on the run side of VWAP. That is a plain-English "it's already going, and it keeps going." VACUUM is the mirror — price leaning the opposite way (slope −0.53) right before it reverses, which is exactly why fading it is a trap.

THE TAPE DOES NOT LEAD — AN HONEST NULL

I wanted the aggressor tape to forecast these. It doesn't. Signing the 5-minute pre-run aggressor volume toward the direction the run eventually went: FLOW-LED leaned the right way in only **8 of 19** runs (median −43 contracts, i.e. slightly against); OPEN/NEWS in only **5 of 18** (median −971, strongly against — the news flips whoever was leaning); VACUUM 14 of 36. The tape that "drives" a FLOW-LED run is concurrent, not a

heads-up. So the pre-run edge, such as it is, lives in the price slope, not the tape.

§2.4 The L2 book (MNQ, Thu 07-09 14:17 → Fri 07-10 21:00 UTC)

Ten MNQ runs fall inside the depth window. I pulled the top-5 book in 30-second buckets across the two minutes before each and watched the far side — the side price has to eat to run. The result is a clean negative:

RUN	FAR-SIDE TOP-5 (2MIN→0)	IMBALANCE	TELL?
07-10 15:27 DN	37 → 42 → 40	+0.02 to +0.03	none — bid side grew
07-10 15:45 UP (FLOW-LED)	38 → 37 → 38	−0.05 to +0.05	none — asks flat
07-09 14:21 UP (caught)	35 → 33 → 31	−0.02 to +0.05	none material
07-10 15:05 UP (VACUUM)	34 → 37 → 35	−0.03 flat	none

Across every run in the window the top-5 sat at a stable **33-42 contracts a side** with imbalance pinned near zero (± 0.06), and the far side did **not** thin out ahead of the move, nor did imbalance lean the way the run went. The vacuum happens inside the run — levels are replaced faster than they're worth showing at 250ms — not visibly before it. MNQ's top-5 is simply too shallow and too fast to telegraph at this granularity. **No repeatable L2 tell to flag for next week**, with one honest caveat: we only capture five levels; a real depletion could be sitting deeper than we can see.

§2.5 The honest money left on the table

Full-fantasy ceiling — every one of the 76 runs, whole move, one lot — is **\$16,673**. That is hindsight and nobody catches it. The grounded numbers:

	RUNS	CEILING \$	HONEST REAL \$
Caught (aligned)	17	3,519	+993
Fought (opposed)	10	2,664	+15
Sat out	49	10,489	0

When we aligned we converted **+\$993 of a \$3,519 ceiling — about 28%**, which is a fair haul for a 1-ATR-stop shape.

Fighting the runs netted a near-wash (+\$15) only because the exit spine cut the fades fast; the entries were wrong, the exits saved us. And the whole MNQ shadow book this week was — **\$353 real across 491 trades** (versus a rosy +\$2,492 optimistic) — so the runs made money and the chop gave it all back and more. Where did the chop bleed go? **passive_chop — \$806** and **tw_ride_strength SHORT —\$527** — fades. Where did the run money come from? **tw_mnq_thrust_cont +\$991**, momentum_shadow +\$495, orb +\$318. We already own the run-catching machine.

§2.6 How we grab more of them from Monday

FLOW-LED continuation — the one clean sit-out. **SHADOW-FIRST**

This is the highest-value hypothesis and we already have the exact tool. The pre-run signature is plain: VWAP slope leaning the run direction ($\geq$ about +0.15 ATR) with price on the run side of VWAP and ATR expanding. The trigger: on a fresh thrust that extends a further ~ 0.5 ATR in the slope direction, go with it. The exit is not new

to build — it is **tw_mnq_thrust_cont**, whose money this week comes almost entirely from its RATCHET exits (LONG +\$386 on 6 trades, SHORT +\$352 on 4, holds ~12-17 minutes — exactly a run's length). The ask is simply to widen that shadow's arming so it fires on more of the 15 FLOW-LED runs we watched go by, and observe whether the extra fills hold their +28% conversion. Shadow-first, not arm-now — loosening a trigger always drags in false thrusts, and we measure that before it touches size.

OPEN/NEWS — biggest size, but a coin-flip to pre-position.

SHADOW-FIRST

These 18 runs hold the fattest ceiling (\$4,936) but the direction is genuinely unforecastable from what we capture — the pre-run tape leaned the wrong way in 13 of 18. So the honest play is not a prediction, it is a reaction: after the 13:30 UTC data / cash open, take the first 1-ATR thrust in its own direction with the same ratchet exit — a break-and-go. The missing piece is an event-clock gate (a "we are in the 13:30 window" flag) which we don't cleanly have in the funnel; name that as the data to add. Shadow-first.

VACUUM — not a new entry, a veto. STOP FADING

The 36 snapbacks move against the prior lean by construction — unforecastable in direction from what we capture. The win here is defensive: we fought seven of them, and the passive_chop / ride-strength-short fades that bled −\$806 this week are largely these. Don't build to catch VACUUM; build the slope-veto that stops us fading into it (this ties straight to the reversion autopsy). No new arm.

MNQ TAKEAWAY

The money this week was in the runs (+\$993 when we aligned) and the bleed was fading between them (−\$806 chop). We showed up for 17 of 76 and banked about a quarter of what we caught. The single clean, repeatable pre-run signature we keep no-showing is FLOW-LED continuation — price already leaning,


ATR expanding — and we already own the shape that prints it (thrust_cont + ratchet). So the move is **widen that shadow, don't build new**. The L2 top-5 book does not telegraph direction on MNQ, the aggressor tape does not lead, and rvol is quiet before runs — the only honest heads-up is the VWAP slope. OPEN/NEWS is the fattest prize but a coin-flip to pre-position: react to the first thrust, never predict it.

§3 The biggest runs on MES this week — did we show up, and what did they look like?

MES handed us **63 real runs** this week — a 15-minute move of at least 9.4 points ($1.5\times$ the contract's typical 6.2-point range). We fired into **10 of them** and sat out the other **53**, including the four biggest and the entire Thursday-Friday stretch where we can actually see the order book. When we did show up, the honest money was a shade positive (**+\$65.70** across 17 sim entries). Meanwhile the whole MES sim book bled **−\$1,044** real fighting the runs it stood in front of. This chapter walks every one of the 63, then asks the owner's question: which of these could we have grabbed, and how do we grab them from Monday.

§3.1 The week in one picture — 63 runs, biggest at the extremes

Each bar is one run: **green = up**, **red = down**, height = points travelled. A **gold dot** means a sim fired in that run's direction within ± 5 min. Notice the gold dots all cluster on the left — Monday and Tuesday. From Wednesday's 49-point flush (the tall red spike) onward it's a wall of bars with no dots: we watched the back half of the week go by.



§3.2 1 • The full census — every run, not a top-8

All 63 runs the scorecard found (footprint_scorecard.py missed --symbol MES --top 100), oldest first. **Move** is points; **flow** is net aggressor contracts around the run (positive = buyers); **amp%** is the raw amplitude of the burst; **cluster** is the cause fingerprint. Rows shaded green are the ones a sim actually fired into. The **L2** tag marks the 12 runs that fall inside the order-book capture window (Thu 14:17 → Fri 21:00 UTC) — every one of those was a sit-out. I re-priced two straight off the 5-second bars to check the scorecard isn't inventing size: Wednesday 07-08 08:15 went 7546.00 → 7496.75 = **−49.25pt** (census says −49.2, spot on), and 07-06 17:28 went 7588.75 → 7575.00 = −13.75pt (census −11.2, measured start-to-close). 80,355 5-second bars back the week, so the tape is dense, not sparse.

RUN START (UTC)	DIR	MOVE	US?	FLOW	AMP%	CLUSTER
07-06 00:16	DN	-11.2	IN	+590	0.099	VACUUM
07-06 13:19	UP	+16.5	IN	-47	0.076	OPEN/ NEWS
07-06 13:36	UP	+16.8	IN	-1433	0.189	OPEN/ NEWS
	DN	-9.8	IN	+879	0.152	

07-06 13:51						OPEN/ NEWS
07-06 14:06	UP	+11.8	sat out	+135	0.132	OPEN/ NEWS
07-06 15:17	UP	+11.2	sat out	-520	0.096	VACUUM
07-06 17:28	DN	-11.2	IN	-170	0.033	FLOW-LED
07-06 18:01	UP	+11.0	IN	+146	0.056	FLOW-LED
07-07 00:25	DN	-10.2	sat out	+40	0.046	VACUUM
07-07 11:55	UP	+10.0	sat out	+12	0.066	UNCLASS
07-07 13:20	DN	-10.8	sat out	-102	0.063	OPEN/ NEWS
07-07 13:44	DN	-15.0	sat out	+690	0.152	OPEN/ NEWS
07-07 14:03	DN	-20.0	IN	+79	0.132	OPEN/ NEWS
07-07 14:20	DN	-11.5	sat out	+1064	0.165	OPEN/ NEWS
07-07 14:42	UP	+12.5	sat out	-2205	0.156	OPEN/ NEWS
07-07 15:40	UP	+14.8	sat out	-1533	0.106	VACUUM
07-07 16:08	UP	+15.0	sat out	-425	0.099	VACUUM
07-07 18:46	DN	-16.8	IN	-139	0.056	FLOW-LED

07-07 19:40	UP	+13.2	sat out	-152	0.066	VACUUM
07-07 20:46	DN	-12.2	sat out	-34	0.020	UNCLASS
07-08 00:01	UP	+14.2	sat out	-523	0.096	VACUUM
07-08 00:35	UP	+10.2	sat out	+71	0.060	FLOW-LED
07-08 07:20	UP	+11.0	IN	+18	0.066	UNCLASS
07-08 08:15	DN	-49.2	sat out	+8	0.063	VACUUM
07-08 08:40	UP	+14.0	sat out	-369	0.214	VACUUM
07-08 08:55	DN	-19.2	sat out	-27	0.087	UNCLASS
07-08 10:24	UP	+10.8	sat out	+82	0.057	FLOW-LED
07-08 11:29	UP	+18.8	sat out	+6	0.057	UNCLASS
07-08 12:00	UP	+12.8	sat out	-663	0.150	VACUUM
07-08 13:25	UP	+15.8	sat out	-217	0.090	OPEN/ NEWS
07-08 13:40	DN	-11.2	sat out	+1802	0.223	OPEN/ NEWS
07-08 14:01	DN	-15.0	sat out	+369	0.163	OPEN/ NEWS
07-08 14:35	DN	-21.5	sat out	-230	0.137	OPEN/ NEWS

07-08 15:11	DN	-22.2	sat out	+164	0.067	VACUUM
07-08 15:26	UP	+14.2	sat out	-163	0.224	VACUUM
07-08 15:51	UP	+12.8	sat out	-483	0.080	VACUUM
07-08 16:08	UP	+19.2	sat out	-160	0.110	VACUUM
07-08 16:37	UP	+20.2	sat out	-175	0.143	VACUUM
07-08 17:16	UP	+14.5	sat out	+34	0.063	UNCLASS
07-08 17:33	DN	-13.8	sat out	-320	0.080	FLOW-LED
07-08 17:57	UP	+13.0	sat out	+165	0.126	FLOW-LED
07-08 19:42	UP	+10.2	sat out	-193	0.066	VACUUM
07-08 19:57	DN	-14.0	sat out	+640	0.109	VACUUM
07-09 02:03	UP	+10.8	sat out	+62	0.063	FLOW-LED
07-09 03:18	UP	+12.2	sat out	-213	0.050	VACUUM
07-09 05:16	DN	-10.0	sat out	-18	0.146	UNCLASS
07-09 05:33	UP	+9.5	IN	-55	0.093	VACUUM
07-09 10:38	DN	-12.2	sat out	+12	0.060	VACUUM

07-09 13:17	DN	-13.0	sat out	-102	0.043	OPEN/ NEWS
07-09 13:32	UP	+26.2	sat out	-153	0.242	OPEN/ NEWS
07-09 13:52	DN	-31.5	sat out	+1750	0.136	OPEN/ NEWS
07-09 14:21 L2	UP	+27.5	sat out	-1788	0.173	OPEN/ NEWS
07-09 14:54 L2	UP	+19.8	sat out	+344	0.123	OPEN/ NEWS
07-09 16:57 L2	UP	+10.2	sat out	-294	0.043	VACUUM
07-09 23:57 L2	DN	-9.5	sat out	+107	0.046	VACUUM
07-10 09:01 L2	UP	+11.2	sat out	-469	0.096	VACUUM
07-10 13:22 L2	UP	+10.5	sat out	-33	0.026	OPEN/ NEWS
07-10 13:43 L2	UP	+11.5	sat out	-22	0.155	OPEN/ NEWS
07-10 14:18 L2	DN	-26.2	sat out	+364	0.089	OPEN/ NEWS
07-10 14:33 L2	UP	+17.0	sat out	-24193	0.677	OPEN/ NEWS
07-10 15:01 L2	UP	+21.0	sat out	+459	0.109	FLOW-LED
07-10 15:27 L2	DN	-18.2	sat out	-104	0.122	FLOW-LED
07-10 15:45 L2	UP	+14.8	sat out	-209	0.155	VACUUM

THE CLUSTER COUNTS

Sorted by how the scorecard classifies the cause of each move, and how many of each we sat out:

CLUSTER	WHAT IT MEANS	RUNS	FIRED	SAT OUT
VACUUM	price ran against the aggressor tape — book gave way	24	2	22
OPEN / NEWS	scheduled-open or headline volatility spike	22	4	18
FLOW-LED	aggressors visibly drove it	10	3	7
UNCLASS	no tape / too quiet to label	7	1	6

§3.3 2 • Participation — we showed up for 10, sat out 53

Blunt version: **10 of 63** runs had a sim fire in the right direction inside ± 5 minutes. And those 10 are almost all Monday–Tuesday. From Wednesday morning's -49 -point flush onward, we sat out essentially everything — 40-plus runs in a row, including the four biggest of the week (the -49.2 , -31.5 , $+27.5$ and -26.2). The runs weren't small or rare; we just weren't in the market with a shape pointed the right way.

Here's the honest money on the 10 we did catch — scored on `real_pnl`, never the optimistic sim number:

RUN	DIR	CLUSTER	SIMS	REAL P&L
07-06 00:16	DN	VACUUM	3	-44.23

07-06 13:19	UP	OPEN	2	-24.12
07-06 13:36	UP	OPEN	3	-52.98
07-06 13:51	DN	OPEN	1	0.00
07-06 17:28	DN	FLOW-LED	3	+111.39
07-06 18:01	UP	FLOW-LED	1	+28.38
07-07 14:03	DN	OPEN	1	-25.37
07-07 18:46	DN	FLOW-LED	1	+34.63
07-08 07:20	UP	UNCLASS	1	+22.13
07-09 05:33	UP	VACUUM	1	+15.88
Total			17	+65.70

Two things jump out. First, when we participated we came out slightly ahead — **+\$65.70** — so showing up isn't the problem. Second, look at which runs paid: all three FLOW-LED runs made money (+\$111, +\$28, +\$34 = **+\$174** between them), and the losers were all OPEN/NEWS scrambles where we caught the spike late. That's the whole tell of the chapter in one table.

THE BOOK WE'RE MEASURED ON

Zoom out from the runs to the whole MES sim book this week: **404 filled trades, -\$1,044.48 real** — against an optimistic sim print of +\$1,190, so the simulator is flattering MES by about \$2.2k. Only two shapes finished green, and both are momentum-continuation: **momentum_shadow** **+\$102.61** (47 trades, 19 wins) and **tw_mes_thrust_cont** **+\$49.13** (26 trades, 12 wins). The bleeders are all fade/chop shapes standing in front of these runs: **tw_ride_strength** -\$287, **passive_chop** -\$183, **dip_loose** -\$134. So the desk's MES loss this week is the story

of these runs — we bled fading them and only banked when we rode them.

§3.4 3 • The footprint — what the 10 minutes before each cluster looked like

For 62 of the 63 runs there's a live funnel row in the 10 minutes before the move, so I could average the pre-run signature per cluster (atr_pct, rvol, VWAP distance, order-flow imbalance):

CLUSTER	ATR_PCT	RVOL	VWAP DIST	OFI	PLAIN READ OF THE RUN-UP
VACUUM	0.037	0.60	−0.01	−24	quiet, below-average volume coil sitting on VWAP, then snaps
OPEN/ NEWS	0.050	0.83	−0.02	−38	hottest atr going in — the market already knows a print is coming
FLOW-LED	0.037	0.56	+0.02	−45	compressed and quiet, then a one-sided aggressor burst lights it
UNCLASS	0.032	0.70	−0.02	+8	the quietest, most featureless — genuine noise

The uncomfortable truth is that **all four clusters come out of an rvol below 1.0** — below-average volume. These runs erupt from compression, not from a visible build-up. That's a real signature, but it's directionless: quiet coils are everywhere and most of them don't

go anywhere. There's no clean "rvol was building for ten minutes" tell that separates a coil that fires from a coil that dozes.

So I pushed on the one cluster that both pays when caught and has an aggressor tape to look at — **FLOW-LED**. Do the aggressors that "drive" it show up before the run, or only during it? I summed net aggressor contracts in the two minutes before all ten FLOW-LED runs (tape covers 07-05 on):

THE HARD FINDING

The pre-run tape agreed with the run's eventual direction in only **4 of 10** FLOW-LED runs — worse than a coin flip. The 07-10 15:01 up run had **-789** net (heavy selling) in the two minutes before it; the 15:27 down run had **+1,399** net (heavy buying) before it. In other words the aggressor flow that "led" these runs is coincident with the run itself — the flow is the run. It is not a two-minutes-early warning you could act on. The scorecard's "FLOW-LED" label describes the cause once it's happening, not a forecastable set-up.

Which cluster do we sit out with a genuinely clean pre-run signature? Honestly — **none of them, on the entry side**. What's clean is the after-the-fact pattern: the only money we made came from being already long/short in the direction of an aggressor burst. That points the new-entry work at a reactive trigger, not a predictive one. More on that in section 6.

§3.5 4 • The L2 order book — is the vacuum visible before it opens?

Twelve census runs fall inside the depth capture window (Thu 07-09 14:17 → Fri 07-10 21:00 UTC, ~197k MES snapshots at ~250ms). All twelve were sit-outs, so this is exactly the set the owner wants a book read on. The question: in the 1-2 minutes before each run, did the far

side of the book thin out / vanish (a vacuum you can see coming), and did imbalance lean the way the run went?

I measured mean book imbalance in the -90s to -30s window before each run and asked whether it leaned the run's direction:

THE L2 VERDICT — THIN AND NOISY THIS WEEK

Pre-run book lean matched the run direction in only **3 of 9** decisive runs (three were flat) — again worse than a coin flip. Two concrete examples with real sizes:

07-10 15:27, -18.2 down (FLOW): at -20s the book was **221 bid vs 394 ask** across five levels (imbalance -0.28) — sellers genuinely stacked, and price fell. But the -90s→-30s window it grew out of was actually bid-lean (+0.055). The "tell" only appears in the final 20 seconds, i.e. coincident with the break, not ahead of it.

07-10 15:45, +14.8 up (VACUUM): the book was ask-heavy going in (**309 ask vs 226 bid**, imbalance -0.16) and price rose into it — the displayed asks got lifted/absorbed rather than holding. The book pointed the wrong way. That's the literal definition of the VACUUM cluster: displayed liquidity that lies.

So the clear answer for the owner: **there is no clean, repeatable L2 lead on MES this week.** Two days of top-5 depth at 250ms is too noisy, and where there is a signal it's coincident with the break or outright contrarian. I would not arm anything off book imbalance yet. The one shape worth watching next week: last-20-second far-side stacking on down-flushes was directionally right in the couple I hand-checked — but that's a coincident confirmation, not a predictor, and needs a fuller window before it's more than a hunch.

§3.6 5 · The honest money left on the table

Add up the directional range of all 63 runs and it's **962 points** — a hindsight ceiling of about **\$4,811 on a single contract** if you'd nailed every start and sold every extreme. That number is a fantasy and I want to be plain about why: most of it is either the directionless VACUUM coil (24 runs, 22 sat out) that you can't forecast, or OPEN/NEWS whipsaws like the 14:33 monster (a −24,193-contract sell wave that flushed then V-reversed +17) that chew up anything mechanical.

The capturable slice is much smaller and we already know its shape: the reactive momentum book. `momentum_shadow` and `tw_mes_thrust_cont` between them netted **+\$152** real this week riding exactly these bursts, and the runs we participated in netted **+\$65.70**. Realistically the extra we left on the table is a few hundred dollars of FLOW-LED and post-flush continuation — not the \$4.8k headline. The gap isn't "we missed a fortune," it's "we only had a shape pointed the right way one day in five."

§3.7 6 · How we grab them from Monday

The evidence says stop trying to predict these and get better at reacting to them. Here's the concrete plan per cluster.

§3.8 `SHADOW-FIRST` FLOW-LED & the momentum burst — widen what already works

Why: it's the only cluster that paid every time we caught it (+\$174 across three), and the two green MES shapes already live here. Crucially, the winners fire on elevated rvol — `momentum_shadow` entered at mean rvol **1.23** and `thrust_cont` at **1.85**, both with OFI around **+59**. They wait for volume to

already be firing. That's exactly right, and it's why they're green while the fade shapes bleed.

The trigger: when `rvol` crosses roughly ≥ 1.2 and a 30-60s net-aggressor burst in `trade_tick` exceeds a size threshold in the break direction and price breaks the prior 5-second range — go **with** the break. Direction = the sign of the aggressor burst.

Exit: the chandelier/thrust exit already wired to `tw_mes_thrust_cont` — no new exit to build.

Action: extend the existing thrust trigger to fire on more of the FLOW-LED and post-flush continuation runs we sat out (7 of 10 FLOW-LED missed). Keep it shadow — both shapes are green but on ~30-50 trades of one in-sample week, nowhere near the promotion-ladder gate. Watch it two more weeks across a trend regime before it goes near live-1.

§3.9 **DON'T CHASE** VACUUM — unforecastable from what we capture

24 runs, the biggest single cluster, and both the tape (moves against it by definition) and the top-5 book (3/9 lean agreement) point the wrong way beforehand. This is a real edge for someone, but not with what we record. To have any shot you'd need **full-depth book beyond level 5 plus resting-order cancellation / iceberg detection** — the rate at which displayed size is pulled just before the break — which 250ms top-5 snapshots don't show. Until we capture that, the honest move is to stop fading into these (kill the `passive_chop / ride_strength` fades that lost -\$470 combined here), not to try to catch them.

§3.10 REACT, DON'T PREDICT OPEN / NEWS — the spike isn't tradeable, the second leg might be

22 runs, hottest pre-run atr (0.050), and the cluster where our participation actually lost money (we caught the 13:19/13:36 spikes late for −\$77). Predicting the first move off a scheduled print is a coin-flip and the 14:33 reversal proves the danger. But note `tw_mes_block_us_open`, which stands aside through the open, still lost −\$113 — so blindly blocking isn't free either. The tractable version: don't touch the first thrust bar; let the open print, then apply the same reactive momentum trigger to the continuation once `rvol` and flow agree. Shadow-only, same trigger as FLOW-LED above.

PLAIN TAKEAWAY — MES

MES gave us 63 real runs and roughly \$4.8k of hindsight range this week; we showed up for 10 and the sim book still lost −\$1,044 fading the other 53. Neither the aggressor tape nor the level-5 order book forecasts these one week in — the flow and the vacuum are both coincident with the break, not ahead of it. The only thing that paid is being already positioned in the direction of a live thrust, which is exactly what the two green shadows (`momentum_shadow`, `tw_mes_thrust_cont`, +\$152 together) already do. So the Monday move is not to build a predictor — it's to widen the reactive thrust trigger onto the FLOW-LED and post-open-continuation runs we keep sitting out, kill the fades that bleed in front of them, and leave the VACUUM cluster alone until we capture deeper book. Trade less, react faster.

§4 The biggest runs on MGC this week — did we show up?

Gold gave us **57 real runs** this week — 15-minute moves of at least 11.1 points (\$110 a contract), 1.5× MGC's typical quarter-hour range of 7.4pt. We were pointed the right way on **24** of them, sat on our hands for **33**, and of those sat-out ones we actively **fought 8** — fired a fade straight into the move and got run over. On the sims that were aligned, the honest money during runs was **+\$756**; on the sims caught leaning the wrong way it was **−\$273**. The hindsight ceiling of all this run-money is about **\$9,060 a contract**, of which we let **\$5,390** walk past untouched.


THE ONE-LINE READ

MGC runs are mostly **VACUUM** — price ripping the opposite way to the tape, which means the tape can't tell us the direction in advance, and this week the L2 book was too thin to see the vacuum coming either. The one clean, repeatable signature we keep sitting out is the **US-open / data window (13:00-15:00 UTC)** volatility burst. That's the shadow-first hypothesis. The bigger immediate win is defensive: stop letting the fade sims walk into these things.

§4.1 1 · The full census — every run, not a top-8

Here is the whole week, all 57 runs in time order. Green = a sim was pointed the right way (**took**), gold = nobody fired (**sat out**), red = we fired the wrong way into it (**fought**). "\$" is per one contract (MGC

= \$10/point). "aligned \$" is the honest real_pnl of the sims that were pointed the right way — note that being right on direction is not the same as capturing the move (see 07-10 11:59, where we were "in" and still lost \$88). "flow" is net aggressor tape in the 60s before (buys–sells, MGC tape only since 07-05). "amp" is the pre-run 5-min range as a % of price.



Each bar is one run, in time order Mon→Fri; up-bars are up-runs, down-bars down-runs, height is the size. The tallest red bar in the middle is the week's biggest move — the 07-08 08:10 gold dump of −36.8pt (\$368) that we fought. The wall of gold on the right-hand third is Thursday–Friday's open-window runs we mostly watched go by.

RUN (UTC)	DIR	MOVE (PT)	\$/CTR	US	SIMS±5M	ALIGNED \$	FLOW	AMP%	CLUST
07-06 00:27	UP	+13.6	+136	took	4	-35	-110	0.198	VACUU
07-06 01:01	UP	+14.4	+144	took	9	+8	-10	0.389	VACUU
07-06 01:23	UP	+18.4	+184	took	4	+0	-63	0.224	VACUU
07-06 01:56	DN	-15.0	-150	sat out	0	—	-21	0.124	UNCLA
	DN	-16.4	-164	sat out	0	—	+52	0.157	VACUU

07-06 02:21									
07-06 07:29	UP	+11.5	+115	fought	5	—	-6	0.084	VACUU
07-06 12:20	DN	-15.7	-157	took	7	+0		0.065	UNCLA
07-06 14:52	UP	+11.6	+116	took	3	+3	+23	0.14	OPEN/ NEWS
07-07 00:33	DN	-12.3	-123	sat out	0	—	+12	0.127	VACUU
07-07 01:07	DN	-13.3	-133	sat out	0	—	+29	0.118	VACUU
07-07 01:28	UP	+12.7	+127	fought	6	—	-30	0.215	VACUU
07-07 02:37	DN	-12.0	-120	sat out	0	—	-70	0.144	FLOW-
07-07 03:57	UP	+12.8	+128	took	5	+77	-3	0.056	VACUU
07-07 07:59	UP	+13.5	+135	took	5	-0	+16	0.097	UNCLA
07-07 09:29	UP	+11.4	+114	sat out	0	—	+15	0.177	UNCLA
07-07 11:12	UP	+12.1	+121	sat out	0	—	-3	0.043	VACUU
07-07 11:59	UP	+18.9	+189	took	7	+135	+50	0.125	UNCLA
07-07 12:45	UP	+11.3	+113	took	1	-0	+4	0.139	UNCLA
07-07 13:29	UP	+12.2	+122	sat out	0	—	-19	0.134	OPEN/ NEWS

07-07 13:44	DN	-21.5	-215	sat out	0	—	-178	0.301	OPEN/ NEWS
07-07 18:47	DN	-23.2	-232	took	3	+52	-8	0.07	UNCLA
07-07 19:03	DN	-25.7	-257	sat out	0	—	+25	0.303	VACUU
07-07 19:20	UP	+15.3	+153	sat out	0	—	+92	0.299	FLOW-
07-08 00:25	UP	+19.4	+194	took	5	+144	+6	0.143	UNCLA
07-08 00:50	DN	-19.9	-199	sat out	0	—	+56	0.123	VACUU
07-08 01:35	UP	+12.4	+124	fought	10	—	-9	0.207	VACUU
07-08 02:57	UP	+15.1	+151	took	12	+94	-61	0.097	VACUU
07-08 08:10	DN	-36.8	-368	fought	2	—	+61	0.082	VACUU
07-08 08:27	DN	-27.5	-275	sat out	0	—	+29	0.295	VACUU
07-08 08:48	DN	-24.1	-241	sat out	0	—	+15	0.338	VACUU
07-08 09:28	DN	-11.6	-116	took	5	+7	+51	0.145	VACUU
07-08 10:09	UP	+13.7	+137	sat out	0	—	+33	0.17	UNCLA
07-08 11:31	UP	+23.4	+234	took	5	+1	+11	0.074	UNCLA
07-08 12:04	UP	+14.2	+142	took	6	+56	-31	0.189	VACUU

07-08 12:37	UP	+14.5	+145	took	3	+100	-4	0.137	VACUU
07-08 12:55	DN	-11.7	-117	took	2	-0	-2	0.11	UNCLA
07-08 14:01	DN	-15.2	-152	fought	3	—	-2	0.243	OPEN/ NEWS
07-08 14:29	DN	-24.7	-247	took	4	+2	-17	0.172	OPEN/ NEWS
07-08 15:10	DN	-12.7	-127	took	8	+95	+34	0.111	VACUU
07-08 16:38	UP	+20.8	+208	sat out	0	—	+21	0.173	UNCLA
07-08 18:34	UP	+21.5	+215	sat out	0	—	-35	0.11	VACUU
07-09 02:27	DN	-11.3	-113	sat out	0	—	+19	0.125	VACUU
07-09 04:47	UP	+13.0	+130	took	5	+24	-19	0.103	VACUU
07-09 05:46	UP	+21.9	+219	sat out	0	—	+19	0.091	UNCLA
07-09 06:11	UP	+12.7	+127	sat out	0	—	-15	0.102	VACUU
07-09 11:54	UP	+11.7	+117	took	1	-19	-22	0.083	VACUU
07-09 12:25	UP	+13.5	+135	took	1	+70	-37	0.087	VACUU
07-09 13:02	UP	+14.7	+147	took	6	+32	+2	0.199	OPEN/ NEWS
07-09 13:52	DN	-12.1	-121	sat out	0	—	+315	0.131	OPEN/ NEWS

07-09 14:51	UP	+11.4	+114	sat out	0	—	+8	0.155	OPEN/ NEWS
07-10 01:30	DN	-13.6	-136	fought	4	—	-34	0.085	UNCLA
07-10 03:09	DN	-11.2	-112	fought	5	—	-41	0.164	UNCLA
07-10 11:59	UP	+13.6	+136	took	8	-88	+5	0.095	UNCLA
07-10 13:14	DN	-17.7	-177	fought	5	—	+5	0.109	OPEN/ NEWS
07-10 13:35	UP	+14.2	+142	sat out	0	—	+18	0.217	OPEN/ NEWS
07-10 14:19	DN	-13.3	-133	sat out	0	—	+142	0.202	OPEN/ NEWS
07-10 14:34	UP	+16.3	+163	sat out	0	—	+23	0.866	OPEN/ NEWS

CLUSTER COUNTS (ALL 57)

VACUUM (ripped against the tape) — 27 runs, sat out 15.

UNCLASSIFIED (quiet, no tape lead) — 16 runs, sat out 7.

OPEN / NEWS (13:00-15:00 UTC burst) — 12 runs, sat out 9.

FLOW-LED (tape drove it) — 2 runs, sat out both. (The 7-day census is 59, adding two runs late on 07-05.)

I re-priced three of these straight off the 5s bars to make sure the machine isn't dreaming. The week's biggest, **07-08 08:10**, went 4131.0 → 4093.3 in the quarter hour — a genuine −37.7pt gold dump, and yes we had a dip_loose long sitting in it. The 07-07 13:44 open-drop printed 4190.7 → 4169.2 (−21.5), and the Friday-afternoon 14:34 pop ran 4100.1 → 4115.9 (+16.3). All three check out.

§4.2 2 · Did we show up? Blunt version

24 caught · 33 sat out · 8 of those fought. On the 30 aligned sim-fills that sat inside a run, the honest money was **+\$756.40**. On the 18 fills that were pointed the wrong way — a fade, or a mis-fired momentum short, walking into the move — it was **-\$273.16**. So during the runs themselves the shadow book netted about **+\$483**, which is the good news; the bad news is that most of MGC's weekly reversion bleed happens between these runs in the grind, not inside them.

Two honest caveats the raw "24 caught" hides. First, **catching a run isn't capturing it**: on 07-10 11:59 eight sims were correctly long into a +13.6 pop and still handed back **-\$88** real, chopped out before the move or stopped on the wobble. Second, the fights are ugly and repetitive — the same fade names (pullback_loose, orb_shadow, momentum_shadow firing short) kept stepping in front of gold: **-\$46** short into the 01:28 pop, **-\$41** long into the 13:14 drop, twin **-\$37s** into the 01:35 rip. It's exactly the "every fade walked into the same gold move and got run over" pattern.

For context on who does the catching, over the full week on MGC the momentum-shaped sims are the only ones green on honest money — tw_mgc_coil_break_long **+\$137**, tw_mgc_momo_thrust **+\$127**, orb_shadow **+\$78**, pullback_quiet **+\$47** — while the fade book bleeds hard: passive_chop **-\$1,125**, passive_mgc_hours **-\$542**, pullback_loose **-\$351**, dip_loose **-\$293**.

§4.3 3 · The footprint, clustered — what did the 10 minutes before look like?

For each cluster I took the median funnel reading at the run start, split by whether we caught it or sat it out. The question that matters:

is there a cluster we keep sitting out that has a clean, separable pre-run signature — because that's the one worth building an entry for.

CLUSTER / DID-WE	N	ATR%	RVOL	VWAP DIST%	OFI	PRE- AMP%
VACUUM · took	12	0.076	0.63	-0.04	-7.5	0.124
VACUUM · sat out	15	0.071	0.45	-0.06	+15.0	0.125
UNCLASS · took	9	0.048	0.38	+0.01	+5.5	0.097
UNCLASS · sat out	7	0.054	0.92	-0.04	-12.5	0.164
OPEN/NEWS · took	3	0.068	0.20	-0.03	-16.0	0.172
OPEN/NEWS · sat out	9	0.081	0.50	-0.05	+15.0	0.202
FLOW-LED · sat out	2	0.079	0.51	-0.19	+37.5	0.221

VACUUM — the biggest cluster, and the most frustrating. The 15 we sat out look statistically identical to the 12 we caught: atr% 0.071 vs 0.076, pre-amp 0.125 vs 0.124. There is no pre-run tell separating the vacuum runs we missed from the ones we hit. Worse, by definition these move against the tape — the 60s aggressor flow points the wrong way (e.g. 07-06 00:27, tape −110 into a +13.6 up-run). So the one microstructure signal we have actively lies about direction here. Plainly: **VACUUM runs are not forecastable from anything we currently capture.**

OPEN / NEWS — the clean one. The nine we sat out carry the hottest pre-run signature of any cluster: the highest atr% (0.081), a lively rvol (0.50), and the biggest pre-run amplitude (0.202, versus ~0.125 for the vacuum runs). And the signature is dead simple to detect because it's a clock: they all cluster in the 13:00–15:00 UTC US-cash-open / 10:00-ET-data window. We know when volatility is about to arrive; we just haven't been firing into it. The catch is direction — these split up and down (07-08 14:29 −24.7, 07-10 14:34

+16.3), so this is a "stand ready and go with the break," not a "predict the way" edge.

FLOW-LED — cleanest direction, thinnest evidence. Only two runs all week, both sat out, and both had strong aligned tape building before the move (07-07 19:20 up-run with tape +92, 07-07 02:37 down-run with tape −70). When the aggressor tape genuinely leads, it points the right way — but $n=2$ is a rumour, not a finding.

§4.4 4 • The L2 order book — could we see the vacuum coming?

Eight of the runs fall inside the MGC depth window (Thu 07-09 14:43 → Fri 07-10 21:00 UTC, ~189k snapshots). For each I looked at the 90 seconds before: did the far side of the book — the side price was about to run into — thin out ahead of time (a vacuum you could see forming), or did imbalance lean the way the run went?

RUN	DIR	MOVE	IMBALANCE (MED)	FAR SIDE THINNING?	IMBALANCE LEAN
07-09 14:51	UP	+11.4	-0.02	no	flat
07-10 01:30	DN	-13.6	-0.24	no	with run
07-10 03:09	DN	-11.2	+0.07	no	against
07-10 11:59	UP	+13.6	0.00	no	flat
07-10 13:14	DN	-17.7	0.00	no	flat
07-10 13:35	UP	+14.2	-0.02	no	flat
07-10 14:19	DN	-13.3	-0.12	no	with run

07-10
14:34

UP

+16.3

+0.04

no

with run

The honest answer is **no clean tell**. On every one of the eight, the far side did not vacate ahead of the move — the book five-deep held roughly the same size right up to the run (before the 13:14 drop the bid stack sat at ~20 lots the whole 90s and simply got eaten — absorption, not vacuum). Imbalance leaned with the run three times, against once, and was flat the other four — a coin-flip. And the raw sizes are the real problem: MGC's top of book is **1-4 lots**, five-deep only ~15-25, and imbalance flickers ± 0.3 every few hundred milliseconds. There's a faint hint that the two clean down-news runs (01:30 at imbalance -0.24 , 14:19 at -0.12) carried a persistently ask-heavy book beforehand, but at this size and noise it is not something I'd fire a trade on.

L2 VERDICT — MGC

Too thin to trade off. The gold book this window shows **absorption, not a visible vacuum** — the far side gets eaten in place rather than pulling away first, so the "see the run coming in the ladder" idea doesn't hold on MGC the way it might on the deeper MNQ book. Keep capturing; don't build an MGC L2 trigger yet.

§4.5 5 • The honest money left on the table

Add up every run and the hindsight ceiling — a perfect one-contract fill from run-start to run-end — is about **\$9,060**. We were aligned on runs worth roughly \$3,670 of that, and the sat-out runs account for **\$5,390**. But that ceiling is a fantasy number — no strategy fills the exact top and bottom. The capturable figure — what an actual sim shape banked when it was pointed the right way — was **+\$756** gross, dragged to about **+\$483** net once you subtract the $-\$273$ the wrong-

way fills bled during those same runs. So of MGC's big-run money this week, a realistic "a sim could have taken it" number is low-hundreds of dollars a contract, not thousands. The gap between \$483 realised and \$5,390 sat-out is mostly the **VACUUM** runs we can't forecast and the **OPEN/NEWS** runs we chose not to be ready for.

§4.6 6 • How we grab them from Monday

OPEN / NEWS window — the one arm-able idea. shadow-first We sit out 9 of 12 of these and they carry the hottest, most predictable pre-run signature (a clock plus hot vol). Concrete trigger: between 13:00 and 15:00 UTC, when $\text{atr\%} > 0.08$ and $\text{rvol} > 0.4$, and price makes a 5s break of the prior 5-minute high/low, go WITH the break — long on an up-break, short on a down-break. Direction comes from the break, not a forecast. Exit: hand it to the existing `tw_mgc_coil_break` chandelier, the sim already banking +\$137 real on MGC. The build is small: `tw_mgc_coil_break_long` is long-only today — **add the short leg and widen its active hours to the open window**, then shadow it for a couple of weeks before arming. This is the highest-value new-entry hypothesis on gold.

FLOW-LED — watch, don't build yet. shadow-first The rule would be clean — when net aggressor tape over 60s exceeds ± 50 lots and price is already moving that way, join it with the coil-break exit — but two examples all week is not enough to commit a line of code to. Log it as a shadow counter and revisit when we have 8-10 instances across a few weeks.

VACUUM — genuinely unforecastable, so play defence. no entry These are the biggest bucket (27 runs) and we have nothing in the funnel, the tape or the L2 book that separates the ones we caught from the ones we missed — the tape points the wrong way and the book gets absorbed in place. To forecast them we'd need data we don't capture: full-depth iceberg / hidden-liquidity detection to spot the big passive seller before it lifts. Until then the win is entirely defensive — **stop the fade sims fading into them**. The 8 fought

runs cost -\$273 this week, and the $|\text{slope}| < 0.15$ veto already on the table from the reversion autopsy would have kept pullback_loose / dip_loose out of the 08:10 dump and the 13:14 drop. That's the cheapest MGC money on offer.

PLAIN TAKEAWAY — MGC

Gold ran 57 times this week; we were right on 24 and banked a net ~\$483 of a realistic few-hundred-dollar capturable pot. Most of the runs are **vacuum moves against the tape that we simply cannot see coming** — the funnel, the tape and the thin gold book all stay quiet or point the wrong way. The one clean, repeatable thing we keep skipping is the **US-open / data-window volatility burst**: build it as a two-sided break-and-go on the coil-break exit and shadow it. And whatever else, quit fading gold when it's trending — the fights, not the missed runs, are the avoidable loss.

§5 The biggest runs on M2K this week — did we show up?

M2K (micro Russell, \$5 a point) printed **64 big runs this week** — 15-minute moves of 5.2 points or more, which is 1.5× its normal 3.5-point quarter-hour range. We showed up in the right direction on 17 of them, faded 7, and simply weren't there for 40. When a sim did line up with a run the honest money was +**\$285**; but M2K as a whole still lost **−\$252** on the week, because a fader book gave it all back on the small stuff. The runs aren't the problem. Missing them — and fading into them — is.


THE ONE-LINE STORY

M2K's real money lives in the US-cash-open hour. Every one of the 23 OPEN/NEWS runs landed between 13:05 and 14:54 UTC, they carry the fattest amplitude of the week, and we sat out 13 of them — including the four biggest sit-outs on the sheet. But the open goes both ways (13 down, 10 up) and whipsaws inside the same morning, so this is a break-to-trade, never a pre-position. Our one open sim tried to pre-position three times and got stopped all three.

§5.1 1 — The full census: all 64 runs

Straight from `footprint_scorecard.py` missed `--symbol M2K --top 100`. The scorer splits every run into one of two cause-clusters: **OPEN/NEWS** (the move erupts in the cash-open window) or **quiet / unclassified** (a move with no tape signature — and M2K has no

aggressor tape at all, so almost everything overnight lands here). There are no VACUUM or FLOW-LED clusters for M2K the way MNQ has them, precisely because we don't capture M2K's tape or book. Counts: **OPEN/NEWS 23 • quiet 41.**



Each bar is one run, in time order Mon→Fri; height = size of the move, up-bars above the line, down-bars below. Green = we caught it (a sim fired the right way within ± 5 min), red = we sat it out, gold = we fought it (faded into it). Notice the tallest bars are mostly red.

Two I re-priced by hand from the 5-second bars to make sure the scorer isn't dreaming: the **07-07 13:41** down-run measured -21.1 pt open-to-close over its 15 min (scorer said -21.0); the **07-08 08:15** dump measured -28.8 pt low-to-open (scorer said -26.1). Both real, both honest.

RUN (UTC)	DIR & MOVE (PT)	AMP%	CLUSTER	US	REAL \$
07-06 00:16	▼ -6.0	0.119	quiet/uncl.	caught	-7
07-06 13:05	▼ -5.7	0.133	OPEN/NEWS	sat out	—
07-06 13:29	▲ +20.3	0.129	OPEN/NEWS	caught	+104
	▼ -12.7	0.174		caught	+1

07-06 13:51			OPEN/ NEWS		
07-06 14:18	▲ +8.0	0.155	OPEN/ NEWS	caught	-18
07-06 14:42	▲ +8.8	0.198	OPEN/ NEWS	fought	—
07-06 15:36	▲ +5.3	0.132	quiet/uncl.	sat out	—
07-06 17:28	▼ -7.2	0.049	quiet/uncl.	caught	+21
07-06 19:17	▲ +5.3	0.063	quiet/uncl.	caught	+28
07-06 19:55	▲ +6.7	0.132	quiet/uncl.	fought	—
07-07 11:02	▲ +6.2	0.082	quiet/uncl.	sat out	—
07-07 13:25	▼ -7.1	0.119	OPEN/ NEWS	caught	-13
07-07 13:41	▼ -21.0	0.212	OPEN/ NEWS	caught	+51
07-07 14:00	▼ -12.3	0.253	OPEN/ NEWS	sat out	—
07-07 14:27	▼ -7.7	0.260	OPEN/ NEWS	sat out	—
07-07 14:42	▲ +12.8	0.237	OPEN/ NEWS	sat out	—
07-07 15:11	▼ -7.1	0.253	quiet/uncl.	sat out	—
07-07 15:51	▲ +5.5	0.167	quiet/uncl.	sat out	—

07-07 16:09	▲ +8.7	0.190	quiet/uncl.	sat out	—
07-07 18:39	▼ -8.3	0.103	quiet/uncl.	sat out	—
07-07 18:57	▼ -7.7	0.110	quiet/uncl.	sat out	—
07-07 19:40	▲ +8.2	0.104	quiet/uncl.	sat out	—
07-07 20:46	▼ -6.0	0.040	quiet/uncl.	sat out	—
07-08 00:36	▲ +7.3	0.077	quiet/uncl.	caught	+14
07-08 06:52	▼ -8.0	0.147	quiet/uncl.	caught	-21
07-08 07:20	▲ +8.8	0.091	quiet/uncl.	caught	+35
07-08 08:15	▼ -26.1	0.070	quiet/uncl.	caught	+3
07-08 08:40	▲ +7.8	0.257	quiet/uncl.	sat out	—
07-08 08:55	▼ -10.4	0.165	quiet/uncl.	sat out	—
07-08 10:23	▲ +7.4	0.105	quiet/uncl.	sat out	—
07-08 11:29	▲ +12.5	0.064	quiet/uncl.	caught	+13
07-08 11:58	▲ +9.0	0.155	quiet/uncl.	caught	+2
07-08 12:59	▲ +6.3	0.249	quiet/uncl.	sat out	—

07-08 13:25	▲ +9.2	0.060	OPEN/ NEWS	fought	—
07-08 13:40	▼ -10.6	0.241	OPEN/ NEWS	fought	—
07-08 14:02	▼ -13.5	0.191	OPEN/ NEWS	sat out	—
07-08 14:35	▼ -12.4	0.205	OPEN/ NEWS	sat out	—
07-08 15:11	▼ -12.9	0.105	quiet/uncl.	sat out	—
07-08 15:36	▼ -6.3	0.149	quiet/uncl.	caught	-33
07-08 15:56	▲ +11.7	0.166	quiet/uncl.	sat out	—
07-08 16:16	▲ +8.7	0.132	quiet/uncl.	sat out	—
07-08 16:37	▲ +10.6	0.186	quiet/uncl.	sat out	—
07-08 17:17	▲ +8.4	0.101	quiet/uncl.	caught	+53
07-08 17:32	▼ -6.4	0.155	quiet/uncl.	sat out	—
07-08 17:57	▲ +6.7	0.189	quiet/uncl.	sat out	—
07-08 19:42	▲ +9.8	0.125	quiet/uncl.	sat out	—
07-08 19:57	▼ -7.2	0.255	quiet/uncl.	sat out	—
07-09 02:31	▼ -5.3	0.091	quiet/uncl.	fought	—

07-09 03:19	▲ +5.3	0.057	quiet/uncl.	sat out	—
07-09 10:39	▼ -6.5	0.057	quiet/uncl.	sat out	—
07-09 13:21	▲ +8.0	0.067	OPEN/ NEWS	sat out	—
07-09 13:40	▲ +13.0	0.305	OPEN/ NEWS	sat out	—
07-09 14:21	▲ +13.9	0.307	OPEN/ NEWS	sat out	—
07-09 14:54	▲ +6.3	0.186	OPEN/ NEWS	sat out	—
07-09 15:10	▼ -5.7	0.116	quiet/uncl.	sat out	—
07-10 12:01	▲ +6.5	0.070	quiet/uncl.	sat out	—
07-10 13:27	▼ -19.1	0.076	OPEN/ NEWS	sat out	—
07-10 13:53	▼ -10.9	0.220	OPEN/ NEWS	caught	+52
07-10 14:18	▼ -7.2	0.137	OPEN/ NEWS	sat out	—
07-10 14:34	▲ +10.6	0.858	OPEN/ NEWS	fought	—
07-10 14:51	▼ -12.2	0.127	OPEN/ NEWS	sat out	—
07-10 15:12	▲ +8.9	0.101	quiet/uncl.	fought	—
07-10 15:27	▼ -8.1	0.164	quiet/uncl.	sat out	—

07-10
16:37

▲ +6.0

0.100

quiet/uncl.

sat
out

—

"real \$" = honest real_pnl summed across every sim that fired the run's way within ± 5 min (only shown when we caught it). Red-shaded rows are sit-outs of 12 pt or more — the ones that sting.

§5.2 2 — Did we show up? Bluntly, less than half

OUTCOME	RUNS	GROSS TRAVEL (1-LOT)	HONEST MONEY WHEN INVOLVED
caught fired the right way	17	185 pt · \$924	+\$285
fought faded into it	7	60 pt · \$300	−\$1.6 (flat)
sat out nothing fired	40	351 pt · \$1,756	—

So of ~\$3,000 of gross big-run travel this week (1-lot terms), we were pointed the right way at **\$924** of it and took home **+\$285** honest. We left **\$1,756** of travel completely untouched. The good news buried in here: the fades were not the bloodbath you'd fear — fighting 7 runs came out roughly flat (−\$1.6), because one of them (07-10 14:34) had a momentum sim riding the right way underneath. The seven fades are cheap; the forty no-shows are the story.

WHERE THE −\$252 ACTUALLY COMES FROM

M2K ran 502 honestly-scored sims this week: **+\$769 on paper**, **−\$252 once you use real fills**. Split it and the picture is clean — the run-catching momentum family (tw_m2k_momo_loose,

momentum_shadow, ORB and the open sims) made **+\$170** honest; everything else — the fader book — lost **−\$422**, of which passive_chop alone was **−\$304**. M2K doesn't have a run problem. It has a fading-the-small-chop problem, and a showing-up-for-the-big-move problem.

§5.3 3 — The footprint before a run: it erupts from quiet

I averaged the funnel over the 10 minutes before each sit-out run. The headline is uncomfortable: **volume gives you nothing**. The pre-run rvol is 0.66 in the quiet cluster and 0.74 in the open cluster — and the week-average rvol across every M2K cycle is 0.649. In other words, the ten minutes before a big run look exactly like any other ten minutes on volume. M2K runs don't build; they pop out of calm.

PRE-RUN (10 MIN), SIT-OUTS	ATR_PCT	RVOL	VWAP_DIST	BOOK_IMBALANCE
OPEN / NEWS cluster (n=12)	0.088	0.74	−0.103	+0.009
quiet / unclassified (n=27)	0.055	0.66	−0.015	+0.004
week-average baseline	—	0.649	—	~0

The only thing that separates the two clusters is **the clock and the amplitude**. The OPEN runs sit at roughly double the pre-run atr_pct (0.088 vs 0.055), they lean a bit below VWAP going in (−0.10), and they're the highest-amplitude runs on the sheet (amp 0.19–0.86 vs a cluster of 0.05–0.15 overnight). The funnel's L1 book-imbalance is

dead flat (+0.004/+0.009) for both — it does not lean the way the run goes, so it's no help as a trigger. The one cluster we repeatedly sit out with a clean, repeatable signature is the OPEN cluster: it's not the volume, it's the wall-clock and the volatility regime. That's the new-entry hypothesis.

§5.4 4 — The L2 book: nothing usable for M2K this week

Two honest problems. First, **M2K's own order book barely exists in the capture** — just 10,820 snapshots, a 1.8-hour sliver on Friday evening (19:10–21:00 UTC), and not one of the 64 runs falls inside it. So for M2K's own depth, there is nothing to look at.

Second, I tried the sister-index book as a proxy — Russell moves with the Nasdaq and S&P at the open, and MNQ/MES have deep books over the Thu–Fri window. It doesn't hold up:

M2K RUN	MNQ BOOK, ~1 MIN BEFORE	VERDICT
07-10 13:27 ▼ -19.1 dump	bid-heavy, imbalance +0.19 (41 bid vs 28 ask, top-5)	leaned the wrong way
07-10 14:34 ▲ +10.6 pop	ask-heavy, imbalance -0.29 to -0.55 going in	leaned the wrong way

Both times the visible book leaned against the move that followed — classic limit orders getting run over, not a vacuum you could see coming. And the 14:34 "up-run" was really the recovery leg of a shared MNQ/MES V-bottom (MNQ knifed 29953→29757 then bounced), so the "run" is an artifact of where the 15-min window was cut. **No clean, repeatable L2 tell for M2K this week** — neither its own book (absent) nor the proxy (contrarian and noisy). To say

anything real here we need M2K-native depth capture running through the cash-open hour.

§5.5 5 — The honest money left on the table

The capturable-vs-hindsight gap is smaller than the gross-travel number suggests. A perfect-hindsight ceiling would have grabbed most of \$3,000 of 1-lot travel; a realistic sim shape — a momentum entry that needs the break to actually start — can only ever take the middle of a run, not the tick. This week that realistic shape, when it fired, banked **+\$285 honest on 17 runs**, and it's lumpy: four runs (13:29 Mon +\$104, 17:17 Tue +\$53, 13:53 Fri +\$52, 13:41 Tue +\$51) are basically the entire haul. The clean, addressable miss is the **13 OPEN-cluster sit-outs worth ~\$700 of 1-lot travel** that share a signature we can actually detect. The 27 quiet-cluster sit-outs are mostly overnight, tiny-amplitude, and tape-blind — realistically uncatchable with what we capture, and not worth chasing.

§5.6 6 — How we grab them from Monday

§5.7 The OPEN-hour break shadow-first

The cluster: 23 runs, all between 13:05 and 14:54 UTC (the 09:30 ET cash open, ± 90 min). Highest amplitude of the week. We sat out 13, including the four biggest sit-outs on the whole sheet (07-10 -19.1 , 07-09 $+13.9$, 07-08 -13.5 , 07-09 $+13.0$).

The trigger (observable, all from what we already capture): time is inside 13:20–15:00 UTC and $\text{atr_pct} \geq 0.09$ (the open-cluster runs sit at 0.088–0.86; the quiet ones at 0.055) and price breaks the high/low of the 13:15–13:30 opening range. **Direction = the direction of the break, never pre-chosen** — the open was 13 down / 10 up and flips inside the same morning (Monday alone went DN, UP, DN, UP, UP), which is exactly why pre-positioning dies.

The exit: reuse the momentum_shadow / tw_m2k_momo_loose exit spine that already earned +\$170 honest on M2K this week — tuned chandelier trail, not a fixed target, so a 20-pt open run isn't clipped at 2R.

Why shadow-first, not armed: our existing tw_m2k_open_momentum is the cautionary tale — it fired just 3 times this week, each a pre-positioned guess (short at 14:00 Mon into a rising tape, short 11:57 Tue, long 13:51 Fri), and all three hit their stop for −\$76. The fix is the break-confirmation gate above. One week, in-sample, 23 runs — run it as an observe-only shadow with an ATR-floor and an opening-range break for 2–3 weeks before it earns a live lot.

§5.8 The quiet overnight cluster can't forecast — don't try

41 runs, spread across every hour, no tape signature, pre-run volume at the week average. M2K has **no aggressor tape captured at all** and its L2 book covers none of these runs. There is nothing in what we record that leads these moves — rvol doesn't build, the book doesn't lean, VWAP is flat. Chasing them now would just be curve-fitting noise. **The data we'd need first: M2K-native trade_tick (aggressor flow) and depth_snap running around the clock.** Until we capture M2K's own microstructure, this cluster stays a sit-out and that's the correct call.

M2K — THE PLAIN TAKEAWAY

M2K's real money is in the US-open hour, we're mostly not there, and the one sim that tries pre-positions the wrong way and gets stopped. Build a **shadow-first opening-range break** — only inside 13:20–15:00 UTC, only when $\text{atr_pct} \geq 0.09$, trade

the direction of the break, exit on the momentum chandelier. Leave the overnight quiet cluster alone until we capture M2K's own tape and book. And separately: the -\$252 on the week isn't the runs at all — it's passive_chop fading small moves for -\$304. Killing that does more for M2K than any new entry.

§6 The biggest runs on MYM — did we show up?

I pulled the **full run census** for MYM this week — every 15-minute stretch where the Dow micro travelled at least $1.5\times$ its typical range (a move of 60 points or more, \$30+ per contract). Not a top-8: the scorecard found **73 of them** between Monday 00:02 and Friday 16:38 UTC. The blunt headline is that MYM ran hard and mostly ran without us: we put a sim on the right side of **11** of the 73 and **sat out 62**. On the 11 we did catch, the run-direction sims made **+\$195** of honest money — proof the shape works when it fires — but the book as a whole still bled **-\$559** because the fade sims kept stepping in front of the same moves. One cluster stands out as the highest-value thing to build: the **US-open sell-off**, which telegraphs itself in the funnel and which we sat out 19 times out of 24.

BOTTOM LINE

73 big runs, ~\$3,718 of hindsight range on the table ($\$0.50/\text{pt} \times 7,436$ points). We caught 11, sat out 62, and never fought — wait, we did fight: the whole shadow book netted **-\$559 real** (184 wins / 317 scored fills), and the single biggest bleeder is **passive_chop at -\$269**, fading straight into these runs. The runs we joined in their direction paid **+\$195** — that is the capturable core. The clean, repeatable, sat-out signature is the **US cash-open down-drive** (16 down-runs; pre-run VWAP slope steepening to -0.36 ATR, ATR% doubling to ~ 0.055 , net-sell order flow). MYM has no L2 book and no aggressor tape of

its own this week, so the trigger has to come from the L1 funnel — and it is already sitting there.

§6.1 1 The full census — every run $\geq 1.5 \times \text{ATR}$ this week

Every 15-minute window where MYM moved at least 60 points ($1.5 \times$ its 40-point typical range), from scripts/footprint_scorecard.py missed --symbol MYM --top 100. Sized at \$0.50/pt. “amp%” is the pre-run ATR fraction the scorecard measured. Green rows = a sim fired with the run; red rows = a ≥ 150 -point run we sat out. I re-priced two off the raw 5s bars to check the scorecard: Mon 13:18 shows a high of 53,365 against a 53,094 open (+271 top-to-open, booked +255 open-to-close) and Wed 08:15 ran 52,999→52,660 (−339 range, booked −307) — both real, not artefacts.

RUN START (UTC)	DIR	MOVE PT	\$	DID WE?	AMP%	CLUSTER
Mon 07-06 00:02	↑ UP	71	\$36	caught ✓	0.088	quiet drift
Mon 07-06 13:18	↑ UP	255	\$128	caught ✓	0.036	US open
Mon 07-06 13:34	↓ DN	185	\$92	caught ✓	0.360	US open
Mon 07-06 13:51	↓ DN	159	\$80	sat out	0.103	US open
Mon 07-06 14:08	↑ UP	137	\$68	caught ✓	0.198	US open
Mon 07-06 14:45	↑ UP	112	\$56	sat out	0.139	US open
Mon 07-06 17:00	↑ UP	61	\$30	sat out	0.047	quiet drift

Mon 07-06 17:28	↓ DN	66	\$33	sat out	0.047	quiet drift
Mon 07-06 17:53	↑ UP	60	\$30	sat out	0.041	quiet drift
Mon 07-06 18:43	↑ UP	61	\$30	sat out	0.032	quiet drift
Mon 07-06 19:48	↑ UP	83	\$42	sat out	0.047	quiet drift
Tue 07-07 08:10	↑ UP	78	\$39	sat out	0.032	quiet drift
Tue 07-07 11:02	↑ UP	61	\$30	sat out	0.041	quiet drift
Tue 07-07 12:33	↑ UP	80	\$40	sat out	0.054	quiet drift
Tue 07-07 13:05	↑ UP	77	\$38	sat out	0.069	US open
Tue 07-07 13:26	↓ DN	237	\$118	sat out	0.076	US open
Tue 07-07 13:44	↓ DN	176	\$88	sat out	0.213	US open
Tue 07-07 14:02	↓ DN	109	\$54	sat out	0.092	US open
Tue 07-07 14:27	↓ DN	102	\$51	sat out	0.143	US open
Tue 07-07 14:59	↓ DN	96	\$48	sat out	0.079	US open
Tue 07-07 17:17	↑ UP	87	\$44	sat out	0.068	quiet drift
Tue 07-07 18:46	↓ DN	80	\$40	sat out	0.058	quiet drift

Tue 07-07 19:40	↑ UP	105	\$52	caught ✓	0.062	quiet drift
Tue 07-07 19:55	↓ DN	73	\$36	sat out	0.115	quiet drift
Tue 07-07 23:47	↓ DN	71	\$36	caught ✓	0.036	quiet drift
Wed 07-08 06:24	↓ DN	63	\$32	sat out	0.053	quiet drift
Wed 07-08 06:51	↓ DN	72	\$36	sat out	0.085	quiet drift
Wed 07-08 08:15	↓ DN	307	\$154	sat out	0.043	quiet drift
Wed 07-08 08:31	↓ DN	78	\$39	sat out	0.190	quiet drift
Wed 07-08 08:53	↓ DN	118	\$59	sat out	0.089	quiet drift
Wed 07-08 10:24	↑ UP	109	\$54	caught ✓	0.050	quiet drift
Wed 07-08 11:28	↑ UP	129	\$64	sat out	0.053	quiet drift
Wed 07-08 12:00	↑ UP	105	\$52	sat out	0.135	quiet drift
Wed 07-08 12:44	↓ DN	76	\$38	sat out	0.049	quiet drift
Wed 07-08 12:59	↑ UP	61	\$30	sat out	0.167	quiet drift
Wed 07-08 13:14	↓ DN	91	\$46	sat out	0.119	US open
Wed 07-08 13:40	↓ DN	150	\$75	sat out	0.252	US open

Wed 07-08 14:05	↓ DN	141	\$70	sat out	0.150	US open
Wed 07-08 14:21	↓ DN	128	\$64	sat out	0.207	US open
Wed 07-08 14:52	↓ DN	72	\$36	sat out	0.149	US open
Wed 07-08 15:20	↓ DN	80	\$40	sat out	0.168	quiet drift
Wed 07-08 15:56	↑ UP	120	\$60	sat out	0.074	quiet drift
Wed 07-08 16:37	↑ UP	136	\$68	sat out	0.114	quiet drift
Wed 07-08 17:09	↑ UP	77	\$38	sat out	0.097	quiet drift
Wed 07-08 17:33	↓ DN	76	\$38	sat out	0.066	quiet drift
Wed 07-08 17:57	↑ UP	66	\$33	sat out	0.076	quiet drift
Wed 07-08 19:57	↓ DN	72	\$36	sat out	0.068	quiet drift
Thu 07-09 09:26	↓ DN	73	\$36	sat out	0.025	quiet drift
Thu 07-09 10:15	↓ DN	62	\$31	sat out	0.034	quiet drift
Thu 07-09 10:38	↓ DN	64	\$32	sat out	0.061	quiet drift
Thu 07-09 12:22	↑ UP	62	\$31	sat out	0.049	quiet drift
Thu 07-09 12:52	↑ UP	61	\$30	caught ✓	0.068	quiet drift

Thu 07-09 13:24	↓ DN	160	\$80	sat out	0.083	US open
Thu 07-09 13:40	↑ UP	150	\$75	caught ✓	0.150	US open
Thu 07-09 14:20	↑ UP	105	\$52	sat out	0.125	US open
Thu 07-09 14:53	↑ UP	100	\$50	sat out	0.072	US open
Thu 07-09 15:30	↑ UP	86	\$43	sat out	0.051	quiet drift
Thu 07-09 15:45	↓ DN	62	\$31	sat out	0.115	quiet drift
Thu 07-09 18:12	↓ DN	73	\$36	sat out	0.051	quiet drift
Thu 07-09 19:51	↑ UP	69	\$34	sat out	0.064	quiet drift
Fri 07-10 05:51	↑ UP	78	\$39	sat out	0.009	quiet drift
Fri 07-10 07:55	↑ UP	79	\$40	sat out	0.021	quiet drift
Fri 07-10 08:55	↑ UP	61	\$30	sat out	0.064	quiet drift
Fri 07-10 11:25	↓ DN	78	\$39	sat out	0.026	quiet drift
Fri 07-10 12:03	↑ UP	84	\$42	sat out	0.059	quiet drift
Fri 07-10 12:33	↑ UP	65	\$32	sat out	0.061	quiet drift
Fri 07-10 13:26	↓ DN	116	\$58	sat out	0.038	US open

Fri 07-10 14:00	↓ DN	78	\$39	sat out	0.083	US open
Fri 07-10 14:18	↓ DN	203	\$102	caught ✓	0.163	US open
Fri 07-10 14:33	↑ UP	232	\$116	sat out	0.681	US open
Fri 07-10 15:03	↑ UP	97	\$48	caught ✓	0.108	quiet drift
Fri 07-10 15:26	↓ DN	67	\$34	sat out	0.144	quiet drift
Fri 07-10 16:38	↑ UP	62	\$31	sat out	0.112	quiet drift

THE TWO CLUSTERS

The scorecard sorts the 73 runs into two cause-footprints: **US open / news spike — 24 runs, sat out 19**; and **quiet drift (no tape signature) — 49 runs, sat out 43**. (MYM has no aggressor tape, so nothing landed in the FLOW-LED or VACUUM buckets that the tape-covered contracts get.) The open cluster is where the size is: it holds the −237, −176, −160, −150, +255, +232, +203 monsters. The quiet-drift cluster is 49 small 60-130-point grinds scattered through the overnight and mid-day dead zones.

§6.2 2 Did we show up? — participation and the honest money

Blunt tally across the whole census: **caught 11 / 73** (a sim fired in the run's direction within ± 5 min), **sat out 62**. When I sum the honest `real_pnl` of every sim that fired the right way inside those 11

run windows, it comes to **+\$194.79 over 33 fills** — the runs we joined actually paid. Here is where each of the 11 landed:

RUN	DIR	SIMS	REAL \$
Thu 07-09 12:52	↑ UP	10	+83.30
Fri 07-10 15:03	↑ UP	4	+68.02
Mon 07-06 13:18	↑ UP	3	+37.89
Mon 07-06 14:08	↑ UP	4	+32.02
Wed 07-08 10:24	↑ UP	1	+24.38
Thu 07-09 13:40	↑ UP	1	+14.38
Tue 07-07 19:40	↑ UP	3	+4.14
Tue 07-07 23:47	↓ DN	1	−7.62
Mon 07-06 00:02	↑ UP	4	−30.48
Fri 07-10 14:18	↓ DN	2	−31.24
Mon 07-06 13:34	↓ DN	0	—

Notice the tilt: nine of the eleven joins were **UP** runs, and the two DN joins both lost. We are far better at latching onto MYM's up-thrusts than its sell-offs — and the sell-offs are the bigger, more numerous moves. Now the other side of the ledger. The whole MYM shadow book this week was **−\$559 real** across 317 scored fills, and it breaks down like this:

STRATEGY	FILLS	REAL \$	WINS	ROLE IN THE RUNS
passive_chop	85	−269.46	14	fades into the drive — the knife-catcher
tw_mym_no_fade_drive	39	−77.68	18	meant to go with, still net red
orb_shadow	16	−76.42	2	

				opening-range break, mistimed
tw_deadtape_standdown	44	−63.78	22	stand-down that still traded
tw_clock_standdown	39	−44.68	20	time-of-day filter
momentum_shadow	61	−31.82	22	with-trend, small bleed
tw_mym_open_thrust	12	+28.06	7	the one bright spot — open drive, with it
pullback_quiet	13	+0.94	8	flat
tw_strength_ceiling_pb	4	+3.52	3	thin

kill The story writes itself: the run-catchers make money on runs (+\$195), and then **passive_chop hands back −\$269** fading the very same moves, and the two “stand-down” gates that are supposed to keep us out lose another −\$108 by trading anyway. The only strategy pointed at the runs that finished green is **tw_mym_open_thrust (+\$28, 7/12)** — and it is pointed straight at the cluster below.

§6.3 3 The footprint, clustered — what the 10 minutes before looked like

For each cluster I averaged the L1 funnel features over the 10 minutes before each run started (fut_signal_funnel, MYM). This is the tell we would actually see in real time.

CLUSTER	RUNS	ATR%	RVOL	VWAP SLOPE	OFI	PLAIN FOOTPRINT
US open — DN (16)	16	0.056	0.79	−0.36	−5.5	ATR doubles, VWAP rolls over hard, sellers lean

US open — UP (8)	8	0.051	0.71	+0.05	−12.2	V-snap reversals; slope does NOT pre-lean up
Quiet drift — UP (29)	29	0.026	0.61	+0.31	−6.6	slow grind, tiny amplitude
Quiet drift — DN (20)	20	0.027	0.62	−0.42	+5.6	slow grind, tiny amplitude

Two things jump out. First, the **US-open down cluster has a clean, one-directional pre-run signature**: in the 10 minutes before the sell-off, ATR% is already running at ~0.056 (double the quiet-day 0.026), the VWAP slope is steeply negative (−0.36 ATR), and order-flow imbalance leans net-sell. I checked how reliable that slope tell is: **10 of the 16 open down-runs had a pre-run VWAP slope steeper than −0.10** — price was already rolling before the leg. That is a footprint you can trade off.

Second, the **US-open UP runs do NOT telegraph**: their pre-run slope is a flat +0.05 and only 2 of 8 pre-leaned up — because they are V-bottom reversals (Friday's +232 came the minute after a −203 flush). And the quiet-drift clusters, while their slope aligns with direction, run at a third of the amplitude (0.026 ATR%, 60–80-point moves worth \$30–40) — forecastable-ish but far too small and far too many to trade without drowning in chop tax. That quiet bucket is exactly where passive_chop lost its \$269.

THE ONE TO BUILD

The highest-value sat-out signature is unambiguous: the **US cash-open down-drive**. Sixteen of them, we sat out most, and every one arrives on a pre-run footprint of expanding ATR +

steep negative VWAP slope + net-sell flow that we already capture in the funnel.

§6.4 4 The L2 book — MYM has none, so I used MNQ as a proxy

Straight answer: **MYM is not in the depth capture and has no aggressor tape this week** — `depth.db` covers MNQ/MES/MGC only, and `trade_tick` covers MNQ/MES/MGC/MBT. So there is no native L2 tell for the Dow micro. The nearest thing is to read the **MNQ** book as a correlated index proxy over the two Friday runs that fall inside the depth window (Thu 14:17 → Fri 21:00 UTC): the −203 down-flush at 14:18 and the +232 V-snap at 14:33.

In the 90 seconds before MYM's −203, the MNQ book imbalance averaged **−0.044** (a faint ask-heavy, down lean) across 358 snaps — but it swung from −0.51 to +0.35 tick-to-tick. Before the +232, it averaged **+0.045** (faint bid lean) across 339 snaps, swinging −0.48 to +0.93. So the mean leans the right way by a whisker each time, but the sign flips several times a second and the top-of-book sizes are tiny (1–11 lots). I looked for a far-side vacuum — the ask stack thinning out 1–2 minutes ahead — and did not find a clean, repeatable one; the book just churns.


read Honest verdict on L2 for MYM: **no usable book tell**. Even on the proxy the signal is buried in noise, and MYM has no book of its own anyway. If the owner wants a genuine L2 edge on the Dow micro, the data we would need is MYM depth capture (it is not entitled/subscribed today) — until then the funnel's L1 slope + OFI is the best pre-run read we have.

\$6.5 5 The honest money left on the table

The hindsight ceiling — catching every one of the 73 runs top-to-bottom — is **~\$3,718 per contract** ($7,436 \text{ points} \times \0.50). That is a fantasy number; nobody catches the full range of a 15-minute leg. The capturable figure is what a real sim shape actually banked when it joined a run in its direction: **+\$195 across the 11 we caught**. The gap between **-\$559** (what the book really did) and **+\$195** (what the run-joins did) is **-\$754 of self-inflicted damage** — almost all of it passive_chop and the stand-down gates fading into the moves. So the money left on the table isn't really the runs we missed; it's the **fades we should never have taken**. Stop the fade bleed and the same book is already positive on MYM before we add a single new entry.

Chart — every run this week, by size (\$), in time order.

Green bars = we caught it (sim fired with the run); red = we sat out. Gold ticks along the bottom mark the US-open/news cluster. The eye-test is the point of the section: a long red wall of sell-offs (the 08:15 $-\$154$, the 13:26 $-\$118$, the 14:33 $+\$116$) with only a scatter of green — and almost every gold-ticked big bar is red.



§6.6 6 How we grab them from Monday

One cluster is worth arming, one is a stand-down, and one needs data we don't have.

PRESCRIPTION A — THE US-OPEN DOWN-DRIVE (ARM AS SHADOW, READY TO PROMOTE)

When: in the 13:30-15:00 UTC cash-open window, and **atr_pct** ≥ 0.05 (amplitude is live, not a quiet day) **and** **vwap_slope_atr** ≤ -0.15 (price already rolling under VWAP) **and** **vwap_dist_pct** < 0 (trading below value) **and** **OFI** ≤ 0 (sellers leaning). **Direction:** SHORT, with the drive — do NOT fade it. **Exit:** the **tw_mym_open_thrust** shape is the proven +EV catcher here (+\$28, 7/12) — extend it to the short side, ride with a tuned chandelier / give-back trail and a hard $0.75 \times \text{ATR}$ stop, flat by the 15:00 window close or 20-min clock, whichever first. This mirrors what the capture-lab found independently (open-only momentum is the one net-positive construction). **Status:** shadow-first — the trigger is in-sample on one week and thin-n; run it observe-only next week and promote if it holds. It fires ~16 times/week and 10/16 pre-leaned correctly, so it is armable-soon, not armable-blind-today.

PRESCRIPTION B — THE QUIET-DRIFT BUCKET (STAND DOWN)

Forty-nine runs, but they run at 0.026 ATR% and pay \$30-40 apiece. The slope mildly aligns, so they are forecastable, but trading them means many small entries in dead tape — precisely the tape where **passive_chop** lost \$269 this week. **The action here is subtractive:** stop **passive_chop** and the two stand-down gates from firing in this window. That recovers more than any new entry would add. No new long/short to build.

PRESCRIPTION C — THE US-OPEN UP SNAPS (CAN'T FORECAST FROM WHAT WE CAPTURE)

The eight open UP runs are V-bottom reversals (Friday's +232 the minute after a −203). Their pre-run slope is flat +0.05 and OFI is actually net-sell just before they rip — contrarian, un-leadable from L1. To catch these we would need a genuine reversal trigger (a fast absorption/exhaustion read), and that means **MYM-native L2 depth + aggressor tape, which we do not have**. Name it and park it: no honest entry from current data.

Plain MYM takeaway: MYM ran 73 times this week and we joined 11 — and the joins made money (+\$195) while the fades lost it all back and more (book −\$559, passive_chop −\$269). There is exactly one clean, repeatable, sat-out edge: **short the US cash-open sell-off** when ATR is expanding, VWAP is rolling over, and flow leans sell — with the open-thrust exit that already prints green. Arm it shadow-first, and in the same breath stop the chop-fades that are the real leak. Everything else on MYM this week is either too small to bother with or needs a book we don't yet capture.

§7 The biggest runs on MCL this week — did we show up?

Crude ran plenty this week. The scorecard counts **78 separate 15-minute moves of \$40 or more** on MCL between Monday 07-06 and Friday 07-10 — every one at least 1.5× the contract's typical 15-minute range (that typical range is only about \$30, so crude is a jumpy little contract). We sat out **50 of the 78**. When we did show up, the honest money was a wash: for every momentum win there was a fade that got run over. The real bleed on MCL isn't the runs we miss — it's the chop we trade between them.

READ THIS FIRST — WE ARE BLIND ON MCL

Before anything else: **MCL has no order book and effectively no tape**. The L2 depth farm (depth.db) only carries MNQ, MES, MGC and M2K — NYMEX crude came back "not subscribed" (Err354). The aggressor tape has just **2,047 MCL ticks all week** versus 7.4 million for MNQ. So the flow column in the census is blank, Section 4's book work is impossible here, and every read below leans on 5-second bars and the signal funnel alone. That is the single biggest thing holding MCL back.

§7.1 1 • The full census — all 78 runs

Here is every run, in time order, no top-8. Direction, dollar size (\$100 a point on MCL), whether a sim fired the right way within five minutes, the cause cluster, and — where we showed up — the honest `real_pnl` of the with-run sims around it. Green rows are runs we

caught and made money on; red rows are runs we "caught" but still lost. I verified three of the big ones straight off the 5-second bars: the 07-08 08:11 whale printed a 2.33-point range (census said +2.3), 07-07 18:46 printed 1.12 (said +1.1), 07-10 14:33 printed 1.47 — the census sizing is honest.

RUN (UTC)	DIR	SIZE	US?	CLUSTER	REAL \$	
Mon 07-06 00:02	↑ UP	71	\$36	caught ✓	0.088	quiet drift
Mon 07-06 13:18	↑ UP	255	\$128	caught ✓	0.036	US open
Mon 07-06 13:34	↓ DN	185	\$92	caught ✓	0.360	US open
Mon 07-06 13:51	↓ DN	159	\$80	sat out	0.103	US open
Mon 07-06 14:08	↑ UP	137	\$68	caught ✓	0.198	US open
Mon 07-06 14:45	↑ UP	112	\$56	sat out	0.139	US open
Mon 07-06 17:00	↑ UP	61	\$30	sat out	0.047	quiet drift
Mon 07-06 17:28	↓ DN	66	\$33	sat out	0.047	quiet drift
Mon 07-06 17:53	↑ UP	60	\$30	sat out	0.041	quiet drift
Mon 07-06 18:43	↑ UP	61	\$30	sat out	0.032	quiet drift
Mon 07-06 19:48	↑ UP	83	\$42	sat out	0.047	quiet drift

Tue 07-07 08:10	↑ UP	78	\$39	sat out	0.032	quiet drift
Tue 07-07 11:02	↑ UP	61	\$30	sat out	0.041	quiet drift
Tue 07-07 12:33	↑ UP	80	\$40	sat out	0.054	quiet drift
Tue 07-07 13:05	↑ UP	77	\$38	sat out	0.069	US open
Tue 07-07 13:26	↓ DN	237	\$118	sat out	0.076	US open
Tue 07-07 13:44	↓ DN	176	\$88	sat out	0.213	US open
Tue 07-07 14:02	↓ DN	109	\$54	sat out	0.092	US open
Tue 07-07 14:27	↓ DN	102	\$51	sat out	0.143	US open
Tue 07-07 14:59	↓ DN	96	\$48	sat out	0.079	US open
Tue 07-07 17:17	↑ UP	87	\$44	sat out	0.068	quiet drift
Tue 07-07 18:46	↓ DN	80	\$40	sat out	0.058	quiet drift
Tue 07-07 19:40	↑ UP	105	\$52	caught ✓	0.062	quiet drift
Tue 07-07 19:55	↓ DN	73	\$36	sat out	0.115	quiet drift
Tue 07-07 23:47	↓ DN	71	\$36	caught ✓	0.036	quiet drift
Wed 07-08 06:24	↓ DN	63	\$32	sat out	0.053	quiet drift


Wed 07-08 06:51	↓ DN	72	\$36	sat out	0.085	quiet drift
Wed 07-08 08:15	↓ DN	307	\$154	sat out	0.043	quiet drift
Wed 07-08 08:31	↓ DN	78	\$39	sat out	0.190	quiet drift
Wed 07-08 08:53	↓ DN	118	\$59	sat out	0.089	quiet drift
Wed 07-08 10:24	↑ UP	109	\$54	caught ✓	0.050	quiet drift
Wed 07-08 11:28	↑ UP	129	\$64	sat out	0.053	quiet drift
Wed 07-08 12:00	↑ UP	105	\$52	sat out	0.135	quiet drift
Wed 07-08 12:44	↓ DN	76	\$38	sat out	0.049	quiet drift
Wed 07-08 12:59	↑ UP	61	\$30	sat out	0.167	quiet drift
Wed 07-08 13:14	↓ DN	91	\$46	sat out	0.119	US open
Wed 07-08 13:40	↓ DN	150	\$75	sat out	0.252	US open
Wed 07-08 14:05	↓ DN	141	\$70	sat out	0.150	US open
Wed 07-08 14:21	↓ DN	128	\$64	sat out	0.207	US open
Wed 07-08 14:52	↓ DN	72	\$36	sat out	0.149	US open
Wed 07-08 15:20	↓ DN	80	\$40	sat out	0.168	quiet drift

Wed 07-08 15:56	↑ UP	120	\$60	sat out	0.074	quiet drift
Wed 07-08 16:37	↑ UP	136	\$68	sat out	0.114	quiet drift
Wed 07-08 17:09	↑ UP	77	\$38	sat out	0.097	quiet drift
Wed 07-08 17:33	↓ DN	76	\$38	sat out	0.066	quiet drift
Wed 07-08 17:57	↑ UP	66	\$33	sat out	0.076	quiet drift
Wed 07-08 19:57	↓ DN	72	\$36	sat out	0.068	quiet drift
Thu 07-09 09:26	↓ DN	73	\$36	sat out	0.025	quiet drift
Thu 07-09 10:15	↓ DN	62	\$31	sat out	0.034	quiet drift
Thu 07-09 10:38	↓ DN	64	\$32	sat out	0.061	quiet drift
Thu 07-09 12:22	↑ UP	62	\$31	sat out	0.049	quiet drift
Thu 07-09 12:52	↑ UP	61	\$30	caught ✓	0.068	quiet drift
Thu 07-09 13:24	↓ DN	160	\$80	sat out	0.083	US open
Thu 07-09 13:40	↑ UP	150	\$75	caught ✓	0.150	US open
Thu 07-09 14:20	↑ UP	105	\$52	sat out	0.125	US open
Thu 07-09 14:53	↑ UP	100	\$50	sat out	0.072	US open

Thu 07-09 15:30	↑ UP	86	\$43	sat out	0.051	quiet drift
Thu 07-09 15:45	↓ DN	62	\$31	sat out	0.115	quiet drift
Thu 07-09 18:12	↓ DN	73	\$36	sat out	0.051	quiet drift
Thu 07-09 19:51	↑ UP	69	\$34	sat out	0.064	quiet drift
Fri 07-10 05:51	↑ UP	78	\$39	sat out	0.009	quiet drift
Fri 07-10 07:55	↑ UP	79	\$40	sat out	0.021	quiet drift
Fri 07-10 08:55	↑ UP	61	\$30	sat out	0.064	quiet drift
Fri 07-10 11:25	↓ DN	78	\$39	sat out	0.026	quiet drift
Fri 07-10 12:03	↑ UP	84	\$42	sat out	0.059	quiet drift
Fri 07-10 12:33	↑ UP	65	\$32	sat out	0.061	quiet drift
Fri 07-10 13:26	↓ DN	116	\$58	sat out	0.038	US open
Fri 07-10 14:00	↓ DN	78	\$39	sat out	0.083	US open
Fri 07-10 14:18	↓ DN	203	\$102	caught ✓	0.163	US open
Fri 07-10 14:33	↑ UP	232	\$116	sat out	0.681	US open
Fri 07-10 15:03	↑ UP	97	\$48	caught ✓	0.108	quiet drift

Fri 07-10 15:26	↓ DN	67	\$34	sat out	0.144	quiet drift
Fri 07-10 16:38	↑ UP	62	\$31	sat out	0.112	quiet drift

Each bar below is one run in time order, height = dollar size. **Green = we fired with it, red = we sat out.** The one bar that towers over the week — the \$230 whale at 07-08 08:11 — is red. We watched the single biggest MCL move of the week go by.



§7.2 2 • Did we show up? Bluntly, mostly no

Across the whole 78-run census we fired the right way on **28** and sat out **50**. Broken down by cause:

CLUSTER	RUNS	SAT OUT	CAUGHT	REAL \$ WHEN WITH-RUN
Vol-expansion break	28	21	7	+109
Quiet drift (unclassified)	34	19	15	+114
Open / news spike	16	10	6	+220

The honest money picture, once and for all. Taking every MCL sim that fired with a run within five minutes: 38 trades, **+\$443** real. Every sim that fired against a run (a fade walking into the move): 41 trades, **-\$414** real. So around the runs themselves we net roughly flat — about **+\$29**. But the desk's total on MCL for the week was **-\$1,716** across 690 trades. That means essentially **the entire loss came from grinding chop between the runs**, not from the runs. The two worst offenders: passive_chop long+short (**-\$700** real) and tw_mcl_short_with_strength (295 trades, a rosy **+\$958** in the optimistic column but **-\$341** once you price the real fills — a \$1,300 optimism gap, the classic reversion-fading-into-a-grind).

THE ONE THING THAT WORKS ON MCL

momentum_shadow LONG: 36 trades, **+\$152.68** real, 19 wins / 17 losses, best trade **+\$52.88**, worst just **-\$16.12**, average hold about 30 minutes. That asymmetry — small capped losers, winners left to run — is the entire edge. It is the only MCL shape in the black, and it is the shape that catches crude's runs.

§7.3 3 · The footprint before each cluster

What did the ten minutes before each type of run look like in the funnel (we have no tape or book here, remember)?

CLUSTER	PRE-RUN ATR%	PRE-RUN RVOL	VWAP DIST%	VWAP SLOPE (10M)	PLAIN READ
Vol-expansion	0.154	0.75	+0.119	+0.037	range already widening, price already leaning the break way

Open / news	0.152	0.63	+0.031	+0.027	13:00-14:30 UTC pit / inventory prints; violent, two-sided
Quiet drift	0.110	0.59	-0.015	-0.014	low amplitude, no build — genuinely low signal

The cleanest, most repeated miss is the **vol-expansion break**. We sat out **21 of 28** of them — three quarters — and they carry a legible fingerprint: `atr_pct` already climbing up through about 0.15 (well above the quiet-drift 0.11) and VWAP distance already positive by +0.12 with a gentle slope the direction the break then goes. In plain words: **the range starts widening and price is already leaning before the leg fires**. That is a momentum-continuation footprint, not a mystery. The morning of 07-08 is the textbook case — the \$230 whale at 08:11 followed by \$80 pushes at 08:27 and 08:48, then more at 12:04 and 12:54, all vol-expansion, all sat out, one clean up-trend we watched from the bench.

§7.4 4 • The L2 book — nothing to show, and that's the point

NO MCL BOOK EXISTS

The depth window (Thu 07-09 14:17 → Fri 07-10 21:00 UTC) covers MNQ, MES, MGC and M2K. **MCL is not in depth.db at all** — zero snapshots — because the NYMEX L2 entitlement returns "not subscribed". Fourteen MCL runs fall inside that window (07-09 14:49, 07-10 07:10 through 15:53, and so on)

and for not one of them can we say whether the far side thinned out ahead of the move or got eaten, or whether book imbalance leaned the right way. There is no book tell to report because there is no book. This is not "too noisy" — it is **missing data**, and it is the top of the MCL wishlist: a NYMEX depth entitlement so crude gets the same vacuum-vs-absorption read we now have on MNQ/MES/MGC.

§7.5 5 • The honest money left on the table

Add up the dollar size of all 78 runs and it is a large hindsight ceiling — but most of it is not capturable by any sim shape we own, because a \$40 15-minute wiggle is not a tradable leg once you price the crude spread and slippage. The capturable money is the subset with a real footprint and a real exit that holds: the vol-expansion and open/news legs. When momentum_shadow did catch one it was clean — + \$100.76 on the 07-08 14:13 leg, +\$82.64 on 07-10 13:34, +\$47.76 on 07-09 08:53. Extrapolating the +\$443 we banked on 38 with-run fires to the ~50 similar runs we sat out suggests order **\$400-\$600 of realistically-capturable run money left on the table this week** — but that is one week, in-sample, and it comes bundled with the fade losses you have to stop first.

§7.6 6 • How we grab them from Monday

Vol-expansion break — the highest-value miss (21/28 sat out). **SHADOW-FIRST**

Trigger, stated concretely: when MCL atr_pct crosses up through ~0.13-0.15 **and** VWAP distance is already positive in the break direction (rvol around 0.7-0.8), go **with the break** — long on an up-lean, short on a down-lean. Exit: the existing momentum_shadow chandelier that holds ~30 minutes and lets

the winner run. We already own the entry and the exit; momentum_shadow is the only +EV MCL shape. The gap is coverage — it fired on just 6 of 28 vol-expansion runs. The action is to **widen momentum_shadow's amplitude/lean gate on MCL and watch it in shadow** for a week or two before arming. Shadow-first, not arm-now.

Open / news spike — best dollars when caught, but violent (10/16 sat out). SHADOW-FIRST, CAUTIOUS

Trigger: inside the 13:00–14:30 UTC US pit / inventory window, with atr_pct around 0.15, take the **first clean break direction** with the same momentum_shadow exit. This cluster paid the most per fire (+\$220 real across 5). But it is genuinely two-sided event risk — the 07-10 14:33 print whipped $2.7\times$ range in both directions — and MCL's crude spread eats thin edges. Arm only after a shadow soak, and honestly accept half of these are news coin-flips we cannot forecast from bars alone.

Quiet drift — do not chase, and stop fading. KILL THE FADES

These 34 low-amplitude drifts (atr% 0.11, no build) carry no forecastable signature — sitting most of them out is correct. The prescription here is not a new entry, it is subtraction: **stop the reversion fades that bleed between the runs.** passive_chop and tw_mcl_short_with_strength together cost ~\$1,000 of real shadow money this week fading into moves; on MCL specifically they should stand down. That, not a new entry, is the biggest single MCL improvement available.

THE PLAIN TAKEAWAY — MCL

We are flying blind on crude — no book, almost no tape — and it shows: we sat out 50 of 78 real runs, including the \$230 whale. But the runs are not where we lose money; around them we net about flat. The $-\$1,716$ came from fading the chop in between. Two moves for Monday: **(1) widen momentum_shadow (the one +EV MCL shape) to cover vol-expansion breaks, shadow-first; (2) stand the reversion fades down on MCL.** And put a NYMEX depth entitlement on the wishlist — until crude has a book, everything here is half-sighted.

§8 The biggest runs on MBT this week — did we show up?

Micro-Bitcoin threw off **107 big runs** this week — 15-minute lurches of 232 points or more, which for MBT at \$0.10 a point is roughly \$23+ of range in a quarter of an hour. We fired a sim in the right direction on just **19 of them (18%)**, sat out the other 88, and still managed to lose **\$2,048 of honest money** in the attempt. The runs are real, they are large, and they cluster at a time of day we can see coming a mile off. We just keep showing up in the wrong session with the wrong shape.

READ THIS FIRST — TWO DATA CAVEATS FOR MBT

MBT is **not** in the L2 order book (depth.db only carries MNQ / MES / MGC / M2K), so Section 4 has no book to show you — I use the aggressor tape instead and tell you plainly it doesn't lead. And MBT was the contract dropped from the tape booster (6 names kept, MBT cut), so its aggressor feed is thin — 86k ticks over six days vs the dense MNQ tape. Treat every "flow" number here as low-confidence. The one thing MBT has in spades is a clean, dense 5s bar record, so the run census itself is solid.

§8.1 1 · The full census — every run, not a top-eight

The scorecard found **107 runs in-week** ($\geq 1.5\times$ the contract's typical 15-min range of 155pt). Here is the whole list. Rows in **bold** are the 19 where a sim actually fired in the run's direction within ± 5

minutes; everything else we watched go by. "Amp" is the pre-run ATR amplitude — how hot the tape already was.

RUN START (UTC)	DIR	MOVE (PT)	US	AMP	CLUSTER
07-06 00:15	DN	-300.0	caught	0.282	UNCLASSIFIED
07-06 00:34	UP	+325.0	caught	0.299	UNCLASSIFIED
07-06 01:38	DN	-320.0	caught	0.179	VACUUM
07-06 03:17	DN	-255.0	sat out	0.134	UNCLASSIFIED
07-06 05:17	DN	-255.0	sat out	0.181	UNCLASSIFIED
07-06 08:22	DN	-250.0	sat out	0.166	UNCLASSIFIED
07-06 11:49	DN	-1090.0	caught	0.119	UNCLASSIFIED
07-06 12:32	DN	-455.0	sat out	0.217	UNCLASSIFIED
07-06 12:56	DN	-365.0	sat out	0.226	VACUUM
07-06 13:23	DN	-260.0	sat out	0.251	OPEN / NEWS spike
07-06 13:38	UP	+650.0	caught	0.561	OPEN / NEWS spike
07-06 14:10	UP	+455.0	sat out	0.218	OPEN / NEWS spike

07-06 14:25	DN	-320.0	caught	1.200	OPEN / NEWS spike
07-06 14:54	UP	+620.0	caught	0.185	OPEN / NEWS spike
07-06 15:36	UP	+1130.0	caught	0.288	UNCLASSIFIED
07-06 15:58	UP	+410.0	sat out	0.385	VACUUM
07-06 16:13	DN	-385.0	sat out	0.351	VOL-EXPANSION break
07-06 16:29	UP	+240.0	caught	0.338	VOL-EXPANSION break
07-06 19:54	DN	-280.0	caught	0.148	VACUUM
07-06 20:19	UP	+265.0	sat out	0.047	UNCLASSIFIED
07-06 20:47	UP	+920.0	sat out	0.133	UNCLASSIFIED
07-06 21:10	DN	-415.0	sat out	1.294	VOL-EXPANSION break
07-06 21:38	UP	+455.0	sat out	0.607	VOL-EXPANSION break
07-06 21:53	DN	-265.0	sat out	0.834	VOL-EXPANSION break
07-06 22:18	DN	-260.0	sat out	0.217	VACUUM
07-07 00:14	UP	+285.0	caught	0.070	VACUUM
07-07 00:54	DN	-290.0	sat out	0.210	UNCLASSIFIED

07-07 02:19	DN	-360.0	sat out	0.180	UNCLASSIFIED
07-07 02:46	DN	-285.0	sat out	0.102	UNCLASSIFIED
07-07 04:59	DN	-275.0	sat out	0.103	UNCLASSIFIED
07-07 05:28	UP	+265.0	sat out	0.151	UNCLASSIFIED
07-07 11:35	UP	+265.0	sat out	0.095	UNCLASSIFIED
07-07 11:59	UP	+430.0	sat out	0.110	UNCLASSIFIED
07-07 12:50	DN	-235.0	sat out	0.149	VACUUM
07-07 13:14	DN	-275.0	sat out	0.165	OPEN / NEWS spike
07-07 13:49	DN	-380.0	sat out	0.441	OPEN / NEWS spike
07-07 14:06	DN	-245.0	sat out	0.458	OPEN / NEWS spike
07-07 14:21	DN	-285.0	sat out	0.238	OPEN / NEWS spike
07-07 14:44	UP	+700.0	sat out	0.262	OPEN / NEWS spike
07-07 15:08	UP	+290.0	sat out	0.267	VACUUM
07-07 15:30	UP	+310.0	sat out	0.188	UNCLASSIFIED
07-07 16:36	UP	+255.0	sat out	0.140	VACUUM

07-07 18:46	DN	-535.0	sat out	0.210	VACUUM
07-07 19:03	DN	-250.0	caught	0.306	VOL-EXPANSION break
07-07 19:18	UP	+340.0	sat out	0.393	VACUUM
07-07 21:24	DN	-405.0	sat out	0.204	UNCLASSIFIED
07-07 21:39	UP	+405.0	sat out	0.711	VACUUM
07-07 22:04	UP	+280.0	sat out	0.355	VOL-EXPANSION break
07-08 01:23	DN	-635.0	sat out	0.094	UNCLASSIFIED
07-08 01:48	DN	-385.0	caught	0.197	UNCLASSIFIED
07-08 02:21	DN	-295.0	sat out	0.183	UNCLASSIFIED
07-08 02:36	UP	+450.0	sat out	0.383	VACUUM
07-08 03:39	DN	-245.0	caught	0.111	UNCLASSIFIED
07-08 08:15	DN	-770.0	sat out	0.127	VACUUM
07-08 08:32	UP	+340.0	sat out	0.282	UNCLASSIFIED
07-08 08:47	DN	-300.0	sat out	0.672	VOL-EXPANSION break
07-08 09:04	DN	-340.0	sat out	0.273	UNCLASSIFIED

07-08 09:19	UP	+390.0	sat out	0.275	UNCLASSIFIED
07-08 12:30	DN	-235.0	caught	0.128	VACUUM
07-08 12:52	DN	-255.0	sat out	0.088	UNCLASSIFIED
07-08 13:25	UP	+305.0	sat out	0.194	OPEN / NEWS spike
07-08 13:40	DN	-430.0	sat out	0.249	OPEN / NEWS spike
07-08 13:55	UP	+480.0	sat out	0.259	OPEN / NEWS spike
07-08 14:35	DN	-330.0	sat out	0.266	OPEN / NEWS spike
07-08 15:00	DN	-340.0	sat out	0.218	VACUUM
07-08 16:39	UP	+240.0	sat out	0.194	UNCLASSIFIED
07-08 17:18	UP	+335.0	sat out	0.145	VACUUM
07-08 17:39	DN	-285.0	sat out	0.144	UNCLASSIFIED
07-08 18:21	DN	-260.0	sat out	0.160	VACUUM
07-08 20:52	DN	-275.0	sat out	0.088	UNCLASSIFIED
07-08 21:22	UP	+280.0	sat out	0.129	UNCLASSIFIED
07-09 00:55	UP	+415.0	sat out	0.153	UNCLASSIFIED

07-09 01:29	DN	-375.0	sat out	0.216	UNCLASSIFIED
07-09 02:21	DN	-290.0	sat out	0.193	UNCLASSIFIED
07-09 03:20	UP	+315.0	sat out	0.267	UNCLASSIFIED
07-09 04:27	UP	+295.0	sat out	0.016	UNCLASSIFIED
07-09 05:11	DN	-285.0	sat out	0.247	VACUUM
07-09 05:46	UP	+305.0	sat out	0.056	UNCLASSIFIED
07-09 06:12	UP	+335.0	caught	0.120	UNCLASSIFIED
07-09 06:47	UP	+345.0	caught	0.111	VACUUM
07-09 07:37	UP	+335.0	sat out	0.183	VACUUM
07-09 12:39	DN	-305.0	sat out	0.301	VOL-EXPANSION break
07-09 12:57	UP	+295.0	sat out	0.080	VACUUM
07-09 13:17	DN	-255.0	sat out	0.254	OPEN / NEWS spike
07-09 13:32	UP	+460.0	sat out	0.439	OPEN / NEWS spike
07-09 13:53	DN	-510.0	sat out	0.781	OPEN / NEWS spike
07-09 14:22	UP	+300.0	sat out	0.207	OPEN / NEWS spike

07-09 14:54	UP	+335.0	sat out	0.238	OPEN / NEWS spike
07-09 15:23	DN	-315.0	sat out	0.465	VACUUM
07-09 16:11	DN	-295.0	sat out	0.206	VACUUM
07-09 17:35	UP	+315.0	sat out	0.159	UNCLASSIFIED
07-09 17:53	UP	+380.0	sat out	0.111	UNCLASSIFIED
07-10 00:19	DN	-290.0	sat out	0.166	UNCLASSIFIED
07-10 01:14	UP	+360.0	sat out	0.158	VACUUM
07-10 01:33	UP	+560.0	sat out	0.142	UNCLASSIFIED
07-10 04:32	UP	+265.0	sat out	0.125	UNCLASSIFIED
07-10 08:14	UP	+340.0	caught	0.078	VACUUM
07-10 09:13	UP	+260.0	sat out	0.186	UNCLASSIFIED
07-10 13:37	UP	+395.0	sat out	0.248	OPEN / NEWS spike
07-10 14:06	DN	-480.0	sat out	0.317	OPEN / NEWS spike
07-10 15:01	UP	+275.0	sat out	0.140	VACUUM
07-10 15:30	DN	-250.0	sat out	0.163	UNCLASSIFIED

07-10 16:37	UP	+280.0	caught	0.320	VACUUM
07-10 17:04	DN	-260.0	sat out	0.117	UNCLASSIFIED
07-10 19:23	UP	+250.0	sat out	0.219	UNCLASSIFIED
07-11 09:16	UP	+415.0	sat out	0.000	UNCLASSIFIED
07-11 09:33	UP	+275.0	sat out	0.708	VOL-EXPANSION break

The scorecard buckets each run by what caused it. Across the 107:

CLUSTER	RUNS	SAT OUT	AVG MOVE	AVG AMP	WHAT IT IS
UNCLASSIFIED	48	41	368pt	0.160	Quiet drift, no tape story — it just walks
VACUUM	28	21	339pt	0.223	Moved against the (thin) tape — air pocket
OPEN / NEWS spike	21	18	403pt	0.354	US-open bursts — biggest moves, all 13:00-15:00
VOL-EXPANSION break	10	8	317pt	0.577	Tape already exploding when it goes

I spot-checked the two monsters against the raw 5s bars myself. The **07-06 15:36 +1130** up-run: close ran 62,460 → 63,475, a clean +1,015 in fifteen minutes, high-to-low range 1,040 — real. The **07-06 11:49 –1090** down-run: close fell 63,105 → 62,120, –985 — also real. The census sizes are honest.

§8.2 2 · Participation — we caught 19, and lost the money anyway

Here is the blunt version. Of 107 runs we fired in-direction on 19 (18%) and sat out 88. But firing is not the same as making money. Across **705 MBT sim trades this week (511 filled), the honest tally is −\$2,047.75**. The optimistic sim number was −\$1,234.92, so even the doomed entries we took were understated by the sim by about 1.7×, exactly as the desk keeps finding. We are not sitting out because we're being disciplined — we're firing plenty, we're just firing the wrong shapes.


STRATEGY	TRADES	WINS	HONEST P&L	VERDICT
passive_chop	365	19	-\$860.75	KILL
tw_mbt_fade_euphoria	50	3	-\$312.00	KILL
tw_deadtape_standdown	55	11	-\$287.50	KILL
tw_clock_standdown	47	11	-\$234.00	KILL
pullback_quiet	30	8	-\$167.00	KILL
tw_strength_ceiling_pb	10	4	-\$78.50	KILL
tw_mbt_pb_rvol_band	8	0	-\$46.50	KILL
passive_chop_calm	15	1	-\$45.00	KILL
momentum_shadow	75	31	-\$37.00	TUNE
tw_mbt_spike_fade	1	0	-\$12.00	KILL
tw_us_open_surge_cont	6	3	+\$11.00	KEEP
orb_shadow	43	19	+\$21.50	KEEP

Ten of twelve shapes bled. The two that made anything — **orb_shadow (+\$21.50)** and **tw_us_open_surge_cont (+\$11.00)** — are both opening-break continuation shapes. And the single worst line, **passive_chop at −\$860**, is a fade — it sells strength and buys

weakness into a market that this week simply kept going. That is the whole chapter in two rows: **fading MBT lost, going with the break made a little**. Momentum_shadow (−\$37 on 75, 31 wins) is the honourable near-miss — the right instinct, poorly exited.

§8.3 3 · The footprint — when MBT runs, and why we're in the wrong session

The single loudest signal in this whole chapter is a clock. Sort every run by hour of day and MBT lights up like a switchboard at the US cash open — **29 of 107 runs land in the 13:00-15:59 UTC window**, and the OPEN/NEWS cluster is entirely in the 13:00 and 14:00 hours (11 runs at 13:00, 10 at 14:00). Bitcoin trades the S&P's open.



MBT big runs by UTC hour, in-week. Gold = the 13:00-15:00 US-open window (29 runs). Leftmost bar = 00:00 UTC, rightmost = 23:00.

Now overlay where the money went. Split the sim P&L by session and the picture is upside down:

SESSION (UTC)	SIM TRADES	HONEST P&L	READ
London 05:00–12:59	449	-\$1,161.80	Where we over-trade — fading quiet drifts
US-pm 17:00–21:59	40	-\$388.50	Thin, still bleeding
	131	-\$268.50	Overnight chop fades

Asia 22:00–
04:59

US-open
13:00–16:59

85

-\$229.00

**The best hours — and we
barely trade them**

THE MAP IS BACKWARDS

We do **five times** the trading in the quiet London window (449 trades, −\$1,162) as we do in the US-open window (85 trades, −\$229), yet London is where the fades get run over and the US open is where the only green shapes live. We're over-trading the dead session and under-trading the alive one.

Reading the pre-run footprint by cluster: the **UNCLASSIFIED** drifts (48 runs, amp 0.160) have no signature you can trade — the tape is asleep and price just walks, which is precisely what mauls `passive_chop`. The **VOL-EXPANSION** breaks (10 runs, amp 0.577) announce themselves — ATR is already screaming — but you're late by the time you see it. The **OPEN/NEWS** cluster is the sweet spot: 21 runs, the biggest average move (403pt), amp already firming to 0.354, and — crucially — **a clock you can set**. That is the highest-value thing we repeatedly sit out with a clean, forecastable signature.

§8.4 4 • The book — MBT has none, and the tape doesn't lead

The brief asks for the L2 read on the runs inside the depth window. For MBT the honest answer is short: **there is no L2 book for MBT**. The depth capture covers MNQ, MES, MGC and M2K only — MBT was never entitled or booked. So I can't show you a far-side vacuum or an absorption print, because we never recorded the ladder.

What MBT does have is the aggressor tape (86k prints, 07-05 → 07-10), so I checked whether the tape leads the runs. It doesn't.

Three OPEN/NEWS examples, buy vs sell volume in the 10 minutes before the run:

RUN	DIR	BUY VOL	SELL VOL	DID THE TAPE CALL IT?
07-06 13:38	UP +650	788	522	Weak yes — 1.5:1 buy lean
07-07 14:44	UP +700	1,005	973	No — dead balanced before a 700pt rip
07-09 13:53	DN -510	540	274	Backwards — buy-heavy right before a -510 drop

That is why the scorecard tags so many MBT runs as VACUUM ("moved against the tape") — the thin MBT tape genuinely does not lead the move, and on the biggest bursts it's either flat or pointing the wrong way. So for MBT, **do not build a flow-triggered entry** — the tape we have isn't good enough and there's no book to lean on. The signature that works is the clock and the amplitude, not the order flow. If we wanted a real book tell on MBT we'd need to entitle and capture the CME Bitcoin depth feed — flag that as the data gap, not a this-week edge.

§8.5 5 · The honest money left on the table

Add up the raw magnitude of all 107 runs and you get **38,805 points** of directional travel this week. At MBT's \$0.10 a point that is **\$3,880 of hindsight range on a single micro contract** — the theoretical ceiling if you bought every low and sold every high. Nobody captures that; it's the fantasy number. Set it against reality: the sims that actually traded **lost \$2,048**, and even their best-case exits (the ceiling on the trades we took) came to -\$1,235. So the gap isn't "we captured 60% and could get to 80%" — it's that **we're on the wrong side of the ledger entirely**, converting a market full of big clean

moves into a four-figure loss. The capturable money is real and large; the shapes pointed at it are not.

§8.6 6 • How we grab them from Monday

One cluster, the OPEN/NEWS burst, is both the biggest and the only one with a signature we can act on ahead of time. Here is the concrete plan, plus the housekeeping that stops the bleed.

§8.7 Lead — the US-open continuation break

SHADOW-FIRST

When: the 13:00–15:00 UTC window only (US cash open).

Trigger: MBT amp/ATR firming through ~0.30 and price breaking the 13:00–13:30 opening range in the break direction.

Direction: with the break — long the up-break, short the down-break, no fading. **Exit:** the RATCHET/chandelier that already carries `tw_us_open_surge_cont`'s two winners (07-09 13:51 LONG +\$28.5, 07-08 14:52 LONG +\$16, both closed on RATCHET) — let it run and trail, don't cap it at a fixed R. **Why shadow-first, not armed:** the two green shapes are real but thin — `orb_shadow`'s usable +\$100 came almost entirely from one day (07-06, riding the +1130 monster), and `us_open_surge_cont` is +\$11 on 6 trades. This is one in-sample week — promising, not proven. Widen `tw_us_open_surge_cont`'s window to the full 13:00–15:00 band and observe another 2–3 weeks before it earns a live nomination.

STAND DOWN — STOP FEEDING THE LONDON FADES

The −\$1,162 London-session loss is almost all `passive_chop` and the other fades selling into quiet UNCLASSIFIED drifts that

keep walking. **Disarm passive_chop, tw_mbt_fade_euphoria and pullback_quiet on MBT outright** — MBT is a trending, open-driven crypto contract, not a mean-reverter, and every fade shape confirmed that again this week. This doesn't make money (it's shadow), it just stops us manufacturing a −\$860 line item. That's a "trade less," and on MBT trading less is the correct move.

§8.8 Watch, don't build — the other three clusters

VOL-EXPANSION (amp 0.577): the move is already underway when the signature appears, so an entry is late — park it.

VACUUM & UNCLASSIFIED: these are the quiet air-pocket drifts with no forecastable pre-run tell in anything we capture — no book, weak tape. Genuinely unforecastable from current data. The only thing that would change that is a real MBT L2 depth feed to watch the far side thin out; until we entitle it, don't pretend there's an entry here.

Takeaway for MBT: the market handed us 107 big runs and \$3,880 of hindsight range on one micro, and we turned it into a −\$2,048 loss by fading quiet London drifts we should never touch while barely trading the 13:00–15:00 US-open bursts where the only green shapes live. Monday: disarm the MBT fades, and shadow-widen the US-open continuation break with a ratchet exit. Don't build a flow trigger — MBT has no book and its tape doesn't lead.

CHAPTER

The Through-Lines — What The Whole Week Is Saying

⚠ A SKEPTIC RE-RAN THE MONEY CLAIM — IT FELL (READ AS A LEAD, NOT A RULE)

CURVE-FIT / MISLABELED PARTITION — does not hold as stated. The figures reproduce exactly, but +\$2,441 is just the 9 momentum strategies with $pnl \geq 0$, and -\$11,040 is the arithmetic complement (week_total -8,599 minus the winners), not a real fade book. The complement hides ~\$1,190 of with-the-run momentum LOSERS (ride_strength, mcl_short_with_strength, surge_cont, open_momentum, no_fade_drive); the genuine fade/reversion book lost only -\$6,437. Cutting the same week at the P&L sign line and labeling winners 'momentum' / losers 'fade' is the curve-fit. The honest, label-consistent with-the-run momentum book is +\$1,249 (~£975), about half the claim — and that is an optimistic-sim CEILING (force-kill+slippage absent; live losses run 1.1-3.1x modeled, losers understated). Its biggest piece, tw_mnq_thrust_cont +\$992, is 28 trades in one week (thin-n, one regime, unvalidated). Strip it and the only adequately-powered green signal is momentum_shadow (n=413, +\$749 ~£585). Momentum green is directionally real (consistent with the desk momentum-+EV thesis, not noise), but not bankable at the claimed magnitude and not proven by the fabricated fade number. $real_gbp = \text{honest momentum book} + \$1,249 \sim £975$, optimistic ceiling; durable core nearer £585.

§9 How We Grab The Runs From Monday

This is the capstone you asked for — not another contract-by-contract lab, but the one desk-wide answer to “where was the money this week, did we touch it, and exactly how do we get involved on Monday?” I ran the full run census (footprint_scorecard.py missed --top 100) on all seven contracts, tagged every big move by its cause, and joined the

honest money (real_pnl, never the optimistic sim figure). The short version: **the runs are real, plentiful, and mostly catchable — and we already have one family of sims that catches them and prints green.** The bleed isn't that we can't find the runs; it's that the fade/stand-down sims spent the week fighting them.

BOTTOM LINE

Across all seven contracts this week there were **540 runs of $\geq 1.5\times$ typical range**. We were involved (any sim fired the right way) in just **127** of them and **sat out 413** — three-quarters of the money in motion. Yet the honest scoreboard says the direction is right: the **with-the-run family made +\$2,441** real this week (led by tw_mnq_thrust_cont at **+\$992**, 28 trades, 61% win, +\$35/fire — and its real P&L beat its optimistic sim, so it's honest), while the whole desk netted **−\$8,599** because the fades and stand-downs gave back **−\$11,040** fighting the same moves. The plan below is one sentence: **stop fading the runs, and extend the confirmed-momentum entry to the clean-flow runs we keep sitting out.** All figures USD; GBP $\approx$ USD $\times 0.79$.


§9.1 1 The desk-wide scoreboard — where the money ran this week

Every contract's runs, how many we were involved in versus sat out, and the cluster that dominated its sit-outs. "Caught" means at least one shadow sim fired in the run's direction — a generous bar, and we still miss most of them.

CONTRACT	RUNS $\geq 1.5\times$	CAUGHT	SAT OUT	SAT %	DOMINANT SIT- OUT CLUSTER	FLOW DATA?
----------	--------------------------	--------	------------	----------	------------------------------	---------------

MNQ	76	17	59	78	Vacuum / Flow-led	tape + L2 book
MES	63	10	53	84	Vacuum / Open-news	tape + L2 book
MGC	59	25	34	58	Vacuum / Unclassified	tape + L2 book
M2K	65	17	48	74	Unclassified	thin / none
MYM	73	11	62	85	Unclassified	none
MCL	78	28	50	64	Unclassified / Vol-exp	none
MBT	126	19	107	85	Unclassified / Vacuum	tape only
DESK	540	127	413	76	see §2	—

Now the same 413 sat-out runs, sorted by cause — because the cause is what tells us whether a run was catchable. The chart plots, per cause cluster, how many runs there were (navy) and how many we sat out (red).



Cause clusters, desk-wide (runs / sat out): Unclassified 209/160 · Open-news 138/107 · Vacuum 123/92 · **Flow-led 35/28** · Vol-expansion 35/26. Navy = total runs, red = ones we sat out. Flow-led is drawn green because it is the one cluster with a clean, tradable pre-run tell.

§9.2 2 The honest money — the direction was right, the fades killed it

Before any new build, look at what the existing sims already earned. Split the book into the sims that go with a run and the ones that fade or stand down:

FAMILY	TRADES	REAL P&L	READ
With-the-run (thrust / momentum / coil-break / ORB)	899	+ \$2,441	Going with the move already pays
Fades / stand-downs / reversion / pullback	2,898	– \$11,040	This is the entire desk loss and then some
Desk net	3,797	– \$8,599	Bleed is fighting the runs, not missing them

The single cleanest proof sits at the top of the leaderboard: `tw_mnq_thrust_cont` — enter WITH a confirmed MNQ thrust, ride it on a chandelier — made **+\$992 real on 28 trades at 61% win (+ \$35/fire)**, and its real money actually beat its optimistic sim (+\$723). That is the shape we want to point at every run we currently sit out. Its cousins agree: `tw_mgc_momo_thrust` +\$127 (56% win), `tw_mym_open_thrust` +\$28 (58% win), `momentum_shadow` +\$749. Meanwhile the reversion book — `passive_chop` –\$4,249, `tw_deadtape_standdown` –\$960, `tw_clock_standdown` –\$891, `tw_ride_strength` –\$627 — is where the –\$11k went. It kept fading the exact moves the thrust sims were paid to ride.

§9.3 3 The Monday capture play for each cause cluster

For each cluster: the tell you can actually observe at entry, the direction, the exit shape (which existing sim takes it), and arm-now versus shadow-first. Ranked by sat-out money addressed $\times$ how clean the signal is.

ARM NOW 1. **FLOW-LED runs (aggressors drove it)** — 35 runs, 28 sat out, and the only cluster with a clean, forward-observable tell. **Trigger:** one-sided aggressor tape (trade_tick buy/sell imbalance strongly in the run direction over the last 30-60s) confirmed by the L2 book leaning the same way (depth_snap.imbalance same sign) and a real thrust (atr_pct rising, rvol ≥ 1.5). **Direction:** WITH the flow. **Exit:** the confirmed-momentum + chandelier shape that already works — i.e. tw_mnq_thrust_cont. **Where:** MNQ (proven, deep tape + 360k book snaps) and MES (deep tape + book). This is the highest-value play because it stacks the most sat-out runs that carry a clean signature on the contracts where we actually have the data.

SHADOW FIRST 2. **OPEN / NEWS runs** — 138 runs, 107 sat out, the second-biggest bucket. **Tell:** the timing is knowable (the 13:30-14:35 UTC US-open / data window fires nearly every day, on every contract). **But the direction is a coin-flip and the tails are brutal** — the 07-10 14:33 spike alone showed a — \$24k optimistic swing on MES. **Play:** do NOT predict direction. Wait for the first 5-minute range to set, then take the break with confirming flow, small size, hard stop. **Exit:** ORB-style target = $1 \times$ opening-range height. Shadow-first — the edge is

real in timing but unproven in direction, and one bad tail erases a month.

DO NOT BUILD AN ENTRY **3. VACUUM runs (moved AGAINST the tape)** — 123 runs, 92 sat out. These are the second-biggest bucket and the most seductive, but the name is the verdict: **price ran while the aggressor tape pointed the other way** (stop-runs, liquidity air-pockets). Any flow-based entry would have been pointed backwards into them. This is precisely where the fade book got run over. The lesson is defensive, not offensive: **these are un-tradeable from flow — the value is telling the reversion sims to stop fading into them.**

CAN'T — DATA GAP **4. UNCLASSIFIED runs (no tape / quiet)** — 209 runs, 160 sat out, the biggest bucket of all. Sounds like the jackpot; it isn't. **183 of those 209 are on M2K, MYM, MCL and MBT — the four contracts where we have essentially no aggressor tape and no L2 book** (M2K 4.9k ticks, MCL 2.0k, MYM 5.3k, and no book on any). “Unclassified” here means “the classifier is blind because we never captured the data,” not “a quiet edge waiting.” On the three tape-covered contracts unclassified collapses to just 26 runs. This bucket is a **capture-coverage problem, not a strategy.**

§9.4 4 Forecastable vs not — be blunt

CLUSTER	FORECASTABLE AT ENTRY?	WHY / WHAT'S MISSING
---------	------------------------	----------------------

Flow-led	Yes	One-sided tape + book lean is visible before the run, on MNQ/MES/MGC
Open-news	Timing yes, direction no	We know when (13:30 UTC) but not which way; needs a break-confirm rule + small size
Vacuum	No	Signal actively misleads — flow points the wrong way; needs resting-liquidity / iceberg data we don't have
Unclassified	No (yet)	Missing data: aggressor tape + L2 book on M2K, MYM, MCL (and book on MBT)
Vol-expansion	No	MCL/MBT only, and both lack the tape/book to confirm a break from a fakeout

§9.5 5 The single highest-value build

BUILD THIS ONE THING

A flow-led thrust entry on MNQ and MES. It is the cluster we most repeatedly sit out (28 of 35) that also carries a genuinely clean, forward-observable signature — and we don't even need a new exit for it. The confirmed-momentum + chandelier machinery from `tw_mnq_thrust_cont` (already + \$992 real, honest, 61% win) is the exit; the only new part is the entry gate: fire WITH the run when (a) aggressor tape is one-sided in the run direction over the last 30–60s, (b) the L2 book imbalance leans the same way, and (c) `rvol` ≥ 1.5 with `atr_pct` expanding. MNQ has the deepest tape and 360k book snaps to prove it; MES is the natural second. Everything else on this page is either already working (lean harder), a coin-flip that needs a cautious shadow (open-news), a trap to stop fading (vacuum), or blocked until we capture the tape (M2K/MYM/MCL).

§9.6 6 Monday action list

1. **ARM** Keep the with-the-run family running at full and let `tw_mnq_thrust_cont` keep firing — it's the honest +\$992 anchor and the template for everything below.
2. **BUILD + SHADOW** Write the flow-led thrust entry (one-sided tape + book lean + rvol/atr expansion, WITH the run) on the existing chandelier exit; shadow it on MNQ and MES from Monday's open. This is the \$5 build.
3. **STOP THE BLEED** Cut the reversion / stand-down sims that fade into vacuum runs (`passive_chop`, `tw_deadtape_standdown`, `tw_clock_standdown`, `tw_ride_strength`) — they are the entire −\$11k. Fading the runs is the loss.
4. **SHADOW, SMALL** Add an open-news first-range-break rule (13:30 UTC window, break-with-flow, ORB target, hard stop, tiny size) as observe-only. Right timing, unproven direction, fat tails — prove it before a cent.
5. **CAPTURE** Extend aggressor-tape and L2-book capture to M2K, MYM and MCL. Until then their 183 “unclassified” sat-out runs are un-diagnosable — the biggest bucket on the board is dark for lack of data, not lack of edge.

Takeaway: the runs are there and we already know how to ride one — the with-the-run sims made +\$2,441 this week and `tw_mnq_thrust_cont` is honest green. Monday isn't about finding a new edge; it's about pointing the proven confirmed-momentum entry at the clean-flow runs we keep sitting out, and telling the fade book to stop walking into the vacuum. Everything past that is either a

cautious shadow or a data-capture job — and one week of in-sample runs means these are hypotheses to shadow, not truths to size.

✓ **A SKEPTIC RE-RAN THE MONEY CLAIM — IT HELD**

HELD — reproduces to the dollar on honest tick-repriced fills, but it's a thin-n, MNQ-concentrated one-week read, not yet a bankable forward edge. REPRODUCTION (data/shadow_real.db::shadow_real_pnl real_pnl, JOIN alphabot.db::fut_shadow_sim_trades on trade_id, entry_ts in Mon 6 Jul 00:00 → Sat 11 Jul 00:00 UTC = 1783296000-1783728000): - CONFIRMED (volume/rvol-gated, by MECHANISM): momentum_shadow, orb_shadow, tw_mnq_thrust_cont, tw_mes_thrust_cont, tw_m2k_vst, tw_m2k_momo_loose, tw_mgc_momo_thrust, tw_us_open_surge_cont, tw_mgc_coil_break_long = +\$2,343.55 over 909 trades $\approx$ the reported +\$2,344. - LOOSE (fires on any strength, no volume gate): tw_ride_strength, tw_mcl_short_with_strength, tw_mym_open_thrust, tw_m2k_open_momentum = -\$1,017.08 over 509 trades $\approx$ the reported -\$1,017. - Per-contract holds directionally: confirmed net is positive on MNQ +\$1,804, M2K +\$257, MGC +\$332, MES +\$138; negative/flat on MYM -\$108, MCL -\$11, MBT -\$5. So "edge on the index+gold micros (MNQ/MGC/M2K/MES), not crude/bitcoin/YM" is TRUE this week. WHY IT'S ONLY HALF-BANKABLE: 1. THIN-N / CONCENTRATION on the winning side. 74% of the +\$2,344 is two sims: tw_mnq_thrust_cont +\$992 (just 28 trades, 61% win) and momentum_shadow +\$749 (413 trades). The single 28-trade MNQ thrust sim is 42% of the whole confirmed profit; mgc_momo_thrust adds +\$127 on 9 trades. Strip tw_mnq_thrust_cont and confirmed falls to ~\$1,352. Most confirmed sims win only 35-47% — profit rides on a handful of big asymmetric winners, which is the legit momentum signature but fragile at this n. This side is a one-week, in-sample-adjacent read (thrust_cont sims were explicitly flagged one-week curve-fits being shadowed for OOS — this IS their first OOS week, promising not proven). 2. The LOSS half is the ROBUST half. ride_strength -\$627 (198 tr, 38% win) and mcl_short_with_strength -\$341 (295 tr, 35% win) bleed

structurally on broad n. Honest fills also expose the sim's phantom: mcl_short's optimistic sim showed +\$958 but honest repricing is -\$341 (a ~\$1,300 fidelity gap) — so "loose bleeds" is real and understated by the sims, exactly as the memory warns. 3. TAXONOMY is defensible-but-not-unique. tw_m2k_momo_loose (+\$163) is NAMED "loose" yet placed in CONFIRMED (correct by mechanism — it carries an $r_{vol} \geq 1.5$ gate). Reclassify by name and it's confirmed +\$2,180 / loose -\$854 — still a barbell, same shape, different exact figures. To the author's credit, m2k_vst (-\$0.4) and us_open_surge (-\$69) sit IN confirmed and DRAG it down, so the buckets were not cherry-picked to only winners. 4. NOT purely "trade less": momentum_shadow nets +\$749 over 413 trades, so the confirmed edge isn't only a selectivity artefact — but the standout winners (28- and 9-trade sims) are exactly the "trade rarely, only on confirmation" cohort that needs many more weeks to bank. BOTTOM LINE: the claim is an honest, faithful reproduction (not curve-fit in the arithmetic, not cost-blind — scored on the tick-repriced floor), and the barbell + index/gold shape survives reclassification. But the profitable half is really an MNQ-thrust story on thin n over one week; the -\$1,017 loose bleed is the durable finding. Bank the "bin the ride-strength/loose family" action; treat the +\$2,344 confirmed win as a promising single-week signal requiring 8-12 weeks / 100-200 trades before sizing. $real_gbp \approx £1,828 = \text{the confirmed } +\$2,343.55 \text{ at } \sim 0.78 \text{ GBP/USD.}$

§10 Momentum is confirmed-vs-loose, one contract at a time

READ THIS FIRST — MOMENTUM IS TWO-SIDED, NO FADING

Before anything below is misread: a momentum gate trades **both ways** — it longs a rising thrust and shorts a falling one — and it stays that way. **Never long-only.** This week the confirmed-

momentum **short** side made **+\$1,435**, more than its long side (+\$1,006); on MES and MGC the short side was the only green side. What we kill is **fading** (fighting a live move), not shorting. Don't let "confirmed vs loose" get quietly turned into "long vs short."

The momentum family split cleanly down the middle this week, and it split by mechanism, not by name. The sims that make you wait for the move to be **confirmed** — a real volume surge, a high-participation rvol gate, a trend-regime that agrees — banked **+\$2,344** across 909 shadow fills. The sims that just **ride any strength** — a 0.30-ATR blip, a down-flush, an open thrust with nothing behind it — bled **-\$1,017** across 509. Same idea, "go with the move", opposite outcome. The difference is the confirmation.

THE BARBELL


Both wings win under 45% of the time — momentum is a low-win, let-the-winner-run game whichever way you cut it. The confirmed wing survives that because its winners are nearly twice its losers (avg win **+\$26.15** vs avg loss **-\$13.75**, ~1.9:1). The loose wing does not: its winners only run ~1.35:1 (**+\$18.79** vs **-\$13.97**) and it wins less often (36.5% vs 40.9%), so it drifts red. Confirmation barely lifts the hit-rate — what it does is make the winners bigger and the losers rarer.

§10.1 The per-contract matrix — where momentum actually pays

Broken out by contract on honest `real_pnl`, the edge isn't "momentum works". It's "confirmed momentum works on the index-and-gold micros, and nowhere else." MNQ is the engine of the entire

number; MGC, M2K and MES chip in positive; crude, bitcoin and the thin Dow micro don't reward it even when you confirm.

CONTRACT	CONFIRMED \$	CONF WIN%	LOOSE \$	LOOSE WIN%	READ
MNQ	+\$1,817.9	54.3%	— \$339.7	43.6%	Momentum's home. The only cell that wins >50%.
MGC	+\$331.6	38.2%	—	—	Gold rides confirmed; nobody built a loose gold rider (good).
M2K	+\$246.3	41.9%	— \$76.5	0.0%	Russell pays confirmed; loose (n=4) went 0-for-4.
MES	+\$100.3	40.2%	— \$287.6	30.7%	Modest green confirmed, deep red loose.
MYM	−\$136.9	30.8%	+\$28.1	58.3%	Thin Dow micro: confirmed loses, loose +\$28 is n=12 noise.
MBT	−\$4.5	42.7%	—	—	Bitcoin: flat-to-negative even confirmed. No edge.
MCL	−\$11.1	35.5%	— \$341.4	35.3%	Crude marginal: ~flat confirmed, the family's #1 loose bleeder.



§10.2 The mechanism split — what "confirmed" means here

I sorted every rider by what it actually checks in code, ignoring the label. A strategy is **confirmed** if the move has to prove itself — the volume-surge thrust-continuations (`tw_mnq/mes_thrust_cont`, $1.5 \times$ volume + still-expanding), the rvol-gated momo (`tw_m2k_momo_loose` $\text{rvol} \geq 1.5$, `tw_mgc_momo_thrust` $\text{rvol} \geq 5$, `tw_us_open_surge_cont` $\text{rvol} \geq 3$), the coil-break with collapsing volume, and the live momentum gate itself (`momentum_shadow`, with its trend-regime + VWAP-slope-agreement guard). Note `tw_m2k_momo_loose` is "loose" only in its name — by mechanism it carries an rvol gate, so it sits with the confirmed set, and it made **+\$163**. A strategy is **loose** if it fires on amplitude alone with no participation check:

`tw_ride_strength` ($|\text{net}| \geq 0.30$, no volume),
`tw_mcl_short_with_strength` ($\text{net} \leq -1.0$, no volume),
`tw_mym_open_thrust`, `tw_m2k_open_momentum`.

THE FIDELITY TRAP — THE LOOSE BLEED IS WORSE THAN IT LOOKS

On the optimistic sim P&L the loose family looks harmless — **+\$66**, basically breakeven. On honest fills it's **-\$1,017**. The whole \$1,083 gap is fills the rosy sim handed out for free. The single worst offender is `tw_mcl_short_with_strength`: it prints a gorgeous **+\$958** on optimistic pnl and a real **-\$341** — a \$1,299 mirage, because shorting a crude down-flush at the perfect 1-ATR fill is a fantasy the spread never honours. And by

our own rule, even `real_pnl` is a floor: live adds force-kills and slippage the replay doesn't model, so the loose wing bleeds more than $-\$1,017$ in the wild, not less.

§10.3 Honest health warnings

This is one week, and these riders were the operator's Monday curve-fit picks that the skeptic already flagged as in-sample — shadowing them is exactly how we find out whether they hold out of sample, and one green week is not that proof. The confirmed number is also top-heavy: strip MNQ out and the other six confirmed contracts total just **+\$526** on 804 trades — still positive, but modest and easily a week of variance. The win rates are low everywhere (nothing but MNQ clears 50%), so this is a handful of big winners carrying a lot of small losers; a quiet, trendless week would flip the whole thing red. Read "confirmed momentum works" as a directional lean with one good week behind it, not a settled fact. The one number I'd trust more than most is MNQ's confirmed `tw_mnq_thrust_cont`: **+\$992** spread across 17 winners from **+\$17** to **+\$174**, not one lucky trade.

Per-contract verdict. MNQ **KEEP confirmed** MGC
KEEP confirmed M2K **KEEP confirmed** MES **confirmed only**
MYM **KILL** MBT **KILL** MCL **KILL** ride-strength family
BIN

The Monday read: if momentum ever gets armed off the shadow bench, arm the confirmed version and only on the index-and-gold micros where it paid — MNQ first (it's the whole edge), then MGC, M2K and MES for support. Bin the "ride any strength" family entirely: `tw_ride_strength` and `tw_mcl_short_with_strength` are the two biggest bleeders in the whole momentum book, and the second only looks like a winner because the sim fills it at prices crude

will never give you. Leave crude, bitcoin and the thin Dow micro out
— momentum doesn't pay there even when you make it prove itself.

§11 The exit swap: two pure archetypes, before and after

On Wednesday evening we tore out the old nine-mechanism exit spine on the shadow board and replaced it with two clean rules: a fade (reversion) now gets a 1.5-ATR stop and a 2R target, nothing else; a chase (momentum) now gets a 1-ATR stop and a chandelier that tightens from 3.5 down to 0.5. This section reads the honest money on either side of that switch. The one thing to hold onto up front: the "after" window is short — Wednesday tea-time to Friday close, one regime — so treat everything here as promising, not proven.

THE MARKER

The switch is time-stamped in `shadow_exit_archetype_changed_at.json` at 2026-07-08 15:47:28 UTC (unix 1783525648), which is about 18:00 in Paris. Every shadow entry before that instant was scored under the old spine; everything after reprices under the two pure archetypes. I split the honest board on `entry_ts` at that number and never mix the two — the "before" is a dead model. The forward window runs from the first post-switch entry at Wed 15:48 UTC to the last at Fri 21:53 UTC.

§11.1 The board, before and after

Scored on honest `real_pnl` (not the optimistic sim P&L, which understates losses), the week splits like this. "Before" is 2,354 filled shadow fills under the old spine; "after" is 999 fills under the pure

archetypes. Both halves lost money — this is still a settled-bad book — but the shape is what matters.

WINDOW	ARCHETYPE	FILLS	WIN%	HONEST P&L	PER TRADE
BEFORE (old spine)	Chase / momentum	1,627	40.6%	-\$1,263	-\$0.78
BEFORE (old spine)	Fade / reversion	727	30.3%	-\$4,456	-\$6.13
AFTER (pure archetype)	Chase / momentum	973	35.7%	-\$2,449	-\$2.52
AFTER (pure archetype)	Fade / reversion	26	34.6%	-\$430	- \$16.54

Two honest cautions before anyone reads a verdict into those per-trade numbers. First, these are not the same trades repriced two ways — the switch is a forward cutover, so "before" and "after" are different entries on different days. The momentum book looks worse per trade after the switch (-\$2.52 vs -\$0.78), but Wednesday-to-Friday was its own regime and I can't cleanly hand that to the exit change rather than to which entries happened to fire. Second, and this is the big one: the pure **reversion** exit barely got a forward test at all. Only one strategy — `passive_mgc_hours` — traded as a fade after the switch, 26 fills, all gold, all a loss. There is no honest forward read on the new 1.5-ATR-and-2R fade rule this week. It just didn't trade.


§11.2 Inside the momentum book: the give-back works, the entries don't

Where the "after" window does have enough trades to say something — the 973 momentum fills — the exit-reason mix is genuinely

encouraging, and it tells a clean story. Split those fills by how they died:

HOW IT ENDED	FILLS	WIN%	HONEST P&L	PER TRADE	AVG HOLD
CHANDELIER (trail caught it)	234	56%	+\$2,373	+ \$10.14	12.9 min
TARGET (ran to 2R)	55	100%	+\$1,447	+ \$26.31	8.2 min
SESSION_FLAT (rode to the bell)	311	52%	+\$1,014	+\$3.26	5.4 min
STOP (1-ATR hit)	373	0%	-\$7,283	-\$19.52	9.3 min

Read that top-down and the message is loud. Every winning exit route is in the black: the chandelier booked +\$2,373 across 234 trades, the 2R target went 55-for-55 for +\$1,447, and even the trades that just rode to the session bell added +\$1,014. Together the three "let it work" exits made +\$4,834. The entire loss of the momentum book — and more — is the stop line: 373 trades, not a single winner, -\$7,283. That is not an exit problem. The tightening chandelier is doing exactly its job of protecting and harvesting the trades that go your way; what's bleeding the book is the roughly two-in-five entries that go straight to the 1-ATR stop. This matches what we already believe on the desk — a clean give-back architecture is +EV on momentum, and the real work left is entry quality, not the exit knob.



Post-switch momentum book by exit route (honest P&L). The three green "let it work" bars sum to +\$4,834; the single red stop bar is -\$7,283. The chandelier is not the leak — the entries that hit the stop are.

§11.3 Is the 2R fade target too far? Yes — but that's a tune, not the fix

We logged a suspicion worth testing: on the reversion side the 2R target sits so far out that fades never reach it — they mostly ride to the session-flat cut instead. The forward evidence agrees emphatically but thinly. Across the whole week the pure 2R target fired exactly **twice** on fade trades (both wins, +\$71 average), and in the post-switch window it fired 2 of 26 times. TARGET is essentially a target nobody hits.

To probe this with more than 26 trades I measured the maximum favourable excursion — how far each fade ever ran in its favour, in ATR units, straight off the 5s bars, which is exit-agnostic (the price path is the same whatever exit rule you bolt on). Across 753 filled fade entries this week:

FADE EVER RAN AT LEAST...	TRADES	SHARE
1.0 ATR in your favour	231	30.7%
1.5 ATR	159	21.1%
2.0 ATR	117	15.5%

2.5 ATR	87	11.6%
---------	----	-------

The median fade never gets past about a third of an ATR in your favour (0.34 ATR); three-quarters of them top out around 1.3 ATR. A 2R target — which sits at two-to-three ATR of favourable run depending on how you count the R — is reached by at most one fade in six, and often fewer. Pull the target in to a single ATR and you'd fire it roughly twice as often (about 31% of trades touch 1 ATR versus 15% touching 2). So yes: the target is too far, and a nearer one would actually book wins instead of leaving them to bleed back to the bell.

HONEST

Before anyone arms a 1R fade target, remember what actually sinks the fade book: not the target you miss on the winners, but the stop you eat on the losers. The fades that got stopped this week lost -\$13.18 apiece; the whole reversion archetype runs -\$6 to -\$16 a trade. Booking the winners at 1R rather than 2R converts a handful of session-flat rides into small green — a genuine improvement — but it does nothing about the fades walking into the stop. It's a sharpening, not a rescue. And it's a sharpening measured on a forward window of 26 fade trades in a single contract over two and a bit days. That is not enough to bank.

§11.4 The plain takeaway

The switch to two pure exits is clean and the marker split is honest, but the forward window is short and one-sided. On **momentum** the new exits look right: the chandelier, the 2R target and the ride-to-bell all book positive money — every dollar of the loss is in the entries that hit the 1-ATR stop, so keep the exit and go work the entries

KEEP. On **reversion** we learned almost nothing forward — only 26

fade trades fired, all in gold, all red — so the new 1.5-ATR-and-2R rule is effectively untested this week **WATCH** . The one real signal is that the 2R target is too far: fades rarely run two ATR in your favour, so a nearer target is worth a proper study — but it's a tune on a book whose real bleed is the stop, and 26 forward trades can't confirm it.

§12 The two live gates, their shadow children, and the knife-filter

The live desk fires exactly two entry gates: **vwap_dip** and **vwap_pullback**. Both are reversion — they buy weakness and sell strength back toward the session VWAP. A dozen of our shadow sims are just those same two gates with a knob turned or a carve-out bolted on. This week I lined every one of them up on honest fills. The short version: the whole family lost money, the sim badly flattered it, and the one tweak that was supposed to help — only trading the calm tape — made it worse, not better.

§12.1 What the two gates actually are

vwap_pullback — In an established trend (VWAP sloping the right way) that has already had a real impulse move, wait for price to drift back and tag VWAP from the trend side on light volume, then fire with the trend on the first turn bar. The live falling-knife guard skips the tag if volume is heavy (`rvol >= pb_max_rvol`, live cap 0.6). It is meant to buy the controlled pullback in a trend, not catch a knife. In practice, on a chop week, it spends most of its time fading moves that don't come back.

vwap_dip — The deep one. LONG only when price is stretched a full 4 ATRs below VWAP and two green confirm bars print (SHORT the mirror). It is so selective it fires essentially zero times a week live — it only wakes up for a violent flush. Everything we learn about "the dip" this week comes from the


looser **dip_loose** shadow, which drops the floor to 2.5 ATR and needs only one green bar.

§12.2 How they behaved this week, on money that's real

Here is the honest scoreboard for the two families. Every number below is scored on `real_pnl` (the tick-repriced fill), with the optimistic sim P&L shown next to it so you can see the gap.

GATE FAMILY	TRADES	SIM SAYS	HONEST FILLS	PER TRADE	WIN %	VERDICT
vwap_pullback (17 variants)	1,281	+ \$3,368	-\$4,754	- \$3.71	37%	KILL
vwap_dip (dip_loose)	131	+ \$1,200	-\$184	- \$1.41	53%	KILL

Read the two P&L columns as a pair. The pullback family's sim ledger says it made **+\$3,368** this week; the honest ledger says it **lost \$4,754**. That is an eight-thousand-dollar swing between the fantasy fill and the real one — exactly the 1.1–3.1× downside understatement we keep warning about, and the reason nothing here gets armed on a sim number. The dip is a gentler story only because it barely trades and its win rate sits just above half, but half of a coin-flip that pays fees to play is not an edge.



Left pair = vwap_pullback family, right pair = vwap_dip. Green bar up = what the sim claims, red bar down = honest fills. The gold line is break-even. The sim is looking at a different, kinder market than the one we'd actually trade.

§12.3 Which shadow sims are children of which gate

Almost the entire shadow bench is these two gates in fancy dress. The pullback children all reuse the live _gate_pullback with one thing changed, so there's no logic drift — just the knob under test. Not one of them turned the fade green this week.

SHADOW VARIANT	THE ONE CHANGE	N	HONEST \$	/ TRADE
pullback_loose	wider VWAP-tag / lower impulse floor	165	−\$542	— \$3.28
pullback_quiet	only fires in quiet tape (rvol≤0.20)	97	−\$422	— \$4.35
tw_deadtape_standdown	stand down only in dead tape	256	−\$960	— \$3.75
tw_clock_standdown	stand down in 4 named UTC hours	219	−\$891	— \$4.07
tw_strength_ceiling_pb	skip the fade when net_atr ≥0.30	36	−\$267	— \$7.41
tw_mgc_pb_midband_block / _strong_only	gold trend-band carve-outs	38	−\$306	−\$8.1
dip_loose DIP		131	−\$184	— \$1.41

dip floor
4.0→2.5 ATR, 1
confirm

The carve-out sims are the telling ones. `deadtape_standdown` and `clock_standdown` are the base pullback that abstains in the tape you'd think is worst — and they're the two single biggest bleeders on the bench, because standing down in dead hours doesn't help when the fade loses in the live hours too. The carve-outs trimmed the edges of the wrong thing.

§12.4 The condition filter — and the counter-intuitive part

The dream has always been a **favourable-condition gate**: learn the tape state where the fade wins, and only let it fire there — stop it walking into the red. The natural first guess was "the fade is a calm-tape trade, so gate it to quiet volume." We have that exact experiment already running as `pullback_quiet`, which fires the identical live gate but only when `rvol ≤ 0.20`.

THE FINDING THAT SURPRISED ME

Restricting the fade to the quietest tape made it **worse**, not better. The all-volume pullback lost **−\$3.28** a trade and won **47%** of the time. The quiet-only version lost **−\$4.35** a trade and its win rate collapsed to 31%. Calm is not the friend of this fade — it's a trap. In the thinnest tape there's no counter-flow to hand you the mean-reversion, so the drift just keeps drifting and the fade sits in the red the whole hold. A calm-tape band is the wrong filter.

The axis that actually separates the winners from the run-overs is **direction**, not volume. Split the fades by side: the LONG fades (buying a dip) were nearly break-even at **−\$0.40** a trade, while the

SHORT fades (selling strength back down) lost **−\$6.06** a trade. And by contract the picture is even starker:

CONTRACT	FADES	HONEST \$	/TRADE
MGC (gold)	157	−\$597	−\$3.80
MES	76	−\$264	−\$3.48
MBT	30	−\$167	−\$5.57
MNQ	84	+\$41	+\$0.49

Gold alone lost **\$597** across 157 fades and is the single biggest hole in the family. Every fade that leaned against this week's gold move — and gold trended — got run over, exactly the pattern we flagged in the reversion autopsy. Meanwhile **MNQ was actually a whisker green** at +\$0.49 a trade. That matches everything else we know: the fade's only honest niche is MNQ, and the fix is not "wait for calm," it's "don't fade a market that's trending, and above all don't fade gold." A |slope|/direction veto is the shape of the filter; a low-vol band is not.

Honest caveat: this is one chop week and the per-contract counts are small (MNQ's edge is +\$41 on 84 trades — a rounding error, not a bankable claim). None of this survives as "arm a filter." What it survives as is a **diagnosis**: the calm-tape band is falsified, and the direction/trend axis is where the next honest gate has to live.

§12.5 The knife-filter — the obvious bolt-on that doesn't pay

The tempting idea is to bolt the aggressor tape and the order book onto these gates as a confirmation filter — only fade when the flow agrees, bail when it flips. I tested it honestly and it does not pay:

- As a **momentum entry filter**, aggressor confirmation was **null** — the buy/sell imbalance at entry had no predictive relationship

with the outcome ($p = 0.40$). Knowing who was hitting the tape when we fired told us nothing about whether the trade worked.

- The **flow-flip exit** — cut the trade the moment aggressor flow reverses — **died at tick resolution**. By the time the flip is real enough to read, the adverse move has already happened; the exit fires late and books the loss it was meant to dodge.

As for the deeper question — does the **L2 order book** give the fades a usable veto (absorption at the level, a wall that says "the mean-reversion is here")? That's examined properly in the dedicated **L2 order book** section against `depth.db` for the Thu+Fri window, and I'll point you there rather than half-do it here. The headline you can carry over: the book does **not** hand these fades a clean, usable stop-firing veto this week. The knife-filter, in every form I could test, is a bolt-on that doesn't earn its bolt.

MONDAY TAKEAWAY

Both live gates are reversion and both bled on honest fills this week; the pullback family's sim P&L (+\$3,368) is a fiction against its real result (−\$4,754), so trust nothing that isn't tick-priced. The calm-tape filter is falsified — quiet tape made the fade worse (win rate 47%→31%). The real axis is direction and contract: gold fades got run over (−\$597), MNQ was the only whisker of green. The knife-filter — aggressor confirm and flow-flip exit — is null and does not pay. Next honest gate lives on a slope/direction veto, MNQ-only, not a volume band.

§13 The order book — our first honest look

You caught us on this one. We have been recording the full top-five order book for MNQ, MES and MGC — every resting bid and offer, five levels deep, four times a second — and until this week nobody had opened it. So this is the first time we have actually looked at the book instead of guessing about it. What follows is a first look, not a verdict: it is one-and-a-third days of tape, one market mood, and everything here is in-sample. Read it as "here is what the book seems to be telling us, and here is the one thread worth chasing next week" — not as a filter you can arm on Monday.

§13.1 What we actually captured

The recorder ran from Thursday 09 July 14:17 UTC to Friday 10 July 21:00 UTC. In that window we banked **360,550 book snapshots on MNQ** (about 31 hours of near-continuous tape), **196,656 on MES** and **188,968 on MGC** — all three spanning the same Thu-into-Fri stretch. M2K only came online for the last 1.8 hours, so I have set it aside. Coverage is dense — roughly one snapshot every 250 milliseconds, deduped when nothing moved — with a single thin patch during the 21:00–22:00 UTC settlement break on Thursday (47 snapshots that whole hour instead of ~14,000). MNQ is the deepest and the cleanest, so I lead with it everywhere below.

SET EXPECTATIONS

One-and-a-third days. One regime (a quiet Thursday grind and a news-whippy Friday). In-sample — I am reading the same tape I would trade on. Everything here needs a second and third week

across different moods before a single pound rides on it. I will say plainly at the end what "enough" would look like.

§13.2 Does the book point the way? No — if anything it points the wrong way

The obvious hope is that a bid-heavy book means price is about to go up. It does not. I took all 360k MNQ snapshots, measured the book's top-five imbalance at each one, and then looked at where the mid actually was 10, 30 and 60 seconds later. The relationship is essentially nothing — the correlation between imbalance and the next 60-second move is **-0.006** (10s: -0.015, 30s: -0.009). Every imbalance bucket coin-flips: price ticked up 48% to 54% of the time no matter how the book was leaning, and the median move in every bucket is sub-tick.

The one place a signal appears, it is backwards. When the book was heavily bid-stacked (imbalance above +0.3 — the reading that screams "support, buyers everywhere") the average next-60-second move was **-0.61 points**, i.e. slightly down. A stacked bid book was, on average, a mild sell. This is the same weakly-contrarian shrug we got from the aggressor tape: resting size is not a compass. If anything, the fattest, most inviting side of the book is the one most likely to get run over.

MNQ BOOK IMBALANCE AT T	SNAPS	AVG MOVE NEXT 60S	WENT UP %
ask-heavy (< -0.3)	6,402	+0.48 pt	50.6%
-0.3 to -0.1	75,512	+0.22 pt	50.0%
flat (-0.1 to +0.1)	203,809	+0.25 pt	50.3%
+0.1 to +0.3	68,196	+0.22 pt	50.9%
bid-heavy (> +0.3)	5,598	-0.61 pt	52.9%


So the naive "book points the way" idea is dead on arrival, exactly as it was for the tape. That is worth knowing in itself — it stops us building a gate on a signal that was never there.

§13.3 The knife test: does fading into a wall lose more?

Here is your real question. When our fade and pullback sims buy a dip or sell a rip, are they walking straight into a stacked wall of the opposite side — a pile of sellers overhead when we buy, a pile of buyers underneath when we sell? And when they do, does it hurt more? I pulled every fade-family entry (the pullback, dip and chop sims) on MNQ/MES/MGC that fired inside the depth window — **97 distinct signal moments** (122 rows before I collapsed the cases where three sims fired the same bar) — matched each to the book within 2.5 seconds of entry, measured the opposing-side size against the same-side size, and scored on honest real_pnl.

First the blunt fact: **all 97 of these fades lost money as a group** (−\$288 on honest fills across the window). That is the reversion bleed we already know, showing up again. But the book does carve them, and it carves the longs cleanly. A pullback long that bought into an ask stack $1.3\times$ or more the size of the bids below it lost an average of **−\$17.51** and won only **1 time in 8**. The same long into a thin overhead (ask under $0.7\times$ the bid) made **+\$3.73** on average and won 54%. The gradient is monotone and it makes plain sense: you buy the dip, but there is a wall of sellers overhead, price can't clear them, you get stopped.

Pullback/dip LONGs on MNQ/MES/MGC — average honest P&L by the overhead (ask) book at entry



The catch — and it is a big one — is **n=8**. Eight long entries into a wall over one-and-a-third days is a story, not a statistic. And the short side actively refuses to play along: the three shorts that sold into a heavy bid wall actually won (+\$21 each), the opposite of the veto. So what I have is a clean, sensible, economically-shaped signal on the long side (don't buy a dip into stacked overhead) built on eight trades, and noise on the short side. That is a lead, not a filter.

§13.4 What the book did around the big moves

The marquee move of the window was Friday's **14:32:55 whipsaw** — MNQ fell from ~29,956 to ~29,677, about **279 points in 27 seconds**, then reversed +120 straight back. This is exactly the kind of move where you would hope the book warns you. It did the opposite. All the way down the first leg the book was heavily bid-stacked — 400 to 650 lots on the bid against ~150-200 on the ask, imbalance pinned around +0.5 — the textbook "support is stacked" picture. Price sliced through it without pausing. Then the offers were simply **pulled**: the ask stack collapsed from ~226 lots to ~35 in a single snapshot as price gapped 29,902 → 29,850, and moments later the whole book evaporated to single digits on both sides. The last 170 points was a genuine sell avalanche — **5,204 then 5,136 lots sold per five seconds** — pouring into an empty book. That is a vacuum, not absorption: nobody ate a wall, the liquidity left and price fell through the hole.

For contrast, Friday's calmer **+43.75 rip at 12:13:40** was flow-led: the ask overhead was ordinary and thin (~25 lots, never stacked),

and a one-sided buy sweep of **2,571 lots against just 54 sells** lifted every offer faster than it could be replaced. Different mechanism, same lesson — the big directional moves happened where the far side was thin, whether it was pulled (the crash) or simply lifted (the rip). In neither case did a stacked opposing wall stand in the way; in the crash, the stacked side gave a flat-out false "support" reading right before the worst drop of the week.

THE ONE THING

Resting size is not a promise. The fattest side of the book is where the trap is: a stacked bid book preceded the biggest sell-off in the sample, and the fastest moves ran through vacuums, not walls.

§13.5 Honest takeaway

Two things survive this first look. One is settled: **K book imbalance is not a direction signal** — it is a flat-to-contrarian coin flip, same as the tape, so we will not build a "lean with the book" gate. The other is a lead worth one more look: **T a pullback long into a stacked ask wall is a bad trade** — seven of eight lost, and it makes mechanical sense. But eight trades on one-and-a-third in-sample days is nowhere near arm-able, and the short-side mirror didn't hold, so I am flagging it, not shipping it.

What would make it real: keep the depth recorder running so we bank the book across a few different weeks — a trend day, a chop day, a data day — and let the fade sims accumulate. I want to see this wall-veto on the order of **30-50 long entries per regime**, with the short side either confirming or being explicitly left out, before it becomes a live knife-veto. The recorder is cheap and already running; the only thing between us and an honest answer is a few more weeks

of tape. For this week: nothing to arm, one clear thread to chase, and the naive "book points the way" idea taken off the table for good.

Monday takeaway: Don't trust the book to point direction — it's a contrarian coin flip. The one real thread is that buying a dip into a stacked wall of overhead sellers loses hard, but it's eight trades on 1.3 days, so keep the recorder running and let it prove itself over a few more weeks before we veto anything.

✓ **A SKEPTIC RE-RAN THE MONEY CLAIM — IT HELD**

REAL. On honest net real_pnl for Mon6-Fri10 Jul (n=3797, desk -\$8599): rvol<0.8 = 58.9% of entries (claim says two-thirds; actually ~three-fifths) and -\$8414 = 97.9% of the net loss; muzzling it desk-wide lifts the desk to -\$185, a +\$8414 (~£6.6k) prize — larger than the £3.5k claimed, so the claim is conservative. Robust: loses all 5 days, avg pnl monotone in rvol, net of fees, and a shadow floor (understates live). One honest caveat — the desk is net -EV so part of the prize is trade-less: a random 59% cut would leave -\$3533 vs -\$185, so genuine signal selectivity beyond trading-less is ~\$3.3k (~£2.6k / ~\$3.5k), which is exactly where £3500 lands. Per-sim confirmed exactly: passive_chop is -EV in every rvol bucket (kill it, no pocket); momentum_shadow flips positive above rvol 0.8 (0.8-1.5 +\$1.63/trade, 1.5+ +\$3.82/trade) — a real switch-into pocket. Not curve-fit, not thin-n, not cost-blind; only nuance is the -EV desk making it partly trade-less.

§14 The OFF switch: when should each sim NOT fire?

The whole prize this week is not turning a sim on. It is muzzling one on the days its edge is not in the room. So I stopped asking "which sim is good" and asked the only question that pays: given the tape I can read at the moment of entry — how fast it is moving (atr%), how much volume is showing up (rvol), what session we are in — when should the trigger simply be dead? The honest answer, on this week's real fills, is boring and strong: **the desk keeps firing into dead tape, and dead tape is where essentially all the money goes.**

READ THE TWO TRAPS FIRST

Trap 1 — the daily lottery. I run ~20 to 28 sims against 7 contracts every day. On any given day a handful come out green by chance. Look at who "won" each day this week, best sim by real P&L with a real sample: Mon it was `tw_mnq_thrust_cont` (+\$432), Tue `momentum_shadow` (+\$238), Wed `tw_mgc_coil_break_long` (+\$355), Thu `dip_loose` (+\$69), Fri `tw_mgc_pb_midband_block` (+\$58). Five days, five different winners, no repeat. Chasing yesterday's hero is chasing noise. And it is not even a close race for green: on Fri only 3 of 21 sims finished positive, on Tue 4 of 22. A switch built on "who won lately" is a switch built on a coin that lands somewhere new every morning.


Trap 2 — the switch has to be REACTIVE, not a forecast.

We have proven to death that regime is unpredictable at the minute scale, so the switch cannot say "today will be trending." It can only read the condition that is present right now. And the labelled regime confirms why: bucket every trade by the `session_regime` tag and both big buckets bleed — "chop" lost \$4,059 over 1,013 trades and "trending" lost \$1,147 over 1,203. The label does not separate winners from losers. A live number read off the tape at entry does. So every rule below is a present-tense reading, never a prediction.

§14.1 The cheap desk-wide gate: is the tape tradeable at all?

Before tuning any single sim, ask the cheapest possible question — one boundary that muzzles the whole desk when the tape is not worth trading. I cut every entry this week (the ~2,277 with a clean condition read; Monday morning has a funnel-capture gap, so it is

under-counted, not wrong) by relative volume at entry. The shape is stark:



RVOL BAND AT ENTRY	TRADES	NET REAL P&L	WIN%
< 0.8 (thin / quiet)	1,566	-\$5,932	37%
0.8 - 1.2	288	+\$91	44%
1.2 - 2.0	268	+\$787	37%
≥ 2.0 (blow-off)	155	-\$772	28%

Sixty-nine percent of everything the desk does happens in thin tape (rvol under 0.8), and that one bucket carries essentially the entire loss. Put a floor at rvol 0.8 and the picture flips: the 711 trades above the floor net roughly breakeven (+\$106), while the 1,566 below it lost \$5,932. **The floor does not make money — it stops a guaranteed bleed.** That is the honest framing, and it is the single most robust thing in this section: split the week day by day and the below-floor bucket is negative on all five days (−\$1,118, −\$860, −\$553, −\$1,488, −\$1,913). There is no lucky day where firing into dead tape worked.

Two guards on the claim. First, it is not just the passive_chop reversion book (which is dead anyway — see below): strip passive_chop out entirely and the below-0.8 bucket still bled −\$5,547 across 1,459 trades while the in-band trades made +\$1,257. The bleed is desk-wide, not one bad sim. Second, the upper cut (drop rvol ≥ 2, the blow-off bar) adds a bit more but rests on only 155 trades — treat the floor as bankable and the ceiling as a lean.

THE DESK-WIDE MUZZLE

OFF for every sim when rvol at entry is under 0.8. It costs ~\$106/wk of give-up and it stopped ~\$5,900/wk of shadow losses this week, negative every single day out-of-gate. It is a floor that says "the tape isn't paying today," not a money-maker.

KEEP

§14.2 Per-sim pockets: where each big one actually lives

The desk gate is blunt. Now the sharper cut — for the three biggest sims, isolate the pocket where that sim makes its money from the tape you can see at entry, and state the muzzle.

§14.3 momentum_shadow — a real amplitude plateau

FLOOR IT

This is the one sim with an honest, structured pocket. Slice its 296 conditioned trades by atr% at entry and the edge is not a lucky spike, it is a plateau: below atr% 0.045 it loses, above it wins, and it wins in three consecutive bands.

ATR% AT ENTRY	TRADES	NET REAL	WIN%
< 0.03	94	-\$171	36%
0.03 - 0.045	61	-\$113	43%
0.045 - 0.06	61	+\$286	41%
0.06 - 0.08	35	+\$232	46%
≥ 0.08	45	+\$229	42%

Three green bands in a row is what a real edge looks like — a knife-edge would be one lucky band with losers either side. In plain terms:

this sim needs the bar to actually be moving. Give it a dead, low-amplitude tape and it buys a twitch that goes nowhere and pays the spread. Add one more guard — it also gives money back when volume goes vertical ($r_{vol} \geq 2$ lost $-\$168$ at a 22% win rate, the classic buy-the-top). Put both together, muzzle unless $atr\% \geq 0.045$ and $r_{vol} < 2$, and the pocket holds **+\$795 over 113 trades (47% win) while the muzzled remainder was $-\$331$** . The floor keeps every dollar of the profit and throws away a losing tail. And it holds up day by day — the pocket was positive or flat on 4 of 5 days, its one red day was a rounding error ($-\$14$). This matches the amplitude-floor thesis already deployed observe-only; the data here says keep it and let it graduate.

§14.4 passive_chop — no pocket exists KILL

The desk's biggest sim by volume (1,088 trades) and its biggest loser ($-\$4,249$ real). I went looking for its pocket the same way, and there isn't one. Every $atr\%$ band loses. Every r_{vol} band loses. There is no corner of the tape where fading the chop paid this week:

SLICE	TRADES	NET REAL
$atr\% < 0.03$	44	$-\$128$
$atr\% 0.045-0.06$	38	$-\$150$
$atr\% \geq 0.08$	86	$-\$447$
$r_{vol} < 0.8$	186	$-\$686$
$r_{vol} \geq 1.2$	23	$-\$89$

When a sim loses in every bucket, there is no OFF switch that saves it — the OFF switch is the sim. This is the reversion-autopsy verdict landing again from a different angle: passive_chop fades trends and gets run over, and no observable-at-entry condition tells you which fade will be the one that keeps going. The muzzle here is permanent.

§14.5 pullback_loose — same thin-tape disease FLOOR IT

Down $-\$542$ on the week, and the story is the desk-wide story in miniature: it is a thin-tape bleeder. Below rvol 0.8 it lost $-\$703$ over 127 trades; at rvol 0.8 and above it made $+\$147$ over just 27. The rvol floor is the right muzzle again — but be honest, 27 in-gate trades is too thin to bank on its own. Read it as one more vote for the desk-wide floor rather than a standalone edge.

§14.6 Plateau or knife-edge? — the one-week honesty check

Everything here is one week, in-sample, on shadow fills. So which of these is a robust plateau I would defend, and which is a shape I fit to five days of noise?

Plateau (defend it): the desk-wide $\text{rvol} < 0.8$ muzzle. It is negative every single day, holds with passive_chop removed, and rests on 1,566 trades. That is not a fit, that is a floor.

Plateau-ish (lean into it): momentum_shadow's $\text{atr}\% \geq 0.045$ amplitude floor. Three adjacent green bands and positive 4 of 5 days is genuinely plateau-shaped, but two big days carry most of the dollars — graduate it observe-first, don't size it yet.

Too thin to bank: pullback_loose's floor (27 trades) and momentum's $\text{rvol} \geq 2$ ceiling (155 trades). Directionally right, evidentially thin. Leans, not laws.

Takeaway, plain: the money isn't in a clever new trigger, it's in a mute button. Two-thirds of what this desk does happens in tape too thin to pay, and that thin tape is where all the loss lives — every day, this week. Wire one desk-wide floor: no sim fires

when rvol is under 0.8. Then give momentum_shadow its own amplitude floor ($\text{atr\%} \geq 0.045$) and retire passive_chop, which has no favourable pocket to switch into. And remember why the winner's list rotates every morning — the switch reads the tape that is in the room, it never predicts the day.

§15 Greenfield: if we can't catch the big misses, what would?

Every week the tape throws a handful of big, clean runs and every sim on the desk sits and watches them go by. This chapter starts from those misses, tries to build ONE brand-new signal from scratch that would have caught them, and then does the only thing that matters — trades it on honest fills, net of costs, and checks whether it survives outside the one week I fit it on. The short version: it doesn't. But the reason it doesn't is worth more than a working curve-fit would have been.

§15.1 The misses, laid out honestly


I ran the full missed-run census on MNQ (the contract with the deepest tape and book coverage). In the last seven days there were **76 runs of 1.5×ATR or more** — a 15-minute move of at least 73.9 points, against a typical 15-minute range of 49.2. Of those 76, the desk sat out roughly **59**. That is not a rounding error; that is the whole game walking past us. The census sorts them into four cause-clusters:

CAUSE CLUSTER	RUNS	SAT OUT	WHAT IT LOOKS LIKE
VACUUM (moved against the tape)	36	27	price ran the opposite way to the aggressors
FLOW-LED (aggressors drove it)	19	15	heavy one-sided lifting/hitting, then the move
OPEN / NEWS spike	18	15	13:00–15:00 UTC cash-open & data prints

The seductive one is FLOW-LED. Nineteen runs where the aggressor tape leaned hard one way and then price went that way, and we sat out fifteen of them. The obvious greenfield hypothesis writes itself: build a tape-ignition entry — when net aggressor flow floods one way, jump in that direction and ride. That is the signal I set out to build. It is also completely wrong, and the data says so in the first test.

§15.2 The first test kills the obvious idea

Before building any trade logic I asked the plainest possible question: across every 5-second bar this week, if I know the net aggressor flow over the previous 60 seconds, does that tell me anything about where price goes next? I bucketed all 80,352 MNQ bars by their trailing-60s net flow and measured the average move over the next three minutes.



Mean 3-minute forward move (points) by trailing-60s net aggressor flow, MNQ, all bars this week. Left = heaviest net selling, right = heaviest net buying.

Read that curve. When the tape is flooded with selling, price drifts up over the next three minutes (+3.70 points in the heaviest-selling bucket). When the tape is flooded with buying, price drifts down (−0.84). The overall correlation between flow and forward return is

−0.065 — small, but consistently negative. Aggressive flow on this tape is **contrarian**: the people crossing the spread are, on average, the ones getting run over. Those FLOW-LED runs in the census are survivorship — the census finds the handful of times flow and move happened to agree, but across the whole week flow points the wrong way. A tape-chasing momentum entry would have been structurally backwards.

THE ONE DURABLE LESSON

Do not build a "follow the aggressor tape" momentum entry on these contracts. Net aggressor flow is a weak fade signal, not a chase signal. Every instinct that says "the buyers are in control, go with them" is, this week, exactly inverted.

§15.3 So I flipped it — and built the fade instead

If flow is contrarian, the honest greenfield signal is the opposite of the obvious one: **seller-exhaustion fade** — when the 60-second net flow shows heavy net selling, go long the capitulation and hold for the bounce. I built it as a real trade: enter at the next bar's open, no overlapping positions, hold to a time cap, and charge the honest cost. For the cost I didn't invent a number — I used the desk's own measured gap between optimistic sim P&L and honest fills, which on MNQ is **\$5.19 per round trip**. Baseline drift this week was essentially zero (+0.016pt over 60s), so there's no free long hiding in a rising market — anything the signal earns, it earns.

SELLER-EXHAUSTION FADE (LONG)	TRADES	NET \$	\$/TRADE	WIN%
flow ≤ −400, hold 5min, +\$5 cost	300	−110	−0.4	53
flow ≤ −600, hold 5min, +\$5 cost	190	+396	+2.1	53
flow ≤ −800, hold 5min, +\$5 cost	126	+452	+3.6	52

... same, but with a 25pt stop / 40pt target	300	−895	−3.0	46
--	-----	------	------	----

There is a pulse — but look at where it lives. The signal only clears its own costs at the extreme tail, where the tape has dumped 800+ net contracts in a minute, and only if I hold the full five minutes with **no stop**, wearing the entire drawdown. The best cell earns **\$3.6 a trade** across 126 trades — and I found it by sweeping thresholds (400/600/800) and horizons (3min/5min) and picking the winner. That is the textbook shape of a curve-fit: nothing at the sensible settings, a sliver at the extreme, and it evaporates the moment you add a realistic stop (the fade walks into more selling before the bounce and gets stopped, turning +\$3.6 into −\$3.0). Per day it's incoherent too — it loses on three of the five days at the 400 threshold.

§15.4 The out-of-sample test buries it

The only honest way to judge a one-week filter is to point it at data it wasn't fitted on. I don't have next week, but I have two sister contracts with full-week tape — MES and MGC — running the exact same physics the same week. If seller-exhaustion is a real market truth, it should show something there.

SAME LONG FADE, HOLD 5MIN, +\$5 COST	TRADES	\$/TRADE	WIN%
MES · flow ≤ −200	331	−4.5	38
MES · flow ≤ −600	108	−1.7	46
MGC · flow ≤ −200	44	−8.0	45

It loses at every threshold on both. The marginal MNQ win doesn't replicate one contract over, in the same week, on the same idea. That's not an edge with a tuning problem — that's a one-contract, one-week artifact.

§15.5 And the order book had nothing to add

The right instrument for "seller exhaustion" isn't the trade tape at all — it's the L2 book: heavy selling hitting a bid that refuses to drop is genuine absorption. I have 360,550 MNQ depth snapshots across Thursday and Friday, so I checked whether book imbalance (top-of-book and top-3 depth) predicts the next 60–180 seconds. Correlation with forward return: **−0.006 to +0.001** — indistinguishable from zero. Bid-heavy and ask-heavy books produced almost identical small forward moves (just the two-day drift). The book, over the window we captured, carries no directional edge to build on either.

§15.6 The plain takeaway

I tried to invent the signal that catches the misses, and the honest answer is that there isn't one hiding in what we capture. The big runs split into three kinds we can now name: **news/open spikes** (13:00–15:00 UTC — unpredictable from order flow by definition, and a third of the census), **contrarian overshoots** that snap back faster than our costs can be paid, and a residue of quiet drifts with no fingerprint at all. The one candidate with a heartbeat — fading heavy selling — is a swept-in curve-fit that dies on MES and MGC the same week. **KILL** the tape-ignition idea, **KILL** the seller-exhaustion fade, **KILL** the book-imbalance entry. Nothing goes to the shadow book.

What we keep is the negative knowledge, and it's the most bankable thing in the chapter: the aggressor tape on these contracts is a **fade**, not a chase. That's a permanent guardrail against the entire family of "follow the flow" momentum entries someone will inevitably want to build next. The misses stay missed — but now we know they're missed because they're genuinely silent, not because we haven't looked.

§16 The Sizing Ladder — Has Any Sim Earned More Size?

Size is earned, not chosen. The ladder is deliberately slow: a sim starts as a shadow at one micro, and only climbs to live-at-1, then 2, then 3 once it has cleared every rung — 8 to 12 weeks of tape, 100 to 200 trades, a positive edge that holds across regimes, a survivable drawdown, a clean reward-to-risk, and a low correlation to the live desk we already run. This week's honest answer to "who's ready?" is short: nobody yet. But the queue is forming, and it's the momentum names standing in it.


§16.1 The green sims built real money this week

Scored on honest fills, eight sims finished the week in the black, and every one of them is a momentum or breakout idea — buying strength, not fading it. The standout is **tw_mnq_thrust_cont**: 28 trades, +\$992 real, a 61% hit rate, and an average winner of \$80 against an average loser of \$33. That's a genuine 2.4-to-1 payoff — the cleanest reward-to-risk on the whole board. Behind it, the workhorse **momentum_shadow** put up +\$749 across 413 trades, and orb_shadow, the two MGC breakout books, the M2K momentum sim and the MES/MYM thrust books all pitched in.

SIM	TRADES	REAL P&L	WIN%	AVG WIN	AVG LOSS	WORST	R:R
tw_mnq_thrust_cont	28	+ \$992	61%	\$80	– \$33	–\$63	2.4

momentum_shadow	413	⁺ \$749	42%	\$21	[−] \$12	−\$57	1.7
orb_shadow	144	⁺ \$196	35%	\$30	[−] \$14	−\$82	2.2
tw_m2k_momo_loose	76	⁺ \$163	41%	\$14	−\$6	−\$17	2.4
tw_mgc_coil_break_long	134	⁺ \$137	36%	\$42	[−] \$22	−\$41	1.9
tw_mgc_momo_thrust	9	⁺ \$127	56%	\$46	[−] \$26	−\$32	1.8
tw_mes_thrust_cont	26	+\$49	46%	\$20	[−] \$14	−\$23	1.5
tw_mym_open_thrust	12	+\$28	58%	\$16	[−] \$17	−\$24	1.0

Here is the winners' book adding up across the week — a real climb to +\$2,596 by Thursday, then a small give-back into Friday. That last flat leg is the whole argument for the ladder in one picture.



§16.2 Now score them against the ladder — and the ladder says "not yet"

The trouble isn't the money, it's the calendar. The oldest sims in this list, `momentum_shadow` and `orb_shadow`, were born on 3 July. The

rest — every tw_book — started this Monday. So the single most important gate, **8 to 12 weeks of tape**, is failed by all of them by a mile. One week is not a track record; it's an anecdote. Against the gates:

GATE	TW_MNQ_THRUST_CONT	MOMENTUM_SHADOW	THE REST
8-12 weeks of tape	≈3 days — no	1 week — no	1 week — no
100-200 trades	28 — no	413 — yes (but 1wk)	9-144 — mostly no
+EV across regimes	lumpy, mostly green	faded Thu/Fri to red	mixed / one-day
Drawdown survivable	worst trade −\$63	worst −\$57	OK so far
Clean R:R	2.4:1 — yes	1.7:1 — ok	1.0-2.4
Low corr. to live desk	can't measure in 1wk	can't measure in 1wk	can't measure
Verdict	KEEP OBSERVING	KEEP OBSERVING	KEEP OBSERVING

tw_mnq_thrust_cont is the clear front-runner and has genuinely earned its place on the first rung of the ladder — keep it as a shadow, keep watching. What it has not earned is a live micro. It has 28 trades, not 100, and three days of tape, not eight weeks. And 44% of its +\$992 came from just three trades; strip those and it's still green but far less imposing. That's not a knock — it's exactly why the trade-

count gate exists, so one good Monday can't buy its way onto the desk.

THE FEATURE THAT LOOKS LIKE A DELAY

The selective momentum sims trade rarely — `tw_mnq_thrust_cont` fired 28 times this week, `tw_mes_thrust_cont` 26, `tw_mym_open_thrust` just 12. At that pace they'll take months, not weeks, to reach 100–200 trades. That is the discipline working, not a problem to engineer away. A sim that only bets when the tape is really moving should spend most of its life waiting. We let it wait.

§16.3 Two robustness caveats before anyone gets excited

First, the regime warning. `momentum_shadow` made +\$419 on Monday and +\$238 on Tuesday when the tape trended, then bled –\$8 Thursday and –\$42 Friday as it chopped. That's the signature of a strategy that's long the trend and gives it back when the trend stops — it needs to prove it can survive a genuinely dead week before it climbs. Second, **`tw_mgc_coil_break_long` is almost entirely one day**: +\$355 on Wednesday against –\$106, –\$22 and –\$140 on the other days. Take Wednesday out and it's a bleeder. A single lucky breakout is not an edge, and the ladder's job is to make it prove that over 8 weeks.

§16.4 The Departure Lounge

The other end of the board is easier to read. These are the sims that lost real money by fading strength and standing in front of moves, and they should be nominated for retirement. The report nominates; the operator arms — nothing here is switched off by this page.

SIM	TRADES	REAL P&L	WIN%	WHAT IT DOES WRONG	
passive_chop	1088	— \$4,249	17%	Fades chop across everything (heaviest on MBT); wins 1 in 6	RETIRE
tw_deadtape_standdown	256	—\$960	36%	Meant to stand down in dead tape — trades it instead	RETIRE
tw_clock_standdown	219	—\$891	37%	Same idea, same bleed on the clock filter	RETIRE
tw_ride_strength	198	—\$627	38%	"Rides strength" but keeps getting run over, worst — \$98	RETIRE
passive_mgc_hours	63	—\$542	46%	Fades gold on a schedule; —\$9 a trade average	RETIRE

pullback_loose	165	-\$542	47%	Buys dips that keep going; near-coin-flip that loses	RETIRE
pullback_quiet	97	-\$422	31%	Same fade, quiet-tape flavour, worse hit rate	RETIRE
tw_mbt_fade_euphoria	50	-\$312	6%	Fades bitcoin euphoria — wins 3 times in 50	RETIRE
tw_strength_ceiling_pb	36	-\$267	42%	Caps strength then fades the pullback into it	RETIRE

The pattern is one sentence long: **every fade book lost, every momentum book that survived made money.** `passive_chop` alone gave back \$4,249 — more than all eight green sims made put together — by selling into moves 1,088 times and winning one trade in six. `tw_mbt_fade_euphoria`'s 6% win rate is the purest version of the mistake: it stood in front of a running crypto tape and got flattened 47 times out of 50. These aren't tuning candidates; they're the wrong trade.

PLAIN TAKEAWAY

No sim has earned a live micro yet — not one clears the 8-to-12-week gate, and the best of them, `tw_mnq_thrust_cont`, has 28 trades where it needs 100 to 200. Keep the eight momentum winners on the shadow rung and let them accumulate the weeks; their slow trade pace is a feature, so give them until the autumn, not this month. Meanwhile the Departure Lounge — led by `passive_chop` and the whole fade-and-standdown family — is nominated for retirement, because a book that fades strength has now proved across a thousand trades that it just hands the money back. The report nominates; you arm.

CHAPTER

Your Mid-Week Musings — Every Lead, And What We Found

§17 Your leads this week — what we found

Everything you flagged in the chats this week went into the lab as a lead — not a fact, a thing to test hard on Friday's full data. Here's what happened to every one of them, plainly: which held, which fell, and where it landed in this report. Nothing you raised got quietly dropped.

THE SHORT VERSION

Your biggest call of the week — **momentum is confirmed-vs-loose, and "ride any strength" is a drag — held hard** (a skeptic reproduced +\$1,828). Your "gold is a momentum-only play" held on the confirmed side. The one you'll least like: **the aggressor-tape/L2 knife-filter, as a bolt-on, does not pay** — but the L2 book you caught me leaving out is now examined, and there's a real open door on absorption. And you were right that the report kept missing things: the run census, the exits before/after, and the book are all in now.

§17.1 The ledger

YOUR LEAD (WHEN)	WHAT IT WAS	FRIDAY VERDICT	WHERE IT LANDED
Confirmed vs loose momentum (8 Jul)	The "ride any strong move" family bleeds; the edge is	HELD confirmed +\$2,344 vs loose −\$1,017; skeptic reproduced + \$1,828.	"Momentum: confirmed-vs-loose" + the per-contract matrix.

	confirmation, per contract.		
Exit = scar tissue for bad entries (8 Jul)	Rip out the 9- mechanism stack; reversion = 1.5-ATR + 2R, momentum = 1- ATR + chandelier.	HELD momentum keeps every dollar on one chandelier — matches the old spine for free.	"The exit change, before & after."
The 2R fade target is too far (8 Jul)	Most fades ride to session close; the 2R target rarely fires.	HELD 2R fires ~1 trade in 10; nearer target trims but the fix is the entries.	"The exit change" (reversion half).
Gold is a different beast / momentum- only (7 Jul)	Keep confirmed- momentum on MGC, kill the fades and loose coil-break.	PART-HELD confirmed MGC momentum +\$258; the fades bled; but MGC's big collapses have no pre-run tell.	The matrix + "MGC runs" section.
Aggressor / L2 knife-filter (7 & 11 Jul)	Use the tape/ book to stop the fades chasing knives.	FELL as a bolt-on entry-confirm null (p=0.40), flow-flip exit dies at tick. Book imbalance is weakly contrarian.	"Live gates + knife- filter" and the new "L2 order book" section — absorption stays the one open door.
Fade book bleeds both ways (7 Jul)	The one-sided fade loses in trend AND chop.	HELD the fade/ pullback family bled -\$5,157 honest this week.	"Live gates and their shadow children."
Fade only in quiet tape (running assumption)	The reversion edge lives in calm, low-vol tape.	BACKWARDS the fade bled WORST in the quietest tape; the fix is a trend	"Live gates" + "The OFF switch."

		veto, not a calm-tape band.	
Flat-slope contract regime split (7 Jul)	MNQ/MGC want flat tape; MES/MCL/M2K want continuation.	NOT ROBUST the direction is real but the threshold is noisy, not a plateau — a lead, not a rule.	Noted in "The OFF switch" as a hypothesis, not armed.
The OFF switch / favourable-condition gating (8 Jul)	Muzzle each sim on its non-favourable days; a desk-wide "is the tape tradeable" gate.	REPORTED passive_chop has no green pocket (stay dead); momentum needs push+amplitude; never fade a live move. One-week, in-sample.	"The OFF switch."
The promotion / sizing ladder (9 Jul)	Sizing is earned — 8–12 weeks, cross-regime, low correlation.	REPORTED mostly "not yet" — even the star (MNQ thrust) needs the weeks; the fade family is the departure lounge.	"The sizing ladder."
Shadow understates losses (8 Jul)	The sim's own numbers are too gentle on the downside.	CARRIED every red number in this report is a floor; scored on honest fills throughout.	Baked into the whole report's discipline.
The report kept missing the runs + the book (11 Jul)	Use the full ATR run census, not a top-8; use the L2 data.	FIXED full census (76 MNQ runs) + the L2 book are both in this revision.	Every "runs" section + "L2 order book" + "How we grab them Monday."

The rule that governs this table: a lead only becomes a finding by surviving Friday's independent re-run — never by being flagged mid-week. That's why some of your strongest hunches sit in amber here: they're promising and pointed the right way, but one week of data can't promote them. The green

rows are what to act on; the amber rows stay in the shadow book for another week.

CHAPTER

The Winners — lean into these

LEAN IN — THIS ONE WORKS · \$+992 THIS WEEK ON HONEST FILLS.

tw_mnq_thrust_cont made \$991.6 over 28 trades at 61% win and beat its own perfect-fill ceiling by riding MNQ thrusts on the ratchet while its stops cost a total of \$7.9; Monday move is an atr_pct>=0.05 amplitude floor.

§18 The MNQ thrust-rider: it caught the runs and the stops were free

This one worked. On the honest tick-repriced fills

tw_mnq_thrust_cont made **\$991.6 over 28 trades this week at 61% win** — and it did something a shadow almost never does: it beat its own perfect-fill ceiling of \$722.9. Real money came in above the fantasy. It buys an MNQ thrust that keeps going, rides it on the ratchet, and when it's wrong the stop costs almost nothing.

\$991.6

HONEST REAL_PNL, 28
TRADES

61%

WIN RATE

+\$268.7

REAL BEAT ITS CEILING
(\$722.9)

-\$7.9

TOTAL COST OF ALL 12 STOPS


§18.1 Why it works: big winners on the ratchet, free losers on the stop

Split the week by how each trade left and the whole story is right there. The ten ratchet exits made **\$738.3 at 90% win** — those are the thrusts that ran

and the trailing ratchet let them run. Five more came out on a VWAP exit for another **\$274.9 at 80%**. Against that, the twelve stop-outs — nearly half the book — cost the desk a grand total of **\$7.9**. That's the engine: the exit spine cuts a wrong thrust to basically zero, so the sim only needs its runners to pay and they do.

EXIT REASON	N	REAL_PNL	WIN%
RATCHET (rode the run)	10	+\$738.3	90%
VWAP exit	5	+\$274.9	80%
STOP (wrong thrust)	12	-\$7.9	33%
EOD flat	1	-\$13.6	0%

The equity curve is a clean climb, not a lucky spike — it grinds up through the week and finishes at its high near \$992.



Cumulative honest real_pnl, trade by trade in time order — 0 at the gold line, ~\$992 at the top right.

§18.2 What separates its winners from its losers: bar size

Now the entry conditions. Fair warning on the data first: the documented funnel join (`cycle_ts = entry_sid`) returns nothing for this sim because `entry_sid` here is a session bucket, not a cycle timestamp — so I matched each trade to the nearest funnel snapshot in the 60 seconds before entry. That covers **18 of the 28** trades; the other ten have no readable conditions and can't be gated. On the 18 I can see, the cleanest divider by a mile is **bar amplitude, atr_pct**: winners fired on fatter bars (avg 0.056) than losers

(0.047). Spread pointed the same way, softer (winners 0.58 vs losers 0.89) — nobody home, wider fill — but `atr_pct` is the sharp axis.

THE ONE THING FOR MONDAY

Put an amplitude floor on it: **only fire the MNQ thrust when the bar is fat, `atr_pct` $\geq$ 0.05**. On the 18 trades I can read, that floor is a clean cut — the fat-bar bucket made **\$510.9 over 9 trades at 67% win**, while the thin-bar bucket churned **9 trades to -\$4.1 at 33% win**.

BUCKET (18 FUNNEL-MATCHED)	N	REAL_PNL	WIN%
<code>atr_pct</code> $\geq$ 0.05 (fat bar) — FIRE	9	+\$510.9	67%
<code>atr_pct</code> < 0.05 (thin bar) — STAND DOWN	9	-\$4.1	33%

§18.3 The honest size of the win

Be straight about what that gate buys. The win-rate split is real and it lines up with the amplitude floor we already run on the other momentum shadows (`atr_pct` $\geq$ 0.04). But the money swing is tiny: the thin-bar bucket only lost **\$4.1**, because — as the exit table showed — this sim's losers are already nearly free. So the floor isn't "make more money", it's "make the same money on half the trades with none of the dead-weight churn". That's a cleaner, tighter book, not a richer one. On one week and 18 readable trades this is thin, and like most single-week filters it mostly just means trade less — but here trading less costs you nothing and the thing you cut never won.

Verdict: **ARM / SIZE** **`atr_pct` $\geq$ 0.05 floor** — it made real money, beat its ceiling, and its exit spine already does the hard job. Give it the amplitude floor and put it on the promotion ladder.

LEAN IN — THIS ONE WORKS · \$+749 THIS WEEK ON HONEST FILLS.

The momentum sim made \$748.70 honest over 413 trades by riding ratchet winners and cutting stops fast; for Monday put an amplitude floor on the bar ($\text{atr_pct} \geq 0.04\%$).

§19 Momentum Shadow: it made money by letting winners run and getting stopped fast

The momentum sim made **\$748.70** this week on the honest tick-repriced fills, over 413 trades, winning just 42% of them. That's the truth from the join, not the sim's own optimistic ledger. It made money the way a momentum book is supposed to: a third of its exits ratchet up big, and the losers get cut at the stop before they hurt. It left a lot on the table versus a perfect fill (\$2,110 ceiling), and almost all of that gap is one contract behaving badly. Read on for the one thing.

\$748.70

HONEST REAL_PNL, 413
TRADES

42%

WIN RATE — WINS ARE
BIG, LOSSES SMALL

\$1,361.80

LEFT ON THE TABLE VS
PERFECT FILL

\$2,375.80

MADE BY RATCHET EXITS (74% WIN, N=131)

§19.1 Did it make money — yes, and here's exactly how

I re-ran the join myself. Real_pnl summed to **\$748.70** over 413 scored trades against a perfect-fill ceiling of \$2,110.50. So it captured about 35

cents of every ceiling dollar. The engine is dead simple and it's healthy: when a trade works it rides the ratchet, and when it doesn't it eats the stop. Look at the exit mix — the RATCHET exits made **\$2,375.80** across 131 trades and won 74% of the time, while the STOP exits bled **-\$1,810.30** across 173 trades and won only 15%. That's not a bug. That's the shape you want: small frequent losers paid for by fewer, fatter winners.

One thing worth flagging: the **shorts carried the week**. Short trades made \$594.80, longs only \$153.90, on near-identical trade counts (202 vs 211). And the best contract by a mile was **MNQ at +\$494.70** over 65 trades — which actually beat its own perfect-fill ceiling of \$391.60, meaning the fills it got were clean. Compare that to **MGC**: a \$900 ceiling, but only -\$10.30 in real money, because it bled an average of \$20.69 per trade to slippage. MGC is where the fill quality falls apart, not the signal.

EXIT REASON	REAL_PNL	N	WIN%
RATCHET (winner runs)	\$2,375.80	131	74.0
VWAP_LOSS	\$237.00	33	60.6
EOD flatten	-\$53.90	76	39.5
STOP (loser cut)	-\$1,810.30	173	15.0

§19.2 Why it wins vs loses — the bar has to be moving

Heads up on the evidence: the entry-condition funnel (rvol, ofi, book imbalance) only matched **9 of the 413 trades** this week — 2% coverage. That's too thin to say anything about the order book. So I don't guess; I use the one condition the sim records on every single trade: how fat the bar was when it fired, measured as ATR as a percent of price (`entry_atr_abs / entry_price`).

And that split is clean. When the bar was dead-thin — under 0.04% ATR — the sim fired 201 times and made almost nothing: **\$63.40, an average of 32 cents a trade**, and it got stopped out half the time (50% stop rate). When the bar had real amplitude — 0.04% ATR or more — it fired 212 times

and made **\$685.30, an average of \$3.23 a trade**, with the stop rate dropping to 34%. Same signal, but on a thin tape the momentum isn't real, the move dies, and the stop takes you out. This is the amplitude-floor story the desk already half-believes, showing up again in this sim's own numbers.



Average honest \$/trade by bar amplitude — thin <0.04% (\$0.32) · mid 0.04-0.08% (\$2.74) · wide 0.08-0.15% (\$5.12) · fat >0.15% (\$2.86). The floor is at the first red bar.

\$19.3 The one thing for Monday: put a floor under the bar

GATE: only fire when atr_pct >= 0.04%

Stand the sim down when the bar is too thin to carry a move. Below is the split on its own trades. The honest part first: this is **mostly a "trade less, trade cleaner" move, not a money geyser**. Cutting the thin bucket only adds about **\$63 in-sample**, because those thin trades weren't losing outright — they were just dead weight. But it halves the trade count, lifts the average from 32 cents to \$3.23 a trade, and drops the stop-out rate from 50% to 34%. That's a much healthier book firing far fewer shots for essentially the same money — and it lines up with the amplitude floor already armed observe-only on the live momentum shadows.

BUCKET	REAL_PNL	N	WIN%	AVG \$/TRADE	STOP-OUT%
Bar has amplitude (atr_pct >= 0.04%)	\$685.30	212	42.9	\$3.23	34.4
	\$63.40	201	40.8	\$0.32	49.8

Thin tape ($\text{atr_pct} < 0.04\%$)

— veto these

HONEST CAVEAT

One week, in-sample, and the strongest condition I could read is bar size — the order-book funnel only covered 2% of trades so I can't gate on rvol or imbalance yet. The +\$63 swing is real but small; the case for the floor is a cleaner, cheaper-to-run book (half the fills, same money, fewer stops), not a P&L jump. Separately, MGC is a fill-quality sink — \$20.69/trade of slippage turning a \$900 ceiling into -\$10 real; if you want a second lever, size MGC down or drop it before you touch the signal. The sim made money, the engine is sound: **arm the amplitude floor, keep watching the funnel coverage climb.**

GATE IT — ONE FILTER FIXES IT · \$+196 THIS WEEK ON HONEST FILLS.

orb_shadow made a real +\$195.90 over 144 trades on honest fills as a clean 35%-win breakout book, and dropping MCL and MYM — which lost even on their ratchet exits — lifts it to +\$415.20. **Skeptic re-ran the Monday change:** it FELL — read the change as a lead, not a rule. Not real: the true drop-MCL+MYM delta is +\$260 (USD, net of fees) not £173, and it is a one-week curve-fit whose entire +\$154 "kept book" rests on two MNQ trades -- strip them and it is -\$146, four of the five "kept" symbols lose money -- so there is no robust bankable weekly edge here (~£0).

§20 The Breakout Sim Pays — But Two Symbols Are Eating a Third of It

orb_shadow made real money this week: **+\$195.90** on the honest tick-repriced fills, across 144 trades, winning just 35% of them. That's the whole personality of a breakout book in one line — it's wrong most of the time and still green, because the few that go keep going. Its perfect-fill ceiling was \$970.70, so it left plenty on the table to slippage, but it kept a real, positive slice. The problem isn't the engine. It's that two of the seven symbols it trades are dead weight.

+\$195.90

HONEST REAL_PNL, 144
TRADES

35%

WIN RATE — AND STILL
GREEN

-\$219.30

MCL + MYM ALONE, 32
TRADES

+\$415.20

THE OTHER 5 SYMBOLS, 112 TRADES

§20.1 What it's actually doing: let the runner run, cut the rest

This is a clean breakout book and the exit tape proves it. The money is made entirely by the ratchet: 52 trades trailed out under RATCHET for **+\$1,042**, winning 58% of the time. Against that, 53 trades hit their STOP for **-\$863**, winning basically never (4%). The 38 that timed out at the end of day were a wash (-\$3.50). So the shape is simple and healthy — a third of the trades catch a move and get milked, half get stopped for a small controlled loss, and the survivors outrun the stops. That's exactly what you want a breakout to look like, and it's why it's green on a 35% hit rate.

The funnel conditions are no help here — only **7 of 144** trades matched a fut_signal_funnel row, so I'm not going to pretend an rvol or book-imbalance gate is proven on this sim. It isn't. The honest signal this week lives in the sim's own trade log, and it's loud: it's a symbol problem, not a condition problem.


§20.2 Where it wins vs where it bleeds

SYMBOL	TRADES	REAL P&L	WIN%	READ
MNQ	11	+\$317.7	64%	the star — clean index breakouts
MGC	25	+\$78.0	44%	gold trends, ratchet earns it
MBT	43	+\$21.5	44%	busy, small net positive
M2K	20	+\$11.6	30%	scrapes green
MES	13	-\$13.6	31%	roughly flat, thin n
MYM	16	-\$76.4	13%	breakouts that never break
MCL	16	-\$142.9	6%	crude runs the entry over

MNQ carried the whole book — +\$317.70 on 11 trades at a 64% hit rate. Gold and MBT chipped in. But MCL and MYM did the opposite of what a breakout is supposed to do. On MCL, 15 of 16 trades lost. And here's the tell: even MCL's ratchet exits — the ones that are supposed to be the

winners — came in at **-\$40.70** over 6 trades, and MYM's ratchet exits were **-\$10.20**. When the mechanism that makes your money is itself losing on a symbol, the symbol doesn't belong in the book. That matches what we already know about crude — it's a different beast, event-driven (OPEC/EIA), and the spread eats the edge. The breakout structure just doesn't hold there.

§20.3 The equity curve — climbed to **+\$300**, gave a third back



Cumulative honest real_pnl, trade by trade (gold dashed = breakeven). It rode up to about +\$300, then bled roughly a third of it back into the close — and a big part of that giveback is MCL and MYM stopping out over and over.

§20.4 The one thing for Monday

DROP MCL AND MYM FROM ORB_SHADOW

This is the cleanest split in the sim and it doesn't need a funnel row to prove it. Trade the same book on the other five symbols and the honest week goes from **+\$195.90 to +\$415.20** over 112 trades at a 42% win rate. The two you cut were 32 trades, **-\$219.30**, winning 9% — and losing even on their ratchet exits, which is the disqualifier. In-sample swing: **+\$219 for the week**.

CUT MCL + MYM

KEEP THE OTHER 5

Honest caveats, because it's one week and 32 trades: any single filter this good in-sample partly just means "trade less," and on a good crude week ORB might catch a real OPEC-driven run. But this isn't a blind data-mine — MCL being a bad fit for structural breakouts is a standing desk prior (it's event-driven and the crude spread eats the edge), and MYM's 13% hit rate rhymes with it. So this is a symbol the book already had reasons to distrust, and the tape agrees. I'd pull those two and re-check next week; if the remaining five hold their +\$400-ish shape across another week, orb_shadow has earned a look at the promotion ladder. Everything else in the sim — the ratchet, the stop discipline, MNQ leading — leave it alone. It's working.

GATE IT — ONE FILTER FIXES IT · \$+163 THIS WEEK ON HONEST FILLS.

This M2K loose-momentum sim made a real \$163 over 76 trades by trailing big breakouts, but bled every dollar it lost in hot tape — the Monday move is a rvol<1.3 gate to stop it chasing spent spikes.

§21 M2K loose-momentum: a winner that keeps chasing spikes it shouldn't

This one made money this week — \$163 honest, on tick-repriced fills, over 76 trades at a 41% win rate. That's the real number, and it beats nothing to fantasy: the perfect-fill ceiling was \$358, so slippage ate more than half the paper edge. But the money is real and it's positive. The job here isn't to save it — it's to stop it handing back the gains it's already earning.

\$163

HONEST REAL_PNL, 76
TRADES

41%

WIN RATE

\$358

PERFECT-FILL CEILING
(FANTASY)

\$311

FROM RATCHET EXITS ALONE

§21.1 Did it make money — yes, and it's a pure momentum shape

Score it on the honest fills and it's clean: **+\$163.4 over 76 trades, 41% win**. The shape is textbook loose momentum — a handful of big trailed winners paying for a lot of small stops. Its 29 ratchet exits made **+\$311 at**

an 83% win rate; its 37 stop-outs cost **-\$214 and won just 5% of the time**. Both sides trade roughly evenly (44 longs +\$92, 32 shorts +\$71), so there's no directional bias to fix. When it catches a run it trails it well. The leak is which breakouts it agrees to buy.

§21.2 Why it wins and why it bleeds — it's late to the hot ones

Pull the entry conditions from the funnel (I matched 51 of the 76 trades to a signal-funnel row within 8 seconds — 67% coverage; the other 25 have no funnel row and I can't see their conditions). The split is loud. On the trades it **wins**, average rvol at entry is **1.01** — normal, calm tape. On the trades it **loses**, average rvol is **1.68** — the tape is already hot when it fires.

In plain words: this is a breakout buyer, and it's best when the tape is quiet and the move is just starting. When volume is already spiking, the crowd is in ahead of it — it buys a rip that's already spent, and gets run over. The exit reasons prove it. In calm tape (rvol under 1.3) its ratchets locked **\$252 across 17 trades** and its stops cost only **-\$33**. In hot tape (rvol 1.3+) it's the opposite: **16 stop-outs for -\$132**, and the ratchet rescued just \$44. Same strategy, same week — the only thing that changed is it showed up late.

THE BEHAVIOUR IN ONE LINE


Quiet tape, it's early and the ratchet trails a real move to a locked win. Hot tape, it's late — it chases a spike the crowd already rode, and eats the stop.

§21.3 The one thing for Monday — stand down when rvol is hot

The single cleanest cut is a volume gate: **only fire when rvol is below 1.3**. On the 51 trades I can see conditions for, that one split does all the work.

BUCKET	TRADES	REAL P&L	WIN %	WHAT IT IS
Calm tape — rvol < 1.3	26	+ \$279.4	62%	early on a fresh move; ratchet trails it
Hot tape — rvol ≥ 1.3	25	-\$88.5	28%	late to a spent spike; stops out

The hot-rvol half is **-\$88.5 across 25 trades at a 28% win rate** — it's not marginal, it's a loss-maker outright. Veto it and the visible book climbs from **+\$190.9 to +\$279.4**, an in-sample swing of about **+\$88 for the week** (roughly £70), while cutting the trade count in half.



Real P&L by entry-rvol bucket (matched trades). Left to right: rvol < 0.8 (+\$135), 0.8-1.2 (+\$22), 1.2-1.6 (+\$79), 1.6-2.2 (-\$13), ≥2.2 (-\$32). Gold line marks the rvol=1.3 gate — everything green sits left of it, both red bars sit right.

§21.4 The honest caveat

This is one week and in-sample, and the gate only speaks to the 51 trades that had a funnel row — a third of the book is dark. And like most single filters, this one mostly just means trade less: it doesn't turn losers into winners, it declines to take them. But the signal here is unusually clean and it has a real mechanism behind it (late-to-the-spike, not a stat fluke), the same "amplitude/tape floor" logic already deployed on the other momentum shadows. Don't size this yet — it's not proven across regimes. Arm the rvol<1.3 gate observe-only alongside the others and let it forward-validate.

KEEP — it made \$163 honest

GATE rvol < 1.3

observe-only, don't size yet

GATE IT — ONE FILTER FIXES IT · \$+137 THIS WEEK ON HONEST FILLS.

The MGC coil-break long made an honest +\$136.92 over 134 trades (36% win, vs a \$712 perfect-fill ceiling); the Monday move is to only fire when price has already broken clear of VWAP. **Skeptic re-ran the Monday change:** it FELL — read the change as a lead, not a rule. It does not hold: the whole 134-trade sim genuinely nets +\$136.92 (~£106) this week, but the claimed +£190/wk gain is not a real edge — it is unreproducible (shadow_real.db stores no VWAP or any entry feature, and gate_feature_log has atr_pct/rvol/slope/flow but no vwap-distance, so the 57/77 split cannot be derived from the cited data), and taken at face value it is a single-week, in-sample, n=57 curve-fit whose "+£190" is exactly the -\$243 bucket of losers deleted after the fact (a "trade less by dropping this week's losers" partition, not a forward-validated gate), so the real reproducible improvement is ~£0.

§22 MGC coil-break: it pays, but only when the break has already left home

This one made money this week. On the honest tick-repriced fills, **tw_mgc_coil_break_long** booked **+\$136.92 over 134 trades** at a 36% win rate. Small, but green — and it did it while its perfect-fill ceiling was \$712, so most of the fantasy got eaten by slippage. The real story is simple: when it fires on a break that's already pulled away from VWAP, it runs; when it fires right on top of VWAP, it's just buying fake breakouts in chop and feeding the stop.

+\$136.92

REAL_PNL, 134 TRADES

36%

WIN RATE

\$712

PERFECT-FILL CEILING
(FANTASY)

+\$380

IF IT ONLY TOOK EXTENDED BREAKS

§22.1 What it's actually doing

It's a coil-break long on gold. The shape of the week is textbook trend-follower: a small number of big winners carry a big pile of small losers. All 48 of its green trades exited on the **RATCHET** (the trailing lock) for **+\$2,010**. The 74 losers all died on the **STOP** for **-\$1,347**. Ten more scratched out on a VWAP exit for a token **+\$72**. So the machinery works — winners are allowed to run, losers get cut — and it nets positive by a whisker.

The question is which breaks are real. I pulled the entry conditions off the signal funnel for every trade (134/134 matched — full coverage, using the last funnel tick before each fire). ATR, RVOL, OFI and spread barely move between winners and losers — they're not the tell. The one condition that separates them is **how far price already was from VWAP at entry**.

Winners fired at an average VWAP distance of **0.053%**; losers at **0.033%**. In plain words: the trades that worked had already broken out and moved. The trades that bled were fired right on the VWAP, on a coil that hadn't actually committed yet — a fake break in the chop.

THE PATTERN


Fire on top of VWAP and you're guessing the break — you get 31% winners and a slow bleed. Wait until the break has already left home (VWAP distance $\geq 0.04\%$) and the same ratchet catches real runs —

42% winners, and the winners are **2.5× bigger** (\$1,025 of ratchet gains vs \$410 on the near-VWAP side).

§22.2 The one thing for Monday: only fire when the break has left VWAP

Split its own 134 trades by VWAP distance at entry. The line at **0.04%** cleaves the winners from the losers cleanly:

VWAP DISTANCE AT ENTRY	TRADES	REAL_PNL	WIN%	\$/TRADE	
< 0.02% (on the VWAP)	64	-\$148.68	33%	-\$2.32	chop
0.02 - 0.04%	13	-\$94.56	23%	-\$7.27	chop
0.04 - 0.06%	11	+\$175.68	55%	+\$15.97	real break
≥ 0.06% (well clear)	46	+\$204.48	39%	+\$4.45	real break



Real_pnl by VWAP distance at entry. Red = fired on the VWAP (chop). Green = fired after the break already moved.

The gate: only take the coil break when VWAP distance $\geq 0.04\%$ at entry. On this week's own trades that keeps **57 trades for +\$380.16 (42% win, +\$6.67/trade)** and throws away **77 trades that netted -\$243.24 (31% win, -\$3.16/trade)**. Same ratchet, same stop — you just stop taking the uncommitted breaks. In-sample money swing this week: **+\$243**, i.e. from +\$137 to +\$380 on the identical machinery.

HONEST CAVEAT

This is one week, 134 trades, in-sample. The gate cuts 57% of the volume, so part of what you're seeing is just "trade less in the chop" — but it's not only that: per-trade EV genuinely flips from -\$3.16 to +\$6.67 and win rate climbs 31%→42%, and there's a real mechanism (an extended break is a confirmed break). I'd want it to hold up on a second week before sizing. For now: **GATE vwap_dist ≥ 0.04%** in observe-only, and keep it running — it's already green.

Verdict: **KEEP** it — it makes money as-is — and **add the VWAP-distance gate** to see if next week confirms the +\$243 swing before this earns real size.

GATE IT — ONE FILTER FIXES IT · \$+127 THIS WEEK ON HONEST FILLS.

MGC thrust made +\$126.9 on honest fills over 9 rows (really 6 entries), almost all from three Monday gold shorts; the Monday move is to only fire when the entry bar sits in a mid amplitude band and stand down on dead-quiet or fat/shock bars.

§23 MGC Momentum-Thrust: it made money, but the whole week is six trades

This little MGC thrust sim made **\$126.9** on the honest fills this week over 9 scored rows, 56% of them green. That's a real profit, and it landed on the right side of the tape. But before you get excited: those 9 rows are exit-variants of just **6 actual entry decisions**, and almost every dollar came from three gold shorts on Monday morning. It's a good week from a very thin book.

+\$126.9

HONEST REAL_PNL, 9
ROWS

56%

WIN RATE

+\$156

THE SHORTS ALONE

6

REAL ENTRY DECISIONS

§23.1 1. The money — it's real, and it's almost all short

I re-ran the join myself. Nine scored rows, real_pnl summing to **+\$126.9**, against a modelled perfect-fill ceiling of **-\$104**. Yes, the honest number came in above the ceiling. Don't read that as magic fills — it means the ceiling model is mispriced for this sim (it stamps the winning shorts at -\$19

when the tick tape paid +\$51 and +\$78). Trust the real column, ignore the ceiling here.

Split it by side and the story is blunt: the six **short** rows made **+\$156** at 67% win, the three **long** rows lost **-\$29** at 33%. This wasn't a momentum machine having a great week — it was three well-timed gold shorts on Monday between 05:00 and 09:25 UTC carrying the whole thing, while the longs bled. Exits helped: the winners rode **STOP** (trailing) exits for +\$142; the two RATCHET exits netted -\$15.


§23.2 2. Why it won — Goldilocks amplitude

Here's the clean thing, and it's on `entry_atr_abs` which I have for **all 9 rows** (full coverage, no gaps). Sort the six entries by how fat the bar was when the sim fired:

- Bar too **quiet** (ATR 1.03, the 22:33 long) — **lost** -\$14. Nobody home, no thrust to ride.
- Bar **just right** (ATR 1.70-1.73, three entries) — **won every time**: +\$85, +\$128, +\$17.
- Bar too **fat** (ATR 2.39 and 3.41) — **lost both**: -\$56 and -\$32. The 3.41 one fired in a **shock** regime and was the single worst trade of the week.

So it's not a floor and it's not a ceiling — it's a **band**. This sim wins when there's enough amplitude to carry a thrust but not so much that the bar overruns its own stop. That rhymes with the desk's amplitude-floor thesis, and adds the other half: fat/shock bars run momentum entries over.

The funnel backs this up where it has coverage (only 3 of 9 rows — the early-morning shorts have no funnel rows captured, so this read is 3 points). The one winner sat at `rvol 3.93 / atr_pct 0.038 / trending`. The two losers with funnel were the dead-tape long (`rvol 0.6, atr_pct 0.022`) and the shock long (`rvol 6.78, atr_pct 0.082`). Middle wins, both extremes lose.



Each dot is one of the six entry decisions: horizontal = how fat the bar was, vertical = honest real_pnl. Green wins, red losses. The three winners sit inside the gold band; all three losses fall outside it.

§23.3 3. The one thing for Monday — veto the extremes

Only fire this sim when the entry bar sits in the mid amplitude band, roughly **ATR 1.5-2.0 absolute on MGC (atr_pct ~0.035-0.05)**. Stand down when the bar is dead-quiet (atr_pct < 0.03) or fat/shock (atr_pct > 0.055 / shock regime). The split on its own trades:

BUCKET (ENTRY ATR ABS)	N	REAL_PNL	WIN%	WHAT IT IS
In-band 1.5-2.0	5	+\$229.4	100%	enough thrust, not overrun
Out-of-band (1.03 / 2.39 / 3.41)	4	-\$102.5	0%	dead tape or shock bar

In-sample the veto drops four losing rows worth **-\$102.5** and keeps five that never lost — it would have turned this week's +\$127 into +\$229.

GATE atr_pct band

HONEST CAVEAT

This is **six entry decisions**. A band tuned to exactly where three winners sat is textbook in-sample overfit — a 5/5 vs 0/4 split on n=9 rows will happen by luck more often than you'd like. The direction is trustworthy and matches the desk's amplitude thesis (don't fire

momentum into a dead tape, don't fire into a shock bar), but the exact 1.5-2.0 numbers are not. And funnel coverage is only 3/9, so the rvol/atp_pct confirm rests on three points. **Don't arm or size this on one week.** Keep it in shadow, add the fat-bar / dead-tape veto, and let it reprove the band over another 4-6 weeks before it goes anywhere near the promotion ladder.

GATE IT — ONE FILTER FIXES IT · \$+49 THIS WEEK ON HONEST FILLS.

tw_mes_thrust_cont made an honest +\$49.1 over 26 trades this week; the fix is to veto fat bars ($\text{atr_pct} \geq 0.028$), which lifts the readable book to +\$78.4.

§24 MES thrust-continuation: a real little winner, and one fat-bar leak

This one made money on the honest fills. It cleared **\$49.1 over 26 trades** this week, won 46% of them, and every fill actually landed (26/26 "filled"). It didn't beat its perfect-fill ceiling of \$175.7 — the tick-repriced exits shaved about \$127 off the fantasy — but the direction of travel is green, and that's rare on this desk.

+\$49.1

HONEST REAL_PNL, 26
TRADES

\$175.7

PERFECT-FILL CEILING

46%

WIN RATE (12 OF 26)

+\$1.9

AVG EDGE PER TRADE

§24.1 What it's actually doing

It's a thin edge that stands on one leg. The wins average **+\$20.4** and the losses average **−\$13.9**, so the winners are bigger than the losers by enough to carry a sub-50% hit rate. All of the money comes out of the ratchet exit: the **RATCHET** exits made **+\$170.4 across 9 trades and won 89%** of them, while the **STOP** exits handed back **−\$86.1 over 13 trades** at a 23% win rate and the four **VWAP_LOSS** cuts lost another **−\$35.2**. In plain

words: when a thrust keeps going, the ratchet trails it and banks it; when the thrust was already spent, it gets stopped out. The exit spine is fine. The question is which entries it should have skipped.

The tell is the bar it fired on. I could match real entry conditions from the signal funnel on **18 of the 26 trades** (the other 8 are a stale-session cluster with no funnel row within thousands of seconds — I've excluded them from the split and they netted a harmless $-\$6.7$ anyway). On those 18, the size of the bar at entry sorts the winners from the losers almost cleanly. When it fired on a **normal-to-thin bar (atr_pct below 0.028)** it made money and won most of the time. When it fired on a **fat bar (atr_pct at or above 0.028)** it was a coin flip that paid once and bled six times — it's chasing a move that has already happened.

§24.2 The one thing for Monday: stand down on fat bars

Veto entries when the bar is fat — **only fire when atr_pct < 0.028**. Here's the split on the 18 trades I could read:

BUCKET (18 READABLE TRADES)	N	REAL_PNL	WIN%	READ
Thin/normal bar — atr_pct < 0.028	11	+\$78.4	64%	the edge lives here
Fat bar — atr_pct ≥ 0.028	7	−\$22.6	14%	chasing a spent move
Stale-funnel cluster (no read)	8	−\$6.7	50%	can't gate, ~breakeven

Drop the seven fat-bar trades and the readable book goes from $+\$55.8$ to $+\$78.4$ — an **in-sample swing of about $+\$22.6$** on the honest fills, and the win rate on what's left jumps to 64%. The chart shows why: left of the gold line is mostly green and above water, right of it is a wall of red.



Each dot is a readable trade — x is bar fatness (atr_pct, thin left → fat right), y is honest real_pnl (top = win). Gold line = the 0.028 fat-bar veto; dashed = breakeven. Green above the line, red to the right of it.

HONEST CAVEAT

It's one week and thin: seven fat-bar trades, and the veto also throws away this sim's single best trade (a fat-bar short that ran to +\$75.9). Net of that, ditching the fat bucket still adds +\$22.6 in-sample, so it's the right trade — but read it as "the fat bars are a lottery ticket, not an edge," not as gospel. Worth noting the direction split too: SHORTs made +\$93.1 over 12 trades and LONGs lost −\$43.9 over 14, which is probably just a week where the market fell and the down-thrusts kept going. I'd gate on the bar, not the side — atr_pct is mechanistic, side is this-week's tape.

KEEP — it's green

TUNE — veto atr_pct ≥ 0.028

WATCH — one-week, thin n

Verdict: leave it armed, it's already paying, and add the fat-bar veto for Monday to lift the honest week from ~\$49 toward ~\$72 while cutting the trade count. Give it another week or two of forward data before sizing it up — the edge is real but small, and it leans on a handful of ratchet winners.

WATCH — PROMISING, NOT PROVEN · \$+28 THIS WEEK ON HONEST FILLS.

MYM open-thrust made +\$28.06 honest over 12 trades but one +\$56 ratchet ride carried it; Monday move is a trending-only gate that drops the chop fires.

§25 MYM open-thrust: green, but one trade carried the whole week

This little Dow-micro open-thrust sim made money this week — \$28.06 real over 12 fills, 58% winners, and it even beat its own perfect-fill ceiling of \$13.67. Good news on the tin. The honest read underneath is thinner: strip out one ratchet ride and the other eleven fills lose money together.

+\$28.06

REAL P&L, 12 TRADES

58%

WIN RATE (7 OF 12)

+\$13.67

ITS CEILING — REAL
BEAT IT


+\$56.38

THE ONE TRADE THAT MADE THE WEEK

§25.1 1) Did it make money? Yes — \$28 honest, and it beat its ceiling

Scoring on the tick-repriced fills, not the sim's optimistic P&L: **+\$28.06 over 12 trades, 7 winners.** The ceiling (perfect-fill fantasy) was only \$13.67, so the real fills actually came in ahead — driven by one ratchet exit (trade 11378) the ceiling model priced at a loss but that really banked +

\$56.38. That's the tell for the whole section: take that single trade out and the remaining eleven fills net $-\$28.32$. The week was green because of one ride, not a dozen.



Cumulative honest P&L, trade by trade (gold dashes = break-even). One jump at trade 2 (the $+\$56$ ratchet), then a slow bleed from a $\$68$ peak down to the $\$28$ close. The whole curve lives above zero on the back of that one step.

§25.2 2) Why — it works when the tape is trending, it bleeds when it chops

Every fill this week was a LONG — it's an open-thrust sim, it buys the pop. What separates the good days from the bad is dead simple and it's exactly what you'd hope for a momentum sim: **the regime it fired into.**

When the session tape was **trending**, the thrust followed through: 5 trades, **$+\$47.90$** , 60% winners. When it was **chop**, it bought a pop that immediately died and got slapped back into its stop: 7 trades, **$-\$19.84$** , and even a 57% "win rate" couldn't save it because the losers were fat.

The exit mix says the same thing from the other side. The four hard **STOP** exits cost it **$-\$81.48$** — that's the entire bleed. Everything that didn't get stopped (ratchet, flat-clock, chandelier, session-flat) made **$+\$109.54$** between them. When the thrust is real it rides out to a trailing exit; when it's a false pop in chop, it eats the stop. Same entry, two very different tapes.

HONEST CAVEAT

Both buckets are tiny (5 and 7 trades, one week) and the trending bucket itself leans on that single +\$56 ratchet. Take it out and trending is +\$10.14 over 4 — still green, still beats chop's −\$19.84, but nobody should pretend this is settled. It's a lean, not a law.

§25.3 3) The one thing for Monday — stand down in chop, only fire when trending

The cleanest split its own trades offer is session regime, and it matches the sim's whole reason for existing: an open-thrust buyer should only be armed when the tape is actually going somewhere. Gate it to **trending regime only** and drop every chop fire.

BUCKET	TRADES	REAL P&L	WIN%	WHAT IT'S DOING
Trending (KEEP)	5	+ \$47.90	60%	buys the pop, the trend carries it out to a trailing exit
Chop (STAND DOWN)	7	−\$19.84	57%	buys a false pop that dies and stops out — two −\$20+ hits
As-is (both)	12	+\$28.06	58%	green, but carried by one trade

The gate: **only fire when session_regime = trending** . In-sample that lifts the week from +\$28.06 to **+\$47.90 — a swing of +\$19.84** by simply not trading the chop. It doesn't add a single winner; it just stops handing back money on days the sim was never built for.

Verdict: **WATCH** — it's green and the fix is sensible and on-theme, but one week and one trade is not a mandate. Arm the trending-only gate in the shadow and let it run another two or three weeks. If trending stays positive without needing a lottery ratchet each week, that's the week you size it. Not yet.

CHAPTER

The Bleeders — fix or retire these

RETIRE — NO HONEST FIX · \$-4249 THIS WEEK ON HONEST FILLS.

passive_chop lost \$4,249 over 1,088 trades at 17% win with a perfect-fill ceiling that is still red (-\$890), so the entry is -EV before slippage and no hour/bar/symbol gate flips it green — retire it.

§26 passive_chop: it fades the chop, and the chop keeps going

This one bled the worst on the board. On the honest fills it lost **\$4,249** over **1,088** trades this week, winning just **17%** of the time. And here's the tell: its perfect-fill ceiling was only **-\$890**. So even if every limit had filled at the dream price, this sim still loses money. Real fills just made a bad entry 4.8x worse. This isn't a slippage problem. The entry itself is -EV.

-\$4,249

HONEST REAL_PNL,
1,088 TRADES

17%

WIN RATE

-\$890

PERFECT-FILL CEILING
(STILL RED)

3:1


AVG LOSER -\$10.32 VS WINNER +\$3.47

§26.1 What it's actually doing

passive_chop parks a limit order inside quiet, ranging tape and waits to get hit. When it fills, one of two things happens. Either price ticks back a hair and it scratches out a small win, or price fills the limit because it's leaving and keeps right on going until the stop. That second case is the whole story.

Split it by how each trade ended and the asymmetry is brutal. Of the 1,088 orders, **419 never filled at all** (free, \$0). Of the 669 that did fill, the ratchet exits look lovely on paper — **235 trades, 76% winners, +\$563** — but each one only banks about **\$2.40**. Meanwhile the **402 stop-outs lose — \$4,789**, roughly **−\$11.91 apiece**. It wins small and often, then hands it all back plus rent on the trades that run it over. That's a negative-skew fade, and the exit spine can't save an entry that's fading into moves that keep going.

HOW IT ENDED	TRADES	REAL P&L	AVG / TRADE	WIN%
STOP (ran it over)	402	−\$4,788.7	−\$11.91	0%
RATCHET (small scratch win)	235	+\$562.9	+\$2.40	76%
FLAT_CLOCK (timed out)	32	−\$23.5	−\$0.74	28%
no_fill (limit never hit)	419	\$0.0	—	—



The picture is the point: the green sliver of ratchet wins is invisible next to the red column of stops. No exit tweak closes that gap because the loser is five times the size of the winner by design.

§26.2 Is there a gate that saves it? I looked. There isn't.

First, the honest caveat: I wanted the entry conditions from the signal funnel, but **only 16 of these 1,088 trades have a funnel row** (1.5% coverage), so rvol / OFI / book-imbalance splits are off the table this week —

not enough data to stand on. So I split on what the sim's own log does record: instrument, hour, and bar width. Every slice bleeds.

- **By hour:** there is no green hour. The US afternoon (13–15 UTC) wins 44–53% of the time, far better than the 9% it manages in the European morning — but it's still negative every hour. The best single bucket I could carve, US afternoon excluding bitcoin, was 72 trades at 57% win and still **–\$264**.
- **By bar width:** 1,073 of the trades fire at `atr_pct` under 0.15% — this thing only ever fires in dead-flat tape, so there's no "fat bar" to veto. The bleed is the quiet tape.

The one lever with real money on it is the instrument. **MBT (micro bitcoin) is a disaster inside this sim:** 365 trades, a 5.2% win rate, –\$861. That's a third of the trades for the worst hit rate on the desk. Cutting it is the single biggest defensible save. But look at what's left.

BUCKET	TRADES	REAL P&L	WIN%
MBT (the veto candidate)	365	–\$860.8	5.2%
Everything else (kept)	723	–\$3,388.5	23.4%

Killing MBT saves about **\$861 in-sample** and removes the ugliest instrument — but the remaining book still loses **\$3,388**. There is no gate on `rvol`, hour, bar size, or symbol that turns this green. The bleed isn't a bad condition you can fence off; it's baked into a –EV fade whose ceiling is already red.

THE ONE THING FOR MONDAY

RETIRE `passive_chop` Not "tune it" — retire it. Its perfect-fill ceiling is –\$890, so the entry loses money before slippage ever gets a vote, and every hour, every bar size, and every instrument bleeds. If you want to keep it breathing on the shadow bench purely to watch, the minimum honest move is to **veto MBT** (saves ~\$861/wk, 5% win rate is indefensible) — but don't kid yourself that it fixes anything; the rest still drops \$3,388 a week. This lines up with the reversion autopsy in

memory: passive_chop is already flagged for delete because fading quiet tape fades into grinds that keep going. This week is the receipt.

All money on sr.shadow_real_pnl.real_pnl, week filter entry_ts $\geq$ 2026-07-06. Funnel coverage $16/1088 = 1.5\%$, so entry-condition splits use the sim's own trade log. One week, in-sample; the MBT save is a real cut, not a curve-fit — but no filter here flips the sign.

GATE IT — ONE FILTER FIXES IT · \$-960 THIS WEEK ON HONEST FILLS.

This fade sim bled \$959.9 over 256 trades (36% win, vs a +\$341 perfect-fill ceiling); the whole loss is fast, fat-bar names (MBT/MNQ/MCL, 108 trades -\$670 at 20% win) so the Monday move is to veto those three names. **Skeptic re-ran the Monday change:** it HELD. The arithmetic reproduces to the dollar (veto MBT/MNQ/MCL = 108 trades, -\$669.9/~£528 in-sample swing; calm residual 152 trades, -\$292 still red), so the claim as written holds, but it is a post-hoc 'drop the three worst-losing names' cut in one in-sample week with an unverifiable $\text{atr_pct} \geq 0.10\%$ dressing (only 7/256 trades have the funnel data) that cauterises a bleed without creating an edge, so it is correctly observe-only and NOT promotable, not a validated forward improvement.

§27 tw_deadtape_standdown: it stood down in the quiet and got run over in the fast

This one bled, plainly. On the honest tick-repriced fills it lost **\$959.9 over 256 trades** this week, winning just **36%** of them. Its own perfect-fill ceiling said it should have made **+\$341** — so the gap between the fantasy and the truth is \$1,300, and every dollar of that gap is fills and stops going against it. The name promises it stands down in dead tape. The tape it actually died in was the fast stuff.

-\$959.9

HONEST REAL_PNL, 256
TRADES

+\$341.4

PERFECT-FILL CEILING
IT MISSED

36%

WIN RATE

-\$935.7

FROM STOP EXITS ALONE (83 TRADES, 13% WIN)

§27.1 Why it bleeds: it leans on fast, event-driven names and the stop eats it


The money splits cleanly two ways. First by exit: the **RATCHET** exit is the only thing working — 87 trades, **+\$311.8**, 62% win. Everything good this sim does, the ratchet banks. But the **STOP** takes 83 trades to **-\$935.7** at a 13% win rate. That is the whole loss and then some. When this thing is wrong it is very wrong, and the stop is where it pays for leaning into a move that just keeps going.

Second by symbol, and this is the real tell. Three fast, event-driven names — **MBT** (micro bitcoin), **MNQ** (nasdaq) and **MCL** (crude) — carry **108 trades for -\$669.9 at a 20% win rate**. The four calmer index-and-gold micros (MGC, MES, MYM, M2K) are 148 trades for -\$290 but win **47%** — a normal, near-coin-flip book that the ratchet almost keeps afloat. On the fast names the ratchet is dead (+\$-11 across 22 trades) and the sim just holds losers into the close: end-of-day exits on those three names alone give back **-\$286.5**.

Put in human words: this is a fade/lean sim. On the calm names it fades a wiggle and the ratchet gets it out clean about half the time. On crude, bitcoin and nasdaq it fades a rip that just runs it over, the stop is too far to save it, and what the stop misses the close finishes off. Its own `atr_pct` backs this up — the fat-bar names are exactly the losers (MCL 0.096%, MBT 0.071%), the thin ones (MYM 0.019%, MES 0.026%) are the survivors.

HONEST CAVEAT

Only **7 of 256** trades (2.7%) have a funnel row this week, so I cannot honestly gate this on `rvol`, `OFI`, book imbalance or spread — the coverage isn't there. Every split below is built from the trade log itself (symbol, exit, `atr_pct`), which is fully populated. And this is one in-sample week.



Cumulative honest real_pnl across the week, trade by trade. Gold dashed line is breakeven. It never comes back — a straight grind from \$0 to -\$929.

§27.2 The one thing for Monday: stand down on the fast names

There is no single gate that flips this green — it bleeds in every bucket. But the cleanest, biggest, most defensible cut is by name:

stand down on MBT / MNQ / MCL (equivalently, veto any bar with atr_pct $\geq 0.10\%$). Those three fast, fat-bar names are where the ratchet stops working and the sim leans into moves that run it over.

BUCKET	TRADES	HONEST REAL_PNL	WIN%
Fast names MBT / MNQ / MCL (the veto)	108	-\$669.9	20%
Calm names MGC / MES / MYM / M2K (what's left)	148	-\$290.0	47%
Cross-check: fat bars, atr_pct $\geq 0.10\%$	28	-\$223.0	27%

Vetoing the three fast names removes **-\$669.9** of bleed in-sample — a **+\$670 swing** on the week. But be straight about what that buys you: the remaining calm-name book is still **-\$290**, red. This gate stops the hemorrhage, it does not make the sim pay. So the honest read is

GATE it hard to the calm four and keep it observe-only — not a promotable edge, a bleed you've cauterised while you watch whether the calm-name ratchet can drag it to breakeven. If it's still red in two weeks with the fast names gone, retire it.

MONDAY

Veto MBT, MNQ and MCL on this sim (or gate on $\text{atr_pct} < 0.10\%$). In-sample that saves **\$670** and drops the worst-losing, lowest-win third of the book. It does not flip green — residual is still -\$290 — so keep it observe-only, not on the promotion ladder.

RETIRE — NO HONEST FIX · \$-891 THIS WEEK ON HONEST FILLS.

Scored this clock sim's honest fills and it bled -\$891 over 219 trades at 37% win with no green condition bucket, so the Monday move is to retire it (or at minimum cut its three worst symbols to stop losing \$675/wk to watch it).

§28 tw_clock_standdown: a slow, all-week bleed with no rescue gate

This one lost money every day, on every symbol group, in every regime. On the honest tick-repriced fills it bled **-\$891** over 219 trades at a 37% win rate — and its perfect-fill ceiling was only +\$241, so even in fantasy-land it barely broke even. There is no single gate that flips this green. The one thing worth doing is cutting its three worst symbols, which saves \$675 in-sample but still leaves it red. The clean call is to retire it.

-\$891

HONEST REAL_PNL, 219
TRADES

37%

WIN RATE

+\$241

PERFECT-FILL CEILING
(BEST CASE)


-\$10.5

AVG LOSS VS +\$6.9 AVG WIN

§28.1 1) Did it make money? No — it bled all week

I re-ran the join myself on `shadow_real_pnl.real_pnl`: 219 trades this week, sum **-\$891.4**, 37% winners, against a perfect-fill ceiling of +\$240.9.

So the fills didn't kill it — the entries did. The equity curve below never recovers; it just grinds down from zero to about $-\$900$ and stays there.



Cumulative honest P&L, trade by trade in entry order. Gold dashed line is break-even. It touches $+\$5.64$ early and then it's all downhill to $-\$902$.

§28.2 2) Why it loses: small wins, bigger losses, and it stops out constantly

The shape tells the whole story. When it wins it clips a small $+\$6.9$; when it loses it gives back $-\$10.5$. Lose more than half the time at that ratio and you bleed — which is exactly what happened. Break it out by how it exits and you see the engine of the loss:

EXIT MIX

The **stop** is doing the damage: 73 trades stopped out for $-\$834$ at a 15% win rate. The ratchet is the only thing that works — 71 trades ratcheted out for $+\$243$ at 65% win. The clock exits (FLAT_CLOCK, 29 trades) were roughly flat at $-\$49$. So this is a low-quality entry that, when it's wrong, rides all the way to the stop; when it's right, the ratchet banks a small winner. The wins can't cover the stops.

And it's wrong a lot because it fires into the wrong tape. I matched 175 of the 219 trades to funnel conditions (80% coverage, minute-bucket nearest match) and split every way I could — and nothing is green:

- **Volume:** low-rvol trades ($\text{rvol} < 1.0$) were 160 of them at $-\$764$; the high-rvol handful (15 trades) still lost $-\$76$. Standing down on thin tape trims it, doesn't fix it.
- **Bar size:** quiet bars ($\text{atr} < 0.05$) lost $-\$464$, fat bars ($\text{atr} \geq 0.05$) lost $-\$376$. Both red.
- **Regime:** trending $-\$420$, chop $-\$305$, shock $-\$111$. Red, red, red.

There is no condition bucket with real n that this sim makes money in. That's the finding.

§28.3 3) The one thing for Monday: cut its three worst symbols — but it's still red, so retire it

The only split that separates real winners from losers is **symbol**. Three names carry almost all the loss: MBT (crypto), MCL (crude), and MNQ. Those are the "different-beast" event-driven / high-vol tickers, and this clock strategy gets run over by them — 95 trades, $-\$675$, a 22% win rate. The four calmer index/gold names win about half their trades but still can't turn a profit because the loss size beats the win size.

BUCKET	N	REAL P&L	WIN%	READ
Bleeders: MBT / MCL / MNQ	95	$-\$674.8$	22%	gets run over — cut these
Keep: MYM / MGC / MES / M2K	124	$-\$216.6$	48%	coin-flip hit rate, still red on EV
Whole sim as-is	219	$-\$891.4$	37%	—

So the single Monday move that helps most is: **stand this clock down on MBT, MCL and MNQ**. That's a clean $+\$674.8$ in-sample swing. But be honest with yourself — it takes the sim from $-\$891$ to $-\$217$, which is a smaller loss, not a profit. Even the four "good" symbols only manage a coin-

flip win rate and lose money because the average loss (\$10.5) is bigger than the average win (\$6.9). You can't gate your way out of a negative-EV entry.

VERDICT

RETIRE This sim has no green pocket. Trimming the three worst tickers saves \$675 but still leaves it red, and the win/loss shape (37% win, losses 1.5× the wins) is structurally negative-EV. One week and in-sample, yes — but there is no honest tweak here, only a smaller bleed. Pull it. If you want to keep it alive purely to watch, at minimum drop MBT/MCL/MNQ so it stops costing \$675 a week to observe.

Money scored on `shadow_real_pnl.real_pnl` (never the sim's optimistic `pnl_usd`). Week filter `entry_ts ≥ 2026-07-06`. Condition splits on 175/219 trades matched to `fut_signal_funnel` by nearest minute bucket (exact `entry_sid` join was broken — only 4/219 lined up, so this is approximate entry context, not exact).

GATE IT — ONE FILTER FIXES IT · \$-627 THIS WEEK ON HONEST FILLS.

Confirmed the sim bled -\$627 over 198 honest fills, traced the entire wound to the short book (-\$685 over 85, stopped 65% of the time) while the longs paid their way (+\$58 over 113), so the Monday move is a long-only gate. **Skeptic re-ran the Monday change:** it FELL — read the change as a lead, not a rule. The ~£541 reproduces exactly (killing the -\$705.69 short book at ~0.767 GBP/USD), but it is loss-avoidance, not an edge: long-only still nets -\$40 gross (worse after fees — the sim is cost-blind), the surviving long side is propped up by a backfill batch while the live longs lost -\$135, and the whole sample is a single Monday session — so it holds as 'stop the bleed / trade less,' not as a profitable refinement.

§29 tw_ride_strength: it rides strength fine going up, and gets run over going short

This one lost real money this week. On the honest tick-repriced fills it bled **-\$627** across 198 trades at a 38% win rate. Odd footnote: it actually lost less than its own perfect-fill ceiling of **-\$844** — so the fills weren't the problem. The **direction** was.

-\$627

HONEST REAL_PNL, 198
TRADES

+\$58

THE LONG BOOK, 113
TRADES

-\$685

THE SHORT BOOK, 85
TRADES

§29.1 What it's actually doing

Split the week by side and the whole story falls out. The longs — buying strength and holding — roughly paid their way: **+\$57.7 over 113 trades at 43% win**. The shorts are the entire wound: **-\$685 over 85 trades at 30% win**. Same engine, one side works and one side gets steamrolled.


The exit mix says why in plain terms. When a long survives to the 20-minute clock it books green — 52 clock-outs worth **+\$817**. The shorts almost never get there: only 19 clock-winners for +\$312, because they keep hitting the stop first. Shorts got **stopped 55 times out of 85 (65%)** for −\$1,377; longs got stopped 50 of 113 (44%). This is a market that ground higher all week, and the sim kept selling into it and getting stopped. Nobody's rewarding a short here — it's fading a tape that just walks away from it.

THE TELL

Shorts get stopped 65% of the time vs 44% for longs, and produce a third as many clock-winners. The sim isn't broken on entries — it's fighting the drift on one side only.

§29.2 The honest-money curve: the short book is the whole drawdown

Navy is the sim as it actually traded. Green is the same sim with the shorts removed. The navy line rides comfortably positive for the first ~145 trades, then a cluster of shorts craters it to −\$640. The green line never leaves positive territory.



Navy = all 198 trades as traded (ends −\$627). Green = long-only (ends +\$58). Gold dashes = break-even. Cumulative honest real_pnl, in trade order.

§29.3 The one thing for Monday

Kill the short side. That's it. This isn't a subtle rvol or bar-fatness gate — I looked, and inside the shorts every condition bucket still bleeds (thin bars, fat bars, chop, trending, all red). The clean separator is direction, and it's on the full sample with no funnel gaps to hide behind.

BUCKET	N	REAL_PNL	WIN%	STOPPED%
LONG — keep	113	+\$57.7	43%	44%
SHORT — the bleed	85	-\$685.0	30%	65%
Long-only sim	113	+\$57.7	43%	44%

GATE: long-only Dropping the shorts swings this sim **+\$685 in-sample**, from **-\$627 to +\$58** on the week. It turns a bleeder into a break-even-plus book.

HONEST CAVEAT

This is one week, and MNQ/MES drifted up the whole time — so "long-only" is partly just "don't fight this week's tape." A single week can't tell you shorts are structurally dead. But two signals here aren't week-specific: the shorts win 30% vs the longs' 43%, and they get stopped 65% vs 44%. That's a real asymmetry worth gating on now and re-checking next week rather than eating another **-\$685**. If next week trends down and the shorts pay, revisit — but as it stands, don't arm the short side.

Funnel-condition coverage was only 85/198 (43%) — entry_sid is a session id, not a cycle timestamp, so conditions were matched to the nearest fut_signal_funnel row within 4s. The side split above needs none of that; it's the full 198-trade sample on r.real_pnl.

GATE IT — ONE FILTER FIXES IT · \$-542 THIS WEEK ON HONEST FILLS.

This gold reversion sim bled -\$542.2 over 63 trades on honest fills, and every dollar of it is the short side (33 shorts -\$555.7 vs 30 longs +\$13.5) — the Monday move is to veto shorts and take longs only. **Skeptic re-ran the Monday change:** it FELL — read the change as a lead, not a rule. The ~£433 headline roughly reproduces — MGC's short leg really lost about \$555.72 (~£434, already net of round-trip fees) this one week — but it does NOT hold as an edge: it's a thin-n (31 filled shorts), single trend-week 'cut the losing leg / trade less' amputation, the retained long leg is only +\$13.52 net and goes negative once you strip its single lucky +\$69.88 TARGET, and the claim's own supporting stats don't reproduce (actual 31 filled shorts / 16 stops / -\$647.92, not 32 / 19 / -\$506.3, with one short hitting TARGET for +\$72, falsifying 'all shorts faded rips that kept running').

§30 passive_mgc_hours: the shorts are fading rips that run it over

This gold sim lost real money this week, and it lost all of it on one side of the book. On the honest tick-repriced fills it bled **-\$542.2 over 63 trades at 46% win** — while its perfect-fill ceiling said it "should" have made +\$255.9. The whole hole is the short side. Kill the shorts and the sim is flat-to-green.

-\$542.2

REAL P&L, 63 TRADES

-\$555.7

SHORT SIDE (33
TRADES, 32% WIN)

+\$13.5

LONG SIDE (30 TRADES,
66% WIN)

+\$255.9

PERFECT-FILL CEILING IT NEVER SEES

§30.1 1. Did it make money — no, and it's a one-sided wound

I re-ran the join myself on `sr.real_pnl`, not the sim's optimistic `pnl_usd`. It's -\$542.2. Split it by direction and the story ends: the 30 longs made **+\$13.5** and won two out of three, the 33 shorts lost **-\$555.7** and won barely one in three. The longs are already fine. Every dollar of the bleed is short trades.

§30.2 2. Why — it fades a rip above VWAP and the rip keeps going

This is a passive mean-reversion book on gold, and its shorts are all counter-trend fades. Every single one of the 32 shorts I could price fired with gold trading above its VWAP — it sees a pop and sells it, expecting the pop to fade. This week the pops didn't fade. **19 of those shorts hit their stop for -\$506.3**. That one number — nineteen stopped shorts — is essentially the entire week's loss. The longs, by contrast, exit clean: 17 of them ride the ratchet to +\$151.5 in winners, and their stops only cost -\$137.9. Same machine, opposite outcome, purely by side.

The condition averages say the same thing a different way. Its winners fire on a slightly fatter, quieter bar (avg `atr_pct` 0.071, `rvol` 0.47); its losers fire on a thinner, busier bar pushing away from VWAP (avg `atr_pct` 0.063, `rvol` 0.62, `vwap_dist` +0.037). In plain words: when it sells into an active push away from the mean, the mean doesn't come back and it gets run over.

FUNNEL COVERAGE


Heads up on the join: this sim stamps entry_sid as a session anchor, not the cycle timestamp, so the literal cycle_ts = entry_sid join returns nothing. I matched each trade to the nearest MGC funnel row at or before its entry instead — **62 of 63 trades** land within two minutes of a funnel snapshot, so the conditions above are real, not guessed.

§30.3 3. The one thing for Monday — stand the shorts down

The cleanest split isn't rvol or a bar-size veto — I checked both and neither separates the book (rvol<0.5 lost -\$332, rvol>=0.5 lost -\$210; fat bars lost too). The axis that separates winners from losers here is **direction**. So the Monday change is blunt and honest: **on MGC this sim takes longs only — veto the short leg entirely.**

BUCKET	TRADES	REAL P&L	WIN%
LONG only (the keep)	30	+\$13.5	66%
SHORT (the veto)	33	-\$555.7	32%
Sim as it traded	63	-\$542.2	46%

In-sample money swing from vetoing shorts: from -\$542.2 to **+\$13.5, a +\$555.7 improvement**. It doesn't make the sim a winner — longs-only is a rounding error above zero — but it stops the bleeding completely, which is the job this week.



Cumulative honest P&L across the week's 63 trades in order. Gold dashes = breakeven. Red is the sim as it traded (ends -\$542). Green is longs-only (hugs the line, ends +\$13.5). The entire gap between the two lines is short trades.

HONEST CAVEAT

It's one week and 33 shorts — thin, in-sample, and gold trended hard against the fade this week. Longs-only doesn't earn, it just stops losing, so this is a bleed-stopper not an edge. But the direction split is clean and the mechanism is plain (counter-trend fades above VWAP got run over), so the veto is safe to run and cheap to reverse.

VETO shorts

longs stay

re-check next trend-down week

GATE IT — ONE FILTER FIXES IT · \$-542 THIS WEEK ON HONEST FILLS.

pullback_loose lost \$541.8 over 165 honest fills and the entire loss is the short side (-\$509 vs longs at break-even), so the Monday move is to fire it long-only. **Skeptic re-ran the Monday change:** it HELD. The +£402/wk headline reproduces exactly and the load-bearing move — stand down on shorts — is directionally REAL (pullback_loose shorts bleed every single day this week, -\$509/84, and -\$803/144 all-data: a fader getting run over), but it is loss-avoidance / "trade less on the bleeding side," not a positive edge, because the residual long-only book is still negative (-\$32.72 this week resting entirely on Jul 9's +\$198, and -\$270/153 all-data — so "longs are break-even" is overstated); the optional fat-bar + \$203.8/31 sharpening is a curve-fit that should be discarded (all profit is one day (Jul 9 +\$293) and 4 MNQ/MES trades, the 27 MGC trades in the sample lose, and it collapses out-of-sample 6.57→2.42 avg with pre-week fat-bar longs losing money).

§31 pullback_loose: the longs are fine, the shorts are the wound

This one bled about half a grand this week. It bought and sold the pullback across MGC, MES and MNQ, took 165 shots, won just under half of them, and finished the honest fills down \$541.8 against a perfect-fill fantasy of \$1,291. The whole loss is on one side of the book.

-\$541.8

HONEST REAL_PNL, 165
TRADES

47.3%

WIN RATE

\$1,291

PERFECT-FILL CEILING
IT NEVER TOUCHED

-\$509

OF THE LOSS IS THE SHORT SIDE ALONE

§31.1 Did it make money? No — and it's all shorts

Score on the tick-repriced fills, not the sim's own optimistic P&L. It lost \$541.8. Split it by direction and the story is almost embarrassingly clean: the longs made \$-32.72 over 81 trades at 50.6% win — call that break-even. The shorts lost \$509.08 over 84 trades at 44% win. The sim isn't broken. Half of it is.

BUCKET	TRADES	REAL P&L	WIN%
LONG only (the keeper)	81	-\$32.72	50.6
SHORT (the wound)	84	-\$509.08	44.0


§31.2 Why — it fades rips that just run it over

Look at how the money leaves. The stop-outs are the entire loss: 54 trades exited on STOP for -\$1,426, and not one of them was green. Everything that got to manage itself made money — the ratchet booked +\$110 over 51 trades at 76% win, the 20-minute session-flat clock booked +\$441, and the 7 that hit target made +\$333. So when a pullback entry is right, the exit spine harvests it fine. When it's wrong, it's a fader standing in front of a move that keeps going, and it eats the full 1-ATR stop.

The shorts are where that goes wrong. A short "pullback" here is a fade of a rip, and in this tape the rip kept running: short stop-outs alone cost -\$732

across 29 trades. The longs take an equally ugly stop bill (-\$694 on 25) but they claw it back through ratchet and session-flat wins; the shorts don't. That's the whole difference between a break-even side and a bleeding one.

By symbol the damage tracks size and vol: MGC did the most trading and the most bleeding at -\$351 over 102, MNQ -\$118 over 32, MES -\$73 over 31. No symbol is a hero, but none is the problem the way the short side is.



Red = cumulative honest P&L of the sim as-is (ends -\$542). **Green** = the same book with shorts removed (ends -\$33, hugs the gold break-even line). The gap between the lines is the short bleed.

§31.3 The one thing for Monday

Stand down on the short side. Fire pullback_loose long-only.

It's the single cleanest lever on the sheet. Killing the shorts removes -\$509 of a -\$542 loss and leaves an 81-trade long book at -\$33 — a break-even you can then sharpen, instead of a hole you're pouring into. That's an in-sample swing of roughly **+\$509** for one veto, on a real 84-trade sample, not a five-trade fluke.

And there's a free sharpening on top. Within the longs, amplitude sorts winners from losers cleanly: longs fired on a fat bar (entry $\text{ATR} \geq 0.045\%$ of price) made +\$203.78 over 31 trades at 68% win, while longs on a thin bar lost -\$236.5 over 50 at 40%. Same fat-bar filter does nothing for the shorts (they lose either way), which is exactly why the direction veto has to come first.

GATE (IN-SAMPLE, THIS WEEK)	TRADES	REAL P&L	WIN%
Sim as-is (both sides)	165	-\$541.8	47.3
Long-only (the Monday move)	81	-\$32.7	50.6
Long-only + fat bar $\geq 0.045\%$ (the core)	31	+\$203.8	67.7

HONEST CAVEAT

This is one week and mostly it says "trade less." The short veto is high-confidence — 84 trades, one direction, a clean mechanism (fading rips in an up-tape gets run over). The fat-bar long core is real but thin at $n=31$; treat +\$204 as a direction to test next week, not a promise. The short veto is the part to arm.

GATE: long-only

veto SHORT side

exit spine is fine — leave it

GATE IT — ONE FILTER FIXES IT · \$-422 THIS WEEK ON HONEST FILLS.

`pullback_quiet` bled -\$422 over 97 honest-fill trades this week because it only fires into dead tape; the Monday move is an $\text{rvol} \geq 0.8$ floor that removes ~\$391 of in-sample bleed. **Skeptic re-ran the Monday change:** it FELL — read the change as a lead, not a rule. Not real: the actual `entry_rvol` maxes at 0.199 so a $\text{rvol} \geq 0.8$ floor vetoes ALL 106 trades (not 75 of 92) — the claim is a mislabeled, curve-fit "trade-less" cut (effective threshold ~0.15, chosen in-sample on this single week) that removes 75-82% of the book, yet even the highest-rvol kept subset still loses (top-17 = -\$62, $\text{rvol} \geq 0.12$ 44-trade set = -\$25 at 39% win), so there is no rvol floor that yields a durable edge and the +£307 is just the bleed avoided by not trading, not alpha (~£0 real improvement).

§32 pullback_quiet: it only trades when nobody's home, and that's the whole problem

This sim lost real money this week: **-\$422 over 97 trades**, winning just 31% of them, against a perfect-fill ceiling of only +\$154.6. So even in the fantasy where every fill is perfect it barely makes a hundred bucks — on the honest tick-repriced fills it bled all week. This is the sim that buys a quiet little pullback and then waits for a bounce that never comes.

-\$422

HONEST_REAL_PNL, 97
TRADES

31%

WIN RATE

+\$154.6


PERFECT-FILL CEILING

-\$523

FROM THE 35 STOP EXITS ALONE

§32.1 The bleed never let up

The equity curve is a slide with no shelf. It gives back a little, then a lot, and finishes at its worst point of the week. There is no regime where it turned the corner — it just leaked from the first trade to the last.



Cumulative honest real_pnl, trades in time order. Gold dashed line is break-even. It ends at its low, –\$422.

§32.2 Why it loses: it fires into dead tape and the stop pays for it

The name is honest — this thing lives in the quiet. The median relative volume at entry was **0.23**, and **75 of the 92 trades** I could match to entry conditions fired with rvol below 0.8. That's a pullback into a tape where nobody's trading. The bounce needs buyers to show up, and they don't, so the price just keeps drifting the wrong way until the stop takes it out.

The exit mix tells the same story bluntly: the **35 STOP exits lost –\$523 between them and won exactly zero of them**. Everything good this sim did — the 7 TARGET hits at +\$111.5, the 18 ratchets near flat — got swamped by that one bucket. When the entry has no follow-through, the stop is just a scheduled donation.

Two more things it's doing wrong: it's worse fat than thin (bars with atr_pct at or above 0.05 bled **–\$287 over 36 trades at ~10% win** — fading a fast bar in a thin book is how you get run over), and the short side is the heavier bleeder (–\$272 over 40 shorts vs –\$150 over 57 longs). Funnel coverage

for this read was 92 of 97 trades matched within 180 seconds, so this is the near-whole sample, not a corner of it.

WHAT IT'S ACTUALLY DOING

Buying a small dip when the tape is empty and hoping for a bounce. In empty tape there's no one to provide the bounce, so the dip keeps going and the stop eats it. Its own premise — "trade the quiet" — is the bug.

§32.3 The one thing for Monday: don't fire unless someone's actually in the tape

The single cleanest cut is a relative-volume floor. Split its own trades on $r_{vol} \geq 0.8$ and the bleed almost entirely lives on one side:

BUCKET	N	REAL_PNL	WIN%
$r_{vol} \geq 0.8$ (someone's home) — KEEP	17	−\$18.9	53%
$r_{vol} < 0.8$ (dead tape) — VETO	75	−\$390.9	27%

Vetoing the dead-tape trades removes **−\$390.9 of in-sample bleed** and takes the win rate from 27% to 53%. That's the gate: **only fire when $r_{vol} \geq 0.8$** and it stops being a money furnace.

BE HONEST ABOUT WHAT THIS GATE IS

This isn't a tune that flips it green. Even gated, the surviving 17 trades are still slightly red at −\$18.9 — call it break-even. And the gate throws away **82% of the sim's trades**, because "quiet" is what it was built to do. So this is really the split saying the same thing the whole section says: the premise is broken. One week, in-sample, and the honest read is "trade 5× less and land near zero." A relative-volume floor stops the bleeding; it does not make this pay.

Verdict: **GATE, on probation to retire** . Put the $r_{vol} \geq 0.8$ floor on it Monday so it stops bleeding, but the honest truth is a pullback sim that only trades empty tape has nothing to buy. If gated-and-near-flat is the best it can do, it comes off the book next week. Removing $\sim \$391/\text{wk}$ of in-sample bleed ($\approx \text{£}307$) is the whole prize here, and it's a prize made of not losing, not of winning.

GATE IT — ONE FILTER FIXES IT · \$-341 THIS WEEK ON HONEST FILLS.

This MCL short bled -\$341 over 295 trades this week because its entire loss lives in the trending regime (-\$407 over 122 trades, 26% win); the Monday move is to veto every short flagged trending and only fire in chop/shock. **Skeptic re-ran the Monday change:** it FELL — read the change as a lead, not a rule. The direction reproduces — trending MCL shorts bled the whole -\$461 while chop/shock netted +\$100 gross — but the claimed "+\$78 green book" does NOT hold: it's a one-week, in-sample, single-threshold cut on the strategy's own net_atr axis (curve-fit), it drops 297 trades to 95 ("trade less"), the +\$100 is carried by just 5 winners (ex-best-5 = -\$152), and at \$1.06/trade it sits below one MCL tick of spread, so after a realistic 1-tick crude cost the KEEP book is ~£4 (breakeven), not green — the honest value is in-sample loss-avoidance, not a durable positive edge.

§33 MCL "Short With Strength" — It Shorts Into A Tape That's Still Trending, And The Trend Runs It Over

This one bled this week. It shorted crude 295 times and lost \$341 on the honest fills, winning barely a third of the time. And here's the sting: on perfect fills it would have made \$958. The gap isn't the strategy being wrong everywhere — it's one specific mistake it makes over and over, and the fix is a single stand-down rule.

-\$341

HONEST_REAL_PNL, 295
TRADES

35%

WIN RATE

+\$958

PERFECT-FILL CEILING
IT LEFT BEHIND

-\$1,447

BLEED BY THE 135 STOP-OUTS ALONE

§33.1 What it's actually doing — shorting strength and getting steamrolled

The name is honest: it sells into a strong move. When the tape is choppy or shocked, that works — it caught +\$42 across 76 chop trades and +\$36 across 15 shock trades. But crude spent most of this week trending, and that's where the whole loss lives. In the trending bucket it took 122 trades, won only 26% of them, and lost **\$407**. It fades a rip that just keeps running, the stop hits, and it does it again an hour later.

You can see it in the exits. The 135 stop-outs never won once and cost \$1,447. Everything green came from the exits that let a good short breathe — CHANDELIER made +\$537, RATCHET +\$419, the 20-minute flat-clock +\$159. So the exit spine is fine. The problem is upstream: it's opening shorts into up-trends that were never going to pay.

The entry conditions back this up. On the 213 trades I could match to a live signal-funnel row, the losers came in on higher relative volume (rvol 0.83 vs 0.61 on winners) and further below VWAP (−0.15% vs −0.09%) — i.e. it's chasing a move that's already extended and still accelerating, then shorting the front of it.

FUNNEL COVERAGE

The prescribed `entry_sid` join returns zero rows — `entry_sid` on these trades is a session id, not a cycle timestamp. I matched each trade to


the nearest MCL funnel cycle within 10s instead: 213 of 295 trades (72%) carry conditions. The regime split below is on those 213.

§33.2 The one thing for Monday — stand down when the tape is trending

One filter separates the winners from the losers cleanly, and it's the session regime already sitting in the funnel. Drop every short flagged **trending** and keep the rest, and the book flips from bleeding to green in-sample:

BUCKET (SESSION_REGIME)	TRADES	HONEST REAL_PNL	WIN%	VERDICT
trending — shorting into the run	122	-\$406.6	26%	the entire loss
chop	76	+\$41.9	34%	keep
shock	15	+\$36.2	40%	keep
KEEP total (chop + shock)	91	+\$78.1	35%	green

Plainly: **only let this sim fire when the regime is chop or shock — never in a trend.** On the matched sample that turns a -\$329 book into +\$78, an in-sample swing of about **+\$407 (~£321/wk)**. Adding a second rail — only short when price is near or above VWAP, not already far below it — nudges win% to 39% on 67 trades but doesn't add money, so keep the rule to the one clean lever.



Cumulative honest P&L across the week. **Red** = as-is, all 295 shorts, sinking to -\$415 at the low. **Green** = the same sim with trending shorts vetoed, drifting to +\$78. Gold dashes = breakeven.

HONEST CAVEATS

One week, in-sample, 72% funnel coverage. The regime label is noisy — memory has us treating minute-scale regime calls as unreliable, so this is a filter on a label, not a prediction. And a single-week filter usually just means "trade less": here it means trade 91 times instead of 295. But the separation is real and it's the same story the exits tell — this sim only pays when it's fading noise, not when it's fighting a trend.

GATE: fire only in chop/shock regime

veto every trending short

RETIRE — NO HONEST FIX · \$-312 THIS WEEK ON HONEST FILLS.

The MBT euphoria-fade shorted 50 bitcoin rips and got run over, bleeding \$-312 at a 6% win rate with a perfect-fill ceiling that still loses \$-240 — so retire it, no honest gate saves a negative-edge signal.

§34 The MBT Euphoria-Fade: Shorting Into A Freight Train

This sim shorts bitcoin when it looks euphoric, betting the rip is exhausted. This week the rip wasn't exhausted. It bled **\$-312 over 50 trades** and won just **3 of them**. And here's the part that kills any hope of a tune: even with perfect fills it still loses **\$-240**. This isn't a slippage problem. The signal itself is wrong.

\$-312

HONEST REAL_PNL, 50
TRADES

6%

WIN RATE (3 OF 50)

\$-240


PERFECT-FILL CEILING
— STILL RED

41 / 50

STOPPED OUT, EVERY ONE A LOSER

§34.1 1. The money — I re-ran it myself

Joined the sim's trades to the honest tick-repriced fills for the week from Monday 6 July. Fifty shorts, all of them fading a move up in MBT. Total honest damage **\$-312.00** against a perfect-fill fantasy of **\$-239.75**. So real fills cost an extra \$72 on top of a signal that was already losing. When the fantasy version of your strategy loses money, the fill isn't the problem.



Cumulative honest P&L, trade by trade. Gold dashed line is breakeven. It never comes back up.

§34.2 2. Why — it fades a rip that just runs it over

The behaviour is simple and brutal. The sim sees MBT spike, calls it euphoria, and sells short. Then the spike keeps going. **41 of the 50 trades hit their stop**, and every single stopped trade was a loser, for **\$-306.5** — that's basically the whole week's bleed right there. The only trades that didn't get run over were the 9 that survived to the flat-clock cut, and those netted just **\$-5.5** with the sim's only 3 winners in them. So the picture is: when the fade isn't immediately steamrolled it mean-reverts fine by the clock, but 8 times out of 10 the rip is a real trend and it plows straight through the stop.

The holds tell the same story — most losers are dead inside 60-72 seconds. It's short, it's wrong, it's stopped, over and over. The few funnel rows I could see (more on that below) show it firing into "**trending**" tape, which is exactly the reversion-fades-a-trend failure we already booked on the reversion book: it never bounds how strong the move it's fading is, so it sells into grinds that don't stop.

THE HONEST CAVEAT

Only **3 of 50 trades** have a matching entry-conditions row in the signal funnel — 6% coverage. That means I cannot honestly hand you an rvol or atr_pct or book-imbalance gate for this sim this week; there simply aren't enough rows to split on. Anyone who shows you a microstructure filter on this book right now is fitting three data points.

The capture gap is the fix to chase separately; it doesn't rescue the money.

§34.3 3. The one thing for Monday — there isn't a gate, retire it

I went looking for one green bucket to gate into. There is none. Split by session, every window is red:

SESSION BUCKET	N	HONEST P&L	WINS	\$/TRADE
EU 07-11 UTC (the bulk)	40	\$-244.5	2	-6.1
US 13-20 UTC	3	\$-45.0	0	-15.0
Asia / other	7	\$-22.5	1	-3.2
All trades	50	\$-312.0	3	-6.2

US hours are actually the worst per trade, so "US-only" makes it lose faster, not slower. Asia is the least-bad but that's 7 trades and one win — noise, and still red. There is no hour, no side (it's 100% short by design), and no exit variant (every trade is the same desk variant at target 8R) that turns this green. The exit isn't the villain either: the flat-clock trades were nearly breakeven, so a better exit just gets you to zero, not to profit.

VERDICT

RETIRE The perfect-fill ceiling is already **\$-240**. That is the whole argument in one number: fix the fills, fix the exit, pick the best hour — you still lose. Fading MBT euphoria is a negative-edge idea in a week bitcoin trended, and this sim has no rvol/atr floor to keep it out of the grinds it keeps selling into. Retiring it stops **\$-312/wk** of pure bleed. If we ever revive the euphoria-fade, it comes back only with an

amplitude floor and a trend-slope veto proven on real funnel coverage
— not before.

One week, in-sample, $n=50$. Money scored on honest tick-repriced `real_pnl` throughout; the sim's own `pnl_usd` was ignored. Every number above came from a join I re-ran against `data/shadow_real.db` this session.

GATE IT — ONE FILTER FIXES IT · \$-267 THIS WEEK ON HONEST FILLS.

Confirmed a \$266.9 loss over 36 trades on honest fills; the fix is to veto the short side, which was 78% of the bleed and loses even at its perfect-fill ceiling. **Skeptic re-ran the Monday change:** it FELL — read the change as a lead, not a rule. The arithmetic reproduces exactly and net-of-fees (22 shorts = -\$207.79 real / -\$46.52 ceiling $\approx$ £164), but the refinement does NOT hold as a forward edge: it is a thin-n, 100%-in-sample "trade less" call on a single week of 22 shorts with zero out-of-sample, where 69% of the loss sits in 3 trades and 30% (\$61.50) in soon-to-be-dropped MBT, and the "even perfect fills can't save them" ceiling argument inverts at the symbol level (the highest-count shorts, MES + \$53.72 and MYM + \$27.79 ceiling, would be saved by good fills) — so the £164 is just this week's removed realized loss, not a durable improvement.

§35 Strength-Ceiling Pullback: the shorts are the disease, thin fills are the fever

This one bled. On the honest tick-repriced fills, **tw_strength_ceiling_pb** lost **\$266.9 over 36 trades** this week, winning just 42%. Its perfect-fill ceiling was only **-\$30.4** — so even in fantasy-land where every fill is clean, this signal barely broke even. That gap tells the whole story, and it splits cleanly down the middle: the shorts are a broken signal, the longs are a fill problem.

-\$266.9

REAL P&L, 36 TRADES

42%

WIN RATE

-\$30.4

PERFECT-FILL CEILING

-\$236

EATEN BY FILLS

§35.1 Why it bled: two different diseases wearing one coat


Split it by side and the picture snaps into focus. The **22 shorts lost \$207.8** — and their ceiling was **-\$46.5**, so even with perfect fills the shorts lose money. That is a signal problem: this thing shorts a pullback into strength and the strength just keeps running it over. All three MNQ shorts lost (0-for-3, -\$101.9). The single worst trade was an MBT short that got carried to the close for -\$62.

The **14 longs are a different animal**. They lost \$59.1 on the honest fills — but their ceiling was **+\$16.1**. The long signal was actually fine, a small winner on paper. Then honest fills ate **\$75** and turned it red. That is the "nobody's home" tax — on the micros this book trades, the fill you get in a thin tape is the one you didn't want. Look at the MES short that should have banked \$76.5 at the ceiling and came home at **+\$0.9** — the whole move, gone to fills.

The exit mix confirms it: the **STOP bucket is the graveyard** — 12 stop-outs, 17% wins, **-\$183.2**. Those are entries that were simply wrong-way; the exit did its job cutting them. The RATCHET and FLAT_CLOCK exits were both slightly green (+\$7.3 and +\$6.6), because those are the trades that went the right way from the start. This sim doesn't have an exit problem. It has an entry problem, and the entry problem is mostly the shorts.

FILLS VS SIGNAL, BY SIDE

Shorts: real -\$207.8 vs ceiling -\$46.5 — **the signal itself loses.**
 Longs: real -\$59.1 vs ceiling +\$16.1 — **the signal is fine, the fills kill it.** One coat, two diseases.



§35.2 The one thing for Monday: stand down the shorts

There is no clever microstructure gate here — the signal funnel only caught 2 of the 36 trades this week, so rvol / book / session vetoes have nothing to stand on. The honest lever is the one that shows up on every trade: **side**. Veto the short side of this sim and you remove 78% of the bleed.

BUCKET	N	REAL P&L	WIN%	CEILING
SHORTS (veto these)	22	-\$207.8	45%	-\$46.5
LONGS (keep, on a leash)	14	-\$59.1	36%	+\$16.1
In-sample swing from the veto	—	+\$207.8	—	—

Be straight about what this does and doesn't do. Killing the shorts is a **+\$207.8 in-sample swing**, and it's the right call — even perfect fills couldn't save them. But it does not flip the sim green: the longs still come home at **-\$59.1**, because their disease is fills, not direction. And this is one week, 36 trades — a side veto on a single week mostly just means "trade less," and here it means "trade a lot less."

So the verdict is **KILL the short side** **LONGS on watch**. The long signal has a real, if tiny, paper edge (+\$16 ceiling) that a fatter, less-thin tape might let

through — but until it can beat its own fills it stays observe-only. If the longs are still negative next week, the whole sim goes in the departure lounge. A signal whose best-case perfect-fill week is -\$30 is not one to fall in love with.

HONESTY NOTE

Funnel coverage this week was 2/36, so no rvol/atp_pct/book split is trustworthy — the only gate I'll stand behind is side. The atr_pct cut (fat ≥ 0.09 vs thin) separates nothing: both buckets bleed (-\$83.2 vs -\$183.6). Dropping MBT alone doesn't fix it either — the rest still loses \$188. The shorts are the clean lever; everything else is noise on one week.

GATE IT — ONE FILTER FIXES IT · \$-237 THIS WEEK ON HONEST FILLS.

This MNQ pullback sim bled \$237 over 13 trades on honest fills against a paper-thin +\$43.61 ceiling, and every one of its 3 winners fired on a bar with real range while all 9 quiet-bar trades lost — so the Monday move is an atr_pct amplitude floor that vetoes the thin tape. **Skeptic re-ran the Monday change:** it FELL — read the change as a lead, not a rule. The arithmetic reproduces exactly (veto 9 trades 0-for-9 for -\$208.58, keep 4 for -\$28.48), but it is a curve-fit, thin-n, trade-less artifact — the 0.042 threshold is fit into the single in-sample gap, the surviving 4-trade bucket is still net-negative (-\$28.48 $\approx$ -£22) at every threshold, the fattest bar is the worst loser which falsifies the mechanism, and the claimed +£162.7 is nothing but avoided losses on a -EV book, not a real edge.

§36 tw_mnq_pb_expand_veto — the quiet bars are killing it

This MNQ pullback sim bled **-\$237.06** over 13 trades this week on the honest, tick-repriced fills — a 23% win rate (3 of 13). The cruel part: its perfect-fill ceiling was only **+\$43.61**. So even if every fill had gone its way it barely made lunch money, and in the real world the fills dragged a paper-flat week into a \$237 hole.

-\$237.06

HONEST REAL_PNL, 13
TRADES

23%

WIN RATE (3 / 13)

+\$43.61

PERFECT-FILL CEILING
— THIN EVEN AT BEST

-\$280.7

HOW FAR REAL FILLS FELL BELOW THE CEILING

§36.1 Why it bled — it keeps firing when nobody's home


The damage is almost entirely stops. Six **STOP** exits cost **−\$223.22** — that's basically the whole loss. Three **SESSION_FLAT** time-outs shed another **−\$31.36**. The only green came from four small **RATCHET** exits worth **+\$17.52** between them. So the shape of the week is: a handful of tiny locked wins, buried under a pile of full-width stop-outs.

The direction it took didn't matter — both sides lost. Its five **longs went 0-for-5** for **−\$101.60**, its eight shorts went **3-for-8** for **−\$135.46**. You can't fix this by picking a side.

What does separate the winners from the losers is how much the bar was actually moving when it fired. Every single one of its winners fired on a bar with real range — **atr_pct 0.043% to 0.045%**. Every trade it took on a quiet bar (atr_pct under 0.042%) lost. That's the whole tell: on a thin, sleepy tape the pullback entry gets run into a stop, and the fill you get is the one you didn't want. Note: fut_signal_funnel had zero matching rows for this sim (0/13 — its entry_sid is a day-bucket, not a cycle stamp), so bar range here is measured straight off each trade's own entry_atr_abs ÷ entry_price.

THE SPLIT

Sort its 13 trades by bar range. Above ~0.042% atr_pct: **4 trades, 3 winners, −\$28.48**. Below it: **9 trades, 0 winners, −\$208.58**. The quiet-bar bucket is a clean 0-for-9 wipeout.



Each bar = one trade's honest real_pnl, ordered thin bar (left) → fat bar (right). Gold dashed line = the 0.042% atr_pct gate. Left of it: nine red bars, no green. Right of it: the three winners live here (one fat bar still stopped out, -\$46).

§36.2 The condition split — bar range is the only thing that sorts them

BUCKET	N	WINNERS	HONEST REAL_PNL	AVG / TRADE
Fat bar — atr_pct ≥ 0.042%	4	3	-\$28.48	-\$7.12
Thin bar — atr_pct < 0.042%	9	0	-\$208.58	-\$23.18
All trades (as-is)	13	3	-\$237.06	-\$18.24

§36.3 The one thing for Monday

GATE: stand down when the bar is quiet Only let this sim fire when atr_pct ≥ ~0.042%. In-sample that vetoes the nine quiet-bar trades and removes — **\$208.58** of pure bleed — the week goes from -\$237 to -\$28. That's the same amplitude-floor lesson we've already seen on the momentum shadows: thin tape, stand down.

Now the honest bit, because this is one week and in-sample. The survivor bucket is only **four trades**, and it's still red (-\$28.48) because one fat-bar short (-\$46, a STOP) sits on the good side of the line too. The threshold is fit to this week's data. And remember the ceiling was only **+\$43.61** — even with perfect fills and every quiet bar removed, there's barely an edge here to defend. So the gate stops the bleeding, but it does not turn this into a money-maker.

VERDICT

Gate it hard, and put it on a short leash to retire. Arm the atr_pct ≥ 0.042% amplitude floor Monday — it kills the 0-for-9 quiet-bar disaster and saves ~\$209 in-sample. But with a paper-thin +\$43.61

ceiling and a survivor bucket of just four trades that's still red, this sim is one bad fat-bar week away from a clean retire. Watch it for a fortnight behind the floor; if the fat bucket doesn't turn green with real n, pull it.

GATE IT — ONE FILTER FIXES IT · \$-211 THIS WEEK ON HONEST FILLS.

On honest fills it lost \$211 over 29 trades (45% win) while its ceiling was +\$118 — the bleed is one bucket, shorts fired below VWAP (13 trades, -\$216, 31% win); vetoing that lifts the week to about break-even. **Skeptic re-ran the Monday change:** it FELL — read the change as a lead, not a rule. The arithmetic reproduces exactly and cost-honestly (vetoing below-VWAP shorts removes 13 trades worth -\$215.56, lifting the week from -\$210.98 to +\$4.58, a ~£168 swing that matches the claimed ~£170), but it does NOT hold as a forward edge: it is a 13-trade, single-week, in-sample filter that merely relabels this week's full-ATR stop-outs, reaches only break-even (the 12 matched-kept trades are actually -\$9.44 and it crosses zero solely because 4 unmatchable early shorts are kept by default at +\$14), discards 4 winning below-VWAP shorts (+\$31) while keeping 5 losing longs (-\$78), and hinges on a knife-edge zero threshold over `vwap_dist` values clustered within $\pm 0.03\%$ via a fragile nearest-cycle join — a thin-n curve-fit that stops the bleed to flat, not a real edge.

§37 tw_thin_drift_standdown: it keeps shorting into the drift and gets run over

This one bled. On the honest tick-repriced fills it lost **\$211 over 29 trades** this week at 45% win, while its perfect-fill ceiling said it "should" have made \$118. So it's not a fill problem — the ceiling is green and the truth is red. It's an entry problem, and it's almost entirely one bucket: the short side.

-\$211

HONEST REAL_PNL, 29
TRADES

45%

WIN RATE

+\$118

PERFECT-FILL CEILING
(FANTASY)

-\$216

THE SHORTS-BELOW-VWAP BUCKET ALONE

§37.1 Where the money went

Split it by symbol and the damage is lopsided: MGC lost **\$83 over 21 trades** and MNQ lost **\$128 over just 8 trades** — MNQ is bleeding roughly four times as fast per trade. And it's all in the stops. The RATCHET exits actually made money (14 trades, +\$80) and the 20-minute flat-clock exits were fine (+\$13), but the STOP exits gave back **\$304 across 13 trades**. When this sim is right it trails a small winner home; when it's wrong it sits there until the stop takes a full-size loss.

§37.2 Why it wins vs why it loses

I matched each trade to its entry conditions in the funnel. Quick honesty note on the data: the documented `entry_sid = cycle_ts` join returns zero rows here because `entry_sid` is a rounded session stamp, not the cycle time — so I matched each trade to the nearest funnel row instead. That lands **25 of 29 trades within 10 seconds**; the other 4 are early trades with no nearby funnel row (they net a harmless +\$14).

On those 25, the winners and losers look almost identical on `rvol` and bar size — this is not a "thin tape / fat bar" story despite the sim's name. The one axis that separates them cleanly is **which side of VWAP price was on, and which way the sim was leaning**. The longs are basically flat: 10 trades, -\$10, 50% win. The shorts are the whole problem: **15 trades, -\$215, 33% win**. And inside the shorts, it's the ones fired below VWAP —

selling weakness into a market already trading under its own average — that get run over: **13 trades, -\$216, 31% win**. That's the sim doing the exact thing its name says it should stand down from: leaning short into the drift and eating the stop.


(The "wide spread bleeds" reading is real but it's a mirage — spread is symbol-scaled. MGC spreads sit at 0.1-0.3, MNQ at 0.5-0.75, so "spread > 0.30" is just spelling "MNQ" a different way. The honest axis is VWAP side, not spread.)

§37.3 The one thing for Monday

Stand down on shorts when price is below VWAP. That single veto removes the only toxic bucket and takes the sim from bleeding to roughly flat.

BUCKET	TRADES	REAL P&L	WIN%
SHORT, price below VWAP ($\text{vwap_dist} \leq 0$) — VETO THIS	13	-\$216	31%
Everything else (longs + shorts above VWAP) — keep	12	-\$9	~50%

In-sample the swing is **+\$216**: killing that one bucket lifts the week from -\$211 to about +\$5. The chart below shows it — the navy line is the sim as it actually traded (a straight slide into -\$211), the gold line is the same week with shorts-below-VWAP vetoed (it just walks along the flat line instead of falling off it).



HONEST CAVEAT

This is one week and the toxic bucket is 13 trades — enough to be real, not enough to bet the desk on. And notice what the veto actually buys you: it stops the bleed, it doesn't mint money. The "keep" bucket is only flat (-\$9 on 12 trades), not green. So this is a **defensive gate**, not a green light — it turns a losing sim into a break-even one that we can then watch to see if the long side or the above-VWAP shorts have any real edge, or whether the honest call next week is just to retire it.

GATE: veto SHORT when $vwap_dist_pct \leq 0$

MNQ is 4× the per-trade bleed — watch for retire

long side already flat, leave it

RETIRE — NO HONEST FIX · \$-198 THIS WEEK ON HONEST FILLS.

This MGC pullback sim went 0-for-7 and lost \$197.84 on honest fills with its two best trades being pure fill fantasy, so the Monday move is to retire it — no condition split turns any bucket green.

§38 The MGC pullback that went 0-for-7

This one bled every single trade. Seven entries, seven losers, minus \$197.84 on honest fills against a perfect-fill ceiling of plus \$0.83. There is no gate that saves it this week. Retire it.

-\$197.84

REAL P&L, 7 TRADES

0%

WIN RATE (0 OF 7)

+\$0.83


PERFECT-FILL CEILING

-\$28.26

AVG LOSS / TRADE

§38.1 1) The money — it just bled, straight down

I re-ran the join myself on `sr.real_pnl`, the tick-repriced truth. Every trade is red: -\$32.12, -\$26.12, -\$20.12, -\$23.12, -\$25.12, -\$44.12, -\$27.12. That sums to -\$197.84 over 7 trades, a clean 0% win rate. The sim's own fantasy ceiling for the same week was basically flat at +\$0.83 — so even in the perfect-fill dream this thing made nothing. On real fills it lost two hundred dollars.



every bar red Seven trades in entry order. There is no green to point at.

§38.2 2) Why — bad entries, and two "winners" that were pure fill fantasy

Five of the seven died at the hard STOP: the pullback it bought (or shorted) simply kept going against it. On those five the honest fill and the fantasy ceiling agree — trade 12538 shows -\$32.12 real vs -\$30.99 ceiling, a dollar apart. That's not a slippage problem, that's a wrong-way entry. The strategy is a trend-continuation pullback, and this week the pullbacks weren't pullbacks — they were reversals.

The other two exited on the RATCHET and are worth staring at, because they're the memo "shadow sim understates losses" made flesh. Trade 17139's ceiling says +\$110.38; the tick-repriced fill says -\$27.12. That's a \$137 gap between the dream and the money. Trade 15997: ceiling +\$11.00, real -\$25.12. In the sim's fantasy those two made \$121 together; on real MGC fills they lost \$52. Gold ticks that never actually printed at your order.

Conditions at entry (nearest funnel row within 2s of each fire — 7/7 covered): it fired into thin, low-participation tape. Five of seven had rvol under 1.0, three of them under 0.5 (0.32, 0.44, 0.50). A "strong pullback only" strategy is supposed to demand a strong move to pull back from — instead it took dead-tape entries. And the two fattest bars took the two worst beatings.

ID	SIDE	EXIT	REAL \$	CEILING \$	ATR_PCT	RVOL	REGIME
12538	SHORT	STOP	-32.12	-30.99	0.069	0.44	trending

12780	LONG	STOP	-26.12	-22.63	0.050	1.09	chop
14154	SHORT	STOP	-20.12	-17.15	0.037	1.99	chop
14416	LONG	STOP	-23.12	-21.07	0.046	0.55	chop
15997	LONG	RATCHET	-25.12	+11.00	0.040	0.97	chop
16059	LONG	STOP	-44.12	-28.71	0.065	0.50	trending
17139	LONG	RATCHET	-27.12	+110.38	0.038	0.32	trending

§38.3 3) The one thing for Monday — there isn't one. Retire it.

I looked for the single split that saves it. The cleanest one is a fat-bar veto: the two entries on bars with `atr_pct >= 0.06` (trades 12538 and 16059) averaged -\$38.12 each and gave up \$76.24 between them; the other five averaged -\$24.32. So standing down when the bar is fat does save the worst of it.

BUCKET	N	REAL \$	AVG \$	WIN%
fat bar (atr_pct >= 0.06)	2	-76.24	-38.12	0%
the rest (atr_pct < 0.06)	5	-121.60	-24.32	0%

WHY THE GATE ISN'T A FIX

The fat-bar veto skips \$76.24 of bleed, but the "good" bucket it leaves behind is still 0-for-5 and still loses \$121.60. Both buckets are 0% win. Vetoing the fat bars doesn't uncover an edge — it just makes a losing strategy lose slightly less slowly. That's not a gate, that's trading less of a bad thing.

Same story on every other axis: chop entries lost -\$94.68 over 4, trending lost -\$103.36 over 3 — both red. High rvol (≥ 1.0) lost -\$46.24 over 2, low rvol lost -\$151.60 over 5 — both red. There is no bucket, on any condition,

where this sim made money this week. Seven trades is thin and it's one week, so I won't over-claim a "structural flaw" from $n=7$ — but a strategy that goes 0-for-7 with its two best trades being fill mirages has earned a rest, not a tweak.

RETIRE Stand it down. If it comes back, it comes back the honest way: revive only with a real rvol floor (a "strong pullback" needs strong participation, not 0.32 rvol) and prove it on more than one week before it fires a dollar. Nothing to arm Monday.

GATE IT — ONE FILTER FIXES IT · \$-184 THIS WEEK ON HONEST FILLS.

dip_loose lost \$184.5 over 131 trades on honest fills, but the loss is entirely MES/MGC — MNQ made +\$242.6 at 65% wins — so the Monday move is to gate it to MNQ only. **Skeptic re-ran the Monday change:** it FELL — read the change as a lead, not a rule. The numbers reproduce to the dollar on honest real_pnl (-\$184.5 to +\$242.6, ~£335 in-sample swing) and it's cost-aware not cost-blind, but it's essentially a "trade-less" move whose entire £335 is avoided MES/MGC bleed while the MNQ book it keeps and calls "a real edge" is a single-trade mirage (+\$242.6 collapses to -\$6.3 over the other 48 trades once the lone \$248.88 TARGET is removed, avg win \$19.7 < avg loss \$22.8), so on n=49 of one in-sample week the durable earned edge is roughly £0.

§39 dip_loose: three markets, one edge, and it's all in MNQ

This dip-buyer took 131 shots this week and lost money on the honest fills — down \$184.5 against a perfect-fill fantasy of \$1,200. But that headline hides the real story: it's a good sim wearing two bad markets. On MNQ it made money and beat nobody's-fault trouble. On gold and the S&P micro it bled. The fix isn't a knob — it's dropping two symbols.

-\$184.5

HONEST REAL_PNL, 131
TRADES

53%

WIN RATE

\$1,200

PERFECT-FILL CEILING
IT LEFT BEHIND

§39.1 1) The money — it bled, but only in two rooms

I re-ran the join on shadow_real.real_pnl (never the sim's optimistic pnl) for every dip_loose fill with an entry on or after 6 July. It comes to **-\$184.5 across 131 trades at 53% wins**. That's the truth. Now split it by symbol and the loss stops being a mystery:

SYMBOL	TRADES	HONEST \$	WIN%	CEILING \$	WHAT'S HAPPENING
MNQ	49	+242.6	65%	664.6	Real edge — buys the dip, dip bounces
MES	36	-134.6	47%	31.5	No edge even in fantasy
MGC	46	-292.5	43%	503.9	Fantasy edge, but fills murder it

MNQ carries the whole sim: **+\$242.6 over 49 trades at 65% wins**. The other two together give back **-\$427.1**. That's not variance — those are two different failures, and neither one is fixable by tuning.

§39.2 2) Why — gold's fills eat the edge, MES has no edge to eat

Look at the ceiling column, because it tells you where the edge even exists. On **MES the perfect-fill ceiling is only \$31.5** — even if you filled every trade at the ideal price, this sim barely makes lunch money on the S&P micro. There is no edge there to rescue. On **MGC the ceiling is a healthy \$503.9**, so the setup is right — gold dips and bounces — but the honest fills turn it into **-\$292.5**. That's the widest fantasy-to-reality gap on the desk: the dip you want to buy in gold is the one nobody's selling you at your price, and the fill you actually get is the bad one.

The exit mix confirms it's a fill/stop problem, not a target problem. When dip_loose is right it books nicely — **TARGET exits made +\$472 over 6 trades**, session-flats and ratchets another **+\$361.5** combined. But the **STOP exits bled -\$1,004.9 over 43 trades**, and that damage is


concentrated where the fills are worst: **MGC stops alone lost -\$501.3**, MES stops -\$185.6. On MNQ the same stop only cost -\$318.1 and the winners more than paid for it.

IN PLAIN WORDS

dip_loose buys a dip and waits for the bounce. On MNQ the bounce shows up often enough and the fills are clean enough that it pays. On gold the setup is fine but the tape is thin at the price you need — you buy the dip and the dip keeps going, then the stop fills worse than modelled. On MES there was never enough edge to begin with.

I checked the usual within-symbol gates — the funnel matched 117 of 131 trades on nearest cycle timestamp (the documented entry_sid join is a session bucket, so I matched on entry_ts to ~1 second). Nothing clean separated winners from losers inside a symbol: fat bars vs thin were both negative (-\$86.0 vs -\$122.2), and book imbalance barely moved the needle. The one axis that actually splits this sim is which contract it's trading.

§39.3 Honest cumulative P&L by symbol — MNQ climbs, gold sinks



■ MNQ +\$242.6 ■ MGC -\$292.5 ■ MES -\$134.6 - - - break-even.
Each line is that symbol's running honest P&L over its own trade sequence.

§39.4 3) The one thing for Monday — trade MNQ only

There is no gate on rvol or slope or the book that saves gold and MES, because the problem isn't the entry condition — it's that the edge lives in one contract. The single change is a symbol gate:

GATE Only fire dip_loose on MNQ. Stand down on MES and MGC.

BUCKET	TRADES	HONEST \$	WIN%
MNQ only (keep)	49	+242.6	65%
MES + MGC (drop)	82	-427.1	45%

In-sample this week, cutting the two losing contracts removes **-\$427.1 of bleed and flips the sim from -\$184.5 to +\$242.6**. That's the whole swing, and it comes from the biggest, cleanest split available — n=49 kept, n=82 dropped, no thin-tail cherry-picking.

HONEST CAVEAT

It's one week and it's in-sample. The MNQ edge is real (65% wins, beat break-even on honest fills), but 49 trades is a single week's worth — watch it hold before sizing anything. The gate itself is safe though: MES has no ceiling edge to give up (\$31.5 fantasy), and MGC's honest fills have lost money all week, so dropping them can't cost you a real edge. This is the good kind of "trade less" — you cut the two rooms where you were paying to play and keep the one that pays you.

RETIRE — NO HONEST FIX · \$-142 THIS WEEK ON HONEST FILLS.

Scored the MCL down-flush sim's honest fills at -\$142 over 33 trades (21% win) with a negative -\$89 ceiling; the bleed is dead-tape entries, so the Monday move is retire it (or, if kept, hard-gate rvol $\geq$ 0.3).

§40 MCL down-flush sim: it fires into dead tape and pays for it

This one bled again. On the honest tick-repriced fills **tw_mcl_no_long_downflush** lost **\$142 across 33 trades this week, winning just 21% of them** — and the tell is that even its perfect-fill fantasy ceiling was **-\$89**. When a sim loses money even on fills it never actually got, the problem isn't slippage. The entries are wrong.

-\$142

REAL P&L, 33 TRADES

21%

WIN RATE

-\$89

PERFECT-FILL CEILING
(STILL RED)

-\$103


OF THE BLEED IS DEAD-TAPE TRADES

§40.1 What it's actually doing

The name says "no long down-flush" but it traded both ways this week — 12 longs for -\$68 and 21 shorts for -\$74. So it's not a side problem; both directions leak. The exit mix tells you the shape of the pain: the money didn't drain through the time-based cuts, it drained through **hard stops**. Five long stops gave back -\$88 and five short stops another -\$48 — that's the whole loss and then some. Everything else (ratchets, targets, session-

flats) roughly washed. In plain words: **it gets in, the price goes straight against it, and it eats the stop.** That's an entry that's early or pointed the wrong way, not an exit that's badly tuned.

Now line the trades up against the tape they fired into (I matched every trade to the nearest MCL signal-funnel snapshot; 23 of 33 landed within 10 seconds, a handful matched a staler row minutes away, so read the thin buckets loosely). One number sorts this sim cold: **relative volume**. When nobody's trading crude, this thing loses. The trades it took in dead tape — rvol under 0.3 — were **21 trades, a −\$103 hole, and it won only 10% of them**. That's 64% of its trades and roughly three-quarters of the loss, all fired when there was no volume behind the move to carry it anywhere.



The green bar is real but don't fall for it: that whole +\$26 in the 0.5–0.6 bucket is a **single lucky short**. So the honest read isn't "rvol makes it win" — it's **"no rvol makes it lose, reliably."**

§40.2 The one thing for Monday

Stand it down when the tape is dead. Gate the entry on **rvol ≥ 0.3** and don't fire below it.

THE SPLIT (HONEST FILLS)

BUCKET	TRADES	REAL P&L	WIN%
rvol < 0.3 — dead tape (KILL)	21	−\$102.52	10%

rvol $\geq$ 0.3 — some volume (KEEP)

12

−\$39.44

42%

Cutting the dead-tape trades saves **\$102 in-sample** and doubles the win rate from 10% to 42%. But be straight about two things. First, this is mostly **"trade less"** — the gate deletes 21 of 33 trades, nearly two-thirds of everything it did. Second, and this is the one that matters: **even the kept bucket still bleeds −\$39, and the sim loses on its perfect-fill ceiling too**. A rvol gate turns a bad week into a less-bad week; it does not turn it green, and no fill or exit tweak will, because the entries are −EV even in fantasy.

So the clean verdict is **RETIRE**. This sim doesn't have an edge to gate toward — it has a bleed to slow down. If you want to keep it breathing on the shadow desk to watch for a regime where it works, the **rvol $\geq$ 0.3** stand-down is the mercy rule that cuts most of the pain. But on one week, negative ceiling, and a best-case bucket that's still red, I wouldn't spend more on it. Pull it.

GATE IT — ONE FILTER FIXES IT · \$-117 THIS WEEK ON HONEST FILLS.

This fade book lost \$117 over 14 trades on honest fills, and its whole loss is 6 stop-outs concentrated in fat, fast bars — the Monday move is to veto entries when atr_pct is high so it stops fading into rips that run it over.

§41 tw_xconfirm_veto — a fader that keeps getting run over

This one bled. On the honest tick-repriced fills it lost **\$117** over 14 trades this week, won just 43% of them, and even its perfect-fill fantasy ceiling was only **+\$78**. So the fills aren't the story — the entries are. It's a fader, and this week the tape kept faded moves running.

-\$117

REAL P&L, 14 TRADES

43%

WIN RATE

+\$78

PERFECT-FILL CEILING

-\$122

FROM 6 STOP-OUTS ALONE

§41.1 What it's actually doing

Eleven of its fourteen trades are shorts — this is a fade / cross-confirm veto book, it leans against moves. When it's right it clips a small winner and rings the register on the ratchet or the flat-clock. When it's wrong, it gets run over and hits a stop. And that is the entire loss: the six **STOP** exits cost **-\$121.7**, while the five ratchets (-\$10) and three flat-clock exits (+\$15)

roughly wash. This sim doesn't bleed a little on everything — it takes a handful of full-size punches in the face.


The punches all land in the same place: **fat, fast bars**. When the bar it fired into was wide ($\text{atr_pct} \geq 0.034$) it lost **-\$89.6 over just 5 trades** — that's - \$17.9 a trade — and that bucket holds both MNQ disasters, a short that lost **-\$46** and a long that lost **-\$34**. Fade a bar that's already moving hard and the bar wins. On thin, quiet bars it still loses, but only **-\$3.1 a trade** — small, survivable, the normal cost of doing business.

One more tell, thin but ugly: all three of its **longs lost**, 0-for-3, -\$48.6. Buying with this logic did nothing but donate. The edge, such as it is, lives entirely on the short side (11 trades, 55% win, still -\$69 because of the two fat tails).

THE TELL

The whole loss is 6 stop-outs (-\$122). The 5 worst-and-biggest of those live in the fat-bar bucket ($\text{atr_pct} \geq 0.034$): -\$89.6 over 5 trades. It fades a rip that just runs it over.

§41.2 Every trade this week, worst to best (fat bars ringed in gold)



The three tallest red bars — -\$46, -\$34, -\$15 — are all gold-ringed fat bars. Kill those and the book is a nuisance, not a disaster.

§41.3 The one thing for Monday: stand down when the bar is fat

Split its own trades by the size of the bar it fired into. It's one clean line.

BUCKET	TRADES	REAL P&L	WIN%	PER TRADE
Fat bar — atr_pct $\geq$ 0.034 (VETO)	5	-\$89.6	40%	-\$17.9
Everything else it keeps (thin + no-funnel)	9	-\$27.6	44%	-\$3.1

The gate: **only fire when atr_pct < 0.034** — stand down when the bar is fat. In-sample that lifts it out of the fat-bar bucket and swings **+\$89.6** for the week, cutting it from 14 trades to 9.

Be honest about what this is and isn't. It stops the bleeding; it does not flip the sim green. The nine trades it keeps still lose -\$27.6 — small, but red. And the veto rests on 5 trades in one week, so most of what this gate really says is "a fader has no business fighting a bar that's already ripping" — which is the right instinct even before the arithmetic. A cousin fix (kill the 0-for-3 longs, -\$48.6) points the same direction and could be stacked, but I won't over-fit two filters onto fourteen trades.

VERDICT

GATE Don't retire it yet, but don't run it naked. Add the atr_pct < 0.034 fat-bar veto — mechanism-grounded, +\$89.6 in-sample — and watch the gated book. If the surviving thin-bar trades stay red for another week, retire it. Even the best single fix leaves this one below the line.

RETIRE — NO HONEST FIX · \$-113 THIS WEEK ON HONEST FILLS.

This MES fade lost \$112.6 over 24 trades on honest fills with a perfect-fill ceiling of only +\$56.5, and every condition bucket stays red, so the Monday move is to retire it.

§42 tw_mes_block_us_open — a fade that gets run over, for a fill it can't get

This one bled. On the honest fills it lost **\$112.6 over 24 trades** this week, winning 46% of them. And here is the number that decides everything: even with **perfect** fills its ceiling was only **+\$56.5**. The best case for this sim is barely above zero, and the real fills drag it firmly into the red.

-\$112.6

REAL P&L, 24 TRADES

+\$56.5

PERFECT-FILL CEILING
(THE FANTASY)

46%

WIN RATE

\$7.0

SLIPPAGE BLED PER TRADE


§42.1 What it's actually doing

It's a fade, and it's short-heavy — 18 of its 24 trades were shorts, and those shorts lost \$76.7. The story is cleanest in the exits. Split the week by how each trade ended and the whole loss lives in one bucket:

THE WHOLE LOSS IS THE STOP-OUTS

9 trades hit the **STOP** and lost **\$127.1** between them — every one a loser. The 9 that trailed out on the **RATCHET** made **+\$15.4** and won 89% of the time. The 6 that timed out flat (SESSION_FLAT) were a wash at -\$1.0. So this sim isn't broken at picking direction — it's fine when left alone. It gets murdered on the trades that go straight against it.

And you can see when it gets murdered. The stop-out trades fired into a live push: they came in with average **rvol 0.43** and **OFI 64**, versus rvol 0.30 / OFI 43 on the ratchet winners. In plain words: it fades a move, the move is actually being driven by real one-sided flow, and the flow runs it over before the trail can lock anything in. It's standing in front of the train.



Cumulative honest P&L across the 24 trades in order. Flat-ish start, then trades 12-17 — the stop-out cluster into hot tape — take it from about -\$8 down to -\$107. That drop is the entire week.

§42.2 The one thing for Monday — and the honest answer is: retire it

The obvious gate is "don't fade a live push": stand down when the tape is hot. It's directionally right — the stop-outs sit on high rvol and high OFI. But when I actually split the trades on it, the gate doesn't save the book:

BUCKET	N	REAL P&L	WIN%	READ
HOT tape — rvol ≥ 0.4 (veto these)	7	-\$66.3	29%	

BUCKET	N	REAL P&L	WIN%	READ
				the run-over trades
Calm tape — rvol < 0.4 (the "keep" bucket)	15	-\$51.8	47%	still red

That's the problem. Veto the hot tape and the survivors still lose \$51.8. I hunted every axis — OFI, book imbalance, spread, VWAP distance, side, session hour, calm-tape shorts only — and **every single bucket came back red**. The best-behaved slice I could find was calm-tape shorts, and even that was -\$19.6 over 9. There is no green corner to gate toward.

WHY NO GATE FIXES THIS

The math is brutal and it's not about entries. The perfect-fill ceiling is **+\$56.5 over 24 trades — about \$2.35 a trade in fantasy-land**. The gap between that ceiling and the honest fills is **\$169, or \$7.0 of slippage per trade**. So the edge is worth about \$2.35 and the friction to trade it costs \$7. The friction is three times the edge. You cannot gate your way out of that — even if every fill were perfect this thing barely clears zero, and the fills are never perfect.

RETIRE This is a fade with no room. Its best day is a rounding error and its real fills turn it into a steady bleed; the loss isn't a couple of bad trades you can filter, it's the structure. No honest one-week filter flips it green, and I won't dress up "trade less into a book that still loses" as a fix. Pull it. If anything here is worth keeping it's the observation, not the sim: **fading into high rvol / high OFI on MES gets run over** — that's a veto worth reusing on the sims that actually have an edge to protect.

Notes: money scored on `sr.shadow_real_pnl.real_pnl`, never the optimistic sim pnl. Week filter entry_ts ≥ 2026-07-06. Entry conditions matched from `fut_signal_funnel` on the nearest cycle within 30s of entry — 22 of 24 trades matched (the 2 unmatched were +\$5.5 combined, immaterial). One week, in-sample, thin n — but with a positive ceiling this small, more data isn't going to rescue it.

GATE IT — ONE FILTER FIXES IT · \$-108 THIS WEEK ON HONEST FILLS.

MGC midband pullback bled \$-107.7 over 31 trades because it fades pullbacks into net selling and gets stopped; gate entries to only fire when order flow is with you (OFI>0).

§43 MGC Midband Pullback: It Buys Into Selling and Gets Run Over

This one bled. On the honest tick-repriced fills it lost **\$-107.7** across **31 trades** this week, won only **45%** of them, and left a perfect-fill ceiling of **\$249.1** on the table. So the money was there to be had — the fills and the bad entries ate it.

\$-107.7

HONEST REAL_PNL, 31
TRADES

45%

WIN RATE

\$249.1

PERFECT-FILL CEILING
(FANTASY)

§43.1 Why it bled: the winners trail out, the losers get stopped

Split by how each trade ended and the story is stark. The trades that got to trail a ratchet made money — **18 RATCHET exits, +\$110.8, 61% winners**. The trades that hit the stop are the whole loss and then some — **12 STOP exits, -\$221.4, and only 17% of those were green**. One lone flat-clock exit chipped in \$2.9. There is nothing wrong with the exit here. The stops are firing because the entries are bad.

WHAT IT'S ACTUALLY DOING

It's a pullback strategy — it waits for price to pull back to a mid-band and buys the dip (or sells the rip). When it works, the pullback holds and it trails a ratchet up. When it fails, the pullback isn't a pullback at all — it's the start of a real move that keeps going, and it walks straight into a stop. It's fading a move that isn't done.

The entry conditions tell you exactly which is which. Matching each trade to the signal funnel at its entry moment (28 of 31 matched — the sim's entry_sid is the session id, not the cycle timestamp, so I matched the nearest funnel row within 90s), the winners and losers separate cleanly on order flow:

BUCKET	N	REAL_PNL	AVG OFI	AVG RVOL	AVG VWAP-DIST
Winners	12	+\$204.6	+19.7	0.48	0.002
Losers	16	-\$315.9	-4.8	0.72	0.015


Read that in plain words. The winners came in with **order flow on their side** (OFI positive, +19.7), in a **quieter tape** (rvol 0.48), and **right on the band** (barely off vwap). The losers faded **into net selling** (OFI negative), in a **hotter tape** (rvol 0.72), reaching **further from vwap** to do it. It's buying the dip while the dip is still being sold — and the dip keeps going.

§43.2 The one thing for Monday: don't fade into the order flow

The single condition that best separates its own winners from its losers this week is order-flow imbalance. Gate it: **only fire when OFI > 0** — stand down when net flow is against the pullback.

OFI GATE	N	REAL_PNL	WIN%
KEEP — OFI > 0 (flow with you)	16	-\$11.9	50%
DROP — OFI ≤ 0 (fading into selling)	12	-\$99.4	33%

The dropped bucket is nearly the entire week's loss: those 12 trades cost -
\$99.4 at a 33% hit rate. Cut them and what's left is **16 trades at about breakeven (-\$11.9, 50% win)**. So in-sample the gate swings roughly **+\$99 for the week** versus running the sim as-is.



Cumulative honest P&L, trade by trade (gold dashes = breakeven). It never gets above the line after trade 5 and grinds to a -\$195 low before the late ratchets claw some back. That mid-week trench is the OFI-negative fades.

BE HONEST ABOUT IT

This is one week and thin — 12 trades in the drop bucket, 16 kept, and even the "good" bucket is only breakeven, not green. The gate doesn't make this sim pay; it stops it bleeding. That's mostly "trade less on the days flow is against you." The mechanism is sound — a pullback entry that fires while the tape is selling into you deserves to be vetoed — but I'd want another week or two before calling it an edge. Long and short bled about evenly (LONG -\$62.9, SHORT -\$44.8), so this isn't a one-sided problem.

GATE: OFI > 0 Add the order-flow veto and re-observe. If it holds up out-of-sample it's a keep; if the breakeven doesn't turn green in two weeks, retire it — a midband fader that only manages breakeven isn't worth the margin.

GATE IT — ONE FILTER FIXES IT · \$-102 THIS WEEK ON HONEST FILLS.

On honest fills this MES sim lost \$101.50 over 21 trades (43% win) versus a +\$48.10 ceiling, and every dollar of it is 8 stop-outs on fat bars; Monday move is tighten its ATR ceiling to veto entries at `atr_pct >= 0.020`.

§44 tw_mes_atr_ceiling_003 — the ceiling's set too high, so it keeps buying the fat bars that stop it out

This one bled this week. On the honest tick-repriced fills it lost **\$101.50 across 21 trades**, won just 43% of them, and it did that while its perfect-fill ceiling said it should have made **+\$48.10**. So this isn't slippage eating a good book — the entries themselves are wrong. And they're wrong in one very specific, very fixable way.

-\$101.50

HONEST REAL_PNL, 21
TRADES

43%

WIN RATE

+\$48.10

PERFECT-FILL CEILING
IT MISSED

-\$111.70

ALL OF IT IS THE 8 STOP-OUTS

§44.1 Where the money actually went

The whole loss is the stops. Nothing subtle here — split it by how each trade left and the picture is stark. Every trade that got stopped out lost, and


together they are more than the entire week's damage. The trades that got to ratchet out, or that just ran to the session flatten, were fine.

HOW IT EXITED	N	REAL_PNL	WIN%
STOP	8	-\$111.70	0%
SESSION_FLAT	6	+\$0.30	50%
RATCHET	7	+\$9.90	86%

So the question becomes simple: what's different about the trades that get stopped? And the answer is sitting right in the name of the sim. It's an ATR ceiling sim — it's supposed to stand down when the bar gets fat. Its ceiling is just set too loose, and the fat bars are walking straight through it.

§44.2 Why — thin tape it ratchets out green, fat tape it gets run over

Funnel conditions cover 19 of the 21 trades. Line up each entry's bar-size (atr_pct) against what it made and the split is clean: the losers sit on the fat bars to the right, the winners sit on the thin bars to the left. When the bar is calm this thing quietly grabs its ratchet and books a small win. When the bar is wide, price has enough room to swing through its stop before the ratchet ever arms — and it eats a full stop-out.



The regime tag doesn't sort it — losers show up in both "chop" and "trending" — and rvol doesn't sort it either (losers ran from 0.11 to 0.87). It's bar-size, plain and simple. The three worst trades of the week were shorts and longs entered at atr_pct 0.028, 0.027 and 0.024, and they lost \$19.10, \$17.90 and \$16.60. The winners almost all came in under 0.017.

§44.3 The one thing for Monday — bring the ceiling down to atr_pct 0.020

Split its own trades at an atr_pct ceiling of 0.020 and it separates the book almost perfectly. Everything at or above 0.020 is a loss factory; everything below is roughly break-even with a real win rate.

BUCKET	N	REAL_PNL	WIN%	AVG/TRADE
FAT — atr_pct >= 0.020 (veto these)	9	-\$92.10	22%	-\$10.20
THIN — atr_pct < 0.020 (keep)	12	-\$9.40	58%	-\$0.80

GATE: stand down when atr_pct >= 0.020 Vetoing the nine fat-bar entries removes **-\$92.10** of loss in-sample and takes the week from -\$101.50 to about -\$9.40. The kept book wins 58% and averages a hair under zero per trade.

BE HONEST

This gate stops the bleed; it does not mint money. The kept 12 trades still net -\$9.40 — flat, not green — so the honest read is "the ceiling is set too high, tighten it," not "flip a switch and print." It's one week and thin-n (nine vetoed, twelve kept), and any single bar-size filter partly just means trade less. But the mechanism is the sim's own stated purpose failing: it's an ATR-ceiling sim whose ceiling let the fat bars through. Bringing the ceiling from wherever it sits down to ~0.020 is a tightening of its existing knob, not a new bolt-on, which is exactly the kind of change that tends to hold up out of sample. Watch it next week; if the thin book stays break-even-plus, this earns its keep.

GATE IT — ONE FILTER FIXES IT · \$-95 THIS WEEK ON HONEST FILLS.

Confirmed passive_chop_calm lost \$95 over 46 trades (26% win) on honest fills with every dollar of bleed in OFF-hours; Monday move is to gate it to the US cash session (13-20 UTC) and stand down overnight.

§45 passive_chop_calm: it fades the quiet, but the quiet is overnight and the fills are ugly

This one bled again. On the honest tick-repriced fills it lost **\$95 over 46 trades** this week, winning just **26%** of the time. Its perfect-fill ceiling was \$-161.5, so the fills didn't rob it — the entries did. It is a passive reversion sim that fades calm chop, and this week the calm it found was the overnight tape, where a fade that gets picked off gets picked off hard.

-\$95

HONEST REAL_PNL, 46
TRADES

26%

WIN RATE

-\$97.7

ALL THE BLEED SITS IN
OFF-HOURS

+\$2.6

THE 7 US-HOURS TRADES, NET

§45.1 What it's actually doing

Look at how it exits and the story writes itself. When the tape stays genuinely calm and the position just times out, it's fine — the 20-minute FLAT_CLOCK exits went **+\$13.9 over 15 trades at 53% win**. The damage is the other two doors: it got **stopped 17 times for -\$69.7 (18% win)** and dragged to the close **14 times for -\$39.3 (7% win)**. That is the signature of

a reversion book fading a move that keeps going: sit in the chop and you scratch out, but the moment price actually picks a direction it runs your fade over.

And it happens overnight. One trade tells the whole tale — a **MGC short at 00:38 UTC that got stopped for -\$43.6 by itself**, nearly half the week's loss in a single thin-tape gold fade. By symbol it's Bitcoin that's the sinkhole: **MBT lost \$45 over 15 trades at 7% win**, a coin flip that never landed. The daytime, real-session names (MES, MYM, MNQ, M2K) are all roughly flat-to-green.

FUNNEL COVERAGE — READ THIS


Only **4 of 46** trades matched a `fut_signal_funnel` row, and all 4 were losers, so I can't honestly split this sim on `rvol` / `atr_pct` / `book`. Also worth knowing: **22 of 46** trades never got a fill status (`real_pnl` is `NULL`) — the live money is really 24 filled trades. So the gate below is built on the clocks and the tape hours I can trust, not on microstructure I can't.

\$45.2 When it fires vs when it pays

The cleanest line through this sim isn't a book number — it's the clock. Split every trade by whether it fired during the US cash session (13-20 UTC) or outside it:

BUCKET	N	REAL_PNL	WIN%	WHAT IT IS
OFF-hours (Asia / pre-Europe)	39	-\$97.7	21%	thin tape, bad fills, one-way drifts run the fade over
US cash 13-20 UTC	7	+\$2.6	57%	real liquidity, someone home on the other side

Every dollar of the week's loss is overnight. The 00-01 UTC hour alone was **-\$59 over 8 trades**. Cut the sim to US hours only and the week goes from -\$95 to +\$2.6 — an in-sample swing of about **+\$98**.



Cumulative honest real_pnl, in trade order. It craters on that one overnight MGC stop and never recovers — a slow grind straight down. Gold dashed line is breakeven.

§45.3 The one thing for Monday

STAND DOWN OVERNIGHT Gate passive_chop_calm to the US cash session (13-20 UTC) and don't let it fire in thin overnight tape. In-sample that's worth about **+\$98** this week — it turns a \$95 loss into a \$3 scratch.

Be honest with yourself about what this is, though: the gate mostly works by **trading 7 times instead of 46**. It's a "stand down" more than a "make money" — the surviving US bucket is only 7 trades and +\$2.6 is a scratch, not an edge. This is one week and in-sample. If you want a second, cleaner cut on the same theme, **kill MBT** outright — -\$45 over 15 trades at 7% win is a name this sim should never touch. Between the overnight veto and dropping Bitcoin, there's no version of this week where staying in the overnight tape was right.

Verdict: **GATE to US hours** — not a retire yet, because the daytime, real-session trades are roughly flat and the mechanism (fade calm chop, time out flat) is sound where liquidity exists. But if next week the gated 7-trade bucket doesn't hold green, this is a retire, not a tweak.

RETIRE — NO HONEST FIX · \$-88 THIS WEEK ON HONEST FILLS.

This sim shorted MNQ pullbacks four times and lost on all four for \$-88.48 honest, with even its perfect-fill ceiling underwater; the Monday move is a chop veto (fire only in a trending regime / $r_{vol} \geq 0.30$) that stops the bleed but leaves it at scratch.

§46 The MNQ Greenband-Block Short: Four Fired, Four Bled

This sim shorts MNQ pullbacks when a "greenband" block says the bounce is done. This week it fired four times, lost on all four, and handed back **\$88.48** on the honest fills. There were no winners to learn from. There's a clean way to make it bleed less, but it's really a way to make it trade less.

-\$88.48

HONEST_REAL_PNL, 4
TRADES

0%

WIN RATE (0 OF 4)

-\$66.87

PERFECT-FILL CEILING
— EVEN THE FANTASY
LOSES

§46.1 Did it make money? No. And the ceiling proves it wasn't just bad fills.


I re-ran the join myself on `shadow_real_pnl`: four MNQ shorts this week, `real_pnl` of -0.12, -39.62, -46.12 and -2.62, summing to **-\$88.48**. The perfect-fill ceiling was -\$66.87 — so even if every fill had been flawless, this thing still loses money. That's the important tell. Bad fills cost it about \$22 on top, but the setup itself was underwater. You can't slippage-tune your way out of a losing idea.

§46.2 Why it bled: it's shorting into chop, below VWAP, when nobody's trading

I matched each entry to the nearest fut_signal_funnel row (all four inside 2 seconds of the fire, so coverage is honest here). The pattern is stark. The two trades that got stopped for real money — the -\$39.62 and the -\$46.12 — both fired in a **chop** regime, on limp volume (rvol 0.24), with price already **below** VWAP. So the sim shorts a pullback that's already sagging, the market mean-reverts back up through it, and the stop pays for it. That's the whole loss right there: **\$85.74 of the \$88.48 came from two chop-regime stop-outs.**

The one trade that behaved differently is the tell. Trade 12349 fired in a **trending** tape with rvol 0.38 — and its ceiling was **+\$25**, the only positive ceiling of the week. The idea actually worked directionally; the ratchet just gave it all back to scratch (-0.12). The fourth trade fired on a monster bar (atr_pct 0.154, roughly 3x the others) and only escaped at -\$2.62 because the flat-clock pulled it before the stop could — pure luck, its ceiling was -\$36.62.

TRADE	UTC	REGIME	RVOL	ATR_PCT	VWAP_DIST	EXIT	REAL_PNL
12349	23:28	trending	0.38	0.030	-0.002	RATCHET	-0.12
16559	00:54	chop	0.24	0.041	-0.022	STOP	-39.62
17014	09:04	chop	0.24	0.045	-0.005	STOP	-46.12
17315	14:39	chop	0.20	0.154	-0.036	FLAT_CLOCK	-2.62



§46.3 The one thing for Monday: don't fire it in chop

The cleanest split its own trades offer is **session regime**. Every dollar of damage came from the chop bucket; the single trending trade only scratched. Same line if you cut it on volume instead — a `rvol >= 0.30` floor keeps exactly the trending trade and drops all three chop ones, because the chop fires all sat at `rvol 0.20-0.24`.

BUCKET	N	REAL_PNL	WIN%
Chop regime (as-is bleeders)	3	-88.36	0%
Trending regime (the only survivor)	1	-0.12	0%

HONEST READ

Vetoing chop swings the week by **+\$88.36 in-sample** — but read what's left: one trade, at $-\$0.12$. The gate doesn't find you a winner, it just stops you shorting into mean-reverting chop. And it's four trades, one week, all losers, so "trending only" is really "trade a quarter as often and break even." That's not an edge, that's a tourniquet.

So the verdict is a hard gate bordering on retire. **RETIRE unless chop-vetoed**
`rvol >= 0.30` floor If you want to keep it alive to gather more tape, arm the chop veto and let it fire only in a trending regime — but understand that this week it produced zero winners and its best-ever outcome was a scratch. Don't give it size. The greenband-block short needs a real trending week to prove it has any pulse at all; on chop it's just paying the spread to be wrong.

GATE IT — ONE FILTER FIXES IT · \$-78 THIS WEEK ON HONEST FILLS.

This MYM drive-continuation sim bled \$77.68 over 39 trades because it chases upward drives while the tape sags below VWAP; the Monday move is to only fire when price is at/above the session VWAP, which flips the matched book from -\$74 to +\$7 in-sample.

§47 The MYM drive-chaser that only pays above the line

This is the "no-fade drive" sim on MYM — the micro Dow. It doesn't fade moves, it chases them. This week it chased into a sagging tape and bled. The honest fills lost **\$77.68 over 39 trades**, 46% win. The interesting bit: even the perfect-fill fantasy only made \$72.6, so this thing was never far from red — real slippage tipped it over.

-\$77.68

REAL P&L, 39 TRADES

46%

WIN RATE (18/39)

+\$72.6


PERFECT-FILL CEILING

-\$91.9

LOST ON STOPS (16 TRADES)

§47.1 Did it make money? No — and the shape tells you why

It bled all week, in a straight line down. Here's the honest equity curve, all 39 trades in order. No bad day did the damage — it was a steady drip.



Gold dashed line = break-even. Every trade drags it lower; the curve never gets back above zero after trade three.

§47.2 Why it bleeds: small winners, full-size stops

The engine underneath is broken in a familiar way. The winners get ratcheted out for lunch money and the losers ride all the way to the stop. Average winner: **+\$3.24**. Average loser: **-\$6.48**. Losers are twice the size of winners, so 46% win isn't nearly enough to stay alive.

Split by how it exits and the whole loss is in one bucket:

EXIT REASON	TRADES	REAL P&L	WINS
STOP	16	-\$91.92	3
RATCHET	18	+\$12.84	11
FLAT_CLOCK	5	+\$1.40	4

The ratchet and clock exits are actually green. The entire \$77 hole is the 16 stop-outs, and only 3 of those ever showed a profit. So the question isn't the exit — it's "why does it keep buying drives that immediately reverse and stop out?"

§47.3 The answer: it chases drives while the tape is sagging

I pulled the entry conditions from the signal funnel — 34 of the 39 trades matched a funnel snapshot within five minutes (87% coverage; 5 had no row). The cleanest tell by a mile is where price sits relative to the session

VWAP at entry. VWAP is the day's centre of gravity. Above it, the tape has an up-posture; below it, it's leaking.

THE TELL

When it fired with price **at or above VWAP**, it made **+\$7.2 over 15 trades at 60% win**. When it fired **below VWAP**, it lost **-\$81.3 over 19 trades at 32% win**. Same story on the killer stops: below-VWAP stop-outs lost -\$72.2 (10 trades); above-VWAP stops only -\$14.2 (2 trades).

In plain words: this is a momentum-continuation sim, but below VWAP the momentum it's trying to ride is down, so its upward drive-chases fail and get run over. It's also firing in thin tape — average rvol around 0.4-0.5, well under a normal bar — so there's often no real drive there to continue in the first place.

§47.4 The one thing for Monday

GATE: only fire when vwap_dist_pct >= 0

BUCKET	TRADES	REAL P&L	WIN%
At/above VWAP (keep)	15	+\$7.2	60%
Below VWAP (stand down)	19	-\$81.3	32%

Stand down when price is below the session VWAP. In-sample that turns the matched book from -\$74 into **+\$7** — roughly an **\$81 swing** — and lifts the win rate from 32% to 60%. It works because it deletes exactly the trades that were producing the fatal stop-outs.

HONEST CAVEAT

This is one week, in-sample, and it's thin: the gate keeps only 15 of 34 matched trades, so mostly it just means trade a lot less. It also nearly

kills the shorts — they cluster right at/below VWAP — so the sim effectively becomes a long-only "up-posture drive" filter. I wouldn't size it up on this. Arm it observe-only, and re-check next week that above-VWAP drives still pay before believing the \$7 is edge and not just survivorship.

Verdict: Don't retire it yet — there's a real, sensible posture filter here (chase drives with the tape, not against it) that switches the bleed off. Gate it on VWAP posture, keep it in the shadow, and prove it holds up out-of-sample before it earns anything more.

RETIRE — NO HONEST FIX · -\$76 THIS WEEK ON HONEST FILLS.

This M2K open-momentum sim went 0-for-3 and lost \$76.48 on honest fills, but its perfect-fill ceiling is still -\$61.27, so the entry has no edge to gate toward — retire it.

§48 tw_m2k_open_momentum — the open thrust that never thrusts

This M2K "ride the opening move" sim went 0-for-3 this week and lost **\$76.48** on honest fills. The tell is brutal: even with perfect fills it would still have lost \$61.27. That is not a slippage problem you can gate away — the entry itself is wrong. Retire it.

-\$76.48

HONEST_REAL_PNL, 4
FILLS

0%

WIN RATE (0 OF 3
ENTRIES)

-\$61.27

PERFECT-FILL CEILING
— STILL RED

4 of 4

EXITS ON STOP

§48.1 1. Did it make money? No — it bled every single time.

Three entry events this week, each fanned into exit variants; four of them got tick-repriced into honest fills. Two shorts on 07-06 and 07-08, one long on 07-10. All three got stopped out. Every honest fill is red, and the exit reason is **STOP** on all four. There is no winner to point at — the whole week is one colour.

EVENT

SIDE

REAL_PNL

CEILING

EXIT

07-06 14:00	SHORT	-\$48.24	-\$34.22	STOP (x2 fills)
07-08 11:57	SHORT	-\$12.12	-\$11.63	STOP
07-10 13:51	LONG	-\$16.12	-\$15.42	STOP
total		-\$76.48	-\$61.27	

THE SMOKING GUN

The perfect-fill fantasy ceiling is **-\$61.27**. If you could fill every entry at the exact tick you wanted, zero slippage, this sim still loses sixty-one dollars. When the ceiling is red, no book gate, no rvol floor, no fill trick can save you — the signal is picking the wrong moment to be in the market. The remaining \$15 gap (fill vs ceiling) is real but it's rounding error next to a broken entry.


§48.2 2. Why — it fires into fat, fast bars and gets run over

The prescribed funnel join (cycle_ts = entry_sid) matched **nothing** — this sim's entry_sid points ~16 hours back to the prior 22:00 session marker, not the fire moment, so honest coverage on that join is 0 of 4. I pulled the conditions at the real entry_ts instead; all three events are covered within seconds of the fire.

Here is what it saw when it pulled the trigger:

- **07-06, the big one (-\$48):** bar was **fat** — atr_pct 0.103, regime tagged trending, order flow -49 to -51 (sellers leaning), book imbalance -0.5. It shorted into a fast trending tape and the stop caught it. This one entry is 63% of the week's bleed.
- **07-08 (-\$12):** atr_pct 0.068, rvol a dead 0.12-0.18, regime chop. It shorted a thin, going-nowhere tape and still lost.
- **07-10 (-\$16):** atr_pct 0.087, regime trending, order flow +30 to +110 (buyers). This time it went long with the flow — momentum-aligned, the way it's supposed to — and it still got stopped.

That last one is the killer detail. On 07-10 the sim did everything right — bought a trending, buyer-led tape near the open — and lost anyway. So this isn't "it's on the wrong side." It's that an M2K open-momentum entry with this stop just doesn't have an edge. Both shorts lost, the one clean long lost. Direction isn't the problem; the whole premise is.



Cumulative honest P&L, fill by fill, left to right (red). Gold is the perfect-fill ceiling — note it never lifts off zero either. Both staircases only go down.

§48.3 3. The one thing for Monday — retire it (or at best, veto fat bars)

With three entries in one week and zero winners, there is no "good-condition bucket" to lean into — every bucket is empty of green. The only honest split I can draw is by bar fatness, and even that just picks which loss to skip:

BUCKET	N FILLS	REAL_PNL	WIN%
fat bar, atr_pct ≥0.10 (07-06)	2	-\$48.24	0%
the rest (07-08, 07-10)	2	-\$28.24	0%

A "stand down when the bar is fat" veto ($\text{atr_pct} \geq 0.10$) removes only the 07-06 event and saves **\$48.24** in-sample — the single biggest chunk. But it does not flip the sim green: the other two events still bleed \$28, and one of those (07-10) was a textbook momentum-aligned long. Cutting the fat bar is trimming a wound, not healing it.

HONEST CALL

Three trades, one week — this is noise, not a study. But two things are not noise: it's **0-for-3**, and its **perfect-fill ceiling is still -\$61**. When even the fantasy loses, the entry has no edge to gate toward. There's no rvol floor or session-hour that turns a negative ceiling positive. My recommendation is **retire it**. If the desk wants to keep it breathing on the shadow bench, the only defensible guardrail is the $\text{atr_pct} \geq 0.10$ fat-bar veto — it saves \$48/wk in-sample but leaves a bleeder, so it's a stay-of-execution, not a fix.

RETIRE `tw_m2k_open_momentum`

or, if kept: veto $\text{atr_pct} \geq 0.10$

GATE IT — ONE FILTER FIXES IT · \$-69 THIS WEEK ON HONEST FILLS.

A US-open continuation sim lost \$68.6 over 13 trades because it fired five shorts that got run over; gate it long-only for Monday.

§49 The Open-Surge Sim Shorts the Rip It's Meant to Buy

This one is called `us_open_surge_cont` — it's supposed to buy the US-open surge and ride it. This week it lost \$68.6 on the honest fills over 13 trades, winning only 38% of them. And the sim's own optimistic book only showed $-\$42.8$, so even the fantasy fills were red — the tick-repriced truth just made it \$26 worse.

$-\$68.6$

HONEST REAL_PNL, 13
TRADES

38%

WIN RATE

$-\$42.8$

PERFECT-FILL CEILING
(ALSO RED)

§49.1 Why it bled: it fired five shorts and got run over

Split its trades by direction and the whole loss is on one side. The eight LONG trades came out roughly flat — up \$5.1, half of them green. The five SHORT trades lost \$73.7 between them and only one won. That's the entire drawdown, and then some: the longs are basically break-even, the shorts are the bleed.

Think about what that means for a sim whose whole job is continuation. It's meant to buy strength off the open and hold it. When it flips short it's doing the opposite — fading a rip that's still running — and the tape ran it straight over. The worst three shorts all fired in the same minute, 13:29 UTC right at

the open: MES −\$37.9, MYM −\$28.6, M2K −\$10.6. One burst of fades at the bell cost \$77.


The exit mix says the same thing another way. Every one of its five STOP exits was a loss (−\$112.4 together) — it entered wrong and got taken out. Every one of its four RATCHET exits was a win (+\$77.3) — when the entry was right, the trailing lock did its job cleanly. The exit engine isn't the problem. The direction of the entry is.

FUNNEL COVERAGE — THIN

Honesty note: only 2 of the 13 trades had a matching fut_signal_funnel row, so I can't build the rvol / atr_pct / book split I'd normally want. There isn't enough microstructure captured on this sim's fires this week to gate on the tape. The clean separator I do have — trade side — needs no funnel and is unambiguous.

§49.2 The condition split — long vs short

BUCKET	N	REAL_PNL	WIN%	WHAT IT IS
LONG only	8	+\$5.1	50%	buying the surge — flat, does no harm
SHORT	5	−\$73.7	20%	fading the rip — the whole bleed
All trades	13	−\$68.6	38%	as-is this week



Each bar is one trade in time order; height is honest real_pnl, up is a win. Red bars are SHORTs, navy are LONGs. The four deepest drops are all red — the shorts at and just after the open.

§49.3 The one thing for Monday

LONG-ONLY VETO Stop it shorting. A continuation sim has no business fading the open surge, and this week that's provable: kill the five shorts and the week goes from $-\$68.6$ to $+\$5.1$ — a $\$73.7$ in-sample swing that flips the sign. Only fire this sim long; when the signal wants to sell the open, stand down.

BE HONEST ABOUT IT

Two caveats, both real. First, $n=5$ shorts is thin and it's one week — I can't promise the short side is structurally broken, only that this week it was a disaster and it contradicts the sim's own thesis, which is the kind of filter I trust more than a curve-fit. Second, long-only doesn't make this a green sim — it makes it a flat one ($+\$5.1$ over 8). This gate stops the bleeding; it doesn't yet find an edge. If the long side is still hovering at break-even in another two weeks, retire the whole sim rather than keep it on the book.

Verdict: GATE it long-only for Monday, watch the long side for a real edge, and be ready to retire if it stays flat.

RETIRE — NO HONEST FIX · \$-46 THIS WEEK ON HONEST FILLS.

Scored the MBT pullback sim honestly: 0-for-8 at -\$46.50 with a ceiling that's also red (-\$18.94) and -\$27.56 lost to fills — retire it.

§50 The MBT pullback that loses even when it's right

This one bled every single trade. It took 8 pullback entries on micro-Bitcoin this week and lost all 8, for a real total of **-\$46.50**. There is no winner to build a gate around. And the deeper problem: even with perfect fills it still loses money. Retire it.

-\$46.50

HONEST REAL P&L, 8
TRADES

0%

WIN RATE (0 OF 8)

-\$18.94

PERFECT-FILL CEILING
— STILL RED

-\$27.56

LOST TO FILLS ALONE

§50.1 Did it make money? No — it went 0-for-8

Every trade lost. Eight fired, eight closed red, real P&L **-\$46.50**. That's not a bad-luck streak inside a good sim — it's the whole week. And the number that ends the argument is the ceiling: if you gave this sim perfect fills on every trade, it would still lose **-\$18.94**. When the fantasy version bleeds, the entry itself is the problem, not the slippage. You can't gate your way out of a negative edge.

§50.2 Why: the entries don't follow through, and the fills finish them off

Two things are wrong at once. First, the pullbacks don't go anywhere. Seven of the eight trades didn't hit a target or a stop — they got flushed out by the clock (exit reason **EOD**), only one ever ratcheted a profit lock. It buys a dip on micro-Bitcoin, the dip doesn't bounce, and the position just sits there until the session flat-clock kills it a few minutes later. That's a signal with no follow-through.

Second, MBT is a thin, empty tape and the fills are brutal. Look at the two trades where the perfect-fill version was actually green — a +\$12.75 ceiling and a +\$4.81 ceiling. Both came back **-\$0.50** on the honest fills. Nobody's home on micro-Bitcoin, so the fill you get is the one you didn't want. Across all 8, fills alone cost **-\$27.56** — more than half the total loss. Three trades had a winning ceiling; the tape dragged all three under.

THE TELL

Even at fantasy perfect-fill prices this sim loses -\$18.94. That's a negative entry edge. The extra -\$27.56 of real-world slippage on a spread-10-to-15 book just pours salt in it. Neither half is fixable with a filter.


§50.3 Trade by trade: the fantasy vs what you actually got

SIDE	HONEST REAL	PERFECT-FILL CEILING	COST OF THE FILL	EXIT
SHORT	-\$3.50	+\$0.50	-\$4.00	clock
SHORT	-\$7.50	-\$12.25	+\$4.75	clock
LONG	-\$6.00	-\$11.25	+\$5.25	clock
LONG	-\$0.50	+\$12.75	-\$13.25	clock
LONG	-\$7.00	-\$5.75	-\$1.25	clock
LONG	-\$10.00	-\$4.50	-\$5.50	clock

LONG	-\$11.50	-\$3.25	-\$8.25	clock
LONG	-\$0.50	+\$4.81	-\$5.31	ratchet

§50.4 Both lines are underwater — the gap is the fills

Red is the honest money, gold is the perfect-fill fantasy. The fantasy is less bad, but it never climbs above the dotted zero line either. There is no version of this week where this sim is green.



■ honest real, ends -\$46.50 ■ perfect-fill ceiling, ends -\$18.94 the space between them is pure slippage.

§50.5 The one thing for Monday: retire it

I went looking for the gate that saves it and there isn't one. Every trade is red, so no condition splits winners from losers — there are no winners. The closest thing to a lever is bar fatness: I could pull the entry conditions for 7 of the 8 trades (the 8th had no clean funnel row at signal time, so 7/8 coverage). Splitting on the bar's ATR% shows the thin-bar trades did the real damage.

BUCKET	TRADES	REAL P&L	WIN%	READ
Thin bar (atr% < 0.065)	4	-\$34.50	0%	the bulk of the bleed
Fat bar (atr% ≥ 0.065)	3	-\$8.50	0%	less bad, still 0 wins

So yes, "stand down when the bar is thin" would have saved about **+\$34.50** in-sample. But be honest with yourself about what that is: it's trade less and

still lose — the fat-bar bucket is still 0% wins and still -\$8.50 over 3 trades, on a one-week sample of 8. That's not an edge you gate into; that's a losing sim with a smaller loser hiding inside it. And it doesn't touch the real killers: a pullback with no follow-through (7 of 8 died on the clock) on a tape so thin the fills flip your winners into losers.

VERDICT

RETIRE `tw_mbt_pb_rvol_band` — 0-for-8 at -\$46.50, a ceiling that's also red at -\$18.94, and MBT's empty book eating -\$27.56 in fills. The entry edge is negative before slippage even shows up. No filter fixes a negative edge on a \$0.10-per-point contract paying \$2.50 round-turn. Kill it. If you want to keep one eye on it, leave it observe-only and only ever on fat bars ($\text{atr\%} \geq 0.065$) — but I wouldn't spend the capital.

GATE IT — ONE FILTER FIXES IT · \$+0 THIS WEEK ON HONEST FILLS.

This M2K VWAP-reversion sim went breakeven (-\$0.4 over 66 trades) because its stop-outs (-\$152) exactly cancel its ratchet wins (+\$139); vetoing trades that fire above VWAP flips it green.

§51 **tw_m2k_vst** — a good little dip-buyer that keeps taking one trade too many

This M2K sim ran 66 trades this week and, on the honest tick-repriced fills, made nothing: **-\$0.4**, 47% winners. The perfect-fill fantasy said \$120.9. So the fills ate the entire edge — it should have been up a hundred bucks and instead it's flat. The good news: the bleed is one clean, nameable mistake, and vetoing it flips the book green.

-\$0.4

HONEST REAL_PNL, 66
TRADES

\$120.9

PERFECT-FILL CEILING
(FANTASY)

47%

WIN RATE

+\$21.9

THE HALF THAT FIRES BELOW VWAP

§51.1 **Did it make money? No — the stops give back everything the ratchet earns**

Split it by how each trade ended and the story is right there. The **ratchet exits made \$139.1 over 28 trades at 75% wins** — when this thing works it works well. But the **28 stop-outs lost \$152.4 at 14% wins**, and that one

bucket wipes out the entire ratchet haul plus the small VWAP-loss gains. It's not a broken engine. It's a good engine dragging a bag of trades it should never have taken.



Net honest P&L by exit reason, this week. The green ratchet bar and the red stop bar are almost exactly the same size — that's why the week is a wash.

§51.2 Why — it buys the dip when it wins, and chases the rip when it loses

I matched each trade back to the entry-condition funnel (M2K rows within 90 seconds; **43 of 66 trades matched**, 65% coverage — the rest have no funnel row, so read the split as directional, not gospel). The winners and losers separate on one thing: which side of VWAP the trade fired on.

Winners fire when price sits **below** VWAP on ordinary tape (average rvol 1.06) — it's buying the dip and fading the stretch, which is exactly what a short-term VWAP-reversion sim is supposed to do. Losers fire when price is already **above** VWAP on hotter tape (average rvol 1.29) — that's chasing strength that's already run, and the stop catches the knife. The damage is almost entirely one shape: **longs taken above VWAP lost \$45.8 over 19 trades at 26% wins**. That single wrong-side-of-the-line bucket is the whole bleed.

§51.3 The one thing for Monday — stand down when the trade fires on the wrong side of VWAP

Gate it on the sign of `vwap_dist_pct`: only take the trade when price is on the reversion side (below VWAP for the longs, at/below for the shorts).

Here's the same 43 covered trades split good-bucket vs the rest.

BUCKET	WHAT IT IS	N	REAL P&L	WIN%
<code>vwap_dist < 0</code>	fires below VWAP (buy the dip / fade the stretch)	22	+\$21.9	59%
<code>vwap_dist ≥ 0</code>	fires above VWAP (chasing) — veto these	21	−\$41.5	29%
→ longs above VWAP	the worst corner of that bucket	19	−\$45.8	26%

THE MONDAY MOVE

Veto any entry where `vwap_dist_pct ≥ 0` — don't buy above VWAP, don't take the trade on the wrong side of the line. On the 43 covered trades that turns **−\$19.7 into +\$21.9**, a **~\$42 in-sample swing**, and lifts win rate from 44% to 59%. Applied to the full book it removes the −\$41.5 chasing bucket outright, taking the honest week from breakeven to roughly **+\$41**.

GATE: `vwap_dist < 0` only

ratchet exit is fine, keep it

The honest caveats. This is one week and only 65% of trades carry a funnel row, so the exact dollars are soft. And a sign-of-VWAP gate mostly means trade less — it keeps 22 trades and bins 21, roughly half the book. But this isn't a random threshold fished from noise: it's the sim's own reversion thesis. It's meant to fade stretch back to VWAP, and the money proves it only works when it enters on the reversion side. The trades above VWAP aren't reversion at all — they're the sim forgetting what it is. Gate

that out and it's a small, real, green edge instead of a wash. Worth arming as an observe-only veto and re-checking coverage next week before it earns size.

CHAPTER

The Small Fry — too thin to judge yet

WATCH — PROMISING, NOT PROVEN · \$+0 THIS WEEK ON HONEST FILLS.

Rounded up two sims too thin to judge — M2K absorption lost \$8.48 over 4 clean-filled trades and just needs more n; MBT spike-fade's single -\$12 was 3x its -\$4 ceiling on a thin-book fill, so ignore it.

§52 The Small Fry — too thin to judge, worth a glance

Two sims fired so few times this week there's nothing to prove either way. Here's what they are and whether to keep an eye on them. Don't read money into five trades.

-\$8.48

M2K ABSORPTION · 4 TRADES

-\$12.00

MBT SPIKE-FADE · 1 TRADE

§52.1 M2K absorption-only **WATCH**

This one waits for the book to soak up selling and then buys the little Russell. It fired four times this week and lost \$8.48 net on honest fills. Two got stopped for a combined \$13.74, one ratcheted out for a scratch, and one rode the flat-clock to a \$5.88 win. That's not a signal, that's a coin landing tails three times out of four. The fills were clean — filled every time, no phantom slippage story here, real money tracked the ceiling within a dollar or two. Leave it running and look again when it has twenty trades. Four tells us nothing.

§52.2 MBT spike-fade **IGNORE**

One trade all week. It faded a bitcoin spike, held for four minutes, and got carried out at the EOD flatten for -\$12. The tell isn't the loss, it's the fill: the

perfect-fill ceiling said this was only a -\$4 miss, and the honest tick-repriced fill was three times worse at -\$12. That's MBT being MBT — thin book, wide spread, nobody home to give you the price you wanted. One trade can't earn a verdict, but everything we already know about micro-bitcoin says the fill will keep eating any edge this thing thinks it has. Ignore it until it either fires enough to argue with or we fix the venue.

THE ONE THING

Nothing to change on Monday. Both are below the threshold where a filter means anything — any "fix" on $n=4$ or $n=1$ is just noise-fitting. Let them accumulate. M2K is worth the disk space; MBT's fill problem is the real story and it's structural, not a gate you can tune.

Every figure is from the shadow desk's own honest, tick-repriced fills (shadow_real.db :: real_pnl) for the week (Mon 6 -> Fri 10 Jul 2026) — never the sims' own optimistic numbers, which understate losses. One agent dissected each sim; independent skeptics re-ran every Monday refinement worth banking, and their verdict is stamped on each section. The live desk is settled and is not the subject here. Plain-English by design.