

A good short-seller wearing a bad long book. That's the desk — and it's the whole job for Monday.

28 sections on the shadow sims, on honest tick-repriced fills, to the locked scope. The edge is proven and one-sided; this report is where it bleeds and how we stop it.

Desk FLAT ·
HEALTHY

MNQ · paper 28 sections · ~30k words

30-agent run, skeptic-verified nothing armed

CONTENTS

- §1 What the board is telling you, and what to do Monday
- §2 Where the money was — the MNQ run census
- §3 Feed the thrust, starve the fade
- §4 The desk only works in one direction — longs bleed, shorts pay, in every family
- §5 The chandelier is the whole edge — and it hands back a third of every winner
- §6 The order book: our honest look at whether the wall tells us anything
- §7 The OFF Switch: Trade the Condition, Not the Gate
- §8 The sizing ladder and the departure lounge: who earned more size, who's booked for the exit
- §9 My Mid-Week Musings, Walked to a Verdict
- §10 grind_fast — the week's best earner, and it's built right
- §11 thrust_cont — the quiet winner that just takes its target
- §12 abs_veto_55s — the absorption-veto that quietly banked a clean week
- §13 thrust_loose — it made money, but only the short half did
- §14 abs_veto_50s — the short-side absorption sim that quietly banked a clean \$443
- §15 thrust_aligned — a short-selling bracket that quietly netted \$194
- §16 abs_veto_60s — the patient absorption-veto that barely traded
- §17 thrust_fast — a small green week carried entirely by its shorts
- §18 chand_k20 — the tightest trail was the only one that stayed green
- §19 capit_loose — a fade that just keeps getting stopped
- §20 rg_long_fast — the fast dip-buyer that mostly buys straight into a stop
- §21 rg_long_fast_v — the long-only grind that stopped out 25 times
- §22 rg_long_20_t20 — a long-only grind gate that only won when the tape was calm
- §23 chand_k35 — the loosest trail, and it gave the most back
- §24 rg_long_20_t15 — a long-only grind that gave it all back on the wild bars
- §25 rg_long_20_flow25 — a long-only reversion that only works in calm tape
- §26 capit_ride — bought two flushes, got stopped on both
- §27 chand_k25 — the middle child of the trail A/B, and it bled
- §28 The small fry: sims too thin to judge, and two housekeeping catches

§1 What the board is telling you, and what to do Monday

+\$21.8k

confirmed momentum +
chandelier, the 2 weeks before
cut-over (honest fills)

-\$15.8k

the fade / reversion / passive
book over the same 2 weeks

-\$1,054

the V7 board's long trades (the
155 shorts made +\$2,884)

The desk is a good short-seller wearing a bad long book around its neck. That one sentence explains the live desk, the V7 board, and the two weeks of history all at once — and it's the whole job for Monday.

The new V7 board is only two days old, so read its per-sim numbers as direction, not gospel. But it says exactly what the fortnight of real, tick-repriced history before it says: the money is in **confirmed momentum ridden by the chandelier** (the MNQ thrust family made **\$22,453** over the two weeks), and it bleeds straight back out through **the fade book** (`passive_chop` alone lost **\$12,528**) and **the long side of everything**. On the current board the shorts made **\$2,884** at 54% while the longs lost **\$1,054** at 30% — same fingerprint, live and shadow.

What's working

The thrust family is the backbone — all four members green, `thrust_cont` the quiet star at **\$594**, every dollar of it the take-profit doing its job. The chandelier is the desk's only real profit centre on the live side (**\$567** at an 85% win rate, its losers a rounding error). And the quiet surprise: the **absorption-veto family is the best-behaved thing on the board** (`abs_veto_55s` **\$499**, green on both sides) — which matters, because the 45-second version of exactly that veto was killed on the live desk this week. We may have cut the wrong thing.

What's bleeding

Everything long, and everything that fades. Capitulation is the worst family on the board (`capit_loose` **-\$622** at a 9% win rate — knifed on entry and stopped inside a minute). The long-only reversal-grabs bleed for the same structural reason: a long-only gate in a short-favourable tape is a losing bet by construction. And the census clears the entry side of blame — we showed up for **77% of the big MNQ runs and 92% of the really big ones**, so we are not missing the trends. The leak is between them: the long branch firing counter-trend, and the stop pile that follows.

What to do Monday (ranked by the money)

1. Attack the long side — it's the entire loss. Get the circuit breaker actually firing (it logged zero trades this week — dormant), and re-open the question of the absorption-veto we killed. **2. Keep the proven core, starve the fades.** Thrust + chandelier stays; do not run reversion / capitulation / passive live — a proven $-\$16\text{k}/2\text{wk}$ drain. **3. Build the vol-adaptive chandelier** — the give-back you watched ($\$224 \rightarrow \88) is the norm on high-ATR days, and an ATR-conditional trail beats a flat one; a flat tighten just clips the runners.

What survived the skeptic

The momentum-beats-fade split and the long-side bleed both held cold. Two things did **not**: the day-type router (it swung from $+\$1,972$ to losing money depending on where you cut the test — two lucky days, not an edge) and grind_fast as a standalone ($-\$7,975$ over 25 days; its big Friday was its single best day). They stay in the lab, not the desk. Nothing here arms on one week.

The one line: feed the confirmed thrust, ride it with the chandelier, and stop paying to be long and to fade. Everything below is the evidence and the per-sim detail.

WHERE THE MONEY WAS

§2 Where the money was — the MNQ run census

Over the 30 archive days in the tape (28 full sessions plus two half-days), MNQ printed **1,338 big runs** — a directional shove of at least 40 points that completed inside 30 minutes. We showed up for **1,030 of them, 77%**. And the bigger the run, the more reliably we were on it: of the 168 monster runs of 120 points or more, we were in **155, that's 92%**. The desk is not missing the moves. It shows up. The money leaks in the flat ground between the runs — and that is the thing to fix on Monday, not the run-catching.

The prize pool is real. Add up the point-size of every one of those 1,338 runs and you get **103,786 points of directional travel** — about **\$207,572** of gross movement at MNQ's \$2 a point across a month. Nobody catches all of it. But that is the pond we are fishing, and it is deep.

How a "run" is defined — and how we checked if we caught it

I ran this on 1-minute MNQ closes, day by day, using the same gate code the live desk runs (`gazbot7.deciders`). A run starts at any bar, scans forward up to 30 minutes tracking the directional extreme, and ends the first time price gives back more than half of the best move it had made. If that best move reached 40 points or more, it counts as a big run. The scanner then jumps past the extreme so runs never overlap and nothing gets double-counted. To score participation, I replayed the two momentum gates — `gate_grind` (slope floor 0.4, fast slope) and `gate_thrust` (thrust 1.5, amplitude floor 0.0004) — bar by bar, and asked one question: did either gate open a position in the run's direction at any point between the run's start and its peak? If yes, we showed up.

The census — 1,338 runs, split down the middle

The board is almost perfectly two-sided: **667 up-runs and 671 down-runs**. Median size is 60 points each way; the median run takes 14 minutes to play out. The up-side high was a 490-point rip on 24 June; the down-side high was a brutal 722-point flush on 25 June. A full-session day carries roughly 45 to 65 qualifying runs — this is a market that moves, and it moves both ways in

equal measure. That symmetry matters: it is exactly why the two-sided momentum gates earn on both sides in the shadow board (grind_fast made +\$747 long and +\$552 short this week), and why a long-only or short-shy stance leaves half the pond untouched.


SLICE OF THE TAPE	RUNS	UP	DOWN	MEDIAN PT	BIGGEST PT
All 30 days	1,338	667	671	60	722
40-70 pt (the small stuff)	821	408	413	51	69
70-120 pt (the meat)	349	177	172	88	119
120 pt+ (the giants)	168	82	86	159	722

The biggest runs — and we caught every one of them

Here are the fifteen largest moves of the month. Every single one had a momentum gate fire in its direction inside the move. When MNQ decides to travel, we are on it — the far right tail is not the problem.

#	DAY	TIME UTC	DIR	SIZE PT	MIN	ON IT?
1	06-25	13:24	DOWN	722	30	YES
2	06-24	19:53	UP	490	30	YES
3	07-02	14:13	DOWN	454	30	YES
4	07-08	08:14	DOWN	441	27	YES
5	07-17	13:44	UP	438	26	YES
6	06-17	19:13	DOWN	420	28	YES
7	06-17	17:55	DOWN	362	10	YES
8	06-29	14:16	UP	335	18	YES
9	07-14	12:11	UP	333	20	YES
10	06-23	13:36	UP	310	28	YES

The ten biggest moves, in dollars per contract, are the prize the desk is built to catch. This is what a single run is worth at \$2 a point:



Participation — where the "on it?" rate slips

The show-up rate is not flat. It climbs steadily with run size, which is exactly what you want: the gates are tuned to real thrust, so the more real a move is, the more surely we are in it.

SIZE BUCKET	RUNS	SHOWED UP	SAT OUT	SHOW-UP %
40-70 pt	821	600	221	73%
70-120 pt	349	275	74	79%
120 pt+	168	155	13	92%

Direction barely matters — we catch **76% of up-runs and 78% of down-runs**, so there is no blind side. Almost all of the 308 sat-out runs live in the smallest bucket: **221 of the 308 misses are 40-70 point runs**. Those are the shallow, twitchy moves that are hardest to tell apart from noise in real time — and, not by accident, they are the same shallow moves that the fade and reversion books try to trade and bleed on. Missing most of that bucket is not really a bug. Chasing it is where the losses come from.

The real leak is not a missed run — it is the ground between runs. Of the tradable tape, about **59% of the clock sits inside a qualifying run and 41% sits between them**. That 41% of flat, chopping tape is where `capit_loose` lost **-\$622** at a 9% win rate this week, and where the old V5 `passive_chop` book bled **-\$12,528** over the prior fortnight. We do not need to catch more runs. We need to stop trading the gaps.

★ How we grab it from Monday — by cluster

The giants (120pt+, 92% caught). Nothing to fix on the entry — grind_fast and thrust_cont are already on 155 of 168. The prize here is the exit. The flat chandelier gives back roughly a third of peak on every run, and on high-ATR winners that is \$95–120 handed back per trade. The graded, ATR-conditional trail beat the flat-3.5 stop by **+\$1,653** on grind and **+\$245** on thrust in replay.

KEEP THE ENTRIES, WIDEN THE TRAIL ON HIGH-ATR GIANTS — this is the single biggest dollar lever in the census.

The meat (70–120pt, 79% caught). Solid and left alone. The 74 misses here are worth a look only if a cheap secondary trigger can lift them without dragging in the 40–70 noise. Low priority.

WATCH — NO CHANGE WITHOUT EVIDENCE IT DOESN'T LEAK INTO THE SMALL BUCKET .

The small stuff (40–70pt, 73% caught, 221 misses). Do not chase these. This bucket is indistinguishable from chop at entry time, and every filter that reaches for it drags the desk into the losing fade trades. The honest read is that a 73% catch here is close to the right number, not a shortfall.

KILL THE INSTINCT TO WIDEN GATES FOR THE SMALL RUNS .

The between-run tape (41% of the clock). This is Monday's actual job. Kill the books that only trade the gaps — capit_loose and the reversion/passive family — and let the desk sit flat through chop. The census proves we already catch the moves that pay; every dollar spent trading the 41% of flat tape is a dollar the momentum book has to earn back.

KILL CAPIT_LOOSE AND THE GAP-TRADING REVERSION BOOK; STAND FLAT BETWEEN RUNS .

The through-line is simple and it should be reassuring: the desk shows up for the money. It is on 77% of all runs, 92% of the giants, and 100% of the fifteen biggest moves of the month, both directions, no blind side. The campaign from here is not "catch more" — it is "give back less on the winners and stop bleeding in the gaps."

§3 Feed the thrust, starve the fade

Across two weeks of V5 and the first two days of the V7 board, the desk tells the same story twice: the money is made when we push a move that is already going, and the money is bled when we lean against a move to fade it. Confirmed thrust is the edge. The fade book is the drain. Nothing else in the shadow data is as clean as this, and it is worth saying plainly before we tune a single knob.

The two-week split is not subtle

Taking every filled shadow trade over the fortnight of V5 and sorting it into "pushing a move" versus "fighting a move," the momentum book made about **+\$36,484** over 1,714 trades while the fade-and-reversion book lost **-\$16,540** over 3,039 trades. Narrow it to the cleanest members of each family and the contrast sharpens: the MNQ thrust family alone made **+\$22,337** over 447 trades, and the passive-chop family alone lost **-\$15,967** over 2,931 trades. One family of strategies, on one contract, out-earned the entire fade book's losses. That is the whole desk in two lines.

The single best line in the V5 book was **tw_mnq_thrust_loose**: **+\$5,858** over 98 trades at an 82% win rate. Of that, 81 chandelier exits carried **+\$5,302** and 10 session-flat exits added **+\$802** — only 7 stops cost it **-\$246**. It wins by getting into a real thrust and letting the trailing exit ride it out. It almost never gets stopped, because a confirmed thrust does not turn around and run it over the way a faded rip does.

The passive-chop line is the mirror image. **passive_chop** lost **-\$11,679** over 2,173 trades at a 45% win rate — the largest single drain on the board, and it is a fade. It sells strength and buys weakness into a tape that keeps going, so it is right just under half the time and the times it is wrong it is wrong in the direction of a trend that keeps paying out against it. Add **passive_chop_calm** at **-\$2,382** and **passive_mgc_hours** at **-\$1,681** and the pattern is a family, not an accident.



V5 filled trades, real tick-repriced pnl. Every thrust line pushes right; every passive line pulls left. The board is drawn along one axis — direction of intent relative to the move.

V7 says the same thing on fresh data

The desk cut over to GAZBOT V7 on 2026-07-15, so the current board is only two days deep — thin, and I will not pretend otherwise. But the shape is already the same. The whole V7 board sits at **+\$1,830** over 414 filled trades, and the thrust family accounts for **+\$1,345** of that over 112 trades. The best line is **grind_fast** at **+\$1,300** over 77 trades, followed by **thrust_cont** at **+\$594**. At the bottom of the board sits **capit_loose** — a capitulation-fade — at **-\$623** over 22 trades at a 9% win rate. Two days in, the worst thing on the desk is again a strategy that tries to catch a falling move, and it is right one time in eleven.

V7 STRATEGY	KIND	TRADES	REAL PNL	WIN%
grind_fast	push	77	+\$1,300	47
thrust_cont	push	29	+\$594	52
thrust_loose	push	43	+\$462	44
rg_long_fast	push, forced-long	51	-\$494	22
capit_loose	fade	22	-\$623	9

The thrust exits explain how it earns. In V7, thrust-family trades that ran to their target booked **+\$3,406** across 50 trades, and the 62 that stopped cost **-\$2,061** — a payoff structure where the winners are bigger than the losers even though there are more losers. That is what a real edge looks like: you are wrong often and paid well when right. It is the opposite of the fade book, which is right often and hurt badly when wrong.

Two-sided — and why the short tilt this fortnight is a trap to read wrong

This is the part that is easy to misread, so I want to be careful. On the V7 board the short side made **+\$2,884** over 155 trades at 54%, while the long side lost **-\$1,055** over 259 trades at 30%. It is tempting to conclude "go short, stop going long." That conclusion is wrong, and the data itself says why.

Look at grind_fast, the best strategy on the board. It made money on both sides — **+\$748** long and **+\$552** short. A genuinely two-sided momentum strategy was green in both directions over the same two days. Meanwhile the strategies that were forced long-only — the rg_long family — lost **-\$1,092** across 117 trades at 26%. So the long book did not bleed because "long is bad." It bled because the fortnight was a down tape, and the strategies that could only buy had no way to take the paying side, while the two-sided strategies simply followed the move down and got paid.

The edge is **momentum**, not **direction**. Killing the fade book is not the same as going long-only. This fortnight momentum happened to pay more on the short side because the tape trended down; a two-sided thrust strategy captured that automatically, and a forced-long one could not. If we respond to a bearish fortnight by biasing the desk long, we would have been long into exactly the move that paid the shorts. The correct response to the short tilt is to keep both sides armed and let the move pick the direction.

The census work backs this up from the other end: over 1,338 big MNQ runs the desk showed up for 77% of them, 92% of the 120-point-plus moves, and 100% of the top fifteen. We are not missing the moves — participation is not the problem. The leak is between the moves, in the trades we take when nothing is thrusting, which is precisely where the fade book lives and precisely where it bleeds.

The amplitude floor, and what to actually do

One refinement earns its place rather than just meaning "trade less." In V7 the thrust trades taken when ATR was above 0.07% of price made **+\$1,322** over 36 trades, while the thrust trades in the quieter 0.04–0.07% band netted essentially flat at **+\$24** over 76 trades. Thrust wants amplitude. A thin tape produces thrust signals that do not carry, and the honest read is to stand down on the small ones rather than dress them up. This is in-sample on two days, so I flag it as a lean, not a law — but it agrees with the V5 amplitude-floor work and it costs nothing to respect.

The prescription is not a new gate. It is a reallocation. Feed the thrust book — grind_fast, thrust_cont, thrust_loose — capital and both-sided arming, with the amplitude floor keeping it off thin tape. Starve the fade book — passive_chop, capit_loose, the coil-break fades — because two independent windows of honest fills say it is a structural drain, not a run of bad luck. And do not let a bearish fortnight talk the desk into a long-only bias: the winner was two-sided, and that is the whole point.

KEEP & FEED the confirmed-thrust book, two-sided, amplitude-floored. **KILL** the passive-chop and capitulation fades. **WATCH** the V7 short tilt as regime, not signal — re-check when the tape turns.

§4 The desk only works in one direction — longs bleed, shorts pay, in every family

If you take one thing from this whole board it is this: the desk is a good short-seller wearing a bad long book around its neck. On the current V7 shadow board the 155 short trades made **\$2,884** at a 54% win rate, while the 259 long trades lost **\$1,054** at a 30% win rate. The shorts didn't just carry the board — they paid for the longs' mistakes and still handed you the entire **\$1,829** net. This is not one strategy misbehaving. It is the same lopsided fingerprint in thrust, in the chandelier A/B, in reversion, and on the live desk. It is the single biggest lever we have.


The one number that says it all: 68% of longs die on a stop, only 38% of shorts do

I split every filled V7 shadow trade by how it ended, and by side. The exit tape is brutal and it is one-sided. Longs get stopped out: 175 of the 259 longs — that's 68% of them — ended on a STOP for a combined **\$5,707** loss, and only 38 longs (15%) ever reached TARGET. Shorts do the opposite: 59 of 155 shorts (38%) walked into TARGET for **\$3,489**, and their stop pile was a third the size of the longs'. The chandelier trail actually helps the longs (+\$2,949) and the targets help too, but it isn't enough to dig out from under that stop pile. Shorts win because the trades they open keep going their way; longs lose because the trades they open turn around and hit the stop.

EXIT REASON	SIDE	TRADES	REAL P&L
STOP	LONG	175	-\$5,707
STOP	SHORT	64	-\$2,133
TARGET	LONG	38	+\$1,910
TARGET	SHORT	59	+\$3,489
CHANDELIER	LONG	43	+\$2,950
CHANDELIER	SHORT	32	+\$1,528

Every family shows the same tilt — even the ones that win overall

I grouped the sims into their families and split each one long-versus-short. The result is almost eerie in how consistent it is. The pure directional and reversion families are catastrophically long-heavy losers and short-heavy winners; the two families that survive on both sides (grind and the absorption-veto) are exactly the two that are built to be bidirectional in the first place. Nothing here is a fluke of one sim.



Read the thrust family line by line and the story writes itself: thrust_loose was **-\$38** long but **+\$500** short; thrust_aligned **-\$99** long, **+\$293** short; thrust_fast **-\$132** long, **+\$227** short. The one thrust sim that stayed green on both sides, thrust_cont (+\$161 long / +\$433 short), only did so because its long entry is the tightest and most trend-confirmed of the bunch — and even it made

nearly three times as much short. The chandelier A/B is the cleanest controlled experiment on the board because the three sims differ only in trail width: every one of chand_k20, k25 and k35 lost money long and made money short. Same exit, same everything, opposite outcome by direction.

The long bleed has a name: reversion firing counter-trend

Where is the actual money going? 135 of the losing longs live in the reversion / mean-reversion family (rg_long_fast, rg_long_fast_v, the rg_long_20 variants, capit_*), and together those long-only reversion sims lost **\$1,563**.

That is the leak, and it is a design leak, not bad luck. These sims are structurally long-only — they buy dips — and buying dips on this instrument means standing in front of moves that keep going. The tell is in the entry conditions: longs on this board fire at an average ATR of 0.064% of price versus 0.058% for shorts, so the long book is systematically stepping in when the tape is faster and more violent, i.e. trying to catch rips that then run it over into the stop. The clock confirms it — during the US afternoon (16:00–18:00 UTC) the long book bled **\$1,735** across those three hours while shorts in the same window made money. The desk buys the panic and gets run over.

The live desk is already living this

This isn't a shadow-only curiosity. The 2-day live V7 desk, which is thrust-only, has taken 24 longs for **-\$733** at a 17% win rate against 25 shorts for **+\$46** at 44%. The exit split is the shadow board in miniature: live CHANDELIER made **\$110** long but **\$457** short, and the live STOP pile was **-\$588** long versus **-\$221** short. Real money, real fills, same tilt. And the two-week V5 history rhymes underneath it — 3,473 longs for **-\$1,316** versus 3,216 shorts for **-\$164**, both at a 44% win rate. The long side has been the heavier loser for the entire through-line; V7 just made the gap loud.

This is NOT "we're missing the trends"

The census kills the easy excuse. Across the 1,338 big MNQ runs we participated in 77% of them, caught 92% of every 120-point-plus move, and were present for 100% of the top-15 runs of the period. We are not sitting out the action. The leak is not *between* the trends we miss — it is the long branch firing *counter-trend*, into rips and afternoon reversals, and then feeding the stop pile. We show up. We just show up on the wrong side of the tape when we go long.

What to actually do — gate the long side, don't kill it

The temptation is to switch off longs entirely, and the raw numbers almost justify it. I'd stop short of that, for two honest reasons. First, grind_fast long (+\$748) and abs_veto long (+\$307) prove a well-conditioned long can win — the problem is the entry condition on the reversion book, not the direction itself. Second, this is one week of V7; a one-side-off rule dressed up as an edge is usually just "trade less," and I'd rather fix the branch than amputate it. So the prescriptions, in order of conviction:

1. Condition-gate the long side — the circuit breaker is sitting

dormant. The cb_thrust and cb_thrust_dropped sims took zero trades all week; the long-side circuit breaker exists in the code and never armed. That is the single highest-value switch on the board. Wire it to veto reversion longs when the tape is trending against them (the afternoon 16:00–18:00 window and the high-ATR 0.064%+ entries are where the damage concentrates). If it had simply suppressed the reversion-long book this week it would have saved on the order of the **\$1,563** that family lost.

2. Reconsider the killed absorption-veto — but only for the long side.

abs_veto was one of only two families green on both sides (long +\$307, short +\$763). It was on the chopping block, but the through-line reframes it: an absorption-aware veto is exactly the tool for the "buying into a rip that gets absorbed and reverses" failure mode. Keep it, and specifically point it at the long entries.

3. Leave the short book alone. It is the profit engine. Every hour you spend "balancing" the desk toward longs is an hour spent diluting the one thing that works. The right shape of this desk, on this evidence, is an unconstrained short book plus a tightly-gated, trend-confirmed long book (grind- and thrust_cont-style), and nothing that buys dips blind.

KILL THE BLIND LONG BOOK the long-only reversion sims (rg_long_*, capit_*) that lost **\$1,563** firing counter-trend. **KEEP & ARM THE CIRCUIT BREAKER** to condition-gate longs — currently dormant at zero trades, the #1 lever. **WATCH THE SHORT BOOK RUN** — it's the engine, don't touch it.

§5 **The chandelier is the whole edge — and it hands back a third of every winner**

If you want to know where the desk actually makes its money, look at how trades end, not how they start. On the live V7 desk the trailing chandelier is the only exit that has ever printed green — **+\$567** over 13 trades at an 85% win rate — while every stop, adverse cut and absorption cut has bled. The shadow board tells the same story with more trades behind it. So this week's through-line is simple: the exit architecture is carrying the desk, the chandelier is the heart of it, and it has one fixable leak — it gives back about a third of every winner's peak, and the dollar size of that give-back scales straight with volatility.

Where the money is actually booked

Across the 28-sim V7 shadow board, sorting every filled trade by how it exited splits the desk cleanly into a profit half and a loss half. The two trailing/target exits book real money; the stop is where it dies.



Two honest caveats before we lean on this. First, on the shadow board the fixed profit **TARGET** actually books slightly more than the chandelier (**+\$5,399** vs **+\$4,477**) — so I won't overclaim that the chandelier is the only profit centre; it's the profit centre on the live desk, where the sims run thrust-only and never hit a fixed target. Second, the **-\$7,840** stop pile is not a give-back problem at all. I checked: the 239 stopped trades had an average peak of **-\$50.7** and not one of them ever showed even \$20 of green before it died. Those are bad entries, full stop. The stop is doing its job. This section is about the winners.

The leak: a third of the peak, every time

Take every filled chandelier winner and compare the best unrealised profit it ever showed (the ceiling) against what actually landed in the book. The gap is the give-back — the money the trail let slip on the way out.

CHANDELIER WINNERS (V7 SHADOW)	N	AVG PEAK	AVG BOOKED	GIVEN BACK
All chandelier wins	74	\$91.0	\$60.4	\$30.6

Ninety-one dollars of peak, sixty booked. That's a **34% give-back** — and it isn't a one-trade fluke. Summed across all chandelier winners the desk showed **\$6,732** of peak profit and kept **\$4,470**, handing back **\$2,262**, which is 0.34 of the total. A third of the peak, in aggregate. When the operator watches a live winner tick to +\$224 and settle at +\$88, that feels like a bug. It isn't — it's the norm the trail was built to produce. The only question is whether a third is the right third to give back, and the answer depends entirely on how fast the instrument is moving.

The give-back scales with volatility — that's the whole prescription

Here's the finding that matters. Bucket every filled winner by its entry ATR-as-percent-of-price and the give-back fraction barely moves — it sits near a third whether the tape is dead quiet or ripping. But the give-back in **dollars** roughly doubles from the quiet bucket to the fast one, because a third of a big peak is a big number.

ATR BUCKET	N	ATR%	AVG PEAK	AVG BOOKED	GIVEN BACK	GIVE-BACK %
Low vol	13	0.028	\$49.8	\$34.6	\$15.2	31%
Mid vol	104	0.053	\$68.6	\$45.4	\$23.2	34%
High vol	54	0.097	\$124.0	\$86.9	\$37.1	30%

Read the last two columns together. The percentage is flat — 31%, 34%, 30% — so a single flat trail is proportionally fair everywhere. But the dollars run \$15 → \$23 → \$37. On the fast tape a high-ATR winner peaks at **\$124** and the trail lets \$37 of it walk out the door. That \$95-to-\$124 peak giving back into the \$85-\$90 range is exactly the operator's +\$224→+\$88 experience, just averaged. A flat-width trail treats a 3-point wiggle in a quiet MNQ the same as a 12-point wiggle in a fast one: in the quiet tape it clips runners early for no reason, and in the fast tape it's still parked a fixed distance back while the price swings a fortune through it.

What the A/B actually says about trail width

The board ran three flat chandelier widths this week, and the raw scores lean tighter-is-better: **k20** at **+\$18.5**, **k25** at **-\$36.5**, **k35** at **-\$101**. But I'm not going to hang a decision on that — it's 9 to 11 trades per arm in a single chop week, and all three are inside coin-flip range on win rate (33-45%). What it

does confirm is the mechanism: the widest arm (k35) is the one that gave back positive dollars on average and scored worst, i.e. a wide flat trail sits too far back and bleeds. That's consistent with the 25-day archive replay, where a fixed 3.5-ATR trail was beaten by an ATR-conditional graded trail by **+\$245** on the thrust family and **+\$1,653** on grind. Same lesson from two directions: the answer is not "pick a tighter number," it's "stop picking one number at all."

The prescription – a vol-adaptive chandelier. Set the trail distance as a function of live ATR, not a fixed point/percent width. Tighten hard in the low-vol bucket (a quiet winner that peaks at \$50 shouldn't need \$15 of room – that's clipping runners for nothing) and give a fast-tape winner proportionally more slack so a normal high-ATR swing doesn't knock it out at the first wobble. The give-back fraction is already a stable third; the goal is to make that third the *deliberate* third rather than an accident of whichever flat width was hard-coded. The census backs the upside: we show up for 77% of big MNQ runs and 100% of the top-15, so the runners exist to be captured – the leak is purely in how we ride them out.

The verdict

The exit spine is the desk's real product this week, and it's working: the trailing chandelier books green on the live desk when nothing else does, and on the shadow board it and the fixed target between them carry roughly \$10k of profit against the stop pile. The one honest leak is the give-back, and it is not a mystery — it's a flat third of peak that turns into real dollars precisely when the tape is fast. A flat trail can't win that fight in both regimes at once.

KEEP THE CHANDELIER — it is the profit centre, live and shadow. **WATCH THE FLAT-WIDTH A/B** — tighter won this week but on 9-11 trades in one chop regime; don't promote a single number off it. **BUILD THE VOL-ADAPTIVE TRAIL** — the give-back scaling with ATR (\$15→\$23→\$37 across the buckets, same 30-34% every time) plus the archive replay's +\$245/+\$1,653 edge for a graded trail is the strongest, most repeatable exit finding on the board. Make the trail a function of ATR, not a constant.

§6 The order book: our honest look at whether the wall tells us anything

For weeks the standing hope has been that the ten levels of resting liquidity on either side of the market would tell us where price is about to go — that a fat wall of bids means "buy", a fat wall of offers means "sell it's about to get rejected". We

finally have enough book to answer it properly, and the answer is the disciplined one: on our home contract the book tells us almost nothing, on the other two it whispers something real but smaller than a single tick, and every folk-theory about "fading the wall" comes out backwards. This is a keep-capturing, don't-trade-on-it chapter.

Here is what we are actually sitting on. The two-week depth store (depth.db, ten levels deep, a fresh snapshot every 250 milliseconds) holds 1,812,648 book snapshots on MNQ, 1,142,749 on MES and 1,073,720 on MGC across July 9-17. M2K has only 10,820 — a single Friday afternoon — so it doesn't get a vote. On top of that the new V7 capture (capture.db) has been recording MNQ since the 15th and already holds 35.2 million book rows across 3.16 million distinct snapshots. So this is not a thin-n complaint about the book itself. There is plenty of book. The problem is what's in it.

Test one: does a lopsided book predict the next move? Barely, and not where we trade.


I sampled each contract down to roughly one reading per second (to stop the 250ms snapshots from flattering themselves with autocorrelation), measured the resting-size imbalance between bids and offers, and asked a plain question: five seconds later, which way did the mid-price actually go? The correlation between the two is the whole ballgame, and it is tiny everywhere.

CONTRACT	BOOK SNAPS	CORRELATION OF IMBALANCE TO NEXT 5S MOVE		READ
MNQ (home)	1,812,648	+0.0024	indistinguishable from zero	
MES	1,142,749	+0.0234	weakly real, still tiny	
MGC	1,073,720	+0.0106	faint	

A correlation of +0.0024 on MNQ is a coin that has heard a rumour about which way it might land. MES at +0.0234 is the strongest of the three and it is genuinely non-zero, but 0.02 is still a signal buried under a mountain of noise. Widening from the top-of-book to the full ten levels didn't move any of these numbers — MNQ stayed at +0.0024, MES at +0.0234 — so stacking deeper levels of resting size buys us nothing. The wall behind the wall is just as mute as the front one.

How far does price actually travel when the book leans hardest?

Correlations are abstract, so here's the concrete version. I took the ten percent of readings where the book was most bid-heavy (the biggest "buy wall") and the ten percent most ask-heavy, and measured the average move over the next five seconds, in ticks. If the wall meant anything tradeable, the strongest signals should move price a meaningful fraction of a tick.



That's the whole story in one picture. The very strongest book lean we can find moves MNQ about eight hundredths of a tick, MES twelve, MGC six. Crossing the spread to act on it costs a full tick or more. The signal is worth roughly a quarter of the toll you'd pay to trade it. The direction is at least honest — bid-heavy nudges price up, ask-heavy nudges it down — but the size is nowhere near the cost of admission.

Test two: "fade the wall" is backwards.

The seductive theory is contrarian: a giant offer overhead is a ceiling, so sell into it. The data says the opposite. When I split by direction, the sign of the relationship is positive on all three contracts — a heavier bid stack goes with an up-move, not against it. The book is momentum-aligned, not a wall to fade. On MES the down-side of this was the sharpest: in the moments the offer stack was heaviest, price continued down 70.4% of the time — continuation, not rejection. If you'd been fading that offer wall you'd have been run over seven times in ten. On MNQ the same bid-heavy readings actually drifted slightly negative on average (–0.39 tick), which is the market quietly mocking anyone who buys the buy-wall. Whatever the book is, it is not a fade signal.

Test three: absorption is a coin flip, not a reversal.

The last hope was absorption — price pushes up into a big resting offer, the offer soaks it up and holds, and the theory says that stalled push snaps back. I

built the event directly: a recent up-push running into a wall three times the median size, then measured whether price continued through or reversed. Across 7,373 such events on MNQ it continued 47.2% of the time and reversed 52.8%; MES 42.7% / 57.3% across 4,310; MGC 47.2% / 52.8% across 6,715. So there is a faint reversal lean — but 53-57% at the five-second horizon is exactly what plain mean-reversion gives you after any sharp push, wall or no wall. The wall adds nothing you couldn't get from the price move alone, and the edge sits under the spread again. "Absorption reverses" is, at best, a coin flip wearing a lab coat.

The three folk-theories, scored: **imbalance predicts direction** — true but a quarter of a tick, buried under a full-tick spread. **Fade the wall** — refuted, the book is momentum-aligned, and fading the MES offer wall loses 70% of the time. **Absorption reverses** — a coin flip (47-53%) indistinguishable from ordinary post-push reversion. None of the three clears the cost of trading it.

Why MNQ — the one we actually trade — is the deafest of the three

This is the awkward part. The book is least informative exactly where we live. MNQ's +0.0024 is the weakest of the three; MES and MGC, which we barely touch, carry what little signal exists. That fits what we already know about MNQ: it's fast, thin at each level, and the resting size churns so quickly that a "wall" is gone before you can react to it. The e-mini (MES) is deeper and slower, so its book holds its shape long enough to mean something — but "something" here is two hundredths of a correlation and a tenth of a tick. It is a tilt, not a trigger.

What the book is actually good for

Not as an entry — on no contract, in no configuration, does the book clear the spread as a standalone reason to fire. That door is closed and I'd stop knocking on it. Where it has a plausible future is as a small exit or veto tilt on MES only: when we're already long and the offer stack ahead is heavy and holding, that 70% continuation read is at least pointing the right way, and it costs nothing to let a runner breathe a beat longer. That's a nudge on a position we already hold, not a new bet — which is the only place a quarter-tick signal can pay, because we're not crossing the spread to use it.

The other honest note: the V7 book is three days old. All the numbers above lean on the two-week depth.db history; the 35-million-row V7 capture has

barely started, and it only covers MNQ so far — the one contract where the book is deafest. Before we build anything, even the MES exit-tilt, we need the V7 book and tape running together across MES and MGC for a couple of weeks so we're testing the world the live desk actually trades in, not last fortnight's.

Verdict. **NO STANDALONE ENTRY** on any contract — the strongest book lean moves price ~0.06–0.12 tick against a ≥ 1 -tick spread.

"FADE THE WALL" REFUTED — the book is momentum-aligned; fading the MES offer wall lost 70% of the time. **KEEP AS AN MES EXIT/VETO TILT ONLY**, and


KEEP CAPTURING — extend V7 book+tape to MES/MGC before we trust anything built on it. The book is real, it's just quieter than the spread.

§7 The OFF Switch: Trade the Condition, Not the Gate

The through-line under everything else this week is a single idea, and it is not "find a better gate." It is: **every gate we own is right some of the time and wrong the rest of the time, and the whole game is turning the right one ON and the wrong one OFF for the condition in front of us.** The board proves it cleanly. When a gate fires into the condition it was built for, it makes money on both sides. When it fires into the opposite condition, it bleeds in every hour of the day, no exceptions. That is not a tuning problem. That is a switching problem.

The board sorts itself by condition-fit, not by cleverness

Score the whole V7 shadow board honestly on **real_pnl** where the fill actually happened and it comes to **+\$1,829** across 414 trades. But that headline hides the real shape. Sort the sims into two piles — the ones that fire with the tape and the ones that fire against it — and the piles do not overlap. The continuation gates (confirmed thrust, the absorption-veto family, grind) are green. The reversal gates that fire into a running move (capitulation, the long-only reversal-grabs) are red. Same market, same week, same fills. The only thing that changed is whether the gate was pointed the right way for the condition.



Read the chart top to bottom and it is the OFF switch drawn out. The top two bars are gates fired into their condition. The bottom three are gates fired against it. Nothing in the middle earns its keep.

Capitulation is the clean proof: it is never the hour, it is the direction

The single most instructive sim on the board is **capit_loose** — the loosened capitulation buyer, meant to grab the snap-back after a flush. It went - **\$622.50** over 22 trades and won **9% of the time**, the worst hit-rate on the board by a distance. Both sides bled (long legs -\$371.50, short legs -\$251). The tempting read is "it fired at bad times." It did not. Bucket it by the hour it entered and every single hour is red:

ENTRY HOUR (UTC)	TRADES	CAPIT_LOOSE P&L	ABS_VETO_55S SAME WINDOW
00-03	11	-\$220	+\$94
05-09	5	-\$159	+\$303
13-16	3	-\$160	+\$46
22-23	3	-\$84	-\$32

There is no good hour to run it. The exit tape says why: **18 of its 22 trades ended on the STOP** for -\$606, and even the 4 that reached TARGET netted -\$16.50. When a "buy the flush" gate gets stopped out four times in five, the flush was not a flush — it was the middle of a continuation, and capitulation was standing in front of the train. This is the definition of a gate firing into the wrong condition. You cannot tune the entry price to fix pointing the wrong way.

The reversal-grabs bleed for the same reason, and they are structurally one-sided

The reversal-grab cluster (**rg_long_fast**, **rg_long_fast_v**, the **rg_long_20_*** variants) is **long-only** by construction — it only ever tries to catch a falling market and buy the bounce. As a family it lost **-\$1,091.50** across 117 trades. The fast variant alone was -\$493.50 at a 22% win rate, and its exit tape is the tell again: **39 STOP exits for -\$1,218**, rescued only partly by 12 chandelier runners that clawed back +\$724.50. It is the capitulation story in a different costume — a gate that only knows how to fade, firing into a tape that keeps going. The one short-side reversal-grab that fired all week (**rg_short_025_raw**) went +\$48.50 on its single trade; too thin to bank, but it is not a coincidence that the only green reversal-grab was the one pointed the other way.

The gates that show up right win on BOTH sides

Now the other pile. The **absorption-veto family** — which does not predict anything, it just declines to enter when the resting book is absorbing the push — is green on both long and short, at all three window lengths:

GATE	LONG	SHORT	TOTAL	WIN%
abs_veto_50s	+\$87.50	+\$356	+\$443.50	55%
abs_veto_55s	+\$143	+\$356	+\$499	57%
abs_veto_60s	+\$76.50	+\$50.50	+\$127	67%

The whole **abs_veto** family banked **+\$1,070**. A gate that is green on both sides is a gate that is reading the condition, not guessing a direction. Confirmed thrust does the same job from the momentum side: **thrust_cont** made **+\$594**, and it is honest about which side is the good one — its short legs (firing with a down-thrust) made +\$433 at 63%, its long legs +\$161 at 38%. Across the whole thrust family the split is stark: **short thrust +\$1,454, long thrust -\$109**. The desk's edge this week was momentum-to-the-downside plus a book-aware veto. Everything that faded a live move gave the money back.

The pattern in one line: gates that fire *with* the condition (short thrust, absorption-veto, grind) made roughly +\$3,000 between them; gates that fire *against* it (capitulation, long reversal-grabs, long thrust) lost roughly -\$1,800. The net +\$1,829 board is that subtraction. Turn the second pile OFF and the board is a different animal.

This is the same lesson the two-week V5 history already taught

None of this is a one-week fluke of framing. Run the same split across the prior fortnight of V5 shadow history and it is the identical shape at 10× the sample: the momentum/thrust/grind family made **+\$23,586** over 1,055 fills while the reversion/fade/capitulation family lost **-\$15,884** over 2,938 fills. Two weeks, two engines, one verdict — the desk gets paid for going with force and gives it back for fading force. The V7 board is the V5 lesson in miniature. That through-line is about as validated as anything we have.

The honest part: the day-type ROUTER is not yet an edge

The obvious next move is to build the switch itself — a day-type router that reads the morning's character and flips reversion ON for a rotational day, momentum ON for a trend day. I want to be straight with you: **we tried to prove that router is an edge and it did not survive.** In-sample it looked great, turning a +\$1,974 stretch into what read like a clean regime call — but split out of sample it collapsed to +\$199, and a dumb "always run the reversion book" control beat it. The grind engine on its own, replayed across 25 archive days, lost -\$7,975 (green on only 7 of 25). The reversion outlier that carried a couple of days (+\$723) was one good session away from being flat. So the router, as a predictor, is refuted for now.

But — and this is the whole point of the section — the refutation is of prediction, not of switching. We already killed regime prediction weeks ago; the market does not tell you at 09:30 what kind of day it will be. What the board above shows works is reactive condition-gating: the absorption-veto does not forecast anything, it reads the book as the trade is forming and declines. Confirmed thrust does not predict a trend, it waits for the thrust to be there. That reactive read is the version of the OFF switch that is actually holding up. The forward-looking day-type router is the version that keeps failing.

Where this leaves us. The direction is right and the mechanism is proven in-week and across the fortnight: fit the gate to the condition and both sides go green; misfit it and it bleeds in every hour. The build order that follows from the evidence is (1) turn OFF the fade-into-force gates outright — capitulation and the long-only reversal-grabs earn a **KILL**, they are pointed the wrong way by construction; (2) keep the reactive condition-readers that win both sides — the absorption-veto family and confirmed thrust earn a **KEEP**; (3) keep researching the day-type router but do not arm it — as a forward predictor it is **REFUTED, KEEP DIGGING**. The switch we want is the reactive one we can already see working, not the fortune-teller the skeptic just shot down.

§8 The sizing ladder and the departure lounge: who earned more size, who's booked for the exit

Short version up front, because it's the honest one: nothing on the V7 board earned a bigger position this week, and it can't have, because the board is exactly two days old. The ladder we agreed to needs eight-to-twelve weeks, a hundred-to-two-hundred trades, and an edge that holds across regimes before a sim moves from shadow to one live micro, let alone to two or three. Two days and seventy-seven trades on the best-populated sim is a rounding error against that bar. So this section does the other useful job: it names who's climbing the first rung, and it walks the losers to the departure lounge so you can decide who gets flattened off the board.

The ladder, and why the top rung is still empty

The rule we set is a promotion ladder, not a leaderboard. A sim goes shadow at one micro, then live at one, then two, then three, and each step has to be earned: eight-to-twelve weeks of runtime, a hundred-to-two-hundred closed trades, positive expectancy in more than one regime, a drawdown you can stomach, low correlation to what's already live, and a clean reward-to-risk. The report nominates; you arm. The point of the gate is to stop a two-day hot streak from talking itself onto the book.

Here's where the current board actually sits against that gate. The whole V7 shadow board spans two calendar days, the 16th and 17th of July. The busiest sim on it, **grind_fast**, has run seventy-seven trades — impressive for forty-eight hours, but it's still under the hundred-trade floor and it's nine-plus weeks

short on runtime. Every other name is thinner: thrust_cont at twenty-nine, the absorption-veto pair at twenty-one and twenty-two, thrust_loose at forty-three. Not one sim has seen a genuine second regime — this was one tape, twice. On the runtime-and-count axis alone, the top rung of the ladder is empty and stays empty. That's not pessimism, it's the gate doing its job.

The honest read on a two-day board: a single week — never mind a single week that's really one repeated tape — mostly just tells you "trade less," not "trade bigger." Anyone who sizes up on this data is sizing up on noise. The frontrunners below are worth *watching*, nothing more.

Who's on the first rung — the watch list


If you set the money aside and just ask "who's behaving like a candidate," three names stand out, and the reason they stand out is the same reason they're only on rung one: they win on both sides of the tape, which is the correlation-and-robustness property the ladder actually cares about.

SIM	TRADES	REAL P&L	WIN%	LONG / SHORT	LADDER READ
grind_fast	77	+\$1,299.5	47%	+\$747.5 / +\$552.0	RUNG 1 – WATCH
thrust_cont	29	+\$594.0	52%	+\$161.0 / +\$433.0	RUNG 1 – WATCH
abs_veto_55s	21	+\$499.0	57%	+\$143.0 / +\$356.0	RUNG 1 – WATCH
abs_veto_50s	22	+\$443.5	55%	+\$87.5 / +\$356.0	RUNG 1 – WATCH
thrust_loose	43	+\$462.0	44%	-\$38.5 / +\$500.5	RUNG 1 – THIN LONG

grind_fast is the only one that clears both a real trade count and a genuine two-sided profit — it made \$1,299.50 over seventy-seven trades, and it did it \$747.50 long and \$552.00 short, so the money isn't coming from one lucky direction. But here's the discipline check, and it matters: the 25-day archive replay had grind_fast at **minus \$7,975 across the month, green on only seven of twenty-five days**. Two green days do not overturn a red month. grind_fast is the poster child for "watch it, don't arm it" — the exact sim a two-day gate would wrongly promote and a proper gate correctly benches. thrust_cont is cleaner in spirit: it took its profit thirteen times through the TARGET exit for \$1,008.50 and gave back \$414.50 through stops, and it's the family that survived the two-week V5 history too. That through-line — thrust green, reversion red — is the only thing on this board with more than one week behind it.

The board, top to bottom

One picture of the whole V7 board, best to worst on real filled P&L, so you can see how far the frontrunners sit above the departure lounge:



The departure lounge — nominated for the exit

These are the names I'd walk off the board. Two clusters, plus one clerical embarrassment.

The capitulation family. capit_loose is the worst sim on the whole board: it lost \$622.50 over twenty-two trades and won only nine percent of them. The autopsy is one line — it took the TARGET four times for a small loss and got STOPPED eighteen times for \$606, which is what "buying the panic and getting run over" looks like as a P&L. Its siblings are the same shape with fewer samples: capit_ride went two-for-two the wrong way at zero percent for minus \$61, and capit_mid was minus \$39 at thirty-three percent. There is no version of the ladder where a nine-percent-win sim advances. The capitulation idea isn't half-formed here, it's actively bleeding on the current tape.

The long-only rg_long_* cluster. Four names, all one-directional, all red: rg_long_fast at minus \$493.50 (twenty-two percent win), rg_long_fast_v at minus \$243.50, rg_long_20_t20 at minus \$203.50, and the rg_long_20 pair at minus \$75.50 each. The common thread is the one the ladder explicitly penalises — they only trade long, so they can't diversify the book and they get taken apart whenever the tape isn't trending their one way. rg_long_fast's exit split says it plainly: it caught twelve chandelier winners for \$724.50 and then handed back \$1,218 across thirty-nine stops. That's not an edge with a bad week, that's a directional bet paying the wrong side of a coin. Note the one flicker of life — rg_short_025_raw, the mirror-image short, is +\$48.50, but it's a single trade, so it means nothing yet; if anything it hints the long-only framing is the problem, not the setup.

The duplicate. rg_long_20_t15 and rg_long_20_flow25 are the same sim wearing two badges. I pulled them trade-by-trade and they are byte-identical — same entries, same eleven fills, same minus \$75.50, down to every individual trade (\$17.5, −\$23.5, −\$53.5, +\$41.5...). Whatever knob was supposed to differ between the "t15" and "flow25" variants isn't wired to anything. That's a wasted board slot and a false sense of two data points where there's one. Kill one of them regardless of what you do with the family.

DEPARTURE LOUNGE	TRADES	REAL P&L	WIN%	WHY IT'S BOOKED
capit_loose	22	−\$622.5	9%	Worst on board; 18 stops for −\$606
rg_long_fast	51	−\$493.5	22%	Long-only; 39 stops undo 12 winners
rg_long_fast_v	35	−\$243.5	23%	Long-only, same shape
rg_long_20_t20	9	−\$203.5	22%	Long-only, thin
rg_long_20_t15 / _flow25	11 ea	−\$75.5 ea	45%	Byte-identical duplicate — kill one
capit_ride	2	−\$61.0	0%	0-for-2 capitulation
capit_mid	3	−\$39.0	33%	Capitulation, thin

The two-week V5 history backs the lounge nominations rather than fighting them: every reversion-shaped family over there bled too — passive_chop alone was minus \$11,679 across 2,173 trades, passive_chop_calm minus \$2,382, passive_mgc_hours minus \$1,099. Fading and buying-the-panic has now lost money on both the old board and the new one. That's the one cross-regime signal we do have, and it points the same way it always has.

What the report nominates (you arm): Promote nobody — the ladder's runtime and trade-count floors aren't met by anyone, and grind_fast's minus-\$7,975 month is the cautionary tale for why that gate exists. Keep grind_fast, thrust_cont and the absorption-veto pair on rung one as **WATCH**, purely because they win two-sided. Walk the capitulation family (capit_loose / capit_ride / capit_mid) and the long-only rg_long_* cluster to the **DEPARTURE LOUNGE**, and delete one half of the rg_long_20 duplicate today — that one's free.

§9 My Mid-Week Musings, Walked to a Verdict

Every idea I floated to you this week, tested against the money and walked all the way to a yes or a no. Nothing got quietly dropped and nothing got armed on the strength of one good week. Eight leads, eight verdicts. Some held, most of the exciting ones didn't, and I'll tell you plainly which is which.


A word on how I scored these before we start. Every dollar below is `real_pnl` on trades that actually filled — the honest number, not the ceiling that flatters losers by anywhere from 1.1 to 3.1 times. And where a musing lived or died on more than this week, I pulled the two weeks of V5 history and the 25-day archive replay so you're not looking at a single lucky Tuesday dressed up as an edge.

1. The chandelier is too loose — HELD, and the fix is ATR-conditional

This was the week's cleanest result. I ran three chandeliers side by side at trail multiples of 20, 25 and 35, same entries, only the trail differs. Tighter won at every step:

THE K20 TRAIL MADE \$18.5 OVER 11 TRADES, K25 LOST \$36.5, AND THE LOOSEST K35 LOST \$101 OVER 9 .

That's a monotonic line — every notch you loosen the trail, you give back more. It confirms what the give-back study kept saying: at any fixed width the runner hands back roughly a third of its peak, and the loose trail just lets it hand back more before the exit fires.



But here's the honest catch, and it's why the verdict is "held" and not "armed": even the tightest of the three barely scraped positive on 11 trades. Flat trailing at any width is a losing game on this sample — tighter just loses less. The real fix isn't picking a number, it's making the trail read the tape: a graded, ATR-conditional trail that rides wide when the move is amplitude-rich and clamps down when it isn't. The deep pass measured that ATR-conditional

trail beating flat-3.5 by about \$245 on thrust and \$1,653 on grind. So the musing holds — the chandelier is too loose — but the answer is a smarter trail, not a smaller constant.

2. A vol-adaptive chandelier — promising, still out-of-sample-pending

This is the natural next step from the one above, and I want to be careful not to oversell it. The in-sample story is genuinely good: grade the trail off ATR and it beats the flat trail by four figures on grind. But that number is fit on the same weeks it's measured on, and you've taught me the hard way that a single-week filter usually just means "trade less" in a flattering disguise. I have not yet run it forward on unseen days. Until I do, this is a promising lead with a pending OOS test, not a result. I'd rather tell you "I don't know yet" than hand you an in-sample curve and call it an edge.

3. Direction × regime gating — HELD, and it's the number-one lever

This is the big one, and it's the same finding from two completely independent angles, which is why I trust it. The whole board splits on which side a strategy is allowed to take in which regime. Look at the two cleanest cases on the V7 board this week:

SIM	LONG SIDE	SHORT SIDE	READ
grind_fast	+\$747.5 (34t)	+\$552.0 (43t)	both sides green — genuinely two-sided
thrust_loose	-\$38.5 (20t)	+\$500.5 (23t)	all the money is on the short side

thrust_loose made \$462 net this week but every dollar of it — and then some — came from shorts; its longs quietly bled. If you'd gated that sim to shorts-only in the prevailing regime you'd have kept the \$500 and skipped the drag. Now widen the lens to the two weeks of V5 history and the same law holds at scale: **MOMENTUM MADE \$23,586 OVER 1,055 TRADES WHILE REVERSION LOST \$15,884 OVER 2,938** . The edge was never "which strategy" — it was "which direction, in which condition." Right idea in the wrong regime is how the desk bleeds between the big runs. This holds, and it's the lever worth building around.

4. A per-side circuit breaker — built, but DORMANT

I built this to catch the exact failure above: if one side of a sim is bleeding, trip a breaker and stop taking that side. It's on the board as `cb_thrust` and `cb_thrust_dropped`. The problem is simple and a little embarrassing —

BOTH BREAKER SIMS TOOK ZERO TRADES THIS WEEK. The trip conditions never armed, so I have exactly no evidence it works. It's not refuted, it's untested-in-the-wild. Either the thresholds are set too tight to ever fire, or this week just didn't hand it the losing streak it's built to catch. Next week I loosen the arming logic enough to actually see it act before I say a word about whether it helps.

5. grind_fast, live if it survives — REFUTED

This one hurt because `grind_fast` was the week's best sim: **+\$1,299.5** over 77 trades, and beautifully two-sided as we just saw. The tempting story writes itself — promote it. So I ran it back over the 25-day archive to see if it survives out of sample, and it does not:

GRIND_FAST LOST \$7,975 OVER 25 DAYS, GREEN ON ONLY 7 OF THEM. This week's win was a good week, not a durable edge. You can even see the fragility inside this week's own trades — its chandelier exits made \$2,606 while its stops gave back \$1,307, so the whole thing lives or dies on how often price lets the trail ride instead of hitting the stop. On a chop-heavy stretch that ratio flips and it bleeds. Refuted as a standalone live candidate. Its value is as a two-sided signal to fold into a regime gate, not as a strategy to arm.

6. Reversion #23, the chop specialist — outlier-carried, not proven

The reversion-in-chop idea keeps flickering because one variant, `rg_long_fast_v`, showed +\$723 over the archive. But when I strip its single best day out, that \$723 collapses toward -\$209 — the whole edge is one outlier session carrying the average. On this week's board the reversion-grid family was uniformly poor: **rg_long_fast lost \$493.5**, `rg_long_fast_v` lost \$243.5, and the `rg_long_20` variants sat between -\$75 and -\$203. A specialist whose edge disappears the moment you remove its best day isn't a specialist, it's a lucky draw. Not proven. I'm not killing the chop-reversion thesis outright — but it needs a real trend-free week to earn a look, and this wasn't it.

7. A day-type router — REFUTED as an edge


The dream here was a switch: run momentum on trend days, reversion on chop days, best of both worlds. The skeptic pass took it apart. In-sample the router looked like +\$1,974; split out of sample it falls to +\$199, and — this is the killer — **IT LOSES TO JUST ALWAYS RUNNING REVERSION'S BETTER COUSIN**. The router's apparent edge was hindsight about which day you were in, and you can't buy hindsight at the open. Once you have to classify the day in real time, the advantage evaporates. Refuted. The regime read is real (see musing 3), but routing whole strategies on a day-label is not the way to cash it — the gate belongs on direction, per trade, not on a morning coin-flip about the day's character.

8. grind is readable early — cliffy, not validated

The hope was that grind days announce themselves in the first hour so you could lean in early. I couldn't make it stand up. The signal is what I'd call cliffy — it looks separable on the days it works and gives you nothing on the days it doesn't, with no stable threshold in between. Given musing 5 already refuted grind as a standalone, an early-read on a strategy that loses \$8k over 25 days isn't worth chasing yet. Filed as not-validated. If the ATR-conditional trail (musings 1-2) proves out and makes grind's runners pay, this early-read question becomes worth reopening. Until then it's a maybe resting on a no.

The through-line. Two things held: the chandelier is too loose and wants an ATR-conditional trail, and direction-by-regime gating is the real lever — both confirmed from independent angles and across the two-week history, not one week. Everything shinier — grind live, the chop specialist, the day-type router — refuted or outlier-carried the moment I pushed it out of sample. The circuit breaker is simply untested because it never fired. Nothing here is armed. The two that held are the two I'd spend next week hardening.

Where the money actually landed this week across the leads, so you can see the spread in one look — the two survivors sit up top, the refuted grind-as-live is the biggest sim but the archive kills it, and the reversion/capitulation family anchors the bottom:



*grind_fast leads the week but loses \$7,975 over the 25-day archive — this week's top sim is the clearest example of why one good week isn't a verdict.

THE WINNERS — LEAN INTO THESE

§10 grind_fast — the week's best earner, and it's built right

grind_fast made **\$1,299.50** over 77 MNQ trades this week at a 47% win rate, and it did it the honest way: small controlled losses, a handful of big chandelier winners, and both the long and short book in the green. This is the cleanest sim on the V7 board — nothing about it looks lucky.

1. The money, in one sentence

Over the two days the V7 board has been live (07-16 and 07-17), grind_fast took 77 filled MNQ trades and cleared **\$1,299.50** of real P&L at a 46.8% win rate — the top earner on the whole shadow board, and it got there by winning bigger than it lost (average win **\$81.40** against an average loss of just **-\$39.80**), not by being right more often.

Every number in this section is scored on **real_pnl** where **fill_status='filled'**, from `/home/alphabot/gazbot7/data/shadow.db` (`shadow_trades JOIN shadow_real`). Never on ceiling — the ceiling understates losses.

2. Why it works — pull it apart

Both sides pay. This is the thing that separates grind_fast from the one-sided sims. The long book made **\$747.50** over 34 trades at 44% and the short book made **\$552.00** over 43 trades at 49%. Neither side is carrying the other; neither side is a hidden bleed. When a strategy is green on both sides in the same week, you're usually looking at a real directional read rather than a lucky tilt in one market. Contrast the live thrust-only desk, where LONG bled and only SHORT held up — grind_fast doesn't have that asymmetry.

SIDE	TRADES	REAL P&L	WIN%
LONG	34	\$747.50	44.1%
SHORT	43	\$552.00	48.8%

The exit is the whole engine. Split by exit_reason and the machine is obvious: the chandelier trail closed 39 trades for **\$2,606.50** at an 82% win

rate, and the stop closed 38 trades for **-\$1,307.00** at a 10% win rate. That is exactly the shape you want from a momentum sim — it lets winners ride out to the trailing chandelier and it cuts losers fast on the fixed stop. The stop isn't broken; it's doing its job of capping the downside at a small, uniform size (worst single loss just **-\$78.50**). The money is made by the winners the chandelier carries, not by avoiding stops.

CHANDELIER (39t, 82% win)

+\$2,606.50

STOP (38t, 10% win)

-\$1,307.00

It likes volatility. Bucket the entries by `atr_pct` (`entry_atr ÷ entry_price`) and the higher-ATR trades earn more: the top band (`atr_pct ≥ 0.08%`) made **\$696.00** over 14 trades at 50%, and every ATR band is green. So there's no "it only works in quiet tape" trap here — if anything it wants a live, moving market, which fits a grind-continuation idea.

ATR_PCT BAND	TRADES	REAL P&L	WIN%
< 0.04%	10	\$205.00	60.0%
0.04-0.06%	28	\$132.50	42.9%
0.06-0.08%	25	\$266.00	44.0%
≥ 0.08%	14	\$696.00	50.0%

In plain words: when `grind_fast` wins, it has caught a move that keeps going and the chandelier walks it up for an \$80-\$335 gain; when it loses, the move fails quickly and the fixed stop takes it out for around \$40. It wins less than half the time and still clears money because the winners are roughly twice the size of the losers. That's a real edge, not a coin flip.

3. The one thing for Monday — cut the dead evening

Split the day by session hour and there is one ugly window. Everything `grind_fast` lost, it lost in the UTC evening. The block from 17:00 to 20:00 UTC took 25 trades for **-\$195.50** at a 32% win rate. The other 20 hours of the day took 52 trades for **\$1,495.00** at a 54% win rate. Hour 18 alone was **-\$191.50** at 17% — nearly the entire evening loss sits in a single thin, chopped-up hour.

WINDOW	TRADES	REAL P&L	WIN%
Evening 17:00-20:00 UTC	25	-\$195.50	32.0%
All other hours	52	\$1,495.00	53.8%

The gate, plainly: don't let grind_fast open new positions between 17:00 and 20:00 UTC. In-sample that lifts the week from **\$1,299.50** to **\$1,495.00**, a swing of about **+\$195.50**, and it does it by cutting a third of the trades — the losing third. The other side of the same coin: the win rate on what's left jumps from 47% to 54%.

Be honest about it: this is two days of data and 25 trades in the flagged window. A single-week hour gate mostly means "this window was chop and you traded it anyway" — it is not a proven time-of-day law, and one strong evening next week would soften it. But it's the cleanest separator in the sim, it lines up with a known thin-liquidity patch (the US-afternoon lull into the pause), and it costs nothing to try because it only ever stands the sim down. Treat it as a veto to watch, not a permanent rule.

What NOT to change: don't touch the sides — both are green, so there's no case to disable long or short. Don't add an ATR floor — every ATR band made money, so a floor would just cut winning trades. And don't loosen the stop; the small uniform losses are exactly what makes the 2:1 win/loss ratio work. The only change worth making is the evening veto.

Verdict: **ARM IT** — grind_fast is the best-constructed sim on the board: both sides pay, the exit logic is doing real work, and it wants volatility rather than fearing it. **WATCH THE 17-20 UTC VETO** for another week before hard-wiring it. If it holds up over more days, this is the first V7 sim I'd nominate for the promotion ladder.

§11 thrust_cont — the quiet winner that just takes its target

thrust_cont made **\$594** across 29 filled trades over the two days it has been on the V7 board (07-16 and 07-17), winning 52% of them. All 29 are MNQ. It is one of only a handful of sims on the new board that finished the week clearly green, and — unusually for this desk — it made money on BOTH sides of the market. There is no bleeding bucket to cut here. The honest question is not "what is broken" but "is this real enough to arm, and does anything make it bigger."

One-sentence money: thrust_cont earned **\$594 over 29 trades at 52% win**, and every dollar of it is the take-profit target doing the work — its 13 TARGET exits banked **\$1,008** while its 16 STOP exits gave back **\$414**.

Why it made money — it wins by getting paid at the target, not by being right often

This sim is mechanically simple and the data shows it plainly. There are exactly two ways it leaves a trade, and they map almost perfectly onto win and loss. When it reaches its profit target it wins 100% of the time (13 of 13, by definition), and those exits carry the whole book at **\$1,008.50**. When it stops out it loses 87% of the time (14 of 16), bleeding **\$414.50**. So the edge is not clever exits or trailing — it is that the average winner (\$77.6) is bigger than the average loser (−\$25.9). It only needs to be right half the time to come out ahead, and this week it was right 52%.

Split by side, the surprise is that the SHORTS carried it. Shorts made **\$433** over 16 trades at 63% win, longs made **\$161** over 13 trades at just 38% win. Both sides are green, which is the thing worth protecting — most sims on this desk lean hard one way. But the longs are clearly the weaker leg: they hit their target only 3 times in 13 tries. The reason the longs still finished positive is that when a long DID run, it ran big — the two largest single trades in the whole sample were both longs on loud tape (\$113.5 and \$176.5). Take those two away and the long book is a net loser. So "longs work" is really "two long runners paid for eleven small long stop-outs." That is thin ice, and I would not read it as a durable long edge on two days of data.



Split by the loudness of the tape at entry ($\text{atr_pct} = \text{entry_atr} \div \text{entry_price}$), a cleaner pattern shows up than the side split does. The trades break into three amplitude bands and the money climbs monotonically with volatility: the quiet 0.04–0.06% band made **\$153** over 14 trades at 50% win, the middle 0.06–0.08% band made **\$210** over 12 at 50%, and the three loudest trades ($\geq 0.08\%$) made **\$231** at 67%. The single biggest winners of the week — the 0.099% long (+\$113.5) and the 0.155% long (+\$176.5) — sit right at the top

of that ladder. In plain words: this thing gets paid when the thrust it is chasing is genuinely big, and it grinds sideways-to-flat when the move is small.

By session hour the picture is noisier and I would not lean on it. The good hours were scattered — the US afternoon runs at 14:00 and 15:00 UTC each produced a single fat winner (\$176.5 and \$113.5), and the 08:00 hour went 2-for-2 for \$166.5. The soft patch is late-US-afternoon into the 16:00–19:00 window: hour 16 lost \$110 on two trades, and the 19:00 hour was the busiest of the day (5 trades) but only 40% winners for a slim \$38. With one or two trades in most hourly buckets this is basically noise, so I am treating hour-of-day as descriptive colour, not a gate.

What it does when it wins vs when it loses, in human terms

When it wins, it has jumped onto a real, loud thrust — often a short into a flush — and the move keeps going far enough to tag a fixed target before it gives back. When it loses, it has jumped onto a smaller or half-hearted thrust, the move stalls, and the fixed stop takes a small, controlled bite. The losses are cheap and uniform (nearly every stop is between $-\$26$ and $-\$52$); the wins scale with how loud the tape was. That is exactly the shape you want from a continuation sim, and it is the opposite of the reversion family that bled the V5 board — this one is paid to be on the right side of momentum, not to fade it.

The one thing for Monday — gate on tape loudness, not on side

The cleanest single split that separates its winners from its losers is amplitude at entry, at a threshold of $\text{atr_pct} \geq 0.07\%$. Above that line it took 9 trades for **\$468** at 67% win. Below it, the other 20 trades made only **\$126** at 45% win — barely above breakeven. Two-thirds of the profit came from under a third of the trades, and the quiet-tape trades did almost nothing but churn.

BUCKET	TRADES	REAL P&L	WIN%
loud tape — atr_pct $\geq 0.07\%$	9	+\$468	67%
quiet tape — atr_pct $< 0.07\%$ (the rest)	20	+\$126	45%

The gate, stated plainly: only take a thrust_cont entry when the ATR is at least 0.07% of price — stand down on quieter tape. In-sample this concentrates the same book into a third of the activity: you keep \$468 of the \$594, forgo \$126, and skip 20 marginal trades. That is not a dollar windfall — the honest in-

sample swing on winrate is the real prize, 67% vs 45%, not the P&L, because the quiet bucket was still slightly green. This is a "trade less, keep the quality" filter, and on a 29-trade, two-day sample I want to be blunt that a nine-trade good bucket is thin. It rhymes with the amplitude-floor finding that already survived on the V5 momentum shadows (arm the thrust only when the move is loud), which is the only reason I would act on so little data.

Verdict: **ARM SMALL & WATCH** – thrust_cont is genuinely +EV on this week, green on both sides, with cheap uniform losses and target-driven wins. Add the atr_pct ≥ 0.07% amplitude floor as an entry veto (in-sample: 67% win on the kept trades vs 45% on the dropped). Do NOT read the long edge as durable – it rests on two runners. Re-check both after another full week before sizing up; 29 trades over two days is not enough to promote.

One caveat I will not paper over: exit_reason here is deterministic (TARGET always wins, STOP mostly loses), so no "exit tweak" is available in this data — the only levers are which entries you allow. The amplitude floor is the honest one. Widening the target or the stop is untestable from these 29 rows and would be a guess. Leave the mechanism alone; gate the entry.

§12 **abs_veto_55s — the absorption-veto that quietly banked a clean week**

On the V7 board this week abs_veto_55s made **\$499** across 21 filled trades, won 57% of them, and made money on BOTH sides of the market — long and short — which is rare and worth taking seriously.

This is the momentum sim that sits on its hands for 55 seconds when it sees the book absorbing (big resting size soaking up the aggressor flow) before it lets an entry through. The sister sims abs_veto_50s (+\$443.5) and the two grind/thrust variants all landed green too, so this isn't a lone lucky print — the whole absorption-veto family had a good week. What follows is only abs_veto_55s, scored on real_pnl from shadow_trades JOIN shadow_real where fill_status='filled', so these are honest fills, not the ceiling.

1. The money, in one sentence

It took 21 trades, won 12 of them (57%), and cleared **\$499** — its winners paid \$661 through the target and its stops cost it \$162, and it was net-green whether it went long or short.

One honest caveat up front: the ceiling on these same 21 trades was \$784, so the desk actually banked \$499 – it gave back about a third of peak (\$285) to slippage and late fills. That give-back is the same one-third-of-peak leak we see across every sim at every trail width, so it's not specific to this one. The edge here is the ENTRY quality; the give-back is a separate, known problem.

2. Why it made money — where the dollars actually came from

Both sides win. This is the headline. Short made **\$356** over 11 trades at a 64% win rate; long made **\$143** over 10 trades at 50%. On the live thrust-only desk longs are bleeding badly (−\$733 at 17%), so a sim whose longs are at least break-even-plus is doing something right — the 55-second absorption wait is filtering out the worst long entries, the ones that buy straight into a wall of offers.

The target is the whole engine; the stop is small and controlled. Split by exit_reason, everything good comes through the TARGET: 13 target exits, 11 of them winners (85%), **\$661** banked. The 8 stop-outs cost **−\$162** between them — small, evenly-sized clips of roughly \$30 each, not a single catastrophic loss. That is a healthy shape: it wins by being right often and taking its profit, and when it's wrong it loses a little. Crossed against side, the cleanest cell is short-into-target — 8 trades, **\$443** — while short stops only cost \$87 and long stops \$75.

SHORT → target (8t)	+\$443
LONG → target (5t)	+\$218
LONG → stop (5t)	−\$75
SHORT → stop (3t)	−\$87

It wants a lively tape. Bucket the entries by amplitude ($\text{atr_pct} = \text{entry_atr} / \text{entry_price}$) and the split is stark. When it entered on a wider, more energetic tape ($\text{atr_pct} \geq 0.06\%$) it made **\$415.5** over 10 trades at 70% wins. When it entered on the thinner 0.04–0.06% tape it scraped just **\$83.5** over 11 trades at 45% wins — barely above water and coin-flip. In plain words: when the market is actually moving, the absorption wait lets it in on a real thrust and it rides it to target; when the tape is quiet, it gets chopped and the small stops nibble it down.

ENTRY AMPLITUDE	TRADES	REAL_PNL	WIN%
atr_pct $\geq$ 0.06% (lively)	10	+\$415.5	70%
atr_pct 0.04-0.06% (quiet)	11	+\$83.5	45%

Session hour is noise here — don't read it. The best hour looks like 08:00 UTC (+\$231, 3/3) and the worst like 11:00 (−\$57.5, 0/2), but those are two- and three-trade cells on a 21-trade week. Carving by hour on this little data is overfitting; I'm flagging it only to say I looked and there's nothing honest to gate on.

3. The one thing for Monday

The single condition that best separates winners from losers is the **amplitude floor**: only take entries when $\text{atr_pct} \geq 0.06\%$. That good bucket did **\$415.5** over 10 trades at 70% wins; the quiet-tape remainder did **\$83.5** over 11 trades at 45%. Applying the floor keeps 83% of the money while cutting more than half the trades and lifting the win rate from 57% to 70%.

Gate, plainly stated: require $\text{entry_atr}/\text{entry_price} \geq 0.06\%$ (i.e. only fire when the tape is genuinely moving). **In-sample effect:** you forgo just \$83.5 of profit but shed 11 low-quality trades and turn a coin-flip book into a 70%-win book.

Be honest about the size of this. The dollar swing is small — the quiet band was still slightly positive, so this isn't rescuing a loser, it's concentrating a winner. And the good bucket is only 10 trades over one week; the sharpest-looking cell inside it (short + lively tape went 5-for-5 for \$382.5) is five trades and I won't hang a rule on that. This is a "trade less, trade the good tape" tweak filter, not a discovered edge — the kind of single-week split that mostly just means the market cooperated on the loud days.

Verdict. The sim itself is the story, not the tweak: a two-sided, target-driven, small-stop profile that made money cleanly. **ARM AT 1 MICRO** and let it run to build out-of-sample trades — and stack the amplitude floor as a soft preference (prefer $\text{atr_pct} \geq 0.06\%$, don't hard-block below it yet). The one real bug to chase is separate from entries: trade 1504 hit its target with a +\$144 ceiling but filled at −\$68, a \$200+ fill gap — that give-back, not the entry gate, is where the next \$285 is hiding.

§13 thrust_loose — it made money, but only the short half did

Over its first honest day and a bit on the V7 board (07-16 18:40 to 07-17 19:28 UTC), thrust_loose took 43 MNQ trades and came out **+\$462** on real fills at a 44% win rate — a green week, but a lopsided one: every dollar of it came from the short side, while the long side was a coin-flip that mostly got stopped out.

The money, in one sentence: thrust_loose made **+\$462** over 43 filled MNQ trades at 44% win because its 20 winners hit the fixed target for **+\$1,221** and its shorts (+\$500) more than paid for its longs (−\$39) — but the whole edge is one direction and one day of data, so read it as "the short side works," not "the sim works."

Why it made money — winners hit target, longs fade

This is close to a binary sim: there are only two exits on the board, a fixed **TARGET** and a fixed **STOP**, and they split almost perfectly by outcome. The 20 target exits made **+\$1,221.5** at a 90% win rate; the 23 stop exits bled **−\$759.5** at a 4% win rate. There's no chandelier here doing clever give-back management — a trade either runs to its target or it dies at its stop. So the whole question of whether thrust_loose is any good collapses into one thing: how often does a "loose" thrust entry actually keep going after you're in?

The answer depends entirely on which way you're facing. Split by side:

SIDE	TRADES	REAL P&L	WIN%
SHORT	23	+\$500.5	57%
LONG	20	−\$38.5	30%

The shorts caught real breaks lower and hit target 14 times out of 23. The longs bought rips that turned around and ran them into the stop — 14 of the 20 longs stopped out for **−\$433**, and the six that worked (**+\$394.5**) only just clawed the side back to roughly flat. In plain words: **when a loose thrust fires to the downside it tends to be a genuine flush that follows through; when it fires to the upside it's usually a spike into a level that fades**. That's the same asymmetry the two-week V5 history hinted at (the thrust family lived on the short-into-weakness trades), and it's the same thing the live V7 desk is showing right now (longs −\$733, shorts +\$46).

SHORT		+\$500.5
LONG	-\$38.5	
TARGET		+\$1,221.5
STOP	-\$759.5	

What it does in different conditions

By **volatility at entry** ($\text{atr_pct} = \text{entry_atr} / \text{entry_price}$, which ranged 0.04% to 0.16%, averaging 0.063%), the sim likes a middle band and gets nothing from the thinnest tape:

ATR_PCT BAND	TRADES	REAL P&L	WIN%
0.04-0.06%	22	+\$70.5	41%
0.06-0.08%	16	+\$283.5	50%
≥0.08%	5	+\$108.0	40%

More than half the sim's whole take came from that 0.06–0.08% ATR band, and the low-vol band (which held half the trades) barely paid at all. That fits the desk's standing amplitude-floor thesis — thin tape doesn't give thrusts room to run to target. But the two axes interact, and it's worth stacking them: split by side and a 0.06% ATR floor and the picture sharpens to a single cell.

SIDE × ATR	TRADES	REAL P&L	WIN%
SHORT, atr ≥0.06%	11	+\$390.5	64%
SHORT, atr <0.06%	12	+\$110.0	50%
LONG, atr ≥0.06%	10	+\$1.0	30%
LONG, atr <0.06%	10	-\$39.5	30%

The longs win 30% of the time regardless of how much vol you demand — turning up the ATR floor doesn't save them, it just gets you a flatter zero. Session hour tells the same story more noisily: the sim's best clumps were the London/early-US hours (08:00 +\$231, 14:00 +\$176.5, 00:00 +\$126), and its worst were the late-US/evening chop (16:00 −\$110, 11:00 −\$87, 18:00 −\$81.5). But with only one to seven trades per hour, that's texture, not a rule — I wouldn't gate on the clock off this.

The one thing for Monday

Drop the long side. Trade thrust_loose short-only. Longs were 20 of the 43 trades and returned **-\$38.5** at 30% win — 14 of them stopped out. Cutting them leaves the 23 shorts at **+\$500.5** and 57% win, i.e. it lifts the sim's take from +\$462 to +\$500.5, raises win rate from 44% to 57%, and removes the side that carries all the -\$433 of stop-out risk. In-sample that's a **+\$38.5** P&L swing plus a much cleaner distribution — you're keeping the whole edge and throwing away a coin-flip.

If you want the sharper version, the single best-separated bucket is **SHORT with atr_pct ≥ 0.06%**: 11 trades, **+\$390.5**, 64% win — that one cell is 84% of the entire sim's profit. The other 32 trades combined made just +\$71.5. So the honest ranking is: the side split is the robust, must-do change; the added ATR floor is the "if it holds up" refinement.

Be honest about the sample. This is one day and a bit — 43 trades, and the headline gate rests on 11. That's far too thin to arm a size change on, and a short-only rule will look brilliant on any dataset where the market happened to sell off, which this window did. What earns confidence is that the short-beats-long asymmetry is not new: it's the same shape in the two-week V5 thrust family and in the live V7 desk this week, three independent looks all pointing the same way. So the call is **WATCH → ARM SHORT-ONLY, KEEP LONGS SHADOW-ONLY**: flip thrust_loose to shorts on the live-armed contracts, leave the long leg running in shadow to see whether it's genuinely broken or just had a down week, and don't touch size until the short book has a couple hundred trades across a mixed-direction fortnight.

§14 **abs_veto_50s — the short-side absorption sim that quietly banked a clean \$443**

This is one of the good-news sims of the week. Over the two V7 days it took 22 trades on MNQ, won 55% of them, and finished **\$443.50** in the black on honest fills. Nothing dramatic, no single hero trade — it just kept its losers tiny and let its targets pay. It sits right in the middle of the V7 leaderboard alongside its sibling abs_veto_55s (+\$499) and the thrust family (thrust_cont +\$594, thrust_loose +\$462), so it is real company, not an outlier.

1. The money, in one sentence

abs_veto_50s made **\$443.50** across 22 filled MNQ trades at a 55% win rate, and it captured about two-thirds of its own best-case ceiling (\$681 of unrealised peak came home as \$443.50) — a give-back that fits the chandelier pattern we see everywhere, not a broken exit.

```
Query: SELECT count(*), sum(real_pnl), avg(real_pnl>0) FROM shadow_trades t JOIN
shadow_real r ON r.trade_id=t.id WHERE strategy='abs_veto_50s' AND fill_status='filled' →
22 trades, $443.5, 54.5% win. Ceiling vs real: $681.0 → $443.5.
```

2. Why it made money — where the dollars actually came from

Three things stand out when you pull the trades apart, and they all point the same way: **this is a short-biased sim with a clean, disciplined exit that shines in a specific volatility band.**

The short book carried it. Split by side, the eleven shorts made **\$356.00** at a 64% win rate, while the eleven longs scraped together **\$87.50** at just 45%. More telling than the totals is the shape: no short in the whole week lost more than \$33, so the downside was tightly capped, while the short winners ran \$44 to \$84 each. The longs were the opposite — they held the two ugliest trades in the book (a −\$68 target-miss and a −\$55.50 stop, both in high volatility) alongside the single best trade (+\$113.50). Longs are a coin-flip here; shorts are the engine.

The exit is doing its job. The exit-reason split is about as clean as you will see: the 13 TARGET exits made **\$661.00** at an 85% win rate, and the 9 STOP exits lost **−\$217.50** at 11%. That is exactly the profile you want — when it's right it rides to target, when it's wrong the stop cuts it fast. Every stop-out was small (worst −\$55.50). One STOP even locked in +\$82 on a long that had already run, so the stop is trailing profit, not just eating losses. There is no third failure mode bleeding money in the background.

It wins in the middle of the volatility range. Bucketed by entry ATR-percent ($\text{entry_atr} \div \text{entry_price}$), the sim is a mid-vol animal. The 0.06–0.08% band made **\$370.00** at 75% over 8 trades — that one band is 83% of the entire week's profit. The quiet 0.04–0.06% band made only **\$83.50** at 45%, and the three high-vol trades above 0.08% netted **−\$10.00**. In plain words: when the tape is dead-quiet it chops around breakeven, when it's wild the longs blow up, and in the sweet middle it prints.



Session hour tells a softer version of the same story. The strong windows were the 08:00 hour (+\$231, 3-for-3), the 00:00 hour (+\$126, 2-for-2) and the 19:00 hour (+\$102). The red came in the US lunch lull and late tape — the 11:00 hour lost −\$57.50 (0-for-2) and 13:00 lost −\$55.50 (0-for-1). It's the right intuition (avoid the dead midday chop), but with one or two trades per hour it's too thin to gate on by itself. I'd treat hour as confirmation of the ATR story, not an independent rule.

CUT	N	REAL_PNL	WIN%
SHORT	11	\$356.00	63.6%
LONG	11	\$87.50	45.5%
Exit = TARGET	13	\$661.00	84.6%
Exit = STOP	9	−\$217.50	11.1%
ATR 0.06-0.08%	8	\$370.00	75.0%
ATR < 0.06%	11	\$83.50	45.5%

3. The one thing for Monday

The cleanest single separator is **entry ATR-percent**, and the nice part is it splits the book almost exactly in half so it isn't a three-trade fluke. Trades entered with ATR at or above 0.06% of price made **\$360.00** over 11 trades at 64%; the low-vol half below 0.06% made only **\$83.50** over the other 11 at 45%.

The gate: only arm abs_veto_50s when entry ATR is $\geq 0.06\%$ of price (skip the dead-quiet tape). In-sample that lifts per-trade take from about \$7.60 to \$32.70, and cuts the trade count in half while keeping 81% of the profit. The dollar swing between the two halves is **\$276.50** — the high-vol half carried essentially all the edge.

Where I add the honesty: this is a filter that says "trade less," not one that conjures new money. You don't make more than \$443 by applying it — you make \$360 on far fewer, better trades and stop paying the low-vol chop tax. And it's one week, 22 trades, in-sample. The 0.06–0.08% band being the exact sweet spot (with high-vol >0.08% actually turning negative) is only three trades wide on the top end, so I would frame it as a floor — "not below 0.06%" — rather than a tight window, until more days come in.

The tempting second move is to cut longs entirely, since shorts made four-fifths of the money and never lost big. I'd hold off. The longs are still net-positive (+\$87.50), and killing a positive-EV side on one week's read is how you overfit. The right posture is: **keep it two-sided, gate on the ATR floor, and know that your conviction sits on the short book.** If longs keep being the choppy, blow-up-prone side over the next two weeks, revisit then.

WATCH → ARM SMALL It made honest money with a clean exit and a defensible volatility edge. I'd arm it at one micro alongside its sibling `abs_veto_55s` (+\$499, same family), apply the $ATR \geq 0.06\%$ floor, and give it two more weeks before touching the long/short balance or sizing up. Thin-n, but the shape is right and the risk was well-controlled all week.

§15 **thrust_aligned — a short-selling bracket that quietly netted \$194**

`thrust_aligned` made money this week, but not the way you'd guess from the name. It took 23 fills on MNQ, won 43% of them, and finished **+\$194** — and essentially all of that green came from the short side. It's a clean fixed-bracket system: enter with the thrust, take a fixed target or eat a fixed stop, nothing fancy in between. It works because its winners are bigger than its losers, not because it's often right.

1. The money, in one sentence

Over the two-day V7 board it took **23 filled MNQ trades, won 43.5% of them, and banked +\$194** (\$8.40 a trade) — a real, if slim, positive week carried entirely by shorts.

Every figure here is `real_pnl` where `fill_status='filled'`, from `shadow_trades JOIN shadow_real` in `/home/alphabot/gazbot7/data/shadow.db`, strategy `thrust_aligned`. No ceiling, no `pnl_usd` — the honest fill number.

2. Why it made money — what it does when it wins vs loses

It's a one-sided desk. Split by side, the whole story falls out. The shorts did the work and the longs quietly bled:

SIDE	TRADES	REAL_PNL	WIN%	PER TRADE
SHORT	17	+\$293.0	52.9%	+\$17.2
LONG	6	-\$99.0	16.7%	-\$16.5

When it shorts a rip that rolls over, it's right more than half the time and pockets it. When it buys a thrust up, it's been right exactly once in six tries this week — it keeps buying tops that immediately turn and stop it out. Same logic, opposite outcome, purely by direction.

The exit mechanism is a pure bracket — there is no middle. Every single trade ended at either its fixed TARGET or its fixed STOP. There's no chandelier, no give-back, no time-clock churn here — you either hit the target or you don't:

EXIT	TRADES	REAL_PNL	WIN%	PER TRADE
TARGET	11	+\$669.0	90.9%	+\$60.8
STOP	12	-\$475.0	0.0%	-\$39.6

This is the engine of the whole result: a target is worth about **\$61** and a stop costs about **\$40**, so the reward-to-risk is roughly 1.5-to-1. At a 48% target-hit rate that nets positive even though it's wrong more often than a coin. The strategy doesn't win by being accurate — it wins by losing small and winning bigger. That's a healthy shape, and it's the reason a 43%-win-rate system is green at all.

SHORT / TARGET	+\$555.5
LONG / TARGET	+\$113.5
LONG / STOP	-\$212.5
SHORT / STOP	-\$262.5

Volatility helps it. Bucketing by entry `atr_pct` (`= entry_atr / entry_price`), the fatter-ATR entries are where the money is. Everything it took was already above 0.04% ATR (the strategy floors there), but inside that:

ENTRY ATR%	TRADES	REAL_PNL	WIN%
≥ 0.07% (fat)	7	+\$209.5	57.1%
0.04-0.07% (thin)	16	-\$15.5	37.5%

In plain words: it needs the tape actually moving. When ATR is fat it wins 57% and makes real money; when the range is compressed it's a coin-flip that nets to roughly zero. The thrust means more when there's genuine energy behind it.

Session hour is noise at this sample. The green is scattered — the US-morning entries (08:00, 15:00 UTC) printed well and the late-Europe 16:00–17:00 slots printed red — but with one or two trades per hour there's nothing to lean on. Don't build an hour gate out of 23 trades; it would just be curve-fitting the calendar.

3. The one thing for Monday

The cleanest separator of winners from losers is **direction**. The long side is the only thing dragging this sim down, and it's not close:

BOOK	TRADES	REAL_PNL	WIN%
Shorts only (proposed)	17	+\$293.0	52.9%
Longs (the drag)	6	-\$99.0	16.7%

Gate: **run thrust_aligned short-only on MNQ**. In-sample that lifts it from **+\$194** to **+\$293** — a **+\$99 swing** — and pushes the win rate from 43% to 53%.

Be honest about the size of this. The whole case against longs rests on six trades over two days. One good long thrust next week and that -\$99 halves. So this is a lean, not a law. What makes me still trust the direction is that it isn't just this sim: the 2-week V5 history had the same fingerprint (thrust_loose was LONG -\$38 / SHORT +\$500, and the live V7 thrust desk is LONG -\$733 @17% vs SHORT +\$46 @44%). The long side of the thrust family has been the weak side everywhere I look, so short-favouring is the through-line, not a two-day fluke.

Secondary, and I'd stack it on top rather than instead: the fat-ATR filter. Shorts that also had $ATR \geq 0.07\%$ were **+\$250** over just 3 trades — great dollars, far too few trades to arm on its own. Treat "fat ATR" as a confidence tilt, not a hard gate, until there are more fills.

Verdict: ARM SHORT-ONLY, KEEP AT 1

thrust_aligned is a genuinely positive fixed-bracket short-seller: 1.5-to-1 reward-to-risk, wins by losing small. The single Monday change is to stop it buying tops — go short-only, worth +\$99 in-sample. But the long-side case is six trades deep, so this earns a watch-and-size, not a promotion. Let the shorts prove the +\$293 holds over another week before adding a contract.

§16 **abs_veto_60s — the patient absorption-veto that barely traded**

This one made money and never really got out of bed. Over the two live V7 days it took just six trades on MNQ, won four of them, and banked **\$127** on honest fills. Every dollar is real — but six trades is not a verdict, it's a postcard. The interesting fact isn't inside abs_veto_60s, it's next to it: the two shorter-window siblings did the same job four times as often and made four times as much.

The money, in one sentence: abs_veto_60s made **\$127** across 6 filled trades at a 67% win rate, with both directions green — LONG **\$76.50** on 3 trades and SHORT **\$50.50** on 3 — because its four target hits (**\$206.50**) more than paid for its two stop-outs (**\$79.50**).

What it actually did — all six trades on the table

With a sample this small there's no hiding behind averages, so here is every fill it took. This is the whole strategy for the week.

WHEN (UTC)	HR	SIDE	ATR_PCT	EXIT	REAL_PNL
07-16 23:40	23	LONG	0.050%	STOP	-\$31.50
07-17 00:51	00	LONG	0.051%	TARGET	+\$59.00
07-17 06:21	06	LONG	0.044%	TARGET	+\$49.00
07-17 06:38	06	SHORT	0.044%	TARGET	+\$49.00
07-17 08:17	08	SHORT	0.076%	STOP	-\$48.00
07-17 12:03	12	SHORT	0.044%	TARGET	+\$49.50

The pattern jumps out immediately. When `abs_veto_60s` wins, it wins the same way every time — it clips a fixed profit target for \$49 to \$59. When it loses, it loses the same way every time — it gets stopped for \$31 to \$48. There is no runner, no big day, no tail. It's a tight, symmetric little scalp that takes roughly a one-to-one payoff and lives or dies on hit rate. Four of six is enough to be green; three of six would have been roughly flat.

Why it made money — the three cuts


By side. Both directions worked, which is the healthiest thing you can say about a momentum-family sim. LONG went 2-of-3 for \$76.50 and SHORT went 2-of-3 for \$50.50. Neither side is carrying the other, and neither side is quietly bleeding under a good headline. That two-sided balance is the same signal the shorter windows show, so I trust the direction-neutrality more than I trust the dollar total.

By exit. This is the whole engine. The four TARGET exits made \$206.50 against a modeled ceiling of \$260 — so on its winners it captured about 79% of the perfect-fill dream, which is a clean, believable haul with no fantasy fills. The two STOP exits cost \$79.50 against a ceiling of \$81 — meaning the stops came in essentially exactly as modeled, no ugly slippage surprise. There is no third exit type. It never used a chandelier, never got flat-clocked, never force-killed. It either reached its target or it hit its stop, full stop.

By ATR and hour. Its winners all fired in a narrow, quiet volatility band — entry `atr_pct` sat around 0.044% to 0.051% on every single profitable trade. Its two losers sit at the edges of that comfort zone: one at the low end (0.050%, a 23:40 overnight long that got stopped) and, more tellingly, the one genuinely loud entry of the week — a 0.076% short at 08:17, its most volatile

setup by half again, which was also its worst single loss at **\$48**. On the clock, the two winning hours were the 06:00 UTC block (both a long and a short scored there) and the 00:00 and 12:00 prints; the two losers were the 23:00 overnight and the 08:17 European-morning spike.

In plain words: `abs_veto_60s` is a quiet-tape target-scalper. When volatility is mild and it can grab its fixed target off an absorption read, it converts cleanly on both sides. When it either reaches into the dead overnight or gets caught in a genuinely fast, high-ATR move, the target never comes and the stop does the talking.



The one thing for Monday

The cleanest in-sample split is the loud-entry cut. Every winner came in under 0.06% `atr_pct`; the one entry above that line — the 0.076% short at 08:17 — was the biggest loss of the week. Fence it off and the picture sharpens:

BUCKET	N	WIN%	REAL_PNL
<code>atr_pct < 0.06%</code> (quiet tape)	5	80%	+\$175.00
<code>atr_pct ≥ 0.06%</code> (loud tape)	1	0%	-\$48.00

The gate, plainly: don't let `abs_veto_60s` take an entry when `atr_pct` is above ~0.06% — when the tape is genuinely fast, its fixed target sits too close and the stop wins. In-sample that lifts the week from **\$127** to **\$175**, a swing of about **\$48**.

Be honest: that swing is one trade. `n=1` on the excluded side is not an edge, it's an anecdote, and the same "skip the loud entry" filter would also just mean "trade even less" on a sim that already only traded six times. I'd hold it as a note, not arm it. The overnight 23:40 stop, by the way, breaks the clean ATR story — it was a low-ATR loser — so even the volatility read isn't airtight.

The real Monday read — look at the siblings

The most useful thing about `abs_veto_60s` is what it says when you stand it next to `abs_veto_55s` and `abs_veto_50s`, which run the identical absorption logic with a shorter veto window. Same family, same week:

SIM	TRADES	WIN%	REAL_PNL
<code>abs_veto_55s</code>	21	57%	+\$499.00
<code>abs_veto_50s</code>	22	55%	+\$443.50
<code>abs_veto_60s</code>	6	67%	+\$127.00

`abs_veto_60s` has the prettiest win rate and by far the least money, because the extra ten seconds of patience made it sit out most of the setups its siblings took — and those setups were, on balance, good ones. The 55s window did nearly four times the work for four times the reward at essentially the same quality. So the honest conclusion isn't a tweak to 60s; it's that the 60-second veto is too patient, and the family's earning member is 55s.

WATCH / DON'T ARM ALONE `abs_veto_60s` is real and green but too thin and too idle to arm on its own. Arm the absorption-veto family through the 55-second window, which does the same thing more often for more money; keep 60s as the observe-only control that confirms both sides work and that the stops fill as modeled. Revisit only after it has 20-plus trades to score.

§17 `thrust_fast` — a small green week carried entirely by its shorts

`thrust_fast` made **\$95** over 17 filled trades this week, winning just 35% of the time — a real profit, but a thin one, and every dollar of it (and then some) came from the short side.

This is a clean, honest little strategy to read because it has no exit machinery to argue about. Every one of its 17 trades ended one of two ways: it hit its **TARGET** or it hit its **STOP**. No chandelier, no absorption cut, no flat-clock. So the whole story is in the entries — which thrusts it chose to chase — and the answer this week is stark: it fades and gets stopped when it buys, and it pays when it sells.

1. The money, in one sentence

Over the two-day V7 board (07-16/17) thrust_fast took 17 MNQ trades and netted **\$95** on real fills at a 35% win rate — profitable only because six winners paid an average of \$84 each while eleven losers cost an average of \$37, a favourable-payoff shape that barely cleared water.


The engine underneath is pure asymmetric R:R. Its 6 **TARGET** hits banked **\$507** at a perfect 100% win rate; its 11 **STOP** hits bled **\$412** at 0% — by construction. The target is worth roughly 1.4× the stop, so the strategy only needs to be right a bit more than 42% of the time to profit. This week it was right exactly 35% — below that line on paper — yet still green, because the trades it did win were its biggest-ATR, highest-payoff ones. That is luck-adjacent. The structural fix isn't the exit; it's refusing the entries that never had a chance.

2. Why it won — and why it nearly didn't

Side is the whole game. The short book did all the work and the long book actively fought it:

SIDE	TRADES	REAL P&L	WIN %
SHORT	10	+\$227	50%
LONG	7	-\$132	14%

Ten shorts split evenly — five stopped, five hit target — and the five short winners (+\$44, +\$48, +\$84, +\$84, +\$134) dwarfed the five short losers. The long book was a disaster: seven longs, **one** winner, six stops. Its single long win (+\$113 on a 0.099% ATR rip at 15:00) is the only thing that kept LONG from being a total write-off, and one trade out of seven is noise, not an edge.



Volatility is the second axis, and it agrees. Bucketed by entry ATR-as-fraction-of-price ($\text{entry_atr} \div \text{entry_price}$):

ENTRY ATR BAND	TRADES	REAL P&L	WIN %
≥ 0.06% (0.064-0.119%)	8	+\$217	50%
0.04-0.06%	9	-\$123	22%

When the tape is genuinely moving — ATR above 0.06% of price — thrust_fast wins half its trades and makes money. When it fires into a thin, low-amplitude thrust (0.04–0.06%), it wins one in five and bleeds. This is the same amplitude-floor lesson the V5 momentum shadows taught: a thrust into quiet tape is a trap, not a signal. The two axes overlap heavily — the biggest winners (+\$134, +\$113, +\$84) are all high-ATR, and four of the six longs sit in the low-ATR trap band.

Session hour is noise here. Spread across twelve different hours with one or two trades each, no hour carries signal — the "good" hours (13:00 +\$94, 15:00 +\$113) are just where the high-ATR short winners happened to land. Don't read an hour gate into 17 trades.

In plain words: thrust_fast wins when it sells a real, fast down-move and loses when it buys a rip into quiet tape and gets faded back through its stop. The long side this week was almost pure give-back.

3. The one thing for Monday

The cleanest single split is **side**. It separates winners from losers better than the ATR band and it's mechanically simple to gate:

GATE	TRADES	REAL P&L	WIN %
SHORT only (the good bucket)	10	+\$227	50%
LONG (the rest — drop it)	7	-\$132	14%

The gate, plainly: disable thrust_fast on the long side; take short thrusts only. In-sample this turns a \$95 week into a \$227 week — a +\$132 swing — and lifts the win rate from 35% to 50%. Second choice, if you'd rather keep both sides: require entry ATR ≥ 0.06% of price, which keeps +\$217 at 50% and drops the low-vol trap band.

The honest caveat has to lead here: this is **seven long trades and two days**. A 14% long win rate on n=7 is one hot short streak away from meaning nothing, and I can't tell you from this week whether thrust_fast's longs are structurally broken or just cold. It also lines up uncomfortably well with the

live V7 desk, where LONG is $-\$733$ at 17% and SHORT is green — so the short-bias read is at least consistent across the shadow board and the real desk. But "consistent over two days" is not "proven."

WATCH – ARM SHORT-ONLY OBSERVE Don't size this yet. Flip thrust_fast to short-only in shadow and let it run another one to two weeks. If the short book stays near 50% and the long book stays underwater, that's an arm-at-1-micro candidate on the short side. If the longs recover, the whole "kill LONG" story was small-sample noise and we keep both sides. Either way the low-ATR floor ($\geq 0.06\%$) is the safer, less side-dependent tweak to carry forward now.

§18 chand_k20 — the tightest trail was the only one that stayed green

chand_k20 made $\$18.5$ over 11 trades at 45% win, all on Friday 2026-07-17 between 15:32 and 19:59 UTC. That is a rounding error in dollars, but it is the meaningful one on the board this week: it ran the exact same 11 entries as its two siblings and was the only chandelier width that finished positive — chand_k25 lost $\$36.5$ and chand_k35 lost $\$101$ on the same signals. The story of this sim is not the entry; it is the trail.

1. The money, in one sentence

Over one Friday session chand_k20 took 11 MNQ trades, won 5 of them, and finished $\$18.5$ ahead on honest fills — and because the three chandelier sims share an entry engine and differ only in how tightly the trail follows price, that $\$18.5$ is a clean like-for-like verdict that the 2.0-ATR trail beat the 2.5 and 3.5 versions this week.

All three siblings fired on the identical entry timestamps. The only thing that moved the P&L was the exit trail width: k20 $+\$18.5$ · k25 $-\$36.5$ · k35 $-\$101$ (k35 skipped two of the eleven). The like-for-like swing from tightening 2.5→2.0 was $+\$55$.

2. Why it made money — it rode shorts to a locked exit and got its longs stopped

Split by side, the whole week is one imbalance. The **shorts carried it** and the **longs bled it**:

SIDE	TRADES	WIN%	REAL_PNL
SHORT	4	75%	+\$108.0
LONG	7	29%	-\$89.5

In plain words: it took seven longs and only two of them worked; the other five were stop-outs in the 16:00–18:30 chop. The four shorts, taken later in the session, went three-for-four and trailed cleanly to profit. Net, the shorts' + \$108 just barely outran the longs' −\$89.5.

Split by **exit reason** the mechanism is even starker, and this is the real "why":

EXIT	TRADES	WIN%	REAL_PNL
CHANDELIER (trail lock)	6	83%	+\$255.5
STOP	5	0%	-\$237.0

Every single stop-out was a loser (0 for 5, −\$237). Every dollar of profit came from a trade the trail carried to a chandelier lock (+\$255.5 across six, five of them green). There is no third category — a chand_k20 trade either trails into money or gets stopped for a loss. That is exactly the behaviour you want to see: the tight trail did its one job, which is to lock a move before it hands the profit back.

CHANDELIER		
STOP		

Exit-reason contribution. The whole book is a tug-of-war between trail-locks and stop-outs; on Friday the locks won by \$18.5.

By **entry ATR%** ($\text{entry_atr} \div \text{entry_price}$) the buckets line up almost perfectly with side, so they tell the same story rather than a new one. The three low-volatility entries (atr% under 0.07) were all shorts and made **+\$62** at 67% win. The mid band (0.07–0.09) was all longs and lost **−\$58.5**. The high band (over 0.09) was a mixed +\$15. Read honestly, atr% is not an independent signal here — it is just re-telling "shorts good, longs stopped" because of when each side fired.

By **session hour**, the middle of the session was the graveyard and the last hour paid: 16:00 UTC lost −\$110 across two longs, while the 19:00 hour made **+\$143.5** across four (three shorts plus one long trail-lock). Again this is the

same longs-early-chop / shorts-late-trend shape, not a clock edge you could bank on from one day.

3. The one thing for Monday — arm the 2.0-ATR trail, it's an exit tweak not an entry bet

The tempting split is "trade the shorts, drop the longs" — shorts +\$108 at 75% versus longs −\$89.5 at 29%, an in-sample swing of about +\$89.5. **Do not do that.** It is four shorts on one afternoon, and every other honest read we have says longs are fine: the 20-sim board's grind_fast wins both sides, and the two-week thrust family was two-sided green. Friday's long-bleed is a chop-day artifact, not a side edge.

The defensible change is the one this sim was actually built to test, and it is an **exit tweak measured like-for-like on identical entries** — the strongest kind of evidence we get. Tighten the chandelier trail to 2.0 ATR (k20) instead of 2.5 (k25) or 3.5 (k35):

TRAIL WIDTH	TRADES	WIN%	REAL_PNL	VS K20
k20 — 2.0 ATR	11	45%	+\$18.5	—
k25 — 2.5 ATR	11	45%	−\$36.5	−\$55.0
k35 — 3.5 ATR	9	33%	−\$101.0	−\$119.5

The gate, plainly: set the chandelier trail multiplier to 2.0 ATR. On the three winning trail-locks the tighter trail simply banked more before give-back — the 15:54 long locked +\$95.5 versus the 2.5's +\$75, the 19:59 short banked +\$46 versus +\$19, and at 19:28 the tight trail caught a chandelier lock at +\$80.5 where both wider trails sagged back into a stop (+\$74.5 and +\$33.5). That is the give-back-a-third pattern we keep finding, and on a day of small winners locking early beat letting them breathe. **In-sample swing: +\$55 versus the 2.5 trail on the same 11 signals.**

Be honest about the size of this. It is one Friday, 11 trades, +\$18.5. And there is a real cost baked into the tight trail: on the two shorts that ran furthest, the *wider* 3.5 trail actually captured more (17:35 gave +\$89 at k35 vs +\$45 at k20; 19:11 gave +\$81.5 vs +\$56.5). That is the classic "a tight flat trail clips a runner" — it just didn't bite this week because nothing ran far. On a genuine trend day the ranking could flip, which is exactly what the standing ATR-conditional-trail finding warns about.

WATCH → ARM THE 2.0 TRAIL

The 2.0-ATR chandelier is the width to run — it is the only green sibling and it won the head-to-head on identical entries by \$55. But this is a single chop session, so arm the trail width, keep it observe-only for sizing, and hold it against a trending day before it earns real capital. Do **not** gate out the longs — that's over-fitting four shorts. The trail width is the honest edge; the side skew is noise.

§19 capit_loose — a fade that just keeps getting stopped

This one had a bad week and there's no soft way to say it: **-\$622.50** over 22 filled MNQ trades, and only 2 of them made money. That's a 9% win rate. It's the worst name on the V7 board this week, and unlike its neighbours there's no corner of the data where it quietly works.

The money, plainly. capit_loose took 22 trades, all on MNQ, and lost on 20 of them. It bled **-\$622.50**. Eighteen of those 22 trades ended on a STOP, and every single stop was a loss — the stops alone cost **-\$606.00**. The four times it actually reached a TARGET, two of those still came in red, so even its "wins" netted **-\$16.50**. There is no filled trade in this sample bigger than a \$68 loss or a \$25 win — it's small, frequent, and grinding the wrong way.

What it actually does — and why it bleeds

capit_loose is a capitulation fade: it waits for price to flush and bets on the snap-back. The problem this week is that the snap-back mostly didn't come. It kept buying the dip (or selling the rip) into moves that just carried on, and the stop caught it 82% of the time. The shape of the loss is not one bad trade — it's death by eighteen \$30 stops.

Split by side, both books are red and neither is hiding a winner:

SIDE	TRADES	REAL_PNL	WIN%
LONG	13	-\$371.50	7.7%
SHORT	9	-\$251.00	11.1%

Longs lost more in dollars, but that's just because it took four more of them — per trade the two sides bleed at almost the same rate (about -\$28 a long, -\$28 a short). This isn't a "kill the longs" story. Both directions of the fade are broken the same way.

Split by exit reason is where the whole thing is written:

EXIT	TRADES	REAL_PNL	WIN%
STOP	18	-\$606.00	0.0%
TARGET	4	-\$16.50	50.0%

Read that STOP row again: 18 stops, zero winners among them, -\$606. The stop isn't a safety net here, it's the outcome — the trade's base case is "get stopped for about \$33." And the TARGET row is almost worse news, because reaching the target didn't reliably pay: two of the four "targets" printed -\$37 and -\$21. When your winning exit still loses money on half its hits, the target sits too close and the stop sits too far — the risk-reward is upside-down before the market even gets a vote.

STOP (18)		-\$606.0
TARGET (4)		-\$16.5

Where the -\$622.50 came from: essentially all of it walked out through the stop. The target exit is a rounding error.

Does any condition save it?

I split the 22 trades every way the brief asks — side, ATR band, session hour — looking for one bucket that's green. There isn't one. By volatility at entry:

ATR_PCT AT ENTRY	TRADES	REAL_PNL	WIN%
< 0.04%	6	-\$106.50	16.7%
0.04% - 0.07%	14	-\$391.50	7.1%
0.07% - 0.10%	2	-\$124.50	0.0%

Every band loses. The two highest-ATR trades (0.074% and 0.084%) were its two biggest single losses at -\$68 and -\$56.50 — fading a fast tape is where it gets hurt most, but the quiet-tape band bleeds too, just slower. There's no volatility floor or ceiling that flips it positive; a floor at 0.07% would only have removed one small winner and kept most of the damage.

By session hour it's the same picture in miniature — thirteen different hours had trades, and **every hour is negative**. The only two green trades in the whole week landed in the 02:00–03:00 UTC window, but that window still netted -\$89.50 across seven trades, so even "only trade the hours that held a

winner" loses money. When your best-looking slice is two trades that happen to sit next to five losers, that's noise, not an edge.

The honest read on "the one thing." The single cleanest separator in the data is the exit itself – STOP (-\$606) versus TARGET (-\$16.50) – but you can't trade on that, because the exit is the result, not a condition you can check at entry. Everything you *can* gate on at entry (side, ATR, hour) is red across the board. There is no atr_pct band, no session window, and no side that turns this positive on the week. With only 2 winners in 22 trades, any "good bucket" I could carve out is one or two lucky prints – not a filter, just curve-fitting to noise.

The one thing for Monday

Retire it. This is a fade in a week where fades didn't work, and it lines up exactly with the two-week V5 history where the whole reversion/fade family bled about -\$15,800 while momentum made money. capit_loose is that same losing bet in a loose-gated V7 wrapper: it fires often, targets small, stops wide, and gets run over. The in-sample swing from just turning it off is the full + **\$622.50** — you keep every dollar it would otherwise have lost, because it kept none.

If you're unwilling to kill it outright and want a study rather than a live line, the only thing worth testing is the payoff geometry, not an entry filter: the stop is doing 97% of the damage and the target pays too little to matter (half its "targets" lost). Any rescue has to come from making the winners bigger or the stops rarer — a wider target with a structure-based stop instead of the current tight-target/wide-stop pairing. But that's a rebuild of the exit, not a gate, and I'd want a trend week and more than 22 trades before believing anything. On this week's evidence the clean call is the honest one.

KILL – RETIRE CAPIT_LOOSE No side, ATR band, or hour is green; 82% stop-out at 0% stop win-rate; +\$622.50 in-sample saved by turning it off. Thin week, but the direction is unambiguous and matches the two-week fade-family bleed.

\$20 rg_long_fast — the fast dip-buyer that mostly buys straight into a stop

rg_long_fast is a long-only, MNQ-only fast reversion grab — it buys pullbacks and tries to catch the snap back. This week it bought 51 times and lost **-\$493.5** at a

22% win rate. The honest read below: it isn't giving back winners, it's buying tops. Three-quarters of its entries go straight to the stop.

1. The money, plainly

Scored on real fills only (`fill_status='filled'`) over the two live V7 days, 2026-07-16 to 07-17: 51 trades, all LONG, all MNQ, **-\$493.5** total, averaging **-\$9.7** a trade, 22% of them green. Its ceiling — the best it could have banked with perfect exits — was only **-\$230.5**, so even a flawless exit engine would not have made this week positive. That tells you the problem is upstream of the exit: the entries themselves are underwater. This is a losing gate as it stands.

2. Why it bled

Side tells us nothing here — the "long" in the name is literal, all 51 trades are longs, so there's no long-vs-short axis to split. The real story is in how the trades ended.

The exit split is the whole autopsy. 39 of the 51 trades (76%) ended on the STOP, and those stops lost **-\$1,218** at a 5% win rate. The other 12 ran into the CHANDELIER trail and made **+\$724.5** at 75%. So the desk works exactly as designed when the entry is right — the trail lets the good ones run — but the entry is right only about a quarter of the time.

STOP · 39t · 5% win

-\$1,218

CHANDELIER · 12t ·
75% win

+\$724.5

Now the tell that settles it. On the 39 stop-outs, the average ceiling was **-\$39.9** — meaning the typical stopped trade never showed a profit at all. It bought and price went against it immediately. By contrast the 12 chandelier trades averaged a **+\$110** ceiling and gave back only about half of it. So this isn't a give-back leak or a too-tight trail — it's an entry that is simply on the wrong side of the next few minutes three times out of four. You can't tune an exit to fix an entry that's wrong on arrival.

EXIT	TRADES	REAL P&L	WIN%	AVG CEILING
STOP	39	-\$1,218	5.1%	-\$39.9
CHANDELIER	12	+\$724.5	75.0%	+\$110.4

By volatility band ($\text{atr_pct} = \text{entry_atr}/\text{entry_price}$). It loses in every ordinary band. The quiet tape it likes to fire in most (0.04–0.06%, 21 trades) is its single worst bucket at **-\$401.5**, 19% win — it's buying dips in dead, drifting tape and getting no snap-back. The one bucket that's green is the loud one, atr_pct at or above 0.08%, but that's just 2 trades for +\$247.5, which is noise, not a signal. Plain words: it fires most where there's no energy to bounce it, and the only place it made money is too thin to trust.

ATR_PCT BAND	TRADES	REAL P&L	WIN%
< 0.04%	13	-\$201.0	15.4%
0.04–0.06%	21	-\$401.5	19.0%
0.06–0.08%	15	-\$138.5	26.7%
≥ 0.08%	2	+\$247.5	50.0%

By session hour the picture is scattered and mostly thin, but one line is clean enough to lean on: the trades bunch up overnight (Europe/Asia hours), and that's where the money goes. The overnight book, everything outside 13:00–20:00 UTC, is 31 trades for **-\$663.5** at just 13% win. The US cash-session window is 20 trades for **+\$170** at 35% win. A reversion long wants a real order-flow crowd to fade into and bounce off; overnight it's fading thin drift and getting run over.

3. The one thing for Monday

The cleanest, least-cherry-picked split is the session clock. Kill the overnight book, keep only the US regular-session window.

BUCKET	TRADES	REAL P&L	WIN%
RTH 13:00–20:00 UTC (keep)	20	+\$170.0	35.0%
Overnight (drop)	31	-\$663.5	12.9%

The gate, stated plainly: only let `rg_long_fast` take a trade between 13:00 and 20:00 UTC; stand it down the rest of the day. In-sample that turns this week's **-\$493.5** into **+\$170**, a swing of about **+\$663**, and lifts the win rate from 22% to 35%.

Be honest about what this is: two days, 20 surviving trades, and +\$170 is barely above break-even. The RTH gate doesn't *find* an edge — it mostly means "trade a lot less, and only when there's a real crowd to fade." That's a real improvement in behaviour (it stops feeding the overnight drift), but it does not turn this sim into a money-maker. It turns a clear bleeder into a coin-flip.

Verdict. As it stands, `rg_long_fast` is a losing gate whose entries are wrong on arrival 76% of the time — no exit tweak can rescue an entry that never goes green. **RETIRE THE OVERNIGHT BOOK NOW**; the only piece worth keeping alive is the RTH-only slice, and that only as **OBSERVE ANOTHER WEEK** to see whether +\$170/35% is a real cash-session dip-buy or just this fortnight's noise. Don't size it. If the RTH slice can't clear break-even with conviction over the next 5–10 sessions, retire the whole thing and give its budget to the momentum family, which is where the two-week ledger says the money actually lives.

§21 `rg_long_fast_v` — the long-only grind that stopped out 25 times

This one had a losing week. On honest fills it made **-\$243.5** across 35 trades, and only 8 of them (23%) were winners. It is a long-only MNQ strategy — every single one of the 35 trades was a buy — so there is no short side to blame. It lost the plain way: it got stopped out over and over, 25 times out of 35, and those stop-outs won just 4% of the time.

1. The money, in one sentence

`rg_long_fast_v` took 35 long trades in MNQ this week and bled **-\$243.5** at a 23% win rate, because it kept buying into moves that immediately turned and hit its stop — 25 stop-outs cost it **-\$906** while its 10 clean chandelier exits banked **+\$662.5**, not quite enough to cover the damage.

The shape of the week is a wide winner, narrow loser: the average winning trade made **+\$101.5**, the average loser cost only **-\$39.1**. That is a healthy 2.6-to-1 payoff. The problem is not the size of the trades, it is the hit rate — at 23% wins, a 2.6:1 payoff still comes out behind. It needs roughly one win in three just to break even, and it isn't getting there.

2. Why it bled — what it does when it wins vs when it loses

By side: there is nothing to split — all 35 trades were LONG, all in MNQ, for **-\$243.5**. This is a long-only grind entry, so the whole story is about which longs worked.

By exit — this is the real split. The two exit doors tell the entire tale. When the trade got room to breathe and rode the chandelier trail out, it won 7 of 10 times for **+\$662.5**. When it got knocked out on the hard stop, it won 1 of 25 times and lost **-\$906**. In plain words: when a long actually goes, this thing captures it well; but most of the time the long doesn't go, it ticks against the entry, and the stop takes a small bite. Twenty-five small bites add up faster than ten good rides can refill.

STOP (25t, 4% win)

-\$906

CHANDELIER (10t, 70% win)

+\$662.5

EXIT REASON	TRADES	REAL P&L	WIN%
STOP	25	-\$906.0	4%
CHANDELIER	10	+\$662.5	70%

That exit split is honest but it is also a bit circular — "the trades that won exited on the chandelier" is almost a definition, and you can't trade on it because you don't know at entry which door a trade will use. So the useful question is: at the moment of entry, what separated the good longs from the stop-fodder? For that we look at volatility and time of day.

By entry ATR (as a % of price). This is where a real pattern shows up. Bucketed by how much range the market was carrying when it entered:

ENTRY ATR BAND	TRADES	REAL P&L	WIN%
0.04-0.06% (quiet)	18	-\$352.5	17%
0.06-0.08% (mid)	15	-\$138.5	27%
≥0.08% (lively)	2	+\$247.5	50%

The quiet-tape band is the sinkhole. When ATR was thin — 0.04 to 0.06% of price — it took 18 trades and lost **-\$352.5** at a 17% win rate. In plain words, it was buying grinds in a dead market, getting no follow-through, and paying the stop. The two trades it took in genuinely lively tape (ATR ≥0.08%) both did

fine. That echoes what the reversion book and the momentum amplitude-floor work have said all along: thin range is a trap, not a filter — a long entry needs the tape to actually be moving or the move never arrives.

By session hour. Spread across the clock it is mostly red — 12 of 14 active hours lost money. The only green came from two lone trades (13:00 UTC made **+\$320**, 16:00 made **+\$106.5**), so this is single-print noise, not a session edge. The busiest windows, the evening 18:00–19:00 block with 12 trades, netted roughly flat-to-slightly-negative. There is no honest hour gate here — the clock isn't the axis.

3. The one thing for Monday

The single cleanest cut is a **volatility floor on entry**: don't take the long unless ATR is at least 0.06% of price. That drops the whole quiet-tape band that did all the damage.

RULE	TRADES	REAL P&L	WIN%
KEEP: entry ATR $\geq$ 0.06%	17	+\$109.0	29%
VETO: entry ATR < 0.06%	18	-\$352.5	17%

Applying that one floor flips the week from **-\$243.5** to **+\$109.0** — an in-sample swing of **+\$352.5** — by simply not taking the 18 quiet-tape longs. The win rate ticks up from 23% to 29%.

Be honest about what this is. The kept side is only **+\$109** across 17 trades over a single week — barely green, thin sample, and chosen in-sample. This is not the discovery of an edge; it is a "trade less" filter that removes the clearly bad half. The remaining half isn't proven good, it just isn't proven bad yet. And the 25-day archive replay backs the caution: `rg_long_fast_v`'s ceiling record was carried almost entirely by one big day (**+\$723** over 25 days, but only **+\$514** with that day removed), so the true baseline is close to flat.

Verdict. **WATCH — WITH THE ATR FLOOR** Keep it on the shadow board, but arm the entry $\text{ATR} \geq 0.06\%$ floor and re-check next week. It does the job on the rare lively long and gives back small on stops, so the payoff shape is not broken — the fix is refusing to fire in dead tape. If a second week with the floor on doesn't hold green, retire it: a long-only grind that needs a volatility floor just to reach break-even is a thin reed to lean on, and the momentum/thrust family is doing this job better on the same tape.

§22 rg_long_20_t20 — a long-only grind gate that only won when the tape was calm

This is one of the V7 "rising-grind, long-only" sims, running a 20-bar grind trigger with a t20 target. It had a losing week: it lost **\$203.50** across 9 filled trades and only won 2 of them (22%). That is thin — nine trades is a single week of noise, not a verdict — but the shape of the nine is clean enough to say one useful thing.

The money, plainly: rg_long_20_t20 took 9 long MNQ trades this week, made money on 2 and lost on 7, and finished **\$203.50** in the red — its two winners (+\$41.50 and +\$47.00) added \$88.50 while the seven losers bled \$292. It actually came in a hair better than its own modelled ceiling (-\$203.50 real vs -\$248.00 ceiling), so this is not a fill-quality problem — the entries themselves just didn't work this week.

Question 1 — confirm the money

Scored on `real_pnl WHERE fill_status='filled'` from the V7 board, it is a clean loss: nine long entries, two winners, seven losers, **-\$203.50** for the week at a 22% hit rate. There is no short side to look at — this sim is long-only by construction (it is the rg_long family), so every one of the nine trades was a long into a rising MNQ tape. The whole "which side pays" question collapses to: the long side paid -\$203.50, and that's the entire book.

Question 2 — why it bled

Three cuts tell the story: the exit it died on, the ATR it entered at, and the hour it fired.

By exit reason — the stop is where the money went. Seven of the nine trades exited on `STOP` and those seven netted -\$176 (one of them a trailing stop that actually locked +\$47, the other six all red). One trade got yanked out on an `ABSORPTION_CUT` for -\$69 — the single worst trade of the week. Only one trade ever reached its `TARGET`, for +\$41.50. So the picture is: it hits its target almost never (1 in 9), and when it's wrong the stop takes a full -\$25 to -\$53 bite each time.



EXIT REASON	TRADES	REAL P&L	WIN%
STOP	7	-\$176.0	14%
ABSORPTION_CUT	1	-\$69.0	0%
TARGET	1	+\$41.5	100%

By entry ATR (the axis that actually separates) — this is the clean line. Both winners came in on the two calmest entries of the week: atr_pct 0.021 (+\$41.50) and atr_pct 0.042 (+\$47.00). Every single entry above atr_pct 0.045 lost — all six of them. The nastiest was the loudest: the atr_pct 0.117 entry (nearly 3x the median volatility) is the one that got absorption-cut for -\$69. In plain words: when it bought a quiet, orderly rise it held and won; when it bought into a wide, fast, volatile push it got stopped or cut out every time.

By session hour — spread too thin to trust. The two winners landed at 20:00 and 11:00 UTC; the losers are scattered across 01, 06, 07, 15, 17, 19 and 22. With one or two trades an hour there is no hour pattern here worth naming — the hour column is noise, and I won't pretend it's a gate.

Question 3 — the one thing for Monday

The single condition that best separates winners from losers is **entry ATR**, not side (there's only one side) and not hour (too sparse). Split the book at atr_pct 0.045:

BUCKET	TRADES	REAL P&L	WIN%
atr_pct ≤ 0.045 (calm tape)	3	+\$65.0	67%
atr_pct > 0.045 (loud tape)	6	-\$268.5	0%

The gate, stated plainly: only take the rg_long_20 entry when atr_pct is at or below ~0.045; stand down on anything louder. In-sample this week that turns a -\$203.50 week into a +\$65 week — a swing of about **+\$268.50** — and drops the six 0%-win loud-tape trades entirely. The honest catch: it's a 3-trade "good" bucket. That is not enough to call it an edge. It rhymes with the amplitude-floor thesis from the momentum family (thin tape / calm entries hold; loud pushes get faded and cut), which is the only reason I'd even flag it rather than bin it outright.

Verdict. **RETIRE AS-IS** — a 22%-win, -\$203.50 long-only week with one target-hit in nine is not a strategy that's working, and the calm-tape rescue rests on

three trades. Don't run it live. The one thing worth carrying forward is the observation itself: this gate only ever paid on quiet entries and lost 6-for-6 on loud ones, so if it earns another look, it comes back with an $\text{atr_pct} \leq 0.045$ floor bolted on from the start — never bare. As it stands it's a duplicate-ish member of a long-only grind family that the 25-day archive already showed bleeds standalone; this week did nothing to argue otherwise.

§23 chand_k35 — the loosest trail, and it gave the most back

This is one of three sims running the same entries through three chandelier trail widths, and it is the widest of the three. Over the honest week it took nine trades on MNQ, won three of them, and finished at **-\$101**. It is the worst of the width A/B: the tight trail (k20) squeaked out **+\$18.50**, the medium (k25) lost **-\$36.50**, and this one lost the most. Same signals, looser leash, less money. That is the whole story in one line, and everything below is just me showing you where the money went.

1. The money, plainly

chand_k35 made nine filled trades between the 15:00 and 19:00 hours, all on MNQ, and ended the week down **\$101** on a 33% win rate. Three winners paid **\$204** between them; six losers took **\$305** back. It is not a blow-up — the biggest single loss was \$68 — it is a slow bleed of small stop-outs that the three winners couldn't cover. Against its two siblings running identical entries, the ranking is clean and monotonic: the wider you let the trail run this week, the worse you did. That is the single most important fact about this sim, because it tells you the loss isn't the entry — the entries are shared — it's the exit width.

Width A/B, same entries, real_pnl: chand_k20 **+\$18.5** (11t, 45%) > chand_k25 **-\$36.5** (11t, 45%) > chand_k35 **-\$101** (9t, 33%). The tightest trail won; this one, the loosest, lost three times as much as the middle. On this week there is no case for the wide leash.

2. Why it bled — where the money actually went


Two cuts explain almost everything. The first is **side**. The longs were a disaster and the shorts were fine:

SIDE	N	REAL_PNL	WIN%
LONG	6	-\$232.0	17%
SHORT	3	+\$131.0	67%

Six longs, one of them green, for a loss of \$232. Three shorts, two of them green, for a gain of \$131. If this sim had simply not taken a single long trade this week it would have finished up \$131 instead of down \$101 — a swing of \$232. Every dollar of the loss and then some came from the long side.

The second cut is **exit reason**, and it maps straight onto the side story. The chandelier trail itself — the thing this sim is A/B-ing — actually made money: three chandelier exits for **+\$102.50**. The bleed is the hard stop: six stop-outs for **-\$203.50**. So the wide trail isn't failing because it gives back profit on runners; it's failing because the trades that never get far enough for the trail to engage just run into the fixed stop and die. When you cross side with exit, the picture sharpens to a single dark corner:

SIDE × EXIT	N	REAL_PNL
LONG · STOP	5	-\$164.0
LONG · CHANDELIER	1	-\$68.0
SHORT · STOP	1	-\$39.5
SHORT · CHANDELIER	2	+\$170.5



The only green cell on the board is short trades that lived long enough for the trail to catch them. Everything else is red. The one long that reached its trail still lost \$68 — the widest chandelier let a long rip give back until it was a loser instead of banking it. In plain words: **this sim went long into moves that immediately reversed and stopped it, and on the rare long that ran, the loose leash handed the profit back.** The shorts, by luck or by tape, were the trades that trended, and the trail did its job on them.

Volatility and time of day don't split it

I checked entry atr_pct ($\text{entry_atr} \div \text{entry_price}$) and session hour so you don't have to wonder. Every one of the nine entries sits in a tight volatility band — 0.061% to 0.115% — so there is no low-vol / high-vol line to draw; the whole book is one bucket. If anything the two green shorts were the two lowest-ATR entries of the week (0.061% and 0.066%) and the single ugliest reading, 0.115%, was the losing chandelier long — but that is one trade each, and I won't build a gate on $n=1$. Session hour is the same: winners and losers are smeared across 15:00 through 19:00 with no clean hour to keep or cut. There is no volatility gate and no time gate hiding here. The axis that separates this week is side, full stop.

3. The one thing for Monday

The cleanest split by a distance is **side**, and it happens to line up with what the width test already told us. The rule I'd write is plain: **on this sim, don't take the wide-trail long.**

Gate: drop chand_k35's long side (or just retire the sim to the tight k20 trail). Good bucket — SHORT only — was **+\$131** over 3 trades at 67% win; the rest — LONG — was **-\$232** over 6 at 17%. In-sample swing of removing the longs: **+\$232**, turning a \$101 loss into a \$131 gain.

Now the honesty, because this is one thin week. Three shorts is not an edge — it's three trades, and "shorts won this week" is exactly the kind of coin-flip that a single Monday can erase. I do not believe the short side has proven anything; I believe the long side has proven it can bleed. The stronger, less-fragile read is the width test sitting right next to it: across the same entries, tighter trail beat looser trail in a clean ladder ($k20 > k25 > k35$), and this loose variant sat at the bottom every way I cut it. That is two independent signals — the side split and the width ladder — both pointing the same direction.

So the call is not "go short next week." The call is: **RETIRE CHAND_K35**. The wide trail is the worst of three widths on shared entries, its longs are a \$232 hole, and the only thing keeping it off the floor is three lucky shorts I don't trust to repeat. Keep the tight k20 trail as the live chandelier width and let this one go. If you want to keep it breathing one more week purely to watch, gate out the long side and it at least stops digging — but the graceful answer is to kill it and stop running a leash we already know is too loose.

§24 rg_long_20_t15 — a long-only grind that gave it all back on the wild bars

This one is a long-only reversion-grind on MNQ. Over its honest first week it took 11 trades and lost **-\$75.5**, winning 45% of them. That is not a blow-up — it is a hair under breakeven — but every dollar of the loss (and then some) came from a handful of trades it took when the tape was too hot for it. Cool the entry down and it is flat-to-green; leave it as-is and it slowly bleeds.

1. The money, plainly

rg_long_20_t15 fired 11 filled trades between 07-15 and 07-17, all of them LONG and all of them MNQ, and finished the week down **-\$75.5** on real fills with a 45% hit rate (5 winners, 6 losers). There is no short book to blame and no second symbol — this is one clean long grind on the Nasdaq micro, and it lost a little money. It also happens to be the exact duplicate of rg_long_20_flow25 on this board, so treat the two as one sim, not two data points.

2. Why it bled — it wins small and clean, loses big and hot

The exit mix tells the whole story in three lines. When it reaches its **TARGET** it is perfect: 5 trades, 5 wins, +\$184.5 banked. Its losses are all stops and one absorption cut: 5 **STOP** exits for -\$191 (zero winners among them) and a single **ABSORPTION_CUT** for -\$69. So the shape is a coin-flip that pays out symmetrically on the wins but the losses are a touch fatter — a -\$75 week is what you get when your 5 stops (-\$191) plus one cut (-\$69) just outweigh your 5 clean targets (+\$184.5).

TARGET (5)		+\$184.5
STOP (5)	-\$191	
ABSORPTION_CUT (1)	-\$69	

Now the reason those stops fire. Bucket the 11 trades by how volatile the bar was at entry ($\text{atr_pct} = \text{entry ATR} / \text{entry price}$) and the split is stark. On the **quiet** bars, atr_pct under 0.04, it went +\$14.5 over 3 trades at a 67% win rate. In the **middle** band, 0.04 to 0.06, it was roughly flat at -\$9 over 4 trades. On the **hot** bars, atr_pct at or above 0.06, it lost -\$81 over 4 trades and won just 25% of them. The single worst trade on the whole sheet is the 0.117 atr_pct entry — more than double any other bar's volatility — which got absorption-cut for -\$69. A long-only grind is a mean-reversion idea: it wants to

buy a dip in an orderly tape and get paid as it grinds back. When the bar is wild, there is no orderly grind to fade — the thing keeps running and stops it out.

Session hour points the same way but is too thin to lean on: the three trades it took in the 20:00 UTC hour were all clean targets for +\$107 combined, while the scattered singles through the overnight and midday (01:00 -\$53.5, 15:00 -\$69, 17:00 -\$38, 19:00 -\$43) are where it stopped out. That is really just the ATR story wearing a clock — the calm 20:00 window happened to hold its quiet bars, and the messy hours held its hot ones. I would not build a time gate off one trade per hour.

In one line: it makes money grinding calm dips (quiet-ATR entries hit target 67% of the time) and hands it back trying to grind the wild ones (hot-ATR entries lose three-to-one).

3. The one thing for Monday — cap the entry ATR

The cleanest single split is a volatility ceiling on the entry. Skip any signal where `atr_pct` is 0.06 or higher and keep the rest:

BUCKET	N	REAL_PNL	WIN%
Keep — <code>atr_pct < 0.06</code>	7	+\$5.5	57%
Veto — <code>atr_pct ≥ 0.06</code>	4	-\$81.0	25%

That one condition is worth an **+\$81 in-sample swing**: it turns the -\$75.5 week into a +\$5.5 week and lifts the hit rate from 45% to 57%. The gate is plain — `rg_long` only takes calm bars; if the ATR is running (`atr_pct ≥ 0.06`) it stands down. It fits the mechanism, not just the numbers: a dip-grind needs a tape that grinds, and it kills the -\$69 absorption-cut outlier that was the biggest single hole in the sheet.

Be honest about what this is, though. The "good" bucket is +\$5.5 — it is breakeven, not an edge. The gate is really just "trade less, and only when it is quiet," and it is drawn on 4 vetoed trades over three days, which is far too thin to promote on. It rhymes with the amplitude-floor work we already know (thin, wild tape = stand down), so it is at least a familiar shape rather than a fresh curve-fit. But nothing here says `rg_long_20_t15` has a live edge — it says the sim is close to flat and its losses are concentrated on hot bars.

WATCH — KEEP IT ON THE BOARD WITH AN ATR_PCT < 0.06 ENTRY CEILING, BUT IT NEEDS ANOTHER WEEK OR TWO OF QUIET-BAR TRADES BEFORE IT EARNS ANYTHING MORE THAN OBSERVE-ONLY. AS IT STANDS IT IS A SLOW LEAK, NOT A PROMOTION CANDIDATE.

§25 rg_long_20_flow25 — a long-only reversion that only works in calm tape

This is a long-only, MNQ-only reversion sim on the new V7 board. Over its first honest week it took 11 trades and lost **\$75.5**, winning just under half of them (45%). It is not blown up and it is not broken — it is sitting right on the line, and one bad high-volatility trade is the difference between red and flat. There is a clean split hiding inside it, but the sample is thin enough that I want to be plain about what it does and doesn't prove.

1. The money, in one sentence

rg_long_20_flow25 made **\$75.5** of losses across 11 filled trades at a 45% win rate — five winners for a combined **+\$184.5**, six losers for **-\$260** — so it bled, but barely, and a single absorption-cut trade is most of the damage.

Worth flagging up front: this sim is an **exact duplicate** of rg_long_20_t15 — same 11 trades, same **-\$75.5** to the cent. Two names, one strategy. Whatever we decide here applies to both, and one of them should just be deleted.


2. Why it bled — what it does when it wins vs when it loses

By side: there is nothing to split. The "long" in the name is literal — all 11 trades were LONG MNQ. There is no short book to blame or credit. So the whole story is in when it goes long, and how those longs exit.

By exit reason is where the shape shows up. Every TARGET hit won and every STOP lost — a textbook symmetric reversion where the outcome is simply "did price mean-revert before it ran me over." The trouble is the third exit:

EXIT REASON	N	REAL P&L	WIN%
STOP	5	-\$191.0	0%
ABSORPTION_CUT	1	-\$69.0	0%
TARGET	5	+\$184.5	100%

The five TARGETs and five STOPS very nearly cancel (+\$184.5 against -\$191, a \$6.5 net loss). The whole negative week is essentially that one ABSORPTION_CUT for -\$69. Take that single trade out and the sim is dead flat at -\$6.5 over 10 trades at a 50% win rate. So it isn't leaking steadily — it caught one long into a rip that got absorbed and stalled, and paid for it.



By entry volatility ($\text{atr_pct} = \text{entry_atr} / \text{entry_price}$): this is the real tell, and it lines up with the exit story. The winners are calm-tape entries; the losers are the hot ones. Sorted by ATR%, the four cleanest winners all sit at or below 0.05% ATR (the 0.021, 0.025, 0.042 and 0.049 entries all hit TARGET), while five of the six entries above 0.05% ATR were losers — including the -\$69 absorption cut, which was the single hottest entry of the week at 0.117% ATR, roughly double the next.

By session hour: too sparse to lean on — 11 trades spread across 11 different hours, so every hour is $n=1$ and there's no cluster to trust. Three of its five wins happen to land in the 20:00 UTC hour, but with one

trade per hour that's coincidence, not a session edge. I'd ignore hour entirely here.

In plain words: **it wins when it fades a quiet, low-volatility drift back to the mean, and it loses when it fades a fast, high-volatility move** — exactly the failure mode reversion has shown all along. Fading into a live thrust runs it into the stop, or gets it absorbed and cut for the worst single loss of the week.

3. The one thing for Monday — gate out the hot tape

The single condition that best separates winners from losers is entry ATR%. Split at 0.05%:

BUCKET	N	REAL P&L	WIN%
calm — atr_pct ≤ 0.05%	5	+\$92.0	80%
hot — atr_pct > 0.05%	6	-\$167.5	17%

The gate, plainly: only let this sim go long when entry ATR% is at or below 0.05% — i.e. when the tape is quiet. Standing down on the six hot-tape entries turns the week from **-\$75.5** into **+\$92**, an in-sample swing of about **+\$167.5**, and lifts the win rate from 45% to 80%.

Be honest about the sample. This is 11 trades over one week, and the split is 5vs 6 — as thin as it gets. The gate is also, mechanically, "trade less": it removes more than half the book. And it isn't perfect — the hot bucket does contain one genuine **+\$69** winner (the 0.082% ATR entry), so an ATR cap would have vetoed a real win too. What gives me some confidence it's not pure noise is that the volatility split and the exit-reason split tell the *same* story, and it's the same story the whole two-week reversion book told: fades die in high-ATR tape. This is a genuinely tiny in-sample read, not a validated edge.

Verdict. **WATCH — GATE TO CALM TAPE, AND DEDUPE** Don't retire it outright: at 50% win rate ex the one absorption cut it's essentially breakeven, and the calm-ATR version is plausibly small-positive. But keep it observe-only under an $\text{atr_pct} \leq 0.05\%$ entry gate and give it another week or two of trades before it earns anything. And since `rg_long_20_flow25` and `rg_long_20_t15` are byte-for-byte the same result, kill one of the two names now — running both just doubles the board noise for zero extra information.

§26 capit_ride — bought two flushes, got stopped on both

This one lost **\$61** over just 2 trades, and neither trade was ever green for a single tick. It is the "capitulation ride" idea — buy the panic flush and ride the snap-back — and this week it bought two quiet dips that were not panics, and both walked

straight into the stop. Zero wins. There is nothing here to tune; the honest read is that the whole capitulation family failed and this sim is the thinnest, weakest corner of it.

1. The money, in one sentence

capit_ride took 2 filled trades on the V7 board, both scored on real_pnl with fill_status='filled', and lost **\$32.50** on the first and **\$28.50** on the second for a total of **-\$61** at a 0% win rate. Both were long, both exited on STOP, and both held for only about two and a half minutes before the stop caught them.

Two trades is not a sample — it is an anecdote. I will not pretend a 2-trade sim has a tunable edge. The useful signal here is not inside capit_ride, it is in the family it belongs to, which I bring in below to give the verdict some weight.

2. Why it bled — it bought calm, not capitulation

Here are both trades in full. The thing that jumps out is the ceiling column: neither trade ever went positive. The best price each one saw was already underwater, so this was not a good entry that gave back a profit — it was a bad entry that never worked for a second.

ENTRY (UTC)	SIDE	HOUR	ATR_PCT	MOVE	CEILING	REAL_PNL	EXIT
07-16 22:12	LONG	22	0.046%	-13.5 pt	-\$28.5	-\$32.5	STOP
07-17 02:17	LONG	02	0.036%	-26.25 pt	-\$54.0	-\$28.5	STOP

By side: both trades were LONG, so the long book is the whole story at **-\$61** over 2 and the short book is empty. capit_ride never took a short this week, which is itself telling — a capitulation-buy strategy only ever leans one way, and one-way books are exactly what got punished elsewhere on the board.

By exit_reason: 100% STOP. Every trade this sim took died at its protective stop, none reached a chandelier trail, a target, or a time-out. When a momentum-style sim shows all-STOP exits and no CHANDELIER exits, it means the entries never earned any room — the trade went the wrong way immediately and the stop did its job. That matches the negative ceilings exactly.

By atr_pct: both entries fired in dead-quiet tape — 0.046% and 0.036% ATR-to-price, which is about as low as the board gets. That is the core mistake. A

genuine capitulation flush is a high-volatility event: price dumps hard, sellers exhaust, and you buy the exhaustion. What capit_ride actually bought was a small drift lower in a calm market and called it a flush. There was no exhaustion to snap back from, so price just kept drifting into the stop.

By session hour: one at 22:00 UTC (the Globex reopen, first hour) and one at 02:00 UTC (Asia hours). Both are thin-liquidity windows where a small down-drift looks like a flush on the indicator but has no real capitulation behind it. Different hours, same failure, so I cannot even blame one bad session.

In plain words: **capit_ride is built to buy panic and ride the bounce, but this week it never saw panic. It saw two slow, quiet slides in illiquid hours, mistook them for capitulation, bought them, and was never once in profit before the stop.**

The family it belongs to — this is the real evidence

capit_ride is one of three capitulation sims on the board, and all three are red. Scored together on real_pnl, the whole capit family took 27 trades, lost **\$722.50**, and won only 3 of them — an 11% hit rate, and 18 of the 27 were long. capit_loose alone is the single worst sim on the entire V7 board at **-\$622.50**. This is not a capit_ride problem; it is a capitulation-buy problem, and capit_ride is just the smallest slice of it.



Every bar points the same way. Across three sims and 27 trades, the capitulation idea did not produce a single positive variant this week. That consistency — thin as any one sim is — is what turns "capit_ride lost \$61" from noise into a verdict.

3. The one thing for Monday

The cleanest condition split inside capit_ride's own two trades is the ATR-floor: both losers fired below 0.05% ATR-to-price. So the "gate" would be only take a capitulation entry when atr_pct is above roughly 0.06% — no flush without real volatility behind it. Applying that this week removes both trades and the in-sample swing is **+\$61**, taking the sim from **-\$61** to flat.

But be honest: a gate that removes 2 of 2 trades is not a filter, it is a delete button with extra steps. There is no surviving good bucket to point at — a "winners vs the rest" table needs winners, and capit_ride has none. The +\$61 swing is real but it is just "stop trading it." That is the truth, not an edge.

Verdict: **RETIRE CAPIT_RIDE**. Two trades, zero wins, never once in profit, and it sits inside a capitulation family that lost \$722.50 across 27 trades with an 11% win rate and no positive variant. The idea it embodies — buy the flush, ride the bounce — is buying quiet down-drifts in thin hours and calling them panic. If the capitulation thesis is ever worth another look, it should come back as one sim with a hard volatility floor (atr_pct well above 0.06%) and a real flush trigger (a sharp multi-ATR drop in the prior minutes), tested fresh from scratch — not by nursing this 2-trade stub. For now, take the money off the table and let the board's proven earners (grind_fast, the thrust family, the absorption-veto pair) carry the load. There is no honest edge to keep here.

§27 chand_k25 — the middle child of the trail A/B, and it bled

chand_k25 is the 25-tick chandelier variant — the momentum entry is identical to its siblings, the only thing that changes is how far the trailing stop sits behind the peak. Over its first honest week on the V7 board it took 11 MNQ trades, won 5 of them (45%), and finished at **-\$36.5**. Small loss, but it's a loss, and it lands exactly where you'd fear on the A/B ladder: the tighter k20 sibling made **+\$18.5**, this k25 lost **-\$36.5**, and the loose k35 lost **-\$101**. Wider trail, worse week. That ordering is the whole story.

The money in one sentence: chand_k25 traded 11 times on MNQ, won 45%, and finished **-\$36.5** — its winners came from short trades that rode the chandelier, and its losses came almost entirely from long trades that got stopped out.

Why it bled — the two cuts that explain it

You don't need a clever filter to see where the money went. Two splits account for the entire loss, and they point at the same trades.


By side. The four short trades made **+\$79.5** and won 3 of 4 (75%). The seven long trades lost **-\$116.0** and won only 2 of 7 (29%). So on a net loss of \$36.5, the shorts were actually carrying the sim and the longs sank it. This is the exact same limp the whole V7 board and the live desk have shown all week —

LONG bleeds, SHORT holds — so it's not a quirk of this variant, it's the room it trades in.

By exit. Only two exit reasons fired. The chandelier — the thing being A/B'd — did its job: 5 chandelier exits made **+\$126.0** at 80% win. The stop, on the other hand, is the bleed: 6 stop-outs lost **-\$162.5** at 17% win. When k25 got to trail a runner it kept most of the money; when the trade never got going and hit its initial stop, it paid full freight. The trailing logic isn't broken — the entries that never work are what feed the stop bucket.

CUT	BUCKET	REAL_PNL	N	WIN%
Side	SHORT	+\$79.5	4	75%
Side	LONG	-\$116.0	7	29%
Exit	CHANDELIER	+\$126.0	5	80%
Exit	STOP	-\$162.5	6	17%

By amplitude and clock — nothing to split on. Every one of the 11 trades entered with atr_pct at or above 0.06 (the range was 0.061 to 0.115), so there's no low-vol trap to filter out here — the whole week fired into genuinely moving tape. The amplitude floor is already effectively satisfied. Session hour tells the same side-story in disguise: the profitable window was 19:00 UTC (+\$129 over 4 trades, 75%), but three of those four were the winning shorts, so the "good hour" is really just "the hour the shorts happened." The 16:00 hour lost -\$110 on two long trades. There's no independent clock edge to lean on — it's side wearing a clock costume.



Read the chart as two views of the same \$-36.5. The top pair is the side cut, the bottom pair is the exit cut. Both times the green bar and the red bar are roughly balanced, and both times the red side is the long / the stop-out. It's not a sim that does one thing badly — it's a sim doing a fine job on half its trades and being dragged under by the long book.

What it looks like when it wins vs loses, in plain words

When k25 wins it's short, or it's a trade that got far enough into profit for the 25-tick trail to arm and then ride the move down — like the short at 19:00 that trailed out for **+\$75**, or the long that ran to **+\$75** before the chandelier clipped it. When it loses it's long and early: it buys the rip, the rip doesn't extend, and price rolls back through the fixed initial stop before the trail ever engages. Six of those, one of them a fluke where the stop caught **+\$74.5**, five of them straight red. The 25-tick trail is wide enough that on the trades that do work it gives back more of the peak than the k20 version — which is precisely why k25 sits below k20 on the ladder.

The one thing for Monday

The single cleanest split is **side**. Trading only the shorts turns this sim's week from **-\$36.5** into **+\$79.5** — an in-sample swing of about **+\$116** — and lifts win rate from 45% to 75%. The long book contributed nothing but drawdown.

Gate, stated plainly: on chand_k25, don't take the longs — trade the short side only. Good bucket: SHORT **+\$79.5**, n=4, 75% win. The rest (LONG): **-\$116.0**, n=7, 29% win. In-sample swing $\approx$ **+\$116**.

Now the honest part. This is 11 trades and the winning bucket is 4 of them. Four trades is not evidence — it's a coin that landed heads three times. I would not arm a "shorts-only" rule off this sim alone. What makes me believe it slightly is that the long-bleed shows up everywhere this week — the live V7 desk ran LONG **-\$733** vs SHORT **+\$46**, thrust_loose ran LONG **-\$38** vs SHORT **+\$500** — so the direction of the finding is corroborated even if this sim's own n is thin.

The stronger Monday move actually isn't a gate on k25 at all — it's the A/B result the sim exists to produce. k25 lost the trail bake-off to k20 (**+\$18.5** vs **-\$36.5**), and both beat k35. The graded-trail work says an ATR-conditional trail beats any flat width by a wide margin, but among the flat widths tested this week, tighter won. So: **RETIRE CHAND_K25** as a live candidate in favour of the tighter k20 trail, and keep k25 running observe-only purely as the A/B control. The shorts-only tilt is a **WATCH** — real direction, far too thin to size on.

§28 The small fry: sims too thin to judge, and two housekeeping catches

Every board has a back row — sims that fired a handful of times, or didn't fire at all, or turned out to be the same idea wearing two name tags. None of these move the needle this week, but they're worth a clean paragraph each so nothing rots quietly in the corner. Two of them are actual bugs, not verdicts.

The rule for this row is simple: with one trade, or three, there is no edge to find yet — there's just a coin that landed. I'm reporting them so you know they're alive and behaving, not because any of them earned a promotion. Where a thin sim is really a slice of a family we already judged, I say so, because a small sample sitting on top of a bad idea is still a bad idea.


rg_short_025_raw — one trade, one win, no verdict

This one took a single short on MNQ at 06:32 on the 17th, at an ATR of about 0.044% of price, and it hit its target for **\$48.5**. That's the whole story. One trade, one winner, 100% win rate — which sounds great until you remember 100% of one is exactly what a fair coin does half the time. I ran the join on `shadow_trades` against `shadow_real` filtered to filled, and there is genuinely just the one row. There is nothing here to keep or kill; it needs weeks of firing before its short-side signal means anything. Worth noting the rest of the raw-short and short-gate family has been quietly the healthier side of the book elsewhere (thrust shorts carried `thrust_loose` to +\$500 while its longs lost \$38), so a short-only reversion probe is a reasonable thing to have running — just leave it running and check back when it has 30 trades, not one.

capit_mid — three trades, two stops, one bounce, still red

The middle-tightness capitulation sim fired three times, all longs, and netted **-\$39** on a 33% win rate. The tape underneath it: a stop for -\$32.5 at 22:12 on the 16th, another stop for -\$28.5 at 02:17, then one target that paid +\$22 at 02:54. So it bleeds two full stops to claw back most of one on the winner — the classic shape of a mean-reversion idea catching a falling knife twice for every clean bounce.

On its own, three trades is a shrug. But `capit_mid` is not really alone — it's the mid setting of the capitulation family, and the family verdict is already ugly. The loose sibling, `capit_loose`, took 22 trades and lost **-\$622.5** at a 9% win rate — the single worst sim on the whole V7 board this week, a wall of STOP exits with barely a target in sight. The ride variant went 0-for-2 for -\$61. So every dial position of this idea is red right now, and `capit_mid`'s tiny -\$39 is just the least-punished slice because it barely traded. I wouldn't read its smallness as safety. The honest call is that capitulation-buying as a family is fading rips that keep running, and the thin sims don't rescue it.



The bottom two bars in that chart are the same result twice — which is the first of the two housekeeping catches.

Catch one: `rg_long_20_t15` and `rg_long_20_flow25` are the same sim

Both of these ran 11 trades and both lost exactly **-\$75.5**. That's not a coincidence — I joined the two sims on entry timestamp and side and walked all 11 trades: every single fill is identical to the cent (+17.5 / -23.5 / -53.5 / +41.5 / -33.0 / +36.0 / -69.0 / -38.0 / -43.0 / +69.0 / +20.5). One is supposed to gate on a t15 trigger and the other on a flow25 filter, but this week they produced a byte-identical tape. Either the two config knobs collapsed to the same condition, or one is silently reading the other's settings. It's not costing money — it's costing a board slot and it's double-counting a bleeder into the reversion column. Point one gate at a genuinely different condition or retire one of them.

Catch two: cb_thrust and cb_thrust_dropped fired zero trades

The two circuit-breaker thrust sims logged **no rows at all** — not zero filled, zero trades. I checked the raw `shadow_trades` table with no fill filter and both come back with a count of 0. They're on the board but they never armed. That's a dormant-sim bug, not a "no signal this week" result: a thrust variant should absolutely have had setups in a week where thrust_cont took 29 trades and thrust_loose took 43. Something in the circuit-breaker gate is refusing everything, or the sims aren't wired into the runner. Until that's fixed they tell us nothing, and worse, a dead sim looks calm on a dashboard while contributing nothing. Fix the wiring, confirm they take trades, then let them observe.

The roundup

SIM	TRADES	NET	READ
rg_short_025_raw	1	+\$48.5	WATCH one lucky short, no verdict
capit_mid	3	-\$39	WATCH thin slice of a red family
rg_long_20_t15	11	-\$75.5	DEDUPE identical to flow25
rg_long_20_flow25	11	-\$75.5	DEDUPE identical to t15
cb_thrust	0	—	FIX dormant, never armed
cb_thrust_dropped	0	—	FIX dormant, never armed

None of these change the week's story. Two are too young to judge and belong on the watch list until they've traded 30-odd times; two are the same sim double-logged and one should be retired or repointed; two are asleep and need their wiring checked before they mean anything. The one genuine warning is capit_mid — don't let its small red number hide that its whole family (loose -\$622.5 at 9%, ride 0-for-2) is bleeding on the same idea of buying capitulation into runs that don't turn.

honest money = shadow_real.db real_pnl · 2-week history + 25-day archive replay + L2 depth.db ·
the report prescribes, you decide