


MNQ pulled apart: the runs we caught, the runs we bled, and the one gate that turns +\$1.8k into +\$5.7k

38 sections, MNQ only, deep — the tape and the shadow board on honest tick-repriced fills. Every run pulled apart, every sim critiqued, the regime router and the L1/L2 hunt worked all the way through.



207 runs • 38 sections • ~35k words skeptic-verified • light

nothing armed

CONTENTS

- §1 Read this first: MNQ handed us the moves, we showed up for most, and we banked barely a third — and the whole leak has one fix
- §2 We showed up to 78% of the week's MNQ runs and only banked 35% of them — firing is not catching, and the small runs are eating us
- §3 The US-open V-reversals are a liquidity sweep read live in one number: net aggressive selling dies at a lower low, then flips — and Friday the desk already stood in the right spot for +320
- §4 The afternoon trend climbed 345 points and our long book still lost money, because we chased every push instead of buying the pullback
- §5 Most of the runs we "sat out" were a dead camera, not a bad decision — and the real money left behind is on the short side in broad daylight
- §6 Both days the shadows traded were down days — the whole game was to know that and lean short, because fading the trend is the only thing that bled every session
- §7 Momentum earns every dollar the desk has; the fade book hands it straight back
- §8 The chandelier is the whole edge — and it hands back a third of every winner
- §9 The whole week says one thing: the edge isn't more gates, it's an OFF switch that reads the room before it fires
- §10 Nothing has earned a bigger size — the board is two days old, and that's the whole answer
- §11 The tape ran 196 times this week — and it told us big MNQ moves are exhaustion reversals off a stop-run, not chases
- §12 Does the order book tell you which way? On MNQ, no — the resting book is worthless to backwards
- §13 Six gates built, six backtested — three dead on arrival, one an illusion, two with a real signal
- §14 Ride-the-flush: the one gate that pays — but it's the momentum you already trade, with a cleaner trigger
- §15 Exhaustion-reversal: the one genuinely new footprint — real direction, but the dollars aren't there yet
- §16 Coiled delta: the +\$1,340 "winner" that I could not reproduce — and why I'm telling you it's dead
- §17 How we grab it from Monday — one sharper trigger to shadow, one footprint to develop, and two habits to stop
- §18 Every dollar we lost this week came from an engine firing into the wrong regime — a router that stands the wrong engine down turns +\$1,830 into +\$5,740
- §19 grind_fast, trade by trade: it enters mid-run and lets the chandelier ride — then hands a third of it back
- §20 thrust_cont — the quiet winner that just takes its target
- §21 thrust_loose made money the same way all week: the short side takes its target, the long side just donates
- §22 thrust_aligned made \$194, but the short side made \$293 and the long side handed almost half of it back

- §23 thrust_fast made \$95 — but only because the short side carried a long side that's pure tax
- §24 abs_veto_55s: the short side and a bit of amplitude do all the work; the quiet-tape longs are dead weight
- §25 abs_veto_50s — the short-side absorption sim that quietly banked a clean \$443
- §26 abs_veto_60s: it hit its target four times out of six and banked \$127 — but it's the slowest, thinnest cut of the absorption-veto family
- §27 chand_k20: dead flat — the tight trail banked \$255, the hard stop bled \$237, and four shorts nudged it green
- §28 chand_k25 — the middle child of the trail A/B, and it bled
- §29 chand_k35: the loosest trail lets longs run into the ground while shorts quietly print
- §30 capit_loose catches a falling knife and gets stopped 4 times out of 5 — the stop is the whole loss
- §31 capit_ride: it bought two flushes, both kept flushing, and it never saw a single green tick
- §32 rg_long_fast keeps buying dips that keep falling — the stop eats it alive
- §33 rg_long_fast_v (#23): it buys the dip, and one time in four the bounce comes — the other three it gets run over
- §34 rg_long_20_t20 — a long-only grind gate that only won when the tape was calm
- §35 rg_long_20_t15 — a long-only grind that gave it all back on the wild bars
- §36 rg_long_20_flow25: a calm-tape long that got dragged under water every time the ATR woke up
- §37 The small fry: sims too thin to judge, and two housekeeping catches
- §38 Eight hunches, walked to a verdict: the exit tweaks are sideshows, direction-by-regime is the only lever that held

§1 Read this first: MNQ handed us the moves, we showed up for most, and we banked barely a third — and the whole leak has one fix

207

real MNQ runs this week — a market moving in size, both ways

78% → 35%

runs we fired into · runs we actually banked

+\$1.8k → +\$5.7k

the book, if a regime gate stood the wrong engine down

This is MNQ, deep — the tape, not the desks. The market printed 207 clean runs across the three days. Our shadow board fired in the right direction on 78% of the ones it was alive for and looked like a machine that shows up. Then you follow the money to honest fills and it banked only 35% of them. The gap between pointing the right way and keeping the money is the whole report — and it has a single cause and a single fix.

The leak. We don't lose by missing the runs — we're in most of them. We lose two ways once we're in: we fire into the small 20-to-40-point runs that are really chop and get stopped before they pay (that bucket is a net loser even pointed the right way), and — the big one — we fade runs that are already moving. In-direction trades made +\$4,900 this week; fading live moves handed back −\$3,405 of it. Nearly three-quarters of the gross edge, eaten by trades that pick a fight with the tape.

The fix. Tag every trade by the regime it fired into and the book splits perfectly in two: every winning bucket is an engine trading with the move, every losing bucket is one trading against it — there is no third case. The same entry logic made +\$293 short and lost −\$99 long in the same week; the strategy was never the variable, the regime was. Stand the wrong-regime engines down and the book goes from its real +\$1,830 to roughly +\$5,740. That is the regime router, and it is the single highest-value thing on the desk.

Where new money would come from. I hunted the L1 aggressor tape and the L2 book for genuinely new gates and backtested six of them. Most died; I killed my own best-looking one when the skeptic couldn't reproduce it. Two survived: a cleaner microstructure trigger for the momentum we already run, and one genuinely new footprint — fading heavy aggression that stalls into a wall — with real direction but sub-cost dollars today. No free gate to arm tomorrow, but a real lead and hard proof of two habits that lose money (chasing delta, trusting the book's lean).

Monday, in one breath: the runs are there and we're in them — stop fading live moves and stop firing into the 20–40-point chop; build the regime router shadow-first (it's the +\$3,900 lever); shadow the flush-follow trigger and keep developing the exhaustion-reversal footprint. Nothing arms on one week. Everything below is the evidence, one run and one sim at a time.

WHERE THE MONEY RAN — EVERY MNQ RUN, PULLED APART

§2 **We showed up to 78% of the week's MNQ runs and only banked 35% of them — firing is not catching, and the small runs are eating us**

Here is the list you asked for: every real MNQ run this week, and whether our shadow board took it. The board fired in the right direction on 93 of the 120 runs it was alive for — that reads like a machine that shows up. But when you follow the money through to honest, tick-repriced fills, we only turned 42 of those 120 runs into a profit. The other side of that coin is the important one: on 51 of the runs we DID fire into, in the right direction, we still got chopped out at a loss before the move paid us. Showing up in-direction and actually banking the move are two completely different things, and the whole week lives in that gap.

What the week actually threw at us

Working from the run list, MNQ printed 207 genuine runs across the two-and-a-bit days of tape (I threw out the three runs tagged longer than 16 minutes — those are overnight-gap artifacts on thin tape with missing minutes, slow drifts dressed up as runs, not something a momentum sim could ever ride). The tape leaned short: 123 of the 207 real runs were DOWN moves and only 84 were UP. It was also a back-loaded week — Wednesday the 15th only gave us 15 runs in the few hours of tape we have, then Thursday exploded to 99 runs and Friday delivered another 93. This was not a quiet week; MNQ was moving in size, over and over, in both directions.

By magnitude, most of what MNQ offered was small. 75 of the 207 runs were in the 20-to-40-point bucket, another 71 were 40-to-60 points, 45 were 60-to-100 points, and 16 were the real monsters at 100 points or more — topping out at a 214-point rip on Friday afternoon. So the shape of the opportunity was a fat pile of small-to-middling moves with a thin tail of huge ones. That shape matters enormously for how we should be trading, and I'll come back to it, because it turns out the pile and the tail behaved like two different markets.

Fair-fight note. The shadow board only came alive at 16:49 UTC on Thursday and ran to 20:57 UTC Friday. Everything before that — all of Wednesday and the first chunk of Thursday's 99 runs — happened with no sims on the field, so scoring us on those would be dishonest. Every participation number below is measured only against the **120 real runs that fell inside the board's live window**. That is the fair scoreboard, and it is the one I'm holding us to.

The participation scoreboard: caught, sat out, and — the honest column — banked

The rule for "caught" is simple and generous: did any shadow sim open a position in the run's direction — LONG for an UP run, SHORT for a DOWN run — at any point between the run's start and five minutes past its peak? That is a wide net; it forgives a late entry and it only asks that we pointed the right way. Even against that generous bar, here is where we landed:

Of the 120 in-window MNQ runs:

- We **caught 93** (78%) — a sim fired in-direction inside the window.
- We **sat out 27** (22%) — nothing fired in-direction at all.
- Of the 93 we caught, we **banked only 42** and got **chopped on 51**. Firing in-direction paid off honest money less than half the time.
- Net across the whole week: we made money on **42 of 120 runs — 35%**.

That 93-vs-42 split is the headline of this section and the thing I most want you to sit with. On paper we are a momentum machine that finds the move 78% of the time. In reality — in honest, tick-repriced, force-kill-and-slippage dollars — we walk away a winner on barely a third of the runs the market handed us. The 51-run gap between "fired the right way" and "actually kept the money" is not bad luck; it is the entry firing into a move that then shakes us out at our stop before continuing, or the exit giving back the whole thing. Either way, pointing the right way is worth very little on its own.

The direction split tells the same story from both sides. On the 46 UP runs we caught 39 and banked 18. On the 74 DOWN runs we caught 54 and banked 24. The board is not lopsided by direction — it shows up roughly equally long and short — and it converts at almost exactly the same weak rate either way (46% of caught UP runs banked, 44% of caught DOWN runs banked). This is

not a "we can't trade shorts" problem or a "we can't trade longs" problem. It is a firing-into-noise problem that hits both directions equally.

The number that changes how we should trade: small runs bleed, big runs pay

When I split the banked money by run size, the whole week snaps into focus. Adding up the honest in-direction P&L of every sim that fired into each run, bucketed by how big the run was:

RUN SIZE	RUNS IN WINDOW	CAUGHT	BANKED	CONVERT %	IN-DIR P&L
20-40 pt	42	34	11	26%	-\$142
40-60 pt	44	32	15	34%	+\$374
60-100 pt	23	16	8	35%	+\$1,202
100+ pt	11	11	8	73%	+\$4,623

Read that bottom row and then the top row. The 20-to-40-point runs — the single biggest pile of opportunity, 42 of them — are a net **-\$142** loser even when we point the right way. We caught 34 of them and only 11 paid. Those small runs are exactly the chop that shakes an entry out and hands it a stop-loss before the tiny move is done. Meanwhile the 100-plus-point monsters are a different animal entirely: we caught all 11 of them, banked 8, and they printed **+\$4,623** — the overwhelming majority of the board's entire in-direction take. The 60-to-100 bucket adds another **+\$1,202**. The edge this week is not spread across the runs; it is concentrated almost entirely in the fat, sustained moves, and it is being partially bled back out by the small ones.

20-40pt **-\$142**

40-60pt **+\$374**

60-100pt **+\$1,202**

100+pt **+\$4,623**

Where the money really leaks: fading the runs

The participation view above only counts sims firing WITH the run. But the board also fires AGAINST runs — a SHORT opening in the middle of an UP rip, a LONG catching a falling knife on a DOWN move. When I tag every filled trade on the board by its run context, the split is brutal and clean:

- **289 trades fired in-direction into a run** → **+\$4,900** honest, 144 winners to 145 losers.
- **111 trades faded a run** (fired counter to a move that was underway) → **-\$3,405**.
- 14 trades fired with no run in sight (pure chop) → a rounding-error **+\$334**.

The whole board netted **+\$1,830** across 415 trades this week (162 winners, 253 losers — a 39% hit rate carried by a handful of big momentum wins). But look at how that number is built: the in-direction momentum trades made **+\$4,900**, and then we handed **\$3,405** of it straight back by fading runs that were already moving. Nearly three-quarters of the board's gross momentum edge is being eaten by trades that pick a fight with a live move. That is the reversion-into-a-grind bleed we have been chasing for weeks, and here it is again in cold numbers: the losses are not random, they cluster precisely where a sim shorts a rally or buys a sell-off.

The runs, and whether we took them — the top 30 by size

This is the list. Sorted biggest first, every one of these fell inside the board's live window. "#Fired" is how many in-direction sims opened during the run; "Banked" is the honest tick-repriced P&L of those fills summed together. Green rows we kept money on, red rows we either sat out or got chopped.

RUN START (UTC)	DIR	SIZE (PT)	MINS	#FIRED	BANKED
Fri 13:57	UP	214.0	13	3	+\$563
Fri 13:25	DOWN	183.2	15	4	+\$143
Fri 13:41	UP	162.8	15	5	+\$1,328
Thu 19:49	UP	142.0	10	4	-\$104
Fri 08:06	DOWN	141.8	14	14	+\$1,065
Fri 15:41	UP	128.5	15	10	+\$1,038
Fri 14:55	DOWN	125.2	12	1	+\$7
Fri 14:32	UP	121.8	8	1	-\$79
Fri 19:52	DOWN	116.0	12	3	+\$6
Fri 14:41	DOWN	111.5	13	2	-\$69
Fri 19:26	UP	101.8	4	10	+\$726
Thu 17:31	DOWN	99.5	14	2	+\$21
Fri 15:31	UP	88.5	9	8	-\$237
Fri 08:43	UP	86.0	15	4	+\$156
Fri 16:08	UP	86.0	9	0	sat out
Fri 17:31	DOWN	85.8	14	8	+\$545
Fri 15:16	UP	83.5	14	6	-\$408
Fri 01:03	DOWN	81.8	15	2	+\$235
Fri 18:56	DOWN	80.2	15	8	+\$628
Fri 06:09	UP	78.0	15	6	+\$331
Fri 16:18	DOWN	78.0	7	0	sat out
Fri 03:56	DOWN	75.0	14	0	sat out
Fri 18:45	UP	73.2	10	3	+\$36
Fri 16:59	UP	72.5	12	2	-\$102
Fri 15:57	DOWN	67.8	10	0	sat out

RUN START (UTC)	DIR	SIZE (PT)	MINS	#FIRED	BANKED
Fri 18:31	DOWN	67.0	13	1	-\$47
Fri 19:31	DOWN	66.8	9	0	sat out
Fri 11:55	DOWN	66.2	15	8	+\$351
Fri 14:11	DOWN	65.8	4	0	sat out
Thu 18:23	DOWN	64.5	13	2	-\$7

The big winners are unmistakable and they carry the week: the 162.8-point Friday-lunchtime UP run banked **+\$1,328**, the 141.8-point overnight DOWN move at 08:06 pulled 14 sims in and banked **+\$1,065**, the 128.5-point 15:41 rip banked **+\$1,038**, and the sharp four-minute 19:26 pop banked **+\$726** off 10 sims. When the move is real and sustained, a crowd of sims piles in and we keep the money. That is the machine working exactly as designed.

But scan the red. The 142-point Thursday-evening UP run — a genuine monster — had four sims fire the right way and still came out **-\$104**: we found it, pointed at it, and got shaken out before it paid. The 83.5-point 15:16 UP run is worse: six sims fired long into a real move and the honest tally is **-\$408**, the single most expensive "we caught it and lost anyway" of the week. And then the pure sit-outs on sizeable moves — the 86-point 16:08 UP run with zero fires, the 78-point 16:18 DOWN, the 75-point overnight 03:56 DOWN, the 67.8-point 15:57 DOWN — several of these clustered on Friday afternoon between roughly 15:50 and 16:20 UTC, a stretch where the board simply went quiet while MNQ kept handing out 60-to-90-point moves. That is a real hole, not noise: a window where our entry conditions stopped triggering right as the tape stayed tradeable.

What this means and the Monday move

The week says three concrete things. First, our momentum edge is real but narrow — it lives almost entirely in the 60-point-plus runs, which we catch well (16 of 23 in the 60-100 bucket, all 11 of the 100-plus) and bank on. Second, the 20-to-40-point runs are a trap: 42 of them, a net **-\$142** even in-direction, small moves that exist mainly to shake an entry into its stop. Third, and biggest, the fade trades — 111 of them at **-\$3,405** — are quietly handing back most of what the momentum trades earn. We are not failing to show up;

we showed up to 78% of runs. We are failing to convert, and we are actively bleeding by betting against live moves.

Monday move. Two changes, both aimed at the convert-not-just-fire gap. (1) **Put an amplitude floor on entries** — stand down when the developing move is small (the run-size buckets say sub-40-point moves are net negative even when we're right, so a live estimate of expected move size or an ATR/thrust-magnitude floor should gate them out and stop us chopping in the small stuff). (2) **Hard-veto counter-move entries** — no new SHORT while an UP run is live and no new LONG while a DOWN run is live; that one rule, applied to this week's tape, would have removed the **-\$3,405** of fade losses and left the **+\$4,900** momentum book almost untouched. Separately, chase the Friday-afternoon dead window (roughly 15:50–16:20 UTC) where we sat out four sizeable moves in a row — the entry gate stopped triggering while the tape was still moving, and that is a missed-money hole worth reproducing and widening the gate for.

VERDICT: WATCH — WE FIND THE RUNS, WE DON'T KEEP THEM. KILL THE FADES AND FLOOR OUT THE SMALL STUFF BEFORE WE CALL THIS AN EDGE.

§3 The US-open V-reversals are a liquidity sweep read live in one number: net aggressive selling dies at a lower low, then flips — and Friday the desk already stood in the right spot for +320

Two mornings running, the 13:30–14:10 UTC window printed the same thing: a hard down-flush that undercut its own low by a few points, and then a violent snap back up worth 150–214 points. This is not noise and it is not two lucky coincidences. Both turns carry an identical, readable footprint in the aggressor tape — sellers keep hitting the bid all the way down, but at the final low they have almost no force left, and a single burst of buying flips the whole thing. Friday one of our reversion-grab longs actually stood on that exact spot and banked +320 points on honest fills. Thursday the same setup printed and we had no shadow board running yet to see it. This section pulls both apart run by run and hands you the one trigger to be there for the next one.

What actually happened on the tape — the shape of the two mornings

Take Friday 07-17 first, because it is the cleaner of the two and it is the one the shadow board saw. Price opened the window chopping around 28,610, spiked up to a high of 28,685 at 13:32:00, and then fell out of bed. From 13:32 to 13:39 it bled straight down to 28,457 — that is the DOWN run the run-list logs as "13:25 DOWN 183pt." It bounced weakly to 28,492, then made a second, lower push and printed the real low of the morning at **28,408.25** at 13:44:00. That second low is the tell: it undercut the 28,457 low by 49 points. From there it never looked back — 28,447 at 13:45, 28,528 at 13:50, 28,614 at 13:56, 28,695 at 14:00, and a peak of 28,776 by 14:03:30. That is the two UP runs stitched together: "13:41 UP 162pt" and "13:57 UP 214pt." A clean V, roughly 320 points of range from the sweep low to the peak in nineteen minutes.

Thursday 07-16 is the same skeleton one day earlier. High of the window was 29,527 at 13:32, then a staged flush: 29,402 by 13:35, a bounce, then lower — 29,265 at 13:50:00 on a big sell burst, a chop back to 29,343, and then a marginally lower low of **29,257.75** at 13:55:00. That undercut the 13:50 low by about 8 points — a shallow sweep, but a sweep. From there: 29,320 at 13:56, 29,405 at 14:02, and a grind up to 29,494 by 14:17. That is the run-list's "13:56 UP 157pt." Same V, same window, same day-of-week rhythm.

Run 1 — 07-17, the down-flush (DOWN 183pt): sellers in full control, cumulative delta craters

This leg was genuine, aggressive, one-directional selling — not a fake-out. I summed the aggressor tape (L1 buy volume minus sell volume) in 30-second buckets and ran it cumulatively from the 28,685 high. It fell off a cliff: -948, then -1,666, -2,455, -3,853, and by 13:34:00 it was at -5,303, with price at 28,554. It kept grinding to a low cumulative reading of **-5,625** at 13:39:00 as price hit 28,463. In plain words: over those seven minutes, net aggressive selling outweighed buying by roughly 5,600 contracts, and price obediently fell about 220 points. Nothing clever to do here except not be long — which, as the scorecard below shows, is exactly where we lost money buying the 28,681 high right before the drop.

Run 2 — 07-17, the sweep + snap (UP 162pt): the divergence that marks the turn

Here is the whole edge in two numbers. At the 13:39 low of 28,457 the cumulative aggressor delta was **-5,625**. Price then made a lower low — 28,408 at 13:44:00, 49 points cheaper — but at that lower low the cumulative delta had climbed all the way back to just **-1,309**. Price went down; net selling force went up (less negative). That is a textbook bullish delta divergence: the second leg down was made on a fraction of the selling pressure of the first. The sellers were spent.

You can watch the absorption happen bucket by bucket. The 13:42:00 bucket sold price from 28,483 to 28,435 — 48 points — on a -690 delta. The next four buckets (13:42:30 to 13:44:00) had positive-ish deltas of +148, +276, -134, +359 while price only slipped another 27 points to the 28,408 low. Sellers were still hitting, but each hit moved price less and less: displacement per unit of selling collapsed. Then the flip: 13:44:30 printed +411, 13:45:00 printed **+783**, and cumulative delta crossed back through zero at 13:45:30. That zero-cross is the snap. Price left 28,408 and did not return.

The L2 book agrees but is the junior partner here — MNQ's book is thin, only about 25-45 contracts stacked across all five levels a side, and it flickers fast. Through the 13:42-13:45 base the book imbalance sat modestly bid-heavy (+0.10 to +0.22), and at 13:44:15 the best bid showed 28,411 x7 holding while price probed the low. There was no giant visible wall to point at; the "wall" was the aggressor tape refusing to make progress. That is the honest version: on micro MNQ, absorption is something you read in the delta, not something you see resting in the book.

Run 3 — 07-17, the continuation (UP 214pt): the desk's easy money, and it took it

Once 28,408 held and delta flipped, the up-leg was a straight trend — 28,528, 28,614, 28,695, 28,735, 28,776 — a 214-point continuation that the run-list logs separately as the 13:57 UP run. This is the part our momentum gates are built for and, as the scorecard shows, they fired repeatedly and banked it. The reversal (Run 2) is the hard, high-value entry; the continuation (Run 3) is the

reward for having been early, plus a second, lower-risk chance to climb aboard once the trend is obvious.

Run 4 — 07-16, the double-capitulation flush (DOWN legs): two sell spikes, the second one hollow

Thursday's flush is the same story told with sharper spikes. There were two obvious capitulation buckets. The first: 13:48:00 printed a $-2,153$ delta and drove price to 29,266. It bounced. The second: 13:55:00 printed a $-2,115$ delta and drove price to 29,257.75 — a lower low by 8 points. But look at the cumulative delta at each: at the 13:48 low it read -697 ; at the 13:55 sweep low it read $+42$ — actually positive on the session-window running total. Price made a lower low; cumulative delta made a clearly higher low. Same divergence as Friday, printed on two clean capitulation spikes instead of a slow grind. The sweep of the 13:48 low ran the stops sitting under it, found no follow-through selling, and reversed.

Run 5 — 07-16, the snap (UP 157pt): the book flip you can see

Thursday is where the L2 book actually earns its keep. Right at the 13:55 low the book snapped from ask-heavy to bid-heavy in fifteen seconds: at 13:55:15 the imbalance was -0.25 (sellers stacked on the offer, best bid 29,276 x2 against best ask x3), and by 13:55:30 it had flipped to $+0.18$ (best bid firming). On the tape, the $-2,115$ capitulation bucket was immediately answered by $+705$ (13:55:30) and $+946$ (13:56:00). Price went 29,257 → 29,320 in ninety seconds and then trended to 29,494. The honest caveat on this run: our shadow board did not start recording MNQ until 16:49 UTC on 07-16, so no sim was running when this printed. This is a tape-only autopsy — we could not have taken it, but it confirms the footprint is repeatable, not a one-day fluke.

Did the shadows catch it? Friday, yes — and it paid

Every number below is honest tick-repriced `real_pnl` from `shadow_real` where `fill_status='filled'`, not optimistic bar marks. The story splits cleanly: we bled buying the top and chopping the pre-flush, made the down-leg short, and then stood exactly on the reversal and banked the up-leg.

SIM	SIDE	ENTRY (UTC)	ENTRY PX	EXIT REASON	
thrust_loose / grind_fast / abs_veto_50s	LONG	13:31:00	28,681	STOP	
capit_loose ×2, thrust_fast	SHORT	13:23-13:29	~28,605	STOP	-3
thrust_fast	SHORT	13:33:00	28,576	TARGET	
grind_fast	SHORT	13:36:00	28,524	CHANDELIER	
rg_long_fast ×2	LONG	13:45:00	28,447.75	CHANDELIER	
grind_fast	LONG	13:53:00	28,589	CHANDELIER	
thrust_loose / thrust_cont	LONG	14:00:00	28,738	TARGET	
grind_fast	LONG	14:04:00	28,718	CHANDELIER	

The headline: `rg_long_fast` , our reversion-grab long, entered at 28,447.75 at 13:45:00 — one minute after the 28,408 sweep low, right as cumulative delta crossed zero — and held it on the chandelier all the way to 28,717 for **+320 points on honest fills**. Two variants took it. That is the desk doing exactly the right thing at exactly the right spot. The chandelier exit is not the weak link here; it carried a 270-point winner without giving it back. The weak links are upstream: we lost roughly –290 points across seven sims before the turn — three of them literally buying the 28,681 high — and only the reversion-grab's own logic, not any conscious "this is an exhaustion reversal" trigger, put us on the low.

The repeatable footprint, in one picture

At both turns, price made a lower low while net aggressive selling made a higher low. The bars show cumulative aggressor delta at the first low versus the swept (lower) low — the gap is the exhaustion.

07-17 first low 28,457  **-5,625**

07-17 sweep low 28,408  **-1,309** → selling force gone

07-16 first low 29,266  **-697**

07-16 sweep low 29,258  **+42** → already net-positive

Say it as a checklist, because that is what makes it armable. One: a multi-minute down-flush drives running aggressor delta deeply negative. Two: price makes a marginal new low that undercuts the prior swing low by a handful of points to a few tens of points — the stop-run. Three: at that new low, running cumulative delta is meaningfully higher (less negative) than it was at the prior low — the second push down carried far less selling. Four: within 20–40 seconds a single fat positive-delta bucket (+700 to +1,900 on these two) flips the tape and running delta reclaims zero. Five: price snaps 150–210 points. On MNQ the L2 book is a supporting witness — when it does something visible it is an ask-heavy imbalance at the low flipping bid-heavy in fifteen seconds (07-16 went -0.25 to $+0.18$) — but the primary read is the delta divergence, because the book is too thin to lean on alone.

The Monday move — arm the divergence, not the knife

Build a US-open exhaustion-long trigger (13:30–14:15 UTC window), gated on aggressor-delta divergence. Track running cumulative aggressor delta in 20–30s buckets and the last swing low. Arm a long only when all four conditions stack: (a) price prints a new low that undercuts the prior swing low, (b) cumulative delta at that new low is at least ~30% less negative than at the prior low, (c) the very next 20–30s bucket prints positive delta / running delta reclaims zero, and (d) optionally the book imbalance flips from ask-heavy to bid-heavy. Enter on the delta-flip reclaim — the 28,447 / 29,265 spot where `rg_long_fast` already stood — stop just under the sweep low, and hand the winner to the existing chandelier, which carried +320 cleanly. Critically, the trigger must NOT fire on the *first* low: that is the knife, and it is where we lost -290 points buying tops. The divergence — the higher-delta second low — is the whole filter. Backtest it across the ~10 sweep-shaped days in the run-list before arming live, and start observe-only.

The prize is concrete. Friday the reversion-grab stumbled onto the right spot and made +320; a purpose-built divergence trigger would have been there deliberately, and Thursday — where we had no coverage at all — it would have caught the 157-point run we completely missed. Two mornings, same footprint, roughly 470 points of reversal we either half-caught by accident or did not see. This is the most repeatable, highest-value pattern in the week's tape.

KEEP & BUILD The exhaustion-reversal footprint is real, quantified, and repeatable: lower price low + higher cumulative-delta low + a positive-delta flip. The desk already banks the up-leg once it is on it (+320 honest). The gap

is a conscious entry trigger so we stand on the sweep low on purpose, not by luck — and never buy the first knife.

§4 The afternoon trend climbed 345 points and our long book still lost money, because we chased every push instead of buying the pullback

After the 14:00 reversal, MNQ walked up all afternoon in a clean staircase: a low of 28,695 at 14:00 became a high of 29,040 by 17:31 — about 345 points of trend in three and a half hours, in a run-then-rest-then-run rhythm. This was the friendliest tape we get. And yet the whole afternoon long book — 54 filled long shadow entries between 14:00 and 17:40 — netted **-390** on honest tick-repriced fills: 14 winners worth +1,884, drowned by 40 losers worth -2,274. The trend paid, and we handed most of it back. This section walks the runs in order, shows what a continuation leg actually looks like on the tape versus the reversal above it, names which sims rode and which gave it back, and points at the one trigger change that turns a losing long book on an up day into a winning one.

The staircase, run by run

Here is the afternoon as the tape drew it. Each up-leg starts with a shallow dip into rising support, then a fresh aggressor push takes out the prior minor high. That is the continuation footprint, and it repeated six or seven times:

RUN	FROM→PEAK	SIZE	
14:32 UP	28,771 → 28,893	+122	only grind_fast, and it entered at 14:41 ne
15:16 / 15:31 UP	28,692 → 28,843	+88	six sims bought the 15:
15:41 UP	28,785 → 28,948	+129	grind_fast bought the 15:42 dip; thrust cluster joined th
16:08 UP	28,859 → 28,951	+86	nobody — zero long entries ope
16:44 UP	28,900 → 29,027	+57	grind_fast rode it from a 16:32 d
16:59 UP	28,970 → 29,047	+73	thrust + grind chased the 16:
17:21 / 17:26 UP	28,984 → 29,040	+58	ten sims bought the 17:26 push at 29,03

Read that column of results top to bottom. The only genuinely green rows are 15:41 and 16:44 — the two legs where a sim was already long from a pullback low before the push happened. Every leg we entered on the push is red. The trend was up the whole time; our timing inside it was upside-down.

What a continuation leg looks like on the tape — and why it is not the reversal

The 14:00 reversal and the afternoon continuations feel the same on a chart — both are up — but the footprint is different, and the difference is the whole game. The reversal off 28,695 was a fresh impulse: heavy two-sided volume (2,245 contracts in the 14:00 bar, 1,800 in 14:01) with buy-aggressor bursts of +15%, +23%, +21% imbalance punching a market that had just bottomed. There was no prior high nearby to lean on; it was a violent turn, and our reversal longs (thrust_loose and thrust_cont at 14:00, grind_fast at 14:04) caught it cleanly for +176 and +209.

A continuation leg is quieter and structurally different. Take the 15:41 run, our best. Price had pushed to 28,833 at 15:40, then pulled back a shallow 48 points to a 28,785 low at 15:42. That dip is not a reversal — it is a breather into rising support. Then the fresh push: through 28,833 by 15:47, through the prior swing high of 28,893 (set back at 14:40) around 15:54, and on to 28,948 by 16:01. The buy aggression that confirmed it did not arrive until the 15:54–15:56 window (+18%, +14%, +15% imbalance) — after the shallow pullback had already resolved. The signature is: shallow dip into support → reclaim → break the prior minor high → then the aggressor burst. The 16:44 leg drew the same shape: a rest around 28,900 through 16:31–16:43 on thinning volume, then a push through 28,974 and the prior 29,008 high up to 29,027, with the confirming buy burst (+31% at 16:51) landing late.

The tell that separates a continuation from a reversal: a continuation does not need a big imbalance to start. It grinds up on mixed, modest tape (the 15:43–15:53 minutes ran between -12% and +16% imbalance while price rose) because it is riding an *existing* order-flow bias, not creating a new one. If we wait for the reversal-sized aggressor burst to confirm, we are always buying the top of the leg, not the base of it.

The give-back: chandelier winners that weren't, and the honest-fill trap at 15:32

The clearest lesson of the afternoon is buried in the 15:32 cluster. Six sims — chand_k20, chand_k25, chand_k35, thrust_loose, abs_veto_50s, abs_veto_55s — all bought the 15:31 push at exactly 28,772. On the one-minute bar-marks they looked like a mixed-to-great outcome: two tagged TARGET with a +144 ceiling, one chand_k35 showed a +234 ceiling. On the honest 250ms tick repriced, every single one scored **-68**. Identical. Why? Their protective stop sat

near 28,738, and the shallow pullback at 15:35 dipped to a 28,741 low that, on the tick tape, pierced the stop — then the market turned and ran to 28,948 without them. The bar-marks never saw the wick; the tick tape did.

That is the afternoon in one trade cluster. The continuation's shallow pullback is exactly the move that shakes out a late or tight entry, and the trend then continues without you. Six "target/chandelier" trades that a bar-backtest would have logged as wins were tick-tape losers for -408 combined. Never score this stuff on minute bars — the 15:32 cluster would lie to us every time.

Now put grind_fast next to them. grind_fast bought the same pullback low — but at 15:42, at 28,785, after the shakeout dip had already happened — and rode to 28,889 for +242, the single best long of the day. Same leg, same support, ~90 seconds and ~13 points of entry difference, and the outcome flips from -68 to +242. The edge was never in the direction; it was in entering after the shake, not into it.

Who rode and who chased — the family scorecard

Split the 54 afternoon longs by family and the pattern is unambiguous:

grind_fast (pullback rider)+396

thrust* (breakout chaser)+118

abs_veto*+2

rg_long* (reversion)-351

chand* (chandelier)-497

grind_fast was the only reliable earner because it is the only family that consistently enters on the dip: 14:04, 15:42, 16:03, 16:32 — all pullback lows, all before the impulse, and it collected +209, +242, +76, +151 on those. The thrust family netted a thin +118 across 19 trades: it caught the 15:54 breakout tail for +113 six times, then gave it all back chasing the 16:27, 16:51 and 17:26 pushes for stop after stop. The chandelier family actually finished negative on an up day (-497) — it either got the 15:32 tick shakeout or trailed so loosely that a leg's normal give-back at the top clipped it.

By exit reason the story is even blunter:

EXIT REASON	TRADES	NET REAL
STOP	29	-1,514
ABSORPTION_CUT	3	-207
CHANDELIER	11	+502
TARGET	11	+830

Twenty-nine stopped longs for -1,514 on a day the market rose 345 points. Those stops are almost entirely the late-entry chases: buy the 16:27 push (seven sims, -58 each), buy the 16:51 push (six sims, -51 each), buy the 17:26 push (ten sims, -44 each). Each time the trigger fired on a burst that had already spent itself, we bought the local high, and the next shallow pullback — the very thing that defines a healthy continuation — took us out. The TARGET and CHANDELIER greens came only from the two legs where we entered early enough to have a cushion.

The legs we sat out entirely

Two legs paid nothing because no sim was positioned for them. The 14:32 run (+122 points, the second-biggest of the afternoon) drew a single long — a grind_fast that arrived at 14:41 already at the 28,884 top and stopped for -78. And the 16:08 run (+86 points) had zero long entries open inside it. On a trend day the cost of standing aside is as real as a loss: 208 points of clean up-leg went by with nobody aboard. The reason we missed them is the same reason we chased the others — our triggers wait for the confirming aggressor burst, and on the quieter continuation legs (16:08 rose on 200–500-contract minutes) that burst is small and late, so we either miss it or arrive at the top.

That thinning tape matters for Monday. Volume fell steadily through the afternoon — 1,000–2,000 contracts a minute at 14:xx, down to 150–350 a minute by 17:xx. The later, thinner legs (16:59, 17:21) are exactly where our chase entries got chopped the worst, because a burst-based trigger in thin tape is buying noise.

The cleaner trigger, and how it grabs more of the trend

The fix is not a new indicator; it is moving the entry from the top of the push to the base of it. A continuation trigger should arm on the pullback, not the breakout: once we have a confirmed higher-low structure (a rising 5s low, a shallow 30–60-point dip that holds above the prior swing low), buy the reclaim

of that dip with a stop under the pullback low — the grind_fast posture — and stop trying to buy the aggressor burst after the prior high is already broken. Every green trade this afternoon fits that posture; every red cluster violates it.

Concretely, on this afternoon that one change does three things. It puts us in the 14:32 and 16:08 legs we sat out, because a pullback trigger does not need the big burst to fire. It saves the 15:32, 16:27, 16:51 and 17:26 clusters — roughly -1,150 of stops — because we would be buying the dip that shook them out rather than the high that trapped them. And it keeps the winners we already have (grind_fast's +242 and +151 came from precisely this entry). A rough honest tally: the four late-chase clusters flip from about -1,150 to break-even-or-better, and the two sat-out legs add positive expectancy where we had zero — turning a -390 long book into a clearly green one, on a day the market handed us 345 points and asked only that we buy the dips.

Monday move: gate continuation longs on a held higher-low pullback, not on the breakout aggressor burst. Arm the entry when a shallow dip (roughly 0.3–0.6 ATR, here 30–60 points) holds above the prior swing low and price reclaims it; place the stop under the pullback low; retire the "buy the push through the prior high" trigger for trend continuation entirely. Watch grind_fast as the live template — it already trades this way and was the only family green on the afternoon.

One honesty note to carry forward: score every one of these on the tick reprice, never the bar-marks. The 15:32 cluster proves the point — six "targets and chandeliers" worth up to +234 on the minute chart were all -68 on the tape. If we tune the continuation trigger against bar-marks we will convince ourselves the chase works. It does not.

KILL THE BREAKOUT-BURST CONTINUATION ENTRY

KEEP GRIND_FAST'S PULLBACK-RIDER POSTURE AS THE TEMPLATE

WATCH THE THIN LATE-AFTERNOON LEGS — PULLBACK TRIGGERS ONLY, NO BURST-CHASING IN 200-CONTRACT TAPE

§5 Most of the runs we "sat out" were a dead camera, not a bad decision — and the real money left behind is on the short side in broad daylight

Of the 207 fast MNQ runs (16 minutes or under) in this window, 114 opened with no shadow taking the trade in their direction. That looks damning until you split it

by cause. 87 of those 114 happened before the shadow board was even recording — the camera was off, so of course nothing fired. Only 27 were genuine passes where the board was live and still declined. When you pull those 27 apart, 11 are thin overnight vacuum drifts you are right to skip, and 16 are liquid daylight moves — mostly clean sell-offs — that we should have been in. That last bucket is the honest sat-out money, and it is smaller than it first appears but pointed squarely at the short side.

The single biggest number on the board is a recording gap, not lost edge

The shadow board's first trade in this data lands at 07-16 16:49 UTC and its last at 07-17 20:57 UTC. Everything before 16:49 on the 16th — the entire back half of the 15th and most of the 16th's day session — happened with the board dark. That is where "most of 07-16" being sat out comes from: it was not a gate turning trades away, it was a blackout. 87 fast runs totalling 4,671 marked points (\$9,342 at MNQ's \$2/point, as a ceiling) fall in that blackout, and they include the three largest moves in the whole window: the 07-16 13:56 UP of 157 points on 122,329 contracts, the 07-15 15:54 DOWN of 104 points, and the 07-16 15:50 DOWN of 102 points. Those are not thin-tape flukes — they are fat, liquid, tight-spread day-session trends. We simply had no camera pointed at them.

Cut that blackout bucket down to the ones that were obviously grabbable — day session, better than 1,500 contracts a minute — and it is 23 runs worth 1,817 points, about \$3,634 of ceiling. That is the biggest lever in this entire section, and it costs nothing to fix: keep the board recording. There is no clever gate to design here. The tape was there, the moves were real, the board was asleep.

Query behind it: shadow board active window from `min/max(entry_ts)` in `shadow.db` = 07-16 16:49 → 07-17 20:57 UTC. 87 of 207 fast runs start before 16:49. Marked points summed from the run list's `size` field; \$ at \$2/point MNQ, quoted as a ceiling because you never catch a run from its exact bottom to its exact top.

Overnight tape is genuinely a different animal — a quarter of the liquidity and a spread half again as wide

Before deciding whether the passes are edge or tax, look at what the off-hours tape actually is. Split every fast run by session and the footprint is night and day. In the US regular window (13:30–20:00 UTC) MNQ traded roughly

3,000 to 8,400 contracts a minute with the top-of-book spread sitting around 0.45 to 0.56 points. Outside that window — the 22:00 through 12:00 UTC stretch — volume collapses to 640 to 1,450 contracts a minute and the spread widens to 0.6 to 0.75 points. That wide spread matters in dollars: 0.7 points on MNQ is about \$1.40 to cross each way, roughly \$2.80 round-trip per contract before you have made a cent, versus about \$1.80 round-trip in the day session.

RTH volume / min~5,400

Overnight volume / min~1,050

RTH spread (pts)0.47

Overnight spread (pts)0.68

There is a second tell that says these overnight moves are hollow. I measured how far price travels per 1,000 contracts of volume — a rough "vacuum score." Overnight runs move 4.48 points per 1,000 contracts; day-session runs need 2.5 times the fuel, moving only 1.82 points per 1,000. In plain words: at night the book is so thin that a modest order walks price a long way, so a 50-point overnight "run" is often price falling through an empty ladder rather than real demand pushing it. That is the classic vacuum drift signature, and it is exactly why the marked points on those runs overstate what you could actually pocket — the print moved, but there was nothing there to fill you at those prints.

The runs are not driven by aggressive one-sided flow — the signal lives in the book, not the tape

Here is the finding that reshapes how we'd ever chase these. I computed, for every run, the aggressor imbalance aligned to the run's direction — buy-initiated minus sell-initiated volume over total, signed so positive means the aggressors were pushing the way price went. Across the board it is essentially zero: +0.019 overnight and +0.029 in the day session. Even on the biggest sat-out moves the number hovers between -0.06 and +0.08. Runs on MNQ, up or down, day or night, are simply not events where market orders pile onto one side. They are the book pulling — resting liquidity thinning or lifting ahead of price — with roughly balanced aggression underneath.

The blunt consequence: any entry trigger we might build to catch these off the aggressor tape is dead on arrival, because the tape is balanced right

through the move. This is the same lesson as the L2-book work — the forward-looking thing is resting liquidity, not who hit whom. If we want to catch runs we sit out, the trigger has to read the book, not the aggressor imbalance.

The 27 genuine passes, sorted into edge and tax

Now the honest part — the runs where the board was live and still took nothing in the run's direction. There are 27 of them. Split by whether the tape could actually support a trade (using 1,500 contracts a minute as the liquidity line), 16 are liquid and 11 are thin.

BUCKET	RUNS	MARKED PTS	\$ CEILING	
Pre-board blackout (camera off)	87	4,671	\$9,342	Coverage gap — turn the came
During-board, liquid, declined flow	16	862	\$1,724	Real edge left on the
During-board, thin vacuum drift	11	484	\$969	A tax — correctly sk

The thin 11 are the overnight drifts you can see in the tape: 07-16 22:39 DOWN 31 points on 597 contracts a minute, 07-17 05:17 DOWN 35 points on 894 a minute, 07-17 08:59 UP 48 points on 910 a minute, 07-17 10:39 DOWN 31 points on 989 a minute — all with spreads of 0.61 to 0.73 points. On tape that thin you cannot get size on without moving the price against yourself, your stop has to be wide because a single order swings the ladder, and the 0.7-point spread eats a chunk of a 31-point move before you start. The marked \$969 on these is a mirage; the fillable number is a fraction of it, and the drawdown risk of standing in a 3am vacuum is real. Skipping these is not a miss, it is discipline. Leave them alone.

The liquid 16 are the miss that matters — and they are overwhelmingly short

The 16 liquid declined-flow runs are a different story, and they cluster with a clear signature. Eleven of the 16 are DOWN moves (614 of the 862 points), only five are UP (247 points). Widen to all 27 during-board passes and the skew is starker still: 20 DOWN versus 7 UP. We sat out shorts at roughly three times the rate we sat out longs, and the biggest ones are a single sustained 07-17 afternoon sell-off in the heart of the day session: 14:11 DOWN 66 points on 11,135 contracts a minute, 14:16 DOWN 52 points on 140,410 contracts, 15:57 DOWN 68 points on 7,144 a minute, 16:18 DOWN 78 points, 19:31 DOWN 67 points — every one of them tight-spread (0.44 to 0.52 points),

deep, liquid, unmistakable trend. This is not thin-tape ambiguity. This is a clean, well-lit downtrend that the board watched go by.

That skew points straight at the short-side entry plumbing. On this desk the reversion shorts are held back by a bear-tape gate and momentum shorts are meant to be the unconstrained two-sided arm — but a live board that fires on longs and passes on clean liquid down-runs is behaving like something is still throttling the short entries. The money here (\$1,724 ceiling, most of it short) is small next to the blackout, but unlike the blackout it is a design signal, not a plumbing outage: the tape gave us grabbable sell-offs and the entries did not engage.

And when we DID catch off-hours runs, they paid — thinly, but positive

One honest check before calling the passes "lost edge": did the runs we caught overnight actually make money, on real fills? Yes. Across the active board window, the 56 overnight runs where an aligned shadow fired netted + \$4,018 of real P&L, +\$71.8 a run, 30 green to 26 red — actually a touch better per run than the 39 caught day-session runs at +\$55.7. At the individual-trade level all 244 filled overnight trades netted +\$839 (about +\$3.4 each) against 170 day trades at +\$990 (+\$5.8 each). So overnight edge is real but thin per trade — the win comes from catching the liquid ones and the losses come from getting chopped in the vacuum ones. That reinforces the split above: the edge is in the liquid off-hours moves, the tax is in the thin ones, and a liquidity floor is what separates them.

The honest sat-out money, in one breath

\$9,342 of ceiling sat out to a dark camera (biggest number, zero-cost fix). \$1,724 sat out to a live board declining real, liquid, mostly-short daylight flow (the actual design miss). \$969 sat out to thin overnight vacuum drifts that we should keep skipping. The blackout dwarfs everything, and it is the one that requires no cleverness at all.

Monday moves

1 — Turn the camera on and leave it on. The shadow board must record 24/5 with no startup blackout. This window lost visibility on 87 runs and \$9,342 of ceiling, including the three biggest moves of the week, purely because the board wasn't up. This is the highest-value, lowest-effort change in the section. [FIX COVERAGE](#)

2 — Chase the short-side daylight passes, not the overnight drifts. The 16 liquid declined-flow runs are 3:1 short and centre on the 07-17 afternoon sell-off. Audit why the live board fires longs and passes clean, liquid, tight-spread down-trends — this smells like the bear-tape gate or a short-entry throttle still choking the momentum-short arm. That is where the recoverable edge is. [INVESTIGATE SHORT ENTRIES](#)

3 — Formalise the tax with a liquidity floor, and read the book not the tape. Make skipping the thin ones a rule, not an accident: no chasing a run under ~1,500 contracts a minute with a spread over ~0.55 points. And since aggressor imbalance is flat (+0.02) right through every run, do not build a tape-aggressor trigger to catch what we miss — the run signal is in resting liquidity, so any Monday entry work belongs on the L2 book. [SKIP THIN, IGNORE AGGRESSOR](#)

Verdict: The scary "we sat out 114 runs" number is mostly a recording blackout — fix the camera and it evaporates. The genuine passes are small (\$1,724 of grabbable edge, \$969 of correctly-skipped tax), but they are informative: the edge we actually left on the table is liquid, in daylight, and short.

THE TAPE, DAY BY DAY AND REGIME BY REGIME

§6 Both days the shadows traded were down days — the whole game was to know that and lean short, because fading the trend is the only thing that bled every session

Three MNQ sessions sit in the tape: a quiet half-day on the 15th, a hard sell-off on the 16th, and a two-sided grind that closed lower on the 17th. The shadow board only came alive at 16:49 UTC on the 16th, so it has an opinion on two of them. Over those two sessions the shorts made **+2,884** and the longs lost **-1,054**, and that single split — with the day versus against it — explains almost everything that paid and everything that bled. The desk did not need a cleverer gate. It needed to read, by lunchtime, which kind of day it was standing in.

The one number that runs the whole section. On both days the shadows traded, price closed lower (07-16 down 594 points, 07-17 down 327 points). Sort every filled trade by direction and the story is brutal and simple: **SHORT — 155 trades, +2,884, 54% win. LONG — 259 trades, -1,054, 30% win.** The shorts were trading with the day's slope. The longs were fading it. Everything below is texture on that one line.

07-15, the quiet half-day: a coin-flip tape the board never had to touch

We only caught the afternoon of the 15th — the tape opens at 15:35 UTC. From there it went almost nowhere with a lot of motion: open 29,622.75, close 29,690.5, a net of just +67.8 points across a 347-point range. Break it into hours and there is no clean trend anywhere. The 16:00 hour dumped 105 points, the 19:00 hour rallied 96 points back, and the evening (20:00–23:00 UTC) did essentially nothing — the 20:00 and 22:00 hours netted +0.5 and +0.2 points on 200-plus points of one-minute path each. That is the textbook signature of chop: range without progress, motion that cancels itself.

The tape confirms it from the aggressor side. Across the whole session buyers lifted 253,691 contracts and sellers hit 253,127 — a buy/sell ratio of 1.002, dead balanced, net aggression of +564 on half a million contracts. Nobody was in control. There is no footprint to read here because there is no story: liquidity was two-sided and neither side could hold ground. The shadow board had not switched on yet, and that is the luckiest thing that happened all week, because a momentum book turned loose on a range-bound afternoon is a stop-machine. The lesson the 15th teaches is a negative one — this is exactly the day the desk must sit on its hands, and the way you'd know it in real time is the same way we know it now: hourly net moves that keep flipping sign and a buy/sell ratio pinned at 1.00.

07-16, the sell-off: sellers owned the tape all day, but the board arrived late and into the messiest part of it

The 16th was the cleanest directional day of the three and the shadows almost entirely missed it. Price opened 29,690.75 and closed 29,096.5 — down 594 points on a 736-point range, the widest of the week. The tape leaned sell the entire session: 1,190,571 buy-aggressor contracts against 1,213,621 sell, a ratio of 0.981 and net aggression of −23,050 contracts. That is the footprint of a real distribution day — offers getting hit, sellers as the aggressor, price making lower lows through the US morning (the 07:00 hour alone dropped 121 points, the 11:00 hour another 83). The run list agrees: 58 genuine down-runs against 41 up-runs on the day, the most sell-side runs of any session.

But the board did not switch on until 16:49 UTC — past the volume peak (the 13:00–14:00 US-open hours cleared 405k and 435k contracts; by 17:00 it had thinned to 176k). So every shadow trade on the 16th was fired into the late-afternoon and overnight grind — a down-trend, yes, but a slow, chopping, low-volume version of it. And in that texture the board bled from both pockets: **77 filled trades, −736.5 total, only 29% winners.**

07-16 (FROM 16:49 UTC)	TRADES	REAL P&L	WIN %
LONG (fading the down-day)	49	−796.5	22%
SHORT (with the down-day)	28	+60.0	39%

Read that carefully, because it is more honest than the tidy "shorts win on down days" headline. The shorts did not really win on the 16th — they

scraped +60 at a 39% hit rate, essentially flat. The longs were the disaster, down 796. What actually happened is that the late-afternoon chop was slow enough that even a with-trend short couldn't reliably reach target before the counter-bounce, while the fade-longs walked straight into every leg down. The worst names were the reversion longs: **rg_long_fast -201** across 15 trades, **thrust_loose -124**, **rg_long_fast_v -88**, **capit_loose -84**. The momentum flagship, **grind_fast**, ran 24 trades and netted +7.5 — a full afternoon of activity for nothing. The 16th's real lesson: a trend day you join in its tired, thin final third is not a trend day you can trade like one. The edge lives in the loud hours, not the quiet tail.

07-17, the day the board came good: a bullish-looking tape that kept getting sold, and the shorts + momentum cleaned up

The 17th is where the shadow board earned its keep: **337 filled trades, +2,566, 42% winners**. And it is the most interesting tape of the three because the surface reading and the outcome point in opposite directions. Price closed lower — open 29,096.25, close 28,768.75, net -327 points on a 716-point range — yet the aggressor tape was net buying all day: 1,422,688 buy contracts against 1,403,355 sell, ratio 1.014, net aggression of +19,333. Buyers were the ones lifting offers, and price fell anyway.

That gap — net-buy aggression, net-down price — is the footprint of the day: **absorption**. Buyers kept hitting the ask, a resting seller kept soaking it up, the buyers got exhausted, and each attempt resolved lower. It is exactly the tape that punishes a naive long (you're buying with the crowd into a wall) and pays a short that waits for the buy-thrust to fail. It is also why the absorption-veto strategies had their best day of the week here — **abs_veto_55s +580**, **abs_veto_50s +443** — they were built to stand aside from, or lean against, precisely this buy-into-supply pattern.

07-17	TRADES	REAL P&L	WIN %
SHORT (with the sell-absorption)	127	+2,824.0	57%
LONG (buying into the wall)	210	-258.0	32%

The winners were the momentum/short complex, and they won big: **grind_fast +1,292** over 53 trades, **thrust_cont +630**, **thrust_loose +586**, **abs_veto_55s +580**. The losers were the same fade-longs that bled on the 16th, just louder: **capit_loose -538 at a 9%-ish hit rate**, **rg_long_fast -292**, **rg_long_20_t20 -221**, **rg_long_fast_v -155**. The difference between

the 16th and the 17th was not the strategies and not the direction bias — it was that the 17th had loud, tradeable legs. Its US session threw genuine runs: a 214-point up-burst at 13:57, a 163-point up-run at 13:41, a 142-point down-run at 08:06, a 128-point up-run at 15:41. When the board participated in those, it collected: the 13:41 up-run drew 5 same-direction shadows for +1,327, the 08:06 down-run drew 14 for +1,064, the 15:41 up-run drew 10 for +1,037, a 19:26 up-burst drew 10 for +726. Momentum that catches a real run gets paid; that is the whole thesis and the 17th is the clean demonstration of it.

The cross-day law: momentum is paid by big runs in the loud hours, and fading the day's slope bleeds in every regime

Stack the three sessions and two laws fall out, one about direction and one about time-of-day.

Law one — trade the day's slope, never fade it. Both traded days closed down; on both, shorts beat longs, and the reversion-long book (capit_loose, rg_long_fast, rg_long_fast_v) lost money on both days without exception. Across everything: shorts +2,884 at 54%, longs −1,054 at 30%. And you can see the mechanism in the exit reasons — 239 trades exited on STOP for −7,840, while the 97 TARGET hits made +5,399 and the 75 CHANDELIER trails made +4,477. The book's edge is entirely in the runners that reach target or trail out with the trend; the stops that fund them are overwhelmingly the fades getting run over. Every point the reversion longs "captured" on a bounce was handed back and then some on the next leg down.

Direction versus the day's slope, both traded sessions combined (real P&L, filled):

SHORT — with the down-trend		+2,884
LONG — fading the down-trend		−1,054
07-16 longs (thin down-chop)		−796
07-17 shorts (loud sell-legs)		+2,824

Law two — momentum needs the loud hours; the quiet tail and the overnight are where it stalls.

The volume tells you when to be awake. On both real days the US cash open (13:00–16:00 UTC) cleared three-to-eight times the contracts of the overnight hours, and that is where the tradeable runs lived — the 214-, 183-, 163-, 142- and 128-point runs on the 17th all printed between 13:00 and 16:00 UTC. The 13:00 hour alone made the board +1,029; the 14:00 hour +415. The overnight and the tired late-afternoon are the opposite: the 16:00–18:00 UTC block on the 17th bled (–457, –333, –594 by hour) even on a winning day, and the entire 07-16 shadow record is one long demonstration of momentum arriving after the volume left. Same strategies, same short bias — the only variable that flipped the result was whether there was a real run to catch. When a top run drew a crowd of participants it paid four-figures; when a run drew one or two lonely shadows it was a wash (the 14:55 down-run: 1 participant, +7; the 14:32 up-run: 1 participant, –78).

So the desk's first question each session is not "which gate" but "which day am I in," and it is answerable from the tape by lunchtime with three cheap reads: the sign and persistence of the hourly net move (does the slope hold or flip?), the aggressor buy/sell ratio (0.98 = sellers own it, lean short; 1.00 = chop, stand down), and the US-open volume (is there enough fuel for a run?). Get those three right and you are short on the 16th, short-and-patient into the absorption on the 17th, and flat on the 15th — which is exactly the P&L-maximising posture the board stumbled into only on the one day it happened to be awake for the runs.

Monday move. Wire a lightweight "day-type" read that runs by 14:00 UTC and gates the book's direction bias: compute the trailing hourly net-move persistence (are the last 3 hourly closes stepping the same way?) plus the session aggressor buy/sell ratio. If the ratio is under 0.99 and the slope is stepping down, bias short and let momentum size up; if it is over 1.01 and stepping up, bias long; if it sits 0.99–1.01 with a flipping slope, cut the reversion-long book to zero and stand the momentum book down — that is a 07-15 chop day and it only donates stops. Concretely for Monday: keep grind_fast / thrust_cont / abs_veto armed only inside 13:00–20:00 UTC when the day-type read agrees with their side, and hard-mute capit_loose and rg_long_fast whenever the day-slope is against them, because those two fade-longs lost on every single session this week.

KEEP the momentum/short complex — it is the entire +2,566 on the 17th and it is paid by real runs in the loud hours. **KILL** the reversion-long book on any day whose slope and tape point down; capit_loose and rg_long_fast bled on both sessions and never earned their stops back. **WATCH** the quiet/overnight window — momentum there is a coin flip that the exit rules cannot rescue, so it should not be firing until the day-type read says a run is live.

THE THROUGH-LINES

§7 Momentum earns every dollar the desk has; the fade book hands it straight back

Step back from the individual sims and one picture is left standing. Over the last two weeks the momentum book made money — real, tick-filled money — and the fade-and-reversion book lost almost exactly as much. The desk is net-green for one reason only: momentum is big enough to carry the fade book on its back. The Monday move writes itself. Run the thing that pays. Bench the thing that bleeds. That is the whole strategy.

Here is the split, scored on honest fills over the two weeks, and it is not close. The momentum book — every thrust and continuation sim — made **+\$21,792**. The fade book — reversion, passive-chop, the "sell the rip" family — lost **-\$15,798**. Inside those two piles the story sharpens further: the MNQ thrust family alone made **+\$22,453**, which is more than the entire momentum book, meaning the odds-and-ends momentum sims roughly cancel out and the clean thrust setups do all the lifting. On the other side, passive_chop by itself lost **-\$12,528** — four-fifths of the entire fade loss came from one idea.

Momentum book	+\$21,792
↳ MNQ thrust family	+\$22,453
Fade / reversion book	-\$15,798
↳ passive_chop alone	-\$12,528

Two-week V5 board, filled trades, real tick-repriced pnl. Every momentum bar pushes right; every fade bar pulls left. The desk is drawn along one axis — whether the trade pushes a move or leans against it.

Re-deriving it straight off the two-week board this morning lands in the same place — the momentum sims net about **+\$23,600**, the fade sims about **-\$14,800**, the MNQ thrust family **+\$22,337**, and every passive_chop variant stacked together **-\$14,286**. The exact dollars wiggle a little depending on how

you bucket the borderline sims and whether you count the tiniest ones, and honest tick fills always run a touch different from the clean bar marks — but the shape does not move. One book is a payer. The other is a payer in reverse.

Why this keeps happening — it is the same trade run backwards

These two books are not two strategies. They are one instinct pointed in opposite directions. Momentum buys the rip and rides it; the fade book sells the same rip and stands in front of it. When MNQ actually goes somewhere — and over these two weeks it went somewhere often enough — the first book collects and the second gets run over. `passive_chop` is the purest example of the losing side: it assumes the tape will keep chopping back to the middle, so it fades every push, and on any day with a real trend it hands back a little on every single entry until the total reads **-\$12,528**. That is not a tuning problem you fix with a tighter stop. It is the setup being wrong about which way the tape wants to go.

The desk has been net-green for a fortnight. That is true, and it is also the trap. Green is not coming from a balanced pair of edges — it is one strong edge dragging a dead weight. Cut the dead weight and you do not just remove a loss, you keep the whole **+\$21,792** instead of netting it down to a few thousand. The fade book is not a hedge. It is a tax on the part that works.

The one thing not to hear

Killing the fade book is not the same as going long-only, and that distinction is the difference between a good Monday and a blown one. Momentum is two-sided — it takes the short whenever the thrust points down, and on the fresh V7 board the short side is where a lot of the money lives: `thrust_loose` made **+\$500** short against just **-\$38** long, and `abs_veto_55s` was green on both sides. Retiring reversion means "stop selling rips you are betting will reverse." It does not mean "stop selling." When a genuine downward thrust fires, momentum shorts it and that is exactly right. Freeze the shorts and you throw away half of the only book that pays. The rule is about the setup, not the direction.

BOOK	TWO-WEEK P&L	VERDICT
Momentum — confirmed thrust, both sides	+\$21,792	Run it live
Fade / reversion / passive-chop	-\$15,798	Do not run live
Net desk (momentum carrying fade)	+\$5,994	The carry is the risk

The Monday move

Run the confirmed thrust book live, two-sided, and do not run reversion, fade, or passive-chop live in any form. The net-green you have been reading is momentum dragging a corpse — unhook the corpse and the same market gives you the full momentum number instead of the diluted one. Keep the shorts armed; the money is on both sides of momentum, and long-only would surrender half the edge for nothing. This is the through-line under every other section: the desk does not need a new idea, it needs to stop paying for the one that loses.

KEEP CONFIRMED TWO-SIDED MOMENTUM — IT IS THE WHOLE EDGE

RETIRE REVERSION / FADE / PASSIVE-CHOP FROM ANY LIVE PATH

§8 The chandelier is the whole edge — and it hands back a third of every winner

If you want to know where the desk actually makes its money, look at how trades end, not how they start. On the live V7 desk the trailing chandelier is the only exit that has ever printed green — **+\$567** over 13 trades at an 85% win rate — while every stop, adverse cut and absorption cut has bled. The shadow board tells the same story with more trades behind it. So this week's through-line is simple: the exit architecture is carrying the desk, the chandelier is the heart of it, and it has one fixable leak — it gives back about a third of every winner's peak, and the dollar size of that give-back scales straight with volatility.

Where the money is actually booked

Across the 28-sim V7 shadow board, sorting every filled trade by how it exited splits the desk cleanly into a profit half and a loss half. The two trailing/target exits book real money; the stop is where it dies.

TARGET (97t, 91%)		+\$5,399
CHANDELIER (75t, 79%)		+\$4,477
ABSORPTION (3t)		\$20.7
STOP (239t, 6%)		-\$7,840

Two honest caveats before we lean on this. First, on the shadow board the fixed profit **TARGET** actually books slightly more than the chandelier (**+\$5,399** vs **+\$4,477**) — so I won't overclaim that the chandelier is the only profit centre; it's the profit centre on the live desk, where the sims run thrust-only and never hit a fixed target. Second, the **-\$7,840** stop pile is not a give-back problem at all. I checked: the 239 stopped trades had an average peak of **-\$50.7** and not one of them ever showed even \$20 of green before it died. Those are bad entries, full stop. The stop is doing its job. This section is about the winners.

The leak: a third of the peak, every time

Take every filled chandelier winner and compare the best unrealised profit it ever showed (the ceiling) against what actually landed in the book. The gap is the give-back — the money the trail let slip on the way out.

CHANDELIER WINNERS (V7 SHADOW)	N	AVG PEAK	AVG BOOKED	GIVEN BACK
All chandelier wins	74	\$91.0	\$60.4	\$30.6

Ninety-one dollars of peak, sixty booked. That's a **34% give-back** — and it isn't a one-trade fluke. Summed across all chandelier winners the desk showed **\$6,732** of peak profit and kept **\$4,470**, handing back **\$2,262**, which is 0.34 of the total. A third of the peak, in aggregate. When the operator watches a live winner tick to +\$224 and settle at +\$88, that feels like a bug. It isn't — it's the norm the trail was built to produce. The only question is whether a third is the right third to give back, and the answer depends entirely on how fast the instrument is moving.

The give-back scales with volatility — that's the whole prescription

Here's the finding that matters. Bucket every filled winner by its entry ATR-as-percent-of-price and the give-back fraction barely moves — it sits near a third whether the tape is dead quiet or ripping. But the give-back in **dollars**

roughly doubles from the quiet bucket to the fast one, because a third of a big peak is a big number.

ATR BUCKET	N	ATR%	AVG PEAK	AVG BOOKED	GIVEN BACK	GIVE-BACK %
Low vol	13	0.028	\$49.8	\$34.6	\$15.2	31%
Mid vol	104	0.053	\$68.6	\$45.4	\$23.2	34%
High vol	54	0.097	\$124.0	\$86.9	\$37.1	30%

Read the last two columns together. The percentage is flat — 31%, 34%, 30% — so a single flat trail is proportionally fair everywhere. But the dollars run \$15 → \$23 → \$37. On the fast tape a high-ATR winner peaks at **\$124** and the trail lets \$37 of it walk out the door. That \$95-to-\$124 peak giving back into the \$85-\$90 range is exactly the operator's +\$224→+\$88 experience, just averaged. A flat-width trail treats a 3-point wiggle in a quiet MNQ the same as a 12-point wiggle in a fast one: in the quiet tape it clips runners early for no reason, and in the fast tape it's still parked a fixed distance back while the price swings a fortune through it.

What the A/B actually says about trail width

The board ran three flat chandelier widths this week, and the raw scores lean tighter-is-better: **k20** at **+\$18.5**, **k25** at **-\$36.5**, **k35** at **-\$101**. But I'm not going to hang a decision on that — it's 9 to 11 trades per arm in a single chop week, and all three are inside coin-flip range on win rate (33-45%). What it does confirm is the mechanism: the widest arm (k35) is the one that gave back positive dollars on average and scored worst, i.e. a wide flat trail sits too far back and bleeds. That's consistent with the 25-day archive replay, where a fixed 3.5-ATR trail was beaten by an ATR-conditional graded trail by **+\$245** on the thrust family and **+\$1,653** on grind. Same lesson from two directions: the answer is not "pick a tighter number," it's "stop picking one number at all."

The prescription – a vol-adaptive chandelier. Set the trail distance as a function of live ATR, not a fixed point/percent width. Tighten hard in the low-vol bucket (a quiet winner that peaks at \$50 shouldn't need \$15 of room – that's clipping runners for nothing) and give a fast-tape winner proportionally more slack so a normal high-ATR swing doesn't knock it out at the first wobble. The give-back fraction is already a stable third; the goal is to make that third the *deliberate* third rather than an accident of whichever flat width was hard-coded. The runs are there to ride – MNQ made a thousand big moves this month, so the runners are never in short supply; the leak is purely in how we ride the ones we're in out to the end.

The verdict

The exit spine is the desk's real product this week, and it's working: the trailing chandelier books green on the live desk when nothing else does, and on the shadow board it and the fixed target between them carry roughly \$10k of profit against the stop pile. The one honest leak is the give-back, and it is not a mystery — it's a flat third of peak that turns into real dollars precisely when the tape is fast. A flat trail can't win that fight in both regimes at once.

KEEP THE CHANDELIER — it is the profit centre, live and shadow. **WATCH THE FLAT-WIDTH A/B**

— tighter won this week but on 9-11 trades in one chop regime; don't promote a single number off it. **BUILD THE VOL-ADAPTIVE TRAIL** — the give-back scaling with ATR (\$15→\$23→\$37 across the buckets, same 30-34% every time) plus the archive replay's +\$245/+\$1,653 edge for a graded trail is the strongest, most repeatable exit finding on the board. Make the trail a function of ATR, not a constant.

§9 The whole week says one thing: the edge isn't more gates, it's an OFF switch that reads the room before it fires

Line every shadow up next to each other and a single rule sorts the winners from the bleeders. It is not which gate is cleverest. It is whether the gate looks at the room before it acts. The two families that check the condition first — confirmed thrust and absorption-veto — made money on both sides. The two that fire on a fixed trigger no matter what the tape is doing — capitulation and the long-only reversal-grabs — bled almost every dollar the board lost. That is the desk's real game, and it has nothing to do with stacking on more gates.

Four families, same week, same tape. I grouped every filled shadow into what it actually is and scored each on honest tick-repriced fills:

FAMILY	WHAT IT DOES BEFORE IT FIRES	TRADES	REAL P&L
Confirmed thrust	Waits for the move to prove itself	112	+\$1,345
Absorption-veto	Stands down when the book is eating the push	49	+\$1,069
Capitulation	Fires on a fixed flush trigger, condition-blind	27	-\$722
Reversal-grab (long-only)	Grabs the falling knife long, always long	117	-\$1,092



The winners aren't better triggers — they're gates that can say "not now"

Look at what the two green families have in common. Absorption-veto's entire job is to refuse: when the resting book is absorbing the push, it stands aside. It went +\$307 long and +\$762 short, 50% and 64% — green in both directions because it isn't picking a side, it's picking a moment. Confirmed thrust is the same idea from the other end: it waits until the move has already committed before it climbs aboard, +\$433 on the short side alone at 63%. Neither of these is a smarter fade or a tighter stop. They are gates with an off switch, and the off switch is what makes them money.

The bleeders all share one flaw: they fire into the wrong room

Now the red. Capitulation lost \$722 over 27 trades and won 11% of them — and when you pull its exits apart the story is brutal: 22 of its trades died on the STOP, -\$460 long and -\$268 short, against a grand total of \$5.50 of TARGET money. It is a fixed flush trigger that keeps firing into continuations that were never going to reverse. The long-only reversal-grabs are the same disease in a different coat: 117 trades, every one of them long, -\$1,092 at 26%. It grabs each dip as if the tape owes it a bounce, and on a week that

trended it just caught knife after knife. Neither gate ever asks whether the condition supports the trade. They just fire.

The one-line honesty note: these dollars are tick-repriced fills, so they run a little tighter than the clean bar-mark points you'd read off a chart. The *ranking* is what matters and it is rock-solid — the condition-reading families are green, the condition-blind ones are red, and nothing about the fill model flips that.

So the obvious move is a day-type router — and it's a trap

If reading the room is the edge, the tempting next step is to read it for the whole day: detect "trend day" versus "chop day" in the morning and route to thrust or reversion accordingly. I built exactly that and the skeptic took it apart. Out of sample it collapsed — the same router that showed +\$1,974 on one split fell to +\$199 when the days were cut a different way, which is the fingerprint of a curve fit, not an edge. The reversion leg it was supposed to feed is +\$723 over 25 days and that whole number is carried by a single outlier session; strip the best day and it's -\$209. The router looks like intelligence and behaves like luck.

The distinction that actually holds. Predicting the day is dead — we proved regime prediction impossible at minute scale months ago, and this router is just that ghost in a new suit. What works is the opposite: reactive, per-condition gating. Don't forecast the room, read it right now and let the gate decline. Absorption-veto reads the book this second and stands down. Confirmed thrust reads commitment this second and waits. That is not prediction — it's an off switch wired to what the tape is doing in the moment, and it is the only thing on the board that made money on both sides.

The Monday move

Stop treating the board as a pile of gates to add to, and start treating it as one question: does each gate have an off switch that reads the condition, or does it fire regardless? Kill the condition-blind bleeders — capitulation and the long-only reversal-grabs, together -\$1,814 this week at a combined 20% win rate — because no amount of tuning fixes a trigger that never looks at the room. Keep leaning on the two families that already do the right thing. And shelve

the day-type router as research only: the honest read is that it did not survive out of sample, and shipping it would be dressing luck up as a signal.

RETIRE THE CONDITION-BLIND GATES (CAPITULATION, LONG-ONLY REVERSAL-GRABS) – THEY FIRE INTO THE WRONG ROOM

KEEP THE OFF-SWITCH GATES (ABS_VETO, CONFIRMED THRUST) – GREEN BOTH SIDES

WATCH-ONLY THE DAY-TYPE ROUTER – REFUTED OUT OF SAMPLE, RESEARCH BENCH

§10 Nothing has earned a bigger size — the board is two days old, and that's the whole answer

This is the part of the report where we ask the honest question: after a week of watching, has any V7 sim proven it deserves real money, or more of it? The answer is no, and it isn't close. Not because the sims look bad — a fair few are green — but because the whole board was born on Thursday. Every number below is drawn from trades between 16 July 16:49 and 17 July 20:57. Two sessions. You don't hand a strategy more contracts on two sessions, no matter how pretty the two sessions were.

The ladder, and why every sim is still on the bottom rung

The promotion ladder we agreed on is deliberately slow: a sim earns its way up (shadow at 1 lot, then live at 1, then 2, then 3) only after roughly eight to twelve weeks and a hundred to two hundred trades, and only if the edge holds across different kinds of week — grind, thrust, chop — with a shallow drawdown and clean reward-for-risk. The V7 board clears none of those bars. The busiest green sim on it, grind_fast, has 77 filled trades. The headline thrust_cont has 29. Most sims have a dozen or two. On a hundred-to-two-hundred-trade rule, the leader is a third of the way to even being eligible, and it has seen exactly one type of week.

So the report nominates nobody for a bigger size this week. That is the correct outcome, not a disappointing one — the ladder exists precisely so a good-looking Thursday doesn't cost us on a bad-looking Monday.

The best number on the board is also the best example of the trap

grind_fast finished the two days at **+\$1,299.50** over 77 trades, 47% winners, and both sides carried their weight (long +\$747.50, short +\$552). On its own

that's the strongest line in the whole shadow fleet. And it's a mirage. We already put the grind idea through an adversarial 25-day walk-forward, and grind standalone bled **-\$7,975** over that stretch; the day-type router that was supposed to rescue it got refuted out-of-sample. What we're seeing here is two days that happened to suit a grinder, not a grinder that has earned trust. If two good sessions could promote a sim, that 25-day result is exactly the bill we'd eventually pay. It stays at shadow-1, and we keep watching.

The Departure Lounge — three names that don't need eight more weeks to judge

The flip side of the ladder is the exit door. Some sims are green-but-unproven and simply need time. Others are already telling us what they are, and we can move them along now. Three groups check out of the lounge:

1. The capitulation family — fading a crash and getting run over

Every capitulation sim is red: capit_loose **-\$622.50** on a 9% win rate, capit_ride **-\$61** at zero wins, capit_mid **-\$39**. As a family that's **-\$722.50** across 27 trades. A 9% win rate isn't a tuning problem, it's the strategy describing itself: it steps in front of a move that's already going and takes the hit. This is the same lesson the reversion autopsy taught us — fading magnitude bleeds. Two days is plenty to see it.

RETIRE THE CAPITULATION FAMILY

2. The long-only rg_long_* sims — missing the side that does all the work

The rg_long_* cluster is long-only by construction, and as a group it lost **-\$1,091.50** over 117 trades. rg_long_fast alone bled **-\$493.50** at a 22% win rate. Look at why: on this same board the green thrust and absorption sims made their money on the short side — thrust_loose was long -\$38.50 but short +\$500.50, thrust_cont was long +\$161 but short +\$433, abs_veto_55s was long +\$143 but short +\$356. The tape rewarded shorts, and a long-only sim structurally can't touch that. It's not that these are badly tuned longs; it's that half the edge is a door they've bolted shut.

RETIRE RG_LONG_FAST AND RG_LONG_FAST_V


3. The duplicate — two names, one strategy

rg_long_20_t15 and rg_long_20_flow25 are not two sims. I checked the actual trade sets and they are byte-identical — same entries, same sides, same exits, both **-\$75.50** over the same 11 trades. One is a wasted board slot pretending to be an independent data point. Kill one, free the slot for something that isn't already being tested. (The cb_thrust and cb_thrust_dropped sims, meanwhile, took zero trades all week — dormant, not dead; leave them armed and see if they ever fire.)

RETIRE THE RG_LONG_20_FLOW25 DUPLICATE

The lounge, in one picture

Here's what's checking out, by what it actually lost on honest tick-repriced fills over the two days. Everything points left because everything here is red — that's the point.



The Monday move: nominate nobody for a bigger size – the board is two days old and the ladder rule is eight-to-twelve weeks plus a hundred-plus trades across regimes, so even the +\$1,299.50 grind_fast leader stays at shadow-1 (and our own 25-day walk-forward says that leader bled -\$7,975 out-of-sample, so treat the two-day number as noise, not signal). At the same time, clear the lounge: retire the capitulation family (-\$722.50, a 9% win rate is the strategy admitting it fades into a move that runs it over), retire the long-only rg_long_fast and rg_long_fast_v (they're locked out of the short side that carried every green sim this week), and delete the rg_long_20_flow25 duplicate to free a board slot. Report nominates, operator arms.

WATCH THE WHOLE GREEN BOARD – A SIZING CONVERSATION IS WEEKS AWAY, NOT SESSIONS

RETIRE THE CAPITULATION FAMILY, LONG-ONLY RG_LONG_FAST/_V, AND THE RG_LONG DUPLICATE

THE L1/L2 TAPE – SIX NEW GATES HUNTED AND BACKTESTED

§11 The tape ran 196 times this week — and it told us big MNQ moves are exhaustion reversals off a stop-run, not chases

Forget what the sims did — you watch that all day. This is the tape itself: three days of MNQ, 3.6 million aggressor ticks and 35 million order-book rows, asked one question — what does the tape do in the seconds before a big move, and can we build a gate to be there for it? The tape printed **196 moves of 30 points or more inside fifteen minutes**, and the fat ones share a shape.

The headline events of the week were not clean trends — they were **V-shaped stop-runs**. On the 17th at 13:31 price flushed 273 points down, then immediately ripped 313 points back up. Same shape on the 16th (223 down, then 235 up). Both fired off the US cash open. That is the signature, and the pre-move tape confirms the mechanism: **in the last 30 to 60 seconds before a big move, the aggressor tape is leaning the wrong way 68-71% of the time**. The crowd hits sells into the low, the sells exhaust against resting bids, and price snaps up. Big MNQ moves are not continuation of the flow you can see — they are the reversal after that flow exhausts.

The single most useful thing the tape said all week: the obvious signal is a trap. When the aggression is loudest in one direction, the next big move is usually the other way. That one fact killed three of the six gates we tried before they started.

BIGGEST MOVES THIS WEEK	DIR	SIZE	MINS
17 Jul 13:47	UP	313 pt	14
17 Jul 13:32	DOWN	273 pt	12
16 Jul 13:55	UP	236 pt	15
16 Jul 13:33	DOWN	224 pt	15

§12 Does the order book tell you which way? On MNQ, no — the resting book is worthless to backwards

The first instinct with an L2 feed is that the book shows you where price wants to go: more bids than offers means it's heading up. On MNQ this week, that is flat wrong, and it's worth being blunt about because it kills two obvious gates.

We tested the static book lean directly — the five-level imbalance as a directional trigger. It came back **breakeven at best and backwards at worst**. The reason is mechanical: a big resting bid is advertised liquidity, and on a liquid contract advertised liquidity gets **lifted**, not respected — the offers stacked above a fat bid are exactly what aggressors reach up and take. Same for the "empty book means a flush is coming" idea (book depletion): when the bid vanishes it usually just **refills and ticks the other way** — the emptying side carries zero signed edge. Both are dead.

The book's *static state* is a losing tell on MNQ. What carries information is the book's *dynamic* — aggression exhausting against a wall that won't move (a reversal), and violent sweeps that eat through levels and keep going (continuation). Everything real we found lives in the dynamic; nothing lives in the snapshot.

§13 Six gates built, six backtested — three dead on arrival, one an illusion, two with a real signal

Here is the whole hunt on one page, every gate built as a one-line armable rule and backtested on the real tape with tick-priced exits. Read the verdict column first; the honest score is two survivors, and neither is a free lunch.

GATE	WHAT IT REALLY IS	BACKTEST
Ride-the-flush	go WITH a violent sweep, not against it	+\$2,200 00S
Exhaustion-reversal	fade heavy aggression that can't move price	+\$0.70 → -\$0.30/tr
Coiled delta	"go with a delta surge while price is pinned"	did not replicate
Book depletion	"bid vanishes → short"	-\$0.50/tr
Static book imbalance	bid-heavy → long	breakeven/backwards
Vacuum-open	thin book + fast tape → momentum	-\$2,600

The three dead ones are not close calls — they lose across every parameter cell on two independent samples. The next three sections take the two survivors and the one illusion each to the end.

§14 Ride-the-flush: the one gate that pays — but it's the momentum you already trade, with a cleaner trigger

The folklore trade is to fade a stop-run: price sweeps the lows, you buy the "reclaim." We backtested that literally and it lost **-\$2,335** over 2,063 signals, in every cell of a 64-cell grid on two samples. The reclaim is a trap. But the mirror — ride the sweep, don't fade it — is the strongest, most robust signal in the whole hunt.

Going WITH a violent down-sweep won **+\$1.10 a trade** and was **positive in all 64 parameter cells across both the 3-day capture and 9 fresh archive days**. That sign-stability is the real robustness test — an overfit gate flips sign when you nudge a parameter; this one never did. **+\$2,200** out of sample. It is real.

The gate, one line. A down-sweep — price drops ≥ 5 pt in ≤ 6 s with $\geq 68\%$ of aggressor volume on the sell side, piercing the prior-20-second swing low — that fails to reclaim that level within 8 s $\rightarrow$ go **SHORT** (ride the flush; mirror for up-sweeps $\rightarrow$ long). Exit: 8-pt stop / 10-pt target / 150 s max hold, tick-priced.

BACKTEST (TICK-PRICED, \$2/PT)	SIGNALS	\$/TRADE	TOTAL	ROBUSTNESS
Fade the reclaim (the folklore trade)	2,063	-\$1.13	-\$2,335	0/64 cells green
RIDE the flush (this gate)	~2,000	+\$1.10	+\$2,200	64/64 cells · 2 samples

The honest catch is novelty, not validity. This is the same continuation edge the desk's thrust gate already harvests — a violent move keeps going for a beat. It's a better, faster, microstructure-clean trigger, not new alpha. And it clears roughly one tick of cost only thinly. So the play is not "arm a new gate."

Monday: shadow it, and A/B it against the live momentum gate — only worth wiring in if it fires cleaner or earlier than what's already there. It is a sharper entry for a trade we already make, not a money machine.

SHADOW • A/B VS THE LIVE MOMENTUM ENTRY

§15 Exhaustion-reversal: the one genuinely new footprint — real direction, but the dollars aren't there yet

This is the honest jewel of the hunt, and it's the same shape the move-foundation flagged: big MNQ moves are exhaustion reversals off a sweep. When one side dumps hard, price refuses to follow, and there's a real wall on the side being hit — the next move is the **flip**. Fading the aggression wins in every sample; following it loses about \$4 a trade more.

What makes me trust the direction: strip out the order-book-wall filter and the whole thing collapses to \$0.13-a-trade noise. The wall is the entire discriminator — which means it's **mechanistic, not fitted** (a curve-fit doesn't have a single physical knob that switches the edge on and off). The sign is confirmed three independent ways: the move-foundation (pre-move aggression is contrarian 68–71%), the absorption agent, and the skeptic's re-run.

The gate, one line. Rolling 20-second net signed tape ≥ 250 contracts one way AND price has moved ≤ 2 pt AND the near-touch resting size on the hit side $\geq$ the far side → **FADE** (enter opposite the aggression). Exit: 8-pt stop / 12-pt target / 120 s.

BACKTEST (OUT-OF-SAMPLE, 9 ARCHIVE DAYS)	\$/TRADE	READ
Fade absorption-into-wall, frictionless	+\$0.70	confirmed 3 ways
...same, after a realistic ½-pt slippage	-\$0.30	and ½pt is optimistic
Strip the book-wall filter	+\$0.13	noise — the wall IS the edge
Harden the filter (wall $\geq 2\times$ far side)	n=29, 52% hit	the one path over cost

The catch is cost. Frictionless it's +\$0.70 a trade out of sample; charge a realistic half-point of slippage — and you're entering exactly as a wall gives

way, so half a point is optimistic — and it turns **-\$0.30 a trade**. The direction is real; the money is under the cost line today. The one path to a live edge is a harder filter — requiring the wall to be $\geq 2\times$ the far side thinned it to 29 trades but lifted the hit rate to 52%.

Monday: do NOT arm. Shadow it, capture more days, and hunt a tighter wall-ratio and stall definition that clears cost. This is the direction worth building the next few weeks around — it's the only new signal in the tape with genuine information.

WATCH • SHADOW + TIGHTEN THE WALL FILTER BEFORE IT EARNS A CENT

§16 Coiled delta: the +\$1,340 “winner” that I could not reproduce — and why I'm telling you it's dead

One agent came back with the hunt's best number: an extreme 30-second delta while price is pinned, go WITH the delta — **+\$1,340** over 642 signals, "every day green," 62% hit. It's exactly the kind of number I'd have led the report with a few tries ago. I re-ran it independently on nine fresh archive days, and it does not hold.

- **The raw move is a coin flip.** At the signal, the with-delta move 30–60 seconds later averages roughly zero points at a 48–50% hit rate — no directional information at all before any exit rule touches it.
- **On the cleanest signals it inverts.** Raise the threshold to the strongest absorption events and the "follow" direction goes consistently negative (43–44% hit) — the fade wins there, not the follow.
- **Nothing in a 24-cell grid saved it.** Across every lookback, ATR window and threshold, not one cell showed the claimed positive continuation edge.

So where did the +\$1,340 come from? **An exit-rule artifact.** The 1-ATR-stop/1-ATR-target bracket quietly harvested session drift and bracket asymmetry — money the exit manufactured, not money the signal predicted. The raw forward move, which an exit can't fake, is flat. Three independent workstreams agree the true sign is the opposite of what was claimed.

Monday: do not shadow "go with the delta" — it's the same 92%-style illusion where a pretty backtest hid a dead signal. If this footprint is traded at all it belongs to the exhaustion-reversal gate above, traded FADE, under the wall filter. I'm flagging this loudly because catching my own hype before it reaches you is exactly the job.

KILL AS SPECIFIED · THE MONEY WAS THE EXIT, NOT THE SIGNAL

§17 **How we grab it from Monday — one sharper trigger to shadow, one footprint to develop, and two habits to stop**

The blunt truth of a deep L1/L2 week: the tape did not hand us a free new money gate. What it handed us is more useful than another curve-fit — a sharper trigger for the trade that already works, one genuinely new direction to develop, and two expensive habits to quit. Ranked by what it's worth.

1 · Shadow the flush-follow trigger — worth it only if it beats the live entry

The ride-the-flush rule is the only thing that pays and survives everything. But it's the momentum edge you already trade, so the question isn't "does it make money" — it's "does it enter cleaner or earlier than the live thrust gate." Wire it in shadow next to the live gate and A/B them on the same tape. If it triggers sharper, swap it in; if not, leave it.

2 · Develop the exhaustion-reversal — the only new footprint with real information

Fade heavy aggression that stalls against a wall. Real direction, sub-cost today. This is the multi-week build: capture more days, tighten the wall-ratio filter until it clears a realistic half-point of slippage, then shadow it. Do not arm it now — but this is where the genuinely new money would come from if it comes from anywhere.

3 · Stop chasing delta surges, and stop trusting the resting book

The most valuable output of the whole hunt is negative and it saves real money: two of the most natural microstructure trades — chase the aggressor flow, and lean on the fat side of the book — both lose on MNQ. The flow you

can see is the exhaustion, not the fuel; the book you can see is bait, not intent. Any future gate that leans on either starts in a hole.

The one line: no new gate to arm this week — but a cleaner trigger for the momentum you run, one new reversal footprint worth the next month of work, and hard proof to stop paying for the two obvious book/flow trades that don't work. Everything here is shadow-first; nothing arms on a week.

§18 **Every dollar we lost this week came from an engine firing into the wrong regime — a router that stands the wrong engine down turns +\$1,830 into +\$5,740**


This week on MNQ the shadow board printed a net **+\$1,829.50** across 414 honestly-repriced, filled sims. That number hides the real story. When you tag every one of those 414 trades with the regime it fired into — a fast run up, a fast run down, or chop — the book cleaves perfectly in two. Every profitable bucket is an engine trading with the prevailing move. Every losing bucket is an engine trading against it. There is no fourth quadrant. The losers are not bad strategies; they are good strategies pointed the wrong way by a book that fires all of them, all the time, regardless of what the tape is actually doing. That is exactly the hole the regime router is built to close.

What the week actually was: a down-biased, trend-heavy tape

First, the ground truth, because the router is reactive and can only ever be as good as its read of the tape. Over the shadow window (07-16 16:49 to 07-17 20:57 UTC) the run list holds 120 fast runs — the real ≤ 16 -minute directional pushes, artifacts stripped out. Of those, 74 were DOWN and 46 were UP. Measured in points travelled, DOWN runs moved 3,956 points against 2,885 for UP. This was not a balanced chop week you could fade both ways; it was a tape that spent most of its energy grinding lower in sustained 12-to-15-minute legs — the 07-17 13:25 down-run of 183 points over 15 minutes, the 08:06 down-run of 142 points over 14, the 19:52 down-run of 116 over 12. In a tape like that, momentum-short is the house edge and buying every dip is a slow bleed. The shadow board proves both halves of that sentence with real money.

The split: winners trade with the trend, losers trade against it

I tagged each filled trade by the regime active at its entry timestamp — using the fast-run windows [run-start, run-peak+5min] as trend periods and everything else as chop — then asked one question: was the trade's side aligned with the dominant run direction, or fighting it? The answer is the whole section.



Read that chart top to bottom and the router almost writes itself. Momentum-short inside a down-trend made **+\$3,170** across 143 trades. Momentum-long inside an up-trend made **+\$1,478** across 93. Reversion-long that happened to fire during an up-run got carried and booked **+\$979.50** on 18 trades. And reversion-long in genuine chop — the one place fading actually belongs — made **+\$171.50**, though on only six trades, so treat that as a hint not a proof. Then the floor drops out. The two red bars are both engines pointed the wrong way: momentum-long buying strength into a falling tape lost **-\$1,056.50** on 29 trades, and reversion-long catching knives in a down-trend lost **-\$2,653** across 109 trades. That single red bar is 68% of everything we bled. Not one losing bucket exists where the engine was aligned with the tape.

Name the bleeders — these are the sims firing into the wrong regime

The abstraction is clean but you deserve the actual strategy names and their actual repriced P&L, because these are the shadows the router would have muzzled this week. The reversion-long complex is the core of the damage: **rg_long_fast** lost **-\$493.50** over 51 longs, **capit_loose** long lost **-\$371.50** over 13, **rg_long_fast_v** lost **-\$243.50** over 35, **rg_long_20_t20** lost **-\$203.50**, and **chand_k35** long lost **-\$232** on six trades. These are dip-buyers and knife-catchers, and this week the dips kept going. On the momentum side the failure mode is the mirror image: **thrust_fast** long lost **-\$132** and

thrust_aligned long lost **-\$99** — momentum entries that bought thrust into what turned out to be down-legs. Even **capit_loose** short lost **-\$251**, a reversion engine fading up-pops that kept popping.

Now the same names on the right side of the tape, because the point is not that these strategies are broken — it is that they were firing in the wrong conditions. **grind_fast** made **+\$747.50** long and **+\$552** short — a momentum engine that wins in both directions when it is aligned. **thrust_loose** short made **+\$500.50**, **thrust_cont** short **+\$433**, **thrust_aligned** short **+\$293**, and the absorption-veto shorts **abs_veto_50s** and **abs_veto_55s** made **+\$356** each. Line up **thrust_aligned**'s two sides and the router thesis is undeniable: the exact same entry logic made **+\$293** short and lost **-\$99** long in the same week. The strategy was never the variable. The regime was.

The counterfactual: what the router would have banked

The router does two things — it stands the wrong engine down, and it routes the capital to the engine the regime warrants. Both are already priced in this week's data, in the same regime windows, so this is a real head-to-head and not a hypothetical.

In the down-trend windows specifically: the reversion-long engine fired 109 times and lost **-\$2,653**. In those same down-trend windows the momentum-short engine fired 143 times and made **+\$3,170**. Same clock, same tape, opposite outcome. The router's job is simply to be running the second engine and not the first when the tape is falling — a swing of roughly **\$5,800** between the right call and the wrong one.

Totalled across the book, the four against-trend buckets — reversion-long into down (**-\$2,653**), momentum-long into down (**-\$1,056.50**), reversion-short into up (**-\$117.50**), momentum-short into up (**-\$83.50**) — bled **-\$3,910.50**. Every dollar of it is an engine that a regime gate would have held flat. Keep only the aligned buckets plus the small chop-reversion bucket and the book goes from its actual **+\$1,829.50** to roughly **+\$5,740**. That is not a tuning gain of a few percent; it is more than tripling the net by deleting trades we already know were structurally doomed the moment they fired.

Actual — every engine fires every regime	+\$1,829.50	—
Router — against-trend engines stood down	+\$5,740	+\$3,910

The honest caveats — why the live number will be smaller than \$5,740

I will not sell you the idealised figure as if it were bankable, because the run list is a hindsight label. Each run's window is drawn from its start to its peak, and the peak is only known after the fact. A live reactive classifier cannot see the peak; it can only measure the tape that has already printed. So in practice the router lags the turn — it will be a bar or two late recognising a down-trend has begun, and a bar or two late releasing when the trend dies. That lag costs some of the edge at both ends. The saving is therefore a ceiling, not a promise. What protects most of it is that this week's bleed lived inside sustained 12-to-15-minute down-legs, not one-bar spikes — a classifier that needs even a few minutes of confirmation still sits inside the meat of a 15-minute run. The money was not made or lost in the first 30 seconds; it was made or lost by being on the right side for ten minutes.

Two more honesty notes. First, every one of these 414 trades is **filled** with a repriced real_pnl — this is not marks, not perfect 1-ATR fills, so the winners are as honest as the losers and the +\$3,170 momentum-short number is real money on the tape, not a simulation flatter. Second, the chop-reversion bucket that makes the router's "run reversion in chop" leg attractive is only six trades for +\$171.50. That is a coin's worth of evidence, not a proven engine. The safe first version of the router does the high-confidence half — stand the wrong engine down in a clear trend — and leaves "which engine to run in chop" as stand-down-or-tiny until we have a real sample. The trend gate is where the −\$3,910 actually lives; the chop question is a rounding error by comparison.

Anti-flicker, because a router that whipsaws is worse than none

The one way this build hurts us is if it flickers. If the classifier flips from "trend" to "chop" to "trend" every time a 4-minute pop crosses its threshold, it will tear the wrong engine off mid-winner and staple the wrong one on mid-loser, and we will pay the switching cost on both. The design has to carry hysteresis: a regime must persist for a minimum confirmation window — think a handful

of consecutive 5-second bars agreeing on displacement-over-path — before the router is allowed to change which engine is live, and a wider band to switch out of a regime than to switch in. Reactive classification of the regime you are standing in, per contract, with a deliberate stickiness so the run list's short 4-minute runs (there was a 102-point one at 19:26) do not yank the whole book around. Predicting the next regime is dead and we are not reviving it; the router only ever answers "what am I in right now, and has it held long enough to trust."

Monday move. Build the router shadow-first, and make it prove this counterfactual before it arms a single live contract. Concretely: (1) write a reactive per-contract regime classifier over the 5s bars — trailing ~15-minute displacement-over-path, three states TREND-UP / TREND-DOWN / CHOP, with N-bar confirmation hysteresis and an asymmetric exit band; (2) stamp every existing shadow trade with the regime live at its entry using only past bars — no peek at the run peak; (3) replay the board with the against-regime engine suppressed and confirm the net climbs from **+\$1,829.50** toward the **+\$5,740** ceiling, minus whatever the lag eats; (4) do not arm until that replay holds clean on next week's fresh tape too, so we know we gated a regime and not a hindsight label. The first live version gates the trend case only — reversion and momentum-long both stand flat in a confirmed down-trend, momentum-short runs — because that is where 68% of the bleed lived in one bar of this chart.

The direction-by-regime split is the single biggest lever on the board right now. It is bigger than any exit tweak, any threshold, any per-contract force-kill, because it is not trimming a losing trade — it is refusing to place it. This week the tape handed us a clean proof: the same strategies that bled buying a falling market were the strategies that printed shorting it. We do not need better engines. We need to stop running them backwards.


BUILD — SHADOW-FIRST ROUTER, GATE THE TREND CASE, PROVE THE +\$3,910 BEFORE IT ARMS

§19 grind_fast, trade by trade: it enters mid-run and lets the chandelier ride — then hands a third of it back


You asked to see how grind_fast actually works, one trade at a time. Here are its ten best on the V7 board, each drawn from the real 5-second tape: the green dot is where it got **in**, the amber dot is where the chandelier let it **out**, and the dashed amber line is the **peak** it reached in between. These ten alone made **+\$1710**. Read them and two things jump out — the good and the tax.

One note on the numbers so nothing looks off: the point-counts on each chart are the bar marks you can see on the tape; the dollar figure is the honest tick-repriced fill — what the desk would actually bank — which runs a little richer than the round marks. We score on the honest number, always.


The good: grind_fast doesn't chase tops. On every one of these it enters *into* a move that's already going and holds while it runs — 8 to 35 minutes — instead of grabbing a few points and bailing. The tax: look at the gap between the amber peak line and the OUT dot on every chart. That gap is money it gave back, and it is not small.




#1 LONG 17 Jul 13:53 UTC · 11 min — it rode a **194-point** peak run and handed back **65** of those points (34%) from the peak before the chandelier let go. Banked **+\$334** on honest fills.




#2 LONG 17 Jul 15:42 UTC · 19 min — it rode a **170-point** peak run and handed back **66 of those points (39%)** from the peak before the chandelier let go. Banked **+\$242** on honest fills.



#3 LONG 17 Jul 14:04 UTC · 8 min — it rode a **145-point** peak run and handed back **48 of those points (33%)** from the peak before the chandelier let go. Banked **+\$210** on honest fills.




#4 SHORT 17 Jul 08:13 UTC · 5 min — it rode a **91-point** peak run and handed back **16 of those points (18%)** from the peak before the chandelier let go. Banked **+\$160** on honest fills.




#5 LONG 17 Jul 16:32 UTC · 21 min — it rode a **119-point** peak run and handed back **27 of those points (22%)** from the peak before the chandelier let go. Banked **+\$152** on honest fills.




#6 SHORT 17 Jul 01:13 UTC · 35 min — it rode a **84-point** peak run and handed back **37 of those points (44%)** from the peak before the chandelier let go. Banked **+\$142** on honest fills.




#7 SHORT 17 Jul 19:06 UTC · 12 min — it rode a **92-point** peak run and handed back **16 of those points (17%)** from the peak before the chandelier let go. Banked **+\$134** on honest fills.



#8 SHORT 17 Jul 13:36 UTC · 12 min — it rode a **116-point** peak run and handed back **97 of those points (84%)** from the peak before the chandelier let go. Banked **+\$116** on honest fills.



#9 SHORT 17 Jul 03:43 UTC · 22 min — it rode a **88-point** peak run and handed back **10 of those points (12%)** from the peak before the chandelier let go. Banked **+\$114** on honest fills.



#10 SHORT 16 Jul 23:13 UTC · 14 min — it rode a **59-point** peak run and handed back **23 of those points (39%)** from the peak before the chandelier let go. Banked **+\$103** on honest fills.

What the ten charts are telling you

Every single one of these winners exited on the chandelier, never a target — that is the whole point of `grind_fast`, and it's why it can bank a \$334 run instead of clipping it at \$80. The entries are genuinely good: it gets in while the move is live and it works both ways (six of the ten are shorts). This is a real momentum-rider, and on a trending day like 17 July it is the best thing on the board.

But the amber gaps are the story for Monday. On #1 it peaked near 28640 and let go at 28590-ish — that's roughly \$100 left on the table on a single trade. Across the ten, the give-back from peak to exit averages a meaningful slice of each run. The chandelier's trail is a fixed multiple of ATR, so on a high-volatility day like this one the trail sits wide and the runner bleeds back before the stop fires. This is exactly the give-back you watched live (+\$224 down to +\$88) — it is not a one-off, it is how the exit behaves on wide days.

What to do — and the honest catch

The Monday move on `grind_fast`'s exit is clear: an ATR-conditional trail — tight on high-ATR days, loose on calm ones — recovers most of these amber gaps and beat the flat trail by **+\$1,653** over the 25-day replay. Build that.

The honest catch you have to hear: these are grind_fast's ten best trades on its best day. Zoom out and the same sim loses **-\$7,975** over 25 archived days — it only prints on days that trend and bleeds every chop day. So this is not "arm grind_fast." It's "grind_fast is a genuine run-day rider whose exit we can sharpen — and it belongs behind a switch that only turns it on when the day is actually running." The charts show you a real edge; the 25-day number shows you it's a specialist, not an everyday gate. **KEEP AS A RUN-DAY SPECIALIST • FIX THE GIVE-BACK • NEVER STANDALONE-LIVE**

§20 **thrust_cont — the quiet winner that just takes its target**

thrust_cont made **\$594** across 29 filled trades over the two days it has been on the V7 board (07-16 and 07-17), winning 52% of them. All 29 are MNQ. It is one of only a handful of sims on the new board that finished the week clearly green, and — unusually for this desk — it made money on BOTH sides of the market. There is no bleeding bucket to cut here. The honest question is not "what is broken" but "is this real enough to arm, and does anything make it bigger."

One-sentence money: thrust_cont earned **\$594 over 29 trades at 52% win**, and every dollar of it is the take-profit target doing the work — its 13 TARGET exits banked **\$1,008** while its 16 STOP exits gave back **\$414**.

Why it made money — it wins by getting paid at the target, not by being right often

This sim is mechanically simple and the data shows it plainly. There are exactly two ways it leaves a trade, and they map almost perfectly onto win and loss. When it reaches its profit target it wins 100% of the time (13 of 13, by definition), and those exits carry the whole book at **\$1,008.50**. When it stops out it loses 87% of the time (14 of 16), bleeding **\$414.50**. So the edge is not clever exits or trailing — it is that the average winner (\$77.6) is bigger than the average loser (−\$25.9). It only needs to be right half the time to come out ahead, and this week it was right 52%.

Split by side, the surprise is that the SHORTS carried it. Shorts made **\$433** over 16 trades at 63% win, longs made **\$161** over 13 trades at just 38% win. Both sides are green, which is the thing worth protecting — most sims on this desk lean hard one way. But the longs are clearly the weaker leg: they hit

their target only 3 times in 13 tries. The reason the longs still finished positive is that when a long DID run, it ran big — the two largest single trades in the whole sample were both longs on loud tape (\$113.5 and \$176.5). Take those two away and the long book is a net loser. So "longs work" is really "two long runners paid for eleven small long stop-outs." That is thin ice, and I would not read it as a durable long edge on two days of data.

SHORT TARGET		+\$663
LONG TARGET		+\$345
LONG STOP	-\$185	
SHORT STOP	-\$230	

Split by the loudness of the tape at entry ($\text{atr_pct} = \text{entry_atr} \div \text{entry_price}$), a cleaner pattern shows up than the side split does. The trades break into three amplitude bands and the money climbs monotonically with volatility: the quiet 0.04–0.06% band made **\$153** over 14 trades at 50% win, the middle 0.06–0.08% band made **\$210** over 12 at 50%, and the three loudest trades ($\geq 0.08\%$) made **\$231** at 67%. The single biggest winners of the week — the 0.099% long (+\$113.5) and the 0.155% long (+\$176.5) — sit right at the top of that ladder. In plain words: this thing gets paid when the thrust it is chasing is genuinely big, and it grinds sideways-to-flat when the move is small.

By session hour the picture is noisier and I would not lean on it. The good hours were scattered — the US afternoon runs at 14:00 and 15:00 UTC each produced a single fat winner (\$176.5 and \$113.5), and the 08:00 hour went 2-for-2 for \$166.5. The soft patch is late-US-afternoon into the 16:00–19:00 window: hour 16 lost \$110 on two trades, and the 19:00 hour was the busiest of the day (5 trades) but only 40% winners for a slim \$38. With one or two trades in most hourly buckets this is basically noise, so I am treating hour-of-day as descriptive colour, not a gate.

What it does when it wins vs when it loses, in human terms

When it wins, it has jumped onto a real, loud thrust — often a short into a flush — and the move keeps going far enough to tag a fixed target before it gives back. When it loses, it has jumped onto a smaller or half-hearted thrust, the move stalls, and the fixed stop takes a small, controlled bite. The losses are cheap and uniform (nearly every stop is between $-\$26$ and $-\$52$); the wins scale with how loud the tape was. That is exactly the shape you want from a continuation sim, and it is the opposite of the reversion family that

bled the V5 board — this one is paid to be on the right side of momentum, not to fade it.

The one thing for Monday — gate on tape loudness, not on side

The cleanest single split that separates its winners from its losers is amplitude at entry, at a threshold of $\text{atr_pct} \geq 0.07\%$. Above that line it took 9 trades for **\$468** at 67% win. Below it, the other 20 trades made only **\$126** at 45% win — barely above breakeven. Two-thirds of the profit came from under a third of the trades, and the quiet-tape trades did almost nothing but churn.

BUCKET	TRADES	REAL P&L	WIN%
loud tape — $\text{atr_pct} \geq 0.07\%$	9	+\$468	67%
quiet tape — $\text{atr_pct} < 0.07\%$ (the rest)	20	+\$126	45%

The gate, stated plainly: only take a `thrust_cont` entry when the ATR is at least 0.07% of price — stand down on quieter tape. In-sample this concentrates the same book into a third of the activity: you keep \$468 of the \$594, forgo \$126, and skip 20 marginal trades. That is not a dollar windfall — the honest in-sample swing on winrate is the real prize, 67% vs 45%, not the P&L, because the quiet bucket was still slightly green. This is a "trade less, keep the quality" filter, and on a 29-trade, two-day sample I want to be blunt that a nine-trade good bucket is thin. It rhymes with the amplitude-floor finding that already survived on the V5 momentum shadows (arm the thrust only when the move is loud), which is the only reason I would act on so little data.

Verdict: **ARM SMALL & WATCH** — `thrust_cont` is genuinely +EV on this week, green on both sides, with cheap uniform losses and target-driven wins. Add the $\text{atr_pct} \geq 0.07\%$ amplitude floor as an entry veto (in-sample: 67% win on the kept trades vs 45% on the dropped). Do NOT read the long edge as durable — it rests on two runners. Re-check both after another full week before sizing up; 29 trades over two days is not enough to promote.

One caveat I will not paper over: `exit_reason` here is deterministic (TARGET always wins, STOP mostly loses), so no "exit tweak" is available in this data — the only levers are which entries you allow. The amplitude floor is the honest one. Widening the target or the stop is untestable from these 29 rows and would be a guess. Leave the mechanism alone; gate the entry.

§21 thrust_loose made money the same way all week: the short side takes its target, the long side just donates

Honest week. thrust_loose booked **+\$462** across 43 filled trades at a 44% win rate. It is a real, positive strategy — but read one layer down and it's really two strategies stapled together: a short book that works and a long book that quietly leaks. Kill the long side and the same week is worth **+\$500**.

The money in one sentence: thrust_loose fires on a loose thrust, sets a fixed target and a fixed stop, and this week the 20 trades that hit their **target** paid **+\$1,221** while the 23 that hit their **stop** gave back **-\$759** — net **+\$462**, and almost all the surviving profit is short trades taking their target.

Why it won: a plain fixed-target exit, carried entirely by shorts


There is nothing clever in the exit. This sim has exactly two ways out — it either reaches its **TARGET** or it hits its **STOP**. No trailing, no clock, no give-back mechanics to autopsy. So the whole question is: which trades reached the target, and which bled to the stop?

Split by side and the answer is stark. The short book made **+\$500** over 23 trades at a 57% win rate. The long book lost **-\$38** over 20 trades at a 30% win rate. Same entry logic, same target-and-stop geometry, opposite outcomes — the shorts are doing 100% of the work and then some, because the longs are a small net drain.

BUCKET	TRADES	REAL P&L
SHORT → TARGET	14	+\$827.0
LONG → TARGET	6	+\$394.5
SHORT → STOP	9	-\$326.5
LONG → STOP	14	-\$433.0

The killer row is **LONG → STOP**: 14 of the 20 long trades never got going and each one paid the stop, for a **-\$433** hole that the 6 long winners (**+\$394.5**) couldn't quite fill. The shorts have the same problem in miniature — 9 stops

for -\$326.5 — but their 14 winners (+\$827) bury it comfortably. That is the entire week in four rows.



Real-P&L contribution by side × exit, honest tick-repriced fills. Zero line marked; the single green SHORT→TARGET bar is bigger than both red stop bars combined.

The amplitude tell: it needs a real thrust, not a twitch

One more thing separates the good trades from the bad, and it's the size of the move it's trading into. Split by ATR-as-a-fraction-of-price, the bigger-thrust trades ($\text{atr_pct} \geq 0.07\%$) made **+\$387** over 13 trades at 54% wins, while the smaller-amplitude majority (0.04–0.07%) made just **+\$75** over 30 trades at 40% wins. And the two reads combine cleanly: the four **SHORT + high-ATR** trades went 4-for-4 for **+\$343**. That n is tiny and I won't oversell a 100% bucket, but it points the same way as everything else — this thing wants a genuine thrust with amplitude behind it, and it does badly on thin low-amplitude longs.

The Monday move

The gate: drop the long side. Trade `thrust_loose` **short-only**.

The long book this week was 20 trades for **-\$38.5** at a 30% win rate — it added risk and cost turnover to net nothing. The short book alone was 23 trades for **+\$500.5** at 57%. Cutting the longs lifts the in-sample week from +\$462 to **+\$500**: a +\$38 swing in dollars, but a jump from 44% to 57% wins and a much cleaner book — you stop paying 14 long stops to collect 6 long targets.

Be honest about what this is and isn't. The dollar lift from the side gate is small because the longs were already near break-even, not a gusher — the real prize is turnover and drawdown, not P&L. The high-ATR read is the sharper edge but rests on 13 trades (four on the short-plus-amplitude corner), so treat it as a lean, not a law: prefer the shorts, prefer the ones firing into a thrust with amplitude, and let the thin low-ATR longs go.

This is a keeper with one obvious trim. It earns its place on the short side and should stop pretending it's two-sided.

§22 thrust_aligned made \$194, but the short side made \$293 and the long side handed almost half of it back

This one finished the week green: \$194 in the till across 23 trades, winning a shade under half the time (43%). Small, honest, real. But the money is not spread evenly — the whole profit came from the short side, and the long side quietly bled a third of it back out.

Here is the plain version. thrust_aligned goes with a thrust once price is lined up the way it wants, and then it lives or dies on two exits: it either reaches its **TARGET** and books a clean win, or it gets run into its **STOP** and takes a fixed clip. There is almost nothing in between. Eleven trades reached target and made \$669 (10 of those 11 were winners — one "target" actually settled a touch negative on honest fills, at $-\$27.50$, which is why it is not a perfect 100%). Twelve trades hit the stop and lost \$475, none of them positive. So the week's \$194 is really just "target dollars minus stop dollars." There is no clever exit management doing work here — the trade takes its number or it gets taken out.

The short side carries the whole thing

Split it by direction and the story is stark. The 17 short trades made **+\$293** and won 53% of the time. The 6 long trades lost **-\$99** and won just 17% — a single winner out of six. Five of those six longs went straight to the stop. In human terms: when it shorted a thrust this week it got paid; when it bought a thrust it mostly got run over.

SHORT (17)

+\$293

LONG (6)

-\$99

The one long that worked (id 1507, $+\$113.50$) was the widest-ATR trade of the entire week — it fired at an atr_pct of 0.099, roughly double the typical entry. The five long losers all fired into more ordinary conditions and got stopped. So it is not that "longs never work" — it is that only a genuinely violent up-move carried a long far enough to pay, and those are rare. Every routine long was a stop.

BUCKET	TRADES	REAL P&L	WIN%
SHORT side	17	+\$293.00	53%
LONG side	6	-\$99.00	17%
TARGET exit	11	+\$669.00	91%
STOP exit	12	-\$475.00	0%
Whole sim	23	+\$194.00	43%

I checked whether the ATR band or the time of day sorted the winners from the losers on the short side, and it doesn't cleanly — short winners span `atr_pct` 0.04 through 0.075, and short losers sit right in the same range. There is no volatility band or session hour that carves the shorts into good and bad. The one axis that separates money from losses is **side**, full stop.

The Monday move

The gate writes itself: **take the short signals, drop the longs**. Shorts alone were +\$293 over 17 trades at 53%; the longs subtracted \$99. Muting the long side is an in-sample swing of **+\$99** and lifts the whole sim from +\$194 to +\$293 with a higher hit rate.

The honest caveat: the long side is only 6 trades, so this is a thin-n read, not a law. But it is not a coincidence either — it rhymes with the whole thrust family this week, where every green sibling (`thrust_cont`, `thrust_loose`, `abs_veto`) was carried by shorts and did little or nothing long. The direction of the bias is the signal even if the exact -\$99 is noisy. And the fix costs nothing: a short-only `thrust_aligned` would have kept the entire \$293 it earned and skipped every long stop.

So it is a keeper, but a lopsided one — a short-side thrust taker wearing a long-side anchor. Cut the anchor, watch the short-only version live to confirm the 6-trade long read holds up, and don't dress this up as anything more than "take its target or take its stop." That is all it does.

KEEP THRUST_ALIGNED — SHORT-SIDE ONLY

WATCH THE LONG SIDE (THIN 6-TRADE SAMPLE) BEFORE PERMANENTLY MUTING

\$23 thrust_fast made \$95 — but only because the short side carried a long side that's pure tax


thrust_fast fired 17 times this week and came out \$95 ahead on honest, tick-repriced fills. It won 6 of them. That's a real, if slim, green week. But the \$95 hides two completely different desks bolted together: a short book that made \$227 and a long book that lost \$132. Nothing else happened — no chandelier, no time-out, no cleverness. Every trade ended one of two ways: it took its full target, or it hit its stop.

The money in one sentence: thrust_fast made \$95 over 17 trades by taking a fixed target 6 times (\$507) against 11 stop-outs (\$412) — and every honest dollar of profit came from shorting; the long side handed back a third of it.

What it does when it wins, and when it loses

This is a clean binary strategy — there's no exit machinery to argue about. When it's right, it holds to a target and books the whole thing; when it's wrong, it eats the stop. The six winners paid \$44.50 to \$133.50 each; the eleven losers cost \$25.50 to \$58.50 each. Win rate is only 35%, but the winners are bigger than the losers, so the maths just barely clears zero. That's the entire personality: it's a low-hit-rate, positive-payoff take-your-target thruster.

The split that matters is direction. Look at the two books side by side.



SIDE	TRADES	WINS	WIN%	REAL_PNL
SHORT	10	5	50%	+\$227.0
LONG	7	1	14%	-\$132.0

Shorting, thrust_fast hit a coin-flip 50% and every winner was a clean target: +\$84, +\$48, +\$83.50, +\$44.50 and the standout +\$133.50. Going long it won exactly once in seven tries — a single +\$113.50 — and the other six were all stops, including its two ugliest losses of the week (-\$54.50 and -\$58.50). The long side isn't just weaker; it's a straight leak.

The ATR angle — same story, second axis

Score the same 17 trades by how much the tape was moving at entry ($\text{atr_pct} = \text{entry_atr} / \text{entry_price}$) and you get a near-mirror of the direction split. When ATR was rich (0.06% or more) it made \$217.50 on 8 trades at 50%; when the tape was quiet (under 0.06%) it lost \$122.50 on 9 trades at 22%. Every one of its big-target wins — the +\$84, +\$83.50, +\$113.50, +\$133.50 — came out of a live, moving tape, because a fast thruster's target is sized off ATR: bigger range, bigger take. In dead air the target is too small to pay for the stops around it.

ATR BAND	TRADES	WIN%	REAL_PNL
0.06%+ (live tape)	8	50%	+\$217.5
under 0.06% (quiet)	9	22%	-\$122.5

The Monday move

Two knobs point the same way, and the cleaner, better-populated one is direction. **Drop the long book.** Trading shorts only, `thrust_fast` is +\$227 on 10 trades at 50% — a genuinely respectable little short-thruster. The long side, all 7 of it, is -\$132 of tax you can cut with one line. In-sample that lifts the week from +\$95 to +\$227, a +\$132 swing, and it does it by removing trades that lost, not by curve-fitting the winners.

The gate, one line: only let `thrust_fast` take shorts. Optionally stack the ATR floor (0.06%+ at entry) on top — that trims the quiet-tape shorts too and concentrates it to +\$261 on the 4 live-tape shorts, though at $n=4$ that's a lean-into, not a promotion.

Be honest about the size of this: 17 trades is one week and one MNQ tape. Shorts-only is 10 trades — enough to act on as a lean, not enough to bet the desk on. But the direction of the read is the same one the whole thrust family keeps showing on this board, and there's zero cost to muting a long side that hit 14%. Keep it observing, muzzle the longs, and let the short book prove itself over another two weeks before anyone talks about arming it.

WATCH — SHORTS ONLY, MUTE THE LONG BOOK (+\$132 IN-SAMPLE), RE-CHECK THE SHORT READ AT N OF 25 OR MORE

§24 **abs_veto_55s: the short side and a bit of amplitude do all the work; the quiet-tape longs are dead weight**

This one had an honest, quiet week: **\$499** in the till over 21 trades, and it won 57% of them. No single fat trade carried it — the biggest winner was \$113.50 and the worst loss was \$68. It just took a lot of small, tidy targets and gave back a handful of small stops. That is exactly the shape you want from a veto gate.

The plain sentence: **abs_veto_55s made its \$499 by shorting well and by only really paying when the tape had some amplitude behind it.** The short book alone put up **\$356** at 64% wins; the long book scraped together **\$143** at a coin-flip 50%. And every dollar of edge lived above a modest volatility floor — the quietest trades made almost nothing.

Where the money came from

Start with the two sides, because that is the cleanest cut. The shorts are the engine here and they are beautifully repetitive: seven winning targets between \$44 and \$84, three small stops around minus-\$29, and one target that filled for a \$32.50 loss. Tight, symmetric, boring — good boring. The longs are the messy half: they broke even on count (5 up, 5 down) and only netted \$143 because one "target" long actually filled for a \$68 loss and dragged the group down.

SIDE	TRADES	REAL P&L	WIN %
SHORT	11	\$356.0	64%
LONG	10	\$143.0	50%

Now the exits, which tell you how it wins. It is a pure two-outcome machine: 13 trades hit the profit target for **\$661** at an 85% win rate, and 8 trades hit the stop for **-\$162**. That is the whole story — take the target most of the time, eat a small stop the rest. (One honesty note: these are tick-repriced fills, so the labels don't map perfectly to sign — one stop closed up \$82 and one target closed down \$68 because the fill landed the other side of the mark. The dollar totals are the honest ones; the point-marks would read a touch differently.)

TARGET

+\$661

STOP

-\$162

The condition that separates winners from losers

Split its trades by how much amplitude was in the tape at entry — the ATR as a percent of price — and the picture snaps into focus. Everything below a 0.06% ATR floor was a coin flip that barely paid its way; everything at or above it did the real earning.

ENTRY AMPLITUDE	TRADES	REAL P&L	WIN %
ATR% $\geq$ 0.06% (has thrust)	10	\$415.5	70%
ATR% < 0.06% (quiet tape)	11	\$83.5	45%

Look at what that does. Just over half the trades — the quiet-tape half — delivered only \$83.50 at a 45% hit rate. That is noise wearing a badge: eleven trades of effort for a tenth of the profit. The other ten trades, taken when there was genuine movement to fade or ride, carried \$415.5 of the \$499 at a clean 70%. The middle band specifically (0.06%–0.09% ATR) was the sweet spot at \$370 over 8 trades and 75% wins.

The Monday move: add an amplitude floor of `atr_pct  $\geq$  0.06%` to `abs_veto_55s` and stand it down in dead-quiet tape. In-sample that lifts win rate from 45% to 70% and concentrates \$415.5 of the week's \$499 into 10 trades instead of 21 — you keep the money and skip eleven coin-flips. Honest caveat: this is a one-week, thin-n read (10 vs 11 trades), and the quiet half still made a small *positive* \$83.50, so it is dead weight, not a bleeder. This is a "trade less, trade cleaner" tweak, not a rescue.

Verdict

`abs_veto_55s` is a real, if modest, contributor: both sides green, a disciplined target-or-small-stop shape, and no reliance on a lucky outlier. The short book is the part to trust; the longs need the amplitude floor to stop coin-flipping. Keep it running as a shadow, add the 0.06% ATR gate, and let another week or two of trades confirm the floor before it earns any size.

KEEP ABS_VETO_55S – ADD ATR_PCT $\geq$ 0.06% FLOOR, LEAN ON THE SHORT SIDE

\$25 `abs_veto_50s` — the short-side absorption sim that quietly banked a clean \$443

This is one of the good-news sims of the week. Over the two V7 days it took 22 trades on MNQ, won 55% of them, and finished \$443.50 in the black on honest

fills. Nothing dramatic, no single hero trade — it just kept its losers tiny and let its targets pay. It sits right in the middle of the V7 leaderboard alongside its sibling `abs_veto_55s` (+\$499) and the thrust family (`thrust_cont` +\$594, `thrust_loose` +\$462), so it is real company, not an outlier.

1. The money, in one sentence

`abs_veto_50s` made **\$443.50** across 22 filled MNQ trades at a 55% win rate, and it captured about two-thirds of its own best-case ceiling (\$681 of unrealised peak came home as \$443.50) — a give-back that fits the chandelier pattern we see everywhere, not a broken exit.

```
Query: SELECT count(*), sum(real_pnl), avg(real_pnl>0) FROM shadow_trades t JOIN
shadow_real r ON r.trade_id=t.id WHERE strategy='abs_veto_50s' AND fill_status='filled' →
22 trades, $443.5, 54.5% win. Ceiling vs real: $681.0 → $443.5.
```

2. Why it made money — where the dollars actually came from

Three things stand out when you pull the trades apart, and they all point the same way: **this is a short-biased sim with a clean, disciplined exit that shines in a specific volatility band.**

The short book carried it. Split by side, the eleven shorts made **\$356.00** at a 64% win rate, while the eleven longs scraped together **\$87.50** at just 45%. More telling than the totals is the shape: no short in the whole week lost more than \$33, so the downside was tightly capped, while the short winners ran \$44 to \$84 each. The longs were the opposite — they held the two ugliest trades in the book (a −\$68 target-miss and a −\$55.50 stop, both in high volatility) alongside the single best trade (+\$113.50). Longs are a coin-flip here; shorts are the engine.

The exit is doing its job. The exit-reason split is about as clean as you will see: the 13 TARGET exits made **\$661.00** at an 85% win rate, and the 9 STOP exits lost **−\$217.50** at 11%. That is exactly the profile you want — when it's right it rides to target, when it's wrong the stop cuts it fast. Every stop-out was small (worst −\$55.50). One STOP even locked in +\$82 on a long that had already run, so the stop is trailing profit, not just eating losses. There is no third failure mode bleeding money in the background.

It wins in the middle of the volatility range. Bucketed by entry ATR-percent (`entry_atr ÷ entry_price`), the sim is a mid-vol animal. The 0.06–0.08% band made **\$370.00** at 75% over 8 trades — that one band is 83% of the

entire week's profit. The quiet 0.04–0.06% band made only **\$83.50** at 45%, and the three high-vol trades above 0.08% netted **−\$10.00**. In plain words: when the tape is dead-quiet it chops around breakeven, when it's wild the longs blow up, and in the sweet middle it prints.

ATR .04–.06%

+\$83.5

ATR .06–.08%

+\$370

ATR > .08%

−\$10

Session hour tells a softer version of the same story. The strong windows were the 08:00 hour (+\$231, 3-for-3), the 00:00 hour (+\$126, 2-for-2) and the 19:00 hour (+\$102). The red came in the US lunch lull and late tape — the 11:00 hour lost −\$57.50 (0-for-2) and 13:00 lost −\$55.50 (0-for-1). It's the right intuition (avoid the dead midday chop), but with one or two trades per hour it's too thin to gate on by itself. I'd treat hour as confirmation of the ATR story, not an independent rule.

CUT	N	REAL_PNL	WIN%
SHORT	11	\$356.00	63.6%
LONG	11	\$87.50	45.5%
Exit = TARGET	13	\$661.00	84.6%
Exit = STOP	9	−\$217.50	11.1%
ATR 0.06-0.08%	8	\$370.00	75.0%
ATR < 0.06%	11	\$83.50	45.5%

3. The one thing for Monday

The cleanest single separator is **entry ATR-percent**, and the nice part is it splits the book almost exactly in half so it isn't a three-trade fluke. Trades entered with ATR at or above 0.06% of price made **\$360.00** over 11 trades at 64%; the low-vol half below 0.06% made only **\$83.50** over the other 11 at 45%.

The gate: only arm abs_veto_50s when entry ATR is $\geq 0.06\%$ of price (skip the dead-quiet tape). In-sample that lifts per-trade take from about \$7.60 to \$32.70, and cuts the trade count in half while keeping 81% of the profit. The dollar swing between the two halves is **\$276.50** — the high-vol half carried essentially all the edge.

Where I add the honesty: this is a filter that says "trade less," not one that conjures new money. You don't make more than \$443 by applying it — you make \$360 on far fewer, better trades and stop paying the low-vol chop tax. And it's one week, 22 trades, in-sample. The 0.06–0.08% band being the exact sweet spot (with high-vol $>0.08\%$ actually turning negative) is only three trades wide on the top end, so I would frame it as a floor — "not below 0.06%" — rather than a tight window, until more days come in.

The tempting second move is to cut longs entirely, since shorts made four-fifths of the money and never lost big. I'd hold off. The longs are still net-positive (+\$87.50), and killing a positive-EV side on one week's read is how you overfit. The right posture is: **keep it two-sided, gate on the ATR floor, and know that your conviction sits on the short book.** If longs keep being the choppy, blow-up-prone side over the next two weeks, revisit then.

WATCH → ARM SMALL It made honest money with a clean exit and a defensible volatility edge. I'd arm it at one micro alongside its sibling abs_veto_55s (+\$499, same family), apply the $\text{ATR} \geq 0.06\%$ floor, and give it two more weeks before touching the long/short balance or sizing up. Thin-n, but the shape is right and the risk was well-controlled all week.

§26 abs_veto_60s: it hit its target four times out of six and banked \$127 — but it's the slowest, thinnest cut of the absorption-veto family

This is the 60-second version of the absorption-veto idea: wait a full minute to see if the tape gets absorbed before you commit, then ride to a fixed target. Over the two-day board it took just six trades, won four of them, and cleared **+\$127** on honest fills. Every trade it managed to hold to its target made money; every trade that got stopped lost money. There was nothing subtle in between.

The money in one sentence

It made **\$127** over six trades at a 67% win rate, and all of it was four clean target-hits (+\$206.5 together) minus two stop-outs (−\$79.5 together) — no give-back, no drama, just a slow gate that either got its target or got clipped.

TRADE	SIDE	EXIT	ATR%	HOUR	REAL_PNL
1275	LONG	STOP	0.050	23:40	−31.5
1293	LONG	TARGET	0.051	00:51	+59.0
1364	LONG	TARGET	0.044	06:21	+49.0
1375	SHORT	TARGET	0.043	06:38	+49.0
1400	SHORT	STOP	0.076	08:17	−48.0
1459	SHORT	TARGET	0.043	12:03	+49.5

Why it made money — the exit reason is the whole story

Split it any way you like and the same thing pops out: this sim is binary. There are exactly two outcomes. When it reaches its fixed target it books a tidy \$49-\$59 win. When it gets stopped it hands back \$31-\$48. There is no partial exit, no chandelier give-back, no clock-out — just target or stop. That's the cleanest possible read, and it means the only question that matters is: how often does the 60-second veto let it into a trade that runs to target?

TARGET ×4

+\$206.5

STOP ×2

−\$79.5

Every dollar of profit was a target being taken; every dollar of loss was a stop being hit. (Figures are tick-repriced fills, so they run a touch different from a clean points-times-multiplier mark, but the direction is exact: four wins, two losses.)

Long vs short: both sides pulled their weight

Unlike some of the momentum sims where one side does all the work, this one was balanced. The three LONGs made **+\$76.5** (2 wins, 1 stop) and the three SHORTs made **+\$50.5** (2 wins, 1 stop). Neither side is the problem.

That's a small point of confidence — the veto isn't just a bull-tape artefact, it took money going both ways.

Where it broke: the one loud trade got stopped

Look at the `atr_pct` column. The four winners all fired when the tape was quiet — `atr_pct` between 0.043 and 0.051. One of the losers (1275) also sat in that quiet band, so quiet alone isn't a guarantee. But the standout is trade 1400: it fired at `atr_pct` **0.076**, by far the loudest print of the week, and it got stopped for −\$48. The 60-second veto is designed to sit out chaos, and on the one genuinely volatile entry it still walked into a stop. That's the tell — this gate does its best work in calm tape, and the single loud entry cost it.

The Monday move

The cleanest separator is an ATR ceiling. Cap entries at `atr_pct` < 0.06 and you drop the loud 1400 stop entirely:

BUCKET	REAL_PNL	N	WIN%
atr% < 0.06 (quiet tape)	+\$175.0	5	80%
atr% ≥ 0.06 (loud tape)	−\$48.0	1	0%

The gate in one line: **only take `abs_veto_60s` entries when `atr_pct` is under 0.06 — stand down when the tape is loud.** In-sample that's a **+\$48** swing, lifting the two days from +\$127 to +\$175 and the win rate from 67% to 80%.

Be honest about the size of this: it's **six trades**. The `atr` ceiling rests on removing a single loser, so it's a one-trade result dressed as a rule — directionally sensible (it matches the whole absorption-veto thesis of sitting out chaos) but nowhere near proven. And the bigger context is unflattering: its own family did far better with a shorter fuse — `abs_veto_55s` cleared +\$499 and `abs_veto_50s` +\$443.5 over the same window, roughly 3.5× this one. The extra 10-15 seconds of waiting isn't buying anything; it's just taking fewer, later trades.

So: it works, it's green, both sides contribute, and the quiet-tape gate makes sense. But it's the weakest cut of the veto family on a six-trade sample. Keep it running with the `atr`<0.06 ceiling to see if the edge holds past ten trades,

but if you're going to trade one absorption-veto for real, the shorter windows are earning more. **WATCH – QUIET-TAPE ONLY (ATR% < 0.06); PREFER THE FASTER 50S/55S COUSINS**

§27 chand_k20: dead flat — the tight trail banked \$255, the hard stop bled \$237, and four shorts nudged it green

chand_k20 made **\$18.5** over 11 trades this week at a 45% win rate. That is a rounding error. Strip it apart and there is no edge in the entry at all — it is a wash where two exits fought to a draw and a handful of shorts tipped the last dollar into the black.

Here is the whole thing in one sentence: the tight chandelier trail booked **\$255.50** across the six trades it managed to profit, the hard stop handed back **\$237** across five trades where it never worked once, and what was left over — eighteen dollars and fifty cents — is the week. This is not a strategy that found something. It is a coin that landed on its edge.

The exit is the whole story, and the two exits cancel

Every winning trade this sim booked came out on the chandelier — the trailing lock. Five of its six chandelier exits were green (the sixth, the very first trade of the day, trailed off a rip and gave back \$68). Every single stop-out lost money: five stops, five losses, not one saved. So the machine is simple — when a trade works, the tight k20 trail rides it and banks it; when a trade is wrong, the hard stop eats a clean \$40-\$60 and there is nothing to soften it.

The trouble is the two piles are the same size. The chandelier made \$255.50. The stop lost \$237. They are within twenty dollars of each other, which is exactly why the sim reads breakeven. A good week is one where the winners' pile dwarfs the losers' pile. Here they are twins.

LONG

-\$89.5

SHORT

+\$108

Real_pnl by side, chand_k20, filled trades only. Longs: 7 trades, 2 winners (29%). Shorts: 4 trades, 3 winners (75%). Dollars are tick-repriced honest fills, so they run a touch different from the bar-mark points you'd eyeball off a chart.

The shorts worked, the longs didn't — and that's the only signal here

Split it by direction and a difference finally appears. The seven longs lost **\$89.5** and only two of them won — the sim spent the session buying into a tape that kept sliding out from under it, and the stops fired one after another (16:27, 16:51, 17:26, 18:30). The four shorts made **\$108** and three of the four won. Every short but one came out on the chandelier for a tidy gain. In plain words: this sim leaned long into a day that wanted to go down, and the only clean money it made was the few times it went with the slide instead of against it.

BUCKET	REAL_PNL	N	WIN%
SHORT only	+\$108.0	4	75%
LONG only	-\$89.5	7	29%
CHANDELIER exits	+\$255.5	6	83%
STOP exits	-\$237.0	5	0%
Everything	+\$18.5	11	45%

ATR didn't sort the winners from the losers — the whole day ran in a narrow 0.06%-0.12% band and the stops were sprinkled right through the middle of it. Nor did the hour: everything landed in one 15:32-to-19:59 session, which is the real caveat screaming underneath these numbers. This is not eleven trades across a week. It is one day. The "shorts win" read rests on four trades.

The Monday move. The only condition that separates winners from losers here is side: **shorts-only would have made \$108 instead of \$18.5**, an in-sample swing of about **+\$89.5** from dropping the seven longs. But n=4 shorts on a single session is far too thin to arm anything on. Do not build a shorts-only gate off this. The honest read is that chand_k20 is a tie, and the one durable thing it confirms is a signal we already trust from the live desk — the tight chandelier trail is the only exit that makes money (live: CHANDELIER +\$567 at 85% is the sole profitable exit while STOP is 0-win), and k20 is the tightest and only green trail variant on the board this week (vs chand_k25 -\$36.5, chand_k35 -\$101).

Verdict

chand_k20 didn't earn its keep and it didn't embarrass itself — it broke even because two exits cancelled and four shorts broke the tie. There is no entry edge to promote here. Keep it on the shadow board purely as the anchor that proves tighter-is-better on the chandelier trail; give it another week or two of live-adjacent days before reading anything into the long/short skew. Trade nothing off it yet.

WATCH CHAND_K20 — FLAT WEEK, SINGLE-SESSION SAMPLE; THE TIGHT K20 TRAIL IS THE ONE KEEPER SIGNAL, THE SHORTS-ONLY SKEW IS TOO THIN (N=4) TO ARM

§28 chand_k25 — the middle child of the trail A/B, and it bled

chand_k25 is the 25-tick chandelier variant — the momentum entry is identical to its siblings, the only thing that changes is how far the trailing stop sits behind the peak. Over its first honest week on the V7 board it took 11 MNQ trades, won 5 of them (45%), and finished at **-\$36.5**. Small loss, but it's a loss, and it lands exactly where you'd fear on the A/B ladder: the tighter k20 sibling made **+\$18.5**, this k25 lost **-\$36.5**, and the loose k35 lost **-\$101**. Wider trail, worse week. That ordering is the whole story.

The money in one sentence: chand_k25 traded 11 times on MNQ, won 45%, and finished **-\$36.5** — its winners came from short trades that rode the chandelier, and its losses came almost entirely from long trades that got stopped out.

Why it bled — the two cuts that explain it


You don't need a clever filter to see where the money went. Two splits account for the entire loss, and they point at the same trades.

By side. The four short trades made **+\$79.5** and won 3 of 4 (75%). The seven long trades lost **-\$116.0** and won only 2 of 7 (29%). So on a net loss of \$36.5, the shorts were actually carrying the sim and the longs sank it. This is the exact same limp the whole V7 board and the live desk have shown all week — LONG bleeds, SHORT holds — so it's not a quirk of this variant, it's the room it trades in.

By exit. Only two exit reasons fired. The chandelier — the thing being A/B'd — did its job: 5 chandelier exits made **+\$126.0** at 80% win. The stop, on the other hand, is the bleed: 6 stop-outs lost **-\$162.5** at 17% win. When k25 got to trail a runner it kept most of the money; when the trade never got going and hit its initial stop, it paid full freight. The trailing logic isn't broken — the entries that never work are what feed the stop bucket.

CUT	BUCKET	REAL_PNL	N	WIN%
Side	SHORT	+\$79.5	4	75%
Side	LONG	-\$116.0	7	29%
Exit	CHANDELIER	+\$126.0	5	80%
Exit	STOP	-\$162.5	6	17%

By amplitude and clock — nothing to split on. Every one of the 11 trades entered with atr_pct at or above 0.06 (the range was 0.061 to 0.115), so there's no low-vol trap to filter out here — the whole week fired into genuinely moving tape. The amplitude floor is already effectively satisfied. Session hour tells the same side-story in disguise: the profitable window was 19:00 UTC (+\$129 over 4 trades, 75%), but three of those four were the winning shorts, so the "good hour" is really just "the hour the shorts happened." The 16:00 hour lost -\$110 on two long trades. There's no independent clock edge to lean on — it's side wearing a clock costume.



Read the chart as two views of the same -\$36.5. The top pair is the side cut, the bottom pair is the exit cut. Both times the green bar and the red bar are roughly balanced, and both times the red side is the long / the stop-out. It's not a sim that does one thing badly — it's a sim doing a fine job on half its trades and being dragged under by the long book.

What it looks like when it wins vs loses, in plain words

When k25 wins it's short, or it's a trade that got far enough into profit for the 25-tick trail to arm and then ride the move down — like the short at 19:00 that trailed out for **+\$75**, or the long that ran to **+\$75** before the chandelier clipped it. When it loses it's long and early: it buys the rip, the rip doesn't extend, and price rolls back through the fixed initial stop before the trail ever engages. Six of those, one of them a fluke where the stop caught +\$74.5, five of them straight red. The 25-tick trail is wide enough that on the trades that do work it gives back more of the peak than the k20 version — which is precisely why k25 sits below k20 on the ladder.

The one thing for Monday

The single cleanest split is **side**. Trading only the shorts turns this sim's week from **-\$36.5** into **+\$79.5** — an in-sample swing of about **+\$116** — and lifts win rate from 45% to 75%. The long book contributed nothing but drawdown.

Gate, stated plainly: on chand_k25, don't take the longs — trade the short side only. Good bucket: SHORT **+\$79.5**, n=4, 75% win. The rest (LONG): **-\$116.0**, n=7, 29% win. In-sample swing $\approx$ **+\$116**.

Now the honest part. This is 11 trades and the winning bucket is 4 of them. Four trades is not evidence — it's a coin that landed heads three times. I would not arm a "shorts-only" rule off this sim alone. What makes me believe it slightly is that the long-bleed shows up everywhere this week — the live V7 desk ran LONG -\$733 vs SHORT +\$46, thrust_loose ran LONG -\$38 vs SHORT +\$500 — so the direction of the finding is corroborated even if this sim's own n is thin.

The stronger Monday move actually isn't a gate on k25 at all — it's the A/B result the sim exists to produce. k25 lost the trail bake-off to k20 (**+\$18.5** vs **-\$36.5**), and both beat k35. The graded-trail work says an ATR-conditional trail beats any flat width by a wide margin, but among the flat widths tested this week, tighter won. So: **RETIRE CHAND_K25** as a live candidate in favour of the tighter k20 trail, and keep k25 running observe-only purely as the A/B control. The shorts-only tilt is a **WATCH** — real direction, far too thin to size on.

§29 chand_k35: the loosest trail lets longs run into the ground while shorts quietly print

Nine trades, a \$-101 week, and only three of them green. On the face of it a losing sim. But the red is entirely on the long side — the shorts made money, and the actual chandelier exit made money. What bled was the wide trail letting bad longs walk all the way to a hard stop.

chand_k35 is the chandelier family run at the loosest trail setting — k35 gives a rip the most room before it cuts. This week that room cost it. It took nine trades and lost **-\$101**, winning just three of them (33%). Its two tighter siblings did better on the same tape: chand_k20 finished **+\$18.5** and chand_k25 **-\$36.5**. The pattern is clean and it points one way — the wider

you let this thing trail, the more it gives back. k35 sat at the wrong end of that line.

The money in one sentence

It lost a hundred bucks because six long trades bled **-\$232** between them, and three short trades clawing back **+\$131** weren't enough to cover the hole.

Why: the longs are the whole problem

Split it by side and the story is almost embarrassing in how obvious it is. The long book won 1 trade out of 6 and lost \$232. The short book won 2 out of 3 and made \$131. Same sim, same week, and the two sides are pointing in opposite directions.

LONG	-\$232 · 6 trades · 17% win
SHORT	+\$131 · 3 trades · 67% win

Real-fill P&L (tick-repriced), chand_k35, filled trades only. Zero line marked; longs sit left of it, shorts right.

Now look at it by exit, because this is a chandelier sim and you'd expect the chandelier exit to be the story. It is — but not the way the name suggests:

EXIT REASON	TRADES	REAL P&L
CHANDELIER (the trail did its job)	3	+\$102.5
STOP (hard stop caught it)	6	-\$203.5

The chandelier trail — the thing this sim exists to test — actually made **+\$102.5** across the three times it fired. Every dollar of loss came from the six trades that never got the chance to trail and were carried all the way to a hard STOP instead. And five of those six stopped-out trades were longs. So the mechanism isn't broken; the trail works when a trade goes its way. The problem is the entries it was trailing on the long side never went its way — they just walked to the stop.

Condition-wise, the two ugliest longs came in at the highest volatility of the week (atr_pct 0.115 and 0.092, a \$-68 and a \$-58.5), both in the 15:00–16:00 hour. It bought into a hot, jumpy tape and got run over. The shorts that worked came in later (17:00–19:00) at a calmer atr_pct around 0.06 and rode the chandelier to profit. In plain words: it wins when it's short into a settling afternoon tape, and it loses when it's long into a hot mid-session rip.

The Monday move

The single split that separates the winners from the losers here is side, full stop. Drop the longs and you don't have a losing sim anymore:

Gate: shorts only. The short book was **+\$131** over 3 trades at 67% win. The long book was **-\$232** over 6 at 17% win. Cutting the longs is an in-sample swing of **+\$232** — it turns this week's \$-101 into roughly +\$131.

Be honest about what that is, though: it's a 3-trade short sample. That's a coin-flip's worth of data, not a proven edge, and I'd never arm a live side off it alone. The more durable read is the one the whole chandelier family already tells us this week: the k35 width is too loose. k20 beat k25 beat k35 in a straight line, so if this branch survives at all it should be tested tighter, not wider. The wide trail's only redeeming trait — that the chandelier exit itself made money — doesn't rescue it, because the losses came from trades that never reached the trail.

So: don't arm chand_k35 as-is. If the desk wants to keep learning from it, keep it observe-only, tighten the trail toward k20, and watch whether the short-side lean holds up past three trades before anyone believes it.

RETIRE CHAND_K35 AT THIS WIDTH

WATCH A TIGHTER-TRAIL SHORTS-ONLY VARIANT

§30 capit_loose catches a falling knife and gets stopped 4 times out of 5 — the stop is the whole loss

This was the worst sim on the board this week. capit_loose ran 22 trades and lost **\$622.50**, winning just 9% of the time. There's no clever split to save here: the idea itself — buy the capitulation, catch the reversal — got run over almost every single time it fired.

Here's the money in one sentence: it took 22 trades, 18 of them hit the protective stop, all 18 stops lost money, and those stops alone bled **\$606** — which is basically the entire week's damage. The other 4 trades reached their target and roughly washed out at **\$16.50**. So the story isn't subtle. When capit_loose is right it barely gets paid, and when it's wrong — which is four times out of five — it eats a full stop.

Both sides bled, so it isn't a long/short problem

The first thing you'd hope for on a red sim is that one side is the culprit and the other is quietly fine. Not here. The 13 longs lost **\$371.50** at an 8% win rate, and the 9 shorts lost **\$251** at 11%. Both sides are doing the same thing: fading a move that keeps going, then stopping out. You can't flip it, gate it to one direction, or trade only the "good" side — there is no good side.

BUCKET	TRADES	REAL P&L	WIN%
LONG	13	-\$371.5	8%
SHORT	9	-\$251.0	11%

The stop did all the damage — and that's the entry's fault, not the exit's

Split it by how each trade ended and the picture is stark. Every one of the 18 stop-outs was a loser, averaging **-\$33.70** each. The 4 targets averaged about breakeven. When 82% of your trades reach the stop before the target, that's not an exit you can widen your way out of — it's an entry that's simply on the wrong side of the move.

STOP · 18

-\$606

TARGET · 4

\$16.5

That chart is the whole autopsy in two bars: the stop line is the loss, and the target line is a rounding error. A capitulation entry is supposed to buy the panic and ride the snap-back. What actually happened is the panic kept going, ran through the entry, and hit the stop.

The fat-vol trades are the worst of it, but calm doesn't save it either

The only condition that sorts the pain at all is volatility. The three trades that fired when ATR was rich (atr_pct at or above 0.06%) lost **\$161.50** between them — an average of **-\$53.80** a trade, more than double the damage of a normal one. That makes sense: a loose capitulation gate in a fast tape is catching the knife at its sharpest. But cutting those three still leaves you with 19 calmer trades that lost **\$461** at an average of **-\$24.30**. Quiet tape doesn't rescue this; it just bleeds slower.

ATR BAND	TRADES	REAL P&L	AVG/TRADE
atr_pct at/above 0.06% (fat vol)	3	-\$161.5	-\$53.8
rest (calmer)	19	-\$461.0	-\$24.3

The two trades it actually won were both low-vol targets — a long at 0.043% and a short at 0.032% — for a combined **\$41.50**. Two winners out of 22 is not a bucket you can build a rule on. It's noise.

The Monday move. There is no condition split that turns this green — the best you can do is cut the fat-vol tail (atr_pct at or above 0.06%) to save \$161.50 in-sample, and you're still sitting on -\$461 over 19 trades at a 9% win rate. Both sides bleed, every stop loses, and the target barely pays. This is a losing entry, not a tuning problem. Retire it. If the capitulation idea is worth revisiting, it needs a completely different trigger that waits for the snap-back to actually start rather than buying into the fall — not a wider gate on the same broken entry.

RETIRE CAPIT_LOOSE — 9% WIN, THE STOP IS THE ENTIRE LOSS, AND NO SPLIT SAVES IT

Note: P&L is scored on tick-repriced honest fills (real_pnl, filled only), so the per-trade cents won't line up with a clean bar-mark stop distance — the direction and the story are what hold.

§31 **capit_ride: it bought two flushes, both kept flushing, and it never saw a single green tick**

This is the "ride the capitulation" idea — wait for a hard flush and buy the bounce. This week it fired exactly twice, both times it bought a falling MNQ, both times price kept going the wrong way, and both times the stop took it out. Two trades, -**\$61** on honest fills, zero winners.

The money in one sentence: capit_ride went LONG into a flush at 22:12 and lost **-\$32.5**, went LONG into another flush at 02:17 and lost **-\$28.5**, and that's the whole week — **-\$61** over 2 trades, 0% win.

Why it bled: it caught the knife, twice, and the knife kept falling

There's no side to split — both trades were LONG. There's no exit to split — both exits were the same STOP. And here's the tell that matters most: look at the ceiling. The best unrealised point either trade ever reached was already

negative. The first trade's peak was **-\$28.5**, the second's was **-\$54**. That means from the moment it clicked buy, neither position was ever in profit for even a tick. It didn't give back a winner — it never had one. It bought, price dropped, the stop caught it. That's the pure signature of a dip-buyer buying a dip that wasn't done dipping.

TRADE	SIDE	WHEN (UTC)	ATR_PCT	EXIT	PEAK SEEN	REAL P&L
#1	LONG	07-16 22:12	0.046%	STOP	-\$28.5	-\$32.5
#2	LONG	07-17 02:17	0.036%	STOP	-\$54.0	-\$28.5

On atr_pct, both entries landed in the quiet-to-middling band (0.046% and 0.036%) — this wasn't a volatility monster running it over, it was a slow grind against a long that had no business being long. On session hour, one fired at the 22:00 reopen and one in the dead 02:00 overnight, so there's no "it only works in the morning" story to rescue either. Every axis you cut it on says the same thing: it was wrong both times, for the same reason.

#1 L			-\$32.5
#2 L			-\$28.5

Both bars are the honest tick-repriced fills. There is no green bar to show — that's the finding.

Note on the numbers: the P&L above is from tick-repriced fills, so it runs a little different from a clean points-times-multiplier mark. But the direction isn't a rounding question — both trades were underwater from entry to exit, and the fills only confirm it.

The Monday move: there's nothing to gate — it's two-for-two wrong on a two-trade sample

Normally here I'd hand you the one condition that splits winners from losers. capit_ride doesn't have that split, because it has no winners. Every bucket — LONG (the only side), STOP (the only exit), quiet ATR, both sessions — is 0-for-2. There's no atr_pct band to keep, no hour to keep, no side to flip to. The best "gate" the data offers is don't take the trade, which on 2 fills is worth the full **-\$61** of avoided loss in-sample and nothing you'd trust out-of-sample.

The honest read: this is a thin-n dud, not a tunable edge. It fired twice, bought two flushes that kept flushing, and never printed green. Its louder cousin **capit_loose** told the same story at scale this week — - **\$622.5** at a 9% win rate — so the family verdict is consistent, not a small-sample fluke. Buying capitulation on MNQ this week was just catching knives.

Monday: pull capit_ride off the board. Don't tune it, don't re-band it, don't wait for more data on a mechanism whose whole family bled — the capitulation-buy idea is fading knives into a market that keeps going, and the two live shadows we got both proved it.

RETIRE CAPIT_RIDE – 0-FOR-2, NEVER GREEN, AND THE FAMILY CONFIRMS IT

§32 **rg_long_fast keeps buying dips that keep falling — the stop eats it alive**

This one lost money the plainest way there is: it goes long, price keeps going the other way, and the stop takes it out. It bled **-\$493.5** over 51 trades this week and won just 22% of them. There's a good desk hiding inside it, but only on the handful of longs that actually work — the rest is a wall of small stop-outs.

Every single trade was a LONG — that's what "rg_long" means, it only buys. So there's no short side to blame and no side-split to make. The whole story is what happens after it buys, and that comes down to which exit it hits.

The money in one sentence

It took 39 stop-outs at a 5% win rate for **-\$1,218**, and clawed back **+\$724.5** on the 12 trades that lived long enough to trail out on the chandelier — net a **-\$493.5** week.

Why it bled — the exit tells the whole tale

Split its trades by how they ended and the problem is obvious. Three out of every four trades die on the STOP, and those trades almost never work — a 5% win rate means when the stop is what closes the position, the dip it

bought just kept dipping. The chandelier trades, on the other hand, are lovely: 75% of them win, averaging **+\$60** apiece. It's not that the winning idea is broken. It's that the entry buys too many falling knives, and 39 of them get cut before the trail ever gets a chance.

HOW IT ENDED	TRADES	REAL P&L	WIN %	AVG/TRADE
STOP	39	-\$1,218.0	5%	-\$31.2
CHANDELIER	12	+\$724.5	75%	+\$60.4



Look at it by how noisy the tape was when it entered and the same message repeats: it does worst in the quiet-to-normal middle where dips just grind against it, and the only green comes from the six loudest-tape entries.

TAPE AT ENTRY (ATR%)	TRADES	REAL P&L	WIN %
Quiet — under 0.04%	13	-\$201.0	15%
Normal — 0.04 to 0.07%	32	-\$440.5	22%
Loud — over 0.07%	6	+\$148.0	33%

The hour-of-day breakdown is just noise at this size — a good hour here, a bad hour there, nothing you could stand a rule on with one or two trades in most buckets. (The \$ figures are honest tick-repriced fills, so they'll run a touch different from what a clean bar-mark of the same entries would show — the direction is what matters.)

The Monday move

There's no side to cut (it's all longs) and no clean hour to lean on. The one split that separates the winners from the losers is the tape: everything under 0.07% atr — 45 of the 51 trades — lost **-\$641.5** between them at ~19% win. Gate it to **only take entries when atr% is above 0.07%** and you'd have kept the **+\$148** bucket and skipped the rest — a roughly **+\$641** in-sample swing from red to green.

Be honest about it though: that green bucket is 6 trades. That's not an edge you fund, it's a hint. The real diagnosis is that the entry itself is -EV — a 76%-stop, 5%-win-on-the-stop long grab is buying dips that aren't done falling. You can't tune your way out of that with an ATR floor on six trades; the chandelier winners are carrying a broken entry. Stand it down and let the entry logic be reworked before it risks a dollar live.

RETIRE RG_LONG_FAST — THE ENTRY BUYS FALLING KNIVES; 76% STOP OUT AT 5% WIN, AND A 6-TRADE ATR FLOOR ISN'T A FIX.


§33 **rg_long_fast_v (#23): it buys the dip, and one time in four the bounce comes — the other three it gets run over**

You wanted this one dug into as a possible chop-reversion gate, so here it is trade by trade. `rg_long_fast_v` only ever buys — it waits for a stretch below the average price to turn up, then goes long. On the V7 board it took **35 trades and lost \$244** at a **23% win rate**. That number hides the real shape, which the charts make obvious: a handful of clean bottom-catches that pay big, sitting on a pile of dips that just kept dipping. The green dot is where it bought; the amber/red dot is where it got out.


Same note as always: the point-counts are the bar marks; the dollars are the honest tick-repriced fills, which run a little different. We score on the dollars.

When it works — it catches a bottom and rides


Its best trade is the tell. At 13:45 on the 17th it bought 28448 as the morning sell-off exhausted, and rode the afternoon reversal for **+\$320** — it got in earlier than `grind_fast` did on the very same move. When `rg_long_fast_v` is right, it is right at the turn, and the chandelier lets it ride:




✓ **LONG 17 Jul 13:45 UTC · 19 min** — it bought the dip and the bounce came: a ~270-point ride the chandelier banked for **+\$320** on honest fills. This is the trade it exists to catch.




✓ **LONG 17 Jul 19:26 UTC · 6 min** — it bought the dip and the bounce came: a ~64-point ride the chandelier banked for **+\$130** on honest fills. This is the trade it exists to catch.



✓ **LONG 16 Jul 16:49 UTC · 11 min** — it bought the dip and the bounce came: a ~13-point ride the chandelier banked for **+\$106** on honest fills. This is the trade it exists to catch.




✓ **LONG 17 Jul 00:37 UTC · 19 min** — it bought the dip and the bounce came: a ~57-point ride the chandelier banked for **+\$106** on honest fills. This is the trade it exists to catch.




✓ **LONG 17 Jul 18:28 UTC · 6 min** — it bought the dip and the bounce came: a ~4-point ride the chandelier banked for **+\$56** on honest fills. This is the trade it exists to catch.

When it fails — it buys a dip that isn't done


And here is why it's red. Three times out of four it buys into a fall that keeps falling, and the 1-ATR stop takes it out for about $-\$45$. There's no reversal to ride because the move wasn't exhausted — it was still trending down:




✗ **LONG 17 Jul 18:13 UTC • 2 min** — it bought the dip and the dip kept going: stopped 2 min later, ~24 points against it, for $-\$46$. This is the trade that happens 3 times out of 4.



✗ **LONG 17 Jul 18:35 UTC • 5 min** — it bought the dip and the dip kept going: stopped 5 min later, ~22 points against it, for $-\$46$. This is the trade that happens 3 times out of 4.



✗ **LONG 17 Jul 08:23 UTC • 8 min** — it bought the dip and the dip kept going: stopped 8 min later, ~29 points against it, for $-\$49$. This is the trade that happens 3 times out of 4.



✗ **LONG 17 Jul 15:02 UTC • 41 min** — it bought the dip and the dip kept going: held a loser for 41 minutes before the trail gave up, for **\$-72**. This is the trade that happens 3 times out of 4.

The Monday move — and the honest verdict

The pattern is clean once you see it: **rg_long_fast_v buys dips, and dips only reverse when the tape isn't trending**. On a trending-down stretch it's picking a knife out of the air. The single change that would help most is a **trend veto** — don't take the long if the short-window slope is still pointing down. That is the classic reversion fix: it fades in quiet, turning tape, never into a live trend.

The honest verdict: this is a genuine chop-reversion idea with one real edge trade and a lot of knife-catching around it. Its +\$723 over the 25-day archive is outlier-carried — take out its single best day and it's **-\$209**. So it is not a gate to arm; it is a lead to sharpen: add the slope veto, test it on a genuine chop day where dips actually reverse, and reprice it honestly before it earns a cent of size.

WATCH • ADD A TREND VETO • TEST ON A CHOP DAY • NEVER LONG-ONLY-LIVE IN A DOWN TAPE

§34 **rg_long_20_t20 — a long-only grind gate that only won when the tape was calm**

This is one of the V7 "rising-grind, long-only" sims, running a 20-bar grind trigger with a t20 target. It had a losing week: it lost **\$203.50** across 9 filled trades and only won 2 of them (22%). That is thin — nine trades is a single week of noise, not a verdict — but the shape of the nine is clean enough to say one useful thing.

The money, plainly: rg_long_20_t20 took 9 long MNQ trades this week, made money on 2 and lost on 7, and finished **\$203.50** in the red — its two winners (+\$41.50 and +\$47.00) added \$88.50 while the seven losers bled \$292. It actually came in a hair better than its own modelled ceiling (-\$203.50 real vs -\$248.00 ceiling), so this is not a fill-quality problem — the entries themselves just didn't work this week.

Question 1 — confirm the money

Scored on `real_pnl WHERE fill_status='filled'` from the V7 board, it is a clean loss: nine long entries, two winners, seven losers, **-\$203.50** for the week at a 22% hit rate. There is no short side to look at — this sim is long-only by construction (it is the rg_long family), so every one of the nine trades was a long into a rising MNQ tape. The whole "which side pays" question collapses to: the long side paid -\$203.50, and that's the entire book.

Question 2 — why it bled

Three cuts tell the story: the exit it died on, the ATR it entered at, and the hour it fired.

By exit reason — the stop is where the money went. Seven of the nine trades exited on `STOP` and those seven netted -\$176 (one of them a trailing stop that actually locked +\$47, the other six all red). One trade got yanked out on an `ABSORPTION_CUT` for -\$69 — the single worst trade of the week. Only one trade ever reached its `TARGET`, for +\$41.50. So the picture is: it hits its target almost never (1 in 9), and when it's wrong the stop takes a full -\$25 to -\$53 bite each time.

TARGET (1t)		+\$41.5
STOP (7t)		-\$176
ABSORPTION_CUT (1t)		-\$69

EXIT REASON	TRADES	REAL P&L	WIN%
STOP	7	-\$176.0	14%
ABSORPTION_CUT	1	-\$69.0	0%
TARGET	1	+\$41.5	100%

By entry ATR (the axis that actually separates) — this is the clean line. Both winners came in on the two calmest entries of the week: atr_pct 0.021 (+\$41.50) and atr_pct 0.042 (+\$47.00). Every single entry above atr_pct 0.045 lost — all six of them. The nastiest was the loudest: the atr_pct 0.117 entry (nearly 3x the median volatility) is the one that got absorption-cut for -\$69. In plain words: when it bought a quiet, orderly rise it held and won; when it bought into a wide, fast, volatile push it got stopped or cut out every time.

By session hour — spread too thin to trust. The two winners landed at 20:00 and 11:00 UTC; the losers are scattered across 01, 06, 07, 15, 17, 19 and 22. With one or two trades an hour there is no hour pattern here worth naming — the hour column is noise, and I won't pretend it's a gate.

Question 3 — the one thing for Monday

The single condition that best separates winners from losers is **entry ATR**, not side (there's only one side) and not hour (too sparse). Split the book at atr_pct 0.045:

BUCKET	TRADES	REAL P&L	WIN%
atr_pct ≤ 0.045 (calm tape)	3	+\$65.0	67%
atr_pct > 0.045 (loud tape)	6	-\$268.5	0%

The gate, stated plainly: only take the rg_long_20 entry when atr_pct is at or below ~0.045; stand down on anything louder. In-sample this week that turns a -\$203.50 week into a +\$65 week — a swing of about **+\$268.50** — and drops the six 0%-win loud-tape trades entirely. The honest catch: it's a 3-trade "good" bucket. That is not enough to call it an edge. It rhymes with the amplitude-floor thesis from the momentum family (thin tape / calm entries hold; loud pushes get faded and cut), which is the only reason I'd even flag it rather than bin it outright.

Verdict. **RETIRE AS-IS** — a 22%-win, -\$203.50 long-only week with one target-hit in nine is not a strategy that's working, and the calm-tape rescue rests on three trades. Don't run it live. The one thing worth carrying forward is the observation itself: this gate only ever paid on quiet entries and lost 6-for-6 on loud ones, so if it earns another look, it comes back with an atr_pct ≤ 0.045 floor bolted on from the start — never bare. As it stands it's a duplicate-ish member of a long-only grind family that the 25-day archive already showed bleeds standalone; this week did nothing to argue otherwise.

§35 rg_long_20_t15 — a long-only grind that gave it all back on the wild bars

This one is a long-only reversion-grind on MNQ. Over its honest first week it took 11 trades and lost **-\$75.5**, winning 45% of them. That is not a blow-up — it is a hair under breakeven — but every dollar of the loss (and then some) came from a handful of trades it took when the tape was too hot for it. Cool the entry down and it is flat-to-green; leave it as-is and it slowly bleeds.

1. The money, plainly

rg_long_20_t15 fired 11 filled trades between 07-15 and 07-17, all of them LONG and all of them MNQ, and finished the week down **-\$75.5** on real fills with a 45% hit rate (5 winners, 6 losers). There is no short book to blame and no second symbol — this is one clean long grind on the Nasdaq micro, and it lost a little money. It also happens to be the exact duplicate of rg_long_20_flow25 on this board, so treat the two as one sim, not two data points.

2. Why it bled — it wins small and clean, loses big and hot

The exit mix tells the whole story in three lines. When it reaches its **TARGET** it is perfect: 5 trades, 5 wins, +\$184.5 banked. Its losses are all stops and one absorption cut: 5 **STOP** exits for -\$191 (zero winners among them) and a single **ABSORPTION_CUT** for -\$69. So the shape is a coin-flip that pays out symmetrically on the wins but the losses are a touch fatter — a -\$75 week is what you get when your 5 stops (-\$191) plus one cut (-\$69) just outweigh your 5 clean targets (+\$184.5).

TARGET (5)		+\$184.5
STOP (5)		-\$191
ABSORPTION_CUT (1)		-\$69

Now the reason those stops fire. Bucket the 11 trades by how volatile the bar was at entry ($\text{atr_pct} = \text{entry ATR} / \text{entry price}$) and the split is stark. On the **quiet** bars, atr_pct under 0.04, it went +\$14.5 over 3 trades at a 67% win rate. In the **middle** band, 0.04 to 0.06, it was roughly flat at -\$9 over 4 trades. On the **hot** bars, atr_pct at or above 0.06, it lost -\$81 over 4 trades and won just 25% of them. The single worst trade on the whole sheet is the 0.117 atr_pct entry — more than double any other bar's volatility — which got absorption-cut for -\$69. A long-only grind is a mean-reversion idea: it wants to

buy a dip in an orderly tape and get paid as it grinds back. When the bar is wild, there is no orderly grind to fade — the thing keeps running and stops it out.

Session hour points the same way but is too thin to lean on: the three trades it took in the 20:00 UTC hour were all clean targets for +\$107 combined, while the scattered singles through the overnight and midday (01:00 -\$53.5, 15:00 -\$69, 17:00 -\$38, 19:00 -\$43) are where it stopped out. That is really just the ATR story wearing a clock — the calm 20:00 window happened to hold its quiet bars, and the messy hours held its hot ones. I would not build a time gate off one trade per hour.

In one line: it makes money grinding calm dips (quiet-ATR entries hit target 67% of the time) and hands it back trying to grind the wild ones (hot-ATR entries lose three-to-one).

3. The one thing for Monday — cap the entry ATR

The cleanest single split is a volatility ceiling on the entry. Skip any signal where `atr_pct` is 0.06 or higher and keep the rest:

BUCKET	N	REAL_PNL	WIN%
Keep — <code>atr_pct < 0.06</code>	7	+\$5.5	57%
Veto — <code>atr_pct ≥ 0.06</code>	4	-\$81.0	25%

That one condition is worth an **+\$81 in-sample swing**: it turns the -\$75.5 week into a +\$5.5 week and lifts the hit rate from 45% to 57%. The gate is plain — `rg_long` only takes calm bars; if the ATR is running (`atr_pct ≥ 0.06`) it stands down. It fits the mechanism, not just the numbers: a dip-grind needs a tape that grinds, and it kills the -\$69 absorption-cut outlier that was the biggest single hole in the sheet.

Be honest about what this is, though. The "good" bucket is +\$5.5 — it is breakeven, not an edge. The gate is really just "trade less, and only when it is quiet," and it is drawn on 4 vetoed trades over three days, which is far too thin to promote on. It rhymes with the amplitude-floor work we already know (thin, wild tape = stand down), so it is at least a familiar shape rather than a fresh curve-fit. But nothing here says `rg_long_20_t15` has a live edge — it says the sim is close to flat and its losses are concentrated on hot bars.

WATCH — KEEP IT ON THE BOARD WITH AN ATR_PCT < 0.06 ENTRY CEILING, BUT IT NEEDS ANOTHER WEEK OR TWO OF QUIET-BAR TRADES BEFORE IT EARNS ANYTHING MORE THAN OBSERVE-ONLY. AS IT STANDS IT IS A SLOW LEAK, NOT A PROMOTION CANDIDATE.

§36 rg_long_20_flow25: a calm-tape long that got dragged under water every time the ATR woke up

This one bled a small, honest \$-75.5 across 11 trades this week, winning 45% of them. It's not a disaster and it's not a keeper — it's a long-only regime buyer whose whole result came down to one thing: how jumpy the tape was when it entered.

The money in one sentence: it took five targets worth **+\$184.5**, gave back five stops worth **-\$191.0**, and then one absorption cut lopped off another **-\$69.0** — so the stops and that single ugly cut just barely outweighed the wins, and it closed the week red by **-\$75.5**.

Every trade was a long, so the story isn't the side — it's the exit

There's nothing to learn from splitting long versus short here: this is a regime-long strategy and all 11 fills were longs. So the honest place to look is what closed each trade. And when you line the exits up, the picture is stark — the targets did all the good work, the stops did all the damage, and one absorption cut was the single worst trade of the week.

TARGET ×5		+\$184.5
STOP ×5		-\$191.0
ABSORB ×1		-\$69.0

The targets and the stops nearly cancel — five wins at **+\$184.5**, five stops at **-\$191.0**. That's a coin-flip strategy on its own. What actually sank it was the one absorption cut at **-\$69.0**: a long that the book stalled and the strategy bailed on before it could hit either the target or the stop. Take that single trade away and the week is dead flat at about -\$6.5. So the coin-flip isn't the problem — the problem is the one trade that never should have been on.

What it does when it wins vs when it loses: it wins in a quiet tape and drowns in a busy one

Sort every trade by how wide the ATR was at entry (that's entry_atr divided by entry price), and the whole result falls cleanly into two piles. When it bought into a calm tape — ATR under 0.05% of price — it won 4 of 5 and made money. When it bought into a livelier tape — ATR at 0.05% or wider — it won 1 of 6 and got carved up. Every one of the mid-range entries (0.050% to 0.069% ATR) was a stop-out, four in a row, and the widest entry of the lot (0.117% ATR) was that -\$69 absorption cut.

ENTRY CONDITION	TRADES	WIN%	REAL P&L
Calm tape — ATR < 0.05% of price	5	80%	+\$92.0
Active tape — ATR ≥ 0.05% of price	6	17%	-\$167.5

That's the entire strategy in two rows. The calm-tape longs made **+\$92.0** at an 80% hit rate; the busy-tape longs lost **-\$167.5** at a 17% hit rate. It's a slow-drift buyer. When the market is grinding quietly it gets carried to its target; the moment volatility expands, its stop is too close to survive the noise and it gets shaken out or absorbed before the move it wanted ever shows up.

One honest note on the numbers: the \$ figures are tick-repriced honest fills, so if you eyeball the entry and exit points and multiply by the MNQ multiplier you'll get slightly different arithmetic than these dollars — the fills ran a touch off the clean bar marks, as they always do live. The direction of the split is what matters, and it's unambiguous.

The Monday move: put an ATR ceiling on it — or just leave it in the shadows

The single condition that separates its winners from its losers is the ATR at entry. Gate it one line: only take the long when entry ATR is under 0.05% of price; stand down when the tape is wider than that. In sample that turns a **-\$75.5** week into a **+\$92.0** one — a swing of about **+\$167.5** — and it does it by simply not taking the six busy-tape trades that lost 5 of 6.

But be honest about the size of this: it's 11 trades, split 5-against-6. That's far too thin to arm on its own, and the whole week's bleed was a rounding-error **-\$75.5**. The ATR-floor finding is real and it lines up with everything else we know — this is a quiet-drift long and quiet-drift longs die when vol expands —

but one week can't carry an arm decision. Leave it observing, apply the calm-tape ATR ceiling in the shadow, and see if the 80% calm-bucket holds up over another few weeks before it earns a nickel of live size.

WATCH RG_LONG_20_FLOW25 – GATE TO CALM TAPE (ATR < 0.05%), KEEP IT IN THE SHADOWS UNTIL THE ATR FLOOR PROVES OUT ON MORE THAN 11 TRADES

§37 The small fry: sims too thin to judge, and two housekeeping catches

Every board has a back row — sims that fired a handful of times, or didn't fire at all, or turned out to be the same idea wearing two name tags. None of these move the needle this week, but they're worth a clean paragraph each so nothing rots quietly in the corner. Two of them are actual bugs, not verdicts.

The rule for this row is simple: with one trade, or three, there is no edge to find yet — there's just a coin that landed. I'm reporting them so you know they're alive and behaving, not because any of them earned a promotion. Where a thin sim is really a slice of a family we already judged, I say so, because a small sample sitting on top of a bad idea is still a bad idea.

rg_short_025_raw — one trade, one win, no verdict

This one took a single short on MNQ at 06:32 on the 17th, at an ATR of about 0.044% of price, and it hit its target for **\$48.5**. That's the whole story. One trade, one winner, 100% win rate — which sounds great until you remember 100% of one is exactly what a fair coin does half the time. I ran the join on `shadow_trades` against `shadow_real` filtered to filled, and there is genuinely just the one row. There is nothing here to keep or kill; it needs weeks of firing before its short-side signal means anything. Worth noting the rest of the raw-short and short-gate family has been quietly the healthier side of the book elsewhere (thrust shorts carried `thrust_loose` to +\$500 while its longs lost \$38), so a short-only reversion probe is a reasonable thing to have running — just leave it running and check back when it has 30 trades, not one.

capit_mid — three trades, two stops, one bounce, still red

The middle-tightness capitulation sim fired three times, all longs, and netted **-\$39** on a 33% win rate. The tape underneath it: a stop for -\$32.5 at 22:12 on the 16th, another stop for -\$28.5 at 02:17, then one target that paid +\$22 at 02:54. So it bleeds two full stops to claw back most of one on the winner — the classic shape of a mean-reversion idea catching a falling knife twice for every clean bounce.

On its own, three trades is a shrug. But `capit_mid` is not really alone — it's the mid setting of the capitulation family, and the family verdict is already ugly. The loose sibling, `capit_loose`, took 22 trades and lost **-\$622.5** at a 9% win rate — the single worst sim on the whole V7 board this week, a wall of STOP exits with barely a target in sight. The ride variant went 0-for-2 for -\$61. So every dial position of this idea is red right now, and `capit_mid`'s tiny -\$39 is just the least-punished slice because it barely traded. I wouldn't read its smallness as safety. The honest call is that capitulation-buying as a family is fading rips that keep running, and the thin sims don't rescue it.



The bottom two bars in that chart are the same result twice — which is the first of the two housekeeping catches.

Catch one: `rg_long_20_t15` and `rg_long_20_flow25` are the same sim

Both of these ran 11 trades and both lost exactly **-\$75.5**. That's not a coincidence — I joined the two sims on entry timestamp and side and walked all 11 trades: every single fill is identical to the cent (+17.5 / -23.5 / -53.5 / +41.5 / -33.0 / +36.0 / -69.0 / -38.0 / -43.0 / +69.0 / +20.5). One is supposed to gate on a t15 trigger and the other on a flow25 filter, but this week they produced a byte-identical tape. Either the two config knobs collapsed to the same condition, or one is silently reading the other's settings. It's not costing money — it's costing a board slot and it's double-counting a bleeder into the reversion column. Point one gate at a genuinely different condition or retire one of them.

Catch two: `cb_thrust` and `cb_thrust_dropped` fired zero trades

The two circuit-breaker thrust sims logged **no rows at all** — not zero filled, zero trades. I checked the raw `shadow_trades` table with no fill filter and both come back with a count of 0. They're on the board but they never armed. That's a dormant-sim bug, not a "no signal this week" result: a thrust variant should absolutely have had setups in a week where `thrust_cont` took 29 trades and `thrust_loose` took 43. Something in the circuit-breaker gate is refusing everything, or the sims aren't wired into the runner. Until that's fixed they tell us nothing, and worse, a dead sim looks calm on a dashboard while contributing nothing. Fix the wiring, confirm they take trades, then let them observe.

The roundup

SIM	TRADES	NET	READ
<code>rg_short_025_raw</code>	1	+\$48.5	WATCH one lucky short, no verdict
<code>capit_mid</code>	3	-\$39	WATCH thin slice of a red family
<code>rg_long_20_t15</code>	11	-\$75.5	DEDUPE identical to flow25
<code>rg_long_20_flow25</code>	11	-\$75.5	DEDUPE identical to t15
<code>cb_thrust</code>	0	—	FIX dormant, never armed
<code>cb_thrust_dropped</code>	0	—	FIX dormant, never armed

None of these change the week's story. Two are too young to judge and belong on the watch list until they've traded 30-odd times; two are the same sim double-logged and one should be retired or repointed; two are asleep and need their wiring checked before they mean anything. The one genuine warning is `capit_mid` — don't let its small red number hide that its whole family (loose -\$622.5 at 9%, ride 0-for-2) is bleeding on the same idea of buying capitulation into runs that don't turn.

§38 **Eight hunches, walked to a verdict: the exit tweaks are sideshows, direction-by-regime is the only lever that held**

Every idea I floated to you this week, tested against the money and walked all the way to a yes or a no. Nothing got quietly dropped and nothing got armed on the strength of one good week. Eight leads, eight verdicts. Some held, most of the exciting ones didn't, and I'll tell you plainly which is which.

A word on how I scored these before we start. Every dollar below is `real_pnl` on trades that actually filled — the honest number, not the ceiling that flatters losers by anywhere from 1.1 to 3.1 times. And where a musing lived or died on more than this week, I pulled the two weeks of V5 history and the 25-day archive replay so you're not looking at a single lucky Tuesday dressed up as an edge.

1. The chandelier is too loose — HELD, and the fix is ATR-conditional

This was the week's cleanest result. I ran three chandeliers side by side at trail multiples of 20, 25 and 35, same entries, only the trail differs. Tighter won at every step, a perfectly straight line:

THE K20 TRAIL MADE \$18.5 OVER 11 TRADES, K25 LOST \$36.5, AND THE LOOSEST K35 LOST \$101 OVER 9 .

Every notch you loosen the trail, you give back more. And the same shape shows from the live desk the other way round — the CHANDELIER exit was the only green exit on the live board, +\$567 across 13 fills at an 85% hit rate, while the plain STOP lost \$810 and never won once. It confirms what the give-back study kept saying: at any fixed width the runner hands back roughly a third of its peak, and the loose trail just lets it hand back more before the exit fires.

But here's the honest catch, and it's why the verdict is "held" and not "armed": even the tightest of the three barely scraped positive on 11 trades. Flat trailing at any width is a losing game on this sample — tighter just loses less. The real fix isn't picking a number, it's making the trail read the tape: a graded, ATR-conditional trail that rides wide when the move is amplitude-rich

and clamps down when it isn't. The deep pass measured that ATR-conditional trail beating flat-3.5 by about \$245 on thrust and \$1,653 on grind. So the musing holds — the chandelier is too loose — but the answer is a smarter trail, not a smaller constant.

2. A vol-adaptive chandelier — promising, still out-of-sample-pending

This is the natural next step from the one above, and I want to be careful not to oversell it. The in-sample story is genuinely good: grade the trail off ATR and it beats the flat trail by four figures on grind. But that number is fit on the same weeks it's measured on, and you've taught me the hard way that a single-week filter usually just means "trade less" in a flattering disguise. I have not yet run it forward on unseen days. Until I do, this is a promising lead with a pending OOS test, not a result. I'd rather tell you "I don't know yet" than hand you an in-sample curve and call it an edge.

3. Direction × regime gating — HELD, and it's the number-one lever

This is the big one, and it's the same finding from two completely independent angles, which is why I trust it. The whole board splits on which side a strategy is allowed to take in which regime. Look at the two cleanest cases on the V7 board this week:

SIM	LONG SIDE	SHORT SIDE	READ
grind_fast	+\$747.5 (34t)	+\$552.0 (43t)	both sides green — genuinely two-sided
thrust_loose	-\$38.5 (20t)	+\$500.5 (23t)	all the money is on the short side

thrust_loose made \$462 net this week but every dollar of it — and then some — came from shorts; its longs quietly bled. If you'd gated that sim to shorts-only in the prevailing regime you'd have kept the \$500 and skipped the drag. Now widen the lens to the two weeks of V5 history and the same law holds at scale: **THE MOMENTUM BOOK MADE \$21,792 WHILE THE REVERSION/FADE BOOK LOST \$15,798** — and the MNQ thrust family alone carried \$22,453 of that green while passive_chop bled \$12,528. The edge was never "which strategy" — it was "which direction, in which condition." Right idea in the wrong regime is how the desk bleeds between the big runs. This holds, and it's the lever worth building around.

4. A per-side circuit breaker — built, but DORMANT

I built this to catch the exact failure above: if one side of a sim is bleeding, trip a breaker and stop taking that side. It's on the board as `cb_thrust` and `cb_thrust_dropped`. The problem is simple and a little embarrassing —

BOTH BREAKER SIMS TOOK ZERO TRADES THIS WEEK. The trip conditions never armed, so I have exactly no evidence it works. It's not refuted, it's untested-in-the-wild. Either the thresholds are set too tight to ever fire, or this week just didn't hand it the losing streak it's built to catch. Next week I loosen the arming logic enough to actually see it act before I say a word about whether it helps.

5. grind_fast, live if it survives — REFUTED

This one hurt because `grind_fast` was the week's best sim: **+\$1,299.5** over 77 trades, and beautifully two-sided as we just saw. The tempting story writes itself — promote it. So I ran it back over the 25-day archive to see if it survives out of sample, and it does not:

GRIND_FAST LOST \$7,975 OVER 25 DAYS, GREEN ON ONLY 7 OF THEM. This week's win was a good week, not a durable edge. You can even see the fragility inside this week's own trades — its chandelier exits made \$2,606 while its stops gave back \$1,307, so the whole thing lives or dies on how often price lets the trail ride instead of hitting the stop. On a chop-heavy stretch that ratio flips and it bleeds. Refuted as a standalone live candidate. Its value is as a two-sided signal to fold into a regime gate, not as a strategy to arm.

6. Reversion #23, the chop specialist — outlier-carried, not proven

The reversion-in-chop idea keeps flickering because one variant, `rg_long_fast_v`, showed +\$723 over the archive. But when I strip its single best day out, that \$723 collapses toward -\$209 — the whole edge is one outlier session carrying the average. On this week's board the reversion-grid family was uniformly poor: **rg_long_fast lost \$493.5**, `rg_long_fast_v` lost \$243.5, and the `rg_long_20` variants sat between -\$75 and -\$203. A specialist whose edge disappears the moment you remove its best day isn't a specialist, it's a lucky draw. Not proven. I'm not killing the chop-reversion thesis outright — but it needs a real trend-free week to earn a look, and this wasn't it.

7. A day-type router — REFUTED as an edge


The dream here was a switch: run momentum on trend days, reversion on chop days, best of both worlds. The skeptic pass took it apart. In-sample the router looked like +\$1,974; split out of sample it falls to +\$199, and — this is the killer — **IT LOSES TO JUST ALWAYS RUNNING REVERSION'S BETTER COUSIN**. The router's apparent edge was hindsight about which day you were in, and you can't buy hindsight at the open. Once you have to classify the day in real time, the advantage evaporates. Refuted. The regime read is real (see musing 3), but routing whole strategies on a day-label is not the way to cash it — the gate belongs on direction, per trade, not on a morning coin-flip about the day's character.

8. grind is readable early — cliffy, not validated

The hope was that grind days announce themselves in the first hour so you could lean in early. I couldn't make it stand up. The signal is what I'd call cliffy — it looks separable on the days it works and gives you nothing on the days it doesn't, with no stable threshold in between. Given musing 5 already refuted grind as a standalone, an early-read on a strategy that loses \$8k over 25 days isn't worth chasing yet. Filed as not-validated. If the ATR-conditional trail (musings 1-2) proves out and makes grind's runners pay, this early-read question becomes worth reopening. Until then it's a maybe resting on a no.

The through-line. Two things held: the chandelier is too loose and wants an ATR-conditional trail, and direction-by-regime gating is the real lever — both confirmed from independent angles and across the two-week history, not one week. Everything shinier — grind live, the chop specialist, the day-type router — refuted or outlier-carried the moment I pushed it out of sample. The circuit breaker is simply untested because it never fired. Nothing here is armed. The two that held are the two I'd spend next week hardening.

Where the money actually landed this week across the leads, so you can see the spread in one look — the two survivors sit up top, the refuted grind-as-live is the biggest sim but the archive kills it, and the reversion/capitulation family anchors the bottom:



*grind_fast leads the week but loses \$7,975 over the 25-day archive — this week's top sim is the clearest example of why one good week isn't a verdict.