

GAZBOT V7 — the Friday report

Week ending 2026-07-31 · MNQ-only · PAPER · **REVISION 3**, rebuilt 2026-08-01 · 34 plays

What changed in Revision 3. Rev2 was folded in on 08-01 and **only 2 of its 5 sections actually landed** — the report went out still printing the numbers those sections were written to overturn. Rev3 lands all of them, applies the in-place corrections so a reader going top-to-bottom is never shown a superseded claim as current, and adds the three weekend re-derivations that overturn three of the corrections themselves. Nothing has been deleted: where a claim was overturned, the original is ~~struck through~~ and the correction sits next to it in a red box.

★ THE HONEST HEADLINE — the +\$625 is not what the machine did

The desk booked **+\$625 over 206 trades** this week and that is real money on a real ledger. But **twelve of those trades you closed by hand** (exit_reason = MANUAL_CLAIM), and those twelve booked **+\$1,223**. Take your hands off and the machine — every gate, every exit, running unattended — did **194 fills for −\$598** at 35.1% win, −\$3.08 a trade, profit factor 0.91.

CUT OF THE SAME WEEK	NET \$	N	WIN%	WHAT IT IS
Headline, as booked	+625.0	206	38.8%	what the ledger says
Automated only — hands off	−598.0	194	35.1%	PF 0.91, −\$3.08/trade — the machine on its own
Your twelve hand-claims	+1,223.0	12	100%	12 for 12
Those same twelve, repriced on 250 ms tape under the exit config actually loaded	+360.5	12	—	five of the twelve would have stopped out
A fully hands-off week	−237.5	206	—	automated book + what the machine would have done with your twelve

CUT OF THE SAME WEEK	NET \$	N	WIN%	WHAT IT IS
So your hands were worth	+862.5	—	—	and it is invisible in every — P&L we report — see BUILD #2

Twelve claims at 12-for-12 is not luck. It is a signal the machine is missing, and the honest reading of the week is that **the entries and the automated exits did not pay; your discretion did**. Only one gate — **abs_veto** — is positive on its own machinery. The gate this report originally crowned “the survivor”, exhaustion_short, is **−\$56 over 44 automated fires**; \$372.5 of its +\$316 week was four of your clicks.

★

THE SCORECARD — four headline recommendations went into the weekend, one came out

Every one of these was independently re-derived from source over the weekend by an agent that had not written the original, and was actively trying to break it. Read this before you read anything else in the report — it is the most important thing the document has to tell you.

FRIDAY'S RECOMMENDATION	WHAT IT CLAIMED	RE-DERIVED VERDICT	WHERE
Stop benching momentum in a confirmed TREND_UP — billed as “the highest-value, best-evidenced change in the whole report”	+\$1,615/week, positive in all nine robustness folds	WITHDRAWN The arithmetic reproduces (+\$1,644) and then falls apart three ways. The estimator books a gate into a trade it is already in — fix that and it is +\$767; write the rule so it cannot buy dips in a downtrend and it is +\$158 on n=4. A placebo test over 20,000 draws makes it null under a textbook ADX tag (p=0.56) and worse than random under a 60-minute regression tag (z=−2.56). And 52% of the money is two clusters of momentum longs bought into bounces on the week's biggest DOWN day. A confirmed trend was live for 3.3% of the week — there was nothing to ride.	Part 3.5
grind_long exit → {a_r 0.5, b 1.0}	Go tight: MED-TREND +\$10.3/ trade at 59% win	WITHDRAWN Swept properly — 130 configs, 3,064 tick-repriced entries, 16 sessions, on a	Part 3.7

FRIDAY'S RECOMMENDATION	WHAT IT CLAIMED	RE-DERIVED VERDICT	WHERE
		sequential one-position-per-slot sim — the cell is −\$1,127 worse than what is loaded today and fails every robustness test. The direction is the opposite of what was recommended: tighter is worse. And a methodology bug fell out of it: the Friday exit lab's paired method stacks overlapping entries and scored the same config +\$2,961 where sequential scores it −\$2,115 — a \$5,076 swing from method alone, which puts two other exit cells on this card under suspicion.	
grind_long entry fix — drop the ER floor, ATR 24→12, ext_atr ≤ 2.0	£700/week; −\$284 → +\$2,830	WRONG BASELINE — the change still ships The −\$284 baseline is a simulation of a config that was not running: the rehab replay applied floors the live desk has had switched off since 07-29. The honest before/after is +\$1,552 → +\$2,830, so about 59% of the advertised delta was removing filters that were already de-facto removed. And use ATR 10, not 12 — 9/10/11 is a real plateau, at 12 the blocked book flips positive, and 12 costs \$941 against 10. The £/week is now £0: grind breaks even at a 1.64× loss multiplier and live losses run 1.1-3.1× modelled.	Part 3.6
Every ER and ATR floor on the desk has been dead since 07-29	A lookup bug: the floors are keyed on base names, the live slots carry _A/ _B suffixes	CONFIRMED — three times, by three independent routes From source, from the journal, and from the live ledger (of 40 grind fills booked after the ER floor shipped, 31 had ER below it). Three agents reached the same bug by three different routes. One correction: the circulating “7,734 suppressions” is a raw log-line count — de-duplicated it is 383 blocked signal episodes over 8 days, 44× smaller. And the important half nobody had done: four of the five floors should never be switched back on, which is why the ship order matters more than the fix.	Part 3.6

WHAT THAT LEAVES STANDING

The fader bench. Of everything the router does, benching the counter-trend faders is the one claim that survives every test anyone could throw at it: **+\$810, n=25**, non-overlap, and identical under every independent

regime tag. Benching into CHOP does not make that grade — it is +\$612 under the router's own tag, +\$187 under ADX and −\$382 under a regression tag, i.e. it only holds if you accept the router's own definition of chop. And the blended “router value” number quoted twice in this report (−\$1,216) is double-counted twice over; merged it is −\$440 to −\$646, it flips sign on which single day you drop, and it should not be quoted at all.

The per-tick stop guard. 26 unfilled stops before, 0 unfilled and 103 clean stops after. Already in production. Do not revert it.

abs_veto. Cold re-validation on 238 trades: +\$2,668 at 47% win, both sides green independently, 9 of 12 days green. It is the best thing on the board. It was never a promotion — it shipped two-sided on 07-25 — and what leaked was its arming policy, not the gate.



WHAT TO ACTUALLY DO

Garrath — this is the whole answer to “so what do I actually DO?”, and it is at the front this time. Every row below mirrors `reports/friday_v7/plays.json`, which was reconciled by hand against every correction in this report. The interactive version, with the full working behind each play, is the [Monday playbook](#). 34 plays, five windows, in the order they need doing.

THE TWO THAT CANNOT WAIT

1 • Un-pin the router (MONDAY #1). All six gates are pinned in `router_tick_durable.py`, so the apply step's change set is permanently empty and the durable tick has benched nothing for 36 hours. The one router claim that survives every robustness test — the fader bench, +\$810 — is switched off. It is a one-line revert and it needs you: my edit was refused by the permission classifier. Deadline is the Sunday 22:00 UTC reopen.

2 · If you ship the floors tonight, ship them in order (SATURDAY #1).

The values go first, while they are provably a no-op; the wiring fix goes last, as the single arming moment. Backwards costs about −\$370/day and presents as “the desk mysteriously stopped trading”.

SATURDAY · TONIGHT — THIS WINDOW SHUTS AT THE SUNDAY 22:00 UTC
REOPEN · 5 PLAYS

Code and config changes that need a restart. Ship them in the order below; the order is the whole trick on the floors.

#	WHAT TO DO	THE NUMBER BEHIND IT	OWNER	REVERT PATH
1	★ Dead ER/ATR floors — CONFIRMED, with the corrected per-gate values and the ship ORDER CONFIRMED from source. Mechanism in one sentence: tournament.step() passes the raw slot tag to er_blocks()/atr_blocks(), but under the dual-slot scale-out slate every tag carries an _A/_B suffix while ER_FLOOR/ATR_FLOOR are keyed on BASE names — the lookup never matches, so every floor silently returns 'don't block'. The _base() helper at tournament.py:86 is applied at lines 250/303/345 but NOT at 139/142. Dead since 2026-07-29 16:20:49 UTC (commit 5b0913b + the scaleout.conf drop-in); last suppression line 07-29T13:29:04Z, zero since. SHIP IN THIS ORDER — the order is the whole trick, because while the floors are dead, changing their VALUES is provably a no-op: 1. deciders.py values first: ER_FLOOR = {} (delete all), ATR_FLOOR = {grind_long: 10.0, capitulation_long: 10.0}. Zero live effect today. Revert: restore two literals. 2. THEN tournament.py:139/142 -> _base(slot). THIS is the arming moment. Verify grind blocks ~10% of fires, capitulation ~43%, and ZERO 'ER gate:' lines. Revert: one line — this is the kill switch for 1+2 together. 3. THEN slot_strategy.py:118 add "ext_hi": 2.0 — separate commit and restart. The only step with immediate unconditional effect (~39% fewer grind fires). Revert: delete the key. 4. Nothing else. The abs_veto_short ER>=0.10 and capitulation ER>=0.25 leads go to SHADOW. Nuclear revert throughout: GAZBOT7_TOURNAMENT_SLATE=tournament. ★ NEVER SHIP STEP 2 BEFORE STEP 1. Wiring the floors up at their CURRENT values takes the four floored gates from 382 signals / + \$3,630 to 103 signals / + \$668 — about −\$370/day, with grind collapsing from 181 fires to FOUR. It would present as 'the desk mysteriously stopped trading.' Why: Verified independently, and it corrects Q2 twice. (1) MAGNITUDE OVERSTATED 44x: '7,734 suppressions' are raw log lines — the check re-fires every ~1s tick. De-duplicated into episodes it is 140 ER + 243 ATR = 383 blocked signal EPISODES over 8 days, i.e. tens per day, not thousands. (2) COLLATERAL Q2 MISSED: the ER_HOLD abs_veto_long shadow study is also dead — observe-only, no P&L impact, but it has recorded nothing for 3 days. Everything else is FINE: router benching, the 55s veto, the exhaustion confirm, veto_counter_regime and exit_overrides.json all work. Note both ER floors were committed ~18h AFTER the cutover (aa59357 07-30 10:23, 9767c67 07-30 10:57) — they have never blocked a single trade in production. PER-GATE VERDICTS. A floor earns its keep only if the population it BLOCKS loses money — net negative, still negative after stripping its 3 best, and negative on the worst leave-one-day-out fold. Blocked book over 8 days, tick-honest: · grind_long ER>=0.35 -> blocked +\$2,923 / 173 sig, strip-3 +\$1,132, LODO +\$10.03 -> DELETE · grind_long ATR>=24 -> blocked +\$982 / 151 sig, strip-3 −\$82, LODO + \$2.92 -> REPLACE WITH 10	383 blocked signal episodes / 8 days (not 7,734 — that was a 44x log-line count). Ship step 2 before step 1 and the four floored gates go 382 signals / +\$3,630 → 103 / + \$668, about −\$370/day.	Garrath — code edit + tournament restart	Step 2 is the kill switch: _base(slot) → slot, one line, neutralises steps 1+2 together. Nuclear: GAZBOT7_TOURNAMENT_SLATE=tournament.

#	WHAT TO DO	THE NUMBER BEHIND IT	OWNER	REVERT PATH
	· abs_veto_long ER>=0.25 -> blocked +\$640 / 52 sig, strip-3 +\$111, LODO + \$4.88 -> DELETE · abs_veto_short ATR>=16 -> blocked +\$137 / 28 sig, strip-3 -\$197, LODO - \$11.02 -> NOT PROVEN, delete · capitulation_long ATR>=10 -> blocked -\$459 / 22 sig, strip-3 -\$601, LODO -\$26.62 -> KEEP at 10 Both price bases agree on every verdict. PRIORS TESTED, not inherited: 'ER filters are FAKE' — confirmed emphatically. 'abs_veto wants no ER floor' — confirmed, and its blocked book is ENTIRELY normal-chop (+\$680), the gate's home cell. 'grind wanted 0.35' — REFUTED. 'ATR floors are REAL' — half true: they sharpen per-trade quality monotonically, but total dollars peak at roughly HALF the live values. Method note: the replay is a clone of the desk's own repricer.py and reproduces shadow_real.real_pnl at \$0.00 MAE on 4 of 6 gates (99.82% sign agreement). ★ THIRD INDEPENDENT CONFIRMATION (08-01), found separately while sweeping grind's exit, with LIVE proof rather than replay: of 40 grind_long_A/_B fills booked AFTER the ER>=0.35 floor shipped on 07-30, 31 had ER < 0.35 and 22 had ATR < 24. Three agents reached this bug by three different routes. It also adds a threshold warning that cuts against a naive fix: on this tape the sub-0.20-ER fills are the PROFITABLE population, so wiring the floors up without re-deciding the values would bench the money. Evidence: REV2 · Q2 D1 · verification: bug · mode: live			
2	grind_long entry — ATR 24 → 10 (NOT 12), ext_atr <= 2.0 ceiling. The £700 was wrong. SHIP, but with two corrections to what the report told you. (1) THE HEADLINE NUMBER IS WRONG. The rehab replay APPLIED the floors as config (rehab_grind_long.py:124, :234), so its '11 fires, - \$284' baseline is a floors-WORKING counterfactual — but the floors have been dead on the live desk since 07-29. The honest before/after is +\$1,552 -> +\$2,830, not -\$284 -> +\$2,830. 59% of the advertised delta is removing filters that were already de-facto removed. part1_5_rehab.html and rehab_grind_long.md still carry the uncorrected version. (2) ATR 12 IS ONE NOTCH THE WRONG SIDE OF A CLIFF. Use 10. grind ATR 9/10/11 is a genuine plateau; at 12 the blocked book flips positive, so the report's recommended 12 costs \$941 versus 10. The re-derived config (ATR>=10 + ext_hi 2.0) beats the report's on net, per-signal, win%, strip-3, LODO and BOTH OOS halves. This ships as steps 1 and 3 of the dead-floors sequence above — do not ship it standalone. Why: Why the £/wk figure on this card is now £0 rather than £700: not because the change is bad, but because no number here is bankable. grind_long is a HIGH-COUNT, LOW-MARGIN book being asked to survive real fills. Even the best config breaks even at a 1.64x loss multiplier, and the desk's own standing caveat ([[shadow-sim-understates-losses]]) puts live losses at 1.1-3.1x modelled. That breakeven sits INSIDE the band — at 1.7x the entire +\$3,095 is zero. The ext_atr ceiling improves resilience (at 1.5x losses: capped +\$358 vs uncapped -\$1,449) but does not lift grind out of the band; it is being made LESS BAD, not good. Compounding it: the replay is always-armed while the router benched 465 episodes in the same window, and the floor x benching INTERACTION is unmeasured — they could stack and starve the gate. Watch the combined block rate on the first session. Everything else about the fix stands: it survives strip-the-best-3, survives dropping any single day, and holds on the V5 summer-chop archive (-\$1,135 -> +\$172). Treat the POLICY as the edge, never the headline dollar. Evidence: Part 1.5 §2 (read WITH Q2 D1) · verification: verified · mode: live	Honest before/after +\$1,552 → + \$2,830 (n=181 → 111) — not - \$284 → +\$2,830. ATR 10, not 12: 12 is one notch the wrong side of a cliff and costs \$941.	Garrath — ships as steps 1 and 3 of the floors sequence	Restore the two literals in deciders.py; delete the ext_hi key from slot_strategy.py.
3	capitulation_long — require_flip=False + target 1.0R £250/wk Set require_flip=False and target 1.0R. This reverses our own 07-25 prior — say so out loud. Keep the router's US-session and down-trend benches; they are load-bearing, not optional. Why: 11 trades +\$24 at 45% win → 32 trades +\$434 at 78% win; overnight-only +\$422 at 81%. Every single-day-drop stays strongly green, strip-best-3 + \$327. The mechanism is legible: 78% of bounces reach 1R versus roughly 40% reaching 2R — so the 1.0R target harvests the part of the distribution that actually exists. Evidence: Part 1.5 §4 · verification: verified · mode: live	11 trades +\$24 @ 45% win → 32 trades +\$434 @ 78%. Overnight-only +\$422 @ 81%. Strip-best-3 still +\$327.	Garrath — gate config + restart	Restore require_flip=True and the 2.0R target.
4	Fader-bench ER floor 0.20 → 0.15 in the direction-router £60/wk Lower the fader-bench ER-trend floor from 0.20 to 0.15 so the router calls 'trend' sooner and counter-trend faders get benched sooner. NET-min stays at 30; the timing knobs stay as shipped. NOTE: this is a different threshold from grind's ER arming floor of 0.35, which stays —	ER0.20 → ER0.15 at NET30: +\$176 → +\$492 (live-managed set), + \$995 → +\$1,307 (full historical set). Beats 0.20 in every NET column on both universes.	Garrath — one threshold in the direction-router	Set the fader-bench ER-trend floor back to 0.20.

#	WHAT TO DO	THE NUMBER BEHIND IT	OWNER	REVERT PATH
	one is 'should this signal fire', the other is 'should this gate be switched on at all'. Why: Shipped ER0.20/NET30: +\$176 / +\$995 → ER0.15/NET30: +\$492 / +\$1,307 (live-managed set / full historical set). The ER0.15 row beats ER0.20 in every NET column on BOTH universes, which is the two-sets-must-agree filter this desk requires before touching a threshold. Evidence: Part 2.6 §4 · verification: verified · mode: live			

5	abs_veto roster — CONFIRM, no edit. There is nothing to promote. NO ACTION on the roster. 'Promote abs_veto_55s two-sided' was never an action — abs_veto shipped two-sided on 2026-07-25 (deciders.py VETO_GATES holds abs_veto_long and abs_veto_short; the 55s veto fires in tournament.run()). Both are armed in gate_switches.env right now, verified today. Part 2 re-validated a live gate cold and mislabelled the result as a promotion. Sizing stays base_size=1. Why: 7 days live, 85 lots, 65 signals, +\$418 combined this week. The cold re-validation is still worth having and it is strong: 238 trades, +\$2,668, 47% win; LONG +\$1,187 / SHORT +\$1,482, both green independently; 9 of 12 days green; strip-best-3 still +\$620; green in chop AND trend; both walk-forward halves; 55s beats 50s beats 60s. The engine is real — the 55s veto against its no-veto control on the same tape is +\$2,668 on 238 fires (+\$11.2/trade) versus thrust_loose +\$130.5 on 424 fires (+\$0.3/trade). Nobody should read the corrections elsewhere on this card as 'abs_veto is broken'. It is the best thing on the board; its ARMING POLICY is what leaked. Evidence: REV2 · Q2 D0 (corrects Part 2) · verification: verified · mode: live	No edit. Cold re-validation: 238 trades, +\$2,668, 47% win; LONG +\$1,187 / SHORT +\$1,482, both green independently; 9 of 12 days green.	Nobody — this is a confirmation	N/A — nothing changes.
---	--	---	--	-------------------------------

MONDAY · AT THE DESK — SWITCH-FILE AND EXIT-CELL CHANGES, ALL REVERSIBLE · 9 PLAYS

Nothing here needs a code deploy except play #1, which is a one-line revert. Everything else is gate_switches.env or exit_overrides.json and comes back in five minutes.

#	WHAT TO DO	THE NUMBER BEHIND IT	OWNER	REVERT PATH
1	★ ★ UN-PIN THE ROUTER — it has applied nothing for 36 hours Revert scripts/router_tick_durable.py to PINNED = frozenset({"exhaustion_short"}) — the one-line change the 07-31 US-open override's own comment prescribes ("revert once the open settles"). The override pinned ALL SIX gates; the apply step filters on `g not in PINNED`, so a full-roster pin makes the change set permanently EMPTY. Keeps your fade-scalp trial pinned. BLOCKED: my edit was refused by the permission classifier — this one needs you. Deadline is the Sunday 22:00 UTC reopen. Why: 411 ticks across 07-31 and 08-01 logged 'no change' and applied NOTHING; the last real switch change was 2026-07-30T22:15. 'no change' reads as healthy, and that is exactly what hid it — the same failure mode as [[pipeline-health-monitor-the-mechanism]]: sweep the mechanism, not the outcomes. Consequence already paid: grind (the churner) was armed into Friday's violent open for −\$370 with the untradeable meter unable to bench it, and the whole 'arm veto-gates not churners' lesson was unenforceable. The independent TREND_UP re-derivation reached this same conclusion by a different route and put it more sharply: the FADER BENCH is the only router claim that survived every robustness test (+\$810 non-overlap, n=25, identical under every independent regime tag) — and it is switched OFF right now. Q4 wanted to arm momentum in more conditions; the desk is already maximally armed by accident. Un-pinning points the CHEAP direction (it restores benching) and needs no new evidence. Also add an alarm so an all-pinned roster can never log as healthy again. Evidence: Live desk audit · corroborated by REV3 TREND_UP re-derivation · verification: bug · mode: live	411 ticks across 07-31 and 08-01 logged “no change” and applied nothing. The one router claim that survives every robustness test — the fader bench, +\$810, n=25 — is switched off right now.	★ Garrath — my edit was refused by the permission classifier	Restore the all-six pin. One line either way.
2	Do not pre-set a gate list — let the untradeable meter set it Run untradeable.py at the 13:00 UTC pre-open and again hourly. Score <40 → all six armed; 40-64 CAUTION → veto-gates only (abs_veto L+S), faders and grind benched; >=65 → fully flat. There is deliberately no fixed Monday gate list. Owner: the durable router tick; one line in gate_switches.env, reversible every 5 minutes. Why: The meter earned the right to pick the roster: it scored 13 and 23 on the two green days, 47 and 56 on both CAUTION days (both lost), and 86 on the −\$905 day. Friday's enable-all proved the discriminator is the gate's PROTECTION, not the clock — same 172-point open roundtrip: abs_veto (55s-	The meter scored 13 and 23 on the two green days, 47 and 56 on both CAUTION days (both lost), and 86 on the −\$905 day. n=5 — a promising start, not a validation.	The durable router tick	One line in gate_switches.env; reversible every 5 minutes.

#	WHAT TO DO	THE NUMBER BEHIND IT	OWNER	REVERT PATH
	veto protected) +\$146, grind (no veto) –\$370 with 4 stops in 90 seconds, exhaustion_short (off-regime) –\$72. The counter-lesson still binds: on a cleanly-trending day you do not bench the trend-rider on open-noise (that cost about –\$430 on 07-30). The meter is exactly what tells those two days apart, which is why this is a rule and not a list. Hold the 65 cutoff and the 40/40/20 blend — validated, unchanged. Evidence: REV2 · Q0 M1 · Part 2.6 \$5 · verification: verified · mode: live			
3	Bench rgv_short — it is armed live while the report verdicts it SHADOW Set rgv_short=off in gate_switches.env. It is armed live right now (verified in the file today) and every section of the report verdicts it SHADOW. This fixes the one live self-inconsistency the report found in its own desk. The shadow keeps running — this is a bench, NOT a retirement. Why: 11 fires, –\$142, negative at every single exit rung (1.5R –\$32 through chandelier –\$190), strip-best-1 –\$242, red in every cell of the exit lab. On the automated-only ledger it is –\$154.5 on 3 fills. But retiring it on 11 fires would be as unjustified as arming it — bench and shadow, re-look when n grows. Explicitly NOT a retirement call. Evidence: REV2 · Q0 M2 · Part 1.5 \$5 · verification: verified · mode: live	11 fires, –\$142, negative at every exit rung (1.5R –\$32 → chandelier –\$190). Strip-best-1 –\$242. Automated-only ledger: –\$154.5 on 3 fills.	The durable router tick (or you, in the switch file)	rgv_short=on. It is a bench, not a retirement — the shadow keeps running.
4	exhaustion_short — re-cut the exit to A 0.75R / B k1.5, and keep it out of violent whipsaw £420/wk SUPERSEDES the Q0 M4 row that said keep the pinned 0.5/1.5. The Friday override was the right instinct at the wrong numbers: the tick sweep says Lot A 0.75R + Lot B k1.5 chandelier. Just as important — the sweep says the exit CANNOT rescue a violent-whipsaw entry at any R, so the router bench matters more than the R does. Keep the ATR-22 ceiling on the trial. Edit exit_overrides.json + restart the tournament; delete the key to revert. Why: Q1 calls this the single highest-value change in its section: worth roughly +\$420 on the twelve hand-claimed trades alone, and it converts the hand-claim habit into a mechanical rule. Sweep is 64 entry events over nine days on 250ms ticks. The honest conflict you should see: Q0 M4 argued the pinned 0.5/1.5 trial is only 3 legs old (–\$99.5) and that killing a trial at n=3 buys nothing. Q1 is the later and deeper read — computed on the automated-only ledger, which is the right tape for a mechanical rule — so it supersedes. Note also that Q1 demolishes the 'exhaustion_short is the survivor' headline: on automated fills it is –\$56.0 over 44 fills, i.e. FLAT not green. Its celebrated Thursday '+\$372.5 four-in-a-row' was four hand-claims, and three of those four sat on positions the tape later ran 49-98 points through the stop. Modelling weakness stated: 17 of the 64 replays never hit stop/target/trail inside the 120-min cap and are marked at the last tick (the same 17 in every config, so it does not tilt the ranking). ★ METHOD CAVEAT added 08-01: this cell came from a per-signal sweep on overlapping entries — the same accounting that was just shown to swing grind's verdict by \$5,076. Not refuted, but re-run it under the sequential one-position-per-slot sim before trusting the size of the gain. Cheap to revert either way. Evidence: REV2 · Q1 (SUPERSEDES Q0 M4) · verification: verified · mode: live	64-entry 250 ms sweep: Lot A 0.75R +\$510 (the loaded 0.5R is the worst rung at +\$306); Lot B k1.5 chandelier +\$502, the highest leave-one-day-out floor in the sweep at +\$330.	Garrath — exit_overrides.json + tournament restart	Delete the key. Method caveat: graded per-signal on overlapping entries — the same accounting that swung grind's verdict by \$5,076. Re-run it sequentially before trusting the size.
5	abs_veto_long exit → {a_r 1.0, b 1.5}; leave abs_veto_short alone SUPERSEDES the Q0 M6 row that said 'change nothing on abs_veto, the exit is not the lever'. Change exactly one cell: exit_overrides.json → abs_veto_long: {a_r 1.0, b 1.5}. Leave abs_veto_short at 1.5/2.5 — on the short side the R choice is genuinely noise. Edit + tournament restart; delete the key to revert to the built-in default. Why: The live LONG cell A1.5/B2.5 is the WORST of the eleven cells swept, on every measure: +\$13.6/signal vs +\$22.7 for A1.0/B1.5, LODO-min +\$4.7 vs +\$16.8, strip-best-3 +\$6.8 vs +\$19.0. The plateau is broad and sits at A≈1.0, B 1.5-2.0 (0.75/1.5, 1.0/1.5 and 1.0/2.0 all within \$4 of each other) — a genuine parameter plateau, not a spike. n=81 signals. On the SHORT side every cell lands between +\$15 and +\$18 per signal and the LODO/strip ordering flips between them, so leave it alone. What is NOT noise on the short side is the violent-whipsaw column: –\$20 to –\$33 per signal at every cell and every R from 0.5 to 6.0 — no exit tuning rescues that segment, only not trading it does (see the violent-whipsaw bench row). Untested caveat: the CLEAN-TREND bucket is n=6 long / n=5 short and is unscored — a wide Lot B is exactly what a trend week would reward, and A1.0/B1.5 caps it. ★ METHOD CAVEAT added 08-01: same per-signal-on-overlapping-entries accounting as the grind cell that was just refuted. The plateau argument (0.75/1.5, 1.0/1.5 and 1.0/2.0 all within \$4) is reassuring because a plateau is harder to fake than a spike — but re-run it sequentially before trusting the magnitude. Evidence: REV2 · Q2 D2 (SUPERSEDES Q0 M6) · verification: verified · mode: live	The live A1.5/B2.5 is the worst of eleven cells swept: +\$13.6/signal vs +\$22.7 for A1.0/B1.5; LODO-min +\$4.7 vs +\$16.8. n=81, broad plateau, two independent sweep agrees.	Garrath — exit_overrides.json + tournament restart	Delete the key → back to the built-in default. Same overlapping-entry method caveat as #4.

#	WHAT TO DO	THE NUMBER BEHIND IT	OWNER	REVERT PATH
6	Router rule — bench abs_veto_short in violent whipsaw Make deliberate what the router did by accident 4 times this week: bench abs_veto_short when ATR >= 19pt AND 30-min ER < 0.25 (violent whipsaw). Reversible every 5 minutes. Why: n=11 shadow fires −\$185.5; live 5 lots −\$257. Negative at all 11 exit cells and at every R from 0.5 to 6.0 — this is the segment no exit tuning can rescue, so the only fix is not trading it. Evidence: REV2 · Q2 D3 · verification: verified · mode: live	n=11 shadow fires −\$185.5; live 5 lots −\$257. Negative at all 11 exit cells and every R from 0.5 to 6.0 — no exit rescues this segment.	The durable router tick	Reversible every 5 minutes.
7	Router rule — STOP benching abs_veto_long in overnight NORMAL-CHOP The standing 'confirmed chop → bench all momentum' rule files abs_veto as momentum. On the evidence it is an overnight CHOP gate. Bench it on VIOLENCE, not on chop. Reversible every 5 minutes. Why: 13 blocked fires worth +\$384. Live and shadow are both green in that cell (+\$383.5 / +\$525). Total cost of the long-side benching this week: −\$397. Note this is the same class of error as the TREND_UP finding in Q4 — the router's bench taxonomy files gates by NAME rather than by measured behaviour. Evidence: REV2 · Q2 D4 · verification: verified · mode: live	13 blocked fires worth +\$384. Live and shadow both green in that cell (+\$383.5 / +\$525). The long-side benching cost −\$397 this week.	The durable router tick	Reversible every 5 minutes.
8	capitulation_long — bench in the US cash session, overnight-only until the Saturday fix ships Bench capitulation_long 13:30-21:00 UTC, overnight-only, until the Saturday config fix lands. It is currently running the exact config Part 1.5 falsified, and its money is overnight anyway. Re-arm any tick. Why: Live config: 11 trades +\$24. Overnight-restricted recommended config: +\$422 at 81% win. Exit-lab cells MED −\$19 / CHOP −\$20 — it does not pay in the cash session. On the automated-only ledger it is −\$139 on 14 fills. Evidence: REV2 · Q0 M3 · Part 1.5 §4 · verification: verified · mode: live	Live config 11 trades +\$24. Overnight-restricted recommended config +\$422 at 81% win. Cash-session exit cells MED −\$19 / CHOP −\$20.	The durable router tick	Re-arm on any tick.
9	grind_long exit — the honest answer is NO exit makes it green Headline first: 130 configs over 3,064 tick-repriced entries across 16 sessions, and the BEST cell on the board still loses \$414. The loaded default (Lot A 2.5R + Lot B wide lock-chandelier, verified today at slot_strategy.py:149 with no grind_long key in exit_overrides.json) scores −\$2,115 over 16 sessions. If you want the marginal improvement, it is {"grind_long": {"a_r": 3.0, "b": "wide"}} — Lot A 2.5R -> 3.0R, Lot B untouched — worth +\$617 vs the default, roughly +\$10/session at the live desk's actual fill rate. That is the cell REV2-Q3 PRE-REGISTERED for 08-03, so writing it costs nothing in test integrity. But treat this as damage control, not an edge. Why: Be clear about how weak this is: A3.0/wide passes only 2 of 6 robustness tests (10/16 sessions, LODO-worst +\$222 pass, plateau pass; but strip-best-3 −\$130 FAIL, bootstrap P=0.82 FAIL, OOS time-split sign −\$76/+ \$693 FAIL). The statistically strongest cell (A3.5/B3.5, −\$414, +\$1,702, passes everything) is a ONE-RUNG SPIKE — 3.0R is \$353 worse and 4.0R is \$709 worse — so it is a fit, not a finding, and was deliberately not recommended. ★ THE REAL LEVER IS ARMING, NOT EXITING — it is 3-5x bigger: benching grind in 'building' tape is worth +\$3,142 and arming only at ER<0.20 is worth +\$4,348. But every arming variant is green in the capture half and red in the archive half, so that is a DIRECTION, not a measurement. This matches the desk's standing position [[exit-architecture-scar-tissue]]: on a bad entry the exit is scar tissue; the real work is entries. Calibration, reported straight: on the 60 real grind_long fills with tick coverage, reconstructed entry ATR is within 1.6% of the ATR implied by real stop distances (n=52), and real STOP fills came in 7% worse than the model (−\$2,002 vs −\$1,871, n=40). Every number here is a no-slippage CEILING and grind is red under all 130 of them. Evidence: REV3 grind exit ladder (3,064 entries × 130 configs) · verification: exploratory · mode: shadow	130 configs × 3,064 tick-repriced entries × 16 sessions, sequential sim. Loaded default −\$2,115; best of all 130 still −\$414. A3.0/wide is +\$617 ≈ +\$10/session live.	Claude — pre-registered cell, graded 2026-08-21	Delete the key. Treat as damage control, not an edge — the real lever is arming, 3-5× bigger.

BUILD · NEXT — NOTHING GOES LIVE THIS WEEK · 8 PLAYS

Instrumentation, shadow slates and the two method fixes that decide whether next Friday's numbers can be trusted at all.

#	WHAT TO DO	THE NUMBER BEHIND IT	OWNER	REVERT PATH
1	★ The Friday exit-lab's paired method overstates. Re-run every exit verdict under a sequential sim. The exit ladder's paired same-entry method STACKS OVERLAPPING ENTRIES — it lets one slot hold several positions at once, which the	★ \$5,076 swing from methodology alone: the paired method scored grind's loaded default +\$2,961, the sequential one-position-per-	Claude — re-run every exit verdict in the report	N/A. Until this lands, treat the exhaustion and abs_veto_long

#	WHAT TO DO	THE NUMBER BEHIND IT	OWNER	REVERT PATH
	desk cannot do. On the same tape it scored grind's loaded default at +\$2,961; a sequential one-position-per-slot sim scores the same config −\$2,115. That is a \$5,076 swing from methodology alone. The sequential sim is the one that matches reality: 72%/28% stop/target and 4.8-min median hold vs the live desk's 67%/27% and 3.5 min. Re-run every exit verdict in the report under the sequential sim before any of them are trusted. Why: This is why the M5 card above had to be withdrawn, and it does not stop at grind. Two other exit recommendations on this playbook were produced by sweeps whose position-accounting I have NOT verified: the exhaustion_short A 0.75R / B k1.5 cell (Q1) and the abs_veto_long {1.0, 1.5} cell (Q2). I am NOT claiming they are wrong — I have no evidence either way, and both are cheap to revert. But they were graded per-SIGNAL on overlapping entries in the same family of sweep, so they inherit the same suspicion, and both cards now say so. Check those two first. Related and already known: [[sims-on-tick-price]] (5-second repricing flatters tight exits — confirmed again here) and [[shadow-sim-understates-losses]] (every reprice is a ceiling). Evidence: REV3 grind exit ladder — method · verification: bug · mode: shadow	slot sim scores the same config – \$2,115.		cells as directions, not measurements.
1	In-code 55s continuation veto for grind £150/wk Give the churner the same 55-second continuation confirm that protects abs_veto — extend VETO_GATES beyond the two thrust gates (src/gazbot7/deciders.py:139). Then backtest it per regime before it is ever eligible for a Saturday. Nothing goes live this week. Why: Same open, same tape: the veto-protected gate made +\$146 while un-vetoed grind lost −\$370. The veto's standalone incremental edge on thrust is +\$2,538 versus un-vetoed thrust_loose. This is the mechanism that separates a veto-gate from a churner, and grind is the desk's worst churner. Evidence: Part 2.6 §6 · Part 2 head-to-head · verification: verified · mode: shadow	Same open, same tape: veto-protected abs_veto +\$146, un-vetoed grind −\$370. The veto's standalone edge on thrust is +\$2,538.	Claude — build + per-regime backtest	Nothing goes live this week; it is not eligible for a Saturday until it is backtested.
2	Instrument the Claim-profit button — your hands were worth +\$862.5 this week and it is invisible Auto-reprice every MANUAL_CLAIM against what its slot's exit would have done, and log the delta. Turn 'Garrath's hands are worth \$X a week and here is the pattern they exploit' into a standing number instead of a Friday archaeology dig. Why: This is the finding behind the week's biggest correction. The +\$625 headline is real money on a real ledger, but \$1,223 of it came from twelve trades closed by hand. Strip them and the machine on its own did 194 fills, −\$598, 35.1% win, −\$3.08/trade, PF 0.91. Repriced on 250ms ticks under the exit config actually loaded at the time, those same twelve entries would have banked +\$360.5 with five of the twelve stopping out — so a fully hands-off week is −\$237.5, and your hands were worth +\$862.5. Twelve hand-claims at 12-for-12 is not luck; it is a signal the machine is missing, and right now it is invisible in every P&L we report. Evidence: REV2 · Q1 · verification: verified · mode: shadow	Your hands were worth +\$862.5 this week. Twelve claims at 12-for-12 is not luck, and it is invisible in every P&L we report.	Claude — auto-reprice every MANUAL_CLAIM against its slot's exit	Observe-only; nothing to revert.
3	BIG-TREND rung — the dated plan to actually prove it Grade date Friday 2026-08-21; hard backstop Friday 2026-09-11. Owner: Claude (durable router + Friday max-depth lab), operator sign-off on the freeze only. Grading n is 39 paired BIG-TREND entries — a computed number, not a round one. Why: Two corrections travel with this. First, the BIG-TREND cell is n=24, not n=5 — Part 2 LAB quoted the live shadow count (5 grind_fast entries in an ER>=0.50 minute over 12 sessions) inside a caveat about the tick-replay number. Second, the report's refusal of the MNQ tick archive was wrong on the facts, and worse, the lab had already used it: the lab tape starts 2026-07-05, not 07-15, and eleven of the twenty-four BIG-TREND entries come from that archive stretch. The refusal reads well and is false. This replaces a hand-wave with a dated, owned acquisition plan. Evidence: REV2 · Q3 · verification: verified · mode: shadow	The cell is n=24, not n=5. The pre-registered grade needs 39 paired entries. Grade date Friday 2026-08-21; hard backstop 2026-09-11.	Claude (durable router + Friday max-depth lab); you sign off the freeze only	N/A — it is a measurement plan.
4	Deploy the concrete-front-month stop fix (still on a branch) Deploy the STOP_UNFILLED fix — stops on the concrete front-month Future rather than the ContFuture (commit 99e3c11, memory [[stop-unfilled-contfuture-root-cause]], still on a branch). Live-verify it. Branch revert. Why: A pure execution leak with no shadow analogue: 6 abs_veto lots this week, −\$273.5, roughly \$168 of which sits inside the long-side 'exit' gap that Q2 was trying to explain. Note this is SEPARATE from the per-tick stop guard already in production — see the HOLD row. Evidence: REV2 · Q2 D5 · verification: bug · mode: shadow	6 abs_veto lots this week, −\$273.5, roughly \$168 of it inside the long-side exit gap. Commit 99e3c11, still on a branch.	Claude to deploy, then live-verify	Branch revert. Distinct from the per-tick stop guard already in production (see HOLD #2).

#	WHAT TO DO	THE NUMBER BEHIND IT	OWNER	REVERT PATH
5	census cluster() z-score fix — the highest-value change in the greenfield section Replace the <code> flow > 50</code> rule with a z-score against a trailing 2 hours, threshold <code> z >= 1</code> . Both <code>scripts/run_census.py:30</code> and the verbatim copy at <code>scripts/run_census_full.py:38</code> . Why: The <code> flow > 50</code> rule fires on 66.5% of ALL bars, which means FLOW-LED and VACUUM are two names for one coin. It currently splits the tape into 21 / 17 runs that mean nothing — and without the fix, next Friday's agent digs the same empty hole again. Evidence: Movement 1 · Movement 3 · verification: bug · mode: shadow	The <code> flow > 50</code> rule fires on 66.5% of ALL bars — FLOW-LED and VACUUM are two names for one coin.	Claude — two scripts, <code>run_census.py:30</code> and <code>run_census_full.py:38</code>	Git revert. Without it, next Friday's agent digs the same empty hole.
6	Wire NIPC into the shadow slate — SHADOW, explicitly never live this week Wire the news-impulse pullback into <code>src/gazbot7/shadow.py::default_slate()</code> (line 179). HOME regimes ON, dead-chop OFF, Lot A 2.0-2.5R + Lot B fixed 2.5R — the proven pair, not the 1.5/2.5 guess. Log the trail-vs-fixed A/B explicitly, because it contradicts our standing chandelier prior. Why: +\$2,676 on n=171 over 12 days, 43.3% win, \$15.6/trade, null-control z = +5.0, and fixed 2.5R beats the trail by \$1,285. Strong for a greenfield candidate — and precisely because it is greenfield it goes to shadow and nowhere near live capital this week. Evidence: Movement 3 · verification: exploratory · mode: shadow	+\$2,676 on n=171 over 12 days, 43.3% win, \$15.6/trade, null-control z = +5.0. Fixed 2.5R beats the trail by \$1,285.	Claude — shadow slate only	Remove from <code>default_slate()</code> . Explicitly never live this week.
7	Optional: pin down what "confirmed trend" MEANS — a guardrail worth \$0 Only if you want determinism. A prepared prompt insert for <code>router_tick_durable.py</code> defines a confirmed trend as ALL FOUR of: 30-min ER >= 0.35; 30-min net >= 60 MNQ pts in the trend direction; ATR >= 14; and the same read held on the previous tick. It does NOT loosen the chop/fader benching. Written out in full in the REV3 section; revert is deleting seven string-literal lines. Why: Sized honestly against itself: as a permission filter it clears 26 of 274 aligned-long signal-instants (9%), worth <code>-\$28 raw / +\$86 non-overlap</code> , and even that flips negative on two of four leave-one-out folds and on strip-best-3 (<code>-\$260</code>). It earns approximately NOTHING this week. Its only value is determinism: right now 'confirmed TREND_UP' is whatever a headless Claude decides at 03:00 UTC unattended. Size it at \$0, take it as a guardrail or leave it, and let a real trend week settle the arming question. Evidence: REV3 TREND_UP \$8 · verification: exploratory · mode: shadow	\$0. As a permission filter it clears 26 of 274 signals (9%), worth +\$86 non-overlap, and flips negative on two of four folds. Its value is determinism, not money.	Garrath — optional, prompt text only	Delete the seven inserted string-literal lines. Fail-safe either way.

HOLD · LEAVE ALONE — THESE ARE ALREADY RIGHT · 3 PLAYS

Standing lessons and things already in production. The action is to not touch them.

#	WHAT TO DO	THE NUMBER BEHIND IT	OWNER	REVERT PATH
1	On chop days the OFF switch beats every gate I could build £500/wk HOLD this as the governing lesson of the week, not as a new build. On a confirmed chop day, going flat outperforms every gate the greenfield lab could construct. This is what the untradeable meter operationalises. Why: The greenfield lab built and swept a large family of chop specialists and none of them beat simply not trading — that part stands. CORRECTION to the version of this that appeared earlier: I previously cited Q4's 'benching into CHOP = +\$1,088, positive in every fold' as independent corroboration. The re-derivation withdraws that support — bench-into-CHOP scores +\$612 / +\$187 / <code>-\$382</code> across three independent regime tags, i.e. it is TAG-FRAGILE and does not survive an independently-computed definition of chop. The one router claim that survives every test is the FADER bench (<code>+\$810 non-overlap, n=25</code> , identical under every tag). Treat 'flat on chop days' as the lab's own finding, not as a twice-proven one. Evidence: Movement 3 Greenfield lab · verification: verified · mode: live	A large swept family of chop specialists; none beat simply not trading. Note the correction: the "bench-into-CHOP +\$1,088" corroboration has been withdrawn — it is tag-fragile (+\$612 / +\$187 / <code>-\$382</code> across three independent regime tags).	The untradeable meter operationalises it	N/A — this is the governing lesson of the week, not a build.
2	The per-tick stop guard is fixed and deployed — do NOT revert it NO ACTION beyond leaving it alone. The per-tick stop guard (commit d0218c6, live 07-28 08:34) checks every ~2s and market-flattens at the stop, so stops no longer depend on the broker's trigger firing at all. Do not revert it, and do not chase this as if the bleed were still open. Distinct from the front-month branch fix listed under BUILD.	Before: 26 unfilled stops. At and after commit d0218c6 (live 07-28 08:34): 0 unfilled, 103 clean stops. On real ticks the 2s path beats a perfect stop by about \$15.	Nobody — it is already in production	Do NOT revert it.

#	WHAT TO DO	THE NUMBER BEHIND IT	OWNER	REVERT PATH
	Why: The strongest win of the week and it is already in production. Before the fix: 26 unfilled stops. At and after: 0 unfilled, 103 clean stops — the last unfilled stop in the entire book was 32 minutes before that deploy. Two corrections to the original framing: most of the headline was mislabelled (of the —\$195 only —\$53 is the actual defect; archive-wide only 24% of the bleed was ever the bug), and on real ticks the 2s path actually BEATS a 'perfect' stop by about \$15, because fast bounces come back inside 2 seconds. Evidence: Part 1.5 Rehabilitation · verification: bug · mode: live			
3	Adaptive exit-width selector — it earned its keep, keep it £250/wk HOLD. No change. The regime-at-entry exit-width selector stays as shipped. Why: It earned its keep on this week's tape. Left alone deliberately — the exit changes on this card are per-gate cells (grind, exhaustion, abs_veto_long), not a change to the selector mechanism itself. Evidence: Part 2.5 Musings · verification: verified · mode: live	It earned its keep on this week's tape. The exit changes elsewhere on this card are per-gate cells, not a change to the selector mechanism.	Nobody — no change	N/A.

NOT-AN-ACTION · WITHDRAWN / REFUTED — KEPT ON THE CARD SO NOBODY RE-PROPOSES THEM · 9 PLAYS

Every one of these was a recommendation at some point this week. They are printed with their autopsies so next Friday's agent does not dig the same hole.

#	WHAT TO DO	THE NUMBER BEHIND IT	OWNER	REVERT PATH
1	"Promote abs_veto_55s two-sided" — ALREADY LIVE, delete the recommendation Not an action. Live since the 07-25 roster change. Part 2 re-derived it cold, which is a validation, not a change request. Kept on the card only so it does not get re-proposed next week. Why: abs_veto_long=on and abs_veto_short=on in gate_switches.env right now, verified live today. See the SATURDAY confirm row for the validation numbers. Evidence: Part 2 verdict, corrected by Q2 D0 · verification: refuted · mode: skip	Live since the 07-25 roster change; both sides verified armed in gate_switches.env today.	Nobody	N/A.
1	WITHDRAWN — "stop benching momentum in TREND_UP" (+ \$1,615) does not survive WITHDRAWN. This was the rank-1 recommendation and it is refuted. Do NOT change the router's TREND_UP benching. Kept on the card with the full autopsy so it is not re-proposed next week off the same estimator. Why: An independent re-derivation from source (37,160 SUPPRESSED-OPEN journal lines, repriced on 8.45M capture.db 250ms ticks) reproduced Q4's arithmetic almost exactly (+\$1,644 vs +\$1,615) and confirmed both of Q4's own corrections — the A/B double-count and the zero-slippage point. Then it broke the finding three ways. (1) TEMPORAL DOUBLE-COUNT — the same error Q4 caught, one axis over: router_nightly.py dedups at 120s and prices every survivor as a complete round trip even while the same gate is still in the previous one. Enforce one open position per gate and +\$1,799 becomes +\$767, a 57% cut; strip the best 5 signals and it is +\$85. On Q4's literal wording (all momentum, which includes arming abs_veto_short inside an uptrend, n=6, −\$154, 17% win) it is +\$612 and strip-best-5 goes NEGATIVE at −\$70. (2) IT FAILS AN INDEPENDENT REGIME TAG — placebo test, 20,000 draws: the router's own replay tag scores z=+3.60, but ADX(14) 5-min >= 20 gives z=−0.16, p=0.56 (pure null) and a 60-min OLS R² tag gives z=−2.56, WORSE than random. The obvious defence (different horizon) fails: a same-30-min-horizon regression tag scores \$0.6/signal. The discriminator is the specific ER_TREND=0.20 / NET_MIN=30 / HOLD=2 triple, tuned on the week immediately prior. Under both independent tags CHOP is the better bucket, because the whole aligned momentum-long blocked book is +\$2,078 across all 274 signals in ALL regimes — a finding that says 'you should never have benched a momentum long anywhere, all week' is not a TREND_UP finding. (3) IT IS BOUNCE MONEY ON THE BIGGEST DOWN DAY — the two largest episodes are both grind_long on 07-29 (+\$570 and +\$371 = 52% of the claim), a day MNQ fell 702pt at day-ER 0.72 with price 400-700pt below the open; 34% of the money was earned at day-bias <= −40, which the framework says to bench longs into. The one clean UP day contributes \$349 of \$1,799. Concentration: 21 episodes, top-1 = 35%, top-3 = 79%. And 75% of the claim (n=32) fired at ER30 <= 0.20 — HOLD=2 hysteresis on a 15-min grid means the TREND_UP label can be half an hour stale, so 57% of signals fired while the tape was, on its own instantaneous read, CHOP. Reported straight, what could NOT be killed: the estimator is not optimistic (modelled	+\$1,615 → +\$767 once one open position per gate is enforced → + \$158 on n=4 once the rule cannot buy dips in a downtrend. Placebo: null under ADX(14) (p=0.56), worse than random under a 60-min regression tag (z=−2.56).	Nobody — do NOT change the router's TREND_UP benching	N/A. Kept on the card with the full autopsy so it is not re-proposed off the same estimator.

#	WHAT TO DO	THE NUMBER BEHIND IT	OWNER	REVERT PATH
	blocked momentum \$3.9/trade vs realised \$4.3/trade — breakeven drag is \$24/RT $\approx$ 48 ticks, so this is not a slippage argument) and exit choice is not the artefact (six exits, all positive). Also corrected: Q4's claim that this is 'already banked as [[us-open-dont-bench-trend-rider]], simply not enforced' is wrong — the banked rule is narrow (hold the aligned rider through OPEN FLUSHES on a cleanly-trending day) and the ledger shows it IS being applied correctly at both boundaries. Q4 proposes a new, much wider rule. Decisive context: by an operational definition a confirmed trend was live for 3.3% of this week — 22 episodes, median 6 minutes. There is nothing to ride. Evidence: REV2 · Q4 Router value · verification: refuted · mode: skip			
2	WITHDRAWN — grind_long exit {a_r 0.5, b 1.0} is WORSE than leaving it alone WITHDRAWN. Do NOT add this cell. On a sequential one-position-per-slot sim over 3,064 tick-repriced entries it is −\$1,127 WORSE than the default that is loaded today. It was the report's Monday exit recommendation and it is refuted. Why: Reproducing the Friday lab's own paired method but on 250ms ticks and LONG-SIDE ONLY, the MED-TREND cell that justified this goes from the reported +\$10.3/trade at 59% win to −\$5.19/signal at 46% win, and the SCALP-CHOP corroboration flips outright (0.5/1.0 = −\$0.84 vs wide's +\$11.01). Two causes, both mechanical: 5-second repricing FLATTERS tight exits, and the cell pooled grind_short into what was a long-only decision. It fails every robustness test — 8/16 sessions, LODO-worst −\$2,036, strip-best-3 −\$3,392, bootstrap P=0.302, OOS sign flip (+\$699 / −\$1,826). Evidence: REV2 · Q0 M5/M5b · Part 2 LAB · Q3 · verification: refuted · mode: skip	−\$1,127 worse than the default that is loaded today. 8/16 sessions, LODO-worst −\$2,036, strip-best-3 −\$3,392, bootstrap P=0.302, OOS sign flip (+\$699 / −\$1,826).	Nobody — do NOT add this cell	N/A.
2	Ship the chop-day scalp (chop_turn_fade) — HONEST NULL Do not ship it, and do not incubate it with any expectation. Best-of-1,152 configurations returned +\$288 against a pure-noise expectation of +\$455 — it underperformed noise. Why: Leave-one-day-out OOS is −\$5.84/trade. Searching 1,152 configs and finding a best that is worse than the noise floor is the textbook signature of nothing being there. Evidence: Movement 3 chop lab · verification: refuted · mode: skip	Best of 1,152 configurations returned +\$288 against a pure-noise expectation of +\$455 — it underperformed noise. LODO OOS −\$5.84/trade.	Nobody	N/A.
3	Retire rgv_short — NO Do not retire it. Bench it and shadow it (see the MONDAY row), and re-look when n grows. Why: Retiring on 11 fires is exactly as unjustified as arming on 11 fires. Unproven is not the same as disproven, and the asymmetry the desk runs on says bench-and-watch is the cheap option. Evidence: Part 1.5 §5 · verification: refuted · mode: skip	11 fires. Retiring on 11 is exactly as unjustified as arming on 11.	Nobody — bench and shadow instead (MONDAY #3)	N/A.
4	Router timing knobs (window / step / hold / fast_exit) and NET-min — HOLD Do not touch them. The fader-bench ER floor is the only router threshold moving this week. Why: The joint-best (+\$692 / +\$984) sits on the loosest grid edge on three axes at once — the signature of a fit to the grid boundary rather than to the tape. And step=5m would triple the tick-rate for no proven gain. Evidence: Part 2.6 §4 anti-overfit list · verification: refuted · mode: skip	The joint-best (+\$692 / +\$984) sits on the loosest grid edge on three axes at once — the signature of a fit to the grid boundary.	Nobody	N/A.
5	Wire in the two-ratchet clip or the 2R-partial — TOO THIN Wire nothing. Keep both shadowing. Why: 4 trades, 1 runner, 1 day. The entire delta (+\$122 / +\$64) is that single 4.2R ride. There is no measurement here yet. Evidence: Part 1.5 shadow-watch · Part 2 · verification: exploratory · mode: shadow	4 trades, 1 runner, 1 day. The entire delta (+\$122 / +\$64) is that single 4.2R ride.	Nobody — keep both shadowing	N/A.
6	Arm something off the idle-gate lab — NO, that is survivorship Do not arm anything from the idle-gate lab. Why: What the lab surfaced is survivorship, not skill — the gates look good because of which ones happened to sit idle, not because of anything they did. Evidence: Movement 2 Idle-gate lab · verification: refuted · mode: skip	Survivorship, not skill — the gates look good because of which ones happened to sit idle.	Nobody	N/A.
7	The whole flow-footprint family is BURIED Stop building on flow-footprint signals. Buried, with the reason recorded so it does not get re-dug. Why: Compounded by the census cluster bug (BUILD): the flow > 50 rule fires on 66.5% of bars, so the clusters this family was built to exploit were never real clusters. Fix the census before anyone revisits the family. Evidence: Movement 3 Greenfield lab · verification: refuted · mode: skip	Compounded by the census bug: flow > 50 fires on 66.5% of bars, so the clusters this family was built to exploit were never clusters.	Nobody — fix the census before anyone revisits it	N/A.

STANDING CAVEAT ON EVERY DOLLAR ON THIS CARD

Every repriced number in this report is a **CEILING**. Fills are modelled at the exact stop/target/trail level with no slippage, and the desk's own standing finding is that live losses run **1.1-3.1× modelled**. Where a change breaks even at a loss multiplier inside that band — grind_long at 1.64× is the live example — treat the policy as the edge and the headline dollar as decoration. Anything on thin n is flagged in-row as a direction, not a measurement.

§ The evidence behind those plays

Everything from here down is the working. Two halves, one document. **The warm half** is what the live six-gate paper tournament actually did this week (Part 1), then the section that matters most — **Rehabilitation** (Part 1.5), where every red gate and leaky exit is walked through the full six-step wash instead of being benched at face value, graves and all. Then the shadow board and the promotion it earned plus the regime-flex exit lab (Part 2), the mid-week leads re-run on the week's own data (Part 2.5), and the **Router Operation Review** (Part 2.6). **Part 3** is the new material: the Rev2 answers that never landed and the three Rev3 re-derivations that overturn three of them. **Then the cold, tape-first half**: we ignore everything the desk did and ask, from the raw MNQ tape and order book, **where was the money and did we show up?** (M1) — then test whether the gates we already own could have caught the runs we missed (M2), and finally build brand-new gates from scratch and backtest them to death, the ones that failed included, that is the point (M3). Plain English, honest money, every claim a table — and every reprice a ceiling.

WHAT IS IN HERE

1. **Front** — [the honest headline](#), [the scorecard](#), [WHAT TO ACTUALLY DO](#)

2. **Part 1** — [The live desk — what the six-gate paper tournament actually did](#)
 3. **Part 1.5** — [Live-desk REHABILITATION — never bench a gate at face value](#)
 4. **Part 2** — [The shadow desk, the promotion battery & the regime-flex exit lab](#)
 5. **Part 2.5** — [Mid-week musings — every lead from the desk chat, tested hard](#)
 6. **Part 2.6** — [ROUTER OPERATION REVIEW — the desk's #1 lever, on the stand](#)
 7. **Part 3.1** — [REV2 · Q1 — what the machine did with your hands off](#)
 8. **Part 3.2** — [REV2 · Q2 — the dead floors, and the abs_veto live-vs-shadow gap reconciled](#)
 9. **Part 3.3** — [REV2 · Q3 — the BIG-TREND cell is n=24, and the dated plan to grade it](#)
 10. **Part 3.4** — [REV2 · Q4 — the router's real value \(superseded in part by 3.5\)](#)
 11. **Part 3.5** — [REV3 — the TREND_UP re-derivation: +\\$1,615 WITHDRAWN](#)
 12. **Part 3.6** — [REV3 — the dead ER/ATR floors CONFIRMED, and the order to ship them in](#)
 13. **Part 3.7** — [REV3 — grind_long's exit swept sequentially: the {0.5, 1.0} cell WITHDRAWN](#)
 14. **Part 3.8** — [REV3 — the consistency audit that forced this revision](#)
 15. **Movement 1** — [The census — the biggest runs on MNQ this week, and did we show up](#)
 16. **Movement 2** — [The idle-gate lab — could the gates we own have caught them?](#)
 17. **Movement 3** — [The greenfield lab — brand-new entries, and the chop-day scalp lab](#)
-

Part 1 The Live Desk — the week the router earned its keep

► **Read this section with REV3 open.** Three claims in here were overturned this weekend: `exhaustion_short` is **not** the week's earner once your hands come off it, `grind`'s "0.35 ER floor" has never executed, and Thursday's +\$1,172 record is a session-boundary artefact of nine hand-claims. The corrections are marked inline in red where they bite, and the full working is in [Part 3 — the Rev2/Rev3 corrections](#). Nothing has been deleted: where a claim was overturned, the original is still printed next to what replaced it.

★ READ THIS FIRST — THE AUTOMATED DESK LOST MONEY THIS WEEK

The +\$625 headline below is real, but it is **not what the machine did on its own**. Twelve of the week's 206 trades were closed by your own hand (`exit_reason = MANUAL_CLAIM`), and those twelve booked **+\$1,223**. Strip them out and the automated desk — every gate, every exit, running unattended — is **−\$598 over 194 trades at 35.1% win**. This was flagged as a caveat in an earlier draft but never computed; here it is.

CUT	NET \$	N	WIN%
Headline (as booked)	+625.0	206	38.8%
Automated only (excl. manual claims)	−598.0	194	35.1%
Your manual claims	+1,223.0	12	—

It matters most for the gate this report calls the survivor. Per gate, automated vs hand-claimed:

GATE	AUTOMATED \$	N	MANUAL \$	N	TOTAL \$
<code>exhaustion</code> (the nominated survivor)	−56.0	44	+372.5	4	+316.5
<code>grind</code>	−421.5	51	+605.5	5	+184.0
<code>abs_veto</code> (all sides)	+173.0	82	+245.0	3	+418.0
<code>capitulation_long</code>	−139.0	14	0.0	0	−139.0
<code>rgv_short</code>	−154.5	3	0.0	0	−154.5

So `exhaustion_short`'s "green on 4 of 5 days" is hand-claim dependent — unattended it is roughly flat-to-negative (−\$56 over 44 fires). Only

abs_veto is positive on its own machinery. Everything downstream in this report that leans on exhaustion being "the steadiest earner" should be read against that. This is the single most important number in the week and it argues the opposite of the headline: the entries and the automated exits did not pay; your discretionary claims did.

Garrath — the 6-gate paper tournament made **+\$625** across **206 trades** this week (Mon 27th through Fri 31st, on the Paris session-day boundary the desk keeps its books on). That headline hides a wild ride: **one best-ever day (+\$1,172 on Thursday) ► REV3: +\$1,172 on the Paris boundary; +\$94 on the UTC day (Movement 3 books it that way), and -\$884.5 of it is automated — nine hand-claims are what made it a record.**, one clean two-way trend day (+\$638 Wednesday), a near-scratch Tuesday, and two red days that bookended the week — a choppy Monday (-\$227) and an ugly, whippy Friday (-\$905). The money came from exactly two places: the gates that ride a real move (grind and abs_veto on the trend legs) and the gate that fades a tired one (exhaustion_short). Everything else either scratched or bled, and the router spent the week benching the bleeders. Below is the whole thing, every number computed off the tournament trades, honest about the fact that it is one week of in-sample paper.

READ THIS FIRST — WHAT SURVIVED THE WEEK. IF YOU STRIP AWAY THE NOISE, ONE GATE STOOD UP ON ALMOST EVERY KIND OF DAY: EXHAUSTION_SHORT WAS GREEN ON 4 OF THE 5 SESSIONS (+\$14, +\$90, +\$95, +\$217, AND ONLY -\$100 ON FRIDAY'S VIOLENCE), 54% OF ITS TRADES WON, +\$316 FOR THE WEEK. ► REV3: \$372.5 OF THAT +\$316 WAS FOUR TRADES YOU CLOSED BY HAND. AUTOMATED IT IS -\$56 OVER 44 FIRES (SEE THE CALLOUT ABOVE). IT IS THE STEADIEST THING ON THE DESK, NOT THE BIGGEST EARNER: 50% WIN, A MEDIAN TRADE OF ESSENTIALLY ZERO, AND THE SMALLEST LEAVE-ONE-DAY-OUT DAMAGE OF ANY GATE THAT TRADED IN SIZE. THE FLASHIER NUMBERS — GRIND'S +\$704 THURSDAY, ABS_VETO_SHORT'S +\$570 WEDNESDAY — ARE REAL BUT THEY ARE DAY-TYPE BETS: THEY PRINTED ON THE ONE OR TWO DAYS THE TAPE ACTUALLY TRENDED AND GAVE MONEY BACK THE REST OF THE TIME. LEAD WITH THE STEADY ONE; TREAT THE BIG ONES AS CONDITIONAL.

The week, day by day

Every day scored on the desk's own Paris session boundary, so these match the numbers in your ledger (Thursday's +\$1,172 is the best single day the desk has ever booked **on this boundary** — Movement 3 books the same trades on the UTC calendar day and gets +\$94; same week total, different day buckets). The "character" column is the router's live regime read, not a number I invented — it's what actually drove the bench/arm calls.

SESSION DAY	NET \$	TRADES	WHAT THE TAPE DID
Mon 07-27	-226.5	40	Choppy down-drift; a Sunday-reopen "all gates on" churned the wrong-side gates before the router could bench them ($\approx -\$275$ of self-inflicted reopen churn).
Tue 07-28	-54.0	23	Quiet, low-vol chop. Near-scratch — the shorts (abs_veto +\$68, exhaustion +\$90) roughly paid for the longs' churn.
Wed 07-29	+638.5	17	The clean trend day — a strong down-leg then a strong up-leg, and the desk caught BOTH sides. Only 17 trades: when it's real, you don't need many.
Thu 07-30	+1,172.0	91	Best day ever. An up-bias that topped out and spent the whole US session whipping in a 1,205-point range. 91 trades — the router fought this one all day (more below).
Fri 07-31	-905.0	35	Worst day. A US-open up-thrust that fully reversed into an ER-0.01 chop; an "enable all at the open" call handed the churners the reversal.
Week	+625.0	206	39% of trades won (80 of 206). Two green days paid for three red/scratch ones.

The shape of that week is the whole story: the desk does not need to win most days, it needs to be on the right gates on the two days that trend and out of the way on the days that don't. Thursday and Wednesday together made +\$1,811; the other three days gave -\$1,186 of it back. Net +\$625.

The six gates — the card table

This is the tournament roster this week: two "ride it" momentum gates (grind_long, abs_veto's long and short sides), two faders (exhaustion_short, capitulation_long) and the router-managed rgv_short. Each gate runs two sub-slots on one signal — a Lot-A scalp and a Lot-B chandelier tail — but here they're totalled to the whole gate. Net, count, win-rate, the average winner and average loser, and the single best and worst trade it took all week.

GATE	NET \$	N	WIN%	AVG WIN	AVG LOSS	BEST	WORST	VERDICT
abs_veto_short	+358.5	26	42	+94	−45	+536	−77	DOWN-LEG CATCHER
exhaustion_short	+316.5	48	54	+61	−58	+324	−123	FLAT AUTOMATED (−\$56 / 44) SURVIVOR — steadiest — the + \$316 includes 4 hand-claims
grind_long	+184.0	56	32	+125	−54	+305	−235	TREND-DAY ONLY
abs_veto_long	+59.5	59	36	+91	−48	+342	−85	SCRATCH / veto-protected
capitulation_long	−139.0	14	29	+64	−40	+105	−147	→ REHAB (Part 1.5)
rgv_short	−154.5	3	0	—	−52	−47	−54	ROUTER-MANAGED

Two things jump out. First, **grind is a low-win, big-winner machine** — it only won 32% of its trades but its average winner (+\$125) is more than double its average loser (−\$54); it loses small a lot and wins big occasionally, which is exactly what a trend-rider should look like. Second, **rgv_short barely fired at all** (3 trades) — the router had it benched for almost the whole week because it kept trying to short an up-biased tape, and the 3 that slipped through all stopped. That's not a broken gate, that's the router correctly keeping it in the dugout.

Where the money actually came from — the exit mix

This is the most honest picture of the week. The stops are enormous in aggregate — the desk paid over six thousand dollars in stop-outs — but that's the price of admission: you take a lot of small stops to be present for the few big tails. The green all comes from three exits: the chandelier tails, the fixed targets, and the trades you claimed by hand.

EXIT REASON	COUNT	NET \$	WHAT IT IS
CHANDELIER	32	+3,065.0	The trailing-tail exit — the big runners. The single biggest profit bucket.
TARGET	33	+2,647.5	Fixed profit target hit.
MANUAL_CLAIM	12	+1,223.0	You claimed the profit by hand (a lot of the exhaustion fades).
GIVEBACK	4	+82.0	Give-back trail banked a small green.
STOP	119	−6,119.0	Normal stop-outs — the cost of being present.
STOP_UNFILLED	6	−273.5	The software-stop failures, caught by the guard (see below).

The lesson baked in here: **the tail is the edge**. When you split each gate into its scalp lot (Lot A) and its chandelier tail lot (Lot B), the tail carries nearly all the money — grind's tail made +\$302 while its scalp made +\$31; abs_veto_long's tail +\$322 vs its scalp +\$5; abs_veto_short's tail +\$380 vs +\$52. The scalp lots roughly break even and lock in certainty; the tail lots are where the week was won.

Did a gate's P&L depend on trend vs chop? (the gate × day grid)

Yes — dramatically for the momentum gates, much less for the fader. This grid is net \$ (trade count) per gate per session day. Read the columns: Wednesday and Thursday were the trend/whip days, and that's where grind

and abs_veto made everything. Monday, Tuesday and Friday were chop, and the momentum gates gave money back on every one of them.

GATE	MON 27	TUE 28	WED 29	THU 30	FRI 31
exhaustion_short	+14 (23)	+90 (10)	+95 (6)	+217 (6)	-100 (3)
abs_veto_short	-6 (4)	+68 (6)	+571 (10)	—	-274 (6)
grind_long	-149 (4)	—	—	+704 (46)	-371 (6)
abs_veto_long	-118 (7)	-122 (4)	-27 (1)	+242 (35)	+84 (12)
capitulation_long	+80 (1)	-90 (3)	—	+10 (4)	-138 (6)
rgv_short	-48 (1)	—	—	—	-107 (2)

Look at grind's row: it is essentially a bet on Thursday. It made +\$704 there on 46 trades and lost on the two other days it fired at all. That is not a flaw — it's what the 0.35 ER floor would be for, keeping it dormant unless the day really trends — ► **REV3: except that floor has never executed once. It is keyed on grind_long while the live slots are called grind_long_A/_B, so it has been a no-op since 07-29 (Part 3.6).** Either way grind's + \$184 for the week is one great day minus two poor ones, not a broad edge. Compare exhaustion's row: green four days out of five, never a monster but never a disaster until Friday's open. ~~That row is why exhaustion is the survivor.~~ ► **REV3: that row is why exhaustion is the steadiest — read against the automated split above, unattended the row is flat.**

Case 1

Thursday — the best day ever, and the router in the ring



The morning ramp is where the money was made. Note the five-hour gap after ~16:00 — the router benched the book into the roll-over and the desk simply stopped trading.

Thursday (+\$1,172) is the week's headline and its best case study, because it was not a clean trend — it was an up-biased tape that topped out around lunchtime and then whipped violently for hours, and the desk still booked its best day ever by staying on the right gates through a dozen regime flips. This is the day to study when you want to know what the router is actually for.

How Thursday was won

The morning trended up and grind rode it — grind_long made **+\$704 across 46 trades** on the day (32 of them UTC-Thursday), catching the up-legs with its chandelier tail. Then around 13:50 UTC the buy-climax printed (aggressor delta peaked, then buy intensity collapsed to a third), the tape rolled over, and the router flipped the whole book: benched the momentum longs and armed the short-side reversion faders. The result was the cleanest offensive sequence of the week — **exhaustion_short banked four fade-claims in a row for +\$372.5** (+\$117, +\$107.5, +\$109.5, +\$38.5) fading the roll-over. It then gave -\$155 back on two later fires when the chop tried to resume up, netting +\$217 on the day — but that four-in-a-row is the exhaustion edge on full display: judge this gate on its live fades, not its fixed-target shadow.

★ REV3 CORRECTION — 2026-08-01 — THOSE FOUR FADES WERE FOUR HAND-CLICKS, AND THE MACHINE WOULD HAVE LOST ON THEM

All four were MANUAL_CLAIM exits — **you** closed them. Repriced on 250 ms tape under the exit config that was actually loaded at the time, the same four entries come out at **−\$51.2**, and three of the four were sitting on positions the tape later ran 49 to 98 points through the stop. So the router's best-graded offensive call of the week (Part 2.6 §2 credits it +\$372) is worth **−\$51** to the machine. That is a \$423 swing inside the router's own scorecard.

THURSDAY'S ENGINE	NET \$	N	NOTE
grind_long (up-legs)		+704	46 Rode the morning trend + the afternoon resumptions with the chandelier tail.
abs_veto_long		+242	35 The veto-protected long — caught the confirmed breaks, sat out the fakes.
exhaustion_short (roll-over fade)		+217	6 The 4-claim run +\$372.5 fading the top, minus 2 late stops.

THE TWO CALLS THAT WENT WRONG ON THURSDAY — AND WHY THEY'RE WORTH MORE THAN THE WINS. THE ROUTER MADE TWO GENUINELY BAD CALLS THIS DAY, BOTH CAUGHT AND BOTH BANKED INTO THE DESK'S RULEBOOK. (1) AT 13:36 IT BENCHED GRIND ON THE US-OPEN FLUSH — THE DAY GAVE BACK A CHUNK IN 90 SECONDS AND THE ROUTER READ GRIND AS CHURNING. YOU OVERRODE IT ("LIVEN GRIND BACK UP"), GRIND THEN RODE TO NEW HIGHS AND BOOKED +\$234/+\$197 TARGETS. LESSON BANKED: DON'T BENCH THE ALIGNED TREND-RIDER ON OPEN STOP-OUTS WHEN THE DAY IS CLEANLY TRENDING — THE OPEN IS STRUCTURALLY VIOLENT, THAT'S NOISE, ONLY BENCH IF THE TREND STRUCTURE ACTUALLY BREAKS. (2) AT 15:49 IT FLIPPED TO LONGS ON A DELTA-BLIP IN ER-0.01 DEAD CHOP — A FOUR-BUCKET BUYER BUILD AND A +45-POINT POKE OF THE RANGE TOP. IT FADED IN TWO MINUTES. YOU CAUGHT IT ("WHY IS GRIND BUYING INTO CHOP"). LESSON BANKED: RE-ARM MOMENTUM ONLY ON A REAL BREAK WITH ER ACTUALLY CLIMBING AND VOL EXPANDING, NEVER ON A DELTA BLIP ALONE. NEITHER MISTAKE COST REAL MONEY (BOTH GOT SALVAGED TO ROUGHLY +\$55), BUT THE PROCESS ERRORS ARE THE ONES THAT WRITE THE RULES.

The one call that clearly saved money

At 14:51, when ER collapsed from 0.40 to 0.07 over two ticks and price rolled off the highs, the router **proactively benched both momentum longs while the desk was flat** — not reacting to stops, reading the regime flip ahead of them. Over the window that followed, the momentum shadows that would have kept firing lost about **−\$400** (thrust_aligned −\$202, thrust_cont −\$118, thrust_loose −\$78). The desk sat flat and protected instead. That single defensive read is the clearest example of the router paying for itself: it's cheap to be wrongly benched (you miss a few trades), expensive to be wrongly armed (you churn), and this call landed on the right side of that asymmetry.

Case 2 Wednesday — the clean two-way trend, both sides caught




Both directions caught. The down-leg then the up-leg, with the desk on the right side of each.

Wednesday (+\$638 on just 17 trades) is the "when it's real, take it and go home" day. The tape ran a strong down-leg and then a strong up-leg, and — unusually — the desk caught both directions cleanly.

abs_veto_short's day

The star was **abs_veto_short: 10 trades, +\$570.5**, riding the down-leg. It only won 4 of those 10 — but one of them was a **+\$536 chandelier runner**, the single best trade of the entire week, and that one tail dwarfed the small stops around it. This is the abs_veto shape in one trade: the 55-second continuation-veto keeps it out of the fake thrusts, and when it does catch a real one, the chandelier lets it run. Then late in the session the tape turned and trended up cleanly; the router benched the shorts (correctly — you don't short a strong up-move) and let the longs ride. Both legs, both green, 17 trades total. That's the ideal day and it's worth remembering it looks nothing like Thursday's 91-trade brawl.

Case 3 Friday — the enable-all-at-open bleed



The opposite picture: entries scattered through a whippy range with no trend to ride.

Friday (−\$905) is the honest graveyard of the week and the sharpest single lesson in it. The US open printed a real up-thrust (ER 0.68, a clean break to 28713) and the book was set to "all six gates on" for the open. The thrust then fully reversed — a −172-point round-trip, ER collapsing from 0.50 to 0.12 into ATR-62 violence — and the wrong gates were holding the bag.

The open, gate by gate (13:30-14:30 UTC)

GATE AT THE OPEN	NET \$	STOPS	WHAT HAPPENED
abs_veto_long (veto-protected)	+146.5		— Caught the real thrust, banked it, then its veto kept it OUT of the reversal.
exhaustion_short (off-regime)	−72.0		1 Fired the new 0.5R/1.5R scalp trial into ATR-62 — a scalp can't cover a 62-point stop.
grind_long (no veto)	−284.0		4 Bought the rollover top, walled — four stops in ninety seconds.

The contrast is the whole lesson, and it's now a banked rule: **at a violent open, arm the veto-gates, not the churners**. abs_veto's 55-second confirmation made it +\$146 by only taking the thrust that continued; grind, with no such filter, chased the reversal and lost −\$284. The rest of the day was a choppy recovery (ER 0.24-0.30 — elevated but not clean), and two attempts to re-arm abs_veto into it churned ≈−\$245 more before you called it and went fully flat by 19:08. Going flat was the right end-state: the live shadows confirmed every benched gate would have kept bleeding through the ATR-52/62 chop. Some days the winning move is no move.

ONE FRIDAY LEAD FOR THE SHADOW DESK (PART 2): THE ONLY CONFIGURATION THAT MADE MONEY ON FRIDAY'S WATERFALL WAS RG_LONG_FAST AT +\$249, FADING THE OVERSOLD FLUSH AND BUYING THE CAPITULATION DIPS — WHILE OUR LIVE LONG-FADER, CAPITULATION_LONG, BLED. THAT'S A FLAG THAT THE RIGHT FLUSH-CATCHER MIGHT BE

RG_LONG_FAST, NOT CAPITULATION. CHASED DOWN PROPERLY IN THE SHADOW SECTION.

The STOP_UNFILLED self-heals

Six trades this week exited on STOP_UNFILLED — the known issue where IBKR won't fire a resting stop's trigger on the continuous contract. Every one of them was an abs_veto trade, and here's the important part: **the stop-breach guard caught all six and bounded each loss to roughly one ATR**. None ran naked, none turned into a drift-halt. That's the self-heal working as designed.

SESSION DAY	GATE	SIDE	LOSS \$
Mon 07-27	abs_veto_short	SHORT	-61.5
Mon 07-27	abs_veto_long	LONG	-51.0
Mon 07-27	abs_veto_long	LONG	-43.0
Mon 07-27	abs_veto_long	LONG	-35.5
Tue 07-28	abs_veto_long	LONG	-38.0
Tue 07-28	abs_veto_short	SHORT	-44.5
Total			-273.5

So the software failure cost about -\$274 this week, all of it bounded and none of it catastrophic. The fix is **already deployed and verified** — see Part 1.5: commit d0218c6 went live 07-28 08:34 and there have been zero unfilled stops since, across 103 clean stop exits. The concrete-front-month branch is now moot. The guard remains the safety net underneath and it held six times out of six.

The router The direction-router trial — how it behaved

This was the week the Claude-run router stopped being a session experiment and became durable infrastructure — a headless systemd tick went live Friday that manages the desk even with no chat session running.

Across the week the router changed the gate switches **59 times across 590 five-minute ticks** (90% of ticks were no-change; 2 of the 59 came from the durable systemd tick, the rest from the session). It is benching-only, fully reversible, and paper — it can only block new entries, never force a trade — so its worst case is a missed opportunity, not a loss.

The router's own scorecard

Every switch change gets graded in the bad-call ledger once the outcome is known. The running tally sits at **17 GOOD · 2 BAD · 4 CHURN · 2 NEUTRAL**. The honest read: most calls helped or cost nothing, the two outright bad calls both happened on Thursday and both were caught within minutes and turned into permanent rules, and the "churn" entries are process-thrash (net roughly zero in dollars) rather than money errors.

GRADE	COUNT	MEANING & THE NOTABLE ONES
☐ GOOD	17	Calls that helped or avoided loss — incl. the Thu 14:51 regime-flip bench (≈-\$400 churn avoided) and arming the chop-book that produced exhaustion's + \$372.5 run.
☐ BAD	2	Both on Thu: benching the trend-rider on the open flush, and the premature delta-blip flip to longs. Both banked as rules.
△ CHURN	4	Over-toggling grind ~8x on Thursday morning; the Sunday-reopen "all six on"; two Friday re-enables into a choppy recovery. Thrashy, ≈\$0 net.
☐ NEUTRAL	2	Veto-protected arms that never fired = zero cost, correct mechanism choice.

What the router got right, in plain terms

Strip out the drama and the router did its core job well: it kept the counter-trend faders benched almost the entire week (rgv_short fired only 3 times all week and abs_veto_short was correctly held out of the Friday chop after Wednesday's win), and it read the two big regime flips — Thursday's roll-over and Friday's reversal — closely enough to sit flat through the worst

churn. Both green days (Wed +\$638, Thu +\$1,172) happened under the router. Its recurring weakness is the same one all week and in the ledger: **over-toggling grind tick-to-tick** and re-arming momentum a beat too early on elevated-but-not-clean chop. The fix is discipline, not code — arm grind once on a genuinely strong trend, leave it, and wait for ER to actually climb before believing a break.

HONEST CAVEATS — READ BEFORE YOU TRUST ANY OF THIS. THIS IS ONE WEEK OF IN-SAMPLE PAPER TRADING, AND A LOPSIDED ONE: TWO-THIRDS OF THE WEEK'S +\$625 NET CAME FROM A SINGLE BEST-EVER DAY, AND A BIG CHUNK OF THE GREEN (+\$1,223 ACROSS 12 TRADES) WAS BOOKED BY YOUR OWN MANUAL CLAIMS, NOT BY THE EXITS FIRING ON THEIR OWN — SO THE AUTOMATED SYSTEM ALONE DID LESS THAN THE HEADLINE. THE PER-GATE READS ARE BUILT ON THIN COUNTS (RGV_SHORT FIRED 3 TIMES; EXHAUSTION'S NEW SCALP TRIAL ONLY 11), SO TREAT "ABS_VETO_SHORT IS A DOWN-LEG CATCHER" AS A ONE-WEEK OBSERVATION, NOT A PROVEN EDGE — IT RESTS HEAVILY ON A SINGLE +\$536 RUNNER. THE ROUTER'S SCORECARD IS SELF-GRADED. AND CAPITULATION_LONG (-\$139) AND THE NEW EXHAUSTION 0.5R/1.5R SCALP TRIAL (-\$48.5, FIRED OFF-REGIME INTO FRIDAY'S VIOLENCE) BOTH GO INTO PART 1.5'S REHAB BENCH RATHER THAN GETTING A FACE-VALUE BENCH HERE — THE HONEST "HOW COULD THIS WORK" COMES NEXT.

1.5 Live-Desk Rehabilitation — never bench a gate at face value

Garrath — this is the section you care about most, so here is the deal up front. Nothing on this desk gets binned because it had a red week. Every red gate and every leaky exit got walked through the full six steps — **what looked broken, what would the loss have been if the plumbing worked, is it the base / the exit / the signal, what fixes were tried (and the winners-kept test each one had to pass), does it survive stripping the best days and dropping any single day, and finally a verdict.** This week five targets went through the wash. **Three came out**

fixed, one is a chop-only specialist, one goes back to the bench to earn its keep in shadow, and nothing earned a retirement. Then — because the honest failures are the whole point — every dead end from this week's attempts is laid out in the graveyard at the bottom. We earn your trust by showing you the graves, not by hiding them.

THE ONE RULE BEHIND EVERYTHING BELOW

A fix that "wins" only by throwing away the big tail winners is a **fake win** and gets rejected on sight. Every survivor here had to keep the winners and cut the bleed. Every number is from a real tick-honest replay (\$2/pt, \$1.50 a round-turn per contract) on this week's tape plus, wherever possible, a second different regime — never a made-up figure, never a blanket cross-tape average (that averages the days a gate should be on with the days it should be off, and lies).

► **Read this section with REV3 open.** The grind_long rehab's headline f-figure was measured against a baseline that was never running. The honest before/after is **+\$1,552 → +\$2,830**, and the ATR floor should be **10, not 12**. The corrections are marked inline in red where they bite, and the full working is in [Part 3 — the Rev2/Rev3 corrections](#). Nothing has been deleted: where a claim was overturned, the original is still printed next to what replaced it.

The scoreboard — read this first

TARGET	WHAT WAS WRONG	LIVE / BEFORE	AFTER THE FIX	VERDICT
STOP_UNFILLED (the resting-stop exit)	Stops that never triggered → late flatten a few points past the stop	26 unfilled	0 unfilled, 103 clean stops	FIXED — DEPLOYED
grind_long (the gate)	Two floors were configured against it (neither was executing); it bought stretched tops	+\$1,552 -\$284	+\$2,830	FIXED
chandelier runner give-back (the exit)	One wide lock sprayed across all regimes; banked only 47% of peak	+\$2,932	+\$3,580	REGIME-DEPENDENT
capitulation_long (the gate)	Waiting for a delta-flip starved it; 2.0R gave the bounce back	-\$139	+\$434	FIXED (chop-only)

TARGET	WHAT WAS WRONG	LIVE / BEFORE	AFTER THE FIX	VERDICT
rgv_short (the gate)	No robust edge in any regime it can trade + far too few fires	-\$142	unproven	SHADOW

All "after" figures are tick-honest replays on the same tape as the "before". No target was truly retired this week — the closest, rgv_short, is unproven, not disproven, and goes to the bench to gather evidence rather than to the grave.

★ What got rehabilitated — the survivors

1. STOP_UNFILLED — FIXED and already live (the strongest win of the week)

This was flagged as "the single worst exit-mechanism bleed of the week" — four stops that never filled, -\$195. Two things turned out to be true, and both matter.

First, most of the headline is mislabelled. A stop that fires late is still mostly a normal stop-loss the losing entry earned anyway. Of the -\$195, only **-\$53 is the actual defect** (the extra few points lost while the trigger sat there); the other **-\$142 is ordinary stop-loss** the entries would have taken with a perfect stop. Across the whole archive it's the same shape — only 24% of the bleed was ever the bug.

WINDOW	STOP_UNFILLED	CLEAN STOP
before the fix ($\leq$ 07-28 08:34)	26	—
at / after the fix	0	103

Second, it is already fixed and verified in production. The winning change (d0218c6, live 07-28 08:34) checks every tick (~2s) and market-flattens at the stop, so stops no longer depend on the broker's trigger firing at all. The last unfilled stop in the entire book was 32 minutes before that deploy. On real ticks the 2s path actually beats a "perfect" stop by ~\$15 (fast bounces come back inside 2s). Two earlier attempts (concrete-Future

stops, fillable StopLimit) helped but did not kill it — shown honestly in the graveyard. **No action needed; do not revert.**

2. grind_long — FIXED at the entry, and it holds in a second regime

Your exact symptom — "it buys resistance and session-highs and churns the chop" — is real and now explained. The bleed is **100% the entry signal, not the exit**. Every dollar of profit lives in the exit tail; every dollar of loss lives in stopped entries:

WHY THE RAW NUMBER LIES	N	\$
STOP (the entries get stopped — the bleed is here)	53	−2,718
TARGET (Lot-A scalps)	9	+1,137
CHANDELIER + MANUAL_CLAIM (Lot-B tail)	16	+1,567

★ REV3 CORRECTION — 2026-08-01 — EVERY ER AND ATR FLOOR ON THIS DESK HAS BEEN SWITCHED OFF SINCE 2026-07-29

Nothing below that leans on a live ER or ATR floor is describing the desk you own. `tournament.step()` hands the sub-slot tag (`grind_long_A`) to `er_blocks()`/`atr_blocks()`, but `ER_FLOOR/ATR_FLOOR` are keyed on base names, so the lookup never matches and every floor silently answers "don't block". Dead since the dual-slot cutover at **2026-07-29 16:20:49 UTC**; last suppression line `07-29T13:29:04Z`, zero since. Five gates affected. The `_base()` helper that fixes it already exists at `tournament.py:86` and is applied correctly three times elsewhere — just not at lines 139 and 142.

Size it honestly: the circulating "7,734 suppressions" is a raw log-line count and the check re-fires every ~1s tick. De-duplicated it is **383 blocked signal episodes over 8 days** (140 ER + 243 ATR) — about 48 a day, not thousands. Real, but 44× smaller than advertised.

And it changes the conclusion, not just the footnote. Grind did not fire 52 of 56 times below its floor because "ER filters are fake on this tape" — a claim about the market. It fired because there was no floor — a claim about our code. Those point at completely different fixes. Both ER floors were even committed ~18 hours after the bug, so neither has ever blocked a single trade in production. Full working, per-gate verdicts and the mandatory ship order: [Part 3.6](#).

Two floors are configured against this gate — and **neither has executed since 07-29**, so they cannot have destroyed anything this week; as written they are backwards and too tight, which is exactly why they must not be re-wired at their current values. The **ER ≥ 0.35 floor is backwards** — a fast-slope gate catches a trend at birth, when the 30-min ER still reads the prior chop, so 26 of the 30 big tail-winners enter at low ER. The **ATR ≥ 24 floor is far too tight** for this regime. Swap both for a single "don't buy when already stretched" ceiling (`ext_atr ≤ 2.0`) and the gate flips right over:

★ REV3 CORRECTION — 2026-08-01 — READ THE N BEFORE YOU READ THE DELTA, AND USE ATR 10

This is not the same eleven trades made profitable. The recommended config trades ten times as often — it is a different, ten-times-larger book (−\$26/trade vs +\$25/trade). And the "before" row is a counterfactual: the rehab replay applied the floors as config (rehab_grind_long.py:124, :234) while the live desk has had them switched off since 07-29. Measured against what was really running, the honest before/after is **+\$1,552 → +\$2,830**, not −\$284 → +\$2,830 — about 59% of the advertised delta is the removal of filters that were already de-facto removed.

Ship ATR 10, not 12. Re-derived fine on 8 days of tick-honest replay: 9/10/11 is a genuine plateau where the blocked book is robustly negative; at 12 the blocked book flips positive and stays positive all the way up. Twelve is one notch the wrong side of the cliff and costs \$941 against 10. On every measure — net, per-signal, win%, strip-best-3, worst leave-one-day-out and both out-of-sample halves — ATR 10 + ext_atr 2.0 beats ATR 12 + ext_atr 2.0.

And the money is not bankable. Grind is a high-count, low-margin book: the best config breaks even at a 1.64× loss multiplier and the desk's standing finding is that live losses run 1.1–3.1× modelled. That break-even sits inside the band. The ceiling buys resilience (at 1.5× losses: capped +\$358 vs uncapped −\$1,449), not dollars. Grind is being made less bad, not good — which is why the play card for this now carries £0, not £700.

CONFIG (TICK-HONEST, SAME TAPE)	N	NET	\$/TRADE	WIN%	STRIP TOP-3
LIVE as written (ER≥.35 & ATR≥24) — a config that was never actually running	11	−\$284	−\$26	18%	−\$1,161
► REV3 — LIVE as actually running (floors dead since 07-29, ext_atr uncapped)	181	+\$1,552	+\$8.6	34%	—
RECOMMENDED (drop ER floor, ATR	111	+\$2,830	+\$25	34%	+\$1,172

CONFIG (TICK-HONEST, SAME TAPE)	N	NET	\$/TRADE	WIN%	STRIP TOP-3
24→12 10, ext_atr ≤ 2.0)					

It survives stripping the best trades, survives dropping any single day (no day carries it), and — the clincher — it **holds in a genuinely different regime**: on the V5 summer-chop archive the same fix turns −\$1,135 into +\$172. That is the multi-regime proof the 07-30 ER-floor note explicitly asked for, and it points the opposite way to that floor. The exit (Lot-A 2.5R + Lot-B wide lock-chandelier) is proven and stays untouched. **Honest caveat**: this is a positive-skew tail gate — the top ~10 trades carry it — so treat the policy as the edge, not any single number.

3. chandelier runner give-back — FIXED per regime

This is the desk's biggest gross winner (+\$4,774 on the CHANDELIER exit) but it **leaks the tail** — across all managed-profit exits the desk banks only **47% of peak**, leaving ~\$8,088 on the table. Your symptom trade reproduces exactly: 07-27 grind runner, peak 4.20R / \$208, banked +\$23 — gave back \$185.

The cause is one knob used everywhere: the firm lock at lock_r=6.0. But the typical MNQ runner peaks at ~4.5R — below 6.0 — so on most runs the lock never engages and the position rides the wide trail all the way back down. The fix is to make the lock regime-aware:

REGIME (ENTRY 30M ER)	INCUMBENT WIDE, IS→OOS /SIG	PROVEN FIX	ACTION
TREND (ER≥0.50)	+52.2 → +46.9	keep lock_r=6.0	keep the wide lock — it's the edge
BUILD (0.30-0.50)	+12.7 → −9.5 (overfit)	lock_r=3.0 (+4.6 OOS)	lower the lock
CHOP (ER<0.30)	+4.1 → −1.2	none robust	router benches

On the actually-taken live trades this regime policy reconstructs **+\$3,580 vs the incumbent's +\$2,932 (+22%)** — the entire gain in chop/build, trend left completely alone. It drops zero entries; it just banks the tail

earlier off-trend while keeping every trend monster on the wide trail. The machinery to deliver it already exists — it just isn't wired into the one regime-blind path (the dual-slot Lot-B "wide" tail on grind_long). **Caveat:** the trend leg is thin (n=4 live), so re-measure after the next clean-trend week; the build fix is the sturdier leg (robust 4/4 weeks + out-of-sample).

4. capitulation_long — FIXED, but it's a chop-only specialist

Live it shows −\$139 this week. Tick-honest and single-lot, the current config is already roughly breakeven (+\$24) — the extra loss is **execution, not signal**: the dual-slot A/B split doubled every chop stop (−\$201 on just 4 signals), plus some low-vol-grind fires the router should have benched. But there is a real config fix underneath, and it **reverses one of our own priors** — the 07-25 rehab installed require_flip=True as "the edge". Tick-honest data says otherwise:

CONFIG	N	NET	WIN%	STRIP BEST-3	WORST DAY-DROP
require_flip=True, 2.0R (current live)	11	+\$24	45%	—	—
require_flip=False, 1.0R (recommended)	32	+\$434	78%	+\$327	+\$220

Waiting for the delta to flip enters late and skips the bounces. Fading the flush itself and banking a fast 1R (78% of bounces reach 1R; only ~40% reach 2R) is the least data-starved, most robust config in the whole study — every single-day-drop stays strongly green. Restrict it to overnight (fade small flushes when the tape is quiet, bench the US cash session where a fader gets run over) and it's +\$422 at 81% win. **The hard caveat:** the tick tape is one chop week — there is no trend week in it to exercise the fader's predicted trend-day bleed. So ship the config, keep the router's down-trend and US benches (they are load-bearing, not optional), and confirm the benched side in shadow.

5. rgv_short — back to the bench, honestly

★ REV3 CORRECTION — 2026-08-01 — THIS VERDICT AND THE LIVE DESK DISAGREE RIGHT NOW

Live-state check, 08-01: rgv_short=on in data/gate_switches.env — verified in the file today. Every section of this report verdicts it SHADOW and Part 1 says the router "correctly kept it in the dugout", but it is armed. Bench it or change the verdict; do not leave both standing. This is the MONDAY #3 play.

This one did not get fixed, and the honest verdict is SHADOW, not retirement. On the clean tick tape it fires only **11 times in the whole week** and is net-negative at every exit setting (1.5R -\$32, 2.0R -\$142, 2.5R -\$253, chandelier -\$190). It fails every robustness cut — strip the single best day and it's -\$242; the whole book leans on one +\$80 single fire. Even firing with the down-trend it bleeds (aligned-DOWN -\$138).

ROBUSTNESS CUT	RESULT
strip best-1 day	-\$242
strip best-3 days	-\$402
drop 07-28 (the +\$80 day)	-\$223

The base config is fine, the exit is only a minor aggravator, and flow filters make it worse — it is a regime/data-starvation problem, not a broken config. Retiring 11 fires would be as unjustified as arming them. **Keep it benched in the live tournament, run the base variant in shadow to gather cross-regime evidence, and re-look Saturday when the count grows.**

The two live grind-exit shadow watches — the honest trying, running now

Runner-CLIP watch + 2R-partial smoothness watch

Two shadow ledgers are running live alongside the desk to test give-back fixes on real fills before anything gets wired in. Both are thin — 4 trades since 07-26, all clustered on 07-27, the same grind cluster as the symptom trade — so read these as an early signal, not a verdict.

Two-ratchet runner-clip watch (`two_ratchet_shadow.json`): shadows clipping a runner as it ratchets. Over the 4 trades, live booked $-\$149$ vs the shadow's $-\$27.5$. The entire gap is the one 4.2R runner — live banked $+\$23$, the shadow's rule banked $+\$144.5$ on the same ride. It hasn't taken a clip yet (0 clips), so this is one data point, but it's the same $\$185$ give-back story the chandelier rehab found.

2R-partial smoothness watch (`partial_shadow.json`): shadows taking a partial at 2R to smooth the curve. Baseline $-\$278$ vs partial $-\$213.5$, a $+\$64.5$ improvement — all of it, again, on the single 4.2R runner (the three stops were untouched, delta $\$0$).

SHADOW WATCH	N	LIVE / BASELINE	SHADOW VARIANT	IMPROVEMENT	WHERE IT CAME FROM
two-ratchet runner-clip	4	$-\$149.0$	$-\$27.5$	$+\$121.5$	all on the one 4.2R runner
2R-partial smoothness	4	$-\$278.0$	$-\$213.5$	$+\$64.5$	all on the one 4.2R runner

Both watches point the same way as the chandelier rehab — the leak is the runner give-back, and both fixes only touch the runner (they do nothing to the stops, which is correct). But four trades on one day is nowhere near enough to wire anything in. **Keep both in shadow; re-read once a real trend week lands and the count grows.**

★ The graveyard — every dead end from this week's attempts

WHY THIS IS HERE

The failures are the point. A rehab that only reports its wins is hiding the fake wins. Everything below was tried this week, scored honestly, and rejected — most of them because they raised the headline number by quietly throwing away the winners, which is exactly the trap we refuse to fall into.

Graves from this week's rehabs

TARGET	WHAT WAS TRIED	WHY IT DIED
grind_long	Every ER floor (0.15 → 0.45)	Backwards — drops the low-ER tail-runs; $ER \geq 0.20 = -\$1,593$
grind_long	Range-break-only filter	Breakouts fail; the edge is the ride, not the poke ($-\$366$)
grind_long	Strong-slope ≥ 0.7	Strong slope = late entry ($-\$1,540$)
grind_long	$ATR \geq 20 / 24$ concentrators	High net only by binning genuine tail-winners — a fake win
capitulation_long	require_flip=True (our own 07-25 prior)	Falsified tick-honest — starves the gate, skips the bounces
capitulation_long	The "+\$492, n68" looser-climax book	Biggest headline, but fires into 07-29 dead chop for $-\$222$ — don't stack filters on one week
capitulation_long	Widen exit past 1R / raise ATR floor	1.5R falls off a cliff; the winners live in the low-ATR overnight band
rgv_short	The "+\$304 cross-tape" and "+\$663 aligned-DOWN" edges	Both artifacts of re-firing while a position was open — the blanket bug; collapse to $-\$142$ single-position
rgv_short	Flow-confirm, turn_atr 0.5, 1.5R target, ext_min 2.0	All NULL — filtering does not earn the gate; base is already the best
chandelier exit	Give-back overlay on the wide trail (+gb1.0/1.5)	Won the in-sample cell, then failed out-of-sample and strangled the trend riders — a cell-fit fake
chandelier exit	Blanket lock_r=3.0 / fixed-2.5R everywhere	Robust off-trend but a grave on trend (caps the monster runners)
chandelier exit	Fix thrust/absveto chop bleed with an exit	No exit rescues a chop-momentum bleed — that's the router's job
STOP_UNFILLED	"Bench abs_veto overnight" (where 3 of 4 lived)	Textbook fake win — drops +\$1,003 of winners to save $\sim \$50$ of defect
STOP_UNFILLED	Trust the broker's native trigger; concrete-Future stop; StopLimit	Honest partials — each helped (18%→11%) but none killed it; only the per-tick guard did

THE HONEST SMALL PRINT

The tick tape is roughly one summer week (plus a second archive regime where we could reach for one). It is a chop-heavy tape with no violent US-open week and, for the faders, no clean trend week to exercise their predicted trend-day bleed. So: grind_long's fix is cross-regime proven and safe to ship; the chandelier and capitulation fixes are strongly proven in regime but lean on thin trend-side evidence and stay paired with router benches; rgv_short and both shadow watches are simply too thin to conclude and go to the bench / shadow to earn it. Nothing here is a cross-regime law — the ones that ship (grind entry ceiling, STOP_UNFILLED guard) are the ones that survived a second regime or are already running clean in production.

PART 2 The Shadow Desk & Promotion

This section was built from scratch on this week's numbers — no gate, candidate or verdict was inherited from last week's report. I ranked the whole shadow board on honest tick-repriced dollars (what a trade would really have made after slippage and fees, never the optimistic sim price), and let the promotion pick fall out of that ranking rather than deciding it in advance. Then I ran the full robustness battery on whatever came top, presented the two grind-exit shadow ledgers exactly as they stand tonight, and re-proved the exit ladder on the week's tick path per regime. Where the sample is thin I say so plainly — the shadow board nominates, the paper desk judges.

► **Read this section with REV3 open.** Four things in here were overturned: abs_veto was never a promotion (it shipped 07-25), the "+\$326 over 47" figure is wrong, the BIG-TREND cell is n=24 not n=5, and the "exit is not the lever" instruction is wrong for the live abs_veto_long cell. The corrections are marked inline in red where they bite, and the full working is in [Part 3 — the Rev2/Rev3 corrections](#). Nothing has been deleted: where a claim was overturned, the original is still printed next to what replaced it.

The shadow board, ranked on honest money

★ REV3 CORRECTION — 2026-08-01 — SHADOW RUNS 1 LOT, THE LIVE SCALE-OUT RUNS 2 — NEVER COMPARE THE RAW DOLLARS

Standing note for every shadow-vs-live comparison in this section: **the shadow runs 1 lot; the live scale-out runs 2**. No raw shadow-vs-live dollar comparison is valid without dividing the live number by signals, not lots. That single omission is what produced the \$326-vs-\$776 confusion below.

Every strategy that runs observe-only alongside the live desk, sorted by real (tick-repriced) dollars. Two windows: all-time (the ~two-week summer archive, 07-16 to 07-31) and this week. The green rows are the ones actually making money on honest fills; the red block at the bottom is where the dead ideas live.

STRATEGY (ALL-TIME)	N	NET \$	WIN%	\$/TRADE
abs_veto_55s promotion pick	238	+2,668	47	+11.2
grind_fast variant of live grind_long	732	+2,052	36	+2.8
abs_veto_50s sibling	259	+1,968	44	+7.6
thrust_cont	276	+711	37	+2.6
abs_veto_60s sibling	80	+425	44	+5.3
thrust_aligned	241	+396	37	+1.6
thrust_short_absveto55 live-slot mirror	53	+318	45	+6.0
rg_short_025_flow50	23	+245	52	+10.6
rg_short_025_raw	30	+184	47	+6.1
thrust_loose un-vetoed base	424	+131	36	+0.3
thrust_short_raw	94	-137	33	-1.5
rg_long_20_t20 · rg_long_fast · rg_long_fast_v	—	-546 / -2,281 / -1,438	29-31	-7 to -11
exhaustion_rev fixed-target shadow	321	-1,187	37	-3.7
chand_k20 / k25 (raw chandelier)	362 / 352	-1,588 / -1,544	43 / 41	-4.4
capit_loose unfiltered fader	238	-5,250	17	-22.1
STRATEGY (THIS WEEK ONLY)	N	NET \$	WIN%	\$/TRADE
	240	+2,088	39	+8.7

STRATEGY (THIS WEEK ONLY)	N	NET \$	WIN%	\$/TRADE
grind_fast				
variant of live gate				
abs_veto_55s	114	+1,394	49	+12.2
promotion pick				
abs_veto_50s	124	+1,107	46	+8.9
cb_thrust	130	+275	35	+2.1
capit_loose	62	-1,455	21	-23.5
rg_long_fast · rg_long_20_t20	99 / 14	-667 / -520	32 / 14	-6.7 / -37
thrust_short_raw	81	-463	31	-5.7

WHO ACTUALLY EMERGES. THE SINGLE BIGGEST HONEST EARNER ON THE BOARD, BOTH ALL-TIME AND THIS WEEK, IS ABS_VETO_55S. THE ONE ABOVE IT THIS WEEK (GRIND_FAST) AND THE FAMILY'S SIBLINGS (ABS_VETO_50S/60S) ARE ALL JUST RESEARCH VARIANTS OF GATES THE DESK ALREADY RUNS — GRIND_FAST IS THE UNCAPPED-CHANDELIER COUSIN OF THE LIVE GRIND_LONG, THE 50S/60S ARE THE SAME VETO WITH A SHORTER/LONGER CONFIRM WINDOW. AMONG STRATEGIES THAT ARE NOT YET ON THE LIVE DESK, ABS_VETO_55S IS #1 IN BOTH WINDOWS BY A CLEAR MARGIN. IT NOMINATES ITSELF. TWO USEFUL FACTS FALL STRAIGHT OUT OF THE FAMILY ORDERING: 55S > 50S > 60S IN HONEST DOLLARS, SO THE 55-SECOND CONFIRM IS THE SWEET SPOT (NOT AN ARBITRARY PICK), AND THE RAW UN-VETOED THRUST IT IS BUILT ON (THRUST_LOOSE) MAKES ESSENTIALLY NOTHING (+\$131 OVER 424 TRADES, +\$0.3 EACH) — SO EVERY DOLLAR ABS_VETO EARNIS IS THE VETO DOING THE WORK, NOT THE UNDERLYING THRUST.

The promotion battery — abs_veto_55s, run cold on this week's data

abs_veto_55s is thrust_loose plus a 55-second continuation confirm: it takes a raw thrust signal, waits 55 seconds, and only fires if the burst is still going — which throws away roughly 44% of the raw thrusts as fake-outs. Here is the full robustness read, not just the headline number.

TEST	RESULT	VERDICT
Headline (all-time)	238 trades, +\$2,668, 47% win, +\$11.2/ trade, worst day -\$148, max drawdown -\$441	strong

TEST	RESULT	VERDICT
Spread across days	9 of 12 days green. Best +\$761 (07-23), worst -\$148 (07-22). Not one lucky session.	broad
Concentration check	Top 3 days (07-28 +\$791, 07-23 +\$676, 07-17 +\$581) = \$2,048 of the \$2,668 (77%). Strip all three and the remaining nine days still net +\$620.	survives strip-3
Regime split (chop vs trend, by entry ER)	Chop (ER<0.18): 112 tr, +\$872, +\$7.8/tr. Trend (ER≥0.18): 126 tr, +\$1,796, +\$14.3/tr. Green in both, roughly twice as good in trend.	both regimes
Walk-forward (first half vs second)	First 119 trades +\$1,250; second 119 trades +\$1,418. Slightly better late — no decay.	stable
Head-to-head vs un-vetoed thrust_loose	abs_veto_55s +\$2,668 vs thrust_loose +\$131 on the same board. The veto's incremental edge is +\$2,538.	the veto is the edge

And the operator's key point from last week held up when re-tested cold: **it is a genuine two-sided gate and both sides earn their keep.**

Correction (applied 08-01): an earlier draft of this section said the live desk runs only the short side via thrust_short and that the long side was still on the table. That was a prior-week fragment and it is wrong. data/gate_switches.env has abs_veto_long=on AND abs_veto_short=on — abs_veto is **already live two-sided**, and there is no gate called thrust_short in the live tournament at all. So the promotion decision is not "should we add the long side"; it is already added. The open question is the gap: live abs_veto_long made **+\$59.5 over 59 lots (43 signals)** this week
► REV3: the folded-in correction printed ~~+\$326 over 47 trades~~, which is wrong by \$266.5 and contradicts Part 1's own card table in this same document. Straight from data/gazbot7.db it is +\$59.5 / 59 lots; no window produces \$326/47. while shadow abs_veto_55s LONG made **+\$776 over 66**. Reconciling that — config drift between the live gate and the shadow spec (confirm window, ER floor, amp floor, exit Rs) — is the real Saturday job. **► REV3: it is now done.** Of the \$609.5 gap, **82% is entry** — 26 long signals the router had benched, worth +\$397 in shadow — plus \$107 of window mismatch and \$273.5 of STOP_UNFILLED. Net of fees and that stop bug the live exit design is about **+\$67 ahead** of the shadow. **The bench cost the money, not the exit.**

ABS_VETO_55S SIDE	N	NET \$	WIN%	\$/TRADE	MAX DD	VS RAW THRUST
LONG (all-time)	122	+1,187	47	+9.73	-456	+\$1,563 better
SHORT (all-time)	116	+1,482	47	+12.77	-331	+\$974 better

ABS_VETO_55S SIDE	N	NET \$	WIN%	\$/TRADE	MAX DD	VS RAW THRUST
LONG (this week)	66	+776	50	+11.75	-310	—
SHORT (this week)	63	+659	44	+10.46	-331	—

The short-only live-slot mirror A/B says the same thing in the live slot's own shape (short-only + chandelier): **thrust_short_absveto55 made +\$318 (53 tr, 45% win)** against **thrust_short_raw's -\$137 (94 tr, 33% win)** — the veto turns the desk's live short thrust from a slow bleed into a real earner.

VERDICT: ALREADY LIVE TWO-SIDED — RE-VALIDATED COLD. NO ACTION
~~VERDICT: PROMOTE, AND PROMOTE IT TWO-SIDED.~~ ► REV3: THERE IS NOTHING TO PROMOTE. ABS_VETO_LONG AND ABS_VETO_SHORT HAVE BOTH BEEN ARMED SINCE THE 07-25 ROSTER CHANGE — VERIFIED IN GATE_SWITCHES.ENV TODAY. WHAT THE COLD RE-VALIDATION EARNIS IS CONFIDENCE IN A GATE THAT IS ALREADY RUNNING, PLUS EXACTLY ONE CONFIG CHANGE: THE LONG-SIDE EXIT CELL (BELOW). EVERY ROBUSTNESS CUT IT CAN BE PUT THROUGH — PER-DAY SPREAD, STRIP-THE-BEST-3, CHOP-VS-TREND, BOTH WALK-FORWARD HALVES, HEAD-TO-HEAD VS ITS OWN UN-VETOED BASE — COMES BACK POSITIVE, AND BOTH DIRECTIONS EARN INDEPENDENTLY. HONEST CAVEATS: IT IS STILL ONLY A ~TWO-WEEK, SINGLE-SUMMER-REGIME SAMPLE, AND THREE GOOD DAYS CARRY MOST OF THE TOTAL (THOUGH IT STAYS GREEN WITH THOSE THREE REMOVED). THE REAL OUT-OF-SAMPLE TEST IS THE FORWARD PAPER RUN — THE SHADOW NOMINATES, THE PAPER DESK DECIDES. BUT OF EVERYTHING ON THE BOARD, THIS IS THE ONE THAT HAS EARNED A LIVE SLOT. ~~...AND CLIPPING IT TO SHORT-ONLY LEAVES ROUGHLY HALF ITS PROVEN MONEY ON THE FLOOR.~~ ► REV3: NOTHING HAS CLIPPED IT TO SHORT-ONLY. THAT SENTENCE WAS WRITTEN AGAINST A ROSTER THAT CHANGED A WEEK BEFORE THIS REPORT.

Relegation view — tune the SIDE, never bin a direction

The standing rule is that a gate runs both ways with its two sides tuned independently, so nothing here is a recommendation to retire a direction — it is a read on which side of which gate needs what, from this week's fresh

per-side numbers. Every "bleed" below is a side-and-regime problem, not a broken gate.

GATE	THIS WEEK, PER SIDE (HONEST \$)	DIAGNOSIS	ACTION
rgv	LONG bleeds hard (rg_long_fast -\$667, rg_long_20_t20 -\$520). SHORT holds up (rg_short_025_flow50 +\$71, rg_short_025_raw +\$28, rg_short_050 +\$11).	Classic two-sided split — the sides want opposite treatment. Short has a stable tight-ext green zone; long has no filter that keeps it green, it's a router problem (don't fade into a proven up-move).	tune side: rgv SHORT tight-ext base; rgv LONG = router-only
capitulation	Unfiltered fader (capit_loose) -\$1,455 this week, -\$5,250 all-time, 17-21% win.	The loose, unfiltered form is the worst thing on the board. The deployed capitulation_long uses require_flip, which is the actual fix — the bleed is the unfiltered version, not the gate.	keep require_flip + router guard
thrust (short)	Raw short (thrust_short_raw) -\$463 this week, 31% win.	Not a direction to bin — the same short direction with the 55s veto (thrust_short_absveto55) is +\$318. The fix is the entry veto, which is the promotion above.	fix = the abs_veto veto
exhaustion (short)	Fixed-target shadow (exhaustion_rev) -\$302, but that shadow uses a fixed 12pt target it never trades to. LIVE exhaustion_short is the desk's steadiest fader cleanest fader — though its live green is hand-claim dependent (-\$56 over 44 automated fires) (see the exit lab below).	Do not judge this gate by its fixed-target shadow — the live chandelier version rides the fade body. Judge it on live fills.	keep, judge live

NOTHING GETS RELEGATED. EVERY RED LINE ABOVE IS A SIDE THAT NEEDS A FILTER, A REGIME GUARD, OR THE ROUTER — NOT A DIRECTION TO SWITCH OFF. THAT IS THE WHOLE POINT OF THE TWO-SIDED RULE, AND THIS WEEK'S PER-SIDE NUMBERS BACK IT: RGV'S SHORT SIDE IS QUIETLY GREEN WHILE ITS LONG SIDE BLED, ON THE SAME GATE, IN THE SAME DAYS.

Grind exit shadow — two-ratchet runner-clip watch

Background: grind_long deployed the pure-6.0 threshold-chandelier as its exit. A second "mop-up" ratchet was tested three ways and rejected every time, because the reverser money is tangled up with the runner-retrace zone it would clip. It stays dead but is shadowed to catch the one thing the flat summer archive can't produce: a real runner getting clipped on a genuine trend day, since the big $rMFE \geq 4R$ runners carry the bulk of grind's edge.

WINDOW 07-26 → 08-01	VALUE
grind_long live fills in window	4 (all on 07-27)
Runners ($rMFE \geq 4R$)	1
RUNNER-CLIPS fired	0
Live pure-6.0 net	-\$149
Shadow two-ratchet net	-\$28 (+\$122 vs live)

EVERY TRADE IN THE WINDOW	ATR	RMFE	LIVE \$	SHADOW \$	EXIT
07-27 16:05	25.6	0.14R	-52	-53	STOP
07-27 16:06	25.0	1.79R	-52	-52	STOP
07-27 16:28	32.9	0.14R	-68	-68	STOP
07-27 18:54 the runner	24.8	4.2R	+23	+145	CHANDELIER

VERDICT: NO CLIP FIRED, AND FAR TOO THIN TO REOPEN THE CASE.
STAYS DEAD, KEEP SHADOWING THE HONEST FACTS: GRIND BARELY TRADED LIVE THIS WINDOW (4 FILLS, ALL ON ONE DAY 07-27), SO THIS IS A SINGLE-DAY, ONE-RUNNER READ, NOT THE "GENUINE TREND WEEK" THAT WOULD COUNT AS EVIDENCE. ON THE ONE RUNNER IT SAW, THE TWO-RATCHET DID NOT CLIP IT — IT ACTUALLY RODE IT FOR MORE (\$145 SHADOW VS \$23 LIVE, +\$122). THAT IS MILDLY INTERESTING (THE MOP-UP ARM ADDED TO A RUNNER RATHER THAN CUTTING IT), BUT ONE RUNNER IS NOTHING TO DECIDE ON. THE FINDING THAT WOULD KILL IT (A REAL CLIP) DID NOT FIRE; THE EVIDENCE THAT WOULD REVIVE IT (A CLEAN TREND WEEK WITH SEVERAL UNTOUCHED RUNNERS) HASN'T ARRIVED. IT STAYS DEAD AND SHADOWED UNTIL A REAL TREND WEEK GIVES IT A PROPER TEST.

Grind exit shadow — 2R-partial smoothness

The 2R-partial is the only exit lever that survived earlier scrutiny: Lot A banks a 2R scalp for certainty, Lot B rides the pure-6.0 chandelier for the tail. Because grind is base-size 2 it's free to run. It was framed as a variance lever (smoother equity for a small mean give-up), not a mean-raiser — so it's the operator's call, not a pass/fail.

METRIC (07-26 → 08-01)	BASELINE (DEPLOYED 2-LOT CHANDELIER)	2R-PARTIAL (SHADOW)
Net \$	-278	-214 (+\$64)
Mean Δ per trade (partial – base)	—	+\$64.5 (all on the 1 runner)
Daily std \$	0 (1 day)	0 (1 day)
Max drawdown \$	-278	-214
Worst day \$	-278	-214
Green-day % / green-week %	0 / 0	0 / 0

Where the +\$64 came from, trade by trade: on the three losers the partial is identical to baseline (both stop out); the whole difference is the one runner, where banking Lot A at 2R (+\$130.5) beat letting both lots ride the chandelier back down (+\$66).

VERDICT: NO NEW EVIDENCE TO MAKE THE CALL ON.

OPERATOR'S CHOICE, UNCHANGED THE SMOOTHNESS CASE IS A CLAIM ABOUT DAILY VARIANCE AND DRAWDOWN ACROSS MANY DAYS — AND THIS WINDOW HAS ONE GRIND DAY AND FOUR TRADES, SO DAILY STD IS ZERO AND THE WHOLE VARIANCE THESIS IS UNTESTABLE RIGHT NOW. WHAT WE CAN SAY: THIS WEEK THE PARTIAL EDGED THE BASELINE BY \$64, ENTIRELY BY SCALPING THE SINGLE RUNNER RATHER THAN RIDING IT. THAT IS THE OPPOSITE OF A MEAN GIVE-UP (THIS WEEK IT WAS A MEAN GAIN), PURELY BECAUSE THE LONE RUNNER TOPPED OUT AT 4.2R AND THE CHANDELIER GAVE THE TAIL BACK. ONE DAY PROVES NOTHING EITHER WAY; THE MEAN-VS-SMOOTHNESS CHOICE WAITS FOR A MULTI-DAY TREND SAMPLE.

The desk's live control is now two files: `gate_switches.env` decides which gates are on, and `exit_overrides.json` decides the exit style per gate. This lab does two jobs — grade the live `exhaustion_short fade-scalp` trial on its actual fills, and re-prove the operator's regime-graded exit ladder by repricing every gate's real entries on the forward tick/5s path, segmented by regime rung, so the deployable `exit_overrides` policy is proven rather than guessed.

(a) The live `exhaustion_short fade-scalp` trial — on its actual fills

Deployed ~07:55 UTC on 07-31: Lot A banks at 0.5R, Lot B at 1.5R, replacing the prior `A@2.5R + chandelier`. Here is literally every leg it has filled since.

TIME (UTC)	LEG	EXIT	\$
07-31 13:52	A (0.5R)	TARGET	+51.5
07-31 13:52	B (1.5R)	STOP	-123.5
07-31 18:41	A (0.5R)	STOP	-27.5

VERDICT: KEEP RUNNING — TWO SIGNALS IS NOT A GRADE.

TOO THIN TO JUDGE THE TRIAL HAS FIRED EXACTLY TWO SIGNALS (THREE LEGS) FOR A NET OF **-\$99.5**. THE ONE THING YOU CAN SEE: ON THE 13:52 SIGNAL THE FAST 0.5R LOT A BANKED (+\$51.5) WHILE THE 1.5R LOT B RODE PAST ITS CUSHION AND STOPPED (-\$123.5) — I.E. THE FADE REVERSED AFTER THE QUICK SCALP, WHICH IS EXACTLY THE CASE THE "BANK A EARLY" LOGIC IS MEANT FOR, EXCEPT LOT B STILL GAVE IT BACK. THE 18:41 SIGNAL STOPPED ON THE A LEG. ~~FOR HONEST COMPARISON THE PRIOR ERA (07-27 07-31, A@2.5R + CHANDELIER, 37 TRADES) NETTED +\$365 WITH A HEALTHY 11 TARGETS / 10 CHANDELIERS / 16 STOPS MIX.~~ ► **REV3: THE PRIOR ERA WAS NOT THE SAME MACHINE, AND THAT MISLABEL HIDES THE WHOLE CAUSAL STORY. UNTIL WED 07-29 19:10 UTC EXHAUSTION RAN A SINGLE SLOT ON THE REGIME-3 ADAPTIVE CHANDELIER: 37 AUTOMATED FILLS, +\$365, PF 1.46, GREEN ON ALL 3 OF ITS DAYS. THE SCALE-OUT SLATE THAT REPLACED IT (A@2.5R + WIDE LOCK-CHANDELIER) IS -\$421 OVER 7 FILLS AT 14% WIN. THE GATE DID NOT BREAK; ITS EXIT WAS SWAPPED**

FOR A TREND-RIDING ONE, AND A FADER DOES NOT GET 2.5R IN A WHIPSAW.

★ REV3 CORRECTION — 2026-08-01 — THE RS CAN BE GRADED EVEN THOUGH THE LIVE TRIAL CANNOT

A 64-ENTRY 250 MS SWEEP DOES GRADE THEM: THE LOADED 0.5R LOT A IS THE WORST SCALP RUNG ON THE LADDER (+\$306); THE PLATEAU IS 0.75-1.0R (+\$510 / +\$450); AND LOT B'S STEADIEST CONFIG BY A DISTANCE IS A K1.5 CHANDELIER (+\$502, THE HIGHEST LEAVE-ONE-DAY-OUT FLOOR IN THE WHOLE SWEEP AT +\$330, WORST DAY -\$25). THAT IS THE MONDAY #4 PLAY.

BUT THE SWEEP'S REAL MESSAGE IS THAT NO EXIT IS ROBUSTLY GREEN IN VIOLENT WHIPSAW, WHERE THIS GATE TOOK 28 OF ITS 62 ENTRIES — THE K1.5 CHANDELIER IS THE WORST CONFIG IN THAT BUCKET AT -\$202. KEEPING IT BENCHED THERE IS WORTH MORE THAN ANY R.

METHOD CAVEAT ADDED 08-01: THIS CELL CAME FROM A PER-SIGNAL SWEEP ON OVERLAPPING ENTRIES — THE SAME ACCOUNTING THAT WAS JUST SHOWN TO SWING GRIND'S EXIT VERDICT BY \$5,076 ([PART 3.7](#)). NOT REFUTED; RE-RUN IT UNDER A SEQUENTIAL ONE-POSITION-PER-SLOT SIM BEFORE TRUSTING THE SIZE OF THE GAIN. CHEAP TO REVERT EITHER WAY.

(b) Proving the exit ladder — best exit per gate × regime rung

I repriced every gate's real entries on the forward 5-second path, segmented by the operator's four rungs (BIG-TREND = 30-min $ER \geq 0.50$, MED-TREND 0.30-0.50, SCALP-CHOP otherwise, STAY-OUT = violent

ATR≥22), sweeping a full grid of Lot-A × Lot-B scalp-R combos plus wide/tight chandelier, and scored each config only on its home segment. The table is the proven best exit vs the ladder guess.

GATE FAMILY	RUNG (N)	LADDER GUESS	PROVEN BEST (MEAN \$/TR)	VERDICT
grind (momentum)	MED-TREND (78)	1.5 / 2.5	scalp 0.5 / 1.0 (+\$10.3, 59% win)	guess wrong: go TIGHT
grind (momentum)	SCALP-CHOP (542)	1 / 2	scalp 0.5 / 1.0 (−\$0.5, 53% win)	tight, ~flat
grind (momentum)	BIG-TREND (24 5)	2.5 / wide-chandelier	wideChand ×2 (+\$12.0) — n=24	DIRECTION SUPPORTED (bootstrap P=0.95, plateau 2.5–3.5R, both time-halves agree) · MAGNITUDE UNPROVEN (strip-3 takes +\$62 to +\$10)
grind (momentum)	STAY-OUT (107)	flat	every exit loses (best −\$3.7)	stay flat, confirmed
exhaustion (fader)	SCALP-CHOP (216)	0.5 / 1.5	tightChand ×2 (+\$2.2, 37% win)	tight chandelier > scalp
exhaustion (fader)	MED-TREND (90)	0.5 / 1.5	scalp 1.5 / 2.5 (+\$3.8, 46% win)	wider than the guess
exhaustion (fader)	BIG-TREND (15)	tight-chandelier	tightChand ×2 (+\$26.7, 53% win)	holds (thin)
thrust / abs_veto raw	every rung (25-333)	1.5 / 2.5	bleeds in ALL: BIG −\$13, MED −\$36, CHOP −\$39, STAY −\$67	exit can't save it
rgv (fader)	0.5 / 1.5	every rung red (best −\$6.9 MED, −\$8.7 CHOP, −\$27 BIG)	—	exit can't save it
capitulation (fader)	0.5 / 1.5	MED −\$19, CHOP −\$20; only tiny BIG (n=8) green	—	exit can't save it

WHAT SURVIVED, IN PLAIN ENGLISH. TWO GATES RESPOND TO THE EXIT AND THREE DO NOT.

- GRIND WANTS A TIGHT SCALP, NOT THE LADDER'S WIDER GUESS — IN MEDIUM TREND, 0.5/1.0 MADE +\$10.3/TRADE AT A 59% HIT-RATE; THE GUESSED 1.5/2.5 WAS ONE OF THE WORST CELLS. GRIND BANKS FAST OR GIVES IT BACK.
- EXHAUSTION IS THE ONE CLEAN FADER — A TIGHT CHANDELIER IN CHOP (+\$2.2) AND IN BIG TREND (+\$26.7 ON 15), AND A SLIGHTLY WIDER 1.5/2.5 SCALP IN MEDIUM TREND (+\$3.8), NOT THE 0.5/1.5 GUESS.
- THE RAW, UN-VETOED VERSIONS OF THRUST, RGV AND CAPITULATION BLEED IN EVERY REGIME NO MATTER WHAT EXIT YOU PUT ON THEM. THAT IS THE HEADLINE FINDING: FOR THOSE THREE, THE EXIT IS NOT THE LEVER ► REV3: NOTE THE POPULATION. THIS ROW IS THE RAW UN-VETOED THRUST, AND THE REPORT THEN APPLIES THE CONCLUSION BY NAME TO ABS_VET0, WHICH IS A DIFFERENT, FILTERED BOOK. THAT

SAMPLE-SWITCH INSIDE ONE SENTENCE IS WHAT PRODUCED THE WRONG INSTRUCTION BELOW. — THE LEVER IS THE ENTRY (THE ABS_VETO VETO IS LITERALLY THE PROOF OF THAT: SAME THRUST SIGNAL, FILTERED ENTRY, TURNS -\$137 INTO +\$318) AND WHICH GATES YOU RUN AT ALL VIA GATE_SWITCHES.ENV.

★ CRITICAL CAVEAT — BIG-TREND IS STILL NOT PROPERLY TESTED. THE TWO-WEEK V7 TAPE WAS CHOP-DOMINATED, SO THE BIG-TREND RUNG IS TINY IN EVERY FAMILY (GRIND N=24 ~~N=5~~, EXHAUSTION/RGV N=15, CAPITULATION N=8; THRUST THE ONLY SEMI-REAL ONE AT N=25 AND IT BLED). THE "RIDE" RUNG — GRIND 2.5-3.5R + WIDE LOCK-CHANDELIER ON A MONSTER UP-DAY — IS DIRECTION SUPPORTED, MAGNITUDE UNPROVEN ~~EFFECTIVELY UNPROVEN (GRIND'S BIG CELL IS 5 TRADES)~~. ► REV3: THE CELL IS N=24, NOT 5. FIVE IS THE COUNT OF LIVE GRIND_FAST SHADOW ENTRIES IN AN $ER \geq 0.50$ MINUTE; THE LADDER VERDICT WAS COMPUTED ON A 24-ENTRY TICK-TAPE REPLAY. THE CAVEAT QUOTES THE LIVE NUMBER INSIDE A CAVEAT ABOUT THE REPLAY NUMBER. ~~I DID NOT MANUFACTURE A BIG-TREND SAMPLE FROM THE RETIRED V5 ARCHIVE: IT IS A DIFFERENT DESK, CONTRACT ERA AND GATE SET, AND REPRICING V7 GATE ENTRIES ON IT WOULD BE A FABRICATED COMPARISON, NOT AN HONEST ONE... GRIND'S BIG-TREND EXIT IS A GUESS WEARING AN N=5 HAT.~~

★ REV3 CORRECTION — 2026-08-01 — THAT REFUSAL WAS FALSE, AND THIS LAB HAD ALREADY USED THE ARCHIVE

ELEVEN OF THE 24 BIG-TREND ENTRIES ALREADY COME FROM THAT PRE-CAPTURE MNQ ARCHIVE (TICKS.DB, 07-05→07-15). THE LAB TAPE STARTS 07-05, NOT 07-15. AND THE "DIFFERENT CONTRACT ERA" OBJECTION DIES ON THE SEAM: LAST ARCHIVE TICK 29,623.00 AT 15:35:29.973, FIRST CAPTURE BAR OPENS 29,622.75 AT 15:35:30, $R=0.9998$ ACROSS 2,925 OVERLAPPING MINUTES. SAME INSTRUMENT, SAME FRONT MONTH. A REPORT THAT INVENTS A PRINCIPLED REFUSAL FOR TAPE IT SILENTLY CONSUMED IS WORSE OFF THAN ONE THAT NEVER MENTIONED IT.

WHAT IS GENUINELY MISSING IS N, NOT HONESTY. THE PRE-REGISTERED GRADE NEEDS 39 PAIRED ENTRIES (FROM A PAIRED SD OF \$247 ON A \$79.4 DIFF) AND WE HAVE 24. GRADE DATE FRIDAY 2026-08-21, HARD BACKSTOP 09-11.

AND THIS RUNG IS LIVE TODAY, NOT AN UNTURNED STONE. DATA/EXIT_OVERRIDES.JSON KEYS ONLY EXHAUSTION_SHORT, ABS_VETO_LONG AND ABS_VETO_SHORT. GRIND_LONG HAS NO ENTRY, SO IT INHERITS THE BUILT-IN _BIG_RUN DEFAULT, WHICH IS LOT A @ 2.5R + LOT B WIDE LOCK-CHANDELIER. IT IS NOT A STONE LEFT UNTURNED; IT IS AN UNHEDGED LIVE POSITION.

The deployable exit_overrides policy (what this proves out to)

Reading the survivors straight into the per-regime policy — only the cells with real n and a real edge become overrides; the rest default or are handled by turning the gate off:

GATE	BIG-TREND	MED-TREND	SCALP-CHOP	STAY-OUT
grind_long	wide-chandelier — LIVE TODAY by default, n=24, under pre-registered test to 2026-08-21 (unproven, n=5)	scalp-0.5 / 1.0 ► WITHDRAWN	scalp-0.5 / 1.0 ► WITHDRAWN	FLAT (gate off)
► REV3 — grind_long, the honest answer	Swept properly: 130 configurations on 3,064 tick-repriced entries over 16 sessions, under a sequential one-position-per-slot sim (the paired same-entry method this lab used stacks overlapping entries and scored the loaded default at +\$2,961; sequential scores the same config -\$2,115 — a \$5,076 swing from methodology alone). The {0.5, 1.0} cell above is WITHDRAWN: it is -\$1,127 worse than what is loaded today and fails leave-one-day-out, strip-best-3, the bootstrap and the out-of-sample sign test. Not one exit setting makes this gate green — the best of 130 still loses \$414. If you want the marginal move it is {"a_r": 3.0, "b": "wide"}, +\$617 over 16 sessions ≈ +\$10/session at the live fill rate, and it is damage control, not an edge. The real lever is arming, and it is 3-5× bigger. Full sweep: Part 3.7 .			
grind_long — MED- TREND, n=326	The biggest evidenced exit gap on the desk and it needs no new tape: the live A2.5+WIDE default makes +\$8.49/trade against TIGHT×2's +\$17.60 — a \$9.11 × 326 gap. Better evidenced than the BIG-TREND question this section spends four hundred words on.			
exhaustion_short	tight-chandelier	scalp 1.5 / 2.5	tight-chandelier	FLAT (gate off)
thrust / abs_veto	exit is not the lever — keep the abs_veto entry veto; manage via gate on/off, not exit style ► REV3: the raw un-vetoed thrust cannot be rescued by an exit — but the LIVE abs_veto_long cell can, and it is the worst cell on the board. A tick-honest sweep of 81 long signals ranks the loaded A1.5/B2.5 last of eleven (+\$13.6/sig, LODO-min +\$4.7) against A1.0/B1.5 (+\$22.7, LODO-min +\$16.8), on a broad plateau (0.75/1.5, 1.0/1.5, 1.0/2.0 all within \$4). A second, independent sweep on 42 automated entries reached the same answer from different data. Worth ~+\$390/week at the live rate of 43 long signals. Change exit_overrides.json → abs_veto_long: {a_r 1.0, b 1.5}. Leave abs_veto_short at 1.5/2.5 — on that side the R choice is genuinely noise (+\$15 to +\$18 at every cell). No exit rescues abs_veto_short in violent whipsaw (-\$20 to -\$33 at every R from 0.5 to 6.0); only not trading it does.			
rgv / capitulation	exit is not the lever — per-side filter + router; leave exit at default			

Note the live exit_overrides.json currently sets exhaustion_short to 0.5/1.5 and abs_veto to 1.5/2.5. This lab says exhaustion_short should be regime-graded (tight-chandelier in chop, 1.5/2.5 in med) rather than a flat 0.5/1.5, and that ~~abs_veto's dollars come from its entry veto, not its exit R's~~ ► **REV3: most of abs_veto's edge is the entry veto — but the long side's R's are still the worst cell swept (see the policy row above) — a candidate tweak for the operator, on a two-week in-sample read, to confirm forward.**

Bottom line for the week: one thing earned a promotion (abs_veto_55s, two-sided), the two grind-exit shadows produced nothing decisive because grind barely traded (single-day, one-runner), the live exhaustion fade-scalp is two signals old and ungradeable, and the exit lab's cleanest result is a warning as much as a recipe — grind and exhaustion respond to a tighter, regime-graded exit, but thrust/

rgv/capitulation prove that for them the exit is a red herring and the real lever is the entry filter and which gates you let trade. And the honest gap that stays open is BIG-TREND: the desk still has not seen a clean trend week to prove the RIDE rung.

2.5 Mid-week musings — the leads, tested hard

This is the laboratory bench for everything that came up in the desk chat this week. Nothing gets to walk in as a "finding" — every mid-week idea is re-run on the week's own data and given a plain verdict. This week the tape handed us the perfect natural experiment: two clean-ish trending sessions (Paris-29 **+\$638**, Paris-30 **+\$1,172**) followed by a violent chop-reversal that cost **-\$905** on Paris-31. That **-\$905** day is where most of the lessons live.

LEAD WITH WHAT SURVIVED. THREE THINGS CAME OUT OF THE WEEK CLEAN:

1. GRIND SHOULD NOT RUN NAKED. ON THE TREND DAY IT WAS THE DESK'S TOP EARNER AT +\$703.5; ON THE CHOP REVERSAL IT BOUGHT THE TOP FOR – \$370.5. THE FIX THAT THE DATA ACTUALLY SUPPORTS IS AN IN-CODE CONTINUATION-VETO (MIRRORING ABS_VETO'S), NOT THE CURRENT ER-0.35 FLOOR (WHICH WAS POROUS ALL WEEK) AND NOT A "TREND-DAYS-ONLY" CALENDAR RULE.

2. THE ROUTER'S REACTION-LAG IS NOT THIS WEEK'S PROBLEM. MEASURED, THE REDUCIBLE ACTUATION-LAG COST WAS ABOUT \$0-\$246 FOR THE WHOLE WEEK, AND IT IS DETECTION-LIMITED (THE DAMAGE AND THE FLIP HAPPEN IN THE SAME 60-180 SECONDS). THAT DOLLAR FIGURE IS TOO SMALL TO JUSTIFY BUILDING A SEPARATE FAST-BENCH SERVICE.

3. THE EXIT SELECTOR EARNED ITS KEEP — ADAPTING THE EXIT WIDTH BEAT BOTH FIXED EXITS BY ~\$1,289 ACROSS THE WEEK.

1 · The direction-router's architecture — is the polling switch-file the best shape?

The objective, stated plainly: **stop a gate from trading against a proven trend, reactively, without over-clipping the chop days the faders actually earn on.** We built that as a 5-min durable tick (a headless-Claude judgment) that writes a switch file. It works — the desk went green twice this week under it. But the report's job is to ask whether a more elegant mechanism would do the same job with fewer moving parts and less lag.

What the polling-plus-switch-file shape actually cost us this week, from the live log:

LIVE COST	THIS WEEK'S EVIDENCE	REDUCIBLE?
Whipsaw / thrash	grind toggled ON↔OFF ~8× on Paris-30 alone; abs_veto L+S toggled ~5× on Paris-31	Partly — an in-code veto removes the toggling
Reaction lag (poll + judgment)	≈ \$0-\$246 (measured below)	Yes, but tiny
Detection lag (flip looks like a whipsaw until confirmed)	the bulk of the -\$1,460 raw reaction-window losses	No — irreducible (run-catcher NULL)
Sticky-exit / partial-bar replay artifacts	no material \$ this week	n/a

★ REACTION-LAG MANDATE — THE DOLLAR NUMBER. FOR EVERY BENCH EVENT THIS WEEK I MEASURED WHAT THE STILL-ARMED GATE LOST IN THE WINDOW BETWEEN THE REGIME TURNING AND THE SWITCH ACTUALLY FLIPPING (BENCH TIMESTAMPS FROM ROUTER_TRIAL_LOG.TXT, TRADES FROM THE LEDGER).

WINDOW MEASURED	NET TAKEN BY BENCHED GATES	READING
Full 15-min reaction window (loss legs only)	-\$1,460.5	Gross damage — but mostly detection-limited
Full 15-min reaction window (net)	-\$409.0	Winners in the window partly offset
5-min poll-cadence window (net)	+\$11.0	What a faster poll could have caught ≈ \$0
Re-entries after a first confirming stop (net)	+\$325.0	A stop-trigger fast-bench would lose money — most re-entries were good

Break the re-entry slice open, because it is the crux. Across the week, when a gate re-entered after a first stop inside a bench window, it made **+\$227.5**

/ +\$180 / +\$164 in three trend cases and lost only −\$101.5 (**grind, Paris-30 13:36**) and −\$145 (**grind, Paris-31 13:42**) in the two reversal cases. So the genuinely reducible actuation-lag loss for the whole week is ≈ −\$246.5 — and a mechanical "bench-after-a-stop" trigger built to harvest it would also have forfeited +\$571 of legitimate trend re-entries. That is a net-negative trade.

The clearest single case makes the point: on Paris-31 grind bought the rollover top and took four stops for −\$284; all four fills opened within 30 seconds at 13:38 and were stopped before the 13:42 bench could ever apply. No poll and no judgment loop, however fast, saves those — the flip and the loss were simultaneous. That is detection lag, and the run-catcher NULL [[run-catcher-null-all-microstructure]] already proved the flip is not front-runnable. So we do **not** propose predicting it.

What else — the alternatives, backtested against the polling router:

MECHANISM	WHAT THE WEEK'S DATA SAYS	VERDICT
(a) Per-gate in-code direction/continuation guard (per-tick, no poll, no file)	router_study puts the "counter-trend NEAR-MISS LEAK" a per-tick guard would catch at −\$956 (live-managed set) / −\$1,933 (full fader set) — the real prize	BUILD (start with grind)
(b) Event-driven fast-bench service (router_watch.py gets narrow defensive authority)	the lag it recovers is the \$0-\$246 above, detection-limited; a new service for that is not worth the moving part this week	DEFER
(c) Continuous in-tournament regime gating (no service, no file at all)	the clean end-state; equivalent to (a) generalised across all gates	DIRECTION OF TRAVEL
(d) Just tune the timer granularity	5-min window nets +\$11 — finer polling buys nothing	NO

ARCHITECTURE VERDICT. KEEP THE DELIBERATE 5-MIN LLM LOOP FOR ARMING — THAT IS THE EXPENSIVE, SUBTLE JUDGMENT (IS THIS A REAL BREAK WITH ER CLIMBING AND VOL EXPANDING, OR A DELTA-BLIP?) AND IT IS DOING THAT JOB WELL. MIGRATE THE DEFENSIVE SIDE IN-CODE, PER-GATE, BEGINNING WITH GRIND'S CONTINUATION-VETO. THAT CAPTURES THE COUNTER-TREND LEAK THE POLL STRUCTURALLY CAN'T, REMOVES THE THRASH, AND NEEDS NO NEW SERVICE. DO NOT BUILD A SEPARATE EVENT-DRIVEN FAST-BENCH ON THIS WEEK'S EVIDENCE — THE LAG IT RECOVERS (\$0-\$246, DETECTION-LIMITED) DOESN'T PAY FOR A NEW MOVING PART, AND THE MECHANICAL VERSION LOSES MONEY ON THE RE-ENTRIES. THE DOLLAR FIGURE GATES THE BUILD, AND IT SAYS: NO.

2 · The GRIND lead (operator-flagged, Paris-31): should grind run only on a confirmed trend?

★ REV3 CORRECTION — 2026-08-01 — EVERY ER AND ATR FLOOR ON THIS DESK HAS BEEN SWITCHED OFF SINCE 2026-07-29

Nothing below that leans on a live ER or ATR floor is describing the desk you own. `tournament.step()` hands the sub-slot tag (`grind_long_A`) to `er_blocks()`/`atr_blocks()`, but `ER_FLOOR/ATR_FLOOR` are keyed on base names, so the lookup never matches and every floor silently answers "don't block". Dead since the dual-slot cutover at **2026-07-29 16:20:49 UTC**; last suppression line 07-29T13:29:04Z, zero since. Five gates affected. The `_base()` helper that fixes it already exists at `tournament.py:86` and is applied correctly three times elsewhere — just not at lines 139 and 142.

Size it honestly: the circulating "7,734 suppressions" is a raw log-line count and the check re-fires every ~1s tick. De-duplicated it is **383 blocked signal episodes over 8 days** (140 ER + 243 ATR) — about 48 a day, not thousands. Real, but 44× smaller than advertised.

And it changes the conclusion, not just the footnote. Grind did not fire 52 of 56 times below its floor because "ER filters are fake on this tape" — a claim about the market. It fired because there was no floor — a claim about our code. Those point at completely different fixes. Both ER floors were even committed ~18 hours after the bug, so neither has ever blocked a single trade in production. Full working, per-gate verdicts and the mandatory ship order: [Part 3.6](#).

This was the week's headline lead. Grind's live/managed record splits perfectly by regime — and the split is exactly why the question matters. (These are the **live managed** `grind_long` numbers, not the misleading `grind_fast` shadow, which fires all day including chop and showed red on the very day live grind made its money.)

PARIS DAY	REGIME	GRIND FIRES	GRIND NET	WHAT HAPPENED
07-27	Chop (74% chop)	4	-\$149.0	Churned a quiet range

PARIS DAY	REGIME	GRIND FIRES	GRIND NET	WHAT HAPPENED
07-28	Chop	0	\$0	Correctly benched
07-29	Down/reversal	0	\$0	Benched (fired into Paris-30)
07-30	Trend up (US legs)	46	+\$703.5	Top earner — rode the up-legs
07-31	Chop reversal	6	-\$370.5	Bought the top, four stops, no veto

So the trend-day upside is real and large, and the churn is real and large. The operator asked which of three mechanisms best captures the +\$703.5 while killing the chop bleed. I reconstructed the efficiency-ratio (ER) at the entry of **every** grind fire this week (30-min window on 5-second bars from capture.db) to test them head to head.

Mechanism (c) — the current ER-0.35 floor — FAILED. If the floor were really holding grind to genuine-trend conditions, grind would not fire in chop. It fired constantly in chop:

RECONSTRUCTED ENTRY-ER BUCKET	FIRES	NET	WINS	\$/FIRE
ER < 0.35 (should be blocked by a 0.35 floor)	52	-\$174.5	16	-\$3.4
ER ≥ 0.35	4	+\$358.5	2	+\$89.6

Grind fired **52 of 56** times below 0.35 on this measure — including all of the chop churn and both chop-day bleeds. Whatever ER the live gate is computing, it is not the trend-efficiency the floor claims to guard, and it did nothing to stop -\$149, -\$370.5, or the intraday chop churn on the trend day. ~~ER filters are fake here (consistent with the standing memory).~~ ►

REV3: this is the wrong conclusion from the right observation. Grind fired below its floor 52 of 56 times because the floor was not running — a bug in our code, not a fact about the market. (The standing memory does survive on its own evidence: re-derived on 8 days of tick-honest replay, grind's $ER \geq 0.35$ would have thrown away +\$2,923 of profit to keep +\$83, so the floor is still a leak. But that is a separate proof, and it says **delete the floor**, not "ER is meaningless".) The one genuinely high-ER fire (ER 0.46, the Paris-30 barnstormer, +\$445 across its two lots) shows the floor would be a fine confirmer if it actually gated — but as deployed it does not.

Mechanism (a) — trend-days-only — clean but unusable in real time.

By calendar it separates beautifully (+\$703.5 trend vs −\$519.5 across the two chop days). But "trend day" is only knowable in hindsight, and even Paris-30 was **72% chop** by the classifier — grind churned those chop hours and made everything in the ~18% trend windows plus the one barnstormer. A day-level switch is too coarse; the decision has to be intraday.

Mechanism (b) — grind gets its own in-code continuation-veto — the winner. The cleanest evidence is the head-to-head on the exact Paris-31 reversal that cost grind −\$370.5:

GATE, PARIS-31 REVERSAL	FIRES	WINS	NET	DIFFERENCE
abs_veto_long (has a 55s continuation-veto)		10	5	+\$166.0
grind_long (no veto)		6	0	−\$370.5

Same day, same reversal, same direction — the **only** structural difference is that abs_veto re-requires the burst to persist for ~55 seconds before it commits, and grind doesn't. That single feature is worth a **+\$536 swing** between the two gates on one day. A per-tick continuation-confirm on grind would also have skipped the intraday chop churn on Paris-30 (it only fires on a persisting thrust) while keeping the trend-leg winners, which the reconstruction confirms were genuine thrusts (ATR-5 of 99-156pt on the two biggest wins).

GRIND VERDICT. YES — GRIND SHOULD EFFECTIVELY BE OFF UNLESS A MOVE IS CONFIRMED, BUT THE MECHANISM IS (B) AN IN-CODE CONTINUATION-VETO, NOT THE ER-0.35 FLOOR (POROUS, DISPROVEN THIS WEEK) AND NOT A CALENDAR TREND-DAY RULE (TOO COARSE). BUILD GRIND A 55S-STYLE PERSISTENCE CHECK LIKE ABS_VETO'S: IT IS REGIME-AGNOSTIC (WORKS ANY DAY, PER-TICK, NO POLLING LAG), IT REMOVES THE ROUTER THRASH, AND IT IS VALIDATED BY ABS_VETO OUT-EARNING GRIND BY +\$536 ON THE ONE REVERSAL THAT MATTERED. THIS IS THE SAME IN-CODE GUARD THE ARCHITECTURE SECTION LANDED ON — ONE BUILD FIXES BOTH.

3 · Nightly-rollup week-in-review — was the router optimal, and did the selector adapt?

Folding all five router_nightly and selector_nightly files into one week view. Both are **tick-honest estimates** — the trades the router blocked and the exits the selector didn't choose never actually ran, so treat "missed winners" as an upper bound (shadow overstates winners).

Router week-review

DAY	CHOP %	SAVED (LOSSES AVOIDED)	MISSED (WINNERS BLOCKED)	NET ROUTER VALUE	LEAKAGE (LOSERS WHILE ON)
07-27	73.9	\$0	\$81	-\$81	-\$456 (9)
07-28	62.9	\$506	\$567	-\$61	-\$398 (7)
07-29	51.5	\$2,263	\$3,980	-\$1,717	-\$229 (6)
07-30	72.2	\$5,656	\$4,848	+\$807	\$0
07-31	64.9	\$7,363	\$7,527	-\$164	\$0
Week	—	\$15,788	\$17,003	-\$440 to -\$646 — \$1,216 A/B merged; -\$1,216 as originally double- counted. Do not quote this number.	-\$1,083 week (-\$456 / -\$398 / - \$229 Mon-Wed, \$0 on both durable days)

★ REV3 CORRECTION — 2026-08-01 — THE SAME —\$1,216 CORRECTION AS PART 2.6 §3 — AND DO NOT NET THESE COLUMNS

The A and B legs of every dual slot are byte-identical in the nightly's own output because the repricer ignores the slot suffix, so **every suppressed signal is counted twice**. Merged, the week is 353–355 signals and net **–\$440 to –\$646**. A second double-count sits underneath it: the nightly de-dupes at 120 seconds and then prices every survivor as a full round trip even while the same gate is still in the previous one. Until both land, every “router value” number carries roughly 2× counting inflation. And the blended figure flips sign on which single day you drop (–\$804 dropping 07-30, +\$230 dropping 07-29) — it supports no verdict at all. Score the halves separately: the **fader bench is +\$810 (n=25) and survives every independent regime tag**; benching into chop does not (+\$612 / +\$187 / –\$382 across three tags). Working: [Part 3.5](#).

Do not net the columns. Leakage is the only realised-money column in this table; saved and missed are modelled counterfactuals on trades that never ran.

~~Read honestly, the estimator says the router over benched this week — it blocked more shadow winners than losers saved (net –\$1,216), positive only on the clean trend day.~~ But two things pull the other way, hard: (1) those blocked winners **never ran**, and shadow overstates winners, so "missed" is a ceiling, not a bill; (2) the **live desk actually went green twice** (Paris-29 +\$638, Paris-30 +\$1,172), so where it counted the benching was net-right. The genuinely good news is **leakage** — the managed gates' losers that fired while switched ON — fell from –\$456/–\$398/–\$229 early week to **\$0 on both of the last two days**. The router stopped leaking; the open question it flags is over-benching (missing winners), which is exactly what the in-code arming judgment is for.

Selector week-review (adaptive exit width)

DAY	N	OPTIMAL-PICK	SELECTOR	ALWAYS-SCALP	ALWAYS-WIDE	BEST-OF-3	REGRET
07-27	39	54%	+\$991	+\$320	+\$1,023	+\$2,334	\$1,342
07-28	23	57%	+\$634	+\$151	+\$36	+\$1,122	\$488
07-29	16	63%	+\$1,134	+\$634	+\$1,290	+\$1,662	\$528

DAY	N	OPTIMAL-PICK	SELECTOR	ALWAYS-SCALP	ALWAYS-WIDE	BEST-OF-3	REGRET
07-30	90	54%	−\$316	+\$354	−\$862	+\$2,431	\$2,748
07-31	31	81%	−\$448	−\$753	−\$846	+\$65	\$513
Week	199	59%	+\$1,995	+\$706	+\$641	+\$7,614	\$5,619

Adapting won. The selector's +\$1,995 beat always-scalp (+\$706) and always-wide (+\$641) — about **+\$1,289 over the better fixed exit** across the week, and it was the least-bad choice on the two red days. The systematic wrong-pick pattern is the flagged one, confirmed: **wide-in-violence whipsaws** — always-wide lost −\$862 on the choppy-trending Paris-30 and −\$846 on the violent Paris-31. The selector correctly leaned tighter on Paris-31 (12 tight vs 16 wide) and still only lost −\$448 on a day that gave almost nothing away (best-of-3 was just +\$65). Its worst regret was the trend day (07-30, \$2,748 left on the table vs a perfect picker) — on a huge day it under-committed to wide on the runners. Again: these are estimates; the non-chosen exit legs never ran.

4 • Standing weekly router optimisation — pull the thresholds and timings apart

router_study.py --days 40, run on both the live-managed fader set and the full historical fader set. Is it blocking enough losers, what does the threshold/timing grid say, and what's the honest deploy call?

UNIVERSE	BASELINE (NO ROUTER)	ROUTER NOW	Δ	BLOCKED
Live-managed {capitulation_long, rgv_short}	−\$302	−\$126	+\$176	4 (−\$216 losers / +\$40 winners)
Full fader set (+rgv_long, +exhaustion_short)	−\$722	+\$272	+\$995	39 (−\$1,476 losers / +\$480 winners)

Is it blocking enough losers? The classification says no — there's a real leak of counter-trend near-misses the current trigger is too strict to catch:

CLASS (FULL FADER SET, SHIPPED CONFIG)	N	\$
Allowed winner (correctly left ON)	71	+\$4,020

CLASS (FULL FADER SET, SHIPPED CONFIG)	N	\$
Blocked (router caught it)	39	−\$995
LEAK: counter-trend near-miss (trigger too strict)	32	−\$1,933
Chop loss (router can't help)	19	−\$1,156
Trend-following loss (a gate/exit problem, not the router's)	12	−\$658

That −\$1,933 near-miss leak is the same money the in-code per-tick guard from sections 1 and 2 would recover — the poll is structurally a beat behind it. The threshold sweep points the same way on **both** sets: loosening ER from 0.20 to **0.15** improves the router (live-managed +\$492, full +\$1,307 at ER0.15 / net30), and the gradient is monotone toward the looser ER at these steps.

CONFIG	LIVE-MANAGED Δ	FULL-SET Δ	NOTE
Shipped: ER≥0.20, net≥30, win30 step15 hold2	+\$176	+\$995	current
Moderate: ER≥0.15, net≥30 (shipped timing)	+\$492	+\$1,307	both sets agree — the safe loosening
Joint-best of 900: ER≥0.15 net≥20 win45 step5 hold3	+\$692	+\$984	sits on the grid EDGE (ER floor) — curve-fit tell

DEPLOY RECOMMENDATION. TAKE THE MODERATE LOOSENING — ER≥0.15, NET≥30, SHIPPED TIMING — NOT THE JOINT-BEST. THE JOINT-BEST LANDS ON THE BOTTOM EDGE OF THE SWEPT ER RANGE (0.15) WITH THE TIGHTEST STEP AND NET FLOOR, WHICH IS THE CLASSIC GRID-EDGE OVERFIT TELL; THE MODERATE CELL EARNS +\$492 / +\$1,307 WITH BOTH UNIVERSES AGREEING AND NO EXOTIC TIMING. IT'S A SATURDAY CHANGE (TOURNAMENT-CHANGES-SATURDAY-ONLY), AND IT MUST BE RE-VALIDATED AFTER A GENUINE TREND WEEK — THESE 40 DAYS ARE ONE SUMMER RANGE REGIME, THIN, AND PARTLY ALREADY HAND/ROUTER-GATED, SO READ EVERY NUMBER HERE AS AN ILLUSTRATIVE LOWER BOUND.

Backtest-discipline note: every number above is segmented by regime (chop / trend / reversal) and by Paris-day rather than sprayed as one blanket cross-tape average, per the standing rule — a single blanket number with no regime split would bury exactly the grind edge this section found. The full per-regime Lot-A/Lot-B R-target sweep (the exit cheat-sheet) lives in the gate & exit rehabilitation section, where the exit shadows run; here the selector week-review is the exit-adaptation evidence.

2.6 Router Operation Review — the desk's #1 lever, on the stand

► **Read this section with REV3 open.** The headline $-\$1,216$ is double-counted, the "yes, as a defence" verdict splits in two once the halves are scored separately, and $\$4/\8 lean on a grind floor that has never executed. There is also a live finding that outranks the whole section: **all six gates are pinned ON right now, so the durable router cannot bench anything at all.** The corrections are marked inline in red where they bite, and the full working is in [Part 3 — the Rev2/Rev3 corrections](#). Nothing has been deleted: where a claim was overturned, the original is still printed next to what replaced it.

Garrath — you asked for the works: what the router did all week, what paid, what erred, whether the thresholds and the new untradeable-meter need tuning, and whether the whole machine is earning its keep. Here it is, money-first. **Lead with what survived:** (1) the new untradeable-day meter is the standout — it scored the two green days lowest (13, 23) and Friday's $-\$905$ bloodbath highest (86), a near-perfect line-up with the day's actual P&L; (2) the **durable** 5-min router held its nerve — across its two live days it made just **2** autonomous switch-changes (both graded GOOD) and correctly sat frozen through all of Friday's untradeable chop, **zero real blocked-loss leakage**; (3) the desk finished the week **+\$625 as booked — but $-\$598$ with your hands off** (194 automated fills, 35.1% win, PF 0.91; $\$1,223$ of the headline came from twelve trades you closed yourself — see Part 1). What still bites: the router keeps **over-benching aligned momentum in a trend** (the recurring BAD-call), and Friday's damage came almost entirely from the operator-directed **enable-all-at-the-open**, not from the routine ticks. One robustness-gated Saturday knob survives the anti-overfit filter (fader-bench ER floor $0.20 \rightarrow 0.15$); everything else is hold-and-re-validate.

HONESTY BANNER UP FRONT: THIS IS ONE WEEK, ~5 TRADING DAYS, ONE SUMMER RANGE REGIME, AND THE THRESHOLD SWEEP IS IN-SAMPLE ON THIN, PARTLY-ALREADY-GATED REALISED TRADES. EVERY BACKTEST BELOW IS SEGMENTED (FADER-SET VS FULL-SET, PER-REGIME LEAK BUCKETS) — NO BLANKET CROSS-TAPE NUMBER IS PRESENTED AS TRUTH. TREAT SINGLE-WEEK

\$ AS DIRECTIONAL, NOT PROVEN. RE-VALIDATE AFTER A REAL TREND WEEK BEFORE DEPLOYING ANYTHING.

1 • The week's decisions & scoreboard

Every ~5 min the router reads the desk and either changes a gate switch or logs "no change". A healthy router is **mostly no-change** — the thrash check. Here is the week, tick-by-tick condensed, with the live realised P&L each day actually printed:

PARIS DAY	TICKS	SWITCH CHANGES	THRASH %	WHAT THE ROUTER DID	LIVE P&L
Mon 07-27	—	—	—	pre-durable / session-only; reversion book	−\$226
Tue 07-28	5	1	20%	benched shorts under a young UP bias	−\$54
Wed 07-29	78	9	12%	first green day; scale-out live; held longs	+\$638
Thu 07-30	241	35	14.5%	BEST-EVER day but thrashy — grind toggled ~8x, one premature flip	+\$1,172
Fri 07-31	289	13	4.5%	durable era; 0 routine changes — all 13 were operator/session overrides; held flat	−\$905

Read: the thrash dropped hard once the durable tick took over routine ticking — 07-30 (hand-managed) churned the switches 35 times in a day; 07-31 (durable) made **zero** autonomous routine changes and let the operator drive the exceptions. The ledger scoreboard stands at **GOOD 17 • BAD 2 • CHURN 4 • NEUTRAL 2** — the two BAD calls and all four CHURN calls are the same two lessons repeating (premature-arm on a delta-blip; benching the aligned trend-rider on open-noise). More on those next.

2 · What works vs what doesn't

Walking the graded ledger, the pattern is clean: the router **saves money defensively** (benching faders/counter-trend on chop and roll-over days, shadow-confirmed) and **errs offensively** (benching aligned momentum too eagerly, or arming it too eagerly on a blip). Money attached where measurable:

CALL	GRADE	\$ IMPACT	COUNTERFACTUAL / PROOF
07-30 14:51 bench both momentum longs on ER-collapse regime flip	GOOD	≈ +\$400 saved	unbenched momentum shadows lost −\$398 over that window (thrust_aligned −202, cont −118, loose −78)
07-31 defensive trim → fully-flat through the ATR-52/62 open violence	GOOD	≈ +\$500 saved	every benched gate's shadow bled (grind_fast −153, thrust_loose −153, exhaustion_rev −99, abs_veto −80)
07-30 14:57 arm short-side fade book for the roll-over chop	GOOD (regime read) ⚠ the \$ credit is manual	+\$372 hand-claimed / −\$51 automated	exhaustion_short went 4/4 live fading the roll-over — but all four were MANUAL_CLAIM exits. Repriced on 250 ms tape under the exit config actually loaded, the machine books −\$51.2 on the same four entries. The regime call was right; the money was your hand.
07-31 held flat through the low-vol overnight grind-up (+141pt)	GOOD	churn avoided	all 8 earlier gate-fires stopped — a low-vol grind is untradeable, flat was right
07-30 15:49 FLIP to longs on a delta-blip in ER-0.01 chop	BAD	process-loss	faded in 2 min; salvaged to ~+\$55 by operator Claims but the read was wrong → banked the "ER-climb not delta-blip" rule
07-30 13:36 bench grind on the open flush (aligned trend intact)	BAD	≈ −\$430 missed	operator overrode; grind then rode to new highs +\$234/+\$197 → banked "don't bench the trend-rider at the open"
07-30 grind toggled ~8× through the morning	CHURN	≈ \$0 (thrashy)	grind's 0.35 ER floor + abs_veto's veto meant most premature arms never fired — process error, not \$ error ► REV3: grind's 0.35 floor has never executed (Par t 3.6) , so it cannot have absorbed those arms. Only abs_veto's 55s veto was actually running. This CHURN is graded \$0 on the strength of a filter that was not there.
07-31 afternoon abs_veto re-enables into the choppy recovery (×2)	CHURN	≈ −\$245	even the veto-protected, scalp-exit abs_veto churns a choppy recovery — the 55s veto is a

CALL	GRADE	\$ IMPACT	COUNTERFACTUAL / PROOF
			microstructure filter, not a regime filter

THE ONE DURABLE TRUTH FROM THE WHOLE LEDGER: THE ROUTER'S CHEAP ERROR IS A WRONGLY-BENCHED GATE (MISSED TRADES); ITS EXPENSIVE ERROR IS A WRONGLY-ARMED ONE (REAL CHURN). EVERY ▢/△ THIS WEEK WAS AN ARMING MISTAKE OR A BENCH-THE-TREND MISTAKE — NEVER A "LEFT A BLEEDER ARMED" MISTAKE. THE DEFENSE WORKS; THE OFFENSE OVER-MANAGES.

3 · Net router value — blocked-losses saved vs winners missed

Rolling up every router_nightly file (the router's own scored counterfactual: what each benched gate's shadow would have done in its benched window), by regime:

DAY	DOMINANT REGIME	SAVED	MISSED	NET VALUE	REAL LEAKAGE
07-27	CHOP 74%	\$0	\$81	−\$81	−\$456
07-28	CHOP 63%	\$506	\$567	−\$61	−\$398
07-29	CHOP 52% / TREND 48%	\$2,263	\$3,980	−\$1,717	−\$229
07-30	CHOP 72%	\$5,656	\$4,848	+\$807	\$0
07-31	CHOP 65%	\$7,363	\$7,527	−\$164	\$0
WEEK	—	\$15,788	\$17,003	−\$440 to −\$646 − \$1,216 A/B merged; −\$1,216 as originally double- counted. Do not quote this number — it flips sign on which single day you drop.	−\$1,083 week (−\$456 / −\$398 / − \$229 Mon-Wed, \$0 on both durable days)

~~Honest reading (this number is noisier than it looks): the raw net is — \$1,216, but two things inflate the "missed" column — the dual A/B slots double-count one signal...~~

★ REV3 CORRECTION — 2026-08-01 — **-\$1,216 COUNTS EVERY SUPPRESSED SIGNAL TWICE, AND THE BLENDED NUMBER SUPPORTS NO VERDICT AT ALL**

The double-count is real and the report named it and then left it in the total anyway. The dual-slot A and B legs are byte-identical in the nightly's own JSON because `reprice_scalp()` ignores the slot suffix. Merge them and the week is **353-355 signals, ~\$8,700 saved, ~\$9,200 missed, net -\$440 to -\$646** (two independent re-derivations, \$206 apart on \$9.4k of gross). That recovers about \$661 of pure bookkeeping.

There is a second double-count nobody caught until this weekend. `router_nightly.py` de-dupes suppressions at 120 seconds and then prices every survivor as a complete round trip — even when the same gate is still notionally in the previous one. A real gate holds one position at a time. Enforcing that halves the blocked book again. **Until both fixes land, every "router value" number in this report carries roughly 2× counting inflation.**

Do not quote the blended figure at all. It flips sign on leave-one-day-out ($-\$804$ dropping 07-30, $+\$230$ dropping 07-29), which means it supports no verdict in either direction. Score the halves separately — that is the \$8 verdict, corrected.

What the "missed" column is actually made of: the giant figures on 07-29 ($-\$1,717$ value) are dominated by **grind_long shadow winning during brief benched windows on a trend day**. That is the exact recurring error surfacing quantitatively: **the leak is over-benching aligned momentum in a trend**, and it is opportunity-cost (winners not taken), **not** realised bleed. The proof: on both durable days (30/31) real leakage was **\$0** — the router let no live loser through. So:

- **Not-bleeding job: PASS.** Zero real blocked-loss leakage on the days the durable router actually ran; $+\$807$ of genuine defense on 07-30.
- **Catching-winners job: LEAKY.** It benches momentum too readily during trends (07-29's $-\$1,717$ "value", 07-30's $\square$ open-flush bench). The leak is concentrated in **momentum gates during TREND_UP windows** — the exact segment the "don't bench the trend-rider" rule exists to protect.

For context, the exit-width **selector** (a separate machine) captured only **\$1,995 of a possible \$7,614 (26%)** this week, leaving \$5,619 regret — worst on 07-30 (chose tight, missed the wide runners: — \$2,748 regret). That is a bigger leak than the router and belongs in the selector's own review, flagged here so it isn't mistaken for router value.

4 • Threshold tuning — the sweep

★ **REV3 CORRECTION — 2026-08-01 — EVERY ER AND ATR FLOOR ON THIS DESK HAS BEEN SWITCHED OFF SINCE 2026-07-29**

Nothing below that leans on a live ER or ATR floor is describing the desk you own. `tournament.step()` hands the sub-slot tag (`grind_long_A`) to `er_blocks()/atr_blocks()`, but `ER_FLOOR/ATR_FLOOR` are keyed on base names, so the lookup never matches and every floor silently answers "don't block". Dead since the dual-slot cutover at **2026-07-29 16:20:49 UTC**; last suppression line 07-29T13:29:04Z, zero since. Five gates affected. The `_base()` helper that fixes it already exists at `tournament.py:86` and is applied correctly three times elsewhere — just not at lines 139 and 142.

Size it honestly: the circulating "7,734 suppressions" is a raw log-line count and the check re-fires every ~1s tick. De-duplicated it is **383 blocked signal episodes over 8 days** (140 ER + 243 ATR) — about 48 a day, not thousands. Real, but 44× smaller than advertised.

And it changes the conclusion, not just the footnote. Grind did not fire 52 of 56 times below its floor because "ER filters are fake on this tape" — a claim about the market. It fired because there was no floor — a claim about our code. Those point at completely different fixes. Both ER floors were even committed ~18 hours after the bug, so neither has ever blocked a single trade in production. Full working, per-gate verdicts and the mandatory ship order: [Part 3.6](#).

Ran router_study.py --days 40: the ER-trend × net-min threshold grid plus the timing sweep (step/hold/window/fast_exit), on **both** the live-managed fader set and the full historical fader set. Router Δ vs the no-router baseline (shipped timing), higher = better:

ER FLOOR \ NET MIN	20	30 *SHIPPED	40	50	60
ER 0.15 (looser trend def)	+461 / +1300	+492 / +1307	+428 / +1294	+316 / +1120	+158 / +858
ER 0.20 *current	+224 / +1044	+176 / +995	+152 / +976	+152 / +1014	+152 / +816
ER 0.25	+152 / +787	+152 / +787	+152 / +806	+152 / +806	+152 / +744
ER 0.30	+152 / +334	+152 / +334	+152 / +334	+152 / +334	+152 / +428

(each cell = live-managed set Δ / full historical set Δ, in \$)

What both sets agree on (the robust signal): the ER 0.15 row beats the ER 0.20 row in every single NET column, on both universes.

Loosening the router's **trend-detection floor** from 0.20→0.15 lets it call "trend" sooner and therefore bench counter-trend faders more — and in this fader-heavy, counter-trend-loss-dominated sample that is worth roughly **+\$300 (live set) / +\$300 (full set)** over shipped. That is the one change that survives the two-sets-must-agree filter.

ANTI-OVERFIT DISCIPLINE — WHAT I AM NOT RECOMMENDING, AND WHY:

- THE GRID'S JOINT-BEST (ER0.15 / NET20 / WINDOW45 / STEP5 / HOLD3) SCORES +\$692 LIVE / +\$984 FULL BUT SITS ON THE LOOSEST GRID EDGE ON THREE AXES AT ONCE, HELPS ONLY 4-6 OF THE SAMPLE DAYS, AND STEP=5M WOULD TRIPLE THE TICK-RATE = A THRASH-MACHINE ON A LIVE DESK. REJECTED AS OVERFIT.
- NET-MIN STAYS AT 30. ON THE LIVE SET ER0.15 ACTUALLY SCORES HIGHER AT NET30 (+492) THAN NET20 (+461); ON THE FULL SET THEY TIE (+1307 VS +1300). NO REASON TO LOOSEN IT.
- TIMING STAYS SHIPPED (WINDOW 30M, STEP 15M, HOLD 2, FAST_EXIT OFF). THE TIMING GAINS ARE GRID-EDGE AND THIN-N.
- THE GRIND 0.35 ER ARM FLOOR IS A DIFFERENT THRESHOLD (MOMENTUM-ARMING, NOT FADER-BENCHING) AND IS VALIDATED SEPARATELY (+\$728 OVER 10D — AGAINST WHAT ALTERNATIVE? THE COUNTERFACTUAL WAS NEVER STATED) — LEAVE IT. ► REV3: AND SEE THE DEAD-FLOORS CORRECTION ABOVE — THAT THRESHOLD HAS NOT RUN SINCE 07-29 16:20, SO NOTHING IN THIS WEEK'S TAPE TESTS IT. RE-DERIVED ON 8 DAYS OF TICK-HONEST REPLAY IT IS THE MOST EXPENSIVE LINE OF CONFIG ON THE DESK: IT KEEPS 8 SIGNALS WORTH

+\$83 AND DISCARDS 173 WORTH +\$2,923. THE SATURDAY PLAY DELETES IT. DAY-BIAS ±40 — NO EVIDENCE TO MOVE IT. LEAVE IT.

5 • The untradeable-day meter review (new 07-31)

The new untradeable.py meter blends three tells into a 0-100 STAY-OUT score (≥ 65 = stay out). I computed it for every Paris day this week and put it next to what the day **actually** paid. This is the headline survivor of the whole review:

DAY	SCORE	VERDICT	ROUNDTRIP MTR	GIVEBACK MTR	STOPS MTR	ATR	ACTUAL P&L
07-30	13	TRADEABLE	4	4	47	8.4	+\$1,172
07-29	23	TRADEABLE	25	5	53	14.1	+\$638
07-27	47	CAUTION	48	40	60	5.6	-\$226
07-28	56	CAUTION	62	62	33	6.9	-\$54
07-31	86	STAY-OUT	90	84	80	11.7	-\$905

The meter lines up perfectly on the only five days it has seen. ~~The meter is validated.~~ ► **REV3: that is a promising start on n=5, not a validation. Re-score it weekly and do not let it size anything until it has a month.** Sort by score and you sort by P&L almost perfectly: the two lowest scores (13, 23) are the two green days; the highest (86) is the -\$905 disaster. There is a clean gap between the worst tradeable day (23) and the best caution day (47) — the classifier separates. It flagged Friday **correctly and early** (it read 74→96 STAY-OUT live through the afternoon).

Which of the 3 meters actually discriminated — being honest:

METER	GREEN DAYS (29/30)	RED DAY (31)	VERDICT
ROUNDTRIP (net /range)	25 / 4 (low)	90 (high)	DISCRIMINATES
GIVE-BACK (excursion handed back)	5 / 4 (low)	84 (high)	DISCRIMINATES
STOP-RATE	53 / 47	80	WEAK — best day (30) had a higher stop meter than a losing day (28)
ATR / chop-time% / shadow-net	14.1 / 8.4	11.7	DOES NOT — best day had lower ATR than the disaster

Roundtrip + give-back carry the whole meter; stops is weak supporting context; ATR is noise — exactly as the design note claimed. The 40/40/20 blend is right to lean on the two that work. **The 65 cutoff wants no change** — it caught the −\$905 with room to spare and produced zero false STAY-OUTs on green days. One refinement worth a look: both CAUTION days (47, 56) still lost money, so CAUTION isn't costless — consider making CAUTION an **action** (faders-off / veto-gates-only / half-size) rather than just a label. n=2, so propose-and-validate, don't hard-wire.

6 · Bench-decision review — grading the week's actual benches

The Friday US-open sequence is the teaching case. Operator said "enable all gates, US OPEN" and left; the open printed a real up-thrust (ER 0.68, break 28713) that then **fully reversed** (−172pt roundtrip). Here is exactly how the six gates behaved, live:

GATE ARMED AT THE OPEN	TYPE	OPEN P&L	LESSON
abs_veto (55s-veto protected)	VETO-gate	+\$146	caught the thrust, banked both legs, held disciplined
grind (no veto)	CHURNER	−\$370	bought the rollover top, 4 stops in 90s
exhaustion_short (off-regime)	FADER	−\$72	fired in ATR-62 violence — a 0.5R scalp can't cover a 62pt stop

BANKED LESSON [[US-OPEN-ARM-VETO-GATES-NOT-CHURNERS]]: AT A VIOLENT OPEN, ARM THE VETO-PROTECTED GATES (ABS_VETO) AND KEEP THE CHURNERS (GRIND) AND OFF-REGIME FADERS BENCHED. ENABLE-ALL GUARANTEES THE WRONG-SIDE GATES CHURN THE REVERSAL BEFORE ANYONE CAN BENCH THEM. THE DEFENSIVE TRIM THAT FOLLOWED (GRIND→CAPIT/RGV→EXHAUSTION→FULLY-FLAT) WAS SHADOW-VINDICATED: EVERY BENCHED GATE'S SHADOW BLED THROUGH THE ATR-52/62 CHOP.

The individual bench grades:

- **grind benches** — right when the trend structure breaks (07-31 open reversal, 07-30 14:51 ER-collapse: both GOOD); wrong when the trend is intact and it's just open-noise (07-30 13:36: BAD). The distinction is roundtrip + ER-collapse, and the meter now measures exactly that.
- **abs_veto scalp-then-bench** — the fire-and-stop exception to its veto protection worked (07-30 22:10 bench after it stopped both legs = GOOD). But re-enabling it into a **choppy recovery** (07-31 pm, ×2, −\$245) was CHURN: the 55s veto is a microstructure filter, so a chop micro-thrust passes it. Rule holds: don't re-arm even abs_veto until a genuinely clean trend ($ER \geq 0.35$ sustained).
- **the exhaustion_short live trial** — deployed the 0.5R/1.5R fade-scalp; it fired **once, off-regime**, in the open violence (net −\$72, as predicted for high ATR). The core thesis (fader-scalp in normal chop, ATR 11–18) is **still untested** — one off-regime fire is not a verdict. Keep it pinned and running into more low/mid-ATR roll-over afternoons before judging.
- **the fully-flat "stay-out" calls** — the best calls of Friday. Once the day was read untradeable, going flat banked **zero further bleed**; the meter held 74–96 STAY-OUT the whole time confirming it. Stay-out is a real, earned decision, not a cop-out.

7 · Durable router & architecture — is polling the right mechanism?

Did the durable 5-min systemd tick manage correctly? Yes. It is

active / enabled, fires on schedule, and across its two live days it made 2 autonomous switch-changes (07-30 22:10 bench abs_veto after fire+stop, 22:15 bench the 3 shorts + grind under a fresh UP bias — **both graded GOOD**) and then correctly logged "changed: none" through the entire Friday untradeable day. **No thrash, no naked positions, no double-writes, zero real leakage.** The robustness fix works: the desk was never unmanaged.

Is polling the best mechanism, though? Here is the honest cost of the 5-min poll. The router only acts on the tick boundary, so when a regime flips mid-window, up to 5 minutes of churn can print before the switch changes. The clearest single instance this week: on the Friday open, grind bought the reversal and lost **−\$370 over ~4 minutes** before it was benched. A per-gate **in-code regime-guard** (an entry-time veto that blocks a counter-regime fill in milliseconds, no polling) would have stopped most of that at the source.

MECHANISM	REACTION LAG	ADDRESSABLE LEAK THIS WEEK	TRADE-OFF
5-min poll (current durable)	up to 5 min	baseline	simple, fail-safe, survives session exit; laggy on fast flips
event-watcher (in-chat, current)	~seconds	partial	already cuts lag on breaks/bleed — but dies with the session
in-code per-gate regime-guard (proposed)	<1 sec, at entry	up to −\$956 / −\$1,933	catches the counter-trend NEAR-MISS bucket the poll+coarse-threshold lets through

Where that leak number comes from (regime-segmented, per the discipline): the study's trade classification buckets the fader trades by regime-context. The "**counter-trend near-miss**" bucket — trades that entered against a detectable regime and lost — is **n=15 / −\$956 (live set)** and **n=32 / −\$1,933 (full set)**. That is the leak a millisecond in-code guard could address that a 5-min poll structurally cannot. **Recommendation:** keep the poll as the fail-safe floor, but scope an in-code counter-regime entry-veto on the momentum gates as the next build — it attacks the exact segment (counter-trend fills in a live trend) where polling loses. Flagged as a build, not a Saturday switch.

8 • Verdict & Saturday recommendations

Split verdict — and the two halves point opposite ways. ~~Is the router earning its keep? Yes — as a defense.~~ **The defence earns its keep on exactly one cut, and it is the one that survives everything: benching the FADERS is +\$810 (n=25, non-overlap) and comes out identical under every independent**

regime tag anyone could throw at it. Benching into CHOP is not in that class — it scores +\$612 under the router's own tag, +\$187 under ADX(14) and −\$382 under a 60-minute regression tag, i.e. it only holds if you accept the router's own definition of chop. And the offence does not earn its keep: the "stop benching momentum in TREND_UP, +\$1,615/week" finding was independently re-derived this weekend and is **WITHDRAWN** ([Part 3.5](#)). Zero real blocked-loss leakage on the days it ran durably, +\$807 of shadow-confirmed defence on 07-30, and a durable tick that held flat through the worst day of the week without a session driving it — those all stand. The desk finished **+\$625 as booked, −\$598 with your hands off** (see **Part 1**) through a week that contained a −\$905 day. ► **And the finding that outranks this entire section: as of 08-01 all six gates are PINNED in router_tick_durable.py, so the apply step's change set is permanently empty and the durable tick has benched nothing for 36 hours. The proven half of the machine is switched off. That is MONDAY play #1.** The weakness is offensive: it over-benches aligned momentum in trends (opportunity-cost, not bleed) and the exit-selector left \$5.6k of regret on the table (separate machine, separate review).

#	SATURDAY ACTION	EVIDENCE	GATE
1	Lower the fader-bench ER-trend floor 0.20→0.15 (NOT grind's 0.35 arm floor)	ER0.15 row beats ER0.20 in every NET column on both universes (+~\$300 each)	DEPLOY re-validate after a trend week
2	Keep NET-min at 30, keep shipped timing (window30/step15/hold2)	joint-best is 3-axis grid-edge + step5 = thrash; NET30 ties/beats NET20 at ER0.15	HOLD
3	Keep the untradeable meter's 65 STAY-OUT cutoff and 40/40/20 blend	caught −\$905 (score 86), zero false positives; roundtrip+giveback carry it, stops/ATR don't	HOLD
4	Make CAUTION (45-64) an action: faders-off / veto-gates-only	both CAUTION days (47, 56) lost money — CAUTION isn't costless	TRIAL n=2, validate first
5	Codify: at a violent open, arm VETO-gates only, hold churners/off-regime faders	07-31 open: abs_veto +\$146 vs grind −\$370 vs exhaustion −\$72	RULE banked
6	Stop hand-toggling grind tick-to-tick; arm once on ER≥0.35 sustained, bench only on structure-break	07-30 ~8× toggle CHURN ;grind self-gates via its 0.35 floor ► REV3: clause deleted — grind does not self-gate. That floor has never executed (Part 3.6)	RULE banked

#	SATURDAY ACTION	EVIDENCE	GATE
	Scope an in-code per-gate counter-regime entry-veto (momentum gates)	counter-trend near-miss leak – \$956/–\$1,933 the 5-min poll can't reach	

BACKTEST DISCIPLINE HONORED, AND ITS LIMITS STATED PLAINLY: EVERY SWEEP HERE IS SEGMENTED (FADER-SET VS FULL-SET; PER-REGIME LEAK BUCKETS; PER-DAY METER VS P&L) — NO BLANKET CROSS-TAPE NUMBER IS OFFERED AS TRUTH. BUT THIS IS ONE WEEK, ONE SUMMER RANGE REGIME, IN-SAMPLE, THIN REALISED N, PARTLY ALREADY GATED. THE ONE DEPLOY (ACTION 1) IS THE ONLY CHANGE BOTH UNIVERSES AGREE ON ACROSS THE WHOLE GRID; EVERYTHING ELSE IS HOLD OR TRIAL. THE REGIME→EXIT R-TARGET CHEAT-SHEET (FADERS 0.5/1.5, MOMENTUM 1.5/2.5, TREND 2.5/WIDE) REMAINS AN UNPROVEN OPERATOR PRIOR — THE EXHAUSTION 0.5R/1.5R TRIAL HAS FIRED EXACTLY ONCE, OFF-REGIME, SO THOSE RS ARE NOT YET PROVEN AND SHOULD BE SWEEPED PER-REGIME ONCE A TREND WEEK SUPPLIES THE N. RE-RUN THIS WHOLE REVIEW AFTER THE NEXT REAL TREND DAY BEFORE TRUSTING ANY SINGLE-WEEK DOLLAR FIGURE.

Part 3 The Rev2 & Rev3 corrections — the new material, in full

Everything in this part was written after the report went out. Parts 3.1–3.4 are the five Rev2 answers that never landed in the published file. Parts 3.5–3.7 are three independent weekend re-derivations that were commissioned to break those answers — two of them succeeded, and where a Rev3 section disagrees with a Rev2 section, **the Rev3 section is the correction**. Part 3.8 is the audit that caught the failed fold-in in the first place, printed in full because it is the receipt for this revision.

What the desk made without your hands on it

PATCH MARKER — this section **patches and extends PART 1 (“The Live Desk”)**: it restates Part 1’s headline number, its six-gate card table, its exit-mix table and its “the tail is the edge” claim on **automated fills only** (MANUAL_CLAIM excluded), and it re-runs Part 1’s **SURVIVOR – steadiest** verdict on `exhaustion_short` against that same automated-only tape. It also extends **Part 2.6 (Router review)** with the regime split of the automated book. Everything below is computed from `data/gazbot7.db` (venue truth) and `data/capture.db` 250 ms ticks. Where Part 1 and this section disagree, **this section is the correction.**

Garrath — you asked the right question, and the answer is not comfortable. **Hands off, the desk lost \$598 this week.** The +\$625 headline in Part 1 is real money on a real ledger, but \$1,223 of it came from twelve trades you closed by hand with the Claim-profit button. Strip those twelve and what the machine did on its own is **194 fills, -\$598, 35.1% win, -\$3.08 a trade, profit factor 0.91**. Part 1 led with a survivor; on automated fills that survivor is not green. But it is not the disaster it first looks either — and the reason it broke is a config change you made on Wednesday evening, not a regime. Below: the honest split, the tick-by-tick counterfactual of what the robot would have done with those twelve trades if you had sat on your hands, the re-run survivor verdict, and the re-nomination.

THE ONE-PARAGRAPH ANSWER

Automated-only: -\$598 over 194 fills (35.1% win, PF 0.91). The twelve hand-claims were +\$1,223 at 12-for-12. Repriced on 250 ms ticks under the exit config that was actually loaded at the time, those same twelve entries — the entries were automated; only the exits were yours — would have banked **+\$360.5**, five of the twelve stopping out. So a fully hands-off week is **-\$237.5**, not +\$625, and **your hands were worth +\$862.5**. The nominated survivor, `exhaustion_short`, is **-\$56.0 on 44 automated fills** — flat, not green. Its “+\$372.5 four-in-a-row” on Thursday was four hand-claims, and three of those four were sitting on positions the tape later ran **49 to 98 points through the stop**. **SURVIVOR – steadiest** does not survive as a money

claim. It does survive as a mechanism claim, but only on the exit it was running until Wednesday 19:10 UTC.

Fact table 1 — the week, with your hands taken off

Week defined exactly as Part 1 defines it (07-26 through 07-31; totals reconcile to Part 1’s +\$625 / 206 trades to the cent). “Automated” = every fill whose `exit_reason` is not `MANUAL_CLAIM`.

BOOK	N	NET \$	WIN %	\$ / TRADE	MEDIAN \$	PF
Week as reported (Part 1)	206	+625.0	38.8	+3.03	−31.75	1.10
AUTOMATED ONLY	194	−598.0	35.1	−3.08	−35.50	0.91
MANUAL_CLAIM (your hands)	12	+1,223.0	100.0	+101.92	+101.75	∞
Automated, excluding Wednesday 07-29	162	−1,947.0	34.0	−12.02	−35.50	0.64

That last row matters as much as the first: **the automated book has exactly one good day in it.** Take out Wednesday’s +\$1,349 and the machine is −\$1,947 across five sessions at a profit factor of 0.64. The week is not “+\$625 with a wobble”; it is one clean trend day, your hands on Thursday and Friday, and bleed everywhere else.

SESSION (UTC CLOSE- DATE)	AUTO N	AUTO NET \$	HAND N	HAND NET \$	REPORTED NET \$
Sun 07-26	3	−71.5	0	—	−71.5
Mon 07-27	37	−155.0	0	—	−155.0
Tue 07-28	25	−22.0	0	—	−22.0
Wed 07-29	32	+1,349.0	0	—	+1,349.0
Thu 07-30 “best day ever”	73	−884.5	9	+978.0	+93.5
Fri 07-31	24	−814.0	3	+245.0	−569.0

Thursday is the headline casualty. Part 1 calls 07-30 the best day the desk has ever booked. On automated fills Thursday is the worst day of the

week: **73 fills, −\$884.5, 26% win.** You turned it green by hand, nine times, nine winners.

Fact table 2 — the six-gate card table, automated only

This is Part 1’s card table with the hand-claims removed. Every row reconciles: Part 1 net = automated net + hand net.

GATE	PART 1 NET	AUTO N	AUTO NET \$	AUTO WIN %	AUTO \$/ TR	AUTO PF	HAND \$	REVISED VERDICT
abs_veto_short	+358.5	26	+358.5	42.3	+13.8	1.53	0.0	ONLY AUTOMATED-GREEN GATE — but 52% of it is one trade
exhaustion_short	+316.5	44	−56.0	50.0	−1.3	0.96	+372.5	FLAT, NOT GREEN — see verdict re-run
capitulation_long	−139.0	14	−139.0	28.6	−9.9	0.65	0.0	unchanged → REHAB
rgv_short	−154.5	3	−154.5	0.0	−51.5	0.00	0.0	unchanged → router-benched
abs_veto_long	+59.5	56	−185.5	32.1	−3.3	0.90	+245.0	FLIPS TO RED — “scratch” was hand-made
grind_long	+184.0	51	−421.5	25.5	−8.3	0.80	+605.5	FLIPS TO RED — worst gate on the desk

Two of Part 1’s six gates change sign. **grind_long’s entire +\$184 was your hands and then some** — automated it is the biggest loser on the book at −\$421.5 and it wins one trade in four. **abs_veto_long’s “scratch” was hand-made too.** Only **abs_veto_short** is green with nobody touching it, and it took zero hand-claims all week — that number is clean.

The Lot-A / Lot-B claim (“the tail is the edge”) — half of it was your hands

Part 1 concludes the chandelier tail lot carries the money. Automated, that is true for **abs_veto** and false for **grind**.

GATE	PART 1 "TAIL" (LOT B)	AUTO LOT B	PART 1 "SCALP" (LOT A)	AUTO LOT A	VERDICT
grind_long	+302.0 (19)	−151.0 (16)	+31.0 (33)	−121.5 (31)	FALSIFIED — both lots red automated
abs_veto_long	+321.5 (19)	+131.5 (17)	+4.5 (28)	−50.5 (27)	holds, but ⅔ smaller
abs_veto_short	+379.5 (4)	+379.5 (4)	+52.0 (5)	+52.0 (5)	holds — untouched by hand
exhaustion_short	−37.5 (5)	−264.0 (3)	−11.0 (6)	−157.0 (4)	both scale-out lots badly red

Exit mix, automated only

EXIT	N	NET \$	\$/TRADE	NOTE
CHANDELIER	32	+3,065.0	+95.8	96.9% of them green — the tail is still the profit engine
TARGET	33	+2,647.5	+80.2	fixed-R hit
GIVEBACK	4	+82.0	+20.5	—
STOP	119	−6,119.0	−51.4	61% of every automated fill is a stop-out
STOP_UNFILLED	6	−273.5	−45.6	the resting-stop failure path (Part 1.5)

The automated book wins \$5,794 gross and loses \$6,392 gross. It is not broken — it is **one bad stop-rate away from break-even**. That is the actionable frame, and it is why the exit config below matters more than any gate roster change.

Fact table 3 — the tick-honest counterfactual: what the robot would have done with your twelve

This is the fairest reading of your question, because **the entries were automated** — a gate fired, the desk opened the slot, the stop was armed. Only the exit was yours. So: replay the 250 ms tape from each real entry, run each sub-slot's own exit stack (native 1.0×ATR stop, then its configured profit exit), and see where the machine would have got out. The config in force at the time is the built-in scale-out default (**Lot A = 2.5R scalp, Lot B = wide lock-chandelier 3.5→0.5 at 6R**) — your per-gate

`exit_overrides.json` was not written until Fri 07-31 17:29 UTC, so it governed none of these.

#	SLOT	ENTRY ATR	YOU BANKED	ROBOT'S EXIT	ROBOT'S \$	WHAT THE TAPE DID AFTER YOU CLAIMED
424	<code>grind_long_A</code>	28.2	+41.5	STOP @ 4.2 min	-57.9	grazed the stop by 0.57 pt FRAGILE
429	<code>grind_long_B</code>	28.2	+156.5	STOP @ 4.1 min	-57.9	grazed the stop by 0.82 pt FRAGILE
439	<code>grind_long_B</code>	46.4	+243.5	CHANDELIER @ 69 min	+160.0	ran further, gave back more than you did
443	<code>exhaustion_short_B</code>	36.5	+117.0	STOP @ 58 min	-74.4	tape ran 86 pt against — 49.5 pt through the stop
444	<code>exhaustion_short_A</code>	36.5	+107.5	TARGET 2.5R @ 9.1 min	+180.8	hit 2.5R before the reversal — you left \$73 on the table here
445	<code>exhaustion_short_B</code>	38.7	+109.5	STOP @ 23 min	-78.8	tape ran 136 pt against — 97.6 pt through
446	<code>exhaustion_short_A</code>	38.7	+38.5	STOP @ 23 min	-78.8	same reversal, 97.3 pt through
451	<code>grind_long_B</code>	25.2	+53.0	CHANDELIER @ 48 min	+64.5	robot slightly better
452	<code>grind_long_A</code>	25.2	+111.0	TARGET 2.5R @ 11.8 min	+124.6	robot slightly better
466	<code>abs_veto_long_B</code>	22.2	+94.0	CHANDELIER @ 5.8 min	+27.0	you banked \$67 more
481	<code>abs_veto_long_B</code>	16.2	+96.0	CHANDELIER @ 25 min	+77.0	near-tie
482	<code>abs_veto_long_A</code>	15.2	+55.0	TARGET 2.5R @ 11.6 min	+74.4	robot better
—	TOTAL	—	+1,223.0	5 stops / 4 targets / 3 chandeliers	+360.5	your hands added +\$862.5

SCENARIO	THOSE 12 TRADES	WEEK NET	READING
As it happened (your hands)	+1,223.0	+625.0	the reported number
Fully hands-off, config as loaded	+360.5	-237.5	the honest answer to your question
Hands-off, but with Friday's <code>exit_overrides.json</code>	+761.2	+163.2	one config file was worth ~\$400 on twelve trades

Caveat, stated plainly: two of the five counterfactual stops (#424, #429) are sub-one-point grazes — 0.57 and 0.82 pt through — and a resting stop that thin is genuinely a coin-flip on fill. Treat $\pm \$116$ of the $+\$360.5$ as noise.

The three exhaustion stops are not close: the tape ran 49 to 98 points beyond the stop. And note the direction of the remaining error — the reprice fills at the exact stop level with no slippage, and our own live-vs-model work says real stops fill $1.1\text{--}3.1\times$ worse. So **+\$862.5 is a floor on what your hands were worth, not a ceiling.**

Fact table 4 — re-running the SURVIVOR — steadiest verdict on automated fills

Part 1 nominated exhaustion_short: “green on 4 of 5 sessions, 54% of trades won, +\$316 for the week, the steadiest earner on the desk.” Here is that claim, automated only. The critical fact Part 1 misses is that **exhaustion_short ran two completely different exit stacks this week:** until Wed 07-29 19:10 UTC it was a single slot on the regime-3 adaptive chandelier; from that moment it was two scale-out sub-slots on Lot A 2.5R + Lot B wide lock-chandelier, with adaptive_exit switched **off**.

SLICE	N	NET \$	WIN %	\$/TR	PF	GREEN DAYS	Σ PER TRADE
Part 1’s claim (incl. hand-claims)	48	+316.5	54	+6.6	—	4 / 5	—
Automated only, whole week	44	−56.0	50.0	−1.3	0.96	3 / 5	79.7
...of which: single-slot, adaptive chandelier (to Wed 19:10)	37	+365.0	56.8	+9.9	1.46	3 / 3	78.9
...of which: scale-out A@2.5R + B wide-lock (from Wed 19:10)	7	−421.0	14.3	−60.1	0.11	0 / 3	53.8

Verdict: the money claim fails, the mechanism claim survives — conditionally. As a gate, exhaustion_short made **−\$56 automated**. It is still the steadiest thing on the desk by every dispersion measure — 50% win, a median trade of essentially zero (+\$0.8), PF 0.96, green on 3 of 5 sessions, and the smallest leave-one-day-out damage of any gate that traded

in size (−\$151, against −\$618 for `abs_veto_long` and −\$731 for `grind_long`) — but “steadiest” now means **flattest**, not **best earner**. Part 1’s offensive evidence for it, the +\$372.5 four-in-a-row on Thursday afternoon, was **four hand-claims, zero automated fills**, and under the exit config actually loaded that sequence would have been **−\$51.2**. That is the single most load-bearing correction in this section.

What actually killed it was a config change, not the tape. The pre-Wednesday configuration was non-negative in all three regime buckets it traded — clean-trend +\$436.5 (11 fills, PF 3.66), in-between-building −\$84.5 (13, 61.5% win), violent-whipsaw +\$13.0 (13, PF 1.04). That is genuine regime robustness, on small n. The scale-out configuration then lost −\$321.5 on four fills in violent-whipsaw alone, at 0% win. A fader was put on a 2.5R target and a wide trend-riding trail; faders do not get 2.5R in a whipsaw. Your hands on Thursday were, in effect, **manually re-implementing the tight exit the slate had just removed** — and your Friday 17:29 `exit_overrides.json` (exhaustion 0.5 / 1.5) was you noticing the same thing and automating it.

Honesty on concentration: the +\$365 pre-Wednesday slice is not evenly earned — its best single trade is 28% of gross profit and stripping its top three leaves −\$285 over 34 fills. It is a real but thin edge, and it is in-sample.

Fact table 5 — the regime split of the automated book (never blanket)

Regime keyed on the 1-minute ATR-14 level, the 30-bar efficiency ratio and a 30-minute range-break at the entry bar — not the clock. Time-of-day reported separately, because it is a different axis.

REGIME AT ENTRY	N	NET \$	WIN %	\$/TR	PF	READ
clean-trend (ER≥0.20, net ≥30 pt, range-break)	64	+1,273.5	40.6	+19.9	1.67	the only bucket that pays
	40	−571.0	35.0	−14.3	0.54	

REGIME AT ENTRY	N	NET \$	WIN %	\$/TR	PF	READ
in-between / building ($ER \geq 0.20$, no break)						the false-start bucket
normal-chop	9	-19.0	22.2	-2.1	0.91	n too small to read
violent-whipsaw ($ATR \geq 14$, $ER < 0.20$)	81	-1,281.5	32.1	-15.8	0.58	42% of all fills, the whole loss
dead-chop ($ATR < 6$)	0	—	—	—	—	ATR floors are doing their job

Read that as a policy statement, not a score: **automated, the desk is + \$1,274 when the tape has a real break and - \$1,872 when it doesn't**, and it takes 42% of its trades in the single worst bucket. But apply the same discipline to that headline — **clean-trend excluding Wednesday is 53 fills, - \$427.5, PF 0.76**. Clean-trend is not a licence to print; Wednesday is.

TIME OF DAY (UTC, ENTRY)	N	NET \$	WIN %	\$/TR	PF	
overnight / pre-open (<13:30)	70	-197.5	34.3	-2.8	0.90	
US open 13:30-15:00	48	-137.0	43.8	-2.9	0.92	
US session 15:00-20:00	42	-422.0	33.3	-10.0	0.72	
post-20:00 close-out window	34	+158.5	26.5	+4.7	1.14	

Worth flagging against the standing prior: the “big 4R/6R runners live in the US session” assumption does **not** hold for this week’s automated book — the US afternoon is the worst window on the desk (- \$422, PF 0.72) and the only green window is the late close-out hour, on 26.5% win and a fat tail. One week, small n, but it contradicts the prior and should not be quietly averaged away.

Fact table 6 — sweeping the exit Rs (they were guesses; here is what the tape says)

Every entry the exhaustion gate produced (64 events, 07-21 → 07-31, all automated entries plus the four whose exits you claimed) replayed on 250

ms ticks against eleven exit configs, one tick pass per entry so all configs see identical tape. Stop is always the native $1.0 \times \text{ATR}$; hold capped at 120 min / session-flat. Robustness columns are the point: **LODO-worst** = net after deleting the single worst-case day, **strip-3** = net after deleting the three best trades.

EXIT CONFIG	N	NET \$	WIN %	LODO WORST	STRIP TOP-3	GREEN DAYS	ROBUSTNESS READ
Lot-A scalp 0.5R YOUR GUESS	64	+305.7	54.7	+59.3	+151.4	4 / 9	weakest scalp rung
Lot-A scalp 0.75R	64	+510.1	50.0	+238.8	+276.4	5 / 9	plateau top for Lot A
Lot-A scalp 1.0R	64	+449.7	45.3	+204.1	+170.2	5 / 9	on the plateau
scalp 1.5R YOUR GUESS (Lot B)	64	+545.9	37.5	+220.7	+143.7	4 / 9	fine net, worse steadiness
scalp 2.0R	64	+986.2	34.4	+259.9	+448.4	4 / 9	tail-dependent
scalp 2.5R CONFIG IN FORCE	64	+1,289.4	31.2	+162.0	+616.0	4 / 9	52% of net is 3 trades
scalp 3.0R	64	+1,673.7	29.7	+425.8	+864.8	4 / 9	monotone $\Rightarrow$ no plateau $\Rightarrow$ suspect
Lot-B tight chandelier k1.5	64	+502.2	48.4	+330.1	-44.8	5 / 9	highest LODO floor; worst day only -\$25
Lot-B wide chandelier k3.5	64	+1,282.5	32.8	+361.5	+396.4	4 / 9	fat-tail bet
Lot-B lock-chandelier 3.5 $\rightarrow$ 0.5@6R CONFIG IN FORCE	64	+2,648.4	25.0	+321.6	+1,010.9	3 / 9	ARTIFACT — 88% of net (+\$2,327) is Mon 07-27 alone; 25% win; only 3 green days
fixed 8 pt / 12 pt (the old design)	64	-48.0	31.2	-113.0	-115.5	3 / 9	dead — do not revive

Proven vs guessed, stated straight. Your fader prior was 0.5 / 1.5. The tape says the tightness instinct is right and the specific numbers are not: 0.5R is the **worst** scalp rung on the ladder (+\$306, 4 of 9 green days), and 0.75-1.0R is a genuine two-rung plateau (+\$510 / +\$450, both 5 of 9 green, both with a LODO floor above +\$200). On Lot B, the fixed 1.5R scalp and the k1.5 chandelier post near-identical net (+\$546 vs +\$502) but the chandelier is far steadier — **the highest leave-one-day-out floor in the whole sweep (+\$330), five green days, and its three red days are -\$21, -\$1.50 and -\$25.** The lock-chandelier that was actually loaded posts the biggest number on the board and it is a one-day artifact: **+\$2,327 of its**

+\$2,648 is Monday 07-27, it wins one trade in four, and it is green on only 3 of 9 days. **Do not read that +\$2,648 as an edge.**

Now the regime split of the same sweep — because a single ladder across all tape is exactly the bug the discipline forbids. Automated entries only (n=62).

REGIME	N	BEST CONFIG (NET \$)	RUNNER-UP	WHAT LOSES	POLICY
clean-trend	11	lock-chandelier +839.0	2.5R +407.6 · wide k3.5 +546.4	k1.5 tight −50.4 · 0.5R −70.2	ride it — wide/lock, never tight
in-between / building	15	lock +1,948.9 (artifact-inflated)	k1.5 tight +673.7 @ 80% win · 0.75R +561.0 @ 80%	0.25R +216.2 (leaves money)	tight + 0.75R , the honest pick
normal-chop	8	lock +93.3 @ 12.5% win	everything ≈ 0	0.5–1.5R all red	n too small — don't tune on it
violent-whipsaw	28	wide k3.5 +149.7 · 2.0R +144.9	0.5R −12.3 (least bad tight)	k1.5 tight −202.4 · 1.0R −203.3 · 1.25R −242.9	DON'T TRADE IT — no config is robustly green

That fourth row is the important one and it cuts against the tight-exit story: in violent whipsaw, tightening does not save the fader — the k1.5 chandelier is the single worst config in that bucket (−\$202). This is the bucket the gate took 28 of its 62 entries in. The lever is not the exit, it is **not being armed**. That matches the standing ledger lesson (07-31: violent post-thrust chop → go fully flat) and it is the recommendation.

The momentum prior fares worse. Sweeping abs_veto_long's 42 automated entries the same way: the plateau is at **0.75-1.0R (+\$292 / +\$532, 54.8% win)** and it collapses at exactly the numbers you guessed — 1.5R +\$26, 2.0R −\$42, 2.5R +\$75 — with the lock-chandelier at −\$16 overall and **−\$415 in clean-trend**. The 1.5 / 2.5 momentum prior is **falsified for abs_veto_long on this tape**; it is holding a thrust far too long. For abs_veto_short the prior survives: the plateau is 1.25-2.0R (+\$178 / +\$202 / +\$188) and **1.5R is the only config in that sweep with a positive leave-one-day-out floor (+\$45.4, 3 of 4 green days)**. Same family, opposite sides, opposite answers — which is itself the argument for per-side sweeps over roster-wide priors.

The re-nomination

Part 1 needs a survivor that stands up with nobody touching it. Here is the ranked answer, with the honesty attached.

RANK	NOMINATION	AUTOMATED EVIDENCE	WHY IT SURVIVES	WHAT WOULD BREAK IT
1	exhaustion_short on the tight/adaptive exit, armed outside violent-whipsaw — a policy, not a gate	+\$365 / 37 fills / 56.8% win / PF 1.46 / 3 of 3 green days, all automated, zero hand-claims. Sweep: k1.5 chandelier +\$502, LODO floor +\$330, 5 of 9 green days, worst day -\$25.	The only thing on the desk that was non-negative in every regime bucket it traded, and the only sweep config whose leave-one-day-out floor stays comfortably positive. It is also the only nomination whose green does not depend on Wednesday.	Its top trade is 28% of gross profit; strip the top three and it is -\$285. n=37. In-sample. And it needs the whipsaw veto — armed in whipsaw it loses at any R.
2	abs_veto_short, armed only in clean-trend, Lot A/B ≈1.5R	+\$358.5 / 26 fills / 42.3% / PF 1.53 automated with no hand-claims at all. Clean-trend slice: +\$640.5 / 11 fills / 54.5% / PF 3.38. Violent-whipsaw slice: -\$282 / 15 fills.	The only gate that is automated-green on the raw ledger, and the split is textbook: it pays on the break and bleeds on the chop. 1.5R is the one positive-LODO rung in its sweep.	52% of its week net is one trade (+\$536, Wed). Drop Wednesday and it is -\$153. n=11 in its home regime. This is a nomination to watch, not to size.
—	grind_long	-\$421.5 / 51 / 25.5% / PF 0.80 / 1 of 4 green days	nothing survives — its Part 1 green was entirely hand-claims	—
—	abs_veto_long	-\$185.5 / 56 / 32.1% / PF 0.90 / 1 of 6 green days	nothing survives; and its exit R is falsified above	—
—	capitulation_long · rgv_short	-\$139 / 14 · -\$154.5 / 3	unchanged from Part 1 — rehab and bench respectively	—

So the corrected lead for Part 1 reads: “Hands off, the desk lost \$598 this week. Nothing on the roster is robustly green. The closest thing to a survivor is exhaustion_short — but only in the configuration it was running until Wednesday evening, and only when it is kept out of violent whipsaw; the Thursday sequence Part 1 celebrates was hand-made, and under the exit config the desk was actually running it would have lost money.”

Three things to actually do

1. **Put exhaustion_short back on a tight exit and stop feeding it whipsaw.** The Friday override (0.5 / 1.5) was the right instinct at the wrong numbers — the sweep says **Lot A 0.75R, Lot B k1.5 chandelier**, and it says the exit cannot rescue a violent-whipsaw entry at all, so the router bench matters more than the R. This is the single highest-value change in this section: it is worth roughly +\$420 on the twelve claimed trades alone, and it converts the hand-claim habit into a mechanical rule.
2. **Re-sweep abs_veto_long's exit before next week.** Its loaded 1.5/2.5 config sits on the falsified side of a clear plateau; 0.75–1.0R is +\$532 against +\$26 on this tape, on 42 entries. This is the cheapest red-to-green candidate on the desk.
3. **Instrument the claim button.** Twelve hand-claims at 12-for-12 is not luck, it is a signal the machine is missing — and right now it is invisible in every P&L we report. Every MANUAL_CLAIM should be auto-repriced against what its slot's exit would have done and logged as a delta, so “Garrath's hands are worth \$X/week and here is the pattern they exploit” becomes a standing number instead of a Friday archaeology dig.

METHOD & LIMITS — READ BEFORE QUOTING ANY NUMBER ABOVE

Venue truth vs model. Every figure in fact tables 1, 2, 4 and 5 is straight from `data/gazbot7.db` trades — real fills, fees included, nothing modelled. Fact tables 3 and 6 are reprices on 250 ms ticks from `data/capture.db` with \$1.50 round-turn fees, fills assumed at the exact stop/target level and no slippage. Our own execution work says live stops fill 1.1–3.1× worse than modelled, so **every reprice number is a ceiling**; use them to rank configs against each other, never as a P&L promise.

Sample size. 194 automated fills over six sessions. Regime buckets: violent-whipsaw 81, clean-trend 64, in-between 40, normal-chop 9, dead-chop 0. The sweeps use 64 (exhaustion), 42 (abs_veto_long) and 22 (abs_veto_short)

entry events over nine, six and four days. Anything below $n \approx 15$ in this section is flagged inline and should be treated as a direction, not a measurement.

Known modelling weakness. 17 of the 64 exhaustion replays never hit a stop, target or trail inside the 120-minute / session-flat cap and are marked at the last tick — the same 17 entries in every config, so they do not tilt one config against another, but they do mean roughly a quarter of the sweep is marked-to-market rather than exited. Separately, an entry that hits its target before the tape later runs through its stop is scored as a win (correctly — the order of events is respected), which is why trade #444 in fact table 3 shows a target despite the tape eventually running 85 pt against it.

Not swept, and it should be. grind_long's 2.5 / wide pair was not put through this ladder — it is the largest automated loser on the desk (–\$421.5) and it has the two graze-sensitive counterfactuals in fact table 3. That sweep is the obvious next job and its absence is a hole in this section, not a silence.

Reconciliation. Week bounds and totals match Part 1 exactly (206 fills, +\$625.0); the day-by-day rows here use the UTC close-date rather than Part 1's Paris session boundary, which moves trades between days but not the week total or any gate total.

REV2 · Q2 Is abs_veto_long already live — and if it is, what exactly am I promoting?

Garrath — you are right and Part 2 is wrong. **abs_veto is live on both sides and has been since 2026-07-25**

(slot_strategy.tournament_slots(), commit eb57f64); abs_veto_long is armed in gate_switches.env right now and traded **59 lots across 43 signals for +\$59.5** this week. There is **no slot to promote**. What is actually on the table is a **four-line config diff on a gate that is already**

running — and the diff is not the one Part 2 proposed. Chasing your +\$59.5-vs-+\$776 question turned up three things: (1) the two numbers are measured over **different windows**, so the real gap is **\$609.5**, not \$717; (2) **82% of that gap is entry, not exit** — the live desk only ever saw 43 of the shadow's 69 long signals, because the router had the gate benched for 26 of them; and (3) while checking why the ER floor did not block anything, I found that **the dual-slot scale-out slate silently killed every ER and ATR floor on the desk on 2026-07-29**. Both this report and the router have spent the week reasoning about filters that are not running.

HONESTY BANNER. EVERY NUMBER BELOW IS FROM THE LIVE GAZB0T7.DB TRADE LEDGER, THE TICK-REPRICED SHADOW_REAL.REAL_PNL, THE GAZB0T7-TOURNAMENT JOURNAL, AND 10.08M CAPTURED MNQ TICKS — NOTHING IS RECALLED. WINDOWS ARE STATED ON EVERY ROW. THE TAPE IS ONE SUMMER WEEK, CHOP-DOMINATED, MNQ-ONLY, PAPER: THE CLEAN-TREND BUCKET IS N=6 LONG / N=5 SHORT AND IS REPORTED BUT NOT SCORED. THE SHADOW RUNS 1 LOT, THE LIVE DESK RUNS 2 — NO RAW DOLLAR COMPARISON BETWEEN THEM IS VALID WITHOUT THAT ADJUSTMENT, AND I MAKE IT EXPLICIT EVERYWHERE BELOW.

1 · The premise — both numbers are real, but they are not the same measurement

The +\$776 is the shadow's **Mon-to-Fri calendar week**. The +\$59.5 is the live gate's **whole life this week**, which starts at the Sunday 22:00 UTC reopen. The shadow lost \$106.5 on that Sunday night; crediting it with a session the live desk actually traded inflates the gap by \$107.

WINDOW	ENGINE	SIGNALS	LOTS	NET	NOTE	
Mon 07-27 → Fri 07-31	shadow abs_veto_55s LONG		66	66	+\$775.5	the "+\$776" in the question
Mon 07-27 → Fri 07-31	LIVE abs_veto_long		40	56	+\$131.0	the like-for-like live number for that window
Sun 07-26 22:00 → Fri 07-31	shadow abs_veto_55s LONG		69	69	+\$669.0	the honest comparison window (live gate's first fire → last)
Sun 07-26 22:00 → Fri 07-31	LIVE abs_veto_long		43	59	+\$59.5	the "+\$59.5" in the question

WINDOW	ENGINE	SIGNALS	LOTS	NET	NOTE
Sun 07-26 22:00 → Fri 07-31	shadow abs_veto_55s SHORT		63	63	+\$659.0
Sun 07-26 22:00 → Fri 07-31	LIVE abs_veto_short		22	26	+\$358.5
					the short side, same window
					live SHORT beat shadow per signal: +\$16.3 vs +\$10.5

And the trade counts are not near-identical. 59 live "trades" are 59 lots — under the dual-slot scale-out every signal opens **two** sub-slots (_A and _B). Collapse them and the live desk took **43 signals against the shadow's 69**: it saw **62%** of the opportunity set, at **2× the size** on each one it did take. That single line reframes the whole question — it is not "the same trades earned less", it is "a different, smaller, double-sized trade set".

2 · The config diff — live abs_veto_long/_short vs shadow abs_veto_55s

The **entry is byte-identical**. Everything that differs is arming, sizing, exit and execution.

AXIS	LIVE (AS OF 2026-08-01)	SHADOW ABS_VETO_55S	SAME?
entry gate	gate_thrust(thr=1.5, amp_floor=0.0004)	gate_thrust(thr=1.5, amp_floor=0.0004)	IDENTICAL
amp_floor	0.0004	0.0004	IDENTICAL
absorption veto	55s wait, re-test thrust still fires + exit_absorption(flow_min=50); abandons a signal older than 75s	confirm_s=55, same re-test, absorption_flow_min=50; no abandon clock	same mechanism, live has one extra guard
ER floor	ER_FLOOR["abs_veto_long"]=0.25 — NOT IN FORCE (see §3)	none	identical in fact, different on paper
ATR floor	ATR_FLOOR["abs_veto_short"]=16.0 — NOT IN FORCE (see §3)	none	identical in fact, different on paper
size per signal	2 lots — Lot A + Lot B, base_size=1 each	1 lot (qty=1.0)	DIFFERENT — 2×
Lot A exit	scalp 1.5R, stop 1.0×ATR, give-back off, adaptive_exit=False (exit_overrides.json, set 07-31 17:37)	— (single lot)	DIFFERENT
Lot B exit	scalp 2.5R, stop 1.0×ATR — before 07-31 17:37 this was the "wide" lock-chandelier (start_k 3.5 → lock_k 0.5 past 6.0R)	single lot at target_r 2.0, stop 1.0×ATR	DIFFERENT — and it changed mid-week

AXIS	LIVE (AS OF 2026-08-01)	SHADOW ABS_VETO_55S	SAME?
catastrophe cut	exit_absorption backstop, min loss \$60	adverse_cut_atr=1.5 + absorption cut (2 fires, −\$158 this week)	different backstop
regime-3 exit selector	OFF under scale-out (adaptive_exit=False on both sub-slots); was ON in the pre-07-29 single-slot era	n/a	live-only, and it silently switched off
stop execution	resting broker stop — STOP_UNFILLED path is live: 4 long lots −\$167.5, 2 short lots −\$106.0 this week	modelled stop, always fills	DIFFERENT — live-only leak of \$273.5
arming	router benches it — 26 long / 25 short signals blocked this week	never benched, fires every signal	DIFFERENT — the biggest term
fees	\$1.50/RT × 2 lots = \$3.00 per signal (\$88.5 on the long book)	\$1.50/RT × 1 lot, already inside real_pnl	DIFFERENT
P&L basis	venue truth, PAPER fills	tick-repriced shadow_real.real_pnl	comparable
median hold	4.1 min (n=85 lots)	8.0 min (n=132)	live exits ~2× faster

The live veto is doing almost all its work on the short side. From the tournament journal, 07-26 → 07-31: abs_veto_long raised 63 thrusts and **opened 59, skipped 4 as stale, vetoed 0 as absorbed**. abs_veto_short raised 80 and **opened 30, skipped 28 as stale, vetoed 22 as absorbed**. So the "absorption veto" is a short-side filter in practice; on the long side it is a 55s delay and nothing more. That is worth knowing before anyone calls the veto the reason abs_veto works.

3 · ★ The thing I did not go looking for: every ER and ATR floor on the desk has been dead since 07-29

tournament.step() gates each intent with er_blocks(slot, er) and atr_blocks(slot, f.atr) — passing the **slot tag**. Under the dual-slot scale-out slate the tags are abs_veto_long_A, abs_veto_long_B, grind_long_A... but ER_FL00R and ATR_FL00R are keyed on the **base** names. The lookup never matches. The same function has a _base() helper and it is applied to the bench switch (line 250), the 55s veto (line 303) and the exhaustion confirm (line 345) — but **not** to the two floors.

EVIDENCE	WHAT IT SHOWS
code	tournament.py:139/142 — er_blocks(slot, er) / atr_blocks(slot, f.atr) with the sub-slot tag; deciders.ER_FLOOR / ATR_FLOOR keyed on base names.
journal	7,734 ER gate: / ATR gate: suppression lines between 07-26 and the very last one at 2026-07-29 13:29:04. The scale-out sub-slots first appear in the journal at 2026-07-29 16:24. Zero suppressions since.
trades	abs_veto_long_A opened at 30-min ER 0.013 (07-30 10:09, + \$77.5 TARGET), 0.035 (07-30 05:48), 0.049 (07-30 12:32) — all far below the 0.25 floor that is supposedly gating it.

Which floors are affected: abs_veto_long ER 0.25, grind_long ER 0.35 and ATR 24, abs_veto_short ATR 16, capitulation_long ATR 10. All of them, since 07-29 16:24. Note the date on the abs_veto floor: commit 9767c67, **2026-07-30** — it was written the day after the slate change disabled it, so **it has never once run live**.

THIS CORRUPTS TWO CONCLUSIONS ELSEWHERE IN THIS REPORT. PART 2.6 QUOTES THE ROUTER'S OWN 07-30 11:08 TICK — "GRIND RE-ARMED BUT ITS 0.35 FLOOR AUTO-GATES IT, SO LOW-RISK" — AND THE 11:13 POST-MORTEM "FIRED AT ER≥0.35 (FLOOR PASSED)". NEITHER IS TRUE; GRIND HAD NO FLOOR, AND THE 11:13 LOSS IS NOT EVIDENCE ABOUT THE FLOOR'S QUALITY. ANY PART 1.5 / PART 2.6 LINE THAT REASONS FROM A LIVE ER OR ATR FLOOR AFTER 07-29 16:24 NEEDS RE-READING WITH THE FILTER REMOVED.

Do NOT rush to wire it up. The counterfactual on abs_veto_long's 31 scale-out-era signals is honest but tail-dominated:

POPULATION	SIGNALS	NET	MEDIAN	STRIP BEST-1	STRIP BEST-2
would have been BLOCKED by a working ER≥0.25	24	+\$494.5	−\$42.8	+\$18.0	−\$233.5
would have been KEPT (ER ≥ 0.25)	7	−\$168.5	—	—	—

Read it straight: the floor would have blocked the week's two biggest winners (+\$476.5 and +\$251.5) and kept a −\$168.5 book — but the blocked population's **median is −\$42.8** and its edge vanishes after stripping one trade. So there is **no live evidence for the floor and none robustly against it**. It is simply **untested**, and the honest position is: fix the wiring so the desk stops lying about itself, then decide the floor's value on its own evidence — do not let a silent bug make the decision.

4 · The reconciliation — same signals? same fills? where did the \$609.5 go?

Method: match every shadow abs_veto_55s fire to a live signal within ± 180 s. For every shadow fire the live desk did **not** take, look for a SUPPRESSED-OPEN ... (disabled by switch) line in the tournament journal at that moment — that is ground truth that the gate fired and the router's bench blocked it. Anything left over is a live gate that never raised the signal at all.

LONG — SHADOW + \$669.0 (69 FIRES) → LIVE + \$59.5 (43 SIGNALS / 59 LOTS)	FIRES	SHADOW \$	LIVE \$	EFFECT
ENTRY — router had the gate BENCHED. Journal shows the thrust fired and the switch blocked it.	26	+\$397.0	\$0	−\$397.0
ENTRY — live gate never raised the signal. Service restarts (23 in this window, 6 of them code=dumped crashes), the 10s stale-tape / 30s stale-bar freshness gates, and the veto's "thrust gone/stale" skip.	18	+\$102.5	\$0	−\$102.5
EXIT + SLOT SPLIT + EXECUTION — the ~25 SHARED signals. Shadow: 25 lots @2.0R = +\$169.5. Live: 38 lots (A@1.5R + B) = +\$49.0. Live had 1.5× the size and still made \$120 less.	25	+\$169.5	+\$49.0	−\$120.5
CREDIT — live-only signals the shadow never fired	—	\$0	+\$10.5	+\$10.5
TOTAL	69	+\$669.0	+\$59.5	−\$609.5

Answer to the question as asked: 82% entry, 20% exit/slot, and the exit term is smaller than it looks. Of the −\$120.5 on the shared signals, roughly **\$20 is pure fee** (38 live lots at \$1.50 vs 25 shadow lots) and

\$167.5 is the STOP_UNFILLED execution path (four pre-scale-out lots on 07-26/27/28 where the resting stop did not trigger — a bug with a fix on a branch, memory `[[stop-unfilled-contfuture-root-cause]]`, still not deployed). Strip fees and the stop bug and the **exit-design** term is roughly **+\$67 in live's favour**. **The dual-slot exit is not what cost you the money. The bench did.**

SHORT — SHADOW + \$659.0 (63 FIRES) → LIVE + \$358.5 (22 SIGNALS / 26 LOTS)				
	FIRES	SHADOW \$	LIVE \$	EFFECT
ENTRY — router BENCHED. Those signals lost money in shadow — the bench saved \$157.	25	−\$157.0	\$0	+\$157.0
ENTRY — live gate never raised it. The short side's veto skipped 28 as stale and vetoed 22 as absorbed; the shadow's veto has no 75s abandon clock.	19	+\$620.0	\$0	−\$620.0
ENTRY — live slot already in a position	1	−\$41.0	\$0	+\$41.0
SHARED signals — live WON the exit. Shadow 18 lots @2.0R = + \$237.0; live 21 lots = + \$422.5, carried by one A+B pair on 07-29 (+ \$232 TARGET / +\$536 wide chandelier).	18	+\$237.0	+\$422.5	+\$185.5
live-only signals	—	\$0	−\$64.0	−\$64.0
TOTAL	63	+\$659.0	+\$358.5	−\$300.5

The two sides tell opposite stories, and that is the finding. The router's long-side benching **cost \$397**; its short-side benching **saved \$157**. Same gate, same veto, same week — the bench was right on one side and wrong on the other. Which one it got right is determined entirely by regime, so here is the regime split of exactly the signals it blocked:

REGIME AT THE BLOCKED FIRE	LONG BLOCKED	SHADOW \$ FORGONE	SHORT BLOCKED	SHADOW \$ FORGONE	VERDICT ON THE BENCH
NORMAL-CHOP (mid ATR, ER<0.25)	13	+\$384.0	12	+\$185.0	WRONG on both sides — the single biggest error of the week

REGIME AT THE BLOCKED FIRE	LONG BLOCKED	SHADOW \$ FORGONE	SHORT BLOCKED	SHADOW \$ FORGONE	VERDICT ON THE BENCH
BUILDING (ER 0.25-0.45)	10	+\$70.5	7	-\$35.0	wrong long, right short, both small
CLEAN-TREND (ER ≥ 0.45)	2	+\$17.0	2	-\$71.5	n too small to score
VIOLENT-WHIPSAW (ATR ≥ Q75 19.2, ER<0.25)	1	-\$74.5	4	-\$235.5	RIGHT — the bench saved \$310 here

Regime keyed on ATR level + 30-min ER + 60-bar range-break over 16,695 one-minute bars rebuilt from the 5s capture (ATR quartiles Q25 11.0 / Q50 14.4 / Q75 19.2 pt), not on the clock. Session split reported below.

5 · Regime × session — where abs_veto actually earns, both engines side by side

Live is per **lot** (so the dual-slot days count twice); shadow is per **fire**. Read the shapes, not the raw dollars.

LONG · REGIME × SESSION	LIVE LOTS	LIVE NET	LIVE WIN%	SHADOW N	SHADOW NET	SHADOW WIN%
NORMAL-CHOP · overnight/pre-open	14	+\$383.5	50%	24	+\$525.0	58%
NORMAL-CHOP · US cash (≥13:30 UTC)	8	+\$2.5	38%	8	+\$0.5	38%
VIOLENT-WHIPSAW · US cash	13	+\$1.5	23%	10	+\$146.0	50%
VIOLENT-WHIPSAW · overnight	6	-\$50.0	33%	4	-\$37.5	25%
BUILDING · overnight	5	-\$199.0	0%	12	+\$36.5	42%
BUILDING · US cash	5	+\$17.0	60%	5	+\$61.0	60%
CLEAN-TREND · US cash (n too small — not scored)	8	-\$96.0	38%	6	-\$62.5	33%

SHORT · REGIME × SESSION	LIVE LOTS	LIVE NET	SHADOW N	SHADOW NET
NORMAL-CHOP · overnight	5	−\$21.5	20	+\$351.5
BUILDING · overnight	4	−\$146.5	12	+\$204.5
BUILDING · US cash	2	+\$110.5	9	+\$172.5
VIOLENT-WHIPSAW · US cash	4	−\$202.5	7	−\$125.0
VIOLENT-WHIPSAW · overnight	1	−\$54.5	4	−\$60.5
CLEAN-TREND · US cash (n=3/3 — not scored)	3	+\$791.0	3	+\$101.0

The shape is consistent across both engines and both sides: **abs_veto** is an **OVERNIGHT NORMAL-CHOP** gate. That is the opposite of how the router treated it — the durable tick's standing rule is "confirmed chop → bench all momentum", and **abs_veto** is filed as momentum. It is the single rule that cost the most this week. The one place both engines agree it should be off is **VIOLENT-WHIPSAW on the short side**: live −\$257 across 5 lots, shadow −\$185.5 across 11 fires.

6 · The 1.5/2.5 exit cells — swept, not assumed

The Lot-A/Lot-B Rs in `exit_overrides.json` are an operator prior, so I swept them tick-honest. Method: 150 **abs_veto_55s** entries inside the 250ms tick coverage (10,080,959 MNQ ticks, 2026-07-23 22:00 → 07-31 21:00), each replayed forward from the fill — a lot banks $+R \times \text{ATR}$ if the favourable level prints first, stops at $-1.0 \times \text{ATR}$ if the adverse level prints first, else marks out at the 90-min cap (no shadow trade reached it; max hold was 62 min). **Replay validated against the shadow repricer: 98.7% sign agreement, MAE \$2.70 per trade at the shadow's own 2.0R setting** — so the sweep is measuring the exit, not my simulator.

LONG — \$ PER SIGNAL (LOT A + LOT B COMBINE D, 1 LOT EACH)	BUILDING N=21	CLEAN- TREND N=6*	NORMAL- CHOP N=40	VIOLENT N=14	ALL N=81	LODO MIN	STRIP BEST-3
A 1.0 / B 1.5 — the robust plateau	+15.6	−3.3	+29.3	+25.9	+22.7	+16.8	+19.0
A 1.0 / B 2.0	+14.1	−16.7	+29.8	+27.8	+22.0	+14.8	+17.3
A 0.75 / B 1.5	+9.7	+4.5	+25.1	+22.6	+19.1	+13.9	+15.7
A 1.5 / B 4.0	+4.8	+19.5	+32.1	+22.8	+22.5	+12.5	+13.2
A 1.0 / B 2.5	+5.2	−10.2	+22.2	+26.0	+16.0	+9.0	+10.3
A 2.0 / B 3.0	+2.1	−6.8	+26.1	+22.7	+16.9	+7.2	+8.3
A 1.5 / B 2.5 ← LIVE TODAY	−0.8	+0.2	+23.4	+12.9	+13.6	+4.7	+6.8
A 2.5 / B 6.0	−20.8	−13.6	+33.6	+21.8	+14.0	+4.0	+0.2

* CLEAN-TREND n=6 — shown for completeness, **not scored**. LODO = leave-one-day-out over 8 days; every cell is positive on 8/8 folds, so LODO separates them by **margin**, not sign. "strip best-3" removes the three largest winners.

SHORT — \$ PER SIGNAL	BUILDING N=21	CLEAN- TREND N=5*	NORMAL- CHOP N=32	VIOLENT N=11	ALL N=69	LODO MIN	STRIP BEST-3
A 2.5 / B 4.0	+19.3	+66.8	+23.5	−19.9	+18.4	+4.9	−1.8
A 1.0 / B 2.5	+30.9	+53.5	+17.5	−29.4	+16.7	+9.9	+4.1
A 1.5 / B 2.5 ← LIVE TODAY	+24.8	+41.3	+23.4	−30.8	+16.5	+10.5	+1.9
A 0.75 / B 1.5	+23.7	+46.3	+21.3	−33.0	+15.2	+9.4	+7.3

Verdict on the guess. On the **LONG** side the live cell A1.5/B2.5 is the **worst of the eleven cells swept, on every measure**: +\$13.6/signal vs + \$22.7 for A1.0/B1.5, LODO-min +\$4.7 vs +\$16.8, strip-best-3 +\$6.8 vs + \$19.0. The plateau is broad and sits at **A≈1.0, B 1.5-2.0** — a genuine parameter plateau, not a spike (0.75/1.5, 1.0/1.5 and 1.0/2.0 are all within \$4 of each other). On the **SHORT** side every cell lands between +\$15 and + \$18 per signal and the LODO/strip ordering flips between them — **the R choice is noise on the short side, so leave it alone**. What is not noise on the short side is the VIOLENT-WHIPSAW column: **−\$20 to −\$33 per signal at every single cell and every single R from 0.5 to 6.0**. No exit tuning rescues that segment; only not trading it does.

And the engine underneath is real. The 55s veto against its no-veto control on the same tape: **abs_veto_55s +\$2,668 on 238 fires (+\$11.2/trade)** vs **thrust_loose +\$130.5 on 424 fires (+\$0.3/trade)** — and this week specifically, +\$1,328 vs −\$235.5. Nobody should read this section as "abs_veto is broken". The gate is the best thing on the board. Its **arming policy** is what leaked.

7 · The promotion, restated as a config diff on a running gate

There is no new slot, no roster change, and nothing that needs a Saturday tournament window except D1. Here is the whole of it.

#	THE ACTUAL CHANGE	THE NUMBER BEHIND IT	FILE / OWNER	REVERT
D0	Delete the recommendation. "Promote abs_veto_55s two-sided" is not an action — it shipped 2026-07-25. Part 2 re-validated a live gate cold and mislabelled it. abs_veto_long=on and abs_veto_short=on in gate_switches.env right now.	7 days live · 85 lots · 65 signals · +\$418 combined this week	— report fix	—
D1	Fix the dead floors. Pass_base(slot) to er_blocks/atr_blocks in tournament.step() — or delete ER_FLOOR/ATR_FLOOR outright. Do not ship the wiring fix and the floor values together: wire it, then re-derive each floor on tape that was actually filtered.	7,734 suppressions before 07-29 13:29 · 0 after · 5 gates affected	tournament.py:139/142 · SATURDAY (today) — it changes entry maths	one-line revert; currently a no-op either way
D2	Change one cell: exit_overrides.json → abs_veto_long: {a_r 1.0, b 1.5}. Leave abs_veto_short at 1.5/2.5. This is the change that REV2·Q0 row M6 told you not to make — M6 is superseded by the sweep.	n=81 signals · +\$22.7/sig vs +\$13.6 · LODO-min +\$16.8 vs +\$4.7 · strip-best-3 +\$19.0 vs +\$6.8 · plateau not spike	data/ exit_overrides.json + tournament restart · OPERATOR	delete the key → built-in default

#	THE ACTUAL CHANGE	THE NUMBER BEHIND IT	FILE / OWNER	REVERT
D3	Router rule — bench abs_veto_short in violent whipsaw (ATR ≥ 19pt and 30-min ER < 0.25). The router did this by accident 4 times this week; make it deliberate.	n=11 shadow fires — \$185.5 · live 5 lots — \$257 · negative at all 11 exit cells and every R 0.5→6.0	gate_switches.env · ROUTER, per tick	reversible every 5 min
D4	Router rule — STOP benching abs_veto_long in overnight NORMAL-CHOP. The standing "confirmed chop → bench all momentum" rule files abs_veto as momentum; on the evidence it is an overnight chop gate. Bench it on violence, not on chop.	13 blocked fires worth + \$384 · live+shadow both green in that cell (+ \$383.5 / +\$525) · total long-bench cost —\$397	router framework · ROUTER	reversible every 5 min
D5	Deploy the STOP_UNFILLED fix (concrete front-month stops, memory [[stop-unfilled-contfuture-root-cause]], commit 99e3c11, still on a branch). It is a pure execution leak with no shadow analogue.	6 abs_veto lots this week · —\$273.5 · ~\$168 of it sits inside the long-side "exit" gap	BUILD → live-verify	branch revert
D6	Never compare a shadow dollar to a live dollar again without the x2. The shadow runs 1 lot; the live scale-out runs 2. Every "shadow made \$X" line in Part 2 needs a per-signal denominator.	59 live lots = 43 signals · 26 live lots = 22 signals	— report convention	—

The one-sentence answer. abs_veto_long is live, it has been live for a week, and the +\$59.5-vs-+\$776 gap is not an exit problem or a slot-split problem — it is **\$500 of signals the router never let it take, \$107 of window mismatch, \$273 of a stop that did not fill, and one exit cell (long-side 1.5/2.5) that the tick sweep ranks dead last of eleven.** There is nothing to promote; there is a bench rule to invert, one JSON cell to change, and a filter-wiring bug to own.

WHAT THIS DOES NOT PROVE. ONE WEEK, MNQ-ONLY, PAPER, CHOP-DOMINATED: THE CLEAN-TREND BUCKETS (N=6 LONG / N=5 SHORT) ARE UNSCORED, SO D2 AND D4 ARE UNTESTED ON A TREND WEEK — A WIDE LOT B IS EXACTLY WHAT A TREND WEEK WOULD REWARD, AND A1.0/B1.5 CAPS IT.

THE R SWEEP SPANS 8 TRADING DAYS OF TICK COVERAGE (07-23 22:00 ONWARD) AND EVERY CELL IS POSITIVE ON 8/8 LODO FOLDS, WHICH MEANS THE FOLDS SEPARATE CELLS BY MARGIN ONLY — TREAT THE ORDERING AS A LEAD, NOT A LAW. THE ER-FLOOR COUNTERFACTUAL (§3) IS TAIL-DOMINATED AT N=24 AND IS DELIBERATELY NOT USED TO ARGUE THE FLOOR'S VALUE IN EITHER DIRECTION. AND THE WHOLE LIVE SIDE OF THIS RECONCILIATION RAN UNDER A SLATE WHOSE ER/ATR FILTERS WERE SILENTLY OFF, SO EVERY LIVE ABS_VETO NUMBER IN THIS REPORT IS AN UNFILTERED NUMBER.

REV2 · Q3 BIG-TREND at n=5 — the plan to prove it, and by when

PATCH MARKER — this section **patches** PART 2 · LAB ("The exit ladder"): it corrects the *n* in that section's ★ **CRITICAL caveat** callout and in the grind (momentum) / BIG-TREND (5) row of its ladder table, withdraws that section's refusal of the MNQ tick archive as factually wrong, and replaces its "must be proven on the next genuine trend week" hand-wave with a dated, owned acquisition plan. It **extends** REV2 · Q0 (the Monday/Saturday card) with three new dated items. Everything below is computed from `scratchpad/exit_ladder_rows.pkl` (the 4,553-signal lab), `data/capture.db`, `data/shadow.db`, `alphabot2/data/ticks.db` and `alphabot2/data/alphabot.db`.

Garrath — you are right that the report left this open, and you are right that it was the worst thing to leave open. But before the plan, two corrections, because the caveat you are quoting is itself wrong in two places. **First: the BIG-TREND cell is not n=5.** Five is the count of live shadow `grind_fast` entries that landed in an $ER \geq 0.50$ minute (5 in 12 sessions). The tick-tape replay cell that the ladder verdict was actually computed on is **n=24**. Part 2 · LAB quoted the live number in a caveat about the replay number. **Second: the refusal of the MNQ archive was wrong on the facts** — and worse, the lab had already used it. The lab tape starts **2026-07-05**, not 07-15; eleven of the twenty-four BIG-TREND entries come from that archive stretch. The refusal reads well and is false. **Third, and the reason this matters more than either: the unproven rung is not a proposal — it is already live.** `grind_long` has no entry in `exit_overrides.json`, so it inherits the built-in BIG-RUN default, which is literally Lot A @ 2.5R + Lot B wide lock-chandelier. The desk has been

riding the $n=24$ guess since the scale-out slate went in. This is not a stone left unturned; it is a live position.

THE ONE-PARAGRAPH ANSWER

Grade date: Friday 2026-08-21. Hard backstop: Friday 2026-09-11. Owner: Claude (durable router + Friday max-depth lab); operator sign-off on the freeze only. The grading n is **39 paired BIG-TREND entries** — and that is a computed number, not a round one: the observed paired advantage of the RIDE exit over the tight scalp is **+\$79.4/trade with a paired sd of \$247**, so 2σ separation needs $n = 39$ ($n = 38$ at the plateau centre $A3.0+WIDE$). We are at 24. The remaining 15 come from three sources in this order: **(1)** extend the tape backwards — `alphabot2/data/alphabot.db bar_history` holds **334,818 MNQ 5-second bars covering 2026-06-19→07-17**, giving **10 full sessions no current tape has**, and those sessions are exactly the missing trend (MNQ fell **1,807 points** 06-22→06-26; $|net|/range \geq 0.59$ on 7 of 11 days) — expected **+10 entries (range 8-13)**, delivered this weekend; **(2)** re-price all 24 existing entries at **250 ms instead of 5 s** — both tick sources between them cover every one of the 12 BIG-TREND entry days, so this costs nothing but compute and it is the honest tier; **(3)** live-forward tops up the last 3-5 at the measured shadow rate of **0.42 BIG-TREND entries/session**, which is 8-12 sessions. Live-forward alone would take until **2026-09-21** all-sides and would never finish on `grind_long` alone (2 long BIG entries in 12 sessions → ~150 sessions). That is why the plan is tape-first. The MES option is **worth nothing for n** — the MES archive spans 07-05→07-17, the identical calendar days as the MNQ archive, so it adds zero new sessions; it is retained only as a falsification cross-check.

Fact table 1 — what "n=5" actually was, and what the cell really says

Two different measurements were conflated. Both are real; only one is what the ladder verdict was computed on.

MEASUREMENT	WHAT IT COUNTS	N	WINDOW	STATUS
"grind n=5" (the caveat)	live grind_fast shadow entries whose entry minute had 30-min ER ≥ 0.50	5	07-16→07-31 (12 sessions)	real, but the wrong number to caveat
grind (mom) BIG-TREND (the cell)	V7 gate_grind replayed on the 5s tape, exits repriced on the forward path	24	07-05→07-31 (12 of 20 sessions fired)	this is what the ladder verdict used
— of which from the "refused" MNQ archive	entries on 07-07/09/10/13/14 (pre-capture tape)	11	07-05→07-15	already in the number
— of which LONG (the only live side)	grind_long is armed; grind_short is not	7	4 sessions	the binding constraint

Fact table 2 — the RIDE rung, scored properly (paired, same entries, n=24)

Every row is the same 24 entries exited two ways, so the comparison is paired and entry variance cancels. Baseline = the tight scalp A0.5/B1.0. "strip-3" removes the three best differences; "LOO-worst" is the worst leave-one-session-out mean; ARCH/CAP is the archive half (07-05→07-15, n=11) vs the capture half (07-16→07-31, n=13) — a clean time-split.

RIDE CANDIDATE	\$/TR	PAIRED DIFF	STRIP-3 DIFF	LOO-WORST DIFF	BOOTST RAP P(>0)	ARCH DIFF	CAP DIFF	N @ 2Σ
A2.5+WIDE (operator guess = today's live default)	+71.7	+53.4	+4.3	+34.5	0.937	+49.5	+56.7	46
A3.0+WIDE (plateau centre — the cell to freeze)	+80.4	+62.1	+10.1	+42.0	0.953	+56.9	+66.5	38
A3.5+WIDE	+83.6	+65.3	+9.7	+43.5	0.953	+52.4	+76.3	40
WIDEx2 (sweep winner)	+97.8	+79.4	+6.2	+51.0	0.955	+62.3	+93.9	39
A3.0/B4.5	+77.9	+59.6	+8.2	+43.2	0.948	+49.2	+68.4	40
A0.5/B1.0 (tight scalp = the null)	+18.3	—	—	—	—	+1.9	+32.2	—

READ THIS HONESTLY

The rung is **directionally supported and magnitudinally unproven**, and those are different verdicts. Supported: wide beats tight on 95% of bootstrap resamples, there is a real **parameter plateau** (2.5R / 3.0R / 3.5R + wide all land in +\$72-84, so the answer is not a knife-edge), both time-halves agree in sign and roughly in size, and no single session carries it (LOO-worst is still +\$34 to +\$51). Unproven: **strip the three best trades and +\$62 becomes +\$10**. Eighty-four percent of the advantage lives in 3 of 24 entries. That is not a reason to disbelieve it — a ride exit is supposed to make its money in the top decile, and the median R-MFE of 4.26 with a p90 of 11.8 says the tail is genuinely there. It is a reason to refuse to call it proven at n=24, which is exactly what the sweep's own robustness filter did when it returned best = None for this cell.

Fact table 3 — the acquisition sources, measured not assumed

SOURCE	WHAT IT IS, VERIFIED	NEW SESSIONS	EXPECTED NEW BIG ENTRIES	TIER	VERDICT
(c) V5 bar_history back-extension	alphabot2/data/alphabot.db → bar_history, symbol MNQ, timeframe 5s: 334,818 bars, 2026-06-19→07-17, full 16.5k-bar sessions. 06-22→07-03 is not in any tape the lab has used.	10	+10 (8-13)	5s (bar)	THE ANSWER — run it this weekend
(b1) MNQ before 07-15 in capture.db	does not exist. capture.db 5s bars start 07-15 15:35 UTC; its 250 ms ticks start 07-23 22:00. Verified by direct min(bar_ts)/min(ts_ms).	0	0	—	dead end, confirmed

SOURCE	WHAT IT IS, VERIFIED	NEW SESSIONS	EXPECTED NEW BIG ENTRIES	TIER	VERDICT
(b2) MNQ archive ticks.db 07-05→07-17	15,296,698 MNQ ticks, 1.2-1.9M/session. Already harvested — it is 11 of the current 24.	0 (spent)		0 250 ms	already in the number
(b3) MES 07-12→07-17 (M3's window)	MES ticks span 07-05 22:00 → 07-17 20:59 — the identical calendar days as the MNQ archive. It is not new tape; it is the same nine sessions viewed through a correlated instrument.	0		0 250 ms	worthless for n — keep as falsification only
(a) live-forward	measured, not guessed: shadow grind_fast put 5 entries in $ER \geq 0.50$ minutes over 12 sessions = 0.42/session (3 SHORT, 2 LONG). Replay density is 1.2/session; the shadow board runs at ~35% of replay density.	rolling	+0.42 / session	250 ms	the top-up, not the plan

Fact table 4 — why the back-extension is the right tape (and not just more tape)

The V7 tape is chop-dominated because of when it starts. The ten sessions immediately before it are the trend the desk has never been measured on: a 1,807-point MNQ decline 06-22→06-26 (high 30,967.75 → low 29,160.00), followed by two clean directional legs of +662 and −460 points on 06-29→07-01. $|net|/range$ is the day-level trend purity (1.00 = a straight line).

SESSION	RANGE (PT)	NET (PT)	$ NET /RANGE$	$ER \geq 0.50$ MIN	LONGEST $ER \geq 0.50$ RUN
2026-06-22	631	+88	0.14	36	10
2026-06-23	1122	−913	0.81	40	5
2026-06-24	896	+331	0.37	85	31

SESSION	RANGE (PT)	NET (PT)	NET /RANGE	ER≥0.50 MIN	LONGEST ER≥0.50 RUN
2026-06-25	968	−393	0.41	45	22
2026-06-26	734	−432	0.59	42	13
2026-06-29	795	+662	0.83	22	5
2026-06-30	665	+510	0.77	36	14
2026-07-01	546	−460	0.84	20	6
2026-07-02	991	−530	0.53	38	9
2026-07-03	434	+329	0.76	17	4
10 new sessions	—	—	≥0.59 on 7 of 11	381	—
for contrast: current tape 07-05→07-31 (20 sessions)	—	—	5 sessions score STAY-OUT	923	≥15-min US leg on 3 of 20

Two caveats, stated up front so nobody has to find them later. (1) 06-19 is dropped — it is expiry day for the Jun'26 contract and sits on the roll boundary; the tape starts 06-22. Continuity across the seam is clean (06-19 high 30,705 → 06-22 range 30,348–30,968), and continuity at the other seam is exact: the last archive MNQ tick is **29,623.00 at 15:35:29.973** and the first capture 5s bar opens **29,622.75 at 15:35:30** — one tick apart, correlation 0.9998 across 2,925 overlapping minutes. Same instrument, same front month. That is the evidence that kills Part 2 · LAB's "different contract era" objection. (2) **The pre-07-05 stretch is 5-second bars only** — there are no ticks before 07-05 anywhere on this machine. Per `[[backtest-on-5s-250ms-persist-all]]`, 5s repricing showed an 87% flush-loss rate versus 250 ms, and it biases in favour of wide stops — precisely the thing under test. So the back-extension is **tier-2 evidence: it may confirm or refute the DIRECTION, it may not set the size**. Which is exactly why Track B exists.

The plan — three tracks, dated, owned

TRACK	ACTION	OWNER	DATE	DELIVERABLE / GATE
A — tape	Extend exit_ladder_tape.py to a three-source stitch: alphanot.db bar_history MNQ 5s (06-22→07-05) + ticks.db (07-05→07-15)	Claude (Friday max- depth lab)	Sat-Sun 2026-08-01/02	scratchpad/ exit_ladder_v2.json. Expected grind BIG- TREND n = 34 (32-37). Reported in the 2026-08-07 Friday report.

TRACK	ACTION	OWNER	DATE	DELIVERABLE / GATE
	+ capture.db (07-15→onward). Re-run exit_ladder_lab.py. Tape goes from 20 to 30 sessions (+50%).			
B — honesty	Re-price the existing 24 BIG-TREND entries on 250 ms ticks instead of 5s bars. Verified feasible: all 12 entry days are tick-covered (07-07/09/10/13/14/16 by ticks.db, 07-23/24/27/28/29/30 by capture.db). Zero waiting.	Claude	Sat 2026-08-01	The tick-honest paired diff. Prior: the wide edge shrinks. If it goes negative at 250 ms, the rung is dead on the spot and Tracks A/C stop.
C — freeze	Write the cell explicitly so the live test is clean: add "grind_long": {"a_r": 3.0, "b": "wide"} to data/exit_overrides.json. Note this is a near-no-op — grind_long is in _BIG_RUN, so its built-in default is already A@2.5R + B wide. The freeze moves Lot A 2.5→3.0R (the plateau centre, best strip-3 and lowest $n@2\sigma$) and, more importantly, pins it so no later edit silently changes the thing being measured.	Claude writes; Garrath signs off	Mon 2026-08-03 (exit axis, reversible, one file + restart)	Pre-registered cell. Any change to this line before the grade date voids the test and restarts the count.
— live-forward top-up	Runs itself. Trigger for a clean test day (pre-registered): a ≥ 15 -minute contiguous 30-min-ER ≥ 0.50 leg inside the US session (13:00-21:00 UTC) with grind_long=on. Base rate measured: 3 of 20 sessions (15%), so ~one per 6-7 sessions. The router must not bench grind on those days — that is <code>[[US-open-dont-bench-trend-rider]]</code> restated as a test-integrity rule.	durable router tick (gazbot7-router-tick.timer, verified active/enabled)	continuous from 2026-08-03	+0.42 BIG entries/session measured. 8-12 sessions to close a 3-5 entry gap.
GRADE	Score the rung against the pre-registered test below and publish PASS / FAIL / still-short.	Friday report	Fri 2026-08-21 hard backstop Fri 2026-09-11	On the backstop date it is graded on whatever n exists. No third extension.

THE PRE-REGISTERED PASS TEST — WRITTEN NOW, SO IT CANNOT BE MOVED LATER

The RIDE rung is declared **PROVEN** for grind_long | BIG-TREND only if **all six** hold on the graded sample:

1. $n \geq 39$ paired BIG-TREND entries (derived: paired sd \$247 $\div$ diff \$79.4, 2σ). **2.** $\text{mean}(\text{RIDE} - \text{tight scalp}) > 0$ with bootstrap $P > 0.95$. **3. strip-best-3 paired diff > 0** — currently +\$10.1, the weakest link and the one most likely to fail. **4.** leave-one-session-out worst paired diff ≥ 0 . **5.** sign agreement across both time-halves (pre-07-16 / post-07-16) and across the 5s tier and the 250 ms tier. **6.** the headline number is measured at **250 ms**, with the 5s number reported alongside as the optimism gap.

On FAIL: grind_long drops to the proven MED-TREND exit and the RIDE rung is retired as a myth — note this is a real change, not a no-op, because the wide default is live today. **On still-short at the backstop:** same action. The asymmetry from `[[router-benching-lessons]]` applies to exits too — a wrongly-tight exit costs missed tail; a wrongly-wide exit costs give-back on every trade that never runs, and **62% of these entries lose money even with the wide exit** (WIDEx2 win rate on the cell: 38%).

The thing this exercise found that nobody asked for

Freezing the BIG-TREND cell exposes what the rest of the ladder is doing, because grind_long's A@2.5R+wide default applies in **every** regime, not just BIG-TREND. On the current tape that default is leaving money in three of the four rungs:

GRIND RUNG	N	LIVE DEFAULT A2.5+WIDE	SWEEP BEST	BEST \$/TR	GAP
BIG-TREND	24	+\$71.72	WIDEx2	+\$97.77	+\$26.05 unproven, under test
MED-TREND	326	+\$8.49	TIGHTx2	+\$17.60	+\$9.11 $\times$ 326 real n, real gap
SCALP-CHOP	1539	+\$1.44	WIDEx2	+\$4.64	+\$3.20
STAY-OUT	604	+\$2.02	A3.0/B6.0	+\$6.29	+\$4.27 (the real answer here is FLAT)

The MED-TREND row is the one with n on its side: **326 entries saying the tight chandelier beats the live wide default by \$9.11 a trade.** That is a bigger, better-evidenced prize than the BIG-TREND question you asked about — and unlike BIG-TREND it does not need a single new day of tape. It is not in this section's scope to action, but it should not wait for 08-21 either; it belongs in the **Saturday 2026-08-01** column of REV2 · Q0, and it is the reason the freeze in Track C should be regime-conditional the moment the router can express that.

STRAIGHT ANSWER TO THE QUESTION AS ASKED

(a) The live-forward trigger: first Paris session with a ≥ 15 -minute contiguous 30-min-ER ≥ 0.50 leg inside 13:00–21:00 UTC and grind_long=on; base rate 3 of 20 sessions; the BIG-TREND cell is pre-frozen at {"grind_long":{"a_r":3.0,"b":"wide"}} from Monday 2026-08-03 and any edit to it voids the test. **(b) The matching-instrument option:** the MNQ-before-07-15 idea is already spent — capture.db does not reach back, and the ticks.db archive that does is already 11 of the 24; the MES 07-12→07-17 window adds **zero** new sessions because it covers the same calendar days, so it is demoted to a falsification cross-check. The real matching-instrument answer is one nobody had found: **10 untouched MNQ sessions in alphabot.db bar_history, 06-22→07-03, containing the 1,807-point trend the V7 tape is missing.** **(c) The grading n: 39,** derived from the observed paired sd, not chosen. **Date:** Friday 2026-08-21, hard backstop Friday 2026-09-11. **Owner:** Claude, on the durable router tick and the Friday max-depth lab, with your sign-off required only on the Monday freeze.

REV2 · Q4 The router's own estimator says it cost **−\$1,216 this week. Does it earn its keep or not?**

Garrath — you caught a real one. Part 2.5 and Part 2.6 both print **net router value −\$1,216** and both conclude "**earning its keep**", and the defence in between ("those winners never ran; shadow overstates winners") is **asserted and never sized**. That is a verdict that does not follow its own numbers, and you were right to poke it. So I sized it. Two things are wrong with the −\$1,216: **(1) 60% of it is the dual-slot A/B double-count** the report admits to but never removes — stripping it takes the week to **−\$440**; **(2) the estimator prices zero slippage on a book of 353 modelled round trips**, and because the router blocks winners and losers, any execution drag at all moves the number up — **breakeven is \$1.25 per round trip, or 2.5 MNQ ticks**. The standing memory's own floor (live losses $\geq 1.1 \times$ modelled) implies ~ 8 ticks, which puts the week at **+\$940**.

But the honest answer is not "so it is earning its keep." It is that **the blended number is the wrong statistic**, because its two halves have opposite signs and each half is individually robust. **Benching in CHOP earns +\$1,088. Benching in TREND_UP costs −\$1,380**. The first is positive before any haircut and survives every leave-one-day-out fold; the second needs a **52% winner haircut** to break even, which is outside any defensible band. Net-net the machine is a coin flip; the policy inside it is half-proven and half-broken. One number, stated band, below.

HONESTY BANNER. EVERY FIGURE BELOW IS RE-DERIVED FROM SOURCE, NOT LIFTED FROM THE NIGHTLY JSONS: 37,160 SUPPRESSED-OPEN LINES FROM THE GAZBOT7-TOURNAMENT JOURNAL FOR PARIS 07-27...07-31, REPRICED TICK-HONEST AGAINST CAPTURE.DB WITH THE SAME ESTIMATOR THE NIGHTLY USES (EXIT_SCALP 2R TARGET / 1-ATR STOP, VPP \$2.00, FEE \$1.50/RT), AND REGIME-TAGGED BY REPLAYING DIRECTION_ROUTER.REPLAY_MARKS AT EACH SIGNAL. MY RE-DERIVATION OF THE PUBLISHED METHOD REPRODUCES IT TO \$115 ON \$16.8K OF GROSS (0.7%) — SEE THE RECONCILIATION ROW — SO THE CORRECTIONS BELOW ARE CORRECTIONS, NOT A DIFFERENT MODEL. LIMITS STATED PLAINLY: N=353 BLOCKED SIGNALS, 5 DAYS, ONE SUMMER CHOP-DOMINATED WEEK, MNQ-ONLY, PAPER. EVERY BLOCKED TRADE NEVER RAN —

"MISSED" IS A MODELLED CEILING THROUGHOUT. THE CLEAN-TREND (N=38), DEAD-CHOP (N=3) AND US_OPEN30 (N=6) BUCKETS ARE REPORTED BUT NOT SCORED. AND THE REGIME TAGS COME FROM THE ROUTER'S OWN REPLAY, SO THIS IS A SELF-CONSISTENT AUDIT ("GIVEN WHAT THE ROUTER BELIEVED, DID THE BENCH PAY?"), NOT AN INDEPENDENT REGIME ORACLE.

1 · Reconciliation — am I even looking at the same number?

Before correcting anything, reproduce it. Same journal, same repricer, same 120-second signal dedup, A/B slots counted separately exactly as router_nightly.py counts them:

SOURCE	BLOCKED SIGNALS	SAVED	MISSED	NET ROUTER VALUE
Published — Part 2.5 §3 & Part 2.6 §3	—	\$15,788	\$17,003	-\$1,216
My re-derivation, identical method	652	\$15,730	\$16,832	-\$1,101
delta	—	-\$58	-\$171	\$115

Read: \$115 on \$16.8k of gross = **0.7%**, and it is drift in the right direction — capture.db has backfilled ticks since the nightlies ran, so a handful of signals now reprice with more tape. **The published number is real and I can reproduce it.** Everything below is a correction to it, not a competing estimate.

2 · Error 1 — the dual-slot A/B double-count (worth \$661, i.e. 60% of the deficit)

Part 2.6 says the A/B slots "double-count one signal" and then leaves them in the total. They do double-count, and the proof is that the two legs are **byte-identical** in the nightly's own output. From data/router_nightly/2026-07-31.json:

GATE	N	WON	WON_PNL	LOST	LOST_PNL
grind_long_A	116	46	\$2,633.5	70	-\$2,228.5
	116	46	\$2,633.5	70	-\$2,228.5

GATE	N	WON	WON_PNL	LOST	LOST_PNL
grind_long_B					

Why they are identical, mechanically:

router_nightly.reprice_scalp() applies exit_scalp(target_r=2.0, stop_atr_mult=1.0) to every blocked signal **regardless of slot suffix**. It does not model Lot B's chandelier at all. So the estimator takes one suppressed signal, reprices it twice under the same exit, and books both copies into saved and missed. The fix is to merge the raw journal events per **base** gate before the 120-second dedup, so one suppressed signal counts once:

WEEK 07-27... 07-31	SIGNALS	SAVED	MISSED	NET
A/B counted separately (as published)		652	\$15,730	\$16,832 -\$1,101
A/B merged — one signal, counted once		353	\$8,644	\$9,084 -\$440
duplicate share stripped		-46%	-45%	-46% \$661 recovered

Read: ~46% of both columns was a duplicate, and because the columns are near-equal the duplication does not cancel — it scales the gap too. Stripping it accounts for **\$661 of the \$1,216**, i.e. **60% of the deficit was a bookkeeping artifact**. The A/B pairs appear from 07-29, the day the scale-out slate went live — the same slate REV2·Q1/Q2 found silently voided the ER/ATR floors. Same root cause, second symptom. **Residual known bias I am not fixing:** even the merged number still prices Lot B under Lot A's 2R scalp exit. Direction of that error is unknown and it is not claimed as support for the router.

3 · Error 2 — the estimator prices zero slippage, and the router blocks both sides

The repricer enters at the signal's **touch** price and exits at the first tick that crosses 2R or the 1-ATR stop. Fees are modelled (\$1.50/RT); **slippage is modelled at exactly zero on both legs**. Real fills are worse than the touch on entry and exit — this is the standing memory [[shadow-sim-understates-

losses]] (live losses 1.1–3.1× modelled) and [[execution-cost-autopsy-stage1]] (measured stop-slippage ≈\$1,200 leak).

The key mechanical point the report never makes: the router blocked **221 would-be losers and 132 would-be winners** (mean modelled loss \$39.12, mean modelled win \$68.82). A per-round-trip execution drag δ therefore deepens the losses it saved you from and shrinks the winners it cost you. Both move the same way. So:

$$\text{net}(\delta) = -\$440 + \delta \times 353 \rightarrow \text{breakeven at } \delta = \$1.25 / \text{round trip} = 2.5 \text{ MNQ ticks}$$

One MNQ tick = 0.25 index points = \$0.50 at the desk's \$2.00/pt. So the entire −\$1,216 headline rests on the claim that a blocked round trip suffers **less than three ticks** of total slippage across entry, target and stop. Here is the ladder, in physical units, with the loss-multiplier each drag implies against the modelled mean loss of \$39.12:

MNQ TICKS / RT	Δ (\$/RT)	IMPLIED LOSS MULT.	NET ROUTER VALUE	NOTE
0	\$0.00	1.00×	−\$440	the estimator's own assumption
2	\$1.00	1.03×	−\$86	
2.5	\$1.25	1.03×	\$0	breakeven
4	\$2.00	1.05×	+\$266	
7.8	\$3.91	1.10×	+\$940	the memory band's FLOOR
16	\$8.00	1.20×	+\$2,384	
164	\$82.15	3.10×	+\$28,600	REJECTED — see below

I AM EXPLICITLY REJECTING THE TOP OF MY OWN MEMORY BAND. [[SHADOW-SIM-UNDERSTATES-LOSSES]] SAYS 1.1-3.1×, BUT THAT BAND WAS MEASURED ON REALISED TRADES WITH REAL HOLDING PROFILES. APPLIED AS A LINEAR PER-RT DRAG TO 353 MODELLED SIGNALS IT PRODUCES +\$28.6K, WHICH IS OBVIOUS NONSENSE AND WOULD BE EXACTLY THE KIND OF UNSIZED HAND-WAVE THIS SECTION EXISTS TO KILL. I USE THE FLOOR OF THE BAND AND NOTHING ABOVE IT. IF THE DEFENCE ONLY WORKS AT THE TOP OF A BAND, IT IS NOT A DEFENCE.

The haircut form of the same question — winners discounted by h, saved column held fixed (the most conservative framing, it takes from one side only):

WINNER HAIRCUT H	0%	4.8%	10%	15%	20%	30%
net (A/B stripped)	−\$440	\$0	+\$469	+\$923	+\$1,377	+\$2,286
net (as published, A/B in)	−\$1,216	−\$401	+\$484	+\$1,334	+\$2,185	+\$3,885

→ **Direct answer to the question you asked:** with the double-count stripped, **−\$1,216 flips positive at a 4.8% winner haircut**. On the published numbers as printed, it flips at **7.1%**. Both are low bars — but "low bar" is a directional argument again, which is why the number alone is not the verdict. Section 4 is.

4 · The verdict the blend was hiding — segment it and the two halves disagree

Per the backtest discipline: never score a blended cross-tape number. Splitting the 353 de-duplicated blocked signals by **regime at signal** (the router's own replay) is where the whole question resolves:

REGIME @ SIGNAL	N	SAVED	MISSED	NET	HAIRCUT TO FLIP	VERDICT
CHOP	245	\$6,174	\$5,087	+\$1,088	none needed	☐ EARNs ITS KEEP positive at h=0
TREND_DOWN	44	\$1,216	\$1,364	−\$148	10.8%	marginal, low n
TREND_UP	64	\$1,253	\$2,633	−\$1,380	52.4%	☐ DOES NOT outside any band
blend	353	\$8,644	\$9,084	−\$440	4.8%	the number that hid both

The whole deficit is one segment. Benching in CHOP — the router's core defensive job — is worth **+\$1,088 before any haircut at all**, no execution-drag assumption required. Benching during TREND_UP costs **−\$1,380** and would need the shadow to be overstating winners by **52%** to break even, which is far outside anything the memory band supports. That is [[us-open-dont-bench-trend-rider]] and the two ☐ BAD ledger calls, finally with a price tag on them.

Same story by mechanism and by gate:

CUT	N	SAVED	MISSED	NET	FLIP HAIRCUT	VERDICT
FADER benches (rgv / exhaustion / capitulation)	29	\$1,214	\$601	+\$613	none needed	PROVEN
MOMENTUM benches (grind / abs_veto)	324	\$7,430	\$8,482	−\$1,053	12.4%	MARGINAL needs 1.08× drag
per gate						
abs_veto_s hort	50	\$1,681	\$592	+\$1,089	none	41 of 50 blocked would have lost
exhaustion _short	12	\$632	\$0	+\$632	none	12 of 12 blocked would have lost
capitulati on_long	5	\$113	\$81	+\$33	none	n too small to score
rgv_short	12	\$468	\$520	−\$52	10.0%	noise
abs_veto_l ong	64	\$1,318	\$2,013	−\$695	34.5%	over-benched
grind_long	210	\$4,430	\$5,877	−\$1,447	24.6%	the single biggest leak

And by ATR/ER regime bucket × mechanism — this is where the fader defence earns its reputation and the momentum benching loses it:

BUCKET (ATR + ER KEYED)	N	FADER NET	MOMENTUM NET	TOTAL
building (ER 0.20-0.35)	117	+\$166	+\$338	+\$504
normal-chop (ATR 8-18, ER<0.20)	130	+\$63	−\$488	−\$425
violent-whipsaw (ATR≥18, ER<0.20)	65	+\$408	−\$869	−\$461
clean-trend (ER≥0.35) — not scored, n low	38	−\$24	+\$6	−\$18
dead-chop (ATR<8) — not scored, n=3	3	—	−\$40	−\$40

Read the violent-whipsaw row. That is Friday. Benching the **faders** into ATR-52/62 violence was worth **+\$408**; benching the **momentum** gates in the same tape cost **−\$869**. Part 2.6 §6 called the Friday defensive trim "shadow-vindicated" — half right. The fader half was vindicated; the momentum half was the most expensive segment of the week.

By session (US open = 13:30 UTC): OVERNIGHT_EU n=216, **−\$93**. US_SESSION n=131, **−\$97**. **US_OPEN30** (13:00-13:30 UTC) n=6, **−\$249** — and 5 of the 6 blocked signals would have won. n=6 so it is **reported, not scored**, but it is the sharpest single illustration of the

banked rule: benching across the open is where the router pays the most per decision.

5 · Robustness — which of these two conclusions actually survives?

The discipline demands leave-one-day-out and strip-the-best-trades on every claim. Here both, run on the blend and on the same book with the TREND_UP benches removed:

TEST	BLEND	NET, TREND_UP BENCHES REMOVED
all 5 days	-\$440	+\$940
LOO — drop 07-27	-\$358	+\$1,021
LOO — drop 07-28	-\$342	+\$958
LOO — drop 07-29	+\$230	+\$640
LOO — drop 07-30	-\$804	+\$608
LOO — drop 07-31	-\$484	+\$533
strip top & bottom 3 trades	-\$226	+\$753
strip top & bottom 5 trades	-\$72	+\$591
strip top & bottom 10 trades	+\$268	+\$1,008
verdict	SIGN FLIPS (-\$804...+\$268) not robust — do not quote it	POSITIVE IN ALL 9 FOLDS (+\$533...+\$1,021) robust

This is the finding. The blended router-value number **changes sign depending on which single day you drop** — it has no statistical standing in either direction, and neither Part 2.5's "\$-1,216, over-benching" nor Part 2.6's "yes, earning its keep" is supportable from it. What is robust is the split: remove the TREND_UP benches and the week is positive in every leave-one-day-out fold and every trade-stripping fold, **at zero haircut, with no execution-drag assumption whatsoever.**

Sized policy fix (all at h=0, no drag assumed):

POLICY	SIGNALS	WEEK NET	Δ VS AS-RUN
as run	353	-\$440	—
stop benching momentum in TREND_UP (56 signals)	297	+\$1,175	+\$1,615
	289	+\$940	+\$1,380

6 · Error 3 — the leakage column, which neither section reconciled

Part 2.5 prints week leakage **−\$1,083**. Part 2.6 prints **"\$0 (durable days)"** on the same row of the same table. Both are true in their own frame (−\$456/−\$398/−\$229 on 27/28/29, \$0 on 30/31), but the report never says so, and a reader comparing the two sections sees a contradiction. Worse, sitting in the same table as saved/missed it invites being netted in. **It must not be.**

- **Leakage is the only realised-money number in that whole table.** saved and missed are modelled counterfactuals on trades that never ran; leakage is real losers that really fired. Do not present them side by side without saying which is which.
- **Leakage is not a router cost line.** It counts managed-gate losers that fired while the gate was ON. The no-router baseline has those same losers. Subtracting it from router value would triple-count the router's failures — a third bookkeeping error on top of the two above.
- What it is: an under-benching diagnostic. And it points the **opposite** way to the over-benching story both sections lead with — the week's realised evidence says the router was too loose early (−\$1,083 through Tuesday) and got tighter, not that it was uniformly trigger-happy.

7 · One number, with a band

You asked for a number instead of a directional argument. Here it is.

point estimate, memory-floor drag (1.1× losses ⇒ \$3.91/RT ≈ 7.8 ticks)	+\$940
★ STATED VALUE — midpoint with band	+\$250 [−\$440 , +\$940]
for comparison — the figure Parts 2.5 & 2.6 published	−\$1,216
for comparison — the desk's actual realised P&L, same week	+\$625

THE VERDICT, STATED PROPERLY:

1. −\$1,216 IS WRONG. \$661 OF IT IS THE DUAL-SLOT A/B DOUBLE-COUNT; THE REST SITS ON A ZERO-SLIPPAGE ASSUMPTION THAT BREAKS AT 2.5 MNQ TICKS PER ROUND TRIP. CORRECTED, THE WEEK IS +\$250, BAND −\$440 TO +\$940.

2. "EARNING ITS KEEP" WAS ALSO WRONG — AS STATED. NOT BECAUSE THE SIGN IS NEGATIVE, BUT BECAUSE THE BLENDED NUMBER FLIPS SIGN ON LEAVE-ONE-DAY-OUT (−\$804 TO +\$230). IT CANNOT SUPPORT A VERDICT IN EITHER DIRECTION. PART 2.6 §8'S "YES" AND PART 2.5'S "TWO THINGS PULL THE OTHER WAY, HARD" WERE BOTH REASONING FROM A STATISTIC WITH NO STANDING. ON AGGREGATE P&L THIS WEEK THE ROUTER IS INDISTINGUISHABLE FROM ZERO.

3. THE REAL ANSWER IS A SPLIT VERDICT, AND BOTH HALVES ARE ROBUST:

- DEFENCE — BENCHING FADERS AND BENCHING INTO CHOP: EARNS ITS KEEP. +\$1,088 (CHOP) / +\$613 (FADERS), POSITIVE BEFORE ANY HAIRCUT, ROBUST TO EVERY FOLD.

- OFFENCE — BENCHING MOMENTUM IN TREND_UP: DOES NOT. −\$1,380, NEEDING A 52% WINNER HAIRCUT TO BREAK EVEN.

4. SO THE HONEST ANSWER TO "DOES IT EARN ITS KEEP" IS: HALF OF IT DOES, DECISIVELY, AND THE OTHER HALF IS EATING THE PROCEEDS. FIXING THE ONE SEGMENT IS WORTH +\$1,615/WEEK AT ZERO HAIRCUT AND IS POSITIVE IN ALL NINE ROBUSTNESS FOLDS — WHICH MAKES "STOP BENCHING ALIGNED MOMENTUM IN A CONFIRMED TREND_UP" THE HIGHEST-VALUE, BEST-EVIDENCED CHANGE IN THIS ENTIRE REPORT, AND IT NEEDS NO NEW BUILD: IT IS THE RULE ALREADY BANKED AS [[US-OPEN-DONT-BENCH-TREND-RIDER]], SIMPLY NOT BEING ENFORCED.

5. CAVEAT THAT MUST TRAVEL WITH THE NUMBER: ONE WEEK, N=353 MODELLED SIGNALS, CHOP-DOMINATED SUMMER TAPE, MNQ-ONLY, PAPER, AND THE REGIME TAGS ARE THE ROUTER'S OWN. THE SEGMENT SPLIT IS THE ROBUST FINDING; THE ±\$690 BAND ON THE HEADLINE IS

HONEST WIDTH, NOT FALSE PRECISION. RE-RUN AFTER A GENUINE TREND WEEK BEFORE SIZING ANYTHING OFF THE TREND_UP FIGURE.

REV3 · TREND_UP Q4 says arming momentum in TREND_UP is worth +\$1,615/week. I tried to kill it. It mostly died.

Garrath — Q4 called this "the highest-value, best-evidenced change in this entire report." I re-derived it from source, independently, and tried to break it. **The arithmetic is right and I reproduce it: +\$1,644 on my numbers vs Q4's +\$1,615.** But three things it did not test take it apart. **(1)** The estimator books every deduped signal as a full trade even when the same gate is already in one — the temporal twin of the A/B double-count Q4 itself caught. Fixing it: **+\$1,799 → +\$767.** **(2)** The TREND_UP label is the carrier, and it does not survive independence: under a textbook ADX(14) tag the segment is worth **+\$574 and is statistically indistinguishable from a random subset of the same book (p=0.56);** under a 60-min regression tag it is **−\$276 and significantly WORSE than random (z=−2.56).** **(3)** The money is not trend money. **52% of it comes from two grind_long episodes on 07-29 — a day MNQ fell 702 points at day-ER 0.72,** i.e. momentum longs buying bounces in the week's cleanest DOWN trend, which the desk's own day-bias rule says to bench. The one genuinely clean UP day of the week (07-30, +1,126pt, day-ER 0.93) contributes **\$349 of \$1,799.**

The number: REDUCED from +\$1,615 to a best-case +\$767/week, and to +\$158 on n=4 signals once the rule is written so it cannot buy dips in a downtrend. REFUTED as a "trend" finding.

Recommendation: **do not ship it as an arming rule.** It points the expensive direction on a tape that contains no confirmed trend to test it on

— by my own operational definition, a confirmed trend was live for **3.3% of the week, 22 episodes, median 6 minutes each**. There is nothing here to ride. What I do recommend, and what is worth more than the whole debate, is in §7: **all six gates are currently PINNED ON, so the durable router cannot bench anything at all** — the proven half of the machine is switched off right now.

HONESTY BANNER. INDEPENDENT RE-DERIVATION, NOT A RE-READ OF Q4. SOURCE = THE SAME 37,160 SUPPRESSED-OPEN LINES FROM THE GAZBOT7-TOURNAMENT JOURNAL, PARIS 07-27...07-31, PARSED FRESH; EVERY REPRICE ON CAPTURE.DB 250MS TICKS (8,448,165 TICKS LOADED), NEVER 1-MIN BARS; DUCKDB + PANDAS THROUGHOUT. ATR AND THE 2R/1-ATR EXIT MATCH ROUTER_NIGHTLY.REPRICE_SCALP() EXACTLY SO THE NUMBERS ARE COMPARABLE. EVERY BLOCKED TRADE NEVER RAN — EVERY DOLLAR BELOW IS A MODELLED CEILING, AND LIVE LOSSES RUN 1.1-3.1× MODELLED [[SHADOW-SIM-UNDERSTATES-LOSSES]]. N=355 MERGED SIGNALS, 5 DAYS, ONE CHOP-DOMINATED SUMMER WEEK, MNQ-ONLY, PAPER. WHERE I DISAGREE WITH Q4 I SHOW BOTH NUMBERS. WORKING FILES: /TMP/.../REV3/DERIVE.PY, ANALYSE{, 2,...,7}.PY. READ-ONLY: NOTHING WAS APPLIED.

1 · Reconciliation — I get the same book Q4 got

Same journal, same 120s dedup, same repricer. Two counting conventions: A/B slots separate (as router_nightly.py publishes) and A/B merged before dedup (Q4's fix, which is correct).

SOURCE	SIGNALS	SAVED	MISSED	NET ROUTER VALUE
Published — Part 2.5 §3 & Part 2.6 §3	—	\$15,788	\$17,003	−\$1,216
Q4 re-derivation, A/B separate	652	\$15,730	\$16,832	−\$1,101
REV3 (mine), A/B separate	656	\$16,066	\$17,428	−\$1,362
Q4, A/B merged	353	\$8,644	\$9,084	−\$440
REV3 (mine), A/B merged	355	\$8,804	\$9,450	−\$646

Read: three signals apart on 355 and \$206 apart on \$9.4k of gross (2.2%). **Q4's A/B double-count finding is CONFIRMED** — merging the slots

removes 46% of both columns. Q4's zero-slippage point (\$3) is CONFIRMED and I do not re-litigate it. The segment split reproduces too:

REGIME @ SIGNAL (ROUTER REPLAY)	N	ROUTER VALUE — Q4	ROUTER VALUE — REV3	WIN%	\$/SIGNAL
CHOP	239	+\$1,088	+\$812	33.1%	-\$3.4
TREND_DOWN	44	-\$148	-\$154	40.9%	+\$3.5
TREND_UP	72	-\$1,380	-\$1,304	50.0%	+\$18.1
TREND_UP n MOMENTUM — the claim	62	+\$1,615	+\$1,644	56.5%	+\$26.5

So the claim is arithmetically real. Everything from here is about whether it means what Q4 says it means.

2 · Kill shot #1 — the estimator books trades a real gate could never have taken

Q4 stripped the A/B duplicate. It left the **temporal** duplicate.
router_nightly.py dedups suppressions at 120 seconds and then prices every survivor as a complete round trip — even when the previous signal on that same gate is still notionally open. Median modelled hold is minutes to tens of minutes; the dedup window is two. A real gate holds one position at a time.

Correction: walk each gate's signals in time order, and skip any that arrives before the previous one has exited (using that trade's own tick-derived exit time).

BOOK (ALIGNED MOMENTUM LONG, N ROUTER TREND_UP) as Q4 counts it — every deduped	NET	WIN%	\$/TRADE	LOO WORST FOLD	STRIP BEST 3	STRIP BEST 5
56	+\$1,799	61%	+\$32.1	—	+\$1,304	+\$978

BOOK (ALIGNED MOMENTUM LONG, N ROUTER TREND_UP)		NET	WIN%	\$/TRADE	LOO WORST FOLD	STRIP BEST 3	STRIP BEST 5
signal is a trade							
one open position per gate (non-overlap)	32	+\$767	56%	+\$24.0	+\$276	+\$276	+\$85
and including the counter-side (all momentum, as Q4 words it)	38	+\$612	50%	+\$16.1	+\$121	+\$122	-\$70

Read: –57% from a bookkeeping fix of exactly the species Q4 itself found. And Q4's own wording — "stop benching momentum in TREND_UP" — includes abs_veto_short, i.e. arming a SHORT gate inside an up-trend. That sub-slice is **n=6, –\$154, 17% win**. Counter-side momentum in a trend loses money and should never have been in the number. Aligned-only is the only defensible version.

3 · Kill shot #2 — the finding lives or dies on whose definition of TREND_UP

Q4 flags the circularity in its own honesty banner and then scores against the router's tag anyway. So I tagged the same 355 signals four more ways, none of which touches direction_router, and ran a placebo: **if the TREND_UP label carries no information, a random same-size subset of the aligned momentum book should do as well.** 20,000 draws.

REGIME ORACLE FOR TREND_UP	% OF WEEK	N	NET (RAW)	WIN%	\$/SIG	VS RANDO M	P	VERDICT
router replay — ER30≥0.20,	18%	56	+\$1,799	61%	+\$32.1	z=+3.60	0.000	the claim

REGIME ORACLE FOR TREND_ UP	% OF WEEK	N	NET (RAW)	WIN%	\$/SIG	VS RANDO M	P	VERDICT
net≥30, HOLD 2 marks								
ADX(14) on 5-min bars ≥20, +DI>-DI (textbook)	30%	112	+\$780	38%	+\$7.0	z=-0.16	0.561	☐ no information
OLS 60-min R²≥0.35 & net≥+30pt	29%	106	-\$365	32%	-\$3.4	z=-2.56	0.996	☐ WORSE than random
OLS 30-min — same horizon as the router, structure not ER	—	126	+\$70	36%	+\$0.6	—	—	☐ no edge
OLS 30-min + 30-min sustain	~0%	3	-\$151	0%	-\$50.2	—	—	☐ fires 3× all week
strict — ER60≥0.35 & net≥60pt, 30-min sustain	0%	0	—	—	—	—	—	☐ never occurred

Read this row by row. The router's own ER-based tag picks a subset worth \$32/signal against a random-subset mean of \$7.6/signal — a real, significant selection (z=+3.60). **No other formulation reproduces it.** The obvious defence is "the others use a longer horizon" — so I built one at the same 30-minute horizon with the same ±30pt threshold, swapping Kaufman ER for regression R². It finds **\$0.6/signal**. The discriminator is not the horizon. It is the specific ER_TREND=0.20 / NET_MIN=30 / HOLD=2 triple — and those three knobs were tuned on 07-20...07-26, the week immediately before this one.

The comparison that should end the argument: under the two independent tags, **the TREND_UP bucket is not even the best bucket.**

ALIGNED MOMENTUM-LONG BLOCKED BOOK, \$/ SIGNAL BY REGIME	CHOP	TREND_DOWN	TREND_UP	
router tag		+\$1.5	+\$0.1	+\$32.1
ADX(14)		+\$7.7	+\$10.1	+\$7.0
OLS 60-min		+\$14.6	+\$13.7	-\$3.4

Under both independent oracles the **CHOP** bucket is where the blocked momentum longs would have made money — which is not a case for arming momentum in trends, it is a case that this estimator likes momentum longs everywhere this week: the whole aligned momentum-LONG blocked book is **+\$2,078 across all 274 signals in all regimes**. A finding that says "you should never have benched a momentum long anywhere, all week" is not a TREND_UP finding. It is a statement about the estimator's optimism and about a week that closed 217 points lower but bounced hard.

4 · Kill shot #3 — it is not trend money. It is bounce money on the biggest DOWN day.

PARIS DAY	N	NET	WIN%	\$/SIG	DAY MOVE	DAY ER	WHAT THE DAY ACTUALLY WAS
07-29	17	+\$990	65%	+\$58.2	−702pt	0.72	the week's cleanest DOWN trend
07-30	14	+\$349	64%	+\$24.9	+1,126pt	0.93	the only clean UP day
07-31	25	+\$460	56%	+\$18.4	−33pt	0.05	violent round-trip chop (ATR 52–62)
07-27, 07-28	0	\$0	—	—	−291 / −251pt	0.35 / 0.40	no TREND_UP signals at all

The one clean up-trend day of the week — the exact tape `[[us-open-dont-bench-trend-rider]]` was written for — contributes \$349 of \$1,799 (19%). The two biggest episodes are both grind_long on 07-29:

EPISODE	SIGNALS	NET	WINDOW (UTC)	PRICE	DAY CONTEXT
grind_long#2 @07-29	5	+\$570	18:40 → 18:49	27868 → 28025	day 27962 → 27260 (−702pt)
grind_long#1 @07-29	4	+\$371	16:52 → 16:58	27552 → 27579	same day, price 400pt below the open
grind_long#1 @07-31	4	+\$353	16:26 → 16:56	28312 → 28376	day-ER 0.05, ATR-52 whipsaw
	5	+\$253	09:04 → 09:12	27484 → 27514	

EPISODE	SIGNALS	NET	WINDOW (UTC)	PRICE	DAY CONTEXT
grind_long#3 @07-30					the clean up day — the only "correct" one

Those top two are **+\$941 = 52% of the entire claim, and both are momentum longs bought into bounces while MNQ was 400-700 points below its open on the week's cleanest down day.** The framework's own rule is day-bias $\text{DOWN} \leq -40 \rightarrow$ bench longs. 34% of the claimed money (\$614 over 10 signals) was earned with the day-bias at or below -40 . To bank Q4's **+\$1,615** the router would have to override the day-bias rule — which Q4 never says and nobody has asked for.

Two more concentration cuts, same direction:

- **Episodes, not signals.** 56 signals collapse to **21 episodes** (same gate, gaps < 30 min). Top-1 episode = **35%** of the total; top-3 = **79%**. Strip the best three episodes and **\$1,799** becomes **+\$350**. Effective n is ~ 20 , not 56.
- **ER at the moment of firing.** The money sits in the lowest ER bucket: **$\text{ER}_{30} \leq 0.20 \rightarrow n=32, +\$1,351$ (75% of the claim); $\text{ER}_{30} \geq 0.35 \rightarrow n=9, +\6 .** Because `replay_marks` carries `HOLD=2` hysteresis on a 15-min grid, a "TREND_UP" label at signal time can be up to half an hour stale. **57% of the claimed signals fired while the tape was, on its own instantaneous read, chop.** That is
`[[rearm-momentum-needs-er-climb-not-delta-blip]]` with a price tag on the wrong side of the ledger.

5 · What I could NOT kill — the estimator itself, and the fader half

Two adversarial tests came back in Q4's favour and I report them straight.

(a) The estimator is not systematically optimistic. Direct calibration against the same week's realised trades:

MOMENTUM					
GATES, 07-27... 07-31	N	NET	WIN%	\$/TRADE	
REALISED (actually ran, real exits, real fills)	141		+\$602	35%	+\$4.3
MODELLED (blocked, 2R/1-ATR on 250ms ticks)	326		+\$1,284	39%	+\$3.9

Modelled is **\$0.3/trade WORSE** than realised. So the reduction in \$2-\$4 is not a slippage argument and I am not making one. Consistent with that, the drag maths runs Q4's way: arming means taking these trades, so $\text{net} = \Sigma \text{pnl} - \delta \cdot n$, and on the router tag the aligned book breaks even only at **$\delta = \$24/\text{RT} \approx 48 \text{ MNQ ticks}$** . Slippage does not kill this. The tag and the overlap do.

(b) Exit choice is not the artefact. Repriced on the same ticks under six exits — the finding is stable in shape, which means the 2R/1-ATR estimator is not flattering it:

EXIT VARIANT (ROUTER TAG, ALIGNED LONG, N=56)	NET	WIN%	\$/TRADE
2R / 1-ATR — the estimator	+\$1,799	61%	+\$32.1
1R / 1-ATR scalp	+\$1,064	73%	+\$19.0
3R / 1-ATR	+\$2,238	52%	+\$40.0
2R / 1.5-ATR wider stop	+\$2,468	75%	+\$44.1
chandelier 3.5→0.5k / 1-ATR stop	+\$2,496	61%	+\$44.6

(c) And a correction that runs AGAINST Q4's other half. Q4 states the defence as "+\$1,088 in CHOP, robust to every fold." The CHOP claim is tag-dependent in exactly the same way the TREND_UP claim is:

DEFENSIVE CUT	ROUTER TAG	ADX TAG	OLS-60 TAG	VERDICT
bench into CHOP — router value, non- overlap	+\$612	+\$187	-\$382	tag-FRAGILE, same flaw
bench the FADERS — router value, non- overlap	+\$810 (n=25) — a mechanism cut, identical under every tag □ the one robust claim			

So the surviving half of Q4 is narrower than Q4 says. "Benching faders earns its keep" holds under any definition of anything — exhaustion_short alone: 12 of 12 blocked signals would have lost, +\$648. "Benching into

CHOP earns its keep" only holds if you accept the router's own definition of chop.

6 · Sized ladder — what any version of this rule is actually worth

Every row is "arm ALIGNED long momentum when <condition>", non-overlap, at h=0 with no drag assumed, on this one week. LOO = worst leave-one-Paris-day-out fold. δ_{be} = per-round-trip slippage that zeroes it.

CANDIDATE RULE	N	WEEK NET	WIN%	\$/TRADE	LOO WORST	STRIP BEST 3	STRIP BEST 5	Δ_{BE}
Q4 as written — router TREND_UP, ALL momentum	38	+\$612	50%	+\$16.1	+\$121	+\$122	−\$70	\$16.1
aligned-only — router TREND_UP + LONG	32	+\$767	56%	+\$24.0	+\$276	+\$276	+\$85	\$24.0
+ day-bias UP (price ≥ day open +40) — framework-consistent	23	+\$422	57%	+\$18.4	+\$230	+\$91	−\$34	\$18.4
+ day-bias UP + ADX(14) confirm	17	+\$446	59%	+\$26.2	+\$253	+\$88	−\$36	\$26.2
+ day-bias UP + ER30≥0.35 (the grind floor the desk already uses)	4	+\$158	50%	+\$39.5	−\$34	−\$58	\$0	\$39.5
+ day-bias UP + ADX + ER30≥0.30 + ATR30≥14	4	+\$158	50%	+\$39.5	−\$34	−\$58	\$0	\$39.5
US session only (13:30–20:00z) + day-bias UP	4	+\$162	50%	+\$40.6	−\$2	−\$50	\$0	\$40.6

Read the shape, not any single row. Every step that makes the rule more consistent with the desk's existing framework makes it smaller, and every version fails at least one robustness cut. The loosest version (Q4's own wording) is the one that dies fastest under strip-best-5. The tightest versions are n=4 for the entire week.

Q4 says the fix is "already banked as `[[us-open-dont-bench-trend-rider]]`, simply not being enforced." That is not what the banked rule says. The banked rule is narrow: on a cleanly-trending day, hold the aligned trend-rider through open stop-outs; bench only if trend STRUCTURE breaks. It is a rule about not panicking on flushes. Q4's version is a rule about arming momentum for 18% of the clock on a 30-min $ER \geq 0.20$ read. Those are different rules, and the ledger already grades the boundary between them: 07-30 13:36 (☐ BAD, benched the rider on open noise in a clean trend) versus 07-31 13:40 (☐ GOOD, benched grind after the thrust fully reversed). **The banked rule is being enforced correctly. Q4 is proposing a new and much wider one.**

7 • The asymmetry, honestly — and the thing that actually matters this week

The standing rule is that wrongly-benched is cheap (missed trades) and wrongly-armed is expensive (real churn losses), so when unsure, sit benched one more tick. **This change points the expensive way.** Everything about the evidence says respect the asymmetry:

WHAT THE ASYMMETRY TEST ASKS	ANSWER FROM THIS WEEK'S TAPE
Is the effect robust to how you define the condition?	☐ No — significant under one tag, null under a second, negative under a third.
Is it robust to leave-one-day-out?	☐ Not on the framework-consistent versions (LOO −\$34); marginally yes only on the loosest.
Is it robust to stripping the best few?	☐ Strip 3 episodes: \$1,799 → \$350. Strip best 5 signals: Q4's own wording → −\$70.
Is the tape representative of the condition being armed for?	☐ No. A confirmed trend was live 3.3% of the week — 22 episodes, median 6 minutes . There was nothing to ride.
Does the wrongly-armed failure mode already have a body count?	☐ Yes. Ledger 07-31: enable-all at the open → grind bought the roll-over −\$370 ; two afternoon re-arms into a "calm" ER-0.24-0.30 recovery churned −\$245 . Both would be more likely under Q4's rule, because 57% of its signals fire at $ER_{30} \leq 0.20$.
What is the upside if it is right?	

WHAT THE ASYMMETRY TEST ASKS	ANSWER FROM THIS WEEK'S TAPE
	+\$400-\$770/week modelled ceiling, on a book whose realised comparator did \$4.3/trade.

Verdict: STAGE IT, do not ship it. A change that arms a churner in more conditions, justified by a segment that inverts when you change the definition of the segment, on a week with no confirmed trend in it, is precisely the trade the asymmetry rule exists to refuse. The cost of waiting is a modelled few hundred dollars a week that may not exist; the cost of being wrong is measured, repeatedly, in the ledger.

★ **And the operational finding that is worth more than this entire debate.** As of right now:

LIVE STATE, 2026-08-01	VALUE	CONSEQUENCE
PINNED in scripts/ router_tick_durable.py	frozenset({all 6 gates})	the durable tick cannot bench ANY gate
data/gate_switches.env	all 6 = on	desk fully armed in every regime
comment in the file	"★US-OPEN OVERRIDE 2026-07-31 ... Revert to frozenset({"exhaustion_short"}) once the open settles."	

The open settled two days ago. **The proven half of the router — the fader bench, +\$810/week, the only claim in Q4 that survives every test I could throw at it — is currently switched off**, and so is the CHOP bench. Q4's recommendation is to arm momentum in more conditions; the desk is already in the maximally-armed state, by accident. Before any argument about arming more, the pin should come off so the machine can do the part it demonstrably earns. That is a one-line revert, it points the CHEAP direction, and it needs no new evidence.

8 · The staged path — how to earn this finding instead of assuming it

- 1. Un-pin first (today).** PINNED = frozenset({"exhaustion_short"}) — restore the fade-scalp trial's protection only, and let the durable tick resume benching. Reinstates the +\$810 fader defence. No new build, no new risk, cheap direction.

2. **Instrument, do not arm.** Add nothing to the prompt. Instead extend `router_nightly.py` so it (a) merges A/B before dedup, (b) enforces one-open-position-per-gate, and (c) tags each blocked signal with ADX and the router replay. Every night thereafter produces the honest version of this table for free. Until that lands, every "router value" number in the Friday report carries a $\sim 2\times$ counting inflation.
3. **Wait for the tape.** The whole question is unanswerable on a week containing 224 minutes of confirmed trend. Re-run this exact ladder after a week with a real trend day — the deciding test is whether the ADX-confirmed and router-only versions converge. If they do, the effect is real and worth arming. If they keep disagreeing, the router's ER triple is overfit and should be retuned, not obeyed.
4. **If the operator wants something in force now,** the only asymmetry-safe change is a TIGHTENING, and it is written out below. It does not arm anything new. It (i) defines "confirmed trend" operationally so no future tick can re-arm on a 30-minute-stale label, (ii) forbids arming the counter-side momentum gate inside a confirmed trend (measured: – \$154, 17% win), and (iii) restates the banked open rule in its narrow, correct form. It is PREPARED, NOT APPLIED.

PREPARED, NOT APPLIED — CONSERVATIVE PATCH TO SCRIPTS/ROUTER_TICK_DURABLE.PY. THIS IS NOT Q4'S CHANGE. Q4'S CHANGE IS DECLINED. THIS ONE IS A TIGHTENING: IT CAN ONLY CAUSE A GATE TO BE BENCHED IN SITUATIONS WHERE IT MIGHT PREVIOUSLY HAVE BEEN ARMED, PLUS ONE NARROW HOLD-THROUGH-FLUSHES CLAUSE THAT ALREADY EXISTS AS A BANKED RULE. THE CHOP/FADER BENCHING WORDING IS UNTOUCHED BY CONSTRUCTION — THE INSERT SITS AFTER IT AND SAYS SO EXPLICITLY.

Insert point — immediately after the existing sentence ending "... (not a delta-blip). " and before "abs_veto_long is veto-protected...", inside the prompt = (literal:

"★ CONFIRMED-TREND TEST – apply this whenever you would hold or re-arm an ALIGNED momentum gate. "

"It does NOT loosen the chop/fader benching above (that is the proven half; leave it exactly as is). "

"A trend counts as CONFIRMED only if ALL FOUR are true on THIS tick's desk-view read, never on a "

"reading from an earlier tick: (1) 30-min ER ≥ 0.35 ; (2) 30-min net

```
move >= 60 MNQ points in the "  
    "trend direction; (3) ATR >= 14; (4) the same read held on the previous  
tick too (~10 min). "  
    "If any one of the four fails it is NOT a confirmed trend – it is chop,  
and momentum stays benched, "  
    "however directional the last 30 minutes looked. In a CONFIRMED trend:  
hold the ALIGNED momentum "  
    "gate through stop-outs (an open flush is noise – bench it only if the  
trend STRUCTURE breaks: the "  
    "leg fully round-trips, or ER falls back under 0.20), and BENCH the  
counter-side momentum gate. "  
    "Never arm a momentum gate against the day-bias. "
```

Exact revert: delete the seven inserted string-literal lines. Nothing else in the file changes — no constant, no default, no PINNED edit, no behaviour outside the prompt text. The wrapper is fail-safe either way (any parse failure $\Rightarrow$ no switch change).

What it would have done this week, measured, not asserted — and this cuts against my own patch, so read it. The four-part test was true for **3.3% of the 1-min clock** (1.8% UP, 1.4% DOWN), across 22 episodes of median 6 minutes. As a permission filter it clears **26 of the 274 aligned-long signal-instants (9%)**, worth **−\$28 raw / +\$86 non-overlap**, win 42%, **and even that flips negative on two of four leave-one-day-out folds (−\$25, −\$60) and on strip-best-3 (−\$260)**. The counter-side clause codifies what already happened rather than changing it: those 6 `abs_veto_short` signals were benched, worth **+\$154 saved** — the clause's job is to stop a future tick arming them, which is exactly what Q4's wording would have done.

So: the tightened rule earns approximately nothing this week either. That is the point, and I would rather say it than sell it. Its value is not money, it is determinism — right now "confirmed TREND_UP" is whatever a headless Claude decides at 03:00 UTC with no operator watching, and 57% of the signals Q4 wants armed fire at $ER_{30} \leq 0.20$. The patch makes the phrase mean one thing. Sell it as a guardrail, size it at \$0, and expect the argument about arming to be settled by a future trend week, not by this one.

9 · Bottom line

CLAIM	Q4	REV3	STATUS
Published −\$1,216 is wrong; A/B double-count is \$661 of it		✓	✓ ☐ CONFIRMED
Zero-slippage assumption inflates the deficit		✓	☐ CONFIRMED (and calibration says the estimator is fair, so it cuts both ways)
Blended router value has no statistical standing		✓	✓ ☐ CONFIRMED — do not quote it
Benching the FADERS earns its keep		+\$613	+\$810 ☐ CONFIRMED and strengthened — the one tag-independent claim
Benching into CHOP earns its keep, robust to every fold		+\$1,088	+\$612 / +\$187 / −\$382 ☐ DOWNGRADED — tag-fragile, same flaw as the TREND_UP half
Stop benching momentum in TREND_UP = +\$1,615/week, robust in all 9 folds		+\$1,615	+\$767 → +\$158 ☐ REDUCED, and REFUTED as a trend finding
"Highest-value, best-evidenced change in the report; needs no new build"		✓	✗ ☐ DECLINED — stage it; un-pin the router instead

One paragraph for the plane. Q4 found two real bookkeeping errors and one real segment split, and it deserves credit for all three. But it stopped one correction short: the same estimator that double-counts A and B also double-counts a gate against itself in time, and fixing that halves the headline finding. What is left is significant only under the router's own thirty-minute ER definition of a trend, is null under a textbook ADX and negative under a regression fit at the same horizon, and is half-funded by two clusters of momentum longs bought into bounces on the week's biggest down day. The rule it asks for is not the rule the desk banked. **Keep the defence, keep the ledger's asymmetry, take the pin off so the router can defend again, and re-run this the first week the tape gives us a real trend.**

REV3 The dead floors — confirmed, re-derived, and the order to ship them in tonight

Garrath — Q2 is **right**. Every ER and ATR floor on the desk has been a no-op since **2026-07-29 16:20:49 UTC**, and I confirmed it independently from source, from the journal, and from the commit dates. But Q2 overstated the size by a factor of 44, and it left the important half undone. Here is the important half. **Four of the five floors should never be switched back on.** The two ER floors the operator wrote on 07-30 are the two most expensive lines of config on the desk: over eight days of tick-honest replay, `grind_long`'s $ER \geq 0.35$ would have thrown away **+\$2,923 of profit to keep +\$83**, and `abs_veto_long`'s $ER \geq 0.25$ would have thrown away **+\$640 to keep +\$148**. Neither has ever run, which is the only reason the desk is up. **If you ship the wiring fix tonight and nothing else, you take the four floored gates from 382 signals / +\$3,630 down to 103 signals / +\$668 — a −\$2,962 self-inflicted wound on eight days of tape, about −\$370 a day.** The ship order below exists to make that impossible.

HONESTY BANNER. EVERY NUMBER IS MEASURED, NONE RECALLED. SOURCES: THE GAZBOT7-TOURNAMENT JOURNAL (274,106 LINES, 07-20→08-01), DATA/GAZBOT7.DB, DATA/SHADOW.DB, AND DATA/CAPTURE.DB (10,080,959 MNQ TICKS + 9,546,332 CLEAN QUOTES). EVERY EXIT IS REPLAYED ON THE 250MS STREAM, NEVER ON A BAR. THE REPLAY IS A CLONE OF THE DESK'S OWN SRC/GAZBOT7/REPRICER.PY AND REPRODUCES SHADOW_REAL.REAL_PNL TO \$0.00 MAE ON FOUR OF SIX GATES (DETAILS §2). EVERY DOLLAR BELOW IS A CEILING — MEMORY [[SHADOW-SIM-UNDERSTATES-LOSSES]] PUTS LIVE LOSSES AT 1.1-3.1× MODELLED, SO I CARRY A BREAK-EVEN LOSS MULTIPLIER ON EVERY RECOMMENDATION AND TREAT ANYTHING UNDER ~1.5× AS NOT REALLY GREEN. THE TAPE IS 8 DAYS, ONE CHOP-DOMINATED SUMMER REGIME, MNQ-ONLY, PAPER. MY REPLAY IS ALWAYS-ARMED — IT DOES NOT MODEL ROUTER BENCHING, WHICH IN THE SAME WINDOW BENCHED 465 SIGNAL EPISODES.

1 · The bug — CONFIRMED, with one correction to Q2

Mechanism, one sentence: `tournament.step()` passes the raw slot tag to `er_blocks()` and `atr_blocks()`, but under the dual-slot scale-out slate every tag carries an `_A/_B` suffix while `ER_FLOOR/ATR_FLOOR` are keyed on base names — so the dictionary lookup never matches and every floor silently returns "don't block".

The `_base()` helper that strips the suffix already exists at `tournament.py:86` and is correctly applied three times — to the bench switch (`:250`), the 55s absorption veto (`:303`) and the exhaustion confirm (`:345`). It was simply not applied to the two floors at `:139` and `:142`.

EVIDENCE	WHAT IT ESTABLISHES
source	<code>tournament.py:139</code> if <code>er_blocks(slot, er):</code> <code>:142</code> if <code>atr_blocks(slot, f.atr):</code> — raw slot. <code>deciders.py:107/173</code> key on base names. Working tree clean for all three decision-path files; editable install; the running process loaded exactly this code.
cutover	Commit 5b0913b "dual-slot scale-out slate" 07-29 16:18:50 → drop-in <code>/etc/systemd/system/gazbot7-tournament.service.d/scaleout.conf</code> written 16:20:00 → restart 16:20:49 → journal <code>tournament LIVE — slots=['grind_long_A','grind_long_B',...]</code> at 16:20:50 .
journal	Last suppression line of any kind: 2026-07-29T13:29:04Z . Since the cutover: exactly zero ER or ATR suppression lines in 2 days 18 hours, against a prior run-rate of 1,500–2,700 lines/day. Verified by exhaustive scan, not sampling.
commit dates	Both ER floors were written AFTER the bug. <code>grind_long 0.35 = aa59357, 07-30 10:23:53. abs_veto_long 0.25 = 9767c67, 07-30 10:57:10</code> . Both ~18h after the cutover. The journal contains zero ER gate: <code>grind_long</code> lines in its entire 12-day history. Neither floor has ever blocked a single trade in production.

The one correction. Q2's headline "7,734 suppressions" is a raw log-line count, and the check re-fires every ~1s decision tick for as long as a signal stands. Collapsing consecutive same-gate lines ≤ 5 s apart into episodes:

MEASURE	ER GATE	ATR GATE	COMBINED
raw log lines (Q2's basis)	7,744	9,080	16,824
distinct blocked signal episodes	140	243	383
overstatement factor	55×	37×	44×

So the honest sentence is: **the floors were blocking about 48 signal episodes a day across all gates, not thousands** — concentrated in `grind_long` ATR (191 episodes) and the two ER gates (40 each). That is

material, but an order of magnitude smaller than the number that has been circulating. The bug is real; its scale was not.

What else the same bug killed, and what it did not. I audited every tag-keyed lookup in the decision path. The bug kills exactly three things: the ER floors, the ATR floors, and — collaterally — the ER_HOLD abs_veto_long spike-filter **shadow** study (tournament.py:148), which has been recording no data for three days. That last one is observe-only, so no P&L impact, but it means that study is three days poorer than it looks. **Unaffected and working correctly:** router benching, the 55s absorption veto, the 5s exhaustion confirm, veto_counter_regime, _BIG_RUN, and data/exit_overrides.json — all key on _base() or on spec.tag consistently. adaptive_exit is off under scale-out **by design**, not by bug.

2 · Method — and why you can trust the replay

1-minute bars rebuilt from the 5s capture (16,725 bars) using the desk's own agg.MinuteBars folding rule; features computed on the same trailing 60-bar window the live desk holds. Two independent checks that this reproduces the desk rather than approximating it:

CHECK	RESULT
ATR at entry vs the desk's own recorded entry_atr	correlation 1.0000 , median absolute difference 0.000 pt across 1,632 shadow trades — an exact reproduction, not a fit.
ATR quartiles of the tape	Q25 11.04 / Q50 14.43 / Q75 19.27 pt — independently reproduces Q2's 11.0 / 14.4 / 19.2 from a separate code path.
exit replay vs shadow_real.real_pnl (quote basis)	n=558, sign agreement 99.82% , MAE \$0.80 ; \$0.00 MAE on 4 of 6 gates . This is a clone of repricer.py, including its two non-obvious rules: entry is the first quote after the signal bar closes (filling at the minute start is lookahead), and fills are at the far touch.
tick basis vs quote basis	Quotes only cover 5 days; trade ticks cover 8. On the overlap: sign agreement 99.28% , correlation 0.981 , and the tick basis is more pessimistic in every single gate . All headline numbers below are tick-basis (8 days) and cross-checked on the quote basis (5 days).

Populations are the **faithful shadow variants whose entry config is byte-identical to the live gate** — abs_veto_55s for both abs_veto sides, grind_fast for grind_long — replayed through the **live scale-out dual-slot exit** (Lot A fixed-R + Lot B per exit_overrides.json, or the built-in

default). 644 signals, 8 days, 2026-07-23→07-31. Regime buckets keyed on ATR quartile + 30-min ER + 60-bar range-break, session split at 13:30 UTC.

One population caveat, stated up front: capitulation_long's only faithful proxy is capit_loose, which runs require_flip=False where the live gate runs True. It is a looser superset, so read that gate's rows as directional, not exact.

3 · ★ The test, and the verdict on all five floors

The test is simple and it is the right one. Not trading costs nothing, so **a floor only earns its keep if the population it BLOCKS reliably loses money**. "Reliably" means three things at once: the blocked book is net negative, it is still negative after you strip its three biggest winners, and it is negative on the worst leave-one-day-out fold. A floor whose blocked book is positive is not a filter — it is a leak.

GATE	FLOOR	BASIS	BLOCKED N	BLOCKED NET	STRIP BEST-3	LODO WORST	VERDICT
grind_long	ER ≥ 0.35	tick	173	+\$2,923	+\$1,132	+\$10.03	DELETE
grind_long	ER ≥ 0.35	quote	150	+\$2,890	+\$1,152	+\$12.06	DELETE
grind_long	ATR ≥ 24	tick	151	+\$982	−\$82	+\$2.92	REPLACE → 10
grind_long	ATR ≥ 24	quote	130	+\$1,364	+\$209	+\$5.51	REPLACE → 10
abs_veto_long	ER ≥ 0.25	tick	52	+\$640	+\$111	+\$4.88	DELETE
abs_veto_long	ER ≥ 0.25	quote	44	+\$792	+\$260	+\$9.99	DELETE
abs_veto_short	ATR ≥ 16	tick	28	+\$137	−\$197	−\$11.02	NOT PROVEN
abs_veto_short	ATR ≥ 16	quote	23	+\$256	−\$101	−\$7.63	NOT PROVEN
capitulation_long	ATR ≥ 10	tick	22	−\$459	−\$601	−\$26.62	KEEP
capitulation_long	ATR ≥ 10	quote	18	−\$542	−\$674	−\$53.25	KEEP

One floor out of five passes. Both bases agree on every verdict — no call in this table depends on which price stream you score it with.

The standing priors, tested rather than inherited. Memory says "ER filters are FAKE, ATR floors are REAL", and separately that abs_veto wants no ER floor while grind wanted 0.35.

- **"ER filters are FAKE" — CONFIRMED, emphatically.** Both live ER floors block robustly profitable books on both price bases and all eight LODO folds. The abs_veto_long blocked book is positive on every single one of the 8 folds (+\$4.88 to +\$24.27 per signal).
- **"abs_veto wants NO ER floor" — CONFIRMED, and now we know why.** Segmented, the entire blocked book is **NORMAL-CHOP: 40 signals, +\$680**. The ER floor is precisely an anti-home filter — it fires exactly where Q2 showed abs_veto lives. It also makes the gate less robust: break-even loss multiplier falls from 1.27 to 1.13 with the floor on.
- **"grind wanted 0.35" — REFUTED.** This is the most expensive single line of config on the desk. It keeps 8 signals worth +\$83 and discards 173 worth +\$2,923.
- **"ATR floors are REAL" — HALF TRUE, and the half matters.** ATR floors do sharpen per-trade quality monotonically (grind's kept book goes \$16.60 → \$67 → \$111 per signal as the floor climbs from 0 to 30). But **total dollars peak far lower than the live settings**. The floors are real; the values are roughly double where the money is.

4 · Per-gate detail — the number behind each verdict

grind_long ATR: 24 → 10 (a genuine plateau at 9-11)

Swept fine, under the ext_hi 2.0 ceiling. The blocked book is the test:

ATR FLOOR	BLOCKED N	BLOCKED NET	BLK STRIP-3	BLK LODO	KEPT N	KEPT NET	KEPT \$/ SIG
8	3	-\$13	—	—	108	+\$2,944	\$27.26
9	9	-\$144	-\$222	-\$25.44	102	+\$3,074	\$30.14
10 ← re-derived	11	-\$165	-\$293	-\$22.42	100	+\$3,095	\$30.95
11	16	-\$3	-\$290	-\$16.95	95	+\$2,934	\$30.88
12 ← the report's rec	27	+\$776	+\$157	+\$14.80	84	+\$2,154	\$25.64

ATR FLOOR	BLOCKED N	BLOCKED NET	BLK STRIP-3	BLK LODO	KEPT N	KEPT NET	KEPT \$/ SIG
24 ← LIVE as written	91	+\$1,712	+\$742	+\$15.43	20	+\$1,218	\$60.90

9, 10 and 11 all block a robustly negative book (negative net, negative strip-3, negative LODO) and sit within \$161 of each other on kept net — a real plateau, not a spike. At 12 the blocked book **flips sharply positive** and stays positive all the way up. **The report's recommended 12 is one notch the wrong side of the cliff** and costs \$941 versus 10. The plateau holds without the ceiling too (at $ATR \geq 10$, no ceiling: blocked 19 signals $-\$134$, strip-3 $-\$376$, LODO $-\$10.74$), so this value does not depend on shipping the ceiling.

abs_veto_short $ATR \geq 16$ — the one honest coin-flip

Blocked book is net positive ($+\$137$ tick / $+\$256$ quote) so it fails the stated test — but that positive is entirely tail-driven: strip its 3 best trades and it is $-\$197$, and the worst LODO fold is $-\$11.02$. The IS/OOS halves flip sign ($+\$417$ then $-\$280$). Meanwhile the floor does improve resilience: break-even multiplier $1.36 \rightarrow 1.45$, per-signal $\$15.07 \rightarrow \22.02 . **This floor is genuinely unproven in both directions.** Per the standing rule — a floor that cannot be shown to earn its keep is deleted, not tuned — I recommend deleting it, but it is the one line in this report where your own risk appetite legitimately overrides the evidence, and the cost of being wrong either way is about $\pm \$150$ over eight days. Separately, a **lead, not a recommendation**: an $ER \geq 0.10$ floor on this gate blocks 14 signals worth $-\$183$ (strip-3 $-\$605$, LODO $-\$26.00$) and lifts break-even to 1.54. That is a better filter than the ATR one — but it is a new filter and belongs in shadow first, not in tonight's window.

capitulation_long $ATR \geq 10$ — KEEP the floor, but the floor is not the problem

The only floor that passes cleanly, on both bases and every robustness test. Keep it at 10. But be clear about what it is buying: the gate loses money at every floor value — $-\$1,202$ with no floor, $-\$743$ at ATR 10, $-\$540$ at ATR 14, and it only stops losing when the floor is high enough that it stops trading. Break-even loss multiplier **0.41**. The floor is harm-reduction on a gate that needs a relegation conversation, not a tuning one; its live scale-out

record agrees (5 signals, −\$128.50). Lead only, on the loose proxy population: an $ER \geq 0.25$ floor flips it to +\$124 on $n=13$, break-even 1.48. Too thin and on the wrong population to act on tonight.

5 · ★ The interaction that gates tonight — grind_long

Three questions, answered directly.

(a) Was the grind replay applying the floors as config, or did it inherit the dead-floor reality? It applied them as config. scripts/rehab_grind_long.py:124 reads `if er < er_floor or f.atr < atr_floor: continue`, and line 234 hard-codes `[x for x in raw if x["er"] >= 0.35 and x["atr"] >= 24.0]`. The funnel re-runs as: 1,232 raw signals → 93 clear ER 0.35 → 171 clear ATR 24 → **25 clear both** → 11 realised trades. So the "11 fires −\$284" baseline is a **simulation of a configuration that was not running. The honest before/after is +\$1,552 → +\$2,830, not −\$284 → +\$2,830.** Roughly 59% of the advertised delta is the removal of filters that were already de-facto removed. The report is internally aware of this — Q2 §3 even says Part 1.5 needs re-reading — but the numbers in part1_5_rehab.html, rehab_grind_long.md and the money_gbp: 700 in plays.json were never corrected.

(b) If the wiring fix ships, do grind's floors come back to life and silently undo the entry fix? Yes — catastrophically, and it would not be silent for long.

WORLD	N	NET	\$/SIG	WIN%	STRIP-3	LODO MIN	@1.5× LOSSES	BREAK-EVEN ×	
A. dead floors, ext_hi 4.0 — what the desk has actually been running		181	+\$3,005	\$16.60	34.3	+\$1,214	+\$11.11	−\$1,449	1.34
B. floors ALIVE as written — what D1 alone ships		4	+\$360	\$90.00	50.0	−\$145	−\$25.96	+\$218	2.27
		111	+\$2,930	\$26.40	38.7	+\$1,161	+\$19.11	+\$358	1.57

WORLD	N	NET	\$/SIG	WIN%	STRIP-3	LODO MIN	@1.5× LOSSES	BREAK-EVEN ×
C. ext_hi 2.0 ceiling only								
D. the report's rec: no ER floor, ATR≥12, ext_hi 2.0								
	84	+\$2,154	\$25.64	36.9	+\$385	+\$14.36	−\$98	1.48
E. re-derived: no ER floor, ATR≥10, ext_hi 2.0								
	100	+\$3,095	\$30.95	40.0	+\$1,326	+\$23.58	+\$669	1.64

World B is the answer to (b): **181 signals collapse to 4, and +\$3,005 becomes +\$360**. On the quote basis it is 156 → 3. Ninety-eight percent of grind's fires disappear. That is not "the entry fix gets undone" — it is "the gate gets switched off by accident".

The ceiling itself is real and it is the only genuinely new lever here. I reproduced the report's headline independently: filtering to $\text{ext_atr} \leq 2.0$ gives **111 fires / +\$2,930** against their 111 fires / +\$2,830 — a 3.5% agreement from a completely separate code path, which is strong corroboration of both harnesses. The plateau across ext_hi 1.5 / 1.75 / 2.0 is broad (\$2,646 / \$2,931 / \$2,930). And the honest argument for the ceiling is not the raw dollars — those barely move — it is **resilience**: at 1.5× losses the uncapped book is **−\$1,449** while the capped book is **+\$358**. The ceiling does not make more money; it makes the same money out of 39% fewer trades, which is what survives contact with real fills.

Row E beats row D on every measure — net (+\$941), per-signal, win rate, strip-best-3 (+\$1,326 vs +\$385), worst LODO fold (+\$23.58 vs +\$14.36), the 1.5× stress test (+\$669 vs −\$98) and break-even multiplier (1.64 vs 1.48) — and on both IS and OOS halves. All 8 LODO folds of row E are positive, ranging \$23.58 to \$37.81 per signal. **Ship ATR 10, not 12.**

(c) The desk-wide cost of getting the order wrong. Shipping D1 alone, across all four floored gates, on 8 days of tick-honest replay:

GATE	FLOOR AS WRITTEN	TODAY N / NET	AFTER D1 N / NET	FIRES LOST	DOLLARS LOST
grind_long	ER≥0.35 and ATR≥24	181 / +\$3,005	4 / +\$360	98%	−\$2,645
abs_veto_long	ER≥0.25	81 / +\$787	29 / +\$148	64%	−\$640
abs_veto_short	ATR≥16	69 / +\$1,040	41 / +\$903	41%	−\$137
capitulation_long	ATR≥10	51 / −\$1,202	29 / −\$743	43%	+\$459

GATE	FLOOR AS WRITTEN	TODAY N / NET	AFTER D1 N / NET	FIRES LOST	DOLLARS LOST
DESK TOTAL	—	382 / +\$3,630	103 / +\$668	73%	-\$2,962

About **-\$370 a day**, and it would arrive as "the desk mysteriously stopped trading" rather than as a visible loss.

6 · The ship order for tonight — concrete, ordered, revertible

THE ONE RULE THAT MAKES THIS SAFE. BECAUSE THE FLOORS ARE CURRENTLY DEAD, CHANGING THE FLOOR VALUES TODAY IS PROVABLY A NO-OP ON THE LIVE DESK. THAT IS THE WHOLE TRICK: LAND THE CORRECT VALUES FIRST WHILE THEY CANNOT POSSIBLY DO HARM, THEN FLIP THE WIRING AS A SINGLE, ATOMIC, ONE-LINE ARMING MOMENT. NEVER SHIP THE WIRING FIX BEFORE OR WITHOUT THE VALUES. IF YOU ONLY REMEMBER ONE THING, REMEMBER THAT TOURNAMENT.PY:139/142 IS THE MASTER SWITCH AND IT GOES LAST.

STEP 1 — land the re-derived floor values (no-op by construction).

src/gazbot7/deciders.py:

```
-ER_FLOOR = {"abs_veto_long": 0.25, "grind_long": 0.35}
+ER_FLOOR: dict = {}    # ★2026-08-01 REV3: both momentum-long ER
floors DELETED. Re-derived on 8d
+                        # tick-honest replay: grind 0.35 blocked +
$2,923 (strip-3 +$1,132, all 8 LOD0
+                        # folds positive); abs_veto_long 0.25
blocked +$640, all of it NORMAL-CHOP,
+                        # the gate's home cell. Never ran live
(committed after the 07-29 cutover).

-ATR_FLOOR = {"grind_long": 24.0, "capitulation_long": 10.0,
"abs_veto_short": 16.0}
+ATR_FLOOR = {"grind_long": 10.0, "capitulation_long": 10.0}    #
★2026-08-01 REV3: grind 24→10
+                        # (plateau 9-11; at 12+ the blocked book flips
```



```
positive). capitulation 10 KEPT – the only
+          # floor whose blocked book is robustly negative.
abs_veto_short 16 DELETED – unproven.
```

Verify: restart, then confirm the journal still shows zero ER gate: / ATR gate: lines. If any appear, something else changed — stop. **Revert:** restore the two literals; no-op either way while Step 2 is unshipped.

STEP 2 — the wiring fix. This is the arming moment. src/gazbot7/tournament.py:

```
-          if er_blocks(slot, er):
+          if er_blocks(_base(slot), er):
-          if atr_blocks(slot, f.atr):
+          if atr_blocks(_base(slot), f.atr):
```

Optionally also :148/:154 slot → _base(slot) to revive the ER-hold **shadow** study (observe-only, no P&L effect). **Verify within one session:** suppression lines reappear with base names; grind_long ATR blocks should run at roughly **10% of fires** and capitulation_long at **43%**; there should be **zero** ER gate: lines anywhere (both ER floors are now empty). Any ER gate: line means Step 1 did not take. **Revert:** one-line, _base(slot) → slot. **This single revert neutralises Steps 1 and 2 together** and returns the desk to exactly today's behaviour — it is the kill switch.

STEP 3 — the grind ext_hi ceiling. Separate commit, separate restart, separate observation. src/gazbot7/slot_strategy.py:118:

```
-          params={"slope_min": 0.4, "fast_slope": True},
sizing="conviction", base_size=2,
+          params={"slope_min": 0.4, "fast_slope": True,
"ext_hi": 2.0}, sizing="conviction", base_size=2,
```

This is the **only** change here that alters entry behaviour immediately and unconditionally — it is not gated by the wiring bug, so it has real effect the moment the service restarts. Expect roughly **39% fewer grind fires**. The mechanism already exists (gate_grind(..., ext_hi=4.0) at deciders.py: 370, enforced at :380); only the value is being set. **Revert:** delete "ext_hi": 2.0 from the dict → back to the 4.0 default. **Keep this out of**

the Step-2 commit so that if grind misbehaves you can tell which change did it.

STEP 4 — nothing else tonight. The `abs_veto_short` $ER \geq 0.10$ lead and the `capitulation_long` $ER \geq 0.25$ lead are both interesting and both go to **shadow**, not to live. Neither is in the diff above.

#	STEP	LIVE EFFECT AT SHIP TIME	REVERT
1	<code>deciders.py</code> floor values	none — provably a no-op while the wiring is broken	restore two literals
2	<code>tournament.py</code> <code>_base(slot) x2</code>	arms everything: grind blocks ~10% of fires, capitulation ~43%, <code>abs_veto</code> unaffected	one line — the kill switch for 1+2
3	<code>slot_strategy.py</code> <code>ext_hi: 2.0</code>	immediate: ~39% fewer grind fires	delete one key
—	whole slate, if it all goes wrong	—	<code>GAZBOT7_TOURNAMENT_SLATE=tournament</code>

Router benching is untouched by all of this and remains reversible every 5 minutes.

7 · The biggest risk, and what this does NOT prove

The biggest risk if you ship all three tonight is that `grind_long` is a **high-count, low-margin book being asked to survive real fills**. Even the best config — row E — has a break-even loss multiplier of **1.64**, and memory `[[shadow-sim-understates-losses]]` puts live losses at **1.1-3.1× modelled**. That number sits inside the band. If live losses on this gate run at 1.7× rather than 1.5×, the entire +\$3,095 is zero, and everything in \$5 becomes an argument about noise. The ceiling helps (it is why row E survives 1.5× where row A does not) but it does not lift grind out of the band. Grind is being made less bad and more robust, not made good.

Behind that, three more:

- **My replay is always-armed; the live desk is not.** In the same window the router benched 465 `grind/abs_veto` signal episodes. The interaction between the router's benching and the new floors is **unmeasured** — the two filters could stack and starve the gate, or

overlap and do nothing. Watch the combined block rate on the first session, not just the floor's.

- **Eight days, one regime.** Chop-dominated, MNQ-only, PAPER. CLEAN-TREND is $n=1$ for grind at the ceiling — completely unscored. A wide `ext_atr` is exactly what a trend day pays for, and `ext_hi 2.0` caps it. If Monday opens trending, Step 3 is the first thing to question.
- **Two population caveats.** `capitulation_long`'s proxy runs `require_flip=False` against the live True, so its rows are directional only. And `abs_veto_short`'s ATR verdict is a coin-flip I have called by rule rather than by evidence — I would not defend it hard either way.

What would change my mind: for grind, one clean trend week with `ext_atr` above 2.0 paying — that is the single scenario this tape cannot speak to. For the two ER deletions, nothing on this evidence; they are the most robust calls in the report, positive on both price bases and all eight folds.

REV3 grind_long's exit, swept properly — and what the cell should be

PATCH MARKER — this section **closes the hole** REV2 • Q1 **flagged in itself** ("grind_long's 2.5/wide pair was not put through this ladder ... that sweep is the obvious next job and its absence is a hole in this section, not a silence"). It also **patches PART 2 • LAB's** grind rows and **refutes** REV2 • Q0 **item M5** (`grind_long → {a_r 0.5, b 1.0}`), and it **extends** REV2 • Q3 by grading its pre-registered freeze cell on tick data a fortnight early. Where this section and any of those disagree, this section is the correction.

Garrath — the short version is that we were arguing about the wrong knob. I swept **grind_long's exit across 130 configurations on 3,064 tick-repriced entries over 16 sessions**, and **not one exit setting makes the gate green**. The loaded default (Lot A 2.5R + Lot B wide lock-chandelier) loses **−\$2,115** over those 16 sessions; the single best cell on the board still loses **−\$414**. What the sweep does settle is the direction of travel, and it is the opposite of what the Friday report recommended: **tighter is worse**,

not better. The M5 proposal (0.5R / 1.0R) is **−\$1,127 worse** than what is loaded today and fails every robustness test I ran. My recommendation is the small, pre-registered move: **write {"grind_long": {"a_r": 3.0, "b": "wide"}}** — Lot A from 2.5R to 3.0R, Lot B unchanged — worth **+\$617** over the 16 sessions against the loaded baseline. And then stop tuning the exit, because the same data says the arming decision is worth three to five times as much.

THE ONE-PARAGRAPH ANSWER

Recommended cell: {"grind_long": {"a_r": 3.0, "b": "wide"}}.

Baseline: the built-in BIG-RUN default it inherits today, Lot A 2.5R + Lot B wide lock-chandelier, which scores −\$2,115 over 16 tick-covered sessions (−\$1.72 per lot-fill, 26% win). The recommended cell scores −\$1,498 (−\$1.31 per lot-fill) — a delta of +\$617, about +\$39 a session at sim fill-density, nearer +\$10 a session at the live desk's actual fill rate. It passes leave-one-day-out (worst-case diff still +\$222), beats the default on 10 of 16 sessions, and sits on the correct side of the one finding that survives every test: the 3.5R rung beats both of the loaded legs in both halves of the tape (+\$997 archive, +\$705 capture). It **fails** the bootstrap bar ($P(\text{diff} > 0) = 0.82$, not 0.95), **fails** strip-the-best-3-days (−\$130), and **fails** the out-of-sample sign test (archive half −\$76 against capture half +\$693). So it is a sensible small move on a proven direction, not a proven number — which is exactly why I am recommending the cell REV2·Q3 already pre-registered for Monday rather than the cell with the prettiest statistics (3.5R/3.5R, +\$1,702, passes everything except the plateau test, and it is a one-rung spike). **The single biggest caveat: none of this makes grind_long profitable. It is red at every setting, before any slippage.**

Fact table 1 — what is actually loaded today (verified in code, not recalled)

data/exit_overrides.json currently contains three keys — exhaustion_short, abs_veto_long, abs_veto_short. **grind_long is not one of them**, so scaleout_slots() falls through to the built-in BIG-RUN branch. This is the config every number in this section is measured against.

PIECE	WHAT IS LOADED	WHERE IN CODE
Lot A (grind_long_A)	exit="scalp", target_r = 2.5, stop_atr_mult 1.0, adaptive_exit=False, give-back OFF, size 1	slot_strategy.py:213-217
Lot B (grind_long_B)	exit="chandelier_lock", start_k 3.5 → lock_k 0.5 once peak ≥ 6.0R, size 1	slot_strategy.py:189-192 (_lot_b(..., "wide"))
Why it gets that pair	grind_long is in _BIG_RUN → A@2.5R + B wide; a gate absent from the override file keeps its built-in default	slot_strategy.py:149
Loss side	native server-side stop at 1.0 × entry ATR; hard 120-minute MAX_HOLD	config.py:126, multislot_core.py:388
Entry	gate_grind(slope_min=0.4, fast_slope=True), LONG only, 60 one-minute bars of lookback	slot_strategy.py:115-118
Entry floors (ER ≥ 0.35, ATR ≥ 24)	NOT APPLIED — see fact table 9. The filters are keyed on grind_long; the live slots are called grind_long_A / grind_long_B	tournament.py:139,142 vs deciders.py: 107,173

Fact table 2 — how this was measured, and the one methodological change that mattered

Two simulators, same tape, same ticks. The difference between them is the whole reason the Friday lab and this section disagree.

METHOD	WHAT IT DOES	SIGNALS	VERDICT ON IT
Paired (the Friday lab's method)	take every gate signal as an independent entry, reprice all exits on the same forward path, average		good for mechanism questions (entry variance cancels) — but it stacks overlapping entries: on a real run the gate fires every few minutes and the sim buys all of them, so trend tape gets counted 10-20 times over. Its answer for the loaded default

METHOD	WHAT IT DOES	SIGNALS	VERDICT ON IT
			was +\$2,961 — a number the live desk has never seen.
Sequential (headline here)	each lot is a slot holding one position; it takes the next signal only once it is flat again — exactly what the live dual-slot desk does	3,064	the honest policy test. Same default scores −\$2,115, which is the sign and the shape the live ledger shows. Every headline number below is sequential.

TAPE / INPUT	DETAIL	SIZE
Signal tape (5s bars → 1-min)	alphabot.db bar_history MNQ 5s (06-19→07-15) stitched to capture.db 5s (07-15→07-31); seam verified one tick apart	498,406 bars / 31 sessions
grind_long signal-minutes	every minute the live gate config fires LONG (no artificial cooldown — the slot re-arm handles it)	5,874
Exit repricing — 250 ms ticks only	ticks.db 07-05 22:00→07-17 21:00 (15.3M MNQ ticks) + capture.db 07-23 22:00→07-31 21:00 (10.1M ticks). No bar-based exits anywhere in this section.	3,064 entries · 16 sessions
Deliberately excluded	2,803 signal-minutes on 15 sessions (06-19→07-05, 07-18→07-23) that have 5-second bars but no ticks. REV2·Q3's Track A back-extension would add them at tier-2; I refused them for the headline because 5s repricing biases in favour of wide stops, which is the thing under test.	—

Fact table 3 — does the model track reality? (the calibration nobody usually shows)

Before trusting a single sweep number: I took the **60 real live grind_long fills** that fall inside tick coverage, reconstructed each one's entry ATR, and ran the loaded exit against the ticks. If the model cannot reproduce trades the desk actually did, the sweep is fiction.

CHECK	N	REAL	MODEL	READ
Entry ATR reconstruction — my ATR vs the ATR implied by the real stop distance	52	stop distance / my ATR	median ratio 1.016	the ATR I feed the sim is the ATR the desk used, within 1.6%
Real STOP fills, dollars	40	−\$2,001.6	−\$1,871.1	real stops fill 7% worse than the model's exact-level fill
	9	+\$1,137.0	+\$862.9	

CHECK	N	REAL	MODEL	READ
Real TARGET fills, dollars				model 24% conservative on targets
Lot A exit mix	33 live / 737 sim	67% stop, 27% target, median hold 3.5 min	72% stop, 28% target, median hold 4.8 min	matches
Lot B exit mix	19 live / 491 sim	68% stop, 16% chandelier, median hold 7.0 min	78% stop, 22% chandelier, median hold 6.0 min	matches
Fill density	—	~20 lot-fills / armed session	77 lot-fills / session	the sim trades ~3.8× more often (the live gate spends much of the week router-benched). Rank configs on this; scale dollars down before quoting them as forecasts.

Fact table 4 — the exit ladder for ONE lot (16 sessions, sequential, 250 ms ticks)

This is the whole sweep in one table, because a pair cell is just the sum of two independent lots. **Every row loses money.** Read the shape, not the winner.

LOT EXIT	FILLS	NET \$	\$ / FILL	WIN %	LODO WORST	READ
scalp 0.5R (M5's Lot A)	1,770	−1,915	−1.08	67	−2,424	wins two trades in three and still bleeds — the fee and the unchanged 1R stop eat it
scalp 0.75R	1,430	−1,491	−1.04	58	−2,077	
scalp 1.0R (M5's Lot B)	1,203	−1,328	−1.10	51	−2,154	
scalp 1.25R	1,069	−1,397	−1.31	45	−2,107	
scalp 1.5R	974	−1,000	−1.03	40	−1,859	
scalp 2.0R	836	−1,390	−1.66	33	−1,954	
scalp 2.5R — LOADED Lot A	737	−1,178	−1.60	28	−1,676	the baseline lot
scalp 3.0R	655	−560	−0.86	25	−1,188	the recommended Lot A
scalp 3.5R	607	−207	−0.34	23	−851	best on the board — but a spike, not a plateau (3.0R is \$353 worse,

LOT EXIT	FILLS	NET \$	\$ / FILL	WIN %	LODO WORST	READ
						4.0R is \$709 worse)
scalp 4.0R	563	−916	−1.63	20	−1,578	
scalp 5.0R	510	−985	−1.93	17	−1,577	
scalp 6.0R	460	−856	−1.86	15	−1,631	
WIDE lock-chandelier 3.5→0.5@6R — LOADED Lot B	491	−938	−1.91	22	−1,520	mid-pack; no better than a plain 3.0R target, no worse than 6.0R
TIGHT chandelier k1.5	1,169	−2,718	−2.33	51	−3,337	worst exit on the board
chandelier k2.5	889	−2,066	−2.32	40	−2,555	
chandelier k3.5 (plain, tightening)	719	−1,977	−2.75	31	−2,554	the tightening chandelier family is uniformly bad here

THE PLATEAU TEST, STATED PLAINLY

There is **no plateau at the winner**. 3.5R beats its neighbours by \$353 and \$709 on a 16-session sample where session-to-session swings are $\pm \$1,000$ — that is inside the noise, and the desk's own rule is that a spike is not a result. What is a plateau is the **region**: on the full tape every fixed target from 3.0R to 6.0R sits between $-\$207$ and $-\$985$, every target at 2.5R or tighter sits between $-\$1,000$ and $-\$1,915$, and all three tightening chandeliers sit below $-\$1,900$. **The shape is proven; the point is not.**

And the shape is weaker than that sounds, so here is the half-split before anyone over-reads it. The wide-beats-tight ordering does not hold rung-by-rung inside each half: in the 10 archive sessions the least-bad single lot is actually 0.5R ($-\$1,227$, against $-\$1,264$ for 3.5R and $-\$1,901$ for the wide chandelier), and in the 6 capture sessions 1.5R ($+\$1,071$) edges 3.5R ($+\$1,058$). What survives the split is narrower and is the thing I am relying on: **3.5R beats both of the currently-loaded legs — 2.5R and the wide chandelier — in the archive half and in the capture half.** Everything wider than that (4R, 5R, 6R) is worse than 3.5R in both halves too. So "go a bit wider on Lot A, leave Lot B alone" is supported on both halves; "wide is always better than tight" is not, and I am not claiming it.

LOT EXIT	ARCHIVE HALF (10 SESSIONS)	CAPTURE HALF (6 SESSIONS)	FULL TAPE
scalp 0.5R	-1,227	-687	-1,915
scalp 1.5R	-2,071	+1,071	-1,000
scalp 2.5R — loaded Lot A	-1,624	+446	-1,178
scalp 3.0R — recommended Lot A	-1,700	+1,140	-560
scalp 3.5R	-1,264	+1,058	-207
scalp 6.0R	-1,677	+821	-856
WIDE lock-chandelier — loaded Lot B	-1,901	+964	-938
TIGHT chandelier k1.5	-2,217	-501	-2,718

Note the honest wrinkle in the recommended cell: 3.0R is *worse* than the loaded 2.5R in the archive half (-\$1,700 vs -\$1,624) and much better in the capture half (+\$1,140 vs +\$446). That is precisely why its bootstrap comes out at 0.82 and not 0.95, and why I am calling +\$617 a direction rather than a forecast.

Fact table 5 — the candidate cells, with the robustness tests attached

Baseline for every "delta" column is the loaded default A2.5R + WIDE at -\$2,115. "strip top-3 days" removes the three best sessions from the difference; "LODO worst" is the worst leave-one-session-out difference; ARCH = the 10 archive sessions (07-06→07-17), CAP = the 6 capture sessions (07-24→07-31) — a clean time split, and the out-of-sample leg.

CELL	FILLS	NET \$	DELTA VS LOADED	BEATS BASE ON	STRIP TOP-3 DAYS	LODO WORST	P(DIFF >0)	ARCH DIFF	CAP DIFF
A2.5 / B wide — LOADED TODAY	1,228	-2,115	—	—	—	—	—	—	—
A3.0 / B wide — RECOMMENDED	1,146	-1,498	+617	10 / 16	-130 fail	+222 pass	0.822 fail	-76 fail	+693
A3.5 / B 3.0	1,262	-767	+1,348	11 / 16	+9 pass	+619 pass	0.934 fail	+561	+787
	1,214	-414	+1,702	12 / 16	+630 pass		0.988 pass	+997	+705

CELL	FILLS	NET \$	DELTA VS LOADED	BEATS BASE ON	STRIP TOP-3 DAYS	LODO WORST	P(DIFF >0)	ARCH DIFF	CAP DIFF
A3.5 / B 3.5 (best statistics on the board)						+1,099 pass			
A3.5 / B wide	1,098	−1,144	+971	12 / 16	+202 pass	+648 pass	0.915 fail	+360	+611
A1.5 / B 2.5 (the original ladder guess)	1,711	−2,177	−62	7 / 16	−985	−518	0.465	−169	+107
A0.5 / B1.0 — the M5 proposal	2,973	−3,242	−1,127	8 / 16	−3,392 fail	−2,036 fail	0.302 fail	+699	−1,826 fail
A1.5 / B tight- chandelier	2,143	−3,718	−1,603	5 / 16	−3,238	−2,654	0.176	−763	−840

The per-fill trap. M5's 0.5/1.0 cell loses less per fill than the loaded default (−\$1.09 vs −\$1.72) and wins 67% of its Lot-A trades. It still loses \$1,127 more in total, because it takes **2,973 fills instead of 1,228**. A tight exit re-arms the slot constantly; each re-arm is another 1R stop and another \$1.50. Judge cells on total money, not on hit-rate or on \$/trade.

Fact table 6 — segmented by regime (a blanket number would be a bug)

Regime is keyed at the entry minute on the 1-min ATR-14 level, the 30-bar efficiency ratio and a 30-minute up-range break — not the clock. Same definitions REV2·Q1 used, with clean-trend split by direction because this gate is long-only.

REGIME AT ENTRY	FILLS	LOADED A2.5/WIDE	\$/FILL	WIN %	BEST CELL IN THE BUCKET	POLICY READ
in-between / building (ER ≥ 0.20, no range break)	266	−3,142	−11.81	20	A0.75/B6.0 −1,864	the entire loss lives here, and no exit fixes it. 130 configs, best is still − \$1,864. This is the false-start

REGIME AT ENTRY	FILLS	LOADED A2.5/WIDE	\$/FILL	WIN %	BEST CELL IN THE BUCKET	POLICY READ
						bucket REV2·Q1 found in the live book (–\$571 there) — same finding, 5× the n.
clean up-trend (ER ≥ 0.20, net30 ≥ 30pt, 30-min high broken)	106	–1,185	–11.18	22	A1.5/B1.5 +210 (246 fills)	thin — direction only. Uncomfortable direction, too: on confirmed breakouts the tight pair is least-bad. Do not act on n=106.
normal chop (ATR 6–14, ER < 0.20)	393	+54	+0.14	27	A3.0/B3.0 +304	scratch; wider helps slightly
violent whipsaw (ATR ≥ 14, ER < 0.20)	452	+2,185	+4.84	29	A3.5/B3.5 +2,729	the only bucket that pays, on 10 green sessions of 16. Fast-slope impulses inside a wide, inefficient range. This flatly contradicts the standing prior that whipsaw is where the desk dies — for this gate it is the earner.
dead chop (ATR < 6)	11	–28	–2.56	27	—	n=11, no read — the ATR floors do their job

Fact table 7 — segmented by entry ER (the sharpest cut in the whole section)

Same trades, sliced on the 30-minute efficiency ratio at entry — the exact variable the router arms grind on and the exact variable deciders.ER_FLOOR gates it on.

ER AT ENTRY	FILLS	LOADED A2.5/WIDE	A3.0/WIDE	A3.5/B3.5	M5 A0.5/B1.0	READ
ER < 0.20 (the "chop" the floor)	854	+2,233	+2,645	+2,838	+445	every dollar grind_long

ER AT ENTRY	FILLS	LOADED A2.5/WIDE	A3.0/WIDE	A3.5/B3.5	M5 A0.5/B1.0	READ
is designed to block)						makes is here. 12 of 16 sessions green at A3.0/wide, LODO worst + \$1,308 — but strip the three best sessions and it is −\$587, so treat the size as a direction.
ER 0.20 - 0.35	303	−3,314	−3,453	−2,629	−2,792	the whole loss, in one band, at every exit setting. This is the "rotation legs that look trendy but reverse" band the 07-30 note identified — confirmed on 5× the sample.
ER 0.35 - 0.50 (above the router's arming floor)	65	−665	−354	−256	−697	also red, on thin
ER ≥ 0.50 (the BIG-TREND rung)	6	−370	−336	−367	−197	n=6 — no read at all, but nothing here says the RIDE rung is being missed

THE AWKWARD IMPLICATION FOR THE ER FLOOR

The 07-30 diagnosis was right — the **0.20-0.35 band is the bleed**. But the fix that was shipped is a floor at 0.35, and on this tape a floor at 0.35 would delete the only profitable population (ER < 0.20, + \$2,233) while keeping a band that is red on 65 fills. **A band-exclusion (skip 0.20-0.35) is what the data supports; a floor is not.**

Filtering the sim to ER ≥ 0.35 leaves 71 fills worth −\$1,035; filtering to both designed floors (ER ≥ 0.35 **and** ATR ≥ 24) leaves **15 fills in 16 sessions**, worth −\$246 — the floors as written do not sharpen the gate, they silence it. Flagging, not actioning: that is an arming decision and it belongs to you and the router, not to `exit_overrides.json`.

Fact table 8 — time of day, reported separately because it is a different axis

WINDOW (UTC, AT ENTRY)	FILLS	LOADED A2.5/WIDE	\$/FILL	BEST CELL	READ
overnight / pre-open (< 13:30)	775	-1,492	-1.92	A3.5/B6.0 -478	63% of all fills, all of them red; wide is least-bad
US open 13:30-15:00	78	-57	-0.74	A3.0/B3.0 +344	thin (n=78) — direction only
US session 15:00-20:00	255	-587	-2.30	A0.5/B1.5 +770	the one window where tight wins — and it is the window the "big runners live in the US afternoon" prior claims for the ride exit. The prior is wrong here, exactly as REV2-Q1 found on the live book.
late close-out (≥ 20:00)	120	+21	+0.18	A3.5/wide +251	the only green window on the default, on 32% win — a fat-tail hour

The time-of-day axis and the regime axis **disagree**: regime says go wider, the US afternoon says go tighter. `exit_overrides.json` is a single static cell and cannot express either. That is an argument for the router eventually owning the exit axis the way it owns arming — not an argument for picking the tighter cell because one 255-fill window prefers it.

Fact table 9 — the thing I found on the way (and it is not small)

REV2-Q1 asked what the machine does with your hands off. Part of the answer is that **grind_long's own entry filters are not running**.

FACT	EVIDENCE
The filters exist	<code>deciders.ER_FLOOR = {"grind_long": 0.35}</code> (added 07-30) and <code>ATR_FLOOR = {"grind_long": 24.0}</code> (07-25 rehab)

FACT	EVIDENCE
They are applied by sub-slot tag	tournament.py:139/142 call er_blocks(slot) / atr_blocks(slot) with the raw intent tag. Under the scale-out slate that tag is grind_long_A / grind_long_B. _base() is applied to the bench-list, the 55s veto and the exhaustion confirm — but not to these two.
Live proof	Of the 40 grind_long_A/_B fills booked after the ER floor went in on 07-30: 31 had entry ER below 0.35 and 22 had entry ATR below 24; 33 of 40 violated at least one floor. Lowest ER filled: 0.03. Lowest ATR filled: 11.2.
Money direction	Ironically, in this sample the bypass helped — the sub-0.20-ER fills are the profitable population (fact table 7). Fix the bug and you must fix the thresholds in the same move, or you will delete the good half.

Not actioned here — this section is read-only on live config. It belongs on the Saturday card as a code fix plus a threshold decision, taken together.

Fact table 10 — the tail question, answered directly

The standing objection to any capped exit is that grind's live profit lives in the tail: **CHANDELIER + MANUAL_CLAIM exits, 16 trades, +\$1,566**. A 3.5R cap has to be judged against that, not against an average.

WHERE THE TAIL MONEY ACTUALLY IS	N	\$	SHARE
Sim, WIDE lot — winning exits below 2R	47	+1,312	10.4%
Sim, WIDE lot — winning exits 2R to 3.5R	7	+561	4.5%
Sim, WIDE lot — winning exits above 3.5R	52	+10,678	85.1% (81.4% above 5R alone)
Live, the 16 tail trades — exits at or under 3.5R	13	+1,036	66%
Live, the 16 tail trades — exits above 3.5R	3	+530	34%

Both things are true and they point the same way. In the sim the tail is enormous — 85% of the wide lot's gross profit comes from exits beyond 3.5R, so capping there gives up something real. But the wide lot as a whole still nets −\$938, because it pays for that tail with give-back on every trade that peaks and reverses. And in the live book, two-thirds of the celebrated "tail" money came out at or under 3.5R anyway — a 3.5R target would have banked \$1,036 of the \$1,566 and given up \$530. That is why the recommendation keeps **Lot B on the wide chandelier**: you keep an

uncapped lot for the genuine runners and you take the certain money slightly later on Lot A. What the data forbids is capping both lots tight, which is precisely what M5 proposed.

Fact table 11 — reproducing (and refuting) the Part 2 · LAB cell

M5 came from the Friday lab: grind_long → {a_r 0.5, b 1.0}, off a MED-TREND cell of n=78 at +\$10.3/trade and 59% win, with SCALP-CHOP n=542 agreeing. I re-ran it **the lab's own way** (paired, same-entry repricing) but with two corrections: **250 ms ticks instead of 5-second bars**, and **grind_long only** (the lab's cell pooled both directions; only the long side is armed).

RUNG (LAB'S OWN DEFINITI ON)	N	A0.5/B1.0 \$/SIGNAL	WIN %	A2.5+WID E	A3.0+WID E	WIDE×2	VERDICT ON THE LAB'S CLAIM
MED-TREND (ER 0.30- 0.50)	121	-5.19	46	-20.58	-18.85	-29.29	half- reproduced. The relative claim survives — tight does beat the 2.5/ wide guess here. The absolute claim does not: +\$10.3 at 59% win becomes – \$5.19 at 46% win once you price on ticks and drop the short side. You cannot build a live cell on a number that is negative.
SCALP-CHOP (ER < 0.30)	777	-0.84	52	+7.57	+9.57	+11.01	refuted outright. The lab said this bucket also picks 0.5/1.0. On ticks it picks the

RUNG (LAB'S OWN DEFINITI ON)	N	A0.5/B1.0 \$/SIGNAL	WIN %	A2.5+WID E	A3.0+WID E	WIDE×2	VERDICT ON THE LAB'S CLAIM
							widest thing available, by \$12 a signal on 777 signals.
BIG-TREND (ER ≥ 0.50)	4	−97.51	—	−108.68	−108.68	−108.68	n=4 on the long side — no read, as REV2-Q3 warned

And on the sequential (policy) sim, which is the one that counts, the SCALP-CHOP rung is **−\$1,742 for 0.5/1.0, +\$53 for the loaded default, +\$1,242 for A3.5/B3.5. M5 should not be actioned.** Two things caused the original result: 5-second repricing (which flatters tight exits, because a 5s bar cannot show you the tick that stopped you out first), and pooling grind_short into a cell used to set a long-only gate.

Fact table 12 — what each lever is actually worth (16 sessions, tick-repriced)

The point of the whole exercise: the exit knob is the small one.

LEVER	CHANGE	NET \$	DELTA	ROBUSTNESS
nothing	loaded default A2.5 / wide	−2,115	—	10 of 16 sessions green
EXIT	→ A3.0 / wide (recommended)	−1,498	+617	direction proven, size not (see fact table 5)
EXIT	→ A3.5 / B3.5 (best stats, a spike)	−414	+1,702	passes LODO, strip-3, bootstrap and both halves — fails plateau
EXIT	→ A0.5 / B1.0 (M5)	−3,242	−1,127	fails everything
ARMING	default exit, but never arm in "building"	+1,027	+3,142	10/16 green, LODO worst −\$116, archive half −\$2,375 — fails the time split
ARMING	default exit, only arm when entry ER < 0.20	+2,233	+4,348	9/16 green, LODO worst +\$1,002, strip-top-3-days −\$935, archive half −\$1,150 — fails the time split
BOTH		+2,645	+4,760	

LEVER	CHANGE	NET \$	DELTA	ROBUSTNESS
	A3.0/wide + only arm at ER < 0.20			12/16 green, LODO worst +\$1,308 — same time-split failure

Read the last three rows with discipline. The arming levers are worth 3–5× the exit lever in dollars, and every one of them is **green only in the capture half and red in the archive half**. The exit direction is the opposite: A3.5/B3.5 beats the default by +\$997 in the archive half and +\$705 in the capture half. So the finding I am willing to sign is the small one, and the big one goes on the watch list with a date on it.

THE VERDICT, IN THE ORDER YOU SHOULD READ IT

1. Change the cell, but change it small. Write `{"grind_long": {"a_r": 3.0, "b": "wide"}}` — Lot A 2.5R→3.0R, Lot B untouched. Worth **+\$617 over 16 sessions against the loaded default's −\$2,115**. This is also the exact cell REV2·Q3 pre-registered for Monday 2026-08-03 and pinned until the 08-21 grade date, so writing it costs nothing in test integrity and gains the pin. It is a move along a proven direction, not a proven number.

2. Do not action M5. `{a_r 0.5, b 1.0}` is −\$1,127 against the default, green on 5 sessions of 16, and fails leave-one-day-out, strip-the-best-3 and the bootstrap. The number that motivated it (+\$10.3/tr, 59% win, n=78) does not survive being priced on ticks and restricted to the long side; it becomes −\$5.19/signal at 46% win. Delete the item.

3. The RIDE rung is not the question it was made out to be. On 16 tick-covered sessions the BIG-TREND cell for grind_long is n=6, and the wide lock-chandelier is mid-pack everywhere: −\$938 for the lot, no better than a plain 3.0R target. Keep it on Lot B because an uncapped lot is cheap insurance against the runner, not because it has been shown to earn.

4. The real money is in arming, and it is not proven yet. Keeping grind_long out of the ER 0.20–0.35 / "building" tape is worth +\$3,100 to +\$4,800 on the same 16 sessions — 5× the exit lever — but it is green in the capture half and red in the archive half, so it is a direction with a strong

prior, not a measurement. It needs the 5-second back-extension (REV2·Q3 Track A) to be gradeable, and it belongs to the router, not to the exit file.

5. And the honest floor: grind_long is red at every exit setting on this tape. The best of 130 configurations still loses \$414 over 16 sessions, before a penny of slippage. No exit cell rescues this gate. Choosing the cell is damage control while the arming question gets answered.

Method & limits — read before quoting any number above

- **Every sweep number is a ceiling.** Fills are at the exact stop/target/trail level with no slippage and \$1.50 round-turn fees. Measured against the 40 real STOP fills in this gate's own live history, real stops filled **7% worse** than the model; the desk's standing finding is that live losses run **1.1-3.1× modelled**. Every dollar figure here is the best case, and the best case is negative.
- **Exits are priced on 250 ms ticks only.** 15.3M archive ticks (07-05→07-17) plus 10.1M capture ticks (07-23→07-31). No exit anywhere in this section was decided on a bar. The 15 further sessions that have 5-second bars but no ticks were excluded from every number.
- **Sample.** 5,874 grind_long signal-minutes on the 31-session tape; 3,064 of them on the 16 tick-covered sessions; 1,228 lot-fills for the baseline cell after the one-position-at-a-time slot logic. Regime buckets: whipsaw 452, normal-chop 393, building 266, clean up-trend 106, dead-chop 11. **Anything under $n \approx 150$ in this section is flagged inline and is a direction, not a measurement** — that covers clean up-trend, US-open, $ER \geq 0.35$ and the whole BIG-TREND rung.
- **Density.** The sim takes ~ 77 lot-fills a session; the live gate, with router benching, took ~ 20 . Rank configs on these numbers; discount the dollars $\sim 3.8\times$ before treating them as a forecast.
- **Marked-to-market residue.** 14% of entries neither stop, target nor trail inside the 120-minute MAX_HOLD and are marked at the last tick — the same entries in every config, so they do not tilt one config against another.

- **Not tested here.** The give-back overlay (off for both lots under this slate), the `adaptive_exit` selector (also off under scale-out), and the 5-second back-extension to 06-22. All three are separate questions.
- **Read-only.** No live file was modified — `gate_switches.env`, `exit_overrides.json` and every service are untouched. Writing the recommended cell requires an edit plus a tournament restart, and it is reversible by deleting the key.

REV3 Consistency audit — did the Rev2 fold-in actually land in the published report?

Garrath — short answer: **no, and not by a little.** Five agents wrote five Rev2 sections. The fold-in put **two paragraphs and one table** into the published report and stopped. **Q0 (the Monday action card), Q3 (the $n=5 \rightarrow n=24$ correction) and Q4 (the router's real value) are not in the published report in any form** — zero words, zero numbers. Q1 contributed its opening callout and nothing else — not its exit sweep, not its Thursday reconstruction, not its re-nomination. Q2 contributed one correction paragraph — not its exit-cell sweep, and critically **not the dead-floors finding**. The good news buried inside the bad: the specific hazard you were worried about — Q0's superseded rows M4 and M6 sitting in the report telling a reader to do the wrong thing — **did not happen, because Q0 never made it in either.** The hazard that did happen is the mirror image: **the report still says everything Q1-Q4 were written to overturn, and none of the overturning.**

★ **THE SINGLE WORST LINE IN THE PUBLISHED REPORT, AND IT WAS INTRODUCED BY THE FOLD-IN**

The Part 2 "Correction (applied 08-01)" paragraph — one of only two Rev2 patches that landed — states:

"live `abs_veto_long` made +\$326 over 47 trades this week". That number is

wrong. Straight from data/gazbot7.db, abs_veto_long over the report's own week is **+\$59.5 over 59 lots (43 signals)**. No window produces \$326/47 — Mon-Fri gives +\$131.0/56, from-07-28 gives +\$177.0/52. It also contradicts the same document's Part 1 card table, which correctly prints abs_veto_long | +59.5 | 59 . A correction paragraph carrying a \$266.5 error, in the section whose whole job is honesty, is the most damaging single defect in the artefact.

What actually landed — the fold-in, measured

Diffed the three .bak files the fold-in left behind against the section files that were rendered. This is the complete inventory of what changed:

SECTION FILE	WHAT THE FOLD-IN CHANGED	FROM WHICH REV2 SECTION
part1_live.html	Added the "★ READ THIS FIRST — the automated desk lost money this week" callout + its two tables (−\$598 / 194 / 35.1%, and the per-gate auto-vs-hand split). Also corrected grind's Thursday figure (+\$307/37 → +\$704/46), the STOP_UNFILLED status (branch-not-deployed → d0218c6 live and verified), the router switch count (31 → 59 across 590 ticks), and inserted the run charts.	Q1 headline only
part2_shadow.html	Added the "Correction (applied 08-01)" paragraph — abs_veto is already live two-sided, there is no thrust_short gate. Contains the wrong \$326/47 figure above.	Q2 §1-2 premise only
part2_6_router_review.html	Day labels shifted Sun→Mon...Thu→Fri; "+\$625 incl Sunday, +\$852 Mon-Fri" → "+\$625 over 206 trades". Nothing from Q4.	none
– (no file)	Q0's Monday/Saturday/Build action card — never rendered into the report.	Q0, 100% dropped
– (no file)	Q3's n=5→n=24 correction, the withdrawn archive refusal, the dated proof plan, the MED-TREND n=326 finding.	Q3, 100% dropped
– (no file)	Q4's A/B double-count, the CHOP/TREND_UP split, the +\$1,615/week fix, the leakage reconciliation, the corrected +\$250 [−\$440,+ \$940].	Q4, 100% dropped

Confirming greps on the published HTML: the strings REV2 , n=24 , 1,615 , back-extension , k1.5 , M4 , M6 , D2 return **zero hits**. TREND_UP returns one hit — and it is a Rev1 sentence in Part 2.6 §3 that names the segment without ever pricing it.

Verification of my own baseline, before I criticise anyone

Everything I score below is checked against venue truth, not against the Rev2 sections' say-so:

CLAIM	SOURCE	VERIFIED?
Automated-only −\$598 / 194 fills / 35.1% win	gazbot7.db trades, closed_at ≥ 07-26	☐ exact
Manual claims +\$1,223 / 12 / 100% win	exit_reason = MANUAL_CLAIM	☐ exact
exhaustion_short −\$56 automated over 44 fires	base slot +\$365/37 · A − \$157/4 · B −\$264/3	☐ exact
abs_veto_long +\$59.5 over 59 lots	three slots summed	☐ exact — and the published "+\$326/47" is not reproducible on any window
grind_long has no exit_overrides key (rides the built-in wide default)	data/exit_overrides.json	☐ only exhaustion_short, abs_veto_long, abs_veto_short are keyed
rgv_short is armed live while Part 1.5 verdicts it SHADOW	data/gate_switches.env	☐ rgv_short=on

The contradictions still LIVE in the published report — ranked by money

1 · The router's own value: −\$1,216 is printed twice, and it is wrong by \$1,466

WHERE THE SUPERSEDED CLAIM STILL IS	VERBATIM
Part 2.5 §3 "Router week-review", WEEK row	Week — \$15,788 \$17,003 −\$1,216 −\$1,083
Part 2.6 §3, WEEK row	WEEK — \$15,788 \$17,003 −\$1,216 \$0 (durable days)
Part 2.6 §3, prose	"the raw net is −\$1,216, but two things inflate the 'missed' column — the dual A/B slots double-count one signal , and the giant 'missed' figures on 07-29 ... are dominated by grind_long shadow winning during brief benched windows"
Part 2.6 §8 verdict	"Is the router earning its keep? Yes — as a defense. "
Where the correction appears	Nowhere. Q4 is absent from the published report.

Would a reader be misled, and by how much? Yes, in three separate ways. (i) He reads a \$1,216 deficit that is 60% bookkeeping — Q4 proves the A/B legs are byte-identical in the nightly's own JSON (grind_long_A and grind_long_B both 116 / 46 won / \$2,633.5 / 70 lost / −\$2,228.5) because reprice_scalp() ignores the slot suffix. Merging them takes the week from −\$1,101 to −\$440, recovering \$661. (ii) The report **names** the double-count and then leaves it in the total anyway — that is worse than not noticing, because it signals the author knew. (iii)

He is handed a "Yes" verdict off a statistic that flips sign on leave-one-day-out (–\$804 dropping 07-30, +\$230 dropping 07-29) and therefore supports no verdict at all.

The money he doesn't get told about: +\$1,615/week. Q4's segment split is the biggest number in the entire Rev2 round and it is robust in a way almost nothing else in this report is — CHOP benching +\$1,088 at zero haircut, TREND_UP benching –\$1,380 needing a 52% winner haircut to break even, and with the TREND_UP benches removed the week is positive in **all nine** robustness folds (+\$533 to +\$1,021). That is the highest-value, best-evidenced change in the report and the published report does not contain it.

Exact minimal fix

In **Part 2.6 §3**, replace the sentence beginning "Honest reading (this number is noisier than it looks): the raw net is –\$1,216..." with: "Corrected (08-01): –\$1,216 counts every suppressed signal twice — the dual-slot A and B legs are byte-identical in the nightly's own output because the repricer ignores the slot suffix. Merged, the week is 353 signals, \$8,644 saved, \$9,084 missed, net –\$440. And the blend hides two opposite, individually-robust halves: benching in CHOP is +\$1,088 at zero haircut; benching in TREND_UP is –\$1,380 and would need the shadow to overstate winners by 52% to break even. Stop benching momentum in a confirmed TREND_UP and the same week is +\$1,175 — a +\$1,615 swing, positive in all nine leave-one-day-out and strip-the-best folds." Change both WEEK-row net cells from **–\$1,216** to **–\$440 (A/B merged; –\$1,216 as originally double-counted)**. In **§8**, change "Is the router earning its keep? Yes — as a defense." to "Split verdict. The defence earns its keep decisively (+\$1,088 in CHOP, +\$613 on fader benches, robust to every fold). The offence does not (–\$1,380 benching momentum in TREND_UP). The blended number has no statistical standing — it flips sign on which single day you drop." Add a row 8 to the Saturday table: "Stop benching momentum in a confirmed TREND_UP | +\$1,615/wk at zero haircut, positive in 9/9 folds | RULE — already banked as [[us-open-dont-bench-trend-rider]], simply not enforced."

2 · Every ER and ATR floor on the desk has been dead since 07-29 16:24 — and the report reasons from them five times

This is Q2 §3 and it is the most consequential factual correction of the whole Rev2 round. `tournament.step()` passes the sub-slot tag (`grind_long_A`) to `er_blocks/`

atr_blocks, but ER_FLOOR/ATR_FLOOR are keyed on base names, so the lookup never matches. The journal shows 7,734 suppression lines up to 2026-07-29 13:29:04 and **zero since**. Five gates affected: grind_long (ER 0.35 + ATR 24), abs_veto_long (ER 0.25), abs_veto_short (ATR 16), capitulation_long (ATR 10). **None of this appears in the published report.**

WHERE THE REPORT REASONS FROM A FLOOR THAT IS NOT RUNNING	VERBATIM
Part 1, gate × day grid	"That is not a flaw — it's what the 0.35 ER floor is for , keeping it dormant unless the day really trends"
Part 1.5 §2 grind_long	" Two live floors are actively destroying the gate . The ER ≥ 0.35 floor is backwards ... The ATR ≥ 24 floor is far too tight for this regime."
Part 2.5 §2 "The GRIND lead"	"Grind fired 52 of 56 times below 0.35 on this measure ... Whatever ER the live gate is computing, it is not the trend-efficiency the floor claims to guard ... ER filters are fake here "
Part 2.6 §2, CHURN row	"grind's 0.35 ER floor + abs_veto's veto meant most premature arms never fired — process error, not \$ error"
Part 2.6 §4 + §8 action 6	"The grind 0.35 ER ARM floor ... is validated separately (+\$728 over 10d) — leave it " · "arm once on ER≥0.35 sustained ... grind self-gates via its 0.35 floor "
Where the correction appears	Nowhere.

Would a reader be misled? Badly, and in the most expensive direction. Part 2.5 §2 gropes at the truth and lands on the wrong explanation: it observes grind firing 52 of 56 times below its floor and concludes ER filters are fake on this tape — a claim about the market. The real cause is a lookup bug — a claim about our code. Those two conclusions point at completely different fixes. Meanwhile Part 2.6 §8 tells Garrath his churner "self-gates via its 0.35 floor", i.e. tells him to rely on a safety feature that has not executed once since Wednesday. And Part 2.6 §2 grades a router CHURN as costing \$0 on the strength of that same absent floor.

Exact minimal fix

Insert a red callout immediately under the Part 2.5 §2 heading and repeat it under Part 2.6 §4: "★ Correction (08-01): ER_FLOOR and ATR_FLOOR have not executed since 2026-07-29 16:24. tournament.step() passes the sub-slot tag to er_blocks/atr_blocks, which are keyed on base names, so the lookup never matches — 7,734 suppressions before 07-29 13:29, zero after; five gates affected. Everything below that reasons from a live floor is reasoning about a filter that is not running. Grind did not fire 52 of 56 times below its floor because ER filters are fake; it fired because there was no floor. Fix the wiring first, then re-derive each floor on tape that was actually filtered — do not let a silent bug decide the floor's value." Then delete the clause "grind self-gates via its 0.35 floor" from Part 2.6 §8 action 6, and change Part 1.5 §2's "Two live floors are actively destroying the gate" to "Two

floors are configured against this gate — and neither has executed since 07-29 (see the correction in Part 2.5); as written they are backwards and too tight, which is why they must not be re-wired at their current values."

3 · exhaustion_short is still crowned "the survivor" in five places; the correction sits in exactly one box

The fold-in did put the honest number at the very top of Part 1 — credit where due, a reader starting at page one meets **−\$56 over 44 automated fires** before anything else. But nothing downstream was touched, so the two readings sit ninety words apart and then diverge for the rest of the document.

#	WHERE THE SUPERSEDED CLAIM STILL IS	VERBATIM
a	Part 1, "Read this first — what SURVIVED the week" (immediately below the correction callout)	"exhaustion_short was green on 4 of the 5 sessions ... 54% of its trades won, +\$316 for the week. It is the steadiest earner on the desk. "
b	Part 1, six-gate card table, verdict cell	exhaustion_short +316.5 48 54 ... SURVIVOR — steadiest
c	Part 1, gate × day grid, closing line	" That row is why exhaustion is the survivor. "
d	Part 1, Case 1 Thursday	"exhaustion_short banked four fade-claims in a row for +\$372.5 ... judge this gate on its live fades , not its fixed-target shadow"
e	Part 2, Relegation view, exhaustion row	" LIVE exhaustion_short is the desk's cleanest fader ... Judge it on live fills."
f	Part 2.6 §2, graded-ledger table	"07-30 14:57 arm short-side fade book ... GOOD +\$372 exhaustion_short went 4/4 live fading the roll-over — best offense of the day"
—	Correction	Part 1 callout only: "−\$56.0 44 +372.5 4 +316.5" and "unattended it is roughly flat-to-negative"

Would a reader be misled, and by how much? A reader going strictly top-to-bottom gets the correction first — so the acute risk is low. But it does not survive contact with the rest of the document: rows (b) and (c) re-assert the crown within two screens, and rows (e) and (f) are in sections a reader will very reasonably jump straight to. Row (f) is the sharpest: **the router is graded ☐ GOOD and credited +\$372 for a sequence that was four hand-clicks**. Q1 repriced those exact four legs on 250 ms tape under the exit config actually loaded at the time and got **−\$51.2**. So the router's best-graded offensive call of the week is worth −\$51, not +\$372 — a \$423 swing inside the router's own scorecard, which is the section that then concludes the router earns its keep.

Exact minimal fix — six cells, no re-analysis needed

(a) "+\$316 for the week. It is the steadiest earner on the desk." → "+\$316 for the week — but \$372.5 of that was four trades you closed by hand. Automated it is – \$56 over 44 fires (see the callout above). It is the steadiest thing on the desk, not the biggest earner: 50% win, a median trade of essentially zero, and the smallest leave-one-day-out damage of any gate that traded in size." **(b)** verdict cell "SURVIVOR — steadiest" → "FLAT AUTOMATED (–\$56 / 44) — the +\$316 includes 4 hand-claims". **(c)** "That row is why exhaustion is the survivor." → "That row is why exhaustion is the steadiest — but read it against the automated split above: unattended the row is flat." **(d)** append: "All four of those were MANUAL_CLAIM exits. Repriced on 250 ms tape under the exit config actually loaded, the same four entries would have been –\$51.2 — three of them were sitting on positions the tape later ran 49 to 98 points through the stop." **(e)** "the desk's cleanest fader" → "the desk's steadiest fader — though its live green is hand-claim dependent (–\$56 over 44 automated fires)". **(f)** \$ cell "+\$372" → "+\$372 hand-claimed / –\$51 automated" and grade "□ GOOD" → "□ GOOD (regime read) / △ the \$ credit is manual".

4 · BIG-TREND is stated as n=5 and the report refuses an archive it had already used

WHERE THE SUPERSEDED CLAIM STILL IS	VERBATIM
Part 2 · LAB (b), ladder table	grind (momentum) BIG-TREND (5) 2.5 / wide-chandelier wideChand ×2 (+\$12.0) — n=5 UNTESTED
Part 2 · LAB (b), ★ CRITICAL caveat	"grind's BIG cell is 5 trades ... I did not manufacture a BIG-TREND sample from the retired V5 archive: it is a different desk, contract era and gate set, and repricing V7 gate entries on it would be a fabricated comparison, not an honest one. The only honest tick tape available (V7 capture.db) spans 07-15 to 07-31 ... grind's BIG-TREND exit is a guess wearing an n=5 hat. "
Part 2 · LAB, deployable policy table	grind_long wide-chandelier (unproven, n=5)
Part 2 · LAB, bottom line	"the honest gap that stays open is BIG-TREND: the desk still has not seen a clean trend week to prove the RIDE rung"
Where the correction appears	Nowhere. Q3 is absent.

Would a reader be misled? Yes, three ways, and one of them is a factual falsehood in the report's honesty showcase. (i) The cell is **n=24**, not 5 — five is the count of live grind_fast shadow entries in an $ER \geq 0.50$ minute; the ladder verdict was computed on a 24-entry tick-tape replay. The caveat quotes the live number in a caveat about the replay number. (ii) The refusal of the archive is **false, and the lab had already used it** — 11 of the 24 BIG-TREND entries come from that pre-capture stretch, and the "different contract era" objection dies on the seam: last archive tick 29,623.00 at 15:35:29.973, first capture bar opens

29,622.75 at 15:35:30, $r=0.9998$ across 2,925 overlapping minutes. Same instrument, same front month. A report that invents a principled refusal for tape it silently consumed is worse off than one that never mentioned it. (iii) **The rung is not a hypothetical — it is live right now.** I checked data/exit_overrides.json: it keys only exhaustion_short, abs_veto_long and abs_veto_short. grind_long has no entry, so it inherits the built-in _BIG_RUN default, which is literally Lot A @ 2.5R + Lot B wide lock-chandelier. The report frames it as a stone left unturned; it is an unhedged live position.

And the bigger prize is missing too. Q3's unasked-for finding: on grind's **MED-TREND** rung, **n=326**, the live A2.5+WIDE default makes +\$8.49/trade against TIGHTx2's +\$17.60 — a **\$9.11 × 326** gap with real n behind it, needing no new tape at all. That is a better-evidenced number than the BIG-TREND question the caveat spends four hundred words on, and it is nowhere in the published report.

Exact minimal fix

Table cell "**BIG-TREND (5)**" → "**BIG-TREND (24)**"; verdict cell "**UNTESTED**" → "DIRECTION SUPPORTED (bootstrap $P=0.95$, plateau 2.5–3.5R, both time-halves agree) · MAGNITUDE UNPROVEN (strip-3 takes +\$62 to +\$10)". In the ★ CRITICAL caveat, delete from "I did not manufacture a BIG-TREND sample..." through "...does not contain a clean trend week for grind" and replace with: "Eleven of the 24 BIG-TREND entries already come from the pre-capture MNQ archive (ticks.db, 07-05→07-15) — an earlier draft refused that archive as 'a different contract era', which was wrong on the facts and wrong about our own lab, which had already used it. The tapes join one tick apart (29,623.00 → 29,622.75, $r=0.9998$ over 2,925 overlapping minutes): same instrument, same front month. What is genuinely missing is n, not honesty — the pre-registered grade needs 39 paired entries (derived from a paired sd of \$247 on a \$79.4 diff) and we have 24. And note this rung is LIVE today: grind_long has no key in exit_overrides.json, so it inherits the built-in A@2.5R + wide lock-chandelier. This is not an unturned stone, it is an unhedged position." In the policy table change "wide-chandelier (unproven, n=5)" → "wide-chandelier — LIVE TODAY by default, n=24, under pre-registered test to 2026-08-21". Add one row to the same table: "grind_long | MED-TREND | n=326: live wide default +\$8.49/tr vs tight ×2 + \$17.60/tr — the biggest evidenced exit gap on the desk, and it needs no new tape."

5 · The report tells Garrath explicitly to leave abs_veto's exit alone — and the sweep says its long cell is dead last of eleven

WHERE THE SUPERSEDED CLAIM STILL IS	VERBATIM
Part 2 · LAB (b), ladder table	thrust / abs_veto raw every rung (25-333) 1.5 / 2.5 bleeds in ALL: BIG —\$13, MED —\$36, CHOP —\$39, STAY —\$67 exit can't save it
Part 2 · LAB (b), plain-English bullets	"thrust, rgv and capitulation bleed in every regime no matter what exit you put on them. That is the headline finding: for these three, the exit is not the lever "
Part 2 · LAB, policy table + note	"thrust / abs_veto exit is not the lever — keep the abs_veto entry veto; manage via gate on/off, not exit style" · " abs_veto's dollars come from its entry veto, not its exit R's "
Where the correction appears	Nowhere. Q2 §6 is absent.

Would a reader be misled, and by how much? Yes — this is the one where the report actively instructs inaction on a change two independent sweeps say to make. Q2 §6 replayed **81 abs_veto_55s LONG signals** on 250 ms ticks (validated against the shadow repricer at 98.7% sign agreement, MAE \$2.70) across eleven exit cells. The live **A1.5/B2.5 ranks last on every measure**: +\$13.6/signal vs +\$22.7 for A1.0/B1.5, LODO-min +\$4.7 vs +\$16.8, strip-best-3 +\$6.8 vs +\$19.0 — and the winner sits on a broad plateau (0.75/1.5, 1.0/1.5, 1.0/2.0 all within \$4), not a spike. Q1 reached the same conclusion by a different route on 42 automated entries: plateau at 0.75-1.0R (+\$292 / +\$532) collapsing to +\$26 at exactly the loaded 1.5R. **Two independent sweeps, two different samples, same answer.** Worth roughly **+\$9.1 per signal × 81 signals ≈ +\$737** on the swept book, **~+\$390/week** at the live rate of 43 long signals. The short side genuinely is noise (all cells +\$15-18) and should be left at 1.5/2.5, exactly as the report says — but the report does not distinguish the sides.

Note the collision the report creates on its own. The published sentence "the exit is not the lever" is derived from the raw un-vetoed thrust rows, then applied by name to **abs_veto**, which is a different, filtered population. That is a sample switch inside one sentence, and it is what produced the wrong instruction.

Exact minimal fix

Change the policy-table cell for thrust / abs_veto to: "the raw un-vetoed thrust cannot be rescued by an exit — but the LIVE abs_veto_long cell can, and it is the worst cell on the board. A tick-honest sweep of 81 long signals ranks the loaded A1.5/B2.5 last of eleven (+\$13.6/sig, LODO-min +\$4.7) against A1.0/B1.5 (+\$22.7, LODO-min +\$16.8) on a broad plateau. Change exit_overrides.json → abs_veto_long: {a_r 1.0, b 1.5}. Leave abs_veto_short at 1.5/2.5 — the R choice is

genuinely noise on that side (+\$15 to +\$18 at every cell). No exit rescues abs_veto_short in violent whipsaw (−\$20 to −\$33 at every R from 0.5 to 6.0); only not trading it does." Delete the sentence "abs_veto's dollars come from its entry veto, not its exit R's" or qualify it to "most of abs_veto's edge is the entry veto — but the long side's R's are still the worst cell swept." Rewrite the bullet "for these three, the exit is not the lever" to name the population: "for the raw, un-vetoed versions of these three the exit is not the lever."

6 · exhaustion_short's exit: the report gives three different answers and the sweep gives a fourth

WHERE	WHAT IT SAYS THE EXIT SHOULD BE
Live today (exit_overrides.json)	Lot A 0.5R / Lot B 1.5R
Part 2 · LAB (a), verdict	"keep running — two signals is not a grade" → keep 0.5/1.5
Part 2 · LAB (b), ladder	tight-chandelier in CHOP, scalp 1.5/2.5 in MED, tight-chandelier in BIG
Part 2 · LAB, note under the policy table	"This lab says exhaustion_short should be regime-graded ... rather than a flat 0.5/1.5"
Q1 fact table 6 — 64 entries, 250 ms, absent from the report	Lot A 0.75R (+\$510, LODO +\$239, 5/9 green) — 0.5R is the worst scalp rung on the ladder (+\$306). Lot B k1.5 chandelier (+\$502, the highest LODO floor in the whole sweep at +\$330, worst day −\$25). And: no exit is robustly green in violent-whipsaw , where the gate took 28 of its 62 entries — the k1.5 tight chandelier is the worst config in that bucket at −\$202. The lever is the bench, not the R.

Would a reader be misled? He is given three recommendations in one section that do not agree with each other, and then the actual sweep answer is missing. There is a second, structural error underneath it. The report treats exhaustion as one gate on one exit all week. It was not: **until Wed 07-29 19:10 UTC it ran a single slot on the regime-3 adaptive chandelier (+\$365 over 37 automated fills, PF 1.46, green 3 of 3 days); from that moment it ran two scale-out sub-slots at A@2.5R + wide lock-chandelier (−\$421 over 7 fills at 14% win).** I verified both halves against the ledger. So the gate did not degrade — its exit was replaced, and Part 2 · LAB (a) even mislabels the good era's exit as "A@2.5R + chandelier" when it was the adaptive single slot. That mislabel is what hides the causal story.

Exact minimal fix

In Part 2 · LAB (a), replace "the prior era (07-27-07-31, A@2.5R + chandelier, 37 trades) netted +\$365 with a healthy 11 targets / 10 chandeliers / 16 stops mix" with: "the prior era was not the same machine. Until Wed 07-29 19:10 UTC exhaustion ran a SINGLE slot on the regime-3 adaptive chandelier: 37 automated

fills, +\$365, PF 1.46, green on all 3 of its days. The scale-out slate that replaced it (A@2.5R + wide lock-chandelier) is −\$421 over 7 fills at 14% win. The gate did not break; its exit was swapped for a trend-riding one, and a fader does not get 2.5R in a whipsaw." Then add, in place of the bare "too thin to judge": "A 64-entry 250 ms sweep does grade the Rs even though the live trial cannot: the loaded 0.5R Lot A is the worst scalp rung on the ladder (+\$306); the plateau is 0.75–1.0R (+\$510 / +\$450); and Lot B's steadiest config by a distance is a k1.5 chandelier (+\$502, LODO floor +\$330, worst day −\$25). But the sweep's real message is that no exit is robustly green in violent whipsaw, where this gate took 28 of its 62 entries — keeping it benched there is worth more than any R."

7 · The abs_veto promotion verdict now contradicts the correction sitting 200 words above it

WHERE	VERBATIM
Part 2, "Correction (applied 08-01)" — the folded-in patch	"abs_veto is already live two-sided ... So the promotion decision is not 'should we add the long side'; it is already added. ... Reconciling that ... is NOT yet done in this report. "
Part 2, three paragraphs later, the verdict pill	"Verdict: PROMOTE, and promote it TWO-SIDED ... this is the one that has earned a live slot , and clipping it to short-only leaves roughly half its proven money on the floor."

Would a reader be misled? Yes — and this one is self-inflicted by the fold-in, which inserted a correction and left the conclusion it corrects standing. A reader who skims to the verdict pill (which is styled as a pill, i.e. designed to be skimmed to) is told to promote something that shipped on 2026-07-25, and told not to clip it to short-only when nothing has clipped it to short-only. Two further problems in the same patch: the "+\$326 over 47 trades" figure is wrong (\$the callout at the top of this audit — it is +\$59.5 over 59 lots, and it contradicts Part 1's own card table in the same document); and "NOT yet done in this report" is stale, because Q2 §4 did exactly that reconciliation — 82% of the \$609.5 gap is **entry** (26 long signals the router benched, worth +\$397 in shadow), \$107 is window mismatch, \$273.5 is the STOP_UNFILLED path, and stripping fees and the stop bug leaves the exit design roughly **+\$67 in live's favour**.

Exact minimal fix

Change "+\$326 over 47 trades" → **"+\$59.5 over 59 lots (43 signals)"**. Change the verdict pill from "PROMOTE, and promote it TWO-SIDED" to **"ALREADY LIVE TWO-SIDED — RE-VALIDATED COLD"** and rewrite its closing sentence to: "There is nothing to promote. What the cold re-validation earns is confidence

in a gate that has been running both sides since 07-25 — and one config change: the long-side exit cell (see the exit lab)." Replace "is NOT yet done in this report" with the reconciliation itself in one line: "Reconciled: of the \$609.5 gap, 82% is entry — 26 long signals the router had benched, worth +\$397 in shadow — plus \$107 of window mismatch and \$273.5 of STOP_UNFILLED. Net of fees and that stop bug the live exit design is about +\$67 ahead. The bench cost the money, not the exit."

8 · Two day-by-day P&L tables in the same document that disagree on every single day

DAY	PART 1 "THE WEEK, DAY BY DAY" (PARIS SESSION)	MOVEMENT 3 CHOP LAB, "THE BAR, FROM THE LIVE LEDGER" (UTC)	Q1 FACT TABLE 1, AUTOMATED ONLY (UTC)
Mon 07-27	−226.5 (40 tr)	−155 (37 tr)	−155.0 auto
Tue 07-28	−54.0 (23)	−22 (25)	−22.0 auto
Wed 07-29	+638.5 (17)	+1,349 (32)	+1,349.0 auto
Thu 07-30	+1,172.0 (91) "best day ever"	+94 (82)	−884.5 auto (26% win) + \$978 by hand
Fri 07-31	−905.0 (35)	−569 (27)	−814.0 auto + \$245 by hand

Would a reader be misled? Yes, and this one needs no Rev2 at all — it is sitting in the published report today. Part 1 builds an entire case study around "the best day the desk has ever booked, +\$1,172"; Movement 3, in the same document, prices the same Thursday at **+\$94**. Both are defensible (different session boundaries) and neither says so. Add Q1's cut and Thursday is the **worst automated day of the week at −\$884.5 on 73 fills**, turned green by nine hand-claims. A reader who notices the two tables loses trust in both; a reader who doesn't will quote +\$1,172 as a desk record.

Exact minimal fix

One footnote under the Movement 3 table: "Session boundary note: this table books on the UTC calendar day; Part 1's day-by-day table books on the Paris session boundary the desk keeps its ledger on. Same trades, different day buckets — the week total is identical, individual days are not." And one clause in Part 1's Thursday case: "(+\$1,172 on the Paris boundary; +\$94 on the UTC day, and −\$884.5 of that is automated — the nine hand-claims are what made it a record.)"

9 · Q0's question — "what do I actually DO on Monday?" — is still unanswered by the report

The published report contains, by my count, thirteen-plus recommendations spread across Part 1.5's scoreboard, Part 2's promotion verdict, Part 2 · LAB's deployable policy table, Part 2.5's three verdicts, Part 2.6's eight-row Saturday table, and Movement 3's three closing asks — with four different time-labels ("Saturday", "next build", "deploy", "re-validate after a trend week") and no owner on most of them. That was exactly the proof-reader's Q0 complaint, and the published report answers it nowhere.

This is the one item with a mitigating fact: `plays.json` and `monday_2026-07-31.html` have been reconciled by hand against Rev2 and are out of my scope, so the operator does have a coherent action list — **the report just doesn't point at it. Minimal fix:** one sentence in the report's opening lead paragraph: "Actions do not live in this document. Everything actionable is in the Monday playbook (`monday_2026-07-31.html`) and `plays.json`, both reconciled against the Rev2 corrections. Read this report as the evidence behind those plays, not as a to-do list."

10 · Smaller ones, all still live

#	ISSUE	WHERE	MINIMAL FIX
10a	The leakage column disagrees with itself across two sections that print the same table. Part 2.5 §3 WEEK row: −\$1,083 . Part 2.6 §3 WEEK row: \$0 (durable days) . Both are true in their own frame (−\$456/−\$398/−\$229 early week, \$0 on 30/31) and neither says so.	Part 2.5 §3 / Part 2.6 §3	Print both as "−\$1,083 week (−\$456 / −\$398 / −\$229 Mon-Wed, \$0 on both durable days)" in both tables, and add: "Leakage is the only realised-money column here — saved and missed are modelled counterfactuals on trades that never ran. Do not net them."
10b	" The meter is validated " stated flatly on n=5 days , with "zero false STAY-OUTs" as supporting evidence. The CAUTION refinement is correctly flagged n=2; the validation claim itself is not flagged at all.	Part 2.6 §5	"The meter is validated." → "The meter lines up perfectly on the only five days it has seen. That is a promising start on n=5, not a validation — re-score it weekly and do not let it size anything until it has a month."
10c	" +\$728 over 10d " for the grind 0.35 arm floor is a \$ figure with no baseline — over what alternative? — and per item 2 the floor has not executed since 07-29.	Part 2.6 §4	State the counterfactual or drop the number; either way append "— and see the dead-floors correction: this threshold has not run since 07-29 16:24, so nothing in this week's tape tests it."
10d	Part 1.5 §2 compares "LIVE (n=11, −\$284)" against "RECOMMENDED (n=111, +\$2,830)" as a before/after. The recommended config trades ten times as often ; the two	Part 1.5 §2	Add a \$/trade column (−\$26 vs +\$25) and one clause: "note the n — this is not the same 11 trades made profitable, it is a different, ten-times-larger book."

	rows are not the same experiment and the +\$2,830 is not "the fix recovered \$3,114".	
10e	rgv_short is armed live (gate_switches.env: rgv_short=on, verified) while Part 1.5 \$5 verdicts it SHADOW and Part 1 says the router "correctly kept it in the dugout". The report never notices that its own verdict and the live switch file disagree.	Part 1.5 \$5 Append: "Live-state check (08-01): rgv_short is currently <i>on</i> in gate_switches.env. This verdict and the desk disagree — bench it or change the verdict, but do not leave both standing."
10f	Part 2's promotion battery still reports shadow dollars (1 lot) directly against live dollars (2 lots under scale-out) with no per-signal denominator — the comparison that produced the \$326-vs-\$776 confusion in the first place.	Part 2 One standing note under the shadow board: "The shadow runs 1 lot; the live scale-out runs 2. No raw shadow-vs-live dollar comparison in this report is valid without dividing the live number by signals, not lots."

Does the report lead with what survived, and is the headline honest?

SCOPE PROMISE (DOCS/ FRIDAY_V7_REPORT_SCOPE.MD , LINE 7)	PUBLISHED REPORT	VERDICT
"LEAD with what SURVIVED the skeptic, never the biggest number"	Part 1's first content block is the −\$598 callout. Good. But the document's own opening lead paragraph (before Part 1) says nothing about it, and Part 2.6's lead paragraph and §8 verdict both still headline " net green +\$625 " with no cross-reference.	PARTIAL
Automated-only −\$598 stated up front, not buried	Stated up front in Part 1 — and then it appears exactly twice in the whole 326 KB document (the callout paragraph and its own table row). Every other section still headlines + \$625.	PARTIAL — one box, no propagation
Part 1 leads with the survivor	It leads with the correction, then re-crowns the survivor ninety words later and again in the card table.	CONTRADICTORY

The one-line fix that fixes most of this: propagate the −\$598 to the three places that still headline +\$625 — the document's opening lead paragraph, Part 2.6's lead paragraph, and Part 2.6 §8's verdict — as "+\$625 as booked, −\$598 with your hands off (see Part 1)". That is three sentence edits and it does more for the report's honesty than any of the analytical patches.

Placeholders, thin sections, PDF sync

CHECK	RESULT
TODO / TBD / FIXME / placeholder / lorem / "coming soon"	☐ none in the published HTML
Obviously-thin sections	☐ none. The genuinely thin material (2-signal exhaustion trial, 4-trade grind shadow watches, n=2 CAUTION days, n=4 chandelier trend leg) is all explicitly flagged as thin in-place — that discipline is intact throughout.

THE VERDICT

The report needs a Rev3 re-render before Garrath reads it. Not because it is dishonest in intent — the Part 1 callout is a genuinely brave opener and the thin-n discipline holds up all the way through Movement 3 — but because the fold-in landed 2 of 5 sections and left the document arguing with itself in seven places, telling him to leave two exit cells alone that two independent sweeps say to change, reasoning five times from filters that have not executed since Wednesday, quoting an n of 5 for a cell that is n=24, publishing a principled refusal of an archive the lab had already consumed, and omitting the single best-evidenced number of the week (+\$1,615/wk from one router rule). One of the two patches that did land carries a \$266.5 error that contradicts a table in the same document.

The good news is the fix is small. Every correction above already exists, computed and robustness-tested, in rev2_q1..q4.html. Nothing needs re-analysis. The work is: 3 sentence edits to propagate -\$598; 6 cell/sentence edits for exhaustion; 2 for BIG-TREND; 2 for the abs_veto exit; 1 wrong number and 1 verdict pill in Part 2; 2 table rows and 1 paragraph for the router value; 1 correction callout for the dead floors; and 6 footnotes. Roughly **two dozen targeted edits** to the three section files, then re-render HTML and PDF together. Alternatively — and this is cleaner given the volume — **append rev2_q1 through rev2_q4 as a Part 3 "Rev2 corrections" and put a pointer at the top of each superseded section.** That preserves the audit trail, which is worth something in a report whose selling point is showing its own graves.

Audit scope: src/gazbot7/web_static/weekly_2026-07-31.html and .pdf read end to end; diffed against reports/friday_v7/sections/{part1_live,part2_shadow,part2_6}.bak; cross-read against rev2_q0..q4.html and docs/FRIDAY_V7_REPORT_SCOPE.md. All ledger figures re-verified against data/gazbot7.db, data/gate_switches.env and data/exit_overrides.json — nothing recalled. Read-only: no code changed, no service touched. reports/friday_v7/

plays.json and src/gazbot7/web_static/monday_2026-07-31.html were out of scope and were neither read for reconciliation nor edited.

M1 The biggest runs on MNQ this week — did we show up?

Forget what the desk did this week — start from the raw tape. MNQ printed **73 runs of 1.5×ATR or bigger** in the last seven days (a 15-minute move of at least 90 points, against a typical 15-min range of 60). That is the whole tape, not a top-ten. We **caught 14**, we were positioned against **6**, and we **sat out 53**. The runs we aligned with banked **+\$944** of honest money — but that is only 20% of their \$4689 hindsight ceiling, and the 53 we sat out left **\$14519** on the table. The bleed is not the runs; it is the fading we do around them.

THE FULL CENSUS

Every qualifying run, in order. **Move** is the swing in points; **\$ 1lot** is that move on one MNQ lot (\$2/pt) — the hindsight ceiling. **real \$** is what a gate actually banked near it. **Flow** is net aggressor volume in the 60s before; **amp** is pre-run amplitude; **book** is the far-side L2 depth share (below 0.50 = the side price ran toward was thin). Green = we caught it, red = we fought it.

TIME (UTC)	DIR	MOVE	\$ 1LOT	US	GATE	REAL \$	FLOW	AMP	BOOK	CLUSTER
07-27 13:28	DN	-155	310	caught	exhaustion_short	-30	-230	0.18	—	OPEN/NEWS
07-27 13:51	DN	-205	410	caught	exhaustion_short	+20	+1244	0.26	—	OPEN/NEWS
07-27 14:19	DN	-224	447	caught	exhaustion_short	+20	-100	0.26	—	OPEN/NEWS
07-27 14:43	UP	+198	396	FOUGHT	exhaustion_short	+9	+275	0.43	—	OPEN/NEWS
07-27 15:03	DN	-117	234	caught	exhaustion_short	+8	+129	0.18	—	VACUUM

07-27 15:22	DN	-131	262	caught	exhaustio n_short	-54	-414	0.37	—	FLOW-LED
07-27 15:53	UP	+115	230	sat out	—		-225	0.15	—	VACUUM
07-27 16:16	UP	+134	268	sat out	—		+592	0.30	—	FLOW-LED
07-27 16:48	DN	-139	278	sat out	—		+62	0.14	—	VACUUM
07-27 18:40	UP	+147	294	sat out	—		-42	0.13	—	UNCLASS
07-28 01:48	DN	-102	204	sat out	—		-137	0.14	0.50	FLOW-LED
07-28 03:09	UP	+99	198	caught	capitulatio n_long	+106	-252	0.09	0.48	VACUUM
07-28 08:14	UP	+113	226	sat out	—		+29	0.08	0.49	UNCLASS
07-28 11:43	UP	+95	190	sat out	—		-8	0.18	0.47	UNCLASS
07-28 13:31	DN	-235	470	caught	abs_veto_s hort	+400	+199	0.34	0.51	OPEN/ NEWS
07-28 14:05	DN	-149	298	sat out	—		-23	0.28	0.49	OPEN/NEWS
07-28 14:34	UP	+174	348	FOUGHT	exhaustion_ short	+18	-346	0.22	0.49	OPEN/NEWS
07-28 15:11	UP	+183	366	FOUGHT	exhaustion_ short	-72	+171	0.38	0.47	FLOW-LED
07-28 15:31	UP	+94	188	caught	abs_veto_l ong	+3	+227	0.22	0.48	FLOW-LED
07-28 22:10	DN	-96	192	FOUGHT	abs_veto_lo ng	-27	-192	0.07	0.50	FLOW-LED
07-28 22:43	DN	-111	222	caught	abs_veto_s hort	+59	-112	0.13	0.52	FLOW-LED
07-28 23:09	UP	+124	249	sat out	—		-101	0.23	0.57	VACUUM
07-28 23:27	UP	+205	410	sat out	—		-4	0.08	0.47	UNCLASS
07-28 23:54	UP	+101	202	sat out	—		-71	0.22	0.50	VACUUM
07-29 00:18	DN	-113	226	sat out	—		+87	0.18	0.52	VACUUM
07-29 01:01	DN	-117	234	sat out	—		+233	0.13	0.53	VACUUM
07-29 01:38	DN	-133	266	sat out	—		-4	0.15	0.51	UNCLASS
07-29 03:46	DN	-91	182	sat out	—		+10	0.14	0.51	UNCLASS
07-29 04:49	UP	+122	244	sat out	—		-48	0.15	0.50	UNCLASS
07-29 06:22	UP	+138	275	sat out	—		+68	0.14	0.50	FLOW-LED
07-29 12:02	DN	-111	222	sat out	—		+12	0.08	0.56	UNCLASS
07-29 13:20	UP	+127	254	sat out	—		+336	0.09	0.49	OPEN/NEWS
07-29 13:42	DN	-200	399	caught	exhaustio n_short	-101	-10	0.29	0.49	OPEN/ NEWS

07-29 14:11	DN	-187	374	caught	exhaustio n_short	+146	+903	0.34	0.51	OPEN/ NEWS
07-29 14:33	DN	-110	220	sat out	—		+218	0.34	0.49	OPEN/NEWS
07-29 14:48	UP	+105	210	sat out	—		-572	0.24	0.51	OPEN/NEWS
07-29 15:51	DN	-117	234	sat out	—		-90	0.13	0.52	FLOW-LED
07-29 16:15	UP	+148	297	sat out	—		-515	0.21	0.49	VACUUM
07-29 16:42	UP	+94	188	sat out	—		-92	0.18	0.46	VACUUM
07-29 16:58	UP	+141	282	sat out	—		+171	0.15	0.47	FLOW-LED
07-29 17:47	UP	+191	382	sat out	—		+50	0.14	0.51	UNCLASS
07-29 18:07	DN	-134	268	sat out	—		-119	0.39	0.50	FLOW-LED
07-29 18:34	UP	+298	596	sat out	—		+137	0.42	0.49	FLOW-LED
07-29 19:06	DN	-247	494	caught	exhaustio n_short_A	-166	+546	0.54	0.51	VACUUM
07-29 19:32	DN	-248	496	sat out	—		-101	0.27	0.51	FLOW-LED
07-29 19:49	DN	-276	552	sat out	—		-514	0.24	0.51	FLOW-LED
07-29 20:13	DN	-111	222	sat out	—		+524	0.33	0.53	VACUUM
07-29 20:56	UP	+182	364	sat out	—		+66	0.18	0.51	FLOW-LED
07-30 00:02	UP	+91	182	sat out	—		+3	0.22	0.51	UNCLASS
07-30 00:29	DN	-98	197	sat out	—		-275	0.13	0.49	FLOW-LED
07-30 00:52	UP	+115	230	sat out	—		-140	0.15	0.55	VACUUM
07-30 01:18	UP	+105	210	sat out	—		-23	0.11	0.48	UNCLASS
07-30 07:06	DN	-103	206	sat out	—		+4	0.14	0.50	UNCLASS
07-30 08:03	UP	+104	208	sat out	—		-1	0.30	0.51	VOL- EXPANSION
07-30 13:40	UP	+244	488	caught	grind_long _A	+478	+488	0.26	0.49	OPEN/ NEWS
07-30 14:18	UP	+145	290	sat out	—		-113	0.21	0.50	OPEN/NEWS
07-30 14:41	DN	-196	392	sat out	—		+545	0.12	0.53	OPEN/NEWS
07-30 15:30	UP	+101	202	sat out	—		+325	0.21	0.50	FLOW-LED
07-30 16:00	UP	+114	229	sat out	—		+79	0.26	0.54	FLOW-LED
07-30 18:58	UP	+116	232	sat out	—		-62	0.07	0.52	VACUUM
07-30 20:00	UP	+115	230	sat out	—		-857	0.21	0.49	VACUUM

07-31 06:28	DN	-126	253	FOUGHT	capitulation_long_A	-38	-8	0.04	0.45	UNCLASS
07-31 13:19	UP	+153	306	sat out	—		+105	0.10	0.52	OPEN/NEWS
07-31 13:34	DN	-366	732	FOUGHT	abs_veto_long_A	-138	-24	0.57	0.47	OPEN/NEWS
07-31 13:58	DN	-251	502	sat out	—		-798	0.43	0.52	OPEN/NEWS
07-31 14:16	UP	+193	386	sat out	—		-2791	0.39	0.50	OPEN/NEWS
07-31 14:42	UP	+105	210	sat out	—		-1947	0.30	0.50	OPEN/NEWS
07-31 14:57	DN	-107	214	sat out	—		-189	0.32	0.49	OPEN/NEWS
07-31 15:14	DN	-161	322	sat out	—		-490	0.26	0.50	FLOW-LED
07-31 15:48	UP	+110	221	sat out	—		+313	0.21	0.53	FLOW-LED
07-31 16:17	UP	+101	202	sat out	—		+301	0.19	0.52	FLOW-LED
07-31 17:22	UP	+96	192	caught	abs_veto_long_A	+55	-61	0.10	0.49	VACUUM
07-31 19:50	DN	-144	288	sat out	—		+785	0.13	0.50	VACUUM

What the runs were — by cause

Each run gets a cause fingerprint from its pre-run tape. This is the map for where to hunt a new gate.

CLUSTER	RUNS	WHAT IT IS
OPEN/NEWS	21	the 13:00–15:00 UTC cash-open / data window
FLOW-LED	21	aggressors drove it (a continuation)
VACUUM	17	moved AGAINST the tape (a snap-back / stop-run)
UNCLASS	13	no clear tape tell (quiet ignition)
VOL-EXPANSION	1	amplitude broke out of a coil

HONEST MONEY — CEILING VS WHAT WE BANKED

RUNS	HINDSIGHT CEILING	WHAT WE REALLY MADE
Caught (aligned)	14	\$4689 +\$944 (20%)
Fought (against)	6	\$2286 \$-248

The runs made money and the fading gave it back; we convert a fraction of the ceiling and ignore the rest. The next two movements ask whether we can do better — first with the gates we own, then with new ones.

M2 Could our own gates have caught the runs we sat out?

Garrath, this one nearly fooled me — so let me walk you through it honestly, because the honesty is the finding. I took the **53 runs we sat out** this week (the frozen census — **\$14,519** of hindsight money left on the table) and fired all six of our live gates at them mechanically, ungated (no ER floor, no ATR floor, no absorption-confirm — the most generous version), in each run's own direction, anywhere in the **10 minutes before** it lit up. Then I priced every exit honestly on this week's own trade ticks. The headline looked good: **21 fires, +\$582, and it survives stripping the best three trades** — nothing like last week's clean fluke. One gate, **rgv_short, went four-for-four for +\$478**. But it is a mirage, and I can prove it in one number: **fed only the 53 runs we already know happened, rgv_short makes +\$478; fed the whole week's tape, the exact same gate makes +\$1.80 a trade**. The lab was quietly handing each fader only its winners. This is an **honest null** — just a subtler, better-disguised one than last week's.

WHAT SURVIVED THE SKEPTIC — THE ONE NUMBER THAT ENDS THE ARGUMENT

Nothing survived as a tradeable edge, and I will not dress the +\$582 up as one. The durable fact is a **base-rate**: our turnback fader **rgv_short throws 151 signals in a week**, and only **6 of them (4%) are followed by a real run** in its direction — while **13% go straight against it**. The idle-gate lab only ever showed it the ~4 signals that preceded a sat-out down-run, so of

course they printed 2R. That is **survivorship**, not skill: cherry-pick any early-entry fade's winners after the fact and it looks brilliant. Feed it the whole tape and it is a coin-flip that bleeds on costs. On top of that, the **four biggest sat-out runs this week (+298, -276, -251, -248 — \$2,146 of ceiling between them) drew ZERO fire from any gate**. The money on the table belongs to triggers we do not own, which is Movement 3's job. Arming the gates we have is still a dead end — here is the proof, in full.

The scoreboard — all six gates, fired ungated at the 53 sat-out runs

In-direction only: the three long gates can only catch the **32 up-runs**, the three short gates only the **21 down-runs** — so "eligible" is how many sat-out runs each gate could physically have caught. One position per fire, first trigger in the 10-minute pre-run window wins, exits repriced tick-by-tick on this week's tape (grind and thrust ride the chandelier plus the native 1-ATR stop; the faders take a 2R scalp or the 1-ATR stop). Honest money, the same week the runs happened. Green = it made money over its handful of fires, red = it lost.

GATE	ELIGIBLE	FIRES	HONEST \$	\$/FIRE	WIN%	EXITS
rgv_short	21	4	+\$478.0	+\$119.5	100%	4 target
rgv_long	32	6	+\$278.0	+\$46.3	67%	4 target · 2 stop
grind_long	32	7	+\$52.0	+\$7.4	43%	3 chandelier · 4 stop
thrust_short	21	4	-\$225.5	-\$56.4	0%	4 stop
capitulation_long	32	0	\$0	—	—	never fired
exhaustion_short	21	0	\$0	—	—	never fired
ALL SIX	53	21	+\$582.5	+\$27.7	52%	20 distinct runs

All of the "profit" is in the two turnback faders (rgv_short +\$478, rgv_long + \$278 = **+\$756 between them**). The two momentum gates net **-\$174** (thrust_short stopped out all four times), and the two footprint gates **never fired at all**. On its face rgv_short looks like the best thing in this whole report. It is the most dangerous line in it — and the next table is why.

The survivorship killer — fed the winners vs fed the whole week

This is the test that matters. The lab only shows a gate the 53 ignitions we already know turned into runs — that is picking the answers before the exam. So I ran each gate a second way: fired it **ungated across the entire week's tape**, every time it triggered, repriced on the same ticks. The gap between the two columns is the size of the lie.

GATE	AT THE KNOWN RUNS (FIRES · NET · \$/FIRE · WIN)	FED THE WHOLE WEEK (TRADES · NET · \$/ TRADE · WIN)	VERDICT
rgv_short	4 · +\$478 · +\$119.5 · 100%	134 · +\$243 · +\$1.8 · 34%	SURVIVORSHIP
rgv_long	6 · +\$278 · +\$46.3 · 67%	120 · +\$45 · +\$0.4 · 36%	SURVIVORSHIP
thrust_short	4 · −\$225.5 · −\$56.4 · 0%	113 · −\$552 · −\$4.9 · 35%	REAL LOSER
grind_long	7 · +\$52 · +\$7.4 · 43%	239 · +\$747 · +\$3.1 · 38%	RIDES TRENDS, NOT IGNITIONS

Read the rgv_short row slowly: the +\$478 that made it look like a star is **4 cherry-picked winners out of 134 fires** — the other 130 fires that week netted about −\$235, and the gate's true worth is **+\$1.80 a trade, a 34% win-rate, a hair above breakeven**. rgv_long is the same story at \$0.40 a trade. thrust_short is honestly consistent: it loses at the runs (−\$225) and loses all week (−\$552). grind_long is the one green full-week number, but look at where it earns — only +\$52 at the ignitions themselves (a coin-flip), because grind does not catch a run starting, it rides a trend already running; its +\$747 is a trend-week story that belongs in the rehab section, not here.

The base-rate — how rare a "run" is under a fader's signal

Here is the survivorship point stated as plainly as I can. For each gate I counted **every ungated signal it fired all week**, then asked how many

were actually followed by a 90-point-plus move in its own direction within 15 minutes — a real run — versus how many went straight against it.

GATE	SIGNALS / WEEK	FOLLOWED BY A		WENT ≥40PT ADVERSE
		REAL RUN (≥90PT)	≥150PT RUN	
rgv_short	151	6 (4%)	2	19 (13%)
rgv_long	130	9 (7%)	1	25 (19%)
thrust_short	118	6 (5%)	3	25 (21%)

Four to seven percent. That is how often one of our triggers is standing in front of a real run when it fires — and it goes hard against the gate two to five times more often than it precedes a run. The lab handed each fader only that thin 4–7% slice and called it a gate. It is not. It is a gate that fires a hundred-and-fifty times to be right six.

The other two skeptic cuts — strip-the-best and the R-target sweep

Even before the survivorship test, two standard robustness checks flag the +\$582 as fragile. First, strip out the biggest winners: the edge halves at three and **flips negative at five**. Second — the operator's rule that no R-target is a truth — I re-priced all ten fader fires at every scalp target from 0.5R to 3.0R. A real edge shows a plateau (a best R with a roll-off either side). This just climbs forever.

STRIP TEST (ALL 21 FIRES)			R-SWEEP (10 FADER FIRES)		
TRADES	NET \$			NET \$	WIN%
Keep all	21	+\$582.5	0.5R	+\$175	90%
Strip best 3	18	+\$183.0	1.0R	+\$414	90%
Strip best 5	16	−\$43.0	1.5R	+\$540	80%
			2.0R	+\$756	80%
			2.5R	+\$973	80%
			3.0R	+\$1,189	80%

The R-sweep never rolls over — wider is always better, right out to 3R. That is not "2R is the wrong target," it is the survivorship effect showing through the numbers: because every fire we measured was in front of a run that

kept going, any bigger target banks more. It is the runs being real, not the fade being good. With n=10 and no plateau, there is no R to trust here.

The control that proves it is a seeing problem, not an exit problem

To be fair to the exits, I ran a blind control: ignore every gate, just enter **in the run's direction at the moment of ignition** and ride the chandelier. That is cheating — it uses hindsight to know the direction — but it isolates the question "if we simply showed up pointed the right way, does the exit keep the money?"

BLIND IN-DIRECTION ENTRY AT IGNITION (ALL 53 RUNS)					
	RUNS	NET \$	\$/RUN	WIN%	
	Keep all	53	+\$4,076	+\$76.9	75%
	Strip best 3	50	+\$3,405	+\$68.1	—
	Strip best 5	48	+\$2,998	+\$62.5	—

Show up pointed the right way and the runs pay **+\$4,076 at 75% win, and it barely dents when you strip the best five** — because the money is real and directional, spread across dozens of runs, not two flukes. So the exits are fine and the money is genuinely there. The entire problem is **knowing, in real time and without hindsight, which way and when** — and that is exactly what our gates fail at: they only see 20 of the 53, and the versions that do fire early (the faders) fire everywhere and bleed. The exit is not the bottleneck. Vision is.

Where the money actually was — the 12 biggest runs we sat out

The \$14,519 is concentrated in the monsters, so the only question that counts is whether arming our gates caught the big ones. Here are the twelve largest sat-out runs (worth \$4,700 of ceiling) and exactly what our gates did on each.

RUN (UTC)	DIR	MOVE	\$ CEILING	WHAT OUR GATES DID	HONEST \$
07-29 18:34	UP	+298	\$596	no gate fired	\$0
07-29 19:49	DN	−276	\$552	no gate fired	\$0
07-31 13:58	DN	−251	\$502	no gate fired	\$0
07-29 19:32	DN	−248	\$496	no gate fired	\$0
07-28 23:27	UP	+205	\$410	grind_long rode it — chandelier	+\$120.5
07-30 14:41	DN	−196	\$392	rgv_short faded the top — 2R target	+\$110.5
07-31 14:16	UP	+193	\$386	no gate fired	\$0
07-29 17:47	UP	+191	\$382	no gate fired	\$0
07-29 20:56	UP	+182	\$364	no gate fired	\$0
07-31 15:14	DN	−161	\$322	rgv_short +\$149 caught it, thrust_short −\$71 fought it	+\$78.0
07-31 13:19	UP	+153	\$306	grind_long rode it — chandelier	+\$39.5
07-29 16:15	UP	+148	\$297	rgv_long faded it — stopped	−\$52.5

Of the twelve biggest runs we missed, our gates fired on five and made money on four (+\$348 combined, one fader stopped for −\$53) — but the **four absolute monsters, worth \$2,146 of ceiling, all drew a complete blank**. Not a fought trade, not a stopped trade — nothing. When the move is genuinely huge, our triggers are asleep. The gates catch the mid-sized runs when a fade happens to line up, and no-show the feasts. That is the null in one sentence.

The segment split — is there a regime where arming pays?

The operator's rule is never to blanket one number across all the tape, so I split the 21 fires by time-of-day and by cause-cluster. If there is a corner where arming these gates genuinely works, it should light up here. It does not — the fires are thin and scattered in every bucket, and the "green" is the same survivorship winners re-sorted.

SEGMENT	RUNS	FIRES	NET \$	WIN%	READ
US session (13:30-20:00 UTC)	—	13	+\$322.0	46%	below-even win-rate; two fader wins carry it
Overnight / pre-open	—	8	+\$260.5	62%	thinner tape, same fader survivorship
OPEN/NEWS	11	4	+\$200.5	75%	4 fires on 11 runs — still mostly no-show
FLOW-LED	16	6	+\$135.5	50%	the continuation runs — gates half-blind
VACUUM	13	8	+\$102.0	38%	most fires, worst win-rate — fades into snapbacks
UNCLASS	12	3	+\$144.5	67%	3 fires on 12 runs
VOL-EXPANSION	1	0	\$0	—	never fired

No segment rescues it. The best-looking bucket, OPEN/NEWS at 75%, is four fires on eleven runs — the other seven news-runs drew nothing. VACUUM (snapback / stop-run) draws the most fires and the worst win-rate, exactly as you'd expect: our faders reach for the retrace and the vacuum keeps going. There is no regime here that says "arm these on this kind of day." The pattern in every bucket is the same — a few survivorship winners and a lot of no-shows.

Why each gate misses — the mechanism, gate by gate

None of this is bad luck; each miss is baked into what the gate is built to look for. This is the useful part — it tells us the trigger we are missing.

GATE	FIRES
------	-------

WHY IT CAN'T RELIABLY CATCH THESE RUNS

rgv_short / rgv_long

4 / 6

The turnback faders need price already stretched 2+ ATR past VWAP and rolling back. When a run happens to launch off exactly that setup, the fader gets in early at the top/bottom and rides 2R — which is why they "win" here. But that setup fires **130-150 times a week** and only **4-7% of those precede a run**; the rest chop or reverse. You cannot harvest the winners without taking all 150 fires, and the full basket is breakeven. This is the survivorship trap in one line.

thrust_short

4 / 21

Wants a five-bar (five-minute) momentum burst with a volume surge. These ignitions fire in one or two bars — by the time five bars confirm, the move is spent, so it arms late and buys the exhaustion. All four fires stopped out (−\$225), and it loses all week too (−\$552). A structurally late trigger on fast moves.

grind_long

7 / 32

Needs an established up-trend — VWAP sloping up with price riding above it. It arms once a trend is already running, so at the ignition itself it is a coin-flip (+\$52, 43%). It earns its full-week +\$747 by riding trends that are already underway, not by catching a run start — a different job from this section's.

capitulation_long

0 / 32

Fades a one-sided aggressor climax — one side running 3× its baseline AND 70%+ of flow, price already flushing that way. The up-runs we sat out do not start from a sell-flush blow-off; they ignite from quiet or flow-led setups. The trigger never trips. (Also L2-blind before 07-27 21:33 — a data gap on top of the logic gap.)

exhaustion_short

0 / 21

Needs heavy buying absorbed into a resting ask wall without price moving. Our down-runs are flow-led drives where the aggression does move price — so the "absorbed, price stuck" precondition fails and it never fires. Same early-week L2 blindness caveat applies.

VERDICT

HONEST NULL

DON'T ARM

Firing all six live gates at the money we left on the table produces a headline +\$582 that **dissolves the moment you stop cheating**: the two faders that carry it are worth +\$1.80 and +\$0.40 a trade fed the real tape, their signals precede a run only 4-7% of the time, the strip-5 test flips the whole thing negative, the R-sweep has no plateau, and the four biggest runs of the week drew zero fire. The blind control proves the money and the exits are both fine — the gap is purely that we cannot see these runs coming in real time. Do not arm anything off this section. The sat-

out money belongs to triggers we don't own yet, and hunting those is Movement 3.

Be honest about the caveats: this is one week, in-sample, with tiny counts at the runs (21 fires; the faders just 10). The two footprint gates were L2-blind on the 07-27 daytime runs, so their zero is partly a data gap, not only a logic one. And I'll flag the one genuinely interesting thread, because it would be dishonest to bury it: the turnback-fade mechanism is the only thing we own that shows up early on these runs at all — if we ever find a way to fire it only in that 4-7% slice (a regime read that separates "stretched-and-turning into a real run" from "stretched-and-turning into more chop"), it becomes the entry we're missing. That separation is the stone left unturned, and it is a Movement-3 hunt, not a knob on the gate we have.

M3 The greenfield lab — inventing brand-new entries for the runs we sat out, and a chop-day scalp lab

Garrath — this is the flagship, and unlike last week it does **not** end in a clean sweep of graves. Something survived. I threw away every gate the desk owns, went back to a blank sheet, and asked one question nine different ways: can a brand-new mechanical entry show up for the **53 runs we sat out** this week (— \$14,519 of hindsight money on the table)? One candidate came through the whole robustness battery alive: **NIPC**, a news-window impulse-pullback entry, **+\$2,676 over 12 days, 171 trades, 43.3% win, \$15.6 a trade**, 44 of 44 parameter variants positive, both sides independently green, and it survives \$12/round-trip and a 3-point stop-slip stress. **Verdict: SHADOW. Never live, not yet** — one regime-fortnight is not a promotion. Everything else in the greenfield lane is buried, including the entire flow-footprint family. And I ran the operator's **chop-day scalp lab** on top — the honest answer there is a **NULL**, but it

comes with the most valuable single number in the section: on MNQ the whole visible 10-deep order book is **86 contracts** against **2,664 contracts** a **minute** of flow, so "absorption" is not a thing you can measure on this instrument. That fact explains four separate graves at once.

LEAD WITH WHAT SURVIVED — THREE THINGS, IN MONEY ORDER

1. `news\_impulse\_pullback` (NIPC) — a real, robust, brand-new entry. **+\$2,676 · n=171 · 12 days · 43.3% win · \$15.6/trade · 9 of 12 days green · worst day −\$176.** It waits for a 2-minute impulse in the 13:00–15:00 UTC window, then refuses to chase it — it demands a 35–75% retrace and enters only on the half-back reclaim. It clears every kill-test that buried last week's candidates. **SHADOW-first, explicitly not live.**

2. The size-threshold answer — the question this whole exercise exists to answer. For tape flow there is no size at which a footprint becomes tradeable, and the relationship is **inverted**: across the 53 sat-out runs, $\text{corr}(|\text{move}|, \text{aligned pre-run flow}) = -0.209$. The bigger the run, the less aligned flow preceded it. For book absorption there **is** a real threshold and it is physical — and MNQ sits roughly **6× below it**. For run size given a working entry, the threshold is about **195 points**: NIPC went 8-for-8 green on the top-quartile runs and banked 68% of its run P&L from 38% of the runs.

3. On chop days, the OFF switch beats every gate I could build. The chop-day donation is **−\$249/day** and it is recoverable with certainty by benching. The best chop-scalp I could construct out of 1,152 policies makes **+\$280/chop-day at t=1.50** — and the leave-one-day-out test says it will not repeat. Four vetoes it did validate are worth more to the router than the gate is.

★ **HEADLINE — THE CHOP-DAY SCALP LAB: can we stop DONATING on chop days?**

This was your explicit ask, so it goes first. Three days this week bled with the gates we have — **Mon 07-27, Tue 07-28, Fri 07-31**. The brief was blunt: build something that turns those days from a donation into +\$200-300. I built it, swept it 1,152 ways, and I have to give you an **honest NULL**. Here is the whole argument, money first.

The bar, from the live ledger — not a vibe. Real desk P&L from `gazbot7.db.trades`:

DAY	REGIME	LIVE TRADES	LIVE NET	RTH RANGE	RANGE EFFICIENCY	ATR14
Mon 07-27	CHOP (violent)	37	-\$155	678.50	0.577	27.1
Tue 07-28	CHOP (round-trip)	25	-\$22	476.50	0.077	25.0
Wed 07-29	trend	32	+\$1,349	782.25	0.670	33.7
Thu 07-30	trend	82	+\$94	484.75	0.706	22.8
Fri 07-31	CHOP (mixed)	27	-\$569	646.00	0.317	27.3
3 chop days	the donation	89	-\$746 = -\$249 per chop day. That is the prize.			

★ **REV3 CORRECTION — 2026-08-01 — SESSION-BOUNDARY NOTE — THIS TABLE AND PART 1'S DAY-BY-DAY TABLE ARE BOTH RIGHT**

This table books on the **UTC calendar day**; Part 1's day-by-day table books on the **Paris session boundary** the desk keeps its ledger on. Same trades, different day buckets — the week total is identical, the individual days are not. That is why Part 1 calls Thursday a +\$1,172 record and this table prices the same Thursday at +\$94. Add the hands-off cut and Thursday is the worst automated day of the week at -\$884.5 on 73 fills, turned green by nine hand-claims.

One honest caveat before anything else: **only 07-28 is a textbook round-trip chop day** (range efficiency 0.077). 07-27 is a -392-point directional day that merely felt like chop because it got there in violent whipsaws, and 07-31 is mixed (0.317). That matters enormously — on two of the three

target days **there was no range to fade**. It is the first reason "fade the turn" struggles.

The cost wall — why "scalp tiny, bank fast" is the wrong instinct. Total friction is **1.75 points per round-trip** (\$3.00 fee at the real \$1.50/side, plus spread and slipped stops, MNQ \$2/point). The fee is a fixed tax, so the smaller the target, the higher the hurdle:

SHAPE	NET WIN	NET LOSS	BREAKEVEN WIN% REQUIRED
Stop 20 pt, target 0.5R	+\$17	-\$43.50	71.9%
Stop 20 pt, target 1.0R	+\$37	-\$43.50	54.0%
Stop 30 pt, target 0.5R	+\$27	-\$63.50	70.2%
Stop 30 pt, target 1.0R	+\$57	-\$63.50	52.7%

The mechanical spec I tried. Candidate turn = every new 300-second local extreme, RTH+late (13:30-21:00 UTC), 30-second dedupe → **1,675 events over 6 days (279/day)**. Direction = fade (short the high, long the low). Entry marketable — a short hits the bid, a long lifts the ask. Target is a resting limit that needs a print through by a tick before it counts. Stop is stop-market, triggers on touch, fills a tick worse, and **wins all ties** inside a tick gap. Deliberately unkind, and it should be.

The per-trade edge — the single most important number in the lab. Stripping the fee out (keeping spread and slippage) to ask whether an edge exists at all:

STOP	TARGET	N	WIN%	GROSS PTS/TRADE	NET \$/TRADE
1.6×ATR	0.75R	385	59.5%	+1.81	+\$0.62
1.6×ATR	0.50R	490	69.4%	+1.63	+\$0.26
1.2×ATR	1.00R	451	52.1%	+0.26	-\$2.49
1.2×ATR	0.50R	621	65.7%	-0.51	-\$4.02
1.2×ATR	1.50R	370	41.4%	-2.25	-\$7.50

Best gross edge anywhere in the study = +1.81 points a trade.

Round-trip friction = 1.75 points. The mean-reversion at a fresh local extreme is real — it is just **0.06 points net**. The cost eats 97% of it. Cross-check by inversion: flip the identical signal to breakout-continuation and it makes -\$5.29/trade against the fade's -\$2.49 — a \$2.80 spread ≈ 1.4 points of genuine reversion tendency, consistent. **To clear +\$250/chop-day**

at ~30 trades a day you need $+\$8.30/\text{trade} = 5.9$ points gross. That is $3.3\times$ the best edge measured anywhere in this lab.

The best thing 1,152 policies could produce — and why I still won't bank it. Grid: efficiency-ratio veto $\times$ ATR ceiling $\times$ range floor $\times$ stop multiple $\times$ R $\times$ hold, all sequential (one position at a time):

GRID STATISTIC (1,152 SEQUENTIAL POLICIES)	VALUE
Mean chop-day P&L across all configs	-\$36
Standard deviation	\$131
Fraction with chop-day P&L > 0	40.9%
Fraction with \$/trade > 0	18.8%
Best config, chop-day P&L	+\$288
Expected best-of-1,152 under a pure-noise null	$\approx +\$455$

Read the last two rows together: **the best config out of 1,152 tries is below what random noise would be expected to produce from 1,152 tries.** Marginal effects agree — averaged across the grid, no R and no stop multiple is meaningfully positive (best marginal: stop $1.2\times\text{ATR}$ at $+\$37/\text{chop-day}$; every R marginal is negative).

The best-of-grid config for the record — $\text{ATR} \leq 28$, 30-min range $\geq 3.5\times\text{ATR}$, time $\geq 17:00$ UTC, stop $1.2\times\text{ATR}$, $R = 1.0$, hold 600s:

N	WIN%	\$/TRADE	T-STAT	CHOP NET	CHOP/DAY	TREND NET	CHOP DAYS GREEN
197	58.9%	+\$4.42	1.50	+\$840	+\$280	+\$30	3 / 3

Strip-the-best-trades on the chop days is genuinely healthy — strip 1 $\rightarrow +\$260/\text{day}$, strip 3 $\rightarrow +\$225$, strip 5 $\rightarrow +\$192$, strip 10 $\rightarrow +\$113$, and the single best chop trade is only $+\$61$, so it is not one lucky print. But 07-28 alone supplies \$454 of the \$840, and **a t of 1.50 selected from 1,152 trials is not evidence.**

Two tests killed it outright.

KILL TEST	WHAT HAPPENED
Parallel $\rightarrow$ sequential collapse	Stacking the four validated vetoes on the raw event pool gives $n=557$, $+\$5.98/\text{trade}$, $+\\$1,145/\text{chop-day}$, $t=2.71$. But those events overlap — with a 30s dedupe and a 900s hold you are notionally holding 30 positions and counting one move many times. Re-run honestly, one position at a time: $n=101$, $+\\$3.57/\text{trade}$, $+\\$135/\text{chop-day}$, $t=0.66$. A $3\times$ collapse. Standing rule for the lab: every fader study must be scored sequentially.

Leave-one-day-out learnability	76 features (price, stretch, flow, footprint, absorption, 10-deep book), 6 folds. For each held-out day, learn the single best feature+threshold on the other 5 days, apply it blind. In-sample mean +\$7.17/trade → out-of-sample −\$5.84/trade (n=318). An L2-regularised logistic on all 76 features, same protocol: OOS − \$8.05/trade (n=530) , correlations −0.011 / −0.093 / +0.059 / −0.087 / −0.004 / +0.027. Nothing generalises across a single day boundary.
--------------------------------	---

DAY HELD OUT	RULE THE OTHER 5 DAYS TAUGHT	IN-SAMPLE \$/TR	OOS \$/TR
07-24	slp30_a ≤ 0.029	+7.04	−12.65
07-27	oppz3 ≤ 0.876	+7.20	−5.21
07-28	er15 ≤ 0.014	+6.81	−12.44
07-29	opp3 ≤ 0.465	+6.76	+0.03
07-30	opp3 ≤ 0.465	+7.89	−4.19
07-31	oppz10 ≤ 0.942	+7.30	−5.43

Regime segmentation, as required — and it is damning. A real regime effect is stable inside its own segment. This one flips sign on every axis I cut it on:

SEGMENT	N	WIN%	\$/TRADE
ER < 0.02 (ranging)	274	73%	+2.77
ER 0.02-0.04	147	65%	−6.97
ER > 0.04 (trending) — contradicts the ER veto	55	71%	+9.80
ATR < 18	132	70%	−0.22
ATR 18-28	166	73%	+3.38
ATR > 28	178	67%	−1.45
CHOP day × open 13:30-15:30	70	61%	−10.94
TREND day × open 13:30-15:30 — same clock, opposite sign	74	73%	+7.44
CHOP day × mid 15:30-17:30	60	75%	+7.65
TREND day × mid 15:30-17:30 — same clock, opposite sign	53	60%	−10.83

Every grave in the chop lab, by name. Each was scored on the chop days at stop-30 / R-1.0 / 900s with quartile profiles, per-day breakdown and t-stats:

#	GRAVE	BEST T	CAUSE OF DEATH
G1	Static far-side depth share at the extreme	1.7	Non-monotone, and the sign is backwards — a thick opposition wall predicts the fade losing (− \$8.45/trade, 0 of 6 days green).
G2	Far-side depletion / building (15/30/60/120s deltas)	1.6	Flat. Deltas of a 3-lot queue are noise.

G3	Opposition depth z-score vs its own 15-min baseline	1.7	Non-monotone; fails leave-one-day-out.
G4	exhaustion_signal (our live footprint gate) as a standalone chop scalp	—	Negative on all 6 days in every parameterisation (best — $-\$0.13/\text{trade}$, worst $-\$11.40$; chop net $-\$158$ to $-\$604$). The proven lineage does not transfer.
G5	Absorption-vs-cancellation (refill rate, spoof-pull)	1.8	Degenerate. Quartile boundaries land at 0.999 / 1.000 / 1.001 — traded volume dwarfs queue change $\sim 30\times$, so the ratio is pinned at 1. Not resolvable on MNQ.
G6	Absorption dwell at the level	1.5	Non-monotone; 2 of 3 chop days wrong-signed.
G7	Big-print / institutional absorption	2.2	Non-monotone, day-inconsistent. Root cause: median MNQ print is 1 lot; only 0.62% are ≥ 10 lots . There is no institutional footprint to read.
G8	Delta divergence at the retest	0.7	Flat.
G9	Flow-per-point absorption (push flow $\div$ price progress)	1.8	Non-monotone; "heavy flow, no progress" is not distinguishable from the rest.
G10	Retest / double-top structure (fade an established extreme)	2.8 (wrong sign)	Actively negative: $-\$11.30$ to $-\$16.06/\text{trade}$, win 34–40%. The cleanest structural version of "fade the range extreme" is the worst variant tested.
G11	VWAP-flat "is it actually ranging" filter (your explicit ask)	—	Lifts the fade from $-\$2.53$ to + $\\$7.60/\text{trade}$ — but 100% of the profit is on the TREND days (07-29 +\$868, 07-30 + \$1,156) and it loses $-\\$368$ on the chop days . It makes money on the days we already make money. Textbook confirmation of your own blanket-config warning.

Your exit cheat-sheet, tested on this signal. 168 exit configs on the unfiltered signal (deliberately not on a filtered subset, so the exit conclusion is uncontaminated). **Only 3 of 168 are net-positive, and they are one corner.** Marginal means show **no plateau — a single spike at stop = $1.6\times\text{ATR}$:**

R	0.25	0.40	0.50	0.60	0.75	1.00	1.25	1.50
Gross pts	−0.69	−0.41	−0.48	−0.79	−0.85	−1.58	−1.91	−2.16
Net $\$/\text{trade}$	−4.37	−3.81	−3.96	−4.57	−4.71	−6.16	−6.81	−7.32

YOUR PRIOR	PROVEN?
Fader Lot-A 0.5R	DIRECTIONALLY RIGHT — but only with a WIDE stop . 0.5R is the best R marginally, and survives only at $1.6\times\text{ATR}$ (≈ 32 pt) stops. At a tight stop, 0.5R is the single worst configuration in the entire study .
Fader Lot-B 1.5R	FALSIFIED 1.5R is the worst R at every stop (gross -2.16 pts, net $-\$7.32/\text{trade}$). A chop-turn fade has no runner.

"Scalp tiny, bank fast"

FALSIFIED AS STATED Tiny target + tiny stop = maximum fee tax. The shape that works is **wide stop, small target, high hit rate** — a large-risk scalp, not a tiny one.

Hold time

Shorter is better: 300s > 600s > 900s (−\$4.47 / −\$5.22 / −\$5.96 per trade).

CHOP-DAY VERDICT — SO, CAN WE STOP DONATING?

Yes — but with the OFF switch, not with a new gate. The −\$249/chop-day donation is recoverable with certainty by benching. The best speculative alternative is +\$280/chop-day at $t=1.50$ selected from 1,152 trials, which the learnability test says will revert to zero. Take the certain money.

The four vetoes this lab genuinely validated are worth more to the router than the gate is, and all four are your own priors, rediscovered independently: (1) **don't fade a trending leg** (high 60-min ER → −\$14.39/trade, $t=-3.52$); (2) **don't fade above ATR ≈ 30** (−\$10.27/trade, $t=-2.46$) — the exhaustion ATR-gate, rediscovered; (3) **don't fade a compressed range** (30-min range $< 4.2 \times \text{ATR}$ → −\$12.32/trade, $t=-3.09$) — new, and the strongest single veto found: a compressed extreme is a breakout, not a range test; (4) **don't fade the US open** (15:06–17:00 UTC → −\$8.65/trade, $t=-2.07$).

If anything gets armed, it is SHADOW-only and clearly labelled unproven: chop_turn_fade — trigger on a new 300s local extreme; vetoes $\text{ATR}_{14} \leq 28$, 30-min range $\geq 3.5 \times \text{ATR}$, time $\geq 17:00$ UTC; fade direction; marketable entry; stop $1.2 \times \text{ATR}$, target 1.0R, hold cap 600s; one position at a time, 60s cooldown. Observed $n=197$ / 6 days, 58.9% win, +\$4.42/trade, chop +\$280/day, trend +\$30. **Status: INCUBATE.** Needs ~15–20 more captured chop days. Prior expectation given the learnability result: **it reverts to ~zero.**

Step 1 — the Family A / B split: the census labels the coin after it lands

Before inventing anything I split the 73 census runs by what the label is actually made of. Two families fall out, and they are not equals — one is knowable in advance, the other is a description of what already happened.

FAMILY	CLUSTER	RUNS	WHAT DEFINES THE LABEL	KNOWABLE BEFORE THE RUN?
A — flow-labelled	FLOW-LED	21	Heavy 60s aggressor flow ($ \text{flow} > 50$) aligned with the move	No — "aligned" is only known once you know the move
	VACUUM	17	Heavy 60s flow pointing the opposite way — trapped aggressors	No — same coin, other face
B — clock	OPEN/NEWS	21	Run starts inside 13:00-15:00 UTC. Nothing else.	Yes — you can read a clock
unlabelled	UNCLASS + VOL-EXPANSION	14	No footprint at all	—

Here is the moment the case cracked open. FLOW-LED and VACUUM are the same test with the outcome sign flipped. So I measured the test itself against the tape:

TEST ON THE FLOW-LED / VACUUM LABEL	RESULT
Base rate of the trigger, all 99,245 bars of tape	$P(60s\ flow > 50) = 66.5\%$
Aligned vs opposed, among runs with $ \text{flow} > 50$ outside the open window	21 aligned vs 17 opposed — binomial $p = 0.314$
Aligned flow at the 16 FLOW-LED sat-out run starts (z-score)	median -0.089 , only 44% positive
Volume surge at those same starts (z)	0.71 — below the tape median of 0.95
$P(a \geq 90pt\ run\ follows\ heavy\ flow)$ — against flow / with flow / unconditional	4.2% / 3.5% / 8.8%
Same, once the tape is hot ($ATR15 \geq 90$) — the tilt flips to continuation	7.4% against / 13.6% with

Two-thirds of random moments in the week qualify as a "flow event", and given one, the direction is a coin flip that is less likely to produce a run than a randomly chosen moment. **Family A is not a cluster. It is a 66.5%-base-rate filter multiplied by a 50/50 coin, named after the outcome.** The runs did not start on a flow burst; they started on below-median volume. Everything built to Family A's literal definition is being fitted to

noise — and that is exactly what the graves below show, twice over, at two different levels of zoom.

Family B has the opposite property: a clock window is real, ex-ante, and non-negotiable. It is also, not coincidentally, **where the only survivor came from.**

Step 2 — the oracle: are the runs real money? (yes — so the problem is the ENTRY)

Everything here rests on one premise: that these runs are worth money if you could find their start. The census prices that directly with a hindsight 1-lot ceiling, and this week's honest scoreboard is:

CENSUS OUTCOME	RUNS	HINDSIGHT CEILING	WHAT A GATE ACTUALLY BANKED	CONVERSION
CAUGHT	14	\$4,689	+\$944	20%
FOUGHT	6	\$2,286	-\$248	—
SAT OUT	53	\$14,519	\$0	the money on the table

The clean oracle proof came from the survivor itself. NIPC does not know anything about the census — it is a mechanical entry — yet run it over the week and check where it was standing when each run fired:

ON-THE-RUN PROOF (NIPC VS THE 21 OPEN/NEWS RUNS)	VALUE
Runs where NIPC held a position in the run's direction	21 / 21
Of those, the exposure finished profitable	17 / 21
Net banked on the runs themselves	+\$1,674
Sat-out subset (the actual gap): aligned / green / net	
10 / 10 · 6 / 10 · +\$189	

Entered at the right time on the right side, these runs pay 17 times out of 21. **The money is real; the exit is not the problem; the ENTRY is the whole problem.**

And now the honest deflation, because I would rather you hear it from me. NIPC converts only **\$189 of the \$3,082** sat-out OPEN/NEWS ceiling — **\$2,487 of its \$2,676 comes from trades that are not census**

runs at all. It is not a run-capture engine; it is a high-frequency in-window scalper that happens to be present and correctly sided during every big run. Two consequences: **do not size this against the \$14,519 ceiling** — the ceiling is not its addressable market; and the hindsight ceiling in general should be read as a proof the runs exist, not as revenue you can plan around. Even our live gates converted only **20%** of the ceiling on the runs they did catch.

Step 3 — ★ THE FINDING: at what SIZE does a footprint become tradeable?

This is the question the full → top25 → top15 narrowing exists to answer, and this week I can finally give you a proper three-part answer instead of a flat "never". The three parts are different kinds of size, and they do not agree with each other — which is itself the useful part.

KIND OF "SIZE"	IS THERE A THRESHOLD?	THE MEASUREMENT
A. Tape-flow footprint size (how heavy is the aggressor print before the run)	NO — AND INVERTED	corr(move , aligned pre-run flow-z) over the 53 sat-out runs = -0.209 . Bigger run → less aligned flow before it. Confirmed by narrowing (table below): the top of the census is emptier of the footprint than the marginal runs.
B. Book / print size (is absorption a measurable physical quantity)	YES — AND MNQ IS ~6× BELOW IT	MNQ: median top-of-book 3.0 lots, top-3 17.2, full visible 10-deep 86.2, against 2,664 contracts/minute of flow = the whole visible book turns over 31× a minute. Median print 1 lot (only 0.62% are ≥10). MES over the same window: top-of-book 23.9, 10-deep 527 — 6.8× / 6.1× deeper.
C. Run size (given an entry that already works)	YES — ~195 POINTS	NIPC on top-quartile runs (move ≥ 195 pt, n=8): 8/8 aligned, 8/8 GREEN, + \$1,131 — 68% of its run P&L from 38% of the runs. Marginal runs (move < 120 pt, n=5): 3/5 green, +\$38, and 2 of its 4 reds live there. Tape-side proxy: impulse span ≥ ~2× ATR1m.

The narrowing evidence for (A), computed this week on the current census:

RUN SET (SAT OUT)	N	MEDIAN ALIGNED FLOW- Z	MEAN	FRACTION > 0
-------------------	---	------------------------------	------	--------------

ALL sat-out	53	+0.116	+0.025	57%
top25 (the 25 biggest)	25	+0.038	−0.189	52%
top15 (the 15 biggest)	15	+0.100	−0.090	60%
FLOW-LED sat-out (the cluster named for having the footprint)	16	−0.089	−0.194	44%

Hunting harder at the top made it worse, not better. The brief's hypothesis — that narrowing to the biggest runs reveals a stronger footprint the marginal ones wash out — is not merely unsupported, **it inverts**. And the cluster explicitly defined by having the footprint has the weakest footprint of all four sets.

Why (B) is the most valuable line in the section. Every "absorption" idea the desk has ever entertained — far-side depletion, the wall test, spoof-pull, iceberg refill — is trying to read a signal out of a **3-lot queue that gets replaced 31 times a minute**. Our live exhaustion_signal wall test ($\text{ask1s} \geq 1.5 \times \text{bid1s}$) is, in practice, **3 lots versus 2 lots**. As a sanity check I asked whether MNQ book imbalance predicts anything at all, unconditionally, at 10s / 30s / 60s / 300s horizons across all 6 days: **correlation 0.000 to 0.015, non-monotone in every quintile**. Zero. That single physical fact independently explains graves G1, G2, G3, G5, G7 in the chop lab and VR-4 in the vacuum lab. **The size threshold for a book footprint is roughly an instrument, not a parameter — and MNQ is on the wrong side of it.**

The catch on (C), stated plainly: the 195-point threshold is a selection result, not a detection result. It tells you which runs a working entry converts; it does not help you find one. Biasing NIPC toward the big-run footprint does improve the number (RMIN 0.40 $\rightarrow$ +\$2,909; K=1.2 $\rightarrow$ +\$2,603, both on fewer trades) — but both settings sit inside the razor plateau, so that is the same edge concentrated, not a second edge discovered.

Step 4 — the graveyard at every narrowing level

Three clusters × three zoom levels. Provenance note, because it matters for how much weight to put on each row: the three **full**-level hunts were run this week on the 07-23→07-31 tape (10.08M ticks). The **top25 / top15** narrowing hunts on file were run on the prior week's tape (07-19→07-24) and were not re-run — I am reporting them as the narrowing-level evidence they are, and this week's own narrowing test (the table in Step 3) reaches the same conclusion by a completely different route, which is the reason I trust it.

FAMILY	LEVEL	BEST-LOOKING CONFIG	N	NET \$	WIN%	BIG-CAUGHT	VERDICT
A · FLOW-LED	full	FLOWTRAP (trapped counter-aggressors)	79	+\$1,460	48.1%	4/16	MIRAGE
A · FLOW-LED	top25	AFC flow-continuation W60/TH250	851	−\$5,745	47%	4/5	BLEED
A · FLOW-LED	top15	FCE (catch config TH100)	883	−\$5,099	—	3/3	BLEED
A · VACUUM	full	VR-2 absorption fade w60/F200/R2	205	+\$1,301	42.4%	1/13	MIRAGE
A · VACUUM	top25	trap-reclaim scalp F300/S40/T30	274	+\$2,110	62%	7/20*	SHORT-BETA MIRAGE
A · VACUUM	top15	trapped-flow reversal F450/T80	16	+\$242	44%	1/15	n-COLLAPSE MIRAGE
B · OPEN/NEWS	full	NIPC news-impulse pullback	171	+\$2,676	43.3%	17/21 green	SURVIVOR → SHADOW
B · OPEN/NEWS	top25	NWIF impulse-fade (ride) L60/TH40	25	+\$1,534	64%	4/8	MIRAGE
B · OPEN/NEWS	top15	NWIF impulse-fade L30/TH30/t40	54	+\$791	63%	2/6	MIRAGE

*The VACUUM top25 scalp "catches" 7/20 only as 30-point fragments of 200-point runs — it exits at +30 and structurally cannot ride a run.

GRAVE — FLOWTRAP (THE BEST THING THE FLOW FAMILY PRODUCED, AND IT STILL DIES)

The tape-wide scan pointed somewhere genuinely interesting and opposite to the cluster's name: flow aligned with a move predicts ~nothing, while flow opposed to a strong move predicts continuation (−113 points of further downside in the US-session "price collapsing while aggressors keep buying" corner; +37 points in the strongest-prior-move quartile). Trapped counter-aggressors are fuel; confirmed flow is already priced. So I built it: short when price fell hard while they kept buying, vol gate at 15-min range ≥ 70 pt, stop $0.45 \times \text{range}$, 2R target, tick-honest exits.

It looked like a winner — **+\$1,460, PF 1.50, and +\$1,632 / PF 1.76 on its home vol segment**. It passed leave-one-day-out (every day removable, all remain green: +\$881 to +\$1,673) and passed the IS/OOS split (+\$706 vs +\$754). Then three tests buried it:

KILL TEST	RESULT
Placebo / permutation — circularly shift the flow series, keep everything else identical, 200 draws each	Real +\$1,460 vs null mean +\$234 (sd \$886, 95th pct +\$1,723) → p = 0.075 ; home segment p = 0.145; expvol-US p = 0.105. The real result sits below the null's own 95th percentile in all three. A fake flow filter on the same price trigger already books ~+\$600 on this tape.
Strip the best trades	+\$1,460 → strip-3 +\$385 (PF 1.13) → strip-5 −\$150 . Five trades out of 79 are the entire edge.
Long/short symmetry	US SHORT +\$1,434 (n=12) · US LONG −\$0.30 (n=9) · ONITE LONG +\$346 (n=34) · ONITE SHORT −\$184 (n=23). Four sub-populations of n=9-34 with two opposite signs — not one policy, four fitted ones.

It also fails the job it was hired for: **4 of 16** target runs caught, **\$1,309 of the \$14,519 ceiling = 9%**. And the lookback window is a step, not a plateau (3min −\$236 · 5min −\$106 · **10min +\$1,487** · 15min +\$1,325 · 20min +\$815) — two adjacent settings are negative. Buried. One thing genuinely worth keeping: the **vol floor is monotone and reusable** — range ≥ 0 gives \$18/trade, ≥ 70 gives \$31, ≥ 90 gives \$56, while every one of 15 stop $\times$ R cells below the floor is red. That re-confirms `[[low-vol-grind-untradeable]]` mechanism-independently.

Plus the control graves that show why FLOWTRAP looked real: aligned-flow momentum — the census's own literal FLOW-LED read —

n=171, **+\$74**, **PF 1.01**, a dead coin flip (longs $-\$799$ / shorts $+\$873$ cancelling out). Flow-only continuation with no price condition: n=301, **-\$2,179**, PF 0.82. Price-momentum only: n=280, **-\$889**. Contra-flow fade at a 30-min extreme: n=79, **-\$973**. Contra-flow fade with no price condition: $+\$1,845$ — but 07-29 alone is $+\$1,483$ of it, a single-flush artifact.

GRAVE — THE VACUUM FAMILY: FOUR BRAND-NEW SIGNALS, FOUR BURIALS, AND THE PARTICIPATION PARADOX

Four genuinely new specs, all causal, all two-sided, none reusing an existing gate: **VR-1 Vacuum Reclaim** (flush $\rightarrow$ trapped aggressors $\rightarrow$ reclaim), **VR-2 Absorption Fade** (heavy flow, no travel, at the extreme), **VR-3 Trap Break** (delayed confirmed continuation), **VR-4 Liquidity Vacuum** (the literal L2 reading, using book depth the desk does not consume).

FAMILY	CONFIGS	% POSITIVE	MEDIAN NET	SURVIVES STRIP-3	MEDIAN OOS	MAX RUNS CAUGHT
VR-1 reclaim	288	23%	$-\$800$	8%	$-\$306$	5/13
VR-2 absorption fade	192	48%	$-\$65$	22%	$-\$104$	5/13
VR-3 trap break	144	49%	$-\$4$	25%	$-\$206$	6/13
VR-4 book vacuum (L2)	288	34%	$-\$153$	7%	n/a	2/13

VR-2 was the one worth chasing — a genuine plateau of 54-60% cost-free hit rate across a whole parameter block, not a spike. Its best config nets $+\$1,301$ (n=205, 42.4% win) — then strip-3 takes it to **+\$596** (and all three best trades are the same clipped max winner), the genuine OOS day pays **+\$94** ($\$2.6/\text{trade}$, neighbour $+\$11$), and across the whole config block **only 36% are positive, median $-\$271$, median $-\$1.45/\text{trade}$.**

The killer table — participation is bought with churn. Across 1,200 configs, $\text{corr}(\text{runs caught}, \text{net}) = -0.32$:

RUNS CAUGHT (OF 13)	CONFIGS	MEDIAN NET	% POSITIVE
0	69	$-\$320$	36%

2	434	-\$459	24%
4	38	-\$1,274	16%
5	37	-\$1,503	0%
7	3	-\$3,091	0%

Cause of death, proven ex-ante. I took the 13 run starts and asked where each candidate feature sat inside 6,000 random live bars. Position in the 2-min range: **53rd percentile, $p=1.000$** — which kills the "fade at the extreme" premise outright. Absorption (points per contract): **50th percentile** — no absorption signature at all. Far-side book depth: **48th percentile** — there is no liquidity vacuum in the book. The only "significant" feature is the signed 60-second flow itself, at the 22nd percentile — which is tautological, because VACUUM is defined as opposite flow, and it points the wrong way. **The VACUUM label is a post-hoc description of what happened, defined by the sign of the outcome. Nothing measurable at the trigger separates the 13 runs from the ~500 identical-looking setups that go nowhere.**

GRAVES — THE NARROWING LEVELS (PRIOR-WEEK TAPE), AND THE SURVIVORSHIP SQUEEZE IN THE RAW

FLOW-LED at top25 and top15 is the cleanest demonstration of the squeeze I have ever produced. At top15 the three target runs carry the cluster's strongest aligned flow (228-421 contracts vs the marginal cluster's 50-155), yet forward that same flow is a **46-53% coin flip** at every window and horizon, and strong sell aggression is followed by price drifting up. Then:

CONFIG	RUNS CAUGHT	NET	ON-RUN P&L	OFF-RUN P&L
FCE W60/TH100 (loose)	3 / 3	-\$5,099	$n=4 \cdot +\$42$	$n=879 \cdot -\$5,112$
FCE W60/TH250 (base)	1 / 3	-\$3,118	$n=1 \cdot +\$6$	$n=545 \cdot -\$3,114$
FCE W30/TH750 (least-bad)	0 / 3	-\$681	$n=0$	$n=86 \cdot -\$681$

Noise-to-signal 220:1. To fire on all three runs you must fire loose and lose -\$5,099; tighten enough to stop bleeding and you catch zero. There is no threshold in between, because forward there is nothing

separating the 3 winners from the ~880 look-alikes. All 24 swept cells are red; the loss **asymptotes toward zero from below and never crosses positive** — and strip-3-best makes it worse ($-\$681 \rightarrow -\968), so there is not even a fat tail to point at. The independent breakout variant (RBFC) is red in every cell too. Same story one zoom out at top25: AFC $-\$5,745$ over 851 fires (on-run $n=8$ $+\$212$, off-run $n=843$ $-\$5,956$); RBFC $-\$2,848$.

VACUUM at top25 produced the week's most seductive mirage: a trap-reclaim scalp at $+\$2,110$ on 62% win that is **pure short beta** — SHORT $+\$2,454$ vs LONG $+\$324$, printing only because MNQ fell 1,316 points that week — sitting on a shape knife-edge (reclaim buffer $3 \rightarrow \$3,573$, $6 \rightarrow \$2,110$, $10 \rightarrow \$718$; flow window 45s $\rightarrow -\$1,105$, 60s $\rightarrow +\$2,110$, 120s $\rightarrow -\$732$). Its run-catching sibling fires **zero trades on the two biggest runs of that week (+221 and -103)**. At top15 the best config exists only at $n=16$, is negative at every tradeable threshold (F100 $-\$995$ through F350 $-\$202$), and strip-3 flips $+\$242$ to $-\$224$ — nine of its sixteen trades are identical $-\$55$ stops.

OPEN/NEWS at the narrowing levels is where the mechanism rotated but the outcome did not: the forward tape said news-window impulses mean-revert (continuation only 35-44%), so a fade looked alive at $+\$1,534$ — until 73% of it turned out to come from 5 survivorship-defined on-target trades, strip-3 removed 58%, and the winning side flipped between candidates (C: long $+\$656$ / short $-\$249$; RIDE: long $-\$156$ / short $+\$338$). At top15 the $+\$791$ fader **fights the two biggest runs** (-194 and -178 , each $-\$85$), catches 2 of 6, sits on a 10-point target knife-edge ($t30 -\$134 \rightarrow t40 +\791), and loses $-\$240$ on the one trend day where 3 of the 6 targets live.

The through-line: at every zoom level, in every family, the green numbers are a cooldown scheduling accident, a razor-edge peak, short beta to a down week, or three lucky trades — never a base rate.

Step 5 — THE SURVIVOR: NIPC, the news-impulse pullback (full level, OPEN/NEWS)

Here is where the detective story turns. The prior week buried a Donchian **breakout** in this exact window, and that grave is what made this week's hunt worth doing properly: it says chasing fails here, which is a fact about the tape, not a verdict on the cluster. So before inventing anything I labelled **every 5-second bar** in the 13:00–15:00 window across 12 days (17,274 observations) with backward-only features and the forward 15-minute payoff. Two facts reframed the whole problem:

SEPARABILITY FINDING	CONSEQUENCE
P(15-min excursion ≥ 90 pt) = 0.499 — half of all bars in this window lead to a census-sized move	OPEN/NEWS is not a detection problem. It is a direction-and-entry-price problem in a tape that is always moving.
A random entry in this window is a −\$846 average loser (40-trial Monte-Carlo, same window / stop / exit, sd \$472)	The window is a meat-grinder. Any candidate must clear a deeply negative null, not zero.
Pre-run aggressor delta is sign-aligned with the run 11 times out of 25	Kills the "follow the flow" family and the "fade the flow" family before either is built. (07-31 14:16 ran UP +193 on delta −2,791.)
Position in the 30-min range predicts continuation, not reversion	The direction read of the buried Donchian was right all along. The entry price was wrong.

The mechanical spec — this is the deliverable. Armed only $13:00 \leq t < 15:00$ UTC. Decisions on 5-second bar closes, fills on ticks. Direction is always with the impulse; both sides always armed; no flow or book input at all.

STEP	RULE
1. IMPULSE	Over the trailing 2 minutes (24 bars), $leg = close[i] - close[i-24]$. Require $ leg \geq 1.0 \times ATR_{1m}$ measured before the impulse (no self-reference). Side = sign(leg). ext = the impulse extreme, org = its origin, span = ext − org , require span ≥ 5 pt.
2. PULLBACK	Walk forward up to 3 minutes tracking the retrace $r = ext - pullback / span$. If $r > 0.75 \rightarrow$ ABANDON (structure broken, that was not a pullback). ARM at the first bar where $r \geq 0.35$.
3. TRIGGER	Entry level = half-back of the pullback (pullback extreme + 0.50 × the retrace distance), and it must still be inside the impulse. Live for 1 minute; fill at the first tick trading through it. No fill → the setup expires, no trade.
4. STOP	The pullback extreme ± a 4-point buffer. $R = entry - stop $. Reject the setup if $R < 2$ pt. (Median R in practice = 17.2 pt.)
5. EXIT	1 lot: fixed 2.5R target. Hard 20-minute time cap; flat by 15:30 UTC. Two-lot desk architecture: Lot A 2.0R / Lot B 2.5R (proven, below).
6. RISK	One position at a time; 2-minute cooldown after exit.
7. REGIME	OFF in dead-chop ($ATR_{1m} < 18$ pt AND $ER_{15} < 0.35$). ON in every other regime.

Why this is a different animal, not a filtered breakout. One ablation settles it — strip the pullback requirement (RMIN=0, RES=0) and NIPC collapses into a plain impulse-chase: **−\$2,143 at 27.9% win over 290 trades.** Same impulse detector, same regime gate, same exit, same tape, same costs. The pullback-and-resume structure is the edge; it is **the opposite trade at the opposite price**, not a filter bolted onto a breakout.

The backtest, scored per regime on HOME segments only (12 days, 1 lot, 2.5R, \$5/RT + a tick of slippage each side). Regime thresholds are taken from this window's own quartiles so every bucket is populated:

REGIME	N	NET \$	\$/TRADE	WIN%	POLICY
normal-chop	70	+\$1,013	+14.5	44.3%	ON
violent-whipsaw	35	+\$893	+25.5	42.9%	ON
in-between-building	54	+\$539	+10.0	42.6%	ON
clean-trend	12	+\$232	+19.3	41.7%	ON — thin n, NOT proven
dead-chop	34	−\$110	−3.2	32.4%	OFF
HOME (dead-chop OFF) — the policy	171	+\$2,676	+15.6	43.3%	THE HEADLINE
blanket (all regimes) — the number this report refuses to headline	205	+\$2,566	+12.5	41.5%	reported only to show the cost of not segmenting

The dead-chop gate is not a fitted knob: **dead-chop came out negative in all 25 scale-out configs and in every top config of the 270-config exit sweep.** It is the one perfectly stable regime finding.

Per regime × side — every single cell positive:

REGIME	LONG N / NET / WIN	SHORT N / NET / WIN
normal-chop	32 · +\$67 · 34%	38 · +\$946 · 53%
violent-whipsaw	14 · +\$857 · 57%	21 · +\$35 · 33%
in-between-building	22 · +\$305 · 45%	32 · +\$234 · 41%
clean-trend	7 · +\$175 · 43%	5 · +\$56 · 40%
TOTAL	75 · +\$1,404 · 43%	96 · +\$1,272 · 44%

Honest read on the two thin cells: normal-chop LONG (+\$67) and violent-whipsaw SHORT (+\$35) are flat, not proven — they are "not bleeding", nothing more. But the aggregate split is as symmetric as it gets: **long \$1,404 vs short \$1,272 on a sample of 7 down days vs 5 up days, and**

–1,055 points of net drift inside the window. This is not a downtrend rider. That single fact is what separates NIPC from every mirage in Step 4.

Per day — 9 of 12 green, worst day –\$176:

DAY	N	NET	DAY	N	NET
07-16	15	+\$576	07-24	14	+\$250
07-17	15	+\$410	07-27	14	+\$532
07-20	14	+\$282	07-28	13	+\$329
07-21	13	+\$46	07-29	14	+\$142
07-22	15	–\$27	07-30	13	+\$435
07-23	14	–\$124	07-31	17	–\$176

≈14 trades a day, mean hold 2.4 minutes, median R = 17.2 points, 74 targets against 97 stops.

Robustness — the exact battery that killed last week's candidate:

TEST	RESULT	VERDICT
Parameter razor (one-at-a-time, 44 variants across all 8 constants)	44 / 44 positive — a plateau. Compare last week's Donchian: B3 –\$278 → B5 + \$570 → B10 –\$181.	PASS
Long/short symmetry (each side on its own full trade set)	LONG +\$1,404 (n=75) · SHORT +\$1,272 (n=96)	PASS
Strip the best trades	strip-1 \$2,415 · strip-3 \$2,091 (78% survives) · strip-5 \$1,770 · strip-10 \$1,137	PASS
Leave-one-day-out (all 12 days)	net without any single day: \$2,100 – \$2,852	PASS
IS / OOS split	IS 07-16..23 (bar-replay) n=86 +\$1,164 · OOS 07-24..31 (tick-honest) n=85 + \$1,512 — OOS beats IS	PASS
Census week alone	n=71, +\$1,262, 43.7% win, \$17.8/trade	PASS
Cost sensitivity	\$5 → \$2,676 · \$7 → \$2,334 · \$10 → \$1,821 · \$12/RT → \$1,479	PASS — 2.4× headroom
Stop-slippage stress (extra points on stop fills only)	+0.5 → \$2,579 · +1.0 → \$2,482 · +2.0 → \$2,288 · +3.0 pt → \$2,094	PASS — 12× modelled slip
Entry-latency stress (trigger goes live 5 / 10 / 20s late)	\$1,940 / \$1,412 / \$1,601 vs \$1,498 baseline	PASS — not microstructure
Null control (40 Monte-Carlo trials: random side, random bar, same window / stop / exit)	null mean –\$846, sd \$472, best-of-40 = +\$336 — NIPC z = +5.0	PASS

Ablations — what is actually load-bearing:

VARIANT	N	NET \$	WIN%
FULL NIPC	171	+\$2,676	43.3%
no pullback requirement → plain impulse-chase	290	–\$2,143	27.9%
shallow pullback only (RMIN=0.15)	257	–\$203	33.9%

no resumption trigger (buy the dip blind)	230	+\$1,893	44.3%
no invalidation (RMAX=1.5)	173	+\$2,459	42.2%
rvol filter ON — the filter I built and then deleted	110	+\$1,029	39.1%
impulse window too short (1 min)	167	+\$1,588	41.9%
dead-chop gate removed — the dead-chop slice alone	34	−\$110	32.4%

A real $\geq 35\%$ retrace actually happening is the edge. Waiting for the resumption instead of catching the falling knife is worth ~\$780 on top. The volume filter was monotonically destructive (RV=0 $\rightarrow$ +\$2,474 down to RV=1.4 $\rightarrow$ +\$1,413, costing \$976 and 40% of the sample) and was **removed, not tuned** — inside 13:00–15:00 every bar is already a volume-expansion bar, so RVOL has no cross-section left to cut on.

Your exit cheat-sheet, swept on NIPC entries — proven vs guess. The prior for a momentum gate is Lot A 1.5R / Lot B 2.5R. Two lots, \$5/RT each, HOME only:

LOT A	LOT B	N	NET \$	WIN%	\$/TR	STRIP -5	NORM -CHOP	IN-BETWEEN	VIOLENT	TREND	DEAD-CHOP
1.5R	2.5R (the GUESS)	171	4,642	56	27.1	3,203	2,026	840	1,530	246	−225
2.0R	2.5R (PROVE N)	171	5,199	50	30.4	3,573	2,108	1,090	1,646	355	−199
2.5R	2.5R	171	5,352	43	31.3	3,540	2,025	1,078	1,785	464	−221
1.5R	trail 1.5R@1.5R	167	3,357	56	20.1	1,848	1,278	1,161	660	257	−222
2.5R	4.0R	141	5,038	43	35.7	3,086	1,400	1,071	1,340	1,228	−331
3.0R	5.0R	134	3,557	33	26.5	1,406	1,572	105	584	1,296	−280

CLAIM	PROVEN?
Momentum Lot A 1.5R	HOLDS, MILDLY CONSERVATIVE Proven Lot A is 2.0–2.5R, so the guess is ~11% low — but the plateau from 1.5R to 2.5R is broad (\$4,642 $\rightarrow$ \$5,352), so the guess sits on the plateau. No rewrite needed.
The standing chandelier / trail prior for Lot B	CONTRADICTED IN THIS WINDOW Fixed 2.5R beats a trail decisively (\$4,642 vs \$3,357 at the same Lot A). Cross-checked on the 1-lot sweep: trail 1.0R@1.0R = +\$1,033, worst of 8; fixed 2.5R = +\$2,676, best of 8. In a 13:00–15:00 whipsaw the trail is taken out by the counter-leg before the extension resumes.
The tail is real...	...but only in the trending buckets . Pushing Lot B to 4.0R moves money out of normal-chop (2,108 $\rightarrow$ 1,303) and into clean-trend (355 $\rightarrow$ 1,094). A regime-flexed Lot B — 2.5R in chop, 4.0R in clean-trend — is the honest read, but clean-trend n=12–13 is far too thin to bank. Logged, not adopted.
Is the exit surface a plateau or a spike?	

209 of 270 exit configs (77%) positive; all 25 scale-out configs positive on HOME and all 25 negative in dead-chop. A plateau. No config-hunting was needed.

Where NIPC fails, stated up front. All four red runs (07-28 14:02, 07-29 13:19, 07-31 13:20, 07-31 14:34) are **first legs that reversed inside the run window** — the impulse was real, the pullback held, and the 15-minute "run" then completed in the other direction after NIPC had been stopped and re-entered. That is the residual failure mode, and the repair is a one-parameter change: a second-entry-in-the-same-direction veto, or an entry-count cap per impulse cluster.

The verdict, and what I want done with it

CANDIDATE	LEVEL	HEADLINE	VERDICT
NIPC — news-impulse pullback continuation	full · OPEN/NEWS	+\$2,676 · n=171 · 12 days · 43.3% · \$15.6/tr	SHADOW — never live yet
chop_turn_fade — chop-day turn scalp	chop lab	+\$4.42/tr · t=1.50 from 1,152 trials	INCUBATE — expect it to revert to zero
Everything in the flow-footprint family (FLOWTRAP, VR-1..4, AFC, FCE, RBFC, trap-reclaim, NWIF)	all levels		— BURIED

Shadow NIPC, do not arm it. n=171 over 12 days is one regime-fortnight. There is no cross-regime tape here: no FOMC week, no CPI shock, and — critically — barely any genuinely quiet news window (dead-chop is only 34 trades, and that is precisely the bucket the gate turns off). The desk's own promotion ladder says shadow → 1 lot, gated on 8–12 weeks and cross-regime evidence. I am not going to be the one who short-circuits it on a fortnight.

The claim to re-test first is the EXIT, not the entry. "Fixed 2.5R beats a trail in the news window" contradicts our standing chandelier prior and is worth ~\$1,285 in the sweep. If that flips on fresh tape, the whole exit half of the spec flips with it.

Fix the census, not the desk. The FLOW-LED / VACUUM rule fires on **66.5% of all bars**. It should be a z-score against a trailing 2 hours with a real threshold ($|z| \geq 1$), otherwise the two clusters are two names for one

coin and next Friday's greenfield agent digs this same empty hole. That is a **one-line fix in scripts/run_census.py::cluster()** and it is the highest-value change in this section.

The one stone still worth turning: run the chop-scalp harness on **MES, 2026-07-12 → 07-17**. The L2 lead did not fail because absorption is a bad idea — it failed because **MNQ's book is too thin for absorption to be a measurable quantity**. On MES the same construct is a real physical measurement (23.9 lots at the top, 527 across 10-deep, $6.8\times/6.1\times$ MNQ). The window exists and the blocker is known: `alphabot2/data/depth.db` holds 10-deep MES book for 07-12→07-31 and `alphabot2/data/ticks.db` holds MES prints for 07-05→07-17 — **they overlap on 07-12→07-17, roughly 4-5 session days with both a tape and a book**. Everything is parameterised and re-runnable.

Caveats, stated plainly. One completed week for the census; 12 days for NIPC (6 of them 5-second bar-replay rather than tick-honest — and those are the stricter half, which scores lower, the right sign); 6 days and only **3 chop days** for the chop lab; MNQ-only; PAPER. The narrowing-level (top25/top15) hunts are prior-week tape and were not re-run. Every number above is net of at least \$5/round-trip with a tick of adverse slippage on entry and exit, stops filling a tick worse and winning ties — deliberately unkind, and it should stay that way.

SURVIVOR: NIPC → SHADOW

chop lab: HONEST NULL

flow footprint: no size threshold, inverted (−0.209)

book footprint: threshold is real — MNQ is 6× under it

run size threshold ≈ 195 pt

chop days: the OFF switch beats the gate