

GAZBOT V7 — the Friday report

Week ending 2026-08-08 · Mon 3 – Fri 7 August 2026 · MNQ-only · PAPER ·
REVISION 2, rebuilt 2026-08-08 · 52 plays · 9 sections, complete

What Revision 2 is. Revision 1 was killed mid-build by the machine running out of memory. It shipped titled “PARTIAL, not proofread”, ended mid-sentence at “That is Movement 3's job.”, carried **no Movement 3** and **no Monday playbook**, and its PDF had been rendered from an earlier copy of the HTML than the one on disk — so the two artefacts disagreed with each other. Revision 2 replaces all of it. All nine sections are stitched, the build now refuses to run at all if any section is missing or empty rather than quietly emitting a short report, and the HTML and the PDF are rendered from one string in one pass so they cannot diverge again.

★ THE HONEST HEADLINE — the desk lost money, and it is not mysterious

The tournament lost **−\$772.50** this week on 110 booked lots. The loss is not spread around: the **long book made +\$345.50** and the **short book lost −\$1,022.50**, and one gate — **abs_veto_short** — lost **−\$559.00 on nine lots without winning a single one**. Take the shorts away and the week is green.

There is no single honest total, so here are all of them, with what separates each from the next.

WHAT YOU ARE LOOKING AT	LOTS	NET \$	WHY IT DIFFERS FROM THE ROW ABOVE
Raw trades rows, 08-03 → 08-07	112	−1,028.00	includes the two corrupt MD_STREAM lots
Clean ledger (data_quality IS NULL)	110	−772.50	the desk's real booked money — this is the number to quote
Tournament gates only	108	−677.00	strips the day-rider's cross-desk pair (−\$95.50), booked to the

WHAT YOU ARE LOOKING AT	LOTS	NET \$	WHY IT DIFFERS FROM THE ROW ABOVE
			tournament by Thursday's flatten
Gate strategy, defects stripped	104	−607.00	strips the ghost stop and the phantom close — −\$165.50 of the loss was three shared- account defects, not a strategy losing money

READ THE WEEK BY VOLATILITY, NOT BY DIRECTION

Every tape number below is recomputed from capture.db 5-second bars rolled up to one minute, not read off a dashboard. **The desk does not lose money on trend days. It loses money on high-volatility days that go nowhere.** Monday walked 482 points up in the US session at ER 0.171 — a genuinely good day to be long — and the desk lost \$280 on it. Wednesday and Thursday, the two days ATR doubled from 12 to 20 points and the tape finished where it started, took **−\$1,122** out of the desk between them. Thursday travelled **9,964 points over the full session to finish eleven points up** — ER 0.001, the flattest day in the record. With n=5 sessions that is an observation, not a law — but it is the same shape the whole week keeps making, and it is why the single largest number in this report is an off switch.

DAY	US NET	US PATH	US ER	US ATR	READ	LOTS	WINS	DESK NET
Mon 08-03	+482	2,815	0.171	12.4	clean up-trend	30	8	−280.00
Tue 08-04	+641	2,799	0.229	12.6	the week's cleanest trend	16	10	+330.50
Wed 08-05	−367	4,095	0.090	20.1	whippy down — ATR doubled	29	4	−817.00
Thu 08-06	+161	4,253	0.038	20.5	pure chop at high vol	14	3	−305.00
Fri 08-07	+120	3,676	0.033	14.2	chop, but a real London run at 12:30	21	12	+299.00

Read that table twice. **The two cleanest trend days produced +\$50.50 between them.**



THE SCORECARD — what survived the week, and what did not

Six claims went into this week carrying real weight. Here is where each one ended up, with the test that decided it. Read this before anything else in the report.

THE CLAIM	WHAT IT CLAIMED	VERDICT, AND THE TEST THAT DECIDED IT	WHERE
abs_veto_long is the one gate that works	The thrust + 55-second confirm earns its place	SURVIVED — the one number worth weight 16 entries, 32 lots, +\$566.00 at 62%, average winner +\$47.60 against an average loser of – \$32.20. Still positive if you delete any single day of the week, and still +\$363 if you delete its three best lots. Everything else on the desk this week is thin, one-week, or explained by a software defect.	Part 1
Never bench a gate at face value	The gates that bled are fixable, not broken	CONFIRMED — four of six rehabilitated Not one of the six was broken in the way it looked broken from the roster. grind_long looked like a bad gate and was a bad ATR floor. abs_veto_short looked like a bleeder and was the strongest signal on the desk on a 3.3% leash. MAX_HOLD looked like the worst exit on the desk and was the fire alarm being blamed for the fire. exhaustion_short looked like an arming problem and is a friction problem. nipc_short looked like a wrong-rung problem and is a direction bet. The chandelier looked like a threshold that stopped being reached and had been switched off in code.	Part 1.5
RUN STATE is the primary discriminator — carried into the week as the desk's headline routing lead	in-run aligned +\$10.34/trade against chop –\$28.43/trade — a \$39/trade spread that dwarfs gate, side and time-of-day	WITHDRAWN AS A ROUTING LEVER The variable is real — it beat 4,000 placebo shuffles twice and held on twelve unseen sessions. The headline cell is not. Re-derived from scratch on the same 89 trades, in-run-aligned TAKEN goes from $n=34 \cdot 50\% \cdot +\\$10.34/\text{tr}$ to $n=46 \cdot 41\% \cdot -\\$1.83/\text{tr}$, and the 34/4/51 split could not be reproduced under any run-span definition tried. Armed inside an aligned run the desk made nothing. And the lever is in the wrong place: arming more was worth ~\$170 this week, but re-pricing the exits on the exact same entries	Part 2.5

THE CLAIM	WHAT IT CLAIMED	VERDICT, AND THE TEST THAT DECIDED IT	WHERE
		— changing nothing about routing — turns the week from –\$772 into +\$1,493. It is an exit find wearing a router's coat.	
The router is earning its keep	The desk's #1 lever, on the stand	YES — its first clearly positive week + \$2,951 of net value on its own nightly repricing, up from – \$1,216 the week before, and without thrashing: 1,393 ticks, 66 switch changes, 95.3% of ticks changed nothing. The independent check agrees — the two gates it kept benched had shadow twins losing \$4,158 over exactly those windows. The desk still lost \$953 on the router's own Paris-day basis, but it lost it through gates the router had ARMED, not gates it failed to bench.	Part 2.6
The day-rider's exit is exit-proof — holding wins	25 exit variants tested, holding beat every one	HALF WRONG Re-run over 38 sessions scoring 114 ways of ending a ride on identical entries: holding still beats every stop, but its whole edge is three days — strip them and \$5,410 becomes \$379. The seven days that lost \$1,000+ lost –\$11,195 between them and the trail we shipped last Friday saves exactly \$0.00 on every one, because price never went our way at all. The exit parameter surface is noise (Spearman –0.30 between in- and out-of-sample rank). The entry is the lever: efficiency floor 0.15 → 0.25 takes the same 38 sessions from +\$5,410 to +\$8,819, 32 of 32 leave-one-out folds positive.	Part 2.7
The greenfield lab will find us a new gate	\$9,917 of hindsight money sat in 54 runs we sat out — go build something that catches them	ONE SURVIVOR, AND IT GOES TO SHADOW Three hunts from a blank sheet; two came back with nothing. The OPEN RIDER — no trigger at all, just show up every five minutes in the open window facing the way the last fifteen minutes went, with a stop twice the house width — made money on 36 days including 19 it was never fitted on (+\$30.66/trade out of sample against +\$32.29 in, a 5% degradation). Three independent skeptics were told to destroy it; the vote was 2-1, not 3-0. It goes to the shadow book. Nothing from this week's greenfield goes live — 17 in-sample days earns SHADOW at most.	Movement 3

THE THREE THINGS THIS WEEK ACTUALLY TAUGHT THE DESK

1 · The arming is worth more than the exits, and the exits are worth more than the entries. `abs_veto_short`'s signal is +\$2,290.50 over 159 fires and the desk had it benched for 91% of the week; the same section shows re-pricing exits on entries we already took is worth +\$2,265 of swing. Every new gate invented from scratch this week is worth less than either.

2 · We are late, not wrong. Entered on the ignition minute with an honest racing exit, the 52 sat-out runs pay **+\$6,846 — 71% of their hindsight ceiling**. Turn up two minutes later and it is +\$6,041; five minutes, +\$4,090; ten, +\$944; fifteen, **-\$2,068**. Every minute of confirmation costs about \$600. That is why the idle-gate lab is a null: our gates fire seven to eight minutes before the money.

3 · Our house stop is the wrong instrument for the open. Three independent routes — a 45-cell grid, a from-scratch skeptic rebuild, and the oracle — agree that every cell at a 1-ATR stop is flat-to-negative and all 27 cells at 2.0×ATR or wider are positive. Even a perfect entry is shaken out of 14 of 52 runs at 1 ATR, against 0 of 52 at 2 ATR. Checking that on the live gates costs nothing.

★ WHAT TO ACTUALLY DO

Garrath — this is the whole answer to “so what do I actually DO?”, and it is at the front. Every row below mirrors `reports/friday_v7/plays.json`, which is the reconciled source of truth for this week; the same file drives the [Monday playbook](#), so the two cannot drift apart. 52 plays, five windows, in the order they need doing. Every play carries its **evidence tier** and its **kill criterion** — if a play has no way to be proven wrong it should not be on the card.

1 · Sit out the chop (MONDAY #1). The largest number anywhere in this document, and it needs no code at all: chop-regime entries cost $-\$5,483$ over 17 days while everything else made $+\$1,482$ — perfect avoidance turns $-\$4,001$ into $+\$1,482$. A perfect chop filter beats every invention in this report by an order of magnitude, and the router already owns the lever. This is an architectural principle, not a play: **the edge is refusing rubbish trades, not finding more of them.**

2 · abs_veto_short: arm it by default — and without the ER30 floor (SATURDAY #1). This is the biggest single leak in the report and three independent sections found it. The gate was benched for 91% of the week while its shadow twin made $+\$928$, and armed for the 9% in which it lost $-\$559.00$ on nine lots without a single winner. The signal is $+\$2,290.50$ over 159 fires and 4-of-4 weeks green. The router had it backwards on both sides. **The ER30 ≥ 0.35 arming floor this play originally carried is WITHDRAWN as refuted** — it keeps 5 of 15 winners and discards $\$1,363$ of profitable trades ([Part 3.1](#), [Part 3.6](#), corroborated independently in [Part 3.3](#)). Arm the gate, ship the BUILDING $\times$ US-SESSION veto, leave ER alone. $\triangle$ Judge the first 15 armed fires on the armed book alone — not blended with the shadow twin. On the live week this gate was **0 for 8 at $-\$523.50$** , the worst short gate on the desk; the entire case for arming it comes from the continuous book, so the 15-fire review is a hard gate and not a formality.

3 · grind_long: revert the 08-01 floors (SATURDAY #2). ATR $10 \rightarrow 22$, and **delete** the ext_hi key rather than pinning it to 3.0 — everything from 3.0 up scores identically, so a new literal just re-fits the thing the revert is undoing ([Part 3.2](#)). One literal changes, one is removed. Quote the number with its shape attached: **$+\$4,151$** is a delta over 11 tick-honest sessions and 165 legs, of which **21 runner events are the entire result**. $\triangle$ The week's six live grind fills all fired at ATR ≥ 22 and are silent on the revert — do not quote $-\$172$ as support for it. $\triangle$ And grind_long is not in reactivate_gates.py: :HOLD, so the 22:00 UTC Sunday timer re-arms it either way: shipping the floor or adding a HOLD entry are the only two states available.

SATURDAY · TONIGHT — THIS WINDOW SHUTS AT THE SUNDAY 22:00 UTC
REOPEN · 10 PLAYS

Code and config changes that need a restart, **in execution order, not in order of headline size. #1 and #2 are forced:** `gazbot7-gate-reactivate.timer` arms every =off gate at 22:00 UTC Sunday, so for both of them “decide later” is not one of the available states. #3–#6 are safety and measurement, and they come before the two experiments at the bottom because those experiments cannot be judged without them.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
1	abs_veto_short: arm it by default LIVE ★ REV2 CORRECTION (REV2 · Q0 + FIX2, two independent derivations) — the ER30 ≥ 0.35 arming floor is WITHDRAWN as REFUTED; every other half of this play ships unchanged. Stop running <code>abs_veto_short</code> as a 22-minute discretionary window: arm it by default, delete the two bad bench rules (the US-session violent-whipsaw bench and 'bench on a stop-out pair'), fix the 22:00 auto re-arm, and add the one new entry veto from the rehab — BUILDING × US-SESSION. Do NOT add an ER floor of any kind, to either twin. The floor was justified as 'mirroring grind's'; <code>deciders.ER_FLOOR</code> is {} — every floor was deleted on 08-01 and the deletion is pinned by <code>tests/test_deciders.py::test_er_floors_all_deleted</code> . There is no floor to mirror. Mechanism: <code>abs_veto_short</code> — thrust + 55-second absorption veto, short side Why: This is the single biggest leak in the report and both halves of it point the same way: the gate's SIGNAL is the strongest thing on the desk, and the desk's ARMING of it was backwards on both sides — benched through the good tape, armed into the bad. Evidence: Part 3.1 (REV2 · Q0) + Part 3.6 (REV2 · FIX2) — supersede Part 2.6 §8 Saturday #1/#2; extend Part 1.5 rehab_abs_veto_short §5 · verification: Track live fills against the shadow twin for 15 fires; they should stop diverging. · mode: live-paper	+\$2,497.50 vs +\$2,290.50 baseline, 15/15 winners kept, +\$913.50 out of sample. The signal is +\$2,290.50 over 159 fires / 17 days / 4-of-4 weeks green / positive in all five regimes, and 53 of 54 filters tested LOSE to just letting it fire. Over the 96.7% of the week it was benched, the identically-configured twin made +\$1,064.50; the router armed it for the 9% in which it lost – \$559.00 on nine lots without a single winner. The BUILDING × US-SESSION veto is worth +\$207.00 · 15/15 winners · strip-3 +\$1,987.00 · net at the 98.5th placebo percentile — but still n=8, re-look at n≥25. ★ REV2 CORRECTION — WITHDRAWN: the ER30 ≥ 0.35 floor keeps 5 of 15 winners, discards +\$1,363 of profitable trades, takes OOS +\$809 → +\$363, fires zero times in wk29, is – \$3,863 pooled across nine exits, sits at the 77th placebo percentile on winner retention (random expects 3.4 of 15), is a spike not a plateau (0.325/0.35/0.375 = +\$849/+\$927/+\$679), its own live evidence dies at strip-1 (+\$550.50 → +\$14.50), and it silently SUBSUMES AND DELETES the BUILDING × US veto rather than adding to it.	Arms: Armed by default. No ER condition of any kind. One per-entry veto: skip when <code>regime=BUILDING</code> and 13:30 ≤ UTC < 20:00. Exit: Unchanged deployed two-lot ladder. (The 55-second entry WAIT is a separate open question — see BUILD: shrink the wait, keep the test.) Kill: Review after 15 fires armed-by-default; re-bench if the armed book is negative over those 15.	Garrath — gate config + <code>reactivate_gates.py</code> + tournament restart Re-bench it in <code>gate_switches.env</code> . Fully reversible on any 5-minute tick; the ER floor fails safe because a floor can only ever reduce arming.
2	grind_long: put the ATR floor back to 22 and DELETE the extension ceiling (not 3.0 — delete it) LIVE ★ REV2 CORRECTION (REV2 · Q1 §10) — the mechanism changed, the direction did not. Revert the 2026-08-01 change: ATR floor 10 → 22 in <code>deciders.py::ATR_FLOOR</code> , and DELETE the "ext_hi": 2.0 key from the <code>grind_long SlotSpec</code> so it falls back to gate_grind's own default of 4.0. The earlier card said 'ext 2.0 → 3.0'; do not pin 3.0. Everything from 3.0 upward scores identically (+\$5,893 / +\$6,081 / +\$5,981 / +\$6,262), so there is no optimum to find and 2.0 is a \$2,210 tax. Two literals: one changes, one is removed. Change nothing else — the exit (start_k 3.5 / lock_r 6.0 / lock_k 0.5), the 1.0×ATR stop, Lot A's 2.5R, base_size=2, slope_min 0.4, ext_lo 0.3 and fast_slope=True all survive their own sweeps. Mechanism: <code>grind_long</code> — trend-continuation gate, entry-side ATR floor and extension ceiling Why: <code>grind_long</code> did not look broken because it is a bad gate — it looked broken because of a bad ATR floor. The 08-01 change cut the floor to 10 and the gate started taking 744 legs to earn what 155 legs earn. Evidence: Part 3.2 (REV2 · Q1) — supersedes Part 1.5 §1/§10, Part 2.5 LEAD 3, Part 2.6 §6/§(b) · verification: Re-score on the next 11 sessions and confirm the leg count stays near 155, not 744. · mode: live-paper	★ REV2 CORRECTION — the headline was a DELTA, not a P&L, and the recommended cell is +\$4,151 (+\$6,081.2 – \$1,929.9), not +\$3,964. Window and shape must travel with it: 11 tick-honest sessions, 165 legs, of which 21 runner events are the ENTIRE result — strip them and the configuration is +\$380.70 over 144 legs (\$2.64/leg, a rounding error above the fee). 77% of the gain (+\$4,708 of +\$6,081) was already available under the pre-08-01 config, so this is a revert to a previously-shipped cell, not a new optimisation. △ THE LIVE FILLS DO NOT SUPPORT IT (REV2 · Q2): all 6 post-card grind legs fired at ATR14 ≥ 22 (23.98, 23.98, 33.50) — the reverted floor would not have blocked one of them, so –\$172.00 is not evidence for this play. The	Arms: ATR1m ≥ 22 at signal; below 22 the gate does not arm at all. No extension-ceiling override — the gate default (4.0) applies. △ OPERATIONAL: <code>grind_long</code> is NOT in <code>reactivate_gates.py::HOLD</code> (which holds only <code>nipc_long/nipc_short</code>), so <code>gazbot7-gate-reactivate.timer</code> flips it back on at 22:00 UTC Sunday regardless of what this report concludes. 'It stays off' is not achievable by doing nothing — either ship the config change and let it re-arm into a gate that self-gates at ATR 22, or add <code>grind_long</code> to HOLD. There is no third state. Exit: Unchanged — the deployed two-lot ladder (Lot A fixed-R scalp, Lot B chandelier), plus the quiet-tape clip below ATR 22 (which the new floor makes almost unreachable). Kill: ★ REV2 CORRECTION — write the review criterion before	Garrath — two literals in <code>deciders.py</code> + tournament restart Set <code>ATR_FLOOR['grind_long']</code> back to 10.0 and re-add "ext_hi": 2.0 to the <code>SlotSpec</code> . One line each; the gate returns to exactly this week's behaviour.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
		case rests entirely on the 11-session replay leg-count (744 → 155). Part 2.5's –\$958 is a different population answering a different question; never quote it in the same sentence.	the first bad week, because P&L alone cannot judge this: 20 admitted legs MINIMUM, and judge on whether 4R events appear at anything near the expected 12.7% rate. A fortnight with no expansion days will produce a flat result that is NOT evidence the change was wrong. Hard kill: fewer than 15 of 19 target-winners retained on fresh tape, or a negative \$/leg over 40 live legs.	
3	★ Build the inverse auditor: does every live STOP have a slot? LIVE Run the orphan sweep on a TIMER, not only at boot. The desk audits 'does every slot have a live stop?' and never asks the inverse — and it is the inverse that fires unattended. Mechanism: order/position reconciliation across the shared account Why: Same class as the MD_STREAM bug: a SHARED resource consumed without checking the tag that says whose it is. There the stream was multi-symbol; here the account is multi-desk. Evidence: Part 1.5 §5 · verification: Plant an orphan stop in a test and confirm the timer alarms on it. · mode: observe-only then alarm	Two orphaned GTC stops opened a naked 2-lot short on 08-06 and it was booked to a gate that had been benched 38 minutes earlier. Separately, the two poisoned stops from 08-04 rested ~2,000pt away for three days. Until it exists, any gate's ledger can be contaminated by another gate's garbage — the nipc rehab would have started from a 6% hit rate on 16 trades instead of 0% on 14 if the order-id namespace had not been checked.	Arms: N/A — safety infrastructure. Exit: N/A. Kill: None. This is the open gap from the 08-06 incident and two separate rehabs found it independently this week.	Claude — build, then live-verify N/A — it is a read-only auditor plus an alarm.
4	MAX_HOLD: keep the 120 minutes, add a ceiling that survives a dead loop LIVE Do not shorten MAX_HOLD — every attempt to is a fake win. Add the ceiling that still fires when the main loop has stopped cycling. Mechanism: MAX_HOLD — the 2-hour time-cap exit Why: MAX_HOLD looked like the worst exit on the desk and was the fire alarm being blamed for the fire. The real defect is that it can be skipped when the loop wedges. Evidence: Part 1.5 §1 · Part 1.5 §5 · verification: Kill the loop in a test and confirm the ceiling still flattens. · mode: live-paper	+\$461.50 on the archive. △ MAX_HOLD has four fires in 21 days and every per-regime 'optimum' in its sweep rests on one to four bound trades — the ceiling is justified by the failure mode, not by the P&L.	Arms: Always on. It is a safety cap, not a strategy. Exit: 120 minutes, unchanged, plus an independent ceiling that does not depend on the main loop being alive. Kill: None — this is a safety net. If the ceiling ever fires spuriously, fix the ceiling, do not remove the cap.	Garrath — code + restart Remove the ceiling. Do NOT touch the 120 minutes: that cap is what limited the MD_STREAM incident to –\$255.50.
5	DAY RIDER: write the trade rows, and pass desk= at the same time LIVE Add a record_trade call on entry and on exit. In the SAME change, pass desk= in web.py and core.py so the day-rider's P&L stops blending into the tournament's. Mechanism: day_rider persistence + P&L attribution Why: The two things protecting us today are accidents: the kill switches are set to 0, and the day-rider writes no rows. The moment either changes, this becomes live. Fix desk= at the same time as the write path, not after it. Evidence: Part 2.7 DR7 #2 and #3 · verification: Confirm rows appear with the right desk tag and that pnl.day() separates the books. · mode: live-paper	The strategy has been live for 3 days and has produced 0 trade rows; we are flying on a journal file that gets overwritten every session. 3 of 4 P&L surfaces blend the two books — including core.py:406, the daily-loss kill switch. A day-rider loss counts against the tournament's headroom, so one bad afternoon on a strategy with a \$1,059 daily standard deviation could halt eight gates that did nothing wrong.	Arms: N/A — infrastructure. Exit: N/A. Kill: None. Without this the strategy is unmeasurable.	Garrath — code + restart N/A — it is additive instrumentation.
6	DAY RIDER: cut to 1 lot until 56 recorded sessions LIVE Halve the size. It costs no information and it is the only thing standing between a bad afternoon and a hole in the desk. Mechanism: day_rider position sizing Why: 38 sessions and one real live trade is not a validation. Size for the sample you have. Evidence: Part 2.7 DR7 #4 · verification: 56 recorded sessions. · mode: live-paper	A day-rider lot carries 15× the swing of a tournament lot; its worst day is 64% of the desk's all-time loss. Halving size costs no information.	Arms: Same detection as above; only the size changes. Exit: Unchanged. Kill: Restore 2 lots only after 56 recorded sessions at the 0.25 floor come back near +\$276/ trade.	Garrath — day_rider config Set it back to 2. One number.
7	★ Commit the working tree — five live behaviours exist only as uncommitted edits LIVE Commit every uncommitted change that is currently live, then make `exit_overrides_uncommitted: true` FAIL the sweep instead of being journalled and ignored. The flag already exists and is already written at every desk startup; nothing has ever acted on it. Mechanism: config integrity / reproducibility Why: An uncommitted live behaviour is a behaviour that cannot be reverted, diffed or reproduced. The desk already journals a flag that	Five live behaviours exist only as uncommitted working-tree edits. `config_journal.jsonl` flags `exit_overrides_uncommitted: true` at every startup and nothing consumes it. Reconstructing a config epoch by hand is what made 07-31–08-02 unrecoverable, and the REV2 audit of last week's card took hours of archaeology	Arms: N/A — infrastructure. Exit: N/A. Kill: None — this is reproducibility, not a hypothesis. If the sweep starts failing on it, that is the mechanism working.	Garrath — git commit, then one clause in sweep.py N/A. Committing cannot change live behaviour; it only records it.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	says this is happening and no instrument reads it — the same shape as `signal_journal.suppressed_by` being NULL in all 596 rows. Evidence: Part 3.3 (REV2 · Q2) §5 — 'The Saturday build, in priority order' #2 · verification: git status clean on the desk tree; sweep goes red if it is not. · mode: infrastructure	that a commit history would have made a query.		
8	★ The play ledger — status / shipped_ts / outcome on every play, plus a Monday grader LIVE The three fields are now written on all 52 rows of plays.json (defaults PROPOSED / null / null). Two things remain: whoever applies a change sets `status` and `shipped_ts` at the time they apply it, and a Monday job fills `outcome` from the trade book on the gate each play names. Mechanism: plays.json ledger + a scheduled grader Why: Three fields turn 'did we do it, and did it work?' from hours of archaeology into a query. It is the cheapest thing on this card and it compounds every week. Evidence: Part 3.3 (REV2 · Q2) §1 and §5 — 'The Saturday build, in priority order' #1 · verification: Next Friday: the card grades itself with no hand-derivation. · mode: infrastructure	Last week's 24-play card had no outcome field, so nothing graded it and nothing grades the new one either. The REV2 audit had to hand-derive all 24 grades: 15 shipped in some form, 4 never built, 3 overtaken, 1 partly done, 1 flatly violated. The card claimed £3,150/wk; the five sessions that followed it came in at −\$772.50 on 110 trades. The single most expensive line was old #19 — the card said SHADOW news-impulse-pullback and do NOT arm it, it was armed live the next day, and it is −\$434.50 over 49 trades at 24.5% win.	Arms: N/A — infrastructure. Exit: N/A. Kill: None. Without it, next Friday repeats this week's archaeology.	Claude — the Monday grader; Garrath — set status/shipped_ts when applying a change Drop the three fields. Nothing reads them yet except the grader.
9	DAY RIDER: raise the efficiency floor 0.15 → 0.25 LIVE One constant in drift.py. This is the only day-rider change this week that touches the money. Mechanism: day_rider — drift detection from the 13:30 UTC cash open Why: Not one of 114 exit rules touches the money on this strategy; the seven days that lost \$1,000+ lost −\$11,195 between them and the trail we shipped saves \$0.00 on every one of them, because price never went our way at all. The entry is the whole lever. Evidence: Part 2.7 (day rider) DR7 #1 · verification: 56 recorded sessions at 1 lot. If +\$276/trade holds anywhere near that it clears t=2. · mode: live-paper	+ \$5,410 → + \$8,819 over the same 38 sessions; + \$142 → + \$276 a trade; plateau across 0.18-0.35; positive in 32 of 32 leave-one-day-out folds; consistent in and out of sample. It cuts time-to-answer from 10.5 months to 2.7.	Arms: efficiency ≥ 0.25 (was 0.15) AND roundtrip ≥ 0.45, measured from the cash open. One entry per session, no entry after 15:00 UTC. Exit: Unchanged — the armed trail, hard flat 20:40 UTC. Do not tune it further (see HOLD). Kill: 56 recorded sessions at the 0.25 floor, at 1 lot, with trade rows actually written. If \$/trade drifts toward zero over that stretch it was three good days in June and July.	Garrath — one constant in drift.py Put 0.15 back. One number.
10	Lot B rung: widen it to a fixed 6R and move the quiet-tape clip's ATR gate 22 → 24 LIVE ★ REV2 CORRECTION (REV2 · FIX0) — the old card ('regime-key the chandelier width') named a mechanism this play does not touch and could not be typed into any file the desk reads. Replaced in full. Edit FIVE of the six non-NIPC rows of data/exit_overrides.json to {"a_r": 3.0, "b": 6.0, "atr_split": 24, "lo": {"a_usd": 40, "b_r": 1.75, "b_floor_usd": 60}}. NIPC keeps its proven {"a_r": 2.0, "b": 2.5} pair. ★ OPERATOR CALL 2026-08-08 — grind_long is EXEMPT from the split move: it takes "a_r": 3.0 and "b": 6.0 like the rest, but KEEPS "atr_split": 22. Three numbers per row, no code change, no new mechanism. Mechanism: Lot B fixed-R rung + the quiet-tape clip's ATR threshold. NOT the chandelier — this edit resolves every sub-slot to exit="scalp" and leaves ZERO chandeliers in the slate. Why: The chandelier looked like a threshold that had stopped being reached. It had actually been switched off in code. Regime-keying is what it should have had all along. Evidence: Part 3.4 (REV2 · FIX0) — replaces Part 2.6 §8 Saturday #3 and the Part 1.5 chandelier lead row · verification: ★ FIRST: re-run FIX0's 20-session archive with grind_long held at atr_split 22 and report the honest delta — the +\$3,024 no longer describes the deployed config. 20 more sessions; re-run the LODO fold and confirm the minimum stays positive. FOLLOW-UP, NOT THIS PLAY: er_split — the ER30 key, a code change on three lines, better on every axis (archive +\$25, week +\$156, LODO floor +\$345, 4/4 vs 2/4 ISO weeks) but margins are small and it needs a deploy. QUEUED, not shipped, and its AND/OR spec must be resolved first (FIX0 §3). · mode: live-paper	Config-only, as shipped: archive — \$2,932 → +\$92 (+\$3,024) over 20 sessions, 20/20 LODO folds positive, min +\$1,259; pre-window −\$1,859 → +\$891; symptom week −\$1,073 → −\$799. ⚠ Better in only 2 of 4 ISO weeks (W31 +4,557 vs +1,683 and W32 −799 vs −1,073 win; W29 −954 vs −910 and W30 −2,712 vs −2,632 lose). NOT strictly dominant. ⚠ ★ OPERATOR CALL 2026-08-08 — the +\$3,024 was measured with ALL SIX rows at 24. With grind exempted the shipped config is NOT the tested config and the direction of that difference is UNMEASURED. This is a judgement that protects a mechanism we understand (runners carry grind) over a cell we measured once — not a free lunch. Do not quote +\$3,024 as the expected value of what is deployed until the archive is re-run.	Arms: ★ OPERATOR CALL 2026-08-08 — grind_long keeps atr_split 22; the other five take 24. WHY: SATURDAY #2 (grind-long-revert-atr22-ext30-0808) sets grind's entry ATR floor to 22, and REV2 Q1 §10 relied on the clip being INERT for the gate at that floor ("leave it in as a fail-safe"). Moving the split to 24 un-inerts it for the ATR 22-24 band — 19.3% of grind-eligible time on this week's tape — and the clip caps Lot B at 1.75R floored \$60, which truncates exactly the 21 runner events that ARE grind's + \$4,151 case. The two plays would have partially cancelled. Entry ATR ≥ 24 only, frozen at entry (slot_strategy.py:422). Below 24 the \$40 / 1.75R · \$60 clip is unchanged — it is rank 1 of 28 in two of the four regime cells and 100% of quiet-tape behaviour is preserved. There is no ER term. Exit selector only; arms and benches nothing. Exit: Lot B: fixed 6.0R rung above the split. Lot A: changes on all six gates to 3.0R — abs_veto_long 1.0→3.0, exhaustion_short 0.75→3.0, capitulation_long / rgv_short / abs_veto_short 1.5→3.0, grind_long 2.5→3.0. Kill: If the ATR ≥ 24 population goes negative over 20 fresh sessions, or the LODO minimum drops below zero on the next recut, revert. ⚠ READ THIS FIRST:	Garrath — six rows of data/exit_overrides.json, then systemctl restart gazbot7-tournament (overrides are read at slate build; confirm the live slate via scaleout_slots(), never source) git checkout data/exit_overrides.json + restart. Partial revert of just the exemption: set grind_long's atr_split to 24 to match the other five. git checkout data/exit_overrides.json + restart. The file is git-tracked and every startup journals its resolved ladder to config_journal.jsonl (scripts/config_at.py --diff).

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
			the parameter-plateau test FAILED (4 of 35 threshold cells positive), 6 of 16 in-cell sessions are green, two days (07-27, 07-23) carry two-thirds of the edge, and strip-best inverts at top-8 (~\$422). The DIRECTION is robust; the KNOB is not. Size accordingly.	

MONDAY · AT THE DESK — SWITCH-FILE AND EXIT-CELL CHANGES, ALL REVERSIBLE · 6 PLAYS

No code deploy in this window at all. Every row is `gate_switches.env`, `exit_overrides.json` or a router constant, and comes back in five minutes. **#1 needs nothing but a decision** — the router already owns the lever.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
1	★ Sit out the chop blocks — this is the largest single number in the report LIVE No new mechanism. The router already owns the lever: when the tape is in a chop block, bench the book. This is the biggest number anywhere in this report and it needs no code. ★ OPERATOR CALL 2026-08-08 — CARVE-OUT: <code>abs_veto_short</code> is EXEMPT from the chop bench for the duration of its 15-fire review (SATURDAY #1, <code>absveto-short-arm-by-default-er035-0808</code>). It is the only gate on this card with per-regime evidence that positively contradicts a blanket bench — + \$2,290.50 over 159 fires across 17 days, POSITIVE IN ALL FIVE REGIMES, chop included, with 53 of 54 filters tested losing to just letting it fire. Benching it in chop would cost money AND corrupt the experiment, because the 15 fires you review would be drawn from a biased subset and you would conclude something about arm-by-default that is really about the chop filter. Mechanism: the router's existing bench, applied to entry-regime chop Why: Three independent sections landed on the same conclusion this week: the desk does not lose money on trend days, it loses money on high-volatility days that go nowhere. Wednesday and Thursday — the two days ATR doubled from 12 to 20 and the tape went nowhere — took −\$1,122 out of the desk between them. Evidence: Movement 3 (master ledger, chop lab) · Part 1 (day-by-day) · Part 2.5 · verification: Grade the chop bench on the non-carved-out book only, so <code>abs_veto_short</code>'s 15 fires stay a clean experiment. Attribute next week's trades to entry regime and re-cut the same table. · mode: live-paper	Live trades attributed to entry regime over 17 days: chop 336 trades −\$5,483; non-chop 263 trades +\$1,482. Perfect avoidance turns −\$4,001 into +\$1,482. A perfect chop filter beats every invention in this report by an order of magnitude.	Arms: Bench the book when the tape is in a confirmed chop block — EXCEPT <code>abs_veto_short</code>, which stays armed through chop until its 15-fire review completes. Reversion gates keep their existing treatment. Inverse: it BENCHES when the entry regime is chop. Re-arm on a real break with ER climbing AND vol expanding — vol expansion is the binding leg. Exit: N/A — benching blocks new entries only. Kill: If a month of chop-benching costs more in missed non-chop trades than it saves, the attribution is wrong and this goes back to the lab. ★ OPERATOR CALL 2026-08-08 — separately: if the carve-out is what makes chop-benching look bad, that is a finding about <code>abs_veto_short</code>, not about chop. Grade the two independently — the carve-out exists precisely so they CAN be graded independently.	The durable router tick Reversible every 5 minutes.
2	★ Fix the midnight reopen: re-arm reversion gates only, make momentum earn its arm LIVE Change <code>reactivate_gates.py</code>'s default set so the 22:00 UTC reopen re-arms the reversion gates only. Momentum gates start benched and the router arms them on evidence. And move the 21:00–22:00Z maintenance gap so the hour before the reopen is supervised. Mechanism: <code>reactivate_gates.py</code> — the Paris-midnight all-on Why: A third of the router's entire work-rate is spent undoing a policy that fires once a night into the worst block on the desk. Evidence: Part 2.6 §8 Saturday #4 · Part 2.6 §7 (session blocks) · verification: Count switch changes in the 22:00–23:59Z block next week; it should collapse. · mode: live-paper	32% of the week's switch changes are spent undoing the 22:00 all-on; it cost −\$275 in one night on 07-31. The overnight block ran 8 trades this week for zero winners and −\$578, and 89 trades over 40 days for −\$906 at a 31% win rate.	Arms: At the 22:00 UTC reopen: reversion gates on, momentum gates off pending a real break with ER climbing AND vol expanding. Exit: N/A. Kill: If the reversion-only reopen costs a genuine overnight trend two weeks running, put the all-on back and solve it in the router instead.	Garrath / the durable router tick — <code>reactivate_gates.py</code> default set Restore the all-on default set. One list.
3	★ Untradeable meter: STAY-OUT cutoff 65 → 45, and no stay-out can trigger before 15:00Z LIVE Two changes to one rule. Drop the cutoff to 45, and add an explicit clause that a reading taken before 15:00 UTC never triggers a stay-out. Mechanism: the untradeable meter, as a stay-out trigger Why: The meter separates green from red days perfectly at end-of-day and is useless before the open. So use it at end-of-day and forbid it from speaking before 15:00Z.	At 45 the meter catches 4 of 6 red days and 0 of 4 green days, worth +\$1,138, with 22 points of headroom above the worst green day. The current 65 cutoff catches 2 of 6. Before the US open the meter correlates +0.13 with the day's P&L — noise — and it read	Arms: Meter ≥ 45 AND the clock is at or past 15:00 UTC → stay out. Before 15:00Z the meter is advisory only. Exit: N/A — it blocks new entries only; opens still exit and stops are untouched. Kill: In-sample on n=10 days. If it stays out of a green day twice, the cutoff goes back up.	The durable router tick One threshold, one clause. Reversible every 5 minutes.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	Evidence: Part 2.6 §8 Saturday #3 · verification: n=10 today. Re-grade at n=25 days before trusting the 45. · mode: live-paper	92 on a +\$638 day. Note the 45 deliberately declines the in-sample optimum of 30.		
4	Promote the vol-expansion caveat to a written rule LIVE Write it into the tick prompt: expanding ATR alone, in ER < 0.15 chop with no range break, is NOT a re-arm trigger. Rule text only, no code. Mechanism: router re-arm policy Why: Vol expansion is the binding leg for a re-arm, but only alongside ER and a break. On its own it is the most reliable way to get armed into chop. Evidence: Part 2.6 §8 Saturday #7 · verification: Count re-arms that decay inside 35 minutes next week. · mode: live-paper	It fooled the desk at least three times this week alone: 08-06 08:15 and 08-06 11:05 (both decayed inside 15 minutes) and 08-07 13:50 (armed on an ATR reading that was wrong by 15 points).	Arms: Re-arm requires a real range break WITH ER climbing AND vol expanding. Vol expansion on its own is not enough. Exit: N/A. Kill: None — it is a fail-safe rule that can only reduce arming.	Garrath — prompt text Delete the sentence. Fail-safe either way.
5	Grade the bad days in the bad-call ledger LIVE The ledger has zero entries for 08-03, 08-04 and 08-05. Backfill them, then keep grading every day, not just the ones that went well. Mechanism: router bad-call ledger discipline Why: A scoreboard that only scores the days somebody was watching is not a scoreboard. Evidence: Part 2.6 §8 Saturday #8 · verification: Ledger coverage should be 5 of 5 sessions next week. · mode: observe-only	30 switch changes and −\$947 of P&L are currently ungraded. The scoreboard samples the well-supervised days, which biases it favourably.	Arms: N/A. Exit: N/A. Kill: None — it is discipline, not code.	Claude — the 4-hourly ledger review N/A.
6	DEFERRED ONE WEEK · Router timing: WINDOW 30 → 45, HOLD 2 → 3, STEP → 5. Leave the thresholds alone. LIVE ★ OPERATOR CALL 2026-08-08 — HOLD until the chop filter has had a clean week. Change the timing only. ER_TREND stays 0.15 and NET_MIN stays 30 — both grids independently pick the incumbent and the surface around it is flat. Correct STEP to 5 to match the live tick. Mechanism: the direction router's regime-detection timing Why: The thresholds do not need touching. The timing does — and unusually for this desk it is an interior peak on four axes at once rather than the edge of a grid. Evidence: Part 2.6 §8 Saturday #2 · verification: Re-run after a two-clean-trend-day week; watch for late benches on aligned momentum. · mode: live-paper	+\$502 on the live-managed set and +\$399 on the full historical set. It is an interior peak on window, step, hold and ER — not a grid edge — and flat plateau on net. Survives LOO on both sets (+ \$574 / +\$1,040), and both universes rank every cell identically.	Arms: N/A — this changes how long the router looks back and how many marks it holds a decision for, not what it arms. Exit: N/A. Kill: Re-validate after the first week containing two or more clean-trend days. A stickier router is a chop optimisation. If it starts benching the aligned rider late, put HOLD back to 2 first.	The durable router tick — config only Put 30 and 2 back. Reversible on any 5-minute tick.

BUILD · NEXT — NOTHING GOES LIVE THIS WEEK · 19 PLAYS

Instrumentation, shadow slates and the method fixes that decide whether next Friday's numbers can be trusted at all. **#1 gates #4**, and gates every future study that concludes “exit earlier” — the shadow repricer leaks past its own stops on 44% of trades, which is exactly the bias that makes cutting early look good.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
1	★ Fix the shadow repricer's stop detection — then re-run everything that concluded 'exit earlier' LIVE The shadow book leaks past its own stops on 44% of trades. Fix the detection, then re-run every study that ever concluded 'exit earlier'. Mechanism: shadow repricer Why: This is a method fix that decides whether next Friday's exit numbers can be trusted at all. It sits in the same family as the standing rule that every shadow loss is a FLOOR. Evidence: Part 1.5 §5 · verification: Re-run the MFE-release rule and confirm shadow and live stop disagreeing. · mode: observe-only	Median overshoot 0.56 × ATR, maximum 32.32 × ATR, on a book that is 99% 1.0 × ATR stops. It is why the MFE-release rule looks brilliant in shadow and negative on live.	Arms: N/A — method fix. Exit: N/A. Kill: None. Until it lands, treat every shadow result that rewards cutting early as inflated.	Claude — repricer + re-run N/A.
2	★ THE OPEN RIDER → the shadow book. Not live. Not this month. SHADOW Add the OPEN RIDER to the shadow slate at its published spec, with the 14:45 cut applied and cadence-3 struck from the language. It does NOT go live on 17 in-sample days, and the skeptic panel voted 2-1, not 3-0.	36 days, n=236, +\$7,433 at the correct \$1.50/RT, 47.0% win, + \$31.50 a trade, +0.288R. The number that matters: on 19 days of TRUE out-of-sample V5 tape it does +\$3,526 / n=115 / +\$30.66 a trade — a 5% degradation on tape no parameter was ever fitted on.	Arms: 13:00-14:45 UTC only (right edge cut from 15:00). Five-minute cadence; re-entry blocked until the open trade closes. No threshold, no efficiency dial, no book reading. Exit: Stop 2.0 × ATR1m — twice the desk's house style, mean 54.6pt ≈ \$109 of risk. Target 2R.	Claude — shadow slate only Remove it from the shadow slate. It is explicitly never live this week.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	Mechanism: OPEN RIDER — no trigger at all. The clock is the trigger: every 5 minutes in 13:00-14:45 UTC, if flat, take a trade in the direction of the last 15 minutes' move. Why: Built as the control that was supposed to lose to the clever inventions, and it beat both of them. But one of three independent skeptics REFUTED it (see NOT-AN-ACTION), and two skeptics with different out-of-sample legs reached opposite verdicts on the same strategy. That is exactly why it is a shadow candidate and not a live one — and the honest way to resolve it is forward tape, not another backtest. Evidence: Movement 3 §4-5 and the master ledger · verification: 200 shadow trades including a non-trending open week. · mode: shadow	36/36 leave-one-day-out positive; reverse walk-forward picks the published cell #1 of 48 and all 48 cells are holdout-positive; placebo z=+2.59, p=0.005. Δ the placebo mean is NOT zero — the wide-stop/2R shape is worth something on its own.	Time cap 45 minutes, then out at market. Kill: PROMOTE ONLY IF all four: +0.15R over 200 shadow trades including a non-trending open week; the 14:45 cut in place; a live 2×ATR stop path (scaleout_slots() forces stop_atr_mult=1.0 today); and position isolation from the day-rider. Fail any one and it stays in shadow.	
3	Wide stops ($\geq 2 \times \text{ATR}_{1m}$) in the open window — and check the live gates, it costs nothing SHADOW Shadow the finding properly, and separately: check the desk's ~1-ATR house stop against the live gates in the open window. That check costs nothing and is the one thing in Movement 3 worth looking at this week. Mechanism: stop width, in the 13:00-14:45 UTC window Why: This is the most transferable thing in the whole greenfield movement, and it is about the gates we already own rather than a new one. Evidence: Movement 3 §4-5 and master ledger · verification: Re-measure on a signal with a non-zero intercept. · mode: shadow	Three independent votes: the survivor's own 45-cell grid, the robustness skeptic's from-scratch rebuild, and the oracle all agree — every cell at a 1-ATR stop is flat-to-negative, all 27 cells at 2.0 × ATR or wider are positive. Mean stop 57.7pt against a mean 1-min ATR of 28.8pt. Even a perfect entry gets shaken out of 14 of 52 runs at 1 ATR against 0 of 52 at 2 ATR.	Arms: Open window only. This is a stop-geometry finding, not an entry. Exit: Stop $\geq 2.0 \times \text{ATR}_{1m}$ instead of the house ~ 1.0. Kill: Revoke/confirm ONLY IF re-measured on a signal whose intercept is not zero — it has only ever been measured inside the OPEN RIDER, whose own intercept is $-\\$23$ ($t=-0.22$).	Claude — shadow + a live-gate check N/A — nothing ships.
4	exhaustion_short: put the exit back to the fixed 8pt / 12pt / 120s it had on 07-25 PARKED ★ BLOCKED on BUILD #1 (fix-shadow-repricer-stop-detection-0808). This is a tighter, earlier exit, and the repricer leaks past its own stops on 44% of trades (median overshoot $0.56 \times \text{ATR}$) — which is precisely the bias that inflates 'cut earlier' results. Re-derive it AFTER the repricer is fixed. The gate's ENTRY stays parked either way. Δ Note this dependency does NOT extend to the Lot B widening: that exits LATER, so the same bias runs against it, and its real objection is its own failed parameter plateau. Revert the exit to the fixed 8pt / 12pt / 120s ladder. Note this is an EXIT fix on a gate whose ENTRY is parked — do both, do not confuse them. Mechanism: exhaustion_short — roll-over fader Why: A two-year-old exit was replaced with a worse one. Reverting costs nothing and removes a confound from any future attempt to revive the gate. Evidence: Part 1.5 §1 · Part 1.5 §6 (microstructure sampling noise) · verification: Re-run at 10 polling phases and quote the mean, never a single run. · mode: live-paper (exit only — entry stays parked)	+\$604 over 6 dense-book days on identical entries, taken as the mean of 10 polling phases. All six fixed-point exits beat the incumbent. Δ the standard deviation across those ten polling phases is \$418 on a $-\\$880$ mean — anyone quoting a single run of that harness is quoting noise.	Arms: PARKED as an entry. The gate is sub-friction by about two dollars a lot; it does not get armed on this evidence. The exit revert is so that when it IS revived it is not revived onto a worse exit. Exit: Fixed 8pt / 12pt / 120s — the 07-25 ladder. Kill: Revival condition: the gate must clear friction — roughly +\$2/ lot of headroom — on a fresh sample before it is armed again. Until then the exit revert ships and the entry stays off.	Garrath — exit_overrides.js n + tournament restart Delete the key.
5	★ The binding constraint on flow research is RETENTION, not ideas LIVE Keep archiving ticks. This is an infrastructure action, not a research one — and it is the single highest-leverage thing on this list for next month's research. Mechanism: tick archive retention (capture.db 5 trading days → Parquet → B2) Why: Growth you notice; a silent gap in the tape you do not. Evidence: Movement 3 §2 and the take-aways · verification: Row-count verification on every mirrored day. · mode: observe-only	Five days of aggressor-tagged ticks caps every flow finding this desk will ever have at $n \approx 40$. Three more weeks would take FLOW-TRAP FADE from $n=44$ to $n \approx 180$ and turn 'not obviously noise' into an answer. VACUUM work needs ~15 sessions against the 5 we have. And the bar-derived flow proxy is REFUTED — aggressor sign is genuinely absent from OHLCV, so there is no way to buy the missing days back.	Arms: N/A. Exit: N/A. Kill: None.	Nobody — it is already running. Do not let the mirror stop. N/A. ★ Remember the interlock: prune_capture.py may not delete a day tape_mirror.py has not exported AND verified by row count.
6	shadow.db needs the data_quality column gazbot7.db already has LIVE Add it, and backfill the tape-consistency screen's strikes. Mechanism: shadow book data hygiene Why: The trade book learned this lesson in August. The shadow book has not. Evidence: Part 1.5 §5 · verification: Re-run the screen and confirm the 166 rows are now flagged and excluded. · mode: observe-only	The tape-consistency screen struck 166 of 7,639 shadow rows across 14 separate days. Nothing flags them. MD contamination in the shadow book is a 14-day condition, not one bad night.	Arms: N/A. Exit: N/A. Kill: None.	Claude — schema + backfill N/A.
7	Populate signal_journal.suppressed_by — it is NULL in all 596 rows	NULL in all 596 rows. Everything in this week's dossiers about what	Arms: N/A. Exit: N/A.	Claude — one write path

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	LIVE Write the field. Until it is populated, every statement about which gate fires were suppressed and why is unsupported by the store. Mechanism: signal journal Why: It is the difference between a reconstruction and a measurement. Evidence: Part 1.5 §5 · verification: Confirm non-NULL rows appear and that they agree with a tape reconstruction. · mode: observe-only	a gate would have done unrouted is RECONSTRUCTED from the tape, not read from a log.	Kill: None.	N/A.
8	Union the two tick sources, and make an empty tick-day RAISE LIVE Neither tick source covers the window on its own. Union them, and make a harness that finds zero ticks for a day raise instead of booking zero trades. Mechanism: tick data access layer Why: A silent zero is the most expensive kind of bug this desk has: it looks like a result. Evidence: Part 1.5 §5 · verification: Point a harness at a day with no ticks and confirm it raises. · mode: observe-only	data/tape/ticks/MNQ/ is missing 2026-08-07 while capture.db has it; capture.db's ticks only reach back to 08-03. A harness reading only the archive silently loses the most recent session and books zero trades with no error — which is exactly what the shipped nipc replay is doing right now.	Arms: N/A. Exit: N/A. Kill: None.	Claude — lake layer + harness guard N/A.
9	Stop-limit orders are filling AT their limit, 20pt through their own trigger LIVE Narrow the limit buffer to ~5 points or use a plain stop-market. Add a fill-plausibility check: if a stop fills more than N points through its trigger, alarm and reprice it against the tape before it is booked. Mechanism: stop order type + fill validation Why: We caught this by hand. Nothing in the desk would have. Evidence: Part 1.5 §5 · verification: Zero fills more than N points through trigger over a month. · mode: observe-only then alarm	2 lots this week, both the second lot of a pair, both within four minutes of the cash open. The other 54 stop fills have a median slippage of 0.00pt. One filled 10.25pt above the highest print in the window. Roughly \$70/week of fake losses. Found independently by two rehabs.	Arms: N/A. Exit: This IS the exit path — it is the stop itself. Kill: None.	Claude — order type + validation Restore the wider buffer.
10	On reconstruct(), re-derive the entry ATR from the tape LIVE Refuse to adopt a slot whose persisted stop is more than a few ATR from entry. Re-derive the ATR from the tape instead of trusting the persisted value. Mechanism: slot reconstruction after a restart Why: A poisoned number survived a restart. Fixing the producer did not fix the consumer. Evidence: Part 1.5 §5 · verification: Restart with a poisoned persisted ATR and confirm the slot is refused. · mode: observe-only	The 08-04 gold-bar leak was fixed in five minutes. Its consequence survived, because reconstruct() re-adopted two open slots with the 1,848pt ATR persisted — leaving the stop naked, both targets unreachable and the quiet clip disabled, for 115 minutes. Cost: \$192.50 in one night.	Arms: N/A. Exit: Protects the stop and target geometry of an adopted slot. Kill: None.	Claude — reconstruct() guard N/A.
11	One log line: the desk cannot tell you why an approved entry never became an order LIVE Add the log line to all three silent return paths. Mechanism: entry → order path Why: Three silent return paths is three ways to lose a day and never know. Evidence: Part 1.5 §5 · verification: Force each path in a test and confirm it logs. · mode: observe-only	On 08-07 at 14:19:36 both lots logged 'confirmed → OPEN' and then nothing at all — no order, no signals row, no error. Three silent return paths, none of which logs. (For the record it saved us \$174 that day, and that has NOT been credited back — you cannot bank a loss you avoided by accident.)	Arms: N/A. Exit: N/A. Kill: None.	Claude — one log line N/A.
12	Fix the nightly's fader repricing — stop scoring blocked fader signals at a fixed scalp-2R LIVE Use the gate's own exit when repricing a blocked fader signal, not a fixed scalp-2R. Mechanism: the nightly router-value repricer Why: The router's own scoreboard is currently marking its best decisions as mistakes. Evidence: Part 2.6 §8 Saturday #5 · verification: Re-run and confirm the nightly and the shadow mirror stop contradicting each other. · mode: observe-only	The nightly says benching capitulation_long cost +\$368 of winners; the shadow mirror says it avoided −\$1,609. Both cannot be right and the fixed-scalp assumption is the suspect. The fader half of the router's '+\$635 of missed winners' is probably the same artefact.	Arms: N/A. Exit: N/A — it is a measurement of exits, not an exit. Kill: None — diagnostic only, no live behaviour changes.	Claude — nightly repricer N/A.
13	abs_veto_short: keep the absorption test, shrink the 55-second wait SHADOW A specific, testable job: the veto is doing its job, the DELAY attached to it is not. Sweep the wait down and re-derive. Mechanism: abs_veto_short — the 55-second confirm clock, separate from the absorption test Why: It lines up with the single strongest finding in Movement 3: every minute of confirmation you wait for costs about \$600 on run tape.	On run tape the wait costs \$596.50 against the \$169 the test saves. The absorption test blocked three raw fires, all three losers, saving \$169. On the two thrusts that survived the test, waiting 55 seconds turned +\$97 into −\$71 and +\$282.50 into −\$146 — same signal, same bar, same exit, only the entry clock	Arms: Same as the gate. This changes only how long it waits after the test passes. Exit: Unchanged. Kill: If a shorter wait does not beat 55 seconds on a population that is not the run tape, leave it at 55. Δ this measurement is on n=2 surviving thrusts — it is a direction, not a measurement.	Claude — sweep then propose N/A — nothing ships this week.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	Evidence: Movement 2 (lead) · Movement 3 (oracle: lateness) · verification: Sweep the wait on a non-run population before proposing anything. · mode: shadow	differs. Net effect on run tape: – \$427.50.		
14	The quiet-tape clip is a chop tool doing chop work in the wrong place SHADOW Design and backtest a regime carve-out so the clip stops firing on run tape. Do not just delete it — it was proven on the whole book. Mechanism: quiet-tape clip (ATR < 22 → Lot A banks \$40, Lot B banks \$60 / 1.75R) Why: It is the right tool in the wrong regime, which is the same shape as the chandelier finding — and the chandelier's answer was to regime-key it, not remove it. Evidence: Movement 2 (lead) · verification: Score the carve-out on the whole book, not the run subset. · mode: shadow	Live on 35 of the 53 testable runs and cost \$1,202 on run tape. Identical entries: clip ON returns +\$1,696.50, clip OFF returns + \$2,898.50.	Arms: Frozen at entry on ATR < 22. The carve-out would add a run/no-run condition on top. Exit: This IS an exit rule. Kill: If the carve-out does not beat the blanket clip on the whole book — not just on run tape — it does not ship. Judging it only on the runs is judging it on exactly the trades it was designed to skip.	Claude — backtest, then a Saturday N/A — nothing ships this week.
15	nipc_short → SHADOW. It is direction-beta wearing a state machine. SHADOW Move it to the shadow slate. It is not a wrong-rung problem, it is a direction bet. Mechanism: nipc_short Why: Re-arming it is a fresh operator decision, not a router decision. Evidence: Part 1.5 §1 · CLAUDE.md 08-06 note · verification: A flat or down week. · mode: shadow	Every train-optimum in the nipc dossier is killed by its own out-of-sample leg. It looked like a wrong-rung problem and is a direction bet.	Arms: Not armed. Shadow only. Exit: Shadow default. Kill: Revive only if it clears its null on a sample where the direction bet is not doing the work — i.e. a flat or down week. Note it already hit the operator's –\$400 kill criterion on 08-06 (n=47, 26% win) and sits in reactivate_gates.py::HOLD.	Claude — shadow slate Remove from the shadow slate.
16	FLOW-BREAK IGNITION and FLOW-TRAP FADE → shadow x2 (check FADE for redundancy first) SHADOW Shadow both. Δ check FLOW-TRAP FADE against our existing abs_veto absorption work BEFORE shadowing what may be a duplicate. Mechanism: FLOW-BREAK IGNITION = sustained flow + a 30-minute range break. FLOW-TRAP FADE = aggressors lean hard and fail to move price. Why: Neither answers the assignment they came from, which is exactly why they are shadow and not live. Evidence: Movement 3 §3 and master ledger · verification: 3 more weeks of ticks. · mode: shadow	IGNITION +\$1,967 on 39 trades, 56.4% win, 45 of 45 sweep cells green — but its out-of-sample half collapses to a coin flip. FADE + \$2,024 on 44 trades, all five days green, leave-one-day-out flat as a board, and its OOS leg holds — the sturdier of the two. Both catch 0 of the 4 runs they were built for. Both are 5-day findings wearing a confident number.	Arms: Both are flow-conditioned and therefore capped at 5 days of aggressor-tagged ticks — which is the real constraint on both. Exit: Shadow defaults. Kill: PROMOTE IF: IGNITION keeps US-session \$/trade above –\$40 at $n \geq 50$ over 3 more weeks of ticks AND its out-of-sample leg stops being a coin flip. FADE holds $\geq \$25/\text{trade}$ over $n \geq 80$.	Claude — shadow slate Remove from the shadow slate.
17	DEAD-CHOP fade, 20:00-21:00 UTC → shadow SHADOW Shadow it. It should never be promoted on this sample. Mechanism: DEAD-CHOP fade — fade the band in the dead hour before the CME halt Why: Thin n is a shadow, never a kill — and never a promotion either. Evidence: Movement 3 (chop lab, master ledger) · verification: 6-8 weeks to n ≥ 60. · mode: shadow	Race n=159, P=0.610, z=+2.78; all 20 costed cells positive. But only n=10 actual trades.	Arms: 20:00-21:00 UTC only, ATR < 10. Exit: As specified in the lab's 20 costed cells. Kill: Revival requires ALL THREE: (1) $n \geq 60$ costed shadow trades, roughly 6-8 more weeks at ~1 trade a day; (2) the race stays at $P \geq 0.57$ with $z \geq 2$ on the FORWARD sample only; (3) the 20-cell plateau still has $\geq 80\%$ of cells positive and the ATR < 10 gate still separates. Clear all three and it earns a 1-lot promotion inside the 20:00–21:00 window only.	Claude — shadow slate Remove from the shadow slate.
18	grind_long's ext_hi ceiling — SUPERSEDED by Part 3.2 (REV2 · Q1). Kept on the card so the reasoning is not re-proposed. REFUTED ★ REV2 CORRECTION — SUPERSEDED. This row said 'Do NOT remove it on this evidence', and on the Movement-2 sat-out-run population that was right: a ceiling judged on exactly the trades it was designed to skip cannot be judged there. REV2 · Q1 §10 re-derived it on the tick-honest book and on the 128 live router-gated lots and removes it on BOTH — everything from 3.0 up scores identically, the live 128 agrees more strongly (the ceiling deletes 12 real lots worth +\$389 at 50% win). Do the SATURDAY #2 (grind-long-revert-atr22-ext30-0808) deletion; do not also run this re-derivation. Mechanism: grind_long extension ceiling ('don't buy what is already stretched') Why: It is the single filter standing between the desk and these runs — but that is an argument for measuring it properly, not for deleting it.	Ungated, grind fired at 7 of the 22 up-runs for +\$742.50; mechanically it fired once, for – \$130. The ceiling blocked six of its seven pre-run fires.	Arms: Today: ext ≤ 2.0 live, going to 3.0 on Saturday. The study would make it conditional on regime rather than a single number. Exit: Unchanged. Kill: It stays as-is unless a regime-conditional version beats the flat ceiling on the WHOLE book. On a sat-out-run population it is being judged on exactly the trades it was designed to skip.	Claude — regime-conditional backtest N/A — nothing ships this week beyond the Saturday 2.0 → 3.0 revert.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	Evidence: Movement 2 §4 — superseded by Part 3.2 (REV2 · Q1 §10) · verification: Whole-book scoring with an out-of-sample leg. · mode: shadow			
19	The census cluster labels are reporting noise as structure — require persist ≥ 3 REFUTED (as written) Either require persist ≥ 3 for a FLOW-LED / VACUUM label, or retire both back into one 'flow event' bucket. Mechanism: Movement 1 census cluster labelling Why: Without it, next Friday's agent digs the same empty hole — two hunts already died in it this week. Evidence: Movement 3 §3 and master ledger (census) · verification: Re-label the census and confirm the two clusters stop being two names for one coin. · mode: observe-only	At the target runs the flow's run-length is one or two minutes and it reverses on the next bar. VACUUM fires on 1,451 minutes to find 8 runs — 0.55%, roughly 300 fires a day to find one. On 08-04 01:40 the flow is NEGATIVE at the start of a +128pt UP run that the census called FLOW-LED.	Arms: N/A — a labelling fix. Exit: N/A. Kill: The labels as written are refuted: one minute of $z \geq 1$ assigns the label, that minute reverses on the next bar, and an independent re-derivation disagrees with the label 3 times out of 4.	Claude — two scripts Git revert.

HOLD · LEAVE ALONE — THESE ARE ALREADY RIGHT · 5 PLAYS

Standing lessons and things already in production. The action is to not touch them.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
1	abs_veto_long — leave it alone this week. 'Do not touch its exit' over-claims; two leads are parked, not absent. LIVE ★ REV2 CORRECTION (REV2 · FIX1 §E) — No config change this week. Do not add an ER floor — the long twin's ER profile is the mirror of the short's. Two leads that earlier drafts proposed are not being deployed and are dispositioned in Part 2 §11: a US-only session guard (REFUTED — it would bench a positive LONDON book) and a wider Lot B (REFUTED as stated; a narrower US-chop-only tightChand x2 variant is SHADOW at n=13). Mechanism: abs_veto_long — thrust-continuation with a 55-second confirm Why: It is the one number in Part 1 worth putting weight on. Everything else on the desk this week is thin, one-week, or explained by a software defect. Evidence: Part 3.5 (REV2 · FIX1) — resolves HOLD #1 against Part 2 §4, §5, §9, §11 · verification: Standing. · mode: live-paper	Sixteen entries, thirty-two lots, + \$566.00 at a 62% strike rate, average winner +\$47.60 against an average loser of −\$32.20. Still positive if you delete ANY single day (+\$235.50 / +\$635.00 / + \$582.00 / +\$245.50) and still + \$363 if you delete its three best lots. ★ REV2 CORRECTION — the '− \$258.50 overnight problem' is an ASIA problem and Asia is already blocked at config level; excluding it the non-US book is +\$0.59/lot all-time. The earlier '−\$0.90 above ER 0.35' figure was a live-book sampling artefact.	Arms: Armed. It only fires when a burst is still going 55 seconds later, which is exactly the behaviour this week rewarded. Do NOT give it an ER30 floor — the long twin is positive in ALL four ER bands (+\$4.04 / +\$9.40 / + \$10.69 / +\$7.77 per trade). ★ REV2 CORRECTION — the earlier card said 'the floor abs_veto_short is getting'; that floor was WITHDRAWN (Part 3.1 / Part 3.6), so neither twin gets one. Exit: Lot A at 1.0R, Lot B at 1.5R, quiet-tape clip on both under ATR 22. Kill: If it goes negative on a month of fresh tape. Nothing shorter.	Nobody — this is a confirmation N/A.
2	★ Do NOT build the event-driven, zero-lag, per-gate regime guard REFUTED Do not build it. This is the most useful negative result in the router review and it is the exact thing that felt most obviously right. Mechanism: an in-code, per-gate regime guard evaluated at entry time instead of a 5-minute poll Why: Kept on the card with its autopsy so it is not re-proposed next quarter. Evidence: Part 2.6 (the one thing that did not survive) · §8 · verification: N/A. · mode: none	Zero trades opened in the five minutes before any bench decision all week — there is no reaction-lag leak to fix, and the theoretical ceiling on chasing it is ~\$300. The hysteresis that makes the poll 'slow' is the thing earning the money.	Arms: N/A — it would replace the poll. Exit: N/A. Kill: Backtested and refuted: + \$304 against the poll's +\$682 on this week, +\$267 against +\$475 over 40 days, and it fails leave-one-day-out.	Nobody N/A.
3	Do not touch ER_TREND, NET_MIN, or the day-bias ±40 LIVE Leave all three. Both grids independently pick the incumbent cell and the surface around it is flat. Mechanism: router thresholds Why: A flat surface is a finding. Tuning into it is how you buy noise. Evidence: Part 2.6 §8 (what I would not do) · verification: N/A. · mode: none	ER 0.15 / net 30 came out top on BOTH test sets. On the day-bias there is no optimum at all: counter-day-bias trades lose \$16–18 each at every threshold from 0 to 150, and raising it makes the aligned book worse.	Arms: ER_TREND 0.15, NET_MIN 30, day-bias ±40 — unchanged. Exit: N/A. Kill: None — the evidence is that there is nothing to find here.	Nobody N/A.
4	DAY RIDER: leave the trail alone, add no stop below 600pt, and stop tuning the exit LIVE Three non-actions on one strategy. The trail's shipped rationale is wrong but the rule is harmless and mildly positive — changing it again on a noise surface is churn. Mechanism: day_rider exit machinery Why: A negatively-correlated rank surface is the definition of a parameter you must not tune. And 'the exit is exit-proof' is now half wrong — holding beats every stop, but its whole edge is three days; strip them and \$5,410 becomes \$379. That is an argument about the ENTRY, which is what Saturday #6 changes.	Every fixed stop from 75 to 400 points is worse than no stop at all; 150/200/250/300/400 all turn the strategy NEGATIVE, and winners routinely draw 250–490pt against them first. On the exit surface, Spearman between in-sample and out-of-sample rank across 114 rules is −0.30, and the best in-sample rule ranked 71st of 114 out of sample and lost to plain holding. The 4×ATR trail is a HARDER bar than the old +150pt	Arms: N/A. Exit: The 4×/2× ATR armed trail, hard flat 20:40 UTC. Unchanged. 600pt venue stop is last-resort insurance only. Kill: None — the finding is that there is nothing here to optimise.	Nobody N/A.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	Evidence: Part 2.7 DR7 #5, #9, #10 · verification: N/A. · mode: none	on 30 of 38 days and rescues exactly zero extra days.		
5	Keep the ASIA block, and keep the 13:30-14:45 window LIVE No change to either. Both are standing findings that survived a full battery. Mechanism: session-time policy (RunConfig.no_open_asia, and the validated US window) Why: Two independent re-confirmations of the 14:45 edge in one week is as good as this desk gets. Evidence: Movement 3 \$5 (lens 2) · standing config · verification: Standing. · mode: live-paper	ASIA shadow n=1840 at −\$3.17/trade, the worst block on the desk. The 13:30-14:45 window is +\$8.93/trade on n=562. ▲ 'don't day-trade until the US open' was TESTED AND REFUTED — LONDON 07-13 is the desk's only positive block.	Arms: No new entries 00:00–07:00 UTC. Exits and every flatten path untouched. Exit: Unchanged. Kill: None. Note Movement 3 independently re-confirmed the 14:45 right edge this week from a completely different direction — the OPEN RIDER's 14:45-15:00 slice is −\$446 on the full tape and −\$453 on the unseen leg alone.	Nobody — already in config no_open_asia=False.

NOT-AN-ACTION · WITHDRAWN / REFUTED — KEPT ON THE CARD SO NOBODY RE-PROPOSES THEM · 12 PLAYS

Every one of these was a recommendation at some point this week. They are printed with their autopsies so next Friday's agent does not dig the same hole.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
1	★ RUN STATE as a ROUTING lever — WITHDRAWN. It is an exit find wearing a router's coat. REFUTED (as a routing lever) / LIVE (as a variable) Do not re-plumb the router around run state. The variable is real; the lever is not the one we were told. Mechanism: run-state detection (runstate.py / run_detect.py) used to arm and bench gates Why: It beat 4,000 placebo shuffles twice and held on twelve sessions it was never fitted on, so the variable survives. But the money is not where the lead said it was. Use these numbers, not those. Evidence: Part 2.5 LEAD 1 · verification: Re-derive on next week's tape before anyone quotes the \$39. · mode: none	What DOES reproduce is the missed side: in-run aligned MISSED n=96 · 59% · +\$2,703, chop MISSED n=194 · 26% · −\$2,089. The genuine spread is between the signals it BLOCKED (+\$27/signal aligned) and the chop it TRADED (−\$19/lot) — about \$17-\$46 a trade depending which pair you compare, not the flat \$39 quoted. And the lever: arming more would have added roughly \$170 this week, benching the 84 no-run trades would have saved \$584, but re-pricing the EXITS on the exact same entries the desk already took — changing nothing about routing — turns the week from −\$772 into +\$1,493.	Arms: N/A — withdrawn as a routing rule. Exit: N/A — but see the number: the money is in the EXITS on the same entries. Kill: The headline cell does not reproduce. Same 89 Mon–Thu trades, re-derived from scratch: in-run aligned TAKEN goes from the lead doc's n=34 · 50% · +\$10.34/tr to n=46 · 41% · −\$1.83/tr. Armed inside an aligned run the desk made NOTHING — it lost \$1.83 a lot. The 34/4/51 split could not be reproduced under ANY run-span definition tried (merge gaps 0s, 120s, 300s, 600s, 900s all give the identical 46/10/33).	Nobody — do not re-plumb the router N/A. Run state stays as a reported STATE that never flips a switch, which is what it already is.
2	The OPEN RIDER as a tradeable EDGE — REFUTED by skeptic lens 1 (the dissent) REFUTED Printed in full, deliberately, because the panel voted 2-1 and this is the 1. The OPEN RIDER still goes to shadow — but nobody should read its headline as an edge. Mechanism: OPEN RIDER, viewed as an exposure rather than a signal Why: Two skeptics with DIFFERENT out-of-sample legs — 17 unseen days versus 19 — reached opposite verdicts on the same strategy. That is not a contradiction to paper over; it is the reason this is a shadow candidate, and the honest way to resolve it is forward tape, not another backtest. Evidence: Movement 3 \$5 lens 1 · verification: Forward tape only. · mode: none	Concentration: 8 of 34 days carry 82% of the net. Cumulative window drift over those 34 days is −1,389.5pt, so a single held SHORT lot books \$2,728 with no strategy at all. What the lens concedes: the clock observation is real and stable (\$1.237/pt holding at 1.207/1.091 across both independent halves) and wide stops genuinely beat the house 1-ATR.	Arms: N/A. Exit: N/A. Kill: Three kills. (1) The intercept is zero: regress each day's P&L on that day's absolute 13:00-15:00 travel over 34 days and the zero-drift day pays −\$23 (t=−0.22) while the drift loading is \$1.237/pt (t=+2.75). (2) The machinery is decorative: a control making ONE direction decision per day earns +\$6,005 of the RIDER's +\$6,810 over 34 days — 88% — so the whole 5-minute cadence is worth \$3.31 a trade, less than the \$1.23 it costs to cross the spread. (3) Out of sample it is inside the null: on 17 unseen V5 days it lands at the 83.2nd percentile of 2,000 random-direction placebos, p=0.168, against the 99.4th percentile (p=0.0065) on the fitted days.	Nobody N/A.
3	The OPEN RIDER as a tournament gate, or running beside the day-rider — REFUTED twice over REFUTED (as a gate) / PARKED (beside the day-rider, pending isolation) Even if it were an edge, it cannot be deployed as configured. Two independent blockers. Mechanism: ODR deployment path	REVIVE the day-rider-collision half IF: it gets its own account or per-strategy position isolation, AND the 'does every live STOP have a slot?' auditor is built (SATURDAY #3, orphan-stop-auditor-on-a-timer-0808).	Arms: N/A. Exit: Its 2×ATR stop is the entire finding — and the live path will not give it one. Kill: As a gate: scaleout_slots() forces stop_atr_mult=1.0 on every live path, and at that stop	Nobody N/A.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	Why: This is the 08-06 cascade waiting to happen again: the account is SHARED and IBKR nets two desks into one number. Evidence: Movement 3 master ledger · verification: N/A until isolation exists. · mode: none		the same entries earn +\$17.59 a trade — a 62% haircut at double the fill count — plus the quiet-tape clip fires on 34–41% of entries. Beside the day-rider: 49% of entries open inside the day-rider's live window, 2.6 direction flips a day, and 34% of colliding entries are OPPOSITE its direction rule, on a netted account.	
4	COIL-CRACK and OPEN EXHAUSTION FADE — both REFUTED, and the fade is the more instructive one REFUTED Neither is reformulable. Both closed. Mechanism: COIL-CRACK = ignition out of a quiet 30 minutes. OPEN EXHAUSTION FADE = fade a straight 30-minute line into the open. Why: A filter that removes trades AND money has nothing to reformulate. And the beautiful monotone table that motivated COIL-CRACK is an artefact of overlapping 1-minute sampling — the same pseudo-replication trap that appears three times in this report. Evidence: Movement 3 \$4 and master ledger · verification: N/A. · mode: none	The lesson is the fade: it has the strongest RAW signal found anywhere in Movement 3 — day-demeaned forward return +54.2pt at 65.8% win, both sides, monotone in the threshold — and it cannot be traded, because the median excursion AGAINST it is 4.76 ATR against 3.00 ATR in its favour. A 15-minute edge delivered only after a 4.8-ATR excursion against you is not an edge you own; it is an edge your stop pays for.	Arms: N/A. Exit: N/A. Kill: COIL-CRACK, by direct ablation — the only test that decides a filter: WITH the efficiency gate +\$2,049 / n=51; WITHOUT it +\$2,238 / n=58. It costs money AND sample. The close-location confirm selected nothing at all (+\$830 / +\$830 / +\$832 at three settings). EXHAUSTION FADE: 51 configurations, every one negative or zero, and it caught 0 of 14 target runs. Both obvious reformulations — widen the stop to 3–4 ATR, delay entry 10 minutes — were tested and both fail, which CLOSES it.	Nobody N/A.
5	The whole VACUUM / FLOW-LED / ABSORPTION-SHELF family — one PARKED, the rest REFUTED REFUTED (FLOW-LED, ABSORPTION-SHELF, FLOW-TRAP, bar-derived flow proxy) / PARKED (VACUUM-BREAK, the raw sell-burst fade, and the L2 book reading) Do not revisit any of it until the tick archive is longer. The clusters they were built to exploit were never clusters. Mechanism: footprint entries built on aggressor flow and L2 depth Why: A label is not a signal. VACUUM fires on 1,451 minutes to find 8 runs (0.55%); FLOW-LED hangs on one minute of flow that reverses on the next bar. Both clusters produced no survivor because the cluster is not a real thing — and finding that out is worth more than a curve-fitted P&L would have been. Evidence: Movement 3 §2–3 and master ledger · verification: ~15 sessions of aggressor-tagged ticks. · mode: none	△ THE FEE CORRECTION MATTERS HERE AND CHANGES NOTHING: every greenfield backtest in this Movement was originally priced at ~\$5/RT against the true \$1.50/RT. The correction is exactly linear (net@1.50 = net@5.00 + 3.50 × n) so it HELPS everything. One row flips sign — FLOW-TRAP T=3.0/ RM=2.0 goes −\$105 → +\$154 on 74 trades — and it is flagged. A loser the correction rescues was never killed by cost in the first place. REVIVE VACUUM IF: (a) a down-trending or genuinely range-bound week arrives so the +1,300pt melt-up stops doing the work and the short side clears zero on its own, AND (b) n ≥ 200 walk-forward trades, needing ~15 sessions of aggressor-tagged ticks against the 5 we have. The one stone worth funding is the L2 book re-cut on QUEUE DEPLETION RATE rather than static average depth — the 82.3M-row book table at ~190 snapshots/second supports it.	Arms: N/A. Exit: N/A. Kill: FLOW-LED, traded as the census's own rule: catch-rate and profitability move in OPPOSITE directions — every cell catching 2 of 4 targets loses \$434–\$1,302 over 111–129 trades. Not thin-n; a mechanism with the wrong sign. ABSORPTION-SHELF: 12 of 12 cells negative and losing on BOTH sides — not even a drift artefact, just wrong. FLOW-TRAP: the cost-free forward-return test — its raw 15-min return (+2.25pt, 52.1% up) is BELOW the tape's own unconditional drift (+2.73pt, 53.9%) at 5, 15 and 30 minutes. Its information content is not small, it is negative. VACUUM-BREAK: strip-the-best kills it (+\$1,069 → +\$342 at strip-3 → +\$26 at strip-5 — five trades are the whole thing) and the walk-forwarded regime policy returns −\$125 over 91 trades, a selection tax of \$1,194.	Nobody — keep archiving ticks instead N/A.
6	Any symmetric chop scalp, at any distance — REFUTED, and here is the bar for the next one REFUTED Do not build a chop scalper. The chop lab was priced correctly at \$1.50/RT from the first line, so there is no fee excuse. Mechanism: CHOP-FADE (band fade) and CHOP-EDGE (failed-break fade) Why: The answer to chop is not a chop strategy. It is the off switch — see MONDAY #1 (chop-avoidance-is-the-biggest-number-0808), which is worth an order of magnitude more than any invention in this report. Evidence: Movement 3 (chop lab) and master ledger · verification: P(revert first) > 0.5 on fresh tape, or it does not get built. · mode: none	Maximum gross expectancy in EITHER direction is \$1.49 a trade against a \$2.24 friction floor, and the bias decays to 0.5006 by 1.5 ATR — so widening to amortise the fixed fee destroys the edge it was meant to pay for. ★ STANDING BAR: any future chop-scalper proposal must FIRST show P(revert first) > 0.5 on fresh tape.	Arms: N/A. Exit: N/A. Kill: Model-free first-touch RACE — not an MFE count: P(revert) = 0.4723 on n=4,249, z = −3.61. Negative BEFORE any cost. 450 frictionless cells: 34 positive at a median −\$1.50 a trade. 77% of random-direction placebos beat the best-of-grid (p = 0.770). CHOP-EDGE: 1 of 15 cells positive, strip-best-1 already negative, out-of-sample sign flip.	Nobody N/A.
7	The idle-gate lab is an HONEST NULL — do not use it to justify arming anything	The useful part is WHY it is a null, with a number rather than a shrug: the money in these runs	Arms: N/A. Exit: Each gate's own deployed two-lot exit, repriced tick-by-tick.	Nobody N/A.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	REFUTED (as arming evidence) Read it for the diagnosis, not for the P&L. Its three leads are already on this card as BUILD items; the lab itself arms nothing. Mechanism: firing the six live gates mechanically and ungated into the 53 testable sat-out runs Why: This corroborates the run-catcher NULL already banked on 22.6M ticks — run STARTS are unpredictable — and adds the complementary fact about latency. The unturned stone is the narrow band in between. Evidence: Movement 2 · verification: N/A. · mode: none	lives in the first ninety seconds after ignition, and our gates fire seven to eight minutes before it. Every gate we own either wants five minutes of history (abs_veto, grind) or wants the move to be over (rgv, capitulation, exhaustion). Being late by ninety seconds is nearly as bad as being wrong.	Kill: Mechanically: −\$474 on 7 fires across 53 testable runs. Ungated: +\$997.50 on 19 fires — but strip the three best trades and that becomes +\$62.50.	
8	The two-ratchet stays dead — but it did NOT clip a runner, and that is evidence FOR reconsidering it PARKED Keep it out. But record the correction: the reason we all believed it was killing runners is not true. Mechanism: two-ratchet partial exit Why: Killing an idea for the wrong reason is how a desk loses an edge for a year. The objection is retracted; the verdict happens to survive anyway. Evidence: Part 2 §4 · verification: N/A. · mode: shadow	26 grind trades, 3 genuine ≥ 4R runners, ZERO clips. The rule says thin n is a shadow and a withdrawn objection is evidence for reconsidering — so it is reconsidered here, and it still fails on its own merits rather than on the myth.	Arms: N/A. Exit: This IS an exit rule. Kill: It still lost on the boring trades, which is why it stays out. REVIVE IF: it can be shown to win on the boring trades, since the runner-clipping objection is now withdrawn.	Nobody — keep it shadowing N/A.
9	The 2026-07-31 exit-ladder guess — wrong in every cell it could be measured in REFUTED Stop using the hand-written ladder as a reference. The DIRECTION of the guess is right; the numbers on it are not. Mechanism: the operator's 07-31 per-gate exit ladder Why: The instinct was right and the arithmetic was not, which is the most common shape on this desk. Evidence: Part 2 §5 · verification: N/A. · mode: none	The direction — ride trends, bank chop — survives and is exactly what the Lot B rung widening (SATURDAY #10, lotb-rung-widen-atrsplit24-0808 — renamed from 'chandelier regime-key' by REV2 FIX0, which found it touches no chandelier) and the quiet-clip carve-out (BUILD #14, quiet-clip-regime-carveout-0808) implement properly.	Arms: N/A. Exit: N/A — it IS an exit spec. Kill: 24 gate × rung cells measured, and the guess never ranked better than 3rd out of 53 policies in any of them.	Nobody N/A.
10	GIVEBACK — TRULY-RETIRED REFUTED It is done. Do not re-propose it. Mechanism: give-back-from-peak exit on the tournament book Why: Named explicitly because the day-rider has a same-named mechanism that is NOT retired. Evidence: Part 1.5 §1 · Part 2.7 DR7 #6 · verification: N/A. · mode: none	Not computed here beyond the dossier's verdict; the rehab section carries the working.	Arms: N/A. Exit: N/A. Kill: Retired in the rehab dossier after the full six-step wash. Δ note this is the TOURNAMENT's give-back. The DAY RIDER's give-back-33%-of-peak is a separate, live SHADOW item (92% days green, half the tail, ~\$1,000 of the total) — do not confuse the two.	Nobody N/A.
11	'+\$4,169 on 102 trades' — PARKED as a quotable figure. Do not quote it. PARKED If you see this number anywhere, it is wrong twice over. Quote the range instead. Mechanism: the OPEN RIDER headline figure Why: A precise-looking number that four harnesses disagree about by \$1,200 is a range wearing a decimal point. Evidence: Movement 3 §5 lenses 1 and 3 · verification: N/A. · mode: none	Three independent harnesses give +\$3,907 / +\$4,620 / +\$5,028 / +\$5,139 at the correct fee on the same 17 days. ★ Quote 'about +\$4.4k ± \$0.9k on 95 verified trades'. A from-scratch harness lands 16% below the published headline — treat it as ±20%, never as a precise figure.	Arms: N/A. Exit: N/A. Kill: The '102' is a different row of the report entirely (the thresholding comparison, +\$4,142). The trade count is 129, confirmed independently by three harnesses.	Nobody N/A.
12	ODR cadence-3 (PARTIALLY REFUTED) and the 14:45-15:00 right edge (REFUTED — cut it) REFUTED Strike cadence-3 from the language and cut the right edge. Both are corrections to the shadow spec, not to anything live. Mechanism: OPEN RIDER cadence and window edges Why: Two corrections to a shadow candidate before it ever trades is the cheapest kind of correction there is. Evidence: Movement 3 §5 lens 2 and master ledger · verification: N/A. · mode: none	The chosen cell sits on the safe side of a cliff the author did not know was there. And the right-edge cut independently re-confirms the desk's own 14:45 boundary from a completely different direction.	Arms: 13:00-14:45 UTC at a 5-minute cadence. Not 3 minutes. Not to 15:00. Exit: Unchanged. Kill: Cadence: on unseen tape cadence-3 is positive in only 4 of 16 cells against 13/16 at 5 minutes and 15/16 at 10. The plateau is real in STOP WIDTH (27/27 positive at stop ≥ 2.0) and one-sided in cadence — slower works, faster does not. Right edge: 14:45-15:00 is −\$446 across all days and −\$453 on the unseen leg alone.	Nobody — spec correction only N/A.

STANDING CAVEATS ON EVERY DOLLAR ON THIS CARD

Every repriced number is a CEILING. Fills are modelled at the exact stop/target/trail level with no slippage, and the desk's own standing finding is that live losses run **1.1-3.1× modelled**. Treat the policy as the edge and the headline dollar as decoration. Anything on thin n is flagged in-row as a direction, not a measurement.

The fee is \$1.50 per round trip. Every greenfield number in Movement 3 was originally priced at ~\$5/RT and has been re-priced; the correction is exactly linear ($\text{net}@1.50 = \text{net}@5.00 + 3.50 \times n$) so it helps everything, winners and losers alike. A loser the correction rescues was never killed by cost in the first place, and the one row where that happens is flagged in NOT-AN-ACTION #5.

Three method wounds are still open and they bound everything above. The shadow book leaks past its own stops on 44% of trades (median overshoot $0.56 \times \text{ATR}$), so every shadow result that rewards cutting early is inflated — BUILD #3. `signal_journal.suppressed_by` is NULL in all 596 rows, so every claim about what a gate would have done unrouted is a reconstruction from the tape, not a measurement — BUILD #5. And microstructure sampling noise is large: on `exhaustion_short` the standard deviation across ten polling phases is \$418 on a -\$880 mean. Anyone quoting a single run of that harness is quoting noise.

§ The evidence behind those plays


Everything from here down is the working. Two halves, one document. **The warm half** is what the live paper tournament actually did this week (Part 1), then the section that matters most — **Rehabilitation** (Part 1.5), where every red gate and leaky exit is walked through the full six-step wash instead of being benched at face value. Then the shadow board and the promotion it did not earn (Part 2), the mid-week leads re-run adversarially on the week's own data (Part 2.5), the **Router Operation Review** (Part 2.6), and **the Day Rider** — the second desk, rebuilt from the raw tape over

38 sessions (Part 2.7). **Then the cold, tape-first half:** we ignore everything the desk did and ask, from the raw MNQ tape and order book, **where was the money and did we show up?** (M1) — then test whether the gates we already own could have caught the runs we missed (M2), and finally build brand-new gates from scratch and backtest them to death, the ones that failed included, that is the point (M3). Plain English, honest money, every claim a table — and every reprice a ceiling.

WHAT IS IN HERE

1. **Front** — [the honest headline, the scorecard, WHAT TO ACTUALLY DO](#)
2. **Part 1** — [The live desk — what the six-gate paper tournament actually did](#)
3. **Part 1.5** — [Live-desk REHABILITATION — never bench a gate at face value](#)
4. **Part 2** — [The shadow desk, the promotion battery & the regime-flex exit lab](#)
5. **Part 2.5** — [Mid-week musings — every lead from the desk chat, tested hard](#)
6. **Part 2.6** — [ROUTER OPERATION REVIEW — the desk's #1 lever, on the stand](#)
7. **Part 2.7** — [THE DAY RIDER — the second desk, and how we make it work](#)
8. **Movement 1** — [The census — the biggest runs on MNQ this week, and did we show up](#)
9. **Movement 2** — [The idle-gate lab — could the gates we own have caught them?](#)
10. **Movement 3** — [The greenfield lab — brand-new entries, and the chop-day scalp lab](#)
11. **Part 3.1** — [REV2 · Q0 — the \$ER30 \geq 0.35\$ floor on abs_veto_short: put it in the table and it dies in the table](#)
12. **Part 3.2** — [REV2 · Q1 — is grind_long on or off on Monday? One population, one number](#)
13. **Part 3.3** — [REV2 · Q2 — last Friday's card: did we do those things, and did they work?](#)
14. **Part 3.4** — [REV2 · FIX0 — Saturday #3 is not a chandelier play, and as written it cannot be typed into any file the desk reads](#)
15. **Part 3.5** — [REV2 · FIX1 — abs_veto_long: resolving HOLD #1 against Part 2 \\$4, \\$5 and \\$9](#)
16. **Part 3.6** — [REV2 · FIX2 — the ER30 floor through the 54-filter columns it never faced: it keeps 5 of 15 winners](#)

P1 The live desk — Monday 3rd to Friday 7th August



Gaz, the tournament lost **-\$772.50** this week across 110 booked lots. But the loss is not spread around and it is not mysterious: **the long book made +\$345.50 and the short book lost -\$1,022.50**, and one gate — **abs_veto_short** — lost **-\$559.00 on nine lots without winning a single one**. Take the shorts away and the week is green. The thing that survived every test I could throw at it is **abs_veto_long**: sixteen entries, thirty-two lots, **+\$566.00 at a 62% strike rate**, still positive if you delete any single day of the week and still **+\$363** if you delete its three best lots. That is the one number in this section I would actually put weight on. Everything else is either thin, one-week, or explained by a software defect.

THE HONEST TOTAL — THREE NUMBERS, NOT ONE

The raw ledger for the week reads **-\$1,028.00** over 112 rows. Two of those rows are the MD_STREAM corruption from Monday night (the two **abs_veto_short** lots opened 08-04 22:02 with 1,848-point stops, closed by MAX_HOLD for **-\$255.50**) and they carry data_quality = EXCLUDE:md_stream_atr_corruption_20260804. Filter them out, as every honest reader must, and the desk booked **-\$772.50 on 110 lots**. Two more of those lots are not the tournament at all — they are the day-rider's position, booked to the tournament by the cross-desk flatten on Thursday

morning (−\$95.50). So **the tournament’s own gates traded 108 lots for −\$677.00**. And of that, **−\$165.50 was not a strategy losing money** — it was three separate shared-account defects (below). The number that actually measures the gates is **−\$607.00**.

WHAT YOU ARE LOOKING AT	LOTS	NET	WHY IT DIFFERS
Raw trades rows, 08-03 → 08-07	112	−\$1,028.00	includes the two corrupt MD_STREAM lots
Clean ledger (data_quality IS NULL)	110	−\$772.50	the desk’s real booked money
Tournament gates only	108	−\$677.00	strips the day-rider cross-desk pair
Gate strategy, defects stripped	104	−\$607.00	strips the ghost stop and the phantom close too

The week, day by day — and what the tape was doing

Every tape number here is recomputed from capture.db 5-second bars rolled up to one minute — not read off a dashboard. **Path** is the total distance price actually travelled minute to minute; **net** is where it ended up. **ER** is net divided by path — how much of the walking got you somewhere. High ER = a trend you can hold; low ER = a treadmill. The US session here is 13:30–20:00 UTC.

DAY	US NET	US PATH	US ER	US ATR	READ	LOTS	WINS	DESK NET
Mon 08-03	+482	2,815	0.171	12.4	clean up-trend	30	8	−\$280.00
Tue 08-04	+641	2,799	0.229	12.6	the week’s cleanest trend	16	10	+\$330.50
Wed 08-05	−367	4,095	0.090	20.1	whippy down — ATR doubled	29	4	−\$817.00
Thu 08-06	+161	4,253	0.038	20.5	pure chop at high vol	14	3	−\$305.00
Fri 08-07	+120	3,676	0.033	14.2	chop, but a real London run at 12:30	21	12	+\$299.00

Read that table twice, because it says something uncomfortable. **The two cleanest trend days produced +\$50.50 between them**. Monday walked 482 points up in the US session at ER 0.171 — a genuinely good day to be long — and

the desk lost \$280 on it. Tuesday was even cleaner and made \$330.50. Then Wednesday and Thursday, the two days where ATR doubled from 12 to 20 points and the tape went nowhere (Thursday travelled 4,253 points in the US session to finish +161), took −\$1,122 out of the desk between them. **The desk does not lose money on trend days. It loses money on high-volatility days that go nowhere.** And with n=5 sessions that is an observation, not a law — but it is the same shape the whole week keeps making.

THE FULL-DAY VIEW, FOR COMPLETENESS

Over the whole 22-hour session, not just the US window: Mon +360pt net on 8,288 of path (ER 0.043) · Tue +864 on 8,465 (ER 0.102) · Wed −252 on 9,085 (ER 0.028) · Thu **+11 points on 9,964 of path** (ER 0.001 — the flattest day in the record) · Fri +247 on 8,349 (ER 0.030). Thursday walked nearly ten thousand points to finish eleven points up. That is what a 445-point range and no direction looks like from the inside.

Lead with what survived: abs_veto_long **SURVIVED**

It is the only gate on the desk that made money and kept making it under pressure. Sixteen entries, thirty-two lots, **+\$566.00**, twenty winning lots out of thirty-two (**62%**), average winner **+\$47.60** against an average loser of **−\$32.20**. It is a thrust-continuation gate with a 55-second confirm, so it only fires when a burst is still going 55 seconds later — and that is exactly the behaviour the week rewarded. Its two lots run different targets (A at 1.0R, B at 1.5R) with the quiet-tape clip on both under ATR 22.

ROBUSTNESS TEST	RESULT	VERDICT
Headline	+\$566.00 · 32 lots · 62% win —	
Drop Tuesday (its best day)	+\$235.50	still positive
Drop Wednesday	+\$635.00	still positive
Drop Thursday	+\$582.00	still positive
Drop Friday	+\$245.50	still positive
Strip its three best lots	+\$363.00	survives
Strip its best whole entry (08-04 13:32, + \$137.50)	+\$428.50	survives

Four leave-one-day-out runs, all positive. Strip-the-best-three, still +\$363. That is as good as a one-week read gets. **The honest caveat is the sample: 16 entries.** Sixteen entries is not a track record, it is a promising month’s worth of evidence compressed into five days, and it is in-sample — I found this by looking at the week after the week happened. It is also worth saying plainly that abs_veto_long only traded on four of the five days, and did nothing at all on Monday. Keep it armed, keep watching it, do not size it up on this.

And what did the damage: the short book **KILLED-ON-SIGHT**

Thirty short lots this week. **Three winners.** Net **−\$1,022.50**, or −\$34.08 a lot. The long book over the same week: 78 lots, +\$345.50, 44% win, +\$4.43 a lot. That is a **\$38.51-per-lot gap** between the two sides of the same desk.

SIDE	LOTS	NET	WIN %	\$ / LOT
LONG	78	+\$345.50	44%	+\$4.43
SHORT	30	−\$1,022.50	10%	−\$34.08

I tried to kill this finding and could not. The short book was **negative on all four days it traded**. Leave any single day out and it is still deep red (−\$519.00 at best, −\$881.50 at worst). Strip its three worst lots entirely and it is still **−\$789.50** on the remaining 27. There is no single bad print carrying this.

SHORT BOOK — ROBUSTNESS	LOTS	NET
All of it	30	−\$1,022.50
Drop Monday	24	−\$838.00
Drop Wednesday	20	−\$519.00
Drop Thursday	22	−\$829.00
Drop Friday	24	−\$881.50
Strip the 3 worst lots	27	−\$789.50

But here is the part that stops it being a simple “bench the shorts” conclusion, and it is the most interesting fact in this section. The shorts were not wrong about direction. On Wednesday the US session fell 367 points and the short book lost −\$503.50 doing it. On Friday the shorts fired at 13:29 into an open that was, five minutes later, back below where they sold. They were early and their stops sat inside the noise. I take that apart properly in the Wednesday case study below.

The per-gate cards

An “entry” is a signal; each entry opens two lots (A and B) with different exits, so the lot count is roughly double the entry count. **Med hold** is the median time a lot stayed open. **Exit mix** is how the lots actually left. Best and worst are whole entries, both lots together.

GATE	ENTRIES	LOTS	NET	WIN %	AVG WIN	AVG LOSS	MED HOLD	EXIT MIX	BEST ENTRY	WORST ENTRY
abs_veto_long	16	32	+\$566.00	62%	+\$47.60	-\$32.20	5m27s	18 TARGET · 12 STOP · 2 CLAIM	08-04 13:32 + \$137.50	08-06 22:15 - \$72.00
capitulation_long	3	6	+\$84.50	33%	+\$103.50	-\$30.60	20s	4 STOP · 2 TARGET	08-07 12:29 + \$207.00	08-05 00:00 - \$70.50
rgv_short	0	0	\$0.00	—	—	—	—	never fired	—	—
nipc_long	17	33	-\$97.50	33%	+\$53.40	-\$31.10	24s	22 STOP · 11 TARGET	08-03 14:10 + \$137.00	08-03 13:38 - \$119.00
exhaustion_short	3	6	-\$162.00	33%	+\$52.00	-\$66.50	1m47s	4 STOP · 1 TARGET · 1 CHANDELIER	08-07 14:48 + \$104.00	08-05 14:31 - \$141.00
grind_long	2	6	-\$172.00	17%	+\$44.50	-\$43.30	11s	5 STOP · 1 TARGET	08-05 13:33 - \$80.00	08-05 13:36 - \$92.00
nipc_short	8	16	-\$337.00	6%	+\$28.50	-\$24.40	27s	15 STOP · 1 TARGET	08-06 13:04 + \$10.00	08-03 13:43 - \$77.50
abs_veto_short	4	9	-\$559.00	0%	—	-\$62.11	1m00s	9 STOP	08-06 13:33 - \$129.50	08-07 13:29 - \$155.50

THE SINGLE MOST TELLING NUMBER ON THE DESK

The median losing lot this week lived 24 seconds. The median winning lot lived 3 minutes 6 seconds. Seventy-one lots exited on a STOP for -\$2,600.50; thirty-four exited on TARGET for +\$1,794.00. When this desk is wrong it finds out almost immediately — which is the stop doing its job — but it also means an enormous share of the book never gets a chance to be right. **Twelve overlapping five-minute windows this week each contained four or more stopped lots** — five on Monday (all nipc), five on Wednesday (the open and the 14:33 short cluster), two on Friday (the 13:30 bell).

Gate by gate, in plain English.

abs_veto_long is the desk. It made money on Tuesday (+\$330.50) and Friday (+\$320.50), lost small on Wednesday (−\$69.00) and Thursday (−\$16.00), and did not fire at all on Monday. Its worst entry all week was −\$72.00 — it never has a disaster, because the 55-second confirm keeps it out of the fake bursts and the 1.0R/1.5R scalp takes money off the table fast.

capitulation_long traded three times and one of them was the week's single best print (+\$207 in ten seconds). On three entries you cannot say anything about a win rate. What you can say is that its mechanism — buy the flush, get paid on the snap-back — only needs to be right occasionally, and its average winner (\$103.50) is over three times its average loser (\$30.60). Note for the record: the router advised benching this gate four separate times on Friday morning and never did; had it been benched, the desk's best trade of the week does not happen.

exhaustion_short is the interesting one. Three entries: −\$141.00 into Wednesday's knife, −\$125.00 at Friday's bell, and then **+\$104.00 at 14:48 Friday** when it was re-armed under the fader carve-out (A took TARGET +\$54, B rode a CHANDELIER to +\$50). The carve-out's argument was that a fader is not a momentum gate and should not be benched on a momentum rationale. Its first live sample paid. $n=1$.

grind_long traded once, for eleven minutes, on Wednesday, and lost \$172 across six lots — five of them stopped. Median hold: eleven seconds. It bought 29,999 at 13:33 and 29,987 at 13:36; the 13:35 one-minute bar topped at 30,023 and the very next minute traded down to 29,913 — a 110-point one-minute bar that closed three of its lots inside that single minute.

abs_veto_short lost every lot it took. Four entries, nine lots, no winners. Its stops were the widest on the desk (median stop distance 32.2 points against a 15.8-point median for its long twin), because all four of its entries landed on the two highest-ATR days of the week — so every one of its losses cost roughly double a normal one.

rgv_short did not exist this week. It was armed for a total of under two minutes across five midnight re-activations, each of which the router reversed within about thirty seconds. It also carries an `atr_min` of 20 points, and the week's median 1-minute ATR was 11-12 — so even fully armed it had almost no eligible tape. It is a gate on paper only.

The roster is not “six gates” any more — and three of them barely exist

The scope calls this a six-gate tournament. It is not. The live switch file carries **eight** names (grind_long, capitulation_long, abs_veto_long, exhaustion_short, abs_veto_short, rgv_short, nipc_long, nipc_short), rgv_long is not in the roster at all, and nipc was killed mid-week. More to the point, here is how much of the 120-hour week each gate was actually armed for.

GATE	ARMED	% OF THE WEEK	ARM WINDOWS	WINDOWS THAT FIRED	LOTS	NET
nipc_long	82.5 h	69%	2	1	31	−\$38.00
nipc_short	77.1 h	64%	2	2	14	−\$302.50
abs_veto_long	56.2 h	47%	19	6	32	+\$566.00
capitulation_long	33.3 h	28%	6	2	6	+\$84.50
exhaustion_short	19.3 h	16%	9	3	6	−\$162.00
abs_veto_short	11.6 h	10%	12	3	9	−\$559.00
grind_long	1.0 h	1%	9	1	6	−\$172.00
rgv_short	2 min	0%	5	0	0	\$0.00

Two things jump out. First, **the two gates that were armed longest were the two that were killed** — nipc sat on for two-thirds of the week because it was armed once by hand on 08-01 and nobody re-judged it until the kill criterion fired. Second, **grind_long was armed for one hour in five days** and every one of its six lots came from a single eleven-minute window. The router has effectively decided grind is not a gate it wants; it just has not said so out loud. That is worth making explicit rather than leaving as an emergent property of 72 separate decisions.

nipc — the kill criterion did exactly what a kill criterion is for **KILLED**

You wrote the criterion yourself: bench immediately if cumulative $\leq -\$400$. Here is the arithmetic, recomputed from the trade rows.

DAY	LOTS	NET	RUNNING TOTAL	WHAT HAPPENED
Mon 08-03	30	-\$280.00	-\$280.00	15 entries in 80 minutes, 13:34-14:53
Tue 08-04	0	\$0.00	-\$280.00	benched
Wed 08-05	13	-\$90.50	-\$370.50	re-armed by operator 12:55, off again 13:5x
Thu 08-06 to 13:08	4	-\$29.50	-\$400.00	criterion trips exactly, n=47
Thu 08-06 13:47	2	-\$34.50	-\$434.50	the ghost stop — not a nipc trade at all

The gate hit **exactly -\$400.00 on its 47th lot** and was benched at 13:09 the same minute, and put into `reactivate_gates.py::HOLD` so the 22:00 re-arm cannot resurrect it. That is the rule executing without argument, which is the whole point of writing one down in advance.

The extra -\$34.50 is not nipc's. At 13:47 — thirty-eight minutes after the gate was switched off — a leftover GTC stop from Thursday morning's shared-account mess fired with no position behind it, opened a naked short, and was attributed to `nipc_short` by order reference. So **nipc's true live record is -\$400.00 on 47 lots at a 26% strike rate**, and its ledger is polluted by a ghost. That distinction matters if you ever revisit it, because the 08-05 regime filter never got a clean out-of-sample day — the short side did the killing (-\$337.00 on 16 lots, one winner) while the long side was almost flat (-\$97.50 on 33 lots, and 11 of those were winners at an average +\$53.40).

Day types, direction, and the hour of the day

Three splits, each computed on the same 108 gate lots.

DAY TYPE (BY US-SESSION ER)	LOTS	NET	WIN %	\$ / LOT
TREND UP (Mon + Tue, ER 0.17-0.23)	46	+\$50.50	39%	+\$1.10
WHIP DOWN (Wed, ER 0.09 at ATR 20)	29	-\$817.00	14%	-\$28.17
CHOP (Thu, ER 0.04 at ATR 20)	12	-\$209.50	25%	-\$17.46
CHOP + a morning run (Fri)	21	+\$299.00	57%	+\$14.24

TRADE DIRECTION VS THE DAY'S OWN NET DIRECTION	LOTS	NET	WIN %	\$ / LOT
Aligned with the day	69	+\$155.50	43%	+\$2.25
Against the day	39	-\$832.50	18%	-\$21.35

A **\$23.60-per-lot gap** between trading with the day and trading against it, on 108 lots. That is the same direction as the RUN-STATE work in the decision record and it is not a re-derivation of it — this is a crude end-of-day-hindsight split, so it tells you the size of the prize, not that the desk can capture it live. Do not read it as evidence run-state detection works; read it as evidence that something that gets direction right is worth roughly \$20 a lot on this desk.

⚠ A DIRECT CONTRADICTION OF STANDING POLICY — REPORT IT, DO NOT ACT ON IT


The standing time finding on this desk (n=562, survived strip-best-3, both halves, +3.27sd) is that **13:30-14:45 is where expectancy lives**, at +\$8.93 a trade. **This week it was the worst window on the desk.**

WINDOW	LOTS	NET	WIN %	\$ / LOT
13:30-14:45 UTC	58	-\$1,005.50	24%	-\$17.34
Every other hour	50	+\$328.50	46%	+\$6.57

And it is not one bad day: the window was negative on **four of the five days** (Mon -\$281.50, Wed -\$647.00, Thu -\$164.00, Fri -\$160.50, Tue the only green at +\$247.50), and leaving out any single day still leaves it between -\$358.50 and -\$1,253.00. **Fifty-eight lots does not overturn five hundred and sixty-two.** But this is exactly the kind of thing that gets quietly forgotten, so it is written down: **WATCH** — if the open block is negative again next week, the standing rule needs re-deriving, not defending.

Case study 1 — Wednesday 5 August: the desk sold a 367-point drop and lost \$817

This is the flagship session of the week, and it is not the biggest loss because the desk picked the wrong side. It is the biggest loss because **the desk picked the right side and its stops sat inside the noise.**



Here is the sequence, in money. The US session opened at 29,975 and was 29,597 by 20:00 — **down 367 points to the 19:59 close** — but it did it by walking 4,095 points, at an ATR of 20 against Monday and Tuesday’s 12.5. The desk was on the correct side of that. It lost \$503.50 being on the correct side.

TIME	GATE	SIDE	IN	OUT	STOP DISTANCE	RESULT	HOW FAR IT WENT ITS OWN WAY WITHIN 60 MIN
13:17	nipc_short	SHORT	29,913.50	29,920.38	6.9 pt	-\$30.50	17.8 pt
13:33	grind_long (4 lots)	LONG	29,999.06	29,989.81	21.3 pt	-\$80.00	74.2 pt
13:36	grind_long	LONG	29,987.50	29,965.25	22.3 pt	-\$92.00	85.8 pt
13:50	nipc_short	SHORT	30,016.00	30,029.75	13.8 pt	-\$58.00	191.3 pt
14:22	abs_veto_short	SHORT	29,944.38	29,976.63	32.3 pt	-\$132.00	182.1 pt
14:31	exhaustion_short	SHORT	29,867.75	29,902.25	34.5 pt	-\$141.00	105.5 pt
14:32	abs_veto_short	SHORT	29,859.50	29,894.25	34.8 pt	-\$142.00	97.3 pt

Look at the 14:22 short. It sold 29,944 and was stopped 32 points higher at 14:25. **Price then fell to 29,762 within the hour and was 29,597 at 20:00 — 347 points below where it sold.** On two lots held to that price that is roughly \$1,390 of hindsight money. It banked −\$132.

The 14:31/14:32 pair is worse in a way, because they were the second attempt. exhaustion_short sold 29,867.75 and abs_veto_short sold 29,859.50, both were stopped inside two minutes by a bounce to 29,914, and the session ended 263–271 points lower at 29,597. Three separate gates, three correct calls, −\$415 between them in about eleven minutes.

THE PATTERN, MEASURED ACROSS THE WHOLE WEEK — AND WHY IT IS NOT AN INSTRUCTION TO WIDEN STOPS

I ran this on every losing lot of the week. **Of the 73 losing lots, 65 (89%) went further in their own direction within 60 minutes of entry than the loss they actually took.** Per gate: nipc_long 22 of 22, abs_veto_short 7 of 9, grind_long 5 of 5, exhaustion_short 4 of 4, abs_veto_long 10 of 12.

That is an MFE statistic and MFE is not a win rate — this file's own hardest-learned lesson. The stop physically came first in every one of those 73 cases. So I re-raced them properly on 5-second bars, target held at each trade's own 2R, 120-minute horizon:

STOP WIDTH	REACHED 2R FIRST	STOPPED FIRST	MODELLER NET ON THESE 71 LOTS
×1.0 (as traded — the control)	1	70	−\$2,542
×1.25	7	64	−\$2,678
×1.5	18	53	−\$1,763
×2.0	30	41	−\$1,047

The ×1.0 control reproduces the real −\$2,600.50 to within 2%, so the model is honest about the tape. And on this week's data a wider stop looks better.

Do not act on this. The 08-04/05 study already tested exactly this on a bigger sample and **refuted 1.5×** (it rescued 2 of 52 stopped trades while making the other 50 lose 50% more), and found 2.0× sits inside an 18% model error. My run here is on losers only, on one week, with a simplified 2R target that is not the live exit ladder. **PARKED** — revive it if the ×1.5 / ×2.0 forward A/B pairs (sw_*_k10 / sw_*_k20) reach $n \geq 40$ and agree with this direction. Until then the honest reading is not “widen the stop”, it is “this


week’s tape had 20-point ATR and 10-15-point stops, and that combination cannot work”.

DAY	MEDIAN STOP DISTANCE	DAY'S MEDIAN 1-MIN ATR	STOP AS A MULTIPLE OF ATR	DESK NET
Mon 08-03	14.2 pt	11.6	1.22×	−\$280.00
Tue 08-04	14.2 pt	11.5	1.23×	+\$330.50
Wed 08-05	17.2 pt	11.4	1.51×	−\$817.00
Thu 08-06	9.2 pt	12.4	0.74×	−\$305.00
Fri 08-07	17.2 pt	10.8	1.59×	+\$299.00

The stop is 1.0 × the ATR read at entry, and on Wednesday and Thursday the US-session ATR was 20 while the day’s median was 11-12. On Thursday the desk was putting on 9-point stops in tape with a 20-point US-session ATR. **That is not a strategy problem, it is a measurement-window problem**, and it belongs in the rehab section rather than here — but the number is on the record now.

Case study 2 — Friday 7 August, +\$299, and every dollar of it explained

Friday is the best-documented session on the desk and the most instructive, because within nine hours it contained the week’s best trade, the week’s worst ninety seconds, a 20-minute halt, and then the cleanest run of gate behaviour all week.



TIME	WHAT HAPPENED	MONEY	RUNNING DAY
12:29-12:30	capitulation_long buys 29,662 and 29,667 — seconds before the 12:30 break — both lots TARGET at 29,717	+\$207.00	+\$207.00
13:29-13:31	Five lots stopped in 90 seconds at the cash open: abs_veto_short both lots, the phantom LONG from a redundant close, exhaustion_short both lots	−\$280.50	−\$73.50
14:13-14:33	Desk HALTED on reconcile drift — flat, healthy, and unable to enter. 14 minute-alerts from the tape watcher over those 20 minutes	\$0.00	−\$73.50
14:48	exhaustion_short, re-armed at 14:25 under the fader carve-out, goes 2-for-2 (A TARGET + \$54, B CHANDELIER +\$50)	+\$104.00	+\$30.50
15:28	abs_veto_long, armed 15:25 on a fresh 2-hour high, both lots TARGET	+\$94.50	+\$125.00
16:19	capitulation_long stopped pair, 20 seconds	−\$52.00	+\$73.00
19:03-19:58	abs_veto_long takes four entries in 55 minutes: +\$101, −\$72, +\$102.50, +\$94.50	+\$226.00	+\$299.00

The 12:29 trade is the one to sit with. The 12:29 one-minute bar swept an **89-point range** — high 29,685.50, low 29,596.00 — and closed back at 29,674.50. capitulation_long bought into the flush: B at 12:29:58 at 29,662, A at 12:30:02 at 29,667. **Both lots were out at 29,717 by 12:30:08.** Ten seconds from the first entry to both exits, +\$207.00 on a \$2-a-point contract. It is the whole reason Friday is green.

The honest other half: **price kept going to 29,867.25 by 12:50.** That is another **150 points beyond where the targets took it**, about \$600 on the two lots. Measured from the average entry of 29,664.50 the hindsight ceiling to that 12:50 high was 202.75 points, or \$811 gross, \$808 after fees — so the gate banked **26% of it**. That is what a scalp exit does; it is the trade-off you bought, not a fault. But the number should be visible rather than hidden inside a happy headline.

The 13:29 ninety seconds are the mirror image. abs_veto_short sold 29,695 and 29,694.50 at 13:29:01. The 13:30 bar then ran from 29,714.50 to a high of 29,813 — a hundred-point vertical — and both lots were stopped at 29,724. exhaustion_short sold 29,715 at 13:30:48 into the same spike and both its lots were stopped too. By 13:35 price was back at 29,696 — 19 points below where exhaustion_short had sold, and within two points of where abs_veto_short had.

Within 60 minutes those shorts had 130-151 points of favourable movement available. They saw none of it.


And there is a fifth lot in that cluster: at 13:30:02 the ledger records an `abs_veto_short_B` position that is **LONG**, in for 29,712.25, out at 29,695.25, – \$35.50. That is the phantom — a redundant market close that was forgotten rather than cancelled when the stop won the race, filling into an opposite position. It is not a trade anyone decided to take.

THE HALT COST NOTHING IN LOSSES AND EVERYTHING IN OPTIONALITY

Between 13:31 and 14:48 the tournament took **zero fills** — 77 minutes flat. The tape watcher fired fourteen minute-alerts between 14:13 and 14:33 saying `halted=True` while the desk was flat and healthy, escalating to **CRITICAL** at 14:27 when it named reconcile **DRIFT** as the cause. The cause was that `reconcile()` compared the tournament's own slots against the whole account net, and the day-rider's legitimate two lots looked like a leak. **A halt is not free even when the desk is flat** — the switches were fiction for at least those 20 logged minutes and everything armed was a no-op. This is now fixed and regression-tested, so it is a **NOT-AN-ACTION** for Saturday; it is recorded because the tape watcher spotting it in one minute and escalating correctly is the best thing that instrument did all week.

Case study 3 — Tuesday 4 August: one arming, eighteen hours, +\$330.50

The durable tick armed `abs_veto_long` at **04:20:21** and did not touch it again until **22:45:46** — eighteen hours and twenty-five minutes, unprompted, with no session running. In that window the gate took eight entries and sixteen lots.



ENTRY	PRICE	RESULT	NOTE
07:28	29,104.63	-\$60.00	both lots stopped in 8 minutes
10:46	29,134.38	-\$65.50	both lots stopped in 4 minutes
11:45	29,204.63	-\$48.00	both lots stopped inside a minute
13:32	29,327.13	+\$137.50	both TARGET in 60 seconds
14:07	29,511.63	+\$110.00	A TARGET, B claimed by hand
15:03	29,578.50	+\$94.50	both TARGET
15:14	29,621.75	+\$98.00	both TARGET
15:40	29,668.00	+\$64.00	A TARGET, B claimed by hand

Three losers for -\$173.50, then five winners for +\$504.00. Not one arm/bench decision in between. Every conventional risk rule on this desk would have benched that gate after the third stop-out — three stopped pairs is a textbook wall-of-stops trigger — and doing so would have cost the entire +\$504.

Friday told the same story from the other end. abs_veto_long was armed at 16:35:33 and left alone for the rest of the session; it took four entries between 19:03 and 19:58 for +\$226.00, and the one loser in that run (-\$72.00 at 19:12) sat between two winners. Across the whole week, **the two windows that produced the desk’s biggest profits were the two longest windows anyone left alone** — 18h25m on Tuesday and 7h24m on Friday.

The router this week — 72 decisions, 101 gate flips, and what they actually bought

The mechanical gazbot7-direction-router.timer is **disabled and inactive** — verified live, not recalled. The router is the durable 5-minute

headless tick plus whatever you and the session did by hand. Across the five days the log records **72 switch-change events touching 101 individual gates** out of roughly 1,390 ticks.

DAY	CHANGE EVENTS	DURABLE TICK	LIVE SESSION	OPERATOR	DESK NET
Mon 08-03	2	6	0	0	−\$280.00
Tue 08-04	12	17	0	0	+\$330.50
Wed 08-05	20	17	5	6	−\$817.00
Thu 08-06	14	13	1	3	−\$305.00
Fri 08-07	24	17	16	0	+\$299.00
Week	72	70	22	9	−\$772.50

(The last three columns count individual gate flips, so they sum to 101, not 72 — one event can flip three gates at once.) Note Monday: **two change events all day**, and the desk lost \$280. Friday: twenty-four, and the desk made \$299. The rate of change tells you nothing about the outcome on its own. What it does tell you is where the decision-making got hard.

HOURL (UTC)	CHANGE EVENTS, WHOLE WEEK
00-06 (Asia block — entries are config-blocked)	11
07-11 (London)	9
12	6
13 (the cash open)	14
14	7
15-16	6
19-20	4
22-23 (reopen housekeeping)	15

Twenty-one of the week’s 72 decisions landed in the two hours 13:00–15:00 — the two hours that also lost the desk \$1,005.50. **Decisions cluster exactly where they are least reliable.** Eleven more happened between midnight and 06:00, inside the Asia block, where `multislot_core._open()` refuses every entry anyway — those are eleven decisions that could not possibly have had an effect.

THE ARM-WINDOW TABLE — THE SINGLE MOST USEFUL THING I COMPUTED THIS WEEK

I reconstructed every period each of the six switch-managed gates was armed, from the trial log plus the 22:00 auto re-activation, and counted what

traded inside it. nipc is excluded because it was armed once by hand and ran for days — it is not a router decision.

HOW LONG THE GATE WAS ARMED	WINDOWS	THAT PRODUCED A TRADE	LOTS	NET
Under 5 minutes	21	0	0	\$0.00
5 to 30 minutes	14	3	10	-\$426.50
30 minutes to 2 hours	14	6	13	-\$168.50
Over 2 hours	11	6	36	+\$352.50
All	60	15	59	-\$242.50

Twenty-one arm-windows lasted under five minutes and produced not one trade between them. Thirty-five of sixty windows — more than half — lasted under half an hour, and between them produced ten lots and – \$426.50. Every dollar the switch-managed gates made this week came from windows longer than two hours. The three biggest earners were an 18h25m window (+\$330.50), a 7h24m window (+\$226.00) and an 18h27m window (+\$155.00).

This is the **churn** metric made concrete. A five-minute arming is not a decision, it is motion. It grades neutral on any per-decision ledger because it costs \$0, and here are twenty-one of them.

The other side of that coin, and I want to be fair to the router: **the six long windows that fired were not uniformly good.** Two of them lost money badly — Wednesday’s 7-hour `abs_veto_short` carve-out (–\$274.00) and Wednesday’s 9-hour `exhaustion_short` carve-out (–\$141.00), both of which came from the operator’s 13:33 do-not-bench instruction on the worst day of the week. So “longer is better” is not the finding. The finding is narrower and safer: **short windows produce nothing at all, so the churn is pure cost in attention and log noise.** **SHADOW** a minimum dwell time — write it, measure how many of the 21 sub-5-minute armings it would have prevented, and check it does not block the 15:25 Friday arm that fired three minutes later for +\$94.50. That one is the counter-example a dwell rule has to survive.

WHO MADE THE CHANGE	WINDOWS OPENED	THAT FIRED	NET OF WHAT FIRED
Durable headless tick	24	7	+\$486.00
22:00 auto re-activation	24	2	+\$83.00

That last row will read badly and it needs its caveat immediately: **the session only intervenes on the hard days**. Every one of its losing windows was Wednesday afternoon or Friday's open — the two most violent windows of the week. The durable tick got the quiet hours, which is where the +\$486 came from. This is selection, not skill difference, and I would not use it to argue for less hand-intervention. What it does argue for is that **hand-intervention at the cash open is the most expensive thing anyone did this week**, which is the same conclusion the 13:30–14:45 table reached from a completely different direction.

STOP_UNFILLED is gone. Three other ways to lose money took its place.

Straight answer on the self-heals: **there were none, because there was nothing to heal**. The desk has 26 STOP_UNFILLED exits in its entire history and the **last one was 2026-07-28 at 08:02**. Zero this week, zero last week. The tape-primary stop guard closed that leak and it has stayed closed for ten days. That is a genuine **FIXED** and it should be said plainly, because it used to be the desk's single biggest source of malfunction losses (nine of them on 07-23 alone, for −\$747).

What replaced it is a different family entirely: **an order or a position outliving the thing that owned it**, on an account both desks share. Three of them this week, all real money.

WHEN	WHAT	BOOKED AS	COST
08-06 07:08	The day-rider's hard-flat branch flattened the account net without checking whose position it was, closing the tournament's two abs_veto_long lots and its own	day_rider_crossdesk x2	−\$95.50
08-06 13:47	A leftover GTC stop from that morning's mess fired with no position behind it and opened a naked short — 38 minutes after nipc had been switched off	nipc_short x2	−\$34.50
08-07 13:30	A redundant market close was forgotten rather than cancelled when the stop won the race, and filled into a phantom opposite position	abs_veto_short_B LONG	−\$35.50

All three are fixed with regression tests, so none of this is a Saturday action. The one that is **still open** and belongs in front of you: the desk audits “does every slot have a live stop?” and never the inverse, “does every live stop have a slot?”. The 13:47 ghost is exactly what that gap produces, and it is invisible to the per-slot auditor. **OPEN** — check TWS for working stops with nothing behind them.

What I could not measure, and why you should be suspicious of some past numbers

Four instruments returned confident wrong answers or no answer at all this week. I re-derived each of these rather than taking them on trust.

INSTRUMENT	WHAT I FOUND	CONSEQUENCE
signal_journal	702 rows for the week. Every single one has taken NULL and suppressed_by NULL , despite both being in the INSERT	There is no way to tell a gate that had no setup from a gate that was blocked. Any past claim of the form “X signals were taken” is unfounded
capture.db retention	Bars go back to 07-15, but ticks, quotes and book start at exactly 2026-08-03 — five trading days	The nightly claim audit priced 2 of 14 claims; the other 12 return no_ticks. Anything that ATTACHes capture.db and looks back more than a week is silently reading nothing
router_trial_log.txt	It uses two different line formats and one entry is literally timestamped 2026-08-05T13:5xZ — the minute is not recorded	A first parse of the log missed nine operator changes entirely (63 vs the true 72). Two nipc lots at 08-05 14:06 appear to fire after their bench, and the log cannot say whether that is true
Day-rider P&L	day_rider.py has no write path to the trade store, and day_rider_state.json is overwritten each session — it already reads session: 2026-08-08, entered: false	The only day_rider* rows that exist anywhere are the two cross-desk flatten rows above. Its P&L is a measurement gap, not \$0 — see its own section

The second one is worth a sentence on its own. **Twelve of fourteen operator claims could not be priced** because the ticks that would price them no longer exist locally. The two that could — both from Tuesday — say your hands took + \$79.50 where the machine’s own exits would have taken +\$128.50: a -\$49.00 delta, or -\$24.50 a claim. On the 14:07 claim you banked 92.7% of everything the trade ever offered and got out in 354 seconds where the machine would have held 1,260. **On n=2 that is a story, not a statistic**, and the honest headline is that the claim audit is currently blind to 86% of its own subject matter.

Dispositions from Part 1

Every lead this section touched, with exactly one verdict. PARKED entries name what would revive them.

LEAD	VERDICT	BASIS / WHAT WOULD CHANGE IT
abs_veto_long is the desk's working gate	LIVE	+\$566 / 32 lots / 62%; positive on all four leave-one-day-out runs; +\$363 after strip-best-3. Already armed — this is a HOLD, not an action. Do not size up on 16 entries
abs_veto_short should not be armed on the current evidence	PARKED	0-for-9, −\$559, widest stops on the desk, and it enters at the worst possible moment. Revive when its entry timing is characterised — on 08-06 it sold 29,253.75 at 13:33, 12.75 points off a day low of 29,241.00 that printed inside that same minute , and four minutes later price was 29,426.50. It is not shorting the break, it is shorting the climax of the flush. Measure that lag and you have a fix; arm it on shadow P&L and you repeat 08-06
nipc	KILLED	Kill criterion tripped at exactly −\$400.00 on lot 47 and was executed the same minute. In <code>reactivate_gates.py::HOLD</code> . Already done — NOT-AN-ACTION . Its true record is −\$400.00/47, not the −\$434.50 in the ledger
rgv_short is a gate that does not exist	PARKED	Zero fires; armed under 2 minutes all week; <code>atr_min</code> 20 against a week whose median 1-min ATR was 11-12. Revive on a genuinely high-vol week where ATR clears 20 for sustained periods — or delete the ATR floor and re-test. It has not been given a chance to fail
grind_long is being benched by attrition	PARKED	Armed 1 hour in 5 days; all 6 lots from one 11-minute window; −\$172. Revive with an explicit decision either way — the current state is 9 arm-windows nobody believes in. Its <code>ext_hi</code> of 2.0 also means it structurally cannot fire on an extended run, so most armings are guaranteed \$0
Arm for periods, not for runs	SHADOW	21 windows under 5 min produced 0 trades; all the profit came from windows over 2 hours. But 2 of the 6 long firing windows lost money, so it is not “longer is better”. Spec a minimum dwell time and check it against the 15:25 Friday arm that fired in 3 minutes for +\$94.50
Wider stops would have converted this week's losers	PARKED	×2.0 turns −\$2,542 into −\$1,047 on the 71 re-raced losers, control reproduces within 2%. But this is losers-only, one week, simplified 2R target, and 1.5× was already refuted on a bigger sample . Revive when the <code>sw_*_k10/sw_*_k20</code> forward A/B reaches $n \geq 40$
13:30-14:45 is where expectancy lives	WATCH	Standing finding ($n=562$, three tests passed). This week the window lost −\$1,005.50 on 58 lots and was negative on 4 of 5 days. 58 does not

overturn 562. Re-derive if it is negative again next week; do not change policy now

First live sample paid: exhaustion_short re-armed 14:25 Friday went 2-for-2 for +\$104.00. n=1. Accumulate

Zero this week; last occurrence 2026-07-28T08:02; 26 all-time. Ten clean days. NOT-AN-ACTION

The 08-06 13:47 ghost proves the inverse audit does not exist. Not a research question — a thing to go and check in TWS

702 rows, 0 populated. Until this is written, nothing on this desk can distinguish “no setup appeared” from “the filter blocked it”, which is the question every bench decision needs answered

The fader carve-out (do not bench a fader on a momentum rationale)

SHADOW

STOP_UNFILLED

FIXED

“Does every live stop have a slot?”

OPEN

signal_journal.taken / .suppressed_by

OPEN

IF YOU READ ONE PARAGRAPH OF PART 1

The desk lost \$772.50 over five sessions. **One gate made \$566, one gate lost \$559, and the difference between them is which side of the market they take.** The long book was +\$345.50 and the short book was – \$1,022.50 with three winners in thirty lots — and the short book was negative on every day, after dropping any day, and after stripping its three worst trades. Meanwhile the two biggest winning sessions both came from leaving one gate armed and not touching it for eighteen hours and seven hours respectively, while twenty-one separate arm-windows lasting under five minutes produced nothing at all. Nothing here is proven — it is five days, 53 entries, and I found all of it after the fact. But the two things I would act on are: **keep abs_veto_long armed and stop touching it**, and **do not arm abs_veto_short again until somebody has measured how late it enters** — on Thursday it sold within thirteen points of the day’s low, in the minute that low printed, and was stopped for –\$129.50 before the minute was out.



P1.5 LIVE-DESK REHABILITATION — never bench a gate at face value

Gaz, this is the section where nothing gets binned because it had a bad week. Six things went on the stand: the four gates that bled (`grind_long`, `abs_veto_short`, `exhaustion_short`, `nipc_short`) and the two exit mechanisms everyone has been complaining about (`MAX_HOLD`, and the whole chandelier/give-back managed-profit path). For each one the question was not “is it working” — obviously not, the desk booked **$-\$1,028$ raw / $-\$772.50$ after the data-quality exclusions** this week. The question was **“HOW can it work, and if it can’t, can I prove it can’t?”**

Four of the six came back rehabilitated and are going live on Saturday. One is a two-number config revert worth more than everything else combined. One is an arming policy that has been strangling a gate whose signal is the strongest thing on the desk. One is a two-year-old exit we replaced with a worse one. One is an exit mechanism nobody has been grading fairly — it was being blamed for damage caused by a broken stop. The other two are honest failures with dates and conditions attached, not shrugs.

And a warning up front, because it is the thread running through all six: **on three separate gates this week the “obvious” fix was a fake win** — a filter that made the number look better by deleting the trades that actually pay. Every table below carries a winners-kept column, and every fix that failed that test is named and rejected in Section 4, not quietly dropped.

THE ONE-PARAGRAPH VERSION

Rehabilitated and shipping: grind_long’s ATR floor goes 10→22 and its extension ceiling 2.0→3.0 — both reverting a 2026-08-01 change — turning **+\$1,930 on 744 legs into +\$5,893 on 155 legs** over eleven tick-honest sessions, and this week specifically **−\$288 → +\$2,178**. abs_veto_short stops being a 22-minute discretionary window and goes **armed by default**; it was armed 3.3% of the week, and in the other 96.7% the identically-configured twin made **+\$1,064.50**. exhaustion_short’s exit goes back to the fixed 8pt/12pt/120s it had on 07-25 — worth **+\$604 over six days on identical entries**. The chandelier width gets regime-keyed — worth **+\$3,049 over 20 sessions, positive in 20 of 20 leave-one-day-out folds**. MAX_HOLD keeps its 120 minutes (every attempt to shorten it is a fake win) but gets a ceiling that survives a dead loop, worth **+\$461.50** on the archive. **The honest failures:** exhaustion_short is sub-friction by about two dollars a lot and PARKED with a concrete revival condition; nipc_short is direction-beta wearing a state machine and goes to SHADOW; GIVEBACK is TRULY-RETIRED; and roughly thirty ideas tried this week are REFUTED with the test that killed them written next to each one.

1 · THE LEAD — what got rehabilitated, and what it is worth

Money first. These are measured on **different windows with different tapes**, so I have put the window next to each one and **I am deliberately not adding them up** — a single total would be a lie. Read each row on its own terms.

WHAT CHANGED	VERDICT	WORTH	OVER WHAT WINDOW	WHY I BELIEVE IT
	FIXED	+\$3,964	11 tick-honest sessions (07-24...08-07). This	Keeps 19 of 19 target winners. Positive in all

WHAT CHANGED	VERDICT	WORTH	OVER WHAT WINDOW	WHY I BELIEVE IT
grind_long: ATR floor 10 → 22, ext ceiling 2.0 → 3.0			week alone: −\$288 → +\$2,178	four ISO weeks (live config is negative in two). Survives deleting its five best trades (+\$3,593). 10 of 11 sessions green.
abs_veto_short: armed by default, one new entry veto, two bad bench rules deleted	FIXED	+\$1,064.50	the 96.7% of this week it was benched, measured on the exact-config twin	The signal is +\$2,290.50 over 159 fires / 17 days / 4-of-4 weeks green / positive in all five regimes. 53 of 54 filters tested LOSE to just letting it fire.
chandelier: regime-key the width (ER30≥0.25 & ATR≥22 → Lot B 6R)	REGIME-DEPENDENT	+\$3,049	20 sessions (07-15... 08-07). This week: −\$1,073 → −\$643	20 of 20 leave-one-day-out folds positive, minimum +\$1,604. Better in all four ISO weeks. Delta survives 2pt of slippage.
exhaustion_short: exit back to fixed 8pt / 12pt / 120s	FIXED	+\$604	6 dense-book days, identical entries, mean of 10 polling phases	All six fixed-point exit cells beat the current ATR-scaled one. The move is a fixed number of points; an ATR-scaled stop on a 40pt-ATR open is risking 40 to make 31.
MAX_HOLD: keep 120 min, but enforce the ceiling at the venue	FIXED	+\$461.50	the whole archive (2 trades). \$0 this week — nothing bound	The cap fired 159 minutes late once and never at all once. Trade id54 was +\$225.50 at the 120-minute mark and closed at −\$141.00.
nipc_long + nipc_short: add the missing lo clip block	FIXED gate stays off	+\$268	the live week, out-of-sample, on the reference book (−\$359 → −\$91)	Keeps 22 of 22 winners. Helps the LONG leg more (+\$463→+\$668, win rate 43.3%→50.9%). Copied verbatim from mature gates — not fitted.
Desk-wide: stop-limit 20pt buffer → phantom fills	SHIP	+\$69.50	this week, 2 lots — recurs every violent open	Found independently by two rehabs. Two stops filled at their limit, 20pt through their own trigger, at prices the tape never printed.

What that does NOT mean. None of these makes the desk profitable on its own. The best exit policy I could prove takes the archive from −\$2,932 to **break-even (+\$117)**, not to profit. The grind_long fix is the only one of the seven that is a genuine money-maker rather than a bleed-stopper, and even that one **does not save 08-05** — in the exact window the desk traded it still loses −\$178.

2 · THE SIX STORIES

grind_long — 1 win in 6, −\$172.00 FIXED

THE SYMPTOM

Six legs, one winner, five stops, −\$172.00. grind_long_B went 0-for-3. Banned by hand at the Paris reopen on 08-07.

MALFUNCTION NORMALIZATION — THERE ARE NONE, AND THAT IS THE POINT

All six legs audited tick-by-tick against capture.db. **No STOP_UNFILLED, no naked rides, no orphan entries, no cross-desk leakage.** Every fill printed on the tape inside ± 1.5 seconds. Longest hold 150 seconds. **The venue's −\$172.00 is the strategy's −\$172.00.** This is the rare week where the software behaved and the strategy simply lost.

But there IS a malfunction and it is a governance one. The entire week's sample is an **11-minute hand-armed window** at the US open on the one day of five the gate loses money. The signal came off the 13:32 bar, which closed at 13:33:00 with price at **29956.50**. The desk did not fill until 13:33:24 at **29999.75** — a **43.25-point chase** into a window whose whole range was 109.75pt. Taking the identical signal at the legal first decision tick banks **+\$41.50 and +\$73.50**; live, 24 seconds later, both lots took −\$44.00.

That single pair is a \$203 swing caused purely by when the switch got flipped — 118% of the week's whole reported loss. I have not normalized it away (a chase you paid for is real money), but it belongs on the arming ledger, not on the gate's card.

ROOT CAUSE — BASE, NOT EXIT, NOT SIGNAL

The whole gate in one table. This is when a runner exists at all:

ENTRY ATR BAND	N	NET \$	MEAN R/TRADE	% REACHING 2R	% REACHING 4R	% REACHING 6R
0-14	289	-36.5	-0.010	20%	0%	0%
14-18	221	-1,428.9	-0.204	4%	0%	0%
18-22	114	+308.8	+0.058	1%	0%	0%
22-24	24	-328.8	-0.284	33%	4%	0%
24-32	53	+272.3	+0.126	36%	15%	4%
32+	43	+3,143.1	+0.903	47%	16%	14%

Below an entry ATR of about 22, not one trade in 624 reached 4R. Above 24, one in six does. grind_long's edge is a >4R runner on Lot B, and a runner is a tape phenomenon, not an exit phenomenon — it needs enough volatility that four risk units of room actually exist inside the next hour. Below ~22pt that room does not exist, so 624 trades pay full friction for a payoff that cannot happen. It survives risk-normalisation too: mean R per trade is roughly zero below ATR 22 and **+0.90 at ATR 32+**. It is not “bigger ATR, bigger dollars, same edge” — the edge itself is different.

And it is definitely not the exit. I ran **seven** different exit configurations against the sub-ATR-18 population, where 73% of the trades and all of the bleed live. **All seven lose.** The best of them still bleeds \$0.88 a trade over 414 trades. You cannot exit your way out of an entry with no runner in it.

THE FIXES TRIED — AND THE KEEP-THE-WINNERS TEST

Nineteen legs cleared +\$150. They are **284% of the book's net** — take them out and the live config is -\$3,549. **Every single one has entry ATR ≥ 24.6.** Eighteen of nineteen are ≥ 27.5. And thirteen of nineteen fired at ER30 below 0.15.

FILTER	N	NET	\$/LEG	STRIP-BEST-3	DAYS GREEN	WINNERS KEPT
LIVE (ATR≥10, ext≤2.0)	744	+\$1,929.9	+\$2.59	+\$472	6/11	19/19
ER floor 0.20 — the fake win	132	-\$1,637	-\$12.40	-\$2,298	2/11	2/19
ER floor 0.35 (the one deleted 08-01)	24	+\$38.5	+\$1.60	-\$623	2/4	2/19
ER15 floor 0.25	364	-\$942	-\$2.59	-\$1,937	4/11	8/19
box ≥ 0.85 (“buy the break”)	338	-\$515	-\$1.52	-\$1,724	—	—
ATR ≥ 18	234	+\$3,395	+\$14.51	+\$1,938	8/11	19/19

FILTER	N	NET	\$/LEG	STRIP-BEST-3	DAYS GREEN	WINNERS KEPT
ATR $\geq$ 20	163	+\$3,522	+\$21.61	+\$2,064	9/11	19/19
ATR $\geq$ 22 (recommended)	120	+\$3,087	+\$25.72	+\$1,629	6/11	19/19
ATR $\geq$ 24 (the pre-08-01 floor)	96	+\$3,415	+\$35.58	+\$1,958	—	19/19
ATR $\geq$ 32	43	+\$3,143	+\$73.09	+\$1,686	7/11	17/19

An ATR floor is the only filter in the whole battery that cuts 87% of the trades and keeps 19 of 19 winners. That is what a real filter looks like: it removes population that had no path to the payoff and touches nothing that did. The ER floor is the textbook fake — it cuts the trade count by 82% and gets there by binbagging **17 of the 19 winners**, keeping the two smallest.

★ **THE FAKE WIN, NAMED — AND WHY THE “OBVIOUS” INVERSION IS ALSO REJECTED**

Somebody will notice that an ER ceiling at 0.20 looks great in-sample (+\$3,567, 17/19 winners kept) and propose that instead. Don't. **The ER filter's sign flips between the two halves of the tape:** in-sample the ceiling wins by \$4,500; out-of-sample the floor wins by \$2,600 (ER $\geq$ 0.20 = +\$1,640 on 61 legs vs ER<0.20 = −\$973 on 321). That is not a parameter with a robust optimum, it is a coin. **ER is not a usable discriminator for this gate in either direction**, and this is the second independent re-derivation to say so. The 08-01 deletion stands permanently — but the reason is not “the floor blocked profitable trades”, it is “ER has no stable sign here at all”.

ROBUSTNESS

TEST	LIVE CONFIG	PROPOSED CONFIG
11 tick-honest sessions	+\$1,930 / 744 legs	+\$5,893 / 155 legs
\$ per leg	+\$2.59	+\$38.02
after deleting its 5 best trades	−\$347	+\$3,593
leave-one-day-out (worst fold)	+\$949	+\$4,656
sessions green	6 of 11	10 of 11
all four ISO weeks	2 negative	4 positive
out-of-sample (10 bar-replay days)	+\$667 / 382	+\$796 / 92
runners captured (MFE $\geq$ 4R)	16	21

TEST	LIVE CONFIG	PROPOSED CONFIG
placebo null, 300 random-entry draws	88th pct — does not clear	100th pct
quadruple fees (\$6/RT) + 3pt stop slip	—	+\$4,626 (strip-3 +\$3,182)

It is a **plateau, not a spike** — everything in the block ATR 18–24 × ext 2.5–6.0 is in-sample $\geq$ \$3,400 with strip-3 $\geq$ \$1,900, and the out-of-sample grid puts its own best in the same neighbourhood. **The live cell sits in the worst corner of both grids.** The exact number is deliberately not load-bearing: 20/24 and 2.5/4.0 all work.

The other finding, which I did not go looking for. Score the fixed gate against the router's actual arming schedule this week: **–\$210.60 in the ~45 minutes it was armed, +\$2,388.10 in the four days it was benched.** The router benched it on all four days the fixed gate would have earned and armed it for the eleven minutes on the day it lost. That is not slightly off — it is **anti-correlated**, because the router arms on ER climbing and the gate's money is at ATR $\geq$ 22 with ER below 0.25. Those are close to opposite conditions. I am **not** asking you to change the router this week; that is its own job.

VERDICT

FIXED Two numbers, both reverting a 2026-08-01 decision:
 ATR_FL00R["grind_long"] = 22.0 (was 10.0) and ext_hi = 3.0 (was 2.0).
Do not touch anything else — the exit, the chandelier, the lock (lock_r 6.0 has the best strip-3 in its own sweep), the 1.0×ATR stop, Lot A's 2.5R, fast_slope=True and slope_min 0.4 all survive their own sweeps. **A fifth of the trades, three times the money, and it stops being a churner.**

Honest limits: 155 legs is 78 signals across 11 sessions and the runners that carry it are 21 events. Three-and-a-half weeks, one instrument, summer tape. The out-of-sample leg is bar-replayed and runs ~17% optimistic. And there is one hole I cannot explain: **grind_long is 0-for-14 in its own namesake regime** (ER30 $\geq$ 0.45, all stops, three episodes). n=14 cannot settle it and I am not pretending otherwise.

abs_veto_short — 11 closes, 11 losers, −\$814.50 FIXED

THE SYMPTOM

The week’s worst bleeder. Eleven closes, zero winners, −\$814.50 across four sessions.

MALFUNCTION NORMALIZATION — 36% OF THE LOSS IS NOT THE STRATEGY

STEP	CLOSES	NET
RAW as booked	11	−\$814.50
less the MD-stream ATR corruption (08-04 22:02) — struck, no valid signal existed	2	+\$255.50
less the phantom LONG opened by an orphaned market close on a gate called short	1	+\$35.50
TRUE STRATEGY BOOK	8	−\$523.50

Eight closes is **four signals × two lots**. On 08-04, gold joined the capture stream and got folded into the MNQ bar deque — ATR read **1,848** against a true **15.11**, a factor of 122, and the “five-bar thrust” the gate saw was a 26,500-point drop from MNQ to gold. It was not a down-thrust, it was a symbol change. The whole shadow desk was up that night and **not one** abs_veto variant fired at 22:02. All eleven rows re-derived from prices: **0 of 11 need restating** — the money is real, it just is not all strategy money.

ROOT CAUSE — NOT THE BASE, NOT THE EXIT, NOT THE SIGNAL

The signal is the strongest thing on the desk. Tick-repriced over 17 trading days:

TEST	RESULT
all 17 days	159 fires, +\$2,290.50, 48.4% win, +\$14.40/trade
strip best 1 / 3 / 5	+\$2,100 / +\$1,780 / +\$1,485
per ISO week	wk29 +\$326 · wk30 +\$497 · wk31 +\$659 · wk32 +\$809 — rising, all green
out-of-sample (fit wk29-31, test wk32)	+\$1,481 → +\$809
per regime	positive in all five

And the brief’s suggested fix was a null before I computed a single number. “Give it the k10 stop width its best shadow uses” — it already has

it. I read the live slate build rather than assuming: live Lot A is 1.5R on a 1.0×ATR stop, Lot B is 2.5R on a 1.0×ATR stop, plus the quiet clip. That is sw_absS_A_k10 and sw_absS_B_k10 exactly. The best shadow in the family is not running a different exit — it is running the same exit at different moments. (And the width is right anyway: on 28 shared entries the live k1.0 pair beats k2.0 by **\$292.50**.)

So the leak is the leash. The gate was armed for **224 minutes out of 6,780 — 3.3% of the week** — in ten windows averaging 22 minutes, the shortest of them three minutes. In the other 96.7% the same signal on the same exit made **+\$1,064.50**.

THE HEADLINE TEST THE BRIEF ASKED FOR — AND IT CAME BACK THE OTHER WAY

EXACT-CONFIG TWIN, WEEK 08-03→08-07	N	NET	WIN%	\$/TRADE
ARMED (gate could trade)	6	+\$289.50	66.7%	+\$48.20
BENCHED (gate stood down)	54	+\$1,064.50	57.4%	+\$19.70
whole week, unrouted	60	+\$1,354.00	58.3%	+\$22.60

I was told the arming policy was selecting the gate’s losing minutes. **It is not true.** The armed minutes were, if anything, the best minutes of the week — four times better per trade than the benched ones, and every armed window that produced a signal was green. n is small (6 twin fires), so the honest claim is the negative one: **there is no evidence the arming policy picked losing minutes.** The money leaked because the leash was 3.3% long and four consecutive stopped signals — roughly a 1-in-14 run on a gate with a +\$14.40 expectancy — was the whole week.

Two more things fell out of the arming ledger. **–\$529.50 of the –\$814.50 was booked in windows the router did not choose.** The 08-04 pair came from the Paris-midnight reactivate_gates cron re-arming every off gate at 22:00 UTC on zero tape (“insufficient bars, range 0pt, day P&L flat” — the router’s own words that night); the corrupt thrust fired **two minutes later** and ran unsupervised for 45 minutes. The 08-05 pair came from an operator DO-NOT-BENCH carve-out that explicitly removed the router’s bench

authority; the router watched the gate bleed, said so in its logs, and could not act.

54 FILTERS TESTED. 53 OF THEM LOSE MONEY.

The 15 biggest winners carry **+\$1,664 of the +\$2,290.50 (73%)**, so winner-retention is the column that decides.

POLICY	KEEP N	KEEP \$	\$/TR	CUT \$ (MONEY THROWN AWAY)	TOP-15 KEPT	STRIP-3
BASELINE — just let it fire	159	+\$2,290.50	+\$14.40	—	15/15	+\$1,780
stacked router rules (what was live)	67	+\$881	+\$13.10	+\$1,409.50	8/15	+\$373
router “arm trio” (extreme + ER rising + vol up)	31	+\$618	+\$19.90	+\$1,672.50	4/15	+\$286
ER30 ≥ 0.30 — best \$/trade in the sweep	51	+\$1,140.50	+\$22.40	+\$1,150	7/15	+\$633
dwelt: ER30≥0.30 in ≥3 of 5 min	27	+\$937	+\$34.70	+\$1,353.50	4/15	+\$503
ATR ceiling < 19pt (“bench the violent tape”)	120	+\$1,254	+\$10.40	+\$1,036.50	0/15	+\$1,037
ATR floor ≥ 15pt	82	+\$1,829	+\$22.30	+\$462	15/15	+\$1,319
★ drop BUILDING × US-SESSION	151	+\$2,497.50	+\$16.50	-\$207	15/15	+\$1,987

The single best policy is no policy. The baseline is beaten exactly once, by a rule that removes **eight fires in seventeen days** and damages not one winner. Every dwell and ER rung wins on \$/trade by amputating the winners — D3 hits the best per-trade number in the entire sweep (\$34.70) by keeping **4 of 15** winners. And the most instructive failure is the ATR ceiling: benching above 19pt ATR keeps **zero** of the fifteen biggest winners. The gate’s whole edge lives in high-ATR tape.

The bench rule that is actively backwards. The router benches this gate whenever $ATR \geq 19$ AND $ER30 < 0.25$ — the “violent-whipsaw” cell. Measured over 17 days and split by the clock: in the **US session that cell is +\$37.00 a trade, 58% win, n=12, and holds 6 of the 15 biggest winners in the book.** Overnight it is $-\$3.20$ on $n=6$. The rule is right

overnight, wrong in the session, and applied without reference to the time of day.

ROBUSTNESS

TEST	BASELINE (CONTINUOUS)	RECOMMENDED	THE POLICY THAT WAS LIVE
n / net	159 / +\$2,290.50	151 / +\$2,497.50	67 / +\$881
strip best 5	+\$1,485	+\$1,692	+\$195
worst leave-one-day-out	+\$1,889	+\$2,096	+\$333
out-of-sample (wk29-31 → wk32)	→ +\$809	→ +\$913.50	→ -\$22.50
top-15 winners kept	15/15	15/15	8/15

The recommended policy dominates the baseline on every line, and both dominate the policy that was actually live — **which is the only one of the three that is negative out of sample.**

VERDICT

FIXED Signal: no change. Exit: no change. **The arming policy is the change.** Arm it by default; add one per-entry veto (skip when ER30 is between 0.15 and 0.30 and it is 13:30–20:00 UTC); delete the US-session violent-whipsaw bench; delete “bench on a stop-out pair” (on 08-07 the 13:31 bench after the first stop-out pair was followed by **+\$311.50** to the twin in the next 60 minutes); and fix the 22:00 auto re-arm so it cannot arm anything on “insufficient bars, range 0pt”.

What I am NOT claiming. I have not proven a longer leash would have made money this week — the four signals it took would still have lost about $-\$256$ at a neutral entry draw. The claim is about the 54 fires it was benched for. The BUILDING×US veto rests on **n=8**; it cross-validates across nine independent exits and costs almost nothing, but it is not a heavyweight finding. And the 61-second live entry lag is measured and real (median 61s over 78 matched entries, $\sim \$6.87/\text{lot}$ on average with a $\pm \$21$ band) but its cause in code is not proven.

chandelier / give-back — the managed-profit exit path went silent REGIME-DEPENDENT TRULY-RETIRE (give-back)

THE SYMPTOM

From 07-10 to 07-30 these two exits booked **145 fires for +\$7,156**. Across 08-03-08-07 the whole book produced **one** chandelier exit (+\$50) out of 110 closes. GIVEBACK last fired **2026-07-27**. The week was binary: 71 STOP / -\$2,600.50 with zero winning stops, 34 TARGET / +\$1,794.

MALFUNCTION NORMALIZATION

New and previously unreported: a 122× ATR poisoning that removed the stop AND both targets. The 08-04 22:02 pair — the largest single loss of the week — has never been attributed. The protective stops were armed **1,848.25 points** above the entries because the desk read $ATR = 1848.25$ against a true 14.73. The desk restarted at 22:07 with the filter in place, but `reconstruct()` re-adopted the two open slots with the poisoned ATR persisted. So for the next 115 minutes the stop was effectively naked, Lot A's target sat 2,772pt away, Lot B's 4,621pt away, and `atr_split 22` compared $1848 < 22$ and said no. **The position had no exit of any kind.** The tape breached the correct stop level three minutes after entry — 117 minutes before the flatten. Normalized swing **+\$192.50**.

And honesty runs both ways: **one stop-guard miss went the desk's way and I normalized it adversely.** Trade 582 printed through its stop 12 times, to 4.14pt below it, in the 08:30 ET release burst; the 1Hz cadence missed it, the desk held, and banked +\$108.50. Restated to the stop it should have taken: **-\$131.20**. Net corrected week: **-\$801.70**.

ROOT CAUSE — THE MACHINERY IS NOT BELOW A THRESHOLD, IT IS SWITCHED OFF

The scope's hypothesis was "the arm threshold is never being reached". It is not a threshold problem. **On 93% of this week's positions the chandelier is not in the code path at all, and the give-back cannot execute under any tape.**

EVERY LOT-B POSITION ROUTES TO...	THIS WEEK (58)	ARCHIVE (495)
quiet-tape clip (atr_split 22)	24 (41%)	179 (36%)
fixed 2.5R scalp	28 (48%)	32
fixed 1.5R scalp	2	15
chandelier "tight" k1.5	2	237
chandelier-lock "wide"	2	32
any chandelier at all	4 of 58 (7%)	269 of 495 (54%)

Four eligible positions, one fired. The quiet-tape clip runs first and returns unconditionally — any gate with atr_split set and ATR below 22 has no chandelier, no give-back, no adaptive selector, whatever its rung says. **And atr_split = 22 keys on the wrong variable:** of the entries with ER30 $\geq$ 0.30, 66% get clipped — and of the entries with ER30 $<$ 0.30, also 66%. The clip's selector is statistically independent of the regime it exists to protect against. Give-back, meanwhile, was killed twice over: the 07-27 adaptive-exit roster deploy guards it out, and the 07-29 scale-out slate hard-sets giveback_enabled=False on all 16 sub-slots. **16 startups $\times$ 16 slots: giveback_enabled is false on 256 of 256 resolved specs.** Its silence is the design, not a symptom.

THE FIX THAT WAS PROPOSED — AND WHY IT FAILED ON ITS OWN TERMS

The rehab request was “restore the pre-08-01 managed-profit shape”. I ran it:

THIS WEEK'S 110 LOTS, REPRICED ON REAL PRINTS	BOOKED	DEPLOYED REPLAY	PRE-08-01 MANAGED PROFIT	REGIME-KEYED FIX
the 34 TARGET lots	+\$1,794.0	+\$1,527.5	+\$1,472.0	+\$1,720.6
the 71 STOP lots	-\$2,600.5	-\$2,510.6	-\$2,260.0	-\$2,403.2
all 110	-\$932.5	-\$848.6	-\$387.7	-\$402.3

Handing this week's 34 winners to the old management makes them worth \$322 LESS, not more. The whole gain of the managed-profit arm comes from the stop side (the trail cuts some losers before 1R) and from the 08-04 normalization — **not** from captured runners. Why: 18 of the 34 winners do eventually reach 6R, but **19 of those 34 round-trip all the way back to the 1R stop** before the two-hour cap, and on a 10-16pt R the trail sits far enough away that it waits through the entire retrace. Slot-

honest over the whole archive, restoring it is **−\$3,433 vs the deployed −\$2,932, losing in 20 of 20 leave-one-day-out folds.**

WHAT THE DESK ACTUALLY FORFEITED — THE WIDTH, IN ONE REGIME

Single-lot net by regime cell, over the archive. Read the last column:

LOT-B EXIT	ER<.30 / ATR<22 (N=213)	ER<.30 / ATR≥22 (N=108)	ER≥.30 / ATR<22 (N=114)	ER≥.30 / ATR≥22 (N=60)
clip \$40 (deployed Lot A shape)	−\$164 (2nd/28)	−\$695	+\$40 (1st)	−\$345 (26th)
clip 1.75R/\$60 (deployed Lot B shape)	−\$238	−\$1,995 (27th)	−\$65	+\$676
fixed 4.0R	−\$699	−\$1,119	−\$680	+\$2,113
fixed 6.0R	−\$761	−\$318 (1st)	−\$631	+\$2,517
chandelier-lock “wide”	−\$1,379 (28th)	−\$1,246	−\$622	+\$1,768
give-back 50/40	−\$1,013	−\$1,490	−\$1,059	−\$41 (24th)

The entire managed-profit edge of the desk lives in ONE cell — high ER and high ATR, 12% of signals. There the deployed \$40 clip ranks 26th of 28 and a wide rung ranks 1st: a \$3,055 swing on 60 entries. In the other three cells the wide trail is 28th, 8th and 14th while the deployed clip is 2nd, 3rd and 1st. **The deployed config is not wrong about clipping — it is wrong about when.** And note which shape wins: holding the regime gate and Lot A constant, a **fixed 6R rung recovers +\$3,049; the lock-chandelier recovers +\$2,144; the tight chandelier +\$761; give-back +\$581.** What we forfeited on 08-01 was the width, not the trail.

ROBUSTNESS — INCLUDING THE TEST THAT FAILED

TEST	DEPLOYED	RECOMMENDED
full archive, 20 sessions	−\$2,932	+\$117 (+\$3,049)
leave-one-day-out, 20 folds	—	20/20 positive, min +\$1,604
per ISO week W29/30/31/32	−910 / −2,632 / +1,683 / −1,073	−840 / −2,587 / +4,187 / −643 — better in 4/4
out-of-sample (fit 07-15...07-31 → test this week)	−\$1,073	−\$643 — still negative
strip top 1 / 3 / 8 / 20 lots	—	+2,781 / +2,374 / +1,546 / +305
parameter plateau test	—	FAILED — only 4 of 35 threshold cells are positive

⚠ **READ THIS BEFORE YOU SIZE IT**

Inside the trend-capable cell, **6 of 16 sessions are green and 2026-07-27 alone is +\$1,452 of the +\$1,768**. Attributed by gate, exhaustion_short contributes +\$3,515 of which 07-27 is +\$3,280, and rgv_short +\$1,030 of which 07-23 is +\$1,374. **Two days carry it**. Strip-the-best inside the cell survives to top-5 (+\$323) and **inverts at top-8 (-\$422)** — the cell's edge rests on 5-7 trades out of 60. The direction is robust (the whole parameter neighbourhood averages ~+\$2,000 better than deployed); the knob is not.

VERDICT

TRULY-RETIRE **GIVEBACK**. Dominated in every regime, in all 42 sweep cells, in 0/20 folds. It earned +\$2,332 in July because it was the only profit exit on trail-less single-lot gates — and **61% of that came from gates that no longer exist** (rgv_long alone was 30 fires / +\$847). Under the scale-out slate every gate already has a fixed-R Lot A floor doing that job. Stop treating its silence as a symptom.

REGIME-DEPENDENT **the WIDTH ships** as a config-only edit (atr_split 22→24, every non-NIPC b→6.0 and a_r→3.0; NIPC keeps its own proven 2.0/2.5 pair). **SHADOW** **the trail machinery itself** — keep exit_chandelier_lock in the code (it is runner-up in the home cell and it is uncapped, which a fixed rung is not) but do not re-arm it on the strength of this report.

And the honest bottom line on cause: against a 200-trial random-side null the deployed desk sits at the **67th percentile** and the fixed desk at the **88th**. Neither clears 95. The exit explains about \$3,000 over 20 sessions and is worth fixing; **the signal explains the other \$3,000 and nothing in this rehab fixes it**.

MAX_HOLD — “the desk’s worst per-trade exit”, −\$255.50 on 2 fires FIXED

THE SYMPTOM — MEASURED ON ROWS THE DESK HAS ALREADY THROWN AWAY

Both of this week’s MAX_HOLD fires carry data_quality = EXCLUDE:md_stream_atr_corruption_20260804.pnl.py filters them at source. **The desk’s own canonical P&L books MAX_HOLD this week at zero fires and \$0.00.** The −\$255.50 comes from reading the raw trades table without the exclusion clause. Worth saying before anything else, because it means the headline “worst mechanism on the desk” is computed on two rows the desk formally does not count.

MALFUNCTION NORMALIZATION — AND THIS IS THE WHOLE FINDING

Widening the window: MAX_HOLD has fired **4 times all-time for −\$642.50**. Two are from the archive and nobody had looked at them. Four fires, four losers, no winners ever. Reconstructed on the 5s tape:

#	GATE	HOLD	BOOKED	MFE	A WORKING 1.0×ATR STOP FILLS...	MALFUNCTION COST
id114	rgv_short 07-21	120.1m	−\$171.00	−6.5pt	0.9 min in → −\$42.11	−\$128.89
id171	grind_long 07-22	278.9m	−\$216.00	+13.5pt	5.1 min in → −\$26.43	−\$189.57
id538	abs_veto_short_A 08-04	120.0m	−\$124.50	+14.3pt	3.0 min in → −\$30.00	−\$94.50
id539	abs_veto_short_B 08-04	120.0m	−\$131.00	+10.8pt	3.0 min in → −\$30.00	−\$101.00
total			−\$642.50		should have cost −\$128.54	−\$513.96

Read the MFE column: the best any of these ever looked was +14.25 points, and one was **never favourable by a single tick**. These are not positions that went good and then bad and the clock caught them late — they were wrong from the opening print and the machine did not close them. Read the

stop column: **0.9 / 5.1 / 3.0 / 3.0 minutes**. The two-hour clock is 24× to 133× later than the exit that should have happened.

Eighty percent of the MAX_HOLD book — \$513.96 — is a stop that did not work, not a clock that held too long. Grading MAX_HOLD on this damage is grading the fire alarm for the fire.

THE BRIEF’S MAIN ANGLE IS A FAKE WIN, AND I AM REJECTING IT

“Shorten the cap.” The blanket sweep looks fantastic. Decomposed, it is not:

CAP	SAVED ON LOSERS	LOST ON WINNERS	NET	TOP-20 WINNERS INTACT	TOP-20 POOL AFTER
5m	+\$4,976	−\$4,382	+\$594	4/20	\$1,864
10m	+\$2,534	−\$2,356	+\$178	7/20	\$2,658
15m	+\$1,908	−\$1,154	+\$754	12/20	\$3,556
30m	+\$955	−\$296	+\$660	13/20	\$3,841
60m	+\$271	−\$126	+\$146	18/20	\$4,168
90m	+\$490	\$0	+\$490	20/20	\$4,356
120m (live)	+\$460	\$0	+\$460	20/20	\$4,356

A 5-minute cap “makes” \$594 by destroying **57% of the desk’s best trades**. The runner audit says the same thing from the other side: at 15 minutes it clips **24 of the 51 runners**, including abs_veto_short_B 07-29 (\$536→\$294) and grind_long_B 07-29 (\$305→\$148.50). **A cap short enough to have caught the four malfunctions would have cost more than the malfunctions did.** Ironically, the 120-minute number you guessed is the only setting in the entire sweep that takes money purely from losers and nothing at all from winners — 20/20 top winners intact across 601 live and 7,473 shadow trades.

THE REAL DEFECT — THE CAP IS NOT ENFORCED

ID	GATE	HOLD	FIRE	BOOKED	P&L AT THE 120M MARK	COST OF LATENESS
id171	grind_long	278.9m	MAX_HOLD at 278.9m — 158.9 min late	−\$216.00	−\$121.00	−\$95.00
id54	thrust	277.1m	never (exited on ABSORPTION_CUT)	−\$141.00	+\$225.50	−\$366.50

ID	GATE	HOLD	FIRED	BOOKED	P&L AT THE 120M MARK	COST OF LATENCY
						-\$461.50

id54 is the sharpest single fact in this rehab. It was a winner at the cap — up \$225.50 at the 120-minute mark — and was allowed to run another 157 minutes and come back as a -\$141.00 loser, because nothing enforced the ceiling. The cap is only alive while the 5-second audit loop is alive. This is a liveness dependency, not a parameter choice.

ROBUSTNESS — AND THE DEFLATER

Enforcing the cap is worth **+\$460, and that advantage is exactly constant across strip-1, strip-3, strip-5 and strip-10.** That is not luck, it is structural: the gain comes entirely from truncating losses, so stripping the best trades cannot erode it. Its whole effect lands in the **overnight** bucket — the cap has never once bound a trade in the US session or pre-open.

The honest deflater. I ran a placebo: keep the same three bound trades but exit at a random moment in each trade's life. Random timing matched or beat the real cap in **161 of 400 draws. The 120-minute cap has no timing skill whatsoever.** Its +\$460 is not "120 minutes was the right moment" — it is "literally any exit would have been better than what actually happened". Which is the same finding as everything above, said a third way. And out of sample the cap is **completely inert — zero trades bound** in the test week. That is not a failure; it is the correct behaviour of insurance in a week where nothing broke. But it means **there is no out-of-sample evidence for any cap length, because there were no out-of-sample events.**

VERDICT

FIXED **Do not change the length.** Ship the enforcement: (1) arm the ceiling at the venue, not just in the loop, so it survives a dead process; (2) make a stale audit loop loud — if any slot is open past $0.9 \times$ the cap and the auditor has not completed a cycle, notify immediately; (3) book any fire more than 5 minutes past the ceiling as MAX_HOLD_LATE, not MAX_HOLD,

because right now a 278.9-minute hold and a 120.1-minute hold carry the same label — which is exactly why nobody noticed.

Two escalations this rehab produced by accident, and they are probably worth more than the rehab. (1) The tape-consistency screen struck **166 of 7,639 shadow rows across 14 separate days**, not just 08-04 — the MD contamination is a recurring condition in the shadow book and nothing flags it there. (2) **44.0% of shadow trades ran more than $0.25 \times \text{ATR}$ past their own stop** (median $0.56 \times$, max $32.32 \times$) on a book that is 99% $1.0 \times \text{ATR}$ stops. **Every shadow study that rewards cutting early is inflated by that leak.**

exhaustion_short — 2 wins in 6, $-\$162.00$, and the most mis-benched gate of the week PARKED exit FIXED

THE SYMPTOM, AND THE OPERATOR'S REAL COMPLAINT

$-\$162.00$ on six lots, with exhaustion_short_B at $-\$98.50$. But the real complaint was the arming: router 8 off / 5 on, with an explicit 08-07 13:24Z carve-out saying a momentum direction rule was being applied to a fader. Reconstructed from the switch log, **the gate was armed 10.0 hours out of 120 — 8.3% of the week**. Your read of the facts is completely correct.

MALFUNCTION NORMALIZATION

A phantom fill. Trade 587 was closed at a price MNQ never traded. Both lots entered at 29715.00 with stops at auxPrice=29733.25, lmtPrice=29753.25; both triggered in the same microsecond. One filled at 29737.75 (4.50pt through the trigger). The other filled at **29753.25 — exactly the limit, 20.00pt through**. The highest trade printed in the window was 29743.00 and the deepest offer across all five levels was 29743.25. **The fill is 10.25 points above the highest print.** That is the paper simulator filling a stop-limit at its limit when the sim book ran out. Restated to its twin's fill: **$-\$47.00$, swing $+\$31.00$** . The other 54 stop fills have a median slippage of 0.00pt, so this is a fat tail, not a distribution — and it happened twice this week (also abs_veto_short_B on 08-06, + $\$38.50$), both within four minutes of the cash open.

A signal that passed every filter and silently never became an order.

On 08-07 at 14:19:36 the journal logs both lots confirmed and “→ OPEN”, and then nothing. No order, no signals row, no error, no warning.

MultiSlotCore._open has three silent return paths and none of them logs.

Priced honestly the bug **made us money** (the pair would have stopped for −\$174.14) and I have not credited that back — you cannot bank a loss you avoided by accident. But **the desk cannot tell you why an approved entry never became an order**, and one log line fixes it.

And the gate’s famous best day is not what the ledger says. The badcall ledger records 07-30 as “4/4 +\$372.50”. It was **4-for-6** — twenty-two minutes after the last winner the same gate took another pair and both stopped, making the honest total **+\$217.00**. And every exit that day was MANUAL_CLAIM — your hand, not the gate’s. Priced on the tick archive, **the gate’s own machinery on the same six entries makes +\$65.10**. Seventy percent of its best-ever day was you on the exit button.

THE BENCH QUESTION, MEASURED WINDOW BY WINDOW

POLICY	LOTS/PHASE	NET	\$/LOT	PHASES GREEN
as actually armed (10 h)	11.2	−\$474	−\$42.34	0/10
always armed (all 120 h)	61.5	−\$614	−\$9.98	1/10

THE PREMISE OF THE REHAB IS REFUTED

Totalling only the trades that fall inside each benched interval: **benching exhaustion_short this week did not cost \$140. It SAVED \$140.** So the 13:24Z carve-out is right about the logic and wrong about the money. Applying “day-bias UP → bench shorts” to a fader IS a category error — it benches the fade exactly when the fade’s setup forms. But this particular fader was losing money in those windows anyway, so the category error was accidentally protective. Both are true and neither cancels the other.

ROOT CAUSE — IT IS THE SIGNAL, AND THE EXIT IS MAKING IT WORSE

The exit first, because it is fixable and worth \$604. Same entries, six days, ten polling phases, only the exit changed:

EXIT	LOTS	NET	\$/LOT	PHASES GREEN
LIVE today (A 0.75R / B tight chandelier / atr_split 22)	71	-\$847	-\$11.93	0/10
8pt stop / 12pt target / 120s (the 07-25 config)	66	-\$243	-\$3.67	2/10
12pt / 8pt / 120s	66	-\$121	-\$1.83	3/10
15pt / 20pt / 300s	66	-\$182	-\$2.76	3/10

The mechanism is obvious once you see the ATRs. This is a snap-back fade and the current exit hangs a **1.0×ATR** stop on it, at entries where ATR is 18–49 points. On 08-07 the desk was risking **41.8 points to make 31** on a move whose entire measured life is a handful of points and a few seconds. The favourable excursion never scales with the stop: at ATR 12–18 the average trade reaches 1.13R in its favour; at ATR 30–40 it reaches **0.60R** while the adverse excursion sits at 0.80R. The move is roughly a fixed number of points, so an ATR-scaled stop and target both become nonsense as volatility rises. **A fixed-point exit is structurally the right shape and we replaced it with an ATR one on 07-29.**

Now the signal, and this is the verdict. Against a 200-trial placebo of random entries in the same sessions:

HORIZON	SIGNAL MEAN (PT)	RANDOM MEAN (PT)	WHERE THE SIGNAL SITS
15 s	−0.64	−0.03	10th pct
30 s	−1.15	−0.00	2nd pct
60 s	−2.56	−0.18	0th pct
120 s	−4.36	−0.75	0th pct

The entry is significantly ANTI-predictive out to two minutes. Not “no edge” — negative edge. Shorting here is measurably worse than shorting at a coin-flip second. And the reason is our own five-second wait: decomposing the trigger’s life, **the fade happens during the five seconds we spend confirming it (+1.88pt) and it is over by the time we are allowed in (0→−1.15pt at every horizon after).** The confirm-veto is not the villain — the fires it rejects drift −5.25pt against you inside the window — but it proves

the shape: **this signal's half-life is under five seconds, shorter than our own confirmation latency.**

Two more nails. **The wall — the mechanism's actual claim — does nothing:** removing the wall test improves 60s and 120s forward drift and 2.5×s the sample; measuring it correctly (there is a genuine off-by-one, the code reads the second price level and calls it the touch) makes it worse; demanding a bigger wall is catastrophic. And **the independent shadow twin agrees:** 467 trades over 19 sessions, short side −\$2.85/trade, long side −\$3.16 — two completely independent implementations agreeing to within a dollar a trade.

42 FILTERS TESTED. NOT ONE IS A REAL WIN.

The eight target winners had ATR 32–49pt and extension below +1.0 ATR. Every cell that gets the number near zero deletes them:

FILTER	LOTS	NET	STRIP-3	WINNERS KEPT	VERDICT
baseline		71	−\$847	−\$1,188	8/8 —
ext floor ≥ 1.5 ATR — the seductive one	15	−\$3	−\$184	0/8	FAKE WIN
ATR ceiling ≤ 18	22	+\$14	−\$169	0/8	FAKE WIN
regime: drop violent-whipsaw	44	−\$729	−\$951	1/8	FAKE WIN
ATR floors 10...35 (7 rungs)	21–66	−\$660 ... −\$865	all negative	8/8	NULL
counter-regime veto (already live)	71	turning it OFF costs −\$311	—	8/8	HOLD

Zero of forty-two clears strip-best-3. The one thing that genuinely earns is the veto we already have — which is, note, precisely the mechanism-based bench predicate the brief asked me to build. It shipped on 07-28.

THE ONE THAT CAME CLOSEST — AND EXACTLY WHY IT FAILS

Enter at the fire instant, skip the five-second confirm entirely (since that is where the move is), fixed-point stop and target, hard time cap:

STOP / TARGET / CAP	LOTS	NET	STRIP-3	\$/LOT	PHASES GREEN
10pt / 5pt / 30s — with our real costs	171	−\$330	−\$355	−\$1.93	1/10
	171	+\$232	+\$202	+\$1.36	8/10

STOP / TARGET / CAP	LOTS	NET	STRIP-3	\$/LOT	PHASES GREEN
------------------------	------	-----	---------	--------	-----------------

the same cell with
ZERO friction

That is the whole story in two rows. Gross, the construct makes +\$1.36 a lot, is green in 8 of 10 sampling phases and survives strip-best-3. Net of a \$1.50 fee, 0.75pt entry slip and 1.0pt stop slip it makes −\$1.93 a lot.

Friction costs \$562 over the window; the edge is worth \$232. It is not broken — it is sub-friction by roughly two dollars a lot, and no arming rule, no regime filter and no exit target closes a gap that large. Even with free trading and perfect fills the live configuration still loses −\$382.

VERDICT

EXIT FIXED — SHIP Revert to the fixed 8/12/120s. **△ And remove exhaustion_short from _BIG_RUN at the same time** — otherwise deleting the override silently hands it the scale-out slate's A@2.5R + B wide-chandelier, which the sweep says is the worst Lot B in every regime. That is the single most important line in this card.

PARKED as a live gate Revive when (a) a zero-latency execution path exists — fill inside the 20-second window, not five seconds after it — **and** (b) round-trip cost falls below about \$0.50. Arming is not the lever. Leave the 13:24Z carve-out exactly as it is; it costs ~nothing to be present for a rare setup, and this report has no evidence a fifth bench rule would help. **The one thing I would ask: with the exit reverted, the next ~20 fires are a clean re-test of the gate as originally rehabilitated. Judge it on those.**

A banked lesson that needs amending. [[exhaustion-short-live-fader-success]] — “judge footprint faders on LIVE P&L, the shadow understates them” — was banked when the shadow was −\$194 on 145 trades. At 467 trades the shadow (−\$2.85/trade) is kinder than the live dual-slot book (−\$19.49/lot). The shadow was not wrong. It was early.

△ The measurement caveat that governs every number in this card. This trigger is a sub-second knife-edge event and the desk only looks once a second. Slide the polling phase by 100ms and you get a different set of trades: the same config on the same six days ranges from −\$365 to −\$1,711 across ten phases, **standard deviation \$418.** Every headline here is the mean of all ten. Any “improvement” smaller than about \$800 over six days is inside the noise of when we happened to look at the book.

nipc_short — 1 win in 16 (6%), −\$337.00 SHADOW

the 10 block FIXED

MALFUNCTION NORMALIZATION — AND IT MAKES THE SYMPTOM WORSE

Two of the sixteen are not nipc_short trades at all. Trades 577/578 carry order ids stp-*, not v7-mnq-*. Of all 49 nipc entry legs, 47 carry the normal namespace and exactly these two carry the protective-stop namespace. They have a filled row with no submitted. No NIPC setup exists at that time on the tape. And nipc had been router-benched 38 minutes earlier. **Two orphaned GTC stops fired with no position behind them and opened a naked 2-lot short.** Remove: +\$34.50 back.

And the gate's only green print should have been a stop. Trade 572 (+ \$28.50, the single winner the whole roster line rests on) had its protective stop trigger at 29390.25. The tape printed **at 29390.25 and never above it**, for about one millisecond. The venue's stop-limit does not trigger on a touch-without-through, and the desk's own breach guard would have fired it but reads one price per second. Its twin, filled 2.75pt lower, had its stop breached in the same burst and was stopped for −\$18.50. Applying the twin's measured 1.00pt stop slip: **−\$18.50. Swing −\$47.00.**

THE HEADLINE OF THIS CARD

Normalized, nipc_short is 0-for-14. Its one green print was a single-millisecond, exactly-at-the-trigger touch that neither the venue's trigger rule nor the desk's one-second breach guard happened to catch, on the lot that got the better of two fills 2.75pt apart.

ROOT CAUSE — NOT THE EXIT, NOT THE RUNG: THE BASE, AND THEN THE SIGNAL

The structural suspect was that 2.0R/2.5R targets on a gate stopping out 94% of the time are never consulted. **Measured, that is false.** Parking the target at 12R and reading how far each trade got before its stop: **31.9% of**

shorts reach 2.5R. The rung is reachable, and the geometric expectancy has a shallow maximum at 2.5–3.0R — the shipped pair sits on that plateau. **The exit is not the disease.**

The disease is arithmetic. Median R on the short book is 13.2pt, so 1R is about \$26.40 a lot:

PER LOT	
gross edge (0.117 R × \$26.40)	+\$3.09
\$1.50 fee + 2 × 1.25pt measured slip	−\$6.50
net per lot	−\$3.41

Even at the lab’s fantasy one-tick slippage the cost is \$2.50 against a \$3.09 gross edge. **The entire edge is the same size as the execution cost**, and a fixed-dollar cost is a regressive tax on a gate whose R is small. The slippage is structural, not a bug: NIPC’s trigger is “price trades through the half-back level”, so you need a stop-entry — a resting limit is geometrically impossible. Stop orders always slip, most in exactly the fast tape this setup selects for.

And then the finding that decides it: the signal is short-delta BETA, not alpha. Per-session P&L against the 13:00–15:00 UTC drift correlates **−0.528**, and the trade count correlates **−0.640** — it even trades more on down days. Down sessions +\$719, up sessions −\$211. 2026-W31’s news window fell on four of five sessions; W32’s rose. **That is the whole story of the “one good week”.**

NULL TEST	RESULT
“Sell one lot at 13:00, cover at 15:00” — same fees, same slip, no signal, no state machine	+\$354 over 17 round trips
the whole NIPC apparatus, same 17 sessions	+\$508 over 126 lot round-trips
random entry, 200 trials, same window, same counts, same exits	nipc sits at the 80th percentile — needs 95th

The entire NIPC apparatus buys \$154 over a one-line rule, at 7× the execution surface and 7× the incident exposure. And a variance source bigger than the effect itself: both nipc legs share one tracker and one position, so which setups the short leg gets depends on what the long leg happened to be holding. **\$1,832 of the answer is decided by slot-contention luck** — larger than any effect in the rehab.

THE ONE THING THAT IS REAL — THE MISSING L0 BLOCK

Your structural suspect was right. `nipc_short` and `nipc_long` are the only families whose override entry has no adaptive `lo` block, running a bare 2.0R/2.5R while every mature gate carries the clip.

CONFIG	N	NET	WIN%	WINNERS KEPT	STRIP-3	STRIP-5	W29 / W30 / W31 / W32
A. bare 2.0R/ 2.5R (shipped)	63	+\$508	41.3%	22/22	-\$90	-\$410	-150 / -148 / +1164 / -359
B. + <code>lo</code> clip block	60	+\$559	43.3%	22/22	+\$118	-\$76	-7 / -34 / +691 / -91
C. B + entry $\geq$ 13:30 UTC	47	+\$649	46.8%	19/22	+\$209	+\$15	-99 / +43 / +709 / -4

B keeps every one of the 22 winners and does not win by cutting anything — it wins by banking more of them (win rate 41.3%→43.3%), which is exactly what a fixed-dollar clip is for on a gate whose R is too small to pay an R-multiple ladder. It collapses the week-concentration (W31 from 229% of net to 124%) and it **generalises to the mirror leg**, which is the strongest evidence it is a mechanism and not a curve-fit: `nipc_long` goes +\$463→+\$668, strip-3 -\$5→+\$332, win rate 43.3%→50.9%. **The `lo` block helps the LONG leg more than the SHORT one.** Its parameters are plateau-stable and \$40 is the same \$40 every mature gate already uses — not a fitted number.

ROBUSTNESS — THE FINDING THAT DECIDES THE CASE

In every single configuration I tested — 9 Lot-B rungs, 7 A×B rows, 6 stop widths, 5 break-even variants, 3 trails, 5 hold caps, 7 slippage assumptions, 23 filters — 2026-W31 is positive and the other three weeks are almost always negative. Even at zero slippage. And out of sample, fitting on W29-W31 and applying unchanged to the live week, **every single train-optimum loses money**: 13 of 13 filters, 8 of 8 rungs, 5 of 5 stop widths. The only two configurations that come back roughly flat are the two that were not fitted on the training window — the `lo` block (copied verbatim from the mature gates) and the 13:30 cut (the US cash open).

VERDICT

SHADOW **Do not re-arm on live capital.** Not FIXED, because the best config I could build makes +\$649 over 18 sessions of which +\$709 is one week, does not clear a random-entry null, beats a one-line rule by \$154 at 7× the trades, and is 0-for-14 live. Not TRULY-RETIRE, because three things are alive and cost nothing to watch. **Re-arm criterion, stated concretely: nipc_short returns to live capital only when it prints ≥ 3 consecutive ISO weeks non-negative, at least one of which has a NET-UP 13:00-15:00 tape, on ≥ 40 shadow fires.** That is a direct test of the one thing that killed it — it has never yet made money in a week the market rose.

SHIP THE 10 BLOCK ANYWAY to both nipc legs. It is a config change with no arming implication; both gates stay off. $\triangle$ **One wrinkle to comment in the code:** NIPC carries `entry_atr = R`, not ATR, so `atr_split: 22` is being compared against R (median 13.2pt) and the clip binds on ~75% of fires. That is fine, but it is not what the field name says.

3 · THE TWO LIVE GRIND-EXIT SHADOW LEDGERS

Two observe-only ledgers have been recording since 2026-07-26 with no involvement from any rehab harness. I re-derived both from the raw JSON myself rather than quoting their summary fields, and **both of the “data defects” flagged against them this week turn out not to be defects** — which matters, because one of them was about to be used to discount a result.

two_ratchet_shadow.json — the runner-CLIP watch

SHADOW, unchanged

The watch asks: would a second ratchet have clipped a runner before the chandelier gave it back? **The answer after 26 live grind_long_B trades: it has never clipped one, and it costs money.**

N	LIVE NET	TWO- RATCHET NET	SHADOW – LIVE	RUNNERS	RUNNER- CLIPS	
whole window (since 07-26)	26	+\$29.50	−\$265.00	−\$294.50	3	0
the 3 runners (rMFE 4.20 / 4.75 / 5.21 R)	3	+\$319.50	+\$555.00	+\$235.50	3	0
the other 23	23	−\$290.00	−\$820.00	−\$530.00	0	0
this week (08-03...08-07)	3	−\$123.50	−\$143.00	−\$19.50	0	0

The shape is clear even at n=26. On the three runners the two-ratchet genuinely does better, by \$235.50. On the twenty-three non-runners it does worse by \$530.00. Net it hands back **\$294.50** — and it is negative in both ATR buckets (ATR≥22: live +\$262.50 vs shadow +\$139.50; ATR<22: live –\$233.00 vs shadow –\$404.50). That is a variance lever with a negative mean, which is the worst kind.

What the scope’s verdict rule says, and why it does not apply. The rule is “zero clips over a genuine trend week is evidence for reconsidering the two-ratchet”. **This was not a trend week for grind — it was 177 seconds of tape.** Three trades, maximum realised excursion **1.17 R**. The ledger did not fail to find a clip; it never got a candidate. And two of the three entries were at ATR 21.0, below the floor this rehab proposes.

But the ledger delivers something it was not built for: it is the independent, live proof of the grind_long ATR finding. Different tool, no harness, its own 26 rows:

TWO_RATCHET LEDGER, LIVE ROWS ONLY	N	LIVE NET	RUNNERS
entry ATR ≥ 22	14	+\$262.50	3 of 3
entry ATR < 22	12	–\$233.00	0

All three runners had entry ATR **24.8, 25.5 and 46.7**. Live data, different tool, same answer as the 744-leg harness.

Disposition: SHADOW, unchanged — and the watch is now conditional. It is only meaningful on trades the new ATR floor would admit; under the proposed config 14 of the 26 rows to date would have existed and

all 3 runners would have. **Revive the question at ≥ 10 runners recorded under an $ATR \geq 22$ entry.**

✓ **LEDGER-INTEGRITY CHECK — THE FLAGGED “\$435 DISAGREEMENT” IS NOT A DEFECT**

This week’s grind dossier flagged that the file’s xcheck field (–\$406.00) disagrees with live_net (+\$29.50) by \$435 with no documented reconciliation, and suggested a look before the ledger is used to settle anything. **I summed the raw rows: xcheck is not a net, it is a delta.** Sum of recon_pure_usd = –\$376.50; sum of live_usd = +\$29.50; the difference is **–\$406.00 to the cent.** The field reconciles exactly — it is measuring how far the pure reconstruction sits from the live book. **No defect. It just needs renaming to recon_minus_live.**

partial_shadow.json — the 2R-partial smoothness watch

SHADOW, prior RESTORED

The claim on the table: the partial is now +\$307.50 better on the mean over the window (–\$1,074.90 vs –\$1,382.40) with a shallower drawdown, which would quietly invert the standing “variance lever, not a mean-raiser” prior. **I re-derived it and the inversion is an artefact of the broken entry gate.**

First, the shape of the thing. Of the 40 trades in the ledger, **only 10 have a non-zero delta at all** — on 30 of 40 trades the partial does literally nothing, because nothing came within reach of 2R. Here is every trade where it acted:

DATE	ENTRY ATR	BASELINE	PARTIAL	DELTA
07-30	21.6	–\$90.00	+\$40.00	+\$130.00
07-30	15.5	–\$65.40	+\$28.10	+\$93.50
07-30	20.7	+\$22.00	+\$92.50	+\$70.50
07-30	39.9	+\$187.00	+\$252.00	+\$65.00

DATE	ENTRY ATR	BASELINE	PARTIAL	DELTA
07-27	24.8	+\$66.00	+\$130.50	+\$64.50
07-30	14.1	+\$28.00	+\$69.50	+\$41.50
07-30	25.5	+\$126.00	+\$163.50	+\$37.50
07-29	18.7	+\$79.00	+\$113.00	+\$34.00
07-29	27.2	+\$356.00	+\$285.50	−\$70.50
07-30	46.7	+\$688.00	+\$529.50	−\$158.50
total				+\$307.50

Now look at where the money comes from and where it goes. **Both negative deltas are the high-ATR runners** (ATR 27.2 and 46.7) and together they are −\$229.00 — the tail forfeit, exactly as the standing prior predicts. **Every positive delta at ATR below 22 is +\$369.50 across five trades.** Split the whole ledger by the floor the grind rehab is proposing:

PARTIAL LEDGER, SPLIT AT THE PROPOSED ATR FLOOR				
N	BASELINE	PARTIAL	DELTA	
entry ATR ≥ 22 — the trades that will REMAIN	19	−\$199.00	−\$261.00	−\$62.00
entry ATR < 22 — the trades being REMOVED	21	−\$1,183.40	−\$813.90	+\$369.50

THE PRIOR IS RESTORED, NOT INVERTED

The 2R-partial's apparent mean-raising is created entirely by the low-ATR population this rehab is removing. Fix the entry and the partial costs about \$62 of mean on the trades that will remain. It is a variance lever after all. This week the partial's delta was exactly **\$0.00 across 3 trades** (baseline −\$286.00 = partial −\$286.00) because nothing came within a mile of 2R. That is a non-observation, not evidence. **Disposition: SHADOW, operator's mean-vs-smoothness choice UNCHANGED.** Re-present after 20+ trades under the new floor.

✓ LEDGER-INTEGRITY CHECK — THE FLAGGED “SUSPICIOUS PAIR OF ZEROS” IS CORRECT, AND THE FINDING STANDS

This week’s grind dossier flagged that `green_day_pct` and `green_week_pct` both read 0.0 for baseline and partial across 40 trades, called it more likely a bug than a result, and said the smoothness half of the file should not be quoted until checked. **I checked it by summing the raw rows per date.** Baseline daily nets: $-\$278.00, -\$204.00, -\$336.40, -\$278.00, -\$286.00$. Partial daily nets: $-\$213.50, -\$240.50, -\$56.90, -\$278.00, -\$286.00$. **Five dates, ten numbers, all ten negative.** Both ISO weeks in the window are negative for both books too. **`green_day_pct = 0.0` is not a bug — it is the honest answer, and it is a brutal one: neither book has had a single green day since 07-26.** The smoothness half of the file can be quoted. The partial does buy real smoothness — worst day $-\$286.00$ against the baseline’s $-\$336.40$, drawdown $-\$1,074.90$ against $-\$1,382.40$ — it just has no green days to smooth.

4 · THE FAILURES — everything tried THIS WEEK that did not work

This is the part of the section that earns the rest of it. Every idea below was tried this week, on this week’s tape, with a named test. Nothing here is imported from a previous week. **REFUTED** means a specific test killed it and I do not expect a reformulation to save it. **PARKED** means it is not dead — it is waiting on a stated condition, and the condition is written next to it.

4A · REFUTED — THE TEST THAT KILLED IT

THE IDEA	ON	THE NAMED TEST THAT KILLED IT
An ER floor as an entry filter (any level, and the ceiling too)	grind_long	Winner-keep: $ER \geq 0.20$ drops 17 of 19 winners and keeps the two smallest. Then the sign flips: in-sample the ceiling wins by \$4,500, out-of-sample the floor wins by \$2,600. Not a parameter with an optimum — a coin.
“Bench after a wall of stops”	grind_long	Entry caps at 2 and 4 per day per lot cost \$3,576 and \$2,740 . The runner arrives after

THE IDEA	ON	THE NAMED TEST THAT KILLED IT
		the stop cluster — 08-05 13:36 is the pattern in miniature: four stops in 20 seconds, then a + \$135/+\$95.50 pair off the same minute.
A DWELL requirement (signal must hold 2-3 bars) as the anti-churn fix	grind_long	DWELL-2: strip-3 −\$1,273 , out-of-sample − \$548. DWELL-3 is a null. The churn fix turned out to be the ATR floor, which cuts trades 79%.
Turning fast_slope off (60-bar slope)	grind_long	In-sample neutral, out-of-sample −\$1,436 . The existing setting is vindicated.
“Give it the k10 stop width its best shadow uses”	abs_veto_short	Read the live slate build: it already has it , exactly. Same entry, same targets, same stop, same clip. The best shadow is the same exit at different moments. A null before a number was computed.
“The arming policy is selecting the losing minutes”	abs_veto_short	Armed windows: +\$48.20/trade . Benchded minutes: +\$19.70/trade. Every armed window that produced a signal was green. The premise is the opposite of true.
An extension cap (“don’t chase an already-spent move”)	abs_veto_short	6-rung sweep over 17 days. Every cap that bites loses money and drops winners . The only net-positive cap removes exactly one trade in seventeen days.
ER floors / dwell rules as an arming rule	abs_veto_short	12-rung + 3-rung sweeps. Best per-trade rung (\$34.70, the best number in the whole sweep) keeps 4 of 15 winners . Fifth gate to reproduce “ER filters fake, ATR floors real”.
An ATR ceiling (bench above 19/22/25pt)	abs_veto_short	5-rung sweep. Benching above 19pt keeps 0 of the 15 biggest winners . The edge lives in high-ATR tape.
The stacked router rules as an arming policy	abs_veto_short	Out-of-sample split: in-sample +\$903.50 → out-of-sample −\$22.50 . Fitted on the weeks that made it look sensible, negative on the week it actually ran.
The router “arm trio” (new extreme + ER rising + volume up)	abs_veto_short	Cuts +\$1,672.50 — 73% of the book — and 11 of 15 winners .
First-fire-of-burst filter	abs_veto_short	Loses \$172 and 3 winners. REFUTED as a filter, KEPT as a diagnosis : first fires are +\$44.40/tr at 81% win vs +\$9.60/tr for later ones — real information about the entry lag.
My own 5s-bar exit simulator	abs_veto_short	Refuted as an instrument. Its stop-width conclusion flips with the detection rule (bar-extreme says double the stop; bar-close says keep it) and the bar-extreme version contradicts the tick truth. It would have recommended doubling the stop on a number produced by my own measurement choice. Do not cite its numbers .
“Benching the fader cost us money”	exhaustion_short	Window-by-window replay of the always-armed counterfactual: benching SAVED \$140 . Always-armed for the full 120 hours loses −\$614 (mean of ten polling phases), 1 of 10 phases green.
A regime/extension bench predicate for faders (4 forms)	exhaustion_short	The best one takes −\$847 to −\$3 by deleting all eight target winners . The most seductive fake win in the whole report.
Every ATR floor / ATR ceiling / ER floor / ER ceiling / break cap — 42 cells	exhaustion_short	Zero of forty-two clears strip-best-3 with winners intact. Every cell near zero deletes the winners; every cell that keeps them keeps the loss.
The resting-wall mechanism — the gate’s own core claim	exhaustion_short	Removing the wall test improves 60s and 120s forward drift and 2.5xs the sample. Measuring it

THE IDEA	ON	THE NAMED TEST THAT KILLED IT
		at the true top of book makes it worse. Demanding wall ≥ 3.0 drifts -13.4pt at 120s.
The fade direction itself	exhaustion_short	200-trial placebo null: the entry sits at the 0th percentile at 60s and 120s — significantly anti-predictive. Following the flow instead is \$900 better.
An R-size floor derived from the cost arithmetic ($R > 27.8\text{pt}$)	nipc_short	The mechanism predicts a cure, so I tested it: n=1 above the threshold and the sign flips twice below it. There is no R region to trade.
The span 35–50pt band (one of only two filters green in all four weeks)	nipc_short	Knife-edge: move either edge by 5 points and it changes sign (35–45 strip-3 $-\$111$, 35–50 $+\$128$, 35–55 $-\$32$). Classic overfit shape, and it keeps 9 of 22 winners.
$a_{115} \leq -1$ (the impulse fights the prior drift — a real news shock)	nipc_short	The most interesting lead in the file: positive in all four ISO weeks including the live one. Dies on two tests — strip-3 $-\\$50$ (all \$354 is three trades) and it keeps 7 of 22 winners . Rejected as a rehab, carried as the shadow’s watch-item.
Break-even stops (at 0.5R / 0.75R / 1.0R / 1.5R)	nipc_short	All four far below the +\$508 baseline , strip-3 in the $-\$333$ to $-\$640$ range. On a gate that reaches 1R half the time, a break-even stop just converts winners into scratches.
Trailing Lot B instead of the fixed rung	nipc_short	0.75R / 1.0R / 1.5R trails give $+\$1$ / $+\$306$ / $+\$429$ against +\$508 for the fixed 2.5R . The 08-01 finding confirmed independently.
Widening the stop to $1.5 \times R$	nipc_short	Looks better in aggregate and makes the live week worse (W32 $-\$537$). Null.
The cheat-sheet’s momentum prior of Lot A 1.5R	nipc_short	Worse than your guess: $+\$193$ against $+\$508$ at $B=2.5R$. The operator’s $A=2.0R$ is the best A at every B. The guess is confirmed, the prior is not.
Shortening the MAX_HOLD cap (5-60 min)	MAX_HOLD	Winner decomposition + runner audit. A 5-minute cap “makes” \$594 by leaving 4 of 20 top winners and taking \$4,382 off the winner pool. At 15 minutes it clips 24 of 51 runners .
Staggering the paired-lot flatten (“the second lot eats the first lot’s impact”)	MAX_HOLD	Full-book same-instant fill census: 18 of 18 groups, all-time cost $-\\$55.00$, of which 56% is one unrelated trade. On the target event the two lots filled one tick apart, for fifty cents .
Restoring the pre-08-01 managed-profit shape — the rehab’s own premise	chandelier	Slot-honest replay: $-\$2,932 \rightarrow -\$3,433$, losing in 0 of 20 leave-one-day-out folds . And repricing this week’s 34 winners under it makes them \$322 worse , not better.
Restoring give-back at 50/40, or at anything	give-back	Worst arm tested ($-\$4,425$, 0/20 folds), then a 42-config sweep (arm $\$30-\$120 \times$ retrace $\$20-\100): 42 of 42 negative in every regime . No configuration exists.
Continuous exit_chandelier at $k = 1.0$ / 1.5 / 2.0 / 2.5 / 3.5	chandelier	All five in the bottom half of the field. The trail is directionally right in one cell but quantitatively second-best; the tightening version recovers $+\$761$ where a plain wide rung recovers $+\$3,049$.
“SHORTs only” — and it PASSED leave-one-day-out	chandelier	Rejected anyway. Correlation between the session’s own net drift and short-minus-long P&L is -0.462 across 20 sessions. The tape fell 182pt over the window; the “edge” is a bet on tape direction. Flagging it rather than banking it .
“No profit exit at all” — which looked like a \$13,540 winner	chandelier	My own bug. Unconstrained scoring lets a wide trail book the same move thirteen times when a

THE IDEA	ON	THE NAMED TEST THAT KILLED IT
		gate fires thirteen times into one leg. Replayed with one position per sub-slot it is −\$1,058 . The artefact was worth \$14,600 and every headline in that dossier is now slot-occupancy-honest.
The cheat-sheet's fader prior of 0.5R / 1.5R	chandelier	In the one regime where faders make money it ranks 131st of 135 configurations and throws away \$3,929 on 25 signals. It is nearly right (rank 9/135) in violent-whipsaw. The prior is not wrong, it is unconditioned.

4B · PARKED — AND EXACTLY WHAT WOULD REVIVE IT

THE IDEA	ON	WHERE IT STANDS	WHAT WOULD REVIVE IT
Stop 1.0 → 1.25 × ATR	grind_long	The only widening that improves in-sample strip-3 and out-of-sample (+ \$2,018 / +\$829 vs +\$1,411 / + \$643). But the surface is noisy — 1.5× dips between 1.25 and 2.0.	SHADOW Holds on ≥ 40 more legs. One parameter change at a time, and the ATR floor is worth ten times as much.
Lot A 2.5R	grind_long	2.0R / 2.5R / 3.0R are indistinguishable at n=44. 2.5 sits in a small local dip. 1.0R and 1.5R are clearly too tight.	A proper answer when the home segment has ≥ 100 legs . Leave it alone rather than chase 3.0R on 42 trades.
US-session-only clock	grind_long	Looked good for a long time. Once the ATR floor is in it earns nothing — the overnight segment under the new floor is +\$986 on 28 legs . It was only ever a proxy for “when is ATR high”.	Only if the overnight segment turns negative over ≥ 30 legs .
Range-position (box 0.5-0.95) filter	grind_long	Whole-tape positive (+\$2,489, strip-3 +\$1,031, OOS +\$1,081) but neutral on the home segment .	As a second-order filter after the ATR floor has 4 weeks of live data .
The clean-trend hole — 0-for-14 in its own namesake regime	grind_long	All 14 legs at ER30 ≥ 0.45 stopped, across three episodes. The one stone left unturned , and I will not explain it away at n=14.	A genuine multi-hour trend week. Nothing smaller can settle it.
The router's arming criteria for grind	grind_long	Anti-correlated this week: − \$211 armed, +\$2,388 benched. The router arms on ER-climb; the gate earns at ATR ≥ 22 with ER < 0.25 — close to opposite conditions.	Needs its own week. Revive as the “ATR band, ignore ER” rule once the config change has a fortnight of live data. Do not change two things at once.
grind_short	grind	Zero live rows in the window. Never measured. The two-sided-gate rule says its thresholds must be derived independently.	A dedicated per-side sweep. Until then, “grind” as a family is half-measured.
Lot A 1.5R → 2.5R in trend-no-break and violent-whipsaw	abs_veto_short	The only two regimes with a real plateau — and the instrument that found it already failed one validation .	SHADOW Run as a shadow arm on tick-repriced fills; promote at n ≥ 30 per regime .
The 61-second entry lag	abs_veto_short	Real and universal (median 61s over 78 matched entries), worth ~\$6.87 a lot on average with a ±\$21 band. This week's four signals drew 2.51 standard errors worse than the population — a bad draw, not the explanation.	Instrumentation: log the SIGNAL bar alongside submitted/filled in the signals table, then compare live against replay per fire.
ATR floor 13-17pt	abs_veto_short		

THE IDEA	ON	WHERE IT STANDS	WHAT WOULD REVIVE IT
		Genuinely real and plateau-shaped — keeps 15/15 winners and lifts \$/trade from \$14.40 to \$26.80. And it still costs \$259-\$760 of net . It buys the same money in fewer trades.	If slot contention becomes binding. It is a slot-contention lever, not a P&L lever — do not deploy it to fix a bleed it does not fix.
The fast=True 2-bar trigger	abs_veto_short	The one stone I did not turn. It changes which bars fire, so it needs a full signal replay from the tape rather than a filter over an existing ledger, and no shadow variant runs it.	A dedicated signal replay. This is the outstanding piece of work on this gate.
exhaustion_short as a live gate	exhaustion_short	Gross edge +\$1.36 a lot , green in 8 of 10 phases, survives strip-3. Friction is \$3.28 a lot . Sub-friction by about two dollars.	(a) a zero-latency execution path — fill inside the 20-second window, not 5 seconds after it — AND (b) round-trip cost below ~\$0.50. Nothing about arming, ER, ATR or regime closes a gap that size.
Dead-chop, Lot A 2.5R / Lot B 2.5R	exhaustion_short	The one positive cell out of 280. A genuine 4x4 plateau — all 16 neighbouring cells positive — sitting on a number that two of six days produce, and it dies on strip-3 (–\$95) and LODO (–\$35). A lucky cell in a 280-comparison search.	n > 60 in that bucket alone , i.e. about four more weeks of dead-chop tape. The story at least coheres: the fade only works when nothing is happening.
Follow-the-flow (flip the side)	exhaustion_short	A \$900 swing from a sign flip (–\$847 → +\$51). But +\$51 ± \$354 at \$0.32 a lot is not an edge, it is the absence of a bleed.	Only if someone wants a cheap two-sided shadow sim. Not worth a slot. Recorded as the cleanest confirmation that heavy buying that has not yet moved price is a continuation tell.
MFE-conditioned early release (cut at T if it never went favourable)	MAX_HOLD	Runner-safe by construction and it sweeps beautifully on the shadow book (+\$9,134, 30/30 winners kept). On the live book it is –\$139 and has no plateau. The decisive test: 84% of its shadow gain comes from rows where the sim’s own stop had already failed.	Re-test once the shadow repricer’s 44% stop leak is fixed, on the stop-honest slice only. Promote if it clears +\$10/fire over 60+ live fires .
Session-boundary cap instead of a fixed 2h clock	MAX_HOLD	Winner-safe and mildly positive — but it reaches the same three malfunction trades , and at 04:00 it arrives later than 120 minutes does for two of them.	It is the correct primitive for any slot that deliberately holds for hours (day_rider holds ~7h20m and already has its own flat-by field) — but not as a replacement for the general backstop.
A reopen cool-down after the CME break	MAX_HOLD	Every MAX_HOLD fire opened overnight and two of them fired two minutes after the 21:00-22:00 UTC reopen , on a 14-bar ATR window holding two minutes of bars. Untested.	Belongs to the entry rehab, not the exit rehab — but it is the single most promising untested idea this reconstruction turned up.
The trail machinery (exit_chandelier_lock)	chandelier	Worth +\$246 across its whole parameter block — noise — and it loses in 3 of the 4 regime cells. It is the runner-up in the home cell and it is uncapped, which a fixed rung is not.	SHADOW Keep it in the code. Do not re-arm it on the strength of this report.
The er_split code change (gate the clip on ER as well as ATR)	chandelier	Beats the config-only version on every axis (+\$25 archive, +\$156 week, +\$345 LODO floor, better in 4 of 4 ISO weeks vs 2 of 4) — but the	Ship the config-only edit today; queue er_split as the follow-up

THE IDEA	ON	WHERE IT STANDS	WHAT WOULD REVIVE IT
		margins are small and it needs a deploy.	and re-measure before promoting it.
The not violent-whipsaw entry filter	chandelier	Keeps 78% of winning lots , saves +\$2,690, drops a population running −\$12.45 a lot, positive in 20/20 folds. But it is an entry change and outside this rehab's mandate.	Recorded for the gate rehabs, not recommended here. Do not let an exit study smuggle in an entry rule.
The adaptive_exit regime-3 selector on grind	grind_long	It never ran (the override file wins). Its would-have behaviour on grind's home segment is +\$45 net, −\$367 strip-3 — the “tight” mode is a disaster on a high-vol expansion.	A warning, not a finding. Do not let it pick “tight” here if it is ever wired in.
The Paris-midnight auto re-arm on zero tape	desk-wide	On 08-04 it re-armed every off gate at 22:00 UTC on “insufficient bars, range 0pt, day P&L flat”, the corrupt thrust fired two minutes later , and nothing benched it for 45 minutes.	Parked → action: add the gate to the re-arm hold list, or gate the whole re-arm on there being enough bars to form a regime read. Arming anything on “range 0pt” is not a decision.

5 · THINGS FOUND ALONG THE WAY THAT ARE NOT ABOUT ANY ONE GATE

Six rehabs run independently turned up the same class of problem more than once. These are the ones that matter beyond their own dossier — and two of them were found by two separate rehabs that did not know about each other, which is the strongest kind of confirmation available here.

WHAT	EVIDENCE	THE FIX
Stop-limit orders filling AT their limit, 20pt through their own trigger (found independently by two rehabs)	2 lots this week, both the second lot of a pair, both within four minutes of the cash open. The other 54 stop fills have a median slippage of 0.00pt . One filled 10.25pt above the highest print in the window. ≈ \$70/week of fake losses.	Narrow the limit buffer to ~5 points or use a plain stop-market. And add a fill-plausibility check: if a stop fills more than N points through its trigger, alarm and reprice it against the tape before it is booked. We caught this by hand; nothing in the desk would have.
Nothing asks “does every live STOP have a slot?” (found independently by two rehabs)	Two orphaned GTC stops opened a naked 2-lot short on 08-06 and it was booked to a gate that had been benched 38 minutes earlier. Separately, the two poisoned stops from 08-04 rested ~2,000pt away for three days .	Run the orphan sweep on a timer, not only at boot . Until it exists, any gate's ledger can be contaminated by another gate's garbage — the nipc rehab would have started from a 6% hit rate on 16 trades instead of 0% on 14 if the order-id namespace had not been checked.
A poisoned number survived a restart	The 08-04 gold-bar leak was fixed in five minutes. Its consequence survived, because <code>reconstruct()</code> re-adopted two open slots with the 1,848pt ATR persisted — leaving the stop naked, both targets unreachable and the quiet clip disabled, for 115 minutes. Cost: \$192.50 in one night .	On <code>reconstruct()</code> , re-derive the entry ATR from the tape and refuse to adopt a slot whose persisted stop is more than a few ATR from entry.
The desk cannot tell you why an approved entry never became an order	On 08-07 at 14:19:36 both lots logged “confirmed → OPEN” and then nothing at all — no order, no signals row, no error. Three silent return paths, none of which logs.	One log line. (For the record it saved us \$174 that day, and I have not credited that back — you cannot bank a loss you avoided by accident.)

WHAT	EVIDENCE	THE FIX
The shadow book leaks past its own stops on 44% of trades	Median overshoot $0.56 \times \text{ATR}$, maximum $32.32 \times \text{ATR}$, on a book that is 99% $1.0 \times \text{ATR}$ stops. Every shadow study that rewards cutting early is inflated by this. It is why the MFE-release rule looks brilliant in shadow and negative on live.	Fix the repricer's stop detection, then re-run anything that ever concluded "exit earlier".
MD contamination in the shadow book is a 14-day condition, not one bad night	The tape-consistency screen struck 166 of 7,639 shadow rows across 14 separate days . Nothing flags them.	shadow.db needs the data_quality column gazbot7.db already has.
Neither tick source covers the window on its own	data/tape/ticks/MNQ/ is missing 2026-08-07 while capture.db has it; capture.db's ticks only reach back to 08-03. A harness reading only the archive silently loses the most recent session and books zero trades with no error — which is exactly what the shipped nipc replay is doing right now.	Union the two sources , and make an empty tick-day raise rather than book zero.
signal_journal.suppressed_by is NULL in all 596 rows	Any statement anywhere about "which gate fires were suppressed and why" is currently unsupported by the store. Everything in these dossiers about what a gate would have done unrouted is reconstructed from the tape, not read from a log.	Populate it. It is the difference between a reconstruction and a measurement.
A replay that was not deterministic	Same-millisecond ticks came back in arbitrary order from the parallel scan, and the one-second downsample keeps the last print per bucket — so tie order changed fills. The identical config measured +\$559 and +\$639 on two runs.	Pinned. But note the size of it: an \$80 swing on 60 trades from tick tie-ordering alone is a useful yardstick for how much of everything is noise.

6 • THE HONEST LIMITS ON ALL OF IT

READ THIS BEFORE YOU SIZE ANYTHING ABOVE

- n is small everywhere.** The grind_long fix rests on 155 legs / 78 signals / 21 runner events across 11 sessions. The abs_veto_short veto rests on **n=8**. The chandelier width rests on **5-7 trades out of 60**, with two sessions carrying two thirds of it. MAX_HOLD has **four fires in 21 days** and every per-regime "optimum" in its sweep rests on one to four bound trades.
- One instrument, one summer, one direction of drift.** Everything is MNQ, three-and-a-half weeks, on a tape that fell 182 points net over the US window. The chandelier dossier caught one result that was purely a bet on that drift and rejected it; there may be others I did not catch.
- Some of it is in-sample by construction.** Where an out-of-sample leg exists I have shown it, and in two cases (abs_veto's router policy, nipc's every train-optimum) the out-of-sample result is the thing that kills the

idea. Where it does not exist — MAX_HOLD, because no out-of-sample events occurred — I have said so rather than manufacture one.

4. **Bar-replay is not tick-honest.** The grind_long out-of-sample leg runs about **17% optimistic in total and up to ±\$370 on a single day.**

Nine of the chandelier dossier's 20 sessions are bar-replay. The exhaustion_short out-of-sample leg uses a different book source that disagrees in sign with the live one on 2 of 3 overlap days, which is why its one green number is not counted.

5. **Microstructure sampling noise is real and it is large.** On exhaustion_short the standard deviation across ten polling phases is **\$418 on a −\$880 mean.** Anyone quoting a single run of that harness is quoting noise.

6. **Three of the fixes are reverts, and reverts deserve suspicion too.**

The grind_long floor, the grind_long ceiling and the exhaustion_short exit all put back something that was changed on 08-01 or 07-29. The evidence says those changes were wrong. It does not say the people who made them were being careless — it says the tape they were fitted on was too short, which is exactly the risk these rehabs also carry.

The one sentence I would keep. Not one of these six was broken in the way it looked broken from the roster. grind_long looked like a bad gate and was a bad ATR floor. abs_veto_short looked like a bleeder and was the strongest signal on the desk on a 3.3% leash. MAX_HOLD looked like the worst exit on the desk and was the fire alarm being blamed for the fire. exhaustion_short looked like an arming problem and is a friction problem. nipc_short looked like a wrong-rung problem and is a direction bet. And the chandelier looked like a threshold that stopped being reached, and was switched off in code. **That is the whole argument for never benching anything at face value.**

P2 The shadow desk & promotion — who earned the right to trade money next week

Short version, Garrath: **the shadow book made money this week and the live desk lost it, and they disagree about which SIDE was the good one.** The shadow board's 53 sims ran on the same tape as the paper desk; the best of them, **abs_veto_55s**, is now at **+\$3,447 over 319 tick-repriced trades** all-time and **+\$779 over 81 trades this week.** The live desk over the same five days went **−\$932.50.** That gap is the section. And the single most useful thing I found is not a new gate at all — it is that **on the SHORT side the live desk earned −\$74.00 a trade while the identical unfiltered signal in shadow earned +\$18.81 a trade, and on the LONG side it is the exact opposite.** Whatever is arming and benching those two gates is worth more money than any exit tweak in this report.

WHAT SURVIVED — READ THESE FIVE LINES AND YOU HAVE THE SECTION

- 1. abs_veto_55s survived the whole battery.** 12 of 17 days green, positive in chop AND trend, positive in both walk-forward halves, and **+\$3,342** ahead of the un-vetoed thrust it is built from. The 55-second confirm is real.
- 2. But it is already LIVE** — **abs_veto_long** and **abs_veto_short** have been paper gates since 2026-07-29. There is **no new gate to promote this week:** every sim in the board's top ten is a variant of something already trading. Saying so is the honest finding, not a failure.
- 3. Two-sided still holds all-time, but only the SHORT side survives a fresh kill-test.** Short: **+\$2,290** all-time, **+\$1,780** after stripping its three best trades, both halves green, 14/17 green days. Long: **+\$1,157** all-time but **−\$30 this week and negative on every robustness cut of this week's data.**
- 4. The two-ratchet stays dead — but it did NOT clip a runner.** 26 grind trades, 3 genuine $\geq 4R$ runners, **zero clips.** That is evidence FOR reconsidering it, exactly as the rule says. It still lost on the boring trades.
- 5. The exit ladder guess is wrong in every single cell I could measure it in** — 24 gate×rung cells, and the operator's 2026-07-31 guess never

ranked better than 3rd out of 53 policies. The direction of the guess (ride trends, bank chop) is right; the numbers on it are not.

1 — THE BOARD, ranked by honest money

This is every sim on data/shadow.db that has closed a trade, scored on tick-repriced `real_pnl` — never the optimistic sim P&L. Two blocks of columns: all-time (since 2026-07-16) and this week only (from 2026-08-02 22:00 UTC). `#` is the stable sim ID from data/sim_registry.json. Green rows are the ones that made money over their whole life; red rows are the ones that lost more than \$500.

#	SIM	FAMILY	N	NET \$ ALL-TIME	WIN%	\$/TRADE	DAYS	N WK	NET \$ THIS WK	WIN% WK	\$/TR WK
6	abs_veto_55s	thrust + 55s veto	319	+3447	46	+10.81	20	81	+779	44	+9.62
5	abs_veto_50s	thrust + 50s veto	347	+2609	43	+7.52	20	88	+641	42	+7.28
49	thrust_cont	thrust	365	+1288	38	+3.53	20	89	+577	43	+6.48
52	thrust_short_absve55	abs_veto short-only	90	+965	49	+10.72	10	37	+647	54	+17.49
42	sw_absS_B_k10	abs_veto short stop-width	29	+847	55	+29.19	3	29	+847	55	+29.19
43	sw_absS_B_k20	abs_veto short stop-width	28	+666	57	+23.77	3	28	+666	57	+23.77
40	sw_absS_A_k10	abs_veto short stop-width	31	+508	61	+16.37	3	31	+508	61	+16.37
19	cx_absSB_clip	abs_veto short clip A/B	25	+418	52	+16.70	3	25	+418	52	+16.70
20	cx_absSB_live	abs_veto short clip A/B	24	+406	46	+16.92	3	24	+406	46	+16.92
48	thrust_aliigned	thrust	326	+395	36	+1.21	20	85	−2	34	−0.02
41	sw_absS_A_k20	abs_veto short stop-width	31	+384	65	+12.37	3	31	+384	65	+12.37
7	abs_veto_60s	thrust + 60s veto	111	+299	41	+2.69	20	31	−126	32	−4.06
53	thrust_short_raw	thrust short raw	159	+283	36	+1.78	12	65	+420	42	+6.45

32	rg_short_025_flow_50	rgv short	41	+159	46	+3.87	13	18	−86	39	−4.78
51	thrust_loose	thrust un-vetoed	564	+105	35	+0.19	20	140	−26	35	−0.19
33	rg_short_025_flow_50_rth	rgv short	7	+100	43	+14.21	6	0	—	—	—
34	rg_short_025_raw	rgv short	49	+80	43	+1.63	15	19	−104	37	−5.47
45	sw_grind_A_k20	grind stop-width	85	+58	52	+0.68	4	85	+58	52	+0.68
46	sw_grind_B_k10	grind stop-width	66	+37	32	+0.55	4	66	+37	32	+0.55
13	cb_thrust_dropped	thrust cb	67	+17	31	+0.25	7	0	—	—	—
30	rg_short_025_flow_25	rgv short	45	+17	42	+0.37	14	18	−87	39	−4.83
3	CL-grind_long	claude-loop verdict	1	+3	100	+3.00	1	1	+3	100	+3.00
1	CL-abs_veto_short	claude-loop verdict	3	−7	0	−2.17	3	3	−7	0	−2.17
4	CL-rgv_short	claude-loop verdict	5	−10	20	−1.90	3	5	−10	20	−1.90
2	CL-capitulation_long	claude-loop verdict	5	−10	0	−2.00	3	5	−10	0	−2.00
35	rg_short_050	rgv short	21	−16	43	−0.74	9	0	—	—	—
31	rg_short_025_flow_25_t15	rgv short	27	−26	44	−0.96	9	0	—	—	—
36	sw_absl_A_k10	abs_veto LONG stop-width	28	−127	39	−4.52	4	28	−127	39	−4.52
39	sw_absl_B_k20	abs_veto LONG stop-width	25	−142	44	−5.68	4	25	−142	44	−5.68
37	sw_absl_A_k20	abs_veto LONG stop-width	27	−162	56	−6.00	4	27	−162	56	−6.00
38	sw_absl_B_k10	abs_veto LONG stop-width	26	−174	27	−6.69	4	26	−174	27	−6.69
44	sw_grind_A_k10	grind stop-width	95	−220	33	−2.32	4	95	−220	33	−2.32
18	cx_absLA_live	abs_veto LONG clip A/B	21	−222	38	−10.55	4	21	−222	38	−10.55
10	capit_mid	capitulation	20	−237	30	−11.85	10	7	−148	14	−21.07
17	cx_absLA_clip	abs_veto LONG clip A/B	21	−247	29	−11.74	4	21	−247	29	−11.74

11	capit_ride	capitulation	19	−274	16	−14.42	10	7	−148	14	−21.14
25	rg_long_20_flow25	rgv long	49	−288	39	−5.87	12	0	—	—	—
47	sw_grind_B_k20	grind stop-width	41	−297	44	−7.23	4	41	−297	44	−7.23
24	grind_fast	grind	1068	−350	33	−0.33	20	336	−2401	26	−7.15
21	cx_grindA_clip	grind clip A/B	82	−403	32	−4.91	4	82	−403	32	−4.91
26	rg_long_20_t15	rgv long	53	−428	36	−8.07	12	0	—	—	—
50	thrust_fast	thrust	181	−431	33	−2.38	14	0	—	—	—
16	chand_k35	chandelier exit	434	−472	33	−1.09	19	115	+100	33	+0.87
27	rg_long_20_t20	rgv long	50	−546	30	−10.91	12	0	—	—	—
22	cx_grindA_live	grind clip A/B	74	−546	23	−7.37	4	74	−546	23	−7.37
12	cb_thrust	thrust cb	258	−878	32	−3.40	12	0	—	—	—
23	exhaustion_rev	exhaustion fade	467	−1400	39	−3.00	16	146	−213	42	−1.46
29	rg_long_fast_v	rgv long	237	−1483	30	−6.26	18	38	−45	29	−1.18
15	chand_k25	chandelier exit	352	−1544	41	−4.39	13	0	—	—	—
14	chand_k20	chandelier exit	362	−1588	43	−4.39	13	0	—	—	—
8	capit_live_mirror	capitulation	76	−1610	12	−21.18	6	76	−1610	12	−21.18
28	rg_long_fast	rgv long	299	−2281	29	−7.63	12	0	—	—	—
9	capit_loose	capitulation	403	−8132	16	−20.18	20	164	−3090	15	−18.84

WHAT THE BOARD ACTUALLY SAYS — THREE THINGS, IN PLAIN WORDS

(a) Every profitable sim on the board is a THRUST. Look at the green block: abs_veto_55s, abs_veto_50s, thrust_cont, thrust_short_absveto55, and the five sw_absS_* / cx_absSB_* short-side arms. Nine of the top eleven. The faders — capitulation, rgv-long, exhaustion — occupy the bottom of the table, and capit_loose alone is **−\$8,132 over 403 trades at −\$20.18 each**. This desk's shadow book has exactly one working idea: a burst that keeps bursting.

(b) The whole top of the board is SHORT. All four sw_absS_* arms are green (+\$384 to +\$847). All four sw_absL_* arms — the identical machinery

pointed the other way — are red (−\$127 to −\$174). Same veto, same stop, same clip ladder, same three days. The only difference is direction.

(c) The 55-second confirm is worth the whole edge. thrust_loose is the same signal with no veto: **564 trades for +\$105**, twenty cents a trade, effectively zero. Put the 55-second confirm on it and you get **319 trades for +\$3,447** at \$10.81 a trade. The veto throws away 43% of the raw fires (245 of 564) and keeps essentially all the money.

2 — SO WHO GETS PROMOTED? Nobody — and that IS the finding

The board's top ten, ranked honestly, contains: three abs_veto confirm-delay variants, one thrust_cont, one short-only abs_veto, four stop-width A/B arms of abs_veto-short, and two clip-ladder arms of abs_veto-short. **Every single one is a variant of a gate that is already trading live.** abs_veto_long and abs_veto_short went into the paper tournament on 2026-07-29 and have been running ever since — 32 and 11 lot-fills this week respectively.

So the honest answer to "what do we promote on Monday" is: **nothing new.** There is no candidate gate sitting in incubation with a fresh mechanism and enough n. What the board IS offering is three config changes to gates already live, and those are worth real money. I have ranked them by how much evidence stands behind them.

RANK	WHAT THE BOARD IS NOMINATING	EVIDENCE	N	VERDICT
1	The 1-ATR stop is right on the SHORT side — keep it, do not widen. sw_absS_B_k10 (1.0×ATR stop) +\$847 vs sw_absS_B_k20 (2.0×) + \$666; Lot A: k10 +\$508 vs k20 +\$384. Targets held in POINTS so only the stop moved.	both lots agree		29+31 HOLD — already live at 1.0
2	abs_veto SHORT is the robust side; abs_veto LONG is decaying. Short survives strip-best-3, leave-one-day-out and both halves on all-time AND this-week data. Long survives all-time	see §4		159 / 160 TUNE THE SIDE

only and fails every cut on this week's data.				
3	abs_veto_50s is a real second arm, not a duplicate. +\$2,609 over 347 trades — 28 more fires than the 55s and \$3.29/trade less. The confirm window is a genuine tuning axis, and 55s is the better of the two.			
	55s beats 50s beats 60s			
	347	SHADOW — keep both running		

The caveat that governs all three: the five sw_* and cx_* arms have only THREE DAYS of history (they were seeded 2026-08-04/05). A \$29-a-trade number off 29 trades on three days is a lead, not a fact. They stay in the shadow book and get re-read next Friday. Shadow nominates; paper judges.

3 — THE PROMOTION BATTERY: scripts/abs_veto_robustness.py, all five tests

This is the full battery output, unedited, on tick-repriced real_pnl. I ran it fresh; nothing here is inherited.

Test 1 — the headline

SIM	N	NET \$	WIN%	AVG WIN	AVG LOSS	WORST TRADE	\$/TRADE	MAX DRAWD OWN
abs_veto_55s	319	+3447	46	+\$67	-\$38	-\$92	+\$10.81	-\$441

A 46% win rate that makes money because the winners are 1.76× the losers. The worst single trade in 319 is −\$92 — the 1-ATR stop is doing its job and there is no fat tail hiding in here. Max drawdown of −\$441 against +\$3,447 of profit is a 7.8:1 ratio; that is a curve you could actually sit through.

Test 2 — per-day spread: is it one lucky day?

DAY	N	NET \$	WIN%	DAY	N	NET \$	WIN%
07-16	1	−50	0	07-30	26	+300	50
07-17	20	+549	60	07-31	24	+74	42
07-20	22	+36	36	08-03	20	−74	35
07-21	10	+83	50	08-04	9	+540	100
07-22	12	−148	25	08-05	12	+156	42
07-23	23	+761	65	08-06	26	−120	31
07-24	18	+110	44	08-07	14	+278	50
07-27	18	−66	28	total	319	+3447	46
07-28	28	+582	57	12 of 17 days green. Best day +\$761, worst −\$148.			
07-29	36	+438	47	The top day is 22% of the total — concentrated, but not a one-day mirage.			

This is a genuinely spread edge. The three biggest days (07-23, 07-28, 08-04) together are \$1,883 — 55% of the total — but take all three out and you are still left with **+\$1,564 over 14 days**. It does not rest on one session, and it does not rest on one week: it made money in ISO week 30, week 31 and week 32.

Test 3 — regime split: chop vs trend

REGIME (30-MIN ER AT ENTRY)	N	NET \$	WIN%	AVG WIN	AVG LOSS	\$/TRADE	MAX DD
chop (ER < 0.18)	150	+1438	45	+\$64	−\$35	+\$9.60	−\$343
trend (ER ≥ 0.18)	169	+2008	47	+\$70	−\$40	+\$11.90	−\$359

This is the most important robustness result in the section. Almost every edge this desk has found turns out to be one regime wearing a disguise — and this one is not. It earns \$9.60 a trade in chop and \$11.90 a trade in trend, with near-identical win rates and near-identical drawdowns, on near-identical sample sizes. A 24% difference in \$/trade between the two

halves of the regime spectrum is noise. **abs_veto_55s is not a regime bet.** That is rare here and it is why it deserves the trust.

Test 4 — walk-forward halves

HALF	N	NET \$	WIN%	\$/TRADE	MAX DD
first half (by trade order)	159	+2236	49	+\$14.10	-\$441
second half	160	+1212	44	+\$7.60	-\$434

Both halves green — but be honest about the shape: **the second half earns 46% less per trade than the first** (\$7.60 vs \$14.10) and wins 5 points less often. That is a decay signature, not a failure. It could be regime (the back half is the flatter summer range), it could be the edge crowding, or it could just be 160 trades of noise. The right read is: still positive out-of-sample, but trending the wrong way — watch it.

Test 5 — head-to-head vs the un-vetoed thrust

SIM	N	NET \$	WIN%	\$/TRADE	MAX DD
abs_veto_55s (with the veto)	319	+3447	46	+\$10.81	-\$441
thrust_loose (no veto)	564	+105	35	+\$0.19	-\$948
the veto's contribution	-245	+\$3,342	+11pts	+\$10.62	halved

The veto drops 245 of 564 raw thrusts — 43% of them — and by doing so it lifts the win rate 11 points, multiplies the per-trade edge by 57×, and **halves the maximum drawdown**. This is the cleanest filter result on the desk. It passes the 15/17-winners test by construction: it is not winning by dropping winners, it is winning by dropping fakeouts, and the surviving book has a HIGHER win rate, which a winner-dropping filter cannot produce.

THE VERDICT, FOLLOWING THE NUMBERS

abs_veto_55s passes the battery on all five tests. It is spread across days, it works in both regimes, it holds in both walk-forward halves, and it beats its own un-vetoed parent by \$3,342. If it were not already live, it would be this week's promotion, two-sided, without argument.

It IS already live. So the deliverable is not a promotion, it is a re-confirmation and a per-side tuning question — which is \$4. And the honest caveats stand: 319 trades over 17 sessions of ONE summer regime, 1-lot sims with no queue position, no partial fills and no slippage beyond the tick reprice. The forward live run is the only real out-of-sample test, and this week it disagreed with the shadow book (\$4).

4 — TWO-SIDED OR NOT? The per-side re-derivation, and an inversion I cannot fully explain

The standing principle is right and I am not touching it: **a gate runs two-sided with independently-tuned per-side thresholds; you tune the SIDE, you never relegate a DIRECTION.** What follows is the per-side arithmetic done fresh on this week's numbers, and it lands somewhere more uncomfortable than "both sides earn".

abs_veto_55s, per side, all-time AND this week — put through every kill-test

SIDE	WINDOW	N	NET \$	\$/TRADE	STRIP BEST 3	WORST LEAVE-ONE-DAY-OUT	HALF 1	HALF 2	GREEN DAYS
SHORT	all-time	159	+2291	+\$14.41	+1780	+1889 (drop 07-29)	+1055	+1236	14/17
SHORT	this week	43	+809	+\$18.81	+420	+461 (drop 08-07)	+102	+707	4/5
LONG	all-time	160	+1157	+\$7.23	+821	+623 (drop 07-28)	+220	+937	10/20

LONG	this week	38	−30	−\$0.79	−278	−429 (drop 08-04)	+328	−358	2/6
------	-----------	----	-----	---------	------	----------------------	------	------	-----

Read the SHORT row and then the LONG row. Short is green on every single cut, in both windows: strip its three best trades and it still makes \$1,780 all-time and \$420 this week; remove its best day and it still makes \$1,889 and \$461; both halves green in both windows; 14 of 17 green days all-time and 4 of 5 this week. **That is as clean as anything gets on this desk.**

Long is a different animal. All-time it passes — +\$1,157, survives strip-3 and leave-one-out, both halves green. **But this week it fails everything.** Flat headline (−\$30), negative after stripping three trades (−\$278), negative if you remove 08-04 (−\$429), and its two halves go +\$328 then −\$358. And note the green-day column all-time: **10 of 20 — a coin flip**, versus short's 14 of 17. Long has always been the weaker side; this week it stopped paying.

The inversion — the live desk and the shadow book disagree about which side is good

Here is the thing I did not expect. This is the same week, the same gate, the same 55-second veto, the same 1-ATR stop, the same A/B target ladder. The only structural difference is that **the live desk is gated by the direction router and the shadow book is not.**

SIDE	LIVE N (LOTS)	LIVE NET \$	LIVE WIN%	LIVE \$/ TRADE	SHADO W N	SHADO W NET \$	SHADO W WIN%	SHADO W \$/ TRADE
LONG	32	+566.0	63	+\$17.70	38	−30.0	37	−\$0.79
SHORT	11	−814.5	0	−\$74.00	43	+809.0	51	+\$18.81

On the long side the live desk beat the unfiltered signal by **\$18.49 a trade**. On the short side it LOST to the unfiltered signal by **\$92.81 a trade**. Live abs_veto_short went **0-for-11** this week — nine stops and two max-holds, not a single target — while the identical ungated signal went 51% and made \$809.

Why I believe this is real and not sim optimism. The obvious objection is that shadow sims fill on clean 1-minute bars and overstate. That is true and it is a constant — it applies to the long side and the short side equally. **A constant bias cannot flip the SIGN of the comparison between two sides.** The long side shows live-better-than-shadow and the short side shows live-much-worse-than-shadow, under the same optimism. Something is systematically selecting good long entries and bad short entries, and the only thing that does any selecting is the arm/bench decision.

Where the short money actually was. Splitting the shadow short book by session tells the story:

BOOK	SESSION	N	NET \$	WIN%
shadow sw_absS_B_k10	overnight / pre-open (before 13:30 UTC)	16	+304.0	56
shadow sw_absS_B_k10	US session (13:30-21:00 UTC)	13	+542.5	54
LIVE abs_veto_short	overnight / pre-open	4	-375.5	0
LIVE abs_veto_short	US session (13:30-21:00 UTC)	7	-439.0	0
LIVE abs_veto_long	overnight / pre-open	12	-258.5	—
LIVE abs_veto_long	US session	20	+824.5	—

The shadow short book earned in **both** sessions. The live short book lost in both, on a fifth of the fires. Meanwhile the live LONG book has a crisp session signature: **-\$258.50 overnight, +\$824.50 in the US session.** That is a clean, actionable fact independent of the router argument: **abs_veto_long should not be armed before 13:30 UTC.**

THE PROMOTION VERDICT — FOLLOWING THE NUMBERS, NOT A PRIOR

PROMOTE: nothing new. There is no unproven gate with a fresh mechanism and enough n on this board.

TWO-SIDED: yes, keep both sides armed — the principle survives, the tuning does not. All-time both sides earn (+\$2,291 short, +\$1,157 long), so relegating the long DIRECTION would be exactly the mistake the rule exists to prevent. But this week short is robust on every cut and long fails every cut, so the two sides must stop sharing a policy. Specifically, from this week's numbers: **abs_veto_long gets a session guard (US only),** and

abs_veto_short's problem is not its exit — see §8: in SCALP-CHOP, where two thirds of its entries fire, every one of 70 exit policies leaves it negative. It is being armed into the wrong tape.

HOLD: the 1-ATR stop, both sides. The stop-width A/B settled it: k=1.0 beats k=2.0 on the short side by \$181 (Lot A) and \$181 (Lot B) over three days. Already deployed; not an action.

The honest hole: I cannot separate "the router picked badly on the short side" from "the live short fills were worse than the sim reprices" with the data I have, because there is no per-decision ledger tying each live short entry to the arming state that permitted it. That is one line of instrumentation and it is the highest-value build on this list.

5 — THE RELEGATION VIEW — per-side, derived fresh, nothing inherited

The rule stands: **tune the side, never relegate a direction.** So this table does not name gates to kill. It names, for each live gate and each side, what this week's numbers say the problem IS — the regime it was in, the exit it used, or the signal itself — and what the per-side fix would be. Live week P&L is the paper tournament 2026-08-03 to 2026-08-07, both lots.

GATE (SIDE)	LIVE N	LIVE NET \$	WIN%	\$/TRADE	EXIT MIX	WHAT THE NUMBERS SAY THE PROBLEM IS	PER-SIDE CALL
abs_veto_long	32	+566.0	63	+\$17.70	18 TARGET / 12 STOP / 2 claim	Not a problem — but overnight cost −\$258.50 while the US session made + \$824.50. And the exit is far too tight: it banked 1.0R/ 1.5R on moves that ran 7-12R (\$8).	KEEP + session guard + widen the exit
capitulation_long	6	+84.5	33	+\$14.10	2 TARGET / 4 STOP	Only 3 entries all week, both winners on 08-07. Its shadow twin capit_loose is	KEEP — router-only, never unfiltered

						–\$8,132 over 403 trades — so the LIVE gate is surviving purely because the router keeps it benched almost all the time. The bench IS the edge here.
nipc_long	33	–97.5	33	–\$3.00	11 TARGET / 22 STOP	EXIT, not signal. Repricing its 17 real entries: the deployed 2.0R/2.5R ladder gives – \$13.40 an entry and ranks 60th of 70 policies. A wide chandelier on the same entries gives + \$22.80. Three 08-03 entries each ran >10R and were banked at 2.5R. FIXABLE — change the exit
exhaustion_short	6	–162.0	33	–\$27.00	1 TARGET / 1 CHAND / 4 STOP	REGIME. Only 3 entries. Over the full 67-entry archive the deployed A0.75+tight ladder makes +\$13.90 an entry — but in MED-TREND it makes +\$14.40 where a wide chandelier makes + \$102.50. It is right in chop and badly wrong in trend. REGIME-DEPENDENT — split the exit
grind_long	6	–172.0	17	–\$28.70	1 TARGET / 5 STOP	REGIME. All 3 entries were on 08-05 in chop. Over 70 live entries: MED-TREND + \$10.20 an entry, SCALP-CHOP –\$29.80, and in chop no exit in the grid is positive (best – \$19.80). Grind REGIME-DEPENDENT — bench in chop

						must not run in chop at all.
nipc_short	16	−337.0	6	−\$21.10	1 TARGET / 15 STOP	SIGNAL. 1 winner in 16. Repriced across all 70 exit policies, the best is −\$35.20 an entry — every single exit loses. Nothing downstream saves it. PARKED — entry is broken
abs_veto_short	11	−814.5	0	−\$74.00	9 STOP / 2 MAX_HOLD	ARMING, not signal and not exit. 0-for-11 live while the ungated shadow twin went 51% for +\$809. Repriced on its 5 entries this week the best of 70 exits is still −\$12.60 (deployed −\$48.20); over its full 27-entry archive a better exit does turn it green (+\$6.10 vs −\$1.80 deployed) — but in SCALP-CHOP, where 18 of those 27 fired, every one of the 70 is negative. It is being armed into chop. KEEP THE SIDE — fix the arming
rgv_long / rgv_short	0	—	—	—	—	Did not fire live at all this week (last rgv fill was 2026-07-30). In shadow this week both sides bled: rg_short_025_r aw −\$104 / 19 trades, rg_short_025_f low50 −\$86, rg_long_fast_v −\$45. Its all-time short book is still the better of the two (+\$159 vs −\$1,483 long). NO DATA THIS WEEK — carry forward

THE RELEGATION SUMMARY IN ONE PARAGRAPH

Not one gate on this desk is a candidate for having a direction removed. Of the seven gates that traded, three have an EXIT problem (nipc_long, exhaustion_short in trend, abs_veto_long's too-tight ladder), two have a REGIME problem (grind_long in chop, capitulation_long unfiltered), one has an ARMING problem (abs_veto_short), and exactly one — **nipc_short** — has a signal that no exit in a 70-policy grid can rescue, and even that is PARKED on a 16-trade week, not refuted. Thin n can never be refuted. The lever this week is overwhelmingly which gate runs when and where it gets out, not which direction it may face.

6 — GRIND EXIT SHADOW: the two-ratchet runner-clip watch

Context in one line: grind_long's chandelier was the thing we spent three studies trying to improve, the two-ratchet mop-up arm was rejected every time, and it was left running in shadow to watch for the one thing a range-bound archive cannot produce — **a runner-clip on a real trend day**. The ~4R-and-bigger runners carry the whole of grind's net, so a second ratchet that banks one of them off a 4R plateau before it ignites would be an expensive mistake.

△ FIRST: THE WATCH WAS BROKEN, AND I FIXED IT BEFORE RUNNING IT

Both grind shadow watchers query WHERE gate='grind_long'. **The desk went multi-slot on 2026-07-29 and the gate column became grind_long_A / grind_long_B.** The exact match silently stopped returning rows — not an error, just four trades where there should have been twenty-six. Every reading these tools produced since 07-29 was computed off the four pre-multislot trades and printed with full confidence. I repointed both at gate IN ('grind_long', 'grind_long_B') (the chandelier lot) and ('grind_long', 'grind_long_A') respectively, and the numbers below are the corrected run. This is the third instrument in three weeks to return a confident wrong number from a silent zero-row join.

The ledger — data/two_ratchet_shadow.json, 2026-07-26
18:00 UTC to now

TRADES	RUNNERS (REACHABLE MFE ≥ 4R)	RUNNER- CLIPS	LIVE PURE-6.0 NET (LOT B)	SHADOW TWO- RATCHET NET	RECONSTRUCTED PURE NET
26	3	0	+\$29.50	-\$265.00	-\$376.50

Every runner, in full — this is the table the watch exists to produce:

OPENED (UTC)	ATR	REACHAB LE MFE	LIVE \$	RECONST RUCTED PURE \$	SHADOW TWO- RATCHET \$	\$ GIVEN UP TO THE CLIP	EXIT REASON
2026-07-27 18:54:06	24.8	4.20R	+23.0	+33.0	+144.5	\$0 — shadow BEAT pure by \$111.50	CHANDELIER
2026-07-30 13:45:06	46.7	5.21R	+243.5	+348.0	+348.0	\$0 — identical	MANUAL_CLAI M
2026-07-30 15:52:27	25.5	4.75R	+53.0	+62.5	+62.5	\$0 — identical	MANUAL_CLAI M

Three genuine runners crossed the watch, including one on 07-30 that reached **5.21R on a 46.7-point ATR** — the biggest single ATR reading in the whole archive, on the day MNQ moved 1,097 points with a 0.80 roundtrip. That is about as close to a real trend day as this summer produced. **The second ratchet clipped none of them.** On one of the three it actually finished \$111.50 AHEAD of the pure chandelier.

THE VERDICT FROM THIS WEEK'S LEDGER — AND IT IS NOT THE ONE I
EXPECTED

Is the shadow beating pure? No. Two-ratchet −\$265.00 against a reconstructed pure −\$376.50 — so on the identical 26 paths the two-ratchet is actually **\$111.50 AHEAD**, and every cent of that came from the single 07-27 runner. On the other 25 trades the two arms are within a few dollars of each other. That is exactly the old finding restated: the second ratchet's money is inseparable from a tiny number of trades.

Did it clip a real runner this week? No. Zero clips, on three runners, one of them a 5.21R monster. By the rule as written, zero clips over a genuine trend week is EVIDENCE FOR reconsidering it. So I will say plainly what the ledger says: **the clip-mirage concern did not materialise, and the two-ratchet's rejection now rests entirely on "it adds nothing on ordinary trades", not on "it will cost you a runner"**. That is a materially weaker case for keeping it dead than the one it was killed on.

It stays dead anyway, for now — but PARKED, not refuted, with a stated revival condition: revive it for a live A/B if a further **5 runners** accumulate with zero clips AND the non-runner trades stop being a wash. $n=3$ runners is not enough to overturn three studies.

△ **The caveat you must weigh this against.** The cross-check line reads — \$406: the reconstructed pure-6.0 result is \$406 away from what the live book actually did. Two of the 26 (07-30 11:11 and 12:43) show live +\$159 and +\$156 against a reconstruction of −\$31 and −\$58. **The reason is that grind_long has not run the pure-6.0 two-lot chandelier since 2026-07-31** — the deployed ladder is Lot A at 2.5R plus Lot B on a wide lock-chandelier, with a quiet-tape clip on top. So the "live" column is a different exit from the "pure" column, and only the pure-vs-two-ratchet comparison (same path, same code, same day) is apples-to-apples. The clip verdict is sound because it is a counterfactual on identical paths; the absolute dollar levels are not the live book.

7 — GRIND EXIT SHADOW: the 2R-partial smoothness watch

The other grind-exit shadow. The idea: Lot A banks certainty at 2R while Lot B rides the chandelier for the tail. Grind runs `base_size=2` so it costs nothing to try. It was never sold as a way to make more money — it was sold as a **variance lever**: a shallower drawdown for a small give-up on the mean.

The ledger — `data/partial_shadow.json`, 40 grind entries over 5 trading days

METRIC

DIFFERENCE

	BASELINE (2 LOTS BOTH ON THE CHANDELIER)	2R-PARTIAL (LOT A SCALP-2R + LOT B CHANDELIER)	
cumulative net \$	-1382.4	-1074.9	+307.5 in the partial's favour
mean give-up (partial – baseline)	—	—	+\$307.50 — a GAIN, not a give-up
daily standard deviation \$	42.3	83.3	+97% — the partial is TWICE as volatile day-to-day
max drawdown \$	-1382.4	-1074.9	+22% shallower
worst day \$	-336.4	-286.0	+15% better
green-day %	0%	0%	no difference — all 5 days red
green-week %	0%	0%	no difference

The two equity curves, side by side. Cumulative \$ at the close of each trading day the watch has data for:

DAY	BASELINE CUMULATIVE	PARTIAL CUMULATIVE	BASELINE THAT DAY	PARTIAL THAT DAY	PARTIAL'S EDGE THAT DAY
2026-07-27	-278.0	-213.5	-278.0	-213.5	+64.5
2026-07-29	-482.0	-454.0	-204.0	-240.5	-36.5
2026-07-30	-818.4	-510.9	-336.4	-56.9	+279.5
2026-07-31	-1096.4	-788.9	-278.0	-278.0	0.0
2026-08-05	-1382.4	-1074.9	-286.0	-286.0	0.0

THE OPERATOR'S CHOICE — AND THE VARIANCE CLAIM DID NOT REPRODUCE THIS WEEK

Read the daily-std row again, because it is the opposite of what this shadow was set up to show. The 2R-partial's daily standard deviation is **83.3 against the baseline's 42.3** — it was twice as jumpy day to day, not smoother. And the reason is visible in the curve table: on 07-30 it saved \$279.50 in a single session and on three of the five days it made no difference at all. A lever that does nothing most days and something large on one day increases dispersion by definition.

So the honest framing is not "smoothness for a mean cost". On this week's data the 2R-partial **won on the mean (+\$307.50), won on max drawdown (22% shallower), won on the worst day (15% better), and LOST on daily volatility.** That is a different trade from the one that was described, and it is a better one: you are being paid to take it.

Your choice, stated plainly: if what keeps you up at night is the depth of the hole — the max drawdown and the worst single day — the partial is strictly better on both and hands you \$307 for the privilege. If what you care about is a quiet daily P&L line, the baseline is calmer. **My read of your own stated preference (worst-day pain, not day-to-day wiggle) says take the partial.**

△ **And the thing that most needs saying: THE 2R-PARTIAL IS ALREADY LIVE, in all but the exact R.** `data/exit_overrides.json` has run `grind_long: {a_r: 2.5, b: "wide"}` since 2026-07-31 — that IS Lot A scalping and Lot B riding the chandelier, at 2.5R instead of 2.0R. So this watch's "BASELINE" arm is a config the desk retired eight days ago, and its "shadow" arm is roughly what is deployed. **This section is measuring the past.** The action is to repoint the watch's baseline at the deployed A2.5+wide ladder and sweep the A leg (2.0 vs 2.5 vs 3.0) against it, which is what would actually inform a decision.

△ **Sample-size honesty:** 40 entries, 5 days, all five red, and the reconstruction is far more pessimistic than the live book — live `grind_long` across all lots since the 07-26 deploy is **+\$12.00**, against this reconstruction's **-\$1,382**. Do not read the absolute levels. Read only the DELTA between the two arms, which is computed on identical paths.

8 — THE REGIME-FLEX EXIT LAB

The desk's live exit control is now two files: `data/gate_switches.env` (which gates are on) and `data/exit_overrides.json` (what exit style each gate uses). The operator's 2026-07-31 proposal was a four-rung regime-graded ladder. This section grades the live trial, then tries to prove the ladder.

8(a) — grading the LIVE `exhaustion_short` fade-scalp trial on its actual fills

FIRST, A CORRECTION: THE TRIAL CONFIG IS NOT WHAT IS DEPLOYED

The brief describes an **A@0.5R / B@1.5R** exhaustion trial. That is not what is in the live file. `data/exit_overrides.json` has read `exhaustion_short: {a_r: 0.75, b: "tight", atr_split: 22, lo: {a_usd: 40, b_r: 1.75,`

b_floor_usd: 60}} since **2026-08-04 21:15 UTC**, confirmed by fourteen consecutive service-startup records in data/config_journal.jsonl all carrying config hash c567c30c6c05. So the trial actually running is **Lot A at 0.75R, Lot B on a tight chandelier**, with a quiet-tape clip. I have graded that.

Every exhaustion_short fill this week, with the regime it fired into and how far the trade could actually have run:

OPENED (UTC)	LOT	ENTRY	30-MIN ER	ATR	RUNG	REACHABLE MFE	P&L	EXIT REASON
08-05 14:31:48	A (0.75R)	29867.75	0.22	23.4	SCALP-CHOP	4.50R	−70.5	STOP
08-05 14:31:48	B (tight chand)	29867.75	0.22	23.4	SCALP-CHOP	4.50R	−70.5	STOP
08-07 13:30:48	A (0.75R)	29715.00	0.31	11.9	MED-TREND	12.66R	−47.0	STOP
08-07 13:30:48	B (tight chand)	29715.00	0.31	11.9	MED-TREND	12.66R	−78.0	STOP
08-07 14:48:15	A (0.75R)	29717.75	0.04	20.4	SCALP-CHOP	2.53R	+54.0	TARGET
08-07 14:48:15	B (tight chand)	29717.25	0.04	20.4	SCALP-CHOP	2.51R	+50.0	CHANDELIER

LOT	ENTRIES	HIT ITS TARGET / TRAIL	STOPPED	HIT-RATE	NET \$
Lot A @ 0.75R	3	1	2	33%	−63.5
Lot B @ tight chandelier	3	1	2	33%	−98.5
trial total, week	3	1	2	33%	−162.0

Three entries is not a grade, it is an anecdote — so I graded the trial config against the alternatives on all 67 live exhaustion entries in the archive instead. Same machinery, same 5-second bars, same 1-ATR stop, stop-first, 60-minute horizon, both lots, costs in:

RUNG	LIVE ENTRIES	DEPLOYED A0.75 + TIGHT CHAND	BEST IN A 70-POLICY GRID	DEPLOYED'S RANK	READ
SCALP-CHOP	49	+\$16.00 / entry	scalp 1.0/1.5 → +\$21.60	15 of 70	The trial is right here. \$5.60 off the grid best on 49 entries is inside the noise.
MED-TREND	15	+\$14.40 / entry	wideChand x2 → +\$102.50	54 of 70	The trial is badly wrong here. It banks 0.75R on trades

					running 12R. This is where the money is.
BIG-TREND	3	−\$23.00 / entry	scalp 0.5/1.5 → + \$14.50	16 of 70	n=3. No call possible.
all rungs	67	+\$13.90 / entry	wideChand ×2 → + \$23.30	23 of 70	Profitable as deployed, and leaving ~\$9.40 an entry on the table.

Trial grade: PASS in chop, FAIL in trend. The one thing this week's three fills do tell you supports it — the two losers fired at ER 0.22 and ER 0.31 and both stopped, and the winner fired at ER 0.04 (dead chop) and paid. The 08-07 13:30 loser is the whole story in one trade: ER 0.31, a MED-TREND rung, the move ran **12.66R** in exhaustion's favour at some point in the hour, and the tight ladder stopped both lots out for −\$125 before it got there. That is not a bad signal. That is a good signal with the wrong exit bolted to it.

8(b) — PROVING THE OPERATOR'S EXIT LADDER: the extended lab

The mid-week first look was chop-dominated, BIG-TREND came out at n=5-25, and the 3.5R-plus-chandelier "ride" rung was never tested. So I went back as far as solid data allows and rebuilt the lab.

What the extended lab actually ran

DIMENSION	FIRST LOOK (MID-WEEK)	THIS RUN — SCRIPTS/ EXIT_LADDER_LAB_V2.PY
window	07-16 to 07-31, chop-dominated	2026-07-15 to 2026-08-07 — 21 trading days, 283,341 five-second bars
entries repriced	a few hundred	3,321 real shadow entries across 8 gate families
policies swept	25 scalp combos + 3 chandeliers	35 scalp A/B combos + 18 chandelier / ride policies = 53 per cell
time-of-day	computed and then never used	every cell split ON (pre-13:30 UTC) vs US (13:30-21:00)
STAY-OUT rung	an ATR≥22 proxy	the deployed untradeable meter (src/ gazbot7/untradeable.py), per day
robustness	none	leave-one-day-out, strip-best-3, walk-forward halves on every winning cell

⚠ **A data-quality bug I had to fix first, and it had already produced a fake result.** shadow_trades.entry_atr contains impossible values on MNQ — **25,809.5, 7,351.7 and 1,846.5 points**, across capit_loose, grind_fast, thrust_loose, thrust_short_raw, rg_long_fast and

chand_k35. Every payoff here is $R \times ATR \times \$2$, so a single poisoned row moves a whole cell. On my first pass the lab reported **capitulation_long / SCALP-CHOP best exit = "scalp 3.5/4.0" at +\$1,114.90 a trade off a 6% win rate** — obvious nonsense, entirely produced by two rows. Guarding ATR to 3–60 points dropped **7 entries of 3,328** and the same cell then read — \$18.60 a trade. Seven bad rows out of three thousand were enough to invert a headline. **Any prior study that used entry_atr unguarded should be re-checked.**

★ **THE BIG-TREND ANSWER — the rung is not refuted, it is UNMEASURABLE on this tape**

Here is why the first look could not test the ride rung, and why this one cannot either. Of the **23,025 minutes** of tape with a valid 30-minute efficiency reading across 21 trading days:

RUNG	MINUTES	SHARE OF TAPE	GATE FAMILIES REACHING $N \geq 12$
SCALP-CHOP (ER < 0.30)	14,035	61.0%	8 of 8
STAY-OUT (untradeable meter)	5,576	24.2%	6 of 8
MED-TREND (ER 0.30–0.50)	2,905	12.6%	8 of 8
BIG-TREND (ER ≥ 0.50)	509	2.2%	3 of 8

Two point two percent. Eight and a half hours of BIG-TREND tape in twenty-one trading days. Even pooling every gate's entries over the whole archive, only three families clear $n=12$ in that rung, and **grind_long — the gate the ride rung was designed for — does not reach $n=12$ in BIG-TREND at all.** On the live fills it is worse: 3 BIG-TREND entries for abs_veto_long, 3 for exhaustion_short, out of 454.

So the answer to "finally prove or refute the 3.5R-plus-chandelier rung" is: **I cannot, and neither can any study run on this archive, and that is itself the finding.** Here is everything the three measurable cells say:

GATE FAMILY	N	DAYS	GRID BEST	\$/TR	WORST LEAVE-	A3.0 + WIDE CHAND	A3.5 + WIDE CHAND	WIDE CHAND $\times 2$
-------------	---	------	-----------	-------	--------------	-------------------	-------------------	-----------------------

				ONE- OUT	(THE LADDER GUESS)			
exhaustion_ short	16	7	midChand ×2	+24.7	+17.4	+14.7	+17.7	+16.3
thrust_raw	16	10	scalp 3.5/4.0	+9.7	−77.1	−23.2	−18.0	−40.4
rgv_long	12	6	scalp 0.5/1.0	−25.5	−32.4	−64.6	−72.9	−85.9

In all three, the ride policies LOSE to something tighter. On exhaustion, A3.5+wide makes \$17.70 where a mid chandelier makes \$24.70. On thrust_raw the whole cell is a mirage — +\$9.70 a trade that becomes **−\$77.10** if you remove one day, which is the textbook one-day artefact. On rgv_long every ride policy is the worst thing in the grid (wide chandelier ×2 = −\$85.90 a trade, dead last of 53).

Verdict on the ride rung: PARKED, not refuted. Where it can be measured at all it loses, but n=12-16 over 2.2% of tape cannot refute anything. **Revival condition: a fortnight containing ≥2,000 minutes of ER≥0.50 tape, or a redefinition of BIG-TREND that captures more than one fortieth of the sessions.** My recommendation is the second one — see the deliverable.

THE LADDER GUESS vs WHAT THE TAPE PROVED — all 24 measurable cells

The guess, as stated: BIG-TREND → A~3.0R + wide lock-chandelier. MED-TREND → A 1.5R / B 2.5R. SCALP-CHOP → A 1.0R / B 2.0R. STAY-OUT → flat. Here is that guess scored against a 53-policy grid on 3,321 real entries. **"rank" is where the guess placed out of 53.**

GATE FAMILY	RUNG	N	PROVEN BEST EXIT	\$/TR	WORST LOO	THE LADDER GUESS	\$/TR	RANK
exhaustion_ short	BIG-TREND	16	midChand ×2	+24.7	+17.4	A3.0 + wideChand	+14.7	33/53
exhaustion_ short	MED-TREND	88	tightChand ×2	+10.2	+5.5	scalp 1.5/2.5	+3.6	29/53
exhaustion_ short	SCALP- CHOP	240	tightChand ×2	+6.2	+2.9	scalp 1.0/2.0	+0.2	33/53
	STAY-OUT	120		+8.8	+3.5	FLAT (\$0)	0.0	—

exhaustion_short			tightChand x2					
grind_long	MED-TREND	77	scalp 0.75/1.0	+10.5	+6.5	scalp 1.5/2.5	-8.0	36/53
grind_long	SCALP-CHOP	707	scalp 0.75/1.0	-2.4	-3.6	scalp 1.0/2.0	-9.1	19/53
grind_long	STAY-OUT	265	scalp 0.75/1.0	+3.0	+0.2	FLAT (\$0)	0.0	—
abs_veto_long	MED-TREND	16	tightChand x2	-22.4	-45.1	scalp 1.5/2.5	-38.9	27/53
abs_veto_long	SCALP-CHOP	87	scalp 0.5/1.0	-43.6	-45.8	scalp 1.0/2.0	-46.0	9/53
abs_veto_long	STAY-OUT	47	wideChand x2	-40.8	-67.8	FLAT (\$0)	0.0	guess WINS
abs_veto_short	MED-TREND	23	A3.0 + wideChand	+10.5	+0.5	scalp 1.5/2.5	-11.4	31/53
abs_veto_short	SCALP-CHOP	71	scalp 3.0/3.5	-33.5	-42.9	scalp 1.0/2.0	-39.5	35/53
abs_veto_short	STAY-OUT	58	scalp 3.5/4.0	-20.8	-29.7	FLAT (\$0)	0.0	guess WINS
thrust_raw	BIG-TREND	16	scalp 3.5/4.0	+9.7	-77.1	A3.0 + wideChand	-23.2	23/53
thrust_raw	MED-TREND	92	A3.0 + tightChand	-31.5	-40.2	scalp 1.5/2.5	-38.3	45/53
thrust_raw	SCALP-CHOP	360	wideChand x2	-40.4	-46.1	scalp 1.0/2.0	-48.5	50/53
thrust_raw	STAY-OUT	240	scalp 3.5/4.0	-36.7	-41.8	FLAT (\$0)	0.0	guess WINS
rgv_long	BIG-TREND	12	scalp 0.5/1.0	-25.5	-32.4	A3.0 + wideChand	-64.6	49/53
rgv_long	MED-TREND	76	wideChand x2	-1.9	-14.4	scalp 1.5/2.5	-18.2	36/53
rgv_long	SCALP-CHOP	126	scalp 0.5/1.0	-9.6	-12.1	scalp 1.0/2.0	-12.9	3/53
rgv_long	STAY-OUT	81	wideChand x2	-0.3	-30.4	FLAT (\$0)	0.0	guess WINS
rgv_short	MED-TREND	13	tightChand x2	+12.5	-13.5	scalp 1.5/2.5	-10.1	50/53
rgv_short	SCALP-CHOP	24	A1.5 + tightChand	+0.4	-4.5	scalp 1.0/2.0	-6.6	13/53
capitulation_long	MED-TREND	41	scalp 0.5/1.5	-15.0	-20.2	scalp 1.5/2.5	-26.2	32/53
capitulation_long	SCALP-CHOP	264	scalp 0.5/1.0	-18.6	-19.5	scalp 1.0/2.0	-31.0	16/53
capitulation_long	STAY-OUT	84	scalp 0.5/1.0	-12.8	-16.1	FLAT (\$0)	0.0	guess WINS

WHAT THE LADDER GOT RIGHT AND WHAT IT GOT WRONG

WRONG — every R number in it. Across the eighteen tradeable-rung cells the guess's specific A/B targets never placed better than **3rd of 53** (rgv_long in chop, and that cell loses money anyway). Its median placing is about 32nd

of 53 — worse than picking a policy at random. The 1.5R/2.5R medium-trend rung in particular is bottom-half in six of the seven cells it appears in.

RIGHT — the STAY-OUT rung, spectacularly. Look at the six STAY-OUT rows where the guess wins: on untradeable days, **going flat beat the best available exit for abs_veto_long (−\$40.80/tr), abs_veto_short (−\$20.80), thrust_raw (−\$36.70), rgv_long (−\$0.30) and capitulation_long (−\$12.80).** The untradeable meter flagged 6 of 21 days as STAY-OUT (07-15, 07-20, 07-28, 07-31, 08-02, 08-06) and on those days no exit rescues a fader or a raw thrust. **The single most valuable rung in the whole ladder is the one that says "do nothing".**

RIGHT with two exceptions. Two gates keep earning on untradeable days and should NOT go flat: **exhaustion_short** (+\$8.80/tr on 120 entries, still +\$3.50 after removing its best day) and **grind_long** (+\$3.00/tr on 265 entries, +\$0.20 worst leave-one-out — thin, but positive). A blanket stay-out would cost you both.

RIGHT in direction on the trend rungs, wrong in instrument. The guess says "ride trends with a big R target". The tape says "ride trends with a CHANDELIER". Look at where the money actually is: exhaustion in MED-TREND on the live entries is +\$14.40 deployed against **+\$102.50** for a wide chandelier; abs_veto_short in MED-TREND is +\$38.30 against **+\$109.40**. A fixed big-R target caps you at exactly the R you guessed; a trailing chandelier takes whatever the move gives. That is the correction the ladder needs.

THE FIRST-LOOK LEADS — confirmed or refuted, one by one

MID-WEEK LEAD	VERDICT	THE EVIDENCE, ON 21 DAYS INSTEAD OF 12
"grind MED-TREND preferred a TIGHT 0.5/1.0 scalp, NOT 1.5/2.5"	CONFIRMED	Best on 77 entries is scalp 0.75/1.0 at +\$10.50/tr (worst leave-one-out still +\$6.50, both halves positive). The 1.5/2.5 guess makes −\$8.00 and ranks 36th of 53. A sign flip, on the same trades, from the exit alone.
	CONFIRMED (and stronger)	

"exhaustion was the only clean-positive fader — tight chandelier in chop, 1.5-2.5 in medium"

Exhaustion is positive in every single rung: BIG +\$24.70, MED +\$10.20, CHOP +\$6.20, STAY-OUT +\$8.80 — and every one of those survives leave-one-day-out (+\$17.40 / +\$5.50 / +\$2.90 / +\$3.50) and has both halves green. Tight chandelier is the best policy in three of four rungs. The "1.5-2.5 in medium" half is refuted: 1.5/2.5 makes +\$3.60 where the tight chandelier makes +\$10.20.

"thrust / rgv / capitulation BLED in every regime regardless of the exit"

CONFIRMED — and say it plainly

thrust_raw: negative in all 53 policies in MED-TREND, SCALP-CHOP and STAY-OUT (best cells -\$31.50, -\$40.40, -\$36.70). capitulation_long: negative in all 53 in all three of its rungs (-\$15.00, -\$18.60, -\$12.80). rgv_long: negative in all 53 in BIG-TREND and SCALP-CHOP. That is 1,700+ entries and roughly 90,000 policy-evaluations that never find a green cell.

So I will say it as plainly as I can: for thrust_raw, capitulation_long and rgv_long, THE EXIT IS NOT THE LEVER. You cannot tune your way out of these with an R target, a chandelier, a time-of-day split or a regime split — I tried all four across 53 policies and 21 days and every combination loses. The lever is the ENTRY, and more specifically it is which gates you allow to run at all. Note carefully that abs_veto_55s is the SAME thrust signal as thrust_raw with a 55-second confirm bolted on the front, and it is +\$3,447 while thrust_raw bleeds in every cell. **That is the entire argument in one comparison: the money was in the entry filter, not the exit.**

9 — THE DELIVERABLE: the per-rung exit policy, ready for data/exit_overrides.json

This is the table the lab was built to produce. Each cell is the policy that won its (gate × rung) segment and then survived leave-one-day-out. **Confidence is graded honestly** — HIGH means $n \geq 60$ and the worst leave-one-out is still positive; MEDIUM means $n \geq 20$; THIN means take it as a lead only. Where the answer is "don't trade this here", I have said so rather than picking the least-bad exit.

GATE	SCALP-CHOP (61% OF TAPE)	MED-TREND (12.6%)	BIG-TREND (2.2%)	STAY-OUT (24.2%)
exhaustion_short	keep A0.75 + tight chandelier — +\$16.00/entry live, +\$6.20/tr	CHANGE → wide chandelier ×2 — +\$102.50/entry vs +	mid chandelier ×2, +\$24.70/tr. THIN (n=16)	DO NOT go flat — tight chandelier keeps +\$8.80/tr on 120 entries. HIGH

	shadow, rank 15/70. HIGH	\$14.40 deployed. MEDIUM (n=15 live, 88 shadow)		
abs_veto_long	CHANGE → scalp 1.5/6.0 or tight chandelier ×2 — +\$11.70 / +\$10.80 per entry vs +\$5.40 for the deployed 1.0/1.5. MEDIUM (n=35 live entries)	CHANGE → tight chandelier ×2 — +\$5.20/entry vs —\$4.70 deployed — a sign flip. MEDIUM (n=12 live)	n=3 live (deployed + \$62.40, tight chand + \$87.40). No call. NO DATA	FLAT. Best of 53 is — \$40.80/tr. HIGH
abs_veto_short	BENCH. Every one of 70 exits is negative (best — \$3.70/entry). This is an arming problem, not an exit problem. HIGH	CHANGE → wide chandelier ×2 — + \$109.40/entry vs + \$38.30 deployed. THIN (n=7 live)	no measurable cell. NO DATA	FLAT. Best of 53 is — \$20.80/tr. HIGH
grind_long	BENCH. Deployed — \$29.80/entry; best of 70 still —\$19.80. No exit is positive. HIGH (n=49 live, 707 shadow)	keep A2.5 + wide chandelier (+\$10.20/entry, rank 13/70) or tighten to mid chandelier ×2 (+\$20.40). MEDIUM	never reaches n=12. NO DATA	DO NOT go flat — scalp 0.75/1.0 keeps +\$3.00/tr on 265 entries (worst LOO +\$0.20 — marginal). MEDIUM
capitulation_long	hold A1.5 + tight (+ \$2.80/entry, rank 10/70; best is tight chandelier ×2 at +\$7.00). Live only because the router benches it. MEDIUM	wide chandelier ×2, + \$74.90/entry vs +\$43.30 deployed. THIN (n=4 live)	no measurable cell. NO DATA	FLAT. Best of 53 is — \$12.80/tr. HIGH
nipc_long	CHANGE → wide chandelier ×2 — a sign flip: +\$34.60/entry against —\$4.80 for the deployed 2.0/2.5, which ranks 47/70. MEDIUM (n=12 live)	all 70 policies negative (best —\$37.50). THIN (n=3)	no measurable cell. NO DATA	no measurable cell. NO DATA
nipc_short	BENCH. All 70 policies negative (best —\$18.10/entry). Signal problem. MEDIUM (n=6 live, 16 lots)	no measurable cell. NO DATA	no measurable cell. NO DATA	no measurable cell. NO DATA
rgv_short	A1.5 + tight chandelier, + \$0.40/tr — basically a scratch. THIN (n=24)	tight chandelier ×2, + \$12.50/tr but worst-LOO —\$13.50 — one day carries it. THIN (n=13)	no measurable cell. NO DATA	no measurable cell. NO DATA
rgv_long	BENCH. All 53 negative (best —\$9.60/tr on 126 entries). HIGH	wide chandelier ×2 — \$1.90/tr — a scratch at best. THIN	all 53 negative, best — \$25.50. THIN (n=12)	FLAT. MEDIUM

THE THREE STRUCTURAL RECOMMENDATIONS THAT FALL OUT OF THAT TABLE

1. Collapse the ladder from four rungs to three. BIG-TREND is 2.2% of the tape and never reaches a testable sample for the gate it was designed for. Redefine the top rung as **TREND = ER ≥ 0.30** (14.8% of tape, n=2,905 minutes, every gate reaches n≥12) and put the RIDE behaviour there. You lose nothing you could measure and you gain a rung you can actually tune.

2. On the trend rung, RIDE WITH A CHANDELIER, not with a big R target. This is the single biggest number in the section: exhaustion + \$102.50 vs +\$14.40, abs_veto_short +\$109.40 vs +\$38.30, capitulation + \$74.90 vs +\$43.30, nipc_long +\$34.60 vs —\$4.80. Four gates, same

direction, and the mechanism is obvious — a fixed target caps you at the R you guessed, a trail takes what the move gives.

3. STAY-OUT is a per-gate call, not a desk-wide one. Flat is right for `abs_veto` (both sides), `thrust_raw`, `rgv_long` and `capitulation_long`. It is **WRONG** for `exhaustion_short` (+\$8.80/tr) and `grind_long` (+\$3.00/tr), which keep earning on untradeable days. Implement it as a per-gate flag in `exit_overrides.json`, not as a global kill.

△ **And the standing caveat on all of it.** These are one summer regime, 21 trading days, and the live-entry cells are $n=3$ to $n=70$. Every "CHANGE" row above is a shadow nomination, not a deployment order. The right next step for each is a paired live A/B on the same feed — the same discipline that settled the stop-width question — not a config edit off a backtest.

10 — INSTRUMENT FAILURES FOUND WHILE WRITING THIS SECTION

Four tools returned confident wrong numbers, or refused to run at all. This matters more than any single result here, because a prior report may have quoted them.

TOOL	WHAT WAS WRONG	BLAST RADIUS	FIXED?
<code>scripts/archive_data.py</code>	Hard-coded <code>ATTACH '/home/alphabot/alphabot2/data/ticks.db'</code> . That database was retired to parquet on 2026-08-05, so the loader crashed outright — and it is the shared loader for both grind shadow watchers and several backtests.	every tick-honest backtest since 08-05	YES — rebuilt as a 3-source union (V5 parquet → nightly tape mirror → <code>capture.db</code> , freshest day wins). Now spans 32,025,269 ticks over 23 trading days, 07-05 to 08-07 .
<code>two_ratchet_shadow_watch.py</code> <code>partial_shadow_watch.py</code>	<code>WHERE gate='grind_long'</code> stopped matching when multi-slot renamed the column to <code>grind_long_A/B</code> on 2026-07-29. Silent zero rows, no error. The watch reported on 4 trades instead of 26.	every grind shadow reading since 07-29	YES — both repointed at the lot-suffixed names, choosing the correct lot for each watch.
<code>shadow_trades.entry_atr</code>	Poisoned values — 25,809.5 / 7,351.7 / 1,846.5 points on MNQ. Seven rows in 3,328. Because every payoff is $R \times ATR \times \$2$, they inverted a whole cell's headline (+ \$1,114.90/trade → −\$18.60/trade once guarded).	any study using <code>entry_atr</code> unguarded	GUARDED in the lab (3–60pt), not fixed at source — the writer that produces them is still live.

`adaptive_exit_bt.py` (mid-week)

the mid-week first look

Computed a time-of-day bucket and then never used it in any output. Also read 5s bars with a window that produced tiny BIG-TREND cells.

SUPERSEDED by exit_ladder_lab_v2.py, which segments by ToD and runs the full archive.

✓ capture.db retention — the GOOD news

The 5-day retention that has broken other tools applies to ticks only. The bars table retains 5-second bars for the whole archive — 283,341 bars from 07-15 to 08-07. Any study that needs 5s resolution over a long window can use it today.

unlocks the extended lab —

11 — DISPOSITION TABLE: every lead this section touched

LEAD	VERDICT	THE NUMBER THAT DECIDED IT	REVIVAL CONDITION (PARKED) / THE TEST THAT KILLED IT (REFUTED)
abs_veto_55s — thrust + 55-second continuation confirm	LIVE — re-confirmed, not a new promotion	+\$3,447 / 319tr; 12/17 green days; chop +\$1,438 AND trend +\$2,008; both halves green; +\$3,342 over un-vetoed thrust	Already deployed 07-29 as abs_veto_long/abs_veto_short. NOT AN ACTION.
abs_veto SHORT side — keep it two-sided	LIVE — the robust side	+\$2,291 all-time / +\$809 this week; strip-3 +\$1,780 / +\$420; worst-LOO +\$1,889 / +\$461; 14/17 and 4/5 green days	Survives every kill test in both windows. The problem is arming, not the side.
abs_veto LONG side — decaying, needs a session guard	SHADOW — tune the side	+\$1,157 all-time but −\$30 this week; strip-3 −\$278; worst-LOO −\$429; 10/20 green days all-time. Live: −\$258.50 overnight vs +\$824.50 US.	Revive as a full-confidence side if a US-session-only guard produces 4 more green weeks. Do NOT relegate the direction — all-time it earns.
1-ATR stop on abs_veto (vs 2-ATR)	HOLD — already live	Short Lot A k1.0 +\$508 vs k2.0 +\$384; Lot B k1.0 +\$847 vs k2.0 +\$666. Targets held in points, so only the stop moved.	Settled by paired live-feed arms. NOT AN ACTION.
the two-ratchet mop-up arm on grind's chandelier	PARKED — case is weaker than it was	26 trades, 3 runners (4.20R / 5.21R / 4.75R), ZERO clips; two-ratchet −\$265 vs reconstructed pure −\$377 — \$111.50 AHEAD, all of it from one runner	Revive for a live A/B if 5 further ≥4R runners accumulate with zero clips AND the non-runner trades stop being a wash. The clip-mirage fear did not materialise; it is now rejected only for adding nothing on ordinary trades.
the 2R-partial (Lot A banks, Lot B rides)	ALREADY GRADUATED	+\$307.50 on the mean, 22% shallower max drawdown, 15% better worst day — but daily std 83.3 vs 42.3, i.e. MORE volatile	Live as grind_long: a_r 2.5 / b wide since 07-31. The watch's "baseline" is a retired config. Action: repoint the watch and sweep the A leg (2.0 / 2.5 / 3.0) against what is deployed.
exhaustion_short fade-scalp trial (A@0.75R + tight)	PASS in chop / FAIL in trend	Week: 1-for-3, −\$162. Archive: +\$16.00/entry in chop (rank 15/70, keep) but +\$14.40 in MED-TREND where a wide chandelier makes +\$102.50 (rank 54/70)	Split the exit by rung. Note the brief's A@0.5/B@1.5 is NOT the deployed config — a_r 0.75 / b tight has been live since 08-04 21:15 UTC.
the BIG-TREND ride rung (A 2.5-3.5R + wide lock-chandelier)	PARKED — unmeasurable, not refuted	ER≥0.50 is 509 minutes = 2.2% of 21 trading days. Only 3 of 8 gate families reach n=12;	Revive on a fortnight with ≥2,000 minutes of ER≥0.50 tape — or, better, redefine TREND as

		grind_long reaches it in none. Where measurable, ride policies lose (thrust +\$9.70 → −\$77.10 on leave-one-out; rgv_long − \$85.90, dead last).	ER≥0.30 (14.8% of tape) and test the ride there, where every gate has n.
the STAY-OUT rung (untradeable meter → flat)	CONFIRMED — the ladder's best idea	On 6 of 21 flagged days, flat beats the BEST of 53 exits for abs_veto_long (−\$40.80/tr), thrust_raw (−\$36.70), abs_veto_short (−\$20.80), capitulation_long (−\$12.80)	Deploy per-gate, not desk-wide: exhaustion_short (+\$8.80/tr) and grind_long (+\$3.00/tr) keep earning on those days and must be exempt.
the ladder's R numbers (0.5/1.5 faders, 1.5/2.5 momentum, 2.5/ wide trend)	REFUTED as an optimum	Named test: full 53-policy grid on 3,321 entries across 24 gate×rung cells, with leave-one-day-out and walk-forward halves. The guess never ranked better than 3rd of 53 and its median placing is ~32nd.	No reformulation of these specific R values survives — but the ladder's structure (regime → exit) is confirmed, and the STAY-OUT rung is confirmed. It is the numbers that are refuted, not the idea.
"a better exit can rescue thrust_raw / capitulation_long / rgv_long"	REFUTED	Named test: 53 policies × 3-4 rungs × 3 time-of-day buckets, ~1,700 entries. Not one green cell. thrust_raw best cells − \$31.50 / −\$40.40 / −\$36.70; capitulation_long −\$15.00 / − \$18.60 / −\$12.80.	No reformulation on the exit axis survives. The same thrust signal WITH a 55-second entry confirm is +\$3,447 — so the lever is provably the entry.
nipc_short	PARKED — entry broken	1 winner in 16 lots, −\$337 live. Best of 70 exit policies on its real entries is −\$35.20/entry — every exit loses.	Revive only with a changed ENTRY condition and n≥40. Thin n can never be refuted; this is 6 entries.
nipc_long	SHADOW — fixable exit	Deployed 2.0/2.5 ladder ranks 60 of 70 at −\$13.40/entry; a wide chandelier on the identical entries makes +\$22.80. Three 08-03 entries each ran >10R and were banked at 2.5R.	A sign flip from the exit alone, on 17 entries. Needs a paired live A/B before deployment.
abs_veto_50s / abs_veto_60s — the confirm-window axis	SHADOW — keep both arms	50s +\$2,609/347tr (\$7.52), 55s +\$3,447/319tr (\$10.81), 60s + \$299/111tr (\$2.69). A clean single peak at 55.	Keep running — a peak with monotone shoulders is a parameter plateau, which is what you want to see. Re-read next Friday.
rgv (both sides)	NO DATA THIS WEEK	Zero live fills since 2026-07-30. Shadow this week: rg_short_025_raw −\$104, rg_short_025_flow50 −\$86, rg_long_fast_v −\$45.	Carry forward. Nothing this week changes its standing either way — and I will not manufacture a verdict from a week it did not trade.

THE ONE STONE STILL UNTURNED

There is no ledger tying each live entry to the arming state that permitted it. That is the missing measurement behind the biggest number in this section — live abs_veto_short at −\$74.00 a trade against the same ungated signal at +\$18.81. I can prove the two books disagree; I cannot prove why, because nothing records "gate X was armed at time T because of reason R, and here is the fire it allowed". Every other question in this section is a tuning question worth tens of dollars a trade. That one is worth **ninety-**

three dollars a trade on the short side alone, and it needs one write path, not a study.

Last word, and it is the uncomfortable one. The shadow book made **+\$779** this week on abs_veto_55s alone while the live desk lost **-\$932.50**. Some of that gap is sim optimism and I have said so repeatedly. But the sign flip between the two sides is not explainable by optimism, and the exit lab says the exits — while genuinely improvable, and worth roughly \$9 to \$88 an entry depending on the gate — are not where the loss came from either. **The desk did not lose money this week because its gates are bad or its exits are wrong. It lost money because of when they were switched on.** Everything in Part 2 points at the same door.

2.5 Mid-week musings — this week's leads, attacked

One thing survived everything I threw at it this week, Garrath, and it is **not** the thing that shipped on Thursday night. **Run state is real.** Trades taken with a directional run, on the run's own side, are worth about **\$27 a signal**; trades taken in chop are worth about **-\$13**. That separation beat 4,000 placebo shuffles, held on twelve sessions of tape it was never fitted to, and held again on Friday. But when I asked where the money actually is, the answer moved. Arming more would have added roughly **\$170** this week. Benching the 84 trades the desk took while no run was on would have saved **\$584**. And re-pricing the exits — on the exact same entries the desk already took, changing nothing about routing — turns the week from **-\$772** into **+\$1,493**. The run-state lead is right about the variable and wrong about the lever. It is not a routing find. It is an **exit** find wearing a router's coat.

HOW EVERY NUMBER BELOW WAS MADE

From scratch, this week, nothing quoted from Thursday's write-up. 1-minute bars rebuilt from the raw 5-second tape in `capture.db` (23,625 minutes, 07-15→08-07). Runs detected causally — trailing bars only, recomputed every minute, no look-ahead. Missed entries are the tournament's own SUPPRESSED-OPEN lines (37,404 raw lines 08-03→08-07), collapsed to distinct episodes and then **raced target-vs-stop on raw ticks, first touch** — stop $1 \times \text{ATR}$ at entry, 90-minute horizon, \$2/pt, \$1.50 round-trip. A race, not an MFE count. Live trades use their real booked P&L, `data_quality IS NULL`. Where a raced figure and a live figure appear side by side they are the same entry, priced two ways.

LEAD 1 — RUN STATE. What reproduced, what didn't, and what it is actually worth.

First: the re-derivation disagrees with Thursday's write-up on the headline cell

The run inventory reproduces almost exactly — I find **26 runs Mon-Thu** (31 across the full week), against the 26 in the lead doc. The bucketing does not. Same 89 Mon-Thu trades, different split, and the cell the whole argument rests on flips sign.

MON 08-03 → THU 08-06	LEAD DOC SAID	RE-DERIVED HERE
in-run aligned, TAKEN	n=34 · 50% · +\$10.34/tr	n=46 · 41% · -\$1.83/tr does not reproduce
in-run counter, TAKEN	n=4 · -\$20.38/tr	n=10 · 0% · -\$34.85/tr same direction, worse
chop, TAKEN	n=51 · 16% · -\$28.43/tr	n=33 · 18% · -\$19.36/tr same direction, smaller
in-run aligned, MISSED (raced)	n=61 · 67% · +\$2,141	n=96 · 59% · +\$2,703 reproduces
in-run counter, MISSED	n=16 · -\$643	n=30 · 3% · -\$1,192 reproduces
chop, MISSED	n=183 · 27% · -\$1,517	n=194 · 26% · -\$2,089 reproduces

What that changes. The claim "armed inside an aligned run the desk makes ~\$10/trade" is **not supported**. Armed inside a run the desk made nothing — it lost \$1.83 a lot. The genuine spread is between the signals it blocked (+\$27/signal aligned) and the chop it traded (−\$19/lot), which is about **\$17-\$46 a trade depending which pair you compare**, not the flat

\$39 quoted. And the fact that the desk could sit inside a run, aligned, and still make zero is the single most important number in this section. Hold that thought.

I could not make the lead's 34/4/51 split come back under any run-span definition I tried (merge gaps 0s, 120s, 300s, 600s, 900s all give the identical 46/10/33 — the labelling is completely insensitive to that knob; only the run count moves, 30→26). So the difference is in a definition that was not written down. **Use these numbers, not those.**

Attack 1 — PLACEBO. It passes, clearly, twice.

Two nulls, 4,000 draws each, on the full week's 114 aligned missed signals (observed **+\$26.97/signal**).

PLACEBO	WHAT IT DESTROYS	NULL MEAN	NULL P95	OBSERVED	BEATEN BY
side shuffle	keeps the run windows, randomises which side is "aligned"	+\$7.88	+\$19.44	+\$26.97	0.00%
duty-matched random	random windows, same 31% duty cycle, same length mix, random side	−\$4.09	+\$9.97	+\$26.97	0.03%

Read the first row carefully. A random side inside the real run windows still earns **+\$7.88**. So roughly **a third** of the effect is simply "these are the busy, wide-range parts of the tape" — and **two thirds is the direction call**. That is the honest decomposition, and it is the first time the desk has separated the two.

PASSED The label carries information. This is not the refuted TREND_UP tag from the 08-01 weekend, which a 20,000-draw placebo said was noise.

Attack 2 — IT IS LATE. Worse than advertised, and shrinking the window is refuted.

The lead doc said the 30-minute window caught 20 of 61 aligned entries (~34%). Re-derived over the full week: it was already armed for **25 of 114 — 22%**. At the moment a run first declares, a median **83 points has already happened and 22 points are left**. Median detection lag **33 minutes**. Median share of the move still available: **20%**.

So: should the window be shorter? I swept 96 configurations (window 5–45 min × net 40–100pt × ER 0.25–0.55) and scored each on **three separate samples** — Mon–Thu missed signals (in-sample), Friday's missed signals (never seen), and 489 taken trades across the twelve sessions 07-16→08-02 (never seen).

CONFIG	IS MON-THU \$/ SIG	OOS FRIDAY \$/ SIG	OOS 12 PRIOR SESSIONS \$/TR	VERDICT
W=15 · net≥40 · ER≥0.25 (best in- sample total, +\$760)	+\$7.60 n=100	−\$9.60	−\$1.64	overfit tell
W=5 · net≥40 · ER≥0.35	+\$17.59 n=34	−\$15.93	+\$26.57	fails Friday
W=30 · net≥80 · ER≥0.35 ← SHIPPED	+\$17.74 n=21	+\$66.67	+\$10.98 n=70	holds everywhere
W=30 · net≥60 · ER≥0.45	+\$39.37 n=8	+\$66.67	+\$30.03 n=35	best all three
W=20 · net≥60 · ER≥0.45	+\$20.85 n=26	+\$51.78	+\$13.72 n=67	best all three

The verdict is the opposite of the hope. Every configuration that catches more of the run catches it by lowering the efficiency bar, and every one of those loses money out of sample. The lever that works is not the window length — it is **ER**. Twenty-four configurations are positive in all three samples and they form a clean plateau at **window 20-30 minutes, ER 0.45**, with the net-points threshold barely mattering (40, 60 and 80 give identical results once ER binds). That is a real gradient, not a grid corner: ER 0.55 falls off the far side into n<5.

RECOMMEND a **moderate** move: keep the 30-minute window, lift RUN_ER_START **0.35 → 0.45** and RUN_ER_HOLD 0.30 → 0.40. It halves how often the detector says RUN (n=21→8 in-sample) and roughly doubles the

quality of what is left. Since run-state reports state and never flips a switch, this is a prompt-input change, not a switch change — cheap to try, cheap to revert. Do **not** shorten the window.

Attack 3 — THE EXITS EAT IT. This is where the section turns.

Same entries, two prices: what the desk actually banked, against what a plain 2R / 1×ATR race would have banked.

BUCKET	N LOTS	LIVE BOOKED	SAME ENTRIES RACED @2R	CAPTURE
in-run aligned — Mon-Thu	46	−\$84.0	+\$777.9	−11%
in-run aligned — full week	54	+\$246.5	+\$862.7	29%
chop	40	−\$650.5	−\$913.9	71%
counter-trend	14	−\$296.5	−\$828.6	36%

Look at the asymmetry in that last column. The live exit ladder takes **71% of the loss** when it is wrong in chop, and **29% of the win** when it is right in a run. On Mon-Thu it took less than nothing. That is not bad luck; it is a ladder tuned to bank small and bank fast, applied indiscriminately to the one regime where the whole point is to stay in. The exit-reason mix on the week's 54 in-run aligned lots: **29 STOP for −\$1,110.50, 23 TARGET for +\$1,277.50**, 2 manual claims +\$79.50. Roughly a coin flip on count — and the targets are too small to pay for the stops.

CONFIRMED The lead doc's attack #3 was right and if anything understated. The \$4,283 two-lot headline is not reachable. See LEAD 2 — this is the actual finding of the week.

Attack 4 — STRIP TUESDAY. One half survives, the other half dies.

SAMPLE	IN-RUN ALIGNED TAKEN	CHOP TAKEN	IN-RUN ALIGNED MISSED (RACED)
all five sessions	n=54 +\$4.56/lot	n=42 -\$17.20/lot	n=114 +\$26.97/sig
Mon-Thu, strip Tue 08-04	n=34 -\$15.88/lot	n=29 -\$17.71/lot	n=71 +\$29.62/sig

Strip Tuesday and the taken in-run book (−\$15.88) becomes **statistically indistinguishable from trading chop** (−\$17.71). Everything good the desk did inside a run this week happened on one day, with one gate. Meanwhile the missed book does not care at all — leave-one-day-out on all five sessions gives **+\$25.40 to +\$28.16 per signal**, a spread of under \$3.

That contrast is the whole diagnosis in two rows. The opportunity inside runs is robust and day-independent. The desk's ability to convert it rested entirely on Tuesday. Routing is not the constraint.

Attack 5 — THE ASYMMETRY. It only inverts if you fix the exit first.

ERROR	N	WHAT IT COST	PER TRADE
wrongly ARMED — traded in chop or counter-trend	56	−\$1,019 realised	−\$18.20
wrongly BENCHED — flat inside an aligned run, at the raced 2R rate	114	−\$3,074 forgone	−\$26.97
wrongly BENCHED — at the desk's own 29% in-run exit capture	114	−\$891 forgone	−\$7.82
rightly BENCHED — the chop and counter signals it blocked	280	+\$4,746 saved	+\$16.95

Here is the quantified answer the lead doc asked for, and it is a "yes, but". Priced at a competent 2R exit, being wrongly benched inside a run (−\$26.97) is worse than being wrongly armed in chop (−\$18.20) — the inversion is real, 1.5×. But priced at **the exit ladder the desk is actually**

running, wrongly benched costs \$7.82 and wrongly armed costs \$18.20 — the old rule **still holds, by 2.3×**.

So the standing guidance in CLAUDE.md — "inside a detected aligned run, BENCHED is the expensive error" — is **a promissory note, not a present fact**. It becomes true the day the in-run exit stops giving back 71% of the move. Until then, "when unsure, sit benched one more tick" is still the money-maximising rule even inside a run. That is uncomfortable and it is what the arithmetic says.

OUT OF SAMPLE — twelve prior sessions and Friday

Friday 08-07 is genuinely untouched: the thresholds were frozen Thursday night. So are the twelve sessions 07-16→08-02 (489 taken trades) which pre-date the idea entirely.

SAMPLE	ORACLE ALIGNED	ORACLE COUNTER	ORACLE CHOP	VERDICT
12 prior sessions, 489 taken trades	n=228 +\$11.46/tr	n=115 -\$29.66/tr	n=146 -\$16.65/tr	HOLDS
Friday 08-07 taken	n=8 +\$41.31/tr	n=4 +\$13.00/tr	n=9 -\$9.28/tr	HOLDS (thin)
Friday 08-07 missed (raced)	n=18 +\$20.61/sig	n=8 -\$45.92/sig	n=48 -\$22.86/sig	HOLDS

The concept is out-of-sample clean. A ~\$28/trade spread between aligned and chop, on 489 trades the idea never saw. That is the strongest thing in this report.

The live detector is a different story. Scored on what it would actually have said in real time:

SAMPLE	DETECTOR SAYS RUN	DETECTOR SAYS NO-RUN
12 prior sessions (taken trades)	n=86 +\$5.42/tr	n=374 -\$8.62/tr
this week (taken trades)	n=26 -\$7.25/tr	n=84 -\$6.95/tr
Friday only (taken trades)	n=6 -\$16.25/tr	n=15 +\$26.43/tr
all 17 sessions, strip-best-day (07-29)	n=105 -\$4.83/tr	—

The detector's arming signal fails strip-best-day. Across 17 sessions its RUN book is +\$277.50; remove 07-29 alone (+\$784.50) and it is -\$507 over 105 trades. On the very week it was built from it produced no

separation on taken trades, and on Friday it inverted. One consolation: that also means it is not simply in-sample flattery — if it were fitted to this week it would have looked good on this week, and it did not.

THE POLICY TEST — so what is arming actually worth?

Policy: arm the aligned side, and only the aligned side, whenever the live detector says RUN; flat otherwise. Priced on this week's real book.

COMPONENT	N	\$
what the desk actually did	110 lots	-\$772.50
— of which taken while the detector said RUN	26	-\$188.50
— of which taken while the detector said NO-RUN ← the bench prize	84	-\$584.00
signals the policy would have ADDED, raced @2R, 1 lot	26	+\$579.20
— at the desk's own 29% in-run capture ← the arm prize	26	+\$168.00
POLICY TOTAL		-\$20.50 (delta +\$752)

The bench is worth 3.5× the arm. \$584 from not trading when nothing is running, against \$168 from trading when something is. Almost all of run-state's value this week was in the negative. That answers the scope's open question directly: **run state is a bench instrument, not an arm instrument** — which is exactly what you would expect from something with a 33-minute detection lag and a 20% capture ratio.

VERDICT · LIVE Keep it exactly as it is — state-only, above the untradeable meter, Claude decides. Lift ER to 0.45. Lead with it as a reason to stay out, and treat "RUN, aligned" as permission rather than instruction.

VERDICT · PARKED Run-state as an arming trigger. It is late (22% capture), it fails strip-best-day on taken trades, and its arm value is \$168/week. **What would revive it:** a detection route that is not a trailing window — and note the run-catcher NULL [[run-catcher-null-all-microstructure]] says front-running the start is dead, so the realistic path is not earlier detection but better conversion once detected, i.e. LEAD 2.

LEAD 2 — THE EXIT LADDER. The money is here, and it needs no routing change at all.

This came out of Attack 3 and it is the biggest number in the section. If run state tells you which regime you are in, the obvious use is not "should I be on" — it is "**how far should I ride**". I raced **every entry the desk had this week — all 511 of them, 110 taken and 401 missed** — at six different targets against the same 1×ATR stop, segmented by regime. The answer is unambiguous and it points the opposite way in each regime.

R-target sweep by regime — \$ per entry, 511 entries, tick first-touch

REGIME	N	R=1.0	R=1.5	R=2.0	R=2.5	R=3.0	R=4.0	R=5.0	R=6.0
in-run aligned	168	+9.52	+17.30	+23.43	+32.10	+40.02	+51.02	+56.70	+56.83
chop / no run	282	−6.92	−10.69	−14.54	−17.64	−19.18	−19.80	—	—
counter-trend	52	−24.75	−33.98	−45.92	−45.63	−45.37	−44.81	—	—

Hit rate on the aligned row falls 57% → 46% → 38% as you widen from 2R to 6R and the money still goes up the whole way. It **plateaus at 4-6R** (+\$51.02, +\$56.70, +\$56.83) — a genuine flat top, not a grid edge. Leave-one-day-out at 4R across the five sessions: **+\$43.62, +\$46.47, +\$53.47, +\$53.67, +\$56.72**. Every single day positive, none of it resting on one session.

The chop and counter rows are monotone the other way: the tighter the better, and no target makes them positive. **There is no chop exit that turns chop into money.** Tight only loses less. The chop answer is the bench, not the ladder — which is exactly what LEAD 1 said.

Re-price the week's exits and nothing else

Same entries. Same order of events. Same stop. Only the target changes, keyed on run state at entry.

THE 54 IN-RUN ALIGNED LOTS THE DESK ACTUALLY TOOK		\$
what the live ladder banked		+\$246.50
the same 54 entries at 2R		+\$862.70
the same 54 entries at 3R		+\$1,975
the same 54 entries at 4R		+\$3,171

WHOLE WEEK, RUN-STATE-KEYED EXIT, SAME 110 ENTRIES	WEEK TOTAL
as it ran	-\$772.50
aligned 2.5R / chop 1.0R	+\$896
aligned 3.0R / chop 1.0R	+\$1,493
aligned 4.0R / chop 1.0R	+\$2,689

A \$2,266 swing on the exact same trades, with the router untouched.

Two honest hedges before you spend it: (1) this is a single fixed-R race against a live 2-lot scale-out with chandeliers and clips, so it is a config-vs-config comparison, not a drop-in; (2) the chop-at-1R leg reads +\$89 on the 40 taken lots but the full 282-entry chop population reads -\$6.92 each, so treat the chop leg as "lose less", not "make money" — on the population rate the same-entries figure lands nearer **+\$870**. Either way the sign and the order of magnitude do not move. And the counter-trend lots are the one place the live exits beat the race (-\$296.50 live vs -\$829 raced) — the clip is doing its job there.

The quiet-tape clip is keyed on the wrong variable

The clip fires on **ATR < 22**. Segment the same 511 entries by ATR-at-entry and that turns out to be nearly orthogonal to where clipping helps.

SEGMENT	N	R=1.0	R=2.0	R=3.0	R=4.0	WHAT IT WANTS
---------	---	-------	-------	-------	-------	------------------

aligned · ATR < 12	17	+10.64	+24.01	+34.32	+41.93	wide
aligned · ATR 12-20	79	+15.44	+26.62	+40.11	+43.32	wide
aligned · ATR ≥ 20	72	+2.77	+19.80	+41.26	+61.61	very wide
chop · ATR < 12	77	−5.17	−4.62	−4.87	−3.47	doesn't matter
chop · ATR 12-20	181	−9.02	−13.51	−17.98	−17.00	tight
chop · ATR ≥ 20	76	−15.88	−48.52	−54.44	−60.11	tight, urgently

The two cells that most need opposite treatment sit on the same side of the ATR-22 line. At $ATR \geq 20$ the aligned book wants 4R (+\$61.61) and the chop book wants 1R (−\$15.88 instead of −\$48.52) — and the clip, keyed on ATR alone, treats them identically. At $ATR < 12$ in chop, where the clip is active, tight versus wide is worth about \$1.70 a trade — nothing. The clip is doing its work in the segment where it barely matters and sitting out the segment where it is worth \$45 a trade.

Time-of-day agrees rather than competes: US session (13:30–21:00Z) aligned entries want the widest target of all (4R = **+\$60.10**, $n=102$), the overnight/London aligned entries top out earlier (4R = +\$36.97, best relative value at 2R = +\$27.23, $n=66$). ATR carries most of that, as the discipline note predicts.

VERDICT · SHADOW → SATURDAY BUILD Re-key the exit ladder on **run state**, not ATR. Concretely: RUN + aligned → Lot A 2.5R, Lot B rides (chandelier, no clip) regardless of ATR; NO-RUN → both lots clip at 1R-ish; counter → leave exactly as it is, the current clip is already the best of the three. Against the cheat-sheet's operator guesses, the proven numbers say **trend Lot A 2.5R is about right but Lot B must not clip**, and **fader 0.5/1.5 is right in chop for the wrong reason** — it wins by shrinking a loss, not by capturing an edge.

△ Caveats stated plainly: one week, 511 entries, 168 of them aligned. First-touch on a fixed $1 \times ATR$ stop with a 90-minute horizon. Both-touched-in-one-tick is resolved as a stop, so this is the conservative side. And 401 of the 511 entries are counterfactual — the desk was benched and they never ran.

LEAD 3 — GRIND: does it run only on a confirmed trend? (prior week, 2026-07-31)

First, the premise doesn't fully reproduce

The lead quoted live managed grind_long at **+\$440 on the 07-30 trend day** and **−\$370 on the 07-31 chop reversal**. Re-derived from the trade table, clean rows only: **07-30 = +\$353.50** over 36 lots, **07-31 = −\$284.00** over 4 lots (0 wins), and 07-29 = +\$263.50. Same story, ~20% smaller both ways. Its full live record across the 17 sessions it has traded: **128 lots, −\$958.00, 34% win**. Two green days out of ten.

The three enable-mechanisms, scored on grind's own live lots

MECHANISM	KEEPS	\$/LOT KEPT	BLOCKS	WORTH
(0) as it runs today — no filter	128 lots −\$958.0	−\$7.48	—	—
(a) trend-day only — session-to-date ER ≥ 0.20 at entry	10 lots −\$174.5	−\$17.45	118 lots −\$783.5	+\$784
(a') trend-day only — session ER ≥ 0.35	0 lots	—	128 lots	deletes the gate
(b) entry confirmation — run-state ALIGNED at entry	36 lots −\$200.5	−\$5.57	92 lots −\$757.5	+\$758
(c) 30-min ER floor 0.25 at entry (either side)	73 lots −\$230.5	−\$3.16	55 lots −\$727.5	+\$728
(c') 30-min ER floor 0.35	16 lots −\$124.0	−\$7.75	112 lots −\$834.0	+\$834

(a) is refuted on mechanism, not on sample size, and this is the clean kill. Grind is a continuation-on-pullback gate: its entries happen precisely when session efficiency is low, because a pullback is what depresses efficiency. On the 07-30 trend day, only **2 of its 36 lots** had session ER ≥ 0.20 at entry, and those two lost −\$120. At ER ≥ 0.35 it fires **zero times all year**. A trend-day gate on grind does not focus it — it switches it off. The

prior week's "grind's trend-day upside is proven" reads the day-level P&L and infers an entry-level condition that does not exist.

(b) blocks the most damage per lot kept but its edge is one day.

Leave-one-day-out on the run-state-aligned book: remove 07-30 and it goes from $-\$200.50$ over 36 lots to **$-\$846.00$ over 28 lots ($-\$30.21/\text{lot}$)**.

Everything good about mechanism (b) is 07-30 and nothing else. Worth noting separately: grind fired **zero** lots during a detected opposite-side run in 17 sessions — its own geometry already prevents counter-trend entries, so there is no counter-veto to add.

No mechanism makes grind positive. The best of them (c) cuts the loss from $-\$958$ to $-\$230$ and leaves a gate with no edge.

VERDICT · REFUTED "Grind trend-day-only" — mechanism (a). Its entries structurally cannot coincide with a high session-ER reading. There is no path; a trend-day filter is a delete button.

VERDICT · PARKED Grind itself. It stays **off**, but the reason is not the enable-mechanism — it is that grind wants to buy pullbacks and then exits them like a scalp. **What would revive it:** the LEAD 2 exit re-key (its aligned lots are exactly the population that wants 4R and gets 0.5R), plus the VWAP-extension gap noted below — it emits nothing at high extension, which is when a run is actually running. Re-test after the exit change with $n \geq 40$ aligned lots. Do not re-test it with another entry filter.

△ As instructed, all of the above is **live managed** grind_long, never the grind_fast shadow.

GATE MECHANICS — reported separately, because arming would have changed nothing

The runs no gate could see

Of the 31 runs this week, **2 produced not one aligned signal from any of the eight gates** — not a suppressed one, not a taken one. Arming the

entire roster would have produced zero trades. That belongs to Friday rehab, not to the router.

WINDOW	NET	RANGE	SIGNALS FROM THE WHOLE ROSTER	TRADES
08-04 15:41→19:40 (the afternoon up-legs)	+264pt	290pt	2, and both SHORT — abs_veto_short 17:27, exhaustion_short 18:20	0
08-05 07:06→07:36	-76pt	98pt	none at all (grind_long fired a counter LONG at 07:05)	0

The first row is worse than "silence". Over a four-hour stretch in which MNQ added 264 points, the roster's only two opinions were **against** the move. That is not a routing failure and it is not bad luck — it is eight gates that between them have no instrument for "price is going up and has been for hours". The related mechanism was flagged on 08-06: grind is a continuation-on-pullback gate and emits nothing at high VWAP extension, so at +6.6 ATR it logged zero suppressed opens for 30+ minutes while an [ACT] watcher demanded five times that it be armed.

Note this is fewer than the four windows named in the mid-week doc — 08-04 16:42 and 17:16 sit inside a longer run that did produce signals elsewhere in its span. I have merged them into the single silent stretch above rather than counting them three times.

HAND OFF

To Friday rehab: the desk needs one long-side continuation gate that fires at extension rather than only on a pullback. That is a build, not a switch.

TWO SMALLER LEADS THE WEEK THREW UP

(i) "Arm for periods" vs "wall of STOP → bench" — settled, and it's not close

These two standing instincts are in direct opposition and 08-04 is the test case. abs_veto_long, armed once at 04:20Z and left alone all day:

ENTRY (UTC)	RESULT	RUNNING	
07:28 10:46 11:45 — three losing entries in a row		−\$173.50	−\$173.50
13:32 14:07 15:03 15:14 15:40 — five winning entries in a row		+\$504.00	+\$330.50

A textbook "three stopped pairs → bench" rule fires at 11:45 and costs the entire +\$504. So I ran the rule properly, per gate, per session, across all 17 sessions:

RULE	17 SESSIONS — WHAT BENCHING WOULD HAVE SAVED	THIS WEEK — WHAT IT WOULD HAVE SAVED
bench after 2 consecutive losing entries	+\$1,905	−\$414
bench after 3	+\$28	−\$610
bench after 4	−\$334	−\$776
bench after 2, only if run-state says NO-RUN	+\$1,752	−\$492

It is an artefact of K. Moving the threshold by one loss swings the seventeen-session verdict from +\$1,905 to −\$334, and this week the rule destroys money at every K including the regime-conditioned version. A rule whose sign depends on an arbitrary integer is not a rule.

VERDICT · REFUTED

"Wall of STOP → bench" as a standing trigger.

VERDICT · LIVE

"Arming is a permission window" — the 08-04 morning losses were the price of the option and the option paid 2.9×. The discriminator is the **regime the losses happened in**, and the honest version of that is simply: bench on chop, never on a loss count.

(ii) CHURN — the desk armed and un-armed 70 times for nine trades

From router_trial_log.txt, 1,400 durable ticks across the five sessions:

DAY	TICKS	GATE FLIPS	ARM FLIPS	ARMS THAT PRODUCED A TRADE	MEDIAN ARM DWELL
2026-08-03	276	6	0	0	—
2026-08-04	276	17	7	2	35m
2026-08-05	276	17	5	1	35m

2026-08-06	276	13	3	1	15m
2026-08-07	288	17	10	5	—
WEEK	1,400	70	25	9 (36%)	35m

Of 19 arm windows that closed, **9 lasted 30 minutes or less and 7 lasted 15 minutes or less**. The five shortest were 5, 5, 5, 10 and 15 minutes. The worst single episode: `abs_veto_long` flipped **on/off/on/off/on/off/on — seven times in 2h25m — between 01:55 and 04:20 on 08-04**, overnight, and traded nothing at all until it was finally left alone at 04:20 and went on to make the week's only green day.

None of those 16 barren flips grades badly on the per-decision scoreboard, because a switch that produces no trade is scored neutral. **The rate is the metric**. A five-minute arm window is arithmetically incapable of catching a signal from a gate that fires a few times a day, so it is a guaranteed \$0 with a non-zero chance of catching exactly the wrong one.

RECOMMEND A **dwel floor**: once a gate is armed it stays armed for a minimum of 30 minutes unless a hard confirmed event fires against it (that exception is the fast-bench path, and it belongs to the router-operation review, not here). Cheap to state, cheap to enforce in the tick prompt, and it would have removed 9 of this week's 19 closed arm windows without touching a single trade — because none of those 9 produced one.

(iii) One instrument still lying — check before quoting

`signal_journal.suppressed_by` is **NULL in all 702 rows**, verified just now, despite being named in the INSERT. Anything computing `COALESCE(suppressed_by, 'TAKEN')` reports every fire as taken, including on benched gates on a zero-fill day. That is why this section reads suppressed entries out of the tournament journal directly rather than out of the journal table. If a past figure in any report was built on `signal_journal`, it is suspect.

WHAT SURVIVED — THE WHOLE SECTION IN ONE TABLE

LEAD	DISPOSITION	THE NUMBER BEHIND IT
Run state separates aligned from chop	LIVE	+\$27/sig vs -\$13/sig; beats 4,000 placebos at $p < 0.001$; holds on 489 out-of-sample trades
Run state is a bench instrument	LIVE	bench prize \$584 vs arm prize \$168 — 3.5×
Run-state-keyed exit ladder	SHADOW → SATURDAY BUILD	same 110 entries: -\$772 → +\$1,493 (aligned 3R / chop 1R); LOO-clean, 5/5 days
Lift RUN_ER_START 0.35 → 0.45	DEPLOY, LOW RISK	+\$39.37/sig in-sample, +\$66.67 Friday, +\$30.03 on 12 prior sessions — positive in all three
Dwell floor of 30 min on any arm	RECOMMEND	9 of the 19 arm windows that closed this week lasted ≤ 30 min and produced zero trades
"Arming is a permission window"	LIVE	08-04: -\$173.50 then +\$504.00, one arming decision
Run state as an arming trigger	PARKED	22% capture, 33-min lag, fails strip-best-day (+\$2.48 → -\$4.83/tr)
"The bench/arm asymmetry inverts inside runs"	PARKED	true at 2R (1.5×), false at today's exits (0.43×) — becomes true when LEAD 2 ships
Grind runs only on a confirmed trend day	REFUTED	2 of 36 lots on the trend day pass the filter, and they lost; at $ER \geq 0.35$ it never fires
Grind itself	PARKED	-\$958 / 128 lots; no filter makes it positive; revive after the exit re-key, $n \geq 40$
"Wall of STOP → bench"	REFUTED	+\$1,905 at $K=2$, -\$334 at $K=4$, negative at every K this week
Shorten the run-state window	REFUTED	best in-sample config (+\$760) loses on both out-of-sample legs
Two runs no gate can see	HAND OFF — FRIDAY REHAB	08-04 15:41-19:40, +264pt, roster's only two opinions were both short

THE HONEST HEALTH WARNING

One week of live book (110 trades, -\$772.50) and 401 counterfactual entries that never ran. The out-of-sample legs are real but thin on the Friday side ($n=3-18$) and use taken trades only on the twelve prior sessions. The exit ladder result is the most robust thing here — five-for-five on days, a genuine plateau, a large effect — and it is still one week, priced on a race rather than on the live two-lot ladder. Nothing in this section should ship without the Saturday re-validation, and the exit re-key in particular wants a forward A/B rather than a switch flip, because the last three exit cells that looked this good on a replay did not survive one.

P2.6 The router operation review — what works, what doesn't, and what to change on Saturday

Gaz, you asked for a full interrogation of the router, so here it is, and the headline is better than the desk's P&L would suggest. **By its own nightly repricing the router was worth +\$2,951 this week** — it blocked \$20,329 of would-be losses and gave up \$17,378 of would-be winners — against **−\$1,216 the week before**. That is its first clearly positive week. The independent check agrees: the gates it kept benched had shadow twins bleeding the whole time (capit_live_mirror −\$1,609, grind_fast −\$2,549 over the exact windows those gates were off). And it did it without thrashing — **1,393 ticks, 66 switch changes, 95.3% of ticks changed nothing**.

Three things need fixing and all three are concrete. **(1)** The thresholds do not need touching — ER 0.15 / net 30 came out top on both test sets — but the timing does: window 30→45 minutes and hold 2→3 marks is worth **+\$502 on the live-managed set and +\$399 on the full historical set**, and it is an interior peak on four of five axes, not a grid edge. **(2)** The untradeable meter separates green days from red days perfectly at end-of-day (every green day ≤ 23 , every red day ≥ 34) but is useless before the US open — its 13:30Z reading correlates +0.13 with the day's P&L, which is noise. The 65 cutoff catches 2 of 6 red days. **(3)** One gate is the whole leak: abs_veto_short was benched 91% of the week while its shadow twin made **+\$928**, and armed for 9% of it during which it lost **−\$559.00 on nine lots without a single winner**. The router had it backwards on both sides.

THE ONE THING THAT DID NOT SURVIVE — AND IT IS THE MOST USEFUL NEGATIVE RESULT IN THIS SECTION

I expected the 5-minute poll to be leaking money to reaction lag, and it is not. **Zero trades opened in the five minutes before any bench decision all week**. More than that: I backtested the alternative you asked about — an in-code, per-gate, zero-lag regime guard evaluated at entry time instead of a poll — and it is **worse**. On this week it would have returned +\$304 of router

value against the poll's +\$682; over 40 days, +\$267 against +\$475. The hysteresis that makes the poll "slow" is the thing earning the money. **Do not build the event-driven version.**

1 · The week's decisions and the scoreboard

Every number in this table is recomputed from data/router_trial_log.txt (one line per 5-minute tick, plus session lines) and data/gazbot7.db. Days are **Paris days** — 22:00 UTC to 22:00 UTC — because that is the unit the router itself works in, and P&L is on the clean ledger (data_quality IS NULL). That basis puts the week at **−\$953.00 on 112 lots**; P1 quotes −\$772.50 on 110 because it uses UTC calendar days, which pushes the two Monday-night lots out of the week entirely. Same trades, different fence.

DAY	TICKS	SWITCH CHANGES	CHANGE RATE	METER (EOD)	DESK P&L	WHAT THE ROUTER ACTUALLY DID
Mon 08-03	276	5	1.8%	34	−\$460.50	All five changes in the 22:00-23:59Z reopen block — benched five of six gates straight after the midnight re-arm. Zero changes for the following 22 hours.
Tue 08-04	277	9	3.2%	22	+\$330.50	Seven of nine changes fell in the 01:00-04:00Z Asia block, cycling abs_veto_long off and on four times each. The gate it kept fiddling with was the one that made the day.
Wed 08-05	276	16	5.8%	44	−\$817.00	The busiest non-Friday. abs_veto_long off 7× / on 5×. The day's damage (−\$474 in the 14:00Z hour) came from

						gates it had armed, not ones it failed to bench.
Thu 08-06	276	11	4.0%	78 STAY-OUT	-\$233.00	abs_veto_short armed and benched three separate times . First two never fired; the third fired and lost -\$129.50. The nipc kill-criterion fired at 13:09 exactly on the written -\$400 line.
Fri 08-07	288	25	8.7%	15	+\$227.00	276 durable ticks + 12 session ticks. 20 of 25 changes fell in 12:00-16:59Z, then the 17:00 and 18:00 hours passed with none at all. Three written expiries fired autonomously and correctly.
WEEK	1,393	66	4.7%	—	-\$953.00	1,326 ticks explicitly logged “changed: none”.

The thrash check passes, and it passes comfortably. Ninety-five percent of ticks changed nothing. But the raw rate hides the shape, and the shape is the finding: the changes are not spread across the week, they are bunched into two blocks. Here is every switch change by the hour it happened.

UTC HOUR BLOCK	WHAT IT IS	CHANGES	SHARE	READ
22:00-01:59	Paris reopen + the midnight all-on re-arm	21	32%	The router spends its first hours every night undoing the automatic re-arm
12:00-15:59	London close through the US cash open	22	33%	The hardest four hours of the day carry a third of all decisions
02:00-11:59	Asia + European morning	17	26%	Mostly Tuesday and Wednesday abs_veto_long cycling
16:00-21:59	US afternoon	6	9%	Quiet — and this is where abs_veto_long made its money on Friday

Two-thirds of every decision the router made this week happened either in the four hours around the US open or in the four hours after the midnight reset. **The reset block is self-inflicted.** gazbot7-gate-reactivate.timer turns everything back on at 22:00 and the router then spends 21 changes turning most of it back off. That is 32% of the week’s decision budget spent correcting a scheduled action, and it is the same “all-6-open reopen churn” the ledger already flagged on

07-31 — when it cost **-\$275** in real money (rgv_short -\$107, abs_veto_long -\$82, grind stops -\$86.50) before the router could get to it.

The graded ledger — where the scoreboard actually stands

data/router_badcall_ledger.md carries the running tally: **GOOD 30 · BAD 5 · CHURN 10 · NEUTRAL 6** as of Friday’s close. Parsed by session, here is how it accumulated — and one thing worth saying out loud.

LEDGER SESSION	▣ GOOD	▣ BAD	△ CHURN	▣ NEUTRAL	DAY P&L
2026-07-30 (Thu)	4	2	1	2	+\$1,172.00
2026-07-31 (Fri)	8	0	4	0	-\$905.00
2026-08-03 → 08-05	—	—	—	—	-\$947.00
2026-08-06 (Thu)	5	1	1	1	-\$233.00
2026-08-07 (Fri, 6 passes)	9	3	4	1	+\$227.00

△ A PROCESS HOLE, AND IT IS IN THE THREE WORST DAYS OF THE WEEK

The ledger has no graded entries at all for Monday, Tuesday or Wednesday. Those three days carry 30 of the week’s 66 switch changes and -\$947.00 of realised P&L, including Wednesday’s -\$817.00 — the worst session of the week. The grading resumes on Thursday and is dense on Friday. So the scoreboard you are reading is **not a sample of the week, it is a sample of the days somebody sat down to grade**, and those days are Thursday and Friday. That biases the tally toward the sessions with the most oversight, which are also the sessions where the router behaved best. Fixing this is a discipline item, not a code item: the 4-hourly review must run on the bad days too, or especially on the bad days.

2 · What paid, what erred — in dollars

Two independent ways to price a bench. **(a)** The gate’s shadow twin over the exact hours the gate was benched — that is the “what it would have done” counterfactual, and it is the honest one because the twin was running the same logic on the same tape with nobody stopping it. **(b)** The nightly’s repricing of every suppressed signal. Both are below. Shadow trades whose entry-to-exit distance exceeds 300 points are dropped as contract-roll contamination (13 rows across the week). **“Hours benched” is reconstructed from the switch-change log**, walking each gate’s on/off events across the 120-hour week; P1 counts arm-windows from the switch file itself and will differ by a few hours per gate. The share is what matters here, not the decimal.

GATE	HOURS BENCHED	% OF WEEK	ITS SHADOW TWIN OVER THE BENCHED WINDOW	LIVE P&L WHILE ARMED	VERDICT
capitulation_long	111.6 h	93%	capit_live_mirror – \$1,609 (70 tr) · capit_loose –\$3,566 (146 tr)	+\$84.50	BEST BENCH
grind_long	119.3 h	99%	grind_fast –\$2,549 (335 tr) · cx_grindA_live – \$202 (74 tr)	–\$172.00	RIGHT CALL
rgv_short	120.0 h	100%	rg_short_025_raw – \$89 (19 tr)	\$0.00	FREE
abs_veto_long	85.8 h	71%	cx_absLA_live –\$77 (18 tr) · sw_absL_A_k20 –\$78 (21 tr)	+\$385.50	WELL TIMED
exhaustion_short	112.7 h	94%	exhaustion_rev + \$218 (124 tr)	–\$162.00	MILD OVER-BENCH
abs_veto_short	109.2 h	91%	abs_veto_55s +\$928 (71 tr) · cx_absSB_live +\$262 (21 tr)	–\$559.00	BACKWARDS

Read the top two rows first, because they are the week’s case for the router existing. It kept capitulation_long off for 93% of the week and grind_long off for 99%, and over exactly those windows their unmanaged twins lost \$1,609 and \$2,549. That is the money that did not leave the account.

Everything else in this review is a rounding error next to those two decisions, and they were both essentially one decision each, held for days without fiddling.

Now read the bottom row, because it is the week’s indictment.

abs_veto_short’s shadow twins made **+\$928** and **+\$262** during the 109 hours it was benched, and the gate itself lost **−\$559.00 across nine lots with zero winners** in the 11 hours it was armed. The router was not merely wrong about this gate — it was anti-correlated. It sat the gate down while the opportunity was there and stood it up for the losses.

WHY — AND THE ANSWER IS NOT “THE ROUTER IS STUPID”

The ledger already caught this on 08-06 and it holds for the whole week: **the live gate and its shadow twin do not enter at the same price.** On 08-06 at 13:33 the live gate sold at ~29253, **66 points below the 29319 break low** its own twin sold; the twin banked +\$108.50 and the live gate stopped for −\$129.50 on the same setup. So the twin’s +\$928 is not a clean measure of what the live gate would have earned — it is a measure of what the setup was worth to a faster fill. The router was reading shadow P&L as a proxy for gate quality and the proxy is broken for this gate. That is banked as `[[shadow-green-does-not-mean-live-green]]` and it is the single most important mechanism lesson of the week.

The  BAD and  CHURN patterns, priced

These are the recurring error shapes from the ledger, with the dollars attached. Each one has a rule already written against it; the question for Saturday is whether the rules are being applied.

PATTERN	WHEN	WHAT HAPPENED	COST	RULE IT VIOLATED
Premature arm on a delta blip	07-30 15:49	Flipped to longs on a 4-bucket buyer-delta build and a 45pt poke of the range top, in ER 0.01 chop . Faded in 2 minutes.	~\$0 net (salvaged by operator Claims)	rearm-momentum-needs-er-climb-not-delta-blip
Benching the aligned trend-rider on open noise	07-30 13:36	Benched grind on the 13:33 false-breakout flush while	−\$431 avoided only by the override	us-open-dont-bench-trend-rider

		the day was cleanly trending up. Operator overrode; grind then booked +\$234 and + \$197.		
All-6-on at the reopen	07-31 21:58	Clean-slate midnight re-arm into a directional reopen; the wrong-side gates churned before the router could reach them.	-\$275	Reset-design cost — still unfixed, still costing 32% of the change budget
Re-arming into a “calm” recovery	07-31 16:41 & 17:37	Two abs_veto re-enables into ER 0.24–0.30 chop. Even veto-protected, even on scalp exits, both churned.	-\$245	Wait for ER≥0.35 sustained, not ER-0.25 elevated chop
Arming on shadow evidence	08-06 13:20	abs_veto_short armed because its shadow family was the day’s one green cell. Live entry was 66pt worse than the twin.	-\$129.50	shadow-green-does-not-mean-live-green (banked from this)
Holding a short through the cash open	08-07 13:20	Armed after a correct operator challenge — the reasoning improved, the outcome did not. Both lots stopped 59 seconds later at the bell, 11pt of slippage through the stops.	-\$155.50	Process was right, ATR-39 tape was the cost
Recommending the bench of the day’s only earner	08-07 morning	The router advised benching capitulation_long four separate times . It was never actioned. It then made +\$207 in ten seconds — the day’s entire P&L.	+\$207 saved by not listening	Fader benches must be on mechanism, never on day-sign
Zero-cost churn	08-07 12:30–16:35	20 switch changes in 4 hours; 16 never fired at all . Every one grades NEUTRAL in isolation. The \$0 was availability, not judgement.	\$0 realised, real risk	count-switch-changes-churn-is-invisible-at-zero-cost

THE CALLS THAT PAID	WHEN	EVIDENCE	VALUE
Bench both momentum longs on a regime flip	07-30 14:51	ER collapsed 0.40→0.20→0.07 off the highs; desk was flat, so this was a regime read not a reaction to stops. The unbenched momentum shadows lost -\$398 over that window.	+\$398
Bench the fader family on the meter	08-06 02:15	Meter 72 STAY-OUT, roundtrip 0.21, 66% given back. Held for 21 hours without thrashing. capit_loose ran -\$551 and capit_live_mirror -\$395.50 that day.	+\$400-550
Execute the written kill criterion	08-06 13:09	nipc hit exactly -\$400.00 on its 47th lot and was benched the same minute, then held in reactivate_gates.py: :HOLD so the 22:00 reopen cannot resurrect it.	rule, not \$
Three written expiries firing headless	08-07	13:05 grind_long on the ER<0.25 line · 14:20 abs_veto_short on the 29700 reclaim · 15:40	mechanism

		exhaustion_short re-armed on NO-RUN. Each condition was written into gate_switches.env before the fact and honoured by a headless tick.	
Re-arm exhaustion_short under the (c2) carve-out	08-07 14:25	The 14:20 bench had applied a momentum rationale to a fader. The tick caught it and re-armed. Went 2-for-2 at 14:48.	+\$104.00
Arm abs_veto_long and then leave it alone	08-07 15:25	Armed on a completed trio, then six entries over 4.5 hours with no further intervention.	+\$248.50
Revert an auto-re-arm built on a false number	08-07 13:52	The tick armed citing "ATR 39.1→27, the spike is over". Measured ATR14 was 42.1 — vol had risen. Codified as (d2): measure vol against the session baseline, never a window containing the spike.	error caught

3 • Net router value — the nightly rollup

Regenerated for every day of this week by running scripts/router_nightly.py and scripts/selector_nightly.py against the tournament journal (they had not been run since 07-31). Each blocked signal is repriced tick-honest under a scalp-2R / 1-ATR-stop exit. **Health warning on the size of these numbers:** every gate reports as an A and a B lot, so “saved” and “missed” are roughly double-counted — the net is unaffected but do not read \$20,329 as a real dollar figure. Read the net, and read the direction.

DAY	CHOP %	TREND UP %	TREND DOWN %	BLOCKED LOSSES SAVED	WINNERS MISSED	NET ROUTER VALUE	LEAKAGE
Mon 08-03	73%	23%	4%	\$4,020	\$3,135	+\$885	0
Tue 08-04	63%	30%	7%	\$2,306	\$3,176	−\$870	0
Wed 08-05	72%	12%	16%	\$4,832	\$3,323	+\$1,509	0
Thu 08-06	63%	20%	18%	\$4,533	\$4,540	−\$7	0
Fri 08-07	66%	27%	7%	\$4,638	\$3,204	+\$1,434	0
WEEK	—	—	—	\$20,329	\$17,378	+\$2,951	0
Previous week 07-27→07-31	—	—	—	\$15,788	\$17,003	−\$1,216	22 tr / −\$1,083

Two things flipped between the weeks and both matter. Net value went from −\$1,216 to +\$2,951, and leakage went to zero. Leakage counts managed gates

that fired while armed and lost in a regime they should have been benched in — the router letting a loser through. Last week that was 22 trades and $-\$1,083$; this week it was nothing, on any day. Whatever else is wrong, **the router is no longer leaving gates on in regimes that kill them.**

The three positive days are all high-CHOP days (73%, 72%, 66%) and the two flat-to-negative days are the two lowest-CHOP days (both 63%). That is the router doing exactly its designed job: in chop it benches the momentum gates and gets paid; in a day with more trend in it, its benches start costing winners. Tuesday is the clean example — 30% TREND_UP, the week's best trend day, and the router's benches cost $-\$870$ of would-be winners while the desk still made $+\$330.50$.

GATE	BLOCKED SIGNALS	WOULD-HAVE-WON	WOULD-HAVE-LOST	NET BLOCKED	VERDICT
grind_long	404	$+\$6,982$	$-\$9,186$	$-\$2,204$	the router's biggest single earner
abs_veto_long	88	$+\$1,623$	$-\$2,168$	$-\$545$	good block on net
nipc_long	22	$+\$568$	$-\$1,010$	$-\$442$	good block
abs_veto_short	172	$+\$3,342$	$-\$3,676$	$-\$334$	marginal — near coin-flip
nipc_short	34	$+\$1,353$	$-\$1,414$	$-\$61$	neutral
rgv_short	10	$+\$400$	$-\$294$	$+\$106$	over-bench
exhaustion_short	70	$+\$2,339$	$-\$2,178$	$+\$161$	over-bench — agrees with the shadow read
capitulation_long	20	$+\$771$	$-\$403$	$+\$368$	over-bench — and it disagrees with the shadow read

THE ONE PLACE THE TWO METHODS DISAGREE — AND YOU SHOULD TRUST THE SHADOW

The nightly says benching capitulation_long cost **$+\$368$ of missed winners**. The shadow twin over the same benched windows says it avoided **$-\$1,609$** . They cannot both be right, and the difference is the exit: the nightly reprices every blocked signal under a fixed scalp-2R, which flatters a flush-buyer that snaps back fast and then gives it all up; the shadow mirror runs the gate's actual exit logic to its actual conclusion. On a gate whose entire mechanism is "buy the flush and get paid on the snap", a 2R scalp on 20 signals is not a fair proxy. **Treat the nightly's per-gate verdicts as directional for momentum gates and unreliable for faders**, and price

fader benches off the shadow mirror instead. This is a real defect in the nightly, worth an hour on Saturday.

The exit selector, same week, same method

The router picks which gates are on; the selector picks which exit each trade gets. Both are “regime→config” policies, so they get judged the same way — against fixed alternatives, not in isolation.

DAY	TRADES	OPTIMAL PICKS	REGRET	SELECTOR	PERFECT HINDSIGHT	ALWAYS-SCALP	ALWAYS-WIDE
Mon 08-03	30	26/30 (87%)	\$357	+\$1,152	+\$1,509	+\$72	+\$1,152
Tue 08-04	16	8/16 (50%)	\$428	+\$1,246	+\$1,674	+\$584	+\$1,510
Wed 08-05	29	23/29 (79%)	\$353	-\$488	-\$134	-\$1,164	-\$1,310
Thu 08-06	12	10/12 (83%)	\$40	-\$294	-\$254	-\$254	-\$294
Fri 08-07	23	14/23 (61%)	\$358	+\$70	+\$427	-\$250	-\$370
WEEK	110	81/110 (74%)	\$1,536	+\$1,686	+\$3,222	-\$1,012	+\$688

Adapting the exit beat both fixed alternatives. +\$1,686 against -\$1,012 for always-scalp (**+\$2,698**) and +\$688 for always-wide (**+\$998**). It picked the best of three exits on 74% of trades and left \$1,536 of regret on the table — but the regret is measured against perfect hindsight, which is not a benchmark you can trade. The one systematic error is visible in the worst-pick lists: **it chose WIDE on longs that needed TIGHT or MID, repeatedly** (08-05 nipc_long chose wide -\$64 when tight made +\$22.50; 08-07 abs_veto_long chose wide -\$38 when mid made +\$35, twice). Wide was picked 70 times out of 110. That is a bias worth sweeping.

4 · Threshold tuning — and the answer is “leave the thresholds alone, fix the timing”

From `scripts/router_study.py --days 40`, extended with an isolated timing sweep and a grid-edge probe I ran on top of it. The fast replay is

asserted identical to the shipped `direction_router.replay_marks` on every day before anything is measured, so there is no drift between what is being swept and what is running. Both universes are reported side by side throughout: **LIVE** = the two gates the router actually manages as faders (`capitulation_long`, `rgv_short`); **FULL** = the historical fader set including `rgv_long` and `exhaustion_short`.

First the thresholds, because this is the short answer. Sweeping `ER_TREND` × `NET_MIN` at the shipped timing, the incumbent wins on both sets.

LIVE MANAGED SET — ROUTER Δ VS NO-ROUTER						
ER \ NET	20	30	40	50	60	
0.15		+\$461	+\$492 ★	+\$428	+\$316	+\$158
0.20		+\$224	+\$176	+\$152	+\$152	+\$152
0.25		+\$152	+\$152	+\$152	+\$152	+\$152
0.30		+\$152	+\$152	+\$152	+\$152	+\$152
0.35		+\$152	+\$152	+\$152	+\$152	+\$152

FULL HISTORICAL FADER SET — ROUTER Δ VS NO-ROUTER						
ER \ NET	20	30	40	50	60	
0.15		+\$1,300	+\$1,307 ★	+\$1,294	+\$1,120	+\$858
0.20		+\$1,044	+\$995	+\$976	+\$1,014	+\$816
0.25		+\$787	+\$787	+\$806	+\$806	+\$744
0.30		+\$334	+\$334	+\$334	+\$334	+\$428
0.35		+\$428	+\$428	+\$428	+\$428	+\$428

The shipped cell **ER 0.15 / net 30** is the maximum on both grids independently, and the surface around it is flat — net 20, 30 and 40 are within \$13 of each other on the full set. **There is nothing to gain here and no case for touching it.** Note also how badly ER 0.35 does (+\$152 on the live set means it essentially never declares a trend, so the router stops doing anything). That is a useful boundary: the router’s ER and grind’s 0.35 entry floor are different animals and must not be confused.

The timing sweep — this is where the money is

Thresholds pinned at shipped (ER 0.15 / net 30); only `WINDOW`, `STEP`, `HOLD` and `fast_exit` move. “Days +/-” counts sessions where the config helped

or hurt. “LOO worst” is leave-one-day-out: the total with the single most favourable day deleted — if that goes negative, the result is one day in a trenchcoat.

CONFIG	LIVE Δ	LIVE DAYS +/-	LIVE LOO WORST	FULL Δ	FULL DAYS +/-	FULL LOO WORST
SHIPPED w30 step15 hold2	+\$492	3 / 0	+\$228	+\$1,307	7 / 0	+\$802
w30 step5 hold2	-\$320	0 / 4	-\$320	+\$308	4 / 3	+\$17
w30 step5 hold3	+\$190	3 / 0	+\$38	+\$744	5 / 3	+\$310
w45 step5 hold3 ★	+\$994	4 / 0	+\$574	+\$1,706	6 / 1	+\$1,040
w45 step5 hold2	+\$838	5 / 0	+\$520	+\$1,430	6 / 1	+\$868
w45 step5 hold4	+\$633	3 / 0	+\$318	+\$1,314	5 / 2	+\$728
w45 step10 hold3	+\$756	3 / 0	+\$440	+\$1,526	5 / 1	+\$954
w45 step15 hold3	+\$400	2 / 0	+\$171	+\$894	3 / 2	+\$370
w60 step5 hold3	+\$832	4 / 0	+\$460	+\$1,285	6 / 1	+\$655
w45 step5 hold3 fast_exit	+\$994	4 / 0	+\$574	+\$1,706	6 / 1	+\$1,040
zero-lag w30 step1 hold1	+\$178	2 / 2	-\$75	+\$856	4 / 3	+\$350

Both sets pick the same winner, and it beats the incumbent by **+\$502 on the live set and +\$399 on the full set**. It survives leave-one-day-out on both (+\$574 and +\$1,040 with the best day deleted — it never goes negative). And fast_exit makes literally zero difference at this setting, so leave it off; it is a knob with no effect on this tape.

The anti-overfit test — is w45/step5/hold3 a peak or a cliff edge?

The backtest discipline says resist the grid edge, so I pushed every axis past where router_study stops and re-measured. A parameter you should trust sits on a plateau with an interior peak; a parameter you should not sits at the boundary of what was tested.

AXIS	LIVE SET Δ / FULL SET Δ — RECOMMENDED VALUE IN BOLD					SHAPE
WINDOW (min)	30 +190/+744	45 +994/+1706	60 +832/+1285	75 +316/+598	90 -39/+308	INTERIOR PEAK

STEP (min)	1	2	3	5	10	15	INTERIOR PEAK
	+364/+1124	+439/+1117	+764/+1454	+994/+1706	+756/+1526	+400/+894	
HOLD (marks)	2	3	4				INTERIOR PEAK
	+838/+1430	+994/+1706	+633/+1314				
ER_TREND	0.10	0.12	0.15	0.20	0.35		INTERIOR PEAK
	+778/+1464	+881/+1527	+994/+1706	+264/+874	−0/−112		
NET_MIN (pt)	10	20	30	40	50		FLAT PLATEAU
	+994/+1706	+994/+1706	+994/+1706	+829/+1658	+764/+1604		

This is as clean as a robustness check gets on this desk. Interior maximum on four axes out of five, and the fifth is a flat plateau — net 10, 20 and 30 give bit-identical results, so router_study’s joint-best answer of “net 20” is not an improvement, it is the grid edge of a region where the parameter does nothing. **Keep net at 30.** Both universes agree on the ranking of every single cell. The recommendation is therefore two numbers, and only two: WINDOW 30→45, HOLD 2→3. STEP is nominally 15 in the module but the live mechanism already ticks every 5 minutes, so that one is a code-comment fix, not a behaviour change.

⚠ THE HONEST CAVEAT, AND IT IS NOT SMALL

This is 40 days of **realised** trades, which is 601 rows across 17 trading days — and those trades are already partly hand-gated and partly router-gated, so it is not a clean A/B. It is also one summer RANGE regime: 62–73% of every day this week read CHOP. **A longer window and a hold of 3 marks makes the router stickier, which is exactly what you want in chop and exactly what will hurt in a real trend week**, because it will hold a stale trend read for 15 minutes instead of 10. Every number above is an argument for chop. Deploy it Saturday, then **re-validate after the first week with two or more clean-trend days in it** before treating it as settled. If the next trend week shows the router benching the aligned rider late, hold is the first thing to put back to 2.

The three live rules you asked me to check specifically

(a) The day-bias ± 40 rule — bench shorts when the day is up more than 40 points, bench longs when it is down more than 40. I tested every threshold from 0 to 300 on 599 clean realised trades over 30 days.

BIAS THRESHOLD	COUNTER-BIAS TRADES	THEIR NET	\$ / TRADE	ALIGNED TRADES	THEIR NET
± 0 pt		207	−\$3,534		−\$467
± 20 pt		194	−\$3,147		−\$854
± 40 pt (live)		186	−\$3,354	−\$18.03	−\$647
± 80 pt		166	−\$2,807		−\$1,194
± 150 pt		136	−\$2,522		−\$1,479
± 300 pt		83	−\$870		−\$3,131

There is no optimum to find here, and that is the answer. Counter-day-bias trades lose \$16–\$18 each at every threshold from zero to 150 points. The rule is not a tuned parameter, it is a directional fact: trading against the day loses money on this desk regardless of how far the day has gone. The ± 40 setting happens to isolate the worst per-trade cohort of the low thresholds, so **leave it at 40** — and note that raising it is actively harmful, because the “aligned” bucket gets steadily worse as you raise it (it starts absorbing counter-trend losers it should be excluding).

(b) The grind ER-0.35 floor — validated hard, and it also tells you what to do next.

GATE	ER < 0.15			ER 0.15-0.25			ER 0.25-0.35			ER $\geq$ 0.35		
grind_long	n=18	+\$118	+\$6.5/tr	n=29	−\$1,160	−\$40.0/tr	n=30	+\$20	+\$0.7/tr	n=21	+\$568	+\$27.1/tr
abs_veto_short	n=10	−\$218	−\$21.9/tr	n=11	−\$452	−\$41.1/tr	n=5	−\$80	−\$16.1/tr	n=9	+\$550	+\$61.2/tr
abs_veto_long	n=27	+\$80	+\$3.0/tr	n=29	+\$552	+\$19.0/tr	n=8	−\$162	−\$20.2/tr	n=29	−\$25	−\$0.9/tr
exhaustion_short	n=33	−\$144	−\$4.3/tr	n=10	−\$170	−\$17.0/tr	n=17	+\$468	+\$27.6/tr	n=15	−\$94	−\$6.3/tr
capitulation_long	n=15	−\$254	−\$17.0/tr	n=8	−\$16	−\$2.1/tr	n=3	+\$56		n=4	+\$177	

grind_long makes **+\$27.10 a trade above ER 0.35** and loses **−\$40.00 a trade in the 0.15-0.25 band**. The floor is exactly in the right place. Do not move it, do not let anyone argue it down on a “close enough” reading.

And now the actionable half of that table. `abs_veto_short` has the identical shape — **+\$61.20/trade above ER 0.35, negative in all three bands below it** — and it has no floor at all. That is this week's biggest single leak with a mechanism attached. Two of the week's four losing entries fired at ER30 of 0.17 and 0.18 for a combined $-\$274$; an ER-0.35 floor blocks both. The other two fired at 0.47 and 0.38, so the floor is not a complete fix — but it is a free one, it uses a threshold already proven on a sibling gate, and it fails safe. Note the mirror-image result on `abs_veto_long`: it earns in the 0.15–0.25 band and does nothing above 0.35. **The two halves of the same gate want opposite ER regimes.** Do not apply one floor to both.

(c) The re-arm test (ER climbing + vol expanding + structure) — the mechanism is sound and the failures are all in the measurement, not the rule. Three separate times this week the test was applied to a false vol number: 08-07 13:17 (ATR60 as denominator, which contains the spike), 13:50 (the tick's "ATR 39.1→27, spike over" when measured ATR was 42.1), and the 08-06 08:15/11:05 pair where ER climbed on ATR 11–12 and both re-arms decayed within 15 minutes. The (d2) codification — measure vol against the SESSION baseline, never a window containing the spike — is the right fix and it is already written. What is not written is the necessary-but-not-sufficient half: **vol expansion alone, in ER-0.08 chop with no range break, is a whipsaw trap**, and it has now fooled the desk at least three times. That deserves promotion from a note to a rule.

5 · The untradeable-day meter — it works, and it works too late

`src/gazbot7/untradeable.py`, shipped 07-31, computed here for every Paris day since 07-27 by calling `untradeable.compute` directly against `capture.db` and `gazbot7.db`. Three meters blend into one 0–100 score: **ROUNDTRIP** (how much of the day's range the day actually kept), **GIVE-BACK** (how much of its biggest excursion it handed back), and **STOPS** (the recent stop-rate, excluded below $n=10$). Score ≥ 65 = STAY-OUT, ≥ 45 = CAUTION.

DAY	SCORE	VERDICT	ROUNDTRIP	GIVE-BACK	STOPS METER	ATR14	DAY NET / RANGE	ACTUAL DESK P&L	RIGHT?
Fri 07-31	86	STAY-OUT	0.10	0.84	80	11.7	−61 / 611	−\$905.00	☐
Thu 08-06	78	STAY-OUT	0.14	0.81	58	4.1	−61 / 422	−\$233.00	☐
Tue 07-28	56	CAUTION	0.38	0.62	33	6.9	−230 / 610	−\$54.00	☐
Mon 07-27	47	CAUTION	0.52	0.40	60	5.6	−417 / 808	−\$226.50	☐
Wed 08-05	44	TRADEABLE	0.45	0.12	87	8.1	−234 / 522	−\$817.00	☐ miss
Mon 08-03	34	TRADEABLE	0.51	0.10	53	3.6	+316 / 624	−\$460.50	☐ miss
Wed 07-29	23	TRADEABLE	0.75	0.05	53	14.1	−724 / 961	+\$638.50	☐
Tue 08-04	22	TRADEABLE	0.77	0.15	33	10.4	+862 / 1115	+\$330.50	☐
Fri 08-07	15	TRADEABLE	0.84	0.05	33	3.4	+339 / 403	+\$227.00	☐
Thu 07-30	13	TRADEABLE	0.96	0.04	47	8.4	+1030 / 1074	+\$1,172.00	☐

★ THE RESULT — AND IT IS BETTER THAN THE VERDICT LABELS SUGGEST

Sort by score and the ordering is **perfect**. **Every green day scores 23 or below. Every red day scores 34 or above.** There is an eleven-point gap with nothing in it. On ten days the meter has never once put a profitable day above an unprofitable one — Spearman −0.76 against realised P&L. It flagged 07-31 (−\$905) at **86**, the highest score in the record, and it put the desk’s two best days at **13** and **15**. **The measurement is right.** What is wrong is the cutoff and the clock.

Which of the three meters actually discriminated? You had a prior on this and it is mostly right, with one correction.

SIGNAL	SPEARMAN VS DAY P&L	VERDICT	NOTE
ROUNDTRIP meter	−0.82	DISCRIMINATES	The strongest single input. Confirms the module’s own claim.
Blended SCORE	−0.76	DISCRIMINATES	Slightly worse than roundtrip alone — the blend is diluting its best input
STOPS meter	−0.72	BETTER THAN BILLED	Stronger than expected, but it is a lagging read of damage already done
GIVE-BACK meter	−0.66	DISCRIMINATES	Real, but weakest of the three — and it pulled the score down on both miss days (0.12 and 0.10)
Router CHOP-time %	−0.38	NOISE	Every day this week read 63-73% CHOP. It cannot separate what it cannot vary.
ATR14	+0.21	NOISE	

			Confirmed dead. The best day (07-30) and a bad day (07-31) sat at ATR 8.4 and 11.7.
Shadow net, all 46 sims	+0.73	COINCIDENT	Correction to the prior: it does track P&L — but only at the close, when it is useless

Two corrections worth writing down. **The stops meter is stronger than we credited it** (−0.72, essentially tied with the blend) — the 08-05 fix that excludes it below n=10 was still right, because n=4 setting a fifth of the score was a coin flip, but the meter itself is not weak. And **shadow-net is not a non-discriminator**; at +0.73 it is one of the better readings in the table. It just cannot be used, because a day’s shadow net is only known once the day has happened — which is exactly the problem the whole meter has.

The clock problem — when is the meter actually right?

The table above uses the end-of-day reading. But the router reads this meter during the session. So I recomputed it at five points in every day and asked, at each one, whether it separates green days from red days.

DAY	12:00Z	13:30Z	15:00Z	17:00Z	20:00Z	END OF DAY	ACTUAL P&L
Mon 07-27	78	66	35	25	44	47	−\$226.50
Tue 07-28	36	34	36	70	50	56	−\$54.00
Wed 07-29	92	82	36	44	20	23	+\$638.50
Thu 07-30	12	12	30	18	16	13	+\$1,172.00
Fri 07-31	44	42	92	92	84	86	−\$905.00
Mon 08-03	21	21	47	43	40	34	−\$460.50
Tue 08-04	14	28	18	11	15	22	+\$330.50
Wed 08-05	44	18	84	96	49	44	−\$817.00
Thu 08-06	52	22	74	86	74	78	−\$233.00
Fri 08-07	43	45	46	42	14	15	+\$227.00

READ TIME	SPEARMAN VS DAY P&L	WORST GREEN DAY SCORES	BEST RED DAY SCORES	CLEANLY SEPARATED?
12:00Z (pre-open)	−0.26	92	21	NO — NOISE
13:30Z (the cash bell)	+0.13	82	18	NO — INVERTED
15:00Z	−0.82	46	35	nearly
17:00Z	−0.72	44	25	nearly
20:00Z	−0.73	20	40	YES
End of day	−0.76	23	34	YES

Before the US open the meter is noise, and on one day it was actively dangerous. On 07-29 it read **92 at 12:00Z** and **82 at 13:30Z** — a screaming STAY-OUT — on a day that went on to make **+\$638.50**. At the cash bell the correlation is +0.13, which is worse than a coin. It only becomes reliable from 15:00Z and only fully separates by 20:00Z. That is because the inputs are cumulative session statistics: before the US open they are measuring eight hours of thin European tape, and a quiet overnight range looks identical to a dead day.

And that timing is what kills its actual value this week. Here is what a stay-flat rule would really have saved, counting only the P&L booked after the meter tripped.

DAY	SCORE @15:00Z	WOULD TRIP AT 65?	DAY P&L	P&L BOOKED AFTER 15:00Z	ACTUALLY SAVED
Mon 08-03	47	no	−\$460.50	\$0.00	\$0
Tue 08-04	18	no	+\$330.50	+\$256.50	—
Wed 08-05	84	yes	−\$817.00	\$0.00	\$0
Thu 08-06	74	yes	−\$233.00	\$0.00	\$0
Fri 08-07	46	no	+\$227.00	+\$268.50	—
Fri 07-31 (reference)	92	yes	−\$905.00	−\$282.00	+\$282

The meter correctly identified Wednesday and Thursday as stay-out days and it saved nothing, because by 15:00Z every dollar had already been lost. Look at where the week’s damage actually landed: −\$434 in the 13:00Z hour on Monday, −\$203 then −\$474 in the 13:00 and 14:00 hours on Wednesday, −\$194 in the 13:00 hour on Thursday. **The losses live in the two hours around the cash open, which is precisely the window in which the meter is a coin flip.** Across the whole record its actionable saving is \$282, all of it on 07-31.

Does the 65 cutoff want tuning? On the end-of-day score, yes — badly. It catches two of six red days.

STAY-OUT CUTOFF	DAYS FLAGGED	RED-DAY LOSS AVOIDED	GREEN-DAY PROFIT MISSED	NET	WHICH DAYS
≥ 30 (in-sample optimum)	6	\$2,696	\$0	+\$2,696	all six red days, zero green days
≥ 45 RECOMMENDED	4	\$1,418	\$0	+\$1,418	07-27, 07-28, 07-31, 08-06
≥ 55	3	\$1,192	\$0	+\$1,192	07-28, 07-31, 08-06
≥ 65 (live)	2	\$1,138	\$0	+\$1,138	07-31, 08-06
≥ 85	1	\$905	\$0	+\$905	07-31

No cutoff anywhere from 30 to 85 has ever flagged a green day. The in-sample optimum is 30, which would have caught all six red days and none of the green — but 30 sits four points above the highest green day in a ten-day sample, and that is curve-fitting to a gap I can see. **Move it to 45, not 30.** That catches four of six red days, still leaves 22 points of headroom above the worst green day, and is the change I would defend if next month's data disagrees. Note also that the CAUTION band at 45 and the STAY-OUT band at 65 currently do the same job on this data — every day in 45–65 was red. Collapsing STAY-OUT to 45 is as much a simplification as a tuning.

6 · Bench-decision review — grading the actual calls

DECISION	WHAT WAS ARMED / BENCHED	RESULT	GRADE	WHAT TO CHANGE
07-31 13:30 US-open enable-all	All six gates armed into a real up-thrust (ER 0.68, break 28713)	abs_veto +\$146 · grind – \$370 · exhaustion – \$72 = ~–\$300	CHURN	The lesson of the week. The veto-protected gate caught the move; the un-vetoed churner bought the rollover top. Arm the veto gates at the open, never the churners.
07-31 13:40→14:22 trim to flat	Escalating bench: grind → capit + rgv_short → exhaustion → abs_veto → fully flat	Shadow confirms every benched gate would have bled: grind_fast –\$153, thrust_loose –\$153, exhaustion_rev –\$99, abs_veto_55s –\$80, thrust_short_raw –\$153 — ~–\$500 avoided	GOOD ×4	Nothing. This is one-directional escalation as the tape genuinely worsened — the opposite of thrash.
07-31 19:08 fully flat	Accepted the no-trade day; zero further exposure	Zero further bleed. The exhaustion trial's one calm-regime fire stopped at –\$27.50 — chop paid nobody.	GOOD	Nothing. some-days-stay-out-fully-flat earned its keep.
07-31 16:41 & 17:37 re-enables	abs_veto re-armed twice into an ER 0.24–0.30 “calm” recovery, second time on scalp exits	–\$131 then –\$114 = –\$245	CHURN ×2	Even a veto-protected gate on scalp exits churns a choppy recovery. The 55s veto is a microstructure filter and a chop thrust passes it. Wait for ER≥0.35 sustained.
The grind benches (all week)	Benched 119.3 of 120 hours	grind_fast lost –\$2,549 over that window; grind's one live hour lost –\$172	GOOD	Make it explicit. The router has effectively retired this gate through 9 separate arm/bench decisions without ever saying so. Either write the retirement down or

			give it a hard ER-0.35 auto-arm.
abs_veto_short scalp-then-bench	Eight separate arm events across 08-05, 08-06 and 08-07 — three arm/bench cycles on Thursday alone	Four armings never fired (\$0 — luck, not judgement). Three fired: -\$559.00 on nine lots, zero winners	Add the ER-0.35 floor. Three arm/bench cycles on one gate in a day is the churn signature and its expected cost is not \$0, it is P(setup) × the loss.
The exhaustion trial (fade-scalp A@0.5R / B@1.5R)	Deployed 07-31 07:55, pinned out of the router	First fire came in ATR-62 open violence, off-regime: A +\$51.50, B -\$123.50 = - \$72 . Later, under the (c2) carve-out on 08-07: + \$104.00, 2-for-2 . Week total -\$162 on 6 lots.	The core thesis — fader-scalp in normal chop, ATR 11-18 — is still untested. The ER table says its habitat is ER 0.25-0.35 (+\$27.60/tr on n=17), not high or low ER. Gate the trial to that band and stop sampling it at the open.
The fully-flat stay-out calls	07-31 19:08, and the overnight low-vol grind holds	Recognised an untradeable trend — price ground +141pt overnight on ATR 13-16 and every one of the 8 earlier gate-fires had stopped out	Nothing. “Not every up-move is tradeable” is the correct and hard-won read.

THE BENCH ASYMMETRY, RESTATED WITH THIS WEEK’S NUMBERS

The ledger’s standing principle is that a wrongly-benched gate costs opportunity while a wrongly-armed gate costs real money, so when unsure, sit benched one more tick. **This week priced that principle exactly.** The over-benches cost **+\$368 + \$161 + \$106 = \$635** of missed winners across capitulation_long, exhaustion_short and rgv_short combined (and the fader half of that is probably illusory — see the nightly defect above). The wrong arms cost **-\$559.00** on abs_veto_short alone, in eleven hours. The asymmetry is real and it is roughly 1:1 in headline dollars but far worse in variance: the missed winners are spread over 100 signals, the losses arrived in three bursts.

7 • The durable tick and the architecture question

The mechanical gazbot7-direction-router.timer is switched off; the live mechanism is gazbot7-router-tick.timer running a Claude judgement every five minutes. So “did the durable router manage correctly” is really

two questions: did the plumbing hold, and is a poll the right shape for this job.

DAY	DURABLE TICKS	MEDIAN GAP	MAX GAP	GAPS > 10 MIN	COVERAGE
Mon 08-03	276	5.0 min	5.2 min	0	23 h
Tue 08-04	276	5.0 min	5.3 min	0	23 h
Wed 08-05	276	5.0 min	5.7 min	0	23 h
Thu 08-06	276	5.0 min	5.8 min	0	23 h
Fri 08-07	276	5.0 min	5.6 min	0	23 h

The plumbing is flawless. Exactly 276 ticks every day, median gap 5.0 minutes to the second, no gap anywhere in the week longer than 5.8 minutes, not one missed window. There is a consistent hole at **21:00-21:59Z** — zero ticks in that hour on all five days — which is the pre-reopen maintenance hour and is by design, but it means the desk is unsupervised for the hour immediately before the 22:00 all-on re-arm. Given that the reopen block carries 32% of the week’s switch changes and cost −\$275 on 07-31, **that gap is in the wrong place.**

Is polling right? I backtested the alternative.

The proposed alternative is an in-code, per-gate regime guard evaluated at signal time — no poll, no lag, no hysteresis, the gate checks the regime the instant it wants to fire. That is exactly step=1 minute, hold=1 mark. I ran it against the shipped poll and the 5-minute tick under the full live policy (DOWN_OFF / UP_OFF / CHOP_OFF plus the 13:15-14:00Z open-window carve-out) on this week’s realised trades and on 40 days.

MECHANISM	THIS WEEK: ROUTER VALUE	SIGNALS BLOCKED	40 DAYS: ROUTER VALUE	40D DAYS HELPED / HURT	40D LOO WORST
SHIPPED poll — 15 min, hold 2	+\$682	26	+\$475	9 / 4	−\$124
Claude tick — 5 min, hold 2	+\$456	24	−\$454	7 / 7	−\$1,052
Recommended — 5 min, hold 3, window 45	+\$402	22	+\$1,382	10 / 4	+\$783
	+\$304	14	+\$267	8 / 5	−\$295

The event-driven version is the worst option on both horizons. +\$304 this week against the poll’s +\$682; +\$267 over 40 days against +\$475, and it fails leave-one-day-out (−\$295 with the best day removed) where the recommended config never does. The reason is visible in the “signals blocked” column: the zero-lag guard blocks only 14 signals this week against the poll’s 26. An instantaneous regime read flickers — it declares a trend on a one-minute burst, un-benches the gate, and the gate fires into the burst that then fails. **The lag is not a bug, it is the filter.** Removing it removes the edge.

Quantifying the actuation lag directly

Two measures, because the mechanism question deserves better than one.

MEASURE	METHOD	RESULT
Observed reaction lag	For every bench decision this week, sum the P&L of trades that opened on that gate in the N minutes before the switch flipped — the money that leaked while the flip was detectable but not yet actioned	N=5 min: 0 trades, \$0.00 N=10 min: 8 trades, −\$301.50 N=15 min: 10 trades, −\$370.50 N=30 min: 12 trades, −\$439.50
Mechanism lag	Router value of a zero-lag guard minus router value of the poll, full live policy, this week	−\$378 — i.e. removing the lag costs \$378

Not one trade opened in the five minutes before a bench, all week. The tick cadence is not leaking anything. The −\$301.50 at ten minutes is the cost of “one tick earlier”, and it is made of four decisions: two abs_veto_long lots 8.3 minutes before the 08-02 22:10 bench (−\$180.50), two abs_veto_long lots 7.5 minutes before the 08-06 07:15 bench (+\$56 — that bench cost money, it did not save it), two exhaustion_short lots 7.2 minutes before the 08-07 13:38 narrowing (−\$125), and two capitulation_long lots 7.6 minutes before the 16:27 flip (−\$52). Stretch the window to fifteen minutes and one more pair appears (−\$69, 11.3 minutes before the 08-05 12:40 bench). Three of those benches were triggered by the stop-out that had just happened — the information did not exist a tick earlier. **The genuinely recoverable lag cost this week is close to zero, and the theoretical ceiling on it is about \$300.** That is not where to spend Saturday.

The regime split — because a blanket number is a bug

Every config above is a policy (regime→config), not a static setting, so it has to be scored on its home segments. Tape is segmented on ATR level, ER and range-break — not the clock — into dead-chop / normal-chop / in-between-building / clean-trend / violent-whipsaw, then separately by session block. 301 router-managed trades over 40 days.

REGIME SEGMENT	N	NET	WIN %	\$ / TRADE	SHIPPED BLOCKED (N / \$)	RECOMMEN DED BLOCKED (N / \$)
violent-whipsaw	19	-\$960	16%	-\$50.50	8 / -\$243	8 / -\$322
normal-chop	110	-\$559	35%	-\$5.10	36 / -\$114	61 / -\$382
in-between- building	110	-\$328	39%	-\$3.00	40 / +\$346	49 / -\$268
clean-trend	55	+\$1,119	49%	+\$20.30	15 / -\$246	19 / -\$182
dead-chop	3	+\$26	33%	+\$8.50	0 / \$0	1 / +\$86

This table explains where the recommended config’s +\$900 comes from, and it is one segment. In in-between-building tape — ER 0.20 to 0.35, the ambiguous zone — the shipped config blocks +\$346 of winners, while the longer window and stickier hold blocks -\$268 of losers instead. That single flip is a \$614 swing and it is the whole story. The shipped 30-minute window is too short to tell a building trend from a chop bounce, so it declares a trend, benches the fader, and the “trend” fails. A 45-minute window with three agreeing marks does not make that mistake.

Note also the segment nobody has a rule for: **violent-whipsaw is -\$50.50 a trade at a 16% win rate**, and neither config blocks more than 8 of its 19 trades. That is the most expensive tape on the desk per trade and the router barely touches it.

SESSION BLOCK	N	NET	WIN %	\$ / TRADE	SHIPPED BLOCKED	RECOMMEN DED BLOCKED
Overnight 21:00- 06:00Z	89	-\$906	31%	-\$10.20	40 / +\$64	57 / +\$21
EU morning 06:00-13:30Z	62	+\$28	32%	+\$0.40	23 / -\$160	27 / -\$485

The overnight block is a \$906 hole over 40 days and the router is not helping there — both configs show it blocking winners overnight (+\$64 and +\$21). This week it was worse: 8 overnight trades, **zero winners, -\$578**. That is the reopen churn showing up in the trade record rather than the switch log. The US session, by contrast, is where the recommended config does its work — 58 blocks worth -\$821 of avoided losses against the shipped 40 blocks worth -\$378.

8 • Verdict, and the Saturday list

IS THE ROUTER EARNING ITS KEEP?

Yes, and by the largest margin it ever has. +\$2,951 of net value on its own accounting, up from -\$1,216 the week before. Zero leakage on every day. Two decisions — keep capitulation_long off and keep grind_long off — held for essentially the whole week with almost no fiddling, and the shadow twins of those two gates lost **\$4,158** over exactly those windows. The desk still lost \$953 this week, but it lost it through gates the router had armed, not through gates it failed to bench. Is it over-benching? Marginally, on the faders (+\$635 of missed winners across three gates), and the fader half of that number is probably an artefact of the nightly's scalp-2R repricing. Where is it leaking? **One gate and one time-block: abs_veto_short, and the overnight/reopen window.**

#	SATURDAY ACTION	EVIDENCE	RISK GATE
1	Add an ER30 ≥ 0.35 arming floor to abs_veto_short, mirroring grind's. Do not apply it to abs_veto_long.	+\$61.20/trade above 0.35 (n=9) vs negative in all three bands below (n=26). Blocks 2 of the week's 4 losing entries, -\$274. Long twin has the opposite profile: +\$19.00/tr in the 0.15-0.25 band, -\$0.90 above 0.35.	Fails safe — a floor can only reduce arming. Review after 15 fires.
2	Router timing: WINDOW 30-45, HOLD 2-3. Leave ER_TREND at 0.15 and NET_MIN at 30. Correct STEP to 5 to match the live tick.	+\$502 on the live set, +\$399 on the full set. Interior peak on window, step, hold and ER; flat plateau on net. Survives LOO on both sets (+\$574 / +\$1,040).	Re-validate after the first week containing two or more clean-trend days. A stickier router is a chop optimisation. If it starts benching the aligned rider late, put hold back to 2 first.

		Both universes rank every cell identically.	
3	Untradeable meter: STAY-OUT cutoff 65→45, and add an explicit rule that readings before 15:00Z do not trigger a stay-out.	At 45: catches 4 of 6 red days, 0 of 4 green, +\$1,418 vs the current +\$1,138 — and 22 points of headroom above the worst green day. Pre-open the meter correlates +0.13 with the day (noise) and read 92 on a + \$638 day.	In-sample on n=10. The 45 choice deliberately declines the in-sample optimum of 30.
4	Fix the midnight reopen. Re-arm reversion gates only; make momentum gates earn their arm. And move the 21:00-22:00Z maintenance gap so the hour before the reopen is supervised.	32% of the week's switch changes are spent undoing the 22:00 all-on. It cost -\$275 in one night on 07-31. Overnight block: 8 trades this week, 0 winners, -\$578; 89 trades over 40 days, -\$906.	Reversible — it is a change to reactivate_gates.py's default set.
5	Fix the nightly's fader repricing. Stop scoring blocked fader signals under a fixed scalp-2R; use the gate's own exit.	The nightly says benching capitulation_long cost +\$368 of winners; the shadow mirror says it avoided -\$1,609. Both cannot be right and the fixed-scalp assumption is the suspect.	Diagnostic only — no live behaviour changes.
6	Write down grind's retirement, or automate its arm. It was benched 119.3 of 120 hours through nine separate decisions.	grind_fast lost -\$2,549 over the benched window; grind's single live hour lost -\$172. The router has already decided; it just costs a decision every time.	Either direction is safe. Prefer the automated ER-0.35 auto-arm so it can still catch a real trend.
7	Promote the vol-expansion caveat to a rule: expanding ATR alone, in ER<0.15 chop with no range break, is not a re-arm trigger.	Fooled the desk at least three times: 08-06 08:15 and 11:05 (both decayed inside 15 min), 08-07 13:50 (armed on an ATR reading that was wrong by 15 points).	Rule text only.
8	Grade the bad days. The ledger has zero entries for 08-03, 08-04 and 08-05 — 30 switch changes and -\$947 ungraded.	The scoreboard currently samples the well-supervised days, which biases it favourably.	Discipline, not code.

WHAT I WOULD NOT DO, AND WHY

Do not build the event-driven per-gate guard. Backtested above: worse on this week (+\$304 vs +\$682), worse over 40 days (+\$267 vs +\$475), and it fails leave-one-day-out. **Do not chase the actuation lag.** Zero trades leaked in the 5-minute window all week and the theoretical ceiling is ~\$300. **Do not touch ER_TREND or NET_MIN.** Both grids independently pick the incumbent cell and the surface around it is flat. **Do not move the day-bias ±40.** There is no optimum — counter-day-bias trades lose \$16-18 each at every threshold from 0 to 150, and raising it makes the aligned book worse.

The sample caveats, stated once and plainly. The threshold and timing work rests on 601 realised trades across 17 trading days, already partly hand-gated and partly router-gated — it is not a clean A/B and it is a lower bound, not a

measurement. The meter work is ten days. The bench grades are one week, in-sample, and drawn from a ledger that skipped three of this week's five sessions. Every regime bucket except normal-chop and in-between-building has n under 60. **What survived everything I threw at it: the router's two long-held benches, the timing change (interior peak on four axes, both universes, both leave-one-out tests), the grind ER-0.35 floor, and the meter's end-of-day ranking.** Everything else in this section is a lead, not a conclusion.

DR The day rider — how do you ride a move for as long as it lasts, and then stop?

Gaz, you framed this one yourself and the framing is right: "today was never a good trade. but future ones could be good for 3 hours and then the tape changes. so maybe we have mechanisms to ride them for as long as we want and then take profit." So I stopped asking whether the day-rider works and went looking for the mechanism. I rebuilt the whole thing from the raw tape over **38 sessions** (15 June → 7 August) and scored **114 different ways of ending a ride** on identical entries. Here is the short version, money first. **Not one exit rule I tested touches the money.** The seven days that lost a thousand dollars or more lost **−\$11,195** between them and the trail we deployed last Friday saves **exactly \$0.00** on every single one of them, because on those days price never went your way at all — there was nothing to trail. What does move the number is the **entry**: raise the efficiency floor from 0.15 to 0.25 and the same 38 sessions go from **+\$5,410 to +\$8,819**, per-trade from **+\$142 to +\$276**, and it survives every robustness test I threw at it. That is the one change I would actually make.

THE FOUR THINGS THAT SURVIVED, AND THE THREE THAT DIED

SURVIVED — (1) **Stops are still impossible.** Every fixed stop from 75 to 400 points is worse than no stop at all, and the sample now contains a genuine disaster. Your old finding holds. (2) **Raising the entry efficiency**

floor to ~0.25 — plateau across 0.18–0.35, positive in 32 of 32 leave-one-day-out folds, and consistent in and out of sample. (3) **Give-back-from-peak** halves the tail for about \$1,000 of the total — the only mechanism that changed the risk shape without breaking. (4) **The direction call is the whole strategy**: when the 13:40 lean holds to the close the desk makes **+\$660 a trade at 89% green**; when it reverses it loses **−\$1,308 a trade at 0% green**.

DIED — (1) **“The exit is exit-proof, holding wins”** is half wrong: holding beats every stop, but its whole edge is three days — strip them and \$5,410 becomes **\$379**. (2) **The rationale we shipped last Friday for the ATR trail is factually backwards** — $4 \times \text{ATR}$ is a harder bar than the old +150pt on 30 of 38 days and rescues exactly zero extra days. (3) **The exit parameter surface is noise**. Rank a rule on the first 19 sessions and its rank on the last 19 is negatively correlated (Spearman -0.30).

First, a correction that makes everything below trustworthy

FIXED

Last Friday’s note carried an honest warning: “the lab’s entries are systematically earlier than live — the lab detected 14:04 @29,619 where the live service detected 14:12 @29,567.” That gap has now been closed and it was a lab bug, not a real effect. The live service ticks at 21 seconds past each minute, and the newest 5-second bar it can possibly see is the one starting at :15. My old harness let it see the :20 bar too — five seconds of the future, every minute, for seven hours.

I calibrated the cutoff against the live journal instead of guessing it. At $m: +15s$ the lab reproduces the live drift read **exactly — all 34 minutes from 13:39 to the 14:12 detection match live’s efficiency and roundtrip to three decimals, including the detection minute itself**. Every other offset ($-10, -5, 0, +5, +10, +20, +25$ seconds) matched at most 3 of 34. So the numbers in this section are what the live service would actually have done, not an optimistic twin.

LAB CUTOFF TESTED	MINUTES MATCHING LIVE EXACTLY	VERDICT
m-10s / m-5s / m+0s	2 / 0 / 0 of 34	wrong
m+5s / m+10s	1 / 3 of 34	wrong
m+15s	34 of 34	this is what live does
m+20s (the old harness) / m+25s	2 / 0 of 34	saw 5s of the future

The proof it works: on 7 August the corrected lab fires at **minute 42 (14:12), DOWN, efficiency 0.183, roundtrip 0.638, ATR 42.9, at 29,566.75**. The live service fired at **14:12:22, DOWN, efficiency 0.183, roundtrip 0.638, entry ATR 42.89, filled 29,567.25**. Half a point apart on the fill, identical everywhere else.

1. The honest ledger — everything the day rider has ever actually done

Two rows. That is the entire recorded history, and one of them is a bug.

DATE	WHAT IT WAS	ROWS IN THE TRADE RECORD	REALISED	IS IT A TRADE?
2026-08-06 07:08	the cross-desk flatten — the day-rider's hard-flat branch sold the tournament's two lots twenty seconds after they opened	2 · gate day_rider_crossdesk	-\$95.50	NO — a defect
2026-08-07 14:12	SHORT 2 lots @ 29,567.25, flat on the 20:40 clock	0 — never written	≈ -\$1,014	yes, but unrecorded
Total in pnl.realized(des k="day_rider")		2 rows	-\$95.50	—

⚠ THE MEASUREMENT GAP IS STILL OPEN — AND IT IS THE TOP BUILD ITEM

day_rider.py contains **zero** references to record_trade or the store. It has never written a trade row and still cannot. I checked the fills table too: between 14:00 and 21:00 on 7 August there are 28 fills and **not one of them is the day-rider's** — no SELL 2 at 14:12, no BUY 2 at 20:40. It trades on clientId 4 and the store only ingests clientId 0.

So pnl.realized(desk="day_rider") correctly returns **-\$95.50 on n=2**, and **every one of those dollars is the bug, none of it is the strategy**. Do

not read that as a result. Until there is a write path, the day-rider can be remembered but not graded.

Where the $-\$1,014$ comes from, since it is not in the record: I reconstructed it from the service's own minute-by-minute journal, which logged 419 state stamps between entry and the flat. Entry 29,567.25; the flatten fired at **20:40:23** with the position **-252.8 points** offside; MNQ is \$2 a point a contract (confirmed against trade 569: 20.5 points cost \$41.00 plus \$1.50 fee) so $252.8 \times \$2 \times 2 \text{ lots} = -\$1,011.20$, plus **$\$1.50$ per contract round trip** = **$-\$1,014.20$** . My tape model of the same trade lands at $-\$1,018$ — within four dollars, which is the entry half-point.

The Friday trade, minute by minute

It ran exactly as designed end to end. The detector rejected the setup for 33 consecutive minutes — efficiency oscillating 0.137, 0.079, 0.119, 0.140, 0.130 — then confirmed at 14:12 on 0.183 / 0.638 and sold two lots. It was 18.8 points offside within a minute, 64 points offside within four, and **never traded back below the entry for the remaining six and a half hours**. The best it ever showed was **+2 points**. The trail needed +150 to arm. It rode to the clock and flattened.

Friday 2026-08-07 — the day-rider's only live trade: SHORT 2 lots, held 6h28m, $-\$1,014$

It sold at 14:12 and price never traded below the entry again. The trail needed 172 points in its favour to arm; the best it ever got was 2.



THE 7 AUGUST TRADE	VALUE	VERIFIED AGAINST
detected	14:12:22Z · efficiency 0.183 · roundtrip 0.638	service journal + tape replay, identical
rejected for	33 minutes before confirming	82 logged drift reads
entered	SHORT 2 @ 29,567.25 · entry ATR 42.89 · venue stop 30,167.25	journal
best it ever got	+2 points	tape
worst it ever got	−294 points	tape
trail	never armed — needed +150pt (or +172pt under the new ATR rule)	journal: 388 consecutive “not armed” ticks
exited	20:40:23Z hard flat · −252.8pt · −\$1,014	journal + venue net went −2 → 0

AND 6 AUGUST WAS NEVER TRADED — THE CROSS-DESK BUG ATE THE WHOLE SESSION

The scope says the day-rider’s first tradeable session was Thursday 6 August. It never got one. When the hard-flat branch sold the tournament’s lots at 07:08 that morning it set `closed = true` on the day-rider’s own state for session 2026-08-06 — and the entry block reads `if entered or closed: no re-entry`. From 13:30 onward the service logged “**already traded this session — no re-entry**” every single minute. It sat out a day it would have made money on: the corrected replay says 6 August detected UP at 13:41 @29,445.50 and would have been **+\$178 held to the clock, +\$494 on the trail**. So the bug cost roughly \$180–\$500 of opportunity on top of the \$95.50 it booked.

DR2 The exit problem, quantified — and an attack on “holding wins”

The old study said the exit was exit-proof: 25 variants, holding beat every one. It was derived on 31 sessions **with no big loser in them**. The sample now has one, so I re-derived the whole thing from scratch on 38 sessions with the corrected entry timing. The verdict is split. **Against stops, holding still wins outright and the finding survives.** Against trailing rules it does not — but neither does anything else, because **the exit parameter surface turns out to be noise**.

All 38 sessions, from the tape — the counterfactual ledger

These are not booked trades. This is the corrected replay: what the live service would have detected and done, session by session, 15 June to 7 August. **best** and **worst** are the furthest the trade ever went in each direction while it was open. Green rows made money holding to the clock; red rows lost.

SESSION	HOW THE DAY ENDED UP	SIDE	EFF AT ENTRY	ATR	BEST (PT)	WORST (PT)	FINAL (PT)	\$ HOLD	\$ ATR TRAIL	LEAN HELD?
06-15	CLEAN TREND	UP	0.29	33	+175	-74	+101	+\$400	+\$429	YES
06-16	CLEAN TREND	UP	0.17	38	+18	-598	-561	-\$2,248	-\$2,248	reversed
06-17	BUILDING TREND	DOWN	0.42	36	+487	-136	+262	+\$1,047	+\$1,056	YES
06-18	BUILDING TREND	DOWN	0.54	47	+59	-332	-278	-\$1,113	-\$1,113	reversed
06-22	BUILDING TREND	DOWN	0.26	46	+286	-178	+135	+\$537	+\$609	YES
06-23	DEAD CHOP	UP	0.34	66	+282	-188	-37	-\$150	+\$596	YES
06-24	BUILDING TREND	UP	0.32	59	+288	-490	+286	+\$1,141	+\$1,141	YES
06-25	BUILDING TREND	DOWN	0.53	40	+722	-10	+242	+\$966	+\$1,639	YES
06-26	VIOLENT WHIPSAW	DOWN	0.24	67	+16	-477	-49	-\$198	-\$198	YES
06-29	BUILDING TREND	UP	0.35	57	+331	-464	+265	+\$1,057	+\$1,057	YES
06-30	CLEAN TREND	UP	0.46	52	+387	-60	+281	+\$1,122	+\$496	YES
07-01	BUILDING TREND	DOWN	0.81	55	+90	-259	+49	+\$194	+\$194	YES
07-02	BUILDING TREND	UP	0.33	54	+116	-848	-624	-\$2,501	-\$2,501	reversed
07-06	DEAD CHOP	UP	0.20	54	+145	-67	+7	+\$25	+\$25	YES
07-07	BUILDING TREND	DOWN	0.49	51	+313	-84	+104	+\$412	+\$839	YES
07-08	DEAD CHOP	UP	0.47	53	+141	-325	+19	+\$73	+\$73	YES
07-09	CLEAN TREND	UP	0.96	38	+139	-240	+88	+\$349	+\$349	YES

07-10	BUILDING TREND	UP	0.20	38	+119	-279	+100	+\$399	+\$399	YES
07-13	CLEAN TREND	DOWN	0.38	42	+252	-114	+202	+\$807	+\$372	YES
07-14	DEAD CHOP	DOWN	0.17	35	+72	-256	-157	-\$631	-\$631	reversed
07-15	BUILDING TREND	DOWN	0.46	38	+484	-11	+167	+\$666	+\$669	YES
07-16	CLEAN TREND	DOWN	0.57	44	+325	-105	+214	+\$852	+\$405	YES
07-17	VIOLENT WHIPSAW	DOWN	0.70	45	+64	-574	-284	-\$1,137	-\$1,137	reversed
07-20	CLEAN TREND	UP	0.53	28	+60	-383	-356	-\$1,429	-\$1,429	reversed
07-21	BUILDING TREND	DOWN	0.24	36	+63	-238	-174	-\$698	-\$698	reversed
07-22	DEAD CHOP	UP	0.26	31	+168	-55	+23	+\$90	+\$423	YES
07-23	DEAD CHOP	UP	0.62	38	+91	-360	-29	-\$119	-\$119	YES
07-24	CLEAN TREND	DOWN	0.28	38	+238	-181	+157	+\$626	+\$301	YES
07-27	BUILDING TREND	DOWN	0.26	46	+569	-18	+288	+\$1,150	+\$1,903	YES
07-28	VIOLENT WHIPSAW	DOWN	0.61	44	+207	-269	-137	-\$551	+\$472	reversed
07-29	CLEAN TREND	DOWN	0.24	44	+533	-366	+408	+\$1,627	+\$424	YES
07-30	CLEAN TREND	UP	0.37	47	+486	-3	+412	+\$1,647	+\$660	YES
07-31	VIOLENT WHIPSAW	DOWN	0.36	50	+393	-52	+82	+\$327	+\$1,170	YES
08-03	CLEAN TREND	UP	0.63	53	+482	-68	+440	+\$1,757	+\$1,757	YES
08-04	CLEAN TREND	UP	0.55	34	+548	-24	+376	+\$1,503	+\$1,687	YES
08-05	CLEAN TREND	UP	0.20	52	+54	-489	-436	-\$1,749	-\$1,749	reversed
08-06	BUILDING TREND	UP	0.36	48	+241	-33	+45	+\$178	+\$494	YES
08-07	DEAD CHOP	DOWN	0.18	43	+2	-294	-254	-\$1,018	-\$1,018	reversed
38 sessions	13 clean · 14 building · 4 whipsaw · 7 dead	19 UP / 19 DOWN	—	—	—	—	—	+\$5,410	+\$6,797	28 of 38

What a no-stop multi-hour hold actually looks like — the full distribution

38 sessions, 2 lots, no stop, flat on the 20:40 clock, fees at \$1.50 a round trip a contract. This is the honest shape of the thing you are proposing to run.

THE HOLD, OVER 38 SESSIONS	VALUE
Net	+\$5,410
Per trade	+\$142
Days green	25 of 38 (66%)
Median day	+\$338
Upper quartile	+\$966
Lower quartile	-\$551
Worst decile (mean of the worst 4 days)	-\$1,982
Worst single day (2 July)	-\$2,501
Best single day (3 Aug)	+\$1,757
Top 3 days as a share of the whole result	\$5,031 of \$5,410 — 93%
Average hold	6h 57m
Standard deviation, per day	\$1,059

⚠ THE SURVIVORSHIP SHAPE YOU ASKED ME TO LOOK FOR — IT IS THERE

Strip the three best days and +\$5,410 becomes +\$379. Three sessions out of thirty-eight carry ninety-three percent of the result. That is not a strategy with an edge you can lean on; that is a lottery ticket that has paid three times. The median day is a perfectly respectable +\$338, but the mean is dragged around entirely by the tails, and **the bad tail is bigger than the good one**: worst decile -\$1,982 against a best decile of +\$1,634.

The old “exit-proof” conclusion was not wrong about stops. It was wrong to conclude anything at all from a sample whose whole signal lives in three days.

Stops: the old finding SURVIVES, and now it survives with a disaster in the sample CONFIRMED

This is the cleanest result in the section. Every fixed stop from 75 to 400 points is worse than running with no stop at all, and three of them turn a profitable strategy into a losing one. The 600-point venue stop is essentially free, which is exactly what it was designed to be.

STOP WIDTH	NET	\$ / TRADE	DAYS GREEN	WORST DAY	STRIP BEST 3	WHAT IT DID ON THE 13 LOSING DAYS
no stop (hold)	+\$5,410	+\$142	25/38	-\$2,501	+\$379	-\$13,542
75 pt	+\$1,953	+\$51	12/38	-\$303	-\$2,954	-\$3,939
100 pt	+\$168	+\$4	13/38	-\$403	-\$4,739	-\$5,239
150 pt	-\$317	-\$8	16/38	-\$603	-\$5,224	-\$7,839
200 pt	-\$1,295	-\$34	18/38	-\$803	-\$6,202	-\$9,786
250 pt	-\$3,438	-\$90	19/38	-\$1,003	-\$8,345	-\$11,881
300 pt	-\$2,430	-\$64	21/38	-\$1,203	-\$7,337	-\$12,672
400 pt	-\$176	-\$5	23/38	-\$1,603	-\$5,207	-\$13,724
600 pt (the live venue stop)	+\$5,508	+\$145	25/38	-\$2,403	+\$477	-\$13,444

Read the last column. A 250-point stop cuts the loss on the bad days from – \$13,542 to –\$11,881 — it saves \$1,661 — and it costs \$8,848 of the good days to do it. **Stops on this strategy are a straight transfer from the winners to nothing.** The reason is in the ledger: the winning sessions routinely draw 250–490 points against you first. On 24 June the trade ended +286 points after being **490 points offside**. On 29 July it ended +408 after being 366 offside. Any stop tight enough to help the losers kills those.

Now the uncomfortable part: the trail we deployed last Friday does not do what we said it does

RATIONALE WRONG

The argument shipped in the code reads: “a FIXED +150pt arm is a threshold a losing trade can never reach... at 4×ATR it arms when the DAY says the move is real.” I checked it against the tape and it is backwards.

CLAIM MADE LAST FRIDAY	WHAT THE TAPE SAYS	VERDICT
4×ATR arms in cases where +150pt cannot	4×ATR is a harder bar on 30 of 38 days — median arm distance 178 points vs a flat 150	backwards
It rescues the losing trades	Both rules fail to arm on exactly the same 16 days — the two never-armed sets are identical	rescues nothing
It would have helped on 7 August	4 × 42.89 = 172 points needed; the trade’s best moment was +2 points . Result under the new rule: −\$1,018, identical	no difference
It beats holding (\$7,341 vs \$6,544)	Re-derived with correct entry timing: \$6,797 vs \$5,410 . Still ahead, but see below	direction right, story wrong

Where the trail’s +\$1,387 over holding actually comes from. On the 13 losing days it improves things by \$1,768 — but **on 11 of those 13 days it changes nothing at all, exactly \$0.00**, because it never arms. The entire improvement is two days: 23 June (+\$746) and 28 July (+\$1,023). Neither was a big loser. On the **seven** days that lost a thousand dollars or more — 2 July, 16 June, 5 Aug, 20 July, 17 July, 18 June, 7 Aug, **−\$11,195 between them** — the trail saves **\$0.00. Not a cent, on any of them.**

LOSING SESSION	BEST IT EVER GOT	POINTS NEEDED TO ARM THE TRAIL	\$ HOLD	\$ ATR TRAIL	TRAIL SAVED
02 Jul	+116	216	−\$2,501	−\$2,501	\$0
16 Jun	+18	151	−\$2,248	−\$2,248	\$0
05 Aug	+54	208	−\$1,749	−\$1,749	\$0
20 Jul	+60	111	−\$1,429	−\$1,429	\$0
17 Jul	+64	179	−\$1,137	−\$1,137	\$0
18 Jun	+59	189	−\$1,113	−\$1,113	\$0
07 Aug	+2	172	−\$1,018	−\$1,018	\$0
21 Jul	+63	146	−\$698	−\$698	\$0
14 Jul	+72	140	−\$631	−\$631	\$0
28 Jul	+207	177	−\$551	+\$472	+\$1,023
26 Jun	+16	270	−\$198	−\$198	\$0

23 Jun	+282	264	−\$150	+\$596	+\$746
23 Jul	+91	153	−\$119	−\$119	\$0

This is the single most important table in the section, so let me say what it means in plain English. **A trailing stop is a mechanism for keeping a profit you already have. On the days this strategy loses money, it never has a profit to keep.** No amount of tuning the trail — not the arm distance, not the trail distance, not switching from points to ATR — can touch those days, because on those days there is nothing to trail. That is why the answer to “how do we stop losing money on days like 7 August” is not in the exit at all.

DR3 Your idea, backtested — four ways to end a ride on evidence instead of a clock

You asked for mechanisms that stop the trade because the tape changed, not because a clock struck. I built and scored four families over every session: an **ATR-scaled trail**, a **regime-change exit** (the day’s own efficiency decaying back below the level that got us in), **give-back-from-peak**, and **time-in-adverse-excursion** including a single dated checkpoint. 114 rules in all. Below is the honest scoreboard, and then the part that matters more than any single row in it.

The scoreboard — best of each family, every session, identical entries

wdec is the worst decile: the average of the four worst days, which is the number that tells you what a bad patch feels like. **sb3** is the total with the three best days deleted — the concentration test. **H1/H2** splits the sample in half by date; a rule that only works in one half is fitting.

NET	MEDIAN	WDEC	SB3	H1	H2
-----	--------	------	-----	----	----

MECHANISM	\$ / TRADE	DAYS GREEN	WORST DAY						
HOLD to the 20:40 clock (baseline)	+\$5,410	+\$142	25/38	+\$338	−\$2,501	−\$2,166	+\$379	+\$2,319	+\$3,091
ATR trail, arm 4× trail 2× (LIVE NOW)	+\$6,797	+\$179	27/38	+\$414	−\$2,501	−\$2,166	+\$1,450	+\$3,215	+\$3,583
the OLD fixed trail, arm +150 trail 100	+\$5,005	+\$132	27/38	+\$302	−\$2,501	−\$2,166	−\$416	+\$2,012	+\$2,993
give back 33% of peak (arms at 1×ATR)	+\$4,413	+\$116	35/38	+\$243	−\$2,248	−\$1,155	+\$1,661	+\$2,409	+\$2,003
regime exit: efficiency < 0.15 or the lean flips	+\$5,508	+\$145	20/38	+\$12	−\$1,110	−\$943	+\$663	−\$503	+\$6,011
adverse: offside >3×ATR for 60 min	+\$4,078	+\$107	24/38	+\$260	−\$1,775	−\$1,427	−\$953	+\$1,110	+\$2,968
checkpoint: at 15:00Z cut if worse than −150pt	+\$8,571	+\$226	25/38	+\$338	−\$1,399	−\$1,183	+\$3,540	+\$4,152	+\$4,419
... same checkpoint 30 minutes later (15:30Z)	+\$2,988	+\$79	23/38	+\$338	−\$1,959	−\$1,616	−\$2,043	+\$679	+\$2,309

⚠ READ THE LAST TWO ROWS TOGETHER — THAT IS THE WHOLE LESSON

The 15:00 checkpoint is the best-looking rule in the study: **+\$8,571**, tail cut nearly in half, both halves positive, survives strip-best-3 at **+\$3,540**. Move it **thirty minutes later** and it collapses to **+\$2,988**. Move the threshold from −150 to −100 and it falls to **+\$6,566**; to −200 and it falls to **+\$7,335**. At 16:00 it is **+\$2,556**, at 17:00 **+\$2,156**, at 18:00 **+\$1,625**.

That is not a mechanism, that is a coincidence with a clock on it. A rule whose value drops by two thirds when you nudge it half an hour is fitted to which specific days happened to be underwater at 15:00 in this particular sample. I am reporting it because it looks great and you would rightly ask

why I did not ship it — and the answer is that it fails the plateau test outright. **REFUTED AS TUNED**

The plateau test, applied to the trail we are actually running

Same test, applied to the ATR trail. Every cell is the same 38 entries, only the arm and trail multiples change. If the rule captured something real, neighbouring cells should look like each other.

ARM ↓ / TRAIL →	0.75×	1.0×	1.5×	2.0×	3.0×
1.0×	+\$3,687	+\$2,946	+\$2,768	+\$7,085	+\$5,261
1.5×	+\$1,253	+\$780	+\$665	+\$4,661	+\$3,510
2.0×	+\$2,595	+\$2,161	+\$2,161	+\$5,759	+\$4,814
2.5×	−\$139	−\$921	−\$1,793	+\$1,200	+\$772
3.0×	+\$277	−\$72	−\$341	+\$3,693	+\$2,206
4.0×	+\$3,909	+\$4,140	+\$3,644	+\$6,797	+\$5,875
5.0×	+\$4,022	+\$4,002	+\$4,687	+\$6,444	+\$5,359
6.0×	+\$5,675	+\$5,060	+\$6,405	+\$7,244	+\$5,825

There is no plateau anywhere in that grid. Arm at 2.5× and three of the five trail widths lose money and none makes more than \$1,200; arm at 4× — one and a half ATR away, about 65 points on a typical day — and every single width makes \$3,600 or more. Walk straight down the 1.5× column and you get +\$2,768, +\$665, +\$2,161, −\$1,793, −\$341, then **+\$3,644** in the very next row. A response surface does not do that. A random field does.

⚠️ THE TEST THAT SETTLES IT: RANK THE RULES ON THE FIRST HALF, CHECK THEM ON THE SECOND

I ranked all **114** rules by their net on the first 19 sessions (15 June → 13 July), then looked at how they did on the last 19 (14 July → 7 August). If exit tuning captured anything real, the good rules would stay good.

TEST

Best rule on the first 19 sessions

RESULT

ATR trail arm 6× / trail 0.75× — +\$3,734 in-sample

What it then did on the last 19	+\$1,941 — worse than simply holding (+\$3,091)
Its rank out of sample	71st of 114
Rank correlation, in-sample vs out-of-sample, all 114 rules	Spearman −0.30 — NEGATIVE

Picking the best exit rule on past data is, on this sample, slightly worse than picking one at random. That is the finding. It applies to the $4\times/2\times$ trail we deployed too — it happens to have held up across the split (+\$3,215 then +\$3,583), but it was chosen by looking at the whole sample, so that is not an out-of-sample result and I am not going to present it as one.

The one exit family that changed something real: give-back-from-peak SHADOW

Arm as soon as the trade is 1 ATR ahead — roughly 45 points, a bar almost every session clears — then exit once it has handed back a fixed share of its best moment. Unlike the trail, this arms early enough to matter.

GIVE-BACK RULE	NET	DAYS GREEN	WORST DECILE	H1	H2	LOSS ON THE 13 BAD DAYS
hold (baseline)	+\$5,410	25/38 (66%)	−\$2,166	+\$2,319	+\$3,091	−\$13,542
give back 25% of peak	+\$2,660	35/38 (92%)	−\$1,155	+\$741	+\$1,919	−\$1,730
give back 33% of peak	+\$4,413	35/38 (92%)	−\$1,155	+\$2,409	+\$2,003	−\$1,918
give back 50% of peak	+\$3,368	35/38 (92%)	−\$1,155	+\$1,657	+\$1,712	−\$2,097
give back 66% of peak	+\$2,691	35/38 (92%)	−\$1,155	+\$1,494	+\$1,197	−\$2,192

Look at what is stable and what is not. **The exact percentage is not resolvable** — 33% is the best cell and I do not believe the gap between 33% and 50% means anything at $n=38$. But **the shape of the family is stable across every setting**: 92% of days green in all four, the same worst decile in all four, both halves positive in all four, and the loss on the bad days cut from −\$13,542 to under −\$2,200 in all four. That is the signature of a mechanism rather than a fit.

The price is real and you should see it plainly: it costs about **\$1,000 of the total (18%)** and it caps the big days — the winners fall from +\$18,952 to +\$6,331. **You are buying a smooth equity curve with your right tail.** Whether that is a good trade is a decision about how you want to feel on a Tuesday, not a statistical question, and I am not going to pretend the data answers it.

The regime exit — the mechanism you actually described, and why I cannot recommend it yet

“Good for 3 hours and then the tape changes” is exactly the regime exit: keep recomputing the same efficiency number that got us in, and leave when it decays back below the floor or the lean flips over. It does what it says on the tin — it produces a completely different risk shape.

REGIME-EXIT THRESHOLD	NET	DAYS GREEN	WORST DAY	LOSS ON THE BAD DAYS	H1	H2
hold (baseline)	+\$5,410	25/38	−\$2,501	−\$13,542	+\$2,319	+\$3,091
leave when efficiency < 0.05	−\$664	12/38	−\$1,110	−\$5,534	−\$4,164	+\$3,500
< 0.08	+\$1,404	15/38	−\$1,110	−\$4,498	−\$3,034	+\$4,438
< 0.10	+\$2,184	17/38	−\$1,110	−\$4,096	−\$2,597	+\$4,781
< 0.12	+\$4,711	18/38	−\$1,110	−\$3,834	−\$1,255	+\$5,966
< 0.15 (leave when it falls back under the entry bar)	+\$5,508	20/38	−\$1,110	−\$3,221	−\$503	+\$6,011

The tail control is genuine and large: worst day **−\$1,110 instead of −\$2,501**, and the bad days cost **−\$3,221 instead of −\$13,542**. But look at the last two columns. **Every single variant is negative or barely positive in the first half of the sample and carries its whole result in the second.** The best one is **−\$503 then +\$6,011**. A rule that lost money for eight weeks and made all of it in four is not something I will put in front of you as ready. **SHADOW** — it is the mechanism you described, it controls the tail, and it needs more tape before it can be trusted.

The run-state flip as an exit — tested, and it is too late by construction

You have `runstate.py` already; the obvious idea is to leave when it declares a run in the opposite direction. I replayed it minute by minute from trailing bars only, no look-ahead, and used it as the exit trigger.

RUN-STATE RULE	NET	\$ / TRADE	DAYS GREEN	WORST DAY	WORST DECILE	H1	H2
hold (baseline)	+\$5,410	+\$142	25/38	−\$2,501	−\$1,982	+\$2,319	+\$3,091
exit when run-state flips against us	+\$1,608	+\$42	19/38	−\$1,347	−\$1,135	−\$204	+\$1,812
exit when run-state says NO-RUN	+\$2,305	+\$61	21/38	−\$1,110	−\$902	−\$256	+\$2,561
run-state flip + ATR trail 4×/2×	+\$4,186	+\$110	21/38	−\$1,347	−\$1,135	+\$105	+\$4,081

Same shape as the efficiency exit: **real tail control, expensive, and lopsided across the halves**. And the reason is measurable. On the 13 losing sessions, the median gap between the day-rider’s entry and the first counter-direction run-state reading is **50 minutes**, ranging from 21 minutes (26 June) to 87 (23 July). On 7 August the entry was 14:12 and run-state first called a counter-run at **15:21 — 69 minutes later**, by which point the trade was already 140 points down. `runstate.py`’s own docstring says it is late by construction — a trailing 30-minute window cannot be otherwise — and this confirms it in the day-rider’s specific context. **PARKED** — revive if a shorter confirmation window is ever built and validated.

DR4

Where the money actually is: the entry, not the exit

Forget exits for a moment. Split all 38 sessions by one question — **did the direction the tape was leaning at 13:40 still hold at 20:40?**

	SESSIONS	NET (HOLD)	\$ / TRADE	DAYS GREEN	NET (ATR TRAIL)
The lean held	28	+\$18,485	+\$660	89%	+\$18,850
The lean reversed	10	-\$13,075	-\$1,308	0%	-\$12,052

Twenty-eight days at +\$660 and ten days at -\$1,308. When the lean holds, it is green nine times out of ten and the exit rule barely matters. When it reverses, it is green zero times out of ten and the exit rule barely matters — the best exit family in the study recovers \$1,023 of the \$13,075. **This is a direction problem wearing an exit problem's clothes.**

The efficiency floor — the one change that survives every test

RECOMMEND

7 August fired at efficiency **0.183**, barely over its 0.15 floor, into a tape that reversed. So I swept the floor across the whole sample. The exit is held constant at plain HOLD throughout, so what you are seeing is the entry in isolation.

EFFICIENCY FLOOR	DETECTIONS	NET	\$ / TRADE	DAYS GREEN	STRIP BEST 3	H1	H2
0.10	38 of 38	+\$5,875	+\$155	25/38	+\$844	+\$2,360	+\$3,515
0.15 — what is live	38 of 38	+\$5,410	+\$142	25/38	+\$379	+\$2,319	+\$3,091
0.18	38 of 38	+\$8,664	+\$228	26/38	+\$3,633	+\$5,802	+\$2,862
0.20	36 of 38	+\$8,447	+\$235	25/36	+\$3,416	+\$4,567	+\$3,880
0.22	34 of 38	+\$8,636	+\$254	23/34	+\$3,605	+\$3,981	+\$4,655
0.25 — the recommendation	32 of 38	+\$8,819	+\$276	23/32	+\$3,912	+\$5,007	+\$3,812
0.28	29 of 38	+\$10,204	+\$352	21/29	+\$5,297	+\$3,804	+\$6,400
0.30	28 of 38	+\$8,643	+\$309	20/28	+\$3,736	+\$3,780	+\$4,863
0.35	26 of 38	+\$9,800	+\$377	19/26	+\$4,893	+\$5,603	+\$4,197
0.40	24 of 38	+\$6,077	+\$253	16/24	+\$1,212	+\$4,143	+\$1,934

That is what a plateau looks like, and it is the exact opposite of the exit grid two sections up. Every floor from 0.18 to 0.35 lands between \$8,600 and \$10,200. The per-trade number rises steadily from \$142 to \$377 and only falls away at 0.40 when the sample gets too thin. Both halves are positive at every setting in the plateau. Compare that to the trail, where 2.5× and 4× are on opposite sides of zero.

ROBUSTNESS TEST ON THE 0.25 FLOOR	RESULT	VERDICT
Leave one day out, all 32 folds	positive in 32 of 32; range +\$7,062 to +\$11,320	survives
Strip the three best days	+\$3,912 (the 0.15 floor manages +\$379)	survives
First half vs second half	+\$289/trade then +\$263/trade	consistent
Out of sample — choose on the first half, score on the second	the 0.25 floor is +\$263/trade out of sample; the 0.15 floor is +\$163	survives
Placebo: 200 shuffles of the entry direction	p = 0.085 (the 0.15 floor scores p = 0.175)	improves, still not significant
Concentration of the benefit	the 6 dropped days total −\$3,671, but 16 June alone is −\$2,248 of it (61%)	thin — the caveat

The six sessions a 0.25 floor would have skipped: 16 June (−\$2,248), 26 June (−\$198), 6 July (+\$25), 10 July (+\$399), 14 July (−\$631) and **7 August (−\$1,018)** — including the one real live trade we have. Strip both the best and the worst of those six and the remaining four still cost −\$1,822 between them, so it is not one print carrying the whole case. But 16 June is 61% of the benefit and I will not pretend otherwise.

△ **ONE HONEST COMPLICATION — THE FLOOR IS PARTLY A TIMING KNOB, NOT JUST A FILTER**

Raising the floor does not only remove days, it changes when the detector fires on days it still takes. Going from 0.15 to 0.18 removes **zero** detections — still 38 of 38 — yet the result moves by **+\$3,254**. That is entirely re-timed entries. Specifically: 2 days fire at a different minute at the 0.18 floor, and on **16 June the higher floor picks the opposite side** and turns −\$2,248 into a winner.

At the 0.25 floor, 3 of the 32 surviving days fire at a different minute and **none** flip direction — so the recommendation is mostly a genuine filter. But the 0.15→0.18 jump is a warning that some of this surface is entry-timing luck, and it is why I am proposing 0.25 (a filter) rather than 0.28 (which scores higher and is closer to the noisy edge).

Two more entry cuts, both pointing the same way

These are causal — both are known at the moment the trade is placed.

EFFICIENCY AT THE MOMENT IT FIRED	SESSIONS	NET (HOLD)	\$ / TRADE	DAYS GREEN
0.15 - 0.25 (includes 7 August at 0.18)	9	−\$4,491	−\$499	33%
0.25 - 0.40	13	+\$5,309	+\$408	85%
0.40 - 0.60	10	+\$4,099	+\$410	80%
0.60 and above	6	+\$493	+\$82	50%

HOW FAST THE DAY DECLARED ITSELF	SESSIONS	NET (HOLD)	\$ / TRADE	DAYS GREEN
minute 9 — 13:39, the earliest it can fire	17	+\$4,814	+\$283	65%
minutes 10-12 — by 13:42	9	+\$3,751	+\$417	78%
minutes 13-19 — 13:43 to 13:49	8	−\$3,532	−\$442	62%
minute 20+ — after 13:50 (7 August fired at minute 42)	4	+\$377	+\$94	50%

Both cuts say the same thing and it is intuitive: **the days worth trading announce themselves immediately**. A day that takes 40 minutes to scrape over a 0.15 bar is a day that has not decided anything. That is precisely 7 August — minute 42, efficiency 0.183, and it spent the previous half hour flickering across the threshold (0.137, 0.079, 0.119, 0.140, 0.130). The strategy’s own detector was telling us the day had not committed; the floor was just set low enough to let it through anyway.

These two cuts are not independent of the floor sweep — they are the same fact seen three ways — so do not add up their apparent benefits. Treat them as three views of one finding: **demand a clearer signal**.

And the floor is fragile in the other direction too — you were right to flag it

The CME-anchor test gave **zero** detections at 0.15 over 33 sessions, because folding in 15 hours of overnight chop makes the path 25× longer and collapses efficiency by an order of magnitude. At the cash-open anchor the same 0.15 fires on **38 of 38 sessions — 100%**. So today’s floor is not a filter at all; **it is a formality that has never once declined to trade**. A threshold that says yes every single time is not selecting anything — it is only choosing a minute and a direction. Raising it to 0.25 makes it decline 6 times in 38, which is the first time it has ever functioned as a gate.

DR5 Regime segmentation — scoring each rule only on its home tape

A single blanket number across the whole tape averages the regimes a config should be ON with the ones it should be OFF. So here is every rule scored per regime. I labelled each session by **how much of its range it converted into net movement** ($|net| \div \text{range over 13:30-20:40}$) plus the size of the range — not by the clock. Median session range over the 38 was 473 points.

REGIME	DEFINITION	SESSIONS
CLEAN TREND	closed with at least 55% of its range as net movement	13
BUILDING TREND	35-55% of range converted	14
VIOLENT WHIPSAW	under 35% converted, on a range at or above the 473pt median	4
DEAD CHOP	under 35% converted, on a range below the median	7

REGIME	N	HOLD	ATR TRAIL 4×/2×	OLD FIXED 150/100	GIVE BACK 33%	REGIME EXIT <0.15	FIXED 300PT STOP
CLEAN TREND	13	+\$5,264 (+ \$405)	+\$1,454 (+ \$112)	+\$1,697 (+ \$131)	—	+\$3,967 (+ \$305)	+\$4,251 (+ \$327)

BUILDING TREND	14	+\$3,435 (+ \$245)	+\$5,688 (+ \$406)	+\$4,803 (+ \$343)	—	+\$2,819 (+ \$201)	+\$39 (+ \$3)
VIOLENT WHIPSAW	4	-\$1,559 (– \$390)	+\$306 (+ \$77)	–\$820 (–\$205)	—	–\$1,110 (– \$278)	–\$2,630 (– \$658)
DEAD CHOP	7	–\$1,730 (– \$247)	–\$651 (–\$93)	–\$675 (–\$96)	—	–\$168 (–\$24)	–\$4,090 (– \$584)

Cells are total \$ (per-trade \$). Give-back is shown blank because with only 4 whipsaw and 7 dead-chop sessions the per-regime split for that family is below any n I would quote.

THE REGIME RESULT IS REAL, AND IT IS NOT DEPLOYABLE. BOTH THINGS ARE TRUE.

The pattern is clean and it makes mechanical sense: **on a clean trend you must HOLD** (+\$405 a trade, and the ATR trail throws away three quarters of it by exiting on a mid-trend pullback). **On a building or a whippy trend you must TRAIL** (+\$406 vs +\$245, and +\$77 vs –\$390). **On dead chop nothing works** — every rule loses, the least-bad being to leave early.

But the regime label is only knowable at 20:40, and the trade is placed at 13:40. You cannot route on it. Anyone who tells you this table is a strategy is selling you hindsight. What it is good for: it explains why the exit grid is noise. The best exit genuinely does depend on the regime, the regime is unknown at entry, so a single fixed exit rule is averaging four different right answers — and which average wins depends on which regimes happened to show up in your sample. That is exactly the –0.30 rank correlation.

So I looked for something knowable at entry that predicts the regime

Entry ATR is the obvious candidate — you have it in hand at 13:40. Terciles across the 38 sessions:

ENTRY ATR BAND	SESSIONS	HOLD NET	\$ / TRADE	ATR TRAIL NET \$ / TRADE
-----------------------	-----------------	-----------------	-------------------	---------------------------------

QUIET — ATR under 38	12	+\$74	+\$6	+\$308	+\$26
MID — ATR 38 to 50	13	+\$3,826	+\$294	+\$3,590	+\$276
BUSY — ATR 50 or more	13	+\$1,510	+\$116	+\$2,900	+\$223

There is a hint here — the quiet band produces essentially nothing either way, which fits the desk’s standing finding that low-vol grind is untradeable — but the ordering is not monotone and 12–13 sessions a band is not enough to act on. **I am not proposing an ATR floor on the day-rider.** Worth a note in the book and a re-check in a month. **SHADOW**

On time-of-day, and why there is no overnight split in this section

The report’s standard split is overnight/pre-open versus the US session after 13:30. **For the day-rider that split does not exist by construction:** the anchor is the 13:30 cash open, the detector needs 9 minutes of tape, and there is a hard no-new-entry cutoff at 15:00. Across all 38 sessions the earliest detection was 13:39 and the latest 14:12 — a 33-minute window. Reporting an “overnight” bucket would be reporting an empty set. The meaningful time cut is the one in the previous section: how early inside that 33-minute window it fired, and that one has real content.

The same reasoning kills the Asia-block question — 00:00–07:00 UTC is permanently benched on the tournament, and the day-rider never trades there anyway.

Does a trade that is deep offside ever come back? **STRIKING**

You asked directly, because it decides everything. Here is the answer, measured on every session.

IF IT IS THIS FAR OFFSIDE AT...	FINISHED GREEN	FINISHED BETTER THAN IT STOOD	MEDIAN FINISH
–50pt at 15:00Z	16	5 of 16	7 of 16
–100pt at 15:00Z	15	4 of 15	6 of 15
–150pt at 15:00Z	8	0 of 8	2 of 8
–200pt at 15:00Z	6	0 of 6	2 of 6
–250pt at 15:00Z	3	0 of 3	1 of 3
–150pt at 17:00Z	10	0 of 10	6 of 10
–250pt at 17:00Z	4	0 of 4	3 of 4

A trade that is 150 points offside at 15:00 has never once finished green in this sample. Not in 8 attempts. Nor at –200 (0 of 6), nor at –250 (0 of 3), nor at –150 by 17:00 (0 of 10). The median such trade finishes around 270–280 points down and the worst finishes 625 down. **The “it might come back” hope is not supported by a single example.**

That is the strongest-looking evidence for a checkpoint in the whole study — and it is exactly the rule that fell apart when I moved the clock thirty minutes. Both things are true, and the resolution is honest: **with 8 observations you cannot tell “never comes back” from “has not come back yet”**. A rule that has 8 supporting cases and collapses under a 30-minute perturbation is a promising hypothesis and not a deployment. **SHADOW** — log the 15:00 mark every session, and revisit at $n = 25$ offside cases.

DR6 Sizing, isolation, and what we cannot yet know

Sizing: one day-rider trade is fifteen tournament scalps

RECOMMEND 1 LOT

Measured, not estimated. Tournament figures are the 597 clean MNQ rows in the trade record (`data_quality IS NULL`, cleanup exits excluded); the day-rider figures are the 38 modelled sessions at its live 2-lot size.

N	MEAN	TYPICAL ABSOLUTE SWING	STD DEV	WORST	BEST	
One tournament lot	597	−\$6.54	\$55.50	\$72.09	−\$235.50	+\$536.00
One day-rider trade at 2 lots	38	+\$142.37	\$855.11	\$1,059.19	−\$2,501	+\$1,757
... the ratio	—	—	15×	15×	11×	3×
One day-rider trade at 1 lot	38	+\$71	\$428	\$530	−\$1,250	+\$879

Put it in the bluntest terms available. **The tournament's entire all-time realised loss across 597 lots is −\$3,905.50. The day-rider's worst single modelled day is −\$2,501 — 64% of that, in one trade, in one afternoon.** On 7 August the tournament made +\$299 across 21 lots and the day-rider lost roughly −\$1,014 on one. That is not a strategy running alongside the desk; at 2 lots it is the desk, and everything else is rounding.

My recommendation is to run it at 1 lot until it has 56 recorded sessions. The reason is not caution for its own sake — it is that we are paying full variance for information we cannot yet read. At 1 lot the same 38 sessions still net +**\$2,705**, the worst day is **−\$1,250** rather than −\$2,501, and every conclusion in this section is unchanged because halving the size halves both sides exactly (fees scale with contracts too). We learn the same amount for half the risk. Size back up when the evidence justifies it, which brings me to...

HOW MUCH TAPE WOULD ACTUALLY SETTLE THIS — THE NUMBER YOU SHOULD HOLD ME TO

N SO FAR	MEAN / TRADE	STD DEV	T-STAT TODAY	SESSIONS NEEDED FOR T = 2.0
Current floor 0.15	38	+\$142	\$1,059	0.83 221 — about 10.5 months
Proposed floor 0.25	32	+\$276	\$1,034	1.51 56 — about 2.7 months

The strategy trades on essentially every session, so sessions and trades are the same thing. **At today's settings we would need ten and a half months to know whether this works.** At the higher floor, under three. That is a second, independent argument for raising it: it does not merely make more money in the backtest, **it makes the experiment answerable inside a quarter instead of a year.**

And the placebo keeps me honest about where we stand right now. Shuffling the entry direction 200 times, a coin flip beats the real detector in 35 of 200 draws at the current floor (**p = 0.175**) and 17 of 200 at the proposed one (**p = 0.085**). **Neither is significant.** On 38 sessions of multi-hour holds we genuinely cannot yet distinguish this from luck — the median coin flip loses \$485 and the real thing makes \$5,410, but the spread on the coin flips runs from −\$10,298 to +\$9,262. That is what a \$1,059 standard deviation does to a 38-day sample.

Isolation: the split exists in exactly one place, and it is not the place you look GAP

You asked me to confirm the two books are separated everywhere you look. They are not. `pnl.realized(desk=...)` was built and works correctly, but I grepped every call site and it is passed in **one file**.

SURFACE	CALL	SPLITS THE BOOKS?
scripts/sweep.py lines 258-259	<code>pnl.day(..., desk="tournament") and desk="day_rider"</code>	YES
web.py:252 — the dashboard P&L strip you actually read	<code>pnl.day(c, "MNQ", now)</code>	no — blended
web.py:253-254, 266 — the 7-day and 30-day figures	<code>pnl.realized(...), no desk</code>	no — blended
core.py:406 — the daily-loss kill switch	<code>pnl.day(self._store, symbol, now)</code>	no — blended

The dashboard one is cosmetic and easily fixed. **The kill-switch one is not cosmetic.** As written, a day-rider loss counts against the tournament's daily-loss headroom, so one bad afternoon on a strategy with a \$1,059 daily standard deviation could halt eight gates that did nothing wrong. Two mitigating facts, both worth stating: it has never fired, because `max_daily_loss_usd` and `loss_streak_halt` are **both set to 0** (the kill switches are off for paper tuning) — and the day-rider does not write trade rows at all yet, so there is currently nothing for it to contaminate. **Both of those protections are accidents.** The moment the write path lands, or the moment you turn a kill switch on, this becomes live. Fix the `desk=` argument at the same time as the write path, not after it.

What I would do, and what I would not

#	CHANGE	EVIDENCE	VERDICT
1	Raise the efficiency floor 0.15 → 0.25. One constant in drift.py.	+\$5,410 → +\$8,819 over 38 sessions; +\$142 → +\$276 a trade; plateau 0.18-0.35; 32/32 leave-one-out; consistent across halves; cuts time-to-answer from 10.5 months to 2.7	DO IT
2	Write the trade rows. A record_trade call on entry and exit.	the strategy has been live for 3 days and has produced 0 rows; we are flying on a journal file that gets overwritten every session	DO IT
3	Pass desk= in web.py and core.py at the same time as #2.	3 of 4 P&L surfaces blend the books, including the kill switch	DO IT
4	Cut to 1 lot until 56 recorded sessions.	15× the swing of a tournament lot; worst day = 64% of the desk's all-time loss; halving size costs no information	DO IT
5	Leave the 4×/2× ATR trail alone.	its shipped rationale is wrong, but the rule itself is harmless and mildly positive; changing it again on a noise surface is churn	LEAVE IT
6	Give-back-33%-of-peak, in shadow.	92% days green and half the tail, at \$1,000 of the total; the family is stable even though the exact % is not	SHADOW
7	Regime exit (efficiency decay) and run-state flip, in shadow.	real tail control — worst day — \$1,110 vs −\$2,501 — but every variant is flat-to-negative in the first half	SHADOW
8	Log the 15:00Z mark-to-market every session.	0 of 8 trades 150pt offside at 15:00 finished green, but the rule collapses under a 30-min perturbation — needs n=25	SHADOW
9	Do not add a stop of any width below 600 points.	every width 75–400pt is worse than none; 150/200/250/300/400 all turn the strategy negative; winners routinely draw 250–490pt first	REFUTED
10	Do not tune the exit further.	Spearman −0.30 between in-sample and out-of-sample rank across 114 rules; best-in-sample rule ranked 71st of 114 out of sample and lost to plain holding	REFUTED

WHAT WOULD SETTLE IT, AND HOW LONG

The question that decides this strategy is not “when do we get out” — it is “is the 13:40 lean better than a coin flip at picking the day’s direction?” Everything else is second-order. In this sample it was right 28 times in 38 (74%), worth +\$660 a trade when right and −\$1,308 when

wrong, and that is not yet statistically distinguishable from luck ($p = 0.175$ at the current floor, 0.085 at the proposed one).

What settles it: 56 recorded sessions at the 0.25 floor, with trade rows actually written, at 1 lot. That is roughly **eleven trading weeks** from the day the write path lands. If the per-trade figure holds anywhere near +\$276 over that stretch it clears $t = 2$ and the strategy is real. If it drifts toward zero it was three good days in June and July.

Two things must be true for that clock to start, and neither is true today: **the trades have to be recorded**, and **the floor has to be the one we are testing**. Right now we are running an experiment that writes nothing down and would need ten months to read.

△ **Everything above rests on 38 sessions and one real live trade.** One loss is not a refutation and an in-sample sweep is not a validation — I have tried hard not to dress either as the other. What I am confident about is narrow and mechanical: **a trailing stop cannot rescue a trade that never went your way, and on the seven days that cost real money the trade never went our way.** Whether the entry filter that follows from that is a genuine edge or three lucky sessions in June is a question 56 more days answers and no amount of further backtesting will.

M1 The biggest runs on MNQ this week — did we show up?

Forget what the desk did this week — start from the raw tape. MNQ printed **67 runs of 1.5×ATR or bigger** in the last seven days (a 15-minute move of at least 65 points, against a typical 15-min range of 44). That is the whole tape, not a top-ten. We **caught 10**, we were positioned against **3**, and we **sat out 54**. The runs we aligned with banked **+\$156** of honest money — but that is only 6% of their \$2416 hindsight ceiling, and the 54

we sat out left **\$9917** on the table. The bleed is not the runs; it is the fading we do around them.

THE FULL CENSUS

Every qualifying run, in order. **Move** is the swing in points; **\$ 1lot** is that move on one MNQ lot (\$2/pt) — the hindsight ceiling. **real \$** is what a gate actually banked near it. **Flow** is net aggressor volume in the 60s before; **amp** is pre-run amplitude; **book** is the far-side L2 depth share (below 0.50 = the side price ran toward was thin). Green = we caught it, red = we fought it.

TIME (UTC)	DIR	MOVE	\$ 1LOT	US	GATE	REAL \$	FLOW	AMP	BOOK	CLUSTER
08-02 23:52	DN	-66	132	sat out	—		+0	0.21	—	UNCLASS
08-03 00:16	DN	-69	138	sat out	—		+6	0.20	0.48	UNCLASS
08-03 08:15	DN	-80	160	sat out	—		+192	0.10	0.50	VACUUM
08-03 08:34	UP	+67	134	sat out	—		+51	0.08	0.50	UNCLASS
08-03 11:46	DN	-76	151	sat out	—		-69	0.09	0.52	UNCLASS
08-03 12:07	UP	+69	138	sat out	—		-28	0.09	0.50	UNCLASS
08-03 13:18	DN	-141	282	sat out	—		-5	0.15	0.51	OPEN/NEWS
08-03 13:33	UP	+155	310	caught	nipc_long_A	-119	+831	0.43	0.47	OPEN/NEWS
08-03 13:48	UP	+105	210	caught	nipc_long_A	-178	+627	0.35	0.53	OPEN/NEWS
08-03 14:15	UP	+115	230	caught	nipc_long_A	+120	+228	0.20	0.51	OPEN/NEWS
08-03 14:45	UP	+88	176	caught	nipc_long_A	+36	-12	0.12	0.49	OPEN/NEWS
08-04 01:16	UP	+65	130	sat out	—		-122	0.08	0.51	UNCLASS
08-04 01:40	UP	+128	256	sat out	—		+222	0.10	0.49	FLOW-LED
08-04 02:01	DN	-68	136	sat out	—		-60	0.07	0.53	UNCLASS
08-04 08:31	DN	-66	133	sat out	—		-13	0.05	0.50	UNCLASS
08-04 11:40	UP	+84	169	caught	abs_veto_long_A	-48	-428	0.08	0.46	VACUUM
08-04 13:27	UP	+216	432	caught	abs_veto_long_A	+138	-368	0.12	0.50	OPEN/NEWS

08-04 13:56	UP	+100	200	sat out	—		-92	0.27	0.51	OPEN/NEWS
08-04 14:58	UP	+88	176	caught	abs_veto_l ong_A	+94	-146	0.09	0.48	OPEN/ NEWS
08-04 17:33	UP	+80	159	sat out	—		-170	0.08	0.47	UNCLASS
08-04 19:44	DN	-78	156	sat out	—		+1034	0.12	0.49	VACUUM
08-04 20:06	DN	-151	302	sat out	—		+193	0.26	0.46	UNCLASS
08-05 08:07	DN	-74	147	sat out	—		-152	0.07	0.51	FLOW-LED
08-05 13:25	UP	+96	193	sat out	—		+601	0.09	0.56	OPEN/NEWS
08-05 13:48	DN	-126	253	caught	nipc_short _A	-58	+638	0.18	0.48	OPEN/ NEWS
08-05 14:16	DN	-129	258	sat out	—		+222	0.15	0.48	OPEN/NEWS
08-05 14:39	DN	-78	156	sat out	—		+376	0.25	0.50	OPEN/NEWS
08-05 14:54	UP	+73	146	sat out	—		-3892	0.40	0.54	OPEN/NEWS
08-05 15:11	DN	-66	133	sat out	—		-92	0.20	0.50	UNCLASS
08-05 15:42	DN	-112	225	sat out	—		+103	0.13	0.50	UNCLASS
08-05 16:14	UP	+108	217	sat out	—		-364	0.10	0.50	UNCLASS
08-05 16:55	DN	-66	132	sat out	—		+144	0.11	0.50	UNCLASS
08-05 19:51	DN	-173	346	sat out	—		+23	0.15	0.52	UNCLASS
08-05 22:19	UP	+74	148	sat out	—		-26	0.10	0.44	UNCLASS
08-06 01:11	DN	-68	136	sat out	—		+55	0.09	0.49	UNCLASS
08-06 02:30	UP	+77	154	sat out	—		-50	0.11	0.51	UNCLASS
08-06 05:59	DN	-74	148	sat out	—		+60	0.07	0.48	VACUUM
08-06 08:58	DN	-72	144	sat out	—		+7	0.05	0.53	UNCLASS
08-06 10:51	DN	-107	214	sat out	—		-43	0.06	0.48	UNCLASS
08-06 11:06	UP	+78	155	sat out	—		-294	0.16	0.50	VACUUM
08-06 12:53	DN	-67	134	sat out	—		+110	0.08	0.50	UNCLASS
08-06 13:17	DN	-82	165	sat out	—		-33	0.09	0.51	OPEN/NEWS
08-06 13:33	UP	+243	486	FOUGHT	abs_veto_s hort_A	-130	+1193	0.44	0.52	OPEN/NEWS
08-06 14:03	UP	+152	304	sat out	—		-29	0.21	0.53	OPEN/NEWS
08-06 14:24	DN	-79	158	sat out	—		-808	0.12	0.48	OPEN/NEWS

08-06 14:46	UP	+120	240	sat out	—		+206	0.28	0.51	OPEN/NEWS
08-06 15:10	DN	-94	189	sat out	—		+1114	0.19	0.53	VACUUM
08-06 15:29	DN	-68	136	sat out	—		+45	0.16	0.52	UNCLASS
08-06 15:48	DN	-106	212	sat out	—		+797	0.25	0.48	VACUUM
08-06 16:46	UP	+88	175	sat out	—		+79	0.15	0.49	UNCLASS
08-06 19:52	UP	+72	144	sat out	—		-209	0.12	0.51	UNCLASS
08-06 22:00	UP	+80	160	sat out	—		+114	0.11	0.52	UNCLASS
08-07 00:08	DN	-106	212	sat out	—		-8	0.10	0.49	UNCLASS
08-07 01:06	DN	-77	154	sat out	—		+59	0.07	0.53	UNCLASS
08-07 01:33	DN	-65	130	sat out	—		-76	0.19	0.50	UNCLASS
08-07 12:28	UP	+142	284	caught	capitulatio n_long_B	+207	-22	0.08	0.47	UNCLASS
08-07 13:00	DN	-105	210	sat out	—		-20	0.06	0.48	OPEN/NEWS
08-07 13:29	UP	+88	176	caught	abs_veto_s hort_B	-36	+51	0.20	0.52	OPEN/ NEWS
08-07 13:50	DN	-143	286	sat out	—		+335	0.26	0.47	OPEN/NEWS
08-07 14:12	UP	+183	366	sat out	—		-2589	0.40	0.49	OPEN/NEWS
08-07 14:29	DN	-88	176	sat out	—		-7	0.27	0.52	OPEN/NEWS
08-07 14:52	UP	+97	194	FOUGHT	exhaustion_ short_A	+104	-329	0.17	0.52	OPEN/NEWS
08-07 15:17	UP	+93	186	sat out	—		+333	0.12	0.50	UNCLASS
08-07 16:17	DN	-132	265	FOUGHT	capitulation _long_A	-52	+207	0.08	0.50	UNCLASS
08-07 16:32	UP	+77	154	sat out	—		+124	0.18	0.51	UNCLASS
08-07 18:56	UP	+113	226	sat out	—		-227	0.06	0.52	UNCLASS
08-07 19:44	UP	+69	138	sat out	—		-266	0.10	0.51	VACUUM

What the runs were — by cause

Each run gets a cause fingerprint from its pre-run tape. This is the map for where to hunt a new gate.

CLUSTER

RUNS

WHAT IT IS

UNCLASS	33	no clear tape tell (quiet ignition)
OPEN/NEWS	24	the 13:00-15:00 UTC cash-open / data window
VACUUM	8	moved AGAINST the tape (a snap-back / stop-run)
FLOW-LED	2	aggressors drove it (a continuation)

HONEST MONEY — CEILING VS WHAT WE BANKED

RUNS	HINDSIGHT CEILING	WHAT WE REALLY MADE
Caught (aligned)	10	\$2416 +\$156 (6%)
Fought (against)	3	\$946 -\$78
Sat out	54	\$9917 \$0 ← the money on the table

The runs made money and the fading gave it back; we convert a fraction of the ceiling and ignore the rest. The next two movements ask whether we can do better — first with the gates we own, then with new ones.

M2 The idle-gate lab — could the gates we already own have caught the 54 misses?

Movement 1 froze the tape: 67 big runs on MNQ this week, **54 of them we sat out**, and those 54 carried a hindsight ceiling of **\$9,917** on a single lot. This movement asks the only question that matters before we go inventing anything new: we already own six gates — if we had simply left all six armed, all week, would any of them have been in those runs? So I took every one of the six live gates out of `deciders.py`, fired them **mechanically** (exactly as the desk runs them — live parameters, live ATR floors, live 55-second veto, live confirm) and again **ungated** (the raw gate shape with every filter stripped off), in-direction only, anywhere in the ten minutes before each run ignited, and repriced every exit tick-by-tick on this week's real `capture.db` trade prints under each gate's deployed two-lot

exit. **The mechanical answer is \$-474 on 7 fires across 53 testable runs.** Ungated it is +\$997.50 on 19 fires, but strip the three best trades and that becomes **+\$62.50**. This is an **HONEST NULL** — and, more usefully, I can now tell you exactly why it is a null, with a number rather than a shrug: **the money in these runs lives in the first ninety seconds after ignition, and our gates fire seven to eight minutes before it.**

LEAD WITH WHAT SURVIVED

Three things came out of this alive and worth acting on, and I want them in front of the null rather than behind it.

1. **The 55-second wait on abs_veto_short is costing more than the veto saves — \$596.50 vs \$169 on this population.** The absorption test is doing its job (it blocked three raw fires, all three losers, saving \$169). The problem is the **55-second delay** attached to it: on the two thrusts that survived the test, waiting 55 seconds turned +\$97 into –\$71 and +\$282.50 into –\$146. Same signal, same bar, same exit — only the entry clock differs. Net effect on run tape: **-\$427.50**. That is a specific, testable Saturday job: keep the test, shrink the wait.
2. **The quiet-tape clip (ATR<22 → Lot A banks \$40, Lot B banks \$60/1.75R) was live on 35 of the 53 runs and cost \$1,202 on run tape.** With identical entries, clip ON returns +\$1,696.50 and clip OFF returns +\$2,898.50. It is a chop tool doing chop work in the wrong place. A regime carve-out is the obvious fix.
3. **grind_long's ext_hi 2.0 ceiling is the single filter standing between us and these runs.** Ungated, grind fired at 7 of the 22 up-runs for +\$742.50; mechanically it fired **once**, for –\$130. The ceiling ("don't buy what is already stretched") blocked six of its seven pre-run fires. That is not automatically an argument to remove it — it was proven on the whole book, and on a sat-out-run population it is being judged on exactly the trades it was designed to skip — but it is the highest-value knob in this section and it deserves a proper regime-conditional re-derivation.

How the test was run — so you can trust the minus sign

Everything below comes from one script, `scripts/idle_gate_backtest_wk32.py`, run once against the frozen census. No number here is carried over from any prior week.

COMPONENT	RULE
Population	The 54 sat-out runs of the frozen Movement-1 census, read straight out of the locked census stdout — not re-derived, not re-filtered.
Testable subset	53 of 54. <code>capture.db</code> keeps five trading days of ticks (first print 08-03 00:00:00 UTC), so the 08-02 23:52 run (−66pt, \$132 ceiling) has no pre-run tape and is excluded rather than faked. The remaining 53 carry \$9,780 of ceiling (the census's \$9,917 total is computed on unrounded moves; \$9,780 is the sum of the 53 printed moves in the frozen table).
Gates	All six live tournament slots: <code>grind_long</code> , <code>capitulation_long</code> , <code>abs_veto_long</code> (long side) and <code>rgv_short</code> , <code>exhaustion_short</code> , <code>abs_veto_short</code> (short side). The actual decider functions are imported and called — there is no re-implementation to drift.
Direction	In-direction only. A long gate is only tested on the 22 up-runs, a short gate only on the 31 down-runs. We are asking "could it have been in the move", not "could it have been anywhere near it".
Window	The ten minutes before ignition . Bar gates are tested on each of the ten completed 1-minute bars; the two tape gates (capitulation, exhaustion) are tested every 5 seconds across the whole window. First fire wins.
Mechanical mode	Live params and live filters: <code>grind</code> $ATR \geq 10$ + <code>ext_hi</code> 2.0; <code>capitulation</code> $ATR \geq 10$ + <code>require_flip</code> ; <code>abs_veto</code> <code>amp_floor</code> 0.0004 + the 55-second absorption veto (and its 55s entry delay); <code>rgv_short</code> $ATR \geq 20$; <code>exhaustion</code> + its 5-second confirm-veto.
Ungated mode	The raw gate shape. <code>ATR</code> floors off, <code>amp</code> floor off, <code>ext</code> ceiling back to its 4.0 default, <code>require_flip</code> off, both vetoes off. The gate's core geometry is kept — this is "filters off", not "a different gate".
Exit	Each gate's deployed two-lot exit from <code>data/exit_overrides.json</code> — Lot A fixed-R, Lot B chandelier or fixed-R, one lot each, plus the quiet-tape clip below 22pt <code>ATR</code> and the native $1 \times ATR$ protective stop on the loss side. Repriced on every real trade print, 60-minute hold cap.
Costs	\$1.50 per lot round-trip, \$2.00 per point. Two lots per fire, so \$3.00 a fire before anything else.
Bars	1-minute bars built from the 5-second bars with integer flooring (<code>bar_ts - bar_ts % 60</code>). This is deliberate — the report's own instrument-failure list flags the float-division version as a silent no-op that reads 5-second bars as minutes.

Data actually consumed: **6,761,816 trade ticks** (08-03 00:00:00 → 08-07 21:00:00 UTC), **16.47 million** level-1 book rows, 6,780 one-minute bars inside the tick window.

The scoreboard — mechanically, the desk cannot get into these runs

This is the headline table. "Eligible" is how many sat-out runs that gate could even have traded (its own direction only). Green rows made money, red rows lost it.

GATE	SIDE	ELIGIBLE RUNS	FIRES	NET \$	\$ / FIRE	WIN %	HOW THEY ENDED
grind_long	LONG	22	1	−130.0	−130.0	0%	both lots stopped
capitulation_long	LONG	22	1	+97.0	+97.0	100%	both lots hit target
abs_veto_long	LONG	22	2	−106.0	−53.0	0%	4 of 4 lots stopped
rgv_short	SHORT	31	0	0.0	—	—	never fired once
exhaustion_short	SHORT	31	1	−118.0	−118.0	0%	both lots stopped
abs_veto_short	SHORT	31	2	−217.0	−108.5	0%	4 of 4 lots stopped
ALL SIX, mechanical	—	53 runs	7	−474.0	−67.7	14%	1 winner in 7

Seven fires across 53 runs. Not seven profitable fires — seven fires at all, from six gates over a whole week of the biggest moves MNQ printed. Six of the seven lost. The one that won, capitulation_long on the 08-06 19:52 up-run, made \$97 on a 72-point move whose hindsight ceiling was \$144, and it fired 100 seconds before ignition, which is the only fire in the whole table that landed anywhere near the money.

The same six with every filter switched off

Now strip the ATR floors, the amplitude floor, the extension ceiling, the flip requirement and both vetoes. This is the "leave everything on, damn the filters" book — the fair upper bound on what the gates we own could physically have done.

GATE	ELIGIBLE	FIRES	NET \$	\$ / FIRE	WIN %	VS MECHANICAL
grind_long	22	7	+742.5	+106.1	57%	+6 fires, + \$872.50
capitulation_long	22	2	+194.0	+97.0	100%	+1 fire, +\$97
abs_veto_long	22	2	−106.0	−53.0	0%	identical — no filter bit
rgv_short	31	1	+97.0	+97.0	100%	+1 fire (the ATR≥20 floor)

exhaustion_short	31	2	−140.5	−70.2	0%	+1 fire, −\$22.50
abs_veto_short	31	5	+210.5	+42.1	40%	+3 fires, +\$427.50
ALL SIX, ungated	53 runs	19	+997.5	+52.5	47%	+12 fires, +\$1,471.50

DO NOT GET EXCITED ABOUT THE +\$997.50

It is three trades. Strip the best **one** and it falls to **+\$601**. Strip the best **three** and it falls to **+\$62.50** on sixteen remaining fires — twenty dollars a day, on a hindsight-selected population, in-sample, with no arming rule attached. Day by day it is +\$97 / +\$430 / +\$318 / −\$80 / +\$234, and leaving out the best day (08-04) still leaves +\$568, so it is not a one-session artefact — but "not a one-session artefact" and "an edge" are different sentences. Cost-stressing to \$3.00 a lot round-trip leaves +\$940.50 and to \$6.00 leaves +\$826.50, so costs are not what is wrong with it. **Nineteen fires is what is wrong with it.**

The money on the table — reconciled

BOOK	RUNS TOUCHED	FIRES	CEILING OF THE RUNS TOUCHED	HONEST \$ BANKED	% OF THE FULL \$9,780	BIGGEST 15 RUNS CAUGHT
Mechanical (as the desk runs)	7 / 53	7	\$1,266	−\$474	−4.8%	2 / 15
Ungated (filters off)	17 / 53	19	\$3,256	+\$997.50	+10.2%	4 / 15
The rest	36 / 53	0	\$6,524	\$0	0%	11 / 15

Read the bottom row twice. **Thirty-six of the 53 runs — carrying two-thirds of the ceiling — drew no fire from any of the six gates even with every filter removed.** Eleven of the fifteen biggest runs of the week are in that group. Whatever we do to the thresholds, most of this money is not reachable by these six mechanisms, because the mechanisms are not looking at the thing that precedes these moves.

Why each gate misses — the mechanism, clause by clause

This is the part I actually care about, and it is the part a scoreboard alone can't tell you. For every eligible run I walked the gate's own rule chain in order and

recorded **how far it ever got** across the whole ten-minute window — which clause was the last one standing between it and a fire. Counts are runs, not bars.

grind_long — the trend was never there to ride

FURTHEST THE GATE GOT (MECHANICAL)	RUNS	UNGATED	RUNS
VWAP slope not up — no established up-trend	12	VWAP slope not up	12
Already stretched (ext > 2.0 ATR) — the ext_hi ceiling	6	—	—
Price not riding above VWAP (ext < 0.3)	2	Price not riding above VWAP	2
FIRES	1	FIRES	7
No bars available	1	No bars available	1

Mechanism: grind is a continuation gate. It needs an established VWAP trend it can join. In 12 of the 22 up-runs there simply was no up-trend in the ten minutes before — these runs ignite out of flat or falling tape, which is precisely why they are runs. Then, on the six occasions where a trend was established, the ext_hi 2.0 ceiling blocked entry because price was already more than two ATR above VWAP. Those six are the interesting ones: ungated they are worth +\$872.50 relative to the mechanical book. But note honestly what that means — the ceiling was added on 08-01 to stop grind buying stretched tape, and a sat-out-big-run population is by construction the population where buying stretched tape happens to work. This needs a full-book regime-conditional re-derivation before anyone touches it, not a Movement-2 headline.

abs_veto_long / abs_veto_short — the thrust just isn't there yet

FURTHEST THE GATE GOT	ABS_VETO_LONG (22 UP-RUNS)	ABS_VETO_SHORT (31 DOWN-RUNS)
No 5-bar thrust of 1.5 ATR	10	13
Thrust was the wrong way	4	3
No volume surge on the thrust bar	3	8

Amplitude floor: tape too thin (atr% < 0.04%)	2	2
55-second absorption VETO	0	3
FIRES	2	2
No bars	1	0

Mechanism: abs_veto is a 5-minute-thrust gate. It needs the move to have already happened at 1.5 ATR over five bars before it will look at it — and by definition, in the ten minutes before ignition, the move has not happened yet. In 23 of 53 eligible runs there was no qualifying thrust at any point in the window; in another 7 the thrust pointed the wrong way (price was moving against the run it was about to make, which is a snap-back, not a continuation). Eleven more got a thrust but no volume behind it. This is a structural mismatch, not a threshold problem: **a five-bar-lookback momentum gate is being asked to be early, and it is built to be late.**

rgv_short — the ATR floor stands it down, and the stretch never appears

FURTHEST THE GATE GOT (MECHANICAL)	RUNS	UNGATED	RUNS
ATR floor of 20pt not met	13	—	—
Not stretched ≥ 1.5 ATR above VWAP	10	Not stretched above VWAP	18
Regime stand-down (vol too high / slope too steep)	7	Regime stand-down	7
No fresh down-turn	1	No fresh down-turn	5
FIRES	0	FIRES	1

Mechanism: rgv is a fade. It wants price stretched above VWAP and rolling over — a spent move it can sell into. Down-runs mostly begin from tape that is not stretched above VWAP (18 of 31 ungated), so the setup never presents. On top of that the $ATR \geq 20$ floor stands the gate down on 13 runs before it even gets to look at the stretch, which is why it fires zero times mechanically. **Zero fires from a gate on 31 eligible runs is the cleanest possible statement that this is the wrong instrument for this job.**

capitulation_long and exhaustion_short — the tape gates barely exist this week

FURTHEST THE GATE GOT	CAPITULATION_LONG (22 UP-RUNS)	EXHAUSTION_SHORT (31 DOWN-RUNS)
Volume climax never reached (sell < 2.5x baseline / flow < 400)	18	18
Aggression DID move price > 2pt — so it wasn't absorption	—	8
One side not dominant enough / no wall on the hit side	1	3
Buyers never flipped at the low (require_flip)	1	—
5-second confirm-veto: price ran against the fade	—	1
FIRES	1	1
No tape	1	0

Mechanism: both of these want a violent one-sided flush that gets absorbed. In eighteen of the runs there was no volume climax at all in the ten minutes before ignition — these moves start quietly. For exhaustion specifically there is a second, harder problem: on eight of the runs there was heavy flow, but it **moved price more than 2 points**, which is exactly the condition the gate uses to say "this is not absorption, this is a real move". The gate correctly identified the thing as a genuine move and correctly declined to fade it — and a genuine move is precisely what we were hoping to catch. **Its rule is doing what it says on the tin; the tin is the wrong shape for run-hunting.**

The killer number: the gates fire at chance, near the runs

Here is the check that turns "they didn't fire much" into "they are not looking at anything predictive". Across the whole tick-covered week I counted every minute (bar gates) or 20-second sample (tape gates) in which each gate's condition was **true**, and then asked what share of those fell inside one of the 53 ten-minute pre-run windows. Those windows cover **8.5%** of the minute bars and **7.8%** of the tape samples — that is the chance rate. If a gate were telling us anything about imminent runs, its share would be well above that.

GATE	MODE	SHARE	VS CHANCE
------	------	-------	-----------

		SIGNAL-MINUTES / SAMPLES ALL WEEK	INSIDE A PRE-RUN WINDOW		
grind_long	mechanical	374	38	10.2%	1.20×
grind_long	ungated	1,180	80	6.8%	0.80×
capitulation_long	mechanical	7	1	14.3%	1.83×
capitulation_long	ungated	49	2	4.1%	0.53×
abs_veto_long	mechanical	81	10	12.3%	1.45×
abs_veto_long	ungated	201	11	5.5%	0.65×
rgv_short	mechanical	5	2	40.0%	4.7× (n=5)
rgv_short	ungated	115	6	5.2%	0.61×
exhaustion_short	either	4	1	25.0%	3.2× (n=4)
abs_veto_short	mechanical	112	13	11.6%	1.36×
abs_veto_short	ungated	221	15	6.8%	0.80×

Every gate with enough signal to mean anything sits between **0.5×** and **1.5×** **chance**. The three cells that look impressive — rgv_short at 4.7×, exhaustion at 3.2× and capitulation at 1.8× — are 2-of-5, 1-of-4 and 1-of-7 respectively, which is noise wearing a percentage sign. **None of our six gates is meaningfully more likely to be live before a big run than at any random moment of the week.** That is the honest null, stated as a mechanism rather than a P&L.

One caveat on that table, stated plainly: these are signal-minutes, not trades. grind_long showing 374 mechanical signal-minutes does not mean 374 entries — the live desk holds one position per slot and was benched by the router for much of the week, which is why it actually took 3 entries (two lots each) on the live book this week. The column is a measure of how often the condition was true, which is the right measure for this question.

Where the money actually is: the timing-decay ladder

This is the table I would pin on the wall. It takes hindsight completely for granted — you know which runs, you know the direction — and asks only **how early you enter**. Same 53 runs, same Lot A 1.5R + Lot B wide-chandelier exit, same tick-honest repricing, entry stepped backwards from the ignition bar.

ENTER THIS EARLY	N	NET \$	\$ / RUN	WIN %	STRIP BEST 3	US SESSION \$	OVERNIGHT \$
at ignition (0s)	53	+4,338.50	+81.9	74%	+2,984.00	+2,511.50	+1,827.00
30 seconds early	53	+2,487.50	+46.9	66%	+1,543.00	+1,334.00	+1,153.50
1 minute early	53	+1,696.50	+32.0	58%	+736.50	+839.00	+857.50

2 minutes early	52	+338.00	+6.5	44%	−628.50	−240.00	+578.00
3 minutes early	52	−307.50	−5.9	42%	−865.00	−905.00	+597.50
4 minutes early	52	−401.50	−7.7	33%	−1,338.00	−87.00	−314.50
5 minutes early	52	+268.00	+5.2	38%	−734.50	+193.00	+75.00
7 minutes early	52	−1,318.50	−25.4	29%	−2,178.00	−600.50	−718.00
9 minutes early	52	−475.50	−9.1	37%	−1,423.00	−139.50	−336.00
10 minutes early	52	+56.00	+1.1	38%	−918.00	+352.00	−296.00

THE WHOLE SECTION IN TWO SENTENCES

Enter at the ignition bar with perfect direction and these 54 runs pay **\$81.90 each**, 74% of them green. Enter **three minutes earlier** with exactly the same perfect direction and the same exit, and they pay **nothing** — and from there out to ten minutes early the curve wanders around zero, never once beating the two-minute mark, with strip-best-3 negative in every single row past two minutes.

Now put the gates on that ladder. When they do fire, they fire a **median of 7.0 minutes early mechanically and 8.0 minutes early ungated**. That is the deepest, deadest part of the curve. **The six gates are not losing at these runs because their exits are wrong or their thresholds are wrong. They are firing in a window where being right about the direction is worth minus twenty-five dollars a trade.**

Regime segmentation — scored on home turf, as the discipline requires

Nothing above should be read as one number sprayed across the whole tape, so here is the same population cut by regime at ignition. Buckets are keyed on the entry ATR and the 30-bar efficiency ratio measured on the last completed bar before the run, plus a 60-minute range-break test — not on the clock. This week's ATR quartiles across the sat-out runs came out at **12.3pt (p25)** and **25.4pt (p75)**, and those are the cut points.

	DEFINITION (AT THE LAST BAR BEFORE IGNITION)		
NORMAL-CHOP	ER < 0.25, ATR between 12.3 and 25.4pt	18	6 / 12
BUILDING	ER 0.25–0.45 — direction forming, not proven	10	5 / 5
DEAD-CHOP	ER < 0.25 and ATR below 12.3pt	10	4 / 6
VIOLENT-WHIPSAW	ER < 0.25 but ATR above 25.4pt — big range, no progress	9	2 / 7
TREND (no break)	ER ≥ 0.45, price still inside the prior hour's range	5	3 / 2
CLEAN-TREND	ER ≥ 0.45 and price outside the prior hour's range	1	1 / 0
UNKNOWN	insufficient bar history to classify	1	1 / 0

The first thing this table says is uncomfortable and important: **only one of the 54 big runs this week ignited out of a clean trend.** Twenty-eight of them ignited out of chop — normal or dead — and nine out of violent whipsaw. The runs are not the continuation of something we could see; they are ruptures in tape that looked like nothing.

REGIME	MECH FIRES	MECH NET \$	UNGATED FIRES	UNGATED NET \$	UNGATED WIN %
NORMAL-CHOP	3	−29.0	6	+278.0	50%
BUILDING	3	−315.0	7	+500.0	43%
DEAD-CHOP	0	0.0	3	+138.0	67%
VIOLENT-WHIPSAW	1	−130.0	2	−152.5	0%
TREND (no break)	0	0.0	1	+234.0	100%
CLEAN-TREND	0	0.0	0	0.0	—

Three honest reads. **(1)** Every regime bucket is between n=0 and n=7 — there is no bucket here with enough trades to make a policy from, and I am not going to pretend otherwise. **(2)** The one consistent sign in the table is the negative one: **violent whipsaw loses in both modes with a 0% win rate** — one mechanical fire at −\$130, two ungated fires at −\$152.50 — and it matches the blind ladder, where whipsaw (−\$265 on 9 runs) and building (−\$156 on 10) are the only two regimes still negative at one minute early. If any regime rule comes out of this section it is a bench rule, not an arm rule. **(3)** The gates never once fired in the two trend buckets, which is where a continuation gate should be earning — because the trend buckets contain six runs in total.

The R-target sweep — and what it says about the cheat-sheet guesses

The discipline says every R is a guess until it is swept, so I swept it. The gates gave me 7 and 19 fires, which is not a population you can sweep anything on — so I swept it on the blind one-minute-early book instead (perfect direction, 53 runs), with the quiet-tape clip switched **off**, because with the clip on the R target is a no-op on any entry under 22pt ATR and a sweep would be measuring nothing. Both lots at the same R so the number is readable.

R TARGET	ALL 53	WIN %	US SESSION (26)	OVERNIG HT (27)	DEAD-CHOP (10)	NORMAL-CHOP (17)	WHIPSA W (9)	BUILDIN G (10)
0.5R	−377	68%	−265	−112	−3	+190	−485	−236
0.75R	−300	62%	−289	−11	+68	+422	−537	−394
1.0R	+165	60%	+60	+105	+144	+537	−435	−326
1.5R	+1,164	58%	+634	+530	+292	+843	−233	−189
2.0R	+2,265	58%	+1,307	+958	+445	+1,248	−31	−55
2.5R	+3,374	58%	+1,985	+1,389	+597	+1,655	+172	+82
3.0R	+3,874	57%	+2,061	+1,813	+748	+2,060	−224	+218
4.0R	+5,491	55%	+3,261	+2,230	+1,052	+2,434	+31	+488
6.0R	+5,609	51%	+2,575	+3,034	+1,657	+1,908	+305	+64

What is PROVEN here, and what is not. This is a hindsight-direction population, so it cannot prove an R for a live gate — nothing here says "set abs_veto_short to 4R". What it can prove, because the entries are the same in every row and only the target moves, is the **shape of the exit that run tape rewards**, and that shape is unambiguous: **the curve rises monotonically from 0.5R to 4.0R and only flattens at 6.0R, with the win rate falling the whole way** (68% → 51%). There is no plateau anywhere below 4R and no turnover. That is a direct challenge to two of the operator cheat-sheet guesses:

CHEAT-SHEET GUESS	WHAT RUN TAPE SAYS	VERDICT ON THIS POPULATION
Faders: Lot A 0.5R / Lot B 1.5R	0.5R is the worst cell in the sweep (−\$377 on 53 runs with a 68% win rate — it wins often and loses money, the classic small-target signature). Even 0.75R is negative.	Wrong shape for run tape
Momentum: Lot A 1.5R / Lot B 2.5R	+\$1,164 and +\$3,374 — both positive, both a long way below the 4R cell.	Safe but leaving money
Trend: Lot A 2.5R / Lot B wide	+\$3,374 rising to +\$5,491 at 4R. The wide half is the correct instinct.	Directionally right

Robustness on the best cell, stated honestly. The 6.0R headline of +\$5,609 does not survive contact: leave out 08-05 and it drops to **+\$2,870**, i.e. a single

session carries \$2,739 of it. Leave-one-day-out across the five days gives \$5,265 / \$4,317 / \$2,870 / \$5,355 / \$4,629 — positive every time, but with a 49% swing. And every cell in the whipsaw column is worthless. **The plateau is at 2.5R-4.0R and it is real in sign but soft in size; the finding is "small targets are wrong on run tape", not "use 6R".**

Two live filters that are costing money on this population

The 55-second absorption veto — the test is fine, the wait is not

abs_veto's veto does two things at once and I was able to separate them, because mechanical and ungated fire from the identical bar and differ only by the 55-second wait plus the absorption test.

RUN	MOVE	UNGATED (ENTER NOW)	MECHANICAL (ENTER 55S LATER)	DELTA
08-03 11:46		−76	+97.0	−71.0
08-05 14:39		−78	+282.5	−146.0
08-05 19:51		−173	−62.0	vetoed
08-06 01:11		−68	−54.0	vetoed
08-06 13:17		−82	−53.0	vetoed
Total		—	+210.5	−217.0

The absorption test is **three-for-three** — every thrust it vetoed was a loser, \$169 saved. The 55-second wait is **nought-for-two** — it took the only two winners in the set and turned them into losers, \$596.50 given up. On the long side the veto blocked nothing and the wait changed nothing (both fires were losers either way, −\$55 and −\$51). **Saturday job: sweep the wait at 0s / 15s / 25s / 40s keeping the absorption test intact, on the full book rather than this population, and see whether the test survives without the delay.** This is thin — five fires — so it is a lead to test, not a change to ship.

The quiet-tape clip — a chop tool standing in the road

Every gate in data/exit_overrides.json carries the 08-02 quiet-tape clip: when entry ATR is below 22 points, Lot A banks a flat \$40 and Lot B banks the larger of 1.75R and \$60, and nothing else in the exit stack gets to run.

MEASURE	VALUE
Sat-out runs where the clip was active at entry (ATR < 22pt)	35 of 53
Blind one-minute-early book, clip ON (as deployed)	+\$1,696.50
Same entries, same exit, clip OFF	+\$2,898.50
Cost of the clip on run tape	-\$1,202.00

Two-thirds of these runs ignite from tape quiet enough to trigger the clip — which is the whole point of the census finding that runs come out of nothing. So on precisely the trades where a run is about to happen, the desk is wearing an exit that caps Lot A at forty dollars. It is doing what it was built to do; it was just never told that a 12-point-ATR tape can produce a 173-point run. **This is a regime carve-out question, not a delete question** — and it wants testing against the whole book, where the clip was justified, not against a population selected for big moves.

The verdict

HONEST NULL — AND HERE IS WHAT IT IS A NULL ABOUT

DON'T ARM. There is no version of "leave the six gates on and let them catch the misses" that makes money this week. Mechanically it is **-\$474 on 7 fires**, one winner. Ungated it is **+\$997.50 on 19 fires**, which is **+\$62.50** once you remove three trades, on a hindsight-selected population, in-sample, with no arming rule and no out-of-sample leg. Neither is a book you would put a lot behind.

But the null is specific, not general, and I want the specific part on the record because it points at next week's work:

- **It is not an exit problem.** Proven by the timing ladder: give the exits perfect entries at the ignition bar and the same 53 runs pay **+\$4,338.50**

at 74% win. The exits are fine. (The one exit finding that is real is the quiet-tape clip, worth \$1,202 on this tape.)

- **It is not a threshold problem.** Proven by the ungated book: strip every filter off all six gates and 36 of 53 runs still draw no fire at all, carrying \$6,524 of the \$9,780 ceiling and eleven of the fifteen biggest moves.
- **It is a timing-and-mechanism problem.** Proven twice: the gates fire at 0.5–1.5× chance in the pre-run windows, and when they do fire it is a median of 7–8 minutes early, where the blind ladder pays –\$25.40 a run. Our six gates are three continuation gates that need the move to have already happened and three fade gates that need a spent move to sell into. **A big run is neither of those things until it has started.**

Disposition — every lead touched in this movement

LEAD	VERDICT	WHY / WHAT WOULD REVIVE IT
"A better EXIT would have caught the misses"	REFUTED	Named test: the timing-decay ladder, n=52-53 per rung, tick-honest. Perfect entries at ignition pay +\$4,338.50 / 74% under the existing exits. The exits are not what is missing.
Arm all six gates mechanically at the misses	PARKED	–\$474 on 7 fires. Revive only if a mechanism changes — more n cannot save 7 fires from 53 runs, so what would revive it is a gate that can fire inside the first 90 seconds of a move.
Run all six ungated (filters off)	PARKED	+\$997.50 → +\$62.50 stripped of 3. Revive at n>60 fires across at least three weeks, with an out-of-sample leg and a regime rule attached; as it stands there is no policy here, only 19 trades.
abs_veto 55-second WAIT (not the test)	SHADOW	Best lead in the section. Test kept –\$169 of losers; wait cost –\$596.50 on two winners. Shadow a 0/15/25/40-second variant on the full book — free, riskless, and it settles a live knob.
Quiet-tape clip (ATR<22) on run tape	SHADOW	Active on 35 of 53 runs, costs \$1,202 with identical entries. Shadow a version that suspends the clip when the ignition context is a range-break, and read it against the whole book where the clip earns.
grind_long ext_hi 2.0 ceiling	PARKED	Blocks 6 of grind's 7 pre-run fires, worth + \$872.50 on this population. Revive with a full-book regime-conditional sweep of ext_hi (2.0 / 3.0 / 4.0 / off) split by regime — a sat-out-run population is structurally biased in favour of removing it, so this number alone must not move it.
rgv_short ATR≥20 floor	PARKED	Stands the gate down on 13 of 31 eligible runs; without it the gate still only fires once. Revive only alongside a stretch-condition rethink — the binding constraint is "never stretched above VWAP" (18 of 31), not the floor.

exhaustion_short as a run-catcher

PARKED

4 signal-samples in the entire week; 8 of its near-misses failed on "the flow DID move price", which is the gate correctly refusing to fade a real move. Revive as a bench instrument (that same clause is a clean "this is a real move" detector) rather than an entry.

VIOLENT-WHIPSAW as a place to arm anything

PARKED

0% win in both modes (3 fires, -\$282.50), and the only regime that is negative on the blind one-minute ladder (-\$265 on 9 runs). Thin, so not refuted — but every sign points the same way. Revive as a bench rule with $n > 25$.

THE ONE STONE STILL UNTURNED

Everything in this movement tested entries **before** ignition, because that is what "could we have caught it" normally means. The decay ladder says the highest-paying rung by a distance is the one at zero — and the second-highest is **thirty seconds in**, at +\$2,487.50 / +\$46.90 a run / 66% green. We have never built a gate whose job is to arm on the first thirty seconds of a move that has already started, on 5-second bars, with no attempt to predict anything. Every gate we own either wants five minutes of history (abs_veto, grind) or wants the move to be over (rgv, capitulation, exhaustion). The desk has already banked "predicting run STARTS" as a null on 22.6 million ticks (carried in as a prior, not re-derived here); this movement adds the complementary fact that **being late by ninety seconds is nearly as bad as being wrong**. The unturned stone is the narrow band in between — and it is exactly the shape NIPC found on the news window. That is Movement 3's job.

M3 GREENFIELD — three hunts from a blank sheet, one survivor, and the panel that tried to kill it

Garrath — the census left **\$9,917 of hindsight money** sitting in 54 runs we sat out. So this week I did the opposite of tuning: I threw away every gate we own and asked three separate hunts to invent an entry from

scratch for three of the clusters. **Two came back with nothing.** The third came back with something almost insultingly simple — show up every five minutes during the US open, face the way the last fifteen minutes went, and give it a stop twice as wide as the desk's house style. That one made money on **36 trading days**, including **19 days it was never fitted on.** I then handed it to three independent skeptics with a mandate to destroy it. **One killed it. Two could not.** It goes to the **shadow book** — not live, not this month, and never on 17 in-sample days. And along the way I proved something about the \$9,917 that I think matters more than any of the three hunts: **the money in those runs is real, and it has a half-life of about seven minutes.**

READ THIS BEFORE ANY NUMBER BELOW — THE FEE WAS WRONG

Every greenfield backtest in this Movement was originally priced at **~\$5 a round trip. The real MNQ fee is \$1.50 a round trip** — venue truth, all 487 closed trades carry fees_usd=1.50. The correction is exactly linear: $net@1.50 = net@5.00 + 3.50 \times n$. It is a **gift of \$3.50 for every trade a strategy ever made**, so it helps everything — winners and losers alike. Two consequences you must hold onto: (1) **the survivor was under-stated, not over-stated**; (2) **a loser that the correction rescues was never killed by cost in the first place**, and I flag the one row in this Movement where that happens. The chop-scalp lab is the only hunt that was priced correctly from the first line.

The re-price, line by line — which numbers moved and by how much

NUMBER AS ORIGINALLY PUBLISHED	N	@ \$5/RT	@ \$1.50/RT (TRUTH)	CHANGE	DOES THE VERDICT CHANGE?
OPEN RIDER headline (author's arithmetic, corrected in place)	129	+\$4,169	+\$4,620	+\$451	No — it gets better
OPEN RIDER, execution-skeptic's	129	+\$4,688	+\$5,139	+\$451	No

own re-derivation, same 17 days					
OPEN RIDER, duty-cycle skeptic's re-derivation, same 17 days	129	+\$4,577	+\$5,028	+\$451	No
OPEN RIDER, robustness skeptic's from-scratch harness, same 17 days	121	+\$3,483	+\$3,907	+\$424	No
VACUUM-BREAK walk-forward policy (the honest deployable answer)	91	-\$443	-\$125	+\$318	No — still a loser
VACUUM-BREAK best static cell (T=3.0 / RM=2.0)	70	+\$824	+\$1,069	+\$245	No — killed by strip-3, not by fees
FLOW-TRAP, the worst cell	220	-\$1,492	-\$722	+\$770	No
△ FLOW-TRAP T=3.0 / RM=2.0 — the one row that flips sign	74	-\$105	+\$154	+\$259	The "24 of 24 configs negative" line no longer holds exactly — at least this one flips. (The other 17 cells were published without trade counts, so I cannot re-price them individually and I will not guess.)
ABSORPTION-SHELF, all three published cells	210 / 146 / 54	-\$1,191 / -\$600 / -\$451	-\$456 / -\$89 / -\$262	—	No — still negative on both sides
FLOW-LED census rule traded directly, the cells that actually catch the cluster (Z 0.75 / 1.00)	129 / 116	-\$1,646 / -\$1,708	-\$1,194 / -\$1,302	+\$452 / +\$406	No — the catch/profit inversion is untouched
FLOW-BREAK IGNITION (shadow) / FLOW-TRAP FADE (shadow)	39 / 44	+\$1,830 / +\$1,870	+\$1,967 / +\$2,024	+\$137 / +\$154	No — both still SHADOW on 5 days of tape
CHOP-DAY SCALP LAB (whole section)	—	already priced at \$1.50/RT throughout		—	No change needed

The honest summary of the fee error: it cost us nothing in decisions. The survivor was under-sold by about \$450, and every loser I could re-price stays a loser except one 74-trade FLOW-TRAP cell that limps to +\$2.08 a trade — and that cell was killed by a cost-free test (its raw forward return is below the tape's own drift), which no fee can rescue. **But it must never happen again:** FEE, VPP = 5.0, 2.0 and VPP, FEE = 2.0, 1.5 look

identical at a glance and are reversed. Grep the fee constant in any harness before you quote its number.

STEP 1 — the oracle: are the sat-out runs actually money, or just hindsight?

Everything in this Movement rests on one premise: **that the 54 runs we sat out were worth real money if only we had found their start**. If they are not — if a real stop and a real target eat them — then the whole greenfield exercise is chasing a mirage and the problem is the exit, not the entry. So I tested it directly rather than assuming it.

THE TEST

Take every sat-out run in the census. Enter **1 lot at the run's start minute, on the run's side** (this is the oracle — it is the entry timing no signal has). Then exit it **honestly**: a real stop and a real 2R target racing each other for first touch on the **raw trade prints** from the lake, stop wins ties, 45-minute cap, **\$1.50/RT**. No MFE anywhere. If the money is real, this books a large fraction of the hindsight ceiling. If the exits are the problem, it does not.

The answer: the money is real, and it is the ENTRY that is missing

52 of the 54 sat-out runs had the tick coverage to be tested — the two that did not are 08-02 23:52 and 08-06 22:00, both sitting on the 22:00 UTC reopen where the tape restarts. Entering each of the other 52 at its start minute, on its side, and letting a real stop and a real target race:

ORACLE ENTRY, HONEST EXIT	N	NET @ \$1.50/RT	\$/TRADE	GREEN	HINDSIGHT CEILING	CONVERTE D
	52	+\$6,846	+\$131.65	52 / 52	\$9,622	71%

stop
2.0×ATR1m,
target 2R

These runs are worth about \$6,850 through a real exit with real fees. Not \$9,917 of fantasy — seven-tenths of it, banked by a stop and a target racing tick by tick with the stop winning ties. Exit mix: 45 targets, 7 time-outs, **zero stops**.

Say plainly what this is and is not. It is an oracle: it is handed the run's start minute and the run's direction, which is exactly the information no signal has. That is why nothing stops out and why it is an upper bound, not a plan. What it settles is the question it was built for — **with the right entry, the exit machinery is not the constraint. The entry is the whole problem.**

And then the oracle produced something I did not go looking for. Hold the entry perfect and vary only the exit geometry:

SAME PERFECT ENTRIES, DIFFERENT GEOMETRY	N	NET	\$/TRADE	WIN%	STOPPED OUT	CONVERTE D
stop 1.0×ATR — the desk's house style, 2R	52	+\$2,321	+\$44.64	73.1%	14 of 52	24%
stop 1.0×ATR, 2.5R	52	+\$3,064	+\$58.92	73.1%	14 of 52	32%
stop 1.5×ATR, 2R	52	+\$5,099	+\$98.06	90.4%	5 of 52	53%
stop 2.0×ATR, 2R	52	+\$6,846	+\$131.65	100%	0 of 52	71%
stop 2.0×ATR, 3R	52	+\$6,826	+\$131.26	96.2%	1 of 52	71%
stop 2.5×ATR, 2.5R	52	+\$6,917	+\$133.02	96.2%	0 of 52	72%

Even when you know exactly where the run starts and which way it goes, the desk's 1-ATR stop throws away two-thirds of it — and gets shaken out of **14 of 52 runs it entered perfectly**. Widen to 2 ATR and the shake-outs go to zero. This is a third, completely independent arrival at the same conclusion the OPEN RIDER's grid and the robustness skeptic both reached from different directions. Note the plateau: 2.0/2R, 2.0/3R and 2.5/2.5R are within \$91 of each other, so this is a region, not a spike.

THE UNCOMFORTABLE COROLLARY — THE 10 RUNS WE DID CATCH

Run the same oracle on the 10 runs a live gate actually caught: an honest $2 \times \text{ATR}/2R$ exit banks **+\$1,630 of their \$2,416 ceiling — 67%**. What the gates actually booked on those same runs was **+\$156 — 6%**. That \$1,474 gap is entry timing and stop width together, not exit logic alone, so do not read it as a pure exit indictment. But it does say the thing this Movement keeps saying from every angle: **we are turning up late and standing too close to the door**. Across all 67 runs the oracle converts \$9,007 of \$12,983 — 69%.

★ And the number that explains this entire Movement: the money has a half-life of about seven minutes

Same runs, same side, same honest exit ($2.0 \times \text{ATR}$ stop, $2R$ target, first-touch race on raw prints, 45-minute cap, \$1.50/RT). The **only** thing that changes is how late you turn up.

ENTRY, RELATIVE TO THE RUN'S START	N	NET	\$/TRADE	GREEN	STOP / TARGET / TIME
on the minute (the oracle)	52	+\$6,846	+\$131.65	100%	0 / 45 / 7
2 minutes late	52	+\$6,041	+\$116.17	90.4%	3 / 43 / 6
5 minutes late	52	+\$4,090	+\$78.66	78.8%	7 / 31 / 14
10 minutes late	52	+\$944	+\$18.14	51.9%	21 / 19 / 12
15 minutes late	52	-\$2,068	-\$39.76	21.2%	37 / 5 / 10

Roughly \$600 a minute walks out of the door for the first ten minutes, and by fifteen you are paying to be there. The win rate falls $100\% \rightarrow 90\% \rightarrow 79\% \rightarrow 52\% \rightarrow 21\%$; the stop-out count goes $0 \rightarrow 3 \rightarrow 7 \rightarrow 21 \rightarrow 37$. Same runs. Same exit. Same side. **Just later.**

THIS ONE TABLE EXPLAINS EVERY RESULT IN THIS MOVEMENT

Why the FLOW-LED hunt could not build anything: its own tape study found that price only confirms the run at **+3, +4 and +5 minutes**, 25–45% of the move in. Anything that waits for price confirmation is buying the back half at 60–70 cents on the dollar — before you have paid a spread.

Why VACUUM-BREAK missed the three biggest runs: the shelf-break trigger arrived far too late on the 08-06 buy-climaxes, and the retrace trigger never fired inside its 10-minute window at all. By then there was nothing left to take.

Why the OPEN RIDER is the only survivor: it does not wait for anything. It has no trigger — the clock is the trigger — so its lateness is structurally near zero: it is already in the market, facing the right way, when the run starts. That is not cleverness, it is the removal of a delay, and this table is why removing the delay was worth more than every footprint we tried to invent.

And it re-frames the whole census gap. "We sat out 54 runs worth \$9,917" is really "**we were between three and fifteen minutes late, and lateness is the most expensive thing on this desk.**"

STEP 2 — cluster one: VACUUM. Three inventions, 90+ cells, no survivor.

What VACUUM is supposed to be: the tape is full of buyers, somebody eats all of it, and then the floor gives way and price falls 80 points. Eight such runs in the census. The hunt built three separate signals against them.

INVENTION	THE IDEA IN ONE LINE	BEST HONEST RESULT	VERDICT
FLOW-TRAP	the burst failed and gave everything back — the people who did it are underwater, sell them	12 of 12 cells losing at the old fee; at the correct one, one cell limps to +\$154 on 74 trades	REFUTED — by a cost-free test
ABSORPTION-SHELF	all that effort and price did not move — it was eaten, take the other side now	12 of 12 cells losing, and losing on both sides; every cell whose trade count was published stays negative re-priced	REFUTED
VACUUM-BREAK	wait for the floor under the trapped side to break — that break is the vacuum	n=91, −\$125, 33% win, −\$1.37/trade	PARKED

The single most useful number in the whole cluster is a base rate. Across the tick tape there are **1,451 minutes** where the flow divergence VACUUM is defined by fires at $|z| \geq 1$. Eight of them preceded a big run. **That is 0.55%**. Any signal keyed on the flow tell alone fires roughly 300 times a day to find one of these. VACUUM is a description of what happened, not a footprint you can see coming — and that, not the P&L, is the finding.

Why VACUUM-BREAK looked alive and then wasn't — four tests, in order

TEST	WHAT IT SHOWED	RESULT
Does the money need n to collapse?	correlation between net P&L and falling trade count = +0.77 . 28 of 77 cells positive, mean cell $-\$468$. Positive cells' median $n=71$; negative cells' median $n=133$. At $T=3.0$ the target knob goes $+\$576 \rightarrow +\$824 \rightarrow -\$147 \rightarrow -\706 in four steps (cells shown at the published \$5 fee — re-pricing lifts every cell but does not change the ridge, because the lift is proportional to each cell's trade count). A knob you cannot be half a step wrong on is not a knob.	FAILED
Strip the three best trades (re-priced at \$1.50/RT)	$+\$1,069 \rightarrow \text{strip-1} +\$776 \rightarrow \text{strip-2} +\$529 \rightarrow \text{strip-3} +\$342 \rightarrow \text{strip-5} +\26 . Three trades are 68% of it; five trades are 98% of it. All six regime buckets go negative after strip-3. (At the wrong \$5 fee this looked worse still — strip-3 $+\$107$, strip-5 $-\$202$ — so the correction softens the kill without changing it: a strategy whose entire net is five trades is not deployable at any fee.)	FAILED
Long/short symmetry (re-priced)	long $+\\$1,083$ on 38, short $-\\14 on 32 — in a week MNQ rose 1,300 points. The short side is a flat zero at best. The always-long control on the same 70 timestamps makes \$367 less than the signal — and that gap is fee-invariant, because both sides get the identical correction. \$367 is smaller than the three trades strip-3 already removed. And six of the eight VACUUM runs are DOWN runs — the side that does not work.	FAILED
Regime $\rightarrow$ config policy, walk-forwarded honestly	pick each regime's best config on four days, trade it on the fifth, rotate: $n=91$, $-\\$125$ at the correct fee, $-\\$1.37$ a trade, 33% win , strip-3 $-\$801$. Best blanket in-sample said $+\$1,069$. The act of choosing cost \$1,194. That gap is what the walk-forward discipline exists to expose.	FAILED
Placebo null	$+\$824$ is $z=+2.31$, $p=0.013$ — but the grid has 77 cells, so $77 \times 0.013 \approx 1$ cell expected to clear that bar by luck. Exactly one did. The placebo passed on the cell that was selected for passing it.	passed, worthless

The operator's own hypothesis — "narrow to the biggest runs and a stronger footprint may appear" — runs backwards here. The three biggest VACUUM runs are missed by every single configuration of all three inventions. The only run any honest config catches is the second-smallest. Narrowing does not sharpen the footprint; it removes the only hits.

THE WALL THIS CLUSTER RAN INTO, AND IT IS INFRASTRUCTURE, NOT RESEARCH

VACUUM is defined by aggressor flow, and `capture.db` keeps aggressor-tagged ticks for **5 trading days**. I tried to buy the other twelve by rebuilding signed flow from 5s bars: bar volume reconstructs faithfully (correlation 0.96), but the **sign is a coin toss** — 50.7% agreement, z-score correlation 0.066. **There is no out-of-sample leg available for any flow-based idea on this desk today.** At ~1.35M ticks a day we need roughly 15 sessions before VACUUM work can carry one. The action is not research, it is **keep archiving ticks**.

STEP 3 — cluster two: FLOW-LED. The assignment itself failed, and that is the finding.

This hunt reported back that it could not build a signal for FLOW-LED, and then did something better: it worked out **why**. The answer is that **FLOW-LED is not a family of runs. It is a labelling accident.** The census hangs the label on one minute of aggressor flow before the run — and at census level "full" that is **2 runs out of 67**.

The tape at the target runs — there is nothing there to trigger on

RUN	WHAT THE FLOW ACTUALLY DID
08-05 08:07 DN -74	label minute $z=-1.75$, then flips to +1.42 the very next minute while price keeps falling. Price only confirms at +3 min, ~27pt late.
08-06 12:52 DN -68	flow had been positive for 7 minutes running , one minute flips to -1.87 , label assigned, and it flips straight back the next bar.

08-07 19:43 UP +74

label minute +1.17, **immediately reverses** to −1.00.

08-04 01:40 UP +128

flow is NEGATIVE at the start of a +128pt UP run. The census called it FLOW-LED; an independent re-derivation called it UNCLASS.

At every one of the four run starts the flow's run-length is **one or two minutes** and it reverses on the next bar. A "sustained one-sided aggression" trigger — the only mechanically sensible reading of "flow-led" — cannot fire, because the sustained aggression does not exist.

Generalised over all 62 big runs in the tick era, the median big MNQ run starts AGAINST the flow: 1-min flow imbalance −0.016 at run start (vs ± 0.078 on all minutes), 3-min net aggressor **−100 contracts**, 1-min displacement −3.25pt, flow persistence −1. What they do start with is a volume expansion (volume z +0.25 vs −0.21) after a counter-move. **That is the exact opposite of the FLOW-LED premise.**

The kill: trade the census's own FLOW-LED rule, exactly as written

The most charitable possible version of the assignment. Read the last two columns against the money column.

FLOW-Z THRESHOLD	N	NET @ \$1.50/ RT	\$/TRADE	STRIP-3 \$/TR	CATCH OF THE 4 TARGET RUNS
0.75	129	−\$1,194	−\$9.26	−\$14.10	2 / 4
1.00	116	−\$1,302	−\$11.22	−\$16.92	2 / 4
1.25	111	−\$434	−\$3.91	−\$9.89	2 / 4
1.50	94	+\$687	+\$7.31	+\$0.65	1 / 4
2.00	83	−\$342	−\$4.12	−\$11.65	0 / 4
2.50	70	−\$154	−\$2.20	−\$10.75	0 / 4

Catch-rate and profitability move in opposite directions. The settings that actually catch the cluster are the settings that lose the most, over 111–129 trades — this is not thin-n. That is not a threshold needing tuning; **it is**

a mechanism with the wrong sign. Note the re-price does not touch this: the losing cells stay losing by \$434-\$1,302.

TWO SPIN-OFFS WORTH KEEPING — NEITHER ANSWERS THE ASSIGNMENT

Digging through the flow data threw off two signals that beat their nulls at $p < 0.01$. **FLOW-BREAK IGNITION** (sustained flow + a 30-min range break): **+\$1,967 on 39 trades**, 56.4% win, 45 of 45 sweep cells green — but its out-of-sample half collapses to a coin flip and it catches **0 of 4** target runs. **FLOW-TRAP FADE** (aggressors lean hard and fail to move price): **+\$2,024 on 44 trades**, all five days green, leave-one-day-out flat as a board, and its OOS leg holds — the sturdier of the two, also **0 of 4**. Both are 5-day findings wearing a confident number, and FLOW-TRAP FADE overlaps conceptually with our own `abs_veto` absorption work — **check for redundancy before shadowing a duplicate**. Both filed SHADOW.

The structural point the hunt makes, and I agree with it: the census's cluster labels should carry a persistence requirement. One minute of $|z| \geq 1$ currently assigns a run to FLOW-LED, that minute reverses on the next bar, and an independent re-derivation disagrees with the label 3 times out of 4. **The cluster column is reporting noise as structure.** Either require `persist` ≥ 3 for the label, or retire FLOW-LED and VACUUM back into one "flow event" bucket.

STEP 4 — cluster three: OPEN/NEWS. The clever ones died. The control lived.

First, the thing that makes this cluster different and is the reason it has a survivor at all: **OPEN/NEWS is not detected from the tape. It is a clock label.** Any big run whose window starts between 13:00 and 15:00 UTC gets it, before anything looks at flow or book. Which means two things — nothing rare is happening, and **the whole 5-second bar tape is usable**, not just the 5 days of ticks.

WHERE A 15-MINUTE SLOT SITS	SLOTS	OF WHICH "BIG RUN" ($\geq 65\text{PT}$)	RATE
13:00-15:00 UTC	134	52	39%

Four out of every ten quarter-hours in that window is a big run. So "catch the OPEN/NEWS runs" honestly reduces to "trade 13:00–15:00 UTC directionally" — which is a much better brief than it sounds, because the base rate is high and the sample is not starved.

The two clever inventions, and how they died

INVENTION	THE PITCH	WHAT KILLED IT	VERDICT
COIL-CRACK	"the last half hour went nowhere and then it snapped — ride the ignition out of balance"	Direct ablation , the only test that decides a filter. With the efficiency gate: +\$2,049 / n=51. Without it: +\$2,238 / n=58. It costs money and sample. The close-location confirm was worse — at three different settings it returned +\$830 / +\$830 / +\$832, i.e. it selected nothing at all.	REFUTED as a filter
OPEN EXHAUSTION FADE	"a straight line for thirty minutes into the open is a finished move — take the other side"	51 configurations, every one negative or zero. And the reason is the path, not the parameters: median excursion against a fade is 4.76 ATR against 3.00 ATR in its favour. It is adverse from the first tick, so every stop you can afford is hit first. Caught 0 of 14 target runs.	REFUTED

The seductive part, and the lesson. The exhaustion fade has the strongest raw signal found anywhere in this Movement — day-demeaned forward return +54.2 pt at 65.8% win, both sides, monotone in the threshold. **And it cannot be traded.** A 15-minute edge delivered only after a 4.8-ATR excursion against you is not an edge you own; it is an edge your stop pays for. The two obvious reformulations (widen the stop to 3–4 ATR, delay entry by 10 minutes) were both tested and both fail, which closes it.

THE OPEN RIDER — the survivor, and the whole spec fits in a box

This was built as the control that was supposed to lose to the clever ones. It didn't.

ELEMENT	RULE
WINDOW	13:00–15:00 UTC
CADENCE	every 5 minutes; if flat, take a trade
TRIGGER	none. The clock is the trigger. No threshold, no efficiency dial, no book reading.
DIRECTION	sign of the last 15 minutes' move
ENTRY	market at the close of that 5-second bar (already past the level — no limit-fill optimism)
STOP	2.0 × ATR1m — twice the desk's house style. Mean distance 54.6 pt = \$109 of risk a trade.
TARGET	2R
TIME CAP	45 minutes, then out at market
RE-ENTRY	blocked until the open trade is closed

Read it as a second desk for two hours a day, not a seventh gate. It is in the market essentially 100% of its window — a five-minute cadence with 45-minute caps means the next trade starts the moment the last one ends. Occupancy measured under honest execution is **101% of the two hours, 7.9 trades a day.**

What it did — and the one number that matters most is the 19 days it never saw

TAPE	DAYS	N	NET @ \$1.50/RT	WIN%	\$/TRADE	R/TRADE	
DESIGN tape (07-16 → 08-07), the days it was built on		17	121	+\$3,907	47.1%	+\$32.29	+0.299R
TRUE OUT-OF- SAMPLE (V5 archive, 06-19 → 07-15) — tape that did not		19	115	+\$3,526	47.0%	+\$30.66	+0.277R

exist for the study						
ALL	36	236	+\$7,433	47.0%	+\$31.50	+0.288R

+\$30.66 a trade out of sample against +\$32.29 in sample — a 5% degradation, on 19 days no parameter was ever fitted on. Green days 15 of 19 OOS against 12 of 17 in-sample. That is the opposite of what a curve-fit produces, and it is the strongest single fact in this Movement.

The robustness battery, on all 36 days

TEST	RESULT	READ
Strip the best day / best 3 days	+\$6,268 / +\$4,486	not one week holding it up
Strip best 1 / 3 / 5 / 10 trades	+\$6,827 / +\$5,901 / +\$5,048 / +\$3,160	concentrated but survives
Leave-one-day-out, all 36 drops	positive 36 / 36, +\$6,268 ... +\$8,334	no single day carries it
Quarters of the tape, scored independently	+\$2,221 / +\$1,059 / +\$868 / +\$3,285 — all four green	no dead stretch
Day-bootstrap, 5,000 draws	5th percentile +\$3,743, $P(\text{net} \leq 0) = 0.001$	—
Day-clustered t-stat (36 clusters)	+3.79	clustered properly, not by overlapping samples
Placebo — same bars, random direction, 200 draws	placebo mean +\$1,129 (sd \$2,432); actual $z = +2.59$, $p = 0.005$	⚠ note the placebo mean is not zero — the wide-stop/2R shape is worth something on its own; the direction rule is worth \$6,300 more
Reversed (fade instead of ride)	−\$3,885	ride and fade are \$11,318 apart on the same bars
Reverse walk-forward: fit 48 cells on V5, score blind on the design tape	48 / 48 cells holdout-positive; the DEV-best cell is the published spec exactly, #1 of 48	the cleanest possible answer to "you tuned on the tape you scored on"
Cost sensitivity	a full point of adverse entry slippage and the wrong \$5 fee still leaves +\$5,297	cost is not where this dies

The stop is the lever, not the entry. Every cell with a stop of $1.0 \times \text{ATR}$ or narrower is flat-to-negative; **all 27 cells with a stop of $2.0 \times \text{ATR}$ or wider are positive.** Mean stop 57.7 pt against a mean 1-min ATR of 28.8 pt. **The desk's ~1-ATR house stop is the wrong instrument for the US open** — that is transferable to the live gates today and costs nothing to check.

Participation — it is on the runs, including the biggest ones

CUT	AUTHOR	INDEPENDENT RE-DERIVATION
The 14 sat-out OPEN/NEWS runs (census week)	11 / 14	—
Top-15 biggest OPEN/NEWS runs, 17-day tape	15 / 15	15 / 15
All OPEN/NEWS runs, 17-day tape	62 / 78 = 79%	67 / 87 = 77%
36-day tape, top-15 / top-25 / top-50	—	93% / 92% / 92%

The three misses in the census week were all runs that started while the RIDER was already holding the other way — **the cooldown missed them, not the signal**. And your narrowing hypothesis — that the biggest runs carry a sharper footprint — is **CONFIRMED here**, which is the exact opposite of what VACUUM found: top-5 caught 80-100%, top-8 caught 75-88%, all-14 caught 79%.

STEP 5 — THE SKEPTIC PANEL. Three independent attacks, and the vote is 2-1 to survive.

A survivor that only its own author has tested is not a survivor. So the OPEN RIDER went to three separate adversarial lenses, each with a mandate to **refute**, each re-deriving everything from raw 5-second bars through `lake.connect()` rather than reusing the author's harness. **Here is the whole vote, including the one that killed it — I am leading with that one on purpose.**

LENS	VERDICT	THE ONE LINE
DUTY-CYCLE / PLACEBO	REFUTED	it is a levered long position on how far the open window happened to move, and on the only days it was never fitted to it sits inside its own placebo
ROBUSTNESS	HOLDS	found 19 more days of unseen tape and it reproduces at 95% of its rate there; two of the author's own stated wounds do not replicate
COST & EXECUTION	HOLDS	every sub-attack was run; the execution bill is \$5.18 a trade against a \$52 gross — it dies on plumbing, not on money

X
LENS 1 — DUTY-CYCLE:
REFUTED

The bookkeeping first. The trade count reproduces to the trade — n=129 exactly on the author's 17 days. Re-priced at the house \$1.50/RT the honest net is **+\$5,028**, not \$4,169 (**the fee error HELPED it**), and **+\$4,870** after crossing the measured 0.615pt entry spread. Also: **the claim's "102 trades" is wrong** — the author's own text and this re-derivation both say 129. Then the controls kill it three ways.

KILL 1 — the drift decomposition. The intercept is zero. Regress each day's P&L on that day's absolute 13:00-15:00 travel, 34 days:

SAMPLE	P&L ON A ZERO-DRIFT DAY	T	\$ PER POINT OF DRIFT	T
All 34 days	−\$23	−0.22	\$1.237/pt	+2.75
17 fitted days	+\$74	+0.45	\$1.091/pt	+1.58
17 unseen days	−\$86	−0.60	\$1.207/pt	+1.91

The only statistically non-zero thing about the OPEN RIDER is **its loading on how far the window happened to move**. Mean absolute drift 180.5pt × \$1.237 = **\$200 a day — exactly the actual \$200 a day**. On a day the open ends where it started, expected P&L is nothing.

KILL 2 — the machinery is decorative. A control that makes **one** direction decision per day — the sign of the window's first 15 minutes, held for the whole two hours, no 5-minute cadence, no rolling lookback — earns **+\$6,005 of the RIDER's +\$6,810 over 34 days. That is 88%**. The entire cadence apparatus is worth **\$3.31 a trade**, less than the \$1.23 a trade you pay just to cross the spread. And a hindsight bet that knows the day's answer in advance earns \$10,778 — so the OPEN RIDER is **a 63%-efficient estimator of "which way is today going"**.

KILL 3 — out of sample it is inside the null. lake.connect() unions the V5 archive, which holds 17 more sessions with full window coverage (06-19 → 07-15) at higher volatility (ATR1m 31.9 vs 29.2) that the author never used:

SAMPLE	N	NET	WIN%	\$/TRADE	DAILY T	PLACEBO PERCENTILE
--------	---	-----	------	----------	---------	--------------------

17 fitted days	129	+\$5,028	48.8%	+\$38.98	3.31	99.4	0.0065
17 unseen V5 days	114	+\$1,782	41.2%	+\$15.63	0.94	83.2	0.168

Inside the null — which is this lens's stated kill condition. On the fitted days it is at the 99.4th percentile ($p=0.0065$) and at the 98.8th under a day-level sign-flip null ($p=0.0125$), **so the significance lives entirely in the design sample.** Other controls on 34 days: always-long $-\$4,401$, always-short $+\$3,233$, reversed $-\$3,798$, yesterday's-direction $-\$3,699$. Cumulative window drift over the 34 days is **$-\mathbf{1,389.5}$ pt**, so a single held short lot books $\$2,728$ with no strategy at all. Concentration: **8 of 34 days carry 82% of the net.**

What this lens concedes: it is not worthless. The clock observation is real ($\$1.237$ per point is stable at $1.207 / 1.091$ across both independent halves) and wide stops genuinely beat the desk's 1-ATR standard. **But the OPEN RIDER is an exposure to open-window realised directionality, not an edge** — and nothing in the spec forecasts that quantity or tells you when it stops paying.

Its leads: **REFUTED** ODR as a tradeable edge (test above). **PARKED** "13:00–15:00 is where MNQ travels" — revive with an out-of-sample forecaster of next-window magnitude, not direction. **SHADOW** wide stops $\geq 2 \times \text{ATR}_{1m}$ in the open window — revive by re-measuring the plateau on a signal whose intercept is not zero.

✓ LENS 2 — ROBUSTNESS: **HOLDS**

This lens set out to prove the thing was 17 in-sample days of the already-known 13:30–14:45 window. It rebuilt the harness from scratch off raw 5s bars, read nothing from the author's working directory, **and found 19 trading days of tape the author had left on the table.**

The killer test: on that unseen leg the RIDER does **$+\$3,526 / n=115 / +\30.66 a trade / $+0.277R$** against the design tape's $+\$32.29 / +0.299R$. **A 5% degradation, 15 of 19 days green.** Reverse walk-forward — fit 48

cells on V5, score blind on the design tape — picks **the published cell as #1 of 48**, and all 48 cells are holdout-positive.

Its own hypothesis, refuted by the control the author never ran. If a wide stop plus a 2R target were just harvesting right-skew, it would print money all day. It does not:

BLOCK (SAME MECHANICS)	N	\$/TRADE	R/TRADE
13:00-15:00 (the spec)	236	+\$31.50	+0.288
ASIA 00-07	678	+\$0.38	—
US-AFTER 15-21	545	+\$0.31	—
LONDON 07-13	645	-\$1.89	—

The spec window ranks #1 of 11 two-hour blocks and is 4× the next-best. The clock is load-bearing; the payoff shape is not a free lunch. And **13:00-13:30 — entirely outside the desk's validated window — is + \$1,427 on 74 trades at +0.390R, the best R of any slice.**

The author's own stated wound also fails. He called the day-demeaning test a failure at $t=-5.35$. That is **pseudo-replication**: 4,224 overlapping 1-minute samples with a 45-minute forward window are perhaps 36 independent days. Clustered by DAY the t-stats are **-0.08 / -0.33 / -0.57** — nowhere near significant. Corroborated twice: forcing **equal longs and shorts within every day** (which removes day direction by construction) leaves **+\$30.89 a trade, unchanged** from +\$31.50; and the literal day-drift bet — take the direction at 13:00 and hold to 15:00 — **LOSES \$3,259 over the 36 days** while the RIDER makes +\$7,433. If it were only harvesting day direction, the pure day-direction bet would not be the losing side of a \$10,692 gap.

THREE THINGS THIS LENS DID BREAK — NONE FATAL, ALL ACTIONABLE

1. The "plateau at every cadence" claim is PARTIALLY REFUTED. On the unseen tape, cadence-3 is positive in only **4 of 16 cells** against 13/16 at 5 minutes and 15/16 at 10. The plateau is real in stop width (27/27 positive at stop ≥ 2.0) and **one-sided in cadence** — slower works, faster does not. The chosen cell is on the safe side of a cliff the author did not know was there. **Strike cadence-3 from the language.**

2. The window's right edge is dead weight. 14:45–15:00 is **−\$446** across all days and **−\$453** on the unseen leg alone. Cut it — and note it independently re-confirms our own 14:45 boundary.

3. The headline does not reproduce precisely. A from-scratch harness gets $n=121$ / $+\$3,483$ at $\$5/\text{RT}$ on the identical 17 days — **16% below** the published $+\$4,169$. Same sign, same order, but treat the headline as $\pm 20\%$, never as a precise figure.

Honest weaknesses this lens still names: $\$115$ of risk a trade (4–5× the live gates); the top 10 of 236 trades carry 57% of the net; and **36 days has never shown this thing a bad month** — the worst day is $-\$901$ and that is also the maximum drawdown, because the curve never had a losing run. Verdict: **SHADOW, upheld and strengthened.** Promote only on the author's own bar — $+0.15R$ over 200 shadow trades including a non-trending open week — plus the 14:45 cut and striking cadence-3.

✓ LENS 3 — COST & EXECUTION: **HOLDS**

Every sub-attack was run and none of them kills the P&L. Re-priced at the real fee and re-run tick-honest with the spread crossed, the stop filled on the tape's own flush and the time-cap crossed out: **+\$4,435 net on 95 sequential trades over the 12 execution-verifiable days — +\$46.69 a trade, +0.386R, 51.6% win, 11 of 12 days green.**

(0) The factual correction. The claim's "102 trades" is wrong: $+\$4,169 = 129 \times \32.32 ; the 102 belongs to a different row of the report (the thresholding comparison, $+\$4,142$). An independent lake re-derivation returns **$n=129$ exactly**, net $+\$4,688$ at $\$5/\text{RT}$ — $\$519$ above the claim, because `rider.py` indexes bars positionally and 07-22 is missing 6 bars; this harness keyed everything by timestamp.

(a) The re-price. Author's headline corrected in place = **+\$4,620 on 129 trades**. Independent re-derivation = **+\$5,139 on 129** (DEV 90 trades +

\$4,235 / HOLDOUT 39 trades +\$904). **The \$5/RT was over-charging it \$451.**

(b) Slippage from REAL fills, not assumed. 367 live fills ASOF-joined to L1 quotes at the fill millisecond:

MEASUREMENT	N	MEAN SLIP VS MID	\$ / LOT
market ENTRY	367	+0.389 pt	\$0.78
(against a measured half-spread of)	367	0.341 pt	—
STOP exits, book-wide	209	+1.550 pt	\$3.10
STOP exits, inside 13:00-15:00 UTC	90	+1.820 pt	\$3.64

Entries pay the touch and about nothing more (0.389 against a 0.341 half-spread). Stops pay 1.55pt book-wide and **1.82pt inside the RIDER's own window** — the open is the worst place on the clock to be stopped. **None of this was assumed in.** The harness fills a stop-market at the worst print inside the next 250ms, and that model came out at 1.624pt against the live-measured 1.820pt — **11% cheap**; charging the live number costs a further \$18 across all 45 stop exits. **The RIDER's own stop-slippage bill is \$146 over 12 days — about \$12 a session.**

(c) A genuine race, not an MFE count. First-touch of the 2R target against the $2.0 \times \text{ATR}$ stop on the raw trade-print sequence, stop wins ties, sequential:

LADDER, SAME 12 DAYS	N	NET @\$1.50	\$/TRADE
5s-bar model, exact fills	94	+\$4,957	+\$52.74
250ms tick race, exact fills	94	+\$4,927	+\$52.42
+ cross the spread on entry	95	+\$4,587	+\$48.28
+ stop filled on the tape's flush	95	+\$4,440	+\$46.74
+ cross out on the time cap = FULL HONEST	95	+\$4,435	+\$46.69

The bar→tick effect is −\$30 on \$4,957 — 0.6%. The 87% flush-loss that 5-second bars caused elsewhere on this desk does not generalise here, and the reason is structural: a 54.6pt stop is far outside a 5s bar's range, so there is almost no intrabar ambiguity to resolve. Total execution bill **\$5.18 a trade** against a \$52.42 exact-fill gross. Exit mix 45 STOP / 40 TARGET / 10 TIME.

And the last optimism is worth exactly \$0. Requiring a print 1 or 2 ticks beyond the resting 2R limit before it fills leaves the result byte-identical — +\$4,435, 40 targets, at 0, 1 and 2 ticks of required penetration. Every target that was reached was blown through. **That directly refutes the author's own "two knife-edge target misses" story.**

(d) Overlapping positions — the attack misses. `rider.py` was already sequential. Re-imposing sequencing on tick-resolved exit times moves `n` from 94 to 95, not 94 to 129. Leave-one-day-out worst +\$3,543, strip-best-3 +\$3,204, strip-best-10 +\$790 (top 10 of 95 carry 82%).

(e) Runnable as a gate? NO — and this is where the attack actually lands. The 120-minute hold cap is fine (longest hold 45.0 min) and the ASIA guard is irrelevant (13-15 UTC). **But it dies on plumbing.** Every live path through `scaleout_slots()` hard-codes `stop_atr_mult=1.0`; a $2.0\times$ stop exists only in the shadow slate, and `exit_scalp` couples the target to the stop width. Re-run at the only stop the live machinery can build:

GEOMETRY	N	NET @\$1.50, FULL HONEST EXECUTION	WIN%	\$/TRADE
as specified — stop 2.0×ATR, 2R	95	+\$4,435	51.6%	+\$46.69
live slate — stop 1.0×ATR, 2R	181	+\$3,183	43.1%	+\$17.59
live Lot-A default — stop 1.0×ATR, 2.5R	157	+\$2,795	37.6%	+\$17.80
stop 1.5×ATR, 2R (already refuted live)	121	+\$2,505	42.1%	+\$20.70

Deploying it through today's slate costs 62% of the per-trade edge and nearly doubles the fill count — it pays the spread and the stop-flush twice as often for a third of the reward. On top of that, the **quiet-tape clip fires on 34-41% of its entries** ($ATR < 22$), replacing a \$222 mean 2R target with \$40 / $\max(1.75R, \$60)$. Whatever that two-lot book earns, **it is not the strategy that was tested.** Any "ODR live" number must be re-derived through `scaleout_slots()`, never from the spec — that trap has already eaten the ER/ATR floors, the ER-hold shadow, the regime-3 selector and capitulation's base target.

△ **AND THE SHARED-ACCOUNT PROBLEM — THIS IS THE 08-06 CASCADE CLASS, VERBATIM**

49% of the RIDER's entries (47 of 95) open inside the day-rider's live 13:38-15:00 window. It flips direction **2.6 times a day**, and **34% of the colliding entries (16 of 47) are opposite the day-rider's own direction rule.** Both desks trade MNQ in DUQ191770 and IBKR nets them into one number. That is exactly what produced 11 minutes halted with the whole safety block skipped, phantom exits and a naked GTC stop firing six hours later.

safety.own_flatten_verdict gates flattens; it does not stop two strategies netting each other's entries. **A second always-in desk in the day-rider's window is not a config change, it is an architecture change** — and the still-open gap ("does every live STOP have a slot?") gets materially worse with 7.9 extra wide stops a day working in the book.

The caveat this lens would not hide. Perturbing only information-free execution nuisances — latency 100ms to 2s, cadence 4/5/6 minutes — spans **\$3,291 to \$5,106, sd \$519 = 41% of the headline**, and it is **non-monotone in latency** (perfect 0ms fills earn \$3,762; sloppy 1,000ms fills earn \$4,938). **Several hundred dollars of the headline is fill-luck.** But the entire band is positive and about 8.5 standard deviations above zero, so this lens cannot refute. **Quote it as "about +\$4.4k ± \$0.9k on 95 verified trades", never as +\$4,169.** And the four days that can never be execution-verified are worth only **+\$182** of the 17-day total, so "you only checked the good days" does not apply.

THE VERDICT ON THE OPEN RIDER

The panel votes 2-1 to survive, so the OPEN RIDER survives — as a SHADOW candidate for incubation, never as a live promotion on 17 in-sample days. But I want you to hold both halves of that at once, because the dissent is the sharpest thing written about this candidate all week.

Does it make money?	Yes, repeatedly and on unseen tape. +\$7,433 over 36 days, 36/36 leave-one-day-out positive, all four quarters green, and + \$30.66 a trade on 19 days it was never fitted on.
Is it a timing edge?	No. All three lenses agree there is no moment-picking skill. Where they split is whether what's left is worth owning.
Is it just a bet on the open being trendy?	Genuinely contested. Duty-cycle says yes: zero intercept, 88% reproduced by one decision a day, and inside the placebo out of sample ($p=0.168$). Robustness says no: equal-longs-and-shorts leaves the per-trade rate unchanged, and the literal day-direction bet loses \$3,259. These two used different out-of-sample legs (17 V5 days vs 19) and reached opposite readings — that disagreement is itself the reason this is not live.
Does execution eat it?	No. \$5.18 a trade against a \$52 gross; its own stop-slippage bill is \$12 a session.
Can we actually run it?	Not today. The live slate has no code path to the $2\times\text{ATR}$ stop the edge depends on (62% of the per-trade edge lost at $1\times$), the quiet-tape clip rewrites a third of its exits, and it collides with the day-rider on a netted account half the time.

What I would do: log it in shadow at the corrected spec — **13:00-14:45** (cut the dead right edge), cadence 5 minutes only (cadence-3 is out-of-sample broken), stop $2.0\times\text{ATR}_{1m}$, 2R, 45-minute cap. Judge it on the author's own bar: **+0.15R over 200 shadow trades including at least one demonstrably non-trending open week** — that is precisely the regime the dissenting lens says will hurt it. Nothing goes live until the plumbing question (a real $2\times\text{ATR}$ stop path) and the shared-account question (position isolation) are both answered.

STEP 6 — THE CHOP-DAY SCALP LAB. A clean, well-proven NULL — and a \$5,483 finding hiding behind it.

Your explicit brief: the trend gates carry the profit, chop days are where we donate, and "untradeable" is only a statement about the current gates — so build a purpose-made chop scalper. **The answer is REFUTED, and not because of costs.** But the lab produced the single largest actionable number in this Movement, and it is not a new mechanism.

First, the premise — confirmed, hard

Every live trade (data_quality IS NULL) attributed to the 30-minute block it opened in, 17 trading days:

REGIME AT ENTRY	N	NET	\$/TRADE	WIN
CLEAN-TREND	120	+\$1,373	+\$11.44	45%
BUILDING	143	+\$109	+\$0.76	42%
DEAD-CHOP	11	-\$125	-\$11.36	36%
NORMAL-CHOP	159	-\$3,413	-\$21.47	31%
VIOLENT-WHIPSAW	166	-\$1,944	-\$11.71	39%
ALL CHOP	336	-\$5,483	-\$16.32	—
ALL NON-CHOP	263	+\$1,482	+\$5.64	—
ALL	599	-\$4,001	-\$6.68	—

The desk is profitable in trend and building blocks and gives all of it back plus \$4,001 in chop. 64% of all 30-minute blocks in the window are chop, so this is a large surface. (Honest caveat: this is hindsight labelling — it answers "where did the money go", not "what should have been armed at 09:14".)

Two mechanisms built, both refuted — and the model-free race that settles it

CHOP-FADE (fade a ≥ 1 -ATR stretch from a 20-minute mean) and CHOP-EDGE (fade a failed poke through the prior 2h range) were both specified mechanically and swept. Then the whole class was tested without any model at all: at every minute whose trailing tape passes the chop gate and whose stretch is ≥ 1 ATR, race $+X \cdot \text{ATR}$ against $-X \cdot \text{ATR}$ on the forward tick path. No stop, no target, no fee, no parameters. Fading needs $P(\text{revert first}) > 0.5$.

X (ATR)	N	P(REVERT FIRST)	AFTER AN UP STRETCH	AFTER A DOWN STRETCH
0.50	4,254	0.4770	0.482	0.472
0.75	4,254	0.4673	0.472	0.462
1.00	4,249	0.4723 (z = -3.61)	0.476	0.469
1.50	4,191	0.5006	0.500	0.501
2.50	3,549	0.4970	—	—

Chop-classified MNQ tape mildly CONTINUES, it does not revert.

Both sides agree, all 11 sessions are at or below 0.506, and every segment agrees. **It is not a cost problem, it is a sign problem** — and the

corroboration is that across **450 frictionless cells** (fee zero, slip zero) only 34 were positive, at a median of **−\$1.50 a trade before a cent of fee**.

And the escape hatch is closed too. The obvious fix — widen the distance so the fixed fee amortises — fails because the directional bias decays to zero by $X=1.5$. The largest gross edge anywhere in either direction is **\$1.49 a trade against a \$2.24 friction floor**. There is no window where they cross.

The final nail is the search-corrected placebo. CF-A's best cell sits at the 97.2nd percentile of a random-side placebo — but that is the p-value of the winner of a 540-cell grid. Give the coin flip the same number of tries and **77.0% of random-direction placebos beat the real best-of-grid. $p = 0.770$** . Also: the ER gate — the variable that defines chop — contributes nothing (removing it entirely performs as well as the tuned value), and CF-A's whole net is **one session**: drop 08-06 and +\$431 becomes **+\$2.40 on 260 trades**.

★ **THE FINDING THE LAB ACTUALLY PRODUCED — AND IT IS A ROUTER JOB, NOT A GATE**

A perfect chop filter is worth about \$5,500 over 17 trading days. A chop scalper is worth nothing. Chop blocks cost the desk −\$5,483 on 336 trades while everything else made +\$1,482. That is **~\$320 a day of pure not-being-there**, and nothing in this greenfield hunt came within an order of magnitude of it. The taxonomy is cheap and reusable — $ER_{30} < 0.22$ with roundtrip > 2.5 and no 2-hour range break, computable from trailing 1-minute bars in milliseconds, the same family of measurement `runstate.py` already runs. **It is a block statistic and will sometimes disagree with a RUN STATE read — the standing rule that RUN STATE WINS on disagreement still applies.**

The two survivors of the chop lab — both thin, both filed honestly

LEAD	NUMBERS	VERDICT
DEAD-CHOP fade in the 20:00-21:00 UTC hour (ATR < 10, the final hour before the CME halt)	model-free race n=159, P(revert) = 0.610, z = +2.78, gross \$3.47 against \$2.19 friction. All 20 costed parameter cells positive — the one genuine plateau anywhere in this hunt. Reference cell: n=10, +\$117, +\$11.69/trade, 9 of 10 won, 5 of 5 days green. Widening the ATR ceiling keeps it alive; removing it kills it, which is the right shape for a real effect.	SHADOW It is ten trades.
Chop-gated WHIPSAW at the US open is a CONTINUATION trade	P(revert) = 0.444 on n=453 (z = -2.40), best-signed gross \$6.79 against \$2.22 friction = +\$4.56 a trade. And it is exactly the block where the desk currently loses — \$1,944 by being armed with churners.	PARKED It is momentum, out of scope here

Revival conditions for the 20:00 fade, all three required: (1) $n \geq 60$ costed shadow trades, roughly 6–8 more weeks at ~1 trade a day; (2) the race stays at $P \geq 0.57$ with $z \geq 2$ on the forward sample only; (3) the 20-cell plateau still has $\geq 80\%$ of cells positive and the $ATR < 10$ gate still separates. Clear all three and it earns a 1-lot promotion inside the 20:00–21:00 window only. **It should never be promoted on this sample.**

Revival for the whipsaw continuation lead: hand it to the trend book as an entry-selection study with its own stop/target sweep and a forward leg. **Do not promote it off this section's numbers.**

THE MASTER LEDGER — every lead in this Movement, by name

Nothing is left dangling. Each lead is LIVE, SHADOW, PARKED (with its revival condition) or REFUTED (with the test that killed it).

CLUSTER	LEAD	VERDICT	THE NAMED TEST / THE REVIVAL CONDITION
---------	------	---------	--

oracle	The sat-out runs are real money, and lateness is what costs it	LIVE finding	52 sat-out runs entered at their start minute with a real racing stop/target at \$1.50/RT: +\$6,846 = 71% of the hindsight ceiling. Delay the entry and it decays +\$6,846 → +\$6,041 (2 min) → +\$4,090 (5) → +\$944 (10) → −\$2,068 (15). The exit machinery is not the constraint; the entry is, and specifically its latency. Corollary, from the same run: at the desk's 1-ATR house stop even a perfect entry converts only 24% and is shaken out of 14 of 52 runs.
chop lab	"Sit out the chop blocks"	LIVE — this is the finding	Live trades attributed to entry regime: chop 336 trades −\$5,483, non-chop 263 trades +\$1,482. Perfect avoidance turns −\$4,001 into +\$1,482 over 17 days. No new mechanism required — the router already owns the lever.
OPEN/NEWS	THE OPEN RIDER — 5-min cadence, direction = last 15 min, 2.0×ATR stop, 2R, 45-min cap	SHADOW (panel 2-1)	True OOS on 19 unseen days: +\$3,526 / n=115 / +\$30.66 a trade, 5% degradation. 36/36 LOO positive; reverse walk-forward picks the published cell #1 of 48; placebo z=+2.59 p=0.005. PROMOTE IF: +0.15R over 200 shadow trades including a non-trending open week, plus the 14:45 cut, plus a live 2×ATR stop path, plus position isolation from the day-rider.
OPEN/NEWS	ODR as a tradeable edge (dissent)	REFUTED by lens 1	The 17 unseen V5 days land at the 83.2nd percentile of 2,000 random-direction placebos, p=0.168 — inside the null; and daily P&L regressed on window drift has intercept −\$23, t=−0.22. One decision a day captures 88% of the net.
OPEN/NEWS	"13:00–15:00 is where MNQ travels"	PARKED	Real and stable (\$1.237 per point of drift, t=+2.75, holding at 1.207/1.091 across both halves) but it is an exposure, not a signal. REVIVE IF: an out-of-sample forecaster of next-window magnitude (not direction) tests positive — then the payoff is worth harvesting deliberately.
OPEN/NEWS	Wide stops (≥2×ATR1m) in the open window	SHADOW	27 of 27 cells at stop ≥ 2.0 positive; every cell at ≤1.0 flat-to-negative; reproduces on 34–36 days independently. REVIVE/CONFIRM IF: re-measured on a signal whose intercept is not zero — it has only ever been measured inside the RIDER.
OPEN/NEWS	Cadence 3 minutes ("the plateau holds at every cadence")	PARTIALLY REFUTED	On unseen tape cadence-3 is positive in 4 of 16 cells vs 13/16 at 5 min and 15/16 at 10. Strike it

			from the language; the chosen cell is on the safe side.
OPEN/NEWS	The window's right edge 14:45-15:00	REFUTED — cut it	–\$446 on the full tape and –\$453 on the unseen leg alone. Independently re-confirms our own 14:45 boundary.
OPEN/NEWS	ODR deployable as a tournament gate at its own geometry	REFUTED	scaleout_slots() forces stop_atr_mult=1.0 on every live path; at that stop the same entries earn +\$17.59 a trade, a 62% haircut at double the fill count ; the quiet-tape clip additionally fires on 34-41% of entries.
OPEN/NEWS	ODR deployable alongside the day-rider	REFUTED as configured — PARKED pending isolation	49% of entries open inside the day-rider's live window, 2.6 direction flips a day, 34% of colliding entries opposite its direction rule — on a netted account. REVIVE IF: its own account or per-strategy position isolation, and the "does every live STOP have a slot?" auditor is built.
OPEN/NEWS	"+\$4,169 on 102 trades" as a quotable figure	PARKED — do not quote it	The "102" is a different row of the report. Three independent harnesses give +\$3,907 / +\$4,620 / +\$5,028 / +\$5,139 at the correct fee on the same 17 days. Quote "about +\$4.4k ± \$0.9k on 95 verified trades".
OPEN/NEWS	COIL-CRACK — momentum ignition out of a quiet 30 minutes	REFUTED as a filter	Direct ablation: with the gate + \$2,049/n=51, without it + \$2,238/n=58 . Costs money and sample. The beautiful monotone table that motivated it is an artefact of overlapping 1-minute sampling. A filter that removes trades and money has nothing to reformulate.
OPEN/NEWS	OPEN EXHAUSTION FADE	REFUTED	51 configurations, every one negative or zero, plus the path measurement: median excursion against the trade 4.76 ATR vs 3.00 ATR in favour. Widening the stop to 3-4 ATR and delaying entry 10 minutes both fail, closing the two obvious reformulations. Caught 0 of 14.
OPEN/NEWS	The 15-minute momentum lookback as a trigger	REFUTED as a trigger	Thresholding always cost money, and a duty-matched control replacing the lookback with "is price above the 13:00 print?" scores the same. The lookback is not the signal; being positioned with the window's direction is.
OPEN/NEWS	Pre-open-range containment filter; segment→R policy; the 5-state regime classifier	SHADOW / SHADOW / PARKED	Containment gives +0.446R vs +0.272R and its mirror loses — but it halves the sample and was picked from a menu of ten, so the multiple-comparison tax is unpaid; promote if it holds +0.40R over 60 shadow trades . The segment policy beat

			blanket out of sample (+\$34.17 vs +\$18.51 a trade) on n=19-48 per bucket; revive at n≥60 per segment. The classifier behaved backwards (CLEAN-TREND second-worst, VIOLENT-WHIPSAW best); revive only if re-cut on realised trendiness over the hold.
VACUUM	FLOW-TRAP — fade a failed aggressor burst on the full retrace	REFUTED	The cost-free forward-return test: the entry's raw 15-min return (+2.25pt, 52.1% up) is below the tape's own unconditional drift (+2.73pt, 53.9%) at 5, 15 and 30 minutes. Its information content is not small, it is negative — no exit rule saves that, and no fee correction touches it. (Δ at the correct \$1.50 fee one of its twelve cells limps to +\$154 on 74 trades; the cost-free test is what kills it, not the P&L.)
VACUUM	ABSORPTION-SHELF — fade a burst that produced no price result	REFUTED	12 of 12 cells negative and losing on both sides (long – \$302 / short –\$298 at the mid threshold), so it is not even a drift artefact — it is just wrong. The "effort without result" split showed no separation in the base-rate study either.
VACUUM	VACUUM-BREAK — enter on the break of the shelf under the trapped side	PARKED	Killed as built by strip-the-best (+ \$1,069 → +\$342 at strip-3 → + \$26 at strip-5, i.e. five trades are the whole thing) and by the walk-forwarded regime policy (– \$125 over 91 trades; selection tax \$1,194). REVIVE IF: (a) we accumulate a down-trending or genuinely range-bound week so the +1,300pt melt-up stops doing the work — the short side must clear zero on its own; and (b) $n \geq 200$ walk-forward trades, which needs ~15 sessions of aggressor-tagged ticks against the 5 we have. Both are free to wait for.
VACUUM	Raw "fade the sell-burst" conditional	PARKED	The one large-n signal in the study (n=519, +6.19pt of excess over drift at 15 min) — but its mirror loses real money and 519 overlapping windows in 4.7 days is perhaps 25 independent observations. REVIVE IF: the short mirror clears zero excess on a flat or down week. Until then it is "buy the dip in a melt-up" in a microstructure costume.
VACUUM	Book / L2 "vacuum" reading	PARKED — the one stone worth funding	Every version tested measured the book as static average depth per minute, and the intuitive version loses. REVIVE BY: re-cutting on queue depletion rate — how fast resting size at the touch disappears and whether it comes back. The

			82.3M-row book table at ~190 snapshots/second supports it. That is a genuinely unturned stone.
VACUUM / FLOW-LED	Bar-derived flow proxy (to extend flow work past 5 days)	REFUTED	Four proxies validated against 6,779 minutes of true aggressor flow: best correlation 0.272, sign agreement 49.9-53.4% , while volume itself reconstructs at 0.96. Aggressor sign is genuinely absent from OHLCV. No revival without a longer tick archive — an infrastructure action, not a research one.
FLOW-LED	FLOW-LED as a tradeable family (the census rule traded directly)	REFUTED	Strip-best-3 negative in 6 of 6 cells, plus the catch/profit inversion across the whole threshold sweep : every cell catching 2 of 4 targets loses \$434-\$1,302 over 111-129 trades. n is not the problem. And the tape shows there is nothing to trigger on — flow run-length at the target runs is 1-2 minutes and reverses on the next bar.
FLOW-LED	Single-minute flow spike; the same on quiet tape; the naive fade; the "flow agrees" rule at all hours	REFUTED × 4	All four die on strip-best-3 plus sign-instability across their sweeps, at n=80-136 — not thin-n. The quiet-tape version is the strongest kill in that section: it loses at every threshold, monotonically worse as you tighten, in its own home segment (win rate 20-29%).
FLOW-LED	FLOW-BREAK IGNITION / FLOW-TRAP FADE	SHADOW × 2	+\$1,967/n=39 and +\$2,024/n=44, both beating their nulls at p<0.01, both 0 of 4 on the cluster they were built for. IGNITION's out-of-sample half collapses to a coin flip; FADE's holds and all five days are green. PROMOTE IF: IGNITION keeps US-session \$/trade above ~\$40 at n≥50 over 3 more weeks of ticks and its OOS leg stops being a coin flip; FADE holds ≥\$25/trade over n≥80. △ check FADE against our existing abs_veto absorption work before shadowing a duplicate.
chop lab	CHOP-FADE (band fade) and CHOP-EDGE (failed-break fade)	REFUTED × 2	Model-free first-touch race P(revert) = 0.4723 on n=4,249, z = -3.61 — negative before any cost . 450 frictionless cells: 34 positive at a median -\$1.50 a trade. 77% of random-direction placebos beat the best-of-grid, p = 0.770. The ER gate contributes nothing. CHOP-EDGE: 1 of 15 cells positive, strip-best-1 already negative, out-of-sample sign flip.
chop lab	Any symmetric chop scalp at any distance	REFUTED	Max gross expectancy in either direction is \$1.49 a trade against a \$2.24 friction floor ,

			and the bias decays to 0.5006 by 1.5 ATR — widening to amortise the fixed fee destroys the edge it was meant to pay for . Any future chop-scalper proposal must first show $P(\text{revert first}) > 0.5$ on fresh tape.
chop lab	DEAD-CHOP fade, 20:00-21:00 UTC	SHADOW	Race n=159, P=0.610, z=+2.78; all 20 costed cells positive; n=10 trades. Three revival conditions above, all required.
chop lab	Chop-gated whipsaw at the US open is a continuation trade	PARKED	$P(\text{revert})=0.444$ on $n=453$, best-signed +\$4.56 a trade net — but it is momentum, out of scope. REVIVE BY: handing it to the trend book with its own sweep and a forward leg.
census	The FLOW-LED / VACUUM cluster labels themselves	REFUTED as written	One minute of $ z \geq 1$ assigns the label; that minute reverses on the next bar and an independent re-derivation disagrees 3 times out of 4. FIX: require $\text{persist} \geq 3$, or merge both into a single "flow event" bucket.

WHAT I TAKE AWAY

Seven things, money first

1. **The runs are real money and we are late, not wrong.** Entered on the minute with an honest exit, the sat-out runs pay **\$6,846 — 71% of their ceiling**. Turn up ten minutes later and that is **\$944**; fifteen and it is **-\$2,068**. **Every minute of confirmation you wait for costs about \$600**. That is why the only survivor in this Movement is the one signal that waits for nothing.
2. **The biggest number in this Movement is not a new gate — it is \$5,483 of not being there.** Chop blocks cost the desk that over 17 days while everything else made +\$1,482. A perfect chop filter beats every invention in this report by an order of magnitude, and the router already owns the lever.
3. **The simplest thing beat both clever things, and it was not close.** Two hand-built footprint signals lost \$2,000–\$5,000 between them; the

control that just shows up every five minutes with a wide stop made money on 36 days and was on 15 of the 15 biggest open-window runs.

4. **The stop is the other lever, and it now has three independent votes.** The survivor's own 45-cell grid, the robustness skeptic's from-scratch rebuild, and my oracle all say the same thing from different directions: every cell at a 1-ATR stop is flat-to-negative, every cell at 2-ATR or wider is positive, and even a perfect entry gets shaken out of 14 of 52 runs at 1 ATR against **0 of 52** at 2 ATR. **Our ~1-ATR house stop is the wrong instrument for the US open** — checking that on the live gates costs nothing and is the one thing here I would look at this week.
5. **A label is not a signal.** VACUUM fires on 1,451 minutes to find 8 runs (0.55%); FLOW-LED hangs on one minute of flow that reverses on the next bar. Both clusters produced no survivor because the cluster is not a real thing, and finding that out is worth more than a curve-fitted P&L would have been.
6. **Skeptics with different out-of-sample legs reached opposite verdicts on the same strategy.** One used 17 unseen days and refuted it; another used 19 and upheld it. That is not a contradiction to paper over — **it is the reason this is a shadow candidate and not a live one**, and the honest way to resolve it is forward tape, not another backtest.
7. **The binding constraint on flow research is retention, not ideas.** Five days of aggressor-tagged ticks caps every flow finding this desk will ever have at $n \approx 40$. Three more weeks would take FLOW-TRAP FADE from $n=44$ to $n \approx 180$ and turn "not obviously noise" into an answer. **Keep archiving ticks.**

HOW HONEST THIS IS — THE LIMITATIONS, IN ONE PLACE

The census week is **5 days** and MNQ rose 1,300 points through it, which flatters every long-biased number; the flow work is capped at **5 days of ticks** and cannot have an out-of-sample leg at all; the survivor has **36 trading days** total and has never been shown a bad month (its worst day, —\$901, is its maximum drawdown); the chop lab is **11 tick days and 17 bar days**, which is enough to refute comfortably at -3.61 sigma and nowhere

near enough to confirm anything; and shadow-simulated losses have historically run **1.1-3.1× modelled**, so every negative number here is a floor and every positive one is a ceiling. Single instrument, single contract month, MNQ Sep'26.

Part 3 The Rev2 corrections — the new material, in full

Everything in this part was written after the report was rendered at 10:17 UTC on Saturday. The self-proofread pass reopened six questions against the sections above, and the build was killed before it could fold any of them back in — so on the first published file, Parts 1 to Movement 3 still argued for things this part refutes. **Where a Part 3 section disagrees with anything above it, Part 3 is the correction**, and the action card at the top of this document has been rebuilt to match — every amended play is tagged ★ **REV2 CORRECTION** in-row. Three of the six (Q0, FIX2 and the ER note in Q2) independently kill the same recommendation from three different books, which is the strongest result in the document and the reason the ER floor is gone from Saturday #2. Two more (Q1, FIX0) change what you type rather than whether to act, and one (Q2) is the first grading anyone has done of the previous week's card.

The ER30 $\geq$ 0.35 floor on abs_veto_short — put it in the table and it dies in the table

PATCHES: Part 2.6 §8 Saturday #1 and plays.json rank 2 (absveto-short-arm-by-default-er035-0808) — the ER floor half of that play is **WITHDRAWN**. **EXTENDS:** Part 1.5 rehab_abs_veto_short §5 (the 54-filter table) with the OOS leg it was missing. The arm-by-default half of the play and the BUILDING $\times$ US-SESSION veto are **unchanged and still ship**.

THE ANSWER IN ONE LINE

Gaz — you are right and I was wrong. Your bar was **13 of 15 winners kept** to survive, **7-9 to be the same fake win the section rejects nine times**. **ER30 $\geq$ 0.35 keeps 5 of 15**. It does not clear the bar, it does not even reach the fake-win band — it is worse than E0.30, which the section already printed and rejected at 7/15. It bins **\$1,064 of the \$1,664 top-15 block**, discards **+\$1,363 of profitable trades**, and goes **+\$809 $\rightarrow$ +\$363 out of sample**. Ship the BUILDING $\times$ US-SESSION veto. Do not ship the floor.

1 · The row you asked for — same columns, same book, plus the OOS leg

Basis: abs_veto_55s SHORT, tick-repriced (shadow_real.real_pnl), 159 fires, 17 trading days 2026-07-16 $\rightarrow$ 08-07, four ISO weeks. \$2.00/point, **\$1.50/round-turn**, 1 lot. Identical ledger and identical strip-3 / top-15 definitions as §5, reproduced byte-for-byte before adding columns (scratchpad/rehab_avs/q0_er035.py). **IS** = ISO weeks 29-31 fitted; **OOS** = week 32 untouched.

POLICY	KEE P N	KEE P \$	\$/TRADE	WIN %	CUT N	CUT \$	CUT \$/TR	TOP -15 KEP T	STRI P-3	STRI P-5	WO RST LOD O	DAY S GRE EN	WK2 9	WK3 0	WK3 1	WK3 2 OOS
00 BASELINE — no filter	159	+2290.5	+14.4	48.4%	0	—	—	15/15	+1780.0	+1484.5	+1889.0	14/17	+325.5	+497.0	+659.0	+809.0
E0.35 — ER30 ≥ 0.35 the play	36	+927.5	+25.8	55.6%	123	+1363.0	+11.1	5/15	+493.0	+327.5	+714.5	10/13	0.0	+244.0	+320.5	+363.0
G2 — drop BUILDING × US-SESSION the verdict	151	+2497.5	+16.5	50.3%	8	-207.0	-25.9	15/15	+1987.0	+1691.5	+2096.0	14/17	+325.5	+542.5	+716.0	+913.5
G2 + E0.35 stacked both, as the play asks	36	+927.5	+25.8	55.6%	123	+1363.0	+11.1	5/15	+493.0	+327.5	+714.5	10/13	0.0	+244.0	+320.5	+363.0
A15 — ATR14 ≥ 15 the REAL-floor control	82	+1829.0	+22.3	52.4%	77	+461.5	+6.0	15/15	+1318.5	+1023.0	+1414.0	11/15	+382.5	+224.5	+579.0	+643.0

Read the row against your own decision rule and there is nothing left to argue about. 5 of 15 winners. The cut pile is **positive** at +\$1,363 — a filter whose discards make money at +\$11.10 a trade is not removing bad trades, it is removing trades. Out of sample it turns +\$809 into +\$363: it does not go negative, it simply hands back 55% of the week it was supposed to protect. And **ISO week 29 is \$0.00 — the floor produces zero fires in an entire week**, and no fires at all on 4 of the 17 days. The gate stops being a gate.

THE LINE THAT SHOULD END THE DEBATE ON ITS OWN — ROW 4

G2 + E0.35 stacked is identical to E0.35 alone. Not similar — identical, to the cent, on all sixteen columns. BUILDING is defined as $ER30 \in [0.15, 0.30)$, which lies entirely below 0.35, so the floor already eats every fire the veto was going to cut and 115 more besides. **The play asks to ship both,**

and shipping both is arithmetically the same as shipping only the floor. Adding the ER floor does not strengthen the one filter this report verdicted — it deletes it and takes **\$1,570** of the reference book out the door with it. That is the answer to “why is play #2 bolting an ER floor onto it”: because nobody stacked the two rows and looked.

2 · Which winners it bins — and the rule it quietly re-imposes

The 15 biggest winners carry +\$1,664.00 of the +\$2,290.50 book (73%). Here are the ten the floor throws away.

#	DAY	UTC	\$	ER30	ATR14	REGIME	TIME-OF-DAY	ER0.35	G2
1	07-29	19:21	+190.0	0.451	47.5	TREND-NOBREAK	US-SESSION	keep	keep
2	07-29	14:12	+164.5	0.305	41.4	CLEAN-TREND	US-SESSION	CUT	keep
3	08-07	13:56	+156.0	0.002	38.9	VIOLENT-WHIPSAW	US-SESSION	CUT	keep
4	07-27	14:30	+153.5	0.395	38.8	TREND-NOBREAK	US-SESSION	keep	keep
5	08-05	14:30	+142.0	0.139	35.6	VIOLENT-WHIPSAW	US-SESSION	CUT	keep
6	08-05	19:59	+91.0	0.557	21.0	TREND-NOBREAK	US-SESSION	keep	keep
7	07-27	16:50	+90.0	0.053	23.0	VIOLENT-WHIPSAW	US-SESSION	CUT	keep
8	08-06	15:18	+88.0	0.125	22.4	VIOLENT-WHIPSAW	US-SESSION	CUT	keep
9	07-28	13:30	+87.0	0.184	22.2	VIOLENT-WHIPSAW	US-SESSION	CUT	keep
10	07-29	01:51	+87.0	0.336	22.2	TREND-NOBREAK	OVERNIGHT	CUT	keep
11	08-06	13:31	+84.0	0.471	21.2	CLEAN-TREND	US-SESSION	keep	keep
12	07-17	08:16	+83.5	0.285	21.4	BUILDING	OVERNIGHT	CUT	keep
13	07-17	08:10	+83.0	0.107	21.0	VIOLENT-WHIPSAW	OVERNIGHT	CUT	keep
14	07-29	01:10	+83.0	0.250	21.1	VIOLENT-WHIPSAW	OVERNIGHT	CUT	keep
15	07-20	18:41	+81.5	0.387	20.5	TREND-NOBREAK	US-SESSION	keep	keep

Seven of the ten binned winners are **VIOLENT-WHIPSAW**. That is not a coincidence and it is the mechanical heart of the contradiction. The same play asks, three bullets earlier, to **delete the violent-whipsaw bench (ATR ≥ 19 AND ER30 < 0.25) because it is backwards in the US session** — and I agree, it is: that cell is **n=12, +\$443.50, +\$37.00/trade, 58% win**, the gate’s second-best pocket. Then the next bullet re-imposes the same rule as a bare number. **ER30 < 0.35 is a strictly wider version of ER30 < 0.25 with the ATR condition removed**, so it cuts **all 12** of those fires instead of some of them. The old bench cut 19 fires and \$357.50; the floor cuts 123 fires and \$1,363.00. It is the refuted rule, 3.8× bigger.

The reason is the one \$5 already gave and this table makes physical: **a thrust worth trading is a thing that has just made the last thirty minutes look messy**. Median ER30 in the **VIOLENT-WHIPSAW × US-SESSION** cell is **0.126**. The gate’s single best winner outside the top slot — 08-07 13:56, +\$156.00 — fired at **ER30 = 0.002**. An ER floor is, definitionally, a filter that removes the gate’s own trigger condition.

3 · Regime-segmented — scored on HOME segments, never blanket

Left: what the floor **cuts**, per regime × time-of-day. A filter is only justified where its cut pile is **negative**. Right: what survives.

REGIME	TIME-OF-DAY	CUT N	CUT \$	CUT \$/TRADE	WIN%	VERDICT ON THE CUT
NORMAL-CHOP	OVERNIGHT	26	+474.0	+18.2	57.7%	throws away the gate’s largest single pocket
VIOLENT-WHIPSAW	US-SESSION	12	+443.5	+37.0	58.3%	the pocket the play’s own bullet says to stop benching
BUILDING	OVERNIGHT	28	+219.5	+7.8	42.9%	positive — G2 correctly leaves this alone
TREND-NOBREAK	OVERNIGHT	8	+139.5	+17.4	50.0%	positive
NORMAL-CHOP	US-SESSION	8	+129.0	+16.1	50.0%	positive
BUILDING	PRE-OPEN	4	+126.0	+31.5	75.0%	positive
CLEAN-TREND		6	+154.0	+25.7	50.0%	

	PRE-OPEN / US / O-N					positive — and these are clean trends it is cutting
NORMAL-CHOP	LATE	2	+95.5	+47.8	100%	positive
BUILDING / V-WHIP	LATE & O-N	9	−25.5	−2.8	33.3%	marginal — roughly free either way
NORMAL-CHOP	PRE-OPEN	10	−38.0	−3.8	30.0%	correctly cut, worth \$38
V-WHIP / TREND-NB	LATE / US	2	−147.5	−73.8	0%	correctly cut — n=2
BUILDING	US-SESSION	8	−207.0	−25.9	12.5%	the only cell worth cutting — this is G2, and G2 alone
TOTAL CUT by E0.35	—	123	+1363.0	+11.1	46.3%	the discard pile is profitable

Ten of the sixteen occupied cells the floor cuts are profitable, and they carry **+\$1,781.00 across 94 fires**. Against that it correctly removes **−\$418.00** across six cells and 29 fires — and **−\$207.00 of that is BUILDING × US-SESSION**, i.e. G2, with a further −\$147.50 sitting in two single-trade cells. **The ER floor is a \$1,570 wrapper around a \$207 finding**. That is the whole story of the recommendation.

The kept book is equally telling. Under the floor `abs_veto_short` survives in exactly **two of five regimes** — TREND-NOBREAK (n=18, +\$597.00, +\$33.20/tr) and CLEAN-TREND (n=18, +\$330.50, +\$18.40/tr) — and nothing else. §3a’s headline claim for this gate is “positive in every one of five regimes”. The floor converts a five-regime gate into a two-regime gate and calls the higher \$/trade an improvement. **That is the “bin the winners and quote the average” failure, stated in regime terms.**

4 · No plateau — a spike, and a \$/trade illusion

Discipline says every threshold is a guess and must show a **plateau**. Here is the full rung sweep, including the half-rungs either side of 0.35.

ER30 RUNG	KEEP N	KEEP \$	\$/TRADE	CUT \$ (MONEY DISCARDED)	TOP-15 KEPT	STRIP-3	WK32 OOS
0.00 (baseline)	159	+2290.5	+14.4	0.0	15/15	+1780.0	+809.0

0.05	140	+2304.5	+16.5	−14.0	14/15	+1796.5	+731.5
0.10	125	+1723.5	+13.8	+567.0	13/15	+1215.5	+582.5
0.15	101	+1146.0	+11.3	+1144.5	10/15	+638.0	+210.5
0.20	81	+1081.0	+13.3	+1209.5	9/15	+573.0	+262.0
0.25	66	+1098.0	+16.6	+1192.5	8/15	+590.0	+285.5
0.30	51	+1140.5	+22.4	+1150.0	7/15	+632.5	+315.5
0.325	41	+849.0	+20.7	+1441.5	6/15	+414.5	+303.0
0.35 ← the play	36	+927.5	+25.8	+1363.0	5/15	+493.0	+363.0
0.375	30	+679.0	+22.6	+1611.5	5/15	+244.5	+290.5
0.40	23	+552.0	+24.0	+1738.5	3/15	+187.0	+317.0
0.45	14	+253.0	+18.1	+2037.5	3/15	−112.0	+173.5
0.50	8	−79.5	−9.9	+2370.0	1/15	−195.0	+89.5

Net across the rungs 0.325 / 0.35 / 0.375 goes +\$849 → +\$927 → +\$679. That is a **spike**, not a plateau — the chosen threshold sits on a local bump between two lower neighbours, and the whole 0.15–0.40 range wanders between +\$552 and +\$1,146 with no structure. Every rung from 0.10 upward is below the baseline on net; the only rung that clears it is **0.05**, which removes 19 dead-flat fires worth −\$14 and is a “don’t trade a frozen tape” cut, not an ER floor. **Winner retention is the only column in this table that behaves monotonically** — it falls from 15/15 to 1/15 as the rung rises, which is exactly what “the filter is selecting on trade count, not on trade quality” looks like.

WHY THE \$/TRADE COLUMN LIES — AND WHY YOUR INSTINCT THAT “EVERY ER RUNG WINS ON \$/TRADE” IS THE TELL

You spotted it without the arithmetic: **on a book that is positive overall, almost any subsample raises \$/trade if it is small enough**, because \$/trade rewards shrinking the denominator. Net does not. So I ran the placebo — keep the same number of fires **at random**, 20,000 draws, and see where the real filter lands.

POLICY	KEEP N	REAL NET	MEAN RANDOM NET (SAME N)	NET PERCENT ILE	TOP-15 KEPT	RANDOM EXPECTS	WINNER PERCENT ILE	VERDICT
E0.35 ER30 ≥ 0.35	36	+927.5	+517.7	91.1%	5/15	3.40	77.1%	FAILS both legs
E0.30 ER30 ≥ 0.30	51	+1140.5	+733.3	88.7%	7/15	4.81	84.1%	fails both legs
D3 dwell ≥3 of 5	27	+937.0	+386.2	97.7%	4/15	2.53	77.4%	fails the winner leg

A15 ATR14 ≥ 15 real filter	82	+1829.0	+1181.3	96.5%	15/15	7.73	100.0%	passes both
G2 drop BUILDING×U S real filter	151	+2497.5	+2175.2	98.5%	15/15	14.24	55.3%	passes on net

This is the cleanest statement of “ER filters FAKE, ATR floors REAL” the desk has produced. **Keeping 5 of 15 winners out of 36 fires is the 77th percentile of a coin flip** — random keeps 3.4 on average, so the ER floor carries essentially no information about which trades are the good ones. The ATR floor at the same test keeps 15/15 where random expects 7.73: the **100th percentile of 20,000 draws**. One of these is a filter. The other is a smaller sample.

5 · So why did it look so convincing? The sampling error, named

The play’s headline number is real and I reproduce it to the cent. It is measured on the wrong book.

“+\$61.20/trade above 0.35 (n=9) against negative in all three bands below (n=26)” comes from the **live** abs_veto_short fill book — all 17 days, data_quality IS NULL. Recomputed: **n=9, +\$550.50, +\$61.17/trade**. Exact. And the bands below really are negative: −\$251.00 / −\$419.50 / −\$80.50 on n=26. Nothing is wrong with the arithmetic.

What is wrong is that this is a 35-fill sample selected by the very policy the play is replacing. The live gate was armed for **224 minutes of 6,780 — 3.3% of the tape**. Play #2’s other half says arm it by default. The moment you do that, the population you are filtering stops being those 35 router-chosen fills and becomes the **159-fire continuous book** — and on that book the floor loses \$1,363. **You cannot fit a filter on the 3.3% and deploy it on the 100%.** That is the entire mechanism of the error, and it is why the two halves of the play point in opposite directions.

Then there is the robustness the play never ran on its own number. Here it is:

TEST ON THE PLAY’S OWN N=9	RESULT	READ
as quoted	n=9 +\$550.50 +\$61.17/tr	reproduces exactly

strip best 1	n=8	+\$14.50	+\$1.81/tr	the recommendation evaporates on one trade
strip best 2	n=7	-\$217.50	-\$31.07/tr	negative
strip best 3	n=6	-\$276.50	-\$46.08/tr	negative
worst leave-one-day-out	drop 07-29 →	-\$167.50	on n=6	one day carries all of it
days green above the floor			3 of 5	and the two most recent (08-06, 08-07) are both negative

Two fills carry the whole thing: id 359 / id 360, both the 07-29 19:23 signal, ATR 47.7 — the widest tape in the sample — for **+\$232.00 and +\$536.00**. Every other fire above the floor sums to **-\$217.50 on n=7**. The report’s own standing method rule is strip-the-best-trades, and the #1 recommendation does not survive **strip-1**. Symmetrically, the “negative in all three bands below” half is **two afternoons**: 07-31 (six fills, -\$274.00) and 08-05 (four fills, -\$274.00) out of -\$751.00 total — and 08-05 is the **operator do-not-bench override window** that \$4 already excludes from router responsibility. Remove those two days and the below-floor book is -\$203.00 on n=14.

The one sub-claim that fully holds: “blocks 2 of the week’s 4 losing entries, -\$274” is **correct** — 08-05 14:22 (ER30 0.160, -\$132.00) and 08-05 14:32 (ER30 0.171, -\$142.00). Verified. But the other two losing signals fired at ER30 **0.451** and **0.355**, and the second of those clears the proposed floor **by 0.005**. A filter whose live justification is decided in the fourth decimal place is not a threshold, it is a coin landing on its edge.

AND THE “MIRROR IMAGE” ON THE LONG SIDE FAILS THE SAME WAY

The play conditions on a striking result — abs_veto_long allegedly earns in 0.15-0.25 and does -\$0.90 above 0.35, “the two halves of the same gate want opposite ER regimes”. That reads off the same thin live book. On the **continuous tick-repriced book** (abs_veto_55s LONG, n=160, +\$1,156.50, 17 days) all four bands are **positive**: <0.15 **+\$4.04**/tr (n=64) · 0.15-0.25 **+\$9.40**/tr (n=45) · 0.25-0.35 **+\$10.69**/tr (n=27) · ≥0.35 **+\$7.77**/tr (n=24). Not -\$0.90 — positive. There is no opposite-regime structure on either side of the gate once you leave the router-selected sample. A finding that two halves of one mechanism want contradictory conditions should always be treated as a sampling artefact until it survives the full book, and this one does not.

6 · The floor it claims to mirror does not exist

Play #2 justifies the threshold as “mirroring grind’s”. **Grind has no ER floor.**

Live source, src/gazbot7/deciders.py:136:

```
ER_FLOOR: dict[str, float] = {} # ★2026-08-01 (operator, Saturday
window): ALL ER floors DELETED.
# They were never in force – the dual-slot scale-out slate broke the
lookup on 07-29 16:20, so no ER
# floor has blocked a single trade in production... Re-derived on the
blocked book: grind_long ER>=0.35
# blocked +$2,923/173sig (strip-3 +$1,132, LOD0 +$10.03) and
abs_veto_long ER>=0.25 blocked +$640/52sig
# – i.e. both floors were blocking PROFITABLE populations. [prev:
{"abs_veto_long": 0.25, "grind_long": 0.35}]
```

Three separate things follow, and each one independently sinks the mirror argument. **(1)** The grind 0.35 floor was **deleted seven days ago** and the deletion is pinned by tests/test_deciders.py::test_er_floors_all_deleted. **(2)** It was deleted for precisely the reason re-derived above — measured on its own blocked book it was throwing away **+\$2,923 across 173 signals**. **(3)** Even before deletion it **never blocked a trade in production**, because the scale-out slate broke the lookup on 07-29 — the exact silent-drop failure mode CLAUDE.md warns about. So the threshold being mirrored is a number that has never once run, was refuted when measured, and has been absent from the codebase for a week. This is the **08-01 ER-floor burn repeating**, and it arrived inside the very report that re-derives the finding independently on grind, exhaustion and abs_veto in the same week.

For completeness on your quoted count: strictly, **two** of the 54 rows beat the baseline, not one — G2 (+\$2,497.50) and X8 (+\$2,319.50, which gets there by deleting a single trade in 17 days). The distinction changes nothing. **G2 remains the only filter in the sweep that improves net, keeps 15/15 winners, and improves out of sample.**

7 · Cross-validation across nine independent exits — the test G2 passed

G2 was cross-validated on nine strategies running the same entry with different exits. The floor gets the same treatment. If a filter is real it travels; if it is fitted to one P&L curve it does not.

STRATEGY (OWN BASE N BASE \$ E0.35 N E0.35 \$ DELTA G2 \$ DELTA EXIT)								
abs_veto_55s	159	+2290.5	36	+927.5	−1363.0	+2497.5	+207.0	
abs_veto_50s	173	+1997.0	35	+879.0	−1118.0	+2204.0	+207.0	
thrust_short_a bsveto55	90	+965.0	16	+460.5	−504.5	+1069.5	+104.5	
thrust_short_ra w	159	+283.0	23	+393.0	+110.0	+322.5	+39.5	
sw_absS_A_k10 (live Lot A)	31	+507.5	9	+127.5	−380.0	+612.0	+104.5	
sw_absS_B_k10 (live Lot B)	29	+846.5	8	+314.5	−532.0	+951.0	+104.5	
cx_absSB_live	24	+406.0	7	+496.0	+90.0	+510.5	+104.5	
cx_absSB_clip	25	+417.5	7	+369.0	−48.5	+522.0	+104.5	
abs_veto_60s	54	+33.5	13	−83.5	−117.0	+120.0	+86.5	
POOLED	—	+7746.5	—	+3883.5	−3863.0	+8809.0	+1062.5	

The floor is worse on 7 of 9 exits and costs −\$3,863 pooled — including both of the exits the live desk actually runs (Lot A −\$380, Lot B −\$532). The two exits it “helps” are thrust_short_raw (the unvetoes control, base +\$283 — a book so weak that removing 85% of it looks like progress) and cx_absSB_live at n=7. G2 improves **all nine**, +\$1,062.50 pooled. Same test, opposite outcomes.

8 · What to ship

ELEMENT OF PLAY #2	STATUS	THE NUMBER THAT DECIDES IT
Arm abs_veto_short by default	SHIP — unchanged	Armed 3.3% of tape; the identically-configured twin made +\$1,064.50 across the 96.7% it was benched. Untouched by this patch.

Veto: BUILDING × US-SESSION (ER30 ∈ [0.15,0.30) AND 13:30-20:00Z)	SHIP — the only filter that earns it	+\$207.00 · 15/15 winners · strip-3 + \$1,987.00 · OOS +\$913.50 vs baseline + \$809.00 · +\$1,062.50 across nine exits · net at the 98.5th placebo percentile. Still n=8 — re-look at n≥25.
Delete the US-session violent-whipsaw bench	SHIP — unchanged	That cell is n=12, +\$443.50, +\$37.00/trade, 58% win, and holds 6 of the top-15 winners.
Delete “bench on a stop-out pair” · fix the 22:00 auto re-arm	SHIP — unchanged	Untouched by this patch.
ER30 ≥ 0.35 arming floor	WITHDRAWN — REFUTED	5/15 winners · discards +\$1,363.00 of profitable trades · OOS +\$809.00→+\$363.00 · zero fires in wk29 · −\$3,863 pooled across nine exits · winner retention at the 77th placebo percentile (random expects 3.4 of 15) · a spike not a plateau (0.325/0.35/0.375 = +849/+927/+679) · its own live evidence dies at strip-1 (+\$550.50 → +\$14.50) · and it subsumes and deletes the BUILDING×US veto rather than adding to it.

Net effect on the play. Everything except the floor survives. The recommended arming policy is the one §9 of the rehab already specified — **default-ON, one entry veto (BUILDING × US-SESSION), two bad bench rules deleted** — running **151 fires for +\$2,497.50 (+\$16.50/trade)** on the 17-day reference book, against the +\$927.50 the floor would produce and the +\$881.00 the live-style router policy produced. Play #2’s topic, play, arms_when, number and kill fields all need the ER clause struck; the rationale stands as written.

THE PROCESS FAILURE, SAID PLAINLY — BECAUSE IT IS THE MORE VALUABLE FINDING

This report contained the refutation of its own #1 recommendation, in its own §5 table, at row E0.35 · 5/15, and in its own disposition table, at the line “ER30 / ER15 floors as an arming rule — **REFUTED**”. The play was assembled from Part 2.6 §8, which measured the **live** book; the refutation lived in Part 1.5, which measured the **continuous** book. Nobody joined them. **The generalisable rule: when a gate is being moved from router-gated to armed-by-default, every filter proposed alongside that change must be re-scored on the continuous book, because the population is about to change from 3.3% of the tape to all of it.** A filter fitted on the arming policy you are deleting is fitted on nothing. I would add one mechanical check to the Friday pipeline — every recommended filter gets a size-matched placebo on winner retention before it reaches *plays.json*. It is twenty lines and it would have caught this on Thursday.

9 · Disposition

LEAD	VERDICT	DETAIL / REVIVAL CONDITION
abs_veto_short ER30 ≥ 0.35 arming floor	REFUTED	Killed by: 13-rung plateau sweep (spike at 0.35), size-matched placebo on winner retention (77th pct, random expects 3.4/15), nine-exit cross-validation (−\$3,863 pooled), and strip-1 on its own live evidence (+\$550.50 → +\$14.50). Not revivable at any rung on this mechanism — the gate’s edge lives in low ER by construction.
ER floors as an arming rule, any gate	REFUTED — 5th gate	Consistent with deciders.ER_FLOOR = {} (08-01) and the +\$2,923/173-signal grind re-derivation. Revives only if a gate is found whose trigger is a clean-tape condition; thrust gates are the opposite by definition.
“The two halves of abs_veto want opposite ER regimes”	REFUTED	Killed by: the same band split on the continuous book. LONG is positive in all four bands (+\$4.04 / +\$9.40 / +\$10.69 / +\$7.77 per trade), not −\$0.90 above 0.35. Live-book sampling artefact.
Veto BUILDING × US-SESSION	LIVE (deploy, low confidence)	Unchanged from \$9. Survives every test in this patch including the placebo. n=8 — re-look at n≥25; kill if winner retention drops below 13/15 on fresh tape.
ATR14 floor 13-17pt on abs_veto_short	PARKED	The genuine winner-preserving filter (15/15 at A13/A15/A17; 100th placebo percentile) but it costs \$259–\$760 of net. Revives if slot contention becomes binding — it buys the same money in fewer trades.
ER30 ≥ 0.05 “frozen tape” cut	SHADOW	The one rung at or above baseline (+\$2,304.50, 14/15 winners, cuts 19 fires worth −\$14.00). Delta is +\$14 — inside noise, and it is a dead-tape cut, not an ER floor. Shadow it; promote only if the delta exceeds \$150 at n≥250.
Placebo gate on the Friday pipeline	PARKED → ACTION	Add a size-matched random-subsample check on winner retention to every filter before it enters plays.json. Revives as a finding the next time a recommendation and its own refutation ship in the same report.

Bottom line, Gaz: your instinct was exactly right and it was right for the exactly right reason — **\$/trade rewards shrinking the denominator, and every ER rung is doing that and nothing else.** The floor keeps 5 of 15 winners, throws away \$1,363 of profit, halves the out-of-sample week, and cannot be shipped alongside the BUILDING × US-SESSION veto because it silently replaces it. **Ship the veto. Arm the gate. Leave ER alone.**

REV2 · Q1 Is grind_long on or off on Monday? — one population, one number

WHAT THIS PATCHES

This **supersedes the deployment block in Part 1.5 §10, overturns the PARKED verdict in Part 2.5 LEAD 3, refuses the retirement recommended in Part 2.6 §6, and withdraws the ER-0.35 endorsement in Part 2.6 §(b)**. Everything below is recomputed from scratch on both disputed populations using the shipped gazbot7.decidors objects; nothing is inherited from the three sections it corrects.

Gaz — you spotted a real contradiction and you were right not to ship into it. Three sections gave three answers because **they scored three different populations and none of them said so**. Part 1.5 scored 744→155 reconstructed legs on the tick tape. Part 2.5 scored 128 lots the desk actually filled. Part 2.6 scored a shadow twin (grind_fast) that has no ATR floor at all and is therefore not this gate. I have now put the proposed config through both real populations, and the honest result is not the one either section reported.

The answer is ON — but with one literal changed and one literal DELETED, not two literals added. Raise ATR_FLOOR["grind_long"] from 10 to 22, and delete the ext_hi key so the gate falls back to its own default of 4.0. That is a clean revert of the 2026-08-01 Saturday window and nothing else. The live 128 — the population Part 2.5 used to park the gate — **agrees with the ATR floor and refutes the ext ceiling**: the floor turns those lots from −\$182.50 to **+\$210.00** at 22 and **+\$605.00** at the pre-08-01 value of 24, while the ext ≤ 3.0 ceiling Part 1.5 wanted to ship deletes 12 of them worth **+\$389.00 at a 50% win rate**. Both populations say the same thing about both numbers once you actually run them side by side.

△ AND THREE OF PART 1.5'S ROBUSTNESS CLAIMS DO NOT REPRODUCE

I re-ran its harness and matched its headline figures to the cent (744 legs / + \$1,929.9; 155 legs / +\$5,893.4; strip-5 -\$346.6 / +\$3,592.9; LODO + \$4,655.9...+\$6,244.8). So the harness is sound. But **(1)** “positive in all four ISO weeks” is **false** — ISO week 29 is **-\$568.2**, and Part 1.5's own weekly row sums to \$5,131 against a real total of \$6,689. **(2)** The out-of-sample leg is **three trades**: +\$796.0 becomes **-\$432.5** after strip-best-3, a test Part 1.5 ran everywhere except there. **(3)** The whole config is **21 events** — the legs that reach 4R net \$5,744.5 of the \$5,893.4, leaving the other 134 legs at **+\$148.9**. The recommendation survives all three. The confidence in Part 1.5 \$10 does not.

1 · The three answers, and the one line that explains all three

Nobody was careless. Each section scored the population it had, and the populations are not comparable — they differ in coverage, in config era and in exit ladder. Here is what each one actually measured.

SECTION	ITS VERDICT	WHAT IT ACTUALLY SCORED	THE PROBLEM WITH USING IT ALONE
Part 1.5 \$10	FIXED — ship $ATR \geq 22$ + ext 3.0	155 reconstructed legs, 11 tick-honest sessions, one config, one exit ladder	Complete signal coverage and one clean config — but 100% counterfactual, and its OOS leg and its ISO-week row do not hold up
Part 2.5 LEAD 3	PARKED — stays off, no more entry filters	128 lots the desk really filled, 10 sessions, four entry-config eras and eight exit mechanisms	It is 23% of the signal population, hand-picked by an arming schedule the same report calls anti-correlated with the gate's edge
Part 2.6 \$6	Write down its retirement	grind_fast, a shadow twin with no ATR floor , over the benched window	It prices the unfixed gate. “The unfloored twin bled \$2,549” is an argument for the floor, not for retirement
Part 2.6 \$(b)	$ER \geq 0.35$ floor validated hard	21 trades above ER 0.35 on a partly hand-gated ledger	Directly contradicts Part 1.5's REFUTED. Settled in \$7 below — the floor loses on both populations

So the report contains four verdicts on one gate, and the disagreement is entirely about which population is admissible. That is answerable, so I

answered it: I scored the proposed config on both, then measured exactly why they diverge.

2 · Population A — the 128 live router-gated lots, rebuilt

data/gazbot7.db, gate LIKE 'grind%', data_quality IS NULL, 10 sessions 2026-07-20...08-05. It reproduces Part 2.5 exactly: **128 lots, –\$958.00, 33.6% win**. To apply an entry filter to it I first had to recover what the tape looked like at each fill, so every lot was matched back to the 1-minute bar on which gate_grind actually fired, using the ext_hi that was shipped that day (4.0 before 08-01, 2.0 after). **126 of 128 matched** — 123 on the immediately-preceding completed bar, 3 one bar earlier. The 2 that never match any firing bar (ids 70, 171, –\$284.00) are excluded from every filter row below rather than silently counted as passes.

CUT APPLIED TO THE LIVE LOTS	N	NET	WIN%	\$/LOT	STRIP-3
all 128 lots — the Part 2.5 number	128	–\$958.00	33.6%	–\$7.48	—
LONG legs only	116	–\$466.50	34.5%	–\$4.02	—
matched LONG legs (the base for every row below)	114	–\$182.50	35.1%	–\$1.60	–\$965.00
Part 1.5's PROPOSED — ATR≥22 and ext≤3.0	32	–\$179.00	31.2%	–\$5.59	–\$838.00
ATR≥22 alone	44	+\$210.00	36.4%	+\$4.77	–\$572.50
ext≤3.0 alone	65	–\$382.00	30.8%	–\$5.88	–\$1,041.00
what the ATR floor removes	70	–\$392.50	34.3%	–\$5.61	—
what the ext ceiling removes	12	+\$389.00	50.0%	+\$32.42	—
the PRE-08-01 shipped cell — ATR≥24, ext≤4.0	36	+\$605.00	44.4%	+\$16.81	–\$177.50
ATR≥27, ext≤4.0	25	+\$870.50	48.0%	+\$34.82	+\$88.00

Read the two highlighted pairs against each other, because they are the whole finding. On the population Part 2.5 used to park the gate, raising the ATR floor works — it turns –\$182.50 into +\$210.00 at 22 and +\$605.00 at the pre-08-01 value of 24, without any help from the tick harness. And the extension

ceiling is the thing that destroys it: the 12 lots it deletes made +\$389.00 at a 50% win rate. That direction holds in every single cell of the live grid — the no-ceiling column is the best of the three at all six floors:

LIVE-128 GRID, NET\$ (STRIP-3)	EXT≤2.0	EXT≤3.0	EXT≤4.0 — NO CEILING
ATR≥18	−\$315.0 (−\$773.5)	−\$236.5 (−\$895.5)	−\$7.5 (−\$790.0)
ATR≥20	−\$219.5 (−\$678.0)	−\$198.0 (−\$857.0)	+\$46.5 (−\$736.0)
ATR≥22	−\$131.0 (−\$580.0)	−\$179.0 (−\$838.0)	+\$210.0 (−\$572.5)
ATR≥24	+\$17.5 (−\$431.5)	+\$76.0 (−\$583.0)	+\$605.0 (−\$177.5)
ATR≥27	+\$50.5 (−\$374.5)	+\$307.5 (−\$351.5)	+\$870.5 (+\$88.0)
ATR≥32	−\$324.0 (−\$284.0)	−\$6.5 (−\$564.0)	+\$111.5 (−\$563.5)

Every strip-3 in that grid except one is negative, which is the honest limit of this population: **32 lots over 8 days is not enough to prove anything on its own**, and 2 days carry it. What it can do is corroborate or contradict a direction derived elsewhere, and it does both — corroborating the floor, contradicting the ceiling.

Why this population cannot govern on its own — three measurements, not an opinion

CONFOUND	MEASURED	EFFECT
It is 23% of the signal population	On the 6 sessions that are both tick-honest and in the live set, the proposed rule admits 105 legs worth +\$2,127.00 . Real fills that the same rule would keep: 24, worth −\$209.50 .	Whatever gated the desk selected a −\$210 subset out of a +\$2,127 population. Per day: 07-24 12%, 07-27 22%, 07-29 0% , 07-30 43%, 07-31 25%, 08-05 12%.
It is not one exit ladder	8 distinct exit mechanisms across the 128: STOP 72, CHANDELIER 17, GIVEBACK 11 , TARGET 10, STOP_UNFILLED 8 , MANUAL_CLAIM 5 , ABSORPTION_CUT 4 , MAX_HOLD 1.	20 lots (+\$717.50) exited on three mechanisms that no longer exist, and 8 (−\$220.50) on a known defect. Re-exiting the 24 overlapping lots under today’s ladder gives −\$598.80 against the −\$209.50 booked — the live result is flattered by \$389 of operator hands.
It is not one entry config	ATR floor 20 → 24 (07-25) → dead through the _base(slot) wiring bug (07-29...08-01) → 10 (08-01); ext_hi unset (=4.0) until 08-01, then 2.0.	“The proposed config on the live 128” is partly a re-scoring of the config that produced those lots, not a counterfactual.
Execution is not a confound	Live fills on the 24 matched lots came in 246.00 points BETTER in total than the harness’s own first legal decision tick (mean −10.25 pt/lot) — worth +\$492 to the live book.	The gap between the two populations is not slippage or chase. Part 1.5’s “43-point chase” is real but measures from the bar-close print, and it lands on two lots the proposed floor rejects anyway (ATR 21.0).

3 · Population B — the 744→155 tick-honest leg set, re-run and extended

11 tick-honest sessions, 2026-07-24...08-07. Same harness, same shipped decider objects, \$2.00/pt and **\$1.50 per lot per round trip**. My re-run matches Part 1.5 to the cent on every headline it published, so the disagreements below are not harness differences — they are tests it did not run. **The PRE-08-01 row is new and it is the most important line in this section.**

CONFIG	N	NET	WIN%	\$/LEG	STRIP-3	STRIP-5	DAYS GREEN	LODO MIN... MAX
LIVE — ATR≥10, ext≤2.0	744	+\$1,929.9	35.6%	+\$2.59	+\$472.4	−\$346.6	6/11	+\$949...+ \$2,696
PRE-08-01 shipped — ATR≥24, ext≤4.0	141	+\$4,708.2	36.9%	+\$33.39	+\$3,250.8	+\$2,431.8	7/11	+\$3,328...+ \$4,967
Part 1.5 PROPOSED — ATR≥22, ext≤3.0	155	+\$5,893.4	38.7%	+\$38.02	+\$4,435.9	+\$3,592.9	10/11	+\$4,656...+ \$6,245
RECOMMEND ED — ATR≥22, ext ceiling DELETED	165	+\$6,081.2	38.8%	+\$36.86	+\$4,623.7	+\$3,804.7	10/11	+\$4,794...+ \$6,433

77% OF THE “FIX” IS A REVERT — AND THAT IS THE STRONGEST THING ANYONE CAN SAY FOR IT

+\$4,708 of the +\$6,081 was already available under the configuration the desk shipped before the 2026-08-01 Saturday window. This is not a new discovery bolted onto a losing gate; it is the undoing of a change that cost **\$2,778 across eleven tick-honest sessions** (+\$4,708 → +\$1,929.9). That matters for how much credence the number deserves: a revert to a previously-shipped config is a far weaker claim than a newly-optimised cell,

and weak claims are the ones that survive. It also means the pre-08-01 cell is a live candidate in its own right — see the last row of §11.

4 · Regime and time-of-day segmentation — scored on its HOME segments only

Regime labels are quantiles of this tape, not guesses: ATR14(1m) p25/p50/p75 = 10.4 / 13.5 / 18.2; ER30 p25/p50/p75 = 0.08 / 0.17 / 0.28. **clean-trend** $ER30 \geq 0.45$ · **building** $0.25 - 0.45$ · **high-vol expansion** (the taxonomy’s “violent-whipsaw”) $ER30 < 0.25$ with $ATR \geq 22$ · **normal-chop** $ER30 < 0.25$, $ATR\ 12 - 22$ · **dead-chop** $ATR < 12$. Session split is the 13:30 UTC cash open. The two tables below are scored on Part 1.5’s exact proposed cell ($ATR \geq 22$, $ext \leq 3.0$) so that the comparison against its published regime split is like-for-like; deleting the ceiling moves every bucket in the same direction and only changes the clean-trend row materially (see §7).

REGIME	LIVE CONFIG ($ATR \geq 10$, $EXT \leq 2.0$)			PROPOSED ($ATR \geq 22$, $EXT \leq 3.0$)		
	N	NET	\$/LEG	N	NET	\$/LEG
high-vol expansion — the home segment	105	+\$2,959.5	+\$28.19	114	+\$4,241.3	+\$37.20
building	50	+\$310.4	+\$6.21	32	+\$2,014.6	+\$62.96
normal-chop	446	−\$723.3	−\$1.62	0	—	—
dead-chop	129	+\$56.6	+\$0.44	0	—	—
clean-trend	14	−\$673.3	−\$48.09	9	−\$362.5	−\$40.28

The floor is not a P&L filter, it is a regime filter — it deletes normal-chop and dead-chop entirely, which is 575 of the live config’s 744 legs, and it does so by construction rather than by fitting. That is the difference between this and every ER filter the report has rejected.

SEGMENT × SESSION, PROPOSED CONFIG	N	NET	WIN%	\$/LEG	STRIP-3
	92	+\$3,950.9	39.1%	+\$42.94	+\$2,493.4

high-vol expansion / US (≥13:30Z) — the money					
building / US	27	+\$1,489.4	44.4%	+\$55.16	+\$487.9
high-vol expansion / overnight-pre-open	22	+\$290.4	31.8%	+\$13.20	−\$361.6
building / overnight-pre-open	5	+\$525.2	80.0%	+\$105.05	+\$5.2
clean-trend / US	8	−\$533.0	0.0%	−\$66.62	−\$339.2
clean-trend / overnight-pre-open	1	+\$170.5	100.0%	—	—

Session totals: **overnight/pre-open +\$986.2 on 28 legs, US +\$4,907.3 on 127**. The overnight block is thin but positive, which is why **no clock filter is recommended** — adding one costs \$986 to remove 28 legs whose only sin is being awake at the wrong hour. The ATR floor already is the clock, indirectly.

5 • Why the two populations disagree — the reconciliation, in dollars

This is the arithmetic that decides which number governs. Same six sessions, same proposed entry rule, two different observers.

SESSION	PROPOSED LEGS AVAILABLE	THEIR NET	LEGS THE DESK WAS ARMED FOR	WHAT THE DESK BOOKED	COVERAGE
2026-07-24	8	+\$680.6	1	−\$90.5	12%
2026-07-27	9	+\$79.2	2	−\$103.5	22%
2026-07-29	21	+\$662.1	0	\$0.00	0%
2026-07-30	35	+\$951.6	15	+\$360.5	43%
2026-07-31	16	+\$104.9	4	−\$284.0	25%
2026-08-05	16	−\$351.4	2	−\$92.0	12%
TOTAL	105	+\$2,127.00	24	−\$209.50	23%

That is the whole disagreement in one row. The gate's signal was worth +\$2,127 over those six sessions under the proposed rule; the desk filled less than a quarter of it, and the quarter it got was the losing quarter. This is the same anti-correlation Part 1.5 found inside a single week (−\$210.60 armed / +\$2,388.10 benched) — measured here over six sessions and three weeks instead, so it is not a one-week artefact. **Scoring an entry filter on that sample measures the router's timing, not the filter.**

THE ANSWER TO “WHICH GOVERNS”

The tick-honest population governs the CONFIG. The live population governs the CONFIDENCE. The tick set is the only one with complete signal coverage, one entry config and one exit ladder, so it is the only one on which an entry-filter comparison means anything — and an entry filter can only ever subtract from the population it is scored on, which is why applying it to 128 lots it did not generate produces a number (−\$179) that answers no question anyone asked. But the live set is real money on real fills, it is out-of-sample in time for the pre-08-01 cell, and it is the reason the second literal is not shipping. Use both, for different jobs. **Do not average them, and do not quote −\$958 or +\$5,893 again without saying which population it came from.**

6 · The parameters, swept — and why only one of them is load-bearing

Both axes on the tick-honest set, with the out-of-sample bar-replay leg beside each. **Every threshold here is an operator guess until it shows a plateau**, so both sweeps are shown whole, including the cells that argue against the recommendation.

ATR FLOOR (AT N EXT≤3.0)		IS NET	\$/LEG	STRIP-3	STRIP-5	DAYS GREEN	OOS NET
≥10 (live)	1050	+\$3,637.2	+\$3.46	+\$2,179.7	+\$1,336.7	6/11	+\$979.9
≥14	637	+\$2,406.1	+\$3.78	+\$948.6	+\$105.6	7/11	+\$462.5
≥18	325	+\$4,712.8	+\$14.50	+\$3,255.3	+\$2,412.3	8/11	+\$221.4
≥20	221	+\$4,764.6	+\$21.56	+\$3,307.1	+\$2,464.1	9/11	+\$579.3
≥22	155	+\$5,893.4	+\$38.02	+\$4,435.9	+\$3,592.9	10/11	+\$796.0
≥24	131	+\$4,543.0	+\$34.68	+\$3,085.5	+\$2,242.5	7/11	+\$1,086.1
≥27	99	+\$4,314.2	+\$43.58	+\$2,856.7	+\$2,013.7	7/11	+\$213.6
≥32	67	+\$3,261.6	+\$48.68	+\$1,804.1	+\$985.1	6/11	+\$354.6

18-27 is a genuine plateau — every cell ≥\$4,300 in-sample with strip-5 ≥\$2,000, against \$3,637 / \$1,337 at the live floor of 10 and a trough at 14. 22 is

its peak and 24 is within \$1,350 of it on a \$6,000 number, so **the exact literal is not load-bearing** and I would not defend 22 over 24 on this evidence. The OOS column has no structure at all and should not be used to pick a cell.

EXTENSION CEILING (AT ATR≥22)	N	IS NET	\$/LEG	STRIP-3	STRIP-5	OOS NET	READ
≤1.5	87	+\$2,852.2	+\$32.78	+\$1,394.7	+\$575.7	+\$571.0	far too tight
≤2.0 (the 08-01 change)	123	+\$3,871.5	+\$31.48	+\$2,414.0	+\$1,595.0	+\$925.2	the tax — costs \$2,210
≤2.5	135	+\$4,208.9	+\$31.18	+\$2,751.4	+\$1,932.4	+\$763.8	still paying
≤3.0 (Part 1.5's proposal)	155	+\$5,893.4	+\$38.02	+\$4,435.9	+\$3,592.9	+\$796.0	fine, but arbitrary
≤4.0 — the gate's own default, i.e. DELETE the key	165	+\$6,081.2	+\$36.86	+\$4,623.7	+\$3,804.7	+\$662.9	a revert, not a literal
≤6.0	191	+\$5,981.4	+\$31.32	+\$4,523.9	+\$3,704.9	+\$862.3	flat
no ceiling at all	197	+\$6,262.0	+\$31.79	+\$4,804.5	+\$3,985.5	+\$470.7	flat

Everything from 3.0 upward is the same number. The ceiling is not a filter with an optimum; it is a tax that is only non-zero below about 2.5. That is why the recommendation is to delete it rather than move it — deleting restores a value that was already in the codebase, needs no justification of its own, and is the only version the live 128 also endorses. Shipping `ext_hi = 3.0` would be inventing a third number to sit between two that both work.

7 · The ER question — Part 2.6 §(b) withdrawn

Part 2.6 says “grind_long makes +\$27.10 a trade above ER 0.35 ... the floor is exactly in the right place, do not let anyone argue it down.” Part 1.5 says ER floors are REFUTED at any level. They cannot both stand, and the ER floor decides whether the gate fires at all, so it has to be settled here. It is settled against the floor, on both populations.

ER30 FLOOR ADDED ON TOP OF ATR≥22	N	IS NET	\$/LEG	STRIP-3	RUNNERS KEPT	OOS NET
-----------------------------------	---	--------	--------	---------	--------------	---------

(TICK-HONEST)						
none	165	+\$6,081.2	+\$36.86	+\$4,623.7	21 of 21	+\$662.9
≥0.20	110	+\$2,002.3	+\$18.20	+\$699.8	12	+\$1,118.5
≥0.25	96	+\$1,940.1	+\$20.21	+\$657.6	11	−\$408.5
≥0.35 — the Part 2.6 endorsement	46	+\$657.4	+\$14.29	−\$346.1	4 of 21	−\$44.9

The 0.35 floor costs **\$5,424**, its strip-3 is negative, and it throws away **17 of the 21 runners the whole config exists to catch** — the identical failure mode Part 1.5 documented for the 0.20 floor. And the ER response is not monotone in either direction, which is the tell that it is not a variable at all here:

ER30 AT ENTRY (TICK-HONEST, ATR≥22, NO CEILING)	N	NET	WIN%	\$/LEG	STRIP-3
0.00 - 0.15	69	+\$3,881.4	44.9%	+\$56.25	+\$2,423.9
0.15 - 0.25	47	+\$664.4	29.8%	+\$14.14	−\$185.6
0.25 - 0.35	22	−\$251.5	27.3%	−\$11.43	−\$795.5
≥ 0.35	27	+\$1,786.9	48.1%	+\$66.18	+\$783.4

Best band, worst band, then best band again. On the live 128 the same slice gives −\$136.00 over 30 lots above ER 0.35, and +\$95.50 over 14 once the ATR floor is applied — i.e. nothing. **Part 2.6's +\$27.10/trade is 21 trades on a partly hand-gated ledger and it does not survive contact with either full population. The ER floor is withdrawn; Part 1.5's REFUTED stands.**

The clean-trend hole — smaller than Part 1.5 thought, and it closes on its own

Part 1.5 flagged “0-for-14 in its own namesake regime” as the one unturned stone. Two things. First, **most of it is the ext ceiling**: at $\text{ext} \leq 3.0$ the clean-trend book is 9 legs / −\$362.5, but with the ceiling deleted it is **11 legs / −\$5.5** — the legs that rescue a clean trend are the stretched ones the ceiling was blocking. Second, an ER30 ceiling at 0.45 does look attractive in-sample (+\$6,255.9, strip-5 +\$3,955.4) but it is a **spike, not a plateau** — 0.50 gives +\$5,844 and 0.40 gives +\$5,630 — and the OOS gradient runs the other way, preferring ever-tighter ceilings (0.30 → +

\$1,140 OOS). That is the same unstable-sign signature that killed the ER filters in the first place. **SHADOW, not a ship.**

8 · Robustness of what is actually shipping — including the parts that fail

TEST	LIVE (ATR≥10, EXT≤2.0)	RECOMMENDED (ATR≥22, CEILING DELETED)	VERDICT
11 tick-honest sessions	+\$1,929.9 / 744 legs	+\$6,081.2 / 165 legs	☐
strip-best-3 / strip-best-5	+\$472.4 / -\$346.6	+\$4,623.7 / +\$3,804.7	☐ survives deleting its 5 best
Leave-one-day-out (min...max)	+\$949...+\$2,696	+\$4,794...+\$6,433	☐ no single day load-bearing
Sessions green	6 of 11	10 of 11 (only 08-05 red)	☐
Placebo null — 300 random-entry draws from the same admitted minutes, same exits, same 120-min horizon	—	real +\$6,081.2 vs placebo mean + \$980.6, sd \$1,674, p95 +\$3,869, max of 300 draws +\$5,391.7	☐ 100th percentile — beats every draw. (Part 1.5's cell: +\$5,893.4 vs mean +\$925.3, max +\$5,290.9, also 100th)
Per ISO week, all 21 sessions	W29 -\$627.6 · W30 +\$1,884.0 · W31 +\$1,628.4 · W32 -\$287.7	W29 -\$565.0 · W30 +\$1,829.1 · W31 +\$3,345.3 · W32 +\$2,134.8	☐ NOT four-for-four — Part 1.5's claim does not reproduce
Out-of-sample, 10 bar-replay days	+\$667.2 / 382 legs, strip-3 — \$561.3	+\$662.9 / 104 legs, strip-3 — \$565.6 , 5 of 7 days green	☐ three trades — the OOS leg supports nothing either way
Concentration	16 legs reaching 4R net \$3,951.5 — 205% of the total; the other 728 legs net -\$2,021.6	21 legs reaching 4R net \$5,700.5; the other 144 legs net +\$380.7	⚠ the config is 21 events

⚠ SAY THIS OUT LOUD BEFORE SHIPPING

Strip the 21 runners and this configuration is **+\$380.70 over 144 legs — \$2.64 a leg, a rounding error above the fee.** That is not a defect: it is what a trend-continuation gate with a locked chandelier is supposed to look like, and Part 1.5's own exit sweep shows the tail beyond 2R is where the money is. It is also a genuine improvement on the live config, whose non-runner book is **-\$2,021.60**. But it means the honest confidence interval is set by **21 events over 11 sessions**, not by 165 legs, and a fortnight with no expansion days will produce a flat result that is not evidence the change was wrong. Write the review criterion now, before the first bad week: **judge it**

on 20 admitted legs minimum, and on whether 4R events appear at anything near the expected 12.7% rate — not on P&L alone.

9 • The retirement question — refused, with the reason

Part 2.6 recommended writing down grind's retirement or automating its arm, on the grounds that the router benched it for 119.3 of 120 hours and its twin grind_fast lost \$2,549 over that window. **That argument prices the wrong gate.** grind_fast is the un-floored shadow; on this tape the un-floored configuration loses in normal-chop and dead-chop, which is exactly the 575 legs the ATR floor deletes. A twin bleeding because it has no floor is the strongest possible argument for the floor.

Part 2.6's second branch — “or give it a hard auto-arm” — is the right one, and it needs no router work at all: **ATR_FLOOR is evaluated in-gate on every decision tick, with no polling lag.** Set it to 22 and the switch can simply be left ON; the gate self-gates and will not fire in tape where it cannot pay. That is what the bad-call ledger has been saying about this gate for a fortnight, and it removes the arm/bench decision entirely rather than automating it. **Do not also add a router arming rule this week** — the in-gate floor and a 5-minute polled ATR band would be two mechanisms doing one job, and the polled one would be strictly worse.

△ THE OPERATIONAL FACT THAT MAKES THIS URGENT

grind_long is **not** in `reactivate_gates.py::HOLD` (which currently holds only `nipc_long` and `nipc_short`). Under the standing operator policy, `gazbot7-gate-reactivate.timer` will therefore flip `grind_long=off` back to on at **22:00 UTC Sunday**, roughly 15 hours before Monday's cash open — regardless of what any section of this report concludes. **So “it stays off” is not a decision that can be taken by doing nothing; it requires a HOLD entry.** The choice in front of you on Saturday is: ship the two-line

config change and let it re-arm into a gate that self-gates at ATR 22, or add grind_long to HOLD. There is no third state where it is off by default.

10 • The exact change

```
# src/gazbot7/deciders.py
ATR_FLOOR: dict[str, float] = {"grind_long": 22.0, "capitulation_long":
10.0}
# ^ was 10.0 (set 2026-08-01). Reverted 2026-08-08 (rev2_q1) on 21
sessions, two
# independent populations. Plateau 18-27; 22 is its peak but 24 is
within $1,373
# on a $6,000 number (and beats it out-of-sample), so the literal is
NOT
# load-bearing. Deletes normal-chop and
# dead-chop entirely (575 of 744 legs). Independently corroborated
on the 128 live
# lots: ATR>=24 turns -$182.50 into +$605.00. Revert: set back to
10.0.

# src/gazbot7/slot_strategy.py, the grind_long SlotSpec
params={"slope_min": 0.4, "fast_slope": True},
# ^ DELETE "ext_hi": 2.0 (added 2026-08-01). Falls back to
gate_grind's own default
# of 4.0. Everything from 3.0 upward scores identically (+$5,893 / +
$6,081 / +$5,981
# / +$6,262) so there is no optimum to find; 2.0 is a $2,210 tax.
The live 128 agrees
# and more strongly: the ceiling deletes 12 real lots worth +$389 at
50% win.
# Revert: re-add "ext_hi": 2.0.
```

Do not change anything else. The exit (start_k 3.5 / lock_r 6.0 / lock_k 0.5), the 1.0×ATR stop, Lot A's 2.5R, base_size=2, slope_min 0.4, ext_lo 0.3, fast_slope=True all survive their own sweeps in Part 1.5 and I found nothing to

overturn there. One free note: with the floor at 22 the atr_split: 22 quiet-tape clip becomes inert for this gate — leave it in as a fail-safe.

11 · Disposition — every thread in this question, closed

LEAD	VERDICT	BASIS / REVIVAL CONDITION
grind_long ATR floor 10 → 22	LIVE	Plateau 18-27, strip-5 +\$3,805, 10/11 sessions green, deletes two whole losing regimes by construction, corroborated on 128 live lots (+\$605 vs −\$182). Ship Saturday.
ext_hi 2.0 → DELETED (default 4.0)	LIVE	2.0 costs \$2,210 in-sample; 3.0/4.0/6.0/ none are indistinguishable; the live 128 says the ceiling deletes +\$389 at 50% win. A revert, not a new literal.
ext_hi = 3.0 as a shipped literal (Part 1.5 §10)	REFUSED	Not by its in-sample number, which is fine, but by parsimony against two populations: it invents a third value between two that already work, and it is the only part of the proposal the live lots contradict.
“Grind stays off; do not re-test it with another entry filter” (Part 2.5)	OVERTURNED	Correct on its own population and for the right instinct, but that population is a 23% arming-schedule sample across 4 entry configs and 8 exit mechanisms. Named test: \$5 — 105 legs / + \$2,127 available, 24 legs / −\$209.50 taken.
“Write down grind’s retirement” (Part 2.6 §6)	REFUSED	Priced on grind_fast, the un-floored twin. Its bleed is the floor’s evidence, not the gate’s obituary. The “automate its arm” branch ships instead — via the in-gate floor, not a router rule.
ER30 ≥ 0.35 arming floor (Part 2.6 §b)	WITHDRAWN	Costs \$5,424, strip-3 −\$346, keeps 4 of 21 runners; the ER response is non-monotone (best band, worst band, best band). Live 128 agrees: −\$136 over 30 lots.
ER30 ceiling at 0.45 (the clean-trend patch)	SHADOW	+\$6,256 IS and +\$918 OOS look good, but it is a spike between 0.50 and 0.40 and the OOS gradient prefers tighter. Revive if the clean-trend book is still negative after 25 legs under the new floor.
Router arming rule for grind_long	PARKED	Not needed — the in-gate floor is the arming rule and has no polling lag. Revive only if the gate proves it needs a bench the floor cannot give.
ATR floor at 24 rather than 22 — the full pre-08-01 revert	PARKED — arguably the better ship	−\$1,373 in-sample (+\$4,708 vs +\$6,081) but + \$415 out-of-sample (+\$1,077.8 vs +\$662.9) and the only cell in this entire section with an OOS strip-3 near zero (−\$150.7), plus the shallowest bad week (W29 −\$36.4 vs −\$565.0). If anyone prefers a pure revert to 24 over a fresh 22, do not argue — on robustness it is the stronger of the two and only the in-sample headline separates them.
Stop 1.0 → 1.25 × ATR	SHADOW (unchanged)	Part 1.5’s finding, untouched here. One parameter change at a time.
grind_short	PARKED — never measured	

THE ONE-LINE ANSWER

grind_long is **ON on Monday**, with `ATR_FLOOR` at 22 and the `ext_hi` key deleted — a straight revert of the 2026-08-01 Saturday window, which is the single change both disputed populations endorse. The switch itself can be left armed because the floor gates it in-code. **Two literals do not ship; one changes and one is removed.** On the headline number: Part 1.5's + **\$3,964** is a delta (+\$5,893.4 − \$1,929.9), not a P&L, and it is fine arithmetic — the recommended cell makes it **+\$4,151** (+\$6,081.2 − \$1,929.9). But quote it with its window and its shape attached, always: **11 tick-honest sessions, 165 legs, of which 21 events are the entire result — against a live sample that never got to see 77% of them.** Part 2.5's −\$958 is a different population answering a different question and the two should never appear in the same sentence again without that label.

Harness: *scratchpad/grind_core.py* (shipped *gazbot7.deciders* objects, 1-second decision grid, tick-level protective stop), driven by *scratchpad/q1_live128.py*, *q1_match.py*, *q1_tick.py*, *q1_final.py*, *q1_er.py*, *q1_rec2.py*. Data: *data/gazbot7.db* (*data_quality* IS NULL throughout), *data/capture.db* 5s bars folded to 1m, ticks from *capture.db* unioned with *data/tape/ticks/MNQ/*.parquet*, *src/gazbot7/slot_strategy.py* + *deciders.py* git history for the config timeline. All P&L at \$2.00/point and \$1.50 per lot per round trip.

REV2 · Q2 Last Friday's card — did we do those things, and did they work?

PATCHES: Scope §Background requirement 3 (“Close the loop on last week's plays”) — the one named scope item the report had left open.

EXTENDS: `plays.json` with the status column it has never had, and Part 1 with a regime-segmented before/after on the five gates the old card

actually touched. **Source of truth:** reports/friday_v7/plays.pre-rev2.json (24 rows — not the 36 the scope says; I audited the 24 that exist), graded against docs/WEEKEND_2026-08-01_CHANGES.md, docs/SESSION_2026-08-04_05.md, data/config_journal.jsonl, the live source tree and the trade book.

THE ANSWER IN ONE LINE

Gaz — **we did most of it, and it did not pay.** Of the 24 plays, **15 shipped in some form, 4 were never built, 3 were overtaken, 1 was partly done and 1 was flatly violated.** The card claimed **£3,150 a week**; the five sessions that followed it (08-03→08-07) came in at **-\$772.50 on 110 trades.** And you are right to poke, because it is worse than a miss: **three of the top four plays on this week's new card are undoing or re-asking things off that card** — new #1 reverts old #1, new #4 reverts the exit trial old #17 asked us to keep running, and new #3 is old #4 asked again because **we never built it.** The single most expensive line is old #19: the card said “SHADOW news-impulse-pullback and do NOT arm it”, we armed it live the next day, and it is **-\$434.50 over 49 trades at 24.5% win.** Nothing in plays.pre-rev2.json records any of this — **there is no outcome field, and every grade below is hand-derived.**

1 · The gap, stated plainly — nothing graded this card, and nothing grades the new one either

plays.pre-rev2.json has exactly eight keys per row: id, topic, play, rationale, section_ref, suggested_mode, verification, money_gbp. **There is no status field, no outcome field, no ship date and no owner.** So a week later there is no machine-readable answer to “did we do it”, let alone “did it work” — the whole audit below had to be reconstructed by hand from two decision docs, the config journal, git log -S on the live source, and the trade book.

This week's plays.json does not fix it. It is much richer — it added rank, tier, number, mechanism, revert, kill, window, arms_when, exit, owner —

but it still carries **no outcome field**. Same hole, one week later, on 50 rows instead of 24. The Saturday slate has 8 items on it and **none of them is the play-ledger**. That is where this belongs: three fields (status, shipped_ts, outcome), written by the same Saturday hand that applies the change, plus a Monday grader that fills outcome from the trade book. Until that exists, every Friday re-derives its own history by hand and some of it will be wrong.

A second, quieter gap found while doing this. Five of the card's shipped changes are still uncommitted working-tree edits, seven days on: slot_strategy.py (the ext_hi 2.0 ceiling), direction_router.py (ER_TREND 0.15), run_census.py (the z-score cluster fix), reactivate_gates.py (the NIPC HOLD) and data/exit_overrides.json. git show HEAD: on each of those files still returns the old value. The desk runs the working tree, so this is live behaviour with no commit behind it — the exact condition SESSION_2026-08-04_05.md blames for the 07-31→08-02 ladder being “permanently unrecoverable”. The 08-08 05:00 startup row in config_journal.jsonl agrees: exit_overrides_uncommitted: true.

2 · All 24 plays, graded

DONE = shipped as written or held as instructed. MODIFIED / AMENDED = shipped with a changed number. REVERTED = shipped then undone. NOT-DONE = never built. SUPERSEDED = overtaken by later evidence. VIOLATED = we did the opposite. “Worked?” is the live trade book 08-03→08-07 unless a source is named; every live figure is data_quality IS NULL, MNQ, \$2.00/point, \$1.50/round-turn.

#	PLAY (OLD CARD)	MODE / £	STATUS	EVIDENCE IT IS / IS NOT LIVE	WORKED?
1	grind_long — kill the ER floor + ATR 24, add an ext_atr ≤ 2.0 ceiling	live · £700	DONE — MODIFIED	deciders.py:203 ATR_FLOOR grind_long = 10.0 (card asked 12;	NO. Live 6 legs, – \$172.00, 16.7% win. Rehab replay over 11 sessions: the live

				weekend cut it to 10, “12 is one notch the wrong side of a cliff, worth \$941”). ER_FL00R = {}. slot_strategy.py: 161 ext_hi 2.0.	config earns + \$1,930 on 744 legs where ATR 22 earns +\$5,893 on 155. New card #1 reverts it.
2	PROMOTE abs_veto_55s two-sided (long and short)	live · £350	SUPERSEDED premise was wrong	Both SlotSpecs have existed since eb57f64, 07-25 16:32 — six days before the card. The play's stated premise (“short-only is what the live desk runs today”) was already false when written.	SPLIT. Long + \$566.00 / 32 / 62.5% — best gate on the desk. Short — \$559.00 / 9 / 0% win. Structure was never the problem; arming was. New card #2 re-asks it.
3	STOP_UNFILLED per-tick guard — leave it alone	live · bug	DONE	Guard untouched; concrete-front-month branch 99e3c11 deliberately not merged.	YES. STOP_UNFILLED still 26 all-time, last 07-28T08:02 — zero in 11 days / 110 trades since.
4	Chandelier runner — make the lock regime-aware (6.0R in TREND, 3.0R in BUILD)	live · £400	NOT-DONE	One flat number everywhere: lock_r=6.0 at slot_strategy.py: 162, :290 and :509. No regime key anywhere in the exit path.	n/a — never built. New card #3 asks for the same mechanism again (+ \$3,049 / 20 sessions, 20-of-20 LOO). Biggest un-shipped item on the card.
5	capitulation_long — require_flip=False + 1.0R	live · £250	DONE, THEN REVERTED <24h	slot_strategy.py: 164: “★2026-08-02 REVERTED to require_flip=True / 2.0R”. Killed by the MFE-is-not-a-win-rate error — 78% claimed, 29% real on the target-vs-stop race.	Revert was right. On the restored config: +\$84.50 / 6 / 33.3%, against −\$122.00 / 24 before. Thin, but the right sign.
6	BUILD grind an in-code 55s-style continuation veto	shadow · £150	NOT-DONE	deciders.py:169 VETO_GATES = {abs_veto_long, abs_veto_short}. Grind is not in it, and no grind branch exists in tournament.py's veto buffer.	n/a — never built.
7	Do NOT build an event-driven fast-bench service	skip	DONE (by inaction)	No such service exists. The 08-04 tape watcher was deliberately built alert-only — “a second writer to gate_switches.env is the two-routers hazard”.	Held. No cost.
8	Router fader-bench ER-trend floor 0.20 → 0.15	live · £60	DONE uncommitted	direction_router.py :47 ER_TREND = 0.15. git show HEAD: still reads 0.20.	Not measurable. Its only live consumer is veto_counter_regime on exhaustion_short, whose whole post-card book is 6

						trades. Not computed.
9	Untradeable-day meter — hold the 65 cutoff and the 40/40/20 blend	live	DONE — AMENDED 08-05	untradeable.py: 133,137 still 40/40/20 and 65. Amended: the stop-rate leg now needs $n \geq 10$ to vote, and the weights renormalise to 50/50 when it cannot.	Weakly yes. Scored predictively over 13 days it reads -0.50 (correct sign); the rival 0-10 gauge read -0.06 and was retired an hour after it was built.	
10	Make the CAUTION band (45-64) an action, not a label	shadow	NOT-DONE	Zero occurrences of the string CAUTION in scripts/router_tick_durable.py. The band is still printed and ignored.	n/a — never trialled.	
11	RULE — at a violent open, arm VETO-gates only, hold the churners	live · £250	DONE (in the prompt)	Codified in router_tick_durable.py:218/239/241 — churner / counter-trend language, plus the 08-04 arm-side rebalance.	Did not rescue the window. 13:30-14:45 post-card: – \$1,005.50 on 58 trades, every class negative (veto-gates –\$191.50/11, churners – \$172.00/6, NIPC – \$376.00/37). Thin n; not a refutation, but no help.	
12	exhaustion_short is the survivor — lead with it, judge on live fills	live · £240	SUPERSEDED	Kept live all week (exhaustion_short=on today). But this week's rehab re-graded the entry to PARKED, sub-friction by ~\$2/lot.	NO. Pre-card + \$222.50 / 69 / 55.1%. Post-card –\$162.00 / 6 / 33.3%. The “steadiest earner” framing did not survive the next five sessions.	
13	Regime-grade exhaustion_short's exit; stop tuning abs_veto's Rs	shadow	SUPERSEDED	What shipped 08-01 was a flat A 0.75R / B tight in exit_overrides.json — not a regime grade. The regime-graded version was never built.	New card #4 goes further back — all the way to the 07-25 fixed 8pt / 12pt / 120s.	
14	Keep the adaptive exit-width selector; do not replace it with a fixed exit	live · £250	DONE	slot_strategy.py: 214 still sets adaptive_exit = True roster-wide, with NIPC explicitly exempt.	Not computed this window — no selector-attributed A/B was run on the post-card book.	
15	rgv_short — bench and shadow, do not retire	shadow	DONE	gate_switches.env: rgv_short=off. The rg_short_* family is still on the slate; only three narrow variants are in RETIRED.	YES. Pre-card it was –\$416.00 / 43 / 44.2%. Post-card: zero fills — the bench held through five sessions, including the nightly re-arms.	
16	Wire nothing from either grind-exit shadow	shadow	DONE (by inaction)	Nothing wired.	Held. No cost.	
17	exhaustion_short 0.5R/1.5R fade-scalp	shadow	DONE	Trial ran (06bdd75, 07-31 09:08), was re-cut to A 0.75R /	Verdict: revert. All six fixed-point exits beat the incumbent,	

	trial — let it run, re-grade with real n				B tight on 08-01, and has now been re-graded — which is exactly what the play asked for.	+\$604 over 6 dense-book days as the mean of 10 polling phases. $\Delta \sigma$ across phases is \$418 — a single run of that harness is noise. New card #4.
18	On chop days the OFF switch beats every gate — bench, don't build	live · £500	PARTLY DONE		No chop scalper was built (correct). The bench itself only half-landed: the desk was near-flat in dead chop (2 trades) but took 44 trades in violent whipsaw.	MIS-AIMED. This week's donation was not in chop: NORMAL-CHOP was +\$118.50 / 27 while VIOLENT-WHIPSAW was −\$688.00 / 44. The avoidance idea is right; the label was wrong. See §4.
19	NIPC — SHADOW it, explicitly do NOT arm it	shadow	VIOLATED		Armed live at 1 lot the next day (slot_strategy.py: 205–210, two-sided, 13:00–15:00 UTC window). Weekend doc records the reason: the acceptance replay had used \$5/RT; at the true \$1.50 it was +\$820, not +\$2,676.	NO — the most expensive line on the card. −\$434.50 / 49 / 24.5% win. Short side −\$337.00 / 16 / 6.2%. Hit your −\$400 kill and was benched 08-06 13:09; it is now the only entry in reactivate_gates.H0 LD.
20	Do not arm chop_turn_fade; log the config and expect zero	skip	DONE (by inaction)		Never armed.	Held. No cost.
21	Stop hunting flow-footprint entries on MNQ — the family is buried	skip	DONE		No footprint gate armed. This week's gf_full_FLOW-LED / VACUUM / OPEN-NEWS runs are the fixed census clusters — i.e. play #22's intended follow-on, not a re-dig of the buried family.	Held.
22	One-line fix to the census FLOW-LED / VACUUM rule	live · bug	DONE uncommitted		run_census.py:35–53 is now a trailing-2h z-score (FLOW_Z, z ≥ 1). git show HEAD: still has the old def cluster(hour, flow, mv, amp).	YES. Fire rate 66.5% → 22.0% of bars. This is the play that made this week's Movement-3 clusters mean anything.
23	NEXT FRIDAY: re-run the chop-scalp harness on MES, 07-12→07-17	skip / next-Friday	NOT-DONE — SUBSTITUTED		The harness ran, on MNQ. The string MES appears 0 times in gf_chop_scalp.md. The named subject was not run.	The MNQ run is a strong result in its own right (fading chop is a 47.2% proposition, $z = -3.61$; 9 of 540 costed cells positive) — but it does not answer the question that was asked.
24	Do NOT arm anything off the idle-gate lab	skip	DONE (by inaction)		Nothing armed. Lab re-run this week and re-confirmed as an honest null.	Held.

2b · The tally, and what the £3,150 bought

STATUS	PLAYS	£ CLAIMED	THE HONEST READ
DONE as written / held as instructed	12	560	Mostly the cheap ones. Two genuinely earned their place: the STOP_UNFILLED hold (#3) and the census fix (#22).
DONE — MODIFIED	1	700	grind (#1). Shipped at ATR 10 not 12, and being reverted this week.
DONE — AMENDED	1	0	The meter (#9), amended for the right reason (n<10 must not vote).
DONE, THEN REVERTED	1	250	capitulation (#5). Shipped Saturday, gone by Sunday.
PARTLY DONE	1	500	The chop-day bench (#18) — aimed at the wrong regime label.
NOT-DONE	4	550	#4 chandelier regime-lock, #6 grind veto, #10 CAUTION-as-action, #23 MES.
SUPERSEDED	3	590	#2 (premise false at write time), #12, #13.
VIOLATED	1	0	#19 NIPC. Cost —\$434.50.
TOTAL	24	3,150	Realised over the five sessions that followed: —\$772.50 on 110 trades.

⚠ **Do not read that last row as “the card lost \$772”.** The card had no attribution mechanism, so no play can be scored in isolation — the router was arming and benching over the top of every one of them all week. What the row honestly says is: **a card that promised £3,150 a week was followed by a week that lost money**, and nothing in the card's own artefact would have told you that.

3 · The uncomfortable part — three of this week's top four plays point backwards

You will spot this the moment you read the new card, so here it is up front rather than buried.

NEW CARD	WHAT IT DOES	TO WHICH OLD PLAY	WHY THE OLD ONE WAS WRONG
#1 grind_long: ATR floor 10→22, ext ceiling 2.0→3.0	REVERT, and past the original — the floor was 24 before the card, the card asked for 12, we shipped 10, and we are now going to 22.	old #1 (£700, “verified”, live)	The old number was a blocked-book argument: “the ATR 24 floor blocked +\$982, therefore drop it.” What it never checked was what the gate does with the permission. At ATR 10 grind takes 744 legs to earn +\$1,930; at ATR 22 it takes 155 legs to earn +\$5,893, keeps 19 of 19 target winners, and is green in 10 of 11 sessions.
#3 chandelier: regime-key the width (ER30 ≥ 0.25 & ATR ≥ 22 → 6R)	RE-ASK — same mechanism, opposite framing (widen the good cell rather than narrow the bad one).	old #4 (£400, “verified”, live)	The old play was not wrong. It was never built. A £400 “verified / live” play sat unbuilt for a week and nothing flagged it — which is the outcome-field gap costing money rather than embarrassment.
#4 exhaustion_short: exit back to fixed 8pt / 12pt / 120s	REVERT to the 07-25 ladder, undoing both the 07-31 fade-scalp trial and the 08-01 re-cut.	old #17 (let the trial run) + old #13 (regime-grade it)	Old #17 was correct process — it said “two signals is not a grade, run it and re-grade with real n”, we did, and the answer is that the replacement was worse. That is a trial working, not a failure. What went wrong is old #13: we shipped a flat 0.75R/tight instead of the regime grade it asked for, so a whole week ran on an exit nobody had proven.
#2 abs_veto_short: arm by default (+ an ER30 floor, since withdrawn in REV2 Q0)	RE-ASK of the arming half.	old #2 (£350)	Old #2 asked for a structural change (two slots) that had already existed since 07-25. The real problem was never structure — it was that the router benched the short side ~91% of the week and armed it into the worst tape of it. See §4.

The pattern across all four: every one of the old plays was scored on a counterfactual population — what a filter blocked, or what a slot would allow — and none was scored on what the gate then did with the trades it

took. Blocked-book arithmetic tells you what you gave up; it never tells you what you bought. That is the method note worth carrying into next Saturday, and it is the same family as the MFE-is-not-a-win-rate error that killed old #5.

4 · Did it work? The regime-segmented before/after

The card shipped over the weekend of 08-01/02, so **PRE = 07-16→08-02** (n=489) and **POST = 08-03→08-07** (n=110, five sessions). Regime is tagged **at each trade's entry minute** from the tape itself, not from the clock: 1-minute MNQ bars rebuilt from the 5s lake (`gazbot7.lake.connect()`, `(bar_ts//60)*60` — integer division, not `/`), ATR14 and a 30-minute efficiency ratio computed on those bars. VIOLENT-WHIPSAW = ATR $\geq$ 22 with ER30 $<$ 0.30; CLEAN-TREND = ER30 $\geq$ 0.50 with ATR $\geq$ 15; BUILDING = ER30 0.30-0.50; DEAD-CHOP = ATR $<$ 10 and ER30 $<$ 0.30; NORMAL-CHOP = the rest. 6 of 599 trades fall outside the tape and are shown as UNCLASS.

REGIME AT ENTRY	PRE — 07-16→08-02				POST — 08-03→08-07			
	N	NET \$	\$/TRADE	WIN%	N	NET \$	\$/TRADE	WIN%
CLEAN-TREND	16	+861.0	+53.81	56.2%	4	+43.5	+10.88	50.0%
NORMAL-CHOP	203	-2154.0	-10.61	39.4%	27	+118.5	+4.39	48.1%
VIOLENT-WHIPSAW	116	-841.5	-7.25	37.9%	44	-688.0	-15.64	27.3%
BUILDING	134	-552.0	-4.12	43.3%	33	-216.0	-6.55	30.3%
DEAD-CHOP	14	-315.5	-22.54	14.3%	2	-30.5	-15.25	0.0%
UNCLASS	6	-226.5	-37.75	33.3%	0	—	—	—
ALL	489	-3228.5	-6.60	39.9%	110	-772.5	-7.02	33.6%

What that says in English. The one thing that is true in both halves and needs no defending: **the desk only makes money in a clean trend**, +\$53.81/trade before the card and +\$10.88 after, and it gets almost no

chances — 20 trades out of 599. Everything else is a leak of a different size. The card's own diagnosis (old #18: “chop days are the donation”) was **true before it shipped and stopped being true after**: normal chop flipped from $-\$10.61/\text{trade}$ to **$+\$4.39$** , while the money moved into **violent whipsaw** — $-\$15.64/\text{trade}$ on 44 trades, 40% of the post-card book. That is a high-ATR, low-efficiency tape, and it is precisely the tape a low ATR floor invites a momentum gate into. The regime split and the grind revert are the same finding seen twice.

Split by time of day, the post-card whipsaw is **entirely a US-session phenomenon**: all 44 of those trades opened after 13:30 UTC. Overnight/pre-open post-card is 24 trades for $-\$325.50$, and 14 of those are normal-chop at $-\$21.04/\text{trade}$.

4b · Gate by gate, the five the card actually touched

GATE	PRE 07-16→08-02				POST 08-03→08-07				READ
	N	NET \$	\$/TR	WIN%	N	NET \$	\$/TR	WIN%	
abs_veto_long	61	-121.0	-1.98	34.4%	32	+566.0	+17.69	62.5%	The desk's best gate post-card, and the only one that improved. All 32 fills in BUILDING (+ \$214.00/14) or NORMAL-CHOP (+ \$352.00/18) — it never fired in whipsaw.
capitulation_long	24	-122.0	-5.08	37.5%	6	+84.5	+14.08	33.3%	Post-revert. All the money is 2 BUILDING trades (+ \$207.00); the 4 chop trades gave back — \$122.50. n=6 — a

									direction, not a result.
grind_long	92	−282.0	−3.07	32.6%	6	−172.0	−28.67	16.7%	All 6 legs fired on 08-05 inside three minutes of the open. See the honesty note below.
exhaustion_ short	69	+222.5	+3.22	55.1%	6	−162.0	−27.00	33.3%	The card's “steadiest earner”. Five sessions later it is the entry the new card parks.
abs_veto_sh ort	26	+358.5	+13.79	42.3%	9	−559.0	−62.11	0.0%	Nine lots, nine losers. Four of them fired on 08-05 at ATR 34-38 with ER30 0.12- 0.13 — the cleanest untradeabl e day in the record.
rgv_short	43	−416.0	−9.67	44.2%	0	—	—	—	Play #15 worked exactly as written: benched, shadowed, not retired, zero fills.
nipc (both sides)	0	—	—	—	49	−434.5	−8.87	24.5%	Should never have been live. Long — \$97.50/33; short — \$337.00/16 at 6.2% win. 30 of the 49 fired in violent whipsaw.

THE HONESTY NOTE ON GRIND — THE LIVE FILLS DO NOT SUPPORT THE REVERT

I have to say this because the table invites the wrong conclusion. All **6** post-card grind legs fired at **ATR14 ≥ 22** by my own tape reconstruction (23.98, 23.98, 33.50). **The reverted ATR-22 floor would not have blocked a**

single one of them. So $-\$172.00$ is not evidence for new card #1 — the case for the revert rests entirely on the 11-session replay leg-count (744 → 155), and this week's six live fills are silent on it. Anyone quoting the live number as support for the revert is quoting the wrong number. What the six fills do show is that a low ER floor let grind buy three minutes into a violent open at ER30 0.05–0.22 — which is an argument for an efficiency condition, not the ATR one being shipped.

Same discipline applied to **abs_veto_short**: an $ER30 \geq 0.35$ floor would have kept four of the nine losers (the 08-06 pair at ER 0.49 for $-\$129.50$ and the 08-07 pair at ER 0.42 for $-\$120.00$) and only blocked the ones the whipsaw already made obvious. That is an independent corroboration of **REV2 Q0**, which withdrew that floor on completely different evidence.

5 • What I could not resolve, and what to build Saturday

Not computed, and I am not going to guess: per-play P&L attribution (the card carries no attribution mechanism and the router was arming over the top of every play all week); the adaptive exit selector's post-card contribution (#14 — no A/B was run on this window); the ER_TREND 0.15 change's effect (#8 — its only live consumer has a 6-trade book). Each of those is **not computed**, not zero.

Two numbers I lifted rather than re-derived, and which are flagged as such: the census fire-rate 66.5%→22.0% (weekend record) and the meter's -0.50 predictive correlation over 13 days (08-05 session record).

The Saturday build, in priority order:

1. **The play ledger.** Three fields on every row of `plays.json` — `status` (PROPOSED / SHIPPED / REVERTED / NOT-DONE / SUPERSEDED), `shipped_ts`, `outcome` — written by whoever applies the change, plus a Monday job that fills `outcome` from the trade book on the gate the play names. It is on none of the eight Saturday items today. This audit took hours of archaeology that three fields would have made a query.

2. **Commit the working tree.** Five live behaviours currently exist only as uncommitted edits. `config_journal.jsonl` already flags `exit_overrides_uncommitted: true` at every startup and nothing acts on it — make that flag fail the sweep.
3. **Build old #4 before re-deriving it a third time.** The chandelier regime-key has now been recommended twice, at £400 and at +\$3,049, and shipped zero times.

Disposition of every lead in this section: old #3, #7, #15, #16, #20, #21, #22, #24 — **LIVE / HELD**, working. Old #1 — **REFUTED** by the 11-session leg-count replay, reverting as new #1. Old #5 — **REFUTED** by the target-vs-stop race (29%, not 78%), already reverted. Old #12 — **PARKED**, revive when the entry clears friction by ~\$2/lot on a fresh sample. Old #13, #17 — **SUPERSEDED** by new #4. Old #18 — **LIVE but re-aimed**: the avoidance target is violent whipsaw, not chop. Old #19 — **PARKED** (in HOLD); revival needs a fresh operator decision plus a working fix for the diagnosed trigger-timing drag, not a re-arm. Old #2, #4, #6, #10 — **SHADOW / unbuilt**, carried forward: #4 as new #3, #2 as new #2, #6 and #10 still owned by nobody. Old #23 — **PARKED**, revival condition: run the harness on MES as written, or formally drop it from the queue — do not let it roll a third week unstated. Old #8, #9, #14 — **LIVE, not computed** this window.

Method

Trades: `data/gazbot7.db`, WHERE `data_quality` IS NULL AND `closed_at` IS NOT NULL — the two MD_STREAM rows (538/539, -\$255.50) are excluded throughout. 599 trades, 07-16→08-07. MNQ, \$2.00/point, **\$1.50 per round trip**. Tape: `gazbot7.lake.connect()` over the Parquet lake (`symbol='MNQ'` filtered explicitly; V7 holds no 1m bars after 07-13, so 1m bars were rebuilt from 5s with $(\text{bar_ts}/60)*60$ — 25,519 minutes, 07-13→08-07). Code state read from the working tree and cross-checked against `git show HEAD:` and `git log -S` per file. Config state from `data/config_journal.jsonl` (17 startup rows) and `data/gate_switches.env`. **Thin n is stated everywhere it occurs:** the post-card window is five

sessions, and every per-gate post-card cell except `abs_veto_long` (32) and `nipc` (49) is under 10 trades. None of the post-card figures survive a robustness battery on their own — they are a status report, not a finding.

REV2 · FIX0 Saturday #3 is not a chandelier play, and as written it cannot be typed into any file the desk reads

PATCHES: Part 2.6 §8 WHAT TO ACTUALLY DO Saturday #3 and `plays.json` rank 3 (`chandelier-regime-key-width-0808`) — the card is **REPLACED** in full. **ALSO PATCHES:** the Part 1.5 lead-table row that reads “chandelier: regime-key the **width** ($ER_{30} \geq 0.25$ & $ATR \geq 22 \rightarrow$ Lot B 6R)” — same wrong key, same wrong mechanism. **EXTENDS:** Part 1.5 `rehab_chandelier` §6 with the resolved `scaleout_slots()` spec before/after, which the dossier prescribed but never printed. The dossier’s verdict is unchanged and still ships; only the card is wrong.

THE ANSWER IN ONE LINE

Gaz — the card describes a mechanism that does not exist, keyed on a variable the config file cannot hold, and quotes the numbers of the other variant. There is no ER key in `exit_overrides.json`: the loader reads **four** fields (`a_r`, `b`, `atr_split`, `lo`) and **silently drops anything else** — I typed `"er_split": 0.25` into it and the loader returned the entry without it, no error, no log line. There is no “6R chandelier” either: a numeric `b` resolves to `exit="scalp"`, `target_r=6.0`, a **fixed rung**; the thing that is a chandelier with a 6 in it (`b: "wide" → lock_r=6.0`) is the runner-up the dossier explicitly declined, and picking it costs **\$905** of the claimed edge. What the dossier actually ships is three config numbers — **`atr_split 22→24`, every non-NIPC `b→6.0`, every non-NIPC `a_r→3.0`** — and the third of those silently

triples abs_veto_long's Lot A from 1.0R to 3.0R. The ER30 key is the er_split code change that \$6 **parks**. Ship the config. Queue the code.

1 • Four claims on the card, checked against the code that would execute them

Basis: src/gazbot7/slot_strategy.py at d0f5bbd-era HEAD — _load_exit_overrides() (L245-282), _lot_b() (L285-296), scaleout_slots() (L299-327), the clip freeze at **L421-422**, and _BIG_RUN at L224. Live slate confirmed GAZBOT7_TOURNAMENT_SLATE=scaleout from systemctl show gazbot7-tournament (a drop-in overrides the unit file's tournament), so the “deployed” column below is what the desk is running right now, not a reading of source.

THE CARD SAYS	WHAT THE CODE DOES	VERDICT
Key the wide cell on ER30 ≥ 0.25 AND ATR ≥ 22, owner “exit_overrides.json + tournament restart”	_load_exit_overrides builds {"a_r", "b"} then optionally {"atr_split", "lo_a_usd", "lo_b_r", "lo_b_floor"}. No ER key of any kind . Unknown keys are not rejected — they are never read. Demonstrated: {"a_r": 2.5, "b": 6.0, "atr_split": 22, "er_split": 0.25, "lo": {...}} → loader returns {'a_r': 2.5, 'b': 6.0, 'atr_split': 22.0, 'lo_a_usd': 40.0, 'lo_b_r': 1.75, 'lo_b_floor': 60.0}. The clip freezes at L422 on f.atr < spec.atr_split — ATR alone .	UNEXECUTABLE . Typing the card into the file is a silent no-op on the ER half — the worst failure mode there is, because the restart succeeds and the desk reports a clean slate build.
“Lot B runs a 6R chandelier”	_lot_b: a numeric b → replace(base, exit="scalp", target_r=float(b), stop_atr_mult=1.0) — a fixed 6R rung, no trail . b: "wide" → exit="chandelier_lock", chandelier_start_k=3.5, lock_r=6.0 , lock_k=0.5 — that lock_r=6.0 is exactly why “6R chandelier” reads plausible.	WRONG MECHANISM, and it costs money . Dossier \$1.3 home-cell shape ladder, regime gate and Lot A held constant: fixed 6R +\$3,049 · fixed 4R +\$2,801 · lock 3.5/6.0/0.5 +\$2,144 · lock 3.5/8/0.5 +\$1,736 · chandelier k1.5 +\$761 · give-back +\$581. Reading “chandelier” and setting b: "wide" forfeits \$905 .
“Everything else keeps today's width” / “Lot A untouched”	The dossier's shipped JSON changes three things on all six non-NIPC gates: atr_split 22→24, b→6.0, a_r→3.0 . Resolved, Lot A moves on 6 of 6 gates — including abs_veto_long 1.0R→3.0R and exhaustion_short 0.75R→3.0R . See §2.	FALSE . “Lot A untouched” is the single most misleading line on the card: it is the change most likely to alter fill behaviour on the gate that trades most.
Numbers: +\$3,049, LODO min +\$1,604, better in all four ISO weeks, week −\$643	Those are the er_split code-change figures (\$6, er_split 0.25 / atr_split 22 / a_r 2.5 / b 6.0). The config-only package the card's owner-line describes measures: archive — \$2,932→+\$92 (+\$3,024) , LODO 20/20 min + \$1,259 , pre-window −\$1,859→ +\$891 , symptom	MIS-ATTRIBUTED . Every headline on the card belongs to the variant \$6 tells you not to ship yet. The honest config-only claim is \$25 smaller, a \$345 weaker LODO floor, and 2 of 4 weeks, not 4 of 4.

2 · The resolved `scaleout_slots()` spec, before and after — the report's own standing rule

“Verify config via `scaleout_slots()`, never source” is a house rule precisely because this slate has silently dropped four things already. So here it is, actually resolved: I called `scaleout_slots()` twice under different `GAZBOT7_EXIT_OVERRIDES` — once against the live data/`exit_overrides.json`, once against the dossier \$6 JSON — and diffed the 16 sub-slot specs field by field. Nothing below is read off source.

SUB-SLOT	BEFORE — LIVE NOW	AFTER — DOSSIER CONFIG-ONLY	LOT A R	ATR_SPLIT
grind_long_A	scalp 2.5R	scalp 3.0R	2.5 → 3.0	22 → 24
grind_long_B	chandelier_lock k3.5 / lock_r 6.0 / lock_k 0.5	scalp 6.0R	—	22 → 24
capitulation_long_A	scalp 1.5R	scalp 3.0R	1.5 → 3.0	22 → 24
capitulation_long_B	chandelier k1.5 (tight)	scalp 6.0R	—	22 → 24
abs_veto_long_A	scalp 1.0R	scalp 3.0R	1.0 → 3.0 (3×)	22 → 24
abs_veto_long_B	scalp 1.5R	scalp 6.0R	—	22 → 24
rgv_short_A	scalp 1.5R	scalp 3.0R	1.5 → 3.0	22 → 24
rgv_short_B	chandelier k1.5 (tight)	scalp 6.0R	—	22 → 24
exhaustion_short_A	scalp 0.75R	scalp 3.0R	0.75 → 3.0 (4×)	22 → 24
exhaustion_short_B	chandelier k1.5 (tight)	scalp 6.0R	—	22 → 24
abs_veto_short_A	scalp 1.5R	scalp 3.0R	1.5 → 3.0	22 → 24
abs_veto_short_B	scalp 2.5R	scalp 6.0R	—	22 → 24
nipc_long_A/B, nipc_short_A/B	scalp 2.0R / 2.5R, no clip	identical	—	off → off

THE THING NOBODY WROTE DOWN — THIS EDIT DELETES THE CHANDELIER FROM THE LIVE SLATE

Count the exit field across the 16 resolved sub-slots. **BEFORE: four chandeliers** (`grind_long_B` lock, plus `capitulation_long_B` / `rgv_short_B` / `exhaustion_short_B` tight). **AFTER: zero.** Every sub-slot on the desk becomes `exit="scalp"`. So the play filed under the topic “chandelier: regime-key the width”, whose owner-line says “the chandelier goes back to a single width”, in fact **retires the chandelier from the**

running slate entirely. That is not a contradiction of the dossier — §6 verdicts the trail machinery **SHADOW**, “keep it in the code, do not re-arm it”, and this is exactly that. It is a contradiction of the card, which sells the same edit as a chandelier improvement. **Rename the play. It is a Lot-B rung play.**

3 · The ER30 key: what it really is, and the one thing to fix before anyone deploys it

The ER key is a **code change**, three touch points, and §6 parks it: “Ship the config-only edit today; queue *er_split* as the follow-up and re-measure before promoting it.” Its cost is a deploy, not a JSON edit.

TOUCH POINT	FILE : LINE	CHANGE
1. new spec field	slot_strategy.py : ~99 (beside atr_split)	er_split: float = 0.0 — 0 = off, so no existing gate moves
2. loader accepts it	slot_strategy.py : ~267 (_load_exit_overrides)	validate and add to entry, inside the same try/except as atr_split so a typo drops only the key
3. the freeze	slot_strategy.py : 421-422	self._exit_lo[spec.tag] = f.atr < spec.atr_split gains the ER leg. bars is already in scope (decide() signature, L397) and efficiency_ratio is already imported (L43) — nothing new to plumb

⚠ BEFORE YOU CODE IT — THE DOSSIER SPECIFIES TWO DIFFERENT POLICIES AND CALLS THEM ONE

§1.3 and §5 define the studied cell as **wide iff (ER30 ≥ 0.25 AND ATR14 ≥ 22)** — 84 of 495 signals, 17%. The §6 code sketch says clip iff *f.er15 < er_split* as well as *f.atr < atr_split*, whose complement is **wide iff (ER ≥ 0.25 OR ATR ≥ 22)**. Those are different populations, not phrasings — and the sketch also reads **er15** where the study cut on **ER30**. Both are reported at the same **+\$3,049**, so at least one label is loose and **I have not re-run the harness to say which** — not computed. **Implement the AND-cell on ER30, and re-measure before promoting.** This is the

same class of error as the card itself: a mechanism described in prose that the code would not produce.

4 • Three collisions with other Part 1.5 cards that the JSON resolves silently

The dossier’s JSON block is a **whole-file replacement**. Applied verbatim it overwrites rows that two other rehab cards also claim.

COLLISION	THE OTHER CARD SAYS	WHAT TO DO
grind_long Lot A 2.5R → 3.0R	Part 1.5 §1 grind_long: “ Do not touch anything else — ... Lot A’s 2.5R ... all survive their own sweeps.” Its own §5.5 grid disagrees with itself: grind_long’s deployed 2.5R/wide-lock ranks 119 of 135 in its home cell at — \$1,308, best pair 1.0R/1.0R — and the dossier flags that as n=21, thin, contradicting a prior study .	Ship 3.0R as part of the measured package. The +\$3,024 was measured with it; carving grind_long out is a config nobody scored. Note it explicitly in the change log so the grind_long card’s “don’t touch” is knowingly overridden, not accidentally.
exhaustion_short Lot A 0.75R → 3.0R, Lot B tight → 6.0R	Part 1.5 exhaustion_short: EXIT FIXED — SHIP revert to the fixed 8pt/12pt/120s, and “remove exhaustion_short from <i>BIG_RUN</i> at the same time” (confirmed <i>BIG_RUN</i> = {grind_long, abs_veto_short, abs_veto_long, exhaustion_short}, L224). The gate is also PARKED as sub-friction. Meanwhile §5.5 says the fader guess ranks 131/135 in that cell and the best pair is 3.0R/6.0R.	These two cannot both be applied. The gate is PARKED, so the money at stake is ~zero — but pick one and write it down. Recommended: honour the exhaustion_short card (fixed 8/12/120s), delete its exit_overrides row, and remove it from <i>BIG_RUN</i> in the same commit — otherwise the default hands it A@2.5R + B wide-chandelier, the worst Lot B in every cell.
nipc_long / nipc_short rows carry no lo block	Part 1.5 nipc_short: SHIP THE lo BLOCK ANYWAY to both nipc legs.	Applying the chandelier JSON verbatim silently declines that card. If both ship, the nipc rows must be {"a_r":2.0,"b":2.5,"atr_split":22,"lo":{"...}}. Resolved-spec check confirms today’s nipc legs have atr_split 0.0 — clip off.

5 • The replacement card — paste this over plays.json rank 3

FIELD	CORRECTED VALUE
id	lotb-rung-widen-atrsplit24-0808 (the old id names a mechanism this play does not touch)
topic	Lot B rung: widen it to a fixed 6R and move the quiet-tape clip’s ATR gate 22→24

play	Edit the six non-NIPC rows of data/exit_overrides.json to {"a_r": 3.0, "b": 6.0, "atr_split": 24, "lo": {"a_usd": 40, "b_r": 1.75, "b_floor_usd": 60}}. NIPC keeps its proven {"a_r": 2.0, "b": 2.5} pair. Three numbers per row, no code change, no new mechanism.
mechanism	Lot B fixed-R rung + the quiet-tape clip's ATR threshold. Not the chandelier — this edit resolves every sub-slot to exit="scalp" and leaves zero chandeliers in the slate.
arms_when	entry ATR ≥ 24 only, frozen at entry (slot_strategy.py:422). Below 24 the \$40 / 1.75R-\$60 clip is unchanged — it is rank 1 of 28 in two of the four regime cells and 100% of quiet-tape behaviour is preserved. There is no ER term. Exit selector only; arms and benches nothing.
exit	Lot B: fixed 6.0R rung above the split. Lot A: changes on all six gates to 3.0R — abs_veto_long 1.0→3.0, exhaustion_short 0.75→3.0, capitulation_long/rgv_short/abs_veto_short 1.5→3.0, grind_long 2.5→3.0.
number	Config-only, as shipped: archive −\$2,932 → +\$92 (+\$3,024) over 20 sessions, 20/20 LODO folds positive, min +\$1,259 ; pre-window −\$1,859 → +\$891; symptom week −\$1,073 → −\$799 . ▲ better in only 2 of 4 ISO weeks (W31 +4,557 vs +1,683 and W32 −799 vs −1,073 win; W29 −954 vs −910 and W30 −2,712 vs −2,632 lose). Not strictly dominant.
owner	Garrath — six rows of data/exit_overrides.json, then systemctl restart gazbot7-tournament (overrides are read at slate build; live slate confirmed scaleout).
revert	git checkout data/exit_overrides.json + restart. The file is git-tracked and every startup journals its resolved ladder to config_journal.jsonl (scripts/config_at.py --diff).
kill	If the ATR ≥ 24 population goes negative over 20 fresh sessions, or the LODO minimum drops below zero on the next re-cut, revert. ▲ Read this first: the parameter-plateau test FAILED (4 of 35 threshold cells positive), 6 of 16 in-cell sessions are green, two days (07-27, 07-23) carry two-thirds of the edge , and strip-best inverts at top-8 (−\$422). The direction is robust; the knob is not. Size accordingly.
follow-up (NOT this play)	er_split — the ER30 key, a code change on three lines. Better on every axis (archive +\$25, week +\$156, LODO floor +\$345, 4/4 vs 2/4 ISO weeks) but margins are small and it needs a deploy. QUEUED, not shipped, and its AND/OR spec must be resolved first (\$3).

6 · Lead status — every thread in this patch, closed

LEAD	STATUS	BASIS / REVIVAL CONDITION
Lot B fixed 6R rung + atr_split 24 + a_r 3.0 (config-only)	LIVE SHIP	+\$3,024 archive, 20/20 LODO min +\$1,259. Ships as three numbers per row. Disclose 2-of-4 ISO weeks and the failed plateau test.
The ER30 × ATR regime key (er_split)	PARKED	Revives when the AND-cell-on-ER30 is implemented at slot_strategy.py:422 and re-measured on fresh tape. Not shippable via config — the loader drops the key silently.
“Lot B runs a 6R chandelier” (b:"wide", lock_r 6.0)	SHADOW	Runner-up in the home cell (+\$2,144 vs +\$3,049) and it is uncapped, which a fixed rung is not. Keep exit_chandelier_lock in the code; do not arm it on this report.
	REFUTED	

The card as written — ER key in JSON, “6R chandelier”, “Lot A untouched”, er_split numbers

exhaustion_short override row vs the 8/12/120s revert

PARKED pending one decision

abs_veto_long Lot A staying at 1.0R

SHADOW

Killed by direct execution: the loader returns the entry with er_split absent; _lot_b maps numeric b to exit="scalp"; the resolved diff moves Lot A on 6 of 6 gates; the four quoted figures are §6's er_split row.

Mutually exclusive; gate is PARKED so the stake is ~\$0. **Revives** the moment either card is applied — and whichever way it goes, _BIG_RUN must be edited in the same commit.

The dossier grants no exception — its JSON lists abs_veto_long at a_r 3.0 and the +\$3,024 was measured with it there. The per-cell grid is the counter-argument: 1.0R/1.0R is the best momentum pair in **2 of 4** cells (+\$534, -\$155), 3.0R in the other two (+\$1,495, +\$397). **The delta of carving it out is not computed.** Ship the package as measured; shadow the carve-out.

Bottom line, Gaz. The dossier is right and the card mistranslated it three ways at once — wrong keying variable, wrong exit mechanism, wrong variant's numbers — and every one of those is the kind that fails quietly: the JSON loads, the tournament restarts, the log says nothing. **What to actually do Saturday: change three numbers on six rows of exit_overrides.json, restart the tournament, and print scaleout_slots() afterwards to confirm all sixteen sub-slots read scalp with atr_split 24.** Expect Lot A to move on every gate, expect the chandelier to vanish from the slate, and expect roughly **+\$3,024** over a 20-session archive whose in-cell edge two days carry two-thirds of. The ER key is next week's deploy, not this weekend's edit.

FIX 1 **abs_veto_long — resolving HOLD #1 against Part 2 §4, §5 and §9**

PATCHES: the 50-play card, **HOLD #1** (hold-absveto-long-0808) · **Part 2 §4** (promotion verdict) · **Part 2 §5** (relegation table row) · **Part 2 §9** (per-rung exit deliverable, abs_veto_long row) · **EXTENDS Part 2 §11** (disposition table, two new rows).

Garrath — you caught a real one. The report told you three different things about the same gate in the same document: the card said touch nothing, §4 and §5 said

add a US-session guard and widen the exit, and §9 put a specific exit change in the deliverable table. Then neither change appeared on the card. I went back to the fills and re-derived both recommendations properly. **The card's ACTION was right and the two recommendations were wrong** — but not for the reason the card gave, and both had been published without the robustness tests §9 claimed for them.

THE RESOLUTION IN FIVE LINES

1. The session guard is REFUTED as specified. The $-\$258.50$ overnight figure is 12 lots in one week. Over the full live history the non-US book is **+\$0.59 a lot**, not a loser — because the all-time overnight loss sits almost entirely in **ASIA 00-07 UTC, which the desk already blocks at config level**. A US-only guard would bench LONDON, which is positive for this gate and is the desk's only positive block.

2. The exit widening is REFUTED as stated. Re-derived with the tests §9 promised: the SCALP-CHOP edge is **98% from 2 of 9 days** and goes **negative on strip-best-3**; the MED-TREND "sign flip" is **one day** (08-04). The deployed ladder already ranks **15th of 70** — top quintile. It is not a broken exit.

3. One narrower lead SURVIVES and is new: in the US session only, in SCALP-CHOP, the deployed ladder ranks **51st of 70** and `tightChand x2` is the single cell in this whole analysis that passes mean, leave-one-out AND strip-3 together ($+\$37.99 / +\$16.70 / +\$8.40$). **n=13 — a SHADOW, not a deployment.**

4. Two instrument defects. The script behind §9 computes no leave-one-out and no strip-3 at all — those JSON fields are `null` — yet §9's preamble claims every cell "survived leave-one-day-out". And §9's `abs_veto_long` row silently mixes two different populations in one row.

5. Net effect on the card: HOLD #1 stays, one clause changes. No config change to `abs_veto_long` this week. But "do not touch its exit" overclaims — it becomes "no change this week; two leads parked/shadowed with named revival tests".

A — THE CONTRADICTION, STATED EXACTLY

WHERE	WHAT IT SAID	IMPLIED ACTION
50-play card, HOLD #1	"No change. Do not add a filter, do not add an ER floor, do not touch its exit."	do nothing
Part 2 §4 (promotion verdict)	"abs_veto_long gets a session guard (US only)" — on $-\$258.50$ overnight vs $+\$824.50$ US	add an arming filter
Part 2 §5 (relegation row)	tag: "KEEP + session guard + widen the exit" — "it banked 1.0R/1.5R on moves that ran 7-12R"	filter AND exit change
Part 2 §9 (deliverable)	"CHANGE → scalp 1.5/6.0 or tight chandelier $\times 2$ — $+\$11.70$ / $+\$10.80$ per entry vs $+\$5.40$ deployed. MEDIUM (n=35)"	edit exit_overrides.json
Part 2 §11 (disposition)	Carried one row only — "abs_veto LONG side ... SHADOW". The exit change had no disposition row at all.	—

Two published recommendations, neither on the card, one with no disposition anywhere. Below, both are re-derived from the fills and given a verdict.

Method — so you can check it

Source: data/gazbot7.db trades, gate LIKE 'abs_veto_long%', data_quality IS NULL, 93 closed lots 2026-07-26→08-07. **Fee verified at the row level: \$1.50/round-trip and pnl_usd is already net of it** (checked against $(\text{exit} - \text{entry}) \times \$2 \times \text{qty}$ on four rows). Entries are lots collapsed to the minute — that reproduces the card's "16 entries / 32 lots / $+\$566.00$ / 62%" and its leave-one-out quartet exactly, so my grouping matches the card's. Exit repricing re-runs scripts/live_entry_exit_reprice.py's grid (70 policies, 5s bars, same-bar stop-first, 1-ATR stop, 60-min horizon, VPP, FEE = 2.0, 1.5) with leave-one-day-out and strip-best-3 added, which the original does not compute.

B — LEAD 1: THE US-SESSION GUARD → REFUTED AS SPECIFIED

B1. The headline number does not survive leaving the week. §4 and §5 generalise from one week. Widen the window and the overnight bucket changes sign:

WINDOW	ON LOTS	ON NET \$	ON \$/LOT	US LOTS	US NET \$	US \$/LOT	READ
this week 08-03→08-07 (what \$4/\$5 quote)	12	−258.50	−21.54	20	+824.50	+41.23	looks decisive
multi-slot era 07-29→08-07 (the deployed config)	45	+122.50	+2.72	36	+589.00	+16.36	overnight is POSITIVE
all live history 07-26→08-07	53	−152.00	−2.87	40	+597.00	+14.93	mildly negative

The direction is stable — US beats non-US in all three windows. The magnitude is not: −\$21.54 a lot becomes +\$2.72 a lot on the very config that is deployed today. And the gap is not statistically distinguishable from zero: **US − ON = +\$17.79/lot, se \$14.01, t = +1.27.**

B2. And this is the one that kills it — "overnight" is not one thing.

Decomposing the non-US book by session block:

BLOCK (UTC)	LOTS	NET \$	\$/LOT	STATUS
ASIA 00-07	14	−175.00	−12.50	ALREADY BLOCKED — no_open_asia=True since 08-05. All 14 lots are 07-27/07-28/07-30. Cannot recur.
LONDON 07-13:30	21	+123.00	+5.86	POSITIVE. A US-only guard benches this.
US 13:30-21	40	+597.00	+14.93	the good block, undisputed
EVENING 21-24	18	−100.00	−5.56	mildly negative, 6 days — the only real residue
⇒ NON-US, excluding the already-blocked Asia window	39	+23.00	+0.59	a scratch, not a loser. This is the population a guard would actually remove.

Read that bottom row against \$4's −\$258.50. The all-time non-US loss the guard was reasoning from is concentrated in ASIA, and the desk stopped opening in Asia on 08-05. What a US-only guard would remove from here forward is a book running at **+\$0.59 a lot** — of which the biggest piece, LONDON, is **positive**. That is the same trap as the "don't trade until the US open" idea that was already tested and refuted at desk level: **LONDON is the desk's only positive block**, and this gate is positive there too. This week's LONDON reading (−\$186.50 on 10 lots) is a sign-flip on tiny n — the 11 LONDON lots before this week made **+\$309.50**.

B3. Robustness and the n that would settle it.

BUCKET (ALL LIVE)	LOTS	NET \$	STRIP-BEST-3	WORST LOO	GREEN DAYS
NON-US (ON)	53	-152.00	-825.00	-485.00 (drop 07-29)	2/9
US	40	+597.00	+300.50	+93.00 (drop 08-04)	5/6

The US side is genuinely robust; the non-US side is dispersed rather than reliably bad (sd \$78.02 a lot vs \$57.09). **To separate a \$17.79/lot gap from zero at 80% power you need ~302 lots per arm.** There are 53 and 40. At this week's fill rate (32 lots) that is roughly four to five months — so this will not be settled by waiting on live P&L, and it should not be deployed on the strength of 12 lots.

Sensitivity to the 08-06 shared-account incident — it does not change the verdict

Trades **567/568** (abs_veto_long_A/B, entered 07:08:01, booked TARGET + \$56.00 at 07:17:22 / 07:19:03) sit inside the documented 08-06 cascade: the day-rider flattened that position at **07:08:21** at 29536.75, booked separately as rows 569/570 CROSS_DESK_FLATTEN -\$95.50. The gate's real outcome on that entry was about **-\$30.50**, not +\$56.00 — an \$86.50 overstatement, and it lands in the LONDON bucket. Repricing it: this week's non-US goes -\$258.50 → **-\$345.00**; the multi-slot-era non-US stays **positive** (+\$122.50 → +\$36.00); non-US-excluding-Asia goes +\$0.59 → **-\$1.63 a lot** — still a scratch, not the -\$21.54 the guard was priced on. **The correction makes the week look worse and the conclusion unchanged**, which is why the verdict does not rest on it. (I checked and discarded a second signal here: rows 567/568 carry no exec ids, but neither do 501 of 599 closed trades, so that is not evidence of anything.)

C — LEAD 2: THE EXIT WIDENING → REFUTED AS STATED, ONE NARROWER LEAD SURVIVES

§9's preamble says each cell "won its (gate × rung) segment **and then survived leave-one-day-out**". For abs_veto_long that is not true — no leave-one-out was ever run (see D1). Here it is, run:

RUNG	POLICY	\$/ENTRY	RANK	WORST LOO	STRIP-3	VERDICT
------	--------	----------	------	-----------	---------	---------

SCALP-CHOP n=35, 9 days	scalp 1.0/1.5 ← DEPLOYED	+5.36	15/70	−0.89	−3.02	top quintile already
	scalp 1.5/6.0 (\$9's pick)	+11.66	1/70	+5.12	−4.53	fails strip-3
	tightChand x2 (\$9's pick)	+10.75	3/70	+1.54	−2.14	fails strip-3
MED-TREND n=12, 6 days	scalp 1.0/1.5 ← DEPLOYED	−4.68	14/70	−18.73	−30.81	—
	tightChand x2 (\$9's "sign flip")	+5.23	1/70	−23.39	−43.80	one day — REFUTED

Three independent things kill it.

(i) Concentration. The SCALP-CHOP edge is $+\$6.30/\text{entry} \times 35 = +\220.50 **total**. Per day it is: 07-30 **+\$80**, 08-07 **+\$136**, and the other seven days sum to **+\$5**. **98% of the edge is 2 of 9 days.**

(ii) Strip-3. Remove the three best entries and every candidate goes negative — challengers and the deployed ladder. The whole cell is carried by three fills.

(iii) The MED-TREND sign flip is a single day. Total edge $+\$118.90$; 08-04 alone contributes **+\$156**. Without 08-04 the "improvement" is negative. Worst-LOO $-\$23.39$, strip-3 $-\$43.80$.

Add the shape of the winner: scalp 1.5/6.0 sits on the **edge of the grid** (B_RS tops out at 6.0), with scalp 3.0/6.0 second — both at B=6.0. That is a boundary solution, the classic overfit signature. And only **14 of 70** policies beat what is deployed. A ladder in the top quintile of its own grid is not the thing costing this gate money.

C2. The interaction — and the one lead that survives. The two recommendations were never independent. Splitting the same repricing by session shows \$9's all-session cell is a blend of two populations that want opposite things:

POPULATION	N	DEPLOYED \$/ENTRY	DEPLOYED RANK	POLICIES BEATING IT	BEST OF 70	WHAT IT MEANS
US session, SCALP-CHOP (5 days)	13	+18.13	51/70	50/70	+38.60	here the exit really is too tight
non-US, SCALP- CHOP (7 days)	22	−2.19	6/70	5/70	+0.33	no exit saves it — best of 70 is a scratch

That bottom row is the better-powered version of the session argument: on 22 non-US entries across 7 days, **the best of 70 exit policies earns \$0.33 an entry**. Nothing downstream rescues that book — the same diagnosis \$5 correctly reached for abs_veto_short. But it argues for arming discipline (which the router

already does daily), not for a hard-coded US-only guard that also deletes a positive LONDON book.

And the one genuinely surviving cell:

CELL	N	\$/ENTRY	RANK	WORST LOO	STRIP-3	VERDICT
tightChand x2 — US session, SCALP-CHOP only	13	+37.99	2/70	+16.70	+8.40	passes all three cuts — the only cell here that does
scalp 3.0/6.0 — same cell	13	+38.60	1/70	+18.68	−11.51	higher mean, fails strip-3

This is **not** the recommendation §9 made — it is narrower (US-only, chop-only) and it prefers the chandelier over the 6.0R target. At n=13 it is a **SHADOW**. It is also the first thing I have seen that makes the "banked 1.0R/1.5R on moves that ran 7–12R" observation in §5 into something testable rather than an anecdote.

D — TWO INSTRUMENT DEFECTS FOUND WHILE DOING THIS

#	DEFECT	EVIDENCE	BLAST RADIUS
D1	scripts/ live_entry_exit_reprice.py computes no leave-one-out and no strip-3 — at all. §9's preamble states every cell "survived leave-one-day-out".	172 lines, no LOO/strip-3 code. In scratchpad/ live_entry_reprice.json the fields loo_worst, strip3 and days are null on every cell. Its per- entry best column is a per-entry oracle (mean \$62.91 vs the cell winner's \$11.66), so it cannot stand in for one either.	Every live-entry cell in §9 — not just abs_veto_long. Every "CHANGE→" row sourced from this file has an unearned robustness claim.
D2	§9's abs_veto_long row mixes two populations across its four columns without saying so.	SCALP-CHOP/MED-TREND/BIG- TREND (n=35/12/3) come from live_entry_reprice.json — live fills, 70 policies. The STAY-OUT cell ("best of 53 is −\$40.80") comes from exit_ladder_v2.json — the ungated shadow book, n=47, 53 policies. The 53 vs 70 policy count is the tell.	The row reads as one study. A reader comparing its STAY-OUT cell to its SCALP-CHOP cell is comparing live fills to shadow fills.
D3	Trades 567/568 are booked as TARGET wins against a position the day-rider had already flattened, and are not data_quality-flagged.	Entered 07:08:01; flatten at 07:08:21 booked separately (569/570, −\$95.50); "exits" stamped 07:17:22 / 07:19:03. Only rows 538/539 carry an EXCLUDE: tag in the whole table.	+\$86.50 of overstated P&L reaching every report number and the router's decision inputs. Small, but it is exactly the class the 08-06 fix was meant to close.

E — THE EDITS THAT REMOVE THE CONTRADICTION

FILE / SECTION	ACTION	REPLACEMENT TEXT
plays.json — hold-absveto-long-0808, play field	EDIT ONE CLAUSE	"No config change this week. Do not add an ER floor — the long twin's ER profile is the mirror of the short's. Two leads that earlier drafts proposed are not being deployed and are dispositioned in Part 2 §11: a US-

		only session guard (REFUTED — it would bench a positive LONDON book) and a wider Lot B (REFUTED as stated; a narrower US-chop-only tightChand x2 variant is SHADOW at n=13)."
Part 2 §4, promotion verdict paragraph	DELETE	Delete "abs_veto_long gets a session guard (US only)". Replace with: "abs_veto_long changes nothing — its apparent overnight problem is an ASIA problem, and Asia is already blocked (see FIX 1 §B2)."
Part 2 §5, abs_veto_long row	RETAG	Tag becomes "KEEP AS DEPLOYED". Notes column becomes: "Not a problem. The – \$258.50 non-US figure is 12 lots; excluding the already-blocked Asia window the non-US book is +\$0.59/lot all-time. The 7-12R give-up is real but only measurable in the US session (FIX 1 §C2)."
Part 2 §9, abs_veto_long row + preamble	DOWNGRADE	SCALP-CHOP cell: "NO CHANGE — deployed ranks 15/70; challengers fail strip-3 and 98% of the edge is 2 of 9 days. THIN". MED-TREND cell: "REFUTED — the sign flip is one day (08-04). THIN". Preamble: strike "and then survived leave-one-day-out" for all live-entry cells, or run the test (D1).
Part 2 §11, disposition table	ADD 2 ROWS	Section F below, drop-in. Also amend the existing "abs_veto LONG side — decaying, needs a session guard" row: its title asserts the conclusion this patch refutes.

F — DISPOSITION ROWS (drop into Part 2 §11)

LEAD	VERDICT	THE NUMBER THAT DECIDED IT	REVIVAL CONDITION / THE TEST THAT KILLED IT
abs_veto_long US-only session guard	REFUTED as specified	Excluding the already-blocked ASIA window, the non-US book is +\$23.00 on 39 lots = +\$0.59/lot, of which LONDON is + \$123.00 / +\$5.86 a lot — positive. The –\$258.50 that justified it is 12 lots in one week; the same bucket is + \$122.50 over the multi-slot era. US–ON = +\$17.79/lot, t=+1.27.	Killed by: session-block decomposition — the all-time non-US loss is concentrated in ASIA 00-07, which no_open_asia=True has blocked since 08-05, so a US-only guard removes a positive LONDON book instead. Revive only as the narrow EVENING form (bench 21-24 UTC: –\$5.56/lot, 18 lots, 6 days) once that block reaches n≥60 lots and is still negative on strip-3 and worst-LOO. Settling the broad US-vs-rest claim on live P&L needs ~302 lots per arm (80% power) against 53/40 today — so run it as a shadow arm, never as a deployment.
abs_veto_long exit widening (§9: scalp 1.5/6.0 / tightChand x2, all sessions)	REFUTED as stated	SCALP-CHOP n=35: deployed + \$5.36, rank 15/70; challengers +\$11.66 / +\$10.80 but strip-3 – \$4.53 / –\$2.14, and 98% of the edge (+\$216 of +\$220.50) is 2 of 9 days. MED-TREND n=12: the sign flip is 08-04 alone (+\$156 of a +\$118.90 total edge), worst-LOO –\$23.39.	Killed by: leave-one-day-out + strip-best-3 + per-day edge concentration — none of which the producing script computes (D1). Also a boundary solution: the winner sits on the grid edge (B=6.0). Not revivable in this form; superseded by the row below.

NEW — tightChand x2 for
abs_veto_long, US session x
SCALP-CHOP only

SHADOW

n=13 over 5 days: +\$37.99/
entry vs +\$18.13 deployed,
rank 2/70, worst-LOO +\$16.70,
strip-3 +\$8.40 — the only cell in
this analysis passing all three
cuts. In the same rung the
deployed ladder ranks 51/70
with 50 of 70 policies beating it.

Thin n is a shadow, never a kill.
Promote to a paired live A/B
when the cell reaches $n \geq 30$ US-
session SCALP-CHOP live
entries with rank still top-10,
worst-LOO positive and strip-3
positive. Run it as an
exit_overrides.json shadow
arm on the same feed — do not
edit the deployed ladder off a
backtest.

\$9's "survived leave-one-day-
out" claim (all live-entry cells)

INSTRUMENT DEFECT

loo_worst, strip3, days are null
on every cell of
live_entry_reprice.json; the
producing script contains no
such computation.

Fix: port the LOO/strip-3 block
from exit_ladder_lab_v2.py into
live_entry_exit_reprice.py and
re-run every \$9 live-entry cell
before any of them is quoted
again. Until then no \$9
"CHANGE→" row is deployment-
grade.

WHAT THIS COSTS YOU IF I AM WRONG

Nothing is being changed, so the only risk is inaction. If the session guard was right, the cost of not deploying it is the non-US book's forward run rate — **+\$0.59 a lot on the tradeable (post-Asia) population**, i.e.

approximately zero. If the exit widening was right, the cost is the SCALP-CHOP edge outside its two carrying days — **+\$5 total across seven days**.

Both are rounding errors, which is the strongest argument for leaving the gate alone: **the two changes the report argued about were worth almost nothing even if true**, while the one cell that survives every cut (US × chop, n=13) is worth ~\$20 an entry and is not yet measurable. That is where the next look should go.

REV2 · FIX2 The $ER30 \geq 0.35$ arming floor, run
through the 54-filter columns it was never made
to face — it keeps 5 of 15 winners

PATCHES: plays.json rank 2 (absveto-short-arm-by-default-
er035-0808) and Part 2.6 §8 Saturday #2 — the ER-floor half of the card is
WITHDRAWN. **EXTENDS:** Part 1.5 rehab_abs_veto_short §5 (the 54-

filter table) with the wk29-31 → wk32 out-of-sample leg it never carried.

CONFIRMS: REV2 · Q0, independently and to the cent — this is a second derivation from the same ledger plus three legs Q0 did not run.

Unchanged and still shipping: arm abs_veto_short by default, and the BUILDING × US-SESSION entry veto.

THE ANSWER IN ONE LINE

Gaz — your bar was **13 of 15 winners kept to survive, 7-9 to be the same fake win the section rejects nine times. ER30 ≥ 0.35 keeps 5 of 15.** It does not clear your bar and it does not even reach the fake-win band — it is worse than the 0.30 rung the section already printed and rejected. It bins **\$1,064 of the \$1,664** top-15 block, throws away **+\$1,363 of profitable trades**, hands back **55% of the out-of-sample week** (+ \$809 → +\$363), and its live justification is **one signal on 07-29. Ship the veto. Do not ship the floor.**

1 · The row you asked for — same book, same columns, plus the OOS leg

Basis: abs_veto_55s SHORT, tick-repriced (shadow_real.real_pnl), **159 fires / 17 trading days** 2026-07-17 → 08-07, four ISO weeks. \$2.00/point, **\$1.50/round-turn**, 1 lot. Identical ledger and identical strip-3 / top-15 definitions as \$5 — the baseline row reproduces byte-for-byte before any column was added. **IS** = ISO weeks 29-31 (the weeks the floor was fitted on); **OOS** = week 32, untouched.

POLICY	KEEP N	KEEP \$	\$/TRADE	WIN %	CUT N	CUT \$	CUT \$/TR	TOP-15 KEPT	STRI P-3	STRI P-5	WORST LOD O	DAYS GREEN	IS WK2 9-31	OOS WK3 2
00 BASELINE — no arming filter	159	+2290.5	+14.4	48.4%	0	—	—	15/15	+1780.0	+1484.5	+1889.0	14/17	+1481.5	+809.0
E0.35 — ER30 ≥ 0.35 the play	36	+927.5	+25.8	55.6%	123	+1363.0	+11.1	5/15	+493.0	+327.5	+714.5	10/13	+564.5	+363.0
	51		+22.4	52.9%	108		+10.6	7/15	+632.5	—	—	—	—	+315.5

E0.30 — ER30 ≥ 0.30 the rung §5 already rejected		+1140. 5				+1150. 0									
G2+E0. 35 stacked, exactly as the card asks	36	+927.5	+25.8	55.6%	123	+1363. 0	+11.1	5/15	+493.0	+327.5	+714.5	10/13	+564.5	+363.0	
G2 — drop BUILDING × US- SESSION the §5 verdict	151	+2497. 5	+16.5	50.3%	8	-207.0	-25.9	15/15	+1987. 0	+1691. 5	+2096. 0	14/17	+1584. 0	+913.5	
A15 — ATR14 ≥ 15 the REAL- floor control	82	+1829. 0	+22.3	52.4%	77	+461.5	+6.0	15/15	+1318. 5	+1023. 0	+1414. 0	11/15	+1186. 0	+643.0	

Read it against your own decision rule and there is nothing left to argue.

Five of fifteen. The cut pile is positive at **+\$1,363 (+\$11.10/trade)** — a filter whose discards make money is not removing bad trades, it is removing trades. Out of sample it does not go negative, it simply hands back **55%** of the week it was meant to protect. And **row 4 is identical to row 2 on all fifteen columns:** BUILDING is defined as $ER30 \in [0.15, 0.30)$, which lies entirely below 0.35, so the floor eats every fire the veto would cut and 115 more besides. Shipping both is arithmetically the same as shipping only the floor — the card does not strengthen the one filter this report verdicted, it deletes it and takes \$1,570 of the reference book with it.

The ten winners it bins are not marginal. They carry **+\$1,064.00 of the + \$1,664.00 top-15 block**, and **seven of the ten are VIOLENT-WHIPS AW** — the exact cell the same card's previous bullet asks to stop benching (n=12, +\$443.50, +\$37.00/trade in session). $ER30 < 0.35$ is a strictly wider $ER30 < 0.25$ with the ATR condition removed: the old bench cut 19 fires and \$357.50, this cuts 123 fires and \$1,363.00. The gate's best non-top fire, 08-07 13:56 for **+\$156.00, fired at ER30 = 0.002**. An ER floor on a thrust gate removes the gate's own trigger condition.

2 · Three legs neither \$5 nor the card ran

(a) The cut pile is profitable in BOTH halves of the tape. \$5 records that the ER filter’s sign flips between halves at the 0.20 rung and calls it a coin. At 0.35 it is not even a coin — it is consistently wrong. Split at the median trading day:

HALF	BASELIN E N	BASELIN E \$	KEPT N	KEPT \$	CUT N	CUT \$	CUT \$/ TR	Δ NET VS BASELIN E	WHAT THE DISCAR DS SAY
H1 · 07-17 → 07-27	60	+1055.0	17	+457.0	43	+598.0	+13.9	−598.0	discards make money — filter is wrong here
H2 · 07-28 → 08-07	99	+1235.5	19	+470.5	80	+765.0	+9.6	−765.0	discards make money — filter is wrong here

(b) The live n=9 headline is one signal. The card’s number is real and I reproduce it exactly — live abs_veto_short fills, data_quality IS NULL, all 17 days: above the floor n=9, +\$550.50, +\$61.17/trade, and the three bands below are indeed negative (−\$251.00 / −\$419.50 / −\$80.50 on n=26). But those nine lots are only six signals, and one of them is the whole result:

TEST ON THE CARD’S OWN LIVE EVIDENCE	N (LOTS)	NET	\$/LOT	VERDICT
as quoted	9		+550.50	+61.17 reproduces to the cent
strip best 1 (id 360, 07-29 19:23, ATR 47.7)	8		+14.50	+1.81 the recommendation evaporates on one trade
remove that signal (ids 359+360, +\$768.00)	7	−217.50		−31.07 every other fire above the floor is negative
whole live book, for scale	35	−200.50		−5.73 31.4% win — a 3.3%-of-tape sample
winners kept — the column the source table never had	—		—	— top-5 live winners: 3/5 · top-10: 4/10

Five of the nine are STOPS, four of them on the two most recent days (08-06 and 08-07, both two-lot signals, both losers). The floor was fitted on the 3.3% of tape the router happened to arm; the card’s other half arms the gate by default, at which point the population it governs becomes the 159-fire continuous book — where it loses \$1,363. You cannot fit a filter on the 3.3% and deploy it on the 100%.

(c) A floor that never arms is a bench, not a filter. Under $ER30 \geq 0.35$ the reference book fires **0 / 13 / 13 / 10** times in ISO weeks 29 / 30 / 31 / 32 — **zero fires in an entire week** — and produces no fire at all on **4 of 17 days** (07-17, 07-19, 07-24, 08-03). On the live book it is silent on 2 of 7 days. That is why “days green 10/13” is not comparable to the baseline’s 14/17: three days stop existing.

3 · What to do

LEAD	DISPOSITION	THE NAMED TEST
abs_veto_short armed by default (delete the two bad bench rules)	LIVE — unchanged, still ships	96.7% of the week benched contained + \$1,064.50 of realised money on the live config; 53 of 54 filters lose to letting it fire.
Entry veto: drop BUILDING × US-SESSION	LIVE — unchanged, still ships	+\$207, 15/15 winners kept, better in every ISO week including the untouched one, +\$913.50 OOS vs +\$809.00. Thin at n=8 — re-look at n≥25.
ER30 ≥ 0.35 arming floor	REFUTED — WITHDRAW from the card	13-rung winner-retention sweep on the 159-fire tick-repriced book (5/15 kept, +\$1,363 of profitable trades discarded, no plateau — 0.325/0.35/0.375 = +\$849/+\$927/+\$679); cut pile positive in both halves; wk29-31 → wk32 + \$809 → +\$363; and strip-1 on its own live n=9 → +\$1.81/lot.
ATR14 ≥ 15 as the arming floor instead	PARKED — slot-contention lever, not a P&L lever	Keeps 15/15 winners at +\$22.30/trade, but still costs \$461 of net. Revive only if slot contention becomes the binding constraint.

The edit, exactly. plays.json rank 2 is still unamended as of this writing (the card, arms_when, number and kill all carry the floor). Replace with: topic → “abs_veto_short: arm it by default”; arms_when → **“Armed by default. No ER condition of any kind.** One per-entry veto: skip when regime=BUILDING and 13:30 ≤ UTC < 20:00”; number → drop the final sentence (“The ER floor is + \$61.20/trade above 0.35...”) and substitute **“+\$2,497.50 vs +\$2,290.50 baseline, 15/15 winners kept, +\$913.50 out of sample”**; kill → “Review after 15 fires armed-by-default; re-bench if the armed book is negative over those 15.” One further note for the file: the floor was justified as “mirroring grind’s”, and ER_FL00R in deciders.py is {} — every floor was deleted on 08-01 and the deletion is pinned by tests/test_deciders.py::test_er_floors_all_deleted. There is no floor to mirror.

Bottom line, Gaz. The reviewer’s instinct was exactly right and the arithmetic backs it without qualification: the one recommendation in the report that never faced the winners-kept column is the one that fails it hardest — **5 of 15, below even the 7/15 rung the same section rejects by name.** This is not a close call to be parked in SHADOW; the sample is

not thin (123 discarded fires, \$1,363 of them profitable), the failure is directional in both halves of the tape and out of sample, and its live support is a single 07-29 signal. **Withdraw the floor, ship the other two halves of Saturday #2 unchanged.** Reproduce: scratchpad/rehab_avs/q0_er035.py (the 15-column row and the rung sweep) and scratchpad/rehab_avs/fix2_er035.py (the both-halves, live-winner-retention and coverage legs).