

# GAZBOT V7 — the Friday report

Week ending 2026-08-14 · Mon 10 – Fri 14 August 2026 · MNQ-led, MGC now captured  
· PAPER · 38 plays · 9 sections

## ★ THE HONEST HEADLINE — the desk made money, and one desk made all of it

The whole desk booked **\$1,292.50** across 113 clean lots, and **three of the five days were green** — which is the thing Part 0 exists to show, because a year ago a green week was one day. But the money is not spread around, and the split is the first thing you should see: **the day rider made \$1,617.50 on 5 trades and the six-gate tournament lost –\$325.00 on 108**. Take the rider out and the week is red. Read every gate card in this report with that in mind — the tournament did not have a good week, it had a week that a second desk paid for.

There is no single honest total, so here are all of them, with what separates each from the next.

| WHAT YOU ARE LOOKING AT             | LOTS | NET \$     | WHY IT DIFFERS FROM THE ROW ABOVE                                                                                                      |
|-------------------------------------|------|------------|----------------------------------------------------------------------------------------------------------------------------------------|
| Raw trades rows, 08-10 → 08-14      | 117  | \$2,843.50 | includes 4 quarantined lots worth \$1,551.00                                                                                           |
| Clean ledger (data_quality IS NULL) | 113  | \$1,292.50 | the desk's real booked money — this is the number to quote                                                                             |
| The day rider alone                 | 5    | \$1,617.50 | the second desk. Five trades, five winners — and it only started writing a ledger at all this week, so treat it as thin, not as proven |
| The six-gate tournament alone       | 108  | –\$325.00  | what the gates this report is mostly about actually did                                                                                |

THE 4 QUARANTINED LOTS, NAMED

They are the 13 August day-rider phantom re-book, flagged EXCLUDE:day\_rider\_phantom\_rebook\_20260813, and they are worth \$1,551.00 — more than the entire week’s clean profit. **That is why the quarantine flag exists and why every query in this report carries data\_quality IS NULL.** A report that quoted the raw table would have led with \$2,843.50 and been wrong by \$1,551.00. Part 1 §6.2 walks the phantom itself.

| DAY        | DESK NET (CLEAN) |
|------------|------------------|
| 2026-08-10 | −\$205.50        |
| 2026-08-11 | \$560.50         |
| 2026-08-12 | −\$435.50        |
| 2026-08-13 | \$768.50         |
| 2026-08-14 | \$604.50         |

★

## THE SCORECARD — what survived the week

Lead with what survived the skeptic, never with the biggest number. Six results carried real weight into Saturday; here is where each one landed and the test that decided it.

| THE CLAIM                                                                                  | VERDICT, AND THE TEST THAT DECIDED IT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            | WHERE                       |
|--------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------|
| The wide profit leg beats the tight one                                                    | <div>SURVIVED — the week’s best-evidenced result</div> + \$331.50 over 50 identical entries — same signal, same price, same stop, the only difference being where profit was taken. Placebo p=0.0040, sign test p=0.0235, all five daily folds positive and all four gate folds positive. The A/B legs are a scale-out, not an experiment, so the instruction is widen Lot A’s target, never drop Lot A.                                                                                                                                                         | <a href="#">Part 1 §2.2</a> |
| Never bench a gate at face value — the standing rule, applied to the one properly red gate | <div>CONFIRMED, and it found the opposite of what was expected</div><br>Every treatment aimed at exhaustion_short’s ENTRY failed against a 4,000-draw placebo — 13 cooldown cells, 15 streak-bench cells and 9 flow floors, best of them at the 85th percentile. The entry was never the problem. Under a wide stop and a far target the same 87 signals are worth +\$3,459.98 against −\$84.50 as traded, and the entries beat 60 of 60 random-entry controls. Under the live scalp exit those same entries sit at the 47th percentile — indistinguishable from | <a href="#">Part 1.5</a>    |

| THE CLAIM                                           | VERDICT, AND THE TEST THAT DECIDED IT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                | WHERE                                                       |
|-----------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------|
|                                                     | random. The edge is invisible at the exit the desk uses.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |                                                             |
| "A wall of stops means bench the gate"              | <div>REFUTED, on two independent populations</div> <p>Mechanically, on 104 signals, every one of fifteen bench rules loses money and the best sits at the 54th percentile of noise — benching after two stops cuts 18 winners to avoid 12 losers. Live, the same rule costs −\$172 across the week and forfeits +\$356 on Friday alone. The discriminator is the REGIME the losses happened in, not the count of losses. "Arm for periods" wins.</p>                                                                                                                                 | <a href="#">Part 1 §2.4</a> , <a href="#">Part 1.5 §3.2</a> |
| The order book tells you which gold breaks will run | <div>REFUTED — and it is backwards</div> <p>On 397 MGC breaks, "the far side is empty" loses −\$2,059.50 at the 4.8th percentile of its own placebo: 1,904 of 2,000 random picks did better. The feature carries real information and the sign is inverted — a gold level that breaks into a vacuum is one nobody is defending, and price comes straight back. The reformulation (fade it) is PARKED, not claimed.</p>                                                                                                                                                               | <a href="#">Movement 3</a>                                  |
| A better exit would have caught the runs we missed  | <div>REFUTED — we are late, not sloppy</div> <p>The timing-decay ladder, n=57 at every rung on this week's tape: perfect entries at the ignition minute pay +\$3,799, at five minutes +\$389.50, and at seven minutes −\$1,073. Our gates fire minutes after the money has gone. No exit fixes being late. The idle-gate lab is an honest null: six mechanical fires for +\$15.50, and −\$94.50 once the single best is stripped.</p>                                                                                                                                                | <a href="#">Movement 2</a>                                  |
| The router is earning its keep                      | <div>YES — and the gain is in its CLOCK, not its levels</div> <p>1,148 ticks, zero service failures, 95% of ticks change nothing. Its thresholds are already at the best cell on both universes independently — there is nothing to win there. The timing is a different story: window 45 / step 5 / hold 3 is worth +\$502 and +\$399 across the two sets. Two honest holes: its leakage detector reported zero leakage while −\$241.00 leaked through a gate outside its universe, and a 10h20m outage on 13 August was indistinguishable from a calm day in every scoreboard.</p> | <a href="#">Part 2.6</a>                                    |

## § Every gate, this week, on the clean ledger

| GATE              | LOTS | NET \$     | WIN %  | \$ PER LOT |
|-------------------|------|------------|--------|------------|
| exhaustion_short  | 71   | −\$218.50  | 42.3%  | −\$3.08    |
| abs_veto_short    | 26   | −\$64.50   | 46.2%  | −\$2.48    |
| capitulation_long | 2    | −\$48.00   | 0.0%   | −\$24.00   |
| abs_veto_long     | 9    | \$6.00     | 44.4%  | \$0.67     |
| day_rider         | 5    | \$1,617.50 | 100.0% | \$323.50   |

Computed at build time from data/gazbot7.db with data\_quality IS NULL, not transcribed. Gates are collapsed across their A and B lots; the per-lot split is in Part 1 §3, and it is where the wide-leg finding lives.

### THE THREE THINGS THIS WEEK ACTUALLY TAUGHT THE DESK

#### 1 • When a gate is red, check the exit before you touch the entry.

Twenty separate entry treatments on exhaustion\_short failed to beat a coin. The exit turned the same signals from  $-\$84.50$  into  $+\$3,459.98$ . The desk has spent months tuning entries; this is the clearest evidence yet that the other half of the trade is where the slack is. And the mirror of it is in Movement 3: when the ENTRY is a coin flip, no exit saves it — median favourable excursion  $+3.28 \times \text{ATR}$  against median adverse  $-3.04 \times \text{ATR}$  is a symmetric distribution with nothing to build a stop around.

2 • A filter that beats a placebo can still be a calendar. The one treatment that cleared its placebo at the 97.6th percentile turned out to be re-discovering a floor that is already deployed, and the 17 signals it cut were 12-from-one-era. The placebo is necessary and it is not sufficient. Ask what the filter is selecting for, every time.

#### 3 • Three instruments reported confidently and wrongly this week.

The router's leakage detector said zero on a week money leaked; the nightly rollup said  $\$0/\$0/\$0$  for a day whose journal had rotated away; selector\_nightly counted a quarantined trade five times and inflated one day's regret by 39%. None of them errored. **A number that arrives without a failure mode is not a measurement.**

---

## ★ WHAT TO ACTUALLY DO

---

Garrath — this is the whole answer to “so what do I actually DO?”, and it is at the front. Every row below mirrors reports/friday\_v7/plays.json, which is the reconciled source of truth for this week; the same file drives the [Monday playbook](#), so the two cannot drift apart. 38 plays, five

windows, in the order they need doing. Every play carries its **evidence tier** and its **kill criterion** — if a play has no way to be proven wrong it should not be on the card.

### THE THREE THAT MATTER MOST

**1 • Router timing: window 45 / step 5 / hold 3 (SATURDAY #1).** Change the direction router's regime window from 30 to 45 minutes, its step to 5 minutes and its confirmation hold to 3 ticks. One constant block in the router service; restart required. +\$502 / +\$399 across the two universes over 40 days; 4-and-6 days helped, 0-and-1 hurt.

**2 • Bench exhaustion\_short against a proven counter-trend (SATURDAY #2).** Add exhaustion\_short to the router's managed set so it is benched when the router reads a trend running against it — the same treatment capitulation\_long and rgv\_short already get. n=26 / −\$398.50 / −\$15.33 per trade over 30 days; red on 7 of 8 days; survives strip-worst and leave-one-out. Separately, −\$241.00 leaked through this gate unseen while the detector reported zero.

**3 • Widen Lot A's profit target — do not drop Lot A (MONDAY #1).** Move the tight leg's target out toward the wide leg's. The A/B legs are a scale-out on ONE signal with one stop, so 'B beat A' is an instruction to widen A, never to delete it. +\$331.50 over 50 identical entries. Placebo p=0.0040, sign test p=0.0235, 5/5 daily folds and 4/4 gate folds positive — the cleanest statistical result in the report.

Generated from the top of the card below, so the two cannot drift apart. Everything else, including every refuted idea and why it died, is in the five windows underneath.

**SATURDAY • TONIGHT — THIS WINDOW SHUTS AT THE SUNDAY 22:00 UTC  
REOPEN • 4 PLAYS**

Code and config changes that need a restart, **in execution order, not in order of headline size. #1 and #2 are forced:** gazbot7-gate-reactivate.timer arms every =off gate at 22:00 UTC Sunday, so for both of them “decide later” is not one of the available states. #3–#6 are safety and measurement, and they come before the two experiments at the bottom because those experiments cannot be judged without them.

| # | WHAT TO DO                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        | THE NUMBER BEHIND IT                                                                                                                                                                                           | ARMS WHEN · EXIT · KILL CRITERION                                                                                                                                                                                                                   | OWNER & REVERT                                                                                                                     |
|---|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------|
| 1 | <b>Router timing: window 45 / step 5 / hold 3</b><br><b>LIVE</b> Change the direction router's regime window from 30 to 45 minutes, its step to 5 minutes and its confirmation hold to 3 ticks. One constant block in the router service; restart required.<br><b>Mechanism:</b> The 40-day threshold sweep found the gain is in the TIMING, not the thresholds. The ER and net levels are already at their best cell on both universes independently, so there is nothing to win there — but the window/step/hold cell is worth +\$502 on the live managed set and +\$399 on the full historical fader set, helping on 4 and 6 days respectively and hurting on 0 and 1.<br>Evidence: Part 2.6 §4 — the 40-day sweep, both sets; corroborated independently by the actuation-lag measurement in §7. · verification:<br><b>Re-run the sweep on next week's tape and confirm the 45/5/3 cell is still the best or on a plateau with it. If it has moved, the surface is noise and this reverts.</b> · mode: deploy | <b>+\$502 / +\$399 across the two universes over 40 days; 4-and-6 days helped, 0-and-1 hurt.</b>                                                                                                               | <b>Arms:</b> Always — it is the router's own clock, not a gate.<br><b>Exit:</b> n/a — this is a routing constant, not a trade.<br><b>Kill:</b> Two consecutive weeks where 45/5/3 scores below the current 30/-/- on the nightly rollout.           | Garrath — router service constants<br><br>Restore the previous window/step/hold and restart; sub-minute revert, no data migration. |
| 2 | <b>Bench exhaustion_short against a proven counter-trend</b><br><b>LIVE</b> Add exhaustion_short to the router's managed set so it is benched when the router reads a trend running against it — the same treatment capitulation_long and rgv_short already get.<br><b>Mechanism:</b> exhaustion_short is currently OUTSIDE the router's managed universe, which is why the leakage detector reported zero leakage while this gate quietly bled. Fading a proven trend is exactly the failure mode the router was built for.<br>Evidence: Part 2.6 §4 and §3 (the leakage the detector cannot see); the gate's own week in Part 1 §3. · verification: <b>n=26 over 30 days is not thin, but review the first 15 benched episodes against what the gate would have done unbenched before treating it as settled.</b> · mode: deploy                                                                                                                                                                                | <b>n=26 / −\$398.50 / −\$15.33 per trade over 30 days; red on 7 of 8 days; survives strip-worst and leave-one-out. Separately, −\$241.00 leaked through this gate unseen while the detector reported zero.</b> | <b>Arms:</b> Router reads a trend opposed to the gate's direction.<br><b>Exit:</b> Unchanged — this changes whether it may enter, not how it leaves.<br><b>Kill:</b> Two weeks where the benched windows would have been net positive for the gate. | Garrath — gate_switches.env + router managed set<br><br>Remove from the managed set; the gate returns to always-armed immediately. |
| 3 | <b>selector_nightly.py must respect data_quality</b><br><b>LIVE</b> Add the `data_quality IS NULL` predicate to scripts/selector_nightly.py so quarantined trades stop being scored. One WHERE clause.<br><b>Mechanism:</b> Quarantined rows are counted as real, and because the selector scores each trade against the best of three exits, one bad row is counted repeatedly.<br>Evidence: Part 2.5 Part C — a broken instrument underneath this report. · verification: <b>Re-run the 13 August rollout after the fix and confirm regret drops from \$748 to \$458.</b> · mode: deploy                                                                                                                                                                                                                                                                                                                                                                                                                        | <b>13 August regret inflated 39% (\$748 → \$458) by a single quarantined trade counted five times. It feeds this report's own selector numbers.</b>                                                            | <b>Arms:</b> n/a — nightly job.<br><b>Exit:</b> n/a<br><b>Kill:</b> n/a — this is a correctness fix, not an experiment. It has no failure mode worth reverting.                                                                                     | Garrath — scripts/selector_nightly.py<br><br>Revert the one-line predicate; the job re-runs nightly.                               |
| 4 | <b>grind_long: veto a long that is fighting a 30-minute fall</b><br><b>SHADOW</b> Add a directional veto inside grind_long's gate function: refuse the entry when the last 30 minutes have fallen more than 1×ATR. In-code, per-tick — no router, no switch file.<br><b>Mechanism:</b> The recurring finding is that the faders' problem is direction-blindness, and the router can only fix it on a 5-minute lag. A veto inside the gate has no lag at all.<br>Evidence: Part 2.5 Part B — the mechanism that does work, and it is a veto. · verification: <b>It is a veto, so the fake-filter guard is the test that matters and it passes: it keeps 10 of the 10 best trades. Confirm that again on the first 30 live fires.</b> · mode: shadow-first                                                                                                                                                                                                                                                          | <b>+\$2,407 → +\$3,134 over 133 tick-honest fires, keeping 10/10 of the best trades, with a plateau across the whole tested range rather than one hot cell.</b>                                                | <b>Arms:</b> Always armed; the veto only ever refuses.<br><b>Exit:</b> Unchanged.<br><b>Kill:</b> It drops any of the gate's three best trades in its first 30 live fires, or the plateau narrows to a single cell on next week's tape.             | Garrath — deciders.py grind_long<br><br>Delete the veto clause; the gate returns to its current shape.                             |

**MONDAY · AT THE DESK — SWITCH-FILE AND EXIT-CELL CHANGES, ALL REVERSIBLE · 3 PLAYS**

No code deploy in this window at all. Every row is `gate_switches.env`, `exit_overrides.json` or a router constant, and comes back in five minutes. **#1 needs nothing but a decision** — the router already owns the lever.

| # | WHAT TO DO                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       | THE NUMBER BEHIND IT                                                                                                                                                           | ARMS WHEN · EXIT · KILL CRITERION                                                                                                                                    | OWNER & REVERT                                                                              |
|---|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------|
| 1 | <b>Widen Lot A's profit target — do not drop Lot A</b><br><b>LIVE</b> Move the tight leg's target out toward the wide leg's. The A/B legs are a scale-out on ONE signal with one stop, so 'B beat A' is an instruction to widen A, never to delete it.<br><b>Mechanism:</b> Both lots enter on the same signal at the same price with the same stop; the only difference is where they take profit. Over 50 identical entries the wide leg beat the tight one, which means the tight target is leaving money behind on trades we are already in. | <b>+\$331.50 over 50 identical entries. Placebo p=0.0040, sign test p=0.0235, 5/5 daily folds and 4/4 gate folds positive — the cleanest statistical result in the report.</b> | <b>Arms:</b> Unchanged.<br><b>Exit:</b> Lot A target widened; Lot B unchanged.<br><b>Kill:</b> Two weeks where the tight leg beats the wide leg on the same entries. | Garrath — exit_overrides.json<br><br>Restore the previous Lot A target; five-minute revert. |

| # | WHAT TO DO                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  | THE NUMBER BEHIND IT                                                                                                                                                                                    | ARMS WHEN · EXIT · KILL CRITERION                                                                                                                                        | OWNER & REVERT                                                       |
|---|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------|
|   | <p>Evidence: Part 1 §2.2 — the week's best-evidenced result, and §3 'By lot — the tight-leg problem in one table'. · verification: <b>Placebo p=0.0040, sign test p=0.0235, all five daily folds positive and all four gate folds positive. Re-check after 30 fresh pairs.</b> · mode: config</p>                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |                                                                                                                                                                                                         |                                                                                                                                                                          |                                                                      |
| 2 | <p><b>Open Rider: stop at 3.0×ATR, not 2.0×</b></p> <p><b>SHADOW</b> Widen the Open Rider's stop from twice ATR to three times.</p> <p><b>Mechanism:</b> The house stop is too tight for the open. A 2×ATR stop is shaken out of rides that subsequently work, and the day rider's edge is in holding.</p> <p>Evidence: Part 2 §5 — the secondary axis, stop width. · verification: <b>+\$8.19 vs −\$15.08 per trade, wins on 4 of 5 days, and the replay agrees across 41 days. But the DAY effect is 2.2× the stop effect, so this is the smaller of the two levers.</b> · mode: config</p>                                                                                                                                                                               | <p><b>+\$8.19/trade against −\$15.08/trade at the current width; 4 of 5 days better; replay agrees over 41 sessions. Flagged SHADOW because the day effect is 2.2× larger than the stop effect.</b></p> | <p><b>Arms:</b> Unchanged.<br/> <b>Exit:</b> Stop 3.0×ATR.<br/> <b>Kill:</b> It loses on 3 of the next 5 days, or the day effect swamps it again on a bigger sample.</p> | <p>Garrath — day rider config</p> <p>Restore 2.0×ATR; immediate.</p> |
| 3 | <p><b>Leave both abs_veto sides armed — tune the side, never relegate a direction</b></p> <p><b>LIVE</b> No change. This is on the card because relegating the short side was proposed and refuted, and the 22:00Z timer will re-arm whatever you switch off anyway.</p> <p><b>Mechanism:</b> A gate runs two-sided with independently-tuned per-side thresholds. A direction that loses in one regime is a tuning problem, not a defective direction.</p> <p>Evidence: Part 2 §4 — relegation, and its disposition row 'Relegate the abs_veto SHORT direction: REFUTED'. · verification: <b>The shadow book says short is the STRONGER side. If that inverts for three consecutive weeks, revisit the per-side thresholds — still not the direction.</b> · mode: no-op</p> | <p><b>Shadow: short +\$2,651 vs long +\$1,371, and +\$15.50 vs +\$8.02 per trade. The side proposed for relegation is the better one.</b></p>                                                           | <p><b>Arms:</b> Both sides, as now.<br/> <b>Exit:</b> Unchanged.<br/> <b>Kill:</b> n/a — this play is an instruction not to act.</p>                                     | <p>Garrath — gate_switches.en v</p> <p>n/a — nothing changes.</p>    |

## BUILD · NEXT — NOTHING GOES LIVE THIS WEEK · 10 PLAYS

Instrumentation, shadow slates and the method fixes that decide whether next Friday's numbers can be trusted at all. **#1 gates #4**, and gates every future study that concludes “exit earlier” — the shadow repricer leaks past its own stops on 44% of trades, which is exactly the bias that makes cutting early look good.

| # | WHAT TO DO                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    | THE NUMBER BEHIND IT                                                                                                                             | ARMS WHEN · EXIT · KILL CRITERION                                                                                                                                                                                                                                                       | OWNER & REVERT                                                                      |
|---|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------|
| 1 | <p><b>Shadow-slate the exhaustion_short wide-stop exit</b></p> <p><b>SHADOW</b> Run the gate's live entries through a 1.5×ATR stop with a far target and a time cap, in the shadow book, alongside the live scalp exit. No live change.</p> <p><b>Mechanism:</b> Every treatment aimed at this gate's ENTRY failed against a placebo — cooldowns, streak-benches and a tighter flow floor all lose to random. The entry is not the problem. Under a wide-stop exit the same entries beat 60 of 60 random-entry controls; under the live scalp exit they land at the 47th percentile, i.e. indistinguishable from firing at random. The edge is invisible at the exit the desk uses.</p> <p>Evidence: Part 1.5 §7 — the exit sweep, its random-entry control and the day-by-day. · verification: <b>Worst leave-one-day-out + \$1,626.49, strip-best-3 +\$2,046.14, 9 of 12 days green, and +\$397.26 on this week's four testable days. Shadow it because the sweep never turns over (the best cell is at the grid boundary) and the repricer takes no slippage.</b> · mode: shadow-first</p> | <p><b>+\$3,459.98 against −\$84.50 live on the same 87 signals. Entry alpha \$3,467.78 at the 100th percentile of 60 random-entry draws.</b></p> | <p><b>Arms:</b> Same entries as the live gate — this changes only the exit.<br/> <b>Exit:</b> Stop 1.5×ATR, far target, 15-minute cap.<br/> <b>Kill:</b> The shadow arm fails to beat the live arm over 40 shadow trades, or its fills show the slippage the repricer assumed away.</p> | <p>Garrath — shadow slate</p> <p>Remove the slate; it never touched live money.</p> |
| 2 | <p><b>Run the OPEN/NEWS cluster greenfield hunt first next week</b></p> <p><b>PARKED</b> Schedule the OPEN/NEWS greenfield hunt at the FRONT of the queue, ahead of any further gold work.</p> <p><b>Mechanism:</b> It is the largest addressable block of missed runs and no gate on this desk has ever been built for it. It was scheduled this week, sat fifth in the queue, and the queue ran out.</p> <p>Evidence: Movement 3 §8 and its disposition — 'MNQ OPEN/NEWS-cluster greenfield hunt: NOT RUN'. · verification: <b>The hunt itself is the test. A null is an acceptable outcome; not running it is not.</b> · mode: build</p>                                                                                                                                                                                                                                                                                                                                                                                                                                                   | <p><b>22 of this week's 68 runs, roughly a third of an \$8,024 single-lot ceiling, never hunted.</b></p>                                         | <p><b>Arms:</b> n/a<br/> <b>Exit:</b> n/a<br/> <b>Kill:</b> n/a — this is a scheduling decision.</p>                                                                                                                                                                                    | <p>Garrath — friday_phases.py phase order</p> <p>n/a — reordering a queue.</p>      |

**Arms:** n/a — nightly instrument.



| # | WHAT TO DO                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 | THE NUMBER BEHIND IT                                                                                                                                                                                                                               | ARMS WHEN · EXIT · KILL CRITERION                                                                                                                                  | OWNER & REVERT                                                                     |
|---|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------|
| 3 | <b>Widen the router's leakage detector to every fader it could manage</b><br><b>PARKED</b> Extend the leakage detector's universe beyond capitulation_long and rgv_short so it can see losses on gates the router is not currently managing.<br><br><b>Mechanism:</b> A detector scoped to the gates that barely trade will report zero leakage on a week when money is leaking through a gate it was never pointed at. Zero leakage currently means an empty universe, not a clean week.<br><br>Evidence: Part 2.6 §3 — the leakage the detector cannot see. · verification: <b>After the fix, re-run this week: it should report the −\$241.00 it missed.</b> · mode: build                                              | <b>Reported zero leakage on all four visible days while −\$241.00 leaked through a gate outside its universe.</b>                                                                                                                                  | <b>Exit:</b> n/a<br><b>Kill:</b> n/a — instrumentation.                                                                                                            | Garrath — scripts/router_nightly.py<br>Restore the narrow universe.                |
| 4 | <b>Specify a dwell time on the router's direction threshold</b><br><b>SHADOW</b> Design and shadow a minimum-dwell requirement before the router is allowed to flip direction, instead of flipping the moment the ±40pt threshold is crossed.<br><br><b>Mechanism:</b> Threshold hysteresis with no dwell lets the direction round-trip inside minutes when the measure is sitting on the line — the flip is decided by noise, not by a regime change.<br><br>Evidence: Part 1 §8.2 (churn) and Part 2.6's disposition row 'A dwell-time requirement: PARKED'. · verification: <b>Score it on the week's actual flips: how many of the 36 would a dwell have suppressed, and what did those cost?</b> · mode: shadow-first | <b>36 flips this week, 78% of them produced no trade, 17 were under 30 minutes, and one round trip was decided by three points. △ It cost \$0 directly this week — this is a build because the mechanism is wrong, not because it is bleeding.</b> | <b>Arms:</b> n/a — routing logic.<br><b>Exit:</b> n/a<br><b>Kill:</b> The churn measurement shows it cost real money and a dwell would not have caught it.         | Garrath — router direction rule<br>Set dwell to zero.                              |
| 5 | <b>Put the two surviving gold gates in the shadow book</b><br><b>SHADOW</b> Shadow-slate the MGC day rider and the MGC coil bouncer. Neither goes live.<br><br><b>Mechanism:</b> Gold is a new instrument for this desk and both survived their own robustness checks, but neither has the sessions behind it to be judged.<br><br>Evidence: Part 2.5 Part A, findings 2 and 3, with their disposition rows. · verification: <b>Day rider: 15 of 15 grid cells positive (+\$552 to + \$1,742) and it beats the best constant on both dollars and hit rate. Coil bouncer: n=154, median trade +\$7.75, survives strip-best-day, beaten 0 of 12 by a same-count placebo.</b> · mode: shadow-first                            | <b>MGC is \$10.00 a point — every gold dollar here is already priced at the gold multiplier, not the Nasdaq one.</b>                                                                                                                               | <b>Arms:</b> Per each gate's own rule.<br><b>Exit:</b> Per each gate's own rule.<br><b>Kill:</b> Either fails to stay positive across 4 untouched weeks of shadow. | Garrath — shadow slate<br>Remove the slates.                                       |
| 6 | <b>Make signal_journal record WHY a signal was blocked</b><br><b>PARKED</b> Fix the write path so `suppressed_by` is actually populated.<br><br><b>Mechanism:</b> It is NULL in every row, so a blocked signal is indistinguishable from a taken one in the table. Every claim about what a gate WOULD have done unrouted is currently a reconstruction from the tape rather than a measurement.<br><br>Evidence: Part 1 §11 disposition — 'Signal journal block reasons: MEASUREMENT GAP'. · verification: <b>After the fix, one day's journal should account for every bench with a named reason.</b> · mode: build                                                                                                      | <b>NULL in all 222 rows this week. It is the same defect class the scope flagged last week, which means it was reported and not fixed.</b>                                                                                                         | <b>Arms:</b> n/a<br><b>Exit:</b> n/a<br><b>Kill:</b> n/a — instrumentation.                                                                                        | Garrath — signal journal write path<br>n/a — restoring a NULL column helps nobody. |
| 7 | <b>Pre-register ONE Open Rider day feature and leave it alone for 60 sessions</b><br><b>PARKED</b> Fix ATR-14 ≥ 9.5pt measured at 13:00 as the single day filter, in advance, and do not touch it for 60 forward sessions.<br><br><b>Mechanism:</b> The eight-feature search was beaten by 40% of pure-noise searches once the placebo was made search-aware. The only way to learn anything is to stop searching and commit to one rule.<br><br>Evidence: Part 2 §5 — 'ODR causal day-classifier: REFUTED as searched' and 'pre-registered single feature: PARKED'. · verification: <b>60 untouched forward sessions. That is the test; there is no shortcut.</b> · mode: build                                           | <b>Best of 272 searched rules was beaten by 40.0% of noise searches. The pre-registration is the only remaining route.</b>                                                                                                                         | <b>Arms:</b> 13:00Z daily read.<br><b>Exit:</b> n/a<br><b>Kill:</b> Fails to separate over 60 sessions — at which point the day-classifier idea is done.           | Garrath — day rider config<br>Remove the filter.                                   |
| 8 | <b>Fix the nightly rollup's journal dependency</b><br><b>PARKED</b> Stop the nightly rollup silently reporting \$0/\$0/\$0 when the systemd journal has rotated past the day it is backfilling.<br><br><b>Mechanism:</b> A rollup that cannot distinguish 'no value' from 'no data' will report a clean day on a day it could not see.<br><br>Evidence: Part 2.6 §3 disposition — 'Monday 08-10 nightly rollup: MEASUREMENT GAP'. · verification: <b>Re-run 08-10 after the fix and confirm it either produces numbers or refuses.</b> · mode: build                                                                                                                                                                       | <b>The journal does not reach past 2026-08-10 23:11Z, so the backfilled file reads \$0/\$0/\$0 because there was nothing to read — not because nothing happened.</b>                                                                               | <b>Arms:</b> n/a<br><b>Exit:</b> n/a<br><b>Kill:</b> n/a — instrumentation.                                                                                        | Garrath — scripts/router_nightly.py<br>n/a                                         |
| 9 | <b>Find the other writer to gate_switches.env</b>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          | <b>\$0 cost this week, which is luck rather than safety. The gate traded twice for −\$48.00.</b>                                                                                                                                                   | <b>Arms:</b> n/a<br><b>Exit:</b> n/a                                                                                                                               | Garrath — gate_switches.env writers                                                |



| # | WHAT TO DO                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                | THE NUMBER BEHIND IT | ARMS WHEN · EXIT · KILL CRITERION       | OWNER & REVERT |
|---|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------|-----------------------------------------|----------------|
|   | <b>PARKED</b> Identify what is writing capitulation_long back to ON.<br>The router wrote it OFF five times and ON never, and the file reads ON.<br><br><b>Mechanism:</b> Something outside the router is editing the switch file. Until it is named, no bench decision on this desk can be trusted to stick.<br><br>Evidence: Part 1 §11 disposition — 'capitulation_long arm/bench tug-of-war: PARKED — unexplained'. · verification: <b>Watch the file's mtime against the router's own decision log for one session.</b> · mode: build |                      | <b>Kill:</b> n/a — this is a diagnosis. | n/a            |

|    |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |                                                                                |                                                                                                                                    |                                                                          |
|----|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------|
| 10 | <b>Shadow a heavier weight on the STOP-RATE meter</b><br><br><b>SHADOW</b> Re-weight the untradeable-day meter to lean more on stop rate, in shadow, before touching the live cutoff.<br><br><b>Mechanism:</b> Of the three meters feeding the untradeable score, stop rate is the best discriminator and carries the least weight.<br><br>Evidence: Part 2.6 §5 — 'which of the three meters actually did the work'. · verification: <b>\$559/day separation is the number to beat. It is also the meter most obviously downstream of the desk's own behaviour, so shadow before deploy.</b> · mode: shadow-first | <b>\$559/day separation, the best of the three, and the lightest-weighted.</b> | <b>Arms:</b> n/a<br><br><b>Exit:</b> n/a<br><br><b>Kill:</b> The re-weighted meter fails to beat the current one over 20 sessions. | Garraath — untradeable meter weights<br><br>Restore the current weights. |
|----|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------|

HOLD · LEAVE ALONE — THESE ARE ALREADY RIGHT · 8 PLAYS

Standing lessons and things already in production. The action is to not touch them.

| # | WHAT TO DO                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               | THE NUMBER BEHIND IT                                                                                                                            | ARMS WHEN · EXIT · KILL CRITERION                                                                                                                         | OWNER & REVERT      |
|---|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------|
| 1 | <b>The stop-fill fix — leave it entirely alone</b><br><br><b>LIVE</b> Nothing. It works.<br><br><b>Mechanism:</b> Resting stops are now placed on a concrete future rather than the continuous contract, so IBKR actually fires them.<br><br>Evidence: Part 1 §2.1 — 'the stop bug is dead, and I can prove it'. · verification: <b>62 of 62 stops filled this week; 393 trades and 17 days since the last STOP_UNFILLED.</b> · mode: no-op                                                                                                                                                                                                                              | <b>62/62 stops filled. This also means no number in this report needs a malfunction adjustment — every dollar lost is a strategy dollar.</b>    | <b>Arms:</b> n/a<br><br><b>Exit:</b> n/a<br><br><b>Kill:</b> A single STOP_UNFILLED reopens it immediately.                                               | Garraath<br><br>n/a |
| 2 | <b>Arm for PERIODS, not runs — and never bench on a count of losses</b><br><br><b>LIVE</b> Keep arming as a permission window. Do not add a wall-of-stops bench rule.<br><br><b>Mechanism:</b> A losing streak tells you the last two trades went badly; it does not predict the next signal. The discriminator is the REGIME the losses happened in.<br><br>Evidence: Part 1 §2.4 (live money) and Part 1.5 §3.2 (the mechanical test on 104 signals). · verification: <b>Two independent populations agree. Fifteen versions of the bench rule were tested on exhaustion_short and every one lost, the best sitting at the 54th percentile of noise.</b> · mode: no-op | <b>Bench-on-2-stops costs –\$172 across this week and forfeits + \$356 on Friday alone. Mechanically it cuts 18 winners to avoid 12 losers.</b> | <b>Arms:</b> n/a<br><br><b>Exit:</b> n/a<br><br><b>Kill:</b> A regime-conditioned bench rule — not a loss-counting one — beats always-armed over a month. | Garraath<br><br>n/a |
| 3 | <b>grind_long's ATR floor at 22 — already deployed, and now confirmed</b><br><br><b>LIVE</b> Nothing. This is the confirmation of the 08-08 revert, not a new proposal.<br><br><b>Mechanism:</b> The floor stands the gate when there is not enough range to pay for the trade.<br><br>Evidence: Movement 2 — 'grind_long ATR floor 10 → 22 (the 08-08 revert): LIVE'. · verification: <b>On run tape the floor blocks four trades that all lose. It is doing its job.</b> · mode: no-op                                                                                                                                                                                 | <b>Already live. Do not re-propose it.</b>                                                                                                      | <b>Arms:</b> n/a<br><br><b>Exit:</b> n/a<br><br><b>Kill:</b> n/a                                                                                          | Garraath<br><br>n/a |
| 4 | <b>abs_veto_55s stays two-sided — already promoted</b><br><br><b>LIVE</b> Nothing. Promoted two-sided on 25 July; the battery re-confirms it.<br><br><b>Mechanism:</b> thrust_loose plus a 55-second continuation confirm, on both sides.<br><br>Evidence: Part 2 §2 — the full promotion battery. · verification: <b>16 of 21 days green, positive in both regimes, both walk-forward halves and both sides.</b> · mode: no-op                                                                                                                                                                                                                                          | <b>Already live. On the card only so nobody re-promotes it as new.</b>                                                                          | <b>Arms:</b> n/a<br><br><b>Exit:</b> n/a<br><br><b>Kill:</b> Three consecutive red weeks on either side.                                                  | Garraath<br><br>n/a |
| 5 | <b>Router ER ≥ 0.15 /  net  ≥ 30 — nothing to gain, leave them</b><br><br><b>LIVE</b> Nothing. The timing change (SATURDAY) is where the gain is; these levels are already right.<br><br><b>Mechanism:</b> The thresholds decide when a regime is declared; the timing decides how fast.                                                                                                                                                                                                                                                                                                                                                                                 | <b>A \$13 plateau. There is nothing here.</b>                                                                                                   | <b>Arms:</b> n/a<br><br><b>Exit:</b> n/a<br><br><b>Kill:</b> n/a                                                                                          | Garraath<br><br>n/a |

| # | WHAT TO DO                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               | THE NUMBER BEHIND IT                                                                                                                              | ARMS WHEN · EXIT · KILL CRITERION                                               | OWNER & REVERT |
|---|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------|----------------|
|   | Evidence: Part 2.6 §4 — 'ER≥0.15 /  net ≥30 thresholds: LIVE — HOLD'. · verification: <b>Best cell on both universes independently, and net 20-40 is a \$13 plateau.</b> · mode: no-op                                                                                                                                                                                                                                                                                                                                                                                                                                   |                                                                                                                                                   |                                                                                 |                |
| 6 | <b>The untradeable meter's 65 cutoff — mid-plateau, leave it</b><br><div>LIVE</div> Nothing.<br><b>Mechanism:</b> The meter scores how untradeable a day looks and 65 is the line.<br>Evidence: Part 2.6 §5 — 'The untradeable meter, and its 65 cutoff: LIVE — HOLD'. · verification: <b>Monotone across 22 days (–\$651 / –\$212 / +\$7 per day) and 50-75 all score identically.</b> · mode: no-op                                                                                                                                                                                                                    | <b>Cutoff sits mid-plateau. Moving it is motion, not improvement.</b>                                                                             | Arms: n/a<br>Exit: n/a<br>Kill: The monotonicity breaks on a fresh 20 sessions. | Garrath<br>n/a |
| 7 | <b>The two grind-exit shadow ledgers — keep watching, do not conclude</b><br><div>SHADOW</div> Nothing. Both remain UNMEASURED and that is the honest state.<br><b>Mechanism:</b> The two-ratchet runner-clip watch and the 2R-partial smoothness ledger both need grind to trade before they can say anything.<br>Evidence: Part 2 §6 and §7. · verification: <b>No grind trades since 5 August; zero clips on three pre-August runners; green-day % is zero for both arms over the five days grind traded, so 'bank early' wins trivially on an all-red sample. None of that is evidence either way.</b> · mode: no-op | <b>△ The two-ratchet tool has a –\$406 cross-check error of its own. Fix before trusting it.</b>                                                  | Arms: n/a<br>Exit: n/a<br>Kill: n/a — they are already suspended pending data.  | Garrath<br>n/a |
| 8 | <b>The router blackout alarm — shipped this week, leave it</b><br><div>LIVE</div> Nothing.<br><b>Mechanism:</b> router_health_check.py alarms on a three-tick ABORT streak.<br>Evidence: Part 2.6 §5 disposition — 'The 08-13 10h20m OAuth outage: FIXED — NOT AN ACTION'. · verification: <b>Live since 15:31Z on 13 August. The next outage is the test.</b> · mode: no-op                                                                                                                                                                                                                                             | <b>The outage it was built for ran 10h20m across 125 consecutive aborted ticks and was indistinguishable from a calm day in every scoreboard.</b> | Arms: n/a<br>Exit: n/a<br>Kill: An outage passes undetected.                    | Garrath<br>n/a |

**NOT-AN-ACTION · WITHDRAWN / REFUTED — KEPT ON THE CARD SO NOBODY RE-PROPOSES THEM · 13 PLAYS**

Every one of these was a recommendation at some point this week. They are printed with their autopsies so next Friday's agent does not dig the same hole.

| # | WHAT TO DO                                                                                                                                                                                                                                                                                                                                                                                                                            | THE NUMBER BEHIND IT                                                                                             | ARMS WHEN · EXIT · KILL CRITERION                                                                                                                 | OWNER & REVERT |
|---|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------|----------------|
| 1 | <b>An entry efficiency floor at 0.10</b><br><div>PARKED</div> Do not ship it.<br><b>Mechanism:</b> Refuse entries when tape efficiency is below 0.10.<br>Evidence: Part 1 §9.1 — 'PARKED, killed by leave-one-day-out'. · verification: <b>Killed twice: Tuesday carries 89% of the gain, and the winners test keeps only 7 of 25.</b> · mode: no-op                                                                                  | <b>Keeps 7 of 25 winners. That is the fake-filter signature.</b>                                                 | Arms: n/a<br>Exit: n/a<br>Kill: Revive if: three or more weeks of data and it still separates, AND it stops discarding two-thirds of the winners. | Garrath<br>n/a |
| 2 | <b>A break-even stop at +8 to +15pt</b><br><div>PARKED</div> Do not ship it.<br><b>Mechanism:</b> Move the stop to break-even once a trade is 8–15 points onside.<br>Evidence: Part 1 §9.2 — 'PARKED, killed by applying it to the winners too'. · verification: <b>It 'saves' \$826 when you only apply it to the losers and costs \$435 when you apply it to every lot, which is the only honest way to apply it.</b> · mode: no-op | <b>–\$435 applied honestly. The +\$826 was an artefact of the test, not a result.</b>                            | Arms: n/a<br>Exit: n/a<br>Kill: Revive if: re-specified so it cannot be evaluated on losers alone.                                                | Garrath<br>n/a |
| 3 | <b>Restoring abs_veto_short's ATR ≥ 16 floor</b><br><div>REFUTED</div> Do not restore it.<br><b>Mechanism:</b> An ATR floor beneath the short side of abs_veto.<br>Evidence: Part 2 §10 disposition. · verification: <b>Named test: score the entries the floor would block. They are +\$233.50 over 38 shadow trades — the floor blocks winners. The timing fitted; the mechanism did not.</b> · mode: no-op                         | <b>+\$233.50 of blocked entries. Only two gates on this desk carry an ATR floor and this is not one of them.</b> | Arms: n/a<br>Exit: n/a<br>Kill: n/a — the entries it blocks make money.                                                                           | Garrath<br>n/a |
| 4 | <b>A grind ER-0.35 floor as an arming mechanism</b><br><div>REFUTED</div> Do not ship it.<br><b>Mechanism:</b> Arm grind_long only above an efficiency of 0.35.<br>Evidence: Part 2.5 Part B. · verification: <b>Three named tests, all failed: the fake-filter guard (it keeps 1 of the 10 best trades),</b>                                                                                                                         | <b>ER filters keep coming back and keep failing the same way. This is the third time.</b>                        | Arms: n/a<br>Exit: n/a<br>Kill: n/a — keeping 1 of the 10 best trades is not a threshold problem.                                                 | Garrath<br>n/a |

| #  | WHAT TO DO                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           | THE NUMBER BEHIND IT                                                                                                                                                    | ARMS WHEN · EXIT · KILL CRITERION                                                                                                                                                                                                                   | OWNER & REVERT |
|----|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------|
|    | strip-3 (–\$1,044) and the placebo (beaten by 18 of 400 shuffles' worth of noise). · mode: no-op                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |                                                                                                                                                                         |                                                                                                                                                                                                                                                     |                |
| 5  | <b>grind entry trend-confirmation (require a 15-minute impulse our way)</b><br><b>REFUTED</b> Do not ship it.<br><b>Mechanism:</b> Require a confirming 15-minute impulse before grind may enter.<br>Evidence: Part 2.5 Part B. · verification: <b>Named tests: placebo (beaten by 391 of 400) and strip-3 (–\$1,862). It fails for a reason rather than by luck — grind buys weakness, so demanding confirmation removes the entry.</b> · mode: no-op                                                                                                                                                                                                                                                               | <b>Beaten by 391 of 400 shuffles.</b>                                                                                                                                   | <b>Arms:</b> n/a<br><b>Exit:</b> n/a<br><b>Kill:</b> n/a — the mechanism contradicts what the gate is for.                                                                                                                                          | Garrath<br>n/a |
| 6  | <b>Cooldowns, streak-benches and a tighter net_min on exhaustion_short</b><br><b>REFUTED</b> Do not ship any of the three.<br><b>Mechanism:</b> Three separate proposals to fix a red gate at the ENTRY.<br>Evidence: Part 1.5 §3, §4 and §6 — the full candidate board. · verification: <b>Named test on all three: a 4,000-draw random-removal placebo. Thirteen cooldown cells and fifteen bench cells are all negative with a best of the 54th percentile. The net_min 'win' is a calendar artefact — the live floor is ALREADY 400, every higher floor is worse, the correlation between flow and outcome is r=0.059, and 12 of the 17 signals the winning floor cuts come from a single era.</b> · mode: no-op | <b>Not one of 20 treatments clears a 95th-percentile placebo. Best is 85th, and that one is the calendar artefact.</b>                                                  | <b>Arms:</b> n/a<br><b>Exit:</b> n/a<br><b>Kill:</b> n/a — the entry was never the problem. See BUILD #1 for what is.                                                                                                                               | Garrath<br>n/a |
| 7  | <b>Exit tuning for the thrust family</b><br><b>REFUTED</b> Do not spend another night on it.<br><b>Mechanism:</b> Sweep the thrust family's exits to rescue it on chop days.<br>Evidence: Part 2 §8 disposition. · verification: <b>Named test: all 58 policies span –\$39.4 to –\$40.6 per signal on SCALP-CHOP. The spread across every exit ever tried is about a dollar.</b> · mode: no-op                                                                                                                                                                                                                                                                                                                       | <b>58 policies, \$1.20 of spread between the best and the worst.</b>                                                                                                    | <b>Arms:</b> n/a<br><b>Exit:</b> n/a<br><b>Kill:</b> n/a — no exit rescues an un-vetoed thrust in chop. The veto is the fix and it is already live.                                                                                                 | Garrath<br>n/a |
| 8  | <b>Moving the untradeable cutoff to 35</b><br><b>REFUTED</b> Do not move it, even though leave-one-out picks 35 in 17 of 22 folds.<br><b>Mechanism:</b> Lower the untradeable cutoff so more days are refused.<br>Evidence: Part 2.6 §5 disposition. · verification: <b>Named test: base-rate contamination. Post-15:00Z trading lost –\$1,712 across all 22 days, so ANY rule that switches the desk off in the afternoon scores well — including rules with no information in them. The fold vote is measuring the time of day.</b> · mode: no-op                                                                                                                                                                  | <b>–\$1,712 of afternoon base rate masquerading as a cutoff result.</b>                                                                                                 | <b>Arms:</b> n/a<br><b>Exit:</b> n/a<br><b>Kill:</b> n/a — this is the most instructive kill in the report: a 17-of-22 fold majority that means nothing.                                                                                            | Garrath<br>n/a |
| 9  | <b>The adaptive exit selector</b><br><b>PARKED</b> Do not act on its picks.<br><b>Mechanism:</b> Choose the exit per trade from three candidates.<br>Evidence: Part 2.6 §5 disposition. · verification: <b>It lost to fixed always-SCALP on 4 of 5 days and by \$1,267 on the week, with 59% optimal picks.</b> · mode: no-op                                                                                                                                                                                                                                                                                                                                                                                        | <b>–\$1,267 on the week against doing the simplest possible thing.</b>                                                                                                  | <b>Arms:</b> n/a<br><b>Exit:</b> n/a<br><b>Kill:</b> Revive after the data_quality fix (SATURDAY #3) — its regret numbers are computed by the same job that was counting quarantined trades.                                                        | Garrath<br>n/a |
| 10 | <b>The MGC far-side-depletion BREAK gate</b><br><b>REFUTED</b> Do not build it. The interesting part is that it is backwards.<br><b>Mechanism:</b> Buy the gold break when the order book ahead of it is empty.<br>Evidence: Movement 3 §2 and §9. · verification: <b>Named test: a 2,000-draw random-subset placebo on n=192. It loses –\$2,059.50 at the 4.8th percentile — 1,904 of 2,000 random picks did BETTER. The feature carries real information and the sign is inverted: on gold, a level that breaks into a vacuum is one nobody is defending, and price comes straight back.</b> · mode: no-op                                                                                                         | <b>–\$2,059.50 at the 4.8th percentile. Also: no exit rescues it — median MFE +3.28 ATR against median MAE –3.04 ATR, with 80.4% going a full ATR against us first.</b> | <b>Arms:</b> n/a<br><b>Exit:</b> n/a<br><b>Kill:</b> The FADE reformulation is PARKED and worth one night: build it as a fade with its own entry timing and stop and test on a held-out window. An inverted backtest is a hypothesis, not a result. | Garrath<br>n/a |
| 11 | <b>Narrowing the gold hunt to the biggest moves</b><br><b>REFUTED</b> Do not narrow further — it goes the wrong way.<br><b>Mechanism:</b> The standing escalation: if nothing works on the full census, hunt the top 25, then the top 15.<br>Evidence: Movement 3 §4. · verification: <b>Named test: re-running all 20 filters at each narrowing. The top-25 best cell is three prints (strip-3 –\$212.50); at the top 15 every single rule is negative and the best loses –\$54.00.</b> · mode: no-op                                                                                                                                                                                                               | <b>The footprint gets worse as the moves get bigger, which is the opposite of the premise.</b>                                                                          | <b>Arms:</b> n/a<br><b>Exit:</b> n/a<br><b>Kill:</b> n/a for gold breaks. The escalation itself remains the right method — it is this population that has no footprint in its monsters.                                                             | Garrath<br>n/a |
| 12 | <b>“A better exit would have caught the runs we missed”</b>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |                                                                                                                                                                         | <b>Arms:</b> n/a                                                                                                                                                                                                                                    | Garrath        |

| #  | WHAT TO DO                                                                                                                                                                                                                                                                                                                                                                                                                     | THE NUMBER BEHIND IT                                                                                                             | ARMS WHEN · EXIT · KILL CRITERION                                                                                        | OWNER & REVERT  |
|----|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------|-----------------|
|    | <b>REFUTED</b> Stop proposing it.<br><b>Mechanism:</b> The recurring claim that the idle-gate null is an exit problem.<br>Evidence: Movement 2 — the timing-decay ladder. · verification:<br><b>Named test:</b> n=57 at every rung, tick-honest, on this week's tape. Perfect entries decay from +\$3,799 at the ignition minute to −\$1,073 seven minutes later. Our gates fire late; no exit fixes being late. · mode: no-op | <b>+\$66.60 per run at zero lead, +\$6.80 at five minutes, −\$18.80 at seven. Every minute of confirmation costs real money.</b> | Exit: n/a<br>Kill: n/a — the finding is that we are late, not that we leave badly.                                       | n/a             |
| 13 | <b>Relegating any gate's short side</b><br><b>REFUTED</b> Never relegate a direction.<br><b>Mechanism:</b> The standing temptation to switch off whichever side lost this week.<br>Evidence: Part 2 §4, and the same finding in Part 1. · verification: <b>The shadow book says the side proposed for relegation is the stronger one: +\$2,651 against +\$1,371, and +\$15.50 against +\$8.02 per trade.</b> · mode: no-op     | <b>This is on the card every week because it gets proposed every week.</b>                                                       | Arms: n/a<br>Exit: n/a<br>Kill: n/a — a gate runs two-sided with independently-tuned per-side thresholds. Tune the side. | Garraath<br>n/a |

## STANDING CAVEATS ON EVERY DOLLAR ON THIS CARD

**Every repriced number is a CEILING.** Fills are modelled at the exact stop/target/trail level with no slippage, and the desk's own standing finding is that live losses run **1.1-3.1× modelled**. Treat the policy as the edge and the headline dollar as decoration. Anything on thin n is flagged in-row as a direction, not a measurement.

**The fee is \$1.50 per round trip and the multipliers are not the same.** MNQ is \$2.00 a point; **MGC is \$10.00**. Every gold number in Movement 3 and in Part 2.5 Part A is priced at the gold multiplier — pricing gold with the Nasdaq one would understate it fivefold, and doing the reverse would flatter it fivefold.

**Four method wounds are open and they bound everything above.** `signal_journal.suppressed_by` is NULL in all 222 rows this week, so every claim about what a gate would have done unrouted is a reconstruction from the tape rather than a measurement — BUILD #6, and it is the same defect the scope flagged last week. The router's leakage detector is scoped to a universe that barely trades, so its zero is an empty set and not a clean week — BUILD #3. The nightly rollup reports \$0/\$0/\$0 when the systemd journal has rotated past the day it is backfilling, which is indistinguishable from a day with no value — BUILD #8. And `selector_nightly` counted a quarantined trade five times, inflating one day's regret by 39% — SATURDAY #3.

**Two sections of this report did not run.** The rehabilitation phase timed out and Movement 3's four MNQ hunts never started. Part 1.5 was rebuilt from the artifacts the timed-out phase had already saved, and Movement 3 is

the one gold hunt that finished; both say so at the top. **Nothing on this card rests on a section that is missing**, and the gaps are named as gaps rather than reported as nulls — BUILD #2 is the reschedule.

## § The evidence behind those plays

Everything from here down is the working. Two halves, one document. It opens with **Part 0**, which answers the only question that survives past Friday — are we getting better? — day by day, green or red, straight off the books. **Then the warm half**: what the live desk actually did this week (Part 1, with **the run charts** inside it — every fill of the week marked on its own session's price path, which is the fastest way to see what the prose is describing). Then the section that matters most — **Rehabilitation** (Part 1.5), where the week's one properly red gate is walked through the full wash instead of being benched at face value, and where twenty failed treatments are printed by name because the failures are the point. Then the shadow board and the promotion battery (Part 2), the mid-week leads re-run adversarially on the week's own data including the whole gold programme (Part 2.5), and the **Router Operation Review** (Part 2.6). **Then the cold, tape-first half**: we ignore everything the desk did and ask, from the raw tape and order book, **where was the money and did we show up?** — on MNQ and on gold (M1) — then test whether the gates we already own could have caught the runs we missed (M2), and finally build gates from scratch and backtest them to death, the ones that failed included, that being the point (M3). **Two of those did not fully run this week and both say so at the top of themselves.** Plain English, honest money, every claim a table — and every reprice a ceiling.


### WHAT IS IN HERE

1. **Front** — [the honest headline](#), [the scorecard](#), [WHAT TO ACTUALLY DO](#)

2. **Part 0** — [Are we getting better? — every day the desk has traded, green or red](#)
3. **Part 1** — [The live desk — what the paper tournament and the day rider actually did](#)
4. **Part 1.5** — [Live-desk REHABILITATION — never bench a gate at face value](#)
5. **Part 2** — [The shadow desk, the promotion battery & the regime-flex exit lab](#)
6. **Part 2.5** — [Mid-week musings — every lead from the desk chat, tested hard](#)
7. **Part 2.6** — [ROUTER OPERATION REVIEW — the desk's #1 lever, on the stand](#)
8. **Movement 1** — [The census — the biggest runs this week, and did we show up](#)
9. **Movement 2** — [The idle-gate lab — could the gates we own have caught them?](#)
10. **Movement 3** — [The greenfield lab — the gold book-break hunt, and the hunts that did not run](#)

## **P0** Are we getting better? — every day the desk has traded

One bar per trading day, green or red, total desk. The question every other section answers is what happened this week. This one answers **are we getting better**, which is the only question that survives past Friday.



**The shape of it.** Week 30 was the floor: **0 green days out of 6** and **-3,083.00** — the stretch the operator remembers as “red for many many weeks”. The most recent week ran **3 green out of 5** at **+1,292.50**. Of 5 weeks on record, **2 closed net positive** — and both of them are in the last three.

| WEEK | DAYS                    | TRADED | GREEN | NET       |
|------|-------------------------|--------|-------|-----------|
| W29  | 2026-07-16 - 2026-07-19 | 3      | 1/3   | -661.50   |
| W30  | 2026-07-20 - 2026-07-26 | 6      | 0/6   | -3,083.00 |
| W31  | 2026-07-27 - 2026-08-02 | 6      | 2/6   | +516.00   |
| W32  | 2026-08-03 - 2026-08-07 | 5      | 2/5   | -772.50   |
| W33  | 2026-08-10 - 2026-08-14 | 5      | 3/5   | +1,292.50 |

**What changed underneath.** The Claude-run router went live 2026-07-30, mid-W31 — the first green week is the one it landed in. Its measured effect is on the tail and the frequency of red days, not their depth: the worst day before it was **-1,513.00**, the worst since is **-817.00**, and green days went from **2 of 12** to **6 of 13**. The day-rider joined on 08-06 as a second, slower engine and is unbeaten so far — on the worst day of the last week it was the only thing in profit.

**Read it honestly.** The cumulative total across all 25 days is still **-2,708.50** — this is a desk that has stopped bleeding wholesale, not one that is paid yet. 5 weeks is a short sample, best day +1,349.00 and worst -1,513.00 are both single sessions, and PAPER fills flatter a real book. The trend is real; the destination is not reached.

## **P1 THE LIVE DESK — the week of 10-14 August 2026**

Gaz, here is the week in one sentence: **the desk made money, but almost none of it came from the gate tournament.** The strategy book closed **+\$1,292.50**, and **+\$1,617.50** of that is the day rider — five trades, five winners. The six-slot tournament that we spend nearly all our time tuning took **58 signals across 108 lots and lost \$325.00**. Take the rider away and this was a flat-to-slightly-red week for the gates.

The good news is that the two biggest things that went wrong this year are now measurably fixed. **The stop bug is dead** — 62 stops fired this week and every single one filled; we have now gone 393 trades and seventeen days without a single STOP\_UNFILLED. And **the wide leg of the scale-out beat the tight leg by \$331.50 on fifty identical entries**, which is the



cleanest, best-evidenced number in this whole report because both lots took the same signal at the same price with the same stop — the only thing that differed was where we took profit. That result survived every robustness test I could throw at it.

The bad news is a bookkeeping one you already ruled on, and I am honouring the ruling exactly. On Wednesday a broken loop in the day rider re-booked one dead trade four times, writing **\$1,551 of profit that never existed** and firing real orders behind it. The four rows are flagged and excluded from every number in this section. The real cost was **-\$794**, and per your instruction it stays **unbooked** — so the strategy book and the account differ by \$794 for 08-13 permanently and on purpose. I have labelled every total below with which book it is. I never show one under the other's name.

#### THE ONE-PARAGRAPH VERSION

**Strategy book +\$1,292.50 / account book +\$498.50.** Day rider + **\$1,617.50 on 5 trades (5-0)**; tournament **-\$325.00 on 58 signals** at **-\$5.60** a signal and a 43% hit rate. **SURVIVED:** the stop fix (62 for 62); the wide profit leg beating the tight one by **+\$331.50** on 50 paired entries (placebo  $p=0.004$ , every leave-one-day-out fold positive, still **+\$163** after deleting its five best); “arm for periods” beating “bench on losses” by **\$172** in hard money; and the router's churn dropping from the 5-an-hour we flagged on 08-07 to **0.80 an hour**. **DIED:** an efficiency floor at the entry that looked like **+\$361** until leave-one-day-out showed 89% of it came from Tuesday alone; and a break-even stop that “saves” **\$826** if you only count the losers it rescues and **costs \$435** once you apply it to the winners as well. **THE REAL PROBLEM: \$1,058** of this week's losses came from trades that were up **39 points** at their best and still went the full distance to the stop. That is an exit fault, not an entry fault.

---

# 1 · THE MONEY — two books, and which is which

There are three numbers you could call “this week’s P&L” and only two of them are legitimate. I am putting all three in one table so nobody ever quotes the wrong one again.

| WHICH NUMBER                                       | VALUE       | ROWS | WHAT IT IS, AND WHEN TO USE IT                                                                                                    |
|----------------------------------------------------|-------------|------|-----------------------------------------------------------------------------------------------------------------------------------|
| Raw sum of the trade table                         | +\$2,843.50 | 117  | <b>WRONG. Never quote this.</b> It includes \$1,551 of phantom profit from the 08-13 re-booking loop.                             |
| STRATEGY book (excludes flagged rows)              | +\$1,292.50 | 113  | <b>What the strategies earned. Use this for every gate, expectancy and robustness question — i.e. everywhere in this section.</b> |
| ACCOUNT book (strategy less the unbooked bug loss) | +\$498.50   | —    | What IBKR thinks happened. Strategy book — <b>\$794.00</b> for 08-13. Use this and only this for any venue reconciliation.        |

Every trade in this book is already net of the real fee — **\$1.50 per round trip**, MNQ at **\$2.00 a point**. Total commission for the week was \$189.00 on 126 lots. I checked a dozen rows by hand against entry and exit price; the arithmetic is right.

## DAY BY DAY

| DAY       | TOURNAME NT | SIGNALS | WIN% | DAY RIDER | STRATEGY DAY | WHAT THE TAPE DID (12:00–20:40Z)                                                                                             |
|-----------|-------------|---------|------|-----------|--------------|------------------------------------------------------------------------------------------------------------------------------|
| Mon 08-10 | –\$244.50   | 10      | 30%  | +\$39.00  | –\$205.50    | Tightest day of the week — 181pt range, drifted 123pt down. Nothing to work with.                                            |
| Tue 08-11 | +\$371.50   | 8       | 75%  | +\$189.00 | +\$560.50    | <b>338pt range, two clean down-legs (13:00Z –117pt, 15:00Z –121pt at ER 0.31). The short book’s day.</b>                     |
| Wed 08-12 | –\$660.00   | 16      | 25%  | +\$224.50 | –\$435.50    | <b>Pure chop.</b> 256pt of range for 25pt of net — a 0.10 net-to-range, lowest 30-min efficiency of the week (0.125 median). |
| Thu 08-13 | +\$37.50    | 5       | 40%  | +\$731.00 | +\$768.50    | <b>The trend day. +355pt, 460pt range, and one hour (13:00Z) that did +286pt at ER 0.44 on its own.</b>                      |

| DAY       | TOURNAME<br>NT | SIGNALS | WIN% | DAY RIDER   | STRATEGY<br>DAY | WHAT THE<br>TAPE DID<br>(12:00-<br>20:40Z)                                                                      |
|-----------|----------------|---------|------|-------------|-----------------|-----------------------------------------------------------------------------------------------------------------|
| Fri 08-14 | +\$170.50      | 19      | 53%  | +\$434.00   | +\$604.50       | Grinding down-<br>drift — 262pt<br>range, 127pt net<br>lower, the 14:00Z<br>and 15:00Z hours<br>doing the work. |
| WEEK      | -\$325.00      | 58      | 43%  | +\$1,617.50 | +\$1,292.50     | Three green days,<br>two red.                                                                                   |

**Read that day-type column carefully, because it is the whole story of the tournament's week.** Our book is overwhelmingly short — 101 of 113 lots. It made money on the two days price ground downwards with some persistence (Tuesday and Friday, +\$542 between them) and it lost money on the day that went nowhere in a very busy way (Wednesday, -\$660) and on the day with no range at all (Monday, -\$244.50). That is not a subtle pattern. The gates need the market to keep going after they fire, and on Wednesday it never did once.

## 2 · WHAT SURVIVED — leading with the things that held up

### 2.1 THE STOP BUG IS DEAD, AND I CAN PROVE IT

For most of July this desk was losing money to stops that did not fill — IBKR would not fire a resting stop's trigger on the continuous contract, so the position rode naked until something else closed it. It cost us **26 trades and \$1,377.50** across seven days in July. The last one was on **28 July at 08:02Z**.

| CHECK                                            | NUMBER     | VERDICT                                                         |
|--------------------------------------------------|------------|-----------------------------------------------------------------|
| Trades closed since the last STOP_UNFILLED       | 393        | Zero recurrences in 17 days.                                    |
| Stops fired this week                            | 62         | 62 filled. 62 of 62.                                            |
| Venue fill records for tournament lots this week | 216 of 216 | Every leg has both an entry and an exit execution on the books. |

**Call it fixed.** This is the single biggest improvement in the desk's honest P&L this quarter and it does not show up as a line item anywhere, because a bug that stops happening is invisible. It is worth saying out loud: had July's failure rate continued, this week's 62 stops would have produced roughly four unfilled ones and somewhere around -\$200 of damage that had nothing to do with whether the strategies were any good.

## 2.2 THE WIDE PROFIT LEG BEAT THE TIGHT ONE BY \$331.50 — THE WEEK'S BEST-EVIDENCED RESULT

Every gate on this desk runs as a **scale-out**: one signal opens two one-lot sub-slots. Lot A takes a tight fixed target and banks certainty; Lot B runs wider or trails, and rides the tail. They share the same entry, the same price, and the same stop. That makes this a genuinely controlled comparison — the kind we almost never get in live trading — because the only variable is where we take profit.

| GATE              | LOT A (TIGHT)                      | LOT B (WIDE)            | PAIRED SIGNALS | A TOOK    | B TOOK   | B – A     |
|-------------------|------------------------------------|-------------------------|----------------|-----------|----------|-----------|
| exhaustion_short  | fixed target at 0.75R              | chandelier trail, k 1.5 | 33             | –\$230.00 | –\$4.50  | +\$225.50 |
| abs_veto_short    | fixed target at 1.5R               | fixed target at 2.5R    | 12             | –\$69.00  | –\$4.00  | +\$65.00  |
| abs_veto_long     | fixed target at 1.0R               | fixed target at 1.5R    | 4              | +\$0.50   | +\$42.50 | +\$42.00  |
| capitulation_long | fixed target at 1.5R               | chandelier, k 1.5       | 1              | –\$23.50  | –\$24.50 | –\$1.00   |
| ALL               | identical entries, identical stops |                         | 50             | –\$322.00 | +\$9.50  | +\$331.50 |

**Where the money comes from is the important bit, and it is exactly where it should be.** On the 37 pairs where the two lots did anything different, the losers were near-identical — both lots hit the same shared stop, typically within a dollar of each other. The entire difference is on the winners: A banks around \$36–\$42 and B banks \$56–\$78 on the same move. Splitting it out: on pairs where A won, B added **+\$11.52** a signal; on pairs where A lost, B added **+\$2.79**. It is a profit-taking effect, not a risk effect.

### THE ROBUSTNESS BATTERY ON THAT \$331.50

| TEST                                                                  | RESULT                                               | PASSED?                              |
|-----------------------------------------------------------------------|------------------------------------------------------|--------------------------------------|
| Delete its best single pair                                           | +\$289.50                                            | Yes                                  |
| Delete its best three pairs                                           | +\$220.00                                            | Yes                                  |
| Delete its best five pairs                                            | +\$163.00 (+\$3.62 a signal)                         | Yes — it is not one lucky trade      |
| Leave-one-day-out, all five folds                                     | worst fold still +\$239.50                           | Yes — every fold positive            |
| Every day taken on its own                                            | +\$29.50 / +\$92.00 / +\$69.50 / +\$81.50 / +\$59.00 | Yes — five days out of five positive |
| Leave-one-gate-out, all four folds                                    | worst fold +\$106.00                                 | Yes — not one gate carrying it       |
| Chronological halves                                                  | +\$190.50 then +\$141.00                             | Yes — no decay                       |
| Placebo: randomly relabel which lot is A and which is B, 20,000 times | null mean –\$0.23, p = 0.0040                        | Yes — 80 of 20,000 shuffles beat it  |
| Sign test on the 37 pairs that differed                               | B won 25 of 37, p = 0.0235                           | Yes                                  |

**Nine tests, nine passes.** I could not kill this one and I tried hard. The honest caveat is that it is still one week and 50 pairs, and the mechanism is a profit-

taking width — the sort of thing that can look great in a week where the winners happened to run. But the per-day column is what convinces me: it is positive on the chop day, the trend day and the drift day alike, which is not what a regime fluke looks like.

WHAT THIS ACTUALLY SAYS — AND WHAT IT DOES NOT SAY

It does **not** say “drop Lot A.” The scale-out exists on purpose, and A’s job is to bank something certain. What it says is that **A’s targets are set too tight for what this tape is offering**. exhaustion\_short’s A leg takes profit at **0.75R** — less than it is risking — and it is the single biggest contributor to the gap. The specific, testable recommendation for Saturday is to widen the A targets one notch and re-measure, and the natural place to run that properly is the rehab section, not here.

2.3 THE DAY RIDER HAD ITS FIRST PROPER WEEK — WITH TWO CAVEATS

Last week the rider had **no ledger at all** — it had never written a trade row in its life, and the report had to say so as a measurement gap rather than report a P&L. That is now fixed. It wrote five trades this week and they are all in the book.

| DAY   | SIDE    | ENTRY     | EXIT      | EXIT REASON                          | NET         | VENUE EXECUTION RECORD? |
|-------|---------|-----------|-----------|--------------------------------------|-------------|-------------------------|
| 08-10 | SHORT 2 | 29,759.50 | 29,749.00 | CLOCK_FLAT_RECON                     | +\$39.00    | NO — reconstructed      |
| 08-11 | SHORT 2 | 29,684.00 | 29,636.00 | MANUAL_CLAIM                         | +\$189.00   | NO                      |
| 08-12 | SHORT 2 | 29,876.13 | 29,819.25 | CLOCK_FLAT_RECON                     | +\$224.50   | NO — reconstructed      |
| 08-13 | LONG 2  | 30,024.25 | 30,207.75 | MANUAL_CLAIM<br>(your claim, 14:26Z) | +\$731.00   | NO                      |
| 08-14 | SHORT 2 | 30,194.25 | 30,085.00 | TRAIL                                | +\$434.00   | YES — first ever        |
| WEEK  | —       | —         | —         | —                                    | +\$1,617.50 | 1 of 5                  |

**Five for five, +\$323.50 a trade.** That is a wonderful week and I am not going to pretend n=5 is evidence. Strip the best trade and it is +\$886.50 on four; strip the best two and it is +\$452.50 on three. It stays positive, but with a sample this small that tells you almost nothing you did not already know. What it does confirm is the shape we already believed: one entry a day near the cash open, hold it, and let a trail take it off.

**Two caveats you should hold in your head when you read that +\$1,617.50.**

**(a) Two of the five are reconstructed, not observed.** The 08-10 and 08-12 rows carry the exit reason `CLOCK_FLAT_RECON` — and there is no code path anywhere in this repo that emits that string. They were backfilled by hand after the `CLOSED_ELSEWHERE` bug meant the rider’s 20:40 flat never reached the ledger. They are our best reconstruction of \$263.50, not a venue record of it.

**(b) Four of the five have no execution record at all.** I checked the fills table day by day against the tournament’s known leg count. On Monday through Thursday the fill count matches the tournament exactly — 36, 28, 60 and 18 fills for 18, 14, 30 and 9 legs — leaving no room for the rider’s executions. On Friday there are 76 fills for 74 tournament fills, and the two extra are rider-73 and rider-78, ingested at 17:56Z. **Friday 08-14 is the first day in this desk’s history that the day rider’s executions reached the venue record.** The fills-gap fix works; it just only covers one of the five trades below it.

**2.4 “ARM FOR PERIODS” BEAT “BENCH ON LOSSES” — THE CONTRADICTION, SETTLED IN MONEY**

There has been an unresolved argument sitting in the scope for two weeks: three stopped trades in a row is a textbook “wall of stop, bench it” trigger, but arming is supposed to be a permission window where early losses are the price of the option. Those two policies are directly opposed. This week gave us four clean instances, so I priced the rule instead of arguing about it.

The rule tested: after K consecutive signals where every lot stopped out, bench that gate for the rest of the day.

| DAY                                                | GATE             | WOULD HAVE BENCHED AT | SIGNALS FORFEITED | \$ FORFEITED      |
|----------------------------------------------------|------------------|-----------------------|-------------------|-------------------|
| 08-10                                              | exhaustion_short | 14:55Z                | 3                 | −\$92.50 (saved)  |
| 08-12                                              | exhaustion_short | 13:47Z                | 11                | −\$103.00 (saved) |
| 08-14                                              | abs_veto_short   | 14:20Z                | 1                 | +\$11.50 (lost)   |
| 08-14                                              | exhaustion_short | 13:42Z                | 12                | +\$356.00 (lost)  |
| NET EFFECT OF THE BENCH-ON-2-STOP RULE, WHOLE WEEK |                  |                       |                   | −\$172.00         |

**The rule costs money, and the reason is one specific session.** On Friday `exhaustion_short` opened with two stops back to back — 13:38Z for −\$87.50 and 13:42Z for −\$107.00 — which is exactly the pattern that screams “bench it.” Bench it there and you forfeit **+\$356.00**, because from 14:19Z onward that same

gate went on to take twelve more signals and finish the day **+\$310.00**. The two opening stops were the price of the option, precisely as the arm-for-periods argument says.

**So what is the discriminator?** It is the tape the losses happen in, not the losses. On Friday the two opening stops happened right before the 14:00Z hour did  $-108\text{pt}$  at ER 0.22 — a genuine down-leg that a short gate is built for. On Wednesday the equivalent pair happened inside an hour doing  $-33\text{pt}$  at ER 0.05, which is nothing at all, and there the bench would have saved \$103. **The instrument to bench on is the regime, not the scoreboard.** At K=3 the rule fires only once all week and saves \$71.50 — but that is a single instance and I would not build policy on it.

### 3 · THE GATE CARDS

Two views. The first counts **signals** — one decision by one gate, whichever way its two lots ended up. That is the honest unit for judging a gate. The second counts **lots**, which is what the ledger holds and what you see on the dashboard. Confusing them is how we have historically double-counted a bad afternoon.

#### BY SIGNAL — THE GATE’S OWN SCORECARD

| GATE              | SIGNALS | NET         | \$ / SIGNAL | WIN% | BEST      | WORST     | VERDICT                 |
|-------------------|---------|-------------|-------------|------|-----------|-----------|-------------------------|
| exhaustion_short  | 38      | −\$218.50   | −\$5.75     | 42%  | +\$123.00 | −\$145.00 | HOLD — regime-dependent |
| abs_veto_short    | 14      | −\$64.50    | −\$4.61     | 50%  | +\$110.00 | −\$111.50 | HOLD                    |
| abs_veto_long     | 5       | +\$6.00     | +\$1.20     | 40%  | +\$132.00 | −\$133.00 | HOLD — thin             |
| capitulation_long | 1       | −\$48.00    | −\$48.00    | 0%   | —         | −\$48.00  | NOT GRADEABLE — n=1     |
| day_rider         | 5       | +\$1,617.50 | +\$323.50   | 100% | +\$731.00 | +\$39.00  | LIVE — thin but clean   |
| TOURNAMENT TOTAL  | 58      | −\$325.00   | −\$5.60     | 43%  | —         | —         | —                       |

**Three gates on the roster never traded at all this week:** grind\_long (armed for about seven hours across two windows, zero fires), rgv\_short and both nipc sides (benched throughout). grind\_long’s silence is not a mystery — its ATR floor is 22 and the tape ran 5.7–12.5 all Friday. The roster is eight base gates now, not the six the scope doc still lists; rgv\_long and thrust\_short were retired earlier and replaced by exhaustion\_short and the two abs\_veto sides.



BY LOT — AND THE TIGHT-LEG PROBLEM IN ONE TABLE

| GATE                | A LOTS | A NET     | B LOTS | B NET    | B – A     |
|---------------------|--------|-----------|--------|----------|-----------|
| exhaustion_short    | 36     | −\$250.00 | 35     | +\$31.50 | +\$281.50 |
| abs_veto_short      | 14     | −\$60.50  | 12     | −\$4.00  | +\$56.50  |
| abs_veto_long       | 5      | −\$36.50  | 4      | +\$42.50 | +\$79.00  |
| capitulation_long   | 1      | −\$23.50  | 1      | −\$24.50 | −\$1.00   |
| ALL TOURNAMENT LOTS | 56     | −\$370.50 | 52     | +\$45.50 | +\$416.00 |

**Look at the bottom line.** The tight leg lost \$370.50 this week. The wide leg made \$45.50. The whole of the tournament’s red number is the certainty leg. That is not a coincidence and it is not noise — section 2.2 shows it survives a placebo test at p=0.004.

HOW THE WEEK’S MONEY ACTUALLY LEFT THE BUILDING

| EXIT REASON      | LOTS | NET         | WHAT IT MEANS                                                                                  |
|------------------|------|-------------|------------------------------------------------------------------------------------------------|
| STOP             | 62   | −\$2,456.50 | The native 1-ATR stop, server-side. All 62 filled.                                             |
| TARGET           | 35   | +\$1,701.50 | A fixed R target hit.                                                                          |
| MANUAL_CLAIM     | 8    | +\$1,061.50 | Your hands on the claim button — 6 tournament lots (+\$141.50) and 2 rider trades (+\$920.00). |
| TRAIL            | 1    | +\$434.00   | The rider’s ATR-scaled trail, Friday.                                                          |
| CHANDELIER       | 5    | +\$288.50   | The B-leg trail on exhaustion_short. Five fires, all green, \$57.70 apiece.                    |
| CLOCK_FLAT_RECON | 2    | +\$263.50   | <div>RECONSTRUCTED</div> — no code emits this string. Backfilled by hand.                      |

**Your claim button was the third-best exit on the desk this week** — eight lots for +\$1,061.50, and that is before you count that the 08-13 claim at 14:26Z is what took the rider off at +\$731.00 near the top of the day’s big hour. Worth noticing, given it is the one exit with no configuration at all.

4 • WHERE THE \$2,430 OF LOSSES ACTUALLY WENT

This is the part I would read twice. I took all 33 losing signals and asked one question of the raw tape: in the fifteen minutes after we entered, how far did the trade go our way before it went wrong? That single number separates an entry problem from an exit problem, and it splits this week almost perfectly into thirds.

| BUCKET                            | SIGNALS | NET         | MEDIAN BEST POINT | MEDIAN WORST POINT | DIAGNOSIS                                                                |
|-----------------------------------|---------|-------------|-------------------|--------------------|--------------------------------------------------------------------------|
| Never went more than 10pt our way | 12      | -\$855.50   | 5.5pt             | 39.4pt             | <b>ENTRY fault.</b> Wrong from the first tick. No exit rule saves these. |
| Went 10-20pt our way              | 10      | -\$516.50   | 14.2pt            | 17.9pt             | Never enough to bank after costs. Marginal signals.                      |
| Went 20pt+ our way, then died     | 11      | -\$1,058.00 | 39.2pt            | 33.5pt             | <b>EXIT fault.</b> We had the money and gave every penny back.           |

The biggest single bucket of losses this week — \$1,058 of it — is trades that were up a median of 39 points and still travelled the whole way to a full stop. On MNQ 39 points is \$78 a lot. Eleven signals, twenty-two lots, and not one of them had any mechanism that could bank a partial or protect an entry price, because giveback is forced off on both lots by the scale-out config.

## AND THE WINNERS ARE LEAVING MONEY ON THE TABLE TOO

Same measurement, other direction. For each of the 25 winning signals I took the best price the tape offered in the fifteen minutes after entry, priced it at \$2 a point less the round trip, and compared it to what each lot actually banked.

| GATE             | WINNING SIGNALS | WHAT THE TAPE OFFERED (PER LOT) | A BANKED  | A CAPTURE | B BANKED    | B CAPTURE |
|------------------|-----------------|---------------------------------|-----------|-----------|-------------|-----------|
| exhaustion_short | 16              | \$1,634.50                      | +\$592.50 | 38%       | +\$795.50   | 50%       |
| abs_veto_short   | 7               | \$889.50                        | +\$250.50 | 28%       | +\$236.50   | 36%       |
| abs_veto_long    | 2               | \$129.50                        | +\$94.00  | 73%       | +\$136.00   | 105%      |
| ALL WINNERS      | 25              | \$2,653.50                      | +\$937.00 | 37%       | +\$1,168.00 | 49%       |

Narrow it to the 22 winning signals where the tight leg hit its fixed target and it is starker still: the tape offered **\$2,448** and Lot A took **\$886.50 — 36% of it**. Put the two findings side by side and the week has a single coherent message. **We are cutting winners at a third of what is there, and letting 39-point winners travel all the way back to a full stop.** Both are exit faults, in opposite directions, and they are the reason a book with a perfectly respectable 43% hit rate finished red.

### ON THE WIN RATE — THE STANDING RULE, HONOURED

43% is not the problem and I am not treating it as one. Nothing in this section is judged on win percentage. Every verdict below rests on

**expectancy** (\$ per signal after the real \$1.50 round trip) and on **robustness** — placebo, strip-the-best, leave-one-day-out, both-halves. Where I kill something, I name the test that killed it. For the record: the desk's two best days this week ran at 75% and 53%, and its worst ran at 25% — but Monday lost money at 30% while Thursday made \$768.50 at 40%. The hit rate is not carrying the information.

## 5 • THE FLAGSHIP CASE STUDY — Wednesday 12 August, the day the desk fought the tape

Wednesday is the most instructive session of the week by a distance. The tournament lost **\$660.00** across 16 signals at a 25% hit rate — the worst day in the ledger this month — and the reason is written all over the tape. **Price travelled 256 points of range and finished 25 points from where it started.** That is a net-to-range of **0.10**, the lowest of the week, on the lowest 30-minute efficiency of the week (0.125 median). It was not a violent day and it was not a quiet day. It was a busy day going nowhere, which is the single worst environment for a book of fade gates.



Every tournament signal and the day rider, marked on the real 1-minute path. Arrows are entries, squares are exits — one square per lot, so a signal with two lots shows two. Green means the signal made money, red means it lost, purple is the rider. The cluster between 13:30Z and 14:55Z is where the \$660 went.

## 5.1 THE MOMENT THAT SAYS EVERYTHING: 10 SECONDS, 4 POINTS, OPPOSITE DIRECTIONS

At **13:50:01Z** `abs_veto_long` bought MNQ at **29,960.25**. At **13:50:11Z** — ten seconds later — `exhaustion_short` sold it at **29,956.00**. Same desk, same instrument, four and a quarter points apart, pointing opposite ways.

| TIME      | GATE                          | SIDE  | ENTRY     | OUT AT    | HOW IT ENDED                    | NET       |
|-----------|-------------------------------|-------|-----------|-----------|---------------------------------|-----------|
| 13:50:01Z | <code>abs_veto_long</code>    | LONG  | 29,960.25 | 13:51:27Z | Both lots stopped in 86 seconds | −\$133.00 |
| 13:50:11Z | <code>exhaustion_short</code> | SHORT | 29,956.00 | 13:53:58Z | A took target, B trailed out    | +\$123.00 |

The short was right and the long was wrong, and they very nearly cancelled out — we paid four round trips to discover which way the next ninety seconds went. **The router armed `abs_veto_long` at 12:35Z**, on what was at the time a genuinely unambiguous reading (60-minute bias +100pt at ER 0.38, sustained across three ticks), and **benched it at 13:55Z — four minutes after the loss**. That four minutes is the honest, measured cost of a polling router: **\$133.00**. The router's own graded ledger marks that bench ✓ GOOD, and it is right to — the gate never traded again. But the arm that preceded it is the reaction-lag bill, and this is the first time this month it has been anything other than free.

## 5.2 THE BLEED ITSELF, SIGNAL BY SIGNAL

| TIME        | GATE                           | ENTRY     | BOTH LOTS ENDED                     | SIGNAL    | RUNNING   |
|-------------|--------------------------------|-----------|-------------------------------------|-----------|-----------|
| 13:31:03    | <code>capitulation_long</code> | 29,962.50 | Stopped in <b>4 seconds</b>         | −\$48.00  | −\$48.00  |
| 13:40:11    | <code>exhaustion_short</code>  | 29,863.75 | Stopped, 2m16s                      | −\$134.00 | −\$182.00 |
| 13:47:27    | <code>exhaustion_short</code>  | 29,911.00 | Stopped, 1m25s                      | −\$130.50 | −\$312.50 |
| 13:50:01    | <code>abs_veto_long</code>     | 29,960.25 | Stopped, 1m26s                      | −\$133.00 | −\$445.50 |
| 13:50:11    | <code>exhaustion_short</code>  | 29,956.00 | Target + chandelier                 | +\$123.00 | −\$322.50 |
| 14:06:49    | <code>exhaustion_short</code>  | 29,903.50 | Stopped                             | −\$121.50 | −\$444.00 |
| 14:11:43    | <code>exhaustion_short</code>  | 29,932.00 | Target + chandelier                 | +\$97.00  | −\$347.00 |
| 14:21:27    | <code>exhaustion_short</code>  | 29,877.75 | Stopped, 48 seconds                 | −\$91.00  | −\$438.00 |
| 14:25:13    | <code>exhaustion_short</code>  | 29,915.50 | Stopped                             | −\$102.00 | −\$540.00 |
| 14:33:01    | <code>abs_veto_short</code>    | 29,884.00 | Stopped                             | −\$111.50 | −\$651.50 |
| 14:40:49    | <code>exhaustion_short</code>  | 29,914.25 | Target + chandelier                 | +\$112.50 | −\$539.00 |
| 14:50-16:48 | five more signals              | —         | Two small greens, three small stops | −\$121.00 | −\$660.00 |

Read the entry prices down that column and you can see the day. 29,863 → 29,911 → 29,960 → 29,956 → 29,903 → 29,932 → 29,877 → 29,915 → 29,884 →

29,914. It is a fifty-point band being sawn back and forth for an hour, and we sold every push up and got stopped by the next one. Five of the nine exhaustion\_short stops that day were over inside three minutes.

**The important thing about Wednesday is that nothing malfunctioned.** Every stop filled. Every alarm fired — the bleed alarm went off at 13:49Z and the router responded to it inside nineteen seconds; the wall-of-stop alarm went off at 14:26Z and again at 14:27Z. The router looked at exhaustion\_short bleeding \$321 over 22 lots, said so in plain words in its own log, and correctly concluded **it had no rule that authorised a bench** — there was no trend running against the fade to point at, because there was no trend at all. The instruments all worked. **The rule set had no answer for “busy and directionless.”** That is the finding, and it is a rules gap, not a bug.

**5.3 AND THAT MORNING, THE SWITCH FILE FLIPPED SEVEN TIMES FOR NOTHING**

Before any of the above, between 08:25Z and 12:35Z, the router flipped abs\_veto\_short on and off **seven times** — five of those inside fifty minutes. It produced **zero trades and zero dollars**. Two adjacent readings show why:

| TICK   | 60-MIN DIRECTION | 60-MIN NET | ER    | DECISION                                                 |
|--------|------------------|------------|-------|----------------------------------------------------------|
| 09:05Z | FLAT             | +39pt      | 0.252 | Re-arm — misses the $\pm 40$ pt line by <b>one point</b> |
| 09:15Z | UP               | +42pt      | 0.268 | Bench — clears the same line by <b>two points</b>        |

**An arm/bench round trip decided by three points of a sixty-minute net** — about one point of MNQ noise every twenty minutes. The direction state has a two-tick confirmation, but the  $\pm 40$ pt threshold underneath it has no hysteresis and no dwell at all, so a tape grinding either side of 40 manufactures churn that is perfectly correct by the rule and pure noise on the tape. It was free on Wednesday morning only because the gate happened to be silent. That afternoon, the 13:40Z re-arm off the same mechanism led directly to the 14:33Z fire that stopped for — **\$111.50. A dwell-time requirement on that threshold is the clearest single process fix in this section.**

---

## 6 · THURSDAY — the trend day, the rider’s \$731, and the phantom

Thursday is the mirror image of Wednesday: **+355 points of net travel on a 460-point range**, and one hour — 13:00Z — that did **+286 points at ER 0.44** all by itself. It is the only genuine trend day of the week.



The day rider goes on at 13:38:21Z at 30,024.25 (purple) and you claim it at 14:26Z at 30,207.75 for +\$731.00 — 183 points, near the top of the big hour. The tournament, in green and red, takes five signals around the edges for +\$37.50 between them.

**The tournament essentially sat this one out and that is the correct behaviour, not a failure.** Five signals all week’s quietest count, +\$37.50 net. Its gates are faders; a 355-point trend day is not their tape and the router had grind\_long — the one gate built for it — benched on an ATR floor of 22 that the market never came near. The money on a day like this belongs to the rider, and the rider took it.

### 6.1 THE BLACKOUT NOBODY NOTICED

Here is something I only found by counting log lines. **The router was blind for ten hours and twenty minutes on Thursday.** From 02:40:03Z to 13:00:02Z, 125 consecutive five-minute ticks returned ABORT: no JSON in output → no change. It did not crash, it did not alert, and it did not appear in any scoreboard — each tick logged as “no change,” which is exactly what a healthy quiet tick looks like.

| ROUTER HEALTH, THIS WEEK | TICKS | NOTE                               |
|--------------------------|-------|------------------------------------|
| Total ticks              | 1,440 | Five-minute cadence, all five days |
| Genuine “no change”      | 1,210 | 84%                                |

| ROUTER HEALTH, THIS WEEK                 | TICKS | NOTE                                    |
|------------------------------------------|-------|-----------------------------------------|
| ABORT — produced no decision at all      | 125   | All on 08-13, one unbroken 10h20m block |
| Scheduled skips (CME halt / maintenance) | 45    | 9 a day, expected                       |
| Switch changes actually applied          | 60    | 67 individual gate flips                |

**We got lucky on the timing.** It recovered at 13:05:28Z, twenty-five minutes before the cash open and before the day’s big hour. Had it stayed down another forty minutes it would have been asleep through the only real trend of the week with `abs_veto_long` benched. **Ten hours of silent failure that reads identically to ten hours of calm is a monitoring gap, and it is the most serious thing I found in the router this week** — more serious than the churn, because churn at least announces itself.

## 6.2 THE PHANTOM — WHAT HAPPENED, AND THE TWO BOOKS

You have already ruled on this one and I am not relitigating it, only recording it accurately. At 15:53Z the tournament opened an `abs_veto_short` position; the shared account net went non-zero; and the rider — whose manage branch did not yet test whether its own trade had already closed — woke back up on the dead 13:38 position and re-booked it once a minute for four minutes.

| TIME          | ROW    | BOOKED PROFIT | REALITY                                                                                                                                                         |
|---------------|--------|---------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 15:53:23Z     | id 673 | +\$372.00     | The 13:38 trade had already closed at 14:26Z on your claim. These four rows re-price a position that no longer existed, against a stale trail and a stale peak. |
| 15:54:23Z     | id 674 | +\$404.00     |                                                                                                                                                                 |
| 15:55:23Z     | id 675 | +\$373.00     |                                                                                                                                                                 |
| 15:56:23Z     | id 678 | +\$402.00     |                                                                                                                                                                 |
| TOTAL FICTION | —      | +\$1,551.00   | All four flagged<br>EXCLUDE:day_rider_phantom_rebook_20260813 and excluded from every number in this section.                                                   |

**The real cost was −\$794.00**, derived by subtraction: IBKR’s figure for the day is −\$25.50, the tournament’s independently venue-verified leg is +\$37.50, and the rider’s one real trade is +\$731.00. Per your ruling of 2026-08-14 — “i dont want to book an 800 loss due to a system bug” — that \$794 stays **unbooked**, because it was lost to a broken order loop rather than a trading decision, and the ledger is a strategy book by design. **The obligation that comes with that is labelling, and I have honoured it everywhere:** Thursday is **+\$768.50 of strategy** and **−\$25.50 of account**, and those two numbers never appear under each other’s name.

**One correction worth banking.** That night’s review concluded “no orders were placed — only the rows” and put the real cost at \$0. That was wrong, and it was wrong for a knowable



reason: the rider’s executions were not being recorded anywhere at the time, so there was nothing to check against. The reconciliation against IBKR on 08-14 is what produced the – \$794. **Where the same-night review and the next-day reconciliation disagree, the reconciliation wins** — and the fills-gap fix that landed Friday is what makes that check possible going forward.

## 7 • FRIDAY — the day the short book was built for

Friday is the counter-example to Wednesday, and putting them side by side is the most useful thing in this section. Same book, same gates, same 43%-ish hit rate. Wednesday lost \$660, Friday made \$604.50. **The difference is entirely in whether the market kept going after we fired.**



The rider goes short at 13:38Z at 30,194.25 and trails out at 16:16Z at 30,085.00 for + \$434.00 (purple). exhaustion\_short takes 16 signals down the same slope in green and red for +\$310.00. Note the two red arrows right at 13:38Z and 13:42Z — the two stops that a bench-on-losses rule would have acted on.

| hour (UTC)  | net move | ER       | what the desk did                                                                                 |
|-------------|----------|----------|---------------------------------------------------------------------------------------------------|
| 12:00       |          | +4.0pt   | One exhaustion_short signal, + 0.016 \$99.50. Best single signal of the day.                      |
| 13:00       |          | −40.8pt  | The open. Four signals, three of them stops. −\$215.00. This is where the bench temptation lives. |
| 14:00       |          | −108.5pt | The real down-leg. Three signals, +\$68.00, and the rider running.                                |
| 15:00       |          | −73.2pt  | The leg continues. Five signals, +\$252.00. The day’s payday.                                     |
| 17:00-19:00 |          | +61.4pt  | Drift back up. Six signals, −\$34.00. Small, choppy, roughly a wash.                              |

**The 14:00Z and 15:00Z hours are the entire day.** Two hours of price going one way with enough persistence to be worth trading produced +\$320.00 between them — more than the tournament’s +\$170.50 for the whole day — plus all of the rider’s +\$434.00. Everything else was noise around it. And crucially — the two stops at 13:38Z and 13:42Z that looked like a wall of stop happened eighteen minutes before the good part started.

The wall-of-stop alarm did fire on Friday, at 18:44Z, correctly reporting two stops in six minutes on both legs of `exhaustion_short`. By then the gate was already +\$700 up on the day and the router left it alone; it was benched at 19:45Z on the stay-out meter after one more small loss. That sequence — alarm fires, router declines to act, gate keeps its profit — is what the “arm for periods” policy looks like when it works, and the router’s own ledger graded it ✓ GOOD for exactly that reason.

## 8 · THE DIRECTION ROUTER — the full trial read

The router made **60 switch changes** this week, touching **67 individual gate keys**. Here is what those changes bought.

### 8.1 DID THE ARMING PRODUCE ANYTHING?

I took every “arm” event, found when that gate was next benched, and counted the trades in between. This is the honest measure of whether the router is doing work or just moving a file.

| METRIC                                | VALUE    | READ                                                                                                      |
|---------------------------------------|----------|-----------------------------------------------------------------------------------------------------------|
| Arm events                            | 28       | —                                                                                                         |
| Arms that produced at least one trade | 11 (39%) | For comparison, on 08-07 only 4 of 20 switch changes led to a trade — a related but not identical measure |
| Arms that produced nothing at all     | 17       | Longest was <b>10.7 hours</b> of <code>abs_veto_short</code> armed for zero fires                         |
| Trades taken inside an arm window     | 41 lots  | —\$188.00                                                                                                 |
| grind_long armed time                 | ~7 hours | <b>Zero trades.</b> ATR floor 22 vs a tape running 5.7–12.5.                                              |

## 8.2 CHURN — BETTER THAN LAST WEEK, STILL CONCENTRATED IN ONE GATE

| DAY                                    | SWITCH CHANGES | PER HOUR | NOTE                                                                                                |
|----------------------------------------|----------------|----------|-----------------------------------------------------------------------------------------------------|
| 08-10                                  |                | 9        | 0.38 —                                                                                              |
| 08-11                                  |                | 20       | 0.83 Busiest day. Four flips on one gate in 55 minutes.                                             |
| 08-12                                  |                | 16       | 0.67 Seven flips on abs_veto_short before noon, decided by three points.                            |
| 08-13                                  |                | 11       | 0.46 Plus the ten-hour blackout.                                                                    |
| 08-14                                  |                | 4        | 0.17 <b>Quietest and best day. Across the 12:13-20:13Z span, 100 of 102 ticks logged no change.</b> |
| WEEK, active hours only (12:00-21:00Z) |                | 36       | 0.80 <b>vs the 5.0/hour flagged on 08-07 — a six-fold improvement</b>                               |

The churn problem is genuinely much better and the fixes that shipped after 08-07 deserve the credit. But it has not gone away, it has concentrated: abs\_veto\_short alone accounts for **35 of the 67 gate flips** and **10 of the 13 fast flips** (a reversal within twenty minutes of the previous one). Thirteen of sixty changes — 22% — were fast flips. Every one of them is correct under the rule as written, which is precisely the point: the rule's edge is the problem, not the router's judgement.

## 8.3 TWO SMALLER FAULTS WORTH A LINE EACH

| FAULT                                                                          | EVIDENCE                                                                                                                                                           | COST THIS WEEK                                                                                                  |
|--------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------|
| capitulation_long is being benched by something and re-armed by something else | The router wrote off for it <b>five times</b> and on <b>zero</b> times — yet the switch file reads on right now. Something outside the router log is re-arming it. | \$0 this week (it took one trade). A tug-of-war worth resolving before it matters.                              |
| The signal journal is blind to everything it blocks                            | 222 signal rows this week; block_reason is <b>NULL in all 222</b> . Every row is submitted, filled or nofill.                                                      | We cannot measure what the benches suppressed from this table at all. Bench value has to be priced another way. |

The router's own graded scoreboard, all-time, stands at **GOOD 41 · BAD 10 · CHURN 22 · NEUTRAL 36**. That ratio is healthy on the decisions it grades. The gap is that a decision it never made — ten hours of aborted ticks — grades as nothing at all.

## 9 · TWO IDEAS THAT LOOKED GREAT AND DIED — with the test that killed each

Both of these are things I genuinely expected to survive. Neither did, and the way they failed is more useful than if they had worked.

## 9.1 AN EFFICIENCY FLOOR AT THE ENTRY — PARKED, KILLED BY LEAVE-ONE-DAY-OUT

The observation that started it is real and stark. I measured the 30-minute efficiency of the tape immediately before every signal. **Forty-eight of this week’s 58 signals — 83% of them — fired into a tape with an efficiency below 0.10**, which is a market going nowhere. And that bucket is where the entire loss lives.

| EFFICIENCY OF THE 30 MIN BEFORE ENTRY | SIGNALS | NET       | \$ / SIGNAL | WIN% |
|---------------------------------------|---------|-----------|-------------|------|
| Below 0.10 — a dead tape              | 48      | −\$686.00 | −\$14.29    | 38%  |
| 0.10 to 0.20                          | 9       | +\$229.00 | +\$25.44    | 67%  |
| Above 0.20                            | 1       | +\$132.00 | +\$132.00   | 100% |

Put a floor at 0.10 and the week goes from **−\$325.00 to +\$361.00** on ten signals. That is a \$686 swing and it is exactly the sort of number that gets deployed on a Saturday and hurts us the following Thursday. Here is why I am not deploying it.

| TEST                                                               | RESULT                                                                   | VERDICT                                                                                      |
|--------------------------------------------------------------------|--------------------------------------------------------------------------|----------------------------------------------------------------------------------------------|
| The winners test — does it keep the trades that pay?               | Keeps <b>7 of 25</b> winning signals, and <b>27%</b> of gross winnings   | <b>FAILS.</b> It wins by deleting the winners.                                               |
| Leave-one-day-out                                                  | Drop Tuesday and it collapses from +\$361.00 to <b>+\$38.50</b>          | <b>FAILS.</b> 89% of it is one session.                                                      |
| Parameter sweep — plateau or step?                                 | 0.06→+\$62 · 0.08→+\$169 · 0.10→+\$361 · 0.12→+\$430 · 0.15→+\$95 on n=2 | <b>FAILS.</b> A step, and the edge only appears as n collapses — the classic curve-fit tell. |
| Placebo — best 10 by efficiency vs 20,000 random 10-signal subsets | p = 0.0476                                                               | Marginal at best, one-sided, on the best subset.                                             |
| Is it efficiency, or is it just “not Wednesday”?                   | Still adds +\$28.12/signal with Wednesday removed                        | Passes — the one test it does pass.                                                          |

**Verdict:** **PARKED**, **cause of death leave-one-day-out**. Tuesday alone carries \$322.50 of the \$361.00, and any rule that keeps only seven of twenty-five winners has to clear a much higher bar than one week and ten surviving signals. **What would revive it:** re-run across at least three weeks and require **n ≥ 40 surviving signals** with every weekly fold positive at a floor somewhere in 0.08-0.12. If it holds there it becomes a shadow candidate, not before. The observation underneath it — that we fire 83% of our signals into a dead tape — stands on its own and is worth acting on some other way.

## 9.2 A BREAK-EVEN STOP — PARKED, KILLED BY APPLYING IT TO THE WINNERS TOO

Section 4 showed \$1,058 of losses on trades that were up a median 39 points. The obvious fix is to move the stop to break-even once a trade is a certain distance onside. I replayed it on the 5-second bars, requiring the favourable move to happen in a bar strictly earlier than the exit bar so there is no cheating inside a bar.

| BREAK-EVEN ARMED AT | LOSING LOTS RESCUED “SAVING” | HONEST RESULT — RULE APPLIED TO EVERY LOT |           |
|---------------------|------------------------------|-------------------------------------------|-----------|
| +8 points           | 20                           | +\$826.00                                 | —         |
| +10 points          | 17                           | +\$733.50                                 | −\$304.50 |
| +12 points          | 15                           | +\$602.50                                 | −\$435.50 |
| +15 points          | 11                           | +\$410.50                                 | −\$489.50 |
| +20 points          | 7                            | +\$314.50                                 | +\$103.00 |

**That middle column is a lie and the right-hand column is the truth.** Count only the losers a break-even stop rescues and it looks like \$826 of free money. Apply the same rule to every lot — including the winners that dipped back through their entry price before running — and it **costs** between \$304 and \$490 at every threshold that matters. The one positive cell, +20 points for +\$103, is non-monotonic against its neighbours, which is what noise looks like.

**Verdict:** **PARKED**, cause of death the apply-to-all-lots test. What would **revive it**: a version that arms on something other than a bare price touch — a dwell requirement (price holds beyond +X for N seconds), or arming only when the 30-minute efficiency clears a floor so it fires on moves that are actually going somewhere. The give-back problem it was aimed at is real and worth \$1,058; this particular instrument just is not the answer, and the proper hunt for one belongs in the rehab section.

---

## 10 • WHAT I AM NOT SURE ABOUT

Standing habit: the caveats go in the section, not in a footnote nobody reads.

- **Everything here is one week and it is in-sample.** 58 signals, 108 lots, five sessions. The A-versus-B result is the only thing with enough paired structure

that I would defend it hard; everything else is an observation awaiting a second week.

- **Per-gate samples are small to the point of being decorative.** `abs_veto_long` is five signals, `capitulation_long` is one. I have given them cards for completeness and graded neither. The day rider is five trades — a lovely five, and still five.
- **Two of the rider's five trades are reconstructed, not observed**, and four of five have no venue execution record. \$263.50 of the \$1,617.50 is our best backfill of what happened.
- **The break-even and wall-of-stop counterfactuals are replays, not trades.** They assume our fills would have been where the bars say. The break-even one in particular does not model queue position at the entry price, which on a fast reversal is exactly where you would not get filled — so if anything it flatters the rule I am already rejecting.
- **The signal journal cannot tell us what the benches saved.** All 222 rows have a NULL block reason, so “what would the benched gate have done” is not answerable from that table and I have not pretended otherwise anywhere in this section.
- **The 15-minute window on the MFE work is a choice.** Most of these trades close inside five minutes, so it is generous to the ceiling and therefore harsh on our capture percentages. A shorter window would make our capture look better; I picked the one that flatters us least.

## 11 · DISPOSITION TABLE

| LEAD                                                  | VERDICT               | THE EVIDENCE, OR THE CONDITION                                                                                                                                                                                                                                 |
|-------------------------------------------------------|-----------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| The stop-fill fix (resting stop on a concrete future) | LIVE — CONFIRMED      | 62 of 62 stops filled this week; 393 trades and 17 days since the last <code>STOP_UNFILLED</code> . Nothing further needed.                                                                                                                                    |
| Wide profit-taking beats tight (Lot B over Lot A)     | SURVIVED THE BATTERY  | +\$331.50 over 50 identical entries. Placebo $p=0.0040$ , sign test $p=0.0235$ , all 5 daily folds positive, all 4 gate folds positive, +\$163 after deleting its best five. Action: widen the A targets one notch and re-measure — hand to the rehab section. |
| “Arm for periods” over “bench on losses”              | HOLD — CONFIRMED LIVE | Bench-on-2-stops costs −\$172 across the week; Friday alone forfeits +\$356. The discriminator is the regime the losses occur in, not the losses.                                                                                                              |
| The day rider                                         | LIVE — thin           |                                                                                                                                                                                                                                                                |

| LEAD                                                       | VERDICT                     | THE EVIDENCE, OR THE CONDITION                                                                                                                                                                                |
|------------------------------------------------------------|-----------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                                                            |                             | <b>5 trades, +\$1,617.50, 5-0. Now writes a ledger; venue fills arrived 08-14. Needs weeks, not tests.</b>                                                                                                    |
| exhaustion_short as a standalone gate                      | HOLD — regime-dependent     | —\$218.50 over 38 signals, but leave-one-day-out swings from +\$149 to —\$528 — it is Friday and Wednesday, nothing else. Not robust either way on one week. Its B leg is already profitable.                 |
| abs_veto both sides                                        | HOLD                        | —\$58.50 over 19 signals combined. Too thin to grade; nothing to fix and nothing to celebrate.                                                                                                                |
| Dwell time on the router's $\pm 40$ pt direction threshold | SHADOW — recommend building | Seven flips on one gate before noon on 08-12, one round trip decided by <b>three points</b> . 22% of all changes were fast flips. The clearest process fix in this section.                                   |
| Entry efficiency floor at 0.10                             | PARKED                      | <b>Killed by leave-one-day-out</b> (Tuesday carries 89%) and the winners test (keeps 7 of 25). Revive if: 3+ weeks, $n \geq 40$ surviving signals, every weekly fold positive at a floor in 0.08-0.12.        |
| Break-even stop at +8 to +15pt                             | PARKED                      | <b>Killed by the apply-to-all-lots test</b> — “saves” \$826, actually costs \$435. Revive if: re-specified with a dwell requirement or an efficiency condition rather than a bare price touch.                |
| Router blackout monitoring                                 | MEASUREMENT GAP             | 125 consecutive aborted ticks over 10h20m on 08-13, indistinguishable from calm in every scoreboard. <b>Needs an alarm on consecutive no-decision ticks</b> . No dollar cost this week purely by timing luck. |
| Signal journal block reasons                               | MEASUREMENT GAP             | NULL in all 222 rows. Bench value is unpriceable from this table. Same defect class the scope flagged last week — still open.                                                                                 |
| Day rider venue fill coverage                              | MEASUREMENT GAP — closing   | 4 of 5 trades have no execution record; 2 are hand-reconstructed. The fix landed and covered 08-14. Re-check next Friday that it covers 5 of 5.                                                               |
| capitulation_long arm/bench tug-of-war                     | PARKED — unexplained        | Router wrote off five times and on never; the file reads on. Revive/resolve by: finding the other writer. \$0 cost this week.                                                                                 |

## THE ONE STONE STILL UNTURNED

We now know where the money goes — \$1,058 of it into trades that were up 39 points and gave it all back. We know the naive fix for that (a break-even stop) is a mirage. **What we have not tested is a partial:** banking one lot of the two at a fixed points-onside and letting the other run its existing exit. The scale-out already gives us two lots for free, so it costs nothing structurally, and unlike a break-even stop it does not require being right about where the pullback ends. That is the first thing I would test next week, and it is the natural bridge from this section into the rehab work.



# THE RUN CHARTS — every fill of the week, marked on the day's own price path




**Wednesday 08-12** 07:00-21:00 UTC · 31 closed lots, 8 green ·  $-\$435.50$  · tape +70pt net on a 274pt range ▲ long in ▼ short in ● out ● win ● loss



**Thursday 08-13** 07:00-21:00 UTC · 10 closed lots, 5 green ·  $+\$768.50$  · tape +328pt net on a 460pt range ▲ long in ▼ short in ● out ● win ● loss



**Friday 08-14** 07:00-21:00 UTC · 38 closed lots, 19 green ·  $+\$604.50$  · tape -18pt net on a 262pt range ▲ long in ▼ short in ● out ● win ● loss



## **P1.5** LIVE-DESK REHABILITATION — **exhaustion\_short**, and the honest failures behind it

---

### ★ READ THIS FIRST — WHERE THIS SECTION CAME FROM

The dedicated rehabilitation phase of tonight's run **timed out at its 90-minute ceiling and never wrote its chapter** (rehab: rc=-1 90.0m artifact=MISSING in data/friday\_durable.log). It did, however, finish and save **the entire battery it was told to run** — 104 live **exhaustion\_short** signals with the tick path after each one, four treatment grids, a 4,000-draw placebo behind every cell, two independent exit sweeps and a random-entry control. **This section is assembled from those finished numbers**, read out of the JSON at build time rather than retyped, so the figures are the study's own. What is missing is the prose that phase would have written, and two rehabs it never got to. Those are named in the gap box at the end rather than quietly skipped.

One gate was properly red this week and it is the same one the scope flagged: **exhaustion\_short**, which lost **-\$218.50** across 71 closed lots on the live desk and has now been red for long enough that “bench it” is the obvious call. The standing rule says we do not get to make the obvious call. So the study took **every one of its 104 signals** back to 24 July, rebuilt each one from the tape, and asked the four questions in order: are we counting malfunctions as strategy, is the true cost right, is the bleed in the BASE, the SIGNAL or the EXIT — and does any proposed fix survive a placebo. **The answer is unusually clean, and it is not the answer anyone expected.** Every treatment aimed at the ENTRY failed — cooldowns, benching after a run of stops, and a tighter flow floor all lose to random. The entry is not the problem. **The exit is:** the same signals, entered identically and held with a wide stop and a far target instead of scalped, are worth **\$3,459.98** against **-\$84.50** as the desk actually traded them, and the entries beat **60 of 60** random-entry controls under the same

exit. The gate is rehabilitated as a SHADOW candidate, not deployed — and the reason it is not deployed is in §7.

## 1 · Normalise the malfunctions first — and this week there are none to strip

The first step of the wash is always to take the software failures out of the strategy verdict, because a gate that was never given a working stop has not been tested. Last week that step mattered enormously. **This week it is a no-op, and that is itself the finding.** Across the whole week the desk filled **62 of 62** stops, with no STOP\_UNFILLED in 393 trades and 17 days — the resting-stop-on-a-concrete-future fix has held. So none of the numbers below need a malfunction adjustment, and every dollar exhaustion\_short lost is a strategy dollar. Of its 104 signals, **52 ended at the stop** — almost exactly half — and all of them were real fills at real prices.

## 2 · The corrected-cost reconstruct — what the gate really is

Priced honestly at **\$1.50 per round trip** and **\$2.00 a point**, over the full signal history rather than just this week, exhaustion\_short is not a disaster. It is something more awkward: **a coin-flip that pays slightly less than nothing.**

| THE GATE, AS TRADED | SIGNALS | NET \$    | WIN % | \$ PER SIGNAL | ENDED AT THE STOP |
|---------------------|---------|-----------|-------|---------------|-------------------|
| LIVE as traded      | 104     | -\$158.00 | 50.0% | -\$1.52       | 52 of 104         |

Every signal since 24 July, tick-honest. A 50% win rate that still loses money is the signature of an exit whose losers are bigger than its winners — hold that thought.

## 3 · Root cause, attempt one: is it the ENTRY? Three treatments, all of them fail

The obvious diagnosis is that the gate keeps re-firing into a move that has already beaten it — it stops out, and then immediately shorts the same push again. That is a testable claim and it produces three concrete fixes, all three of which the operator has asked for by name at some point: **cool off after a stop, bench after a run of stops, and demand more aggressor flow before firing at all.** Every cell below is scored against **4,000 random draws that remove the same number of signals at random.** That column is the one that matters: a filter that

drops 30% of a losing book will usually look like it helped, and the placebo is what tells you whether it actually did anything a coin could not.

3.1 Cool off after a stop — every cell loses, and the best one is a coin flip

| COOLDOWN        | SIGNALS LEFT | NET \$   | WIN % | \$/SIGNAL | SIGNALS CUT     | PLACEBO %ILE | VERDICT |
|-----------------|--------------|----------|-------|-----------|-----------------|--------------|---------|
| cooldown 1min   | 97           | −\$454.5 | 47.4% | −\$4.69   | 7 (6 winners)   | 9.0          | no      |
| cooldown 2min   | 94           | −\$536.5 | 47.9% | −\$5.71   | 10 (7 winners)  | 7.0          | no      |
| cooldown 3min   | 90           | −\$295.5 | 48.9% | −\$3.28   | 14 (8 winners)  | 28.8         | no      |
| cooldown 5min   | 87           | −\$94.5  | 51.7% | −\$1.09   | 17 (7 winners)  | 54.0         | no      |
| cooldown 8min   | 81           | −\$338.0 | 48.1% | −\$4.17   | 23 (13 winners) | 28.5         | no      |
| cooldown 10min  | 76           | −\$314.5 | 46.1% | −\$4.14   | 28 (17 winners) | 30.6         | no      |
| cooldown 15min  | 74           | −\$210.5 | 47.3% | −\$2.84   | 30 (17 winners) | 40.2         | no      |
| cooldown 20min  | 73           | −\$192.0 | 47.9% | −\$2.63   | 31 (17 winners) | 41.8         | no      |
| cooldown 30min  | 67           | −\$345.0 | 46.3% | −\$5.15   | 37 (21 winners) | 29.0         | no      |
| cooldown 45min  | 62           | −\$289.0 | 46.8% | −\$4.66   | 42 (23 winners) | 32.3         | no      |
| cooldown 60min  | 54           | −\$541.5 | 40.7% | −\$10.03  | 50 (30 winners) | 15.2         | no      |
| cooldown 90min  | 46           | −\$551.0 | 41.3% | −\$11.98  | 58 (33 winners) | 14.3         | no      |
| cooldown 120min | 43           | −\$545.0 | 44.2% | −\$12.67  | 61 (33 winners) | 13.3         | no      |

Thirteen cells. **Not one is positive, and not one clears its placebo.** The best cell (5 minutes) still loses money and sits at the 54th percentile of pure noise — i.e. 46% of random cuts of the same size did better. Longer cooldowns get steadily worse because they throw away winners: at 60 minutes it has binned 30 winning signals to avoid 20 losing ones.

3.2 Bench after a run of stops — the “wall of STOP” rule, priced

This one is worth dwelling on, because the scope explicitly asks the report to settle it. “Three stopped trades in a row, bench the gate” is a textbook desk rule and it is directly opposed to the “arm for periods” finding Part 1 confirmed in live money. Here it is applied mechanically to all 104 signals:

| BENCH RULE                  | SIGNALS LEFT | NET \$   | WIN % | \$/SIGNAL | SIGNALS CUT     | PLACEBO %ILE | VERDICT |
|-----------------------------|--------------|----------|-------|-----------|-----------------|--------------|---------|
| bench after 2 stops for 15  | 91           | −\$370.5 | 47.3% | −\$4.07   | 13 (9 winners)  | 20.6         | no      |
| bench after 2 stops for 30  | 84           | −\$376.0 | 47.6% | −\$4.48   | 20 (12 winners) | 23.1         | no      |
| bench after 2 stops for 60  | 74           | −\$306.5 | 45.9% | −\$4.14   | 30 (18 winners) | 30.4         | no      |
| bench after 2 stops for 120 | 61           | −\$480.5 | 44.3% | −\$7.88   | 43 (25 winners) | 19.0         | no      |
|                             | 42           | −\$210.5 | 47.6% | −\$5.01   | 62 (32 winners) | 36.9         | no      |

| BENCH RULE                      | SIGNALS LEFT | NET \$   | WIN % | \$/SIGNAL | SIGNALS CUT        | PLACEBO %ILE | VERDICT   |
|---------------------------------|--------------|----------|-------|-----------|--------------------|--------------|-----------|
| bench after 2 stops for session |              |          |       |           |                    |              |           |
| bench after 3 stops for 15      | 100          | −\$362.0 | 48.0% | −\$3.62   | 4 (4 winners)      | 11.1         | no        |
| bench after 3 stops for 30      | 96           | −\$287.0 | 49.0% | −\$2.99   | 8 (5 winners)      | 27.1         | no        |
| bench after 3 stops for 60      | 90           | −\$282.0 | 47.8% | −\$3.13   | 14 (9 winners)     | 30.7         | no        |
| bench after 3 stops for 120     | 80           | −\$239.0 | 48.8% | −\$2.99   | 24 (13 winners)    | 37.2         | no        |
| bench after 3 stops for session | 75           | −\$78.0  | 50.7% | −\$1.04   | 29 (14 winners)    | 54.2         | no        |
| bench after 4 stops for 15      | 104          | −\$158.0 | 50.0% | −\$1.52   | 0 — never triggers | —            | no effect |
| bench after 4 stops for 30      | 104          | −\$158.0 | 50.0% | −\$1.52   | 0 — never triggers | —            | no effect |
| bench after 4 stops for 60      | 104          | −\$158.0 | 50.0% | −\$1.52   | 0 — never triggers | —            | no effect |
| bench after 4 stops for 120     | 104          | −\$158.0 | 50.0% | −\$1.52   | 0 — never triggers | —            | no effect |
| bench after 4 stops for session | 104          | −\$158.0 | 50.0% | −\$1.52   | 0 — never triggers | —            | no effect |

Fifteen cells, every one of them worse than doing nothing, and the best (bench for the rest of the session after three stops) is the 54th percentile of noise. Note the bottom five rows: a four-stop trigger never fires at all in 104 signals, so the rule people imagine is protecting them is not even switched on.

#### ★ THE CONTRADICTION THE SCOPE ASKED US TO SETTLE

“Bench on a wall of stops” and “arm for periods” cannot both be right, and on this gate **the wall-of-stops rule is simply wrong**. Every version of it loses money here, and the reason is visible in the “signals cut” column: benching after two stops for an hour removes **18 winners to avoid 12 losers**. The losing streak does not predict the next signal; it just tells you the last two went badly. Part 1 reached the same verdict from the live book — bench-on-2-stops would have cost the desk −\$172 across this week and forfeited +\$356 on Friday alone. Two independent populations, same answer. **The discriminator is the REGIME the losses happened in, never the count of losses.**

3.3 Demand more flow — the treatment that looked like the fix

The third entry treatment is the gate’s own net\_min knob: the minimum net aggressor flow the tape must show before it is allowed to fade. Raise it and you fire less often but supposedly better. **This is the one cell in the entire battery that beat its placebo**, and the next section is entirely about why it still is not a fix.

| FLOW FLOOR  | SIGNALS LEFT | NET \$  | WIN % | \$/SIGNAL | SIGNALS CUT     | PLACEBO %ILE | VERDICT      |
|-------------|--------------|---------|-------|-----------|-----------------|--------------|--------------|
| net_min 400 | 87           | \$483.0 | 52.9% | \$5.55    | 17 (6 winners)  | 97.6         | beats random |
| net_min 450 | 74           | \$233.5 | 52.7% | \$3.16    | 30 (13 winners) | 80.8         | no           |
| net_min 500 | 62           | \$362.5 | 54.8% | \$5.85    | 42 (18 winners) | 84.9         | no           |
| net_min 550 | 49           | \$109.5 | 53.1% | \$2.23    | 55 (26 winners) | 66.0         | no           |
| net_min 600 | 44           | \$315.5 | 54.5% | \$7.17    | 60 (28 winners) | 80.8         | no           |
| net_min 650 | 38           | \$164.0 | 52.6% | \$4.32    | 66 (32 winners) | 71.2         | no           |
| net_min 700 | 32           | −\$3.5  | 50.0% | −\$0.11   | 72 (36 winners) | 55.0         | no           |
| net_min 800 | 24           | −\$61.5 | 50.0% | −\$2.56   | 80 (40 winners) | 47.8         | no           |
| net_min 900 | 21           | −\$60.5 | 47.6% | −\$2.88   | 83 (42 winners) | 48.0         | no           |

net\_min 400 turns a −\$158 book into +\$483 and sits at the 97.6th percentile of its placebo. On this table it is the answer. It is not the answer.

4 · ★ The autopsy on the only winner — net\_min 400 is already switched on

Here is the trap, and it is a good one, because the number was real and the conclusion drawn from it would have been completely wrong. **The live gate already runs a net\_min of 400.** The 104-signal population the first grid scored includes signals the live gate would never have taken, so “apply net\_min 400” was not a proposed change at all — it was the study re-discovering the filter that is already deployed. Once you take the live floor as the baseline and ask the only question that is actually actionable — should the floor go HIGHER? — the effect evaporates:

| FLOOR                     | SIGNALS | NET \$  | WIN % | \$/SIGNAL | CUT        | WINNERS FORGONE | PLACEBO %ILE | VS THE LIVE FLOOR      |
|---------------------------|---------|---------|-------|-----------|------------|-----------------|--------------|------------------------|
| live 400 floor (baseline) | 70      | \$556.5 | 52.9% | \$7.95    | —          | —               | —            | this is the desk today |
| net_min 450               | 57      | \$307.0 | 52.6% | \$5.39    | 30 (34.5%) | \$640.0         | 81.7         | worse than live        |
| net_min 500               | 45      | \$436.0 | 55.6% | \$9.69    | 42 (48.3%) | \$945.5         | 87.8         | worse than live        |
| net_min 550               | 32      | \$183.0 | 53.1% | \$5.72    | 55 (63.2%) | \$1,592.0       | 69.9         | worse than live        |



| FLOOR       | SIGNALS | NET \$  | WIN % | \$/SIGNAL | CUT        | WINNERS FORGONE | PLACEBO %ILE | VS THE LIVE FLOOR |
|-------------|---------|---------|-------|-----------|------------|-----------------|--------------|-------------------|
| net_min 600 | 27      | \$389.0 | 55.6% | \$14.41   | 60 (69.0%) | \$1,729.5       | 85.1         | worse than live   |
| net_min 650 | 21      | \$237.5 | 52.4% | \$11.31   | 66 (75.9%) | \$2,025.0       | 75.2         | worse than live   |
| net_min 700 | 15      | \$70.0  | 46.7% | \$4.67    | 72 (82.8%) | \$2,319.0       | 60.9         | worse than live   |
| net_min 800 | 7       | \$12.0  | 42.9% | \$1.71    | 80 (92.0%) | \$2,655.0       | 53.8         | worse than live   |
| net_min 900 | 4       | \$13.0  | 25.0% | \$3.25    | 83 (95.4%) | \$2,761.0       | 55.5         | worse than live   |

Nine floors above the deployed 400. **Every single one is worse than leaving it alone**, none of them clears a 95th-percentile placebo, and the damage grows monotonically with the winners you forgo — by net\_min 900 the floor has thrown away \$2,761 of winning signals to dodge \$2,858 of losing ones, on n=4.

Three more checks, all pointing the same way. **First, the correlation.** Above the live floor, the relationship between how much flow a signal had and how much money it made is  $r = 0.059$  on  $n = 70$  ( $t = 0.49$ ). That is nothing. More flow does not mean a better trade. **Second, the quintiles.** If the knob worked at all, the money would climb as you go up the flow ranking. It does not — it zig-zags:

| QUINTILE OF FLOW | NET RANGE | N  | NET \$   | \$/SIGNAL | WIN % | STOPPED % |
|------------------|-----------|----|----------|-----------|-------|-----------|
| Q1               | 405-451   | 14 | \$104.5  | \$7.46    | 50.0% | 50.0%     |
| Q2               | 463-511   | 14 | \$295.5  | \$21.11   | 57.1% | 42.9%     |
| Q3               | 513-555   | 14 | -\$340.5 | -\$24.32  | 42.9% | 57.1%     |
| Q4               | 563-701   | 14 | \$575.0  | \$41.07   | 71.4% | 28.6%     |
| Q5               | 706-1217  | 14 | -\$78.0  | -\$5.57   | 42.9% | 57.1%     |

The best quintile (Q4) and the worst (Q3) are adjacent, and the very highest-flow quintile is a loser. There is no monotone anything here — this is noise with five buckets drawn on it.

**Third, and this is what actually kills it: the era confound.** The 17 signals that the “winning” floor removes are worth -\$641.0 and contain 6 winners — but **12 of those 17 come from a single era**. The filter is not selecting for weak flow; it is selecting for a stretch of calendar. That is a date filter wearing a flow filter’s clothes, and it will not generalise by construction.

**Honest instrument note.** Of the 104 signals, 87 have usable tape; 17 fall on 2026-07-21, 2026-07-22, 2026-07-23, days the rolling capture window no longer holds, and those carry -\$73.5. Every exit number in §6 and §7 is therefore on  $n = 87$ , not 104, and the live comparison is restated on the same 87 so the two sides are the same population.

5 • The absorption veto — a fake win, caught by the winners test

The fourth treatment delays the entry a few seconds and refuses it if price has already gone against us by more than a couple of points — the same shape as the 55-second veto that works on the abs\_veto family. On this gate it fails, and it fails in the most instructive way: **by winning through the destruction of winners.**

| DELAY | ADVERSE CAP | SIGNALS LEFT | SIGNALS CUT     | NET IF SHIFTED | NET AS DESIGNED | PLACEBO %ILE |
|-------|-------------|--------------|-----------------|----------------|-----------------|--------------|
| 5s    | 2.0pt       | 39           | 48 (23 winners) | −\$508.0       | −\$96.0         | 40.1         |
| 5s    | 3.0pt       | 51           | 36 (19 winners) | −\$954.0       | −\$30.0         | 9.2          |
| 5s    | 5.0pt       | 67           | 20 (12 winners) | −\$919.0       | −\$72.5         | 5.2          |
| 5s    | 8.0pt       | 81           | 6 (4 winners)   | −\$517.8       | \$30.0          | 7.2          |
| 10s   | 2.0pt       | 28           | 59 (29 winners) | −\$763.5       | −\$136.0        | 21.2         |
| 10s   | 3.0pt       | 40           | 47 (26 winners) | −\$1,292.0     | −\$87.0         | 3.0          |
| 10s   | 5.0pt       | 52           | 35 (19 winners) | −\$947.5       | −\$144.5        | 18.5         |
| 10s   | 8.0pt       | 66           | 21 (8 winners)  | −\$157.0       | \$83.0          | 86.1         |
| 15s   | 2.0pt       | 24           | 63 (31 winners) | −\$810.5       | −\$38.0         | 24.7         |
| 15s   | 3.0pt       | 33           | 54 (28 winners) | −\$1,171.5     | −\$58.5         | 9.9          |
| 15s   | 5.0pt       | 49           | 38 (19 winners) | −\$865.0       | −\$203.5        | 35.0         |
| 15s   | 8.0pt       | 61           | 26 (10 winners) | −\$463.0       | −\$306.0        | 85.1         |
| 20s   | 2.0pt       | 24           | 63 (31 winners) | −\$811.0       | −\$143.0        | 23.3         |
| 20s   | 3.0pt       | 29           | 58 (28 winners) | −\$844.0       | −\$120.5        | 28.9         |
| 20s   | 5.0pt       | 44           | 43 (20 winners) | −\$961.5       | −\$212.5        | 31.6         |
| 20s   | 8.0pt       | 58           | 29 (12 winners) | −\$888.0       | −\$225.0        | 48.3         |
| 30s   | 2.0pt       | 18           | 69 (32 winners) | −\$465.5       | \$8.0           | 59.7         |
| 30s   | 3.0pt       | 23           | 64 (29 winners) | −\$553.5       | −\$20.5         | 63.4         |
| 30s   | 5.0pt       | 37           | 50 (21 winners) | −\$511.5       | −\$78.5         | 78.3         |
| 30s   | 8.0pt       | 50           | 37 (13 winners) | −\$462.5       | −\$51.0         | 88.6         |
| 45s   | 2.0pt       | 16           | 71 (32 winners) | −\$151.0       | \$148.0         | 79.0         |
| 45s   | 3.0pt       | 21           | 66 (30 winners) | −\$322.0       | \$104.5         | 72.4         |
| 45s   | 5.0pt       | 33           | 54 (22 winners) | −\$118.0       | \$37.5          | 92.8         |
| 45s   | 8.0pt       | 46           | 41 (15 winners) | −\$475.5       | −\$55.0         | 90.9         |
| 60s   | 2.0pt       | 14           | 73 (33 winners) | −\$228.0       | −\$21.5         | 81.5         |
| 60s   | 3.0pt       | 18           | 69 (31 winners) | −\$308.5       | −\$109.0        | 80.5         |
| 60s   | 5.0pt       | 28           | 59 (24 winners) | −\$157.5       | −\$86.0         | 92.7         |
| 60s   | 8.0pt       | 37           | 50 (18 winners) | −\$374.5       | −\$56.0         | 94.9         |
| 90s   | 2.0pt       | 10           | 77 (35 winners) | −\$156.0       | −\$51.5         | 92.7         |
| 90s   | 3.0pt       | 11           | 76 (35 winners) | −\$181.5       | −\$43.5         | 89.5         |
| 90s   | 5.0pt       | 23           | 64 (27 winners) | −\$236.5       | −\$92.0         | 94.7         |
| 90s   | 8.0pt       | 34           | 53 (20 winners) | −\$616.5       | −\$151.5        | 94.8         |

Baseline on the same 87 tape-backed signals: live −\$84.5. Read the fourth column. The cell that cuts the most (10s / 2pt) removes **59 signals of which 29 are winners** — it is binning half the good trades to bin half the bad ones, which is

the 15-of-17-winners rule failing in plain sight. Every cell that cuts aggressively is negative, and the one cell with a high placebo percentile (10s / 8pt, 86.1) barely cuts anything and is inside the noise.

6 · The full candidate board — every treatment on one page, with the kills

This is the table to keep. Every treatment proposed for this gate, scored the same way, with the two robustness kills applied to each: **leave-one-day-out** (drop the single best day — how much of the gain was one session?) and **strip-the-best-3**.

| TREATMENT                           | WHERE IT CAME FROM    | N   | NET \$   | Δ VS LIVE | \$/ SIGNAL | WINNER S CUT | WORST LODO Δ | STRIP-3 Δ | PLACEBO %ILE |
|-------------------------------------|-----------------------|-----|----------|-----------|------------|--------------|--------------|-----------|--------------|
| NO FILTER<br>(live as traded)       | baseline              | 104 | −\$158.0 | \$0.0     | −\$1.52    | 0            | \$0.0        | \$0.0     | —            |
| cooldown 5min after a stop          | operator test 1       | 87  | −\$94.5  | \$63.5    | −\$1.09    | 7            | −\$63.0      | \$243.5   | 53.4         |
| cooldown 15min after a stop         | operator test 1       | 74  | −\$210.5 | −\$52.5   | −\$2.84    | 17           | −\$147.0     | \$127.5   | 41.0         |
| cooldown 30min after a stop         | operator test 1       | 67  | −\$345.0 | −\$187.0  | −\$5.15    | 21           | −\$342.5     | −\$7.0    | 28.3         |
| cooldown 60min after a stop         | operator test 1       | 54  | −\$541.5 | −\$383.5  | −\$10.03   | 30           | −\$533.5     | −\$203.5  | 14.3         |
| bench 60min after 2 stops           | operator test 2       | 74  | −\$306.5 | −\$148.5  | −\$4.14    | 18           | −\$307.5     | −\$148.5  | 31.9         |
| bench rest-of-session after 2 stops | operator test 2       | 42  | −\$210.5 | −\$52.5   | −\$5.01    | 32           | −\$286.0     | −\$52.5   | 37.0         |
| bench rest-of-session after 3 stops | operator test 2       | 75  | −\$78.0  | \$80.0    | −\$1.04    | 14           | −\$24.0      | \$80.0    | 52.9         |
| net_min 500 (from 400)              | operator test 4       | 62  | \$362.5  | \$520.5   | \$5.85     | 18           | \$251.5      | \$520.5   | 85.4         |
| net_min 600 (from 400)              | operator test 4       | 44  | \$315.5  | \$473.5   | \$7.17     | 28           | \$65.0       | \$878.0   | 81.6         |
| net_min 700 (from 400)              | operator test 4       | 32  | −\$3.5   | \$154.5   | −\$0.11    | 36           | −\$131.0     | \$559.0   | 53.6         |
| skip when entry ATR >= 37pt         | from the regime split | 77  | \$239.5  | \$397.5   | \$3.11     | 13           | \$218.5      | \$577.5   | 81.7         |
| skip when entry ATR >= 27pt         | from the regime split | 53  | \$100.5  | \$258.5   | \$1.90     | 25           | −\$4.5       | \$987.5   | 65.2         |
| skip when entry ATR < 14pt          | from the regime split | 78  | −\$190.0 | −\$32.0   | −\$2.44    | 12           | −\$170.5     | −\$32.0   | 43.7         |

| TREATMENT                  | WHERE IT CAME FROM         | N   | NET \$   | Δ VS LIVE | \$/ SIGNAL | WINNER S CUT | WORST LODO Δ | STRIP-3 Δ | PLACEBO %ILE |
|----------------------------|----------------------------|-----|----------|-----------|------------|--------------|--------------|-----------|--------------|
| US session only (>=13:00Z) | from the time-of-day split | 101 | −\$382.5 | −\$224.5  | −\$3.79    | 3            | −\$224.5     | −\$224.5  | 6.7          |

**Not one row clears 95.** The best is net\_min 500 at the 85th percentile — and \$4 has already shown that the whole net\_min family is a calendar artefact. The bottom row is worth a line on its own: restricting the gate to the US session (≥13:00Z) makes it materially WORSE (−\$382.50, 6.7th percentile), which means whatever this gate is doing right, it is doing some of it overnight.

## 7 · ★★ Root cause, attempt two: it was never the entry. It is the EXIT.

Every entry treatment failed. That is not a dead end — it is a result, and it points straight at the other half of the trade. If the signals themselves were worthless, then how they are exited should not matter much; you would lose money slowly whatever you did. So the study swept the exit on the same 87 tape-backed signals, holding the entries exactly where the live gate put them, walking the real tick path in order, checking the stop before the target on every tick, at \$1.50 a round trip and \$2.00 a point.

**The result is not subtle.**

| STOP     | TARGET  | N  | NET \$     | \$/ SIGNAL | WIN % | WORST LEAVE-ONE-DAY-OUT | STRIP THE BEST 3 | DAYS GREEN | WORST DAY |
|----------|---------|----|------------|------------|-------|-------------------------|------------------|------------|-----------|
| 2.5×ATR  | 3.0×ATR | 87 | \$5,242.96 | \$60.26    | 62.1% | \$3,116.11              | \$3,630.52       | 9/12       | −\$248    |
| 2.5×ATR  | 2.5×ATR | 87 | \$4,915.15 | \$56.50    | 62.1% | \$2,978.80              | \$3,513.65       | 9/12       | −\$248    |
| 2.0×ATR  | 3.0×ATR | 87 | \$4,627.16 | \$53.19    | 60.9% | \$2,427.06              | \$3,213.32       | 9/12       | −\$346    |
| 2.0×ATR  | 2.5×ATR | 87 | \$4,387.45 | \$50.43    | 60.9% | \$2,377.85              | \$3,132.85       | 9/12       | −\$346    |
| 2.5×ATR  | 2.0×ATR | 87 | \$4,184.96 | \$48.10    | 63.2% | \$2,547.95              | \$3,048.44       | 9/12       | −\$248    |
| 1.75×ATR | 3.0×ATR | 87 | \$3,872.05 | \$44.51    | 56.3% | \$2,122.32              | \$2,458.21       | 9/12       | −\$314    |
| 2.0×ATR  | 2.0×ATR | 87 | \$3,745.36 | \$43.05    | 62.1% | \$2,035.10              | \$2,678.92       | 9/12       | −\$346    |
| 1.75×ATR | 2.5×ATR | 87 | \$3,650.11 | \$41.96    | 56.3% | \$2,073.11              | \$2,395.51       | 9/12       | −\$314    |
| 1.5×ATR  | 3.0×ATR | 87 | \$3,459.98 | \$39.77    | 54.0% | \$1,626.49              | \$2,046.14       | 9/12       | −\$392    |
| 1.5×ATR  | 2.5×ATR | 87 | \$3,289.74 | \$37.81    | 54.0% | \$1,628.98              | \$2,035.14       | 8/12       | −\$392    |
| 2.5×ATR  | 1.5×ATR | 87 | \$3,145.61 | \$36.16    | 67.8% | \$2,054.07              | \$2,295.47       | 7/12       | −\$92     |
| 1.75×ATR | 2.0×ATR | 87 | \$2,986.95 | \$34.33    | 57.5% | \$1,730.36              | \$1,920.51       | 9/12       | −\$314    |

The live gate on these same 87 signals is **−\$84.50**. The top twelve cells of a 64-cell decoupled sweep, both legs in ATR units so widening the stop does not silently widen the target too. Every one survives dropping its best day and stripping its best three trades with thousands still on the table.

## 7.1 Is it one hot square? No — it is a ramp, and that is the important part

A single brilliant cell surrounded by rubble is the classic curve-fit tell, and the FIRST exit sweep produced exactly that shape — which is why a second, decoupled sweep was run. Here is the stop held at 1.5×ATR and the target walked out, so you can see the whole line rather than the winner:

| TARGET   | NET \$     | \$/SIGNAL | WIN % | TARGET / STOP / TIME-OUT | WORST LODO | STRIP-3   |
|----------|------------|-----------|-------|--------------------------|------------|-----------|
| 0.5×ATR  | \$672.21   | \$7.73    | 85.1% | 74 / 13 / 0              | \$271.1    | \$408.1   |
| 0.75×ATR | \$1,307.34 | \$15.03   | 77.0% | 67 / 20 / 0              | \$743.6    | \$906.6   |
| 1.0×ATR  | \$2,065.29 | \$23.74   | 73.6% | 64 / 23 / 0              | \$1,375.9  | \$1,528.0 |
| 1.25×ATR | \$1,743.75 | \$20.04   | 64.4% | 55 / 27 / 5              | \$1,048.5  | \$1,080.6 |
| 1.5×ATR  | \$1,802.89 | \$20.72   | 59.8% | 49 / 30 / 8              | \$938.8    | \$1,005.3 |
| 2.0×ATR  | \$2,678.28 | \$30.78   | 55.2% | 39 / 31 / 17             | \$1,337.9  | \$1,611.8 |
| 2.5×ATR  | \$3,289.74 | \$37.81   | 54.0% | 31 / 31 / 25             | \$1,629.0  | \$2,035.1 |
| 3.0×ATR  | \$3,459.98 | \$39.77   | 54.0% | 22 / 31 / 34             | \$1,626.5  | \$2,046.1 |

The line climbs almost the whole way and the win rate falls the whole way — from 85% at a half-ATR target down to 54% at three — while the money goes UP fivefold. That is the shape of a gate whose edge lives in a few big continuations and whose current exit is clipping them. Note what happens to the exit mix: at the far target, 34 of 87 trades end on the clock rather than at either price, so this is as much “stop scalping and hold” as it is “widen the stop”.

## 7.2 The control that makes this believable — random entries under the same exit

Here is the objection, and it is the right objection: a wide stop and a far target will flatter ANY entry on a trending tape — you are not measuring the gate, you are measuring the exit policy. That is testable. The study fired **60 random entries** through each exit policy and compared:

| EXIT POLICY | THE REAL ENTRIES | TARGET / STOP / TIME | RANDOM MEAN | RANDOM P95 | RANDOM BEST | REAL %ILE | ENTRY ALPHA |
|-------------|------------------|----------------------|-------------|------------|-------------|-----------|-------------|
|             | \$3,465.48       | 22 / 31 / 34         | −\$2.30     | \$2,350.83 | \$3,038.82  | 100.0     | \$3,467.78  |

| EXIT POLICY                             | THE REAL ENTRIES | TARGET / STOP / TIME | RANDOM MEAN | RANDOM P95 | RANDOM BEST | REAL %ILE | ENTRY ALPHA |
|-----------------------------------------|------------------|----------------------|-------------|------------|-------------|-----------|-------------|
| stop 1.5xATR / target 2.0R / 15min cap  |                  |                      |             |            |             |           |             |
| stop 1.5xATR / target 0.75R / 15min cap | \$1,860.92       | 59 / 25 / 3          | -\$294.31   | \$1,297.36 | \$1,524.57  | 100.0     | \$2,155.23  |
| stop 1.0xATR / target 0.75R / 15min cap | -\$186.04        | 50 / 37 / 0          | -\$420.67   | \$738.58   | \$883.76    | 65.0      | \$234.63    |
| stop 20pt / target 24pt / hold 300      | \$336.00         | 36 / 32 / 19         | -\$253.66   | \$910.50   | \$1,204.00  | 85.0      | \$589.66    |
| stop 8pt / target 12pt / hold 120       | -\$245.50        | 31 / 46 / 10         | -\$251.53   | \$154.50   | \$182.00    | 46.7      | \$6.03      |

This is the load-bearing table of the section. Under the wide-stop policies the real entries beat **all 60** random draws, and the random draws themselves average about zero — so the exit policy is NOT what is making the money. The entry is. Note the last row: under the tight 8pt/12pt scalp the real entries land at the 47th percentile, i.e. **indistinguishable from firing at random**. The gate's edge is invisible at the exit the desk actually uses.

### 7.3 Where the money is, day by day — and the concentration to be honest about

| SESSION      | NET \$ UNDER THE WIDE-STOP EXIT |
|--------------|---------------------------------|
| 2026-07-24   | \$181.53                        |
| 2026-07-27   | \$1,833.49                      |
| 2026-07-28   | \$17.87                         |
| 2026-07-29   | \$160.40                        |
| 2026-07-30   | \$351.50                        |
| 2026-07-31   | \$698.19                        |
| 2026-08-05   | -\$17.00                        |
| 2026-08-07   | -\$163.26                       |
| 2026-08-10   | -\$391.88                       |
| 2026-08-11   | \$389.88                        |
| 2026-08-12   | \$293.42                        |
| 2026-08-14   | \$105.84                        |
| <b>total</b> | <b>\$3,459.98</b>               |

Twelve sessions, 9 of them green. **One day carries more than half of it:** 27 July alone is \$1,833.49 of the \$3,459.98. Drop it and the policy still makes \$1,626.49; strip the best three trades and it still makes \$2,046.14. Both are real, both are

positive, and both are a long way below the headline — which is the honest way to hold this result.

And on THIS week’s four testable days the same policy makes \$397.26 (08-10 –\$392, 08-11 \$390, 08-12 \$293, 08-14 \$106) against the live gate’s –\$218.50 on 71 lots. So this is not purely a July result being projected onto August — it is positive on the week the report covers, on the days the report can actually see.

## 8 · The verdict pills

WHY THIS IS SHADOW ON MONDAY AND NOT A DEPLOY

Three reasons, and I would rather write them down than have them found later. **(1) The sweep never turns over.** The best cell is the widest stop and the furthest target in the grid —  $2.5\times\text{ATR}$  and  $3\times\text{ATR}$  are both at the boundary, so the true optimum is somewhere outside what was tested and the shipped number would be a guess. **(2) The repricer fills at the exact stop and the exact target with no slippage,** and a  $2.5\times\text{ATR}$  stop on this instrument is a wide, thin price to be filled at in the moves where it gets hit. **(3) Part 2 graded the live fade-scalp trial on this same gate and it lost –\$480 over 80 legs,** with the damage concentrated on big-trend rungs. That is a different exit family, but it is a live warning that this gate’s exit behaves differently in the wild than in a repricer. A shadow slate costs nothing and settles all three.

## 9 · DISPOSITION — every lead touched in this rehabilitation

| LEAD                                         | VERDICT                         | THE EVIDENCE, OR THE CONDITION THAT REVIVES IT                                                                                                                                                                                                              |
|----------------------------------------------|---------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| exhaustion_short entry (the signal itself)   | REHABILITATED — SHADOW          | Beats 60 of 60 random-entry controls under a wide-stop exit (entry alpha \$3,467.78, percentile 100). The signal is real; it has never been the problem.                                                                                                    |
| exhaustion_short wide-stop / far-target exit | SHADOW — the week’s deliverable | +\$3,459.98 vs –\$84.50 live on the same 87 signals; worst leave-one-day-out +\$1,626.49, strip-best-3 +\$2,046.14, 9 of 12 days green, +\$397.26 on this week’s four days. Shadow it because the grid never turns over and the repricer takes no slippage. |
| Cooldown after a stop (13 cells)             | REFUTED                         | Named test: 4,000-draw random-removal placebo. Every cell loses money; best cell sits at the 54th percentile of noise. No reformulation                                                                                                                     |

| LEAD                                    | VERDICT                   | THE EVIDENCE, OR THE CONDITION THAT REVIVES IT                                                                                                                                                                                                                          |
|-----------------------------------------|---------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                                         |                           | survives — the mechanism assumes the last stop predicts the next signal, and it does not.                                                                                                                                                                               |
| Bench after a run of stops (15 cells)   | REFUTED                   | Named test: same placebo. Best cell 54.2nd percentile and still negative; the rule cuts 18 winners to avoid 12 losers. Corroborated live in Part 1 (−\$172 across the week).                                                                                            |
| Raising net_min above the live 400      | REFUTED                   | Named test: era-confound check plus correlation. All nine higher floors are worse than the deployed floor, $r = 0.059$ between flow and outcome, quintiles non-monotone, and 12 of the 17 signals the ‘winning’ floor cuts come from a single era. It is a date filter. |
| The absorption veto on exhaustion_short | REFUTED                   | Named test: the keep-the-winners rule. Its most aggressive cell removes 59 signals of which 29 are winners; every aggressive cell is negative. Works on abs_veto, does not transfer here.                                                                               |
| US-session-only restriction             | REFUTED                   | Makes the gate worse (−\$382.50, 6.7th percentile). Some of whatever this gate does right, it does overnight — do not clip the session.                                                                                                                                 |
| Bench exhaustion_short outright         | REJECTED as premature     | The obvious call, and the study says it is wrong: the entry carries measurable alpha and the exit is a fixable, testable thing. Bench the EXIT, not the gate.                                                                                                           |
| abs_veto (both sides) rehabilitation    | NOT RUN — measurement gap | The rehab phase timed out before reaching it. — \$58.50 over 19 signals this week is too thin to grade anyway. Carry to next week with the same battery.                                                                                                                |
| capitulation_long rehabilitation        | NOT RUN — measurement gap | Two lots, −\$48.00, 0% win this week. The phase never reached it. Movement 2 does cover its require_flip clause on run tape (SHADOW) — that is the only reading it got.                                                                                                 |

### ★ WHAT IS MISSING FROM THIS SECTION, NAMED

The rehabilitation phase timed out at 90 minutes having completed the exhaustion\_short battery and nothing else. **Three things the scope asks for are therefore absent and are not to be read as “nothing found”:** (1) the **abs\_veto** exhaustion-chase rehab the scope flagged; (2) a **capitulation\_long** rehab; and (3) the two grind-exit shadow ledgers — those are not lost, Part 2 §6 and §7 carry them and both report UNMEASURED for the honest reason that grind has not traded since 5 August.

Dossier note: 17 rehab dossier file(s) are on disk from prior cycles and 1 path(s) matched the dossier glob but are DIRECTORIES, not files — a filename containing a “/” that became a real directory. The glob in this builder skips non-files rather than dying on them.



---

## **P2 THE SHADOW DESK & PROMOTION — week of 10-14 August 2026**

---

Gaz, the headline this week is a correction to my own brief, and it is worth more than any new candidate. **The thing I was told to put forward for promotion on Monday is already live, and has been since 25 July.** abs\_veto\_55s is top of the shadow board all-time (+\$4,022 over 342 trades) and top of it this week (+\$575 over 23) — but the live desk has been running it two-sided as abs\_veto\_long and abs\_veto\_short for three weeks. So there is nothing to promote there. Writing it up as a promotion candidate would have been a wasted play.

What I found instead when I checked is the real story: **the live short side and its shadow twin have come apart by \$2,349 over the same fifteen days.** Same gate, same 55-second veto, same tape. Shadow abs\_veto\_55s SHORT made **+\$1,828.50**; live abs\_veto\_short lost **-\$520.50**. The long side tracks fine (+\$853.50 shadow vs +\$451 live — same sign, live about half). It is only the short side that has broken away, and I can tell you exactly where it goes: our stops cost us **-\$48.80 a leg live against -\$29.20 in shadow**, while our winners are the same size. We are not picking worse trades. We are losing more on the ones that go wrong.

### **THE ONE-PARAGRAPH VERSION**

**SURVIVED:** abs\_veto\_55s as a mechanism — 16 of 21 days green, positive in BOTH regimes (chop +\$1,746 / trend +\$2,276), positive in BOTH walk-forward halves, positive on BOTH sides (LONG +\$1,371 / SHORT +\$2,651), and +\$11.76 a trade against -\$0.07 for the un-vetoed thrust it is built on. It is already live; this is confirmation, **not an action.** **THE ACTION:** the live SHORT side's exit ladder (1.5R/2.5R) is too wide — it converts winners into stop-outs (36.5% win live vs 49.2% shadow) and its average loss is 33% bigger. **Tune the SIDE, do not relegate the direction** — the shadow says

short is the STRONGER of the two. **DIED:** the Open Rider day-classifier. Eight pre-session features, 272 candidate rules, and the best one is beaten by **40% of pure-noise searches** once you price in the hunting. Also died: the idea that ONE day-detector could serve both the Open Rider and the MGC run-catcher — the two instruments' good days agree **38% of the time against 50% by chance.** **THE HONEST GAP:** both grind-exit shadow ledgers are frozen. grind\_long has not placed a trade since 5 August, so the two-ratchet clip watch and the 2R-partial smoothness watch have **no new data at all** this week and cannot say anything.

## 1 • THE SHADOW BOARD, RANKED ON HONEST MONEY

Every number below is shadow\_real.real\_pnl — the tick-repriced figure, never the optimistic ceiling\_pnl the sims write for themselves. 63 sims, 9,955 trades, 16 July to 14 August. I am showing the top and the bottom, because the bottom is where the lesson is.

### TOP TWELVE, ALL-TIME (16 JUL - 14 AUG)

| #  | SIM                     | N   | NET \$      | WIN % | \$/TRADE | WHAT IT IS                                             |
|----|-------------------------|-----|-------------|-------|----------|--------------------------------------------------------|
| 1  | abs_veto_55s            | 342 | +\$4,022.00 | 47.4  | +\$11.76 | thrust + 55s absorption veto — already live both sides |
| 2  | abs_veto_50s            | 372 | +\$3,034.00 | 43.8  | +\$8.16  | same, 50s veto window                                  |
| 3  | thrust_short_abs_veto55 | 101 | +\$1,330.00 | 51.5  | +\$13.17 | the short-only live-mirror A/B arm                     |
| 4  | sw_absS_B_k10           | 42  | +\$1,267.50 | 59.5  | +\$30.18 | abs_veto short, wide-leg stop-width arm                |
| 5  | thrust_cont             | 394 | +\$1,202.00 | 37.8  | +\$3.05  | thrust continuation                                    |
| 6  | sw_absS_B_k20           | 41  | +\$1,030.50 | 61.0  | +\$25.13 | ditto, k=2.0                                           |
| 7  | cx_absSB_live           | 37  | +\$855.00   | 51.4  | +\$23.11 | abs_veto short clip A/B, live-exit arm                 |
| 8  | cx_absSB_clip           | 38  | +\$838.50   | 57.9  | +\$22.07 | ditto, clip arm                                        |
| 9  | sw_absS_A_k10           | 44  | +\$732.50   | 63.6  | +\$16.65 | abs_veto short, tight leg                              |
| 10 | abs_veto_60s            | 119 | +\$620.00   | 42.9  | +\$5.21  | 60s veto window                                        |
| 11 | thrust_short_raw        | 184 | +\$362.00   | 37.5  | +\$1.97  |                                                        |

| # | SIM               | N   | NET \$    | WIN % | \$/TRADE | WHAT IT IS                   |
|---|-------------------|-----|-----------|-------|----------|------------------------------|
|   |                   |     |           |       |          | the un-vetoed control for #3 |
|   | 12 thrust_aligned | 349 | +\$351.00 | 36.1  | +\$1.01  | thrust, trend-aligned only   |

**Nine of the top nine are the same family.** Every single one is a thrust with the absorption veto on it, or a stop-width / clip arm sitting on top of one. That is not twelve independent discoveries — it is one discovery measured twelve ways, and I want that said plainly before anyone counts the board as twelve winners. The board has found exactly one thing that works, and we already run it.

### BOTTOM SIX, ALL-TIME — WHERE THE MONEY ACTUALLY WENT

| SIM               | N     | NET \$       | WIN % | \$/TRADE | WHY IT BLEEDS                                             |
|-------------------|-------|--------------|-------|----------|-----------------------------------------------------------|
| capit_loose       | 655   | −\$11,291.50 | 16.2  | −\$17.24 | the no-flip capitulation control — catches falling knives |
| capit_live_mirror | 192   | −\$2,839.00  | 15.1  | −\$14.79 | the live capitulation entry, mirrored                     |
| rg_long_fast      | 299   | −\$2,280.50  | 29.1  | −\$7.63  | fading longs into a down-tape                             |
| exhaustion_rev    | 664   | −\$1,807.00  | 39.5  | −\$2.72  | the exhaustion fade, unfiltered                           |
| rg_long_fast_v    | 257   | −\$1,776.00  | 29.6  | −\$6.91  | same, with the veto variant                               |
| grind_fast        | 1,407 | −\$1,725.00  | 31.7  | −\$1.23  | the loose grind entry — huge N, small bleed               |

A note on capit\_loose that matters for anyone reading the raw table: its ceiling\_pnl is **+\$37,045** while its honest repriced P&L is **−\$11,291.50**. That is a \$48,000 gap on one sim. It is the single best argument on this desk for never quoting ceiling numbers.

### THIS WEEK ONLY (9 - 14 AUGUST), TOP TEN

| # | SIM                                     | N       | NET \$    | WIN %       | \$/TRADE            |
|---|-----------------------------------------|---------|-----------|-------------|---------------------|
| 1 | abs_veto_55s                            | 23      | +\$575.00 | 60.9        | +\$25.00            |
| 2 | cx_absSB_live                           | 13      | +\$449.00 | 61.5        | +\$34.54            |
| 3 | abs_veto_50s                            | 25      | +\$425.50 | 52.0        | +\$17.02            |
| 4 | cx_absSB_clip / sw_absS_B_k10           | 13      | +\$421.00 | 69.2        | +\$32.38            |
| 5 | sw_absS_B_k20 / thrust_short_absvet o55 | 13 / 11 | +\$365.00 | 69.2 / 72.7 | +\$28.08 / +\$33.18 |
| 6 | sw_absL_B_k10                           | 11      | +\$358.00 | 63.6        | +\$32.55            |
| 7 | abs_veto_60s                            | 8       | +\$321.50 | 75.0        | +\$40.19            |
| 8 | odr_c5_s30                              | 20      | +\$262.50 | 50.0        | +\$13.12            |
| 9 | sw_absS_A_k10                           | 13      | +\$225.00 | 69.2        | +\$17.31            |

| # | SIM                   | N   | NET \$      | WIN % | \$/TRADE |
|---|-----------------------|-----|-------------|-------|----------|
|   | — capit_loose (worst) | 254 | -\$3,289.50 | 16.1  | -\$12.95 |

The promotion candidate that emerges from the ranking is the one we already promoted. That is the honest read of the board and I am not going to dress it up as something else. Every sim above +\$200 this week is an abs\_veto derivative except odr\_c5\_s30, which is one cell of the Open Rider and which I deal with in section 4 — and which, taken as a family, actually lost money this week.

THE OVER-TRADING CAVEAT, BAKED IN

The shadow board **overstates** and always will. A sim takes every signal its rule generates with no position limit, no benching, no router, no slippage beyond the repricing, and no operator watching. The live desk takes a fraction of them. thrust\_loose fired **611 times** in four weeks; the live tournament took 58 signals in the whole of this week. So the board **nominates** — it tells you which mechanism is worth a live slot — and the **paper desk judges**. A shadow sim's dollar total is never the dollar total you would have banked, and section 3 below is a \$2,349 demonstration of exactly that.

## 2 · THE PROMOTION BATTERY — abs\_veto\_55s, run in full

I ran scripts/abs\_veto\_robustness.py as instructed and I am presenting all five legs plus the two-sided split the script does not do. The verdict up front: **the mechanism survives everything I threw at it**. But read section 3 before you conclude that the live version is fine, because it is not.

### LEG 1 — HEADLINE

| SIM          | N   | NET \$   | WIN % | AVG WIN | AVG LOSS | WORST | \$/TRADE | MAX DD |
|--------------|-----|----------|-------|---------|----------|-------|----------|--------|
| abs_veto_55s | 342 | +\$4,022 | 47    | +\$67   | -\$38    | -\$92 | +\$11.76 | -\$441 |

A 47% win rate with wins nearly twice the size of losses. Per your standing rule I am judging this on expectancy, not hit rate — and \$11.76 a trade with a worst-

ever trade of −\$92 is a tight, well-behaved distribution. The max drawdown of −\$441 against a +\$4,022 total is the number I like most here.

LEG 2 — PER-DAY SPREAD (IS IT ONE LUCKY DAY?)

**18 of 25 days green.** Best day +\$791, worst day −\$198. The seven red days total −\$826.50 against +\$4,848.50 from the eighteen green ones. There is no single day whose removal changes the verdict: take out the best day and it still makes **+\$3,231**. The worst day in four weeks costs less than 5% of the total — that is a very flat distribution for a gate with a 47% hit rate.

| DAY   | N  | NET \$    | DAY   | N  | NET \$    | DAY   | N   | NET \$    |
|-------|----|-----------|-------|----|-----------|-------|-----|-----------|
| 07-16 | 2  | −\$81.50  | 07-26 | 3  | −\$106.50 | 08-05 | 12  | +\$240.00 |
| 07-17 | 19 | +\$580.50 | 07-27 | 15 | +\$41.00  | 08-06 | 26  | −\$198.00 |
| 07-19 | 4  | −\$143.50 | 07-28 | 33 | +\$791.00 | 08-07 | 13  | +\$307.00 |
| 07-20 | 20 | +\$272.00 | 07-29 | 36 | +\$407.50 | 08-09 | 2   | +\$20.00  |
| 07-21 | 8  | −\$10.00  | 07-30 | 22 | +\$83.50  | 08-10 | 7   | +\$232.00 |
| 07-22 | 14 | −\$118.50 | 07-31 | 23 | +\$111.50 | 08-11 | 6   | +\$40.00  |
| 07-23 | 23 | +\$676.00 | 08-02 | 3  | −\$168.50 | 08-13 | 4   | +\$215.00 |
| 07-24 | 16 | +\$165.00 | 08-03 | 17 | +\$94.00  | 08-14 | 4   | +\$68.00  |
| —     | —  | —         | 08-04 | 10 | +\$504.50 | Total | 342 | +\$4,022  |

⚠ **These are my own figures, not the battery script's, and they disagree.** `abs_veto_robustness.py` reports "16/21 days green, top day +\$761, worst −\$148". Its day totals sum to the right grand total, but its buckets do not reconcile trade-for-trade with `shadow.db` — it omits the four Sunday-evening sessions (07-19, 07-26, 08-02, 08-09) entirely and its remaining days do not equal mine even after folding those in (it has 07-28 at 28 trades where the table shows 33). The table above is computed straight from `shadow_trades` joined to `shadow_real` and reconciles exactly to 342 trades and +\$4,022. **Use these; the script's per-day leg needs a fix.** The verdict is unaffected — both versions say the edge is spread across many days — but the exact day figures in any prior report that quoted that leg are suspect.

LEG 3 — REGIME SPLIT (THIS IS THE ONE THAT MATTERS)

| REGIME AT ENTRY   | N   | NET \$   | WIN % | \$/TRADE | MAX DD |
|-------------------|-----|----------|-------|----------|--------|
| Chop (ER < 0.18)  | 163 | +\$1,746 | 47    | +\$10.70 | −\$343 |
| Trend (ER ≥ 0.18) | 179 | +\$2,276 | 48    | +\$12.70 | −\$359 |

**This is the strongest single result in the section.** Almost every edge this desk has found turns out to be a regime accident — it works in trend and dies in chop, or the reverse, and then a week of the other regime arrives and wipes it out. This one is flat across both: +\$10.70 a trade in chop, +\$12.70 in trend, near-identical

win rates and near-identical drawdowns. That is what a structural edge looks like as opposed to a lucky tape, and it is why the veto deserved its live slot.

## LEG 4 — WALK-FORWARD HALVES

| HALF              | N   | NET \$   | WIN % | \$/TRADE |  |
|-------------------|-----|----------|-------|----------|--|
| First 171 trades  | 171 | +\$2,416 | 49    | +\$14.10 |  |
| Second 171 trades | 171 | +\$1,606 | 46    | +\$9.40  |  |

Both halves positive, which is the test passing. But be honest about the shape: the second half earns **a third less per trade** than the first. That is either normal decay of a small edge, or the start of something going stale. One more week tells us which; I would not act on it yet, but I would watch it.

## LEG 5 — HEAD-TO-HEAD AGAINST THE UN-VETOED THRUST

| SIM                   | N   | NET \$   | WIN % | \$/TRADE |  | MAX DD   |
|-----------------------|-----|----------|-------|----------|--|----------|
| abs_veto_55s (vetoed) | 342 | +\$4,022 | 47    | +\$11.76 |  | -\$441   |
| thrust_loose (raw)    | 611 | -\$42.50 | 35    | -\$0.07  |  | -\$1,030 |

The script prints the difference as “+\$4,064 incremental edge”. **Do not quote that number** — it subtracts a 342-trade total from a 611-trade total, so most of it is just a different trade count. The honest comparison is per trade: **+\$11.76 against -\$0.07**, and the drawdown falling from -\$1,030 to -\$441. The veto drops roughly 44% of raw thrusts and what it drops is, on the whole, the losing half. That is the mechanism working exactly as advertised.

## BOTH SIDES — THE SPLIT THE BATTERY SCRIPT DOES NOT PRINT

| SIDE               | N   | NET \$      | WIN % | \$/TRADE |  | SAME, UN-VETOED (THRUST_LOOSE)    |
|--------------------|-----|-------------|-------|----------|--|-----------------------------------|
| abs_veto_55s LONG  | 171 | +\$1,371.00 | 45.0  | +\$8.02  |  | 293 tr, -\$862.00, 32.4%, -\$2.94 |
| abs_veto_55s SHORT | 171 | +\$2,651.00 | 49.7  | +\$15.50 |  | 318 tr, +\$819.50, 38.1%, +\$2.58 |

**Both sides earn, and the short side earns nearly twice as much as the long.** This week the same pattern holds on a small sample: LONG 11 trades + \$214.50 (54.5%), SHORT 12 trades +\$360.50 (66.7%). Hold that thought, because the live desk's numbers are the exact opposite way round.

**Caveats, stated plainly.** 1-lot sims with no position limit. A single summer regime — four weeks, one tape, no autumn, no event day. The exit mix is concentrated: 160 TARGET exits carry +\$8,979 and 177 STOPS give back −\$4,567, so the whole edge is “the target gets hit slightly more often than the stop”, which is a thin thing to rest on. And five ABSORPTION\_CUT exits lost −\$390 between them (0% win) — a small, ugly tail worth watching.

### 3 • ★ THE FINDING — live and shadow have come apart on the short side

Because the promotion was already done, I spent the time on a question nobody had asked: **is the live version actually behaving like the shadow one it was promoted from?** The long side does. The short side does not, and the gap is large enough that it should change what we do on Saturday.

Window is 26 July onward — the day the live gate started filling.

| BOOK                          | N   | NET \$      | WIN % | \$/TRADE OR LEG |
|-------------------------------|-----|-------------|-------|-----------------|
| SHADOW abs_veto_55s<br>LONG   | 118 | +\$853.50   | 44.9  | +\$7.20         |
| SHADOW abs_veto_55s<br>SHORT  | 118 | +\$1,828.50 | 49.2  | +\$15.50        |
| LIVE abs_veto_Long (legs)     | 102 | +\$451.00   | 44.1  | +\$4.40         |
| LIVE abs_veto_short<br>(legs) | 63  | −\$520.50   | 36.5  | −\$8.30         |

The long side is fine — live earns about half what shadow does, which is exactly the over-trading haircut you would expect. **The short side is upside down.** Shadow says short is the better of the two sides; live says it is the only losing one.

### WHERE THE MONEY GOES — EXIT-REASON MIX, SAME WINDOW

| EXIT                 | LIVE LEGS | LIVE \$     | LIVE \$/LEG | SHADOW N | SHADOW \$   | SHADOW \$/<br>TR |
|----------------------|-----------|-------------|-------------|----------|-------------|------------------|
| STOP                 | 35        | −\$1,707.00 | −\$48.80    | 57       | −\$1,663.00 | −\$29.20         |
| TARGET               | 11        | +\$711.50   | +\$64.70    | 59       | +\$3,652.50 | +\$61.90         |
| CHANDELIER           | 9         | +\$759.00   | +\$84.30    | —        | —           | —                |
| MAX_HOLD             | 2         | −\$255.50   | −\$127.80   | —        | —           | —                |
| STOP_UNFILLED        | 2         | −\$106.00   | −\$53.00    | —        | —           | —                |
| GIVEBACK /<br>MANUAL | 4         | +\$77.50    | +\$19.40    | —        | —           | —                |

**Our winners are the same size as shadow's. Our losers are a third bigger.** Live TARGET pays +\$64.70 against shadow's +\$61.90 — identical for practical

purposes. Live STOP costs –\$48.80 against shadow's –\$29.20. And the win rate is 36.5% live against 49.2% shadow, on the same signals. Both facts point the same way and it is not the entry:

THE DIAGNOSIS

The live short side runs a **wider profit ladder than the long side** — from data/exit\_overrides.json, abs\_veto\_short is **A 1.5R / B 2.5R** while abs\_veto\_long is **A 1.0R / B 1.5R**. A wider target does not widen the stop, but it does change which trades reach it: positions that would have banked at 1.0R are still open at 1.5R, and a chunk of them travel all the way back to the stop instead. That converts winners into full-size losers, which is precisely the signature in the table — same winner size, more of them turning into losers, bigger average loss. **The short side's entry is the best thing on the shadow board. Its exit ladder is the widest on the live desk. Those two facts do not belong together.**

NORMALISING THE SOFTWARE FAILURES FIRST, PER THE STANDING REHAB RULE

Before blaming strategy, strip the malfunctions. Two STOP\_UNFILLED legs (–\$106) and two MAX\_HOLD legs (–\$255.50) are handling faults, not the gate's opinion about the market:

| LIVE ABS_VETO_SHORT             | NET \$    |
|---------------------------------|-----------|
| As booked                       | –\$520.50 |
| Add back STOP_UNFILLED (2 legs) | +\$106.00 |
| Add back MAX_HOLD (2 legs)      | +\$255.50 |
| Strategy-only result            | –\$159.00 |

So the short side is a **–\$159 strategy problem wearing a –\$520 coat**. Seventy percent of the headline loss is four legs of plumbing. That does not make it fine — –\$159 against a shadow twin's +\$1,828.50 is still a wide gap — but it changes the size of the fix we should be looking for.

**One hypothesis I tested and rejected.** The abs\_veto\_short ATR ≥ 16 floor was deleted on 1 August, and the live short's bleed does start right after: +\$13.80 a leg before (26 legs, 4 days) against –\$23.80 a leg after (37 legs, 9 days). Tempting. But the shadow counterfactual does not support it — since 1 August the entries that floor would have blocked (ATR < 16) are **positive in shadow: +\$233.50 over 38 trades**. The deleted floor



let in a weaker population, not a losing one. So the floor is a dilution story at most, and the exit ladder remains the better explanation. Restoring the floor is **not** my recommendation.

## 4 • RELEGATION — tune the SIDE, never relegate a DIRECTION

The standing principle says a gate runs two-sided with independently-tuned per-side thresholds, and that you never relegate a direction because it had a bad week — you tune the side. This week gives that rule its cleanest demonstration yet, from fresh numbers.

| LIVE GATE         | SIDE  | LEGS | NET \$ (ALL-TIME) | WIN % | \$/LEG   | READ                                          |
|-------------------|-------|------|-------------------|-------|----------|-----------------------------------------------|
| abs_veto_long     | LONG  | 102  | +\$451.00         | 44.1  | +\$4.40  | working — hold                                |
| abs_veto_short    | SHORT | 63   | −\$520.50         | 36.5  | −\$8.30  | tune the exit, do not drop the side           |
| exhaustion_short  | SHORT | 146  | −\$158.00         | 47.9  | −\$1.10  | near flat — the exit trial is the issue (\$7) |
| rgv_long          | LONG  | 60   | −\$691.00         | 46.7  | −\$11.50 | router-only, per the standing call            |
| rgv_short         | SHORT | 43   | −\$416.00         | 44.2  | −\$9.70  | regime-dependent, not defective               |
| grind_long        | LONG  | 98   | −\$454.00         | 31.6  | −\$4.60  | silent since 5 Aug — see §6                   |
| capitulation_long | LONG  | 32   | −\$85.50          | 34.4  | −\$2.70  | barely trades now                             |
| nipc_short        | SHORT | 16   | −\$337.00         | 6.2   | −\$21.10 | 1 winner in 16 — thin, but ugly               |

**If we judged by direction we would relegate the short side of abs\_veto this weekend, and it would be the worst decision we made all month.** The shadow board says that same short side, on the same signals, is the best-performing thing we own (+\$2,651 all-time, +\$15.50 a trade, 49.7% win). The direction is not the problem. The side's exit ladder is. That is the two-sided-gate rule doing real work rather than being quoted.

What I would actually do, stated as a proposal and not a deploy: **bring abs\_veto\_short's ladder in from 1.5R/2.5R towards the long side's 1.0R/1.5R**, and run it as a forward A/B against the current wide ladder rather than switching outright. The evidence for the tighter ladder is a diagnosis (winner-size identical, loser-size 33% worse, win rate 13 points worse), not a backtest of the narrower setting, and I would rather prove it forward on paired legs than assert it now.

△ A counting note that applies to this whole table: the live desk runs a two-lot A/B scale-out, so a single signal writes **two rows**. Everything above is counted in LEGS, not signals. The A and B legs are a scale-out of one decision, not an experiment with two opinions — so “63 short legs” is roughly 32 short decisions, and every per-leg figure is about half a per-signal figure.

5 • ★★ THE OPEN RIDER — is a juicy day callable BEFORE the session commits?

Your framing: “open rider best shadow again ... but i know they bleed on other days. thats where im saying we need detection some how to block open rider if conditions arent juicy.” You are right on both halves and I verified both from `real_pnl` before going anywhere near a detector.

FIRST — CONFIRMING THE SHAPE YOU DESCRIBED

| DAY        | N  | NET \$      | WIN % | \$/TRADE  | CHARACTER           |
|------------|----|-------------|-------|-----------|---------------------|
| 2026-08-10 | 16 | −\$1,076.50 | 12.5  | −\$67.28  | bled                |
| 2026-08-11 | 18 | +\$59.00    | 44.4  | +\$3.28   | scratch-green       |
| 2026-08-12 | 19 | −\$1,147.50 | 26.3  | −\$60.39  | bled                |
| 2026-08-13 | 15 | +\$1,688.50 | 73.3  | +\$112.57 | owned the board     |
| 2026-08-14 | 19 | −\$2.50     | 42.1  | −\$0.13   | flat                |
| Week       | 87 | −\$479.00   | 39.1  | −\$5.51   | ungated cells 54-57 |

Two green days paid **+\$1,747.50** across 33 trades. Three red days lost **−\$2,226.50** across 54. Add the 19 gated trades and the family is **−\$481.50 over 106**, which is the **−\$482** in your note. **The average hides the strategy completely** — a **−\$5.51-a-trade** line conceals a strategy that pays **+\$112** a trade on its day and loses **\$60-67** a trade on the wrong one.

**The oracle.** Trading only the green days is **+\$1,747.50 a week**, several times over your “few hundred a week” bar. It cheats outright — the labels only exist after the close — so it is a ceiling, not a plan. But it does settle the important question: **the prize is real and the whole problem is day selection**, not entry timing and not exit tuning.

## ESCAPING N=5 — THE REPLAY

Five days cannot test a day-classifier. But the Open Rider is a **pure clock rider**: no signal, no gate. Every N minutes between 13:00 and 14:45 UTC it goes with the sign of the last 15 minutes, stop at  $2.0\times/3.0\times\text{ATR}$ , target coupled at 2R, 45-minute cap. That is fully deterministic, so I replayed it over **every day of tape we hold** — 41 trading days, 19 June to 14 August, on 5-second bars from the V5 archive + lake + hot union (scripts/odr\_day\_classifier.py). That turns 5 day-labels into 41.

Before using it I checked it against the five real forward days:

| DAY   | REAL N | REAL \$     | REPLAY N | REPLAY \$   | SIGN AGREES?           |
|-------|--------|-------------|----------|-------------|------------------------|
| 08-10 | 16     | −\$1,076.50 | 22       | −\$1,180.50 | yes                    |
| 08-11 | 18     | +\$59.00    | 18       | −\$285.10   | no (both near zero)    |
| 08-12 | 19     | −\$1,147.50 | 21       | −\$916.20   | yes                    |
| 08-13 | 15     | +\$1,688.50 | 13       | +\$1,420.40 | yes                    |
| 08-14 | 19     | −\$2.50     | 18       | +\$301.10   | no (both near zero)    |
| Total | 87     | −\$479      | 92       | −\$660      | −\$5.51 vs −\$7.18 /tr |

**Honest read of the validation: good on the days that matter, poor on the days that don't.** The three days with real money in them all match in sign and roughly in size. The two mismatches are both days where the true answer is “about zero” and the replay says “about zero the other way”. Overall level is close and slightly pessimistic. Good enough to hunt a day-classifier on; not good enough to quote a replay dollar total as if it were money. Across all 41 days the replay makes +\$19,247 on 753 trades with 24 green days — a number I am deliberately not putting forward as an edge, because it is in-sample on an optimistic 5-second-bar fill model.

## THE HUNT — EIGHT CAUSAL FEATURES, AND THE TRAP

Everything below is measured **strictly before 13:00 UTC**, the instant the rider's first entry fires. That is the causality line and I held to it. I searched every feature against a threshold in both directions — 272 candidate day-rules in total.

| PRE-13:00 FEATURE      | RANK CORR. WITH DAY \$ | BEST RULE FOUND     | DAYS KEPT | RULE \$   | NAIVE PLACEBO |
|------------------------|------------------------|---------------------|-----------|-----------|---------------|
| Gap from prior close   | +0.057                 | keep $\geq$ 4.475pt | 36        | +\$23,279 | 1.4%          |
| ATR-14 at 13:00        | +0.140                 | keep $\geq$ 9.48pt  | 33        | +\$22,104 | 3.5%          |
| Run-in efficiency (2h) | −0.249                 | keep $\leq$ 0.179   | 37        | +\$21,071 | 8.3%          |
| Overnight roundtrip    | +0.232                 | keep $\geq$ 11.24   | 29        | +\$20,047 | 5.3%          |

| PRE-13:00<br>FEATURE | RANK CORR.<br>WITH DAY \$ | BEST RULE<br>FOUND  | DAYS KEPT | RULE \$   | NAIVE<br>PLACEBO |
|----------------------|---------------------------|---------------------|-----------|-----------|------------------|
| Overnight efficiency | -0.171                    | keep $\leq$ 0.047   | 25        | +\$18,872 | 5.1%             |
| Overnight range      | -0.084                    | keep $\geq$ 224.8pt | 35        | +\$19,531 | 16.5%            |
| Overnight net move   | -0.143                    | keep $\leq$ 277.8pt | 33        | +\$18,826 | 17.9%            |
| Run-in range (2h)    | +0.056                    | keep $\leq$ 207pt   | 35        | +\$21,448 | 5.3%             |

Read the “naive placebo” column and it looks like we have something: the gap rule is beaten by only 1.4% of random day-removals of the same size, ATR by 3.5%. **That column is a trap and I nearly walked into it.** It tests the winning rule as if I had picked it in advance. I did not — I searched 272 of them and kept the best. With 272 tries, a 1-in-70 result is roughly what noise produces on its own.

### THE SEARCH-AWARE PLACEBO — AND IT KILLS THE CLASSIFIER

The correct control is to run **the entire search** against shuffled day labels and ask how often the hunt finds something this good on pure noise. I did that 2,000 times (scripts/odr\_day\_placebo2.py):

| BEST RULE FOUND                                  |           |
|--------------------------------------------------|-----------|
| The real search (gap $\geq$ 4.475)               | +\$23,279 |
| Same search on shuffled labels — median          | +\$22,819 |
| Same search on shuffled labels — 90th percentile | +\$25,137 |
| Same search on shuffled labels — max of 2,000    | +\$28,887 |

**The real search is beaten by 40.0% of pure-noise searches.** It sits almost exactly on the median of what random data produces. There is no day-classifier here. Not a weak one — none.

### WHY THIS WAS NEVER PROVABLE THIS WEEK ANYWAY — THE POWER LIMIT

The day-to-day standard deviation of Open Rider P&L across the 41 replayed days is **\$1,384**. Against that noise:

| TO DETECT AN EDGE OF... | DAYS NEEDED (~80% POWER) | WEEKS |
|-------------------------|--------------------------|-------|
| \$300 a day             | 167                      | 33    |
| \$500 a day             | 60                       | 12    |
| \$800 a day             | 23                       | 5     |

**The forward book has five days. That is 8% of what a \$500-a-day effect needs.** So the deliverable here was always going to be the method plus what would prove it — not a shipped rule — and I want to be clear that the negative result above is a genuine negative on 41 days, not a shrug about five.

## WHAT WOULD PROVE IT

1. **Pre-register one feature, not eight.** The search is what killed this. Pick the single most mechanically plausible one — my vote is ATR-14 at 13:00, the only feature with both a positive rank correlation and a physical story (a dead tape cannot pay a momentum rider) — fix the threshold at 9.5pt now, and let it run forward untouched. One pre-registered rule needs no multiple-comparison discount.
2. **Give it 60 sessions.** Twelve weeks at the current trade rate. That is the honest cost of an answer at the effect size that would matter.
3. **Judge it on the days it removes, not the total.** Log, per skipped day, what the rider would have made. A day-filter that removes eight days and misses two of the best is a bad filter even if the total goes up.

## ★★ COULD ONE DAY-CLASSIFIER SERVE BOTH THE OPEN RIDER AND THE MGC RUN-CATCHER?

This was the most valuable question in the brief, because the same shape turned up twice this week on two instruments: the Open Rider lives on a minority of days, and docs/MGC\_LEADS\_2026-08-14.md found gold's run-catcher does too (run-days pay, the rest bleed). If the same days are juicy on both, one detector is worth more than either gate. I tested it on the 16 days where we hold both MNQ and MGC tape (scripts/odr\_mgc\_shared\_day.py), pricing gold correctly at **\$10.00/point**.

| TEST                                        | RESULT    | MEANING                                                      |
|---------------------------------------------|-----------|--------------------------------------------------------------|
| Green odr day = big MGC run-day?            | 38% agree | chance agreement is 50% — <b>12 points WORSE than a coin</b> |
| Rank corr. (odr day \$, MGC run \$)         | −0.129    | nil, and faintly negative                                    |
| Rank corr. (odr day \$, MGC day range)      | −0.129    | nil                                                          |
| Rank corr. (odr day \$, MGC day efficiency) | −0.341    | gold's cleanest days are mildly BAD for the rider            |
| MNQ ATR at 13:00 → MGC run size             | +0.213    | the only faint positive; nowhere near usable                 |

**The answer is no.** A good day for the MNQ Open Rider is not a good day for the MGC run-catcher; if anything they lean slightly opposite. Two strategies with the same shape of problem do not have the same solution. Gold needs its own detector built on gold's tape, exactly as the MGC notes already insist about

thresholds. This is 16 days, so it is a PARK and not a refutation — but it is 16 days pointing consistently in one direction, and it should stop us building a shared detector on the assumption the days coincide.

## THE DRIFT GATE (SIMS 60-63) — LIVE ONE DAY, AND IT HAS NOT FIRED

| CELL                      | UNGATED \$     | UNGATED N | GATED \$       | GATED N   |
|---------------------------|----------------|-----------|----------------|-----------|
| odr_c10_s20               | −\$141.50      | 5         | −\$141.50      | 5         |
| odr_c10_s30               | +\$226.00      | 4         | +\$226.00      | 4         |
| odr_c5_s20                | −\$8.50        | 5         | −\$8.50        | 5         |
| odr_c5_s30                | −\$78.50       | 5         | −\$78.50       | 5         |
| <b>Total (08-14 only)</b> | <b>−\$2.50</b> | <b>19</b> | <b>−\$2.50</b> | <b>19</b> |

**Identical, cell by cell, trade for trade. The gate has rejected exactly zero entries.** The gated arms have existed for one day (14 August) and on that day the drift detector never disagreed with the rider's own direction, so nothing was blocked. That is not evidence the gate works and it is not evidence it doesn't — it is one day of no information.

⚠ **Never compare gated against ungated on all-time totals.** The gated arms cover one day and the ungated five; an all-time comparison would measure the calendar and report it as a gate effect. The only legitimate comparison is the overlapping window, which is the table above. **What would prove it:** a day where the drift detector confirms a direction and the rider's 15-minute momentum points the other way. Until that day happens the gate is untested by construction.

## THE SECONDARY AXIS — STOP WIDTH, AND WHY I AM NOT LEADING WITH IT

| STOP           | N  | NET \$    | WIN % | \$/TRADE |
|----------------|----|-----------|-------|----------|
| <b>2.0×ATR</b> | 58 | −\$874.50 | 31.0  | −\$15.08 |
| <b>3.0×ATR</b> | 48 | +\$393.00 | 50.0  | +\$8.19  |

Real, and it points the same way on the replay (41 days: 3.0× makes +\$29.59 a trade against +\$22.61, at a 50.3% win rate against 41.8%). The wider stop also beats the tighter one on **4 of the 5 days**. But look at the sizes before acting on it:

| EFFECT                            | SPREAD  |
|-----------------------------------|---------|
| DAY effect (best day – worst day) | \$2,836 |
| STOP-WIDTH effect (3.0× – 2.0×)   | \$1,268 |

**The day effect is 2.2× the stop-width effect.** Tuning the stop while the day problem is unsolved is polishing the smaller lever, and I would rather say that than lead with the tuning because the tuning is the part that has a tidy number attached.

△ Every Open Rider arm is flagged **n < 40**, below this desk's threshold for a forward A/B to claim anything. **No promotion follows from this week**, on any cell, in either direction.

## 6 · GRIND EXIT SHADOW — the two-ratchet runner-clip watch

I ran scripts/two\_ratchet\_shadow\_watch.py without --ping as instructed. The result is dominated by one fact that has nothing to do with the two-ratchet:

△ **THE LEDGER IS FROZEN — GRIND\_LONG HAS NOT TRADED SINCE 5 AUGUST**

The grind family's last live fill was **2026-08-05**. Not a single grind trade this week, or last. So both grind-exit shadow ledgers contain **zero new data** since the previous report, and neither can say anything new. I checked this was not a tooling artefact: the 5-second tape is complete through 14 August (the loader unions V5 archive + lake + capture, and every weekday 10-14 August is present), so the silence is real — the gate is not firing, the watch is not broken.

| TWO-RATCHET WATCH (26 JUL - 14 AUG) | VALUE     |
|-------------------------------------|-----------|
| Trades repriced                     | 26        |
| Runners (rMFE ≥ 4R)                 | 3         |
| RUNNER-CLIPS                        | 0         |
| Live pure-6.0 net                   | +\$29.50  |
| Shadow two-ratchet net              | -\$265.00 |
| Shadow – live                       | -\$294.00 |

On the verdict logic in the scope, zero clips over a genuine trend week would be evidence for reconsidering the two-ratchet. **I cannot give you that, because there was no trend week — there was no week at all.** Three runners appeared, all before 5 August, and on all three the two-ratchet actually did better than the live exit (+\$144 vs +\$23, +\$348 vs +\$244, +\$62 vs +\$53). Yet it still finished \$294 behind overall. Where it lost is instructive: two trades on 30 July

where the live book records +\$159 and +\$156 while the reconstruction says – \$31 and –\$58.

**⚠ INSTRUMENT WARNING — DO NOT ACT ON THE –\$294**

The watch prints its own cross-check: **reconstructed-pure minus actual-live = –\$406**, where zero is correct. **The instrument's error is larger than the signal it is measuring.** Five of 26 trades carry the whole discrepancy, and two of them exited MANUAL\_CLAIM — you closed them by hand, which no tape reconstruction can ever reproduce. A third is worse: the 30 July 11:11 trade books **+\$159 live on a trade whose maximum favourable excursion was 0.81R**. At ATR 14.1 that is about \$23 a lot at best. That P&L is not physically reachable on the tape, which points at the known fill mis-attribution mechanism rather than at a real exit difference. **Verdict: the two-ratchet comparison is UNMEASURED this window, not lost.**

7 · GRIND EXIT SHADOW — the 2R-partial smoothness choice

scripts/partial\_shadow\_watch.py, same frozen window, same caveat — nothing after 5 August. It covers 40 legs across the 5 days grind actually traded.

| METRIC          | BASELINE (DEPLOYED 2-LOT CHANDELIER) | 2R-PARTIAL (SHADOW) | DIFFERENCE       |
|-----------------|--------------------------------------|---------------------|------------------|
| Net \$          | –\$1,382.40                          | –\$1,074.90         | +\$307.50        |
| Daily std \$    | 42.3                                 | 83.3                | worse (+97%)     |
| Max drawdown \$ | –\$1,382.40                          | –\$1,074.90         | +\$307.50 better |
| Worst day \$    | –\$336.40                            | –\$286.00           | +\$50.40 better  |
| Green-day %     | 0%                                   | 0%                  | no difference    |

The choice you were being asked to make is not available on this data, and I would rather say so than manufacture a recommendation. The 2R-partial was framed as a variance lever — give up a little mean to get a smoother curve. Over these five days it does the opposite on the headline variance measure: the partial's daily standard deviation is nearly **double** the baseline's, not lower. It does show a shallower drawdown and a better worst day, and it gives up no mean



at all here (it is \$307.50 ahead, because banking Lot A early on five straight losing days is simply less bad).

That last clause is the whole problem. **Grind lost money on all five days it traded — green-day percentage is zero for both arms.** On a book with no winning days, "bank half the position early" wins automatically and tells you nothing about what happens on a runner day, which is the only day the question is actually about. The mean give-up that the original backtest measured (~\$665 over 3 weeks) is entirely a tail forfeit on runners, and there are no runners in this sample.

**My recommendation: no decision this week.** Not "adopt", not "reject" — the data cannot separate them. The 2R-partial stays in shadow. What would make it decidable: **grind\_long trading again at all**, and then a week containing at least a few  $rMFE \geq 4R$  runners.

## 8 · ★ THE REGIME-FLEX EXIT LAB

### 8A · GRADING THE LIVE EXHAUSTION\_SHORT FADE-SCALP TRIAL ON ACTUAL FILLS

The trial (Lot A at 0.5R, Lot B at 1.5R) went live on 31 July. Graded on real fills, split by the exit config that was genuinely running at the time:

| CONFIG                              | LEGS | NET \$    | WIN % | TARGET | STOP | CHAND |  |
|-------------------------------------|------|-----------|-------|--------|------|-------|--|
| PRE — single slot                   | 58   | +\$271.00 | 57    | 14     | 25   | 10    |  |
| DEFAULT — A2.5<br>+ wide chandelier | 8    | +\$51.00  | 50    | 0      | 4    | 0     |  |
| TRIAL — A0.5 / B1.5                 | 80   | -\$480.00 | 41    | 23     | 47   | 6     |  |

**The live trial is losing: −\$480 over 80 legs** (roughly 40 signals, since these are A/B legs of one decision). But a raw before/after comparison is not fair — the trial ran across different days and different regimes than the period before it. So I repriced **every exhaustion entry we hold** (717 of them: 105 live signals plus the 612-strong exhaustion\_rev shadow board) on the forward tape under both configs, so the only thing that differs is the exit:

| RUNG       | N   | 0.5R HIT | 1.5R HIT | STOP FIRST | TRIAL \$/SIG | PRIOR \$/SIG | DELTA    |
|------------|-----|----------|----------|------------|--------------|--------------|----------|
| BIG-TREND  | 28  | 79%      | 64%      | 21%        | +\$22.24     | +\$42.80     | −\$20.56 |
| MED-TREND  | 169 | 62%      | 35%      | 38%        | −\$16.73     | −\$15.15     | −\$1.58  |
| SCALP-CHOP | 519 | 68%      | 42%      | 32%        | −\$0.57      | −\$1.54      | +\$0.97  |
| ALL        | 716 | 67%      | 41%      | 33%        | −\$3.49      | −\$3.02      | −\$0.48  |

**This is the ladder argument in one table, and it says you were right.**

Overall the trial is near-neutral (−\$0.48 a signal). But it is not uniform: it wins slightly in chop (+\$0.97) and it gives up **−\$20.56 a signal on BIG-TREND**, cutting that rung's return roughly in half. Scalping at 0.5R on a day that is genuinely trending throws away the only trades that pay for the rest. A single exit config sprayed across all regimes averages the tape it should be ON with the tape it should be different on — which is exactly the failure the backtest discipline exists to prevent.

### 8B · THE PER-RUNG R GRID — AND THE WINDOW EXTENSION

I built scripts/exit\_ladder\_grid.py to sweep Lot-A × Lot-B across **8 × 7 R-combos plus two chandeliers** per gate family per rung. Two deliberate departures from the existing adaptive\_exit\_bt.py, both aimed at the BIG-TREND n problem you flagged:

- The tape.** The old lab reads capture.db, whose 5-second bars begin 15 July. This one runs on the full **45-day union** (19 June – 14 August).
- The entries.** Added the V5 archive's fut\_shadow\_sim\_trades — **4,270 MNQ entries from 29 June to 15 July**, thirteen trading days the V7 board does not contain. I verified these reprice correctly before using them: V5 entry prices match our own tape to a **median 1-2.5 points** on every day but one, and the ATR units are comparable (V5 median 13.0pt vs V7 13.4pt).

Robustness on every cell: parameter plateau (I rank by the neighbourhood mean, so a lone spike cannot win), strip-the-best-3-days, leave-one-day-out worst case, and the V5 archive as an out-of-sample leg.

| FAMILY            | RUNG      | N   | BEST POLICY         | \$/SIG  | STRIP-3 | LOO WORST | V5 LEG        | VERDICT |
|-------------------|-----------|-----|---------------------|---------|---------|-----------|---------------|---------|
| exhaustion (fade) | BIG-TREND | 27  | tight chandelier x2 | +\$22.8 | +\$178  | +\$13.1   | −\$8.0 (n22)  | SHADOW  |
| exhaustion (fade) | MED-TREND | 161 | tight chandelier x2 | +\$8.3  | +\$160  | +\$5.5    | −\$6.7 (n128) | SHADOW  |

| FAMILY              | RUNG       | N     | BEST<br>POLICY       | \$/SIG  | STRIP-3   | LOO<br>WORST | V5 LEG         | VERDICT              |
|---------------------|------------|-------|----------------------|---------|-----------|--------------|----------------|----------------------|
| grind (mom)         | MED-TREND  | 163   | A 0.5R / B 1.0R      | +\$9.4  | +\$731    | +\$8.0       | -\$2.9 (n53)   | SHADOW               |
| exhaustion (fade)   | SCALP-CHOP | 475   | tight chandelier x2  | +\$2.1  | -\$498    | +\$0.7       | +\$0.6 (n573)  | HOLD dies on strip-3 |
| grind (mom)         | SCALP-CHOP | 1,067 | A 0.5R / B 1.0R      | -\$3.4  | -\$4,980  | -\$4.1       | +\$9.4 (n141)  | HOLD                 |
| thrust/absveto      | BIG-TREND  | 26    | A 2.0R / B 2.5R      | -\$60.3 | -\$1,613  | -\$75.4      | —              | HOLD                 |
| thrust/absveto      | MED-TREND  | 138   | A 0.5R / B 1.0R      | -\$40.9 | -\$5,493  | -\$42.9      | +\$101.4 (n35) | HOLD                 |
| thrust/absveto      | SCALP-CHOP | 392   | A 3.0R / B 4.0R      | -\$39.4 | -\$15,558 | -\$41.8      | +\$94.2 (n123) | HOLD                 |
| rgv (fade)          | MED-TREND  | 126   | wide chandelier x2   | -\$8.3  | -\$2,316  | -\$15.7      | +\$34.0 (n22)  | HOLD                 |
| rgv (fade)          | SCALP-CHOP | 196   | A 0.5R / B 1.0R      | -\$10.0 | -\$2,450  | -\$12.4      | +\$6.0 (n155)  | HOLD                 |
| capitulation (fade) | MED-TREND  | 91    | A 0.5R / tight chand | -\$7.0  | -\$1,438  | -\$13.9      | —              | HOLD                 |
| capitulation (fade) | SCALP-CHOP | 522   | A 0.5R / B 1.0R      | -\$14.2 | -\$7,234  | -\$14.8      | —              | HOLD                 |

**The single most important line in that table is what it says about thrust.** Across the whole grid the thrust family's SCALP-CHOP cell ranges from  $-\$39.4$  to  $-\$40.6$  a signal — that is the entire spread across 58 different exit policies. **No exit rescues an un-vetoed thrust.** The fix for that family was never the exit; it was the 55-second veto, and section 2 is the proof. That is a satisfying piece of consistency between two independent studies in this report.

**And the shape of the exhaustion result is the ladder you described.** On both trend rungs the winner is a tight chandelier on both lots — ride it, trail it, do not scalp it — and the neighbourhood confirms it is a plateau rather than a spike (BIG-TREND:  $\text{tightChand} \times 2$   $+\$22.8$ ,  $A1.5/\text{tightChand}$   $+\$20.6$ ,  $A1.25/\text{tightChand}$   $+\$19.3$ ,  $A1.0/\text{tightChand}$   $+\$18.8$ ). That is the same conclusion 8a reached from a completely different direction, and it is the opposite of the 0.5R fade-scalp currently deployed.

## 8C · ★ BIG-TREND — THE HONEST ANSWER AFTER EXTENDING THE WINDOW

**I EXTENDED THE WINDOW AS INSTRUCTED. THE RUNG IS STILL NOT PROVEN.**

I want to be exact about this because the brief asked me to finally settle it. I added 13 trading days and 4,270 archive entries and ran on 45 days of tape instead of 22. It was not enough, for a reason that will not go away: **BIG-TREND is only 2.96% of all minutes on this tape** ( $ER \geq 0.50$ ; MED-

TREND is 20.2%). The rung is structurally rare, so entry populations land in it a few dozen times a month no matter how much history you load. And the two independent windows **disagree**:

| EXHAUSTION, BIG-TREND        | N  | BEST POLICY         |  | \$/SIG  |
|------------------------------|----|---------------------|--|---------|
| V7 board (16 Jul - 14 Aug)   | 27 | tight chandelier x2 |  | +\$22.8 |
| V5 archive (29 Jun - 15 Jul) | 22 | A 0.5R / B 1.0R     |  | -\$8.0  |

Both cells are about n=25. One says ride, the other says scalp. **That is not a proof and it is not a refutation — it is an unresolved rung**, and the honest thing is to say so rather than pick the half that agrees with the story. **What would settle it:** roughly 120 BIG-TREND entries, which at 3% of minutes and current fire rates is about 8-10 weeks — or a deliberate archive pull of the 6-8 highest-ER days from the wider B2 history, scored on their own.

One methodological caveat I owe you on the V5 leg: the archive's strategies are not our gates. I mapped them by mechanism (thrust-family to thrust, dip/pullback to rgv, passive-chop to exhaustion) and that mapping is a judgement call, which is exactly why the V5 column is reported beside the V7 one and never merged into it. Some V5 cells are plainly not comparable — thrust STAY-OUT shows +\$238.8 a signal at a 100% win rate on n=19, which is a tiny-n artefact of a violent day, not a finding.

8D · DELIVERABLE — THE PER-RUNG POLICY

I built the policy file with a rule that refuses to change anything that did not earn it: a cell only gets a change if it is positive on its home rung and positive after stripping its best 3 days and positive leave-one-day-out. Written to **data/exit\_overrides\_proposed.json** — a proposal. **I did not touch the live data/exit\_overrides.json**; deploying is a Saturday decision, not a side-effect of writing a report.

| GATE             | RUNG      | PROPOSED                   | CURRENTLY LIVE  | STATUS | WHY NOT STRONGER                                 |
|------------------|-----------|----------------------------|-----------------|--------|--------------------------------------------------|
| exhaustion_short | BIG-TREND | tight chandelier both lots | A 0.75R / tight | SHADOW | n=27 and the V5 leg disagrees                    |
| exhaustion_short | MED-TREND | tight chandelier both lots | A 0.75R / tight | SHADOW | survives strip-3 and LOO; V5 leg disagrees       |
| grind_long       | MED-TREND | A 0.5R / B 1.0R            | A 2.5R / wide   | SHADOW | survives both; V5 leg disagrees; gate is dormant |

| GATE            | RUNG      | PROPOSED                             | CURRENTLY LIVE | STATUS          | WHY NOT STRONGER                   |
|-----------------|-----------|--------------------------------------|----------------|-----------------|------------------------------------|
| everything else | all rungs | NO CHANGE — holds at the live config |                | <div>HOLD</div> | failed at least one robustness leg |

**Three shadow changes, zero adoptions.** That is a thinner deliverable than the brief hoped for and I am not going to inflate it. Every cell that survived the V7 robustness battery was contradicted by the V5 out-of-sample leg, and under this desk's own rules a candidate that fails its out-of-sample leg goes to shadow, not to production. The one change I would argue hardest for is `exhaustion_short` on the trend rungs, because it is the only proposal supported by two independent studies — the live trial grade in 8a ( $-\$20.56$  a signal from scalping BIG-TREND) and the grid in 8b (tight chandelier wins that rung on a plateau).

## 9 · INSTRUMENT NOTES — what I had to fix to trust these numbers

| TOOL                                            | PROBLEM                                                                                                                                                                        | WHAT I DID                                                                                                                       |
|-------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------|
| <code>scratchpad/tape5s.npz</code>              | Cached 31 July — it would have graded the exhaustion trial with <b>all of August missing</b>                                                                                   | Rebuilt full-window from the lake (45 days, 660,071 bars); old file kept as <code>.bak-20260814</code>                           |
| <code>two_ratchet_shadow_watch.py</code>        | Cross-check $-\$406$ against a $\pm\$294$ signal; one trade books $+\$159$ on a 0.81R excursion                                                                                | Reported as UNMEASURED; the $-\$294$ is not quoted as a result                                                                   |
| <code>abs_veto_robustness.py</code>             | Prints a “ $+\$4,064$ incremental edge” that subtracts a 342-trade total from a 611-trade one                                                                                  | Quoted per-trade instead ( $+\$11.76$ vs $-\$0.07$ )                                                                             |
| <code>abs_veto_robustness.py</code> per-day leg | Reports 21 days where the sim traded <b>25</b> — drops all four Sunday-evening sessions, and its remaining buckets still do not match trade-for-trade (07-28: 28 trades vs 33) | Recomputed the whole per-day table from <code>shadow.db</code> ; <b>prior reports quoting this leg's day figures are suspect</b> |
| My own first placebo                            | Tested the winning rule as if pre-registered, after searching 272                                                                                                              | Replaced with the search-aware shuffle — which reversed the conclusion entirely                                                  |
| My own first grid run                           | An $n \geq 25$ floor silently deleted the BIG-TREND rung — the one thing the study was for                                                                                     | Lowered to 15, thin cells printed and flagged                                                                                    |
| MGC run-day threshold                           | A flat $\geq \$300$ bar is only 30 points at $\$10/\text{pt}$ — all 16 days qualified, crosstab carried zero information                                                       | Switched to a median split on gold's own distribution                                                                            |

The `scipy` package is not installed on this box, so every rank correlation here is computed as Pearson-on-ranks by hand. That is the same quantity, but worth knowing if anyone re-runs these scripts and finds the import missing.

10 · DISPOSITION — every lead touched in this section

| LEAD                                                | VERDICT                        | EVIDENCE / REVIVAL<br>CONDITION                                                                                                                                                              |
|-----------------------------------------------------|--------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| abs_veto_55s mechanism                              | <div>LIVENOT-AN-ACTION</div>   | Already promoted two-sided 25 Jul. Battery re-confirms: 16/21 days green, both regimes, both halves, both sides. Nothing to deploy.                                                          |
| abs_veto_short exit ladder 1.5R/2.5R → 1.0R/1.5R    | <div>SHADOW</div>              | The week's main proposal. Diagnosis: identical winner size, 33% bigger losses, 13pt worse win rate vs its shadow twin. Run as a paired forward A/B, not a switch.                            |
| Restore abs_veto_short ATR ≥ 16 floor               | <div>REFUTED</div>             | Timing fits, mechanism does not: the entries the floor would block are +\$233.50 over 38 trades in shadow since 1 Aug. Not the cause.                                                        |
| Relegate the abs_veto SHORT direction               | <div>REFUTED</div>             | Shadow says short is the STRONGER side (+ \$2,651 vs +\$1,371, +\$15.50 vs +\$8.02/tr). Tune the side, never the direction.                                                                  |
| ODR causal day-classifier (8 pre-13:00 features)    | <div>REFUTED</div> as searched | <b>Named test: search-aware placebo, 2,000 shuffles.</b> Best of 272 rules beaten by 40.0% of pure-noise searches. Not a thin-n verdict — a negative on 41 days.                             |
| ODR day-classifier, pre-registered single feature   | <div>PARKED</div>              | <b>Revive with:</b> ATR-14 ≥ 9.5pt at 13:00 fixed in advance and 60 forward sessions untouched. Per-day sd is \$1,384; 5 days is 8% of what a \$500/day effect needs.                        |
| ONE shared day-classifier for odr + MGC run-catcher | <div>PARKED</div>              | 16 overlapping days: agreement 38% vs 50% by chance, rank corr −0.129. <b>Revive only if</b> a mechanism is proposed first — do not re-test by search. Gold needs its own detector.          |
| ODR drift gate (sims 60-63)                         | <div>SHADOWUNTESTED</div>      | Identical to ungated, cell by cell — <b>zero entries rejected</b> in its one day. <b>Proves out when:</b> drift confirms a direction opposite the rider's 15-min momentum.                   |
| ODR stop width 3.0× over 2.0×ATR                    | <div>SHADOW</div>              | Real (+\$8.19 vs −\$15.08/tr; wins 4 of 5 days; replay agrees on 41 days) but the DAY effect is 2.2× larger. All arms n < 40 — <b>no promotion this week.</b>                                |
| Two-ratchet runner-clip watch                       | <div>UNMEASURED</div>          | No grind trades since 5 Aug; 0 clips on 3 pre-Aug runners; and the tool's own cross-check error (− \$406) exceeds the signal (−\$294). <b>Revive when</b> grind trades a genuine trend week. |
| 2R-partial smoothness                               | <div>UNMEASURED</div>          | Green-day % is <b>zero for both arms</b> over the 5 days grind traded — "bank early" wins trivially on an all-red book. <b>Revive when</b> the sample contains rMFE ≥ 4R runners.            |
| Live exhaustion_short fade-scalp trial (A0.5/B1.5)  | <div>PARKED</div>              | −\$480 on 80 live legs; repriced −\$0.48/sig overall but −\$20.56/sig on <b>BIG-TREND</b> . <b>Revive as</b> a chop-only config — it is +\$0.97/sig there.                                   |
| exhaustion_short tight chandelier on trend rungs    | <div>SHADOW</div>              | <b>Backed by two independent studies (trial grade + grid plateau). Blocked from ADOPT only by the V5 out-of-sample leg disagreeing.</b>                                                      |
| grind_long MED-TREND A0.5/B1.0                      | <div>SHADOW</div>              | +\$9.4/sig, strip-3 +\$731, LOO +\$8.0 — the cleanest robustness in the grid. Moot until the gate trades again.                                                                              |
| Exit tuning for the thrust family                   | <div>REFUTED</div>             | <b>Named test:</b> all 58 policies span −\$39.4 to −\$40.6/sig on SCALP-CHOP. No exit rescues an un-vetoed thrust; the veto is the fix.                                                      |

| LEAD                     | VERDICT | EVIDENCE / REVIVAL<br>CONDITION                                                                                                                                                                                                                                   |
|--------------------------|---------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| BIG-TREND rung R-targets | PARKED  | Window extended to 45 days + 4,270 archive entries as instructed and <b>still not provable</b> — the rung is 2.96% of minutes and the two windows disagree at $n \approx 25$ each. <b>Revive at</b> ~120 BIG-TREND entries, or a deliberate high-ER archive pull. |

**The one stone still unturned.** Every study in this section measured the exit or the day and treated the entry as given. But the two clearest results both point at the entry population instead: no exit saves the un-vetoed thrust, and no pre-session feature calls the Open Rider's day. The Open Rider takes its direction from a raw 15-minute momentum sign with no quality test at all — and I never asked whether the losing days are days when that particular signal is unusable, as opposed to days that are simply bad. That is a per-signal question, not a per-day one, and it has roughly twenty times the sample size: **106 trades instead of 5 days**. If the Open Rider is rescuable, I now think that is where it happens, and it is the first thing I would run next week.

---

## **P2.5** MID-WEEK MUSINGS — the leads, tested hard · week of 10-14 August 2026

---

Gaz, three things came up mid-week that were worth a night each: **gold**, **grind's on/off switch**, and **a measuring instrument this very report leans on**. I re-derived every number below from the tape and the ledger tonight — nothing here is inherited from the mid-week chat notes, and where my numbers disagree with those notes I say so and show you which one is right.

**What survived.** Two things, and both are gold. First, the **MGC day rider**: pick a direction an hour into the US session and hold it to the 20:40 flat. Its threshold grid is a plateau — all fifteen cells I swept make money, from +\$552 to +\$1,742 over sixteen sessions — and at the loose end it beats the



best always-on constant on both hit rate and dollars while trading half as often. Second, the **coil bouncer**: 154 trades, a median trade of **+\$7.75**, beaten by **0 of 12** placebo runs, and its mirror image loses \$924. Small money, but real shape. On the MNQ side one clean thing survived too, and it is a negative: **don't buy a grind long that is fighting the last thirty minutes**. That single veto turns grind's tick-honest book from **+\$2,407** into **+\$3,674** on the same 26 days, and it does it while keeping every one of the ten best trades.

#### THE ONE-PARAGRAPH VERSION

**GOLD**: its big moves are slow grinds, not thrusts — median efficiency **0.20**, four of the top five ran for **3 to 24 hours**. Boarding them is already solved; the money question is which days to show up for. **The day rider is the best shape gold has ever shown us** (plateau, not a peak) and goes to shadow. **GRIND**: the "ER 0.35 floor" in the brief **does not exist** — it was deleted on 1 August and had been broken since 29 July. What actually decides whether grind trades is the open-hour watcher's break-trio, and on 26 days of tape that trio would have let grind fire **four times**. The mechanism that works is not a trend confirmation at all, it is a counter-trend veto. **DIED**: the ER floor as an enable mechanism — at 0.35 it throws away nine of the ten best trades, its strip-3 is **-\$1,044**, and 182 of 400 random subsets of the same size beat it. **BROKEN INSTRUMENT**: `selector_nightly.py` counted four quarantined phantom trades on 13 August, inflating that day's regret by **39%** — and that file feeds the router-value figure in this report.

Housekeeping for the reader: the full router operation review is its own section (P2.6) and I have not repeated it here. Money framing for the week — the desk booked **+\$1,292.50** across 113 legs from 10-14 August (day rider **+\$1,617.50** on 5 legs, tournament **-\$325.00** on 108), so the leads below are being judged against a desk whose engine is currently the rider, not the gates.

---



## PART A — GOLD. Everything you have not seen yet, from the start.

You have never been shown a gold section, so this one starts from nothing. We have been capturing **MGC** (COMEX gold, \$10.00 a point — five times MNQ's \$2.00) since 4 August, plus an older block recovered from the retired V5 desk. That gives me **18 trading days of gold bars and 20 of ticks**, spanning 7 July to 14 August with a hole from 18 July to 4 August where nothing was captured. Sixteen of those days have a full US session. **Sixteen sessions is thin and I will keep saying so.** Nothing in this part is live or proposed for live; it is all shadow-or-nothing.

### FIRST, WHY THIS IS NOT JUST ANOTHER SWEEP — SIX GOLD ATTACKS ARE ALREADY DEAD

On 5 August six separate attempts to build a gold gate were closed out, each with a named test. That matters here because it tells you what not to re-run, and because the two survivors below are a different shape of idea, not a re-tune of a dead one.

| THE ATTACK                    | HOW IT DIED — THE NAMED TEST                                                                                                 |
|-------------------------------|------------------------------------------------------------------------------------------------------------------------------|
| momentum / direction gate     | 5 of 72 cells positive; the best strips to \$16; plain always <b>SHORT</b> (−\$532) <b>BEAT</b> the directed gate (−\$1,274) |
| arm-then-confirm delay        | the confirmed side is a coinflip — 48-50% across 24 cells                                                                    |
| router-filtered (live meters) | placebo-controlled and beaten by 31-80% of random removals                                                                   |
| L2 book direction / veto      | <b>NULL</b> and well powered — the edge flips sign across 5 disjoint samples                                                 |
| fader on extension + clock    | Hurst on our own gold tape is 0.489 — a random walk, so there is nothing to fade                                             |
| thrust-continuation           | catches 22 of 22 runs, but the losers are indistinguishable from the winners (ATR effect size +0.21)                         |

The one survivor of that whole round was an oracle day-filter — given foresight of which days contained a run, +\$779 with strip-best-day +\$356. That is exactly your "let the router switch it off", and it is still the live thread. Everything below is downstream of it.

### FINDING 1 — GOLD'S BIG MOVES ARE GRINDS, NOT THRUSTS

I ran the run-catcher over the last nine days of gold tape tonight. It finds every directional leg worth  $\geq \$300$  on one lot that never gave back more than a third

of itself while it was alive. Fifteen qualified. Here are the five biggest, and the shape of them is the finding:

| START<br>(UTC) | DIR | POINTS | \$ 1 LOT | MINUTES | EFFICIENCY | PRE-60M<br>RANGE | MINS TO<br>20%<br>MARK |
|----------------|-----|--------|----------|---------|------------|------------------|------------------------|
| 08-05 01:33    | UP  | 225.4  | \$2,254  | 1,474   | 0.12       | 7.1×ATR          | 68                     |
| 08-10 14:09    | UP  | 110.4  | \$1,104  | 726     | 0.16       | 8.6×ATR          | 19                     |
| 08-14 07:31    | UP  | 65.2   | \$652    | 367     | 0.20       | 6.6×ATR          | 29                     |
| 08-11 04:33    | DN  | 63.3   | \$633    | 206     | 0.23       | 5.1×ATR          | 16                     |
| 08-07 12:29    | UP  | 62.8   | \$628    | 18      | 0.74       | 9.1×ATR          | 1                      |

**Read the efficiency column.** A move with efficiency 0.20 travels five times its net distance getting where it goes — up, back, up, back, up. Four of these five ran for between three and twenty-four hours. Only one (08-07) is a thrust, and it is the smallest of the five. The median of the top five is **efficiency 0.20**, and they all start from tape that is already moving — the hour before ignition is 7× a normal ATR, not a quiet coil.

**Boarding them is not the problem.** The gate fires inside **5 of the 5**, and price reaches the 20% mark of the eventual move in a median of **19 minutes**, leaving roughly \$520 gross still on the table. We are not missing these because we cannot see them.

## AND HERE IS WHERE I HAVE TO CORRECT THE MID-WEEK NOTE

The mid-week note said the exit was the cap: that switching from a 2.5×ATR trail to a wide 5×ATR stop held eight hours moves the ceiling from about \$83 a week to about \$346 a week. **I can reproduce the direction of that, but only inside the oracle**, and the note did not make that boundary loud enough. Here is the matrix, on identical entries, with both columns shown — ALL is every signal a live gate would take, ORACLE is only the signals that landed inside one of the week's top-five runs (labels you can only know afterwards):

| EXIT                                                                            | N  | WIN% | ALL NET | MEDIAN<br>TRADE | ORACLE N | ORACLE<br>NET |
|---------------------------------------------------------------------------------|----|------|---------|-----------------|----------|---------------|
| window A — 7-14 August (the mid-week note's own window), 49 signals, ~1.2 weeks |    |      |         |                 |          |               |
| 2.5×ATR trail ·<br>1.5 stop · 3h<br>(baseline)                                  | 49 | 36.7 | −\$275  | −\$19.80        | 13       | +\$143        |
| no trail · 5×ATR<br>stop · 3h                                                   | 49 | 34.7 | −\$131  | −\$56.50        | 13       | +\$382        |
| no trail · 5×ATR<br>stop · 8h                                                   | 49 | 22.4 | −\$609  | −\$75.80        | 13       | +\$589        |

| EXIT                                                                | N  | WIN% | ALL NET  | MEDIAN<br>TRADE | ORACLE N | ORACLE<br>NET |
|---------------------------------------------------------------------|----|------|----------|-----------------|----------|---------------|
| no trail · 5×ATR<br>stop · 12h                                      | 49 | 14.3 | −\$2,076 | −\$87.90        | 13       | −\$464        |
| window B — 5–14 August (two days more tape), 69 signals, ~1.6 weeks |    |      |          |                 |          |               |
| 2.5×ATR trail ·<br>1.5 stop · 3h<br>(baseline)                      | 69 | 34.8 | −\$352   | −\$23.70        | 21       | +\$220        |
| no trail · 5×ATR<br>stop · 8h                                       | 69 | 26.1 | +\$1,336 | −\$81.90        | 21       | +\$3,051      |
| no trail · 5×ATR<br>stop · 12h                                      | 69 | 20.3 | +\$2,611 | −\$89.70        | 21       | +\$4,357      |

### Three honest readings of that table.

1. **The exit really does matter, and a trail really is wrong for gold.** On the oracle column, in both windows, moving off the trail and onto a wide stop held for hours multiplies the take — +\$143 to +\$589 in window A, +\$220 to +\$3,051 in window B. That part of the mid-week claim stands.
2. **But the exit alone does not make it a strategy.** On the ALL column in window A, every single cell loses money. The wide stop without the day filter takes bigger losses on the ordinary days than it wins on the run days. **The exit was a cap, not the cap** — you still need the filter, so the order of the gold problem is unchanged: day selection first.
3. **The hold cap is not identified at this sample size.** Window A says 12 hours destroys it (−\$464 even on the oracle); window B says 12 hours is the best cell (+\$4,357). The whole difference is one run — the 08-05 move that ran 225 points over 24 hours. **Adding two days of tape flipped the conclusion.** That is the clearest statement I can make about how thin this is.

The hold-time curve on window B makes the same point — it is not a smooth curve, it is a staircase built out of a handful of trades:

| HOLD CAP                | 2H     | 4H     | 6H     | 8H       | 10H      | 12H      |
|-------------------------|--------|--------|--------|----------|----------|----------|
| net (all 69<br>signals) | +\$283 | +\$517 | +\$762 | +\$1,336 | +\$2,139 | +\$2,611 |
| win rate                | 39.1%  | 34.8%  | 29.0%  | 26.1%    | 23.2%    | 20.3%    |

Note what holding longer does to the feel of it: the win rate falls from 39% to 20% as the money rises. That is a strategy where four trades in five lose and the fifth pays for everything. On our own judging rule that is not a disqualification — expectancy is what counts — but it is exactly the shape you will find hardest to sit through, and you should know that before it ever goes near a live switch.

## FINDING 2 — THE MGC DAY RIDER, AND WHY A PLATEAU IS THE EVIDENCE

This is the best-shaped thing gold has produced. The rule is simple enough to say in one line: **at 13:30 UTC start watching; an hour in, ask the live drift detector which way the day has efficiently already moved; if it is confirmed, take two lots that way and hold to the 20:40 flat.** One trade a day, never overnight, and it finishes inside the session you watch.

**The thresholds were derived from gold, not borrowed from MNQ.** You asked for that specifically. The method was to find what percentile MNQ's live constants occupy in MNQ's own distribution, then read gold's number at the same percentile — transferring the selectivity, not the number. Tonight's derivation:

|     | EFFICIENCY P50 | ROUNDRIP P50 | ATR P50 | SESSIONS |
|-----|----------------|--------------|---------|----------|
| MNQ | 0.110          | 0.418        | 37.8 pt | 49       |
| MGC | 0.124          | 0.411        | 3.6 pt  | 16       |

MNQ's live ER floor of 0.15 sits at MNQ's 65th percentile; its roundtrip floor of 0.45 sits at the 53rd. Read off gold's own distribution those percentiles give **ER 0.155 and RT 0.417** — within a whisker of MNQ's numbers. **Gold's ATR is ten times smaller and the thresholds barely move**, because efficiency and roundtrip are ratios and the size cancels out. That had to be measured rather than assumed, and it is the answer to "why are you using MNQ thresholds": we are not — they simply landed in the same place.

**Now the grid.** I swept five efficiency floors against three roundtrip floors around that derived pair, scoring every cell on the same sixteen sessions, always against the best always-on constant:

| ER FLOOR | RT FLOOR | DAYS FIRED | DIRECTION HELD TO THE FLAT | NET      | MEDIAN DAY | MFE (WATCHING) | STRIP-BEST-DAY |
|----------|----------|------------|----------------------------|----------|------------|----------------|----------------|
| 0.093    | 0.417    | 8 / 16     | 87.5%                      | +\$1,742 | +\$280     | +\$3,852       | +\$899         |
| 0.109    | 0.417    | 8 / 16     | 87.5%                      | +\$1,742 | +\$280     | +\$3,852       | +\$899         |
| 0.124    | 0.417    | 8 / 16     | 87.5%                      | +\$1,742 | +\$280     | +\$3,852       | +\$899         |
| 0.093    | 0.350    | 9 / 16     | 77.8%                      | +\$1,259 | +\$279     | +\$3,923       | +\$416         |
| 0.140    | 0.350    | 7 / 16     | 85.7%                      | +\$1,395 | +\$279     | +\$3,327       | +\$552         |
| 0.093    | 0.480    | 7 / 16     | 85.7%                      | +\$899   | +\$279     | +\$2,921       | +\$552         |
| 0.155    | 0.480    | 6 / 16     | 83.3%                      | +\$552   | +\$223     | +\$2,396       | +\$225         |

| ER FLOOR                                                                                                                                                                             | RT FLOOR | DAYS FIRED | DIRECTION HELD TO THE FLAT | NET      | MEDIAN DAY | MFE (WATCHING) | STRIP-BEST-DAY |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------|------------|----------------------------|----------|------------|----------------|----------------|
| ...and the eight cells not shown, which sit between these. All 15 of 15 cells are positive. Range: +\$552 to +\$1,742. Held: 77.8% to 87.5%. Worst single day in every cell: -\$619. |          |            |                            |          |            |                |                |
| CONTROL — best constant (always short)                                                                                                                                               | —        | 16 / 16    | 56.2%                      | +\$1,520 | —          | —              | —              |

**Why I am making a fuss about the plateau.** When a swept grid has one good cell surrounded by bad ones, you have found a fluke and fitted it. When every cell in the neighbourhood pays, you have found a structure and the exact number does not matter much. This is the second kind. **The evidence is the plateau, and the peak cell is not the recommendation** — if this ever ships, it ships from the middle of the plateau, not from +\$1,742.

And against the control: gold drifted down over this sample, so always short scored +\$1,520 with a 56.2% day-hit-rate. That is the bar, not a coinflip. The loose cells clear it on both counts — **+\$1,742 against +\$1,520, at 87.5% against 56.2%, while trading on half the days.** Trading half as often for more money is the part that makes me take it seriously.

**How long to wait before calling the direction** is its own axis, and it has a shape too:

| DECIDE AT | DAYS FIRED | HELD TO THE FLAT | NET      | MFE      | WORST ADVERSE | VS BEST CONSTANT |
|-----------|------------|------------------|----------|----------|---------------|------------------|
| +30 min   | 9 / 16     | 77.8%            | +\$1,585 | +\$5,175 | -\$1,440      | +21.5pp          |
| +45 min   | 10 / 16    | 60.0%            | +\$826   | +\$4,820 | -\$860        | +3.8pp           |
| +60 min   | 6 / 16     | 83.3%            | +\$552   | +\$2,396 | -\$936        | +27.1pp          |
| +90 min   | 5 / 16     | 40.0%            | -\$845   | +\$1,839 | -\$1,068      | -28.8pp          |
| +120 min  | 4 / 16     | 25.0%            | -\$950   | +\$1,006 | -\$1,360      | -50.0pp          |

Waiting too long is fatal — by 90 minutes in you are buying the end of the move and the direction holds only 40% of the time. Between 30 and 60 minutes it works. **The +45 minute cell is the honest warning sign:** it dips to 60% held while its neighbours either side are 77.8% and 83.3%. With sixteen sessions, one day landing differently moves a cell by six percentage points, so do not read that dip as structure — read it as the noise floor of this sample.

**The number you will care about most** is the gap between what the strategy books and what you could claim. Median MFE per fired day is **+\$453** against a median adverse excursion of **-\$400**. The profit is genuinely there to be taken,

but only behind a wide stop and only if someone is watching — which is what you said you do. An unattended tight version of this gets stopped out of the same days.

### FINDING 3 — THE COIL BOUNCER (THE STRONGEST STANDALONE GOLD GATE, AND IT IS SMALL)

Different idea entirely: when gold's 20-minute range squeezes to 0.6× its own two-hour median, fade whichever edge of the coil price is sitting on, targeting the middle. It trades compression, which is why it is not the same thing as the extension-fader that died in the earlier round.

|                                        | N   | DAYS | GREEN DAYS | WIN%  | \$/TRADE | NET       | STRIP-BEST-DAY | MEDIAN TRADE |
|----------------------------------------|-----|------|------------|-------|----------|-----------|----------------|--------------|
| coil bounce (fade the edge)            | 154 | 17   | 10         | 51.9% | +\$3.05  | +\$470.42 | +\$289.42      | +\$7.75      |
| mirror control (break OUT of the coil) | 154 | 17   | 4          | 33.1% | -\$6.00  | -\$923.75 | -\$1,059.61    | -\$12.54     |

Three things make this the most trustworthy gold result I have, even though it is the smallest: **the median trade is positive** (+\$7.75 — it is not one lucky day carrying 153 losers); **it survives stripping its best day** (+\$289 of the +\$470 remains); and **a placebo of twelve random-time trade sets with the same count per day has a median of −\$214 and beat it 0 times out of 12**. On top of that, doing the exact opposite loses \$924 — the mirror control is the cleanest evidence that it is the coil doing the work and not the clock. At roughly **\$137 a week** it is below the bar you have set for going live, and it stays in the shadow book.

### FINDING 4 — THE LONDON-OPEN BREAK FADE (SURVIVES ITS PLACEBO, STILL FAILS AS A TRADE)

Gold's opening-range breaks fail: taking the continuation wins 6.2% / 25.0% / 37.5% of the time across the three anchors I tested. So the trade is to fade them. Only one of the three anchors holds up:

| ANCHOR      | N  | GREEN DAYS | WIN%  | \$/TRADE | NET       | STRIP-BEST-DAY |
|-------------|----|------------|-------|----------|-----------|----------------|
| LONDON open | 16 | 7          | 43.8% | +\$31.40 | +\$502.42 | +\$125.01      |
| RTH open    | 16 | 6          | 37.5% | -\$7.32  | -\$117.04 | -\$346.83      |

| ANCHOR | N | GREEN DAYS | WIN% | \$/TRADE | NET      | STRIP-BEST-DAY |
|--------|---|------------|------|----------|----------|----------------|
| PM fix |   | 16         | 3    | 18.8%    | -\$31.88 | -\$510.07      |
|        |   |            |      |          |          | -\$587.71      |

The London cell is beaten by **0 of 7** shifted anchors (the seven fakes score – \$325, –\$112, –\$305, +\$99, –\$156, –\$228, –\$100) — so the anchor is carrying it, not the general shape of the day. That is a real result. But look at the honest side: **nine of the sixteen trades lose, so the median trade is a loss**, and stripping the best day takes +\$502 down to +\$125. Sixteen trades is one trade a day for three weeks. It is a lead, not a strategy.

## THE THREE TRAPS THAT BIT DURING THIS WORK — CHECK ANYTHING NEW AGAINST THEM

| TRAP                                   | WHAT IT DOES                                                                                                                                  | COST WHEN IT BIT                                                                    |
|----------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------|
| resample labels a bar by its LEFT edge | the 10:00-10:01 bar is stamped 10:00, but its close is only known at 10:01 — racing from the stamp replays the decision minute with hindsight | the coil gate read +\$906.55; honest, it is +\$470.42                               |
| DriftRead.ok is not "fired"            | ok means "enough data to say anything"; the field that means the signal fired is confirmed                                                    | produced a fake "16 of 16 days fired, +12.5pp edge" on rows showing efficiency 0.01 |
| lake.connect() defaults to MNQ         | query gold without symbol="MGC" and you get zero rows — no error, just an empty answer                                                        | a confident NULL on a question never actually asked                                 |

## WHAT IT WOULD COST TO FIND OUT — AND THE ANSWER IS DAYS, NOT CONFIGURATIONS

Every gold result above is limited by the same thing, and it is not cleverness. It is **sixteen sessions**. Look at what that does: adding two days of tape flipped the exit matrix's conclusion about hold time; the day rider's grid moves by a whole cell when one day lands differently; the London fade rests on sixteen trades. We have already swept the configurations — fifteen threshold cells, five decision times, twelve exit rules, three anchors. **More sweeping on this sample cannot tell us anything new; it can only fit it harder.**

The good news is that the missing ingredient is the cheap one. Capture is already running for both symbols and costs nothing extra. At five sessions a week, gold reaches **n = 40 sessions around the third week of September**, which is roughly where a day-rider hit rate of 87.5% stops being a story about eight days. Everything above goes in the shadow book now — free, riskless, and accumulating exactly the thing it lacks. The one thing I would build rather than wait for is the causal day filter for the run catcher, because that is the piece the



whole gold problem still turns on, and it needs to be placebo-controlled against discarding the same number of firings at random.

---

## **PART B — THE STANDING GRIND LEAD: which enable-mechanism actually works?**

The question on the table: what best captures grind's trend-day upside while killing its chop-churn — **(a)** trend-day-only arming, **(b)** its own entry trend-confirmation, or **(c)** the current ER-0.35 floor? I built the book to answer it properly: every grind\_long signal on 26 days of tick-covered tape, fired exactly as the live gate fires it (slope 0.4, fast-slope,  $ATR \geq 22$  — the floor that is genuinely deployed), then **exited tick by tick as the live managed slate exits**: Lot A a 2.5R scalp, Lot B the loose-then-lock chandelier ( $3.5 \rightarrow 6.0R \rightarrow 0.5$ ), both behind the native 1-ATR stop, \$1.50 a round trip per lot. That is the LIVE/MANAGED grind, not the unmanaged grind\_fast shadow, which flatters itself and would have sent me the wrong way.

### **FIRST, A CORRECTION: OPTION (C) DOES NOT EXIST**

There is no ER-0.35 floor on grind. `deciders.ER_FL00R` is an **empty dictionary**. The floors were deleted on 1 August, and the audit note in the code says they had not blocked a single live trade since 29 July anyway, because the dual-slot cutover broke the lookup. When they were re-derived on the blocked book, grind's  $ER \geq 0.35$  turned out to have been blocking a profitable population (+\$2,923 over 173 signals). So the brief's option (c) is a ghost — and tonight's book says deleting it was right, which I will show in a moment.

**What actually decides whether grind trades** is two things: the  $ATR \geq 22$  floor inside the gate, and the on/off switch outside it. The 0.35 number does still exist — but in `open_hour_watch.py`, as one leg of the **break trio** that is allowed to arm a benched momentum gate: last-hour  $ER \geq 0.35$  AND last-hour  $ATR \geq 18$  AND a new session extreme is taken. So the real question (c) is "is the break trio the right arming test", and that is answerable.



THE BOOK ITSELF — WHERE GRIND'S MONEY ACTUALLY IS

| N                                      | NET | \$/TRADE | WIN%     | STRIP-3-BEST | WORST LEAVE-ONE-DAY-OUT |          |
|----------------------------------------|-----|----------|----------|--------------|-------------------------|----------|
| grind_long, live-managed, 26 tick-days | 133 | +\$2,407 | +\$18.10 | 31.6%        | +\$346                  | +\$1,368 |

Two facts about that book decide everything else. **It wins under a third of the time and it is carried by runners:** 28 trades reached 4R or better and those 28 carry **+\$11,991**, meaning the other 105 bleed **−\$9,584**. Strip the three best trades and **+\$2,407** becomes **+\$346**. Any filter that touches the runners kills the gate, which is why every mechanism below is scored on "how many of the ten best trades does it keep".

Split by regime at entry — ATR level, efficiency and range position, exactly as the discipline requires — the shape is not what the folklore says:

| REGIME AT ENTRY                               | N  | NET      | \$/TRADE  |
|-----------------------------------------------|----|----------|-----------|
| dead-chop (ER < 0.10)                         | 40 | −\$802   | −\$20.00  |
| normal-chop (ER 0.10-0.20)                    | 31 | −\$358   | −\$11.60  |
| in-between / building (ER 0.20-0.35)          | 35 | +\$897   | +\$25.60  |
| clean-trend (ER ≥ 0.35)                       | 14 | +\$262   | +\$18.70  |
| violent-whipsaw (top-quartile ATR, ER < 0.15) | 13 | +\$2,408 | +\$185.20 |

**Grind's money is not in "clean trend".** The 14 clean-trend entries made **+\$262**. The 13 high-volatility, low-efficiency entries made **+\$2,408** — the entire book. That is the chandelier catching a big swing on a violent day, and it is the single most important thing in this table, because **every mechanism in the brief is a trendiness test, and trendiness is not where the money is.**

THE THREE MECHANISMS, SCORED AND SWEPT

Every threshold is a guess until it is swept, so each mechanism was scanned across its range rather than tested at one number. Placebo column = how many of 400 random subsets of the same size beat the real one; if a mechanism cannot beat coin-flip removal it has no claim.

| MECHANISM | N KEPT | NET      | \$/TRADE | STRIP-3 | TOP-10 WINNERS KEPT | PLACEBO |
|-----------|--------|----------|----------|---------|---------------------|---------|
|           | 133    | +\$2,407 | +\$18.10 | +\$346  | 10 / 10             | —       |

| MECHANISM                                                             | N KEPT | NET      | \$/TRADE | STRIP-3  | TOP-10 WINNERS KEPT | PLACEBO   |
|-----------------------------------------------------------------------|--------|----------|----------|----------|---------------------|-----------|
| BASE —<br>deployed,<br>ATR ≥ 22 only                                  |        |          |          |          |                     |           |
| (a) trend-day-only — arm once the session so far has proven efficient |        |          |          |          |                     |           |
| session-ER ≥ 0.02                                                     | 96     | +\$2,780 | +\$29.00 | +\$718   | 8 / 10              | 97 / 400  |
| session-ER ≥ 0.05                                                     | 45     | +\$2,630 | +\$58.50 | +\$758   | 3 / 10              | 40 / 400  |
| session-ER ≥ 0.08                                                     | 13     | +\$10    | +\$0.70  | −\$834   | 0 / 10              | 221 / 400 |
| (b) entry trend-confirmation — require a 15-minute impulse our way    |        |          |          |          |                     |           |
| net15 ≥ 0.20×ATR                                                      | 118    | +\$198   | +\$1.70  | −\$1,862 | 8 / 10              | 391 / 400 |
| net15 ≥ 0.50×ATR                                                      | 109    | +\$299   | +\$2.70  | −\$1,762 | 7 / 10              | 374 / 400 |
| net15 ≥ 1.00×ATR                                                      | 96     | +\$1,380 | +\$14.40 | −\$681   | 7 / 10              | 249 / 400 |
| (c) the ER floor at entry — the mechanism the brief names             |        |          |          |          |                     |           |
| ER30 ≥ 0.25                                                           | 33     | −\$558   | −\$16.90 | −\$2,118 | 2 / 10              | 331 / 400 |
| ER30 ≥ 0.35                                                           | 14     | +\$262   | +\$18.70 | −\$1,044 | 1 / 10              | 182 / 400 |
| the deployed lever for comparison — raise the ATR floor               |        |          |          |          |                     |           |
| ATR ≥ 24                                                              | 107    | +\$2,518 | +\$23.50 | +\$458   | 10 / 10             | 116 / 400 |
| ATR ≥ 34                                                              | 33     | +\$1,766 | +\$53.50 | −\$294   | 7 / 10              | 52 / 400  |

**(c) is dead, and I will name the tests.** At  $ER \geq 0.35$  it throws away **nine of the ten best trades in the book** — that is the fake-filter rule failing outright. Its strip-3 is −\$1,044, meaning what is left after removing three trades is a loser. And 182 of 400 random removals of the same size beat it, so it is not distinguishable from deleting trades at random. The  $ER \geq 0.25$  cell is worse still (−\$558, beaten by 331 of 400). This independently confirms the 1 August decision to delete the floor, on a different population and a different exit model.

**(b) is worse than dead — it is backwards.** Requiring a 15-minute impulse our way takes +\$2,407 down to +\$198 and 391 of 400 random subsets beat it. The reason is visible in the trades themselves: the ten worst fires had 15-minute impulses of +\$113, +\$139, +\$147, +\$192 and +\$263 points, which are bigger than the winners'. Confirming the impulse means arriving later into the same move and buying the part that reverses. This is the "the losers are indistinguishable" pattern showing up on MNQ exactly as it did on gold.

**(a) is directionally real but noisy.** Requiring the session to have shown some efficiency improves things at 0.02 and 0.05 (+\$2,780 and +\$2,630 against the base's +\$2,407, on 96 and 45 trades of the 133), but the curve is not monotone — 0.04 dips to +\$1,944, 0.06 falls to +\$1,204, 0.08 collapses to +\$10. That is not a

plateau, it is a wobble. And the 0.05 cell keeps only 3 of the 10 best trades, which is the fake filter tell again.

AND THE ARMING TEST WE ACTUALLY RUN: THE BREAK TRIO ADMITS ALMOST NOTHING

Scored the way it is used — scan the session, the first minute all three legs pass arms the gate for the rest of the day — here is what the trio does to 26 days of grind signals:

| N                                       | NET | \$/TRADE | WIN%            |
|-----------------------------------------|-----|----------|-----------------|
| every grind fire                        | 133 | +\$2,407 | + \$18.10 31.6% |
| armed by the trio (fires after the arm) | 4   | +\$100   | + \$24.90 25.0% |
| never armed — refused                   | 129 | +\$2,308 | + \$17.90 31.8% |

My replication of the trio fired on **3 of 45 tape-days**, and on just one of the 26 days grind actually signalled. (The live watcher's version armed once more than mine did — on 13 August at 13:58 — because it reads its efficiency and its session extremes over slightly different windows than I can reconstruct; so treat the "4 fires" as approximate and the order of magnitude as the finding.) The order of magnitude is the finding: **the arming test in force is so strict that grind is, in practice, a gate that does not trade.** The live ledger agrees without any modelling — **grind\_long has not taken a single live trade since 5 August**, seven sessions ago, and the Paris-midnight service now start-benches it every night by design ("momentum start-benched per MONDAY #2 — router arms on evidence").

What that bench has cost, priced honestly:

| WINDOW                      | GRIND SIGNALS | TICK-HONEST VALUE IF TAKEN | WITH THE VETO BELOW |
|-----------------------------|---------------|----------------------------|---------------------|
| this report week, 10-14 Aug | 10            | -\$66                      | +\$64               |
| since the bench, 6-14 Aug   | 20            | +\$490                     | +\$912              |
| router era, 30 Jul onward   | 51            | +\$814                     | +\$1,688            |

So: **this week the bench cost nothing** — it saved \$66. Over the nine sessions since the bench began it cost about **\$490**, and almost all of that is two trades: 6 August 13:38 (+\$664) and 13 August 13:41 (+\$447). Both are chandelier runners on high-ATR entries. Notice the timing on the second one: the open-hour watcher armed grind on 13 August at 13:58 — **seventeen minutes after that signal fired.** The trio was not wrong about the day; it was late. That is grind's real

problem in one line: the gate lives or dies on two or three runner entries a month, and an arming test that needs a proven trend cannot be early enough to catch the entry that starts it.

THE MECHANISM THAT DOES WORK, AND IT IS A VETO

Since every positive confirmation failed, I swept the negative: rather than requiring proof of trend, refuse the entries that are fighting the recent tape. Two cheap conditions at signal time:

| ARM                                                         | N   | NET      | \$/TRADE | WIN%  | STRIP-3  | WORST LODO | TOP-10 KEPT | PLACEBO  |
|-------------------------------------------------------------|-----|----------|----------|-------|----------|------------|-------------|----------|
| BASE                                                        | 133 | +\$2,407 | +\$18.10 | 31.6% | +\$346   | +\$1,368   | 10 / 10     | —        |
| skip if the 30-min net is worse than $-1.0\times\text{ATR}$ | 125 | +\$3,134 | +\$25.10 | 32.8% | +\$1,072 | +\$2,094   | 10 / 10     | 27 / 400 |
| ...and only in the top 40% of the hour's range              | 97  | +\$3,674 | +\$37.90 | 35.1% | +\$1,613 | +\$2,751   | 9 / 10      | 16 / 400 |

The first line is the one I trust. It removes **eight trades spread across seven different days** — every one of them stopped out, seven of them losers — for a saving of **+\$727** — and it keeps every single one of the ten best trades. It reads in plain English as: don't buy a continuation long while price has just fallen more than an ATR over the last half hour. The sweep behind it is a genuine plateau — every threshold from  $-2.0\times\text{ATR}$  to  $-0.5\times\text{ATR}$  improves the book, so nothing hangs on the literal:

| 30-MIN NET FLOOR         | $-2.0\times\text{ATR}$ | $-1.5\times$ | $-1.25\times$ | $-1.0\times$ | $-0.75\times$ | $-0.5\times$ |
|--------------------------|------------------------|--------------|---------------|--------------|---------------|--------------|
| net (no range condition) | +\$2,636               | +\$3,014     | +\$3,144      | +\$3,362     | +\$3,362      | +\$3,282     |

**The second condition is a peak cell and I am flagging it as one.** Requiring price in the top 40% of the hour's range scores +\$3,674 at exactly 0.60, but 0.65 falls to +\$2,266 and 0.70 to +\$1,008. That is a spike with cliffs, i.e. a fitted literal. **Ship the veto, do not ship the 0.60.**

Robustness on the veto pair, stated plainly: it improves both halves of the sample when split by date (first thirteen days +\$965  $\rightarrow$  +\$1,063; last thirteen +\$1,442  $\rightarrow$  +\$2,611), its worst leave-one-day-out is +\$2,094 (i.e. no single day is holding it up), and 27 of 400 same-sized random removals beat it — about a 7% chance of arising by luck, which is suggestive rather than conclusive. And the honest caveat

that matters most: **the whole effect is eight trades**. Eight. That is why this is a shadow arm and not a Saturday deploy.

**THE ANSWER TO THE GRIND QUESTION, IN ONE PLACE**

**(c) the ER-0.35 floor:** it does not exist, and tonight's book says good — as a filter it is **REFUTED** (kills 9 of the 10 best trades; strip-3 −\$1,044; beaten by 182 of 400 random subsets). **(b) entry trend-confirmation:** **REFUTED** — it is backwards; the losers have bigger impulses than the winners, and 391 of 400 random subsets beat it. **(a) trend-day-only:** **PARKED** — real but wobbly, and it drops the runners; revive it if a genuinely trend-heavy month gives the session-ER measure enough days to show a plateau instead of a wobble.

**What to do instead:** keep the  $ATR \geq 22$  floor, keep the router's on/off, and add the **counter-move veto** as a shadow arm **SHADOW**. And take seriously that the current arming test is so strict grind has been a dead gate for nine sessions — that may be the right call on a chop fortnight, but it is a decision being made by a threshold nobody chose for that purpose.

## PART C — A BROKEN INSTRUMENT UNDERNEATH THIS REPORT

One more lead, and it is the least glamorous but the most immediately actionable: **the nightly exit-selector rollup counts quarantined trades**, and this report reads that rollup to produce its router-value figure.

On 13 August the day rider's re-booking bug wrote four phantom trades that were flagged EXCLUDE:day\_rider\_phantom\_rebook\_20260813 and deliberately never deleted. The selector's own query (scripts/selector\_nightly.py) filters on exit\_reason but **never filters data\_quality** — the one filter this desk's P&L rules mandate. So it counted all five day-rider rows for that day, and because it reprices every trade off bars they all collapsed to the same synthetic row: five identical entries of chosen −\$31.50, regret \$72.50.

| 13 AUGUST SELECTOR FILE | AS PUBLISHED | DEDUPED | ERROR              |
|-------------------------|--------------|---------|--------------------|
| trades counted          |              | 14      | 10 +4 phantoms     |
| regret                  |              | \$748   | \$458 inflated 39% |

| 13 AUGUST SELECTOR FILE | AS PUBLISHED | DEDUPED | ERROR           |
|-------------------------|--------------|---------|-----------------|
| selector P&L            | −\$228       | −\$102  | overstated 2.2× |

I checked every one of the eleven selector files against the ledger: **13 August is the only contaminated one** (the two quarantined rows from 4 August fall outside that file's window), so the week's total regret of \$7,903 is overstated by \$290 rather than by anything dramatic. But **48% of that single day's regret is one trade counted five times**, and any selector tuning done off that day is chasing a phantom. Two other things I would fix while in there: the file set is gappy (27–31 July, 3–7 August, 13 August) because there is no systemd timer for it at all despite the name "nightly", and its only consumer is this report.

Related and worth stating as a rule rather than an incident, because it recurs: a shadow variant is only as trustworthy as its pure-reconstruction column. The two-ratchet ledger ships recon\_pure\_usd, so its stop-leak could be measured and turned out to be one trade; the partial ledger ships none, so its "exits earlier" delta cannot be falsified at all. The grind-exit ledgers themselves belong to the shadow section and I have not re-run them here.

## DISPOSITION — every lead touched tonight

| LEAD                                                                                    | VERDICT | THE EVIDENCE, OR WHAT WOULD REVIVE IT                                                                                                                                                                                                                                                              |
|-----------------------------------------------------------------------------------------|---------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| MGC day rider — direction called +30/60 min into the US session, held to the 20:40 flat | SHADOW  | 15 of 15 grid cells positive (+\$552 to +\$1,742); beats the best constant (+\$1,520 / 56.2%) on both dollars and hit rate while trading half the days; median MFE +\$453. Thin at 16 sessions — ship from the middle of the plateau, never the peak cell.                                         |
| MGC coil bouncer — fade the edge of a compressed 20-min range                           | SHADOW  | n=154, median trade +\$7.75, survives strip-best-day (+\$289 of +\$470), beaten 0/12 by a same-count placebo, mirror control −\$924. About \$137/wk — below the bar for live, ideal for incubation.                                                                                                |
| MGC run-catcher exit — wide stop held for hours instead of a 2.5×ATR trail              | PARKED  | Holds inside the oracle in both windows (+\$143→+\$589 on 7–14 Aug), but every ALL-signals cell in that window loses, and shifting the window by two days flips the hold-cap conclusion entirely. Revive when a causal day filter exists, or at $n \geq 40$ gold sessions — whichever comes first. |
| MGC London-open break fade                                                              | PARKED  | Beaten 0/7 by shifted anchors, so the anchor is real — but 9 of 16 trades lose, the median trade is a loss, and strip-best-day takes +\$502 to +\$125. Revive at $n \geq 40$ trades, i.e. roughly two more months of London opens.                                                                 |
| grind counter-move veto — skip a long fighting a 30-min fall > 1×ATR                    | SHADOW  | +\$2,407 → +\$3,134 over 133 tick-honest fires; keeps 10/10 of the best trades; plateau across the whole −2.0 to −0.5×ATR range; both date-halves improve; beaten                                                                                                                                  |

| LEAD                                                              | VERDICT         | THE EVIDENCE, OR WHAT WOULD REVIVE IT                                                                                                                                                                                                                                                                                                    |
|-------------------------------------------------------------------|-----------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                                                                   |                 | by 27/400 placebos. But the effect is 8 trades — incubate, do not deploy.                                                                                                                                                                                                                                                                |
| grind ER-0.35 floor as an enable mechanism                        | REFUTED         | Named tests: fake-filter guard (keeps 1 of the 10 best trades), strip-3 (−\$1,044), placebo (beaten by 182/400 same-size random removals). Also does not exist in production — ER_FL00R is empty.                                                                                                                                        |
| grind entry trend-confirmation (require a 15-min impulse our way) | REFUTED         | Named tests: placebo (beaten by 391/400), strip-3 (−\$1,862). Mechanism fails for a reason, not by luck — the losing entries have larger 15-minute impulses than the winning ones.                                                                                                                                                       |
| grind trend-day-only arming (session efficiency)                  | PARKED          | Best cells beat base (+\$2,780 / +\$2,630) but the curve wobbles rather than plateaus and the tight cells drop the runners. Revive on a trend-heavy month, or if the session-ER measure can be replaced with one that keeps 8+ of the top 10.                                                                                            |
| the break trio as grind's arming test                             | MEASUREMENT GAP | Admits ~4 of 133 signals in replication; live grind has taken zero trades in nine sessions; and on 13 Aug it armed 17 minutes after the day's only paying signal. Cannot be graded properly until the watcher logs its own ER/ATR/extreme reads in a machine-readable form so the test can be replayed exactly rather than approximated. |
| selector_nightly.py ignores data_quality                          | FIX ON SATURDAY | 13 Aug regret inflated 39% (\$748 → \$458) by one quarantined trade counted five times; feeds this report's router-value figure. One-line filter, then regenerate the historic files. Also: it has no timer despite being called "nightly".                                                                                              |

**The one stone still unturned**, and it is the same stone on both instruments: a causal day classifier. Gold's run-catcher and MNQ's grind both live or die on a minority of days, and both currently rely on an arming test that either fires too rarely (the break trio) or requires foresight (the oracle). Everything else in this section is either incubating or answered.

## **P2.6** THE ROUTER OPERATION REVIEW — week of 10-14 August 2026

Gaz, the router had a quiet, disciplined, slightly profitable week, and then I found two things underneath it that matter more than the week itself. **The decisions it actually made were worth +\$51.50** — its benches blocked \$611 of losers and gave up \$410 of winners (+\$201), its armings paid



\$716.50 and cost \$866 (−\$149.50). That is a rounding error, and on a week where the tournament book closed **−\$325.00 over 108 trades** a rounding error is not a bad result. The nightly rollup, which prices what the benched gates would have done, is more generous: **+\$400 net over the four days it can actually see.**

**The first thing underneath.** The router is managing an empty room. Its managed universe is `capitulation_long` on the down side and `rgv_short` on the up side. Those two gates took **2 trades between them all week.** Meanwhile `exhaustion_short` — which is an up-fader, exactly the shape the router exists to bench — took **71 of the week's 108 trades**, and **12 of them fired straight into a confirmed TREND\_UP for −\$241.00, at −\$20.08 a trade.** The nightly leakage detector reported **zero leakage every single day**, because it is only allowed to look at the two gates that never traded. That is not a small number and it is not a one-week fluke: over 30 days the same gate has 26 counter-trend fires for −\$398.50, negative on **7 of the 8 days it happened**, and it survives both strip-the-worst-day and leave-one-day-out. **That is the Saturday play, and it is the best-evidenced thing in this section.**

**The second thing underneath.** I do not fully trust the ruler the nightly uses, and I can prove it with the desk's own trades. The nightly prices every blocked signal under a plain scalp-2R / 1-ATR exit. So I took the week's **108 real live trades** — same entries, same tape — and repriced them under that identical exit. They go from **−\$325.00 to +\$459.50**. The ruler is \$7.26 a trade more generous than what the desk actually achieves, which means every “wins we missed” figure in the nightly is inflated. But — and this is why I segmented it before saying anything — **the gap is almost exactly zero in chop (−\$0.28 a trade over 65 trades) and huge in trend (+\$25.02 in TREND\_UP, +\$11.38 in TREND\_DOWN).** The router benches in trend. So the over-benching it gets accused of is measured with a ruler that is at its most generous precisely where the router does its work.

**SURVIVED:** the durable router as a mechanism. 5-minute systemd tick ran **288/288 on three days and 284/284 on the fourth, zero service failures; 95% of ticks changed nothing** and Friday changed almost nothing at all (6 transitions in 24 hours) on the day the desk made its second-best number. The **untradeable meter genuinely works** over 22 days — STAY-OUT days averaged −\$651, CAUTION days −\$212, TRADEABLE days +\$7, cleanly monotone — and the 65 cutoff sits in the middle of a flat plateau from 50 to 75, so it does **not** want tuning. **THE ACTION:** put `exhaustion_short` into the router's UP-side managed set. 26 counter-trend fires over 30 days, −\$398.50, −\$15.33 a trade, red on 7 of 8 days, survives strip-the-worst. **THE SECOND ACTION:** the 40-day sweep says leave ER 0.15 / net 30 exactly where they are and change the timing — window 30→45 min, step 15→5 min, hold 2→3 marks. That is worth +\$502 on the live set and +\$399 on the full set, and it independently arrives at the two fixes the ledger identified by hand: a longer window kills the boundary flicker, a shorter step cuts the actuation lag. **WHAT ERRED: 62% of the week's 64 gate transitions were dead** — no trade, no blocked signal, nothing. `abs_veto_short` alone was flipped **36 times and 28 of those flips produced literally nothing**. Actuation lag cost −\$242.00, all of it on the bench side. **THE HONEST GAP:** on Thursday the router was **brain-dead for 10 hours 20 minutes** — 125 consecutive ticks aborted on an expired OAuth session while systemd reported every one of them as a success. It cost \$0, and that was luck, not design.

## 1 · THE WEEK'S DECISIONS AND THE SCOREBOARD

Here is the week, day by day. The tournament book and the day-rider are kept apart because one day-rider trade swings more than twenty tournament scalps — blending them answers neither question.

| PARIS DAY | TOURNAMENT \$ | TRADES | WIN% | DAY-RIDER \$ | GATE TRANSITIONS | OF WHICH DEAD | WHAT THE ROUTER ACTUALLY DID                                                |
|-----------|---------------|--------|------|--------------|------------------|---------------|-----------------------------------------------------------------------------|
| Mon 08-10 | −244.50       | 18     | 33%  | +39.00       | 9*               | n/a           | Two stay-outs triggered and both lapsed. The lapsed-bench bug fired for the |

| PARIS DAY | TOURNAMENT \$ | TRADES | WIN%  | DAY-RIDER \$ | GATE TRANSITIONS | OF WHICH DEAD | WHAT THE ROUTER ACTUALLY DID                                                                                                                                     |
|-----------|---------------|--------|-------|--------------|------------------|---------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|           |               |        |       |              |                  |               | third time in a day.                                                                                                                                             |
| Tue 08-11 | +371.50       | 14     | 79%   | +189.00      | 21               | 14            | Busiest day. Benched grind_long at 13:40 for 48 hours — the week’s single best call. Also the operator carve-out on abs_veto_long, honoured and expired on time. |
| Wed 08-12 | −660.00       | 30     | 23%   | +224.50      | 19               | 12            | The bad day. Seven flips on one gate decided by three points of net. exhaustion_short lost — \$367.50 on 24 lots with 17 stops and no rule could touch it.       |
| Thu 08-13 | +37.50        | 9      | 44%   | +731.00      | 18               | 9             | Router dead 02:40-13:00 on expired credentials. Recovered and had its best afternoon window of the week.                                                         |
| Fri 08-14 | +170.50       | 37     | 49%   | +434.00      | 6                | 5             | Near-total restraint. 100 of 102 ticks logged “changed: none” and the desk made money by being left alone.                                                       |
| WEEK      | −325.00       | 108    | 42.6% | +1,617.50    | 64               | 40 (62%)      | All-in across both desks: + \$1,292.50.                                                                                                                          |

\* Monday’s transition count is from the router’s own log only. The systemd journal does not reach back past 2026-08-10 23:11 UTC, so Monday cannot be verified the second way. Everything else in this section that depends on the journal starts Tuesday.

## COUNTING THE SWITCH CHANGES — AND WHY IT HAS TO BE DONE THREE WAYS

The ledger learned this the hard way on Thursday morning and it is worth restating, because a churn metric read off the wrong instrument is worse than no churn metric. **The router's own log under-reports:** its changed: field only records what the router changed, so the 22:00Z reactivation timer moves gates while every router tick that hour says "changed: none". **The systemd journal over-reports:** tournament.py initialises its last-known set to empty, so the first tick after any restart logs a phantom change. Here is the reconciliation.

| PARIS DAY      | ROUTER LOG SAYS | JOURNAL RAW | RESTART ARTIFACTS | 22:00Z TIMER | TRUE SWITCH EVENTS | GATE TRANSITIONS |
|----------------|-----------------|-------------|-------------------|--------------|--------------------|------------------|
| Tue 08-11      |                 | 21          | 22                | 1            | 0                  | 21               |
| Wed 08-12      |                 | 16          | 19                | 2            | 1                  | 19               |
| Thu 08-13      |                 | 12          | 15                | 1            | 1                  | 18               |
| Fri 08-14      |                 | 4           | 5                 | 1            | 1                  | 6                |
| <b>Tue-Fri</b> |                 | <b>53</b>   | <b>61</b>         | <b>5</b>     | <b>3</b>           | <b>64</b>        |

The three instruments disagree by up to 25% on any given day and they agree within one event on the week. **The arbiter is the modification time on gate\_switches.env** — the file is what the tournament actually reads. Fifty-six real switch events, sixty-four individual gate flips, over five days.

## THE THRASH CHECK — AND IT FAILS ON ONE GATE

Sixty-four flips sounds like a lot. Spread across 1,148 ticks it is not: **the router changed nothing on 95% of its ticks**. The rate per day was 0.88, 0.79, 0.75 and 0.25 changes an hour — well under the 5-an-hour that got flagged on 08-07. **But the rate is not the problem this week. The concentration is.**

| GATE              | TRANSITIONS | BENCHES   | ARMS      | DEAD (\$0 EITHER WAY) | DEAD%      | BLOCKED SIGNALS | LIVE TRADES |
|-------------------|-------------|-----------|-----------|-----------------------|------------|-----------------|-------------|
| abs_veto_short    | 36          | 18        | 18        | 28                    | 78%        | 3               | 20          |
| abs_veto_long     | 9           | 5         | 4         | 2                     | 22%        | 10              | 7           |
| exhaustion_short  | 7           | 4         | 3         | 2                     | 29%        | 15              | 55          |
| grind_long        | 6           | 3         | 3         | 4                     | 67%        | 14              | 0           |
| capitulation_long | 6           | 3         | 3         | 4                     | 67%        | 1               | 2           |
| <b>TOTAL</b>      | <b>64</b>   | <b>33</b> | <b>31</b> | <b>40</b>             | <b>62%</b> | <b>43</b>       | <b>84</b>   |

**abs\_veto\_short** accounts for **56% of every flip the router made this week, and 78% of its own flips produced nothing whatsoever** — no trade taken, no signal blocked. Eighteen benches and eighteen arms, and twenty-eight of those thirty-six were pure motion. Sixteen of them lasted under thirty minutes. This is the same disease the 08-07 review named — churn is invisible at \$0 because each half of a round trip grades NEUTRAL in isolation — and it is now concentrated in one gate rather than spread across three.

The cause is not judgement. It is the threshold. The direction rule requires 60-minute net above 40pt and ER at or above 0.20, and this week’s tape sat on top of both bars for hours at a time. Wednesday morning is the clean example, straight from the log: UP +63pt/ER 0.385 → FLAT +38pt/ER 0.222 (net leg fails by 2 points) → UP +60pt/ER 0.324 → FLAT +39pt/ER 0.252 (fails by 1 point) → UP +42pt/ER 0.268 → FLAT +5pt. **Five flips in fifty minutes, and three of them turned on a single point of net.** Each individual read was correct.

THE GRADED SCOREBOARD

The bad-call ledger graded 61 items across the week’s 25 review windows.

| VERDICT   | THIS WEEK | ALL-TIME RUNNING | READING                                                                                                                 |
|-----------|-----------|------------------|-------------------------------------------------------------------------------------------------------------------------|
| ☐ GOOD    | 11        | 41               | 18% of the week’s grades against 32% all-time — a quieter week with fewer chances to be right.                          |
| ☐ BAD     | 2         | 10               | Both on Monday, both the same lapsed-bench bug, one of them —\$92.50.                                                   |
| △ CHURN   | 15        | 22               | ★ The week booked 15 of the 22 CHURN grades this ledger has ever issued. That is the honest headline of the scoreboard. |
| ☐ NEUTRAL | 32        | 36               | Mostly \$0 flips on gates that never got a setup — availability, not judgement.                                         |

Read plainly: **the router got materially better at not being wrong and materially worse at not being busy.** Two bad calls in five days is a good ratio. Fifteen churn grades in five days, against seven in all prior history, is not.

2 · WHAT PAID AND WHAT ERRED, IN DOLLARS

Every one of the 64 flips is graded here on its own terms. A **bench** is judged on what it blocked — each SUPPRESSED-OPEN in the benched window, repriced tick-honest at \$1.50 a round trip and \$2 a point. A negative blocked number means it

stopped a loser, which is the point. An **arm** is judged on the gate’s real live P&L over the window it was armed.

THE BENCHES

| OUTCOME                          | N  | SIGNALS | \$            | WHAT IT MEANS                                      |
|----------------------------------|----|---------|---------------|----------------------------------------------------|
| GOOD — blocked net losers        | 8  | 27      | saved +611.00 | The bench earned its keep.                         |
| OVER-BENCH — blocked net winners | 4  | 16      | cost −410.00  | Money left on the table.                           |
| FREE — nothing wanted in         | 21 | 0       | 0.00          | 64% of all benches cost and saved exactly nothing. |
| NET BENCH VALUE                  | 33 | 43      | +201.00       |                                                    |

| WHEN (UTC)  | GATE              | WINDOW | SIGNALS | BLOCKED \$ | VERDICT                                                                                                                             |
|-------------|-------------------|--------|---------|------------|-------------------------------------------------------------------------------------------------------------------------------------|
| 08-11 13:40 | grind_long        | 48.3h  | 12      | −247.0     | The week’s best call. Benchd into the US pre-open on a confirmed DOWN segment (−150pt / ER 0.381) and left alone for two full days. |
| 08-13 15:35 | grind_long        | 31.4h  | 2       | −81.0      | Same gate, same treatment, same result.                                                                                             |
| 08-11 13:40 | abs_veto_long     | 0.5h   | 1       | −65.5      | Half an hour, one signal, a loser stopped.                                                                                          |
| 08-13 15:35 | abs_veto_long     | 31.4h  | 2       | −57.0      |                                                                                                                                     |
| 08-13 13:15 | abs_veto_short    | 1.8h   | 2       | −50.0      | The only one of eighteen abs_veto_short benches that blocked anything at all.                                                       |
| 08-12 13:55 | abs_veto_long     | 23.7h  | 4       | −40.0      |                                                                                                                                     |
| 08-11 16:00 | abs_veto_long     | 20.6h  | 3       | −37.0      | The operator carve-out expiring on schedule — and it was right to expire.                                                           |
| 08-12 23:10 | capitulation_long | 22.8h  | 1       | −33.5      |                                                                                                                                     |
| 08-12 16:50 | exhaustion_short  | 5.2h   | 4       | +21.5      | Small over-bench.                                                                                                                   |
| 08-12 13:15 | abs_veto_short    | 0.4h   | 1       | +49.5      | Benchd for 25 minutes and the one signal in that window would have won.                                                             |
| 08-12 23:10 | exhaustion_short  | 22.8h  | 6       | +156.0     |                                                                                                                                     |
| 08-11 15:00 | exhaustion_short  | 7.0h   | 5       | +183.0     | The week’s worst bench. The 15:00Z stay-out fired on a genuinely untradeable read (239pt range for −6pt net) and cost               |

| WHEN (UTC) | GATE | WINDOW | SIGNALS | BLOCKED \$ | VERDICT                                       |
|------------|------|--------|---------|------------|-----------------------------------------------|
|            |      |        |         |            | seven hours of the one gate that was working. |

**Note what is missing from that table: twenty-one benches.** Nothing wanted to open during them, so they are neither good nor bad — they were free. The router’s standing asymmetry says a wrong bench costs opportunity while a wrong arm costs real money, so when unsure, sit benched. That policy is why two-thirds of its benches are free, and it is the right policy. But it also means the bench half of the scoreboard is being judged on 12 events, not 33.

### THE ARMINGS

| OUTCOME       | N | LIVE TRADES | \$ | WHAT IT MEANS |
|---------------|---|-------------|----|---------------|
| PAID          |   | 6           | 46 | +716.50       |
| ERRED         |   | 6           | 38 | −866.00       |
| NEVER FIRED   |   | 19          | 0  | 0.00          |
| NET ARM VALUE |   | 31          | 84 | −149.50       |

| WHEN (UTC)  | GATE             | WINDOW | TRADES | \$      | READING                                                                                                                |
|-------------|------------------|--------|--------|---------|------------------------------------------------------------------------------------------------------------------------|
| 08-11 22:00 | exhaustion_short | 18.8h  | 24     | −367.50 | The 22:00Z reactivation timer, not the router. It armed the gate into Wednesday and Wednesday is where the money went. |
| 08-13 22:00 | exhaustion_short | 21.8h  | 31     | +310.00 | The same timer, the same gate, two days later — and it paid. That is the honest counterweight.                         |
| 08-13 17:50 | abs_veto_short   | 26.2h  | 6      | −139.50 | A restore-on-lapse: the direction bench stopped applying, so the gate came back. Nothing argued for it.                |
| 08-12 12:35 | abs_veto_long    | 1.3h   | 2      | −133.00 | Armed on a confirmed UP read that lasted 20 minutes.                                                                   |
| 08-11 11:25 | abs_veto_short   | 3.5h   | 3      | +131.50 |                                                                                                                        |
| 08-12 13:40 | abs_veto_short   | 3.2h   | 2      | −111.50 |                                                                                                                        |
| 08-11 06:10 | abs_veto_long    | 7.5h   | 2      | +98.00  | Two TARGET wins. Quietly the cleanest arm of the week.                                                                 |
| 08-13 13:35 | abs_veto_long    | 2.0h   | 3      | +95.00  | Armed on the 29922 ceiling breaking. Two targets and one regime-normal stop.                                           |
| 08-11 15:05 | abs_veto_short   | 4.5h   | 3      | +73.00  |                                                                                                                        |



| WHEN (UTC)  | GATE              | WINDOW | TRADES | \$     | READING                         |
|-------------|-------------------|--------|--------|--------|---------------------------------|
| 08-13 16:10 | abs_veto_short    | 1.1h   | 2      | −66.50 |                                 |
| 08-11 22:00 | capitulation_long | 21.8h  | 2      | −48.00 | Timer again. Both lots stopped. |
| 08-13 15:05 | abs_veto_short    | 0.9h   | 4      | +9.00  |                                 |

The single most important line in that table is that the two biggest armings were not the router’s. They were the 22:00Z reactivation timer, which brings gates back on schedule regardless of what anyone thinks. It armed exhaustion\_short into Wednesday for −\$367.50 and into Friday for +\$310.00. Net across the week the timer is roughly a wash — but it is the largest single lever on this desk and nothing grades it, because it is not a decision, it is a clock.

### 3 • NET ROUTER VALUE — THE NIGHTLY ROLLUP

Every data/router\_nightly/\*.json and data/selector\_nightly/\*.json for the week. Two things to say before the numbers. **First, none of these files existed for this week when I started** — the last router nightly on disk was 2026-08-07. I backfilled all five days. **Second, Monday is a measurement gap, not a zero.** The nightly reads blocked signals out of the systemd journal, and the journal does not reach back past 2026-08-10 23:11 UTC. Monday’s file reads \$0 / \$0 / \$0 because there was nothing to read. It is not a quiet day; it is a blind spot, and I am reporting it as one.

| PARIS DAY             | CHOP% | TREND_UP % | TREND_DO WN% | LOSSES SAVED | WINS MISSED | NET ROUTER VALUE | LEAKAGE REPORTED |
|-----------------------|-------|------------|--------------|--------------|-------------|------------------|------------------|
| Mon 08-10             | 69.3  | 10.2       | 20.5         | —            | —           | —                | —                |
| Tue 08-11             | 69.1  | 11.3       | 19.6         | +1,123       | 967         | +156             | 0                |
| Wed 08-12             | 91.8  | 5.2        | 3.1          | +1,192       | 666         | +526             | 0                |
| Thu 08-13             | 82.5  | 13.4       | 4.1          | +588         | 1,064       | −476             | 0                |
| Fri 08-14             | 92.8  | 0.0        | 7.2          | +484         | 290         | +194             | 0                |
| WEEK (4 visible days) | 84.1  | 7.5        | 8.5          | +3,387       | 2,987       | +400             | 0                |

Net positive, but thinly so, and one bad Thursday nearly wipes it out. Three days green, one red. And note the regime column — **this week was 84% chop**. A router built to bench counter-trend faders has very little to do on a tape that is barely trending, which is exactly why the week’s numbers are small in both directions.

## WHERE THE BENCHING LANDS, BY GATE

| GATE              | BLOCKED SIGNALS | WOULD-WIN | \$   | WOULD-LOSE | \$   | NET BLOCKED (2-LOT) | VERDICT                                                                                                                                   |
|-------------------|-----------------|-----------|------|------------|------|---------------------|-------------------------------------------------------------------------------------------------------------------------------------------|
| grind_long        | 14              | 3         | +234 | 11         | −562 | −656                | Best-blocked gate on the desk. It has an ATR floor of 22pt and the tape ran 2.6-4.4 — it could not have fired anyway on most days.        |
| nipc_short        | 12              | 2         | +137 | 10         | −345 | −416                | Held under the −\$400 kill criterion. Correctly.                                                                                          |
| abs_veto_long     | 11              | 3         | +135 | 8          | −290 | −310                | ★ The one long gate with no ATR floor, so it is the one that actually had signals to suppress. Its bench was not free and it priced well. |
| nipc_long         | 11              | 4         | +240 | 7          | −243 | −7                  | A coin-flip.                                                                                                                              |
| abs_veto_short    | 3               | 1         | +50  | 2          | −50  | −1                  | Thirty-six flips produced three blocked signals worth one dollar.                                                                         |
| capitulation_long | 6               | 3         | +124 | 3          | −88  | +70                 | Mild over-bench.                                                                                                                          |
| rgv_short         | 1               | 1         | +100 | 0          | 0    | +199                | n=1. Ignore it.                                                                                                                           |
| exhaustion_short  | 15              | 10        | +476 | 5          | −116 | +721                | The single biggest over-bench figure of the week — and the one I do not believe. See below.                                               |

## ★ THE RULER PROBLEM — AND WHY I RE-MEASURED IT BEFORE BELIEVING THE OVER-BENCH

That exhaustion\_short line says the router benched \$721 of winners. But exhaustion\_short was the desk's worst gate this week — 71 live trades for −\$218.50, minus three dollars a trade. Both things cannot be true of the same gate on the same tape unless the difference is the **exit**. The nightly prices every blocked signal under a plain scalp-2R with a 1-ATR stop. The live desk does not use that exit.

So I ran the control. Same 108 live trades, same entries, same tape, repriced under the nightly’s exit.

| GATE                  | N   | LIVE EXIT<br>\$ | \$/TRADE | SCALP-2R<br>\$ | \$/TRADE | DIFFEREN<br>CE | SCALP<br>BETTER<br>ON |
|-----------------------|-----|-----------------|----------|----------------|----------|----------------|-----------------------|
| exhaustion_sho<br>rt  | 71  | −218.50         | −3.08    | +177.00        | +2.49    | +395.50        | 42/71                 |
| abs_veto_long         | 9   | +6.00           | +0.67    | +203.50        | +22.61   | +197.50        | 5/9                   |
| abs_veto_short        | 26  | −64.50          | −2.48    | +128.00        | +4.92    | +192.50        | 20/26                 |
| capitulation_l<br>ong | 2   | −48.00          | −24.00   | −49.00         | −24.50   | −1.00          | 0/2                   |
| TOTAL                 | 108 | −325.00         | −3.01    | +459.50        | +4.25    | +784.50        | 67/108                |

**The nightly’s ruler is \$7.26 a trade more generous than the desk.** Every “wins we missed” number in this section is measured with it, so every one of them is too big — which means the router’s true net value is higher than the +\$400 reported. But before recommending anything off that, I segmented it, because a blanket number across mixed regimes is the exact bug the standing backtest rule exists to prevent.

| SEGMENT AT<br>ENTRY                                                              | N  | LIVE \$/TRADE | SCALP-2R \$/<br>TRADE | GAP PER<br>TRADE | READING                                                               |
|----------------------------------------------------------------------------------|----|---------------|-----------------------|------------------|-----------------------------------------------------------------------|
| CHOP (router label)                                                              | 65 | +0.65         | +0.36                 | −0.28            | <b>Zero gap. In chop our live exits are as good as the ruler.</b>     |
| TREND_UP                                                                         | 23 | −13.98        | +11.04                | +25.02           | <b>The whole gap lives here.</b>                                      |
| TREND_DOWN                                                                       | 20 | −2.27         | +9.10                 | +11.38           |                                                                       |
| and the same cut on ATR/ER rather than the router’s own label, as a cross-check: |    |               |                       |                  |                                                                       |
| violent-whipsaw<br>(ER<0.20, ATR≥12)                                             | 47 | −10.40        | +1.99                 | +12.39           |                                                                       |
| clean-trend (ER≥0.35)                                                            | 22 | +9.98         | +22.95                | +12.98           | <b>Live makes money here — just less than half what was on offer.</b> |
| normal-chop (ER<0.20,<br>ATR 6–12)                                               | 25 | +4.72         | +2.10                 | −2.62            | <b>Live BEATS the ruler here.</b>                                     |
| in-between (ER 0.20–<br>0.35)                                                    | 14 | −12.39        | −13.68                | −1.29            | <b>Both lose. Nobody’s regime.</b>                                    |

**So the honest statement is narrow and useful: our live exits are fine in ordinary chop and leak badly once the tape moves.** That is an exit finding, not a router finding, and I am handing it to the rehab section rather than acting on it here. What it means for the router is precise: **the wins-missed figure is inflated exactly in trend, which is the only regime where the router does**

**anything.** The \$721 exhaustion\_short over-bench is mostly the ruler. **Do not tighten the router on the strength of it.**

★★ THE LEAKAGE THE DETECTOR CANNOT SEE

Every nightly this week reported **leakage\_n = 0**. Four days, zero leaked losers. That reads as a router with nothing getting past it. It is not. **The detector is only allowed to look at capitulation\_long and rgv\_short — the two gates in its managed set — and those two took two trades between them all week.** It is reporting zero because it is looking at an empty room.

Here is every live trade this week, laid out against the router’s own regime label at the moment of entry.

| GATE              | CHOP N / \$ | TREND_UP N / \$ | TREND_DOWN N / \$ | ROUTER-MANAGED?                                       |
|-------------------|-------------|-----------------|-------------------|-------------------------------------------------------|
| abs_veto_long     | 4 / +44.00  | 5 / -38.00      | 0 / 0.00          | no (momentum, not a fader)                            |
| abs_veto_short    | 14 / +52.00 | 4 / +5.50       | 8 / -122.00       | no (momentum, not a fader)                            |
| capitulation_long | 0 / 0.00    | 2 / -48.00      | 0 / 0.00          | YES — and it took 2 trades all week                   |
| exhaustion_short  | 47 / -54.00 | 12 / -241.00    | 12 / +76.50       | NO — and it is an UP-fader that took 71 of 108 trades |

Read the exhaustion\_short row across. **Aligned with a down-trend it made + \$76.50. In chop it lost a trivial -\$54.00. Fired into a confirmed up-trend — fading a move that was proven to be going the other way — it lost -\$241.00 on twelve trades, twenty dollars a trade.** That is textbook counter-trend fader bleed, it is precisely what this router was built to prevent, and the router is not allowed to touch it.

Then I checked whether that is a one-week accident, across the full 30 days of tape.

| FADER             | COUNTER-TREND FIRES | \$      | \$/TRADE | ALIGNED FIRES | \$/TRADE | CHOP FIRES | \$      | MANAGE D TODAY?     |
|-------------------|---------------------|---------|----------|---------------|----------|------------|---------|---------------------|
| rgv_long          | 36                  | -698.50 | -19.40   | 2             | +8.00    | 22         | -8.50   | no (dormant gate)   |
| exhaustion_s hort | 26                  | -398.50 | -15.33   | 48            | -6.98    | 72         | +575.50 | no — and it is LIVE |
| rgv_short         | 12                  | -380.00 | -31.67   | 6             | +30.17   | 25         | -217.00 | yes (dormant gate)  |
| capitulation_long | 2                   | -112.00 | -56.00   | 9             | -14.22   | 21         | +154.50 | yes                 |

| FADER    | COUNTER<br>-TREND<br>FIRES | \$        | \$/TRADE | ALIGNED<br>FIRES | \$/TRADE | CHOP<br>FIRES | \$      | MANAGE<br>D<br>TODAY? |
|----------|----------------------------|-----------|----------|------------------|----------|---------------|---------|-----------------------|
| ALL FOUR | 76                         | −1,589.00 | −20.91   | 65               |          | 140           | +504.50 |                       |

Every single fader on this desk loses money firing counter-trend, and all four lose by a similar amount — fifteen to fifty-six dollars a trade. Seventy-six fires, −\$1,589. The router currently manages fourteen of those seventy-six. The other sixty-two, worth −\$1,097, belong to two gates it is not allowed to see — and one of those two is the busiest live gate on the desk.

The robustness on the actionable one:

| TEST                                                                     | RESULT                           | READING                                                                                                                                 |
|--------------------------------------------------------------------------|----------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------|
| headline                                                                 | 26 fires, −\$398.50, −\$15.33/tr | over 8 separate days in 30 days of tape                                                                                                 |
| days where the bench would have helped                                   | 7 of 8                           | Not one day carrying it.                                                                                                                |
| strip-the-worst-day (drop 08-12, −\$166)                                 | −\$232.50                        | Still clearly negative — survives.                                                                                                      |
| leave-one-day-out, worst fold (drop the only green day, 07-30 +\$224.50) | −\$623.00                        | Gets stronger, not weaker.                                                                                                              |
| the honest caveat                                                        | n=26                             | Thin. This is a SHADOW-grade sample proposed for a LIVE change, and the change is a bench, which is the cheap direction to be wrong in. |

### THE SELECTOR, SAME ROLLUP

| PARIS DAY | N   | OPTIMAL<br>PICKS | REGRET \$ | SELECTOR<br>\$ | PERFECT<br>HINDSIGHT<br>T \$ | ALWAYS-<br>SCALP \$ | ALWAYS-<br>WIDE \$ |
|-----------|-----|------------------|-----------|----------------|------------------------------|---------------------|--------------------|
| Mon 08-10 | 19  | 79%              | 311       | −170           | +141                         | −296                | −375               |
| Tue 08-11 | 15  | 40%              | 544       | −79            | +466                         | +231                | +180               |
| Wed 08-12 | 31  | 68%              | 336       | −397           | −61                          | −158                | −660               |
| Thu 08-13 | 14  | 29%              | 748       | −228           | +520                         | +144                | −165               |
| Fri 08-14 | 38  | 61%              | 1,140     | −252           | +888                         | +220                | +125               |
| WEEK      | 117 | 59%              | 3,079     | −1,126         | +1,954                       | +141                | −895               |

The selector lost to a fixed always-SCALP rule on four of five days and by \$1,267 on the week. It beat fixed-WIDE, which is something, but the thing it is supposed to be — an adaptive chooser — is currently worse than picking one exit and never thinking again. And the wrong-pick pattern is completely systematic: on Wednesday all five of the worst picks were exhaustion\_short choosing TIGHT when SCALP or WIDE was best, and on Friday all five were the same gate doing the same thing. One gate, one mode, one direction of error. That lines up exactly with the exit finding above — the live exits are too tight for the tape once it moves.

### ⚠ THURSDAY'S SELECTOR NUMBERS ARE POLLUTED — THE STANDING BUG CONFIRMED AGAIN

Thursday shows 14 trades and \$748 of regret. **Five of those fourteen rows are the same day-rider phantom counted five times**, at \$72.50 of regret each — \$362.50, or **48% of the day's entire regret figure**, from one trade that did not exist. Thursday also carries \$1,551 of data\_quality-flagged rows that the P&L module correctly excludes and selector\_nightly.py does not. This is the known selector-nightly-ignores-data-quality defect, unfixed, and it is still inflating regret. Read Thursday's selector row as roughly \$385 of regret on 9 real trades, not \$748 on 14.

## 4 • THRESHOLD TUNING — THE 40-DAY SWEEP

scripts/router\_study.py --days 40, both universes, thresholds and timings. The shipped config is  $ER \geq 0.15$ ,  $|net| \geq 30pt$ , 30-minute window, 15-minute poll, 2-mark hysteresis, fast-exit off. The fast replay validated identical to the live direction\_router.replay\_marks before anything was swept, so this is the real router being measured, not an approximation of it.

### SET A — THE LIVE MANAGED SET (CAPITULATION\_LONG, RGV\_SHORT)

Baseline with no router at all:  $-\$302$ . Router as shipped: **+\$190, a swing of +\$492**, from blocking 14 trades (losers  $-\$630$ , winners  $+\$138$ ).

| ROUTER Δ VS NO-ROUTER | NET $\geq 20$ | NET $\geq 30$ | NET $\geq 40$ | NET $\geq 50$ | NET $\geq 60$ |
|-----------------------|---------------|---------------|---------------|---------------|---------------|
| ER $\geq 0.15$        | +461          | +492 ★shipped | +428          | +316          | +158          |
| ER $\geq 0.20$        | +224          | +176          | +152          | +152          | +152          |
| ER $\geq 0.25$        | +152          | +152          | +152          | +152          | +152          |
| ER $\geq 0.30$        | +152          | +152          | +152          | +152          | +152          |
| ER $\geq 0.35$        | +152          | +152          | +152          | +152          | +152          |

### SET B — THE FULL HISTORICAL FADER SET (+ RGV\_LONG, + EXHAUSTION\_SHORT)

Baseline  $-\$722$ . Router as shipped: **+\$584, a swing of +\$1,307**, blocking 57 trades (losers  $-\$1,994$ , winners  $+\$687$ ). Note this set is  $2.7\times$  the size of the live

one — because it includes the two faders the router is not currently allowed to manage.

| ROUTER Δ VS NO-ROUTER | NET≥20 | NET≥30          | NET≥40 | NET≥50 | NET≥60 |
|-----------------------|--------|-----------------|--------|--------|--------|
| ER≥0.15               | +1,300 | +1,307 ★shipped | +1,294 | +1,120 | +858   |
| ER≥0.20               | +1,044 | +995            | +976   | +1,014 | +816   |
| ER≥0.25               | +787   | +787            | +806   | +806   | +744   |
| ER≥0.30               | +334   | +334            | +334   | +334   | +428   |
| ER≥0.35               | +428   | +428            | +428   | +428   | +428   |

Both grids say the same three things. One: **ER 0.15 is not a peak, it is a cliff edge — and we are on the right side of it.** Everything at ER 0.20 and above is worth a third to a half as much, on both sets independently. Two: **the net threshold barely matters between 20 and 40** — on the full set those three cells are +\$1,300 / +\$1,307 / +\$1,294, thirteen dollars apart across the whole range. That is a genuine plateau and 30 sits in the middle of it. Three: **the shipped cell is the best cell on both sets.** Whoever set ER 0.15 / net 30 got it right and there is nothing to gain by moving either.

THE TIMING SWEEP — WHERE THE ACTUAL GAIN IS

900 configurations, thresholds crossed with timings. The joint winner is **the same config on both universes**, which is the first thing that made me take it seriously:

| KNOB            | SHIPPED | JOINT BEST | WHAT THE CHANGE DOES IN PLAIN TERMS                                                                                               |
|-----------------|---------|------------|-----------------------------------------------------------------------------------------------------------------------------------|
| ER threshold    |         | 0.15       | 0.15 unchanged                                                                                                                    |
| net threshold   |         | 30pt       | 20pt inside the flat plateau — a \$13 difference on the full set. I would not take this one.                                      |
| lookback window |         | 30 min     | 45 min A longer window is a smoother read. This is the direct fix for the boundary flicker that produced 36 abs_veto_short flips. |
| poll step       |         | 15 min     | 5 min Matches the actual systemd tick. The router already runs every 5 minutes; its regime replay pretends it runs every 15.      |
| hysteresis hold |         | 2 marks    | 3 marks 3×5min = 15 minutes of confirmation instead of 30. Faster to act, on a steadier signal.                                   |
| fast_exit       |         | off        | off unchanged — still dormant, as it has been since it backtested + \$0                                                           |

|                                           | LIVE MANAGED SET | FULL HISTORICAL SET |
|-------------------------------------------|------------------|---------------------|
| router value, shipped config              | +\$190           | +\$584              |
| <b>router value, joint-best config</b>    | <b>+\$692</b>    | <b>+\$984</b>       |
| <b>improvement</b>                        | <b>+\$502</b>    | <b>+\$399</b>       |
| trades blocked (shipped → best)           | 14 → 27          | 57 → 70             |
| days helped / days hurt                   | 4 / 0            | 6 / 1               |
| worst single day                          | \$+0             | −\$10               |
| <b>days out of 40 that changed at all</b> | <b>4</b>         | <b>7</b>            |

**Now the honesty.** That +\$502 rests on **four days out of forty**. Thirty-six days are completely unaffected. It never hurt a single day on the live set and hurt one day by ten dollars on the full set, which is as clean a robustness profile as you can ask for — but it is clean because it is rare, not because it is reliable. Four positive events is not a proven edge; it is four positive events.

#### THE MODERATE RECOMMENDATION — AND WHAT I DELIBERATELY LEFT OUT

**TAKE:** window 30→45 min, step 15→5 min, hold 2→3 marks. Three reasons, in order of weight. **(a)** It is the joint optimum on both universes independently — the live set and a set 2.7× larger with two different gates in it. **(b)** It never hurt a day on the live set and cost \$10 on the worst day of the full set. **(c)** Most importantly, it is not a number found by searching — it is the mechanical fix for two problems this section documents from live evidence and not from the grid: a 45-minute window smooths the exact one-point boundary flicker that generated 36 flips on one gate, and 5-minute steps close the actuation lag that cost −\$242 this week. **When a grid search and a ledger written by hand arrive at the same answer from opposite directions, that is the strongest signal available here.**

**DO NOT TAKE:** net 30→20. It is worth \$13 on the full set and \$31 on a live set of 14 trades. That is noise, it moves us to the edge of the plateau rather than the middle, and taking a free-looking knob change for no reason is how a config gets overfitted one dollar at a time. **Leave ER 0.15 and net 30 exactly where they are.**

#### △ SATURDAY, NOT NOW, AND WITH THE CAVEAT WRITTEN DOWN.

The market is shut, so a Saturday deploy carries no live risk and gets a full session to be watched from a clean start. And the caveat that must go in the commit message: **this was fitted on 40 days that are overwhelmingly ONE regime.** This week alone ran 84% CHOP. The router is a trend



instrument being tuned on a range tape. Re-run this study after the first genuine trend week and be ready to revert — the revert is a three-constant edit.

## IS IT BLOCKING ENOUGH? THE LEAK CLASSIFICATION

| EVERY<br>MANAGED-<br>FADER TRADE,<br>CLASSIFIED | LIVE SET N | \$     | FULL SET N | \$     | WHOSE<br>PROBLEM IS<br>IT?                                          |
|-------------------------------------------------|------------|--------|------------|--------|---------------------------------------------------------------------|
| BLOCKED — the router caught it                  | 14         | −492   | 57         | −1,307 | working as designed                                                 |
| allowed WINNER — correctly left alone           | 21         | +1,456 | 63         | +3,813 | the over-benching test: it is NOT strangling the book               |
| LEAK — counter-trend NEAR-MISS                  | 15         | −911   | 30         | −1,876 | the router's own problem — the threshold just failed to catch these |
| chop loss                                       | 4          | −208   | 7          | −355   | nothing a direction router can do                                   |
| trend-FOLLOWING loss                            | 1          | −147   | 16         | −997   | a gate or exit problem, not the router's                            |

**The single biggest bucket of remaining loss is the near-miss leak: 30 trades, −\$1,876 on the full set.** These are counter-trend fader fires that the threshold missed by a whisker. That is the argument for loosening, and the grid answers it directly — ER 0.15 / net 20 blocks 13 more trades on each set and captures part of it, which is where the joint-best config's gain actually comes from. It is also the argument I am declining, because the same grid says the money in that change is thirteen dollars on 40 days. **The bigger prize is not a lower threshold. It is pointing the existing threshold at exhaustion\_short.**

## 5 • THE UNTRADEABLE-DAY METER, REVIEWED

src/gazbot7/untradeable.py. Three meters — ROUND-TRIP (does the day go anywhere), GIVE-BACK (how much of its own move it hands back), STOP-RATE (what fraction of recent trades stopped) — blended into a 0–100 score. 65 and above is STAY-OUT, 45 and above CAUTION. **Five days is far too few to judge it, so I scored all 22 trading days the tape supports** and put each against the money.

| DATE         | SCORE     | VERDICT          | CHOP      | GIVE-BACK | STOP-RATE | ROUND-TRIP  | DAY \$           | N         |
|--------------|-----------|------------------|-----------|-----------|-----------|-------------|------------------|-----------|
| 07-16        | 29        | TRADEABLE        | 27        | 19        | 53        | 0.73        | −535.00          | 22        |
| 07-17        | 48        | CAUTION          | 45        | 44        | 60        | 0.55        | −152.00          | 27        |
| 07-20        | 94        | STAY-OUT         | 99        | 99        | 73        | 0.01        | −1,732.50        | 57        |
| <b>07-21</b> | <b>22</b> | <b>TRADEABLE</b> | <b>19</b> | <b>10</b> | <b>53</b> | <b>0.81</b> | <b>+52.50</b>    | <b>59</b> |
| 07-22        | 53        | CAUTION          | 55        | 50        | 53        | 0.45        | −329.50          | 23        |
| 07-23        | 46        | CAUTION          | 51        | 38        | 53        | 0.49        | −398.00          | 46        |
| 07-24        | 26        | TRADEABLE        | 26        | 19        | 40        | 0.74        | −578.50          | 47        |
| 07-27        | 47        | CAUTION          | 48        | 40        | 60        | 0.52        | −226.50          | 40        |
| 07-28        | 56        | CAUTION          | 62        | 62        | 33        | 0.38        | −54.00           | 23        |
| <b>07-29</b> | <b>23</b> | <b>TRADEABLE</b> | <b>25</b> | <b>5</b>  | <b>53</b> | <b>0.75</b> | <b>+638.50</b>   | <b>17</b> |
| <b>07-30</b> | <b>13</b> | <b>TRADEABLE</b> | <b>4</b>  | <b>4</b>  | <b>47</b> | <b>0.96</b> | <b>+1,172.00</b> | <b>91</b> |
| 07-31        | 86        | STAY-OUT         | 90        | 84        | 80        | 0.10        | −905.00          | 35        |
| 08-03        | 34        | TRADEABLE        | 49        | 10        | 53        | 0.51        | −460.50          | 32        |
| <b>08-04</b> | <b>22</b> | <b>TRADEABLE</b> | <b>23</b> | <b>15</b> | <b>33</b> | <b>0.77</b> | <b>+330.50</b>   | <b>16</b> |
| 08-05        | 44        | TRADEABLE        | 55        | 12        | 87        | 0.45        | −817.00          | 29        |
| 08-06        | 78        | STAY-OUT         | 86        | 81        | 58        | 0.14        | −137.50          | 10        |
| <b>08-07</b> | <b>15</b> | <b>TRADEABLE</b> | <b>16</b> | <b>5</b>  | <b>33</b> | <b>0.84</b> | <b>+227.00</b>   | <b>23</b> |
| 08-10        | 55        | CAUTION          | 62        | 36        | 80        | 0.38        | −244.50          | 18        |
| 08-11        | 50        | CAUTION          | 65        | 50        | 20        | 0.35        | +371.50          | 14        |
| 08-12        | 53        | CAUTION          | 49        | 46        | 73        | 0.51        | −660.00          | 30        |
| 08-13        | 21        | TRADEABLE        | 20        | 14        | 36        | 0.80        | +37.50           | 9         |
| 08-14        | 77        | STAY-OUT         | 84        | 75        | 67        | 0.16        | +170.50          | 37        |

## DOES IT DISCRIMINATE? YES — CLEANLY, AND MONOTONICALLY

| VERDICT                   | DAYS      | TOTAL \$      | \$ PER DAY   | TRADES     | \$ PER TRADE | GREEN DAYS     |
|---------------------------|-----------|---------------|--------------|------------|--------------|----------------|
| STAY-OUT (≥65)            | 4         | −2,604.50     | −651.12      | 139        | −18.74       | 1 of 4         |
| CAUTION (45-64)           | 8         | −1,693.00     | −211.62      | 221        | −7.66        | 1 of 8         |
| <b>TRADEABLE (&lt;45)</b> | <b>10</b> | <b>+67.00</b> | <b>+6.70</b> | <b>345</b> | <b>+0.19</b> | <b>6 of 10</b> |

**That is a real instrument.** Every step down the score is a step up in money, on every measure — dollars per day, dollars per trade, and share of green days. A STAY-OUT day costs nineteen dollars a trade; a TRADEABLE day breaks even. On the five days of this week alone the relationship looked like pure noise, which is exactly why five days was never going to answer the question.

## WHICH OF THE THREE METERS ACTUALLY DID THE WORK?

Each meter split at its own median across the 22 days, high half against low half:

| METER     | WEIGHT IN THE SCORE | HIGH HALF \$ PER DAY | LOW HALF \$ PER DAY | SEPARATION | VERDICT |
|-----------|---------------------|----------------------|---------------------|------------|---------|
| STOP-RATE |                     | −522.72              | +36.46              | \$559.18   |         |

| METER             | WEIGHT IN THE SCORE          | HIGH HALF \$ PER DAY | LOW HALF \$ PER DAY | SEPARATION | VERDICT                                                                 |
|-------------------|------------------------------|----------------------|---------------------|------------|-------------------------------------------------------------------------|
|                   | 20%, and excluded below n=10 |                      |                     |            | ★ The best discriminator of the three — and the one the module demotes. |
| ROUND-TRIP / chop | 40%                          | −407.60              | −12.88              | \$394.73   | Works. Correctly weighted heavily.                                      |
| GIVE-BACK         | 40%                          | −368.45              | −16.14              | \$352.32   | Works, slightly the weakest of the three.                               |

All three discriminate, which is the first real finding — none of them is dead weight. The second is more awkward: the **STOP-RATE** meter separates best and it is the one carrying the least weight. It gets 20% when it votes at all, and since the 2026-08-05 change it is excluded entirely below ten trades — which on this week's volumes means it was excluded on the days with the fewest trades, and the LOW half of its split is the only bucket in this entire table that is positive.

I am **not** recommending a re-weight, and the reason is the same n=22 discipline I applied everywhere else. The stop-rate meter is the only one of the three that is endogenous — it measures what the desk just did, not what the tape did. On a day the router has already benched everything, the stop-rate reads off a tiny sample of whatever was left armed. The 2026-08-05 change that demoted it was made for exactly that reason and it was right to make it. What this table justifies is **re-testing** the weight, not changing it. [SHADOW](#)

## DOES THE 65 CUTOFF WANT TUNING? — THE ACTIONABLE TEST

The full-day score is hindsight. The question that matters is the one the router faces live: **at 15:00Z, does the meter predict the money still to come?** The router only allows a stay-out past 15:00Z, so this is the exact decision. I swept the cutoff against post-15:00Z P&L over all 22 days.

| CUTOFF    | DAYS BENCHED | \$ AVOIDED (POST-15:00Z, ON THE BENCHED DAYS) | DAYS TRADED | \$ EARNED ON THEM |
|-----------|--------------|-----------------------------------------------|-------------|-------------------|
| 35        | 17 of 22     | +1,771.00                                     | 5           | +59.00            |
| 40        | 12           | +1,562.50                                     | 10          | −149.50           |
| 45        | 10           | +1,642.00                                     | 12          | −70.00            |
| <b>50</b> | <b>6</b>     | <b>+1,279.00</b>                              | <b>16</b>   | <b>−433.00</b>    |
| <b>55</b> | <b>6</b>     | <b>+1,279.00</b>                              | <b>16</b>   | <b>−433.00</b>    |

| CUTOFF       | DAYS BENCHED | \$ AVOIDED<br>(POST-15:00Z, ON<br>THE BENCHED<br>DAYS) | DAYS TRADED | \$ EARNED ON<br>THEM |
|--------------|--------------|--------------------------------------------------------|-------------|----------------------|
| 60           | 6            | +1,279.00                                              | 16          | -433.00              |
| 65 ★ shipped | 6            | +1,279.00                                              | 16          | -433.00              |
| 70           | 6            | +1,279.00                                              | 16          | -433.00              |
| 75           | 5            | +1,279.00                                              | 17          | -433.00              |
| 80           | 3            | +209.00                                                | 19          | -1,503.00            |
| 85-90        | 1            | +282.00                                                | 21          | -1,430.00            |

Baseline — trade every day after 15:00Z, no meter at all: **-\$1,712.00** over 22 days.

**The answer is no, and it is a clean no. Everything from 50 to 75 gives an identical result** — the same six days benched, the same \$1,279 avoided. Sixty-five sits in the dead centre of a 25-point plateau. That is what a well-set threshold looks like: you can be off by ten in either direction and nothing happens. **Leave it.**

The grid does show cutoff 35 avoiding more — \$1,771 — and leave-one-day-out picks 35 in 17 of 22 folds, which looks like a strong signal until you read what it is actually saying. **Cutoff 35 benches seventeen of twenty-two days.** That is not a stay-out rule, that is closing the afternoon desk. And it wins for a reason that has nothing to do with the meter: **trading after 15:00Z lost -\$1,712 across these 22 days full stop.** On a sample where the activity is a net loser, any rule that does less of it scores well. The LOO test is picking up the base rate, not the instrument. **I am refusing that recommendation and naming why.**

#### TWO HONEST HOLES IN THE METER, NEITHER OF WHICH CHANGES THE CUTOFF

**(1) It rests on one day.** Of the \$1,279 the 65 cutoff avoids, **\$1,009 — 79% — is 2026-07-20 alone.** Strip that single day and the rule is worth \$270 across the remaining five. The meter is real (the 22-day verdict table is monotone and that is not one day), but the dollar case for the cutoff is one day. Say it plainly and do not build on it.

**(2) The module and the router use different numbers.** `untradeable.py` defines STAY-OUT as  $\text{score} \geq 65$ . The router's live operating trigger, in its own words all week, is "two consecutive  $\geq 55$ , past the 15:00Z cutoff". Two thresholds, ten points apart, for one decision. Because 50-75 is a flat plateau **neither is wrong and this has cost nothing** — but it means the dashboard and the router can disagree about whether today is a stay-out

day, and nobody would notice. **Pick one number and put it in one place.**

**SATURDAY, 10-MINUTE FIX**

**This week's two misses, named.** Friday scored 77, **STAY-OUT** — **and the desk made +\$170.50**, its second-best day. Wednesday scored 53, **only CAUTION** — **and lost -\$660.00**, the worst day of the week. Two clean misses in five days, in opposite directions. Against the 22-day record I read those as variance rather than a broken instrument, but they are exactly the pair that would make you distrust it if this week were all you had.

## **6 · THE BENCH-DECISION REVIEW — grading the week's actual benches**

Thirty-three benches. Eight blocked money (+\$611), four gave money away (−\$410), twenty-one were free. Beyond the arithmetic, four patterns are worth naming because each one is a rule.

### **What made a bench GOOD**

**Duration.** Look at the two best calls — grind\_long benched at 08-11 13:40 for **48.3 hours** and again at 08-13 15:35 for **31.4 hours**. Between them: 14 blocked signals, −\$328. **Every single one of the eight good benches ran for at least half an hour and six of them ran for over 20 hours.** The good benches are the ones nobody touched. Compare the churn column: twelve benches lasted under thirty minutes and exactly **one** of those twelve blocked a single signal.

### **What made a bench BAD — and it is one gate, again**

**Three of the four over-benches were exhaustion\_short, and together they cost \$360.50.** The worst was the 08-11 15:00Z stay-out: the meter read 80 after 78, past the cutoff, on a genuinely awful day (239pt of range

for -6pt of net), and it was a correct call by the written rule — and it benched the one gate that was working for seven hours and gave up \$183. That is the meter and the gate disagreeing, and the meter winning. **Which is the right outcome to keep, because the same instrument is what stopped Wednesday being worse.** I am recording this as the cost of a rule I want kept, not as an error to fix.

## The bug that produced both ☐ BAD grades

Both bad calls this week are the same mechanism, both on Monday, and it has now fired **three times in one day: arming a gate because its bench lapsed.** The router's own words at 17:45Z — “the 17:15 stay-out that benched exhaustion\_short has now lapsed for three ticks, so I restore it.” Nothing made a case for the gate. The meter merely stopped making the case against it. **A lapsed bench returns a gate to benched-by-default, not to armed.** The gate was at that moment the day's worst performer at - \$137.50 with 4 of 6 lots stopped; it then took four more lots and stopped all four, for **-\$92.50**. Had the 17:15 bench simply been left alone the window is flat. The same shape produced the 08-13 17:50 restore of abs\_veto\_short for -\$139.50. **Across the week this one pattern is worth -\$232.00.**

## The bench that finally had a price — and it was positive

Two consecutive ledger entries this week assert that the long-side bench cost “exactly \$0, zero SUPPRESSED-OPEN”. Friday afternoon produced **240 of them**, all abs\_veto\_long, in two bursts at 15:46 and 16:01. The structural reason is worth carrying forward because it decides which benches can ever be free: deciders.py holds an ATR floor for grind\_long (22.0) and capitulation\_long (10.0) **and for nothing else.** On a tape running 2.6-12.5pt ATR that floor refused grind\_long 1,932 times all by itself — its bench was free because the gate could not have fired anyway.

**abs\_veto\_long** has no floor, so it was eligible the entire time, and its **bench** was the only long-side bench that could have a price. It priced GOOD: −\$310 blocked across the week, and roughly \$180 of that on Friday alone. **Never assume a bench is free — price it from SUPPRESSED-OPEN per window.**

## 7 · THE DURABLE ROUTER AND ITS ARCHITECTURE

### DID THE 5-MINUTE TICK MANAGE CORRECTLY?

| UTC DAY | TICKS FIRED | EXPECTED | SERVICE FAILURES | TICKS THAT MADE NO DECISION | %   | READING                                                      |
|---------|-------------|----------|------------------|-----------------------------|-----|--------------------------------------------------------------|
| 08-11   | 288         | 288      | 0                | 9                           | 3%  | The 9 are the scheduled CME maintenance-halt skips. Perfect. |
| 08-12   | 288         | 288      | 0                | 9                           | 3%  | Perfect.                                                     |
| 08-13   | 288         | 288      | 0                | 134                         | 47% | See below.                                                   |
| 08-14   | 284         | 284      | 0                | 9                           | 3%  | Perfect (partial day to 23:35Z).                             |

On the narrow question the operator asked — did the timer manage correctly, without thrash — the answer is an unqualified yes. 1,148 ticks, every single one fired on schedule, **zero service failures, zero timeouts, zero non-zero exits**. Median tick 16 seconds against a 280-second timeout. The plumbing is not the problem and has not been for weeks.

#### △△ AND THAT IS EXACTLY WHY THURSDAY IS THE WEEK’S MOST IMPORTANT FINDING

On 2026-08-13, from 02:40:03Z to 13:00:02Z — ten hours and twenty minutes — the router made no decisions at all. One hundred and twenty-five consecutive ticks aborted on the same line:

```
ABORT: no JSON in output -> no change | raw='Failed to
authenticate: 0Auth session expired
and could not be refreshed'
```

Every one of those 125 ticks was reported by systemd as a success. “Starting... Deactivated successfully... Finished... Consumed 2.5s CPU time.”

The service health was flawless. The service was doing nothing. A 2-second tick against a 16-second median is the only tell there was, and nothing was watching for it.

**What it cost: \$0 realised.** Zero trades fired in the whole window. The switches froze in the state the 02:30Z tick left them, and the only thing that state blocked was **one** `abs_veto_long` signal at 12:31:05Z, which reprices at  $-\$24.50$  a lot — so the frozen bench made  $\$49.00$  on the two-lot gate.

△ **That is luck and must be recorded as luck.** The outage sat almost entirely inside the Asia block, where new entries are refused whatever the switches say, and the rest of it was a dead London morning. Run the same ten-hour blackout starting at 13:00Z instead of 02:40Z and it covers the US open. **The correct reading is not “\$0, no harm” — it is “a ten-hour unmonitored failure of the desk’s number-one lever, which happened to land in the only ten hours of the day where it was survivable.”**

**ALREADY FIXED — NOT AN ACTION** `scripts/router_health_check.py` shipped the **same day** (commit 07cf6ef, “alarm the one thing nothing was watching”), with its timer first firing at 15:31:33Z — two and a half hours after the outage cleared. It alarms on an ABORT streak of 3 ( $\approx 15$  minutes of a frozen router) and on a silent log, pages once on entry and re-nags hourly. It has run every two minutes since. **Do not re-recommend this; it is done.** The open question it raises is the generalising one: **what else on this desk reports success while doing nothing?**

## THE ACTUATION-LAG COST, QUANTIFIED

The router decides on a 5-minute tick and needs two consecutive confirming marks, so a genuine regime change takes up to 30 minutes to reach `gate_switches.env`. Here is what that delay actually cost, measured as what happened in the ten minutes immediately before each flip took effect.

| WHEN (UTC)  | GATE                        | FLIP  | WHAT SLIPPED THROUGH | \$        | READING                                                    |
|-------------|-----------------------------|-------|----------------------|-----------|------------------------------------------------------------|
| 08-12 13:55 | <code>abs_veto_long</code>  | BENCH | 2 live trades        | $-133.00$ | Both taken in the last 10 minutes before the bench landed. |
| 08-13 15:58 | <code>abs_veto_short</code> | BENCH | 2 live trades        | $-101.00$ | The execution-pathology bench —                            |



| WHEN (UTC)                              | GATE             | FLIP  | WHAT SLIPPED THROUGH | \$      | READING                                          |
|-----------------------------------------|------------------|-------|----------------------|---------|--------------------------------------------------|
|                                         |                  |       |                      |         | the trades were already on by the time it fired. |
| 08-14 19:45                             | exhaustion_short | BENCH | 2 live trades        | −31.00  |                                                  |
| 08-11 14:10                             | abs_veto_long    | ARM   | 1 missed signal      | +65.50  | Being late DODGED a loser.                       |
| 08-12 12:35                             | abs_veto_long    | ARM   | 1 missed signal      | +55.00  | Dodged another.                                  |
| 08-12 13:40                             | abs_veto_short   | ARM   | 1 missed signal      | −49.50  | Missed a winner.                                 |
| 08-13 13:35                             | abs_veto_long    | ARM   | 1 missed signal      | −48.00  | Missed a winner.                                 |
| ★ TOTAL ACTUATION-LAG COST FOR THE WEEK |                  |       |                      | −242.00 | bench-late −\$265.00<br>· arm-late +\$23.00      |

**The lag is asymmetric and the asymmetry is exactly the one the router's own policy assumes.** Being late to BENCH cost real money every single time — three flips, six trades, −\$265, no exceptions. Being late to ARM was free: four flips, two missed winners worth \$97.50 and two dodged losers worth \$120.50, netting +\$23 in the desk's favour. **So the standing rule — when unsure, sit benched one more tick, and re-arm the moment evidence is real — is now measured, not assumed. It is worth about \$242 a week to make the bench side faster and roughly nothing to make the arm side faster.** That is a direct, independent argument for the step 15→5 change in section 4, arrived at from live money rather than from a grid.

## THE ARCHITECTURE QUESTION THE WEEK KEEPS ASKING

Two of the week's three structural problems are not solvable by tuning the router, because they are not the router's decisions. **The 22:00Z reactivation timer** made the two largest armings of the week (−\$367.50 and +\$310.00) and it is a clock, not a judgement — nothing grades it, and it re-arms gates the router benched hours earlier with no memory of why. **The lapsed-bench unwind** (−\$232.00 this week) is the same defect in a different costume: a bench expiring is treated as an argument for arming. **Both are the same bug — a permission that expires into the wrong default.** The fix is one line of policy rather than any new mechanism: **a bench that lapses returns a gate to BENCHED-BY-DEFAULT, and an arm needs a positive case made fresh.** The router's prompt already says this in words; the 22:00Z timer overrides it in code.

## 8 · VERDICT AND SATURDAY RECOMMENDATIONS

The verdict on the week: the router is doing its job, its job is currently very small, and the biggest gains available are in what it is pointed at rather than how it is tuned. It was worth +\$400 by the nightly, +\$51.50 by direct grading of its 64 decisions, and it did that on a tape that was 84% chop — the regime where a direction router has least to do. It thrashed one gate badly and left the rest alone. Its thresholds came out best-in-grid on both test sets. And it spent ten hours dead on Thursday without anyone knowing.

| # | RECOMMENDATION                                                                                                                                                                                                                                  | THE EVIDENCE                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 | CONFIDENCE    |
|---|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------|
| 1 | Add exhaustion_short to the router's UP-side managed set (bench it in confirmed TREND_UP).                                                                                                                                                      | 26 counter-trend fires over 30 days for −\$398.50 at −\$15.33/trade; red on 7 of the 8 days it occurred; survives strip-the-worst-day (−\$232.50) and gets stronger dropping the best day (−\$623.00). This week alone it leaked 12 fires for −\$241.00 while the leakage detector reported zero, because the gate is outside its universe. The same gate makes +\$76.50 aligned and is roughly flat in chop — the bench is surgical, not a relegation of the short side.                                    | SHIP SATURDAY |
| 2 | Router timing: window 30→45 min, step 15→5 min, hold 2→3 marks. Leave ER 0.15 and net 30 alone.                                                                                                                                                 | Joint optimum on both universes independently: +\$502 on the live set, +\$399 on the full set. 4 days helped / 0 hurt (live), 6 / 1 (full), worst day −\$10. Independently corroborated by the live ledger: a 45-min window is the fix for the 1-point boundary flicker behind 36 flips on one gate, and 5-min steps address the −\$242 bench-lag measured in section 7. △ Rests on 4 of 40 days, and fitted on one predominantly RANGE regime — re-run after a trend week, revert is a three-constant edit. | SHIP SATURDAY |
| 3 | Fix the lapsed-bench unwind in code, not in the prompt. A bench that expires returns a gate to benched-by-default; the 22:00Z reactivate timer must not re-arm a gate the router benched inside the last N hours without a fresh positive case. | Both of the week's □ BAD grades, −\$232.00 across two incidents (08-10 17:45 −\$92.50, 08-13 17:50 −\$139.50), and the pattern has now fired three times in a single day. The 22:00Z timer separately made the week's two largest armings (−\$367.50 / +\$310.00) with nothing grading it.                                                                                                                                                                                                                   | SHIP SATURDAY |
| 4 | Add a dwell requirement to the direction rule — a confirmed direction must hold for a minimum time before the opposite read can flip it back.                                                                                                   | 36 flips on abs_veto_short, 78% of them producing nothing, 16 lasting under 30 minutes; five flips in 50 minutes on Wednesday morning with three                                                                                                                                                                                                                                                                                                                                                             | TEST FIRST    |

| # | RECOMMENDATION                                                                                            | THE EVIDENCE                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        | CONFIDENCE           |
|---|-----------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------|
|   |                                                                                                           | of them turning on a single point of net. $\Delta$ No dollar figure attached — the churn cost \$0 directly this week, so this is a process fix, not a proven money fix. The 45-minute window in #2 may resolve it on its own; measure before adding a second mechanism.                                                                                                                                                                                                                                             |                      |
| 5 | Make the untradeable cutoff one number in one place.                                                      | untradeable.py says STAY-OUT at $\geq 65$ ; the router operates on “two consecutive $\geq 55$ ”. The 50-75 plateau means neither is wrong and this has cost nothing — but the dashboard and the router can silently disagree about the same day. Do NOT change the value; only unify it.                                                                                                                                                                                                                            | 10-MIN FIX           |
| 6 | Give selector_nightly.py the data_quality IS NULL filter and de-duplicate the day-rider rows.             | Thursday’s file counts one phantom day-rider trade five times at \$72.50 of regret each — \$362.50, 48% of the day’s regret, from a trade that did not exist — and ignores \$1,551 of flagged rows that the P&L module correctly excludes. Third consecutive report to carry this. Every regret figure in section 3 is overstated until it lands.                                                                                                                                                                   | SHIP SATURDAY        |
| 7 | Do NOT lower the net threshold to 20, and do NOT tighten the router on the exhaustion_short “over-bench”. | Net 20→30→40 differ by \$13 across 40 days — noise, and 30 is mid-plateau. The \$721 over-bench figure is largely the measuring instrument: repricing 108 real live trades under the nightly’s own scalp-2R exit turns −\$325.00 into +\$459.50, and the gap is +\$25.02/trade in TREND_UP against −\$0.28 in CHOP. The nightly’s ruler is at its most generous exactly where the router works.                                                                                                                     | HOLD — NOT AN ACTION |
| 8 | The exits, not the router, are where this week’s money went. Handed to the rehab section.                 | Same 108 entries, same tape: live exits −\$325.00, scalp-2R + \$459.50. Zero gap in ordinary chop (−\$2.62/tr, live actually better) and +\$12-13/tr in both clean-trend and violent-whipsaw. The selector agrees from the other side: fixed always-SCALP beat the adaptive selector by \$1,267 on the week, and every one of the top-5 worst picks on both Wednesday and Friday was exhaustion_short choosing TIGHT. Not a router action — but it is the largest number in this section and it should not be lost. | TO REHAB             |

# DISPOSITION TABLE

| LEAD                                                           | VERDICT                | BASIS — AND FOR PARKED, WHAT WOULD REVIVE IT                                                                                                                                                                                                                                           |
|----------------------------------------------------------------|------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| The durable router as a mechanism (5-min systemd tick)         | LIVE                   | 1,148 ticks, zero service failures, 95% of ticks change nothing, +\$400 nightly / + \$51.50 graded on an 84%-chop week. Keep.                                                                                                                                                          |
| ER≥0.15 /  net ≥30 thresholds                                  | LIVE — HOLD            | Best cell on both universes independently; net 20-40 is a \$13 plateau. Nothing to gain. Not an action.                                                                                                                                                                                |
| Timing change: window 45 / step 5 / hold 3                     | LIVE (Saturday deploy) | +\$502 / +\$399 on the two sets, 4-and-6 days helped, 0-and-1 hurt. Corroborated independently by the −\$242 measured actuation lag and the 36-flip churn. Fitted on one range regime — re-run after a trend week.                                                                     |
| exhaustion_short counter-trend bench                           | LIVE (Saturday deploy) | n=26 / −\$398.50 / −\$15.33 per trade over 30 days; red on 7 of 8 days; survives strip-worst and LOO. Thin n, but the error direction is a bench, which is the cheap way to be wrong.                                                                                                  |
| The untradeable meter, and its 65 cutoff                       | LIVE — HOLD            | Monotone across 22 days (−\$651 / −\$212 / +\$7 per day). Cutoff sits mid-plateau, 50-75 identical. Do not tune. △ 79% of the cutoff's dollar case is one day (07-20).                                                                                                                 |
| Re-weighting the STOP-RATE meter upward                        | SHADOW                 | It is the best of the three discriminators (\$559/day separation) and carries the least weight. But it is the only endogenous meter — it measures the desk, not the tape — and the 08-05 demotion was made for that reason. Accumulate n; re-test the weight, do not change it now.    |
| Cutoff 35 (what LOO picks in 17 of 22 folds)                   | REFUTED                | Named test: base-rate contamination. Post-15:00Z trading lost −\$1,712 across all 22 days, so any rule that trades less scores well. Cutoff 35 benches 17 of 22 days — it is not a stay-out rule, it is closing the afternoon desk. The LOO is measuring the base rate, not the meter. |
| Tightening the router on the \$721 exhaustion_short over-bench | REFUTED                | Named test: the scalp-2R control on 108 live trades. The nightly's exit is +\$7.26/ trade more generous than the desk's, and +\$25.02 in TREND_UP against −\$0.28 in CHOP. The over-bench is the ruler, not the router.                                                                |
| The router's leakage detector as currently scoped              | PARKED                 | Reported zero leakage on all four visible days while −\$241.00 leaked through a gate outside its universe. Revives the moment recommendation #1 lands — widen LIVE_UP/ LIVE_DOWN to every live fader and the detector starts telling the truth on its own.                             |
| A dwell-time requirement on the direction rule                 | PARKED                 | The churn is real (36 flips, 78% dead, 17 under 30 min) but it cost \$0 directly this week, so there is no money case yet. Revive if the 45-minute window in #2 fails to cut the flip count below ~10/day, or if a churn round-trip is ever measured costing real dollars.             |
| Monday 08-10 nightly rollop                                    | MEASUREMENT GAP        | The systemd journal does not reach past 2026-08-10 23:11Z, so the backfilled file reads \$0/\$0/\$0 because there was nothing to read. Not a quiet day — a blind day.                                                                                                                  |

| LEAD                                                    | VERDICT               | BASIS — AND FOR PARKED, WHAT WOULD REVIVE IT                                                                                                                                                                                                              |
|---------------------------------------------------------|-----------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                                                         |                       | Revives only by persisting SUPPRESSED-OPEN to disk rather than relying on journal retention.                                                                                                                                                              |
| The 08-13 10h20m OAuth outage                           | FIXED — NOT AN ACTION | router_health_check.py shipped the same day (07cf6ef), live since 15:31Z, alarms on a 3-tick ABORT streak. Cost \$0, and that was luck — the blackout landed in the Asia block. Do not re-recommend.                                                      |
| The 22:00Z reactivate timer as an ungraded arming lever | PARKED                | It made the week's two biggest armings (–\$367.50 and +\$310.00) and nothing scores it because it is a clock, not a decision. Revive as recommendation #3 — it needs the benched-by-default rule applying to it in code.                                  |
| The adaptive exit selector                              | PARKED                | Lost to fixed always-SCALP on 4 of 5 days and by \$1,267 on the week; 59% optimal picks. Revive after the data_quality fix (#6) — its numbers are provably polluted and it deserves re-judging on clean input before any verdict on the mechanism itself. |

**The honest frame on all of it.** This is one week, 108 tournament trades, 64 router decisions, on a tape that was 84% chop. Every threshold sweep here is fitted on 40 days of overwhelmingly one regime and is in-sample to that regime. The two backfilled nightly rollups are tick-honest estimates — a benched gate's trade never ran, so its P&L is a reconstruction, and section 3 shows the reconstruction is systematically optimistic in trend. Where a number rests on a single day I have said which day and what it is worth without it. **The one recommendation I would still make if you told me half of this was wrong is #1** — a live up-fader firing into confirmed up-trends, red on 7 of 8 days, invisible to the detector built to catch exactly that.

## **M1** The biggest runs on MNQ this week — did we show up?

Forget what the desk did this week — start from the raw tape. MNQ printed **68 runs of 1.5×ATR or bigger** in the last seven days (a 15-minute move of at least 44 points, against a typical 15-min range of 29). That is the whole tape, not a top-ten. We **caught 3**, we were positioned against **5**, and we **sat out 60**. The runs we aligned with banked **+\$237** of honest money — but that is only 50% of their \$470 hindsight ceiling, and the 60 we sat

out left **\$8024** on the table. The bleed is not the runs; it is the fading we do around them.

THE FULL CENSUS

Every qualifying run, in order. **Move** is the swing in points; **\$ 1lot** is that move on one MNQ lot (\$2/pt) — the hindsight ceiling. **real \$** is what a gate actually banked near it. **Flow** is net aggressor volume in the 60s before; **amp** is pre-run amplitude; **book** is the far-side L2 depth share (below 0.50 = the side price ran toward was thin). Green = we caught it, red = we fought it.

| TIME (UTC)  | DIR | MOVE | \$ 1LOT | US      | GATE               | REAL \$ | FLOW  | AMP  | BOOK | CLUSTER   |
|-------------|-----|------|---------|---------|--------------------|---------|-------|------|------|-----------|
| 08-09 22:02 | DN  | -59  | 118     | sat out | —                  |         | +0    | 0.16 | —    | UNCLASS   |
| 08-09 23:47 | DN  | -62  | 123     | sat out | —                  |         | +0    | 0.07 | —    | UNCLASS   |
| 08-10 00:02 | UP  | +82  | 165     | sat out | —                  |         | -154  | 0.13 | 0.49 | UNCLASS   |
| 08-10 00:22 | DN  | -59  | 118     | sat out | —                  |         | +19   | 0.09 | 0.51 | UNCLASS   |
| 08-10 01:01 | UP  | +54  | 108     | sat out | —                  |         | -145  | 0.09 | 0.53 | UNCLASS   |
| 08-10 06:13 | DN  | -50  | 100     | sat out | —                  |         | +166  | 0.07 | 0.44 | VACUUM    |
| 08-10 07:11 | UP  | +54  | 108     | sat out | —                  |         | -659  | 0.11 | 0.47 | VACUUM    |
| 08-10 10:57 | DN  | -66  | 132     | sat out | —                  |         | +39   | 0.06 | 0.49 | UNCLASS   |
| 08-10 11:35 | DN  | -56  | 112     | sat out | —                  |         | -57   | 0.03 | 0.57 | UNCLASS   |
| 08-10 12:09 | DN  | -67  | 134     | sat out | —                  |         | +79   | 0.06 | 0.50 | UNCLASS   |
| 08-10 13:20 | DN  | -74  | 147     | sat out | —                  |         | +6    | 0.09 | 0.50 | OPEN/NEWS |
| 08-10 13:35 | UP  | +84  | 167     | FOUGHT  | exhaustion_short_A | +82     | -2910 | 0.37 | 0.52 | OPEN/NEWS |
| 08-10 13:56 | UP  | +109 | 218     | sat out | —                  |         | -118  | 0.21 | 0.51 | OPEN/NEWS |
| 08-10 14:25 | UP  | +46  | 92      | sat out | —                  |         | -438  | 0.22 | 0.50 | OPEN/NEWS |
| 08-10 14:45 | DN  | -90  | 179     | sat out | —                  |         | +101  | 0.10 | 0.48 | OPEN/NEWS |
| 08-10 15:20 | DN  | -62  | 125     | sat out | —                  |         | -190  | 0.09 | 0.50 | UNCLASS   |
| 08-10 15:40 | UP  | +76  | 152     | sat out | —                  |         | -274  | 0.10 | 0.47 | UNCLASS   |

|                |    |      |     |         |                         |     |       |      |      |               |
|----------------|----|------|-----|---------|-------------------------|-----|-------|------|------|---------------|
| 08-10<br>16:29 | DN | -84  | 168 | sat out | —                       |     | -132  | 0.12 | 0.49 | UNCLASS       |
| 08-10<br>16:54 | UP | +51  | 102 | sat out | —                       |     | +52   | 0.09 | 0.52 | UNCLASS       |
| 08-10<br>19:44 | DN | -52  | 104 | sat out | —                       |     | +12   | 0.08 | 0.50 | UNCLASS       |
| 08-10<br>23:47 | DN | -48  | 96  | sat out | —                       |     | +351  | 0.07 | 0.53 | VACUUM        |
| 08-11<br>00:02 | UP | +74  | 149 | sat out | —                       |     | -377  | 0.14 | 0.48 | VACUUM        |
| 08-11<br>01:09 | UP | +63  | 126 | sat out | —                       |     | -82   | 0.05 | 0.48 | UNCLASS       |
| 08-11<br>05:16 | DN | -46  | 92  | sat out | —                       |     | -109  | 0.04 | 0.57 | FLOW-LED      |
| 08-11<br>06:18 | DN | -66  | 132 | sat out | —                       |     | -38   | 0.06 | 0.53 | UNCLASS       |
| 08-11<br>08:03 | DN | -72  | 144 | sat out | —                       |     | +59   | 0.12 | 0.48 | UNCLASS       |
| 08-11<br>10:13 | UP | +70  | 141 | sat out | —                       |     | -33   | 0.04 | 0.51 | UNCLASS       |
| 08-11<br>10:53 | UP | +64  | 128 | sat out | —                       |     | -74   | 0.04 | 0.51 | UNCLASS       |
| 08-11<br>11:08 | DN | -51  | 102 | sat out | —                       |     | +50   | 0.07 | 0.50 | UNCLASS       |
| 08-11<br>12:01 | UP | +47  | 94  | sat out | —                       |     | -305  | 0.05 | 0.49 | VACUUM        |
| 08-11<br>12:16 | DN | -44  | 89  | sat out | —                       |     | +52   | 0.05 | 0.53 | UNCLASS       |
| 08-11<br>13:24 | DN | -137 | 274 | sat out | —                       |     | +48   | 0.05 | 0.51 | OPEN/NEWS     |
| 08-11<br>13:51 | UP | +124 | 249 | sat out | —                       |     | -1145 | 0.16 | 0.50 | OPEN/NEWS     |
| 08-11<br>14:16 | DN | -49  | 98  | sat out | —                       |     | +1720 | 0.15 | 0.50 | OPEN/NEWS     |
| 08-11<br>14:40 | UP | +55  | 110 | sat out | —                       |     | +543  | 0.13 | 0.51 | OPEN/NEWS     |
| 08-11<br>14:55 | DN | -89  | 178 | sat out | —                       |     | +601  | 0.13 | 0.51 | OPEN/NEWS     |
| 08-11<br>15:22 | DN | -54  | 108 | sat out | —                       |     | +539  | 0.11 | 0.51 | UNCLASS       |
| 08-11<br>16:59 | DN | -45  | 90  | sat out | —                       |     | +105  | 0.07 | 0.46 | UNCLASS       |
| 08-11<br>18:15 | DN | -59  | 118 | sat out | —                       |     | -28   | 0.04 | 0.49 | UNCLASS       |
| 08-11<br>20:10 | UP | +47  | 94  | sat out | —                       |     | +197  | 0.07 | 0.52 | UNCLASS       |
| 08-12<br>00:08 | UP | +44  | 89  | sat out | —                       |     | -200  | 0.08 | 0.54 | VACUUM        |
| 08-12<br>12:25 | UP | +87  | 174 | sat out | —                       |     | -161  | 0.07 | 0.53 | VACUUM        |
| 08-12<br>13:08 | UP | +52  | 105 | sat out | —                       |     | +0    | 0.08 | 0.45 | OPEN/NEWS     |
| 08-12<br>13:28 | DN | -92  | 183 | FOUGHT  | capitulation<br>_long_A | -48 | +694  | 0.09 | 0.47 | OPEN/NEWS     |
| 08-12<br>13:50 | DN | -96  | 192 | caught  | exhaustio<br>n_short_A  | -8  | +401  | 0.33 | 0.49 | OPEN/<br>NEWS |

|                |    |      |     |         |                        |      |      |      |      |               |
|----------------|----|------|-----|---------|------------------------|------|------|------|------|---------------|
| 08-12<br>14:40 | DN | -60  | 120 | caught  | exhaustio<br>n_short_A | +112 | -75  | 0.22 | 0.48 | OPEN/<br>NEWS |
| 08-12<br>15:30 | DN | -47  | 94  | sat out | —                      |      | +588 | 0.10 | 0.48 | UNCLASS       |
| 08-12<br>16:36 | UP | +59  | 118 | FOUGHT  | exhaustion_<br>short_A | -46  | -340 | 0.08 | 0.52 | UNCLASS       |
| 08-12<br>20:05 | DN | -59  | 118 | sat out | —                      |      | +392 | 0.11 | 0.52 | VACUUM        |
| 08-12<br>23:59 | UP | +58  | 115 | sat out | —                      |      | +5   | 0.04 | 0.51 | UNCLASS       |
| 08-13<br>11:49 | DN | -48  | 95  | sat out | —                      |      | +60  | 0.04 | 0.53 | UNCLASS       |
| 08-13<br>12:50 | UP | +57  | 114 | sat out | —                      |      | -60  | 0.08 | 0.51 | UNCLASS       |
| 08-13<br>13:32 | UP | +236 | 472 | sat out | —                      |      | +73  | 0.19 | 0.50 | OPEN/NEWS     |
| 08-13<br>13:56 | UP | +79  | 158 | caught  | abs_veto_l<br>ong_A    | +132 | -119 | 0.10 | 0.51 | OPEN/<br>NEWS |
| 08-13<br>14:20 | UP | +75  | 150 | sat out | —                      |      | -711 | 0.11 | 0.42 | OPEN/NEWS     |
| 08-13<br>14:41 | DN | -45  | 90  | sat out | —                      |      | +432 | 0.08 | 0.53 | OPEN/NEWS     |
| 08-13<br>15:22 | DN | -98  | 196 | sat out | —                      |      | +32  | 0.13 | 0.51 | UNCLASS       |
| 08-13<br>15:43 | DN | -61  | 122 | sat out | —                      |      | -227 | 0.15 | 0.50 | UNCLASS       |
| 08-13<br>16:18 | UP | +51  | 102 | FOUGHT  | abs_veto_s<br>hort_A   | -66  | +211 | 0.07 | 0.51 | UNCLASS       |
| 08-13<br>16:55 | UP | +49  | 98  | sat out | —                      |      | -361 | 0.09 | 0.49 | UNCLASS       |
| 08-13<br>18:27 | UP | +45  | 90  | sat out | —                      |      | -6   | 0.07 | 0.53 | UNCLASS       |
| 08-14<br>00:03 | DN | -47  | 94  | sat out | —                      |      | -173 | 0.12 | 0.54 | FLOW-LED      |
| 08-14<br>00:58 | DN | -47  | 94  | sat out | —                      |      | +60  | 0.03 | 0.44 | UNCLASS       |
| 08-14<br>13:22 | DN | -80  | 159 | sat out | —                      |      | +144 | 0.05 | 0.48 | OPEN/NEWS     |
| 08-14<br>13:37 | UP | +79  | 158 | FOUGHT  | abs_veto_s<br>hort_A   | +170 | +408 | 0.24 | 0.51 | OPEN/NEWS     |
| 08-14<br>14:13 | DN | -98  | 197 | sat out | —                      |      | -411 | 0.13 | 0.48 | OPEN/NEWS     |
| 08-14<br>14:48 | DN | -48  | 96  | sat out | —                      |      | +160 | 0.07 | 0.48 | OPEN/NEWS     |
| 08-14<br>15:11 | DN | -73  | 146 | sat out | —                      |      | -92  | 0.10 | 0.48 | UNCLASS       |



## What the runs were — by cause

Each run gets a cause fingerprint from its pre-run tape. This is the map for where to hunt a new gate.

| CLUSTER   | RUNS | WHAT IT IS                                      |
|-----------|------|-------------------------------------------------|
| UNCLASS   | 36   | no clear tape tell (quiet ignition)             |
| OPEN/NEWS | 22   | the 13:00-15:00 UTC cash-open / data window     |
| VACUUM    | 8    | moved AGAINST the tape (a snap-back / stop-run) |
| FLOW-LED  | 2    | aggressors drove it (a continuation)            |

## HONEST MONEY — CEILING VS WHAT WE BANKED

| RUNS             | HINDSIGHT CEILING | WHAT WE REALLY MADE                      |
|------------------|-------------------|------------------------------------------|
| Caught (aligned) | 3                 | \$470      +\$237 (50%)                  |
| Fought (against) | 5                 | \$728      \$+92                         |
| Sat out          | 60                | \$8024      \$0 ← the money on the table |

The runs made money and the fading gave it back; we convert a fraction of the ceiling and ignore the rest. The next two movements ask whether we can do better — first with the gates we own, then with new ones.

## THE SAME CENSUS ON GOLD — MGC, the second instrument

### M1 The biggest runs on MGC this week — did we show up?

Forget what the desk did this week — start from the raw tape. MGC printed **74 runs of 1.5×ATR or bigger** in the last seven days (a 15-minute

move of at least 10 points, against a typical 15-min range of 7). That is the whole tape, not a top-ten. We **caught 0**, we were positioned against **0**, and we **sat out 74**. The runs we aligned with banked **+\$0** of honest money — but that is only 0% of their \$0 hindsight ceiling, and the 74 we sat out left **\$10649** on the table. The bleed is not the runs; it is the fading we do around them.

THE FULL CENSUS

Every qualifying run, in order. **Move** is the swing in points; **\$ 1lot** is that move on one MGC lot (\$10/pt) — the hindsight ceiling. **real \$** is what a gate actually banked near it. **Flow** is net aggressor volume in the 60s before; **amp** is pre-run amplitude; **book** is the far-side L2 depth share (below 0.50 = the side price ran toward was thin). Green = we caught it, red = we fought it.

| TIME (UTC)  | DIR | MOVE | \$ 1LOT | US      | GATE | REAL \$ | FLOW | AMP  | BOOK | CLUSTER   |
|-------------|-----|------|---------|---------|------|---------|------|------|------|-----------|
| 08-09 23:59 | UP  | +12  | 124     | sat out | —    |         | +0   | 0.25 | —    | UNCLASS   |
| 08-10 01:03 | DN  | -15  | 153     | sat out | —    |         | +12  | 0.28 | —    | VACUUM    |
| 08-10 01:26 | DN  | -12  | 118     | sat out | —    |         | +6   | 0.15 | —    | UNCLASS   |
| 08-10 04:08 | UP  | +14  | 138     | sat out | —    |         | +0   | 0.07 | —    | UNCLASS   |
| 08-10 05:38 | UP  | +11  | 108     | sat out | —    |         | +2   | 0.12 | —    | UNCLASS   |
| 08-10 06:34 | UP  | +11  | 108     | sat out | —    |         | +4   | 0.06 | —    | UNCLASS   |
| 08-10 13:42 | DN  | -16  | 163     | sat out | —    |         | -13  | 0.11 | —    | OPEN/NEWS |
| 08-10 14:12 | UP  | +20  | 200     | sat out | —    |         | +12  | 0.08 | —    | OPEN/NEWS |
| 08-10 15:48 | UP  | +11  | 108     | sat out | —    |         | +78  | 0.11 | —    | FLOW-LED  |
| 08-10 16:29 | DN  | -11  | 110     | sat out | —    |         | -3   | 0.05 | —    | UNCLASS   |
| 08-10 18:23 | UP  | +18  | 175     | sat out | —    |         | +3   | 0.04 | —    | UNCLASS   |
| 08-10 19:15 | UP  | +11  | 106     | sat out | —    |         | +11  | 0.05 | —    | UNCLASS   |
| 08-11 00:47 | DN  | -18  | 180     | sat out | —    |         | +50  | 0.25 | —    | VACUUM    |
|             | UP  | +20  | 201     | sat out | —    |         | -8   | 0.35 | —    |           |

|                |    |     |     |         |   |  |      |      |                   |
|----------------|----|-----|-----|---------|---|--|------|------|-------------------|
| 08-11<br>01:02 |    |     |     |         |   |  |      |      | VOL-<br>EXPANSION |
| 08-11<br>01:41 | UP | +12 | 124 | sat out | — |  | +2   | 0.14 | — UNCLASS         |
| 08-11<br>02:24 | DN | -14 | 141 | sat out | — |  | -4   | 0.12 | — UNCLASS         |
| 08-11<br>05:17 | DN | -26 | 264 | sat out | — |  | -11  | 0.09 | — UNCLASS         |
| 08-11<br>05:36 | DN | -14 | 136 | sat out | — |  | -2   | 0.29 | — UNCLASS         |
| 08-11<br>06:18 | DN | -13 | 130 | sat out | — |  | +2   | 0.17 | — UNCLASS         |
| 08-11<br>08:21 | UP | +12 | 125 | sat out | — |  | -8   | 0.11 | — UNCLASS         |
| 08-11<br>10:58 | UP | +12 | 118 | sat out | — |  | +1   | 0.05 | — UNCLASS         |
| 08-11<br>12:01 | UP | +11 | 106 | sat out | — |  | -3   | 0.07 | — UNCLASS         |
| 08-11<br>13:45 | DN | -16 | 162 | sat out | — |  | +19  | 0.13 | — OPEN/NEWS       |
| 08-11<br>15:04 | DN | -12 | 124 | sat out | — |  | +45  | 0.12 | — VACUUM          |
| 08-11<br>15:23 | DN | -14 | 141 | sat out | — |  | +7   | 0.17 | — UNCLASS         |
| 08-11<br>15:41 | UP | +12 | 116 | sat out | — |  | -12  | 0.19 | — UNCLASS         |
| 08-11<br>17:53 | DN | -13 | 131 | sat out | — |  | -4   | 0.09 | — UNCLASS         |
| 08-11<br>19:17 | UP | +12 | 118 | sat out | — |  | +1   | 0.08 | — UNCLASS         |
| 08-12<br>00:04 | UP | +11 | 112 | sat out | — |  | -2   | 0.12 | — UNCLASS         |
| 08-12<br>00:57 | UP | +11 | 110 | sat out | — |  | +0   | 0.08 | — UNCLASS         |
| 08-12<br>01:35 | UP | +11 | 110 | sat out | — |  | -14  | 0.13 | — UNCLASS         |
| 08-12<br>02:39 | UP | +14 | 144 | sat out | — |  | -1   | 0.07 | — UNCLASS         |
| 08-12<br>05:15 | DN | -11 | 112 | sat out | — |  | -30  | 0.10 | — FLOW-LED        |
| 08-12<br>06:37 | UP | +12 | 124 | sat out | — |  | -4   | 0.03 | — UNCLASS         |
| 08-12<br>09:01 | UP | +15 | 147 | sat out | — |  | +0   | 0.03 | — UNCLASS         |
| 08-12<br>12:15 | UP | +20 | 202 | sat out | — |  | +1   | 0.07 | — UNCLASS         |
| 08-12<br>12:31 | UP | +29 | 293 | sat out | — |  | +200 | 0.94 | — FLOW-LED        |
| 08-12<br>13:00 | DN | -18 | 176 | sat out | — |  | +42  | 0.28 | — OPEN/NEWS       |
| 08-12<br>13:18 | DN | -13 | 127 | sat out | — |  | +20  | 0.17 | — OPEN/NEWS       |
| 08-12<br>13:40 | UP | +21 | 207 | sat out | — |  | -83  | 0.21 | — OPEN/NEWS       |
| 08-12<br>14:03 | UP | +14 | 138 | sat out | — |  | +6   | 0.26 | — OPEN/NEWS       |

|                |    |     |     |         |   |      |      |                     |
|----------------|----|-----|-----|---------|---|------|------|---------------------|
| 08-12<br>14:21 | DN | -19 | 187 | sat out | — | -12  | 0.10 | — OPEN/NEWS         |
| 08-12<br>16:20 | DN | -13 | 130 | sat out | — | +6   | 0.08 | — UNCLASS           |
| 08-12<br>16:42 | DN | -16 | 155 | sat out | — | +22  | 0.06 | — VACUUM            |
| 08-13<br>00:04 | UP | +13 | 133 | sat out | — | +16  | 0.17 | — FLOW-LED          |
| 08-13<br>00:19 | UP | +17 | 174 | sat out | — | +5   | 0.21 | — UNCLASS           |
| 08-13<br>00:46 | DN | -17 | 168 | sat out | — | -21  | 0.09 | — FLOW-LED          |
| 08-13<br>01:10 | DN | -14 | 136 | sat out | — | +15  | 0.11 | — VACUUM            |
| 08-13<br>01:37 | DN | -17 | 168 | sat out | — | -26  | 0.09 | — FLOW-LED          |
| 08-13<br>02:21 | DN | -12 | 121 | sat out | — | +3   | 0.05 | — UNCLASS           |
| 08-13<br>03:49 | DN | -11 | 114 | sat out | — | +4   | 0.06 | — UNCLASS           |
| 08-13<br>05:29 | DN | -20 | 196 | sat out | — | -1   | 0.10 | — UNCLASS           |
| 08-13<br>05:52 | DN | -11 | 112 | sat out | — | +8   | 0.10 | — UNCLASS           |
| 08-13<br>12:32 | UP | +13 | 126 | sat out | — | -30  | 0.47 | — VACUUM            |
| 08-13<br>13:22 | DN | -17 | 172 | sat out | — | +2   | 0.12 | — OPEN/NEWS         |
| 08-13<br>14:15 | DN | -12 | 115 | sat out | — | -33  | 0.17 | — OPEN/NEWS         |
| 08-13<br>14:30 | UP | +21 | 211 | sat out | — | -308 | 0.53 | — OPEN/NEWS         |
| 08-13<br>15:00 | DN | -18 | 183 | sat out | — | -3   | 0.08 | — UNCLASS           |
| 08-13<br>15:33 | UP | +11 | 113 | sat out | — | -31  | 0.34 | — VOL-<br>EXPANSION |
| 08-13<br>16:03 | DN | -14 | 140 | sat out | — | +12  | 0.10 | — UNCLASS           |
| 08-13<br>16:20 | UP | +14 | 139 | sat out | — | -8   | 0.11 | — UNCLASS           |
| 08-13<br>19:00 | DN | -12 | 121 | sat out | — | +1   | 0.03 | — UNCLASS           |
| 08-14<br>00:14 | DN | -12 | 123 | sat out | — | +12  | 0.16 | — UNCLASS           |
| 08-14<br>00:35 | DN | -12 | 118 | sat out | — | +28  | 0.20 | — VACUUM            |
| 08-14<br>00:56 | DN | -16 | 156 | sat out | — | +3   | 0.14 | — UNCLASS           |
| 08-14<br>05:44 | UP | +17 | 167 | sat out | — | -27  | 0.16 | — VACUUM            |
| 08-14<br>11:49 | UP | +13 | 133 | sat out | — | +1   | 0.05 | — UNCLASS           |
| 08-14<br>12:16 | UP | +16 | 165 | sat out | — | -2   | 0.14 | — UNCLASS           |
| 08-14<br>12:31 | DN | -14 | 140 | sat out | — | +48  | 0.25 | — VACUUM            |

|                |    |     |     |         |   |     |      |             |
|----------------|----|-----|-----|---------|---|-----|------|-------------|
| 08-14<br>13:22 | UP | +13 | 129 | sat out | — | +6  | 0.09 | — OPEN/NEWS |
| 08-14<br>13:38 | DN | -12 | 124 | sat out | — | +15 | 0.21 | — OPEN/NEWS |
| 08-14<br>14:27 | DN | -12 | 115 | sat out | — | +1  | 0.17 | — OPEN/NEWS |
| 08-14<br>14:58 | UP | +13 | 128 | sat out | — | +10 | 0.07 | — OPEN/NEWS |
| 08-14<br>15:13 | DN | -11 | 107 | sat out | — | +18 | 0.17 | — VACUUM    |

## What the runs were — by cause

Each run gets a cause fingerprint from its pre-run tape. This is the map for where to hunt a new gate.

| CLUSTER       | RUNS | WHAT IT IS                                      |
|---------------|------|-------------------------------------------------|
| UNCLASS       | 41   | no clear tape tell (quiet ignition)             |
| OPEN/NEWS     | 15   | the 13:00–15:00 UTC cash-open / data window     |
| VACUUM        | 10   | moved AGAINST the tape (a snap-back / stop-run) |
| FLOW-LED      | 6    | aggressors drove it (a continuation)            |
| VOL-EXPANSION | 2    | amplitude broke out of a coil                   |

## HONEST MONEY — CEILING VS WHAT WE BANKED

| RUNS             | HINDSIGHT CEILING | WHAT WE REALLY MADE                       |
|------------------|-------------------|-------------------------------------------|
| Caught (aligned) | 0                 | \$0      +\$0 (0%)                        |
| Fought (against) | 0                 | \$0      \$+0                             |
| Sat out          | 74                | \$10649      \$0 ← the money on the table |

The runs made money and the fading gave it back; we convert a fraction of the ceiling and ignore the rest. The next two movements ask whether we can do better — first with the gates we own, then with new ones.

## **M2** The idle-gate lab — could the six gates we already own have caught the 60 misses?

---

Movement 1 froze the tape: **68 big runs on MNQ this week**, and we **sat out 60 of them**, carrying a hindsight ceiling of **\$8,024** on a single lot. Before anyone invents a new gate, this movement asks the cheaper question: we already own six gates — if we had simply left all six armed all week, would any of them have been in those runs? So I took all six live tournament gates straight out of `deciders.py`, fired them **mechanically** (exactly as the desk runs them — live parameters, live ATR floors, the live 55-second absorption veto, the live 5-second exhaustion confirm and its counter-regime veto) and again **ungated** (the raw gate shape with every filter stripped off), in-direction only, anywhere in the ten minutes before each run ignited, and repriced every exit tick-by-tick on this week's real `capture.db` trade prints under each gate's deployed two-lot exit. **The mechanical answer is +\$15.50 on six fires across 57 testable runs** — and strip the single best trade and that becomes **-\$94.50**. Ungated it is +\$379.50 on 27 fires, +\$69.50 once you remove the best three. That is an **HONEST NULL**, and this week I can tell you the mechanism in one line with a number behind it: **this week's runs ignited out of unusually quiet tape (median peak ATR in the ten minutes before ignition: 12.0 points), and the desk's two biggest magnitude floors are set at 22 and 20 points — they are calibrated to stand down in exactly the tape the runs come out of.**

### **LEAD WITH WHAT SURVIVED**

Four things came out of this alive. Three of them are small and one of them is a correction to what this section told you last week, so I want all four in front of the null rather than buried behind it.

1. **The two short FADERS were the only gates that made money, and they were perfect: 3 fires, 3 winners, +\$302.50.** `exhaustion_short` fired twice for **+\$192.50** and `rgv_short` once for **+\$110.00**, all three mechanically, all three with both lots reaching target. Every other gate

lost. That is three trades, so it is a **SHADOW** lead and nothing more — but it is the first time in this section that the mechanical book has had a green column at all, and it is worth saying that the fade gates fired 9–10 minutes ahead of ignition and still won, which the timing ladder says should be nearly impossible.

2. **The 08-08 grind\_long ATR 10 → 22 revert is VINDICATED on run tape.** With no floor at all grind fires 5 times for **−\$178**; with the deployed  $\text{ATR} \geq 22$  floor it fires once for **−\$120**. The floor blocked four runs that were collectively losing money. This section is usually where floors get accused of blocking the good stuff; this week the floor is doing its job.
3. **Last week's headline — "grind\_long's ext\_hi 2.0 ceiling is the single filter standing between us and these runs" — does NOT hold on this week's tape.** The ceiling was deleted on 08-08 (it now falls back to 4.0). I re-fired grind under the old cell, the new cell and a 3.0 cell on the same 23 up-runs: **all three produce the identical single fire and the identical −\$120**. The binding clause is not the extension ceiling at all — it is "VWAP slope not up", which killed grind on **14 of its 23** eligible runs. Removing ext\_hi bought us nothing here.
4. **capitulation\_long's require\_flip is the one live filter that cost real money on this population: +\$158.50.** Ungated the gate fired six times for +\$116.50; mechanically it fired **zero** times. Three of those six were blocked by the  $\text{ATR} \geq 10$  floor (and that floor saved \$42 — two losers and a small winner), but the other three were blocked by require\_flip, and two of them were \$97 winners. Thin — three trades — and require\_flip was re-derived on the whole book on 08-02 with an ordering-correct win rate, so this number must not be allowed to overturn it. It is a shadow question, not a change.

## How the test was run — so you can trust the sign

Everything below comes from two scripts run once each against the frozen census: scripts/idle\_gate\_backtest\_wk33.py and its companion scripts/idle\_gate\_timing\_decay\_wk33.py. No number in this section is

carried over from any prior week; where I quote last week it is to contradict it, and the contradiction is computed here.

| COMPONENT                       | RULE                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
|---------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Population                      | The <b>60 sat-out runs</b> of the frozen Movement-1 census, read straight out of the locked census stdout — not re-derived, not re-filtered, not re-run (run_census.py is a multi-minute tick crunch and the whole point of freezing it is that every movement reads the same snapshot).                                                                                                                                                                                                                                                                                                                                                                                       |
| Testable subset                 | <b>57 of 60.</b> capture.db holds five trading days of ticks and its first print this week is <b>08-10 00:00:00 UTC</b> . Three runs sit before that or need a pre-window that does — 08-09 22:02, 08-09 23:47 and 08-10 00:02 — and they are excluded rather than faked. I checked the Parquet lake for them too: it has no 2026-08-09 tick file, so the Sunday-evening reopen genuinely has no tick tape anywhere on this desk. The 57 testable runs carry <b>\$7,619</b> of the sat-out ceiling and the three dropped runs carry <b>\$406</b> . (The census headline is \$8,024 and the printed per-run ceilings sum to \$8,025 — a dollar of rounding, not a discrepancy.) |
| Why capture.db and not the lake | The census window is <b>08-09 → 08-14</b> , i.e. this week, which is strictly inside capture.db's five-trading-day window — and capture.db is the only source that has Friday 08-14 at all (the lake's last mirrored tick day is 08-13). For a one-week study it is both legal and the more complete source. Anything multi-week would have to go to the lake.                                                                                                                                                                                                                                                                                                                 |
| Gates                           | All six live tournament slots: grind_long, capitulation_long, abs_veto_long on the long side; rgv_short, exhaustion_short, abs_veto_short on the short side. The actual decider functions are imported and called — there is no re-implementation to drift out of sync. (nipc_long/nipc_short ship benched and are not part of the live six.)                                                                                                                                                                                                                                                                                                                                  |
| Direction                       | <b>In-direction only.</b> A long gate is tested only on the 23 up-runs, a short gate only on the 34 down-runs. We are asking "could it have been in the move", not "could it have been near it".                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
| Window                          | The <b>ten minutes before ignition</b> . The three bar gates are tested on each of the ten completed 1-minute bars; the two tape gates (capitulation, exhaustion) are tested every 5 seconds across the whole window. First fire wins.                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| Mechanical mode                 | Live params and live filters, as deployed today: grind ATR≥ <b>22</b> and no ext ceiling (both changed on 08-08); capitulation ATR≥10 + require_flip; abs_veto amp_floor 0.0004 plus the <b>55-second absorption veto</b> and its 55-second entry delay; rgv_short ATR≥20; exhaustion_short plus its 5-second confirm-veto and its veto_counter_regime entry veto (that last one was missing from last week's harness — it is in now, and it blocked nothing).                                                                                                                                                                                                                 |
| Ungated mode                    | The raw gate shape. ATR floors off, amplitude floor off, extension ceiling off entirely, require_flip off, all three vetoes off. The gate's core geometry is kept — this is "filters off", not "a different gate".                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
| Exit                            | Each gate's <b>deployed</b> two-lot exit from data/exit_overrides.json — Lot A fixed-R, Lot B chandelier or fixed-R, one lot each — plus the quiet-tape clip below 22pt entry ATR and the native 1×ATR protective stop on the loss side. Repriced on every real trade print, 60-minute hold cap.                                                                                                                                                                                                                                                                                                                                                                               |
| Costs                           | <b>\$1.50 per lot round trip, \$2.00 per point</b> (MNQ). Two lots a fire, so \$3.00 a fire before anything happens.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| Bars                            | 1-minute bars built from the 5-second bars with <b>integer</b> flooring (bar_ts - bar_ts % 60) — deliberately, because the report's own instrument-failure list flags the float-division version as a silent no-op that reads 5-second bars as minutes.                                                                                                                                                                                                                                                                                                                                                                                                                        |



Data actually consumed: **5,172,092 trade ticks** (08-10 00:00:00 → 08-14 21:00:00 UTC), **13.96 million** level-1 book rows, 6,780 one-minute bars inside the tick window.

The scoreboard — mechanically, six gates managed six fires all week

This is the headline table. "Eligible" is how many sat-out runs that gate could even have traded, its own direction only. Green rows made money, red rows lost it.

| GATE                | SIDE  | ELIGIBLE RUNS | FIRES | NET \$  | \$ / FIRE | WIN % | HOW THEY ENDED                 |
|---------------------|-------|---------------|-------|---------|-----------|-------|--------------------------------|
| grind_long          | LONG  | 23            | 1     | −120.00 | −120.00   | 0%    | both lots stopped              |
| capitulation_long   | LONG  | 23            | 0     | 0.00    | —         | —     | never fired once               |
| abs_veto_long       | LONG  | 23            | 1     | −76.00  | −76.00    | 0%    | both lots stopped              |
| rgv_short           | SHORT | 34            | 1     | +110.00 | +110.00   | 100%  | both lots hit target           |
| exhaustion_short    | SHORT | 34            | 2     | +192.50 | +96.25    | 100%  | target + target/<br>chandelier |
| abs_veto_short      | SHORT | 34            | 1     | −91.00  | −91.00    | 0%    | both lots stopped              |
| ALL SIX, mechanical | —     | 57 runs       | 6     | +15.50  | +2.58     | 50%   | 3 winners, 3 losers            |

Six fires. Not six profitable fires — six fires **at all**, from six gates, over a whole week of the biggest moves MNQ printed. Here is every one of them, because six is few enough to look at individually and this is exactly the kind of table where an average hides the story.

| RUN (UTC)   | MOVE   | GATE             | ENTRY ATR | FIRE<br>THIS<br>EARLY | REGIME AT<br>IGNITION | NET \$  | ENDED                  |
|-------------|--------|------------------|-----------|-----------------------|-----------------------|---------|------------------------|
| 08-11 14:16 | −49pt  | rgv_short        | 20.8      | 3.0 min               | BUILDING              | +110.00 | both target            |
| 08-11 15:22 | −54pt  | exhaustion_short | 16.1      | 9.6 min               | BUILDING              | +97.00  | both target            |
| 08-10 14:45 | −90pt  | exhaustion_short | 22.9      | 9.0 min               | VIOLENT-<br>WHIPSAW   | +95.50  | target /<br>chandelier |
| 08-13 14:20 | +75pt  | abs_veto_long    | 18.2      | 7.1 min               | NORMAL-CHOP           | −76.00  | both stopped           |
| 08-10 15:20 | −62pt  | abs_veto_short   | 21.8      | 6.1 min               | VIOLENT-<br>WHIPSAW   | −91.00  | both stopped           |
| 08-10 13:56 | +109pt | grind_long       | 29.1      | 6.0 min               | VIOLENT-<br>WHIPSAW   | −120.00 | both stopped           |

Read that table and the pattern is not subtle. **All three winners are SHORT fades. All three losers are MOMENTUM entries taken in whipsaw or chop.** The three winners also fired 3.0, 9.0 and 9.6 minutes ahead of ignition, which by the decay ladder further down should be dead ground — a fade that is right about the level does not care much when you took it, because the level is where the money is. The three losers fired 6–7 minutes early, chasing something that had not started.

## The same six with every filter switched off

Now strip the ATR floors, the amplitude floor, the extension ceiling, the flip requirement and all three vetoes. This is the "leave everything on, damn the filters" book — the fair upper bound on what the gates we own could physically have done.

| GATE              | ELIGIBLE | FIRES | NET \$  | \$ / FIRE | WIN % | VS MECHANICAL                                     |
|-------------------|----------|-------|---------|-----------|-------|---------------------------------------------------|
| grind_long        | 23       | 7     | −152.00 | −21.70    | 29%   | +6 fires, −\$32 — the filters were helping        |
| capitulation_long | 23       | 6     | +116.50 | +19.40    | 50%   | +6 fires, +\$116.50 — the filters took it to zero |
| abs_veto_long     | 23       | 3     | −7.00   | −2.30     | 33%   | +2 fires, +\$69                                   |
| rgv_short         | 34       | 2     | +212.50 | +106.25   | 100%  | +1 fire, +\$102.50 (the ATR≥20 floor)             |
| exhaustion_short  | 34       | 2     | +193.50 | +96.75    | 100%  | identical — no veto bit at all                    |
| abs_veto_short    | 34       | 7     | +16.00  | +2.30     | 43%   | +6 fires, +\$107                                  |
| ALL SIX, ungated  | 57 runs  | 27    | +379.50 | +14.06    | 48%   | +21 fires, +\$364                                 |

### DO NOT GET EXCITED ABOUT THE +\$379.50

It is a handful of \$97 target-hits. Strip the best **one** and it falls to **+\$269.50**; strip the best **three** and it falls to **+\$69.50** across 24 remaining fires — about fourteen dollars a day, on a hindsight-selected population, in-sample, with no arming rule attached to it. Day by day it is +\$96.50 / +\$373.00 / +\$191.00 / −\$281.00 / **\$0.00**, and leaving out the best day (08-11) drops the whole book to **+\$6.50** — so one Tuesday is the book. Costs are not the problem: stress the fee to \$3.00 a lot round trip and it is still +\$298.50, to \$6.00 and it is still +\$136.50. **Twenty-seven fires and a single load-**

**bearing day is the problem.** And note the last day in that list: Friday 08-14 had **six** sat-out runs and drew **zero** fires from six gates with every filter removed.

## The money on the table — reconciled

Movement 1 said the sat-out runs left \$8,024 of one-lot hindsight ceiling lying there. Here is how much of that the gates we already own could actually have touched.

| MEASURE                                        | MECHANICAL | UNGATED  |
|------------------------------------------------|------------|----------|
| Sat-out runs in the census                     | 60         | 60       |
| Testable (tick tape exists for the pre-window) | 57         | 57       |
| Ceiling on those 57 runs                       | \$7,619    | \$7,619  |
| Runs that drew ANY in-direction fire           | 6 of 57    | 24 of 57 |
| Ceiling of the runs we touched                 | \$878      | \$3,452  |
| Of the 15 biggest runs, how many touched       | 3 of 15    | 6 of 15  |
| Of the 25 biggest runs, how many touched       | 3 of 25    | 9 of 25  |
| Fires                                          | 6          | 27       |
| Net \$                                         | +15.50     | +379.50  |
| Strip the best 1 trade                         | −94.50     | +269.50  |
| Strip the best 3 trades                        | −287.00    | +69.50   |
| Conversion of the ceiling we touched           | 1.8%       | 11.0%    |
| Conversion of the whole \$7,619                | 0.2%       | 5.0%     |

The line that matters is the second-to-last pair. Even with **every filter removed**, 33 of the 57 testable runs draw no fire at all — and those 33 carry **\$4,167** of the \$7,619. The biggest run of the entire week, the **+236pt monster at 08-13 13:32** with a \$472 ceiling, was touched exactly once in this whole study: ungated capitulation\_long fired 115 seconds ahead of it and banked **+\$18.00**. Eighteen dollars of four hundred and seventy-two.

## Why each gate misses — the mechanism, clause by clause

For every eligible run I recorded how far down its own rule chain each gate ever got in the ten-minute window, and kept the deepest reach. This turns "it didn't fire much" into "here is the exact sentence that stopped it".

## grind\_long — there was no established up-trend to ride, and that is the whole story

| WHERE THE GATE STOPPED                      | RUNS (OF 23) | MODE              |
|---------------------------------------------|--------------|-------------------|
| VWAP slope not up — no established up-trend | 14           | both              |
| ATR floor 22pt not met                      | 4            | mechanical only   |
| Price not riding above VWAP (ext < 0.3 ATR) | 2            | both              |
| Already stretched (ext > 4.0 ATR)           | 2            | mechanical only   |
| <b>FIRES</b>                                | <b>1</b>     | <b>mechanical</b> |

Sixty-one percent of the up-runs never even reached grind's first test. A continuation gate needs a trend to continue; these runs did not have one, they were the trend arriving. That is the census finding restated from inside the gate: big moves rupture flat tape.

## abs\_veto\_long / abs\_veto\_short — the five-bar thrust simply isn't there yet

| WHERE THE GATE STOPPED                        | ABS_VETO_LONG (23 UP-RUNS) | ABS_VETO_SHORT (34 DOWN-RUNS) |
|-----------------------------------------------|----------------------------|-------------------------------|
| No 5-bar thrust $\geq 1.5$ ATR                | 16                         | 18                            |
| Amplitude floor: tape too thin (atr% < 0.04%) | 5                          | 6                             |
| Thrust is the wrong way                       | 0                          | 4                             |
| No volume surge on the thrust bar             | 1                          | 2                             |
| 55-second absorption VETO (burst absorbed)    | 0                          | 3                             |
| <b>FIRES</b>                                  | <b>1</b>                   | <b>1</b>                      |

Seven in ten up-runs (16 of 23) and just over half the down-runs (18 of 34) never produce a 1.5-ATR five-minute burst in their own direction in the ten minutes before they go. Of course they don't — the burst is the run. The amplitude floor then removes another 5-6 each, correctly: the pre-run tape is genuinely thin.

rgv\_short — the ATR floor stands it down on two-thirds of the tape

| WHERE THE GATE STOPPED                             | MECHANICAL (34 DOWN-RUNS) | UNGATED (34 DOWN-RUNS) |
|----------------------------------------------------|---------------------------|------------------------|
| ATR floor 20pt not met                             | 22                        | —                      |
| Regime stand-down (vol too high / slope too steep) | 9                         | 9                      |
| Not stretched $\geq 1.5$ ATR above VWAP            | 2                         | 16                     |
| No fresh down-turn                                 | 0                         | 7                      |
| FIRES                                              | 1                         | 2                      |

This is the clearest single number in the section. **The ATR $\geq$ 20 floor stands rgv\_short down on 22 of its 34 eligible runs** — and independently of the gate, I measured the peak ATR reached anywhere in each run's ten-minute pre-window: only **6 of the 34 down-runs** ever get to 20 points at all. Take the floor away and the binding clause simply moves to "not stretched above VWAP" on 16 runs, and the gate still only fires twice. **The floor is not what is stopping rgv\_short from catching runs; the geometry is.** Both fires it did get were winners, so nothing here argues for touching it.

capitulation\_long — the only gate whose live filters cost it money this week

| WHERE THE GATE STOPPED                          | MECHANICAL (23 UP-RUNS) | UNGATED (23 UP-RUNS) |
|-------------------------------------------------|-------------------------|----------------------|
| Sell volume < 2.5 $\times$ baseline — no climax | 15                      | 15                   |
| Sellers not 60% dominant                        | 2                       | 2                    |
| Buyers never flipped at the low (require_flip)  | 3                       | —                    |
| ATR floor 10pt not met                          | 3                       | —                    |
| FIRES                                           | 0                       | 6                    |

Six ungated fires, +\$116.50, three of them winners. Mechanically: nothing. The two live filters split the six cleanly and pull in opposite directions:

| BLOCKED BY   | RUN         | UNGATED RESULT | VERDICT ON THE FILTER |
|--------------|-------------|----------------|-----------------------|
| require_flip | 08-10 01:01 | +97.50         | blocked a winner      |

|                    |             |                |                                             |
|--------------------|-------------|----------------|---------------------------------------------|
| require_flip       | 08-12 23:59 | +97.00         | blocked a winner                            |
| require_flip       | 08-13 18:27 | -36.00         | blocked a loser                             |
| ATR ≥ 10           | 08-11 00:02 | -31.00         | blocked a loser                             |
| ATR ≥ 10           | 08-11 10:13 | -29.00         | blocked a loser                             |
| ATR ≥ 10           | 08-13 13:32 | +18.00         | blocked a small winner (the +236pt monster) |
| require_flip total | —           | -\$158.50 cost | net harmful here                            |
| ATR ≥ 10 total     | —           | +\$42.00 saved | net helpful here                            |

Three trades either side. I am not going to pretend that overturns the 08-02 re-derivation, which measured `require_flip` on the whole book with an ordering-correct win rate and found the loose version (`capit_loose`) at -\$5,250 over 238 fires. What I will say is that on run tape specifically, the flip requirement arrives too late twice out of three, and that is a shadow-book question with a clean shape.

## exhaustion\_short — barely exists, but everything it does is right

| WHERE THE GATE STOPPED                                       | RUNS (OF 34) |
|--------------------------------------------------------------|--------------|
| 20-second net aggressor flow < 400 — no aggression to fade   | 16           |
| The aggression DID move price > 2pt — that is not absorption | 10           |
| No ask wall (hit side not 1.5x)                              | 6            |
| FIRES                                                        | 2            |

Mechanical and ungated are **identical** — neither the 5-second confirm-veto nor the `veto_counter_regime` entry veto blocked a single thing on this population, so the gate's two live safety catches cost nothing here. It fired twice and won twice, +\$192.50. Note the second row: on ten runs the gate correctly refused because the flow did move price. That is a fade gate declining to fade a real move, which is exactly right — and it is also the most interesting negative signal in this whole section, because "heavy flow that IS moving price" is a clean, cheap, already-implemented detector of a move that is genuinely underway.

# The killer number: near a run, our gates fire at roughly chance

Here is the check that turns "they didn't fire much" into "they are not looking at anything predictive". Across the whole tick-covered week I counted every minute (bar gates) or 20-second sample (tape gates) in which each gate's condition was **true**, then asked what share of those fell inside one of the 57 ten-minute pre-run windows. Those windows cover **9.2%** of the minute bars and **8.4%** of the tape samples — that is the chance rate. A gate that knew something about imminent runs would sit well above it.

| GATE              | MODE       | SIGNAL-MINUTES / SAMPLES ALL WEEK | INSIDE A PRE-RUN WINDOW | SHARE | VS CHANCE                 |
|-------------------|------------|-----------------------------------|-------------------------|-------|---------------------------|
| grind_long        | mechanical | 44                                | 8                       | 18.2% | 1.97× (n=44)              |
| grind_long        | ungated    | 1,118                             | 92                      | 8.2%  | 0.89×                     |
| capitulation_long | mechanical | 19                                | 3                       | 15.8% | 1.88× (n=19)              |
| capitulation_long | ungated    | 79                                | 8                       | 10.1% | 1.21×                     |
| abs_veto_long     | mechanical | 25                                | 4                       | 16.0% | 1.73× (n=25)              |
| abs_veto_long     | ungated    | 231                               | 20                      | 8.7%  | 0.94×                     |
| rgv_short         | mechanical | 1                                 | 1                       | 100%  | 10.8× (n=1 — meaningless) |
| rgv_short         | ungated    | 122                               | 5                       | 4.1%  | 0.44×                     |
| exhaustion_short  | either     | 8                                 | 1                       | 12.5% | 1.49× (n=8)               |
| abs_veto_short    | mechanical | 44                                | 7                       | 15.9% | 1.72× (n=44)              |
| abs_veto_short    | ungated    | 252                               | 15                      | 6.0%  | 0.64×                     |

Every ungated cell — the only ones with enough signal to mean anything — sits between **0.44× and 1.21× chance**, and four of the five are below chance. The mechanical rows all read 1.7-2.0×, which looks encouraging until you notice they are built on 19-44 signal-minutes across a whole week — that is 8, 3, 4 and 7 raw observations landing in the windows. There is a real and slightly hopeful hint in there (the filtered versions are more concentrated near runs than the raw versions, consistently, across four independent gates), and I am flagging it rather than dismissing it — but at these counts it is one or two coincidences wide. **The honest statement is that none of our six gates is meaningfully more likely to be live before a big run than at a random moment of the week.**

One caveat stated plainly: these are signal-minutes, not trades. *grind\_long* showing 1,118 ungated signal-minutes does not mean 1,118 entries — the live desk holds one position per slot and was router-benched for stretches of the week.

The column measures how often the condition was TRUE, which is the right measure for this particular question and the wrong one for any other.

## Where the money actually is: the timing-decay ladder

This is the table I would pin on the wall. It takes hindsight completely for granted — you know which runs, you know the direction — and asks only **how early you enter**. Same 57 runs, same Lot A 1.5R + Lot B wide-chandelier exit, same tick-honest repricing, entry stepped backwards from the ignition bar.

| ENTER THIS EARLY | N  | NET \$    | \$ / RUN | WIN % | STRIP BEST 3 | US SESSION \$ | OVERNIGHT \$ |
|------------------|----|-----------|----------|-------|--------------|---------------|--------------|
| at ignition (0s) | 57 | +3,799.00 | +66.60   | 75%   | +3,162.50    | +2,371.50     | +1,427.50    |
| 30 seconds early | 57 | +1,678.00 | +29.40   | 54%   | +1,315.50    | +1,063.50     | +614.50      |
| 1 minute early   | 57 | +1,377.00 | +24.20   | 47%   | +796.50      | +1,065.00     | +312.00      |
| 2 minutes early  | 57 | +1,603.50 | +28.10   | 49%   | +1,005.00    | +1,121.00     | +482.50      |
| 3 minutes early  | 57 | +1,010.50 | +17.70   | 42%   | +472.00      | +733.00       | +277.50      |
| 4 minutes early  | 57 | +1,235.50 | +21.70   | 47%   | +722.50      | +439.00       | +796.50      |
| 5 minutes early  | 57 | +389.50   | +6.80    | 40%   | +84.50       | −254.00       | +643.50      |
| 7 minutes early  | 57 | −1,073.00 | −18.80   | 21%   | −1,387.00    | −625.00       | −448.00      |
| 9 minutes early  | 57 | −140.50   | −2.50    | 33%   | −452.50      | +93.00        | −233.50      |
| 10 minutes early | 57 | −267.50   | −4.70    | 32%   | −776.50      | +47.50        | −315.00      |

### THE WHOLE SECTION IN THREE SENTENCES

Enter at the ignition bar with perfect direction and these 57 runs pay **\$66.60 each**, 75% of them green. The money halves within thirty seconds, holds a soft shelf worth **\$18-\$29 a run out to four minutes**, falls to **+\$6.80** at five minutes, and is **negative from seven minutes out** — with strip-best-3 negative in every row past five minutes.

Now put the gates on that ladder. When they do fire, they fire a **median of 6.6 minutes early mechanically and 6.7 minutes early ungated** (mechanical range 3.0–9.6 minutes). That is the exact rung where the ladder crosses zero and starts paying **−\$18.80 a run**. **The six gates are not losing at these runs because their exits are wrong or their thresholds are**



**wrong. They are firing in the one part of the window where being right about the direction is worth nothing.**

One honest difference from the last time this ladder was run: **the shelf out to four minutes is new.** Last week the curve was through zero by two minutes; this week there is real money at two, three and four minutes early (+\$28.10, +\$17.70, +\$21.70 a run). I do not think that is a discovery, I think it is a quieter tape with slower ignitions, and the fact that it moved this much on a one-week sample is itself the caveat — do not build anything on the exact shape of this curve. What has replicated across both weeks is the two ends: **entering at the ignition bar is worth several times entering a minute before it, and entering seven minutes before it is worth less than nothing.**

## **Regime segmentation — scored on home turf, as the discipline requires**

Nothing above should be read as one number sprayed across the whole tape, so here is the same population cut by regime at ignition. Buckets key on entry ATR and the 30-bar efficiency ratio measured on the last completed bar before the run, plus a 60-minute range-break test — not on the clock. This week's ATR quartiles across the sat-out runs came out at **8.3pt (p25)** and **15.1pt (p75)**, and those are the cut points. Both are dramatically lower than last week's 12.3 / 25.4, which tells you the headline fact about this week's tape on its own.

| REGIME           | DEFINITION (AT THE LAST BAR BEFORE IGNITION)            | SAT-OUT RUNS | UP / DN |
|------------------|---------------------------------------------------------|--------------|---------|
| NORMAL-CHOP      | ER < 0.25, ATR between 8.3 and 15.1pt                   | 24           | 11 / 13 |
| VIOLENT-WHIPSAW  | ER < 0.25 but ATR above 15.1pt — big range, no progress | 11           | 3 / 8   |
| DEAD-CHOP        | ER < 0.25 and ATR below 8.3pt                           | 10           | 5 / 5   |
| BUILDING         | ER 0.25–0.45 — direction forming, not proven            | 9            | 3 / 6   |
| TREND (no break) | ER ≥ 0.45, price still inside the prior hour's range    | 2            | 1 / 1   |
| CLEAN-TREND      | ER ≥ 0.45 and price outside the prior hour's range      | 1            | 0 / 1   |

**Only 3 of the 57 big runs ignited out of a proven trend (ER ≥ 0.45), and exactly one out of a clean one.** Forty-five ignited out of chop — normal or dead

— or violent whipsaw, and the remaining nine out of BUILDING, which is direction forming but not yet proven. That is the same uncomfortable finding as last week, and it is now two weeks old rather than one: the runs are not the continuation of something we could see. They are ruptures in tape that looked like nothing.

| REGIME           | RUNS | MECH FIRES | MECH NET \$ | MECH WIN % | UNGATED FIRES | UNGATED NET \$ | UNGATED WIN % |
|------------------|------|------------|-------------|------------|---------------|----------------|---------------|
| NORMAL-CHOP      | 24   | 1          | −76.00      | 0%         | 14            | +255.50        | 50%           |
| VIOLENT-WHIPSAW  | 11   | 3          | −115.50     | 33%        | 4             | −192.50        | 25%           |
| DEAD-CHOP        | 10   | 0          | 0.00        | —          | 5             | +98.00         | 40%           |
| BUILDING         | 9    | 2          | +207.00     | 100%       | 4             | +218.50        | 75%           |
| TREND (no break) | 2    | 0          | 0.00        | —          | 0             | 0.00           | —             |
| CLEAN-TREND      | 1    | 0          | 0.00        | —          | 0             | 0.00           | —             |

Three honest reads. **(1)** Every bucket is between n=0 and n=14, and I am not going to pretend a policy can be built from that. **(2)** The one consistent sign is the negative one: **VIOLENT-WHIPSAW loses in both modes** — three mechanical fires at −\$115.50 and four ungated at −\$192.50, and it holds the two worst individual trades in the study (grind −\$120 and abs\_veto\_short −\$91). If any regime rule survives this section it is a bench rule, not an arm rule. **(3)** Whipsaw was also the losing bucket last week, which is the one cross-week agreement in the regime table — but note it is the opposite of what the blind ladder says (below), where whipsaw is the best regime. Those two facts together say something specific: **whipsaw runs are worth a lot if you are in them and expensive if you are guessing at them.** The gates are guessing.

The blind control — perfect direction, one minute early, same exit — cut the same way:

| REGIME                   | N  | NET \$    | \$ / RUN | WIN % |
|--------------------------|----|-----------|----------|-------|
| VIOLENT-WHIPSAW          | 11 | +620.00   | +56.40   | 73%   |
| TREND (no break)         | 2  | +282.00   | +141.00  | 50%   |
| DEAD-CHOP                | 10 | +211.00   | +21.10   | 40%   |
| NORMAL-CHOP              | 24 | +198.50   | +8.30    | 42%   |
| BUILDING                 | 9  | +98.50    | +10.90   | 44%   |
| CLEAN-TREND              | 1  | −33.00    | −33.00   | 0%    |
| US session (13:30-20:00) | 26 | +1,065.00 | +41.00   | 58%   |
| Overnight / pre-open     | 31 | +312.00   | +10.10   | 39%   |

Time-of-day is the cleaner split than regime: with perfect direction a US-session run is worth **four times** an overnight one (\$41.00 vs \$10.10 a run) at more than

half again the win rate. That is a real and stable-looking asymmetry and it belongs in Movement 3's design brief.

## The R-target sweep — and what it says about the cheat-sheet guesses

The discipline says every R is a guess until it is swept, so I swept it. Six and 27 fires is not a population you can sweep anything on, so I swept on the blind one-minute-early book instead (perfect direction, 57 runs), with the quiet-tape clip switched **off** — with the clip on, the R target is a no-op on any entry under 22pt ATR and the sweep would be measuring nothing. Both lots at the same R so the number reads cleanly.

| R TARGET | ALL 57 | WIN % | US (26) | OVERNIG HT (31) | DEAD-CHOP (10) | NORMAL-CHOP (24) | WHIPSA W (11) | BUILDIN G (9) |
|----------|--------|-------|---------|-----------------|----------------|------------------|---------------|---------------|
| 0.5R     | −182   | 63%   | +27     | −209            | −67            | −175             | +132          | −48           |
| 0.75R    | +28    | 56%   | +293    | −265            | −110           | −162             | +266          | +24           |
| 1.0R     | +229   | 53%   | +411    | −182            | −84            | −19              | +403          | −115          |
| 1.5R     | +751   | 49%   | +790    | −39             | −33            | +20              | +675          | −24           |
| 2.0R     | +1,475 | 49%   | +1,274  | +201            | +18            | +251             | +956          | +69           |
| 2.5R     | +2,093 | 47%   | +1,781  | +312            | +68            | +355             | +1,258        | +162          |
| 3.0R     | +2,324 | 44%   | +1,953  | +371            | +119           | +405             | +1,229        | +253          |
| 4.0R     | +3,585 | 44%   | +2,811  | +774            | +220           | +791             | +1,751        | +438          |
| 6.0R     | +4,349 | 44%   | +3,071  | +1,278          | +424           | +1,103           | +1,928        | +509          |

**What is PROVEN here, and what is not.** This is a hindsight-direction population, so it cannot prove an R for a live gate — nothing here says "set abs\_veto\_short to 4R". What it can prove, because the entries are identical in every row and only the target moves, is **the shape of exit that run tape rewards**, and that shape is unambiguous: **the curve rises monotonically from 0.5R all the way to 6.0R while the win rate falls from 63% to 44%**. There is no plateau below 4R and no turnover anywhere. Every single regime column is negative at 0.5R except whipsaw, and every single one is positive at 4R. This is the identical finding as last week on a completely different tape, which is the closest thing to replication this section has produced.

| CHEAT-SHEET GUESS               | WHAT RUN TAPE SAYS                                                                                                                                                                                               | VERDICT ON THIS POPULATION |
|---------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------|
| Faders: Lot A 0.5R / Lot B 1.5R | 0.5R is the <b>worst cell in the sweep</b> — −\$182 on 57 runs at a 63% win rate. It wins often and loses money, the classic small-target signature. This is the second consecutive week 0.5R is the worst cell. | Wrong shape for run tape   |

|                                   |                                                                                              |                        |
|-----------------------------------|----------------------------------------------------------------------------------------------|------------------------|
| Momentum: Lot A 1.5R / Lot B 2.5R | +\$751 and +\$2,093 — both positive, both a long way below the 4R cell.                      | Safe but leaving money |
| Trend: Lot A 2.5R / Lot B wide    | +\$2,093 rising to +\$3,585 at 4R and +\$4,349 at 6R. The wide half is the correct instinct. | Directionally right    |

**Robustness on the best cell, and this one holds up better than last week's.** Leave-one-day-out on the 6.0R headline of +\$4,349: without 08-10 **+\$2,882**, without 08-11 **+\$3,662**, without 08-12 **+\$4,204**, without 08-13 **+\$3,239**, without 08-14 **+\$3,409**. **Positive on all five folds, worst fold still 66% of the headline** — no single session carries it, which is a genuinely better result than last week's 49% swing. The finding remains "small targets are wrong on run tape", not "use 6R" — a 6R target on a real gate with a real win rate is a different animal from a 6R target on hindsight-perfect direction.

## The three live filters that moved money on this population

### grind\_long's 08-08 config change — re-fired under every cell we have shipped

Last week this section named `ext_hi 2.0` as "the single filter standing between us and these runs". The ceiling was deleted on 08-08 and the ATR floor went  $10 \rightarrow 22$  at the same time, so I re-fired grind on this week's 23 up-runs under four cells — identical entries logic, identical exits — to see which knob actually binds.

| CELL                 | EXT CEILING        | ATR FLOOR | FIRES | NET \$  | WIN % | RUNS BLOCKED BY THE ATR FLOOR |
|----------------------|--------------------|-----------|-------|---------|-------|-------------------------------|
| deployed (08-08)     | none (4.0 default) | 22        | 1     | −120.00 | 0%    | 4                             |
| wk32 cell (08-01)    | 2.0                | 10        | 1     | −120.00 | 0%    | 1                             |
| ext_hi 3.0 candidate | 3.0                | 22        | 1     | −120.00 | 0%    | 1                             |
| no ATR floor at all  | none (4.0 default) | 0         | 5     | −178.00 | 20%   | 0                             |

Two clean conclusions, and they both go against last week's write-up.

- **The extension ceiling is not the binding constraint on run tape.** 2.0, 3.0 and "none" all produce the same single fire and the same – \$120. Deleting it bought us nothing on this population. The clause that actually kills grind is "VWAP slope not up", on 14 of 23 runs, and no setting of ext\_hi touches that.
- **The ATR≥22 floor is the only knob doing work, and it is doing GOOD work.** Remove it and grind takes four extra trades — –\$86, –\$31, +\$97, –\$38 — for a net –\$58. The four runs it blocked had entry ATRs of 20.6, 6.8, 9.5 and 8.7 points. That is the floor stopping a trend-continuation gate from buying in dead tape, which is precisely what it was reverted to do on 08-08.

The one fire that survives every cell is 08-10 13:56 — a +109pt up-run out of VIOLENT-WHIPSAW at 29.1pt ATR. Grind got in six minutes early and both lots stopped for –\$120. Being allowed to trade the biggest whipsaw of the week was not the same as being able to trade it.

### The 55-second absorption veto — a small net positive this week

Mechanical and ungated fire from the identical bar and differ only by the 55-second wait plus the absorption test, so the two can be separated.

| RUN         | MOVE | UNGATED<br>(ENTER NOW) | MECHANICAL | EFFECT OF THE<br>VETO                  |
|-------------|------|------------------------|------------|----------------------------------------|
| 08-10 00:22 | –59  | +97.00                 | vetoed     | –97.00                                 |
| 08-12 15:30 | –47  | –69.00                 | vetoed     | +69.00                                 |
| 08-13 15:43 | –61  | –78.00                 | vetoed     | +78.00                                 |
| 08-10 15:20 | –62  | –91.00                 | –91.00     | 0.00 — the 55s wait<br>changed nothing |
| Net effect  | –    | –                      | –          | +\$50.00                               |

Three blocks: two losers and one winner, netting **+\$50 saved**. And on the one thrust that survived the test, the 55-second delay cost **exactly nothing** — both entries stopped for the same –\$91. That is a different verdict from last week, where the wait was the villain (it cost \$596.50 on two winners). I

am flagging the disagreement rather than papering over it: **on four fires in one week the veto is mildly positive, on five fires the week before it was strongly negative, and neither sample can settle it.** The Saturday job is unchanged — sweep the wait at 0/15/25/40 seconds keeping the absorption test intact, on the FULL book, not on run tape.

### The quiet-tape clip — a chop tool standing in the road, again

Every gate in `data/exit_overrides.json` carries the 08-02 quiet-tape clip: when entry ATR is below 22 points, Lot A banks a flat \$40 and Lot B banks the larger of 1.75R and \$60, and nothing else in the exit stack gets to run. I priced identical entries with the clip on and off.

| MEASURE                                              | AT IGNITION | 1 MINUTE EARLY | 5 MINUTES EARLY |
|------------------------------------------------------|-------------|----------------|-----------------|
| Runs where the clip was active at entry (ATR < 22pt) | 55 of 57    | 55 of 57       | 55 of 57        |
| Clip ON (as deployed)                                | +\$3,799.00 | +\$1,377.00    | +\$389.50       |
| Same entries, same exit, clip OFF                    | +\$4,929.50 | +\$2,289.00    | +\$1,222.00     |
| Cost of the clip on run tape                         | -\$1,130.50 | -\$912.00      | -\$832.50       |

**Fifty-five of fifty-seven runs this week ignited from tape quiet enough to trigger the clip** — up from 35 of 53 last week, because this was a much quieter week (median peak pre-run ATR 12.0pt, against a clip threshold of 22). So on precisely the trades where a run is about to happen, the desk is wearing an exit that caps Lot A at forty dollars. The clip is doing exactly what it was built to do; nobody ever told it that a 12-point-ATR tape can produce a 236-point run. **This is a regime carve-out question, not a delete question** — the clip earns its keep on the whole book, which is where it was justified, and this population is selected for the one situation that makes it look worst. But it is now the second consecutive week it has cost roughly a thousand dollars on run tape, and that is worth a shadow.

# The tape shape — the number that explains the whole null

Separately from any gate, I measured what the ten minutes before each run actually looked like: the peak ATR reached, the furthest price ever got from VWAP, the biggest five-bar thrust. This is the honest physical answer to "why didn't anything fire".

| WHAT THE PRE-RUN TAPE LOOKED LIKE                                                     | VALUE          |
|---------------------------------------------------------------------------------------|----------------|
| Peak ATR in the 10 minutes before ignition (median of 57 runs)                        | 12.0pt         |
| — lowest / highest across the 57                                                      | 5.5pt / 37.7pt |
| Up-runs that EVER reach grind_long's ATR≥22 floor                                     | 3 of 23        |
| Down-runs that EVER reach rgv_short's ATR≥20 floor                                    | 6 of 34        |
| Runs that EVER reach capitulation_long's ATR≥10 floor                                 | 35 of 57       |
| Runs where a 5-bar thrust EVER reaches 1.5 ATR (abs_veto's trigger, either direction) | 42 of 57       |
| Peak distance from VWAP (median)                                                      | 2.8 ATR        |

There it is in one line. **Three of the 23 up-runs and six of the 34 down-runs ever get anywhere near the magnitude floors our two floor-carrying gates require.** Those floors were derived on the whole book, where they are right — the ATR≥22 revert alone was worth thousands on the live book, and it blocked losers here too. But on the specific population of "moves we sat out", they are structurally incompatible with the tape: **the runs come out of quiet, and the floors exist to keep us out of quiet.** That is not a bug and it is not a bug to fix by lowering the floors — the no-floor grind test above shows lowering them loses money. It is the reason a NEW instrument is needed, which is Movement 3's brief.

## The verdict

**HONEST NULL — AND HERE IS WHAT IT IS A NULL ABOUT**

**DON'T ARM.** There is no version of "leave the six gates on and let them catch the misses" that makes money this week. Mechanically it is **+\$15.50 on six fires** — nominally green, but strip the single best trade and it is –\$94.50, and it rests entirely on one Tuesday (leave 08-11 out and it is –\$191.50). Ungated it is **+\$379.50 on 27 fires**, which is **+\$69.50** once three trades come out and **+\$6.50** once 08-11 comes out. Neither is a book you would put a lot behind, and both are in-sample on a hindsight-selected population with no arming rule attached.

But the null is specific, not general, and the specific part is where next week's work is:

- **It is not an exit problem.** Proven by the timing ladder: give the same exits perfect entries at the ignition bar and these 57 runs pay **+\$3,799 at 75% win**. The exits are fine. (The one exit finding that is real is the quiet-tape clip, worth \$832-\$1,130 on this tape depending on entry timing.)
- **It is not a threshold problem.** Proven by the ungated book: strip every filter off all six gates and **33 of 57 runs still draw no fire at all**, carrying \$4,167 of the \$7,619 ceiling and nine of the fifteen biggest moves. And the filters we do have were, on balance, helping this week — grind's ATR floor saved \$58, the 55s veto saved \$50, capitulation's ATR floor saved \$42; only `require_flip` cost anything (\$158.50).
- **It is a timing-and-mechanism problem.** Proven three ways: the gates fire at 0.44-1.21× chance in the pre-run windows once you have enough signal to measure, four of five below chance; when they do fire it is a median of 6.6-6.7 minutes early, where the blind ladder pays — \$18.80 a run; and the tape they would have to fire in has a median peak ATR of 12 points against floors of 20 and 22. **Our six gates are three continuation gates that need the move to have already happened and three fade gates that need a spent move to sell into. A big run is neither of those things until it has started.**
- **The one green thread worth pulling.** The three mechanical winners were all SHORT FADES, all with both lots reaching target, two of them fired 9-10 minutes early. Fades do not care about timing the way momentum does, because they are aimed at a price rather than at a moment. That is three trades and I will not dress it up — but it is the only positive mechanism in the section and it points somewhere.

## Disposition — every lead touched in this movement

| LEAD                                         | VERDICT | WHY / WHAT WOULD REVIVE IT                                                                                                                                                                                                                                             |
|----------------------------------------------|---------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| "A better EXIT would have caught the misses" | REFUTED | Named test: the timing-decay ladder, n=57 at every rung, tick-honest, on this week's tape. Perfect entries at ignition pay +\$3,799 / 75% under the existing exits. Second consecutive week this test has returned the same answer. The exits are not what is missing. |



|                                                                                            |         |                                                                                                                                                                                                                                                                                                                                                                                   |
|--------------------------------------------------------------------------------------------|---------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| "grind_long's ext_hi ceiling is what stands between us and these runs" (last week's claim) | REFUTED | Named test: the four-cell grind A/B on the same 23 up-runs. ext_hi 2.0, 3.0 and none all give the identical single fire and the identical –\$120. The binding clause is "VWAP slope not up" on 14 of 23 runs, which no ext setting touches. Last week's number came from a cell where the ATR floor was 10; with the floor at 22 the ceiling is not reachable in the first place. |
| grind_long ATR floor 10 → 22 (the 08-08 revert)                                            | LIVE    | <b>Not an action — it is already deployed and this is the confirmation. On run tape the floor blocks four trades worth –\$58 net at entry ATRs of 20.6 / 6.8 / 9.5 / 8.7pt. HOLD. Note honestly that this is a \$58 confirmation on 4 trades, not independent support for the \$4,151 full-book number.</b>                                                                       |
| The two SHORT FADERS (exhaustion_short, rgv_short) as run-catchers                         | SHADOW  | <b>3 fires, 3 winners, +\$302.50, both lots to target on all three, two of them fired 9-10 minutes ahead and still won. n=3, so it proves nothing — but it is free to watch. Shadow: log every mechanical fade fire near a census run and accumulate n. Revive as a real candidate at n≥20.</b>                                                                                   |
| Arm all six gates mechanically at the misses                                               | PARKED  | +\$15.50 on 6 fires; –\$94.50 stripped of one; –\$191.50 without 08-11. Revive only if a MECHANISM changes — more n cannot save six fires from 57 runs. What would revive it: a gate that can fire inside the first two minutes of a move.                                                                                                                                        |
| Run all six ungated (filters off)                                                          | PARKED  | +\$379.50 → +\$69.50 stripped of 3 → +\$6.50 without 08-11. Revive at n>60 fires across at least three weeks with an out-of-sample leg and a regime rule attached. As it stands there is no policy here, only 27 trades and one good Tuesday.                                                                                                                                     |
| capitulation_long require_flip on run tape                                                 | SHADOW  | <b>Blocked 3 of the gate's 6 ungated pre-run fires, two of them \$97 winners — cost \$158.50 here. Does NOT overturn the 08-02 full-book re-derivation (capit_loose – \$5,250 / 238 fires, ordering-correct). Shadow a "flip OR 90-second timeout" variant so the gate can still take the flush when buyers are slow to show, and read it against the whole book.</b>             |
| Quiet-tape clip (ATR<22) on run tape                                                       | SHADOW  | <b>Active on 55 of 57 runs, costs \$832-\$1,130 with identical entries depending on entry lead. Second consecutive week at roughly a grand. Shadow a version that suspends the clip when the ignition context is a range-break, and read it against the whole book where the clip earns.</b>                                                                                      |
| abs_veto 55-second WAIT (as distinct from the test)                                        | PARKED  | Last week: the wait cost \$596.50 on two winners. This week: it cost exactly \$0 on its one survivor, and the veto as a whole SAVED \$50. The two weeks disagree, which means neither is evidence. Revive via the proper test — sweep 0/15/25/40s on the FULL book, not on run tape.                                                                                              |
| rgv_short ATR≥20 floor                                                                     | PARKED  | Stands the gate down on 22 of 34 eligible runs, and only 6 of those 34 ever reach 20pt ATR at all. But removing it moves the binding clause to "never stretched above VWAP" (16 runs) and the gate still fires only twice, both winners. Revive only alongside a stretch-condition rethink — the floor is not the constraint, the geometry is.                                    |
|                                                                                            | SHADOW  |                                                                                                                                                                                                                                                                                                                                                                                   |

|                                                                                  |                                                                                                                                                                                                                                                                                                                                                                                  |
|----------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <p>exhaustion_short's "the flow DID move price" clause as a BENCH instrument</p> | <p>It refused 10 of 34 runs on exactly that test — heavy aggressor flow that IS moving price. That is a cheap, already-implemented, 20-second detector of "a real move is underway". Useless as an entry (it says don't fade); potentially valuable inverted, as an arm/bench signal. Free to shadow, and it is the same idea Movement 3 needs.</p>                              |
| <p>VIOLENT-WHIPSAW as a place to arm anything</p>                                | <p>LOSES in both modes (−\$115.50 mech / −\$192.50 ungated) and owns the two worst trades in the study. Second consecutive week it is the losing bucket. But the blind ladder says whipsaw is the BEST regime (+\$56.40/run, 73%), so the mechanism is "worth being in, not worth guessing at". Thin (11 runs), so parked not refuted. Revive as a bench rule at n&gt;25.</p>    |
| <p>Small R targets (0.5R Lot A) on run tape</p>                                  | <p>Named test: the 9-cell R sweep on 57 blind entries, clip off, LODO across five days. 0.5R is the worst cell in the sweep at −\$182 with a 63% win rate, negative overall and in three of the four regime columns (whipsaw is the only exception). It was also the worst cell last week on a different tape. Wins often, loses money — the classic small-target signature.</p> |

## THE ONE STONE STILL UNTURNED

Everything in this movement tested entries **before** ignition, because that is what "could we have caught it" normally means. The decay ladder says the highest-paying rung by a distance is the one at zero (+\$66.60 a run, 75% green) and that the money is still real at thirty seconds (+\$29.40) and holds a shelf out to four minutes. We have never built a gate whose job is to arm on the **first thirty seconds of a move that has already started**, on 5-second bars, with no attempt to predict anything.

Two things this week added to that. First, **the three winners here were all fades**, and fades were the only entries whose money survived being early — because a fade aims at a price, not at a moment. Second, exhaustion\_short handed us a ready-made ignition detector by accident: it refused ten runs because the aggressive flow WAS moving price. Inverted, that clause is a twenty-second "a real move is underway" test that already runs on this desk's tape, costs nothing, and fires exactly where the ladder says the money is. Combine that with the US-session skew (+\$41.00 a run vs +\$10.10 overnight) and you have the shape of the thing Movement 3 should be trying to build.

Honest scope on all of it: **one week, five trading days, 57 runs, in-sample, hindsight-selected, on the quietest tape of the summer.** Nothing in this section is a Saturday deploy.

## **M3** The greenfield lab — the gold book-break hunt, and the four hunts that never ran

### ★ READ THIS FIRST — WHAT IS HERE AND WHAT IS NOT

This movement was scheduled as **five separate hunts**: one per MNQ cause-cluster (VACUUM, FLOW-LED, OPEN-NEWS), a chop-day scalp lab, and a dedicated gold hunt. **Four of the five never started.** The serial runner ran out of budget after the rehabilitation phase overran — BUDGET: stopping section builds to protect the tail (0m left) — and the synthesis phase that reads all five never ran either. So **there is no MNQ greenfield hunt in this report.** That is a gap, not a null: nobody looked, so nobody found nothing. It is named in the disposition table and it is the first thing to reschedule.

The gold hunt did get far enough to save its signal file before it died, and that file is a real dataset: **397 MGC book breaks across 21 sessions** from 16 July, each with the order-book features it was built to test and the forward outcome. It is scored below to the full discipline — placebo, strip-the-best-three, leave-one-day-out, the size escalation, the excursion statistic and the book-separation table. **Additional gold work that DID complete is in Part 2.5 Part A** (the MGC day rider, the coil bouncer, the London-open break fade) — that is where the two gold SHADOW candidates of the week live. This section is the fourth gold idea, and it is the one that failed.

The idea was a good one and it is the natural thing to try with L2 on a new instrument. Gold breaks a level; look at the order book at the moment it

breaks; if the far side is empty — nothing resting in the way — the move should run, and if the far side is stacked it should stall. That is the **far-side-depletion** hypothesis, it is exactly what Movement 1 asks of the book on MNQ, and on 397 gold breaks **it is not just absent, it is backwards**. The raw signal makes  $-\$291.00$  at  $-\$0.73$  a go. The book features it was built on select the LOSERS. And the one feature that does separate — resting book imbalance — only works in a middle band and evaporates the moment you narrow to the biggest moves, which is the opposite of what the escalation was supposed to reveal.

## 1 · The raw gate, before anything is filtered

| THE POPULATION                | SIGNALS | SESSIONS      | NET \$           | \$/SIGNAL | WIN % | STRIP THE BEST 3 | WORST LEAVE-ONE-DAY-OUT |
|-------------------------------|---------|---------------|------------------|-----------|-------|------------------|-------------------------|
| every book break, as detected | 397     | 22 (12 green) | <b>-\$291.00</b> | $-\$0.73$ | 48.4% | $-\$1,496.00$    | $-\$1,594.00$           |

Every break the detector found on MGC from 16 July to 14 August, priced at \$10.00 a point with \$1.50 off each round trip. It is a coin flip that pays the fee — and strip its best three trades and it is  $-\$1,496.00$ , so even the coin flip is carried by a handful of prints.

## 2 · Twenty candidate filters, and the two that are worth reading

Each row is a single mechanical rule applied to the same 397 breaks. The **placebo percentile** is against 2,000 random subsets of the same size: below 95 means a coin could have picked as well, and **below about 5 means the filter is actively selecting the wrong signals**, which turns out to be the interesting column here.

| CANDIDATE FILTER                    | N   | NET \$     | \$/SIGNAL | WIN % | STRIP-3     | WORST LODO  | PLACEBO %ILE |
|-------------------------------------|-----|------------|-----------|-------|-------------|-------------|--------------|
| book imbalance with us (imb > 0.10) | 134 | \$1,795.00 | \$13.40   | 50.0% | \$590.00    | \$970.00    | 96.2         |
| ASIA session only                   | 146 | \$596.00   | \$4.08    | 51.4% | $-\$218.50$ | $-\$306.00$ | 74.7         |
| empty far side AND not Asia         | 123 | \$232.50   | \$1.89    | 48.0% | $-\$972.50$ | $-\$198.50$ | 61.0         |
| efficiency >= 0.50 at the break     | 21  | \$198.50   | \$9.45    | 47.6% | $-\$396.50$ | $-\$140.50$ | 67.0         |

| CANDIDATE FILTER                                  | N   | NET \$      | \$/SIGNAL | WIN % | STRIP-3     | WORST LODO  | PLACEBO %ILE |
|---------------------------------------------------|-----|-------------|-----------|-------|-------------|-------------|--------------|
| LONDON session only                               | 129 | \$163.00    | \$1.26    | 49.6% | −\$835.00   | −\$537.50   | <b>59.6</b>  |
| book imbalance strongly with us (imb > 0.25)      | 41  | \$51.00     | \$1.24    | 43.9% | −\$880.50   | −\$438.50   | <b>56.4</b>  |
| far side depleted vs pre (obst < 50% of obst_pre) | 200 | \$9.00      | \$0.04    | 49.0% | −\$1,157.50 | −\$525.50   | <b>55.8</b>  |
| support behind us (support > obstacle)            | 146 | −\$45.00    | −\$0.31   | 49.3% | −\$1,041.00 | −\$550.50   | <b>53.6</b>  |
| ATR >= 3.0pt                                      | 45  | −\$353.50   | −\$7.86   | 51.1% | −\$1,353.00 | −\$815.00   | <b>31.5</b>  |
| not DEAD_CHOP                                     | 327 | −\$463.00   | −\$1.42   | 48.0% | −\$1,668.00 | −\$1,206.50 | <b>38.8</b>  |
| LONDON or US (skip Asia)                          | 251 | −\$887.00   | −\$3.53   | 46.6% | −\$2,092.00 | −\$1,491.00 | <b>26.9</b>  |
| US session only                                   | 122 | −\$1,050.00 | −\$8.61   | 43.4% | −\$2,115.50 | −\$1,573.50 | <b>17.6</b>  |
| depletion ratio = 1.0 (fully cleared)             | 141 | −\$1,175.00 | −\$8.33   | 44.7% | −\$2,341.50 | −\$1,464.50 | <b>16.2</b>  |
| empty far side AND trend                          | 49  | −\$1,367.50 | −\$27.91  | 38.8% | −\$2,241.50 | −\$1,668.00 | <b>3.6</b>   |
| efficiency >= 0.35 at the break                   | 114 | −\$1,526.00 | −\$13.39  | 40.4% | −\$2,375.50 | −\$2,206.50 | <b>7.5</b>   |
| CLEAN_TREND only                                  | 73  | −\$1,605.00 | −\$21.99  | 42.5% | −\$2,530.50 | −\$2,025.50 | <b>2.8</b>   |
| empty far side AND ER>=0.35                       | 65  | −\$1,621.00 | −\$24.94  | 35.4% | −\$2,410.00 | −\$2,220.00 | <b>2.8</b>   |
| far side is EMPTY (obstacle = 0)                  | 192 | −\$2,059.50 | −\$10.73  | 43.8% | −\$3,264.50 | −\$2,664.00 | <b>4.8</b>   |
| far side thin (obstacle < 5 lots)                 | 228 | −\$2,091.50 | −\$9.17   | 45.6% | −\$3,296.50 | −\$2,698.00 | <b>3.7</b>   |
| ATR >= 2.0pt (the move has room)                  | 128 | −\$2,262.50 | −\$17.68  | 47.7% | −\$3,328.00 | −\$2,756.00 | <b>1.6</b>   |

One row clears 95 and five rows are below 10. The rows below 10 are the hypothesis itself.

★ **THE FINDING — AN EMPTY BOOK AHEAD OF A GOLD BREAK IS A WARNING, NOT A GREEN LIGHT**

Read the bottom of that table, not the top. **“Far side is empty” loses −\$2,059.50 over 192 signals and sits at the 4.8th percentile of its placebo.** “Far side thin” is the 3.7th. “Empty far side AND a clean trend” is the 3.6th and loses −\$27.91 a go. Being in the 4th percentile is not noise — it means that of 2,000 random ways to pick 192 of these breaks, 1,904 of them did BETTER. The feature carries real information and **the sign is the**

**other way round:** on gold, a level that breaks into a vacuum is a level nobody is defending because nobody is there, and price comes straight back. That is a refutation of the gate as specified, and simultaneously the most promising thing in this section — it is a fade lead, not a break lead, and it is PARKED as one.

### 3 • The one filter that separates — and why it is still not a gate

Resting book imbalance is the exception. Skew the book our way at the moment of the break and the population turns from  $-\$291$  to  $+\$1,795$  over 134 signals, at the 96th percentile of its placebo, still  $+\$590$  after stripping its best three and still  $+\$970$  after dropping its best day. That is the only row in the study that survives all four checks. Then you cut it finer and it falls over:

| BOOK<br>IMBALANCE<br>AT THE<br>BREAK | N   | NET \$      | \$/SIGNAL | WIN % | STRIP-3     | WORST<br>LODO |
|--------------------------------------|-----|-------------|-----------|-------|-------------|---------------|
| book AGAINST us<br>(imb < -0.10)     | 101 | -\$315.00   | -\$3.12   | 45.5% | -\$1,077.50 | -\$771.00     |
| book neutral (-0.10<br>to +0.10)     | 161 | -\$1,703.00 | -\$10.58  | 49.1% | -\$2,459.00 | -\$2,271.50   |
| book with us<br>(+0.10 to +0.25)     | 91  | \$1,758.50  | \$19.32   | 51.6% | \$691.00    | \$902.00      |
| book strongly with<br>us (> +0.25)   | 44  | -\$31.50    | -\$0.72   | 45.5% | -\$963.00   | -\$521.00     |

**This is the killer column and it is the fourth row.** If book imbalance were a real mechanism, more of it would be better. It is not: the +0.10 to +0.25 band makes  $\$19.32$  a signal and the band ABOVE it — the book most strongly on our side — loses money. An edge that lives in one middle bucket and dies on either side of it is a bucket, not a mechanism.

### 4 • The escalation — full census, then the top 25, then the top 15

The scope's instruction is explicit and it exists because marginal moves wash out the footprint the monsters might carry: if nothing works on everything, narrow to the biggest moves and hunt again. **Here the escalation runs the wrong way,** which is a result in itself.

FULL population

| RULE                                | N   | NET \$     | \$/SIGNAL | WIN % | STRIP-3     | PLACEBO %ILE |
|-------------------------------------|-----|------------|-----------|-------|-------------|--------------|
| no filter — FULL population         | 397 | −\$291.00  | −\$0.73   | 48.4% | −\$1,496.00 | —            |
| book imbalance with us (imb > 0.10) | 134 | \$1,795.00 | \$13.40   | 50.0% | \$590.00    | 97.0         |
| ASIA session only                   | 146 | \$596.00   | \$4.08    | 51.4% | −\$218.50   | 74.7         |
| empty far side AND not Asia         | 123 | \$232.50   | \$1.89    | 48.0% | −\$972.50   | 62.2         |
| efficiency >= 0.50 at the break     | 21  | \$198.50   | \$9.45    | 47.6% | −\$396.50   | 67.7         |
| LONDON session only                 | 129 | \$163.00   | \$1.26    | 49.6% | −\$835.00   | 61.2         |

TOP 25 by move size

| RULE                                  | N  | NET \$    | \$/SIGNAL | WIN % | STRIP-3     | PLACEBO %ILE |
|---------------------------------------|----|-----------|-----------|-------|-------------|--------------|
| no filter — TOP 25 by move size       | 25 | −\$696.50 | −\$27.86  | 44.0% | −\$1,901.50 | —            |
| book imbalance with us (imb > 0.10)   | 10 | \$992.50  | \$99.25   | 60.0% | −\$212.50   | 94.7         |
| empty far side AND not Asia           | 10 | \$288.00  | \$28.80   | 50.0% | −\$917.00   | 79.8         |
| LONDON session only                   | 9  | −\$287.00 | −\$31.89  | 44.4% | −\$1,285.00 | 46.5         |
| depletion ratio = 1.0 (fully cleared) | 10 | −\$288.50 | −\$28.85  | 40.0% | −\$1,455.00 | 51.0         |
| ASIA session only                     | 9  | −\$333.00 | −\$37.00  | 44.4% | −\$1,147.50 | 43.4         |

TOP 15 by move size

| RULE                              | N  | NET \$      | \$/SIGNAL | WIN % | STRIP-3     | PLACEBO %ILE |
|-----------------------------------|----|-------------|-----------|-------|-------------|--------------|
| no filter — TOP 15 by move size   | 15 | −\$630.00   | −\$42.00  | 40.0% | −\$1,835.00 | —            |
| LONDON or US (skip Asia)          | 11 | −\$54.00    | −\$4.91   | 45.5% | −\$1,259.00 | 74.5         |
| far side is EMPTY (obstacle = 0)  | 10 | −\$267.50   | −\$26.75  | 40.0% | −\$1,472.50 | 70.8         |
| far side thin (obstacle < 5 lots) | 12 | −\$321.00   | −\$26.75  | 41.7% | −\$1,526.00 | 67.7         |
| not DEAD_CHOP                     | 13 | −\$659.50   | −\$50.73  | 38.5% | −\$1,864.50 | 36.9         |
| ATR >= 2.0pt (the move has room)  | 10 | −\$1,063.50 | −\$106.35 | 30.0% | −\$2,129.00 | 16.7         |

**Read the three tables together.** On the full population the imbalance rule makes +\$1,795 and survives strip-3 at +\$590. Narrow to the 25 biggest moves and it still shows +\$992.50 — but strip its best three and it is **−\$212.50**, i.e. the whole thing is three prints. Narrow again to the top 15 and **every single rule is negative**; the best of them loses −\$54.00. The footprint does not get cleaner as the moves get bigger. It gets worse. Whatever separates gold breaks, it is not something that concentrates in the monsters — and that is a genuinely useful thing to know before anyone spends another night on it.

## 5 • The excursion statistic — why no stop can save this

| WHAT WAS MEASURED                    | VALUE            | WHAT IT MEANS                                                        |
|--------------------------------------|------------------|----------------------------------------------------------------------|
| median favourable excursion          | <b>3.28×ATR</b>  | the typical break does eventually go our way by about three ATR      |
| median adverse excursion             | <b>−3.04×ATR</b> | ... and goes against us by almost exactly the same amount first      |
| signals reaching 2×ATR our way       | 66.0%            | two-thirds do offer a real move at some point                        |
| signals going 1×ATR against us first | <b>80.4%</b>     | four-fifths hurt before they help — this is the number that kills it |
| signals that 'ran'                   | 48.4%            | a coin flip, by the detector's own definition                        |

**Median MFE 3.28 ATR against median MAE −3.04 ATR is a symmetric distribution.** There is no asymmetry to build a stop around: any stop tight enough to make the losers small takes out 80% of the population before it works, and any stop wide enough to survive costs more than the winners pay. This single table is why the exit rehabilitation that rescued `exhaustion_short` in Part 1.5 cannot rescue this — there, the entry beat 60 of 60 random controls; here, the entry is the coin flip.

## 6 • Day by day — the whole sample, nothing hidden

| SESSION    | BREAKS | NET \$    |
|------------|--------|-----------|
| 2026-07-16 | 18     | −\$366.00 |
| 2026-07-17 | 18     | −\$720.00 |
| 2026-07-20 | 19     | −\$211.50 |
| 2026-07-21 | 17     | −\$75.00  |
| 2026-07-22 | 18     | −\$9.00   |
| 2026-07-23 | 16     | \$72.50   |
| 2026-07-24 | 14     | \$159.00  |
| 2026-07-27 | 19     | \$160.00  |
| 2026-07-28 | 19     | \$174.00  |
| 2026-07-29 | 21     | \$848.00  |
| 2026-07-30 | 19     | −\$754.50 |



| SESSION    | BREAKS | NET \$      |
|------------|--------|-------------|
| 2026-07-31 | 15     | \$115.50    |
| 2026-08-03 | 19     | −\$366.00   |
| 2026-08-04 | 18     | \$108.00    |
| 2026-08-05 | 21     | \$1,303.00  |
| 2026-08-06 | 19     | −\$298.50   |
| 2026-08-07 | 17     | −\$1,076.50 |
| 2026-08-10 | 18     | \$227.00    |
| 2026-08-11 | 16     | \$87.00     |
| 2026-08-12 | 21     | \$42.00     |
| 2026-08-13 | 19     | −\$281.00   |
| 2026-08-14 | 16     | \$571.00    |

12 green sessions of 22. No week, no session and no regime rescues it.

## 7 • Regime and session splits

| CUT             | N   | NET \$      | \$/SIGNAL | WIN % | STRIP-3     |
|-----------------|-----|-------------|-----------|-------|-------------|
| NORMAL_CHOP     | 181 | \$2,156.50  | \$11.91   | 54.1% | \$1,030.50  |
| ASIA            | 146 | \$596.00    | \$4.08    | 51.4% | −\$218.50   |
| DEAD_CHOP       | 70  | \$172.00    | \$2.46    | 50.0% | −\$487.00   |
| LONDON          | 129 | \$163.00    | \$1.26    | 49.6% | −\$835.00   |
| BUILDING        | 61  | −\$497.00   | −\$8.15   | 37.7% | −\$1,220.00 |
| VIOLENT_WHIPSAW | 12  | −\$517.50   | −\$43.12  | 41.7% | −\$772.00   |
| US              | 122 | −\$1,050.00 | −\$8.61   | 43.4% | −\$2,115.50 |
| CLEAN_TREND     | 73  | −\$1,605.00 | −\$21.99  | 42.5% | −\$2,530.50 |

CLEAN\_TREND — the regime the whole idea assumed — is the worst cut in the study at −\$21.99 a signal. Asia is the least-bad, which is a liquidity story rather than a mechanism, and it does not survive strip-3.

## 8 • Against the census — the big-moves-caught column

Movement 1 froze **68 big runs on MNQ** this week of which we **sat out 60**, carrying a hindsight ceiling of **\$8,024** on a single lot. **This section catches none of them, and it was never going to** — it is a gold gate and that ceiling is a Nasdaq number. The honest big-moves-caught line for Movement 3 this week is **0 of 60, because the four hunts aimed at those 60 runs did not run.**

Movement 2 is the only part of the report that engaged with them, and its verdict was an honest null: six mechanical fires for +\$15.50, −\$94.50 once the best is stripped.

| CAUSE CLUSTER | RUNS THIS WEEK | SHARE OF THE CEILING |
|---------------|----------------|----------------------|
| UNCLASS       | 36             | \$4,248 (pro-rata)   |
| OPEN/NEWS     | 22             | \$2,596 (pro-rata)   |
| VACUUM        | 8              | \$944 (pro-rata)     |
| FLOW-LED      | 2              | \$236 (pro-rata)     |

The cluster-by-cluster hunt these were frozen for is the work that did not happen. OPEN/NEWS at 22 runs is the largest addressable block and has never had a dedicated hunt on this week's tape.

## 9 · DISPOSITION — every lead touched in this movement

| LEAD                                                | VERDICT                         | THE NAMED TEST, OR WHAT WOULD REVIVE IT                                                                                                                                                                                                                                                                        |
|-----------------------------------------------------|---------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| MGC far-side depletion as a BREAK-CONTINUATION gate | REFUTED                         | Named test: 2,000-draw random-subset placebo on n=192. −\$2,059.50 at the <b>4.8th percentile</b> — 1,904 of 2,000 random picks did better, so the feature carries information and the sign is inverted. No reformulation saves it AS A BREAK SIGNAL; the reformulation is the next row.                       |
| Fade the gold break when the far side is empty      | PARKED — the reformulation      | The inverse of the refuted rule is +\$2,059.50 gross on 192 signals before it is a strategy. Revive by: building it as a FADE with its own entry timing and stop, and re-testing on a held-out window — an inverted backtest is a hypothesis, not a result.                                                    |
| MGC book imbalance (+0.10 to +0.25) at the break    | PARKED                          | The only cut that clears its placebo (96.2nd percentile, +\$1,795, strip-3 +\$590, LODO +\$970). Parked not shadowed because the band above it LOSES money — a non-monotone response is a bucket, not a mechanism. Revive if: it stays positive on 4+ untouched weeks AND the >0.25 band stops being negative. |
| The size escalation (top 25 / top 15)               | REFUTED as a route to this edge | Named test: re-running all 20 filters at each narrowing. Top-25 best cell is three prints (strip-3 −\$212.50); at top-15 every rule is negative. The footprint does not concentrate in the monsters, so narrowing further is not the answer here.                                                              |
| An exit fix for the MGC break gate                  | REFUTED                         | Named test: the excursion table. Median MFE +3.28 ATR against median MAE −3.04 ATR with 80.4% going a full ATR against us first. The distribution is symmetric — there is no stop that makes it work, unlike exhaustion_short in Part 1.5.                                                                     |
| MNQ VACUUM-cluster greenfield hunt                  | NOT RUN — measurement gap       | Phase never started (budget exhausted). 8 runs this week. First thing to reschedule.                                                                                                                                                                                                                           |
| MNQ FLOW-LED-cluster greenfield hunt                | NOT RUN — measurement gap       | Phase never started. Only 2 runs this week, so it is also the lowest-value of the four.                                                                                                                                                                                                                        |
| MNQ OPEN/NEWS-cluster greenfield hunt               | NOT RUN — measurement gap       | Phase never started. <b>22 runs — the largest addressable cluster of the week</b> and the highest-value hunt not done.                                                                                                                                                                                         |
| The chop-day scalp lab                              | NOT RUN — measurement gap       | Phase never started. Relevant: Movement 2 found this week's runs ignited out of unusually quiet tape, which is exactly the population this lab was for.                                                                                                                                                        |
|                                                     | SEE PART 2.5                    |                                                                                                                                                                                                                                                                                                                |

LEAD

VERDICT

THE NAMED TEST, OR WHAT  
WOULD REVIVE IT

MGC day rider / coil bouncer / London break  
fade

Not missing — relocated. Both SHADOW gold  
candidates of the week are dispositioned in Part  
2.5 Part A, which completed.

★ THE ONE STONE STILL UNTURNED

Both of this movement's honest verdicts point at the same missing piece and it is not another feature. **It is the OPEN/NEWS cluster on MNQ — 22 of this week's 68 runs, roughly a third of an \$8,024 ceiling, and no gate on this desk has ever been built for it.** That hunt was scheduled tonight, was fifth in the queue, and the queue ran out. It costs one phase to find out and it should be first next week, ahead of any further gold work.