

GAZBOT V7 — the Friday report

Week ending 2026-08-22 · Mon 17 – Fri 21 August 2026 · MNQ-led, MGC captured · PAPER · 30 plays · 10 sections

★★ READ THIS BEFORE ANYTHING ELSE — THE DESK IS NOT FLAT

The day rider is long 4 MNQU6 from 29,491.62, opened 2026-08-21T13:13:06Z, and the position is still open. Marked at the last tick the tape printed before the CME shut (29,370.00) that is **-\$973.00**. It is larger than everything else in this report put together.

Every automatic path that exists to close it failed on the same call — the IB gateway entered a partial wedge at about 19:57Z and the rider's own hard-flat, the watchdog, the EOD flatten and the reconciler all share one connectAsync. The EOD flatten's timer is **Mon-Fri**, so it does not retry over a weekend; its next elapse is Monday 20:53Z, and the tape reopens **Sunday 22:00Z**. That is a live-market window of about 23 hours with an unmanaged four-lot position in it.

What is already done: gazbot7-weekend-flatten.timer is installed and enabled, firing at the Sunday 22:00Z reopen and then every minute for two hours, self-disabling once flat. **What still needs you:** confirm the gateway is answering before Sunday 22:00Z — nothing on this box can heal a wedge, and if it is still wedged at the reopen the retry loop will fail sixty times in a row exactly as designed and the position will still be there. That is SATURDAY #1.



THE HONEST HEADLINE — a small red week, and the loss is not where you would look for it

The desk booked **-\$247.50** across 43 clean legs (49 contracts — the day rider's legs are 1-3 lots each, so a row count is not a lot count), **3 of the five days green**. Three sessions made \$1,120.50 between them and two lost -\$1,368.00, and essentially all of that loss is **one day-rider position on Thursday** that ran for six hours with no stop in the water. The tournament — the eight base gates filling six live slots that most of this report is about — traded only 36 legs all week for -\$58.50, because four of its six slots were stood down and never fired at all.

The three things worth carrying into Saturday. One: the standdown was right. Firing all six gates mechanically through this week's own tape loses **-\$1,905.36** on 117 fires (Movement 2), so the four silent slots are the best-performing thing on the desk. Two: your hands made \$550.00 on 16 claims and the machine's own profit exits made almost nothing — but the study that tried to automate your hand comes back **-\$4 out-of-sample**, so the answer is not to copy it. Three: the week's real damage was done by things that are not strategies at all — a session-roll bug, an off-tape fill, a missing stop and a wedged gateway. ★ **REV2 sharpens that last one:** the four lots still open are a missing-stop incident that a wedge then made unfixable, not a wedge incident. Every stop width tested — all seven, from 0.5× to 3×ATR — closes that position within nineteen minutes of entry, six and a half hours before the gateway stopped answering.

There is no single honest total, so here are all of them, with what separates each from the next.

WHAT YOU ARE LOOKING AT	LEGS	NET \$	WHY IT DIFFERS FROM THE ROW ABOVE
Raw trades rows, 08-17 → 08-21	44	-\$330.50	everything the desk booked, flags and all — the CASH book
Clean ledger (data_quality IS NULL)	43	-\$247.50	drops 1 flagged lot worth - \$83.00 — this is the STRATEGY book and it is the number to judge a gate on

WHAT YOU ARE LOOKING AT	LEGS	NET \$	WHY IT DIFFERS FROM THE ROW ABOVE
The day rider alone	7	-\$189.00	the second desk, on the same IB account. It made the week's best trades and its single worst one
The tournament alone (eight base gates, six live slots)	36	-\$58.50	what the gates this report is mostly about actually did — on two slots, because the other four were stood down
... plus the position still open (4 lots long from 29,491.62)	4	-\$973.00	not in any row above, because the trades table only enumerates what CLOSED. The week's real economic position is -\$1,303.50.

THE ONE FLAGGED LOT, NAMED

It is a single day-rider lot from Friday morning, flagged BADFILL:stale_price_30pt_outside_book_20260821, worth -\$83.00. It filled at 29388.25, about 26 points below a 30-minute tape low, roughly 92 ms after its 3-lot sibling filled on-tape. **A BADFILL flag marks a defective fill PRICE, not a phantom trade** — the money left the account, and IB's own realizedPNL of -\$149.38 for Friday reconciles to the FLAGGED figure, not the clean one. So: clean book for judging a gate, cash book for counting cash, and the venue position line above both.

DAY	DESK NET (CLEAN)
2026-08-17	\$458.00
2026-08-18	\$162.50
2026-08-19	\$500.00
2026-08-20	-\$1,294.50
2026-08-21	-\$73.50

★ THE SCORECARD — what survived the week

Lead with what survived the skeptic, never with the biggest number. Seven results carried real weight into Saturday; here is where each one landed and the test that decided it.

THE CLAIM

VERDICT, AND THE TEST THAT DECIDED IT

WHERE

The desk is better off with four of its six slots switched off

SURVIVED — the week's best-evidenced result All six gates fired mechanically in their exact live config across the whole of this week's captured tape, every fire counted rather than only the ones near a big move: –\$1,905.36 on 117 fires, and –\$517.79 of that is the four slots that were stood down. Fired into the 65 big runs the desk sat out, they lose –\$1,333.85; strip every filter off and they lose more, not less. The base-rate control is what makes this different from the last three idle-gate labs, which measured only windows already known to contain a run.

The honest limit: one week, mostly stood down. It licenses leaving them benched, not retiring them — and ★ REV2 puts a bootstrap round the headline: the whole-week total's 95% interval straddles zero, so –\$1,905.36 is a point estimate, not a measured saving. What IS statistically solid is the comparison Rev1 got backwards: the ten minutes before a big run are 3.6× worse per fire than the base rate (12.9× once the overlapping fires are removed), 95% CI excluding zero on both readings. Rev1 called those two rates “close”; they are not, and the difference strengthens the bench case rather than blurring it.

[Movement 2](#)

“abs_veto_55s is the promotion candidate, +\$332 this week”

WITHDRAWN — \$308.50 of it is untradeable The permanent Asia bench (00:00-07:00Z) sits DOWNSTREAM of where a shadow arm records its fire, and the desk has filled nothing in that window since 5 August. Void it and the family goes from +\$332 to +\$23.50 on 34 fires, 2 of 5 days green, and stripping its single best fire turns it –\$133.50. Its 50-second sibling goes negative. Across all 38 families, –\$3,631 of the week's board was booked in that window. The family's real case is its all-time record (+\$3,078 on 263 tradable-hour fires), which is a reason to keep incubating and not a reason to promote on Monday.

[Part 2](#)

“Only run grind on a clean trend day” — the week's headline mid-week lead

PARKED — it works and it cannot be run On 128 live grind fires the rule turns –\$958 into +\$580.50 and keeps 25 of 43 winners, so it is not a fake win — but strip its best three fires and it is –\$202. Then the killer: of 240 classified sessions, 97 are clean-trend and exactly ONE is both clean-trend and above grind's 22-point ATR floor. The two conditions are very nearly mutually exclusive. Arming was never the binding constraint. The real question is the floor, which sits at the ninetyeth percentile of one-minute ATR and has never been swept.

[Part 2.5](#)

“abs_veto_short is chasing exhausted moves”

REFUTED, and the sweep gives a real verdict Measured inside each trade's own life: 0 of its 15 losing lots were chases. Every one went at least 6 points onside, median +11.25. The difference between a winner and a loser is how far it ran, not whether it ran. Re-raced at five stop widths twice over

[Part 1.5](#)

THE CLAIM

VERDICT, AND THE TEST THAT DECIDED IT

WHERE

— once scaling the whole trade, once moving only the stop — all ten cells are negative, so the sign survives and this IS a verdict rather than a knife-edge. The gate reads −\$203 live only because your hand took +\$314.74 out of it that the machine would not have taken (★ REV2 — recomputed off the exact venue stop orders rather than an estimated ATR, and with the quiet-tape clip modelled; the sign and the verdict are unchanged, the figure was \$412.93).

★ REV2 also finally ran the EXIT ladder §2.1 diagnosed and §2.2 never tested. It does not fail: four rungs beat the as-traded book while keeping 7 of 7 winners — a scalp at 1.0R (+\$187.00), at 0.5R (+\$178.50), a trail 8pt behind the peak (+\$149.00), and a scalp at the losers' own median peak of +11.25pt (+\$141.00) — each positive on strip-best-1 and on the worst leave-one-day-out. But against a 400-draw random-entry control on the same tape they sit only at the 80th-93rd percentile, so it is a SHADOW row, not a Monday change.

“Thursday was a bad trade”

REFUTED — it was a missing stop One rider long, four lots, six hours, no protective exit in the water at any point. Raced on the real tape from its own entry, all seven stop widths from 0.5× to 3×ATR beat what happened, the tightest by more than a thousand dollars. The entry is the rider's ordinary open-hour entry and it made money on three other days this week.

★ REV2 CHANGED THE WIDTH, on this report's own winners-retention rule. Rev1 priced the stop on Thursday alone — the one losing day — and picked 1.5×ATR. Raced across all six rider entries, 1.5×ATR (41.4pt on Monday, 46.4pt on Tuesday) stops out BOTH winners and takes the week's closed rider book from −\$272.00 to −\$388.90: worse than doing nothing. 3.0×ATR is the only width that is positive (+\$195.70) and keeps all three winners, and it still caps Thursday at −\$826.80. The honest week-level figure is +\$468, not \$1,000.

And the finding neither revision expected: every width tested closes Friday's 13:13Z entry within nineteen minutes — six and a half hours before the gateway wedged. The four lots at the top of this report would not exist under any stop.

[Part 1.5](#), [Part 1.6](#)

“A standdown means no gate can be armed”

FALSE — and it happened on Friday The standdown is a variable inside one process. open_hour_watch.py is a second writer to gate_switches.env with off→on authority and no awareness of it, and at 14:03Z on Friday it armed abs_veto_short straight through the standdown. The router caught it on its next tick and re-benched it 2m15s later, so the cost was \$0 by luck, not by design — the exposure is bounded by a 5-minute poll, not by a guard. Also on the

[Part 2.6](#)

THE CLAIM	VERDICT, AND THE TEST THAT DECIDED IT	WHERE
	stand: five expired carve-out comments are still spliced into every router prompt labelled ACTIVE, the oldest thirteen days old and asserting that this very gate is "ARMED BY DEFAULT".	
Four instruments reported confidently and wrongly this week	CONFIRMED — none of them errored core_health.flat said the desk was flat through 137 of the last 141 router ticks while the venue was long four lots. hour_watch read Friday at −\$150.50 against the clean book's −\$69.00 (it has no data_quality filter — and it is the one that reconciles to IB). router_nightly and selector_nightly, the two files the weekly router review is built from, silently stopped being written on 08-14 and 08-18; neither has a timer. And 7 of the 38 shadow arms are byte-identical to another arm, so three A/B experiments have been reporting perfect nulls from a broken arm rather than from evidence. A number that arrives without a failure mode is not a measurement.	Part 1 , Part 2 , Part 2.6

§ Every gate, this week, on the clean ledger

GATE	LEGS	NET \$	WIN %	\$ PER LEG
abs_veto_short	22	−\$203.00	31.8%	−\$9.23
day_rider	7	−\$189.00	42.9%	−\$27.00
exhaustion_short	14	\$144.50	57.1%	\$10.32

Computed at build time from data/gazbot7.db with data_quality IS NULL and net of the \$1.50 round-trip fee, not transcribed. Gates are collapsed across their A and B lots; the per-lot split is in Part 1 §3. **Four of the six live slots are absent from this table because they fired nothing** — Movement 2 prices what that was worth.

THE THREE THINGS THIS WEEK ACTUALLY TAUGHT THE DESK

1 • The trades table cannot tell you what you are carrying. It enumerates what CLOSED, so it can never contradict a benign story. This week the router spent 137 consecutive ticks reciting closed rows and the word “flat” while updatePortfolio in the same journal read position=4.0. Attribution has to terminate at the venue position line, and “not mine to close” obliges escalation, not dismissal.

2 · A filter that cannot fire is not a filter, it is an off-switch. grind's 22-point ATR floor blocked 24,110 of 24,121 raw signals this week and sits at the p90 of one-minute ATR. Every grind result in this report — the unrunnable trend-day rule, the two empty exit-shadow watches, the zero live fires since 5 August — is downstream of that one number, and nobody has ever swept it.

3 · Where you measure decides what you find. Three separate results this week flipped on the measurement window rather than on the data: the shadow board reorders completely once the untradeable Asia hours are voided; the idle-gate lab only says something because a whole-week base rate was fired alongside the run windows; and the claim detector makes + \$1,069 in-sample and −\$4 out of it. **Ask what the denominator is, every time.**

★ WHAT TO ACTUALLY DO

Garrath — this is the whole answer to “so what do I actually DO?”, and it is at the front. Every row below mirrors `reports/friday_v7/plays.json`, which is the reconciled source of truth for this week; the same file drives the [Monday playbook](#), so the two cannot drift apart. ★ **This week's card is deliberately SHORT.** The tournament traded two of its six slots for 22 lots, which is not enough evidence to re-tune a threshold, and four separate sections say so independently. So most of what is here is either a fault to fix or a measurement to build — and the single most urgent row is not a strategy at all, it is a position that is still open. 30 plays, five windows, in the order they need doing. Every play carries its **evidence tier** and its **kill criterion** — if a play has no way to be proven wrong it should not be on the card.

THE THREE THAT MATTER MOST

1 · The four lots that are still open — confirm the gateway before Sunday 22:00Z (SATURDAY #1). Before the Sunday 22:00 UTC reopen, check that the IB gateway is answering on 127.0.0.1:4002 (restart it if it is not — nothing on this box can heal a wedge). `gazbot7-weekend-flatten.timer` is already installed and enabled and will fire at 22:00Z then retry every minute for two hours, self-disabling once flat. Your job is the precondition, not the flatten. 4 lots long from 29,491.625, marked $-\$973.00$ at the last print before the halt (29,370.00). Larger than the whole week's trading.

2 · Put a protective stop on the day rider's own book — at $3.0\times\text{ATR}$, not $1.5\times$ (SATURDAY #2). Give the rider a protective stop sized off its own entry ATR at $3.0\times$, armed at entry and never removed. $\triangle$ REV2 CHANGED THE WIDTH. Rev1 priced $1.5\times\text{ATR}$ on Thursday alone — the one losing day — and quoted $\approx \$1,000$. Raced across all six of the week's rider entries, $1.5\times\text{ATR}$ stops Monday's $+\$453$ winner (in-life MAE -75.25pt against a 41.4pt stop) and Tuesday's $+\$185$ winner (MAE -64.50pt against 46.4pt), and makes the closed book WORSE than doing nothing. Closed rider book: $-\$272.00$ as traded, $+\$195.70$ at $3.0\times\text{ATR}$, $-\$388.90$ at $1.5\times$. Including the still-open carry marked at the last print: $-\$1,251.00 / -\$547.10 / -\$763.30$. The week-level value of the change is $+\$468$ to $+\$704$, not $\$1,000$.

3 · Branch the rider on `venue_net`, never on `entered` (MONDAY #1). One line in the day rider: `gate the flatten path on venue_net != 0` instead of `on entered`. $-\$73.50$ clean plus a $-\$83.00$ BADFILL lot = $-\$156.50$, in sixty seconds, on 2026-08-21.

Generated from the top of the card below, so the two cannot drift apart. Everything else, including every refuted idea and why it died, is in the five windows underneath.

SATURDAY · BEFORE THE SUNDAY 22:00 UTC REOPEN — ONE OF THESE IS A REAL DEADLINE · 4 PLAYS

The first row is not a strategy change, it is an open four-lot position that no automatic path on this box was able to close on Friday. Everything under it is code and config that needs a restart. **Note what is NOT forced any more:** `gazbot7-gate-reactivate.timer` no longer arms anything at 22:00Z — its hold set is the whole roster and its last two runs both logged nothing to

change — so the gate decisions on this card can genuinely wait if you want them to. The position cannot.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
1	The four lots that are still open — confirm the gateway before Sunday 22:00Z LIVE Before the Sunday 22:00 UTC reopen, check that the IB gateway is answering on 127.0.0.1:4002 (restart it if it is not — nothing on this box can heal a wedge). gazbot7-weekend-flatten.timer is already installed and enabled and will fire at 22:00Z then retry every minute for two hours, self-disabling once flat. Your job is the precondition, not the flatten. Mechanism: The day rider opened 4 MNQU6 long at 29,491.625 (the fill-weighted price: 2 @ 29,477.00 + 1 @ 29,506.25 + 1 @ 29,506.25; day_rider_state.json's entry field says 29,491.93 and is wrong by 0.305pt) on Friday at 13:13:06Z. At about 19:57Z the gateway entered the documented fourth-class partial wedge — it accepts nothing new while serving everything old — and all four independent closing paths share one ib.connectAsync: the rider's 20:40 hard flat, the flatten-only watchdog, the EOD flatten and desk_reconcile. All four failed. The EOD flatten's timer is Mon-Fri with Persistent=false, so a Friday failure gets no catch-up at all; its next elapse is Monday 20:53Z and the tape reopens Sunday 22:00Z. ★ REV2, from the unit itself: gazbot7-weekend-flatten has NEVER run — empty ExecMainStartTimestamp, empty ActiveEnterTimestamp, no journal entries; next elapse Sun 2026-08-23 22:00Z. Its 120 OnCalendar firings cover 22:00–23:59Z ONLY, and Persistent=false means a box that is down across that window gets no catch-up. Both gaps end at the same place: the position stays open until the EOD flatten's next elapse, Monday 20:53Z. Why: This is the only row on the card with a hard deadline and it is the largest single number in the report. Everything else can wait a week; a live-market window of about 23 hours with an unmanaged four-lot position in it cannot. Evidence: Front page and Part 1 §0; the mechanism is Part 1.6 §5 and the nightly review of 2026-08-21 22:43Z, which read the position live off updatePortfolio. · verification: After the reopen: systemctl is-failed gazbot7-weekend-flatten and day_rider_state.json showing venue_net: 0. Check the state file, not core_health.fl at — that flag said the desk was flat throughout. · mode: operate	4 lots long from 29,491.625, marked –\$973.00 at the last print before the halt (29,370.00). Larger than the whole week's trading.	Arms: Once, at the Sunday 22:00Z reopen. Exit: Flat. The unit disables itself. Kill: n/a — this is a position, not an experiment.	Garrath — gateway health check. The flatten itself is automated. n/a.
2	Put a protective stop on the day rider's own book — at 3.0×ATR, not 1.5× LIVE Give the rider a protective stop sized off its own entry ATR at 3.0×, armed at entry and never removed. △ REV2 CHANGED THE WIDTH. Rev1 priced 1.5×ATR on Thursday alone — the one losing day — and quoted ≈\$1,000. Raced across all six of the week's rider entries, 1.5×ATR stops Monday's +\$453 winner (in-life MAE –75.25pt against a 41.4pt stop) and Tuesday's +\$185 winner (MAE –64.50pt against 46.4pt), and makes the closed book WORSE than doing nothing. Mechanism: One stop for the whole position, first tick at or through it closes every leg still open, raced on capture.db MNQ ticks from each entry using the ATR the detector itself read at the confirming tick (27.6 / 30.9 / 35.3 / 34.2 / 30.7). Closed book, net of \$1.50 per round trip: as-traded –\$272.00; 0.5× –\$495.90; 0.75× –\$658.10; 1.0× –\$135.10; 1.25× –\$262.00; 1.5× –\$388.90; 2.0× –\$642.70; 3.0× +\$195.70. Winners kept, of three: 0 / 0 / 1 / 1 / 1 / 1 / 3. 3.0×ATR is the only width that is both positive and keeps every winner — it is 82.8 / 92.7 / 105.9pt on the three winning days, all wider than their worst excursions, and 102.6pt on Thursday, which caps that day at –\$826.80 instead of –\$1,294.50. Why: The rider ran naked for 6.7 hours on Thursday and is still naked on the four open lots. And the width matters as much as the stop: a 1.5×ATR stop would have taken –\$116.90 out of this week's closed rider book, which is the winners-retention failure this report refuses to accept anywhere else. Every width, including 1.5×, closes Friday's 13:13Z entry within 19 minutes — 3.0×ATR at 13:31:55Z for –\$742.80 — so any stop at all would have removed the weekend carry before the gateway wedged. Evidence: Part 1.6 §3b — the seven-width × six-entry grid with the winners-kept row, which is the verdict-bearing table. Part 1.5 §3 is the Thursday-only race, rebuilt in REV2 on the detector's own ATR of 34.2 (Rev1 used an estimated 36.1); and Part 1.6 §5. · verification: Re-race all rider entries at all seven widths every Friday and	Closed rider book: –\$272.00 as traded, +\$195.70 at 3.0×ATR, –\$388.90 at 1.5×. Including the still-open carry marked at the last print: –\$1,251.00 / –\$547.10 / –\$763.30. The week-level value of the change is +\$468 to +\$704, not \$1,000.	Arms: Every rider entry. Exit: The stop is the exit being added. Kill: Two consecutive weeks in which the 3.0×ATR stop cuts more winner dollars than it saves on losers, measured on the widths×days grid in Part 1.6 §3b, not on one day.	Garrath — day_rider config + restart Remove the stop key. It is additive; nothing else changes.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	publish the grid with the winners-kept column. A width picked from the losing day alone is not a width. · mode: deploy			
3	Make the tournament standdown a FILE every writer checks LIVE Move TOURNAMENT_STOOD_DOWN out of router_tick_durable.py and into a small state file that every writer to gate_switches.env reads before it writes. Six processes can write that file; two respect the lock. Mechanism: On Friday at 14:03Z, with the tournament stood down and all six gates reading off, open_hour_watch.py armed abs_veto_short. It has off-on authority and zero standdown awareness, and it did exactly what it was built to do. The router caught it on its very next tick and re-benched it at 14:05:18Z. Why: The gate was live for 2 minutes 15 seconds and fired nothing, so this cost \$0 by luck rather than by design. The exposure is bounded by the router's five-minute poll, not by any guard — arm three minutes later and the gate is live through the whole US open window. Evidence: Part 2.6 §4 — the three log lines, unedited. · verification: Set the standdown, then run open_hour_watch by hand against a confirmed break and assert it declines to write. · mode: deploy	One incident this week, \$0 by luck. Six writers to one file, two of which respect the lock, and all six share a single .tmp path.	Arms: Always. Exit: n/a. Kill: n/a — it is a guard, not a strategy. If it ever blocks a legitimate arm, that is the signal that the standdown itself was wrong, not the guard.	Garrath — code edit across the switch writers Delete the file check; the variable stays where it is.
4	Delete the five expired carve-out comments from gate_switches.env LIVE Delete the expired instruction blocks from the top of gate_switches.env. switch_notes() splices the top 45 comment lines into every router prompt labelled ACTIVE and nothing prunes them. Mechanism: Five instructions are currently in that header: 2026-08-21 (fresh), 2026-08-19 ("PINNED until 21:55Z", that deadline passed on Wednesday), 2026-08-13, 2026-08-11 ("expires 16:00Z", eleven days ago) and 2026-08-09 — which states that abs_veto_short IS ARMED BY DEFAULT, directly contradicting the standdown the router is enforcing. Why: Only 172 of 1,425 ticks restated any of them this week and all 172 were the two recent ones, which looks like the problem has gone away. It has not, and no mention-rate streak can show that it has: 1,424 of the 1,425 reason fields are pinned at exactly the 1,200-character cap, so the mention rate is measuring spare room, not belief. Any quieter week and the thirteen-day-old "armed by default" line is one tick from being read as policy. Evidence: Part 2.6 §5 — the header table with each line's age and its own mention count. · verification: Re-read the header after the edit and confirm only instructions with a live expiry remain. Revoking a carve-out in the ledger does not revoke it in the file. · mode: operate	5 carve-outs spliced as ACTIVE; ages 0d, 2d, 8d, 11d and 13d. Obeying the 08-09 one would arm a gate the operator has stood down.	Arms: n/a. Exit: n/a. Kill: n/a — housekeeping.	Garrath — text edit, no restart Git restore the file's header.

MONDAY · AT THE DESK — SWITCH-FILE AND EXIT-CELL CHANGES, ALL REVERSIBLE · 4 PLAYS

No code deploy in this window. Every row is gate_switches.env, exit_overrides.json or a router constant, and comes back in five minutes. This week's list is short on purpose: the tournament traded two slots and 22 lots, which is not enough evidence to re-tune anything, and the report says so in three separate places.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
1	Branch the rider on venue_net, never on entered LIVE One line in the day rider: gate the flatten path on venue_net != 0 instead of on entered. Mechanism: At 00:00Z the session roll resets day_rider_state.json to entered: false with note: "switch off" while venue_net still reads the true position. On Friday morning the rider read its own reset state, concluded it was flat, and flattened a position it had opened sixty seconds earlier. Why: It stops being silent and starts asserting FLAT, which is worse. The 3-lot leg went out on-tape; the 1-lot remainder filled 26 points below a 30-minute low about 92ms later and had to be flagged BADFILL. Evidence: Part 1.5 §1 — both lots named, with the timestamps one timer tick apart. · verification: Watch the next 00:00Z roll with a	−\$73.50 clean plus a −\$83.00 BADFILL lot = −\$156.50, in sixty seconds, on 2026-08-21.	Arms: Always. Exit: n/a. Kill: n/a — it is a defect fix.	Garrath — one line in day_rider.py Restore the entered branch.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	position open and confirm no <code>CLOCK_FLAT_BUG</code> row appears. · mode: deploy			
2	Leave the four silent slots benched LIVE No change. Keep <code>grind_long</code> , <code>capitulation_long</code> , <code>abs_veto_long</code> and <code>rgv_short</code> off, and note that <code>reactivate_gates</code> no longer re-arms them at 22:00Z, so this needs no action to hold. Mechanism: All six gates were fired mechanically in their exact live config across the whole of this week's captured tape — every fire counted, not only the ones near a big move. Fired into the 65 sat-out runs they lose —\$1,333.85; stripped of every filter they lose more, not less; and joining a run after ignition is roughly free where predicting one is not. Why: The base-rate control is what makes this different from the last three idle-gate labs, which measured only windows already known to contain a run and reported survivorship as skill. Evidence: Movement 2 \$2–\$4, with the base-rate control in \$3. · verification: Re-run the base rate on the first genuine trend week. This is one quiet week and the finding is scoped to it. · mode: hold	All six mechanically: –\$1,905.36 on 117 fires across the week's tape. The four silent slots alone: –\$517.79 on 48 fires. Aligned fires — the ones pointing the same way the run then went — still lost \$602.60 on 11.	Arms: Not until the base rate turns on a trend week. Exit: n/a. Kill: One trend week where the mechanical base rate for these four is positive.	Nobody — this is the current state n/a.
3	Bench the automated half of <code>abs_veto_short</code>, keep the signal shadowed SHADOW Leave the live slot off (it already is, since Friday 14:05Z). Keep the signal running in shadow through the 55-second confirm and the ladder arms, which is where its case now lives. Mechanism: Re-raced at five stop widths, twice over — once scaling the whole trade, once moving only the stop with the target pinned — all ten cells are negative . The sign survives, so this is a verdict rather than a knife-edge. The live book reads –\$203 only because manual claims took +\$314.74 out of it that the machine would not have taken; left to itself the machine reads –\$446.24. ★ REV2 — those are the VENUE-STOP figures and they are the ones to quote: R comes from each lot's own <code>StopLimitOrder.auxPrice</code> in the tournament journal and the <code>atr_split=22</code> quiet-tape clip is modelled, so 18 of 18 machine-closed lots reproduce their live exit reason. \$2.2's –\$548.73 / +\$412.93 used an estimated ATR and no clip. Same sign, same verdict. Why: The root cause is not what it looks like. Zero of its 15 losing lots were chases — every one went at least 6 points onside, median +11.25 — so the entries work briefly and then stop working. That is exactly what a continuation-confirm addresses, which is why the case moves to shadow rather than to a bench-and-forget. Evidence: Part 1.5 \$2.1–\$2.4 — \$2.4 is the REV2 exit ladder on the venue stops — and Part 2 \$4b for the short-side case the signal now rests on. · verification: The 55s mirror needs a positive TRADABLE-HOURS week with a positive leave-one-day-out before this comes back as a promotion question. · mode: hold	Live –\$203.00 on 22 lots. Machine-only on the exact venue stops –\$446.24 (–\$548.73 on \$2.2's estimated ATR). Ten sweep cells, all negative, worst leave-one-day-out negative at every width.	Arms: Not this week. Exit: n/a. Kill: Two consecutive tradable-hours-positive shadow weeks would re-open it.	Nobody — the slot is already off <code>abs_veto_short=0</code> n. Reversible in five minutes.
4	Do not re-tune any exit cell or router threshold on this week LIVE No change to <code>exit_overrides.json</code> and no change to the router's thresholds or timing constants. Mechanism: The tournament traded two slots and 22 lots. The A/B leg comparison has 16 paired signals and the sign flips on one pair (B–A –\$76.50, –\$96.00 with the best pair stripped). The router wrote 20 switches all week, 5 of them worth \$0 by construction because they were in the Asia block. Why: Four separate sections reach the same conclusion independently. Tuning on a week the desk was stood down is fitting to the standdown. Evidence: Part 1 \$3 (the A/B legs), Part 2.6 \$8 (the router verdict). · verification: n/a — the action is inaction. · mode: hold	16 paired A/B signals, B–A –\$76.50, strip-best-1 –\$96.00. 20 router switches, 5 of them \$0 by construction.	Arms: n/a. Exit: n/a. Kill: A week with ≥40 paired signals would re-open the exit-ladder question.	Nobody — no change n/a.

BUILD · NEXT — NOTHING GOES LIVE THIS WEEK · 9 PLAYS

Instrumentation, shadow slates and the method fixes that decide whether next Friday's numbers can be trusted at all. **Four of these are measurement faults found this week**, not ideas: byte-identical shadow arms reporting perfect nulls, two nightly rollups that silently stopped, a

health flag that says flat while the desk is long, and gold results confounded with a large short drift nobody is reporting.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
1	Row-level duplicate check on the shadow board, and re-seed three broken arms LIVE Add a duplicate check at shadow-write time that compares an arm's rows against every other arm's, and re-seed the three A/B pairs that are currently writing one configuration's trades under two names. Mechanism: Seven of the 38 families with $n \geq 8$ this week are byte-identical to another family on every field including qty, target_r, stop multiple and real_pnl: capit_flip_live == capit_flip_t90, cx_clip_brk_live == cx_clip_brk_standdown == cx_grindA_clip, and odr_c5_s30 == odr_c5_s30_g. Why: Every one of those pairs exists to measure a difference and every one reports a perfect null — not because the configurations agree but because only one is running. A broken arm and a genuinely neutral change look identical in the board's summary view, and the summary view is what a report reads. Note especially the clip-vs-standdown pair, which exists specifically to test whether standing an exit down changes anything. Evidence: Part 2 §5 — the four groups, with the hash comparison that found them. · verification: The check should fail the three known pairs today. After re-seeding, they should diverge within a week. · mode: build	7 of 38 arms duplicated; 3 A/B experiments invalidated. Until this lands, every shadow A/B delta on the board is unverified.	Arms: n/a. Exit: n/a. Kill: n/a — it is an instrument fix.	Claude — shadow writer + a one-off re-seed Remove the check. The re-seed is not revertible and does not need to be.
2	grind's 22-point ATR floor — swept tonight on the 128 live fires; the ladder is now a confirmation, not a discovery SHADOW Keep the shadow arms at 14 / 16 / 18 / 22, but stop calling the floor untested. ★ REV2: the 128 live grind fires already on disk carry their own entry ATR and their own real exits, so the floor is a filter on a real book and was swept tonight. Nothing goes live. Mechanism: Every live grind fire, NET of the \$1.50 round trip (the — \$958.00 in circulation is gross; the same fires are —\$1,150.00 net), subset by entry ATR: no floor —\$1,150.00 on 128 · 14pt —\$1,143.50 on 97 · 16pt —\$799.50 on 90 · 18pt —\$554.50 on 82 · 20pt —\$514.50 on 70 · 22pt —\$280.50 on 51 · 24pt +\$21.00 on 45 · 26pt +\$347.50 on 33 · 28pt +\$469.00 on 28. Monotone upward from 14, which is the shape of a real effect — but strip the best three fires and EVERY rung is negative (—\$309.00 even at 28pt), and the worst leave-one-day-out is negative at every rung (—\$633.50 at 28pt). Against 5,000 random same-size subsets of the same 128, the 22pt rung sits at the 65th percentile and the 24pt rung at the 83rd. Why: The direction of the floor is right and 22 is not obviously the wrong number — it is simply not a high enough one to make the gate positive, and no rung between 10 and 28 survives strip-best-3 or leave-one-day-out. That reframes the question: the floor is not what is wrong with grind. ★ And the 2.3x disagreement about reachability is SETTLED — see the number field. Evidence: Part 2.5 §1–§2, Part 1.5 §4, Movement 2 §3. · verification: Report each rung's fires, real_pnl and winner retention at $n \geq 40$. Do not promote a rung on fire count alone. · mode: shadow	Live floor 22pt: —\$280.50 on 51 fires, keeps 18 of 43 winners, strip-best-3 —\$1,058.50, worst leave-one-day-out —\$1,174.50. ★ REACHABILITY, RESOLVED: 11.0% and 25.3% are the same measurement on different DENOMINATORS. All 24 hours, 33 sessions of capture.db: 11.0% of minutes clear 22pt, median 11.68, p90 22.62 — reproduced to the second decimal. 13:00–21:00Z only, 239 days of the lake: 25.1%, median 14.64, p90 31.34 — the p25 lab's 25.3% / 14.66 / 31.34 to within rounding. The lake and capture.db AGREE at ~10–11% on the same all-hours basis. So '22pt is the p90' is true of the all-hours denominator and FALSE of the session one, where it is the p75.	Arms: Shadow only. Exit: Same exits as the live grind slot, so the floor is the only variable. Kill: Already fired for the live-fire population: no rung produced a strip-best-3-positive or leave-one-day-out-positive result, so the floor is not the problem and grind is back to being an entry question. The shadow arms now only test whether the same shape holds forward; they cannot revive the floor as a lever on their own.	Claude — shadow slate only Remove the arms.
3	Point core_health.flat at the venue position, or rename it LIVE Make core_health.json's flat field reflect the venue position across both books, or rename it tournament_flat so nothing downstream can read it as a desk-wide safety signal. Mechanism: It is tournament-scoped today. The router reads it and 137 of the 141 ticks since the rider's Friday entry contain the word “flat” while the venue was long four lots. The last of those ticks describes the rider as “clientId 4, a desk I neither route nor bench”. Why: The trades table only enumerates what CLOSED, so it can never contradict a benign story. The venue position line can. Citing closed rows and declining responsibility for an open one is an attribution that terminates in the wrong place — “not mine to close” obliges escalation, not dismissal. Evidence: Part 2.6 §6, Part 1 §7. · verification: Open a position on either book and assert the flag goes false. · mode: build	137 of the last 141 router ticks said flat against a 4-lot venue position worth about \$235k of notional.	Arms: n/a. Exit: n/a. Kill: n/a — instrument fix.	Claude — core_health writer Restore the tournament-scoped field.
4	Put timers on router_nightly and selector_nightly, and fix selector's two defects		Arms: n/a. Exit: n/a.	

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	LIVE Add systemd timers for both nightly rollups, and before believing selector_nightly again, give it a data_quality filter and make it reprice at each gate's actual stop multiple instead of a hard-coded 1.0. Mechanism: The newest router rollup is 2026-08-14 and the newest selector rollup 2026-08-18. Neither has a timer, so both simply stopped. The whole weekly router review had to be rebuilt from raw logs this week. Why: A monitoring file that stops being written is worse than one that errors, because its absence looks like a quiet week. And the selector rollup is wrong even when it runs: quarantined rows inflate its regret, and two of the live gates do not use a 1.0 stop multiple. Evidence: Part 2.6 §7, and Part 1 §7 for the instrument list. · verification: Both directories should gain a dated file every night for a week. · mode: build	Router rollup 8 days stale, selector rollup 4 days stale, at this report.	Kill: n/a.	Claude — two timer units plus the selector fixes Disable the timers.
5	LIVE Add a long/short split and a directional-control column to every MGC study, retrospectively for the ones in this report. Mechanism: On the same out-of-sample window the gold rider is scored against, holding gold LONG every session loses \$21,738 and holding it SHORT every session makes \$20,382 over 226 sessions. Any gold arm with even a slight short bias will read as profitable without carrying any edge. Why: Two arms in this report were already reading that drift as edge before the control was run: the derived-trigger gold rider (+\$655 out-of-sample) and the stay-out meter's direction call, whose UP/DOWN hit-rate gap turns out to be the same drift. Evidence: Part 2.5 §4d and §5. · verification: Every gold table in next week's report carries the split. Any arm whose edge vanishes against the control is re-dispositioned. · mode: build	\$21,738 long / +\$20,382 short over 226 out-of-sample sessions. MGC is \$10.00 a point.	Arms: n/a. Exit: n/a. Kill: n/a — method fix.	Claude — every MGC harness n/a — it is an added column.
6	SHADOW Seed a shadow arm that fires the existing six gates only AFTER a run has ignited, and let it accumulate n against the pre-ignition arm. Mechanism: Movement 2's four arms: fired in the 10 minutes before ignition with live filters on, the gates lose −\$1,333.85 on 23 fires; fired in the 5 minutes AFTER ignition with the same filters, they lose −\$45.00 on 17. That is the only asymmetry in the whole lab and it points the same way as the greenfield lab's timing-decay ladder. Why: Predicting a run costs money; joining one is roughly free. Free is not profitable, so this is an incubation, not a proposal — but it is the one place in the idle-gate work where the sign is not clearly against us. Evidence: Movement 2 §2 — the four-arm table. · verification: n≥60 fires before it is discussed again, and it must beat the post-ignition arm's own base rate, not just the pre-ignition one. · mode: shadow	Pre-ignition −\$1,333.85 on 23 fires; post-ignition −\$45.00 on 17.	Arms: Shadow only. Exit: Live scaleout exits. Kill: If it is still not positive at n=60 it joins the refuted pile with the rest of the idle-gate programme.	Claude — shadow slate only Remove the arm.
7	SHADOW No change — keep lad_absS_B_15, lad_absS_A_10, sw_absS_A_k10 and sw_absS_B_k10 running, and do not promote any of them on this week. Mechanism: On the tradable-hours board these are the four most robust things this week: +\$244.00, +\$203.50, +\$183.00 and +\$186.00, four of five days green on the ladder pair, and leave-one-day-out positive on all four. But the ladder arms have n=20 in their first week and no all-time record at all. Why: They also cross-confirm Part 1.5 from a different direction: the k=1.0 stop-width arms beat the k=2.0 ones (+\$183.00 / +\$186.00 against +\$61.50 / +\$103.50), which is the same “tighter is better” ordering the live gate's own stop sweep produced. Evidence: Part 2 §3 — the promotion battery, tradable hours only. · verification: n≥60 tradable-hour fires with a positive leave-one-day-out. · mode: shadow	+\$244.00 / +\$203.50 / +\$186.00 / +\$183.00 tradable-hours this week. Best-3-stripped is negative on all four, which is what n=20 looks like.	Arms: Shadow only. Exit: As configured per arm. Kill: Leave-one-day-out negative for two consecutive weeks.	Claude — shadow slate only n/a — no change.
8	SHADOW Build the entry-side study the refuted claim detector relocated the lead to: on the 671 exits nobody watched, what distinguishes the 378 that ended in a STOP from the rest, measured at ENTRY. Shadow only, no live change, no exit rule.	−\$18,672 across 378 unwatched STOP exits, out of 671 unwatched exits.	Arms: Never — this is a study, not an arm. Exit: n/a Kill: If no entry-time feature reaches AUC 0.60 with a permutation p below 0.01 on the unwatched population, the —	Claude — study only Nothing to revert; it writes no config.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	Mechanism: The peak-watch claim detector is REFUTED (+\$1,069 in-sample, −\$4 out-of-sample on 126 held-out trades, −\$436 as the live config stands). The finding underneath it is that scoring any exit rule against the trades the operator chose to be present for was always going to lose, because his presence IS the filter. The unwatched population is where an automated rule could win, and inside it the STOP bucket is −\$18,672 across 378 trades. Why: On this report's own framing this is the single largest identified pool of money in the document, and in Rev1 it left the report with no play, no kill criterion and no owner. ★ REV2 adds it. It is deliberately a MEASUREMENT and not a rule: 'which trades are stop-bound' is an entry question, and the desk has just spent a week learning what happens when an exit answer is fitted to an entry problem. Evidence: Part 1 §6 (the claim replay); reports/claim_replay/. · verification: Report the discriminator's AUC against a matched control drawn from the same unwatched population, with a permutation p-value, before any rule is written. An AUC quoted without its control is not a measurement. · mode: shadow		\$18,672 is variance rather than a selectable bucket and the lead is closed rather than parked.	

9	Populate signal_journal.suppressed_by — the field that turns every counterfactual into a reconstruction LIVE Write the suppressing rule's name into signal_journal.suppressed_by at the moment a gate is refused, so the desk can answer 'what would this gate have done unrouted?' from a record instead of from a replay. Mechanism: The column is NULL in all 222 rows this week. Every 'if it had been armed' number in this report — Movement 2's −\$1,905.36, Part 2.6 §3's per-gate counterfactual, every bench price — is therefore a re-simulation off the tape rather than a measurement of what the live process actually refused. ★ REV2 shows exactly how much that costs: the mechanical sim reproduces 11 of 12 live abs_veto_short signals within 6 seconds and 5 points, but only 1 of 7 live exhaustion_short signals — three of the seven have no sim fire within TEN MINUTES, and they are the three that made money. A counterfactual that does not reproduce the gate's own fires is not that gate's counterfactual. Why: This is the wound that bounds every other number in the report, and the scope flagged it last week too. It is also cheap: the refusal reason already exists in the decider at the moment of refusal; nothing writes it down. Evidence: Front page standing caveats; Part 2.6 §3; Movement 2 §1. · verification: After one week, the count of non-NULL suppressed_by rows should equal the count of refused intents in the same window. If it does not, the field is being written on some paths and not others, which is worse than NULL. · mode: live (instrumentation only — writes no order)	222 of 222 rows NULL this week. 1 of 7 live exhaustion_short signals reproduced by the sim.	Arms: n/a — it is a log field. Exit: n/a Kill: n/a — an instrument, not a bet. Remove it only if it measurably slows the decider path.	Claude Revert the commit; the column goes back to NULL and nothing else changes.
---	---	---	--	--

HOLD · LEAVE ALONE — THESE ARE ALREADY RIGHT · 3 PLAYS

Standing lessons and things already in production. The action is to not touch them.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
1	The weekend-flatten unit is installed — do not disable it LIVE Nothing. gazbot7-weekend-flatten.timer is installed and enabled, fires at Sunday 22:00Z and retries every minute for two hours, and self-disables once flat. Mechanism: It exists because the EOD flatten's timer is Mon-Fri with Persistent=false, so a Friday failure gets no catch-up: next elapse Monday 20:53Z against a Sunday 22:00Z reopen. The retry loop is there because the failure being fixed was a gateway that would not answer, and a single shot at 22:00 would reproduce the original bug. ★★REV2 — THE UNIT WAS PROBED DIRECTLY AND HERE IS WHAT IT DOES AND DOES NOT CLOSE. (1) IT HAS NEVER RUN: empty ExecMainStartTimestamp, empty ActiveEnterTimestamp, `journalctl -u gazbot7-weekend-flatten` returns 'No entries'. Its next elapse is Sun 2026-08-23 22:00Z, so Sunday night is the first execution of an untested path. (2) THE SELF-DISABLE CALL WILL FAIL, AND IT DOES NOT MATTER. `systemctl disable --now` as User=alphabot returns 'Interactive authentication required' — confirmed by running it. But scripts/weekend_flatten.py removes data/ weekend_flatten_pending.json BEFORE it calls systemctl, and main() returns 0 immediately when that marker is absent, so the job is inert from the instant it succeeds. The systemctl call is `check=False`, so the failure is silent and the TIMER stays enabled — cosmetic, not	Both of Friday's EOD attempts (20:53Z and 20:57Z) died on the same gateway TimeoutError, four minutes apart. Four minutes buys nothing; two hours of retries across the reopen might. ★REV2: 0 executions to date; 120 retries covering 22:00-23:59Z only; Persistent=false.	Arms: Sundays, until it is no longer needed. Exit: n/a. Kill: n/a.	Nobody — already in production Do not.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
---	------------	----------------------	-----------------------------------	----------------

dangerous: it fires 120 no-ops next Sunday. The real defect is the silence, not the privilege. (3) IT CLOSES A WEDGE AND NOT AN OUTAGE. The three OnCalendar lines expand to 22:00:00 plus every minute of hours 22 and 23 — 120 attempts, all inside 22:00-23:59Z. Persistent=false means a box that is down across that whole window gets NO catch-up, and there is no OnCalendar line after midnight, so a wedge that outlives 23:59Z is uncovered too. Both gaps end in the same place: the position stays open until the EOD flatten's next elapse, Monday 20:53Z — the same ~21-hour live-market hole the Friday failure created. So it closes the failure mode that actually happened (a gateway that will not answer for a while) and not the one that would happen if the box itself were down.

Why: Note what it does NOT fix: if the gateway is still wedged at the reopen the loop fails sixty times exactly as designed. It is a retry, not a heal. That is why SATURDAY #1 is a human check.

Evidence: Front page; the failure it addresses is in Part 1.6 §5. · verification: **★REV2: after Sunday, check three things and not just the position — `systemctl is-failed gazbot7-weekend-flatten.service`, whether data/weekend_flatten_pending.json is gone (that is the real self-disable), and whether the timer is still enabled (it will be, because the systemctl call cannot succeed as alphabot; that is expected and harmless).** · mode: hold

2	Keep voiding the Asia block on every shadow number LIVE Nothing to build — apply the rule. Any shadow figure quoted without the 00:00–07:00Z entries removed is wrong, in either direction. Mechanism: The permanent Asia bench sits downstream of where a shadow arm records its fire, and the desk has filled nothing in that window since 5 August. Across the 38 families with $n \geq 8$ this week, —\$3,631.00 of the board was booked there. Why: It reorders the board rather than shifting it: the family second on the raw board falls to eleventh, and one promotion candidate crosses from positive to negative. Evidence: Part 2 §1–§2. · verification: Every shadow table in this report already carries both columns. · mode: hold	–\$3,631.00 of –\$13,780.50 booked in untradeable hours; abs_veto_55s +\$332 → +\$23.50.	Arms: n/a. Exit: n/a. Kill: n/a.	Nobody — a standing rule n/a.
3	Keep claiming by hand — and do not try to automate it LIVE No change. The Claim button stays as it is, and the peak/give-back detector does not go live in any form. Mechanism: Your 16 claims made +\$550.00 this week against +\$314.74 on abs_veto_short alone (★ REV2 — the venue-stop figure; §2.2's estimated-ATR version reads +\$412.93) that the machine would not have taken, at a median capture of 80% of each trade's own in-life peak. The detector calibrated against those same trades makes +\$1,069 in-sample on 341 trades and –\$4 out-of-sample on 126; applied exactly as it stands live it fires on 10 of 467 trades for –\$436. Why: Scoring a claim rule against the trades you chose to be present for was always going to lose, because your presence IS the filter. The money an automated claim could win is in the 671 exits nobody watched, and in that population the STOP bucket is –\$18,672 over 378 trades — which is an entry-side question, not an exit-side one. Evidence: Part 1 §6 and the claim-replay study; the OOS table is in reports/claim_replay/REPORT_UNWATCHED.md. · verification: n/a — the action is inaction. · mode: hold	+\$550.00 on 16 claims, 80% median peak capture. Detector: +\$1,069 in-sample, –\$4 out.	Arms: n/a. Exit: n/a. Kill: An entry-side 'is this trade stop-bound?' study coming back positive would re-open it as a different question.	Nobody — no change n/a.

NOT-AN-ACTION · WITHDRAWN / REFUTED — KEPT ON THE CARD SO NOBODY RE-PROPOSES THEM · 10 PLAYS

Every one of these was a recommendation at some point. They are printed with their autopsies so next Friday's agent does not dig the same hole.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
---	------------	----------------------	-----------------------------------	----------------

1	Promote abs_veto_55s to a live two-sided slot REFUTED Do not. It was the obvious read off the raw board and it does not survive the ruler. Mechanism: \$308.50 of its \$332.00 was booked between 00:00 and 07:00Z, a window with zero live fills since 5 August. Tradable hours: +\$23.50 on 34 fires, 35% win, 2 of 5 days green, and –\$133.50 once the single best fire is stripped. Its 50-second sibling goes negative. Why: The family is NOT dead — all-time tradable-hours it is +\$3,078.00 on 263 fires, which is the best-evidenced thing on the board	+\$332 raw → +\$23.50 tradable. All-time tradable +\$3,078.00 on 263.	Arms: n/a. Exit: n/a. Kill: A positive tradable-hours week with a positive leave-one-day-out re-opens it.	Nobody n/a.
---	---	--	---	----------------

and the reason it is still incubating. The refuted claim is specifically “promote it on this week’s numbers”.

Evidence: Part 2 §2–§3. · verification: **n/a.** · mode: hold

2	Arm grind only on clean-trend days PARKED Do not deploy it — but do not throw it away either. It is parked behind the floor question in BUILD #2. Mechanism: On 128 live grind fires the rule turns —\$958.00 GROSS into +\$580.50 GROSS and keeps 25 of 43 winners, so it is not a fake win. ★ REV2, net of the \$1.50 round trip: +\$548.00 — and 2026-07-29 alone is +\$634.50 of it on 46 fires, so drop that one session and the other eight days are —\$1,784.50. “The best rule on the board” and “arm on 2026-07-29” are nearly the same instruction on this sample. Strip its best three fires and it is —\$202; worst leave-one-out —\$123. And of 240 classified sessions, exactly ONE is both clean-trend and above the 22-point ATR floor. Why: Arming was never the binding constraint. ★ REV2 also scored the BUILDING-block version of the same idea on the same 128 fires with Movement 3 §8’s own classifier — —\$748.50 on 51 fires keeping 17 of 43 winners, worse per fire than no rule at all — so it is not parked behind the wrong condition either. In production this rule would have armed the gate on about one day in 240, and on the other 239 the floor would have refused every signal regardless. Evidence: Part 2.5 §1 — the mechanism table and the day-type conjunction. · verification: Revive it the moment the floor ladder returns a workable rung. · mode: hold	+\$580.50 gross (+\$548.00 net) vs a —\$958.00 gross control; strip-3 —\$202.00; one session, 2026-07-29, is +\$634.50 of it; 1 of 240 sessions satisfies both conditions.	Arms: n/a. Exit: n/a. Kill: If no floor rung works, this dies with the floor.	Nobody n/a.
3	Fix grind with an efficiency-ratio floor or a directional-net requirement REFUTED Do not. This is the third time it has been refuted and it keeps coming back. Mechanism: Nine cells across two families, scored on the same 128 live grind fires: every one negative. The ER-0.35 floor that is already deployed keeps 35 fires and still loses \$417.50. Why: Kept on the card with its autopsy so next Friday’s agent does not re-derive it. Evidence: Part 2.5 §1, the b- and c-family rows. · verification: n/a. · mode: hold	9 of 9 cells negative. The live ER-0.35 floor: 35 fires, —\$417.50.	Arms: n/a. Exit: n/a. Kill: n/a.	Nobody n/a.
4	Stand a gate down after two consecutive stops REFUTED Do not. THIRD independent population to refuse it. ★ REV2 fixes the count — this card said “fourth” by counting the run-state work as a population. It is not one: it refuted the same PRINCIPLE on a different question and never tested this rule on a sample of its own, so it is supporting reasoning, not evidence. Mechanism: Applied to abs_veto_short’s 22 lots this week it cuts winners along with losers and fails the 80%-winner-retention rule. The canonical list of three, which every other mention now points at: (1) mechanically on 104 signals last week; (2) live across last week’s tape; (3) abs_veto_short’s 22 lots this week. Why: The discriminator is the REGIME the losses happened in, not the count of them. Evidence: Part 1.5 §2.3 — the filter table with the winners-kept column. · verification: n/a. · mode: hold	No variant keeps 80% of the winners while beating the as-traded total.	Arms: n/a. Exit: n/a. Kill: n/a.	Nobody n/a.
5	Tighten abs_veto_short’s stop to 0.5× or 0.75×ATR REFUTED Do not — even though the sweep’s ordering makes it look attractive. Mechanism: Both tight cells are less bad than the live width but still negative: —\$124.83 and —\$57.94 on the shape sweep, —\$386.89 and —\$498.24 with the target pinned. Ten cells across two sweeps, all ten negative. Why: ‘Less negative’ is not a reason to deploy. The instruction the two sweeps support is to stop running the automated slot, not to re-size its stop. Evidence: Part 1.5 §2.2 and §2.2b. · verification: n/a. · mode: hold	10 sweep cells, all negative, at both target treatments.	Arms: n/a. Exit: n/a. Kill: n/a.	Nobody n/a.
6	Describe the 2R-partial as a variance / smoothness lever REFUTED Stop calling it that. On its own current window it does the opposite. Mechanism: Baseline —\$1,382.40 against partial —\$1,074.90, so it improves the mean by \$307.50 — while daily standard deviation goes from 42.3 to 83.3 and green-day percentage is 0% for both. Why: It was adopted as a variance lever on a backtest that showed daily vol —29%. On the live shadow window it is a mean lever that	std 42.3 → 83.3; mean +\$307.50; newest trade in the watch 2026-08-05.	Arms: n/a. Exit: n/a. Kill: A fresh window where std falls would revive the variance claim.	Nobody n/a.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	increases variance, which is a different product. Also note the watch has had no new trade since 5 August, because grind has not fired. Evidence: Part 2 §6. · verification: Re-measure on the first week grind actually trades. · mode: hold			
7	The open-rider 'inside the window' session filter REFUTED Do not deploy it. Three sessions are the entire lead. Mechanism: Across 229 replayed sessions at two lots the open rider loses \$20,739. Split by the window, 'inside' makes +\$825 on 60 sessions and 'outside' loses \$21,564 on 165 — but strip the best three inside sessions and it is −\$7,745. Why: Kept on the card because the split looks compelling in isolation and will be re-proposed otherwise. Evidence: Part 2.5 §5. · verification: n/a. · mode: hold	+\$825 on 60 sessions → −\$7,745 with the best three stripped.	Arms: n/a. Exit: n/a. Kill: n/a.	Nobody n/a.
8	Read the stay-out meter as a DIRECTION signal REFUTED Do not. It stays in service as a stay-out instrument, which is what it was built to be. Mechanism: When both of its reads agree it is right 51.7% of the time over 178 sessions. Its UP calls hit 56.7% and its DOWN calls 45.7%, and that gap is the same large short drift the gold control exposed, not skill. Why: The label on the instrument already says this. It keeps getting forgotten because the up/down asymmetry looks like signal. Evidence: Part 2.5 §5. · verification: n/a. · mode: hold	51.7% over 178 agreeing sessions — a coin.	Arms: n/a. Exit: n/a. Kill: n/a.	Nobody n/a.
9	Trade the MGC coil edge, in either direction REFUTED Do not. Both directions lose and the placebo closes it. Mechanism: Fading the coil edge loses \$1,828.93 over 2,487 trades across 251 sessions; the mirror control — breaking out of it — loses \$4,767.86 on the same trades. Zero of the placebo sets did worse, against a placebo median of −\$3,354.96. Why: The tightest-compression cell is separately PARKED (+\$616.32 on n=666) with a named revival condition: run the same placebo, hold out a year. That is a different claim from this one. Evidence: Part 2.5 §4a-§4b. MGC priced at \$10.00/point. · verification: n/a. · mode: hold	−\$1,828.93 / −\$4,767.86 on 2,487 trades; 0 placebo sets beaten.	Arms: n/a. Exit: n/a. Kill: n/a.	Nobody n/a.
10	Arm the idle gates to catch the runs we sat out REFUTED Do not. This is the fourth idle-gate lab and the first one with an honest denominator. Mechanism: 65 sat-out runs, four arms, plus a whole-week base-rate control. Live config into the pre-ignition window: −\$1,333.85 on 23 fires. Ungated: −\$1,620.40 on 60. Whole week, all six gates, every fire counted: −\$1,905.36 on 117. Even the fires that pointed the right way lost \$602.60 on 11. Why: Previous idle-gate labs reported a small positive and it was survivorship — they measured only windows already known to contain a run. The base-rate control is what makes this one different, and it cuts both ways. ★ REV2 CORRECTS THE READING REV1 TOOK OFF IT: the run windows and the whole week do NOT lose at about the same rate per fire. It is −\$57.99 a fire against −\$16.29 — a 3.6× gap, and 12.9× once the 21 overlapping fires are removed (−\$57.99 vs −\$4.49 on 96 disjoint fires), with a 20,000-resample bootstrap CI excluding zero on both readings. The ten minutes before a big run are substantially WORSE for these gates than ordinary tape, which strengthens the bench case rather than blurring it. What does NOT survive is the size of the prize — the −\$1,905.36 headline's own per-fire CI straddles zero. The comparison is solid; the total is a point estimate. Evidence: Movement 2, all sections. · verification: Re-run the base rate on a genuine trend week — this is one quiet week and the finding is scoped to it. · mode: hold	−\$1,905.36 on 117 fires across the week's tape — a point estimate whose own bootstrap CI straddles zero; −\$517.79 of it the four silent slots. What IS solid: −\$57.99 a fire in the pre-ignition windows against −\$16.29 across the week, 12.9× disjoint, CI excluding zero on both readings.	Arms: n/a. Exit: n/a. Kill: A trend week where the base rate is positive.	Nobody n/a.

STANDING CAVEATS ON EVERY DOLLAR ON THIS CARD

Every repriced number is a CEILING. Fills are modelled at the exact stop/target/trail level with no slippage, and the desk's own standing finding is that live losses run **1.1-3.1× modelled**. Treat the policy as the edge and the

headline dollar as decoration. Anything on thin n is flagged in-row as a direction, not a measurement.

The fee is \$1.50 per round trip and the multipliers are not the same. MNQ is \$2.00 a point; **MGC is \$10.00**. Every gold number in Movement 3 and in Part 2.5 Part A is priced at the gold multiplier — pricing gold with the Nasdaq one would understate it fivefold, and doing the reverse would flatter it fivefold.

Four method wounds are open and they bound everything above.

signal_journal.suppressed_by is NULL in all 222 rows this week, so every claim about what a gate would have done unrouted is a reconstruction from the tape rather than a measurement — and it is the same defect the scope flagged last week — **BUILD #9**. The router's leakage detector is scoped to a universe that barely trades (capitulation_long and rgv_short), so its zero is an empty set and not a clean week — exhaustion_short leaked -\$241 unseen — **BUILD #4**. The nightly rollup reports \$0/\$0/\$0 when the systemd journal has rotated past the day it is backfilling, which is indistinguishable from a day with no value — **BUILD #4**. selector_nightly counted a quarantined trade four extra times, inflating one day's regret by 39% (\$747.50 → \$457.50 on 10 real trades) — the phantom is FOUR rows worth \$290.00, not five worth \$362.50; the fifth day-rider row that day is the genuine claimed trade and must survive the fix — **BUILD #4**.

What did not run this cycle, named exactly. The rehabilitation phase timed out, and Part 1.5 was rebuilt from the artifacts it had already saved. In Movement 3 the four per-cause-cluster hunts did not run — gf_full_FLOW-LED.md, gf_full_OPEN-NEWS.md, gf_full_UNCLASS.md and gf_full_RIDER_ALL.md on disk are all dated 15 August, i.e. last week's artifacts, and nothing stitched them in. **The rest of Movement 3 is fresh and ran tonight:** the VACUUM cluster hunt, the 686-run lake study, the three forward tests, BIGBOARD, SMALLFADE, LOTDIV and the chop lab all ran 2026-08-22 00:00-00:39. So this is not "you got no MNQ greenfield work this week" — it is four cluster hunts missing from a movement that otherwise delivered. **Nothing on this card rests on a section that is missing.**

§ The evidence behind those plays

Everything from here down is the working. Two halves, one document. It opens with **Part 0**, which answers the only question that survives past Friday — are we getting better? — day by day, green or red, straight off the books. **Then the warm half**: what the live desk actually did (Part 1, with **the run charts** inside it — every fill of the week marked on its own session's price path, which is the fastest way to see what the prose is describing). Then the section that matters most — **Rehabilitation** (Part 1.5), where every red thing this week is walked through the full treatment instead of benched at face value, and where the failed treatments are printed by name because the failures are the point. Then **the day rider** (Part 1.6), which this week is not a footnote to the tournament but the bigger of the two books. Then the shadow board and the promotion battery (Part 2), the mid-week leads re-run adversarially including the whole gold programme (Part 2.5), and the **Router Operation Review** (Part 2.6). **Then the cold, tape-first half**: we ignore everything the desk did and ask, from the raw tape and order book, **where was the money and did we show up?** — on MNQ and on gold (M1) — then test whether the gates we already own could have caught the runs we missed (M2), and finally build gates from scratch and backtest them to death, the ones that failed included, that being the point (M3). Plain English, honest money, every claim a table — and every reprice a ceiling.

★ **Where work did not run this cycle, it says so in red at the point where it should have been**, naming the phase and what you are not getting. A fragment that is missing or stale is replaced by a dated red placeholder automatically; a phase that ran partially gets a red banner at the top of its section, and there are two of those this week — Part 1.5 and Movement 3. Nothing is silently omitted, and a gap is never a null result: “we looked and found nothing” and “nobody looked” are opposite instructions for next week. (★ REV2: in Rev1 this sentence was true of the build mechanism and false of the document — no fragment was missing, so

no banner ever rendered, and the two partial phases went unmarked. The banners below are the fix.)

WHAT IS IN HERE

1. **Front** — [the honest headline, the scorecard, WHAT TO ACTUALLY DO](#)
2. **Part 0** — [Are we getting better? — every day the desk has traded, green or red](#)
3. **Part 1** — [The live desk — what the paper tournament and the day rider actually did](#)
4. **Part 1.5** — [Live-desk REHABILITATION — never bench a gate at face value](#)
5. **Part 1.6** — [THE DAY RIDER — the second desk, its week, and the four lots it left open](#)
6. **Part 2** — [The shadow desk, the promotion battery, and the ruler that reordered the board](#)
7. **Part 2.5** — [Mid-week musings — every lead from the desk chat, tested hard](#)
8. **Part 2.6** — [ROUTER OPERATION REVIEW — the desk's #1 lever, on the stand](#)
9. **Movement 1** — [The census — the biggest runs this week, and did we show up](#)
10. **Movement 2** — [The idle-gate lab — could the gates we own have caught them? \(with a base rate\)](#)
11. **Movement 3** — [The greenfield lab — detection is solved, selection is not, and the size threshold](#)

P0 Are we getting better? — every day the desk has traded

One bar per trading day, green or red, total desk. The question every other section answers is what happened this week. This one answers **are we getting better**, which is the only question that survives past Friday.



The shape of it. Week 30 was the floor: **0 green days out of 6** and **-3,434.00** — the stretch the operator remembers as “red for many many weeks”. The most recent week ran **3 green out of 5** at **-247.50**. Of 6 weeks on record, **2 closed net positive** — and both of them are in the last three.

WEEK	DAYS	TRADED	GREEN	GROSS	FEES	NET
W29	2026-07-16 – 2026-07-19	3	1/3	-661.50	-75.00	-736.50
W30	2026-07-20 – 2026-07-26	6	0/6	-3,083.00	-351.00	-3,434.00
W31	2026-07-27 – 2026-08-02	6	1/6	+516.00	-307.50	+208.50
W32	2026-08-03 – 2026-08-07	5	2/5	-772.50	-165.00	-937.50
W33	2026-08-10 – 2026-08-14	5	3/5	+1,292.50	-177.00	+1,115.50
W34	2026-08-17 – 2026-08-21	5	3/5	-174.00	-73.50	-247.50

Gross is `sum(pnl_usd)`; NET subtracts `fees_usd` at \$1.50 per round trip per lot. **NET is the column to read** — it is what the account moved by, and it is the convention the rest of the report uses.

What changed underneath. The Claude-run router went live 2026-07-30, mid-W31 — the first green week is the one it landed in. Its measured effect is on the tail and the frequency of red days, not their depth: the worst day before it was **-1,603.00**, the worst since is **-1,294.50**, and green days went from **2 of 12** to **8 of 18**. The day-rider joined on 08-06 as a second, slower engine, and this is the week it stopped being unbeaten: **-189.00** across 7 legs, with a single Thursday long at $-\$1,294.50$ that is larger than the other four sessions made between them. It is still the engine that produced the desk’s best days; it is no longer the one that cannot lose.

Read it honestly. The cumulative total across all 30 days is still **-4,031.50** — this is a desk that has stopped bleeding wholesale, not one that is paid yet. 6 weeks is a short sample, best day +1,301.00 and worst -1,603.00 are both single sessions, and PAPER fills flatter a real book. The trend is real; the destination is not reached.

★ **Every number on this page is NET of the \$1.50 round-trip fee** — and until this revision it was not. The chart used `sum(pnl_usd)` with no fee subtraction, which made the desk's flagship progress page the one section measured with a different ruler from everything in front of it. Across 30 trading days the desk has paid **1,149.00** in fees, so the all-time book is **-4,031.50** net where the old page said **-2,882.50**. Nothing about the shape changes — the green-day counts are identical and both net-positive weeks are still net-positive — but the depth does, and the honest all-time figure is the larger one. Where you see this page quoted elsewhere at -2,882.50, that is the pre-fee number.

P1 THE LIVE DESK — the week of 17-21 August 2026

Garrath — the short version is that the desk lost \$247.50 on the clean book across five sessions, and that number is the least important one on this page. Three sessions were green and two were red, but the two red ones were not the gates being wrong about the market. Thursday was one position held for six hours through a fall it should have been stopped out of, and Friday was a software fault plus a gateway that stopped answering — and that second one is not finished. **You are still long four lots right now.** They were opened at 13:13 UTC on Friday, every automatic path that exists to close them failed on the same call, and at the last price the tape printed before the CME shut they were -\$973.00 down. That is bigger than the whole week's trading, it is not in any P&L on this page, and it is the first thing on the action card.

★ READ THE VENUE LINE BEFORE YOU READ ANY OF THESE TABLES

Every table in this section counts trades that **CLOSED**. The desk is carrying **4 MNQU6 long from 29,491.625**, opened 2026-08-21T13:13:06Z by the day rider and never closed. Marked at the last tick before the weekend (21:00:00Z, 29,370.00) that is **−121.625 points → −\$973.00**. ★ **REV2 — one price, and it is the FILLS**. Rev1 printed this position two ways and never reconciled them: 29,491.93 on the front page, in this section and in plays.json, and 29,491.63 in Part 1.6. The fills settle it — 2 @ 29,477.00 + 1 @ 29,506.25 + 1 @ 29,506.25 at 13:13:06Z is a fill-weighted **29,491.625** exactly, so −\$973.00 is the number. 29,491.93 is the entry field of day_rider_state.json, which disagrees with its own fills by 0.305pt (\$2.44 on four lots). data/weekend_flatten_pending.json carries the same wrong figure; it flattens by quantity so nothing depends on it there, but do not re-import it. The nightly review of 22:43Z Friday read it live off updatePortfolio at −\$950.27 unrealised, which is the same position a few ticks earlier.

A trades table can never contradict a benign story, because it only enumerates what closed. The one instrument that can is the venue position line, and it says four lots. Everything below is the story of the money that finished; this is the money that has not.

1 • The money, session by session

SESSION (UTC CLOSE)	LEGS	CONTRACTS	CLEAN BOOK	FLAGGED LOTS	CASH BOOK
Mon 17 Aug	9	10	\$458.00	—	\$458.00
Tue 18 Aug	21	22	\$162.50	—	\$162.50
Wed 19 Aug	9	10	\$500.00	—	\$500.00
Thu 20 Aug	3	4	−\$1,294.50	—	−\$1,294.50
Fri 21 Aug	1	3	−\$73.50	−\$83.00 (1)	−\$156.50
THE WEEK	43	49	−\$247.50	−\$83.00 (1)	−\$330.50

★ **REV2 — “Legs”, not “Lots”**. This column was labelled Lots and was counting rows. The clean week is **43 rows but 49.0 contracts**, because the day rider’s legs are 1-3 lots each — so the old label understated the week’s size by 12% and disagreed with §2’s slot table, which had it right. Both counts are printed now.

The **clean book** is data_quality IS NULL and it is the strategy verdict — what the gates did when the machinery worked. The **cash book** adds the flagged lot back, and it is what reconciles to the broker: one Friday day-rider lot filled at 29388.25, about 30 points outside anything the book was showing, flagged BADFILL:stale_price_30pt_outside_book_20260821. **A BADFILL flag marks a defective fill PRICE, not a phantom trade** — the money left the account either way, so quote the clean book when you are judging a gate and the cash book when you are counting cash. Add the open carry and the week's real economic position is **−\$1,303.50**.

2 · Every slot that fired, and the four that did not

SLOT	SIGNALS	LOTS	NET	WIN %	AVG WIN	AVG LOSS	\$ / SIGNAL	BEST	WORST	
abs_veto_short_A		11	11	−\$89.00	36%	\$47.50	−\$39.86	−\$8.09	\$86.50	−\$79.50
abs_veto_short_B		11	11	−\$114.00	27%	\$66.83	−\$39.31	−\$10.36	\$92.50	−\$79.50
day_rider		7	13	−\$189.00	43%	\$393.00	−\$342.00	−\$27.00	\$541.00	−\$730.00
exhaustion_short_A		7	7	\$88.50	57%	\$82.25	−\$80.17	\$12.64	\$117.00	−\$105.50
exhaustion_short_B		7	7	\$56.00	57%	\$74.25	−\$80.33	\$8.00	\$129.50	−\$105.50

Two base gates traded all week: abs_veto_short and exhaustion_short, both of them short-side, plus the day rider on its own book. The other four — grind_long, capitulation_long, abs_veto_long, rgv_short — fired **nothing at all**, because they were benched for the whole week under the standing tournament standdown. That is not an accident and it is not a gate failure; it is the desk's current posture, and Movement 2 prices what it saved. The short answer, in advance: firing all six mechanically across this week's tape would have lost — **\$1,905.36** on 117 fires, so the four silent gates are the best-performing thing on this page.

3 · The two legs of the same signal — is the wide one still winning?

Every tournament signal opens two lots at the same price with the same stop, and the only difference is where profit is taken: Lot A takes the tight rung, Lot B rides for the wider one. **They are a scale-out, not an experiment** — if B beats A the instruction is to widen A's target, never to drop A. Paired on the signal, this week:

GATE	PAIRED SIGNALS	LOT A (TIGHT)	LOT B (WIDE)	B – A	B AHEAD ON	B – A, BEST PAIR STRIPPED
abs_veto_short	9	–\$93.50	–\$137.50	–\$44.00	4 / 9	–\$63.50
exhaustion_short	7	\$88.50	\$56.00	–\$32.50	2 / 7	–\$45.00
BOTH GATES	16	–\$5.00	–\$81.50	–\$76.50	6 / 16	–\$96.00

–\$76.50 across 16 paired signals, with B ahead on 6 of them. **That is not a result.** On this n a single pair moves the sign, and stripping the best one leaves –\$96.00. Last week's wide-leg finding stood on 51 paired entries with all five daily folds positive; this week's 16 pairs neither confirm nor contradict it, and the honest statement is that the week was too quiet to re-test it. Do not re-rank the exit ladder on this table.

4 • How the week's lots actually ended

EXIT REASON	LEGS	CONTRACTS	NET	\$ / CONTRACT	WIN %
CLOCK_FLAT_BUG	2	4	–\$156.50	–\$39.13	0%
MANUAL_CLAIM	16	22	\$550.00	\$25.00	81%
STOP	21	21	–\$1,075.00	–\$51.19	0%
TARGET	4	4	\$166.00	\$41.50	100%
TRAIL	1	2	\$185.00	\$92.50	100%

★ **REV2 — this table is the CASH book** (44 legs / 50 contracts, –\$330.50), because the flagged BADFILL lot has to end somewhere and its exit reason is CLOCK_FLAT_BUG. Legs and contracts are both shown now: the two CLOCK_FLAT_BUG rows are a 3-lot and a 1-lot fill from one 09:11Z entry, which is how §5 and Part 1.5 §1 describe them. The CLOCK_FLAT_BUG rows are not an exit style, they are the Friday fault — Part 1.5 takes them apart. Note the shape of everything else: 21 stops for –\$1,075.00 against 16 manual claims for \$550.00 and four targets for \$166.00. **The desk's automated profit-taking barely appears this week** — one TRAIL exit for \$185.00 and four targets. Almost every dollar the desk made, your hand made.

5 • THURSDAY — the week's whole loss, in one position

Thursday 20 August cost –\$1,294.50, which is more than the other four sessions made between them. It was not the gates. The tournament did not fire a single lot that day. It was one day-rider long, four lots, opened at 13:38 UTC and closed by your own hand in three pieces between 19:33 and 20:17 as it kept going the wrong way:

CLOSED (UTC)	SLOT	SIDE	LOTS	IN	OUT	HELD	EXIT	NET	
19:33:54	day_rider	LONG		1	29,447.75	29,301.75	356 min	MANUAL_CLAIM	−\$295.00
19:43:23	day_rider	LONG		1	29,447.75	29,314.50	365 min	MANUAL_CLAIM	−\$269.50
20:17:55	day_rider	LONG		2	29,447.75	29,266.75	400 min	MANUAL_CLAIM	−\$730.00
				4					−\$1,294.50

One entry, one direction, no stop that bound, and just under six hours of holding. The rider had no protective exit in the water for any of it. Part 1.6 walks the rider's own week in detail and Part 1.5 asks the rehabilitation question — whether the entry was wrong, or whether an entry that was fine was ruined by having no stop — and the answer is not the obvious one.

6 • Your hands, scored honestly

16 of the week's 43 clean lots were closed by you pressing Claim, and they account for \$550.00 of the P&L, 13 of 16 of them green. Everything the machine closed by itself came to -\$797.50. Stated that way it reads like your hand carried the week, and in cash terms it did — but the honest version needs the counterfactual, so here is the peak that existed inside each claimed trade's own life against what you actually booked:

DAY	SLOT	LOTS	PEAK INSIDE THE TRADE	HE BOOKED	KEPT	WORST IT WENT AGAINST HIM
08-17	day_rider	2	\$498.00	\$453.00	91%	-75.25 pt
08-17	exhaustion_short_A	1	\$109.00	\$91.00	83%	-39.00 pt
08-17	exhaustion_short_B	1	\$108.50	\$74.50	69%	-39.25 pt
08-18	exhaustion_short_A	1	\$96.50	\$84.00	87%	-40.00 pt
08-18	exhaustion_short_B	1	\$81.00	\$55.50	69%	-40.50 pt
08-18	exhaustion_short_A	1	\$146.50	\$117.00	80%	-10.50 pt
08-18	exhaustion_short_B	1	\$135.00	\$129.50	96%	-10.25 pt
08-18	abs_veto_short_A	1	\$39.50	\$31.50	80%	-10.00 pt
08-18	abs_veto_short_B	1	\$69.00	\$51.00	74%	-13.75 pt
08-18	exhaustion_short_B	1	\$62.00	\$37.50	60%	-7.75 pt
08-19	day_rider	2	\$606.00	\$541.00	89%	-28.75 pt
08-19	abs_veto_short_A	1	\$116.50	\$86.50	74%	-26.00 pt
08-19	abs_veto_short_B	1	\$100.50	\$92.50	92%	-26.00 pt

DAY	SLOT	LOTS	PEAK INSIDE THE TRADE	HE BOOKED	KEPT	WORST IT WENT AGAINST HIM
08-20	day_rider	1	\$37.00	-\$295.00	-797%	-245.25 pt
08-20	day_rider	1	\$37.00	-\$269.50	-728%	-245.25 pt
08-20	day_rider	2	\$74.00	-\$730.00	-986%	-245.25 pt

Median capture of the in-trade peak: **80%**. The peak column counts only ticks between the entry and the exit — an excursion after the trade closed is money the desk was flat for, and counting it would let a chart accuse your own book of leaving money behind that was never available.

THE CLAIM REPLAY: YOUR HAND IS REAL, AND A PEAK-AND-GIVE-BACK RULE DOES NOT REPRODUCE IT

This week's claim-replay study (reports/claim_replay/) rebuilt the 1 Hz peak-watch ladder against all 35 booked manual claims on the tape. Where the watcher is actually in scope — the rider's two-lot book — an exit-shaped signal fired inside the hold on 4 of 4 trades and fired BEFORE you pulled every time, median lead 1,383 seconds. It also booked less money on every one of them: \$165 vs \$189, \$382 vs \$731, \$137 vs \$456, \$454 vs \$544.

Then it was calibrated properly — 75 cells on 341 training trades, best cell taken to 126 held-out ones. In-sample it makes **+\$1,069**. Out-of-sample it makes **-\$4**. Applied exactly as it stands live it fires on 10 of 467 trades for **-\$436**. **The give-back exit shape is REFUTED as a way to reproduce your hand**, over 44 + 75 calibrations on two independent populations.

The finding underneath it is the useful part: scoring any rule against the trades you chose to be present for was always going to lose, because your presence is the filter. The money an automated claim could win is in the 671 exits nobody watched — and in that population the STOP bucket is **-\$18,672** over 378 trades. That is an ENTRY-side question (which trades are stop-bound?), not an exit-side one, and the lead has been moved there rather than killed — **BUILD #8** carries it, and it is a study rather than a rule, because the desk has just spent a week learning what happens when an exit answer gets fitted to an entry problem.

7 · Three instruments that reported confidently and wrongly this week

INSTRUMENT	WHAT IT SAID	WHAT WAS TRUE	COST OF BELIEVING IT
core_health.json · the router's own health line	"flat": true, "healthy": true — still says it as this report renders	The venue has been long 4 lots since 13:13Z Friday	It is the flag every other tool trusts. 132 of the 137 router ticks since the entry repeated the word "flat".
hour_watch.py · the hourly desk read	Friday day net −\$150.50	−\$69.00 on the clean book — it has no data_quality filter, so the BADFILL lot lands in its total	Neither number is wrong; they answer different questions. IB's own realizedPNL −\$149.38 reconciles to the FLAGGED figure, so hour_watch is right about CASH and wrong about STRATEGY.
router_nightly / selector_nightly	Nothing at all for this week	The newest router rollup is 2026-08-14 and the newest selector rollup 2026-08-18 — no timer runs either of them	The weekly router review in Part 2.6 had to be rebuilt from raw logs. A monitoring file that silently stops being written is worse than one that errors.

8 · What actually happened, in five sentences

Monday was the week's clean session: the short book worked in a falling tape, the rider took \$453.00 out of one claimed short, and the desk closed \$458.00. **Tuesday** was the busiest day by far at 21 lots for \$162.50 — a genuine chop day where exhaustion_short faded the oscillation well and abs_veto_short gave most of it back. **Wednesday** was the best at \$500.00, essentially one rider short claimed for \$541.00 inside eleven minutes, after which the short book stopped out four times in a row for −\$153.00 and the operator benched it at 18:15Z. **Thursday** was −\$1,294.50 in the single rider long described above. **Friday** booked −\$156.50 on a software fault at 09:12Z, then opened the four-lot long at 13:13Z that is still open, and at 19:57Z the IB gateway entered the partial wedge that stopped every closing path on the box.

Disposition — every lead this section touched

LEAD	VERDICT	FOR PARKED: WHAT WOULD REVIVE IT. FOR REFUTED: THE TEST THAT KILLED IT.
The four silent gates (grind_long, capitulation_long, abs_veto_long, rgv_short)	HOLD	They traded nothing and that was correct: the mechanical counterfactual in Movement 2 is −\$1,905.36 across this week's tape. Keep them benched; revisit only on a measured trend week, not on a quiet one.
abs_veto_short as a live slot	SHADOW	−\$203.00 over 22 lots this week, and its 55-second shadow mirror is the week's best shadow family (+\$332 over 46). The

LEAD	VERDICT	FOR PARKED: WHAT WOULD REVIVE IT. FOR REFUTED: THE TEST THAT KILLED IT.
		live/shadow gap is the whole question — Part 1.5 takes it apart and Part 2 tests whether the mirror is promotable.
exhaustion_short as a live slot	LIVE	The only live slot green on the week: + \$144.50 over 14 lots at 57% win. ★ REV2 — it has now been interrogated, and the disposition needs a caveat. Repriced machine-only on its own 14 lots (venue stops, so R is exact; 7 of 7 machine-closed lots reproduce their live exit reason), the machine alone books +\$84.00 and your hand adds +\$60.50 — so unlike abs_veto_short this gate is NOT hand-carried. But strip its best single lot and the machine is −\$91.50; leave 08-18 out and the machine is −\$317.00; and its live 1.5×ATR stop is a spike, not a plateau (0.75× −\$22.00, 1.0× −\$163.68, 1.25× −\$305.36, 2.0× −\$140.36). It also ran the whole week on a quiet-tape clip the operator reverted on 08-18 as UNMEASURED. Keep it live — nothing here says bench it — but stop calling the + \$144.50 evidence. Part 1.5 §2.4 has the full battery.
The wide-leg (Lot B) exit finding	PARKED	16 paired signals is too thin to re-test it and the sign flips on one pair. (★ REV2: this row said 22, which is the abs_veto_short LOT count, not a pair count — so the revival condition $n \geq 40$ pairs was being measured against a different denominator from §3's own table.) Revive it next week on $n \geq 40$ pairs, or sooner if the shadow ladder arms produce their own pairs.
The automated claim detector (peak / give-back)	REFUTED	Out-of-sample −\$4 on 126 held-out trades against +\$1,069 in-sample, and −\$436 as the live config stands. Two independent populations, 119 calibrations. Relocated to the entry-side question rather than killed — and ★ REV2 gives that relocation a play: BUILD #8, with an AUC-0.60 / permutation-p kill criterion. In Rev1 the largest identified pool of money in the report (−\$18,672 over 378 unwatched STOP exits) left the document with no play, no kill criterion and no owner.
core_health.flat as a safety signal	FAULT	It is tournament-scoped and it says the desk is flat while the venue is long four lots. BUILD item: make it read the venue position line, or rename it.


THE RUN CHARTS — every fill of the week, marked on the day's own price path



One trade is the whole day: the day_rider held a short from 13:38 to 18:04 for +\$453 net, while the seven gate legs around it netted +\$6 between them.




The busiest session of the week - sixteen exhaustion/abs_veto legs machine-gunned the 13:37-15:10 chop and gave back \$50.50 net, so once again the day_rider's single short (+\$185 net) is the entire day.



The best day of the week was over by 13:52 - the desk caught the 13:36 crash for +\$720 net in thirteen minutes, then handed \$220 back to six shorts that were stopped in the afternoon drift back up.



Three markers, one position: the day_rider bought 29,447.75 at 13:38 and scaled out of a single long over the next six and a half hours as price slid 180 points - the only trade the desk took all day, and the worst of the week.



An empty chart is the finding: the desk did not trade the Friday session at all, so a 158-point drop through the cash open and a 115-point afternoon rally both printed with the book flat (the day's only fills were a 09:11 day_rider bug, hours outside this window).




The one the desk got right: the day_rider was short three minutes after the ignition and banked +\$541 net out of a 207-point flush, and two abs_veto legs added +\$179 on the last leg down - \$720 net against a \$414 one-lot ceiling.




The largest sat-out run of the week: price left the 29,376 low and never looked back, and the desk's only response was to sell 29,543 at 14:51 and be stopped 35 seconds later for -\$87 net while the rally carried on to 29,630.



Friday's cash open dropped 158 points in fifteen minutes and kept going to 29,257 - the single cleanest short of the week, in the direction every armed gate trades, and not one of them fired.



A 150-point rally into the 29,759 pre-open high, untouched: the desk does trade this hour (it did on Monday and Tuesday) but every gate that fired this week was a short, so an up-run like this has nothing pointing at it.



Three hours after sitting out the Friday flush, the desk sat out the bounce as well - 115 points up off 29,331 with no position on either side of it.

★ GAP — THE REHABILITATION PHASE TIMED OUT — THIS SECTION WAS REBUILT FROM WHAT IT HAD SAVED

The phase did not finish. Everything here was reconstructed from the artifacts it had already written to disk before it was killed, plus the REV2 work done afterwards (§2.4, §2.5 and the corrected rider-stop grid). **What that means for you:** the rehab treatments listed are the ones that got run, not the ones that were planned — so an absent treatment here means nobody looked, not we looked and found nothing. Nothing on the action card rests on a treatment that is missing.

P1.5 REHABILITATION — never bench a gate at face value

Garrath — the rule is that nothing gets written off on its P&L alone. If a gate is red, the job is to find out how it could work, and to print the treatments that failed by name so nobody digs the same hole next Friday. Three things were red this week and they are red for three completely different reasons, which is the point: one is a software fault, one is a missing stop, and only one of them is actually a gate. **The headline is that the gate everyone would have benched is the one with the least wrong with it, and the loss that looks like bad luck was structural.**

1 · First, strip the malfunctions out of the strategy verdict

Two of the week's lots are not trading decisions at all. They have to come out before any gate is judged, and they have to be reported, because a fault that gets quietly absorbed into a gate's P&L is a fault nobody fixes:

LOT	WHAT IT WAS	BOOKED	IS IT THE STRATEGY'S FAULT?
#761 day_rider LONG 3 @ 29428.25 → 29417.50	The 00:00Z session roll flipped day_rider_state.json to entered: false / note: "switch off" while venue_net still read 4.0. The rider then flattened a position it had just told itself it	-\$73.50	No. The entry was 09:11:06Z and the flat was 09:12:06Z — one timer tick. Nothing about the market happened in that minute.
-\$239.50			

LOT	WHAT IT WAS	BOOKED	IS IT THE STRATEGY'S FAULT?
	did not have, 60 seconds after opening it.		
#762 day_rider LONG 1 @ 29428.25 → 29388.25	The 1-lot remainder of the same 3+1 exit. It filled 26 points below a 30-minute tape low, ~92 ms after the 3-lot filled on-tape. Flagged BADFILL:stale_price_30pt_outside_book_20260821.	-\$83.00	No, and note it is a SELL — the 'off-tape fills are always BUYs' reading is dead. Same fault, worse fill.
		-\$239.50	

With both removed, the week's strategy book is **-\$174.00** rather than -\$330.50. Note what the fault actually was: the rider branched on entered, and at the UTC midnight session roll entered gets reset to false while venue_net still reads the true position. It is a one-line fix — **branch on venue_net != 0, never on entered** — and it is on the Saturday card.

2 · REHAB TARGET ONE — abs_veto_short, -\$203.00 over 22 lots

This is the only red gate that is actually a gate. It fired 22 lots at 32% win for -\$203.00, and the reflex is to bench it. Before that, the treatment.

2.1 — What KIND of loser is it? (the root-cause step)

There are only two kinds of losing trade and they need opposite fixes. A **chase** never goes green at all — it is adverse from the first tick, and that is an ENTRY problem. A **give-back** goes green and hands it back, and that is an EXIT problem. Measured inside each trade's own life, on the tick tape:

CLOSED (UTC)	LEG	ENTRY	BEST IT GOT (PT)	WORST IT GOT (PT)	NET	EXIT	PEAK, IN MONEY
08-17 12:43	A	30,208.50	+21.25	-4.25	\$34.50	TARGET	\$41.00
08-17 13:25	B	30,208.50	+28.50	-19.00	-\$35.50	STOP	\$55.50
08-17 15:51	A	30,223.50	+11.25	-13.00	-\$26.50	STOP	\$21.00
08-17 15:51	B	30,223.50	+11.25	-13.00	-\$26.50	STOP	\$21.00
08-18 09:33	A	29,738.25	+20.25	-10.75	\$37.50	TARGET	\$39.00
08-18 09:41	B	29,738.75	+30.25	-10.25	\$57.00	TARGET	\$59.00
08-18 12:20	A	29,734.00	+8.50	-15.50	-\$33.00	STOP	\$15.50
08-18 12:20	B	29,734.00	+8.50	-15.50	-\$33.50	STOP	\$15.50
08-18 14:02	A	29,659.00	+38.25	-38.25	-\$79.50	STOP	\$75.00
08-18 14:02	B	29,659.00	+38.25	-33.50	-\$79.50	STOP	\$75.00
08-18 14:59	A	29,557.50	+20.50	-10.00	\$31.50	MANUAL_CLAIM	\$39.50
08-18 15:10	B	29,557.50	+35.25	-13.75	\$51.00	MANUAL_CLAIM	\$69.00

CLOSED (UTC)	LEG	ENTRY	BEST IT GOT (PT)	WORST IT GOT (PT)	NET	EXIT	PEAK, IN MONEY
08-18 16:32	A	29,609.50	+6.00	-12.25	-\$30.00	STOP	\$10.50
08-18 16:33	B	29,609.75	+6.25	-12.00	-\$29.50	STOP	\$11.00
08-19 13:52	A	29,466.75	+59.00	-26.00	\$86.50	MANUAL_CLAIM	\$116.50
08-19 13:51	B	29,466.75	+51.00	-26.00	\$92.50	MANUAL_CLAIM	\$100.50
08-19 14:51	A	29,543.00	+6.00	-19.75	-\$43.50	STOP	\$10.50
08-19 14:51	B	29,543.00	+6.00	-20.50	-\$43.50	STOP	\$10.50
08-19 17:28	A	29,527.50	+17.25	-12.75	-\$28.50	STOP	\$33.00
08-19 17:28	B	29,527.25	+17.00	-13.00	-\$29.00	STOP	\$32.50
08-19 18:12	A	29,514.00	+15.25	-15.75	-\$38.00	STOP	\$29.00
08-19 18:12	B	29,514.75	+16.00	-17.00	-\$37.50	STOP	\$30.50

0 of the 15 losing lots were chases. Not one of them was adverse from the first tick — the worst-behaved loser still got 6 points onside and the median loser reached **+11.25 points**, which is \$21.00 a lot sitting on the screen before it turned round. So the reflex diagnosis is wrong. This is not a gate buying tops; every one of these entries worked, briefly.

The difference between a winner and a loser here is not whether the trade went onside, it is how far: the winners' median peak is **+30.25 points** against the losers' +11.25. That is a give-back shape, and give-back shapes are an EXIT question — so the exit is where the treatment has to start, which is the opposite of where this desk usually looks.

2.2 — The stop-width sweep, run twice, because the obvious version cheats

Every lot re-raced from its own entry price on the real tick path at five widths, first touch wins, stop wins a tie inside a gap, three-hour cap. The first pass is the usual one: risk = $k \times \text{ATR}$ and the target moves with it, so the whole trade keeps its shape.

STOP WIDTH	LOTS	NET	WIN %	\$ / LOT	WORST LEAVE-ONE-DAY-OUT	BEST LOT STRIPPED	STOP / TARGET / 3H-CAP
0.50×ATR	22	-\$124.83	32%	-\$5.67	-\$189.47	-\$170.07	15 / 7 / 0
0.75×ATR	22	-\$57.94	32%	-\$2.63	-\$118.53	-\$149.45	15 / 7 / 0
1.00×ATR ← live	22	-\$548.73	18%	-\$24.94	-\$507.28	-\$607.02	18 / 4 / 0
1.25×ATR	22	-\$778.01	14%	-\$35.36	-\$627.37	-\$851.24	19 / 3 / 0
1.50×ATR	22	-\$1,076.45	9%	-\$48.93	-\$896.89	-\$1,162.81	20 / 2 / 0

Calibration first, because a sweep you cannot check is decoration. Of the 22 lots, 18 were closed by the machine on a stop or a target. Re-raced at the live 1.00× width they come to −\$397.30 against −\$464.50 actually booked, and **18 of 18 land on the same exit reason.** The remaining 4 lots you claimed by hand: you booked \$261.50 on them, the simulator rides them to their own stop or target for −\$151.43. **That gap — \$412.93 — is your hand on this one gate,** and it is the whole reason the live book reads −\$203.00 where the machine alone reads −\$548.73.

The verdict: the sign survives all five widths. Every cell is negative, from −\$124.83 at the tightest to −\$1,076.45 at the widest, and the worst leave-one-day-out is negative at every width too. That is a real verdict rather than a knife-edge, and it is bankable: left entirely to the machine, this gate lost money this week at every stop width tested.

2.2b — But is that the STOP, or is it the target riding along with it?

The sweep above moves the stop and the target together, because R is defined on risk. That means “tighter is better” could just be “a nearer target gets hit more often”, which is a completely different instruction. So here is the same sweep with the target pinned where the live config actually puts it, moving only the stop:

STOP WIDTH	LOTS	NET	WIN %	\$ / LOT	WORST LEAVE-ONE-DAY-OUT	BEST LOT STRIPPED	STOP / TARGET / 3H-CAP
0.50×ATR	22	−\$386.89	5%	−\$17.59	−\$387.24	−\$433.55	21 / 1 / 0
0.75×ATR	22	−\$498.24	9%	−\$22.65	−\$477.70	−\$556.53	20 / 2 / 0
1.00×ATR ← live	22	−\$548.73	18%	−\$24.94	−\$507.28	−\$607.02	18 / 4 / 0
1.25×ATR	22	−\$728.09	18%	−\$33.09	−\$665.74	−\$786.38	18 / 4 / 0
1.50×ATR	22	−\$659.37	23%	−\$29.97	−\$576.13	−\$781.89	17 / 5 / 0

Pinned target, moving stop: 0.50× −\$386.89 · 0.75× −\$498.24 · 1.00× −\$548.73 · 1.25× −\$728.09 · 1.50× −\$659.37. It still gets worse as the stop widens, so the effect is genuinely the stop and not the target. Either way both tables agree on the thing that matters: **every cell in both sweeps is negative.** There is no stop width at which this gate, run by the machine alone, made money this week.

2.3 — Filters, judged by whether they keep the winners

The rule here is old and it is the one that catches most bad ideas: a filter that improves the P&L by dropping the trades that WON is a fake win, however good the total looks.

TREATMENT	THE IDEA	LOTS KEPT	NET	Δ VS AS-TRADED	WINNERS KEPT	VERDICT
Drop the 12:00-14:00Z open window	the loudest two hours of the day	16	−\$314.50	−\$111.50	4 / 7	no better
US cash session only (13:30-20:00Z)	trade only when the depth is real	16	−\$230.00	−\$27.00	4 / 7	no better
Europe/pre-open only (before 13:30Z)	the mirror of the row above	6	\$27.00	\$230.00	3 / 7	FAKE
Stand down after two consecutive stops	the wall-of-stops rule, again	12	−\$40.00	\$163.00	4 / 7	FAKE

Nothing clears the bar. The two that beat the as-traded number do it by cutting winners along with losers, and the stand-down-after-two-stops rule is the “wall of stops” idea. ★ **REV2 fixes its count, which the report gave as three here and four on the action card.** It is **THREE** independent populations, and this is the canonical list — every other mention of the rule points here: **(1)** mechanically on 104 signals last week; **(2)** live across last week’s tape; **(3)** abs_veto_short’s 22 lots this week, in the table above. The card said “fourth” by counting the run-state work as a population; it is not one — that work refuted the same principle on a different question and never tested this rule on a sample of its own, so it is supporting reasoning, not evidence. Three is the number. **Post it as REFUTED and stop proposing it.**

SO WHAT IS THE TREATMENT FOR ABS_VETO_SHORT?

Three things are now closed off. It is NOT a stop-width change — ten cells across two sweeps, every one negative. It is NOT a filter — none keeps the winners. And it is NOT a bench on a knife-edge, because the sign survived, so a bench here would at least be honest.

What the excursion table actually says is that the entries work and then stop working: median peak +30.25 points on the winners against +11.25 on the losers, with nothing adverse from the first tick. The live experiment that addresses exactly that is the 55-second continuation-confirm — it re-requires

the burst to PERSIST before entering, which is a filter on the difference between an 11-point poke and a 30-point move. This week its mirrors are the top of the shadow board (+\$332 over 46 fires) while the live gate lost \$203 over 22. **Part 2 is where that promotion is argued, and it is the one real proposal this report makes about this gate.**

The honest caveat, stated before Part 2 rather than after it: the shadow mirror does NOT model the live gate's default-veto on thin tape, so it over-fires, and it enters on the 1-minute bar open where the live decider quotes the tick it fired on. Both biases run in the mirror's favour. It is an upper bound, and it is presented as one.

2.4 — ★ REV2: THE EXIT LADDER — the treatment §2.1 diagnosed and §2.2-§2.3 never ran

THE DANGLING DIAGNOSIS

§2.1 ends on a sentence that commits this section to something: “that is a give-back shape, and give-back shapes are an EXIT question — so the exit is where the treatment has to start.” What followed was §2.2 sweeping the **stop**, §2.2b sweeping the **stop** again, §2.3 sweeping time-of-day and stand-down **filters**, and a recommendation of the 55-second confirm, which is an **entry** filter. Not one rung of an exit ladder was ever tested on the gate this report says has an exit-shaped loss. Here they are, on the same 22 lots, with the winners-kept column carried all the way through.

Every lot raced from its own entry on the tick tape, three-hour cap, stop wins a tie. **The stop and the target come from the VENUE, not from an ATR estimate** — every lot's protective StopLimitOrder auxPrice is in the tournament journal, so $R = |\text{entry} - \text{stop}|$ is exact, and the live target model (the `atr_split=22` quiet-tape clip: Lot A banks \$40, Lot B takes $\max(1.75 \times \text{ATR}, \$60)$ below a 22-point ATR) is confirmed to the tick. **Calibration: 18 of 18 machine-closed lots reproduce their own live exit reason.**

EXIT TREATMENT	NET	\$ / LOT	WIN %	WINNERS KEPT	WORST LEAVE- ONE-DAY- OUT	BEST LOT STRIPPED	WHAT IT IS
AS TRADED (your hand + the machine)	-\$203.00	-\$9.23	31.8%	7 / 7	-\$162.00	-\$295.50	the live book
MACHINE, live exit (baseline)	-\$446.24	-\$20.28	22.7%	5 / 7	-\$420.00	-\$516.50	live stop, live target
Scalp at +11.25pt (losers' median peak)	\$141.00	\$6.41	72.7%	7 / 7	\$57.00	\$120.00	bank the poke
Scalp at 0.5R	\$178.50	\$8.11	81.8%	7 / 7	\$99.00	\$140.25	half the risk
Scalp at 1.0R	\$187.00	\$8.50	54.5%	7 / 7	\$60.00	\$109.00	one-for-one
Trail after +15pt, 15pt behind the peak	\$107.50	\$4.89	59.1%	7 / 7	\$12.50	\$21.00	give-back cap
Trail after +15pt, 8pt behind the peak	\$149.00	\$6.77	63.6%	7 / 7	\$15.00	\$71.00	tight give-back cap
Breakeven stop once +15pt	\$107.50	\$4.89	59.1%	7 / 7	\$12.50	\$21.00	same as trail-15
Lot A scalp +11.25pt / Lot B rides	-\$162.74	-\$7.40	45.5%	6 / 7	-\$182.50	-\$233.00	split the pair

The ladder does not fail. It is the first treatment this gate has been given that keeps every winner. Four rungs beat the as-traded book while holding 7 of 7 winners, stay positive on the worst leave-one-day-out fold, and stay positive with their best lot stripped: a scalp at **1.0R** (+\$187.00), at **0.5R** (+\$178.50), a **trail 8pt behind the peak once +15pt** (+\$149.00), and a scalp at the losers' own median peak of **+11.25pt** (+\$141.00). Against a machine baseline of -\$446.24 and an as-traded -\$203.00, that is a swing of six hundred dollars from changing nothing but where profit is taken.

★ AND THE PLACEBO, WHICH IS WHY THIS IS A SHADOW ROW AND NOT A MONDAY CHANGE

A tight target flatters any entry in a simulator, so the rung has to beat a random-entry control before it counts. 400 draws, same 22-lot shape, same R distribution, same tape, same fee:

RUNG	REAL	PLACEBO MEDIAN	PLACEBO 5-95%	WHERE THE REAL RESULT SITS
Scalp at 0.5R	\$178.50	-\$6.75	-\$219.00 ... \$192.75	93rd pct
Scalp at 1.0R	\$187.00	\$14.00	-\$305.00 ... \$344.00	80th pct
Scalp at +11.25pt	\$141.00	\$20.00	-\$232.00 ... \$258.00	80th pct
Trail +15pt / 8pt behind	\$149.00	\$1.00	-\$284.00 ... \$259.00	83rd pct

The entry is adding something, and not much. A random entry on this week's tape, scalped at 0.5R against a 1R stop, lands at -\$6.75 — which is what a fair coin minus the fee should look like, and is a good sign the harness is honest. The real signals land at the 80th to 93rd percentile of that distribution. On n=22 that is a direction, not a measurement: none of these rungs clears its own placebo band. So the ladder goes to **SHADOW**, beside the 55s arms, and nothing changes live on Monday.

What this does close. The diagnosis in §2.1 no longer dangles. The exit IS where this gate's money is — every rung that banks earlier than the live 1.5R/2.5R pair improves it, and the losers' median peak of +11.25pt was real money sitting on the screen on 22 of 22 lots. It also sharpens what §2.2 actually proved: the stop sweep found no width that rescues the gate because the stop was never the problem. And it makes the machine-only baseline honest — on the exact venue stops the machine books **-\$446.24** and your hand is worth **+\$314.74**, not the -\$548.73 / +\$412.93 in §2.2, which used an estimated ATR and did not model the quiet-tape clip. Both readings agree on the sign and on the verdict; the venue-stop version is the one to quote.

2.5 • ★ REV2 — REHAB TARGET ONE-AND-A-HALF: exhaustion_short, the green one nobody interrogated

WHY A GREEN GATE GETS THE SAME TREATMENT AS A RED ONE

This section's standing rule is that nothing gets written off on its P&L alone. The mirror of that rule is that nothing gets waved through on its P&L alone either, and Rev1 did exactly that: exhaustion_short is the only live slot green on the week (+\$144.50, 14 lots, 57%) and it received no rehab pass, no stop sweep, no leave-one-day-out, no strip-best and no hand-versus-machine split — while Part 2.6 §3 quietly attributed it -\$638.02

mechanically on the same week. A \$782.52 contradiction sat in the report unexamined. Here is the same battery abs_veto_short got.

2.5a — Is the +\$144.50 the gate, or your hand?

Every one of the 14 lots re-raced from its own entry to its own configured stop or target, three-hour cap. R comes from the venue: each lot's protective stop order is in the tournament journal, and exhaustion_short runs stop_k=1.5, so entry ATR = R / 1.5 exactly. **7 of the 7 machine-closed lots reproduce their own live exit reason**, repricing to −\$435.50 against −\$444.50 booked — \$9.00 apart on seven lots.

BOOK	NET	WIN %	STRIP BEST 1	STRIP BEST 3
LIVE — your hand plus the machine	\$144.50	57.1%	\$15.00	−\$193.00
THE MACHINE ALONE	\$84.00	42.9%	−\$91.50	−\$422.50
... so the hand is worth	\$60.50	—	—	—

The answer is: mostly the gate. That is the opposite of abs_veto_short, where the machine alone loses −\$446.24 and the hand is the entire live result. Left to itself exhaustion_short still books **+\$84.00**. On the report's own logic that is a real difference between two gates that look superficially similar, and it is the reason this one stays live.

2.5b — And then the robustness battery, which it does not pass

TEST	LIVE BOOK	MACHINE ALONE	VERDICT
Strip the single best lot	\$144.50 → \$15.00	\$84.00 → −\$91.50	One lot is 90% of the live result and all of the machine's
Strip the best three	→ −\$193.00	→ −\$422.50	Fails on both books
Leave one day out (it only traded two)	drop 08-18 → \$59.00 · drop 08-17 → \$85.50	drop 08-18 → −\$317.00 · drop 08-17 → \$401.00	The machine result is entirely 08-18. One of two folds is deeply negative
Parameter plateau on the stop	see the sweep below — the live 1.5× is a spike, not a plateau		Fails

STOP WIDTH (MACHINE ALONE)	NET	WIN %	\$ / LOT	WORST LEAVE-ONE- DAY-OUT	BEST LOT STRIPPED	STOP / TARGET / 3H-CAP
0.5×ATR	\$29.66	21.4%	\$2.12	−\$109.68	−\$125.84	11 / 3 / 0
0.75×ATR	−\$22.00	28.6%	−\$1.57	−\$161.50	−\$177.50	10 / 4 / 0
1×ATR	−\$163.68	28.6%	−\$11.69	−\$213.36	−\$319.18	10 / 4 / 0
1.25×ATR	−\$305.36	28.6%	−\$21.81	−\$265.22	−\$460.86	10 / 4 / 0
1.5×ATR ← live	\$83.98	42.9%	\$6.00	−\$317.04	−\$91.52	8 / 6 / 0

STOP WIDTH (MACHINE ALONE)	NET	WIN %	\$ / LOT	WORST LEAVE-ONE- DAY-OUT	BEST LOT STRIPPED	STOP / TARGET / 3H-CAP
2×ATR	-\$140.36	42.9%	-\$10.03	-\$420.72	-\$315.86	8 / 6 / 0

The live 1.5×ATR width is an isolated positive cell surrounded by negative ones — 0.75× -\$22.00, 1.0× -\$163.68, 1.25× -\$305.36, 2.0× -\$140.36. A real effect gives a plateau; this is a step. And 0.5× is positive too (+\$29.66), which makes the shape a pair of spikes rather than a slope. On 14 lots across two sessions that is what you would expect from noise.

★ AND THE CONFIG THAT PRODUCED THE GREEN WEEK WAS REVERTED MID-WEEK AS UNMEASURED

Every exhaustion_short fire this week ran with atr_split=22 active. That clip was put back on 08-16 at 21:15Z by revert_clip_lift's daily restore job against a stale roster, and the operator removed it again on **08-18 at 15:48Z** (commit 465c943) with the words "the desk then ran an UNMEASURED cell live on its only armed short". Below a 22-point ATR it makes Lot A bank at \$40 and Lot B at max(1.75×ATR, \$60). Two of the seven signals fired under it, and it is confirmed to the tick — the 14:55 Lot A exited TARGET at exactly 20.00pt, which is \$40. So the gate's only green week was booked on a configuration that is no longer live and was never graded.

The disposition, restated honestly. Keep exhaustion_short live — nothing here says bench it, and it is the one gate on the desk whose machine half is positive without the operator. But "the only live slot green on the week, no change proposed" needs the same caveat abs_veto_short got: **+\$144.50 on 14 lots across two sessions, under a config that was reverted mid-week, does not survive strip-best-1 on the machine book and has a negative leave-one-day-out fold.** It is a green number, not evidence. Do not re-tune anything on it and do not cite it as a promotion precedent.

3 · REHAB TARGET TWO — the day rider's Thursday, -\$1,294.50

One long, four lots, entered 13:38:06Z at 29,447.75 on the US open, closed by your hand in three pieces between 19:33 and 20:17. Inside its own life it got +19.25 points onside at best and -245.25 points offside at worst. **It never went**

green in any meaningful way. So the same root-cause question: was it the entry or the exit?

The entry is defensible — it is the rider's normal open-hour entry and it has made money all week doing exactly that. What is not defensible is that there was no stop in the water at any point. Here is what one would have done, raced on the real tape from the rider's own entry, with the ATR the live detector itself read at the confirming tick, **34.2** points (★ REV2 — Rev1 used an estimated 36.1 here while Part 1.6 §3 printed the detector's own 34.2 for the same trade; a stop the rider actually carried would be sized off arm_atr, which is the detector's read, so 34.2 is the number and every cell below is recomputed on it, net of the \$1.50 round trip):

IF THE RIDER HAD CARRIED THIS STOP	IT WOULD HAVE FILLED AT	AT	4 LOTS WOULD HAVE COST (NET)	VS WHAT ACTUALLY HAPPENED	... AND ON THE THREE WINNING DAYS
0.50×ATR = 17.1 pt	29,430.65	13:40:03Z	−\$142.80	\$1,151.70	kills all three (−\$291.60)
0.75×ATR = 25.7 pt	29,422.10	13:40:06Z	−\$211.20	\$1,083.30	kills all three (−\$290.40)
1.00×ATR = 34.2 pt	29,413.55	13:40:13Z	−\$279.60	\$1,014.90	keeps 1 of 3 (\$301.00)
1.25×ATR = 42.8 pt	29,405.00	13:40:20Z	−\$348.00	\$946.50	keeps 1 of 3 (\$242.50)
1.50×ATR = 51.3 pt ← the width Rev1 picked	29,396.45	13:40:40Z	−\$416.40	\$878.10	keeps 1 of 3 (\$184.00)
2.00×ATR = 68.4 pt	29,379.35	13:47:32Z	−\$553.20	\$741.30	keeps 1 of 3 (\$67.00)
3.00×ATR = 102.6 pt ← REV2 picks this one	29,345.15	13:51:24Z	−\$826.80	\$467.70	keeps 3 of 3 (\$1,179.00)

★ **REV2 — READ THE LAST COLUMN, WHICH REV1 DID NOT HAVE.** Every width beats what happened on Thursday, and the tightest ones beat it by more than a thousand dollars — but Thursday is the one day the rider lost money, and pricing a stop only there is the fake-win failure this same section refuses to accept two screens earlier. The full widths×days grid is in [Part 1.6 §3b](#). Its verdict: **1.5×ATR takes the week's closed rider book from −\$272.00 to −\$388.90,** because 51.3pt sits inside Monday's −75.25pt excursion and 46.4pt inside Tuesday's −64.50pt, so it stops both winners out. **3.0×ATR is the only width that is positive on the closed book (+\$195.70) AND keeps all three winners.** The honest week-level value of the change is **+\$468 to +\$704**, not \$1,000.

What is NOT in dispute: having a stop at all. Every width beats what happened on Thursday, **This is not hindsight about direction** — it is not asking the rider to have known the tape was going down. It is asking it to have had any exit at all, at any of six widths, all of which are cheaper than holding to your hand at 19:33. The rider ran naked for 6.7 hours.

★ AND A LOSING RIDER IS STRUCTURALLY SILENT — THAT IS THE REAL DEFECT

The reason this ran for six hours is that nothing was shouting. The rider pages on peak give-back and on stalls; below its arm threshold there is no stop, no trail and no alert, so a position that is simply losing produces no output at all. Both of the tools you would naturally check print “flat, \$0” while it happens, because the trade has not CLOSED yet and they count closed rows.

The rehabilitation of the rider is not an entry change and it is not a threshold change. It is a stop. A 3.0×ATR protective stop on the rider's own book, sized off its entry ATR, is the single highest-value change on this week's card — and ★ REV2 corrects both halves of what Rev1 said here: not 1.5×, and not “\$1,000 on one day” but +\$468 on the week’s closed book. It is also the change that removes the report’s headline: every width tested, including 3.0×, closes Friday’s 13:13Z entry by 13:31:55Z — six and a half hours before the gateway wedged.

4 • REHAB TARGET THREE — the four gates that never fired

The standing rule cuts the other way too: a gate that is benched is not being rehabilitated, it is being ignored. So Movement 2 fired all four of them mechanically, in their exact live config, across this whole week's captured tape — the honest denominator, every fire counted, not just the ones near a big move:

GATE (BENCHED ALL WEEK)	RAW SIGNALS ON THE TAPE	WOULD HAVE FIRED	WOULD HAVE MADE	\$ / FIRE	WHAT ITS OWN FILTERS STOPPED
grind_long	24,121	11	−\$319.17	−\$29.02	ATR floor 22pt ×24,110
capitulation_long	132	10	−\$97.66	−\$9.77	ATR floor 10pt ×122
abs_veto_long	38	26	−\$83.42	−\$3.21	55s veto: absorbed ×5; 55s veto: thrust gone ×7
rgv_short	1	1	−\$17.54	−\$17.54	—

Total for the four: **−\$517.79. The bench is the rehabilitation.** Look particularly at grind_long: 24,121 raw signals on the week's tape, of which its own ATR floor stopped 24,110. The gate is not benched by the router in any meaningful sense — it is benched by its own 22-point floor, which this week's tape cleared on a small minority of minutes. The mid-week work in Part 2.5 puts a number on how often that floor is even reachable, and it is about a quarter of sessions.

5 · The treatments that FAILED this week, by name

TREATMENT	APPLIED TO	RESULT	DISPOSITION
Tighter stop (0.50× / 0.75×ATR)	abs_veto_short	−\$124.83 and −\$57.94 — both worse than as-traded	REFUTED for this gate this week
Time-of-day carve-outs (three variants)	abs_veto_short	None keeps 80% of the winners while beating the total	REFUTED — fake wins
Stand down after two consecutive stops	abs_veto_short	Cuts winners with losers; THIRD independent population to refuse it (★ REV2 — the count is settled at three and the populations are named in §2.3)	REFUTED
“It is chasing exhausted moves” (the entry diagnosis)	abs_veto_short	0 of 15 losing lots were chases. Median loser reached +11.25 pt onside first — the entries work, briefly	REFUTED on this week's tape
Blame the entry	the day rider's Thursday	The entry is its normal open-hour entry and made money on three other days; six stop widths all beat the outcome	REFUTED — it is the missing stop

Disposition — every lead this section touched

LEAD	VERDICT	FOR PARKED: WHAT WOULD REVIVE IT. FOR REFUTED: THE TEST THAT KILLED IT.
abs_veto_short — bench the machine's half of it	SHADOW	Ten cells across two stop sweeps, every one negative, worst-leave-one-day-out negative throughout: the sign survives, so this IS a verdict rather than a knife-edge. The gate is only live-positive because your hand took +\$314.74 out of it — the venue-stop figure from §2.4, which is the one to quote; §2.2's estimated-ATR version of the same split reads +\$412.93. Bench the automated slot, keep the signal in shadow, and let the 55s mirror carry the case.
abs_veto_short — fix the ENTRY via the 55s confirm	SHADOW	The mirror is the top of this week's shadow board. Promote only on the Part 2 battery, and only with the mirror's two known upward biases stated on the same page.
abs_veto_short — a stop-width or filter change	REFUTED	Named tests: two stop sweeps of five cells each (all ten negative) and the winner-retention rule (no filter keeps 80% of the winners while improving the total).
	LIVE	

LEAD	VERDICT	FOR PARKED: WHAT WOULD REVIVE IT. FOR REFUTED: THE TEST THAT KILLED IT.
The day rider — a protective stop on its own book		★ REV2 re-raced ALL SIX rider entries at SEVEN widths, not just Thursday. Closed book: as-traded −\$272.00; the picked width is now 3.0×ATR at +\$195.70, the only positive cell and the only one that keeps all three winners. 1.5×ATR — Rev1's pick — is −\$388.90 and keeps one. Cause of death for 1.5×: winners-retention, on the grid in Part 1.6 §3b. Still the week's highest-value change, at +\$468 rather than ~\$1,000.
The day rider — branch on venue_net, not entered	FIXED	One line. The 00:00Z session roll orphaned an open position and cost \$156.50 in two lots on Friday morning.
The four benched gates — un-bench any of them	REFUTED	Fired mechanically across the whole week's tape in their exact live config: −\$517.79 on 48 fires. Revisit on a measured trend week, not on this one.
grind_long's 22-point ATR floor	PARKED	It blocked 24,110 of 24,121 raw signals this week, which makes the gate un-testable rather than benched. Revive the question with the Part 2.5 floor-reality numbers: how often is the floor reachable at all, and is 22 the right number or just an untested one?

P1.6

THE DAY RIDER — the week of 17-21 August 2026

Gaz, the two books swapped places. Last week the rider made **+\$1,617.50 on five trades** while the tournament lost \$325.00; this week the tournament is a rounding error at **−\$58.50 on 36 lots** and **the rider is the loss**. It closed seven legs for **−\$189.00** in the strategy book, **−\$272.00** in cash.

But that headline hides the thing worth knowing. Split the book by who pressed the button and it separates cleanly: the **automatic rider went 3-1 for −\$115.50** across four sessions, and **every other dollar of damage came from a hand** — two manual entries on Friday, one killed by a bug sixty seconds after it was placed, the other **still open right now** because the IB gateway wedged at 19:58Z and the 20:40 hard flat never ran.

THE ONE-PARAGRAPH VERSION

Realised –\$189.00 (strategy) / –\$272.00 (cash), plus a 4-lot long marked –\$973.00 that is UNRESOLVED as I write. ENTRIES: all four automatic confirmations landed inside the validated 13:38–14:09 window; the two entries outside it were both manual and both are under water. **THRESHOLDS:** nothing to change. The margins ran +0.006 to +0.490 over the ER floor, and I probed the floors directly — raising ER to 0.20 and RT to 0.55 removes **zero** fires this week, it only makes them **later and worse**, by –\$116 and –\$156 on the two it touches. **EXITS:** the operator's claims took **91.0% and 89.3%** of everything available on the two winners he claimed and beat the machine trail by **+\$180.29**; the one mechanical TRAIL took **36.3%**; on the losing day the claims cost **–\$120.50** against just letting the clock do it. **THE WEAK DAY IS THURSDAY 08-20 and it is a FAULT, not a drought** — the tape offered **\$1,158** on the short side and the rider bought the high of the day **37 seconds after it printed**.

1 • THE MONEY — and which hand did it

BOOK	NET	LEGS	LOTS	WHAT IT IS
Rider, STRATEGY book	–\$189.00	7	13	Clean rows only. Use for every threshold and exit question below.
Rider, CASH book	–\$272.00	8	14	Adds the 08-21 BADFILL lot — a defective fill price, not a phantom trade. This is the figure that reconciles to IBKR.
Open 4-lot long (unresolved)	–\$973.00	—	4	Marked at the last print before the 21:00 halt (29,370.00). Not a result. See section 5.
— of which entered AUTOMATICALLY	–\$115.50	6	10	Four detector-confirmed sessions, 3 winners 1 loser . Essentially flat.
— of which entered BY HAND	–\$156.50	2	4	Friday's 09:11 press, killed by a bug a minute later. The 13:13 press is the still-open lot above.
Tournament, same week	–\$58.50	36	36	For scale. Last week it was –\$325.00 on 108 and the rider carried the desk; this week that is exactly inverted.

MNQ at **\$2.00 a point**, fee **\$1.50 per round trip**; every figure in this section is now net of it. ★ **REV2 — IT WAS NOT**. This section rendered at 22:32 while every other section was rebuilt between 00:48 and 01:17, so it never picked up the fee convention the rest of the report adopted (friday_v7_build.py:717): it claimed to be net and printed `pnl_usd`, which is gross. Corrected throughout — strategy $-\$169.50 \rightarrow -\189.00 , cash $-\$251.00 \rightarrow -\272.00 , automatic $-\$100.50 \rightarrow -\115.50 , Thursday $-\$1,288.50 \rightarrow -\$1,294.50$ (which is what Part 1 said all along), tournament $-\$4.50 \rightarrow -\58.50 . Two things deliberately did NOT move. The capture ratios barely shift, because an exit at the peak pays the same round trip as the exit you took — Monday is 91.0% on either ruler. And the hand-vs-machine deltas do not move at all: $+\$59.79$ for the week and $+\$180.29$ on the winners are differences between two figures that each pay the same fee, so they are fee-invariant and stand exactly as first printed. The rider went from **2 lots to 4** on 08-20 alongside the no-stop / per-lot-ladder redesign — so the week's only loss is also its first day at double size. At the old 2-lot size the same Thursday is about **$-\$647$** net.

2 • DID WE ENTER OK?

Every row below is the **live detector's own read at the tick it fired on**, lifted from the service journal — not a replay. “Travelled” is how far the move had already gone from the 13:30 open reference the detector itself measures against, and “of extension” is that distance as a share of how far the move eventually ran in the same direction. MFE and MAE are measured from entry, on `capture.db` ticks, to the 20:40 flat.

DAY	CONFIR MED	+MIN	SIDE	ENTRY	13:30 REF	TRAVEL LED	OF EXTEN SION	MFE AFTER	MAE AFTER	WINDO W
Mon 08-17	13:38:07	+8.1	SHORT	30,211.50	30,253.00	41.50 pt	21.5%	+151.50	-75.25	IN
Tue 08-18	13:54:07	+24.1	SHORT	29,642.00	29,713.50	71.50 pt	35.8%	+128.00	-64.50	IN
Wed 08-19	13:39:07	+9.1	SHORT	29,607.50	29,672.75	65.25 pt	22.0%	+231.75	-59.75	IN
Thu 08-20	13:38:07	+8.1	LONG	29,447.75	29,347.25	100.50 pt	81.7%	+19.25	-245.25	IN
Fri 08-21	09:11:06	-258.9	LONG	29,428.25	—	—	—	+1.00	-10.50	MANUAL — OUT
Fri 08-21	13:13:06	-16.9	LONG	29,491.625	—	—	—	-3.12	-271.38	MANUAL — OUT

The automatic entries are fine and there is nothing to fix in them. All four confirmed inside the validated **13:38-14:09** band — three of them within nine minutes of the open, and Tuesday's 13:54 is comfortably inside. On the three shorts the rider boarded with **21.5%, 35.8% and 22.0%** of the eventual move

consumed and left itself **128 to 232 points** of room. That is exactly what this detector is supposed to do.

THE ONE BAD ENTRY, AND IT IS BAD IN A SPECIFIC WAY

Thursday's long boarded with **81.7% of the up-move already spent** — and the session high did not arrive after the entry, it arrived **at 13:37:30, thirty-seven seconds before it**. The rider bought the high of the day. What remained was **19.25 points** of favourable excursion against **245.25 points** of adverse. That single row is the whole week's loss.

Both entries outside the window were the operator's own manual presses, and manual entry is deliberately live across the whole CME session — so “outside the window” is a description here, not a rule violation. It is still worth stating plainly: **the four entries the detector chose are 3-1; of the two a hand chose, one lost and one is still open and red.**

3 · ARE THE THRESHOLDS OK?

Live floors are **ER ≥ 0.15** and **RT ≥ 0.45**. These are the numbers the detector actually read at the confirming tick, with the margin it cleared by.

DAY	ER READ	MARGIN	RT READ	MARGIN	ATR	RESULT
Mon 08-17	0.290	+0.140	0.565	+0.115	27.6	+\$453.00
Tue 08-18	0.156	+0.006	0.520	+0.070	30.9	+\$185.00
Wed 08-19	0.326	+0.176	0.480	+0.030	35.3	+\$541.00
Thu 08-20	0.640	+0.490	0.748	+0.298	34.2	−\$1,294.50

Both scrapers were winners, and the runaway best read was the loser. Tuesday cleared the efficiency floor by **six thousandths** and paid **+\$185**; Wednesday cleared the round-trip floor by **0.030** and paid **+\$541**. Thursday cleared both by the widest margin of the week — ER 0.640, RT 0.748, the most confident read the detector produced — and lost **−\$1,294.50**. On four fires that is an anecdote, not a relationship. It is logged, not acted on.

THE FLOOR PROBE — RAISING A THRESHOLD DOES NOT DELETE A FIRE, IT DELAYS IT

Rather than guess, I re-ran the live detector minute-by-minute over this week's tape with the floors moved. This is the result that matters:

FLOORS	FIRES	08-17	08-18	08-19	08-20	08-21*
LIVE — ER .15 / RT .45	5	+9m 30,187.50	+23m 29,642.75	+9m 29,611.50	+9m 29,436.50	+9m 29,340.50
ER .20 / RT .45	5	+9m 30,187.50	+26m 29,613.75	+9m 29,611.50	+9m 29,436.50	+9m 29,340.50
ER .15 / RT .55	5	+9m 30,187.50	+26m 29,613.75	+12m 29,572.50	+9m 29,436.50	+9m 29,340.50
ER .20 / RT .55	5	+9m 30,187.50	+26m 29,613.75	+12m 29,572.50	+9m 29,436.50	+9m 29,340.50

Not one fire is removed by any of it. Raising both floors as far as ER 0.20 and RT 0.55 leaves Monday, Thursday and Friday untouched to the tick, and on the two days it does bite it simply makes the rider board later, further into a move that has already run:

- **08-18**, ER 0.15→0.20: confirmation slides +23m→+26m, entry 29,642.75→29,613.75. Held to the 20:40 clock that is **\$227 → \$111**, a cost of **−\$116**, and the MFE it can play for drops by the same \$116.
- **08-19**, RT 0.45→0.55: +9m→+12m, entry 29,611.50→29,572.50. **\$351 → \$195**, a cost of **−\$156**.

VERDICT ON THE THRESHOLDS: NOTHING TO CHANGE THIS WEEK

Margins this week: ER **+0.006 / +0.070 / +0.140 / +0.176 / +0.298 / +0.490** across the two metrics. One genuine scraper (ER +0.006) and it was a **winner**. Tightening buys nothing here and costs **−\$272** in delayed entries on a week where it removes no trade at all. That is a mechanical property, not a fluke of five sessions: a higher floor can only be met later, and later means more of the move is already gone. **The margin distribution goes in the book and accumulates. No change recommended on four fires.**

* 08-21 is a replay, not a live fire — the automatic path never ran on Friday because a hand-opened position already occupied the book at 13:30. What the detector would have read is in section 5, and it is not flattering to the hand.

3b · ★ REV2 — THE PROTECTIVE STOP, RACED ON EVERY DAY AND NOT JUST THE LOSING ONE

WHY THIS SECTION EXISTS

Rev1 called a 1.5×ATR rider stop “the highest-value change on this week’s card” and priced it at about \$1,000 — **on Thursday alone**, the one day the rider lost money. That is exactly the winners-retention failure this report applies rigorously to abs_veto_short in Part 1.5 §2.3: a filter that improves the P&L by dropping the trades that WON is a fake win. Rev1’s own tables contained the refutation and nobody joined them up — Monday’s in-life worst excursion is −75.25pt (Part 1 §6) against a 41.4pt stop (§3’s ATR of 27.6). So here is every rider entry of the week, at every width.

One stop for the whole position; the first tick at or through it closes every leg still open, and legs already closed keep their real exit. Raced on capture.db MNQ ticks from each entry, sized off the ATR the live detector itself read at the confirming tick (§3), net of \$1.50 per round trip per lot. **Bold = better than what actually happened.**

RIDER POSITI ON	ATR	WORST EXCUR SION	AS TRADE D	0.5×	0.75×	1×	1.25×	1.5×	2×	3×
08-17 SHORT 2L	27.6	-75.25	\$453.00	−\$58.20	−\$85.80	−\$113.40	−\$141.00	−\$168.60	−\$223.80	\$453.00
08-18 SHORT 2L	30.9	-64.50	\$185.00	−\$64.80	−\$95.70	−\$126.60	−\$157.50	−\$188.40	−\$250.20	\$185.00
08-19 SHORT 2L	35.3	-28.75	\$541.00	−\$73.60	−\$108.90	\$541.00	\$541.00	\$541.00	\$541.00	\$541.00
	34.2	-245.25	−\$1,294.50	−\$142.80	−\$211.20	−\$279.60	−\$348.00	−\$416.40	−\$553.20	−\$826.80
CLOSED BOOK — the verdict- bearing rows	—	—	−\$272.00	−\$495.90	−\$658.10	−\$135.10	−\$262.00	−\$388.90	−\$642.70	\$195.70
... plus the still-open carry (PROVISIO NAL)	—	—	−\$1,251.00	−\$624.70	−\$848.30	−\$386.70	−\$575.00	−\$763.30	−\$1,139.90	−\$547.10
Winners kept	—	—	3 / 3	0 / 3	0 / 3	1 / 3	1 / 3	1 / 3	1 / 3	3 / 3
Winner dollars kept	—	—	\$1,179.00	−\$196.60	−\$290.40	\$301.00	\$242.50	\$184.00	\$67.00	\$1,179.00

RIDER POSITI ON	ATR	WORST EXCUR SION	AS TRADE D	0.5×	0.75×	1×	1.25×	1.5×	2×	3×
08-20 LONG 4L										
08-21 LONG 4L	30.7	-10.50	-\$156.50	-\$156.50	-\$156.50	-\$156.50	-\$156.50	-\$156.50	-\$156.50	-\$156.50
08-21 LONG 4L (still open)	30.7	-271.38	-\$979.00	-\$128.80	-\$190.20	-\$251.60	-\$313.00	-\$374.40	-\$497.20	-\$742.80
CLOSED BOOK — the verdict- bearing rows	—	—	-\$272.00	-\$495.90	-\$658.10	-\$135.10	-\$262.00	-\$388.90	-\$642.70	\$195.70
... plus the still-open carry (PROVISIO NAL)	—	—	-\$1,251.00	-\$624.70	-\$848.30	-\$386.70	-\$575.00	-\$763.30	-\$1,139.90	-\$547.10
Winners kept	—	—	3 / 3	0 / 3	0 / 3	1 / 3	1 / 3	1 / 3	1 / 3	3 / 3
Winner dollars kept	—	—	\$1,179.00	-\$196.60	-\$290.40	\$301.00	\$242.50	\$184.00	\$67.00	\$1,179.00

READ THE CLOSED-BOOK ROW, AND READ WINNERS-KEPT BESIDE IT

1.5×ATR is the fourth-best of seven widths and it LOSES money on the closed book. It takes the rider from −\$272.00 to −\$388.90 — \$116.90 worse than doing nothing — because 41.4pt is inside Monday's −75.25pt excursion and 46.4pt is inside Tuesday's −64.50pt, so it converts a +\$453 winner into −\$168.60 and a +\$185 winner into −\$188.40. It keeps **one of the three winners**.

3.0×ATR is the pick, and it is the only width that is both. On the closed book it is **+\$195.70** — the only positive cell in the row — and it is the only width that keeps **all three** winners, because 82.8 / 92.7 / 105.9pt all sit outside Monday, Tuesday and Wednesday's worst excursions. It still caps Thursday at −\$826.80 against −\$1,294.50. Against as-traded that is **+\$467.70 on the closed book**, and **+\$703.90** once the open carry is marked in.

★ AND THE THING NEITHER REVISION EXPECTED — ANY STOP CLOSES THE WEEKEND CARRY

Every one of the seven widths stops Friday's 13:13Z entry, and all of them do it within **nineteen minutes**: 0.5× at 13:13:06, 1.0× at 13:16:54, 1.5× at 13:20:49, 3.0× at 13:31:55 for −\$742.80. The gateway did not wedge until

19:57Z. So the four lots that are the headline of this report, the ones that need a Sunday-night gateway check and a weekend-flatten timer, would not exist under any stop this section tested. The carry is a missing-stop incident that a wedge then made unfixable — not a wedge incident.

One honest artefact: at $0.5 \times \text{ATR}$ the stop (29,476.28) sits inside the position's own fill spread — the fills were 2 @ 29,477.00 and 2 @ 29,506.25, a 29.25pt spread around a 29,491.625 VWAP — so it triggers in the entry second. That is a real property of stopping a badly-filled position off its VWAP, not a simulator bug, and it is one more reason not to pick the tightest cell.

WHAT THIS DOES NOT LICENSE

Six rider entries is not a sample. Three winners decide the winners-kept column and one of them (Wednesday) survives every width from $1.0 \times$ up, so the ranking between $1.0 \times$ and $3.0 \times$ rests on Monday and Tuesday alone. The closed-book row is the verdict-bearing one precisely because every leg in it is resolved; the row beneath it contains a position that is **still open**, and a marked-to-market leg can only be scored at its current excursion, never at its eventual exit — so treat that row as provisional and re-cut it the moment the carry closes. The instruction that survives all of this is not “ $3.0 \times$ ”, it is “**have a stop, and never pick its width from the losing day**”.

4 • WHICH EXIT CAPTURED IT?

MFE-available is measured on ticks **inside each leg's own life** — entry to its own exit, never a window we were flat for. The clock column is what a hold to 20:40 would have offered, for contrast.

DAY	EXIT	LOTS	HELD	TAKEN	MFE IN LIFE	CAPTURE	VS 20:40 HOLD	\$
08-17	MANUAL_CLAIM	2	4h26m	+125.25 pt	\$498.00	91.0%	75.2%	+\$453.00
08-18	TRAIL	2	1h33m	+128.00 pt	\$509.00	36.3%	36.3%	+\$185.00
08-19	MANUAL_CLAIM	2	11m	+152.25 pt	\$606.00	89.3%	58.7%	+\$541.00
08-20	MANUAL_CLAIM	1	5h56m	−146.00 pt	\$38.50	never green	—	−\$293.50
08-20	MANUAL_CLAIM	1	6h05m	−133.25 pt	\$38.50	never green	—	−\$268.00
08-20	MANUAL_CLAIM	2	6h40m	−181.00 pt	\$77.00	never green	—	−\$727.00
08-21	CLOCK_FLAT_BUG	3	1m	−10.75 pt	\$6.00	n/a	—	−\$69.00
08-21	CLOCK_FLAT_BUG	1	1m	−40.00 pt	\$2.00	n/a	—	−\$81.50

GROUPED BY EXIT REASON

EXIT REASON	LEGS	LOTS	TAKEN	MFE AVAILABLE	CAPTURE
MANUAL_CLAIM — on the winners	2	4	+\$994.00	\$1,104.00	90.0%
TRAIL (armed chandelier)	1	2	+\$185.00	\$509.00	36.3%
MANUAL_CLAIM — on the loser	3	4	−\$1,294.50	\$148.00	never green
CLOCK_FLAT_BUG	2	4	−\$156.50	\$2.00	a defect, not an exit
CLOCK_FLAT / operator claim on the open lot	—	4	unresolved	—	see section 5

THE OPERATOR'S HAND, SCORED AGAINST THE MACHINE

The desk's own `claim_audit.json` (regenerated 21:31Z today) prices each claim against what the configured mechanical exit would have taken on the same entry. It is a ceiling for the machine — exact levels, no slippage — so it flatters the machine side, and the hand still wins:

LEG	HAND	MACHINE	DELTA	MACHINE'S EXIT
08-17 SHORT 2	+\$453.00	+\$303.14	+\$149.86	TRAIL @ 30,134.2 (armed)
08-19 SHORT 2	+\$541.00	+\$510.57	+\$30.43	TRAIL @ 29,478.4 (armed)
08-20 LONG 1	−\$295.00	−\$293.50	−\$1.50	CLOCK_FLAT @ 29,302.5 (never armed)
08-20 LONG 1	−\$269.50	−\$293.50	+\$24.00	CLOCK_FLAT @ 29,302.5 (never armed)
08-20 LONG 2	−\$730.00	−\$587.00	−\$143.00	CLOCK_FLAT @ 29,302.5 (never armed)
Week	−\$300.50	−\$360.29	+\$59.79	5 claims

What the exits did this week, plainly. The claim is a first-class exit and it earned that status again: **90.0% capture** on the two winners it was used on, and **+\$180.29** ahead of the machine on those same two. It is on the loser that it went the other way — three claims spread over 45 minutes, six hours into a position that was never more than 19 points green, each one worse than the 20:40 clock would have been, costing **−\$120.50** in aggregate. That is the honest shape of a discretionary exit: it beats the machine when there is something to take and it bleeds when there is not.

The single mechanical **TRAIL** took **36.3%**. That is a low number in isolation and I am not going to draw a conclusion from one leg — and per the standing instruction, the exit battery is settled and is not re-run here. What is worth

flagging is that on the four losing lots the trail **never armed at all**, which is by design and is exactly what “no stop” means in practice: on a day that never goes green there is no mechanical exit before 20:40.

5 · THE WEAK DAY — Thursday 08-20, and it is a fault

The test for “bad strategy” against “bad tape” is whether the tape had money in it. It did.

DAY	RANGE	NET FROM OPEN	LONG CEILING	SHORT CEILING	RIDER TOOK
08-17		226.8	−188.0	\$135	\$772 +\$453.00
08-18		256.0	−128.2	\$226	\$798 +\$185.00
08-19		366.2	−149.8	\$277	\$1,188 +\$541.00
08-20		267.8	−44.8	\$984	\$1,158 −\$1,294.50
08-21		268.2	−44.8	\$506	\$1,640 −\$156.50 + open

Ceilings and “net” are measured from the same 13:30 reference the detector uses, to the session extreme, at the size the rider was actually carrying that day (2 lots Mon–Wed, 4 lots Thu–Fri). They are oracles — nobody catches an extreme — and they are here only to answer “was there anything to catch”.

WHY 08-20 IS A FAULT AND NOT A QUIET TAPE

Thursday had a **267.8-point range** and **\$1,158** sitting on the short side. This is nothing like 08-10, which took \$39 of a \$147 opportunity and was simply a dead session. The rider went **long**, at **13:38:07**, **thirty-seven seconds after the session high printed at 13:37:30**, with **81.7%** of the up-move already behind it. It then held **245 points** of adverse excursion with no stop, and the operator liquidated in three pieces between 19:33 and 20:17. **The opportunity was there and we took the wrong end of it. That is an entry fault.**

The mitigating and the aggravating fact, both stated. Mitigating: this was the **first session at 4 lots**, doubled that morning alongside the naked / per-lot-ladder redesign — the same day at 2 lots is roughly **−\$647** net, and the desk-level damage is a size decision as much as a signal one. Aggravating: it is the only day

of the week the detector leaned **UP**, every other session drifted **down** from the open, and it leaned up with the strongest numbers it printed all week.

FRIDAY 08-21 — A SEPARATE FAILURE, AND IT IS STILL RUNNING

Friday does not compete for “weak day” because it is not a strategy result at all. Three things happened, in order:

1. **09:11:06Z** — a manual BUY of 4 lots at 29,428.25, hours outside the automatic window. The next tick hard-flatted it at **09:12** because the flatten branch keyed on `mod < OPEN_UTC_MIN`, true all morning; the tick then died on an `UnboundLocalError: px` after the orders had reached the venue. Booked `CLOCK_FLAT_BUG`, **-\$156.50**, one lot filled 30 points outside the book and is flagged `BADFILL`. A regression test now pins the management window (`tests/test_day_rider_hold_window.py`).
2. **13:13:07Z** — a second manual BUY of 4 lots at 29,491.625, seventeen minutes before the open. Because the book was then occupied, **the automatic detector never ran on Friday at all.**
3. **19:58:20Z** — the IB gateway wedged. Last good venue read **19:56:37Z** (“position 4.0 managed”); every tick since is `TimeoutError`, the watchdog has logged **73 consecutive blind cycles** as of 22:20Z and is still counting, and there are **48 CLOSE-WAIT sockets on :4002** — the documented wedge fingerprint. **The 20:40 hard flat never executed.**

UNRESOLVED 4 LOTS LONG, NOT CONFIRMED FLAT

As of **22:20Z** the rider's book claims **4 lots long from 29,491.625** and `desk_reconcile` reads venue ? = tournament +0 + rider +4 → unaccounted ?. The position marks **-\$887.00** at the 20:40 price it should have been flattened at and **-\$973.00** at the last print before the 21:00 halt. **Neither number is a result and I am not booking one.** The standing rule is NEVER HOLD OVERNIGHT; the book is currently in breach of it and the venue cannot be read to confirm either way. This is the first thing to check when the gateway comes back, ahead of anything else in this report.

WHAT THE DETECTOR WOULD HAVE DONE ON FRIDAY

Worth knowing, because a hand took the opposite side. Replaying the live detector on Friday's tape it confirms at **+9m, DOWN, at 29,340.50**, on the strongest efficiency read of the week (**ER 0.854 / RT 0.553**). That short would have gone **+120.25 points (+\$962)** in its favour and **−148.00 (−\$1,184)** against, and held to the 20:40 clock it finishes **−\$326.00** — the trail never arming. So the automatic rider had the direction right and would still have lost money holding; it is not a clean “the machine would have saved us”. But it is **−\$326 instead of −\$973-and-counting**, and it would have been flat before the gateway wedged.

6 · ONE DEFECT FOUND WHILE DOING THIS

A MANUAL ENTRY OUTSIDE 13:30-21:00 SILENTLY DISABLES THE ATR TRAIL

`do_manual_entry()` writes `arm_atr = drift_read(...).atr`. Outside 13:30-21:00 `drift.read()` returns early with `detail='before the 13:30 UTC open'` and `atr = 0.0`. I confirmed that directly: both of Friday's manual entries stored `arm_atr: 0.0`, and `trail_level()` therefore falls back to the fixed **ARM_PT = 150 pt** instead of **4×ATR**. The live ticks say it in as many words — “riding · trail not armed (need +150 [fixed], at −97)”. At Friday's 30.7-point ATR, **4×ATR is 122.8 pt**, so the fallback is **27 points harder to arm**. It cost **\$0 this week** — neither position ever went green — but on a hand-opened winner it silently removes the only mechanical exit the rider has. Three independent confirmations: the code path, the state file, and the rider's own tick note.

7 · THE VERDICT

QUESTION	ANSWER	ACTION
Did we enter OK?	Yes, on all four automatic entries — every one inside 13:38–14:09, three within nine minutes of the open, boarding 21.5–35.8% into the move on	NO CHANGE

QUESTION	ANSWER	ACTION
	the three winners. The one bad entry bought the day's high 37 seconds after it printed.	
Are the thresholds OK?	Yes. Margins +0.006 to +0.490; the one genuine scraper was a winner. Raising ER to 0.20 and RT to 0.55 removes zero fires and costs −\$272 in later, worse entries.	NO CHANGE
Which exit captured it?	The operator's claim, decisively, on the two winners it was used on — 90.0% capture and +\$180.29 over the machine. It cost −\$120.50 on the loser. The one TRAIL took 36.3%.	NO CHANGE
Name the weak day	Thursday 08-20, and it is a fault. A 267.8-point range with \$1,158 on the short side; the rider went long at 81.7% of the up-move, 37 seconds past the high, at double size on its first day at double size.	ENTRY
Anything else	4 lots long are unresolved into the weekend because the gateway wedged at 19:58Z and the 20:40 flat never ran. Plus: a pre-open manual entry stores <code>arm_atr = 0.0</code> and loses the ATR trail.	OPEN

So: nothing to change in the detector or the thresholds this week. The automatic rider is roughly flat at $-\$115.50$ on 3-1 and its entries are behaving exactly as designed. Every dollar of real damage this week came from either a hand-timed entry or an infrastructure failure, and both of those are outside the strategy the operator asked me to tune. The margin distribution is logged and will accumulate; on four fires I am recommending nothing.

Sources, all recomputed for this section: live detector reads from the gazbot7-day-rider journal (6,953 parsed ticks, 17-21 Aug); trades from data/gazbot7.db; tape from data/capture.db ticks, whose 5-trading-day rolling window is exactly this week; floor probe by re-running gazbot7.drift.compute minute-by-minute over the same 5s bars the live service reads; hand-vs-machine from data/claim_audit.json; venue state from day_rider_state.json, day_rider_watchdog.log and desk_reconcile_state.json. MNQ \$2.00/point, \$1.50 per round trip.

P2 THE SHADOW DESK — and the promotion that did not survive the ruler

Garrath — the shadow book ran 38 families with eight or more fires this week, and if you rank them the way the board is usually ranked, the answer is easy and it is wrong. **Once the two filters this desk has already learned to apply are applied, the family sitting second on**

the raw board falls to eleventh, the promotion case for it collapses from +\$332 to +\$23.50, and its 50-second sibling goes negative. The candidates that survive are different ones, they are newer, and the honest recommendation this week is smaller than last week's.

1 • The ruler, stated before any number

Score `real_pnl`, never `ceiling_pnl`. The ceiling is what the trade would have made with a perfect exit; real is what the tape would have paid. On this week's board they disagree by hundreds of dollars per family and the ceiling is always the friendlier of the two.

Void the Asia block. This is the one that reorders everything. The desk carries a permanent bench on 00:00–07:00 UTC entries, and that bench sits downstream of where a shadow arm records its fire — so an arm happily books money in a window where the live desk has filled **nothing since 5 August**. It is not a small correction:

FAMILY	RAW NET	BOOKED IN 00:00-07:00Z	FIRES THERE	WHAT IS LEFT
capit_loose	−\$2,457.50	−\$835.00	84 of 189	−\$1,622.50
bank20_grindA_ctl	−\$1,193.50	−\$520.50	53 of 128	−\$673.00
cx_grindA_live	−\$1,099.50	−\$520.50	53 of 123	−\$579.00
bank20_grindA	−\$1,212.50	−\$486.50	52 of 134	−\$726.00
cx_clip_brk_live	−\$1,166.50	−\$486.50	52 of 125	−\$680.00
cx_clip_brk_standdown	−\$1,166.50	−\$486.50	52 of 125	−\$680.00
ALL 38 FAMILIES	−\$13,780.50	−\$3,631.00		−\$10,149.50

Across all 38 families, **−\$3,631.00 of the week's −\$13,780.50 was booked in a window the desk cannot trade**. Every number from here down is tradable-hours only.

And dedupe before quoting any delta — see §5, where 7 of the 38 arms turn out to be the same rows under different names.

2 • The board, raw and then honest

SHADOW FAMILY	FIRES	RAW REAL_PNL	TRADABLE FIRES	TRADABLE REAL_PNL	WIN %	\$ / FIRE
odr_c10_s30	18	\$516.00	18	\$516.00	50%	\$28.67
odr_c10_s30_g	18	\$326.00	18	\$326.00	50%	\$18.11
lad_absS_B_15	27	\$255.50	20	\$244.00	45%	\$12.20

SHADOW FAMILY	FIRES	RAW REAL_PNL	TRADABLE FIRES	TRADABLE REAL_PNL	WIN %	\$ / FIRE
lad_absS_A_10	29	\$242.00	21	\$203.50	52%	\$9.69
sw_absS_B_k10	27	\$197.50	20	\$186.00	40%	\$9.30
sw_absS_A_k10	28	\$221.50	20	\$183.00	50%	\$9.15
rg_long_fast_v	36	\$329.50	31	\$118.50	29%	\$3.82
sw_absS_B_k20	27	\$100.50	20	\$103.50	50%	\$5.17
sw_absS_A_k20	28	\$9.50	20	\$61.50	60%	\$3.08
thrust_short_abs_veto55	26	\$114.50	19	\$56.50	42%	\$2.97
abs_veto_55s	46	\$332.00	34	\$23.50	35%	\$0.69
abs_veto_50s	47	\$299.00	35	−\$9.50	34%	−\$0.27
CL-capitulation_long	11	−\$22.00	7	−\$12.50	29%	−\$1.79
rung_grindA_med_05	8	−\$46.00	6	−\$55.50	33%	−\$9.25
... 14 families omitted, all between these two blocks ...						
cx_grindA_clip	125	−\$1,166.50	73	−\$680.00	23%	−\$9.32
bank20_grindA	134	−\$1,212.50	82	−\$726.00	27%	−\$8.85
capit_live_mirror	77	−\$1,316.00	42	−\$857.50	7%	−\$20.42
rider_w5	36	−\$1,136.50	36	−\$1,136.50	17%	−\$31.57
capit_loose	189	−\$2,457.50	105	−\$1,622.50	20%	−\$15.45

Read the two money columns side by side. abs_veto_55s is second on the raw board and eleventh on the honest one. abs_veto_50s crosses from \$299.00 to −\$9.50 — from a promotion candidate to a losing arm — purely by removing trades that could never have been filled. The arms that barely move are the ones worth looking at, because their money was made in hours the desk was actually open.

3 · The promotion battery — the seven families that clear the ruler, plus the one Rev1 left out

FAMILY (TRADABLE HOURS ONLY)	N	NET	WIN %	BEST FIRE STRIPPE D	BEST 3 STRIPPE D	WORST LEAVE-ONE-DAY-OUT	GREEN DAYS	ALL-TIME, TRADABLE HOURS
odr_c10_s30	18	\$516.00	50%	\$208.50	−\$255.50	\$13.50	3 / 5	37 · \$499.00
odr_c10_s30_g	18	\$326.00	50%	\$18.50	−\$340.00	−\$176.50	3 / 5	22 · \$552.00
lad_absS_B_15	20	\$244.00	45%	\$125.00	−\$26.50	\$92.50	4 / 5	20 · \$244.00
	21	\$203.50	52%	\$118.50	−\$15.00	\$64.00	4 / 5	21 · \$203.50

FAMILY (TRADABLE HOURS ONLY)	N	NET	WIN %	BEST FIRE STRIPPED	BEST 3 STRIPPED	WORST LEAVE- ONE-DAY- OUT	GREEN DAYS	ALL- TIME, TRADABLE HOURS
lad_absS_A_10								
sw_absS_B_k10	20	\$186.00	40%	−\$10.00	−\$144.00	\$85.50	4 / 5	52 · \$1,232.50
sw_absS_A_k10	20	\$183.00	50%	\$64.00	−\$60.00	\$65.00	4 / 5	53 · \$636.50
rg_long_fast_v ★ REV2	31	\$118.50	29%	−\$172.50	−\$459.00	−\$186.50	2 / 5	211 · − \$1,020.50
abs_veto_55s	34	\$23.50	35%	−\$133.50	−\$280.50	−\$64.50	2 / 5	263 · \$3,078.00

★ **REV2 — two fixes to this table.** The heading said “the four that survive the ruler” over seven rows; it now says the real count. And rg_long_fast_v was **silently absent** despite sitting 7th on \$2's own tradable board at +\$118.50 on 31 fires — more n and more money than sw_absS_A_k20 or thrust_short_absveto55, and more n than either ladder arm that WAS in the battery, with no exclusion reason given anywhere. It is in the table now, and it has a cause of death rather than an omission: **strip its best three fires and it is −\$459.00, its worst leave-one-day-out is −\$186.50, and 2 of 5 days are green** — against an all-time tradable-hours record of **−\$1,020.50 on 211 fires at 30.3% win**, whose own strip-best-3 is −\$2,084.00 and worst fold −\$1,528.00. This week is its best week against a deeply negative body of evidence, which is what a lucky week looks like. It is also a long-side rgv arm rather than an abs-veto or open-rider one, so it is not a candidate for either live slot under discussion — but it is excluded on the battery, not on the category.

This is the part that decides Monday, so read the last four columns rather than the second one. odr_c10_s30 has the biggest number on the board and the worst robustness on it: stripping its best three fires takes −\$255.50, and two of its five days are red. That is one or two trades carrying a family, on n=18. The lad_absS_* and sw_absS_* arms — ladder and stop-width variants on the abs-veto short signal — are smaller and steadier: four of five days green, leave-one-day-out still positive, and the k=1.0 stop-width arms carry real all-time n behind them (\$1,232.50 on sw_absS_B_k10 across all its tradable-hour fires).

★ THE CASE FOR ABS_VETO_55S IS AN ALL-TIME CASE, NOT A THIS-WEEK CASE — SAY IT THAT WAY

This week, tradable hours: **\$23.50 on 34 fires**, 35% win, two of five days green, and stripping the single best fire turns it −\$133.50. On this week alone it is nothing.

All-time, tradable hours: **\$3,078.00 on 263 fires**. That is a real body of evidence and it is why the family is still the lead candidate at all.

So the honest sentence is: the 55-second confirm has a strong multi-week record and a flat week. That is a perfectly good reason to keep incubating it and a bad reason to promote it on Monday off this report's numbers. Anyone quoting ‘+\$332 this week’ is quoting the Asia block.

4 • Two-sided, or short only?

FAMILY	SIDE	FIRES	NET	\$ / FIRE	WIN %
abs_veto_55s	LONG	14	−\$109.50	−\$7.82	29%
abs_veto_55s	SHORT	20	\$133.00	\$6.65	40%
abs_veto_50s	LONG	14	−\$109.50	−\$7.82	29%
abs_veto_50s	SHORT	21	\$100.00	\$4.76	38%
abs_veto_60s	LONG	4	−\$157.00	−\$39.25	0%
abs_veto_60s	SHORT	3	−\$62.00	−\$20.67	33%
thrust_loose	LONG	20	−\$189.50	−\$9.47	25%
thrust_loose	SHORT	42	−\$89.50	−\$2.13	33%

The long side of the abs-veto family is small and positive, the short side is bigger and positive, and neither is decisive on this week's n. The standing principle still holds — you tune a SIDE, you never relegate a DIRECTION — and nothing here argues for dropping either. It also does not argue for adding the long side to a live slot: 14 tradable-hour long fires is not a sample.

4b • ★ REV2 — SO WHAT IS THE SURVIVING CASE FOR THE ENTRY TREATMENT?

THE CIRCLE THIS CLOSES

Part 1.5 concludes that abs_veto_short's treatment is the 55-second confirm and hands the case to Part 2. Part 2 §3 withdraws the promotion and hands it back to “the all-time record”. Neither states the record of **the thing**

that would actually replace the live gate — the SHORT side alone. The live slot is short-only. Its long sibling is a different gate. Scoring a short-side replacement on a two-sided family mixes in 126 fires that would never be taken.

POPULATION (TRADABLE HOURS, REAL_PNL ONLY)	N	NET	\$ / FIRE	WIN %	STRIP BEST 1	STRIP BEST 3	WORST LEAVE- ONE-DAY- OUT	GREEN DAYS
abs_veto_55s — all time, both sides	263	\$3,078.00	\$11.70	46.4%	\$2,888.00	\$2,566.50	\$2,361.00	19 / 30
abs_veto_55s SHORT — all time	137	\$2,207.50	\$16.11	48.9%	\$2,017.50	\$1,696.00	\$1,898.50	21 / 25
abs_veto_55s LONG — all time	126	\$870.50	\$6.91	43.7%	\$752.50	\$534.50	\$423.50	13 / 29
abs_veto_55s SHORT — this week	20	\$133.00	\$6.65	40.0%	−\$24.00	−\$170.50	\$70.00	4 / 5
abs_veto_55s — this week, both sides	34	\$23.50	\$0.69	35.3%	−\$133.50	−\$280.50	−\$64.50	2 / 5
abs_veto_50s SHORT — all time	148	\$1,997.50	\$13.50	46.6%	\$1,807.50	\$1,486.00	\$1,651.50	21 / 25

The surviving case is the short side, and it is the strongest single result in this report. abs_veto_55s SHORT over all tradable-hour fires is **+\$2,207.50 on 137 fires** at \$16.11 a fire and 48.9% win, **21 of 25 days green** — and it stays positive through every test the desk owns: strip the best fire (+\$2,017.50), strip the best three (+\$1,696.00), and the worst leave-one-day-out fold (+\$1,898.50). Nothing else on the board does that.

And it changes what §3 concluded. Part 2 sets the revival condition as “a positive tradable-hours week with a positive leave-one-day-out”. On the short side this week already meets it: **+\$133.00 on 20 fires, 4 of 5 days green, worst leave-one-day-out +\$70.00**. What fails is the pooled family, and it fails because of the LONG side (−\$109.50 on 14 fires this week), which is not the side under discussion.

THE HONEST CAVEATS, BEFORE ANYONE PROMOTES ANYTHING ON MONDAY

One: this week's short-side 20 fires do not carry the case on their own — strip the best one and they are $-\$24.00$, strip three and $-\$170.50$. The case is the 137-fire all-time record; the week merely fails to contradict it. **Two:** a shadow mirror is still an upper bound. It does not model the live gate's default-veto on thin tape so it over-fires, and it enters on the 1-minute bar open where the live decider quotes the tick it fired on — both biases run in the mirror's favour. **Three:** the LONG side is not refuted either. All-time it is $+\$870.50$ on 126 fires, and it survives strip-best-3 ($+\$534.50$) and leave-one-day-out ($+\$423.50$) — on only 13 of 29 days green. "Tune a SIDE, never relegate a DIRECTION" still holds; what this table says is that the two sides have very different records and should stop being scored as one number.

So the recommendation is unchanged and its reason is now stated: keep incubating, and the specific thing to watch is the SHORT arm's tradable-hours week and its leave-one-day-out — both of which were positive this week for the first time. Two consecutive weeks of that, on the short side, re-opens the promotion question with a real case behind it.

5 • ★ 7 of the 38 shadow arms are the same rows under different names

ARMS THAT ARE THE SAME ROWS	FIRES EACH	WHAT THE PAIR WAS SUPPOSED TO MEASURE	THE DELTA IT REPORTS
<code>capit_flip_live == capit_flip_t90</code>	18	a 90-second flip window against the live one	exactly \$0.00, on every field, by construction
<code>cx_clip_brk_live == cx_clip_brk_standdown == cx_grindA_clip</code>	125	a clip-on-break exit, its stand-down control, and the grind-A clip arm — three supposedly different exits	exactly \$0.00, on every field, by construction
<code>odr_c5_s30 == odr_c5_s30_g</code>	19	the 'g' variant of the open-rider claim ladder against the plain one	exactly \$0.00, on every field, by construction

Every one of these pairs was created to measure a difference, and every one of them reports a perfect null — not because the configurations agree, but because **only one configuration is running and its trades are being written twice**. A broken arm and a genuinely neutral change are indistinguishable in the board's summary view, and the summary view is what a report reads. **Diff the rows**

before quoting any shadow A/B delta; the check is four lines of SQL and it is now in scripts/fv7_p2_shadow.py. Note especially cx_clip_brk_live == cx_clip_brk_standdown: that pair exists specifically to test whether standing an exit down changes anything, and it has been answering “no” from a broken arm.

6 · The two grind-exit watches — and why they say nothing this week

GRIND EXIT WATCH	N	LIVE / BASELINE	SHADOW / PARTIAL	DELTA	NEWEST TRADE IN IT
Two-ratchet runner-clip watch	26	\$29.50	−\$265.00	−\$294.50	2026-08-05
2R-partial smoothness	40	−\$1,382.40	−\$1,074.90	\$307.50	2026-08-05

The newest trade in either watch is 2026-08-05. Grind has not fired a live lot since then, because its own 22-point ATR floor blocked 24,110 of 24,121 raw signals on this week's tape (Part 1.5 §4). So both watches are reporting the same fortnight-old window they reported last Friday, and neither has anything to add.

Two things are still worth saying about them. The two-ratchet has now watched 26 trades containing 3 genuine runners and has clipped **0** of them — but its own net is −\$265.00 against the live pure-6.0's \$29.50, so the “zero clips means reconsider it” test cannot fire while the thing is losing anyway. And the 2R-partial's headline is now the opposite of its rationale: it improves the mean by \$307.50 while its daily standard deviation goes from 42.3 to **83.3**. It was adopted as a variance lever. On this window it is a mean lever that increases variance, which is a different product, and it should not be described as smoothing anything until it is re-measured on a week where grind actually trades.

THE OVER-TRADING CAVEAT, BAKED IN

The shadow board took 2,236 fires this week across 38 families. The live desk took 55 lots. A shadow arm has no capital, no queue position, no thin-tape default-veto and no operator, so it nominates far more trades than the paper desk could ever hold, and its per-fire numbers are systematically friendlier than the desk's. **Shadow nominates; paper judges.** Nothing on this page is a P&L forecast — it is a ranking of which signal to let the paper desk judge next.

Disposition — every lead this section touched

LEAD	VERDICT	FOR PARKED: WHAT WOULD REVIVE IT. FOR REFUTED: THE TEST THAT KILLED IT.
abs_veto_55s — promote to a live two-sided slot on Monday	REFUTED	Named test: the Asia void. \$308.50 of its \$332.00 was booked in a window with zero live fills since 08-05, leaving \$23.50 on 34 fires, 2 of 5 days green. Refuted as a THIS-WEEK case only.
abs_veto_55s — keep incubating on its all-time record	SHADOW	It is still the best-evidenced family on the book across all tradable-hour fires. Revive the promotion question when it has a positive tradable-hours week with a positive leave-one-day-out — not before.
lad_absS_B_15 / lad_absS_A_10 (ladder arms on abs-veto short)	SHADOW	The most robust things on this week's honest board: 4 of 5 days green and positive leave-one-day-out on both. But n=20 in their first week and no all-time record at all. Revive as promotion candidates at n≥60 tradable fires.
odr_c10_s30 (open-rider claim ladder)	SHADOW	Biggest tradable number on the board and the least robust: strip-best-3 turns it negative and 2 of 5 days are red on n=18. Keep it running; do not act on it.
abs_veto_50s and abs_veto_60s	PARKED	Both negative in tradable hours this week. Revive only if the 55s arm is promoted and a window-length sweep is wanted around it — there is no independent case for either.
The 7 byte-identical shadow arms	FAULT	BUILD item: a row-level duplicate check at shadow-write time, and re-seed the three broken arms. Until then every A/B delta on this board is unverified.
The two-ratchet grind exit	PARKED	Unchanged since 08-05 because grind has not fired. Revive the clip test on the first week grind trades at all — the ATR-floor question in Part 2.5 gates it.
The 2R-partial as a VARIANCE lever	REFUTED	Named measurement: daily standard deviation 42.3 → 83.3 on the same window it improves the mean. It is not smoothing. Re-open only as a mean lever, and only on fresh grind trades.

P2.5 MID-WEEK MUSINGS — every lead from the week, tested until it broke

Garrath — four leads came out of the week's chats and mid-week labs, and this section exists to test them hard rather than repeat them. **One of them dies in a way that is genuinely useful:** the “only run grind on a clean

trend day” idea works on paper, makes +\$580 where the control loses \$958, and then turns out to be un-runnable, because the gate's own ATR floor and the clean-trend condition almost never happen on the same day. That is a better answer than either a yes or a no, and it changes what the Saturday question is.

1 · LEAD ONE — should grind only run on a confirmed trend?

The mechanism lab took every live grind fire in the record — 128 of them, losing \$958.00 between them — and asked which permission rule would have kept the winners while dropping the churn. Four families of rule, each scored on the same fires, with the winner-retention column carried all the way through because a rule that improves the total by throwing away winners is a fake:

MECHANIS M	FIRES KEPT	NET	\$ / FIRE	WIN %	WINNERS KEPT	BEST 3 STRIPPED	LEAVE- ONE-OUT
Every live grind fire (the control)	128	−\$958.00	−\$7.48	34%	43 / 43	−\$1,740.50	−\$1,661.50
arm only on clean-trend days day-type arming	67	\$580.50	\$8.66	37%	25 / 43	−\$202.00	−\$123.00
arm on clean- trend OR in- between day-type arming	81	\$245.00	\$3.02	36%	29 / 43	−\$537.50	−\$458.50
arm when the day's roundtrip >= 0.55 day-type arming	67	\$580.50	\$8.66	37%	25 / 43	−\$202.00	−\$123.00
arm when the day's roundtrip >= 0.45 day-type arming	71	\$431.50	\$6.08	37%	26 / 43	−\$351.00	−\$272.00
require side*net30 >= 0.5 x ATR directional-net requirement	124	−\$889.50	−\$7.17	34%	42 / 43	−\$1,672.00	−\$1,593.50
require side*net30 >= 1.0 x ATR directional-net requirement	117	−\$675.00	−\$5.77	35%	41 / 43	−\$1,457.50	−\$1,513.50
require side*net30 >= 1.5 x ATR directional-net requirement	113	−\$1,037.00	−\$9.18	33%	37 / 43	−\$1,819.50	−\$1,513.50
require er30 >= 0.25 efficiency-ratio floor	94	−\$791.50	−\$8.42	35%	33 / 43	−\$1,574.00	−\$1,403.00

MECHANIS M	FIRES KEPT	NET	\$ / FIRE	WIN %	WINNERS KEPT	BEST 3 STRIPPED	LEAVE- ONE-OUT
require er30 >= 0.35 efficiency-ratio floor	35	-\$417.50	-\$11.93	31%	11 / 43	-\$1,092.50	-\$1,043.00
require er30 >= 0.45 efficiency-ratio floor	11	-\$636.00	-\$57.82	18%	2 / 43	-\$804.50	-\$833.50
clean-trend day AND side*net30>=1 .0xATR the two stacked	61	\$715.50	\$11.73	39%	24 / 43	-\$67.00	-\$123.00
clean-trend day AND er30>=0.35 the two stacked	12	\$636.50	\$53.04	50%	6 / 43	-\$38.50	\$11.00

Read the last three columns, not the third. The best day-type rule (arm only on clean-trend days) turns -\$958.00 into **\$580.50** — and then strip its best three fires and it is -\$202.00, leave one day out and the worst fold is -\$123.00. It keeps 25 of the 43 winners, so it is not a fake win; it is a real effect with no robustness underneath it. The stacked version (clean-trend day AND side*net30>=1.0xATR) is nominally the best cell in the whole table at \$715.50 on n=61, and it keeps only 24 of 43 winners. Neither is deployable.

Note also what the b- and c-family rows say, because they close a recurring proposal: **no directional-net requirement and no ER floor rescues grind**. Every single one of them stays negative. The ER-0.35 floor that is already live keeps 35 fires and still loses -\$417.50. This desk has now refuted “fix grind with an efficiency filter” on three separate occasions and it should stop being proposed.

★ AND THEN THE FINDING THAT ACTUALLY MATTERS — THE RULE CANNOT BE RUN

The day-type lab classified 240 sessions. **97 of them are clean-trend days.** The number that are clean-trend AND clear grind's 22-point ATR floor is **1**.

So “arm grind on clean-trend days” and “grind's ATR floor lets it fire” are very nearly mutually exclusive conditions. A clean trend is a day that goes somewhere smoothly; the 22-point floor wants a day that is violent. The rule that scores +\$580 in the table above would, in production, have armed the gate on about one day in two hundred and forty — and on the other 239 the arming decision would have been irrelevant because the floor would have refused every signal anyway.

This is why the mid-week lead is neither adopted nor refuted, but reframed. The question is not ‘when should grind be armed?’. Nothing about arming is binding. The question is whether 22 is the right floor, and that has never been tested — it is a number the gate was born with.

1b · ★ REV2 — “ARM ON CLEAN-TREND DAYS” vs “THE DESK EARNS IN BUILDING BLOCKS”: WHICH IS IT?

THE CONTRADICTION, STATED

This section parks arm grind only on clean-trend days as the best rule on the board. Movement 3 §8 attributes every live MNQ tournament trade to the 30-minute regime block it opened in and finds the desk **loses** in CLEAN-TREND blocks (−\$181.00 on 24 trades) and makes all its money in BUILDING blocks (+\$1,836.50 on 121). Read side by side those look like opposite instructions from the same desk.

They are not in conflict, because they are not the same classifier, the same unit, or the same population — and the first of those is the one that matters.

	PART 2.5'S CLASSIFIER	MOVEMENT 3 §8'S CLASSIFIER
Unit	a whole SESSION	a 30-MINUTE BLOCK
Rule	terciles of the session round-trip ratio (7.18 / 10.99)	ER≥0.45 + a breakout → CLEAN-TREND; high-ATR & ER<0.25 → VIOLENT-WHIPSAB; ER<0.22 & low-ATR → DEAD-CHOP; ER<0.30 → NORMAL-CHOP; otherwise BUILDING

	PART 2.5'S CLASSIFIER	MOVEMENT 3 §8'S CLASSIFIER
Population scored	128 live grind fires, 07-19 → 08-04	699 live tournament trades, all gates
How often the label fires	97 of 240 sessions (40.4%) are clean-trend	17 of 991 blocks (1.7%) are CLEAN-TREND

A “clean-trend day” is not a day made of clean-trend blocks. Across the 22 days with a block census on file, the CLEAN-TREND share of a day’s blocks runs from 0.0% to **4.8%**, median 2.2% — it never exceeds one block in twenty. The two labels share a name and almost nothing else, which is the whole answer to “which is it?": neither claim is about the other’s object.

1c · So re-score the grind rule against BUILDING blocks, which is the real question

Each of the 128 live grind fires classified by its own 30-minute block, using Movement 3 §8’s exact rule and its own ATR tercile boundaries (10.17 / 15.09 from regime_blocks.csv). Net of the \$1.50 round trip throughout.

ARMING CONDITIO N	FIRES	NET	\$ / FIRE	WIN %	WINNERS KEPT	BEST 3 STRIPPED	WORST LEAVE- ONE-OUT
none (the control)	128	-\$1,150.00	-\$8.98	33.6%	43 / 43	-\$1,928.00	-\$1,784.50
arm only in BUILDING blocks	51	-\$748.50	-\$14.68	33.3%	17 / 43	-\$1,419.00	-\$1,245.00
arm except in VIOLENT-WHIPSAW blocks	99	-\$829.00	-\$8.37	35.4%	35 / 43	-\$1,607.00	-\$1,514.00

Arming grind on BUILDING blocks does not rescue it — it makes it worse per fire (-\$14.68 against the control’s -\$8.98) while throwing away 26 of the 43 winners. So the parked rule is not parked behind the wrong condition; the block-level version of the condition is the worse of the two. Movement 3 §8’s BUILDING result is a statement about the tournament as a whole — six gates, 121 trades — and it does not transfer to this gate.

★ AND ONE THING THAT FALLS OUT OF DOING THIS: THE DAY-TYPE RULE IS ONE SESSION

Net of fees, the clean-trend-day rule’s +\$580.50 (gross) is **+\$548.00 on 63 of the 67 fires from just two sessions**, and 2026-07-29 alone is **+\$634.50 on 46 fires**. Drop that one day and the remaining eight days are **-\$1,784.50 on 82 fires**. The report’s own leave-one-out column said –

\$123.00 and that is the same fact in gross terms; it is worth stating in plain English because “the best rule on the board” and “arm on 2026-07-29” are, on this sample, very nearly the same instruction.

2 · So how reachable IS a 22-point ATR? Two independent measurements

MEASUREMENT	WINDOW	MEDIAN 1-MIN ATR-14	P90	% OF MINUTES ≥ 22PT	DAYS WITH ANY SUCH MINUTE
The p25 lab (via the parquet lake)	239 sessions	14.66	31.34	25.3%	213 of 239
This report, recomputed off capture.db's own 5s bars	33 sessions, 2026-07-15 → 2026-08-21	11.69	22.64	11.0%	31 of 33

★ REV2 — RESOLVED, AND IT IS NOT A WINDOW, IT IS A DENOMINATOR. Rev1 printed both and picked neither, which left every proposed ladder rung sized off an unresolved 2.3×. Recomputed on ONE method (ATR-14 on 1-min bars, halt-aware, exactly deciders._atr) with only the population varied:

POPULATION	MINUTES	MEDIAN	P90	≥ 22PT	22PT IS THE...
capture.db, 07-15→08-21, all 24 hours	37,411	11.68	22.62	11.0%	p89
the lake, last 239 days, all 24 hours	231,391	10.21	22.12	10.2%	p90
the lake, last 239 days, 13:00-21:00Z only	79,411	14.64	31.34	25.1%	p75
capture.db, THIS WEEK, all hours	6,780	9.66	16.88	5.4%	p95
capture.db, THIS WEEK, 13:00-20:00Z only	11,436	16.46	34.64	29.5%	p71

The p25 lab measured the CME session block; this report measured all 24 hours. The session row reproduces the p25 lab to within rounding on all three of its numbers (14.64 vs 14.66, 31.34 vs 31.34, 25.1% vs 25.3%), and the all-hours row reproduces this report’s capture.db recompute to the second decimal (11.68 vs 11.69, 22.62 vs 22.64, 11.0% vs 11.0%). Neither is wrong and they never disagreed about the tape. The lake and capture.db agree at ~10–11% once put on the same basis, so the “two independent measurements” framing was never the

issue either. **Which one sizes a ladder: the SESSION one**, because a floor can only bind in hours the gate can fire, and the desk refuses new entries 00:00–07:00Z anyway.

★ **REV2: “the floor sits at the ninetieth percentile” is true of the ALL-HOURS denominator and false of the session one, where it is the p75** — and on this week’s session tape it was reachable in 29.5% of minutes. So it is not a p90 off-switch; it is a p75 filter that behaves like an off-switch, and the reason is not the height of the floor. It is that grind’s own trigger and its own floor are anti-correlated: 29.5% of this week’s session minutes cleared 22pt, and yet only **11 of 24,121** raw grind signals did. A gate that triggers on quiet absorption looks for exactly the minutes its floor excludes. Lowering the floor does not fix that — §2b shows dropping it to 14pt adds 46 fires and makes the book worse. A floor at p90 is not a filter, it is an off-switch that opens during violence, and it explains every grind number in this report: 24,110 signals blocked out of 24,121 this week, no live fire since 5 August, and two grind-exit shadow watches with nothing in them.

One more number worth having: **31 of 33 days contain at least one minute above the floor, but only 1 of 33 spend a quarter of the session there**. So “the floor was reachable today” is true almost every day and means nothing; the floor is reachable for a few minutes of spike and then gone.

2b · ★ REV2 — THE FLOOR, SWEEPED TONIGHT, ON THE 128 LIVE FIRES

WHY THIS DID NOT NEED TO WAIT FOR A SHADOW LADDER

Rev1’s answer to “is 22 the right floor?” was a shadow ladder that needs $n \geq 40$. **A gate that has fired zero live lots since 5 August cannot deliver $n \geq 40$** . But the 128 fires it made before the floor bound are on disk with their own entry ATR and their own real exits, so the floor is a filter on a real book and can be swept now, tick-honest by construction. Everything below is NET of the \$1.50 round trip — the –\$958.00 control in circulation is gross; the same 128 fires are **–\$1,150.00** net.

ATR FLOOR	FIRES KEPT	NET	\$ / FIRE	WIN %	WINNERS KEPT	BEST 3 STRIPPED	WORST LEAVE-ONE-OUT	VS RANDOM SUBSET
no floor (the control)	128	−\$1,150.00	−\$8.98	33.6%	43 / 43	−\$1,928.00	−\$1,784.50	—
ATR ≥ 10pt	118	−\$1,108.00	−\$9.39	33.9%	40 / 43	−\$1,886.00	−\$1,742.50	—
ATR ≥ 12pt	107	−\$1,111.00	−\$10.38	32.7%	35 / 43	−\$1,889.00	−\$1,796.00	—
ATR ≥ 14pt	97	−\$1,143.50	−\$11.79	32.0%	31 / 43	−\$1,921.50	−\$1,727.00	25th pct
ATR ≥ 16pt	90	−\$799.50	−\$8.88	33.3%	30 / 43	−\$1,577.50	−\$1,613.50	51th pct
ATR ≥ 18pt	82	−\$554.50	−\$6.76	35.4%	29 / 43	−\$1,332.50	−\$1,444.00	65th pct
ATR ≥ 20pt	70	−\$514.50	−\$7.35	34.3%	24 / 43	−\$1,292.50	−\$1,453.00	61th pct
ATR ≥ 22pt ← live	51	−\$280.50	−\$5.50	35.3%	18 / 43	−\$1,058.50	−\$1,174.50	65th pct
ATR ≥ 24pt	45	\$21.00	\$0.47	40.0%	18 / 43	−\$757.00	−\$982.50	83th pct
ATR ≥ 26pt	33	\$347.50	\$10.53	42.4%	14 / 43	−\$430.50	−\$643.00	—
ATR ≥ 28pt	28	\$469.00	\$16.75	46.4%	13 / 43	−\$309.00	−\$633.50	—

The ladder is monotone upward from 14pt, which is the shape of a real effect — −\$1,143.50 → −\$799.50 → −\$554.50 → −\$514.50 → −\$280.50 → +\$21.00 → +\$347.50 → +\$469.00. Three things follow, and none of them is the one the report expected.

One: 22 is not obviously wrong, and it is not enough. The live floor still leaves the gate at **−\$280.50 on 51 fires**. It does not cross zero until 24pt and does not clear it properly until 26–28. So the honest reading of the direction is “higher would have been better”, not “22 is a mistake”.

Two: no rung survives the battery. Strip the best three fires and every rung is negative — −\$309.00 even at 28pt. The worst leave-one-day-out fold is negative at every rung, −\$633.50 even at 28pt. And against 5,000 random same-size subsets of the same 128 fires, the 22pt rung sits at the 65th percentile and the 24pt rung at the 83rd — the floor is selecting a little, and not much. **Cause of death: strip-best-3 and leave-one-day-out, at every rung tested.**

Three, and this is the part that changes the Saturday question: the floor is not what is wrong with grind. Rev1 framed it as the un-asked question on which the whole grind programme turns. It has now been asked, and the answer is that raising the floor cannot make this gate positive on its own live record. That does not make the shadow ladder pointless — it can still test whether the same monotone shape holds forward — but it demotes it from “the real grind question” to a confirmation, and it sends grind back to being an entry question, which is where §1’s b- and c-family rows already pointed.

3 · What the week's days actually looked like

DAY TYPE	DAYS	SHARE OF 240
clean-trend	97	40.4%
in-between building	57	23.8%
normal-chop	52	21.7%
dead-chop	27	11.2%
violent-whipsaw	7	2.9%

Two-fifths of sessions classify as clean trend, which is a lot more than the desk's folklore assumes, and only 7 of 240 are the violent-whipsaw type the desk spends most of its worry on. The terciles behind the classification are a session round-trip ratio of 7.18 and 10.99.

4 · THE GOLD PROGRAMME — four arms, and gold is \$10 a point

Before any number: the multiplier check. The gold tables below show a median session of about −\$133.14. At \$10 a point that is roughly 13 points, which is an ordinary MGC day-trade excursion. At the MNQ multiplier it would be 67 points, which gold does not do intraday. **The lab is priced correctly.**

4a — The coil, and its mirror control

GOLD COIL ARM	N	SESSIONS	NET	\$ / TRADE	WIN %	BEST 3 STRIPPED	BEST DAY STRIPPED
REVERSION (fade the coil edge → mid)	2487	251	−\$1,828.93	−\$0.74	43.7%	−\$2,398.93	−\$2,244.93
CONTINUATION (break out) [mirror control]	2487	251	−\$4,767.86	−\$1.92	39.3%	−\$5,470.86	−\$5,519.21

Both directions lose. Fading the coil edge loses \$1,828.93 over 2,487 trades and breaking out of it loses \$4,767.86 — and the placebo is the sentence that closes it: **0 of the random sets did worse**, against a placebo median of −\$3,354.96. A coil edge on gold carries no directional information in either direction at this fidelity.

4b — Except in the tightest coils, which is a PARKED lead, not a finding

COIL TIGHTNESS	TARGET	N	NET	\$ / TRADE	WIN %
compression ≤ 0.45	1.0×	666	\$307.29	\$0.46	47.1%

COIL TIGHTNESS	TARGET	N	NET	\$ / TRADE	WIN %
compression ≤ 0.45	1.5×	666	\$616.32	\$0.93	37.8%
compression ≤ 0.45	2.0×	666	\$280.36	\$0.42	31.1%
compression ≤ 0.6	1.0×	2487	−\$1,828.93	−\$0.74	43.7%
compression ≤ 0.6	1.5×	2487	−\$231.36	−\$0.09	35.1%
compression ≤ 0.6	2.0×	2487	\$962.00	\$0.39	29.8%
compression ≤ 0.75	1.0×	4726	−\$4,430.00	−\$0.94	41.9%
compression ≤ 0.75	1.5×	4726	−\$4,084.29	−\$0.86	32.7%
compression ≤ 0.75	2.0×	4726	−\$4,978.36	−\$1.05	26.8%

Restrict to the tightest compression and the sign flips: \$962.00 at 2.0× on n=2,487. That is a real number on real n and it is **still not a result**, because it is the best cell of nine on the same data that produced the null above, with no placebo run against it and no out-of-sample. It is PARKED with a named revival condition: run the same placebo the full-sample arm got, and hold out a year.

4c — The gold exit and stop lab

GOLD STOP WIDTH	STOPPED OUT	NET	\$ / SESSION	WIN %	BEST 3 STRIPPED	LEAVE-ONE- OUT
2.0x	67	−\$6,304.86	−\$78.81	15.0%	−\$9,251.86	−\$6,879.57
2.5x	62	−\$6,112.36	−\$76.40	20.0%	−\$9,413.36	−\$6,252.50
3.0x	56	−\$4,669.00	−\$58.36	26.2%	−\$8,366.00	−\$5,233.57
3.5x	52	−\$3,437.00	−\$42.96	31.2%	−\$7,134.00	−\$4,006.50
4.0x	45	−\$376.29	−\$4.70	38.8%	−\$4,285.29	−\$2,465.00
4.5x	41	\$1,884.57	\$23.56	43.8%	−\$2,024.43	\$66.64
5.0x	37	\$3,352.71	\$41.91	46.2%	−\$556.29	\$246.43
5.5x	37	\$1,985.79	\$24.82	46.2%	−\$1,923.21	−\$905.43
6.0x	35	\$1,240.29	\$15.50	47.5%	−\$2,668.71	−\$1,435.86
7.0x	30	\$3,264.00	\$40.80	52.5%	−\$645.00	−\$2,084.00
8.0x	24	\$2,776.29	\$34.70	53.8%	−\$1,132.71	−\$2,363.29
10.0x	17	\$2,055.14	\$25.69	55.0%	−\$1,853.86	−\$2,667.57
naked	0	−\$1,570.00	−\$19.63	56.2%	−\$5,479.00	−\$6,825.00

GOLD EXIT	NET	\$ / SESSION	WIN %	WORST SESSION	BEST 3 STRIPPED	WORST LEAVE- ONE-OUT	HOW THEY ENDED
5xATR stop, hold to clock	\$3,352.71	\$41.91	46.2%	−\$1,045.14	−\$556.29	\$246.43	CLOCK×43, STOP×37
2.5xATR trail, 3h cap	−\$156.43	−\$1.96	28.7%	−\$284.43	−\$3,185.14	−\$1,940.00	TRAIL×79, CAP×1
hold naked to the 20:40 clock	−\$1,570.00	−\$19.63	56.2%	−\$3,847.00	−\$5,479.00	−\$6,825.00	CLOCK×80
2xATR stop + 4xATR trail	−\$3,440.29	−\$43.00	22.5%	−\$419.86	−\$5,980.43	−\$3,835.00	STOP×46, TRAIL×32, CLOCK×2

GOLD EXIT	NET	\$ / SESSION	WIN %	WORST SESSION	BEST 3 STRIPPED	WORST LEAVE-ONE-OUT	HOW THEY ENDED
1xATR stop, hold to clock	−\$4,263.14	−\$53.29	5.0%	−\$242.57	−\$5,962.14	−\$4,233.14	STOP×76, CLOCK×4
1.5xATR stop, hold to clock	−\$4,500.43	−\$56.26	11.2%	−\$362.36	−\$7,447.43	−\$5,348.36	STOP×71, CLOCK×9
3xATR stop, hold to clock	−\$4,669.00	−\$58.36	26.2%	−\$688.71	−\$8,366.00	−\$5,233.57	STOP×56, CLOCK×24
2xATR stop, hold to clock	−\$6,304.86	−\$78.81	15.0%	−\$460.14	−\$9,251.86	−\$6,879.57	STOP×67, CLOCK×13

This is the most interesting gold result and also the one most likely to be a mirage. There is a genuine plateau — the stop has to be enormous before gold pays, and the progression from 2.0× to 5.0× is monotone rather than jumpy, which is what a real effect looks like rather than a grid artefact. But look at the two robustness columns on the best cell: strip its best three sessions and \$3,352.71 becomes −\$556.29, and the worst leave-one-out fold is \$246.43.

Nearly the whole result is three sessions. And a 5×ATR stop on gold is not really a stop; it is a disaster brake on a naked hold, which the top row of the exit table already shows losing money.

4d — The gold rider, in and out of sample — and a control that undoes it

GOLD RIDER ARM	FIRES	SESSIONS	FIRE RATE	NET	\$ / FIRE	MEDIAN MFE / MAE (\$)
loose trigger — IN-SAMPLE (the 16-ish)		12	27	44.4%	−\$328.00	−\$27.33 253 / -302
loose trigger — OUT-OF-SAMPLE		106	226	46.9%	−\$3,306.00	−\$31.19 307 / -264
loose trigger — ALL 253		118	253	46.6%	−\$3,634.00	−\$30.80 299 / -268
derived trigger — IN-SAMPLE		11	27	40.7%	−\$147.00	−\$13.36 293 / -298
derived trigger — OUT-OF-SAMPLE		79	226	35.0%	\$655.00	\$8.29 319 / -260
derived trigger — ALL 253		90	253	35.6%	\$508.00	\$5.64 314 / -261

The loose trigger is negative in and out of sample, which is clean. The derived trigger is the interesting one: −\$147.00 in-sample and **\$655.00 out-of-sample on 79 fires over 226 sessions**. An arm that does better out of sample than in is usually a good sign.

★ BUT THE DIRECTIONAL CONTROL KILLS IT, AND IT HAS TO BE SAID ON THE SAME PAGE

The same out-of-sample window carries a plain directional control: hold gold LONG every session and you lose **−\$21,738.00**; hold it SHORT every session and you make **\$20,382.00**, over 226 sessions.

The gold rider's out-of-sample \$655.00 is a rounding error against a \$20,382 short drift sitting underneath it. Any gold arm that is even slightly short-biased will look profitable on this window without carrying one gram of edge. **Every gold result in this report needs its long/short mix reported next to it, and none of them currently do.** That is a method fix, it is on the BUILD card, and until it lands no gold arm should be promoted to anything.

5 • The MNQ open-rider leads, and the stay-out meter

OPEN-RIDER POPULATION	SESSIONS	HELD TO THE CLOCK %	NET AT 2 LOTS	\$ / SESSION	MEDIAN SESSION	BEST 3 STRIPPED
every session it fired	225	47.1%	−\$20,739.00	−\$92.17	−\$78.00	−\$29,666.00
inside the window	60	53.3%	\$825.00	\$13.75	\$124.50	−\$7,745.00
outside it	165	44.8%	−\$21,564.00	−\$130.69	−\$112.00	−\$27,477.00

The open rider, replayed across 229 sessions at two lots, loses \$20,739.00. Split it by the window lead and the “inside” population makes \$825.00 on 60 sessions while “outside” loses \$21,564.00 on 165. That looks like the window is the whole story until the strip-best-3 column: the inside population goes to −\$7,745.00.

Three sessions are the entire lead. REFUTED as a session filter on this evidence.

The stay-out meter's direction call, on the same population: when its two reads agree it is right 51.7% of the time over 178 sessions, which is a coin. Split by call, its UP calls hit 56.7% and its DOWN calls 45.7% — and that gap is the same short drift the gold control just exposed, not skill. The 51 mixed-signal sessions average −\$20.79. The corrected series (231 sessions) says the same thing to within a point. **The meter is a stay-out instrument and it should not be read as a direction instrument**, which is exactly what it says on its own label and exactly what keeps getting forgotten.

Disposition — every lead this section touched

LEAD	VERDICT	FOR PARKED: WHAT WOULD REVIVE IT. FOR REFUTED: THE TEST THAT KILLED IT.
Arm grind only on clean-trend days	PARKED	It is not refuted — it makes \$580.50 gross against a −\$958.00 gross control and keeps 25 of 43 winners. It is un-runnable: only 1 of 240 sessions is both clean-trend and above the 22-point floor. ★ REV2 adds two things and neither revives it. Net of fees the rule is +\$548.00 and 2026-07-29 alone is +\$634.50 of it on 46 fires — drop that one session and the other eight days are −\$1,784.50, so “the best rule on the board” and “arm on 2026-07-29” are nearly the same instruction on this sample. And the BUILDING-block version of the idea, scored on the same 128 fires with Movement 3 §8's own classifier, is −\$748.50 on 51 fires keeping 17 of 43 winners — worse per fire than no rule at all, so it is not parked behind the wrong condition. The floor question below is answered too, and it does not revive this either.
grind's 22-point ATR floor as a number	SHADOW	★ REV2 — ASKED AND ANSWERED, tonight, on the 128 live fires (§2b). The ladder is monotone upward — 14pt −\$1,143.50 → 22pt −\$280.50 → 28pt +\$469.00 — but strip-best-3 is negative at every rung and the worst leave-one-day-out is negative at every rung, and against 5,000 random same-size subsets the live 22pt rung sits at only the 65th percentile. Cause of death: strip-best-3 and leave-one-day-out. Raising the floor cannot make grind positive on its own live record, so this is no longer “the real grind question” — grind is an entry question. Keep the shadow arms as a forward confirmation of the monotone shape, nothing more. Reachability is also resolved: 22pt is the p75 of SESSION minutes and the p89 of ALL minutes; the two published figures were one denominator apart, not two measurements.
Fix grind with an ER floor or a directional-net requirement	REFUTED	Named test: nine cells across two families on 128 live fires, every one negative, including the ER-0.35 floor that is already deployed.
MGC coil — fade the edge (or break it)	REFUTED	Named test: a mirror control that also loses, and a placebo that 0 random sets failed to beat from below on 2,487 trades.
MGC coil, tightest-compression cell only	PARKED	\$962.00 on n=2,487. Revive by running the SAME placebo the full-sample arm got and holding out a year of tape. Until then it is the best of nine cells on data that produced a null.
MGC 5×ATR stop / hold-to-clock	PARKED	Monotone plateau, which is the shape of a real effect, but strip-best-3 takes \$3,352.71 to −\$556.29. Revive at n≥200 sessions or on a second instrument.
The MGC derived-trigger rider	PARKED	Out-of-sample \$655.00, which is real, against a \$20,382 always-short drift on the

LEAD	VERDICT	FOR PARKED: WHAT WOULD REVIVE IT. FOR REFUTED: THE TEST THAT KILLED IT.
Report the long/short mix on every gold arm	LIVE	same window. Revive only with the long/short mix reported beside it.
The open-rider &inside the window’ session filter	REFUTED	A method fix, not a strategy. Without it every gold number on this desk is confounded with a large short drift, and this section found two arms already reading it as edge.
The stay-out meter as a DIRECTION signal	REFUTED	Named test: strip-best-3 takes \$825.00 to -\$7,745.00 on 60 sessions. Three sessions are the whole lead.
		51.7% hit rate over 178 sessions when both reads agree. Its up/down asymmetry is the market's short drift, not the meter's skill. It stays in service as a stay-out instrument, which is what it was built to be.

P2.6 ROUTER OPERATION REVIEW — the desk's #1 lever, on the stand

Garrath — the router ran 1,425 ticks this week and wrote a switch on 20 of them. That is the headline and it is not a criticism: the tournament has been stood down all week, so most of what the router had to decide was whether to keep doing nothing, and it kept doing nothing correctly. **What this section is actually about is the three things that can write the switch file or claim to know the desk's state, and the fact that on the two occasions this week where those disagreed, the router was right and the instruments around it were wrong.** It also has to say plainly that the two files built to answer “was the router optimal this week” both stopped being written, so everything here is rebuilt from raw logs.

1 • The duty cycle

SESSION	TICKS	TICKS THAT CHANGED A SWITCH	CHANGE RATE
Mon 17 Aug	279	5	1.8%
Tue 18 Aug	283	1	0.4%
Wed 19 Aug	291	11	3.8%
THE WEEK	1425	20	1.4%

SESSION	TICKS	TICKS THAT CHANGED A SWITCH	CHANGE RATE
Thu 20 Aug	288	2	0.7%
Fri 21 Aug	284	1	0.4%
THE WEEK	1425	20	1.4%

1.4% of ticks change anything. The audit's number for the fortnight was 97.2% no-change; this week is 98.6% no-change, so if anything the router has got quieter, which under a standdown is the correct direction. There were no systemd failures and no aborted decisions this week — both worth stating, because last cycle had a 10-hour block of aborts that no scoreboard noticed.

2 · Every switch it wrote, graded

WHEN (UTC)	WHAT IT WROTE	GRADE	WHY
08-17 08:10Z	{'abs_veto_short': 'off'}	☐ GOOD	Benched into the pre-open. The gate took two stops the previous hour and the tape was going nowhere.
08-17 08:25Z	{'abs_veto_short': 'on'}	⚠ CHURN	Re-armed 15 minutes later. Nothing changed in between; this is the toggle pattern the ledger keeps flagging.
08-17 10:30Z	{'abs_veto_long': 'on'}	☐ GOOD	Armed the long fader into a genuine bid.
08-17 12:40Z	{'abs_veto_long': 'off'}	☐ GOOD	The fail-safe bench that avoided the 12:40 reversal. A fail-safe bench is exempt from the confirmation rule, and this is the example of why.
08-17 15:00Z	{'capitulation_long': 'off', 'exhaustion_short': 'off'}	☐ GOOD	Benched both faders into the afternoon drift. Neither fired again that day.
08-18 01:55Z	{'capitulation_long': 'off'}	— \$0	An Asia-hours bench. Worth exactly nothing in either direction: entries are refused in that block anyway.
08-19 07:00Z	{'abs_veto_short': 'off'}	☐ GOOD	Benched the short book at the European open before the 07:50 squeeze.
08-19 07:50Z	{'abs_veto_short': 'on'}	☐ GOOD	Re-armed for the down-leg it then caught (+\$179 in claims at 13:50).
08-19 12:40Z	{'abs_veto_long': 'on', 'abs_veto_short': 'off', 'exhaustion_short': 'off'}	☐ BAD	Benched abs_veto_short and exhaustion_short ten minutes before the 13:39 down-leg that made the rider \$541, and armed the LONG fader into it. Wrong side, wrong ten minutes.
08-19 12:51Z	{'grind_long': 'on'}	☐ BAD	Armed grind_long into the same down-leg. It could not fire — the ATR floor saw to that — so the cost is \$0 and the process is still wrong.

WHEN (UTC)	WHAT IT WROTE	GRADE	WHY
08-19 13:40Z	<code>{'grind_long': 'off', 'abs_veto_short': 'on'}</code>	▢ GOOD	Corrected itself within the hour: grind off, short book back on, one minute after the rider's entry.
08-19 13:41Z	<code>{'grind_long': 'off', 'abs_veto_long': 'off', 'abs_veto_short': 'on'}</code>	△ CHURN	A second write 30 seconds after the first, same content plus <code>abs_veto_long</code> . Two writes, one decision.
08-19 13:50Z	<code>{'capitulation_long': 'off'}</code>	— \$0	Housekeeping on a gate that had not fired.
08-19 15:05Z	<code>{'abs_veto_short': 'off'}</code>	▢ GOOD	Benched the short book after the 14:51 stop.
08-19 15:35Z	<code>{'abs_veto_short': 'on'}</code>	▢ BAD	Re-armed 30 minutes later with no positive case. It then took three more stops for −\$133.
08-19 18:16Z	<code>{'abs_veto_short': 'off'}</code>	▢ GOOD	Applied the operator's 18:15Z bench instruction within 90 seconds.
08-19 22:00Z	<code>{'abs_veto_short': 'off'}</code>	— \$0	Re-asserted the bench at the reopen. No-op.
08-20 06:40Z	<code>{'exhaustion_short': 'off'}</code>	— \$0	Asia-hours bench.
08-20 08:05Z	<code>{'capitulation_long': 'off'}</code>	— \$0	Housekeeping under the standdown.
08-21 14:05Z	<code>{'abs_veto_short': 'off'}</code>	▢ GOOD	Re-benched <code>abs_veto_short</code> 2m15s after the OPEN-HOUR WATCHER auto-armed it through a standing operator standdown. See §4 — the router was the only thing that caught it.

10 good, 3 bad, 2 churn, 5 worth exactly \$0. The \$0 rows are Asia-hours writes: new entries are refused in that block regardless, so a bench there is neither a save nor a cost, and it should not be counted as either. The two genuinely bad calls are both on Wednesday and they are the same mistake made twice — benching the short book at 12:40 ten minutes before the day's best down-leg, then re-arming at 15:35 with no positive case, straight into three more stops. **Re-arming without a positive live case is now this gate's most expensive recurring pattern**, and it has cost money on three separate occasions that the ledger records.

3 · Is the router earning its keep? The mechanical counterfactual

The honest way to ask is not “did the benched gates lose money”, because a benched gate has no P&L. Movement 2 fires all six mechanically in their exact live config across the whole of this week's captured tape — every fire counted, not just the ones near a big move — which is the counterfactual this question needs:

GATE	LIVE THIS WEEK	IF IT HAD BEEN ARMED ALL WEEK, MECHANICALLY	DIFFERENCE
grind_long	benched, \$0	–\$319.17 on 11 fires	\$319.17
capitulation_long	benched, \$0	–\$97.66 on 10 fires	\$97.66
abs_veto_long	benched, \$0	–\$83.42 on 26 fires	\$83.42
rgv_short	benched, \$0	–\$17.54 on 1 fires	\$17.54
exhaustion_short	\$144.50	–\$638.02 on 18 fires	\$782.52
abs_veto_short	–\$203.00	–\$749.55 on 51 fires	\$546.55
ALL SIX	–\$58.50	–\$1,905.36 on 117 fires	\$1,846.86

The tournament's live book was –\$58.50 this week. Firing the same six gates mechanically through the same tape is **–\$1,905.36 on 117 fires**. On that comparison the standdown plus the router's benching is worth about **\$1,846.86** this week, and \$517.79 of it is the four gates that never fired at all.

★ REV2 — THE **-\$638.02 IS NOT EXHAUSTION_SHORT'S COUNTERFACTUAL, AND THE \$782.52 IS NOT YOUR HAND**

That row invited a reading the data does not support: that the operator turned a **-\$638.02** gate into a **+\$144.50** one. Two things are wrong with it, and the second is a method wound rather than an arithmetic slip.

One: repriced on the SAME 14 lots, the machine alone books +\$84.00. Taking each lot's protective stop straight off the venue (so R is exact, and 7 of 7 machine-closed lots reproduce their own live exit reason) and racing it to its own configured stop or target, the machine gets **+\$84.00** and your hand adds **+\$60.50**. Not \$782.52. This gate is not hand-carried, which is the opposite of what the same test says about `abs_veto_short`.

Two: only ONE of the 18 mechanical fires is a live signal. Matching the sim's fires against the gate's seven live entries: one lines up within 3 seconds (08-18 13:37, and even that one enters 8.75 points away), two more line up within 3 minutes at 17-33 points of difference, and **three live signals have no sim fire within TEN MINUTES** — 14:02, 14:20 and 14:37 on 08-18, which are precisely the three that made money. Meanwhile the sim takes 17 fires the live gate never took. The contrast with `abs_veto_short` is the tell: for that gate the sim reproduces **11 of 12** live signals within 6 seconds and about 5 points, so its **-\$749.55** is a counterfactual. `exhaustion_short` fires on a 5-second footprint confirmation that does not replay faithfully off bars, so its **-\$638.02** is a different strategy's P&L wearing the same name.

What the row can honestly claim, then: most of the **-\$638.02** is the value of the router's benches (the gate was benched at 15:00Z Monday, 12:40Z Wednesday and 06:40Z Thursday, and 16 of those 17 never-live fires happened while it was), and none of it is a verdict on the seven trades the gate actually took. Quote the ALL SIX total as a rough estimate of what the standdown was worth, exactly as the box below already says — but do not read any single gate's row as that gate's alternative history without first checking that the sim reproduces its fires. **BUILD #9 is the fix:** populate `signal_journal.suppressed_by` so a refusal is recorded rather than re-simulated.

★ AND HERE IS WHY THAT NUMBER IS AN ESTIMATE AND NOT A RESULT

The mechanical sim has no operator. Two of the live gates' best trades this week were closed by hand at a price the machine would not have taken, and the sim rides every one of those to its own stop. It also has no queue, no partial fills and no thin-tape veto. Every one of those biases makes the sim look WORSE than the live desk would have been, which means \$1,846.86 is an upper bound on the value of benching, not a measurement of it.

The audit ran the same question on the fortnight with a different instrument and got a range from +\$2,062 to -\$673 — it straddles zero. **So the correct statement is: the standdown was probably worth something this week, it was certainly not worth thousands, and no instrument on this desk can currently put a confident number on it.** That is a measurement problem and it is on the BUILD card.

4 • ★ THE STANDDOWN IS A ROUTER-LOCAL VARIABLE, AND SOMETHING ELSE ARMED A GATE THROUGH IT

On Friday at 14:03Z, with the tournament stood down and all six gates reading off, `abs_veto_short` was switched **on**. The router did not do it. Here is the sequence, unedited, out of the trial log:

```
2026-08-21T14:00:14Z | tune | DURABLE | changed: none | 14:00Z US-OPEN,
tournament STOOD DOWN and all six gates already read off: no-change
tick. Flat, healthy=True, halted=False;
```

```
2026-08-21T14:03Z APPLIED {abs_veto_short: on} [OPEN-HOUR WATCHER,
automatic]
```

```
2026-08-21T14:05:18Z | tune | DURABLE | changed: {'abs_veto_short':
'off'} | 14:05Z US-OPEN, tournament STOOD DOWN – abs_veto_short reads on
because the OPEN-HOUR WATCHER auto-armed
```

`open_hour_watch.py` is a second writer to `gate_switches.env` with off→on authority and **no awareness of the standdown at all**. It saw a confirmed break (ER 0.36, ATR 22, a new extreme) and did exactly what it was built to do, which is the standing “if a huge trend develops, arm immediately” instruction. The problem is not the watcher's logic. The problem is that a standing operator standdown lives in a variable inside one process, and the other five writers to that file cannot see it.

The gate was live for 2 minutes 15 seconds and fired nothing, so this cost \$0 — by luck, not by design. The exposure is bounded by the router's 5-minute tick, not by any guard. If the watcher had armed at 14:06 instead of 14:03 the gate would have been live for the whole 5-minute window into the US open. The standdown needs to be a file the writers check, not a variable one of them holds. That is a Saturday item.

The one genuinely reassuring thing here: the router caught it on its very next tick and re-benched it, calling it “housekeeping, not thrash”. That is the behaviour you want from the slow deliberate loop.

5 · The switch file's comment header is still being read as live instruction

switch_notes() splices the top 45 comment lines of gate_switches.env into every router prompt, labelled ACTIVE. Nothing prunes them. Here is what is currently in there:

INSTRUCTION SITTING IN THE SWITCH FILE'S HEADER	AGE AT THIS REPORT	STILL BEING READ TO THE ROUTER?	TICKS THAT RESTATED IT
2026-08-21T14:03Z OPEN-HOUR WATCHER ARMED abs_veto_short on a CONFIRMED	0d 10h	YES — spliced into every prompt as ACTIVE	106 of 1425
2026-08-19T18:15Z OPERATOR BENCH — abs_veto_short OFF.	2d 6h	YES — spliced into every prompt as ACTIVE	66 of 1425
2026-08-13T13:58Z OPEN-HOUR WATCHER ARMED grind_long on a CONFIRMED	8d 11h	YES — spliced into every prompt as ACTIVE	0 of 1425
2026-08-11T14:1xZ — OPERATOR CARVE-OUT (expires 16:00Z, then normal rules resume).	11d 1h	YES — spliced into every prompt as ACTIVE	0 of 1425
2026-08-09T11:0xZ — SATURDAY #1 + MONDAY #1 SHIPPED. abs_veto_short IS ARMED BY DEFAULT.	13d 1h	YES — spliced into every prompt as ACTIVE	0 of 1425

Five instructions, the oldest **thirteen days old**. The 19 August operator bench says “PINNED until 21:55Z” — that deadline passed on Wednesday and the line is still labelled ACTIVE. The 11 August one carries its own expiry (“expires 16:00Z”) and outlived it by eleven days. And the 9 August one is the dangerous one: it states that **abs_veto_short IS ARMED BY DEFAULT**, which directly contradicts the standing operator standdown the router is currently enforcing. A carve-out

comment never expires, and revoking one in the ledger does not revoke it in the file.

★ **THE THREE DEAD ONES GOT ZERO MENTIONS — AND THAT IS NOT A FIX**

Of the 1,425 tick reasons this week, 172 restated one of these five lines. But the split is the whole point: the two LIVE-ish ones (Friday's watcher note and Wednesday's operator bench) account for every single mention, and the three genuinely dead ones — 8, 11 and 13 days old — were restated **zero times**.

A zero-mention streak reads like the problem has gone away. It has not, and no streak ever can show that it has. **1,424 of the 1,425 reasons this week are pinned at exactly the 1,200-character cap** — median length 1200, maximum 1200. The mention rate is measuring how much room was left in the reason field after the live situation was described, and this week the live situation was a standdown, an Asia block and a four-lot carry, all of which are more urgent to narrate than a nine-day-old carve-out.

Read the desk's six previous samples (42/54 → 49/49 → 9/40 → 0/48 → 3/48 → 9/48) and this one the same way: as readings of spare capacity, not as a trend. **The only thing that closes this is deleting the lines.** Any week the reason field is less crowded, a thirteen-day-old 'armed by default' instruction is one tick away from being read as policy again.

6 • ★ **The router has spent 141 consecutive ticks calling a four-lot position flat**

The day rider opened four lots long at 13:13Z on Friday and never closed them. Since that moment the router has ticked 141 times, and **137 of those ticks contain the word “flat”**. It is not lying and it is not broken: it reads `core_health.json`, whose `flat` field is tournament-scoped, and the tournament genuinely is flat. But the sentence a human reads is “Flat, healthy=True, halted=False”, and it is being printed while the desk carries roughly \$235,000 of notional into a weekend gap.

The ticks are not silent about the rider — most of them name its closed rows, — \$727, —\$268, —\$293.50 — and the last one describes it as “clientId 4, a desk I neither route nor bench”. **That is the failure in one sentence.** Citing what CLOSED and declining responsibility for what is OPEN is an attribution that

terminates in the wrong place. The trades table can only ever enumerate closed rows, so it can never contradict a benign story; the venue position line can, and it says four. “Not mine to close” obliges escalation, not dismissal.

7 • The two instruments that should have written this section

ROLLUP	NEWEST FILE	DAYS STALE AT THIS REPORT	WHAT RUNS IT
data/router_nightly/	2026-08-14	8	Nothing. No timer.
data/selector_nightly/	2026-08-18	4	Nothing. No timer.

These are the files the standing weekly review is supposed to roll up: net router value, leakage by regime, selector optimal-pick rate and regret. Neither has a timer, so both simply stopped. **There is no week-in-review to present**, and rather than dress that up, the honest version is the raw-log rebuild above. Two known defects also make the selector rollup wrong when it does run: it has no data_quality filter, so quarantined rows inflate its regret, and all four of its modes reprice at a hard-coded 1.0 stop multiple, which is not what two of the live gates use.

8 • Verdict, and the Saturday actions

Is the router earning its keep? On this week, yes, and for a reason that is not about its thresholds. It made 10 good calls and 3 bad ones, it applied your 18:15Z bench inside 90 seconds, and it was the only component on the box that noticed a gate had been armed through your standdown. Its thresholds were not the lever this week and nothing here argues for touching them. Its clock was not the lever either — there were only 20 writes to be timed. **The router's problem this week is not its judgement, it is that it is surrounded by instruments that give it wrong inputs:** a health flag that says flat when the desk is long, a comment header full of expired instructions, a second writer that can undo a standdown, and two rollups that quietly stopped.

Disposition — every lead this section touched

LEAD	VERDICT	FOR PARKED: WHAT WOULD REVIVE IT. FOR REFUTED: THE TEST THAT KILLED IT.
Make the standdown a FILE that every writer checks	LIVE	It is currently a variable inside router_tick_durable.py. open_hour_watch.py armed a gate through

LEAD

VERDICT

FOR PARKED: WHAT WOULD
REVIVE IT. FOR REFUTED: THE
TEST THAT KILLED IT.

it on Friday and the only reason it cost \$0 is that the tape did not produce a signal in the 2m15s window. Saturday item.

FIVE instructions are spliced into every prompt as ACTIVE — ★ REV2: this row said THREE, which contradicted \$5's own table and the Saturday card, and it understated the exposure. The five, by header timestamp in data/gate_switches.env: 2026-08-21T14:03Z, 2026-08-19T18:15Z, 2026-08-13T13:58Z, 2026-08-11T14:1xZ and 2026-08-09T11:0xZ. Four of the five are past their own stated expiry or their situation, and the oldest is the dangerous one: the 08-09 line asserts that abs_veto_short IS ARMED BY DEFAULT, which directly contradicts the standing operator standdown the router is currently enforcing — on the same gate that open_hour_watch.py armed through that standdown on Friday. Nothing prunes them and no mention-rate streak closes it; only deleting the lines does.

Delete the expired carve-outs from gate_switches.env's header

LIVE

Point core_health.flat at the venue position

LIVE

137 of the last 141 router ticks said 'flat' while the venue was long four lots. Either make the field desk-wide or rename it tournament_flat.

Change the router's thresholds or timing this week

REFUTED

There were 20 switch writes all week, 5 of them worth \$0 by construction. There is no sample here to tune a threshold on, and tuning on this week would be fitting to a standdown.

Re-arm a gate on a lapse rather than a positive case

REFUTED

Named instance: 2026-08-19 15:35Z re-armed abs_veto_short 30 minutes after benching it, with no new evidence, straight into three stops for −\$133. Third occasion on this gate.

Give router_watch narrow fast-BENCH authority

PARKED

The asymmetry argument still holds — a wrong bench is cheap, a wrong arm is expensive — but this week produced no measurable actuation-lag cost to justify the build, because the desk was stood down. Revive it with a lag number from a week the tournament actually trades.

Put a timer on router_nightly and selector_nightly

LIVE

Both stopped silently. A monitoring file that stops being written is worse than one that errors, because its absence looks like a quiet week.

Fix selector_nightly's two known defects before believing it again

PARKED

No data_quality filter (quarantined rows inflate regret) and a hard-coded 1.0 stop multiple for all four modes. Revive when both are fixed; until then do not quote its regret.

M1

The biggest runs on MNQ this week — did we show up?

Forget what the desk did this week — start from the raw tape. MNQ printed **70 runs of 1.5×ATR or bigger** between **2026-08-16 22:00** and **2026-08-21 20:59 UTC** (a 15-minute move of at least 55 points, against a typical 15-min range of 36). That is the whole tape, not a top-ten. We **caught 4**, we were positioned against **1**, and we **sat out 65**. The runs we aligned with banked **+\$618** of honest money — but that is only 60% of their \$1,038 hindsight ceiling, and the 65 we sat out left **\$9,986** on the table. The bleed is not the runs; it is the fading we do around them.

THE FULL CENSUS

Every qualifying run, in order. **Move** is the swing in points; **\$ 1lot** is that move on one MNQ lot (\$2/pt) — the hindsight ceiling. **real \$** is what a gate actually banked near it. **Flow** is net aggressor volume in the 60s before; **amp** is pre-run amplitude; **book** is the far-side L2 depth share (below 0.50 = the side price ran toward was thin), read from `capture.db.book`, 3 deep, 30s. Green = we caught it, red = we fought it.

TIME (UTC)	DIR	MOVE	\$ 1LOT	US	GATE	REAL \$	FLOW	AMP	BOOK	CLUSTER
08-17 07:01	DN	-56	113	sat out	—		+16	0.04	0.53	UNCLASS
08-17 07:27	UP	+57	114	sat out	—		-216	0.05	0.42	VACUUM
08-17 11:50	DN	-55	110	sat out	—		+113	0.02	0.60	UNCLASS
08-17 12:24	DN	-76	152	sat out	—		+23	0.04	0.57	UNCLASS
08-17 13:53	DN	-70	139	sat out	—		-583	0.20	0.55	FLOW-LED
08-17 14:53	UP	+68	136	sat out	—		-326	0.09	0.43	OPEN/NEWS
08-17 15:28	DN	-57	114	sat out	—		+216	0.07	0.60	UNCLASS
08-17 17:09	DN	-59	118	sat out	—		-1564	0.07	0.54	FLOW-LED
08-18 01:38	DN	-101	202	sat out	—		+458	0.07	0.49	VACUUM

08-18 02:10	DN	-72	144	sat out	—		+17	0.08	0.49	UNCLASS
08-18 02:40	DN	-57	114	sat out	—		+173	0.11	0.49	UNCLASS
08-18 05:52	DN	-82	163	sat out	—		+2	0.07	0.48	UNCLASS
08-18 06:08	UP	+61	122	sat out	—		-39	0.07	0.55	UNCLASS
08-18 07:52	DN	-95	190	sat out	—		-25	0.07	0.52	UNCLASS
08-18 12:22	DN	-65	130	sat out	—		-119	0.12	0.53	UNCLASS
08-18 13:25	UP	+73	146	sat out	—		+261	0.09	0.51	FLOW-LED
08-18 13:40	DN	-138	276	caught	exhaustio n_short_A	-208	+485	0.29	0.49	VACUUM
08-18 13:56	DN	-70	140	caught	day_rider	+32	+1850	0.21	0.34	VACUUM
08-18 14:11	UP	+100	200	sat out	—		-1738	0.32	0.52	VACUUM
08-18 14:39	DN	-103	206	caught	exhaustio n_short_A	+250	-33	0.12	0.50	OPEN/ NEWS
08-18 15:19	UP	+56	112	sat out	—		-158	0.08	0.49	UNCLASS
08-18 15:39	UP	+63	126	sat out	—		-52	0.09	0.47	UNCLASS
08-18 23:58	DN	-75	150	sat out	—		-34	0.08	0.61	UNCLASS
08-19 00:16	UP	+92	184	sat out	—		-90	0.09	0.54	UNCLASS
08-19 05:04	DN	-60	121	sat out	—		-106	0.13	0.47	UNCLASS
08-19 06:10	UP	+63	126	sat out	—		+136	0.08	0.52	FLOW-LED
08-19 06:43	UP	+86	173	sat out	—		-68	0.08	0.55	UNCLASS
08-19 07:54	UP	+61	122	sat out	—		+76	0.05	0.50	UNCLASS
08-19 12:32	UP	+150	300	sat out	—		-131	0.12	0.56	VACUUM
08-19 13:21	DN	-86	173	sat out	—		+8	0.09	0.52	OPEN/NEWS
08-19 13:36	DN	-207	414	caught	day_rider	+544	-1719	0.34	0.51	FLOW-LED
08-19 13:59	DN	-108	215	sat out	—		+741	0.23	0.52	VACUUM
08-19 14:25	UP	+159	318	sat out	—		-427	0.24	0.51	OPEN/NEWS
08-19 14:42	UP	+80	160	sat out	—		-157	0.15	0.47	OPEN/NEWS
08-19 15:55	UP	+74	147	sat out	—		-494	0.09	0.51	UNCLASS
08-19 16:21	DN	-62	125	sat out	—		-133	0.12	0.52	UNCLASS
08-19 16:42	DN	-57	114	sat out	—		-182	0.08	0.52	UNCLASS

08-19 18:03	DN	-89	178	sat out	—		+552	0.19	0.54	VACUUM
08-19 18:18	UP	+62	124	sat out	—		-337	0.17	0.52	UNCLASS
08-19 19:34	DN	-70	140	sat out	—		+118	0.13	0.51	UNCLASS
08-19 20:49	UP	+78	155	sat out	—		+5	0.04	0.50	UNCLASS
08-19 23:47	DN	-57	114	sat out	—		-110	0.05	0.53	UNCLASS
08-20 00:07	UP	+66	133	sat out	—		-14	0.11	0.47	UNCLASS
08-20 02:06	DN	-60	120	sat out	—		-5	0.09	0.51	UNCLASS
08-20 05:54	UP	+56	112	sat out	—		-128	0.05	0.44	VACUUM
08-20 07:50	DN	-71	142	sat out	—		-40	0.10	0.51	UNCLASS
08-20 08:17	DN	-70	140	sat out	—		+7	0.07	0.49	UNCLASS
08-20 08:55	UP	+65	130	sat out	—		-61	0.09	0.54	UNCLASS
08-20 10:58	DN	-81	162	sat out	—		+4	0.05	0.51	UNCLASS
08-20 11:16	DN	-55	110	sat out	—		+16	0.12	0.49	UNCLASS
08-20 11:42	DN	-87	174	sat out	—		-146	0.13	0.50	FLOW-LED
08-20 12:02	DN	-91	182	sat out	—		-26	0.10	0.52	UNCLASS
08-20 12:39	UP	+88	175	sat out	—		-137	0.11	0.48	UNCLASS
08-20 13:20	UP	+128	256	sat out	—		+61	0.08	0.50	OPEN/NEWS
08-20 13:39	DN	-112	224	FOUGHT	day_rider	-1288	+56	0.34	0.50	VOL- EXPANSION
08-20 13:58	UP	+56	112	sat out	—		+352	0.29	0.51	OPEN/NEWS
08-20 14:27	UP	+76	152	sat out	—		+59	0.25	0.52	OPEN/NEWS
08-20 14:42	DN	-66	132	sat out	—		+306	0.18	0.50	OPEN/NEWS
08-20 15:02	DN	-80	161	sat out	—		+1542	0.37	0.49	VACUUM
08-20 15:29	DN	-63	126	sat out	—		+200	0.16	0.49	UNCLASS
08-20 16:32	UP	+71	142	sat out	—		+170	0.13	0.50	UNCLASS
08-20 17:53	DN	-84	168	sat out	—		-211	0.08	0.52	UNCLASS
08-20 23:59	UP	+68	135	sat out	—		+22	0.05	0.48	UNCLASS
08-21 13:10	DN	-62	123	sat out	—		-283	0.07	0.47	FLOW-LED
08-21 13:29	DN	-158	315	sat out	—		-99	0.23	0.48	OPEN/NEWS

08-21 13:47	DN	-58	116	sat out	—	+233	0.18	0.49	OPEN/NEWS
08-21 14:06	DN	-94	188	sat out	—	+210	0.21	0.52	OPEN/NEWS
08-21 14:42	UP	+56	113	sat out	—	-819	0.14	0.49	VACUUM
08-21 15:26	UP	+115	230	sat out	—	+165	0.10	0.51	UNCLASS
08-21 15:55	DN	-77	154	sat out	—	+20	0.16	0.51	UNCLASS

What the runs were — by cause

Each run gets a cause fingerprint from its pre-run tape. This is the map for where to hunt a new gate.

CLUSTER	RUNS	WHAT IT IS
UNCLASS	39	no clear tape tell (quiet ignition)
OPEN/NEWS	12	the 13:00–15:00 UTC cash-open / data window
VACUUM	11	moved AGAINST the tape (a snap-back / stop-run)
FLOW-LED	7	aggressors drove it (a continuation)
VOL-EXPANSION	1	amplitude broke out of a coil

HONEST MONEY — CEILING VS WHAT WE BANKED

RUNS	HINDSIGHT CEILING	WHAT WE REALLY MADE
Caught (aligned)	4	\$1,038 +\$618 (60%)
Fought (against)	1	\$224 \$-1288
Sat out	65	\$9,986 \$0 ← the money on the table

The runs made money and the fading gave it back; we convert a fraction of the ceiling and ignore the rest. The next two movements ask whether we can do better — first with the gates we own, then with new ones.

THE SAME CENSUS ON GOLD — MGC, the second instrument

M1 The biggest runs on MGC this week — did we show up?

Forget what the desk did this week — start from the raw tape. MGC printed **68 runs of 1.5×ATR or bigger** between **2026-08-16 22:00** and **2026-08-21 20:59 UTC** (a 15-minute move of at least 11 points, against a typical 15-min range of 7). That is the whole tape, not a top-ten. We **caught 0**, we were positioned against **0**, and we **sat out 68**. **We have no gate on MGC**, so all 68 runs and **\$11,111** of hindsight ceiling were sat out by construction, not by a decision that went wrong. Nothing here is a bleed — it is an empty seat, and the question the next movements ask is what could sit in it.

THE FULL CENSUS

Every qualifying run, in order. **Move** is the swing in points; **\$ 1lot** is that move on one MGC lot (\$10/pt) — the hindsight ceiling. **real \$** is what a gate actually banked near it. **Flow** is net aggressor volume in the 60s before; **amp** is pre-run amplitude; **book** is the far-side L2 depth share (below 0.50 = the side price ran toward was thin), read from `depth.db.depth_snap`, 5 deep, 60s. Green = we caught it, red = we fought it.

TIME (UTC)	DIR	MOVE	\$ 1LOT	US	GATE	REAL \$	FLOW	AMP	BOOK	CLUSTER
08-16 23:22	DN	-12	121	sat out	—		+0	0.07	0.50	UNCLASS
08-17 00:06	UP	+15	149	sat out	—		-12	0.18	0.48	UNCLASS
08-17 00:52	UP	+14	141	sat out	—		+0	0.06	0.37	UNCLASS
08-17 01:24	DN	-14	145	sat out	—		+22	0.20	0.39	VACUUM
08-17 05:15	UP	+11	112	sat out	—		-10	0.11	0.46	UNCLASS
08-17 07:22	UP	+13	133	sat out	—		+1	0.07	0.45	UNCLASS

08-17 12:34	DN	-11	111	sat out	—	+0	0.12	0.58	UNCLASS
08-17 13:42	UP	+26	262	sat out	—	-35	0.16	0.54	VACUUM
08-17 14:10	UP	+22	225	sat out	—	-6	0.10	0.41	OPEN/NEWS
08-17 18:19	DN	-17	170	sat out	—	-8	0.08	0.59	UNCLASS
08-17 19:06	UP	+15	151	sat out	—	-9	0.11	0.45	UNCLASS
08-18 00:13	UP	+11	107	sat out	—	-6	0.08	0.46	UNCLASS
08-18 01:03	DN	-12	115	sat out	—	-8	0.16	0.55	UNCLASS
08-18 01:33	DN	-17	173	sat out	—	+1	0.09	0.48	UNCLASS
08-18 13:15	DN	-11	112	sat out	—	+0	0.13	0.36	OPEN/NEWS
08-18 13:30	UP	+18	181	sat out	—	-283	0.40	0.51	VACUUM
08-18 13:48	UP	+12	115	sat out	—	+3	0.26	0.57	OPEN/NEWS
08-18 14:28	DN	-16	158	sat out	—	+13	0.13	0.51	OPEN/NEWS
08-18 14:47	DN	-17	171	sat out	—	-10	0.17	0.48	OPEN/NEWS
08-18 15:06	DN	-22	218	sat out	—	+0	0.23	0.41	UNCLASS
08-18 19:42	DN	-16	160	sat out	—	+1	0.10	0.33	UNCLASS
08-18 23:59	UP	+11	112	sat out	—	+10	0.06	0.58	UNCLASS
08-19 00:47	DN	-16	160	sat out	—	+2	0.11	0.49	UNCLASS
08-19 01:25	UP	+19	192	sat out	—	-5	0.16	0.63	UNCLASS
08-19 05:01	DN	-12	116	sat out	—	+49	0.14	0.47	VACUUM
08-19 06:08	UP	+11	111	sat out	—	-1	0.09	0.59	UNCLASS
08-19 07:21	UP	+14	142	sat out	—	+2	0.08	0.53	UNCLASS
08-19 12:29	UP	+50	502	sat out	—	-1	0.08	0.22	UNCLASS
08-19 12:46	UP	+25	246	sat out	—	+67	0.36	0.51	FLOW-LED
08-19 13:04	UP	+34	335	sat out	—	+4	0.33	0.40	VOL- EXPANSION
08-19 13:33	DN	-13	132	sat out	—	-3	0.24	0.50	OPEN/NEWS
08-19 13:56	UP	+20	198	sat out	—	+22	0.25	0.53	OPEN/NEWS
08-19 14:12	UP	+17	171	sat out	—	-7	0.15	0.38	OPEN/NEWS
08-19 14:28	UP	+12	119	sat out	—	+2	0.27	0.51	OPEN/NEWS

08-19 15:31	UP	+11	111	sat out	—	-62	0.11	0.56	VACUUM
08-19 18:13	DN	-16	163	sat out	—	-1	0.10	0.69	UNCLASS
08-19 18:28	UP	+18	175	sat out	—	-19	0.21	0.47	VACUUM
08-19 20:13	UP	+14	144	sat out	—	-5	0.10	0.44	UNCLASS
08-20 00:50	DN	-17	167	sat out	—	+6	0.09	0.58	UNCLASS
08-20 01:58	UP	+12	123	sat out	—	-4	0.08	0.49	UNCLASS
08-20 05:56	UP	+13	134	sat out	—	+1	0.06	0.35	UNCLASS
08-20 11:33	DN	-16	163	sat out	—	+1	0.05	0.53	UNCLASS
08-20 11:51	DN	-20	196	sat out	—	-18	0.16	0.47	FLOW-LED
08-20 12:19	DN	-14	143	sat out	—	+5	0.10	0.65	UNCLASS
08-20 12:34	UP	+18	183	sat out	—	-25	0.21	0.64	VACUUM
08-20 13:29	UP	+15	152	sat out	—	+6	0.10	0.61	OPEN/NEWS
08-20 14:19	UP	+13	126	sat out	—	-9	0.18	0.49	OPEN/NEWS
08-20 14:48	UP	+30	305	sat out	—	+12	0.16	0.43	OPEN/NEWS
08-20 15:29	DN	-21	209	sat out	—	+5	0.18	0.72	UNCLASS
08-20 16:01	DN	-15	149	sat out	—	-1	0.09	0.58	UNCLASS
08-20 18:17	UP	+11	111	sat out	—	+8	0.09	0.65	UNCLASS
08-21 01:02	UP	+13	126	sat out	—	+4	0.10	0.31	UNCLASS
08-21 01:20	UP	+15	148	sat out	—	+17	0.14	0.62	FLOW-LED
08-21 01:58	DN	-12	123	sat out	—	+3	0.09	0.34	UNCLASS
08-21 02:37	UP	+15	148	sat out	—	-1	0.13	0.60	UNCLASS
08-21 05:41	UP	+16	156	sat out	—	+5	0.12	0.59	UNCLASS
08-21 06:24	UP	+15	147	sat out	—	-6	0.14	0.47	UNCLASS
08-21 08:09	UP	+18	179	sat out	—	-5	0.06	0.55	UNCLASS
08-21 09:01	UP	+17	169	sat out	—	-3	0.10	0.47	UNCLASS
08-21 09:16	DN	-17	174	sat out	—	+25	0.12	0.48	VACUUM
08-21 12:06	DN	-13	131	sat out	—	-17	0.17	0.49	FLOW-LED
08-21 12:24	DN	-19	190	sat out	—	+7	0.16	0.33	UNCLASS

08-21 12:39	UP	+15	149	sat out	—	-7	0.27	0.52	UNCLASS
08-21 13:35	UP	+14	144	sat out	—	-5	0.29	0.51	OPEN/NEWS
08-21 14:15	UP	+14	145	sat out	—	-13	0.10	0.31	OPEN/NEWS
08-21 14:30	DN	-21	212	sat out	—	-2	0.11	0.58	OPEN/NEWS
08-21 15:21	UP	+11	113	sat out	—	+2	0.08	0.49	UNCLASS
08-21 17:11	DN	-14	137	sat out	—	-52	0.10	0.42	FLOW-LED

★ REV2 — the book column, and the first thing it says

In Rev1 this column was — **on all 68 gold rows**. That was not missing data: the census read capture.db.book, which is **MNQ-only** (IBKR allows three depth subscriptions and gold's went elsewhere), while MGC's L2 sits in depth.db.depth_snap — **918,431 MGC snapshots inside this window**. The scope's mandated per-run L2 read was simply absent for the instrument the desk is hunting gates for. It is now read from depth.db, 5 deep over the 60 seconds before each run. △ The two instruments' book columns are **not the same measurement**: MNQ is 3-deep over 30s from a 41ms event stream, MGC is 5-deep over 60s from a 250ms sample. Same sign convention — below 0.50 means the side price ran toward was thin — but do not compare the numbers across instruments.

WHAT THE GOLD BOOK SAYS	VALUE	READ
Runs with a book read	68 / 68	the column is complete; nothing is imputed
Median far-side share	0.49	a coin, across the whole census
Runs where the far side was thin (<0.50)	35 / 68 = 51%	also a coin — as a direction tell it is worth nothing
Median on the 9 runs with a ceiling ≥ \$200	0.430	against 0.490 on the other 59 — the big ones look thinner
Correlation (far-side share, ceiling \$)	-0.262	n=68, permutation p = 0.020 over 20,000 shuffles
... with the single biggest run removed	-0.039	the relationship was that one run
... with the top three removed	+0.086	it flips sign
Spearman rank correlation (outlier-proof)	-0.033	essentially zero

Verdict: the gold book column is populated and it says nothing yet.

The apparent size relationship is one run — 08-19 12:29, +50pt for a \$502 ceiling, on the thinnest far-side read in the census (0.22). Strip it and the

correlation goes to -0.039 ; strip the top three and it turns positive; the rank correlation, which cannot be moved by one outlier, is -0.033 . **Cause of death: strip-the-best, and the rank test.** That is the honest state of the one genuinely new input the gold programme has, and it is worth having on the record precisely so nobody re-derives the -0.262 next Friday and calls it an edge.

It does rhyme with Movement 3’s MNQ finding — that size is detectable and direction is not — but a rhyme on $n=68$ with one dominant point is a reason to keep the column, not a result. The way to make it one is more gold tape: `depth_snap` has held MGC since 2026-07-26, so the next census can run this on four times the runs without waiting for anything.

What the runs were — by cause

Each run gets a cause fingerprint from its pre-run tape. This is the map for where to hunt a new gate.

CLUSTER	RUNS	WHAT IT IS
UNCLASS	39	no clear tape tell (quiet ignition)
OPEN/NEWS	15	the 13:00–15:00 UTC cash-open / data window
VACUUM	8	moved AGAINST the tape (a snap-back / stop-run)
FLOW-LED	5	aggressors drove it (a continuation)
VOL-EXPANSION	1	amplitude broke out of a coil

HONEST MONEY — CEILING VS WHAT WE BANKED

	RUNS	HINDSIGHT CEILING	WHAT WE REALLY MADE
Caught (aligned)	0	\$0	+\$0
Fought (against)	0	\$0	+\$0
Sat out	68	\$11,111	\$0 ← the money on the table

Every row is a sat-out row because there is no gate on this instrument to catch or fight anything. The ceiling is what one lot would have made on a

perfect entry and a perfect exit, so treat it as the size of the seat, never as money forgone. The next two movements ask what could fill it.

M2 THE IDLE-GATE LAB — could the gates we already own have caught them?

Garrath — Movement 1 found 65 big runs this week that the desk sat out completely: 4,994 points of movement, worth \$9,987.00 if they had been caught perfectly. The obvious question is whether the six gates we already own would have caught any of them if they had simply been left armed.

The answer is a clean and rather uncomfortable no. Fired mechanically into those runs, in their exact live configuration, they lose \$1,333.85. Fired across the whole week rather than only near the runs, they lose \$1,905.36. **Not arming them was worth more than any entry idea in this report.**

1 • The method, and the one thing that makes it honest

Each of the 65 sat-out runs gets four independent simulations: the six gates fired mechanically in the **10 minutes before** the run ignited (which is the interesting one — catching a run means being in before it goes), and in the **5 minutes after**; each of those both in the exact live config and **ungated**, meaning every ATR floor, veto and confirmation stripped out. Exits are the live scaleout slate plus `exit_overrides.json`, raced on real ticks at \$1.50 per round trip and \$2.00 a point.

The thing that makes it honest is the **base-rate control**. Measuring a gate only in windows we already know contained a big run is hindsight selection — it counts the fires where the move existed and never counts the ones where it did not. So the same six gates are also fired across the entire captured week, every fire counted. That control is the reason this section can say something the last three idle-gate labs could not.

2 · The four arms

ARM	RUNS IT SHOWED UP FOR	FIRES	NET	\$ / FIRE	WIN %	BEST FIRE STRIPPED
10 min BEFORE ignition · live config	18 / 65	23	−\$1,333.85	−\$57.99	13%	−\$1,684.85
10 min BEFORE · every suppressor removed	33 / 65	60	−\$1,620.40	−\$27.01	23%	−\$1,971.40
5 min AFTER ignition · live config	13 / 65	17	−\$45.00	−\$2.65	41%	−\$317.00
5 min AFTER · ungated	24 / 65	49	−\$626.75	−\$12.79	31%	−\$898.75

Every arm loses. Read the first two rows together: with their live filters on, the gates show up for 18 of the 65 runs and lose \$1,333.85; strip every filter off and they show up for 33 runs, take 60 fires, and lose \$1,620.40. **Removing the suppressors triples the activity and makes the result worse**, which is the strongest defence of the gates' own filters this desk has produced.

★ **REV2 — THE NEARLY-FLAT ROW IS ONE GATE.** Split the post-ignition live arm by gate and `abs_veto_short` takes 5 fires for **+\$504.50 at a 100% win rate** (+\$97.00, +\$97.00, +\$97.00, +\$102.50, +\$111.00, on five distinct runs) while the other five gates take 12 fires for **−\$549.50**. Strip `abs_veto_short`'s single best fire and the whole arm goes from **−\$45.00 to −\$156.00**. So “joining a run is roughly free” is not a property of the roster; it is one gate on five fires, and every other gate is as bad after ignition as before. **BUILD #6’s $n \geq 60$ kill criterion has to be PER GATE, not pooled** — pooled, five lucky `abs_veto_short` fires would carry a shadow arm that five other gates are losing money in. The original observation, kept because the asymmetry is real for that one gate: fired after ignition with live filters on, the gates lose only \$45.00 across 17 fires. Joining a run once it has already started is roughly free; predicting one costs money. That is the same shape the greenfield lab reports from a completely different direction in Movement 3.

3 · Gate by gate, with the base rate beside it

GATE	BEFORE IGNITION	AFTER IGNITION	WHOLE WEEK (BASE RATE)	\$ / FIRE, BASE RATE	RAW SIGNALS	WHAT ITS OWN FILTERS BLOCKED
grind_long	3 · \$86.00	6 · −\$254.92	11 · −\$319.17	−\$29.02	24,121	ATR floor 22pt ×24,110
capitulation_long	2 · −\$113.66	1 · \$97.00	10 · −\$97.66	−\$9.77	132	ATR floor 10pt ×122
abs_veto_long	3 · −\$169.58	3 · −\$173.36	26 · −\$83.42	−\$3.21	38	55s veto: absorbed ×5; 55s veto: thrust gone ×7
rgv_short	0 · \$0.00	1 · −\$17.54	1 · −\$17.54	−\$17.54	1	—
exhaustion_short	7 · −\$612.42	1 · −\$200.68	18 · −\$638.02	−\$35.45	46	counter-regime veto ×26; 5s confirm: went adverse ×2
abs_veto_short	8 · −\$524.19	5 · \$504.50	51 · −\$749.55	−\$14.70	67	55s veto: absorbed ×8; 55s veto: thrust gone ×8

Whole week, all six gates, every fire counted: **−\$1,905.36 on 117 fires**. That is the number to quote when anyone asks what the standdown was worth. And look at the raw-signal column, because it explains the shape of this entire report: grind_long generated 24,121 raw signals this week and its own 22-point ATR floor blocked 24,110 of them. The gate is not benched by anybody's decision; it is benched by its own arithmetic.

4 · Which kinds of run, and did direction even matter?

CAUSE CLUSTER	RUNS	POINTS IN THEM	CEILING IF PERFECTLY CAUGHT	MECHANICAL FIRES	WHAT THOSE FIRES MADE
UNCLASS		39	2,754 pt	\$5,508.00	13 −\$782.45
OPEN/NEWS		11	1,029 pt	\$2,058.00	4 −\$266.40
VACUUM		9	797 pt	\$1,595.00	6 −\$285.00
FLOW-LED		6	414 pt	\$826.00	0 \$0.00

THE MECHANICAL FIRE WAS...	FIRES	NET	\$ / FIRE	WIN %
pointing the SAME way the run then went		11	−\$602.60	−\$54.78 9%
pointing the OPPOSITE way		12	−\$731.25	−\$60.94 17%

This is the part that closes the “just leave them armed” argument for good. The mechanical fires that happened to point the same way the run then went made – \$602.60 across 11 fires. **Being right about the direction was not enough.** The gates enter on their own triggers, which fire during the chop before the move rather than at the ignition, so an aligned fire still gets stopped out in the noise and is not there when the run happens.

★ **THE SURVIVORSHIP TRAP THIS LAB EXISTS TO AVOID — AND IT CUTS BOTH WAYS**

Last cycle's idle-gate lab reported a small positive and it was survivorship: the gates looked good because of which ones happened to be idle in the windows chosen. This week the base-rate control was built specifically to close that, and it does — but note that it closes the flattering reading AND the damning one.

In the run windows the gates make **−\$57.99 a fire**. Across the whole week they make **−\$16.29 a fire**. ★ **REV2 — Rev1 called those “close” and concluded there is nothing special about the tape near a big run. That is the one reading the numbers do not support, and it is wrong in the direction that makes the bench case STRONGER.** The gap is **3.56×** and the run windows are the worse side. They are also not independent samples: **21 of the 23 run-window fires sit inside the 117-fire base-rate set** (same gate, same second), so the base rate mostly contains the thing it is being compared against. Remove them and the rest of the week is **−\$4.49 a fire on 96 fires** — a **12.9×** gap.

Bootstrapped over 20,000 resamples, the difference in \$/fire is **−\$41.71, 95% CI [−\$72.15, −\$2.52]** against the full base rate (98.0% of resamples put the run windows worse), and **95% CI [−\$82.41, −\$15.44]** against the disjoint remainder (99.6%). Both intervals exclude zero. So the sentence Rev1 explicitly ruled out is the true one: **these gates are substantially worse in the ten minutes before a big run than on ordinary tape.** They fire into the chop that precedes an ignition and get stopped in it — which is exactly what §4's direction table says from the other side. That is not a blur, it is the damning case, and it is the strongest form of the argument for leaving them benched.

The honest limit on the headline, stated plainly: the **−\$1,905.36** whole-week total has a bootstrap 95% CI on its per-fire mean of **[−\$35.30, +\$3.94]**, which scaled to 117 fires is roughly **[−\$4,130, +\$461]** — it straddles zero. The TOTAL is confidently negative only as a point estimate; what is statistically solid here is the comparison (run windows worse than base rate, both CIs excluding zero), not the size of the prize. Quote **−\$1,905.36** as the point estimate it is, never as a measured saving. So the conclusion is narrow and should not be over-read: these six gates, at these settings, on THIS week's tape, do not contain a run-catcher. It is one week, the desk was stood down for most of it, and a genuine trend week could look different. What it

does license is leaving them benched for now, which is what is already happening.

5 • What this section does NOT say

It does not say the \$9,987.00 in those 65 runs is unreachable — Movement 3 takes a different route at it and finds a size threshold that pays. It does not say the gates are worthless; two of them traded live this week and one of them was green. And it emphatically does not say “arm nothing forever”: it says that this particular six-gate roster, fired mechanically, is not the tool for catching runs, and that the tool will have to be built rather than switched on.

Disposition — every lead this section touched

LEAD	VERDICT	FOR PARKED: WHAT WOULD REVIVE IT. FOR REFUTED: THE TEST THAT KILLED IT.
Leave the six live gates armed to catch big runs	REFUTED	Named test: 65 sat-out runs, four arms, plus a whole-week base-rate control. Every arm negative; live config –\$1,333.85 in the run windows and –\$1,905.36 across the week.
Strip the gates' filters so they show up more often	REFUTED	Ungated triples the fire count and makes the result worse in both windows. The ATR floors and vetoes are earning their keep.
Join a run AFTER ignition rather than predicting it	SHADOW	The post-ignition live arm is nearly flat at –\$45.00 on 17 fires, against a clearly negative pre-ignition arm. That is the only asymmetry in the whole lab. Shadow an ignition-confirmed entry and let it accumulate n. ★ REV2: per-gate. The –\$45.00 is abs_veto_short's +\$504.50 on 5 fires against –\$549.50 on the other five gates' 12; strip one fire and the arm is –\$156.00. Score and kill each gate on its own n, never on the pool.
Direction alignment as a permission rule	REFUTED	Aligned mechanical fires made –\$602.60 on 11 fires. Knowing which way the run would go was not enough, because the gates fire before the ignition and get stopped in the chop.
The 22-point ATR floor on grind_long	PARKED	It blocked 24,110 of 24,121 raw signals this week. Revive with the floor ladder proposed in Part 2.5 — this lab cannot say whether 22 is right, only that it binds almost everything.

★ GAP — FOUR OF MOVEMENT 3'S PER-CAUSE-CLUSTER HUNTS DID NOT RUN — THE REST OF IT DID, TONIGHT

The FLOW-LED, OPEN-NEWS, UNCLASS and RIDER_ALL cluster hunts did not run this cycle: gf_full_FLOW-LED.md, gf_full_OPEN-NEWS.md, gf_full_UNCLASS.md and gf_full_RIDER_ALL.md on disk are all dated 15 August, i.e. LAST week's artifacts, and none of them is stitched into this document. **Everything else in this movement is fresh and ran 2026-08-22 00:00-00:39** — the VACUUM cluster hunt, the 686-run lake study, the three forward tests, BIGBOARD, SMALLFADE, LOTDIV and the chop lab. ★ REV2: Rev1's front page said "Movement 3's four MNQ hunts never started", which reads as you got no MNQ greenfield work this week. That is false. Four cluster hunts are missing from a movement that otherwise delivered.

M3 The greenfield lab — we can see the runs coming, we still cannot tell which way they go

Movement 1 froze the tape: **70 big runs on MNQ this week, 65 of them sat out, carrying \$9,986 of one-lot hindsight ceiling.** This movement ignores every gate we own and asks whether a brand-new entry could have taken any of it. The answer this week is a **NULL on MNQ and one thin survivor on gold** — but it is a null with three hard numbers behind it, and they say something different from "nothing works". **First: the money is real and the exit is not the problem.** Board all 65 sat-out runs at their own start, in the right direction, with a purely mechanical $3 \times \text{ATR}$ stop and $6 \times \text{ATR}$ target, and it banks **+\$4,009 over 45 trades at +\$89.09 each, 84.4% wins — 40.1% of the entire ceiling, converted by the exit alone.** **Second: we can board them.** A dumb causal detector — "the last ten minutes covered two ATRs one way" — boards **50 of the 65 in the right direction, a median six minutes in, with a median 42 points still on the table.** **Third, and this is the whole problem: the same detector fires 155 times in the same five days, and only 50 of those fires are runs.** The 50 that are runs made +\$483. The other 105 lost —

\$1,651. **Detection is solved. Selection is not.** And the size threshold the mandate asked us to find is real, sharp and sits at **120 points** — below it, boarding a run loses money at every band; at and above it, boarding pays **+\$69.38 a trade over 51 trades and survives strip-3, strip-5, leave-one-day-out and long/short symmetry**. It is a hindsight band, and closing that gap is next week's entire job.

LEAD WITH WHAT SURVIVED

Five things walked out of this movement alive. Everything else is in the graves table, by name, with the test that killed it.

1. ★ **THE SIZE THRESHOLD IS 120 POINTS, and it is the deliverable the mandate asked for.** Rebuilding the census picker over the whole lake gives **686 runs on 33 trading days** — real n, not fifteen. Board each one at its own first fire and the answer is not gradual, it is a step: **55-70pt loses \$12.10 a trade, 70-90pt loses \$6.35, 90-120pt loses \$9.48 — and then 120-160pt makes + \$54.04 and 160pt+ makes +\$86.64.** Pooled at 120pt+: **n=51, + \$3,538, +\$69.38 a trade, 52.9% wins**, strip-3 still +\$1,940, strip-5 still +\$1,051, worst leave-one-day-out +\$2,357, and both sides pay (long +\$74.82, short +\$64.91). That is the most robust thing this movement has produced in three Fridays.
2. **The exit is exonerated.** The desk keeps re-litigating exits. On this population the exit does 40.1% of the ceiling on its own. Whatever is wrong, it is not the stop and it is not the target.
3. ★ **The one stone last Friday named for "next week" — the large-lot / small-lot aggressor split — is TURNED, and it is not empty.** It is a null on direction like everything else (large-lot flow agrees with the next 15 minutes **48.71%** of the time on unselected big moves, against price's 48.66% and L2's 49.44%) — but **SMALL**-lot flow is inversely informative and it strengthens monotonically with size: **49.56% on all minutes → 48.50% at 30pt → 46.57% at 55pt → 44.31% at 90pt → 42.95% at 120pt.** Fading the retail tape reads 57.05% on the biggest moves. And **|large-lot flow| detects size better than anything VACUUM found: AUC 0.721 (p<0.0001) on the 120pt+ band,**

against volume-z's 0.690. Both findings are real. Neither trades — see the graves.

- 4. **One gold cell survived its forward test.** Last week's mgc_hole_break_fade double-hole refinement made **+\$159 over 27 trades (+\$5.88 each) on the five days it had never seen**, green on 3 of 5. Everything else in gold died, including its own simpler primary arm. It stays SHADOW on thin n.
- 5. **Every candidate that survived last Friday on MNQ is now dead, and they were killed by the same test.** BREAKFLOW, the pooled rider mechanism and VAC-ABS all went forward onto 08-17-08-21 and all three lost. That is not three coincidences; it is the selection problem showing up on schedule.

The ruler — stated before any number, because two of last week's headlines died on it

COMPONENT	RULE
Census	The frozen Movement-1 snapshot, read from the locked stdout — 70 runs, 65 sat out, ceiling \$9,986 (the printed per-run figures sum to \$9,987; that is a dollar of rounding, not a discrepancy). Never re-derived.
Tape	The Parquet lake , never a bare capture.db ATTACH. 42,143 tick-derived minute bars, 2026-07-05 → 2026-08-21 , sitting on 43,186,167 MNQ trade prints . Gold adds the 1-minute backfill: 357,692 bars, 2025-07-27 → 2026-08-18, 184 clean days after two defects were cut out.
Bars	Built from the trade ticks , so price and the aggressor split are the same object. ATR is ATR(1m,14) over contiguous minutes only.
Entry	The first tick at or after the trigger minute's close . No same-bar fill, no lookahead. Every stop and target is walked tick by tick, and the stop is checked before the target on every tick .
Costs	\$1.50 per round trip — the true venue fee, not \$5, not \$2, not per side — plus 0.25pt slippage each way . MNQ at \$2.00/point . Gold at \$10.00/point and a \$7.50 all-in round trip, because gold crosses a ~0.30pt spread on both legs.
★ Accounting	ONE POSITION AT A TIME , everywhere, without exception. This is not cosmetic: VACUUM's first pass let overlapping entries run concurrently and five separate shorts booked the same 2026-07-31 impulse, worth \$3.9k of a \$10.3k headline . Under honest serial accounting that candidate fell from +\$10,316 to +\$911.55 — 91% of the apparent edge was one impulse counted five times . Any number below produced without the constraint is labelled as such and is never the verdict.
Working	reports/friday_v7/gf3/lab{,2,3,4,5,6,7,8}.py (this movement's fresh runs, stages A-I), reports/friday_v7/vacuum/

1. The oracle — the money is real, and the exit is not what is broken

Before inventing anything, the cheapest possible question: if we knew everything, what would the exit hand back? So I took each of the 65 sat-out runs, entered at its own start minute in its own direction, and let the desk's most ordinary wide exit run — stop $3\times\text{ATR}$, target $6\times\text{ATR}$, two-hour cap, one position at a time, full costs. No filter, no cleverness, no hindsight in the exit at all. Just: could a mechanical exit convert a hindsight ceiling into money?

MEASURE	VALUE
Perfect entry, mechanical exit (know the start and the side; $3/6\times\text{ATR}$, 120-min cap)	$n=45 \cdot +\$4,009.14 \cdot +\$89.09/\text{trade} \cdot 84.4\% \text{ wins}$
Census hindsight ceiling on those runs	\$9,987
Converted by the exit alone	40.1%

Forty per cent of a hindsight ceiling, taken by the plainest exit on the board. (n=45 rather than 65 because the serial constraint drops runs that overlap an open position — that is the honest count a desk with one lot could actually have traded.) **Read this next to the desk's own conversion: Movement 1 says we caught 4 runs and banked +\$618 of a \$1,038 ceiling, 60% — and FOUGHT a fifth for −\$1,288, so the desk's actual net near the week's runs was −\$670.** So the exits on this desk are fine. The gap between \$618 and \$4,009 is not exit quality, it is how many runs we were in. Every remaining page of this movement is about that gap.

2. Family A / Family B — which of the 65 were catchable at all?

The mandate asks for a catchable/uncatchable split. Last Friday's VACUUM finding rules out doing it the obvious way: **every pre-run price statistic on a census run list is an artefact of the census's own run-picker**, which dedupes "keep-strongest" and therefore always lands on the local extremum of the move. Prior-5-minute price direction agrees with the run 48.66% of the time on unselected big moves and **22.55%** after the dedup — a 26-point shift created entirely by the selection rule. So a Family A/B split built on "what did the ten minutes before look like" would be measuring our own arithmetic.

The causal version instead asks: **could a detector that reads only the move that has already happened get on board, in the right direction, while there was still money left?** The detector is deliberately stupid — at the close of every minute, has the last ten minutes covered at least 2×ATR in one direction? Take that direction. That is the whole rule, and it cannot see the future.

OUTCOME	RUNS	CEILING	READ
FAMILY A — boarded, right way	50 / 65	\$7,901	catchable in principle
FAMILY B1 — boarded the WRONG way	13	\$1,858	the detector fired against the run
FAMILY B2 — never boarded at all	2	\$228	never put 2×ATR together in ten minutes

Seventy-seven per cent of the runs and 79% of the money are in Family A. And boarding is not late or marginal:

FAMILY A, AT THE MOMENT WE COULD BOARD	VALUE
median minutes into the run	6
median points already gone	32.5
median points still left, to the run's own end	42.1
median R still left	3.16
median R the tape takes off you first	1.12
p75 / p90 R against first	1.67 / 2.00
runs with ≥20pt still left	46 / 50
runs with ≥50pt still left	19 / 50

That "R against first" row is the one that decides stop width, and it is why every tight stop in this movement dies: **a 1×ATR stop is shaken out of the top half of this distribution before the move it correctly identified starts paying.** The p90 is two full ATRs of adverse travel on a trade that goes on to work.

CLUSTER	SAT-OUT RUNS	FAMILY A	CLUSTER CEILING	FAMILY A CEILING
UNCLASS	39	34	\$5,508	\$4,844
OPEN/NEWS	11	8	\$2,058	\$1,698
VACUUM	9	5	\$1,595	\$944
FLOW-LED	6	3	\$826	\$415

The two runs that were never boardable at all are **08-20 13:58 (+56pt, OPEN/NEWS)** and **08-21 13:47 (−58pt, OPEN/NEWS)** — both 56–58 point moves that never assembled ten minutes of one-directional travel worth two ATRs. Both are also just over the census's 55-point bar. **Family B is not a mysterious class of uncatchable event; it is thirteen runs where a momentum detector was pointing the other way when the move began, and two runs that were barely runs.**

3. And here is the whole problem, in two rows

Family A says we can board 50 of 65 runs. So run the identical detector across the whole of the same five days, take every fire, and see what a desk would actually have booked.

THE SAME DETECTOR, 08-17 → 08-21, ALL HOURS, ONE POSITION AT A TIME					
	N	NET	\$/TRADE	WIN %	
every fire	155	-\$1,168.11	-\$7.54	31.0%	
fires that landed INSIDE a sat-out run	50	+\$483.00	+\$9.66	46.0%	
fires that landed anywhere else	105	-\$1,651.11	-\$15.72	23.8%	
ORACLE on the same 155 fires (keep only the winners — cheats)	155	+\$5,272.07	—	—	—

The run trades work. The non-run trades pay for them three times over. The detector boarded **41 of the 65 sat-out runs**, whose combined ceiling is \$6,321 — and still lost money on the week, because two thirds of what it does is not a run. Causal-over-oracle is **negative** here, where last Friday's pooled rider managed 34% on its own entries over a longer tape. That is the number this movement exists to move, and nothing invented this week moved it.

4. ★ The size threshold — the answer the mandate asked for, on real n

The escalation instruction says to narrow to the biggest runs and report the size at which a footprint becomes tradeable. This week's census gives 65 sat-out runs and the size bands inside it are n=4 to n=21, which cannot answer the question. So the census's own picker was rebuilt over the entire lake — every minute with a 15-minute move of 55pt or more, keep-strongest dedup at 15-minute spacing — giving **686 runs across 33 trading days**. Then the identical causal board-and-ride was run against every one of them.

RUN SIZE	N	NET	\$/TRADE	WIN %	CEILING	CONVERSI ON	BOARDED RIGHT WAY
55-70 pt	157	-\$1,899.75	-\$12.10	34.4%	\$19,379	-9.8%	72.6%
70-90 pt	114	-\$724.29	-\$6.35	41.2%	\$18,043	-4.0%	68.4%

90-120 pt	66	-\$625.96	-\$9.48	39.4%	\$13,517	-4.6%	78.8%
120-160 pt	27	+\$1,458.96	+\$54.04	55.6%	\$7,396	+19.7%	77.8%
160 pt +	24	+\$2,079.29	+\$86.64	50.0%	\$10,356	+20.1%	87.5%
all 388 boarded runs pooled	388	+\$288.25	+\$0.74	39.7%	—	—	—

That is a step, not a slope, and it is at 120 points. Three consecutive bands below it lose money — and note they lose it while boarding the right way 68-79% of the time, so this is not a direction failure, it is that a 55-120 point move does not leave enough behind a six-minute-late entry to clear a 3×ATR stop. Above 120 points it does, and the conversion jumps from negative to a fifth of the ceiling.

The 120pt+ band pooled, put through the battery:

TEST ON THE ≥120PT BAND	RESULT	VERDICT
headline	n=51 · +\$3,538.25 · +\$69.38/trade · 52.9% wins	—
strip the best 3 trades	+\$1,939.68	passes
strip the best 5 trades	+\$1,050.89	passes
leave-one-day-out, worst	+\$2,356.82 (drop 2026-07-29)	passes — still green with any day removed
long / short symmetry	LONG +\$74.82/tr (n=23) · SHORT +\$64.91/tr (n=28)	passes, cleanly
days green	15 / 26	—
US session only (13:00-20:00Z)	n=44 · +\$2,683.21 · +\$60.98/trade	holds under the clock cut too
t-statistic on \$/trade	1.83	not significant — n=51 with right-skewed returns
★ THE KILL: the size is only known afterwards	board every qualifying minute → +\$0.23/trade over 860 fires	not causal

So the honest answer to "at what size does a footprint become tradeable" is: 120 points, and it is a hindsight band. The 388 run-boarding fires are only 45.1% of the 860 fires the same rule produces, and they are picked by knowing the move was a run. What would close the gap is a **causal size predictor** — something that says "this one is going to be big" before the fact. Two independent measurements this week say such a thing exists but is too blunt: VACUUM's volume-z separates 120pt+ runs from ordinary tape at **AUC 0.690, p<0.0001** (rising to 0.867 on ATR), and this movement's large-lot flow does better at **AUC 0.721, p<0.0001**. Both were turned into gates. Both died — see §7.

5. The forward tests — last Friday's three MNQ survivors, judged on tape they had never seen

This is the part of last week's deliverable coming due. Three candidates were shipped to SHADOW on 2026-08-15. Five trading days have happened since and

nothing in those studies could have seen them. **No refitting: the specs are frozen exactly as they were written.**

⚠ **An honesty note about two of the three, read this before the tables**

BREAKFLOW and the pooled rider were originally built on their own engines — BREAKFLOW on a 28-tick-day feature table with a 07-18-07-23 hole, the rider on 5-second bars back to 06-19. I re-implemented both on this movement's single ruler (tick-built 1-minute bars, 07-05 onward) so that in-sample and forward are measured with the same instrument. **My in-sample figures therefore do not reproduce theirs to the cent** — I get +\$17.57/trade for BREAKFLOW where the original reported +\$24.44, and +\$13.75 for the rider where it reported +\$33.44. What is valid, and what the verdict rests on, is the **in-sample-versus-forward comparison on one consistent ruler**. I am not claiming to have reproduced last week's engine; I am claiming the mechanism reverses sign on fresh tape, and that claim is measured the same way on both legs.

CANDIDATE	LEG	N	NET	\$/TRADE	WIN %	VERDICT
BREAKFLOW (15-min break + $\geq 2\sigma$ flow event, US-routed, 1.5/3 $\times$ ATR, 60-min clock)	in-sample 07-06→08-14	120	+\$2,108.70	+\$17.57	41.7%	—
	FORWARD 08-17→08-21	24	−\$421.25	−\$17.55	29.2%	REFUTED
The pooled rider mechanism (board 10min/2 $\times$ ATR, US-routed, 3/6 $\times$ ATR, 120-min cap)	in-sample 07-06→08-14	232	+\$3,188.93	+\$13.75	41.8%	—
	FORWARD 08-17→08-21	52	−\$929.46	−\$17.87	28.8%	REFUTED
...with its rvol / atr-percentile floors on	FORWARD	24	−\$160.00	−\$6.67	37.5%	floors help, sign unchanged
VAC-ABS (fade a $\geq 3\sigma$ flow event that failed to move price)	in-sample 07-24→08-15	53	+\$1,409.57	+\$26.60	—	—
		27	−\$529.50	−\$19.61	28.6%	REFUTED

Three for three, and each one is a proper kill rather than a bad week. The proof that it is not one unlucky cell:

SURFACE	CELLS	POSITIVE IN-SAMPLE	POSITIVE FORWARD
rider entry grid (w × k, 4×4)	16	16 / 16	1 / 16
rider exit grid (stop × target, 4×4)	16	12 / 16	6 / 16
BREAKFLOW z × stop × target	12	11 / 12	0 / 12

Sixteen of sixteen entry cells positive in-sample, one of sixteen forward. Zero of twelve for BREAKFLOW. And the controls make it worse rather than better: on the forward leg the rider's **sign-flip control earns +\$9.10 a trade** against the rule's −\$17.87, and BREAKFLOW's sign-flip earns **+\$10.36** against its −\$17.55. **Both rules were worse than their own mirrors on the fresh week.** VAC-ABS is killed independently by strip-5 (−\$427) and t=1.39.

The one thing that did survive on both is the thing this movement keeps finding: **they still show up for the runs.** BREAKFLOW's blanket book caught **24 of 65 sat-out runs and 9 of the top 15** in the fresh week, banking +\$262 on those trades — and lost −\$491 overall. Same shape, again.

6. The gold arc — a year of tape opened, and the trigger did not survive it

Gold got its own second hunt this week (scripts/gf2_mgc_*.py), and it did three things nobody had done: forward-tested last week's gold survivors, put the price trigger in front of **twelve months** instead of 22 days, and turned the cross-asset stone last week's section named. Two data defects were found and cut before any of it: **the backfill pads non-trading minutes as forward-filled flat bars** (2025-10-07 has nineteen flat hours — a rolling 60-minute high over that is a constant and a break gate fires on the first real tick every single day), and **the series rolls to a different contract on 2026-07-30** (mean 1-minute range collapses 0.48 → 0.03pt and price sits 16–43pt below our own captured tape). After cutting: 184 clean days, correlation 0.994 against our quote tape, median absolute difference one tick.

6a. The forward test — one cell lives

GOLD CANDIDATE (FROZEN 2026-08-15)	IN-SAMPLE	\$/TR	FORWARD 08-17→08-21	\$/TR	VERDICT
double-hole reversion (obstacle==0 AND support==0)	n=157, +\$2,648	+\$16.87	n=27, +\$159	+\$5.88	SHADOW holds
primary vacuum fade (far side empty only) — the arm that was to be shadowed first	n=219, +\$2,730	+\$12.47	n=45, −\$107	−\$2.38	REFUTED
mgc_wall_break_go momentum pair	n=63, +\$746	+\$11.84	n=33, +\$47	+\$1.43	PARKED — a wash
CONTROL: the same trigger with no book gate at all	n=449, +\$1,926	+\$4.29	n=110, −\$470	−\$4.27	control behaves

The double-hole cell is **the only candidate in the whole of Movement 3 that came through a forward test positive**. Green on 3 of its 5 days (+\$100, +\$34, −\$74, +\$263, −\$164). Two cautions, said plainly: **n=27**, and the sides split badly forward — **SHORT +\$19.88/trade (n=13), LONG −\$7.11/trade (n=14)**, where in-sample they were within \$4.43 of each other. It stays SHADOW. It is not a deploy on 27 trades and half a symmetry.

6b. The year tape — the price trigger, alone, is a disaster

The 08-15 gold section built its whole 2×2 on one price trigger (close beyond the 60-minute extreme) and let the book decide fade-or-follow. The book cannot be tested before 2026-07-16 — there is no depth. **The price part can, and this is the first time it has been.**

OVER 184 DAYS / 3,716 BREAKS	N	NET	\$/TRADE	WIN %
FADE the break (chandelier 3/2/2)	3,716	−\$11,055	−\$2.97	53.4%
FOLLOW the break	3,716	−\$25,599	−\$6.89	51.6%
the best CONSTANT on the same rows (always long / always short)	3,716	−\$17,786 / −\$18,868	−\$4.79 / −\$5.08	—

Both directions lose. The fade is better than random — a 30-draw random-minute placebo averages −\$19,364 and the real trigger beats 28 of 30 — but "loses less than random" is not an edge. Quarterly it is negative in all five quarters (−\$0.17

to −\$5.79 a trade), monthly green in **2 of 13**, days green 38%, strip-best-3 −\$13,718, both halves negative.

Then the exit was blamed and the exit was tested. **The full matrix — 25 exits on the fade, 25 on the follow, all on the identical 3,716 entries: scalps 0.5R to 3R, chandeliers arm×trail, the desk's dual slot, time caps 15 minutes to 8 hours, break-even-then-trail, no-stop-EOD-flat.**

EXIT FAMILY	BEST CELL	\$/TRADE
FADE — 25 cells	chandelier sl5.0/arm3.0/trail3.0	−\$0.70
FOLLOW — 25 cells	no stop at all, flat at EOD	−\$0.82
dual slot (Lot A scalp 1.5R + Lot B chandelier)	combined	−\$11,387 total
positive cells, out of 50	0	—

Zero of fifty. The excursion statistic says why and it is brutal: median MFE 8.10pt against a cost of 0.383R, and only 2.8% of trades never clear the cost — so the trades do go green, they just do not go green enough more often than they go red. That is a signal problem, not an exit problem. **The 08-15 gold gate's entire edge must therefore live in the book condition, because the price trigger underneath it is worth −\$2.97 a trade over a year.** That is a much narrower claim than the section made, and it is the honest one.

6c. The volatility-shaped entry, and a look-ahead worth \$86,000

The brief's explicit question — if direction is unpredictable but size is predictable, does a breakout-either-way entry beat a directional one? — was tested properly, and the first pass was **wrong** in a way worth showing because nothing about the output revealed it.

FILL MODEL ON THE IDENTICAL STRADDLE	N	NET	\$/TRADE	WIN %
TRIGGER-OPEN — the bug: filled at the open of the bar that went on to reach the level	4,131	+\$53,216	+\$12.88	65.8%
RESTING STOP — fill at the level + 1 tick + half a spread (what a real bracket does)	4,131	−\$32,710	−\$7.92	51.8%
NEXT OPEN — fill at the open of the bar after the trigger	4,131	−\$29,907	−\$7.24	51.6%

An off-by-one in slice-offset arithmetic — `np.nonzero` returns the index inside the slice, so the triggering bar is $i+1+j$ and it was stamped $i+j$ — bought the bottom of the up-bar that triggered the buy. **It is an \$86,000 swing and the output looked perfect: plausible size, 65% wins, works both sides, stable across the year.** Only re-reading the index arithmetic found it. With honest fills the best surviving cell (band $0.25 \times \text{ATR}$, stop 5/trail 3) reads +\$4.45 a trade over 4,135 overlapping positions — and overlapping is a lie, because up to eight run at once. One position at a time it is **−\$8.96 a trade over 2,793**. Re-raced on the 250ms quote tape where the fill order is observed rather than assumed, it is **−\$18.24 a trade (n=445, 27 days)**, and on the fresh five days **−\$28.38 (n=89)**. The mirror loses too. Dead three separate ways.

And the premise itself is false on gold: quintile the 20-minute range and the tightest quintile's next-hour range is 23.5bp against the widest quintile's 65.9bp. Compression expands relatively ($1.29 \times$ vs $0.74 \times$) but the absolute next-hour range still ranks with the prior one, so an ATR-sized straddle gets nothing. **Compression does not predict expansion in a way you can trade.**

6d. ★ The cross-asset stone — named last week, turned this week, closed

Last week's gold section ended: "the one genuinely different input we have never used is the relationship between MGC and MNQ on the same clock." It was unturnable then (27 days of depth). With the 1-minute backfills there are **207,253 paired minutes across 155 days**. MNQ is used only as an exogenous macro input — the trade is in gold, nothing of MNQ's own gates comes near it.

TEST	RESULT	READ
lead-lag, MNQ back 5/15/30/60min → MGC forward 5/15/30/60min (16 cells)		max corr 0.023 no lead-lag at any horizon
NQ 60-min move quintile → gold up% next hour	51.2 / 48.4 / 50.5 / 52.0 / 52.3	no monotone bias
...the same, US session only	53.2 / 48.6 / 50.5 / 47.8 / 48.9	it inverts — so it is noise
HAVEN gate (buy gold when NQ breaks down) and its MIRROR, 3 thresholds × 2 exits	12 of 12 cells negative , −\$1.66 to −\$13.89/ trade	REFUTED
NQ state as a filter on the gold fade (4 cuts, 200 random keeps each)	best cut −\$1,468, beaten by 21/200	better than random, still negative

The stone is turned and it is empty. Do not re-dig it.

7. ★ Every candidate invented this week, by name, with the test that killed it

The standing rule is that the honest failures are the deliverable. Nine new candidates reached a backtest this week across the two hunts. None survived. Each cause of death is a named robustness test — **not one of them is a win-rate kill**, per the operator's judging rule.

CANDIDATE	THE IDEA, IN ONE LINE	N	NET	\$/TR	WIN %	CAUSE OF DEATH
VAC-Z1	trade the VACUUM label literally	1,768	−\$4,150.79	−\$2.35	35.7%	The label has no edge at all. Unconstrained it loses −\$32,184 over 8,589 trades. REFUTED
VAC-SNAP	heavy one-sided aggression that failed to move price, then a structural break	410	+\$370.64	+\$0.90	34.9%	Strip-the-3-best → −\$1,436 (top 5 trades = 738% of net); OOS leg −\$18.02/tr; LOO goes negative; 12/33 days green. REFUTED
VAC-CUMDIV z1.5 / z2.0	cumulative-delta divergence against price	506 / 264	−\$1,814 / −\$1,669	−\$3.59 / −\$6.32	~32%	Negative at every threshold and in both directions of the sweep. REFUTED
VAC-IGN-BRK	the detection bracket: volume-z ≥2 + ATR expansion ≥1.15, then break the 5-bar range	284	+\$911.55	+\$3.21	32.4%	SIGN-FLIP CONTROL: fading the break earns + \$3.08/tr against the signal's + \$3.21. Both directions win → volatility harvesting, not edge. Also strip-3 −\$915; the parameter surface is a step, not a plateau (4 of 15 cells negative). REFUTED
BIGBOARD (new, this movement)	the combination the two hunts jointly asked for: use VACUUM's size detector only to decide when to look, and the move itself to decide which way	281	+\$171.79	+\$0.61	37.0%	Strip-the-3-best → −\$1,497; leave-one-day-out → −\$1,019 (drop 07-30); in-sample −\$2.10/tr against forward + \$12.84/tr — opposite signs; REFUTED

							the arm is not load-bearing (boarding with no arm at all makes +\$0.23/tr on 860 fires, so the arm buys \$0.38 a trade for a 67% volume cut); random-arm bootstrap beats it in 9 of 40 draws; t=0.06; long – \$5.87 vs short + \$7.33.
SMALLFADE (new, this movement)	fade the retail tape: 5-min cumulative small-lot (<5) aggressor flow z-scored, fade it	1,009	–\$683.14	–\$0.68	51.8%	nothing	The parameter sweep — 0 of 16 cells positive , at any threshold, stop or target; the sign-flip loses in all 16 too, so nothing directional clears costs; strip-3 – \$1,351; LOO – \$1,580; all four time-shift placebos also negative; t=–0.27.
LOTDIV (new, this movement)	when large lots and small lots point opposite ways, take the large lots' side	1,668	–\$292.07	–\$0.18	35.5%		A flat null: in-sample –\$0.22/tr, forward + \$0.34/tr. The 53.59% directional read it is built on (n=1,506) does not survive costs at any exit.
mgc_straddle (breakout-either-way)	volatility-shaped entry on gold	2,793	–\$25,026	–\$8.96	54.4%		The look-ahead audit (an \$86k swing), then the honest fill model, then the serial constraint, then the 250ms quote tape (–\$18.24/tr). The mirror also loses. Placebo beats it in 25 of 30 draws.
MGC HAVEN / MIRROR (cross-asset)	trade gold off what the Nasdaq just did	415–1,279	–\$1,572 to –\$16,554	–\$1.66 to –\$13.89	50–61%		12 of 12 cells negative , on 155 days; the underlying conditional bias is 51.2–52.3% and it inverts in the US session.

Two nulls that never reached a spec, recorded so nobody re-walks them: **flow-sign continuation** (48.95% agreement on unselected big moves) and **the coil ratio** (AUC 0.501 on the sat-out set — no detection at all).

7a. The two real findings inside the dead candidates

SMALLFADE and BIGBOARD both died, and both left something behind that is worth more than they were.

DIRECTION DISCRIMINATOR, ON UNSELECTED BIG MOVES (NO PICKER, NO DEDUP)	≥55PT	≥90PT	≥120PT
prior 5-min price change	48.66%	—	—
total aggressor flow z	48.92%	—	—
L2 far-side depletion (246.9M book rows)	49.44%	—	—
LARGE-lot flow (≥5 lots)	48.71% ±1.39	46.20% ±2.54	47.75% ±3.86
SMALL-lot flow (<5 lots) — inverse, and monotone in size	46.57% ±1.39	44.31% ±2.54	42.95% ±3.86
large vs small DIVERGENCE — large lots win	53.59% ±2.53	53.45% ±4.86	—

Fading the small-lot tape is **53.43% at 55pt and 57.05% at 120pt**. That is the first pre-run direction reading on this desk that is outside its own error bar — against three that are refuted at tight bars. **And it still does not trade, and the reason is important:** on all minutes, small-lot flow reads **49.56%**. The separation only exists once you already know a big move happened. It is the same conditioning trap as the census picker, one level up — and it is exactly why 0 of 16 SMALLFADE cells make money. I am reporting it because it is a real measurement and because it is the seed of next week's work, not because it is an edge.

SIZE DETECTOR, AUC AGAINST A RANDOM LAKE MINUTE (4,000- DRAW PERMUTATION)	55-70PT	70-90PT	90-120PT	120PT+
large-lot flow , 5-min cumulative	0.510 (p=0.26)	0.558 (p=0.002)	0.659 (p<0.0001)	0.721 (p<0.0001)
volume-z (VACUUM's detector)	0.557	0.542	0.567	0.690
ATR(1m,14)	0.616	0.670	0.801	0.867
large-lot SHARE of volume — runs from small lots, not big ones	0.452	0.475	0.427	0.368

ER15 — the only feature below 0.5, and it falls as runs get bigger	—	—	0.442 (sat-out 65)	0.403 (top-15)
--	---	---	--------------------	----------------

Two things in that table belong in the router section rather than here. **The biggest runs come out of chop** — ER15 is the only measured feature whose AUC sits below 0.5, and it falls further as the runs get bigger. Every ER-floor arm condition the live router owns is standing gates down at the moment the largest moves begin. And **the big runs are made of small lots, not large ones** — large-lot share reads AUC 0.368 on the 120pt+ band, which is a strong inverse reading. Runs are a crowd, not a whale.

8. ★ THE CHOP-DAY SCALP LAB — can we stop donating on chop days?

You asked for something specific: the trend days carry the profit, the chop days donate, so build a purpose-built chop-turn scalper, catch the oscillation turns, bank tiny, and **lead with the order book**. It was built — nine named candidates, **290,700 five-second decision points across eighteen trading days, 25.2 million trade ticks, 5.06 million ten-deep book snapshots and 33.9 million level-1 book updates at 41 milliseconds**, everything raced tick by tick at \$1.50 the round trip. **The answer is no, and it is a clean no with two named causes of death. But the premise needs correcting, and the correction is worth about \$5,400.**

★ THE PREMISE IS WRONG — WE ARE NOT DONATING IN THE CHOP

Before building anything, every live MNQ tournament trade was attributed to the 30-minute regime block it opened in. The result changes the target.

BLOCK REGIME	N	NET	WIN %	\$/TRADE
BUILDING	121	+\$1,836.50	48.8%	+\$15.18
NORMAL-CHOP	165	-\$46.00	46.1%	-\$0.28
CLEAN-TREND	24	-\$181.00	33.3%	-\$7.54
DEAD-CHOP	41	-\$199.50	36.6%	-\$4.87
VIOLENT-WHIPSAB	348	-\$5,414.00	33.9%	-\$15.56

-\$5,414 on 348 trades opened in violent-whipsaw blocks, against -\$246 across all 206 trades opened in the genuinely quiet blocks

combined. On the census week of that study, half the tournament's trades opened inside whipsaw blocks and that half was the entire loss (−\$842.50 on 54 trades) — while the quiet tape, the tape a chop scalper would live in, was mildly profitable (+\$146.50 dead-chop, −\$46.50 normal-chop, +\$100 net on 44 trades). **We do not donate in the chop. We donate in the high-ATR thrash inside chop days.** A scalper aimed at the quiet tape is aimed at the wrong target — the money here is a bench, not a gate.

8a. The model-free race — before any gate exists

No parameters, no fitting, no P&L. At every 5-second point where price sits in the top or bottom decile of the trailing 30-minute range, put symmetric barriers at ± 0.75 ATR and let the real tick path decide which is touched first. One event per side per minute. A fade with no information is 50%.

SEGMENT	N	P(REVERT FIRST)	Z
all extremes	8,386	51.2%	+2.16
chop blocks (ER30 < 0.30) — the target population	5,607	50.6%	+0.84
chop + US session	1,642	52.5%	+2.02
chop + overnight	3,965	49.8%	−0.30

Fading a range extreme in chop is a 50.6% proposition. With \$2.25 of friction against a ± 0.75 -ATR payoff — about \$12 a side on an 8-point ATR — you need roughly **55%** to break even. The confidence interval comfortably contains a coin. (Last week's independent study, fading a stretch from VWAP rather than a range extreme, got 47.2% over 4,249 races. Two event definitions, two harnesses, same verdict.)

8b. The book — the core hypothesis, refuted at two fidelities

MEASURE	TAIL	N	P(REVERT)	Z
far-side 5-deep / near-side 5-deep (250ms)	thin 10%	564	49.3%	−0.34
far-side 5-deep / near-side 5-deep (250ms)	thick 10%	562	50.0%	+0.00
41ms: a DEFENDED wall (added/removed ≥ 1.0 on the continuation side) + price stalled ≤ 2 pt — your literal brief	—	481	49.3%	−0.32
41ms level-1 refills, top quintile	—	648	55.1%	+2.59

41ms level-1 pulls, top quintile	—	650	56.2%	+3.14
41ms total churn, top quintile	—	654	55.8%	+2.97

The wall ratio is 49.3% thin and 50.0% thick — the measure the whole hypothesis rests on has no sign at all, and the defended-wall cell, the literal thing the brief asked for, is below a coin. The only quintiles that lift are refills, pulls and total churn together — all in the same direction, which means the signal is "the book is busy", not "the book is defended". A measure that reads the same whether size is being added or removed is measuring activity. Two fidelities, 33.9M level-1 updates, same answer.

8c. The nine candidates — and the two instrument failures found on the way

Two leaks were found in the harness before any verdict. The second one was the entire edge: the feature builder joined each 5-second row to the 1-minute bar it sits inside, so the ATR and efficiency reading at 09:14:05 came from a minute that had not finished — up to 55 seconds of future, straight into the variable everything else was conditioned on.

MEASUREMENT	WITH THE LEAK	LEAK FIXED
P(revert) at range extremes, chop blocks	54.2% (z=+6.16)	50.6% (z=+0.84)
P(revert), dead-chop blocks — "the great segment"	60.2% (z=+6.63)	51.6% (z=+1.24)
first-pass configs positive (6 candidates × 4 exits)	23 of 24	5 of 24
best first-pass config	+\$3,250.83	−\$823.01

Twenty-three green cells out of twenty-four is what tipped it off. Six candidates that share a mechanism are not six independent edges. The tell is always the same: when the null refuses to appear anywhere, the harness is lying.

CANDIDATE	N	NET	WIN %	\$/TRADE	GREEN DAYS	TRADES/DAY	IN-SAMPLE	OOS (THE 5 CHOP DAYS)
CT-9 STALL+PUSH	230	+\$232.05	69.1%	+\$1.01	11/18	12.8	+\$399.49	−\$167.44
CT-3 EXHAUST-CT — your literal brief	78	+\$230.97	70.5%	+\$2.96	12/18	4.3	+\$279.92	−\$48.95
CT-6 THINWALL+PUSH	167	−\$88.99	65.9%	−\$0.53	11/18	9.3	+\$43.10	−\$132.09
CT-7 STALL	404	−\$968.23	63.6%	−\$2.40	7/18	22.4	−\$505.24	−\$462.98

CT-4 PUSHFADE	911	−\$1,127.01	65.3%	−\$1.24	5/18	50.6	−\$823.01	−\$304.00
CT-1 FLATLINE (control, no book input)	838	−\$1,807.38	64.0%	−\$2.16	5/18	46.6	−\$1,494.86	−\$312.52
CT-5 DEplete	689	−\$2,460.69	61.7%	−\$3.57	3/18	38.3	−\$1,809.20	−\$651.49
CT-2 THINWALL	675	−\$3,188.04	60.9%	−\$4.72	3/18	37.5	−\$2,368.56	−\$819.48
CT-8 STALL- NO-BOOK	1,113	−\$4,130.16	62.2%	−\$3.71	1/18	61.8	−\$3,647.10	−\$483.06

Look at the win-rate column and then ignore it, exactly as your judging rule says. Every candidate wins between 61% and 71% of its trades, including the one that loses \$4,130. That is geometry, not quality: a 1.5-ATR stop against a 0.75-ATR target wins twice as often as it loses by construction and needs about 67% just to stand still. CT-8 wins 62% of 1,113 trades and is the worst gate in the table. This is the clearest illustration yet of why win% is off the kill list — and why it must never go on the promotion list either.

CT-8 versus CT-7 is the L2 verdict in two rows. CT-7 is CT-8 plus the wall term. Adding the book takes 1,113 trades to 404 and −\$4,130 to −\$968 — but it did that by removing trades, not by finding good ones, and both numbers are still deep negative. A filter that improves a loser by cutting exposure is a volume cut wearing a filter's clothes. This desk has seen that exact shape before, in the rgv absorption-confirm.

8d. Where the two survivors died

TEST	CT-9 STALL+PUSH	CT-3 EXHAUST-CT
in-sample (13 days)	n=148, +\$399.49, +\$2.70/tr	n=41, +\$279.92, +\$6.83/tr
OUT-OF-SAMPLE — the 5 chop days it exists for	n=82, −\$167.44, −\$2.04/tr, green 2 of 5	n=37, −\$48.95, −\$1.32/tr
strip the best 3	+\$27.77 → 12 cents a trade	+\$99.81
side-shuffled placebo, 300 draws	real +\$399 vs placebo mean +\$197 (sd \$261), p=0.203	real +\$280 vs placebo mean +\$112 (sd \$181), p=0.207
mirror test (every side inverted)	−\$643 in-sample — passes, the sign is real	−\$174 in-sample — passes
leave-one-day-out	0 of 18 removals negative	0 of 18 removals negative
side symmetry	LONG +\$110 / SHORT +\$122 — passes	LONG +\$134 / SHORT +\$97 — passes
cost stress +\$1.00/RT	+\$2.05 — dead	+\$152.97 — survives

CT-9's exit mix says the rest: **63 stops at −\$33.82 each against 153 targets at +\$16.06 each. It needs 67.8% just to tread water and delivers 69.1% — the whole "edge" is 1.3 percentage points of win rate above its own geometric break-even**, which on 230 trades is a rounding error. CT-3's placebo is the subtle

one and deserves the fair reading: **a random-direction gate makes +\$111.69 on those same entries**, because the 1.5-ATR-stop / 0.75-ATR-target geometry pays out that way regardless of side. CT-3's real +\$279.92 sits 0.93 standard deviations above that. But the mirror test says its direction is emphatically not arbitrary — flipping every side costs about \$11 a trade. **What the placebo actually says is narrower: at 41 in-sample trades the spread between the real gate and a coin is smaller than the noise of 41 trades.** That is a thin-n verdict, not a mechanism verdict, so CT-3 is PARKED, not REFUTED.

8e. Against your \$200-300 bar

CANDIDATE, BEST HONEST CONFIGURATION, ON THE FIVE CHOP DAYS	WEEK TOTAL	PER CHOP DAY	VS THE BAR
CT-3 EXHAUST-CT (37 trades)	−\$48.95	−\$9.79	misses by ~\$250/day
CT-9 STALL+PUSH (82 trades)	−\$167.44	−\$33.49	misses by ~\$270/day
CT-1 FLATLINE (control)	−\$312.52	−\$62.50	misses
the kindest possible number — CT-9's in-sample rate, which does not replicate, × its trade rate	—	+\$30.78	still misses by ~\$220/day

To clear \$250 on a chop day at CT-3's in-sample \$6.83 a trade you would need 37 qualifying turns a session. The gate produces three to five. The gap is not a tuning gap, it is an order of magnitude, and no threshold in the 108-cell sweep closes it — every parameter that lifts median n above ~80 also puts median dollar-per-trade below zero.

8f. Two pieces of engineering worth keeping even though the gate died

FINDING	EVIDENCE	USE
The exit ridge: stop 1.5–2.0×ATR with target 0.5–1.0×ATR (0.33R–0.67R)	53 of 60 cells positive, a broad plateau; peak +\$6.90/tr. A tight 0.75-ATR stop on identical entries makes \$1.30 where a 1.5–2.0 stop makes \$6.36–\$6.90.	Reusable by any fading gate on the roster. Your 0.5R instinct was right; the axis you did not guess is that a tight stop kills it.
"Price refused to follow" travels; the wall does not	Tighten the stall from 4pt to 2pt and the 108-cell grid goes from 10-of-54 positive to 40-of-54. The wall term spikes at exactly 1.5× and collapses at 1.2× and 2.0×.	If anyone reaches for the exhaustion_short footprint again, take the stall condition and leave the wall ratio behind.
Your VWAP-flat filter	Three settings — 0.5 ATR / 1.0 ATR / switched off — give −\$0.70 / −\$0.01 / +\$0.00 median \$/trade across 108 cells.	REFUTED as redundant. The 30-minute efficiency read already does that job; the flatness test measures it twice.

8g. Chop-lab disposition

LEAD	VERDICT
------	---------

		CAUSE OF DEATH / WHAT WOULD BRING IT BACK
CT-3 EXHAUST-CT (your literal brief)	PARKED	Killed by its side-shuffled placebo (p=0.207) and the OOS leg . The mechanism is not disproven — the sample is hopeless at 41 in-sample trades. Revive at ~250 real fires (about 12 weeks at 4.3/day) with placebo p<0.05 on that n. Cheap to accumulate; a shadow slot costs nothing.
CT-9 STALL+PUSH	PARKED	Killed by the OOS leg (green on 2 of 5) and strip-the-best-3 (+\$27.77 total). Revive only if the US-session pocket (n=38, +\$7.52/tr) holds on independent days.
CT-6 THINWALL+PUSH	PARKED	Best placebo showing in the study (p=0.090) but +\$43 in-sample and −\$89 across the window. Revive if the placebo result replicates on fresh days only.
CT-1 FLATLINE / CT-8 STALL-NO-BOOK (controls)	REFUTED	Named test: the model-free first-touch race — 50.6% over 5,607 events. That is the mechanism failing, not the sample. CT-8 is green on 1 of 18 days.
CT-2 THINWALL / CT-5 DEplete (thin/depleting far side)	REFUTED	Worse than their own placebos (p=0.875, p=0.820). Both were built from what the leaked race pointed at and did not survive the leak's removal.
The far-side WALL as a turn separator (the core L2 hypothesis)	REFUTED	Named tests: the 250ms ten-deep race (49.3% thin / 50.0% thick) and the 41ms event-stream race (defended wall + stall = 49.3%, z=−0.32, n=481). No reformulation is left — the ratio was tested at level 1, 5 deep, 10 deep, its 30-second change, refill count, pull count and added/removed balance.
★ The VIOLENT-WHIPSAW block bench	SHADOW → hand to the router	The finding this lab actually produced: – \$5,414 on 348 trades at −\$15.56 each in high-ATR / low-efficiency blocks against roughly breakeven in all quiet blocks. The classifier is cheap and strictly trailing. Not claimed as an edge — the missed-winner counterfactual has not been run, so \$5,414 is gross, not net.

Chop-lab verdict: DO NOT BUILD A CHOP-TURN SCALPER. Not because chop is untradeable in principle, but because the specific thing that would have to be true — that the resting book separates a turn from a break at a range extreme — is measurably false at two fidelities across 8,386 model-free races, and because the desk is not losing money in the chop anyway. **The real lead is a bench, not a gate, and it is worth roughly \$5,400 gross over eighteen days.** The single most valuable backtest anyone could run off this lab is the counterfactual: what would benching those blocks have cost in winners forgone?

9. The escalation trail — where does an edge start to appear as we narrow?

The mandate asks for every attempt at every narrowing level, so the loop is laid out end to end. This is the trail, and the trail is the answer.

LEVEL	POPULATION	WHAT WAS HUNTED	RESULT
1 — full sat-out census	65 runs, \$9,986	VAC-Z1, VAC-SNAP, VAC-CUMDIV, VAC-ABS forward, SMALLFADE, LOTDIV, BIGBOARD	All dead. Detectors are indistinguishable from noise here: volume-z AUC 0.597, ATR 0.553 (p=0.071), flow z 0.522 (p=0.27).
2 — top 25 by size	76pt and up	the same detector battery, permutation-tested at n=25	Still noise: volume-z 0.637 (p=0.0089), ATR 0.583 (p=0.072), flow z 0.525 (p=0.34).
3 — top 15 by size	87pt and up	the same battery	★ ALL FOUR detectors become significant at once: volume-z 0.730 (p=0.0004), ATR 0.676 (p=0.0081), flow z 0.679 (p=0.0088), ATR-expansion 0.674 (p=0.0093). The footprint appears exactly where the mandate hoped.
4 — the lake-wide check, 686 runs	33 days, real n	does the footprint keep climbing?	Yes, monotonically. Volume-z 0.557 → 0.542 → 0.567 → 0.690; ATR 0.616 → 0.670 → 0.801 → 0.867; large-lot flow 0.510 → 0.558 → 0.659 → 0.721. And boarding pays: +\$54.04/tr at 120-160pt, +\$86.64 above 160pt.
5 — convert it		VAC-IGN-BRK and BIGBOARD, both — built directly on the level-3/4 footprint	Both dead. VAC-IGN-BRK caught 9 of 65 sat-out and 5 of 15 top-15 on 48 fires — 3.6× the random-fire baseline — and its sign-flip control earns +\$3.08 against its +\$3.21. BIGBOARD caught 20 of 65 and 7 of 15 on 51 fires and died on strip-3 and LOO.

So the escalation did exactly what the mandate hoped it might, and then hit a wall that is now precisely located. The footprint is real, it is size-graded, it becomes statistically unambiguous at 87 points and economically positive at 120, and it converts into catches at 3.6× the random baseline. What it does not do is supply **direction**, and every instrument this desk owns has now been ruled out on that question at tight error bars:

DIRECTION DISCRIMINATOR	POPULATION	AGREEMENT	STATUS
prior 5-min price change	n=4,930	48.66%	REFUTED
prior 15-min price change	n=4,932	50.39%	REFUTED
aggressor flow z / 5-min cumulative delta z	n=4,928 / 4,887	48.92% / 48.95%	REFUTED
position in the 30-min range	n=4,930	49.89%	REFUTED

L2 far-side depletion — last week's "one stone still unturned"	21,249 book-minutes, 246.9M rows	49.72% ±0.67	REFUTED — the null is not n-limited
large-lot aggressor flow — this week's stone	n=4,939	48.71% ±1.39	REFUTED
cross-instrument (MNQ → MGC lead-lag)	207,253 paired minutes	max corr 0.023	REFUTED
SMALL-lot aggressor flow, faded	n=4,937	53.43% ±1.39 (57.05% at 120pt+)	real separation, PARKED — does not trade, and only exists conditional on a big move

10. Robustness battery — every surviving-ish candidate, test by test

TEST	VAC-ABS	VAC-SNAP	VAC-IGN-BRK	BIGBOARD	≥120PT BOARD
sign-flip control	−\$23.86/tr ☐	−\$1.46/tr ☐	+\$3.08 vs +\$3.21 📉	−\$0.35/tr ☐	n/a (hindsight)
unconditional control (arm every minute)	—	—	−\$1.94/tr ☐	+\$0.23/tr 📉 arm not load-bearing	+\$0.23/tr 📉
random-arm bootstrap (40 draws)	—	—	beaten 0/40 ☐	beaten 9/40 📉	—
time-shift placebo	+240m earns +\$2.60 ⚠	—	+90m +\$1.06 ⚠	—	—
parameter plateau	—	edge only as n collapses 📉	step, not plateau 📉	no plateau; flip wins in 2 of 5 cells 📉	—
strip the best 3	+\$118 ⚠	−\$1,436 📉	−\$915 📉	−\$1,497 📉	+\$1,940 ☐
strip the best 5	−\$427 📉	−\$2,364 📉	−\$1,685 📉	−\$2,493 📉	+\$1,051 ☐
leave-one-day-out, worst	+\$9.53 ☐	−\$2.55 📉	−\$0.26 📉	−\$1,019 📉	+\$2,357 ☐
out-of-sample leg 08-17→08-21	−\$19.61/tr 📉	−\$18.02/tr 📉	+\$3.85/tr ☐	+\$12.84/tr (in-sample −\$2.10 — signs inverted 📉)	—
v5-archive leg 07-05→07-17	+\$439 ☐	−\$758 📉	−\$19.44/tr 📉	—	—
long / short symmetry	+\$27.60 vs −\$2.58 📉	+\$8.76 vs −\$6.57 📉	+\$6.74 vs −\$0.47 ⚠	−\$5.87 vs +\$7.33 📉	+\$74.82 vs +\$64.91 ☐
big-moves-caught	0/65, 0/15 📉	6/65, 2/15	9/65, 5/15 — 3.6× baseline ☐	20/65, 7/15 ☐	—
t-statistic on \$/trade	1.39	0.17	0.45	0.06	1.83
days green	15/28	12/33	15/32	14/32	15/26

Two rows deserve reading rather than skimming. **VAC-IGN-BRK** and **BIGBOARD** are the only things in three Fridays of greenfield work that pass the placebo family, the unconditional control and the bootstrap while catching the big runs at 3.6× baseline — and each is killed stone dead by a single control. And the **≥120pt board column is the only one with a clean sweep of green ticks on strip-3, strip-5, LOO and symmetry**, and its cause of death is not a robustness test at all: it is that the band is defined by hindsight.

11. The router — what would be armed, and the honest answer that nothing is

Nothing survived, so what follows is the segmentation evidence and the rule that would have been written, stated so it is falsifiable rather than aspirational. The regime split is consistent across all three VACUUM candidates and is the one genuine policy finding: **every one of them is an ATR-high, US-session animal and every one bleeds in low ATR.**

SEGMENT	VAC-ABS	VAC-SNAP	VAC-IGN-BRK
ATR low tercile	−\$5.26/tr (n=33)	−\$8.10/tr (n=137)	−\$11.64/tr (n=95)
ATR mid tercile	−\$5.99/tr (n=31)	−\$10.09/tr (n=136)	+\$3.24/tr (n=94)
ATR high tercile	+\$50.86/tr (n=33)	+\$20.82/tr (n=137)	+\$18.03/tr (n=95)
US session 13:00-21:00Z	+\$51.24/tr (n=26)	+\$22.89/tr (n=104)	+\$16.33/tr (n=110)
overnight / pre-open	−\$0.18/tr (n=71)	−\$6.57/tr (n=306)	−\$5.08/tr (n=174)

The rule that would have been written, had the sign-flip control not killed it

ARM vac_ign_brk when $ATR(1m, 14) \geq 21.0pt$ on MNQ **and** $13:00 \leq t < 21:00Z$. **BENCH** below 19.0pt (2-point hysteresis, matching the existing convention) with a **≥ 5 -minute dwell on the arm side** — ATR crossing a bare threshold on a partial bar is exactly the oscillation that produced 20 switch changes for 4 trades on 08-07. **No ER condition**, for the reason below. Expectancy on that home regime only: **+\$18.03/trade over 95 trades**. Blanket, as context and never as the verdict: +\$3.21 over 284.

Do not arm it. The home-regime number is the same number the fade control earns. Conditioning on ATR-high selects trades with more range in both directions, and a gate that makes money whichever way you point it is a volatility harvester with a stop, not an edge.

The one thing here that is a genuine router finding, and it belongs to the router section rather than to any gate: the biggest runs start from chop. ER15 is the only feature measured whose AUC sits below 0.5, and it falls further as the runs get bigger — **0.442 on the sat-out 65, 0.403 on the top-15**. Every ER-floor arm condition the live router owns is standing gates down at the moment

the largest moves begin. Two independent phases reached that conclusion this week and a third reached it last week from the filter side (ER cuts fail their placebo while ATR floors beat theirs).

ER is also not a discriminator for the gates themselves, in either direction.

Serially, VAC-IGN-BRK books $-\$0.49/\text{tr}$ in dead chop ($n=83$), $+\$17.85/\text{tr}$ in the 0.15–0.35 middle ($n=93$) and $-\$6.55/\text{tr}$ in clean trend ($n=108$). A middle bucket carrying the whole of a $\$911$ net across three buckets is noise, not a regime. $\triangle$ Worth recording: the unconstrained run said $+\$22.28$ in chop and $+\$13.79$ in trend — that was the overlap-inflated ledger, exactly the number this movement warns against quoting.

12. What could still be wrong with all of this

- **The forward legs are five days.** They are the right five days — frozen specs, tape nothing had seen — but five is five. What makes them a kill rather than a bad week is that the whole parameter surface reversed ($16/16 \rightarrow 1/16$, $11/12 \rightarrow 0/12$) and both rules were beaten by their own mirrors.
- **My re-implementations are not last week's engines.** Stated in full in §5. The comparison is internally consistent; the absolute in-sample figures are not reproductions.
- **The $\geq 120\text{pt}$ band is $n=51$ with $t=1.83$.** It passes every robustness test I can run and it is still not statistically significant. It is a shape to build toward, not a number to trust to the cent.
- **The gold year-tape has no book and no spread.** The $\$7.50$ round trip is modelled, not raced, and stressed at 0.15/0.50pt per leg — the sign never changes. But it means the year tape can only refute the price trigger, and it does not touch the book condition the surviving gold cell actually rests on.
- **The gold year-tape cannot resolve intrabar ordering.** On 1-minute bars the straddle reads $+\$3.90/\text{trade}$ before the tiebreak and $-\$7.96$ after it — three times the apparent edge is one assumption a minute bar physically cannot answer. That is why it was re-raced on the 250ms quote tape, and that is the number quoted.
- **Three separate look-ahead / accounting bugs were found and fixed this week** — the gold straddle's one-bar offset ($\$86\text{k}$), VACUUM's overlapping-position inflation ($\$9.4\text{k}$ of a $\$10.3\text{k}$ headline), and the chop lab's 55-second feature join. I cannot rule out a fourth. What I can say is that every conclusion here is a refutation, and leaks manufacture edges rather than

destroy them — so an undiscovered leak would make this verdict more negative, not less.

- **MGC gates were invented fresh, never ported.** No MNQ gate was transferred to gold and no gold result is priced at \$2.00/point.

13. Verdict

MOVEMENT 3 IS A NULL ON MNQ THIS WEEK, WITH ONE THIN GOLD SURVIVOR — and for the first time the null has a precise address.

Nine new candidates reached a backtest and all nine are graves, each with a named cause of death, none of them a win-rate kill: **VAC-Z1** (refuted outright, −\$4,151 over 1,768), **VAC-SNAP** (strip-3 → −\$1,436), **VAC-CUMDIV** (negative at every threshold), **VAC-IGN-BRK** (sign-flip control earns +\$3.08 against its +\$3.21), **BIGBOARD** (strip-3 → −\$1,497, LOO → −\$1,019, arm not load-bearing), **SMALLFADE** (0 of 16 sweep cells positive), **LOTDIV** (a flat null), **mgc_straddle** (a look-ahead worth \$86k, then dead three ways), **MGC HAVEN/MIRROR** (12 of 12 cells negative). And all three of last Friday's MNQ SHADOW survivors — **BREAKFLOW**, **the pooled rider mechanism** and **VAC-ABS** — were forward-tested on frozen specs and all three lost, each with the whole surface reversing rather than one cell.

What is not a null: **the money is real and the exit converts 40.1% of it** (perfect entry, mechanical exit, +\$4,009 on 45 trades). **Fifty of the 65 sat-out runs are boardable in the right direction a median six minutes in with 42 points still on the table. The size threshold the mandate asked for is 120 points, it is sharp, and above it boarding pays +\$69.38 a trade and passes strip-3, strip-5, leave-one-day-out and long/short symmetry.** And detection keeps improving as the runs get bigger, in three independent instruments, up to AUC 0.867.

The single sentence: we can see the runs coming, we can get on board with money left, the exit does its job — and we cannot tell which way they go, so two thirds of what any detector fires on is not a run and those trades eat the whole prize. Every direction instrument this desk owns has now been ruled out at tight error bars: price (48.66%), aggressor flow (48.92%), L2 depth (49.72% ±0.67 on 21,249 book-minutes), large-lot flow (48.71%) and cross-instrument lead-lag (max |corr| 0.023).

★ THE ONE STONE STILL UNTURNED FOR NEXT WEEK

A causal SIZE predictor sharp enough to be an arm — not a direction predictor. This is a change of target and it is what this week's work earned.

Every hunt for three Fridays has gone looking for direction and every one has been refuted. But the ≥ 120 pt band does not need direction: the board rule supplies it causally, from the move that has already happened, and it makes +\$69.38 a trade up there and passes every robustness test. **The only thing missing is knowing, before the fact, that this is a 120-point move and not a 70-point one.**

Three measurements say that predictor is close but not yet sharp enough: | **large-lot flow** | AUC **0.721**, ATR AUC **0.867**, volume-z AUC **0.690** on the 120pt+ band, all $p < 0.0001$ on real n. BIGBOARD tried the crudest possible version of this — a two-condition arm at the 2.88% most extreme minutes — and it was too blunt: it caught 7 of the top 15 but died on strip-3. **The concrete next step is to stop hand-picking two thresholds and fit the arm properly against the 686-run lake-wide label**, scoring it on catch rate at fixed fire count rather than on P&L, and only then bolting the board rule on. That is a well-posed supervised problem with 686 positives, three measured features, a clean holdout in the fresh week, and an honest baseline ($3.6\times$ random) already established.

Second, cheaper, and worth running alongside: the **whipsaw-block bench counterfactual** from the chop lab — what benching the high-ATR / low-ER blocks would have cost in winners forgone. Everything in §8 says that is where the desk's money actually goes; nothing yet proves the bench is free, and \$5,414 is a gross number until someone runs it.

Nothing from this movement goes live. Nothing from this movement even goes to the paper tournament. The gold double-hole cell stays SHADOW on 27 forward trades and a broken symmetry. Everything else is PARKED or REFUTED, by name, above.