

GAZBOT V7 — the Friday report

Week ending **Fri 2026-09-18** · Mon 14 – Fri 18 September 2026 · published 2026-09-18
· MNQ, with MGC captured · 22 plays · front matter + 4 body sections

★★★ THIS REPORT WAS REFOCUSED ON 2026-09-18 — IT IS NOW ABOUT ONE THING

Every previous edition of this report was a review of the automated desk — the paper tournament, the gate rehab bench, the shadow board, the router, the greenfield labs. **This one is about you.** One subject, start to finish: the trades you make with your own hands, what the tape looked like when you made them, and how much of it a machine could do instead. Four sections and nothing else.

★ **What that means for anything you remember from last month's report.** The tournament sections are not in this document. They are not summarised, not appended and not “covered elsewhere” — they are absent, by instruction. Their fragments are all still sitting on disk in reports/friday_v7/sections, every one of them dated 2026-08-29 or 2026-09-15, and the build refuses to name them: stitching one in would publish last month's tournament review under this week's masthead. **Last week's 58-play action card is UNGRADED for the same reason**, and that gap is stated in full rather than papered over — see the ledger below the card.

★ **THE WEEKEND CARRY CHECK — THE DESK IS FLAT, AND HERE IS WHAT SAYS SO**

Nothing is riding over this weekend. The last leg on the book closed at **18:33:23Z on Friday 18 September** (\$284.00 at 29,775.50) and no position was opened after it. data/day_rider_state.json reads closed true; its venue read at 23:01:06Z returned **net 0** with venue_ok true — taken after that final claim, which is what makes it evidence rather than an assumption. core_health at 23:01:23Z reads flat true, halted false, and

desk_reconcile at 23:01:02Z reads venue +0 = tournament +0 + rider +0 → unaccounted +0.

⚠ **One field in that file will mislead you and it is named here on purpose:** lots_open reads **4**, not 0. It is the size of the LAST position, not the open size, and it is never cleared. closed is the field that decides, and the venue line is what confirms it. (weekend_carry_box() in this build checks lots_open first and is only saved from shouting by the closed test beside it.)

Related: LAST WEEK’S CARD below grades the flatten row this outcome speaks to — and grades it **ungraded**, because the final position was closed by hand more than two hours before either eod-flatten shot, so the flatten path was never exercised this week and nothing here says it works.

◆

THE WEEK THAT WAS, AND THE WEEK COMING

Everything in this section is read from data/event_calendar.json (regenerated 2026-09-18 21:28 UTC) and the times are computed into **Paris** from the New York wall-clock the calendar stores. **The impact column is OUR measurement on OUR tape**, not a general claim about these releases — and where we have not measured one, it says so.

The week that was — 14 to 18 September 2026

One event, and it was the week’s worst day.

DATE	TIME	EVENT	INSTRUMENTS	WHAT IT DID TO OUR TAPE
Wed 16 Sep 2026	20:00 Paris 18:00Z · 14:00 ET	FOMC FOMC decision + projections & press conference	MNQ, MGC	MNQ ran 2.7-3.5× a normal 30-minute window on both FOMC days in our tape (98th/100th percentile); MGC 6.6× its median and 2.8× its previous MAXIMUM

The week coming — 21 to 25 September 2026

★ **NOTHING.** There is not one FOMC, CPI, NFP or tracked earnings release in the week of 21 September. The next scheduled event of any kind is **NFP on Friday 2 October, 14:30 Paris** — 14 days away. An empty calendar week is a real and useful fact rather than a missing table: it means the single measured, repeatable thing this desk knows about the tape — that event windows run several times a normal window — does not apply to a single minute of next week. ⚠ It is not a forecast of a quiet tape. Unscheduled news is not in this file and never will be.

And the seven weeks after that — so nothing arrives as a surprise

Every confirmed date the calendar holds out to 6 November. The **28 October FOMC prints at 19:00 Paris, not 20:00** — the US and EU clocks change three days apart and the table below computes each row rather than offsetting it.

DATE	TIME	EVENT	INSTRUMENTS	WHAT IT DID TO OUR TAPE
Fri 2 Oct 2026	14:30 Paris 12:30Z · 08:30 ET	NFP Employment Situation (Sep)	MNQ, MGC	MNQ 218-220pt in the 30-minute window (100th percentile both times); MGC 3-6×
Wed 14 Oct 2026	14:30 Paris 12:30Z · 08:30 ET	CPI CPI (Sep)	MNQ, MGC	MNQ 205-272pt in the 30-minute window (95th-100th percentile against a 69pt median); MGC 3-6×
Tue 20 Oct 2026	22:05 Paris 20:05Z · 16:05 ET	EARNINGS NFLX earnings	MNQ	NOT MEASURED on our tape. Carried because it moves the index, not because we have tested it — treat as unquantified
Wed 28 Oct 2026	19:00 Paris 18:00Z · 14:00 ET	FOMC FOMC decision	MNQ, MGC	MNQ ran 2.7-3.5× a normal 30-minute window on both FOMC days in our tape (98th/100th percentile); MGC 6.6× its median and 2.8× its previous MAXIMUM
Wed 28 Oct 2026	21:05 Paris 20:05Z · 16:05 ET	EARNINGS TIME UNCONFIRMED TSLA earnings (time UNCONFIRMED — assumed after close)	MNQ	NOT MEASURED on our tape. Carried because it moves the index, not because we have tested it — treat as unquantified
Thu 29 Oct 2026			MNQ	

DATE	TIME	EVENT	INSTRUMENTS	WHAT IT DID TO OUR TAPE
		EARNINGS AAPL earnings 21:05 Paris 20:05Z · 16:05 ET (time UNCONFIRMED — assumed after close) TIME UNCONFIRMED		NOT MEASURED on our tape. Carried because it moves the index, not because we have tested it — treat as unquantified
Thu 29 Oct 2026		EARNINGS AMZN earnings 21:05 Paris 20:05Z · 16:05 ET (time UNCONFIRMED — assumed after close) TIME UNCONFIRMED	MNQ	NOT MEASURED on our tape. Carried because it moves the index, not because we have tested it — treat as unquantified
Wed 4 Nov 2026		EARNINGS GOOG earnings 22:05 Paris 21:05Z · 16:05 ET (time UNCONFIRMED — assumed after close) TIME UNCONFIRMED	MNQ	NOT MEASURED on our tape. Carried because it moves the index, not because we have tested it — treat as unquantified
Wed 4 Nov 2026		EARNINGS GOOGL earnings (time UNCONFIRMED — assumed after close) TIME UNCONFIRMED	MNQ	NOT MEASURED on our tape. Carried because it moves the index, not because we have tested it — treat as unquantified
Wed 4 Nov 2026		EARNINGS MSFT earnings 22:05 Paris 21:05Z · 16:05 ET (time UNCONFIRMED — assumed after close) TIME UNCONFIRMED	MNQ	NOT MEASURED on our tape. Carried because it moves the index, not because we have tested it — treat as unquantified
Wed 4 Nov 2026		EARNINGS META earnings 22:05 Paris 21:05Z · 16:05 ET (time UNCONFIRMED — assumed after close) TIME UNCONFIRMED	MNQ	NOT MEASURED on our tape. Carried because it moves the index, not because we have tested it — treat as unquantified
Fri 6 Nov 2026		NFP Employment Situation (Oct) 14:30 Paris 13:30Z · 08:30 ET	MNQ, MGC	MNQ 218-220pt in the 30-minute window (100th percentile both times); MGC 3-6x

The macro backdrop — and what it is NOT for

★ **READ THIS BEFORE THE TABLE.** These four were measured against our own instruments on **2026-09-18** and the result is a warning label, not a signal: **all four CO-MOVE with the tape and NONE of them predicts the next day.** Same-day, VIX against MNQ is $r = -0.669$ and DXY against MGC is -0.408 ; oil is unrelated to either. Next-day: nothing. **So this table EXPLAINS the week you just had. It does not forecast the one coming,** and no play on the card above is sourced from it. `scripts/macro_watch.py` says the same thing in its own docstring — “IT MEASURES NO EDGE”.

⚠ **And the window is one day, not the week.** `data/macro_watch.jsonl` was started on the morning of 2026-09-18 and holds **14 hourly reads across 1 day(s)** — 07:22Z to 22:52Z. Below is Friday. It is not a week of

context and it must not be read as one; next week's report gets the first full five days. The two same-day correlations quoted above were measured separately on 2026-09-18 and are carried here as a stated result — **this build does not re-derive them, and thirteen hourly rows could not.**

SERIES	LAST	VS PREV CLOSE	RANGE ON THE DAY
Oil CLX6 — front WTI future	97.23	+0.00%	95.06 - 97.88
The dollar DXZ6 — dollar index future	99.94	-0.03%	99.92 - 100.23
The 10-year ZNZ6 — note future — Δ PRICE, so DOWN = yields UP	106.17	+0.00%	105.80 - 106.28
VIX VIX — Δ DELAYED ~15 min on this account	14.81	-4.08%	14.81 - 15.50

★★ IF YOU READ ONLY THIS — FOUR THINGS, IN ORDER

1 — You had the best week this book has ever had, and five trades are the week. \$2,916.49 across 39 entries, every one of them four lots, every one opened by your hand. The previous best was +\$1,617.50 and the week before this one was −\$3,049.49. It also turned the whole book around: the rider was −\$2,591.43 over 53 entries for its entire life and now stands at **+\$325.06 over 92**. ★ **And the same sentence has to carry the caveat:** the best five entries made **+\$3,162.48** — more than the whole week — and the other 34 net −\$245.99. Every per-entry statistic in this report describes a distribution whose mass sits in five draws.

2 — The one scheduled event of the week was the one day that lost real money. FOMC, Wednesday 16 September, **20:00 Paris**. The day booked **−\$833.00** over 6 entries. You were flat through the release itself — and then bought at **T+52, with the slide still running**, for −\$1,272.00. In our own tape an FOMC 30-minute window runs **2.7-3.5×** a normal one on MNQ, and on gold **6.5×** **the median of the same half-hour of the day** (31 prior 18:00-18:30Z windows in capture.db's 5-second MGC bars) — **6.1×** **against all 1,423 prior 30-minute windows**, which is the comparison most readers will assume. Section 4 names the bar source and the baseline for both, and withdraws Rev 1's "largest event response in either tape": NFP on 4 September was bigger. The stand-aside stays on the card as a page and is explicitly refused as a bench, on four event days of evidence — and §3.4(b)

now prices it three ways, because “+\$166.60 a day” is the whole-DAY unit and the whole day is not where the damage was.

3 — Half of the automation already exists and is switched off by one number. leg_watch saw a live leg at **33 of your 34** entries in its window and took your side at 26. **It paged you ahead of 3 of those 34 within half an hour, and 8 within the hour** — every coverage figure on this page is on the **≤60-minute** window unless it says otherwise, and the two figures are one sentence apart here on purpose, because Rev 1 printed them with no window named and the arithmetic looked broken. Its first page fires at 60 minutes of leg age; the median leg dies at 22 and the median leg you press into is 28 minutes old. Move the band to 15 and the same detector precedes **82% of your entries within the hour** against a chance rate of 56% (**P = 0.0009**, and **0.0036** once corrected for winning a four-band sweep — it still clears), where **today's shipped policy reads 24% on that same ≤60-minute window against 18%** and does not separate from chance at all. Apples to apples, both on the hour. $\triangle$ The median 15-minute page lands **26 minutes** before you press, so this buys an early warning, not a punctual one — which is why the row also moves it to the dashboard. That is a config change. MONDAY #2.

4 — ★ And the structural finding, which outranks all three. There are **36 recorded presses and ZERO recorded passes**. The PASS button shipped on Wednesday and has never been pressed. This desk's own fitting threshold is 30 labelled examples and it was written for positives and negatives — so against the threshold that matters the count is **zero**, and no amount of waiting fixes it, because a press only ever produces a positive. **Everything in Sections 1 and 2 is therefore a description, not a rule**, and this report deliberately hands you none. Thirty passes over two weeks changes that. MONDAY #1.

★ THE HONEST HEADLINE — the best week on record, on a sample that cannot yet teach anything

You made **\$2,916.49** this week across **39 entries** and **156 contracts** — **\$583.30 a day** over five sessions. **3 days green, 2 red**, and the worst of the two was FOMC Wednesday.

Three things to hold while reading the rest. One — rows are not entries. The clean week is **61 rows** in trades and only **39 decisions**, because a scale-out books one row per lot leaving. Counting rows would inflate your trade count by **1.56×**. Every number on this page counts decisions where it says entries and rows where it says legs. Two — one fill this week was not real. A 2026-09-15 entry is tagged `BADFILL:paper_engine_0.1pct_fabrication_3of4_lots_20260915`: the paper engine invented the price on three of its four lots. It is **-\$304.50 over 1 row(s)** and it is excluded from every total in this report — but it is not excluded by the live \$250 loss limit, which has no `data_quality` filter and tripped on it (SATURDAY #4). Three — three sections, three denominators, and they are three different instruments. Section 1 counts **39 entries over 5 sessions** from the trade record; Section 2 counts **36 of 78 presses over 4 sessions** from `operator_reads.jsonl`, which only began recording mid-week; Section 3 counts **34 entries over 4 days** because `gazbot7-leg-watch` started at 05:55:43Z on 15 September. ★ **And the one footnote that lets you reconcile them: §2's 36 INCLUDES the BADFILL entry that §1's 39 excludes** — deliberately, because the tape at that moment was real even though the fill price was not. It is why §2.1 lists 7 presses on 09-15 where §1's day table says 6. Every dollar cell in §2.4 is computed on **35** and says so. (78 presses counted by timestamp; 77 records under the *(kind, request.raw)* key, which cannot separate two real unlabelled presses 83 seconds apart. Both are printed in §2.1.) Each section names its own window. **Do not carry a rate across them.**

★ **PASSES RECORDED: 0.** The PASS button shipped on Wednesday 2026-09-18 and data/operator_reads.jsonl holds no pass row. **At an observed rate of zero a week, 30 never lands** — there is no date to put here, and that is the honest form of the answer rather than an extrapolation from nothing. This line is generated from the file at render time and will appear in the same place every Friday, so the count is a fact on the page rather than something somebody has to remember to check. **It is the kill criterion for MONDAY #1, the unblock for BUILD #4’s 78 narratives, and the ceiling on Sections 1 and 2.**

There is no single honest total, so here are all of them, with what separates each from the next.

WHAT YOU ARE LOOKING AT	ROWS	NET \$	WHY IT DIFFERS FROM THE ROW ABOVE
Raw trades rows, 2026-09-14 → 2026-09-18	62	\$2,611.99	everything the book recorded — including the fabricated one
... of which: the BADFILL entry of 15 September	1	−\$304.50	a real decision whose PRICE came from the paper engine on 3 of its 4 lots. Dropped from every total below, and from Section 1's day table
Clean ledger (data_quality IS NULL) — the number this report means	61	\$2,916.49	the week's money. 39 decisions, 156 contracts
... opened by hand (entry_source = 'manual')	61	\$2,916.49	identical — and that IS the finding. Every single clean row this week carries your stamp
... opened by any automated gate	0	\$0.00	the automated desk took nothing this week. That is why this report is about you and not about it

Day by day

DAY	LEGS	NET \$
Mon 14 Sep	9	\$839.49
Tue 15 Sep	12	−\$114.00
Wed 16 Sep FOMC	10	−\$833.00
Thu 17 Sep	12	\$1,477.50
Fri 18 Sep	18	\$1,546.50

⚠ **Legs, not entries** — Section 1's day table counts decisions and reads 4/6/6/8/15. Both are right about different things and the gap is the scale-out. **Three days green, two red** — Tuesday is red at −\$114.00, which is one scratch trade wide, and Wednesday is the day that cost money. (Rev 1 of Section 1 said

“four green days and one red” in its summary callout while its own table said three and two; the callout has been corrected, so the amber banner that used to stand on that section is gone.)

How the money left the book

EXIT	LEGS	NET \$	PER LEG
MANUAL_CLAIM	44	\$2,247.62	\$51.08
TARGET_100	14	\$1,528.62	\$109.19
TARGET_200	2	\$428.25	\$214.12
CLOCK_FLAT	1	-\$1,288.00	-\$1,288.00

MANUAL_CLAIM is your own hand on the button; TARGET_100 and TARGET_200 are the per-lot profit ladder; CLOCK_FLAT is the 20:40Z timer, and it appears exactly once — Tuesday's seven-hour long, -\$1,288.00, the single worst decision of the week and the only position all week that was closed by a machine rather than by you.

★ WHAT TO ACTUALLY DO

Garrath — this is the whole answer to “so what do I actually DO?”, and it is at the front. Every row below mirrors reports/friday_v7/plays.json, which is the reconciled source of truth for this week; the same file drives the [Monday playbook](#), so the two cannot drift apart. ★ **Almost nothing on this card changes how you trade, and that is the honest answer to this week.** You had the best week this book has ever had — and a week is 39 entries, five of which ARE the week. The automated desk took **nothing**, so there is no gate to re-tune and no threshold that this week has earned the right to move. What the card is instead: **one habit, one config change, six instrument faults, six measurements and four refutations. 4 of the rows are NOT-AN-ACTION:** refutations printed with their autopsies, because on a week like this the most valuable output is a list of things that have now been ruled out. ★ The first of them is the idea you are most likely to want after a week like this — only trade the

European morning — and it dies on its own sample: +\$2,445.51 over 20 entries this week, −\$779.48 over the record, one trade from changing sign. 22 plays, five windows, in the order they need doing. Every play carries its **evidence tier** and its **kill criterion** — if a play has no way to be proven wrong it should not be on the card.

THE THREE THAT MATTER MOST

1 · ★ drift is blind before 13:30Z — and that is two thirds of his trading (SATURDAY #1). Anchor the drift detector's day to the CME session (22:00Z), or fall back to a trailing 24 hours — which is what `tape_reader` already does for its bars. Until this is fixed the automated path cannot even LOOK at most of his entries. **51 of 78 presses / 24 of 36 entries** returned no rows. 0 of 78 returned a confirmed drift read.

2 · `tape_reader`'s “day open/high/low” splices yesterday afternoon onto this morning (SATURDAY #2). Anchor `tape_reader`'s session window to 22:00Z, the same anchor SATURDAY #1 asks for. This is about what you are SHOWN, not about any conclusion in this report. **203 bars on 2026-09-14** labelled “session”, spanning an 11-hour gap. No conclusion in §2 changes when corrected (checked).

3 · ★★ PRESS THE PASS BUTTON. This is the highest-value thing on the desk and it costs you one tap (MONDAY #1). Every time you look at the tape and do NOT trade, press PASS. Not sometimes — every time, including the times it feels obvious. Roughly 30 passes over the next two weeks and this desk can, for the first time in eleven months, learn what separates a moment you take from a moment you leave. **36 positives, 0 negatives.** 78 presses across 4 sessions. Negatives from the alternative source (declined *leg_watch* pages, BUILD #1) accrue at ~21/week against ~7/week for positives, so positives are the binding side either way and the pool reaches 30/30 in about **3.6 weeks** — or roughly one week if MONDAY #2 (drop-the-page-band-to-15-0918) ships alongside it.

Generated from the top of the card below, so the two cannot drift apart. Everything else, including every refuted idea and why it died, is in the five windows underneath.

MONDAY · MONDAY 2026-09-21, AT THE DESK — ★ THE FIRST ROW IS A HABIT, NOT A TRADE, AND IT IS THE WHOLE REPORT · 3 PLAYS

★★ ALMOST NOTHING ON THIS CARD CHANGES HOW YOU TRADE, and that is the honest answer to this week. You had the best week this book has ever had — and a week is 39 entries, five of which ARE the week. Nothing in that is enough to hand you a rule. So the three rows here are: press a button you are not pressing, change one number in an alert that is already running, and keep doing the event-day thing you already do. ★ MONDAY #1 is ranked first with no dollar figure attached, deliberately. There are 36 recorded presses and ZERO recorded passes on this desk. Every behavioural finding in Sections 1 to 3 is therefore one side of a decision boundary, and a description is not a rule. Thirty passes changes that and nothing else does.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
1	★★ PRESS THE PASS BUTTON. This is the highest-value thing on the desk and it costs you one tap BLOCKER Every time you look at the tape and do NOT trade, press PASS. Not sometimes — every time, including the times it feels obvious. Roughly 30 passes over the next two weeks and this desk can, for the first time in eleven months, learn what separates a moment you take from a moment you leave. Mechanism: data/operator_reads.jsonl holds 78 recorded presses, 36 of them entries, and 0 passes. The PASS button shipped on Wednesday and has never been pressed. Every threshold, band and description in Sections 1, 2 and 3 is therefore fitted on positives only. Why: The desk's own written fitting threshold is 30 labelled examples, and it was written for positives AND negatives. On that count the real figure is zero, not 36 — and no amount of waiting fixes it, because a press only ever generates a positive. This is the one row on the card that is genuinely blocking: Section 2's continuation description cannot become a filter without it, Section 1 cannot promote a single observation to a rule without it, and Section 3's decision half stays at "no" until it exists. Evidence: \$2.1 (the sample), \$2.5 lead 1 (HIGHEST VALUE ON THE DESK); \$1 \$8 row 1 (BLOCKER); \$3.3 (what is not automatable, and exactly how far off). · verification: <code>python3 -c "import json;print(sum(1 for l in open('data/operator_reads.jsonl') if json.loads(l)['kind']=='pass'))"</code> — it reads 0 today. Any number above zero next Friday is this row working. · mode: do it	36 positives, 0 negatives. 78 presses across 4 sessions. Negatives from the alternative source (declined leg_watch pages, BUILD #1) accrue at ~21/week against ~7/week for positives, so positives are the binding side either way and the pool reaches 30/30 in about 3.6 weeks — or roughly one week if MONDAY #2 (drop-the-page-band-to-15-0918) ships alongside it.	Arms: Nothing arms. It is a habit, and it is the cheapest item on this card. Exit: n/a. Kill: If 30 passes land and the press/pass boundary STILL does not separate on any of the features in \$2, that is a real answer and this desk stops trying to model entry selection. That outcome would be worth knowing and is currently unreachable.	Garrath — at the desk, every session n/a. ★ Nothing else in this report is unblocked without it, which is why it is #1 despite having no dollar figure attached.
2	Drop leg_watch's first page from 60 minutes of leg age to 15 — and route it to the dashboard, not the phone MEASURED Change one number and one destination. The detector is already right; it is speaking forty minutes after you have acted. At a 15-minute band the same detector, unchanged in every other respect, precedes 82% of your entries. Mechanism: leg_watch saw a live 1×ATR leg at 33 of your 34 entries in its window and took your side at 26 of them. It told you about 3. Its first page fires at 60 minutes of leg age; the median leg dies at 22 minutes and the median leg you press into is 28 minutes old. The machine is watching the right thing at the wrong time. Why: This is the only row on the card with a significance test behind it and no model in it. Scored against a session-blocked permutation of your own presses: the policy running today precedes 24% of entries	82% coverage at P = 0.0009 — 0.0036 sweep-corrected, against 24% at P = 0.24 today, both on the ≤60-minute window. Price: your phone goes from 7 to 32 pages a day — which is the entire reason for the dashboard half of this row. Δ Median page lands 26.5 minutes ahead and only 10 of 29 inside twenty; demand a page inside 20 min and coverage falls to 29%.	Arms: On the config change. Nothing about the detector's definition moves. Exit: n/a — it is an alert, never a filter and never an order. Kill: PRE-REGISTERED BEFORE IT SHIPS — band 15 min, every leg, no travel floor, destination dashboard. Measured quantity: fraction of entries preceded by a SAME-SIDE page within 60 minutes, on entries opened after the change lands. Target ≥56% (the shuffled chance rate at this band, not the 82% in-sample point). Kill:	Claude — one config change to the leg-watch band + dashboard route One number in the leg-watch config. ★ No fitting on your judgement, no negatives required, and it is reversible in a minute.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	within the hour against a chance rate of 18%, P = 0.24 — no separation at all. At 15 minutes it is 82% against 56%, P = 0.0009. &#9733; That P is the winning band's and pays nothing for having won a four-band sweep — corrected over the four (60/30/15/10) it is 0.0036, which still clears comfortably. 10 minutes reads 0.015 raw and 0.060 corrected, so it does NOT survive; 30 minutes clears neither way (0.43). That is an edge between 15 and 30, not a plateau, so the band matters and 15 is the measured end of it. The live leg-watch journal reaches 85% independently of the replay. Evidence: §3.1 (delivery, not detection), §3.2 (the page band, priced, with the null and the sweep-wide correction), §3.2b (how far ahead a page lands and whether the leg is still alive); §3.5 row 1. · verification: <code>scripts/leg_survival.py</code> for the leg census; the permutation is re-runnable against data/operator_reads.jsonl and the gazbot7-leg-watch journal. Freeze the band on Monday and count coverage forward. · mode: ship it Monday		forward coverage below 50% over the first 30 forward entries, or you stop reading the pages — either puts the band back to 60 and the row dies. Re-scored Friday 2026-10-16. The band, the 60-minute window, the same-side rule and the target may NOT be re-chosen; changing any of them turns the forward count back into a second fit. Full text in &sect;3.2 under PRE-REGISTRATION. Uncertainty that remains: 34 presses, four days, and the 15 was chosen off this week's own curve.	
3	Keep standing aside on event days — as a PAGE, never as a bench MEASURED Leave the calendar and the T-60/T-15 pages exactly as they are and keep taking them seriously. They are already built and already running. What this row asks for is that they stay a prompt to you and never become an automatic bench. Mechanism: data/event_calendar.json plus the T-60/T-15 pager. The report week held exactly one calendar event — FOMC, Wednesday 2026-09-16, 20:00 Paris — and it was the week's worst day. Why: FOMC Wednesday was −\$833.00 over 6 entries and the only day in the week that lost more than a hundred dollars. Over the book's life event days run −\$1,518.49 across 12 entries against +\$1,843.55 over 80 ordinary ones. And the run charts show the mechanism plainly: he was flat through the release itself and then bought at T+52 with the slide still running for −\$1,272.00. Evidence: §1 §5 (by event day) and §1 §8 row 3; §3.4(b), where it is now priced three ways; §3.5 row 2; §4.1 row 2, which reconciles this row with Section 4's gold lead. · verification: <code>data/event_calendar.json</code> is regenerated and human-verified; the event-day split is re-derivable from trades against it. · mode: carry on	+\$166.60 a day this week for the whole-DAY unit — but that is one of three answers and it prices only the upside. T&plusmn;60 around the release: +\$254.40/day (it removes the −\$1,272.00 and keeps the +\$819.00). After-the-release only: +\$164.40/day. Over the book the order reverses — +\$52.36 / +\$43.86 / +\$20.60 a session. FOMC ran MNQ at 2.7-3.5× a normal 30-minute window (98th/100th percentile) and MGC at 6.5× the median of the same half-hour of the day, 6.1× against all windows — capture.db 5-second bars, named because a gold number without its bar source is not a number.	Arms: Already armed. The pages fire on the calendar. Exit: n/a. Kill: Four event days, one of them positive at +\$224.50, and the split was chosen after seeing the losses. The Wednesday damage was ONE entry at T+52, not the session — and the entire difference between the three units above is that same single trade, so none of them is a window study. &#9733; And the cost of being wrong is now priced: a whole-day stand-aside deletes \$2,267.51 of winners across the four event days, including the week's biggest (+\$819.00) and 12 August outright — a winning day whose only entry was +\$224.50. Nowhere near strong enough to become a bench. If the next four event days are net green, it dies.	Garrath — it is a page, and the decision stays yours n/a. ★ Note what this row deliberately does NOT say: it does not say block the book, because a bench can only ever remove trades and this book's event-day sample is four days deep.

SATURDAY · SAT 2026-09-19 / SUN 2026-09-20 — BEFORE THE SUNDAY 22:00Z REOPEN. ★ NOTHING HERE HAS A DEADLINE, BECAUSE THE DESK IS FLAT · 6 PLAYS

★ **No row in this window has a clock on it.** The last position closed at 18:33:23Z on Friday and the desk went into the weekend flat — the front page states the evidence for that rather than asserting it. Contrast the 2026-09-04 card, whose top SATURDAY row (rank 0, last card) was four lots riding an unreachable gateway across a weekend. **What is here is six instrument faults**, and not one of them moves a threshold. They are all the same species: reasons a number in this report, or a number on your screen at the desk, may be wrong. Two are a session-anchor bug in different clothes (a drift detector that is switched off for two thirds of your trading, and a “day high” that includes yesterday afternoon). Three are in the \$250 loss limit, which is listed on the “already automated” side of Section 3 and turns out to page you a claim its own journal contradicts, count fabricated fills towards its trigger, and be worth — \$67.09 a session on your book — see NOT-AN-ACTION #4.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
1	★ drift is blind before 13:30Z — and that is two thirds of his trading FIX — DEFECT Anchor the drift detector's day to the CME session (22:00Z), or fall back to a trailing 24 hours — which is what tape_reader already does for its bars. Until this is fixed the automated path cannot even LOOK at most of his entries.	51 of 78 presses / 24 of 36 entries returned no rows. 0 of 78 returned a confirmed drift read.	Arms: On the code change. No threshold moves. Exit: n/a. Kill: None needed — it is a defect, not a hypothesis. △ But note what fixing it does NOT do: it does not make drift predictive. It	Claude — one session anchor in the drift reader Revert the anchor. ★ Related, same species, same fix: SATURDAY #2.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	Mechanism: <code>drift.read</code> is gated to 13:30–21:00Z. Outside that window it returns <code>no 5s rows since 13:30Z today</code> and the caller carries on. Why: This is the biggest single instrument fault in the report, measured on his own record: 51 of 78 presses and 24 of 36 entries returned the blind reading. The other 12 all read <code>confirmed=false</code>. So for the majority of his trading the desk's drift signal is not wrong — it is absent, and absent reads as neutral to everything downstream. He trades the European morning heavily (20 of 39 entries this week were before 13:00Z); the detector is switched off for all of it. Evidence: §2.5 lead 3 (FIX — DEFECT). · verification: Call <code>drift.read()</code> at 08:00Z on Monday. Today it returns <code>no 5s rows since 13:30Z today</code>; after the fix it returns a real read. · mode: fix it		makes drift measurable outside RTH, which is the precondition for ever finding out.	
2	tape_reader's “day open/high/low” splices yesterday afternoon onto this morning FIX — INSTRUMENT Anchor <code>tape_reader</code>'s session window to 22:00Z, the same anchor SATURDAY #1 asks for. This is about what you are SHOWN, not about any conclusion in this report. Mechanism: The “session” window does not start at the CME reopen, so it carries the previous afternoon's bars into the current morning's day-open, day-high and day-low. On 2026-09-14 its 203 “session bars” were Sunday night plus Monday morning spanning an eleven-hour hole. Why: Position-in-range is REFUTED as an entry condition (NOT-ACTION #2), so no conclusion in this document moves when this is fixed — that was checked, and it is why this is #2 and not #1. What it changes is the number on his screen: a day-high that includes yesterday afternoon is not the day's high, and he reads it live. Evidence: §2.5 lead 8 (FIX — INSTRUMENT). · verification: Print <code>tape_reader</code>'s session bar count and first bar timestamp at 07:00Z Monday; the first bar should be 22:00Z Sunday, not Friday afternoon. · mode: fix it	203 bars on 2026-09-14 labelled “session”, spanning an 11-hour gap. No conclusion in §2 changes when corrected (checked).	Arms: On the code change. Exit: n/a. Kill: None — it is a defect. ★ It is ranked below SATURDAY #1 precisely because nothing in this report depends on it; that is the honest reason, not a judgement about severity.	Claude — one session anchor in <code>tape_reader</code> Revert the anchor.
3	★ The \$250 daily loss limit pages you a claim its own journal contradicts FIX — DEFECT Drive the notification sentence from the changed list the guard actually wrote, not from a hard-coded string. A guard that reports an action it did not take is worse than a guard that stays quiet. Mechanism: The notify string reads “All six gates benched”. The journal for the same firing reads “benched”: <code>[]</code>. Every gate the limit can reach was already <code>=off</code>, so the guard is structurally incapable of doing the thing it announces. Why: It fired twice this week, benched zero gates both times, and paged him both times saying it had benched six. This is the instrument-reports-healthy failure in its purest form: the page is the only evidence a human sees, and it is manufactured. Evidence: §3.4(c) and §3.5 row 6. · verification: Grep the notify call site for the literal string, then compare against the benched/changed key written to the same journal record. · mode: fix it	Fired 2x, benched 0, paged 2x claiming six.	Arms: On the code change. Exit: n/a. Kill: None — it is a defect.	Claude — one f-string driven from the guard's own return value Revert the string. ★ Same subsystem as SATURDAY #4 and NOT-ACTION #4; fix them together.
4	★ The loss limit's trigger can be moved by a price that never printed FIX — DEFECT Put a <code>data_quality IS NULL</code> filter in <code>session_pnl()</code>. It is the desk's standing convention everywhere else and this reader was missed. Mechanism: <code>session_pnl()</code> sums every closed row for the session with no <code>data_quality</code> filter, so a row the desk has already marked known-bad counts towards a guard that can halt trading. Why: Tuesday's trip read −\$418.50. That is −\$114.00 real and −\$304.50 fabricated by the paper engine, which invented the price on 3 of that entry's 4 lots. On the real number the session was below the limit and the guard should not have fired at all. Evidence: §3.4(c) and §3.5 row 7; the BADFILL row is named in §1 §1 and excluded from every total in this report. · verification: <code>sqlite3 data/gazbot7.db "SELECT data_quality FROM trades WHERE id=895" → the BADFILL:paper_engine_0.1pct_fabrication_3of4_1ots_20260915 tag. Then re-run <code>session_pnl()</code> for that date.</code> · mode: fix it	Tuesday 2026-09-15: trigger read −\$418.50 = −\$114.00 real + −\$304.50 fabricated.	Arms: On the code change. Exit: n/a. Kill: None — it is a defect.	Claude — one WHERE clause Revert the clause.
5	A hand-run of <code>event_alert.py</code> wrote a live PAGED line into the record	1 fabricated PAGED line; would have inflated the §3.2 alerting score by 25%.	Arms: On the flag. Exit: n/a. Kill: None — it is a defect.	Claude — one <code>argparse</code> flag Drop the flag.

	FIX — DEFECT Give event_alert.py a --dry-run that logs nowhere, or stamp manual runs so they can be filtered out. Mechanism: 2026-09-18T07:04:03Z PAGED T-15 FOMC sits in the log — off the *:0/5:30 timer cadence, two days after the only FOMC in the calendar, with the timer's own run at 07:05:30 logging “nothing in a band” thirty seconds later. Why: It is small and it is exactly the kind of small that corrupts a score. It cost this report one false LED and would have inflated the alerting measurement in §3.2 by 25% had it not been caught — and §3.2 is the evidence behind MONDAY #2. Evidence: §3.5 row 8. · verification: Diff the PAGED timestamps against the timer's *:0/5:30 cadence — a line that is not on the cadence was hand-run. · mode: fix it			
6	DEFECT A TARGET_100 rung booked a fill 79.5 points BELOW the entry it was supposed to take \$100 a lot above Read trade id 868 and find out what price that exit actually saw. Small money; a target exit reporting a price its own trigger cannot have produced is not small. Mechanism: Trade id 868, 2026-09-11: long at 29,320.25, out at 29,240.75 on a TARGET_100 exit — −\$160.50. A \$100-a-lot rung on MNQ is +50 points; this booked −79.5. Why: It is the only one. There are 25 TARGET_100 exits in the whole clean record and 24 of them booked between +\$97.50 and +\$132.25 — including all fourteen in the report week. This one landed 129.5 points from the level it names. One row, and it is outside the report week, which is why it is last in this window. Evidence: §1 §8 row 6 (DEFECT); §1 §6 for how the week's exits ended. · verification: sqlite3 data/gazbot7.db "SELECT * FROM trades WHERE id=868", then the rider journal for that timestamp. · mode: investigate	1 row of 25, −\$160.50, against 24 that booked +\$97.50 to +\$132.25 (14 of them this week).	Arms: Nothing arms. It is an investigation. Exit: n/a. Kill: If it turns out to be a mislabelled manual claim, close it and move on. n=1.	Claude — one journal read n/a.

BUILD · THE MEASUREMENT QUEUE — SEVEN ROWS, AND THE FIRST ONE IS THE ONLY THING THAT UNBLOCKS MONDAY #1 WITHOUT WAITING · 7 PLAYS

This window buys information, not P&L. BUILD #1 is the important one: every leg_watch page you see and decline is a negative example, which is the one thing this desk has none of. Five positives and fifteen negatives are already recoverable from pages already sent. ★ **It supports MONDAY #1 rather than replacing it** — negatives accrue at ~21/week and positives at ~7/week, so positives are the binding side either way and the PASS button is still the fastest route. Below it sit a stopwatch that would have earned **\$0.00 this week** and is on the card anyway (BUILD #2, and the row says so in its own kill line), a standing hour-of-day figure that has already dispersed, 78 captured narratives that are deliberately left UNSPENT until passes exist, and two instrument fixes. ★ BUILD #7 is new this revision and is a PARK, not a build: Section 4 floated trading the event window on GOLD and never priced it, so it is priced here — 1.97× MNQ's dollar range to a 1-lot and a coin-flip direction on four events — and parked with the condition that would revive it, so it stops being re-floated every time an event day is red.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
1	★★ Turn the paged legs into LABELS — every page you decline is a negative example ENABLER Record the outcome of every leg_watch page while your phone is awake. A page you see and do not act on is a negative, and it is free. This is the only item on the card that can unblock MONDAY #1 without waiting on the PASS button. Mechanism: leg_watch already pages. What is missing is the disposition — the join between a page and whether an entry followed it inside a window. Why: The dataset already half exists: 5 positives and 15 negatives are recoverable from the 25 pages actually sent, plus 5 more that nobody saw. Negatives accrue at 21/week and positives at 7/week,	5 positives / 15 negatives already held. 30/30 in ~3.6 weeks at the current band, ~1 week at a 15-minute band. Direct P&L: \$0/day — this buys a dataset, not a trade.	Arms: Nothing arms. It is a log. Exit: n/a. Kill: If the negatives turn out to be indistinguishable from the positives on every feature in \$2, that is the same real answer MONDAY #1's kill line describes, reached sooner.	Claude — one disposition log joined to the leg_watch pages Delete the log. ★ It writes nothing to the order path.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	so positives are the binding side — which is why this row supports MONDAY #1 rather than replacing it. At the current 60-minute band the pool reaches 30/30 in about 3.6 weeks; at MONDAY #2 (drop-the-page-band-to-15-0918)'s 15-minute band, in roughly one. Evidence: §3.5 row 3 (ENABLER); §3.3 for what the missing negatives actually block. · verification: Count rows in the disposition store next Friday. It does not exist today. · mode: build it			
2	The eight-hour stopwatch — a hard flat-out at 8h, which needs no model and no judgement MEASURED Build the stopwatch, but ship it knowing this week says it would have earned nothing. It is a clock, not a rule about your judgement: no fitting, no negatives, no threshold on the tape. Mechanism: A hard flat at eight hours of hold time, on top of the existing 20:40Z clock. Why: Over the book's life 4 entries held past eight hours, 0 winners, −\$6,612.44 — and every other entry ever made adds to +\$6,937.50. That is a −\$228.02 drag per session traded. It did not recur this week: 0 of 39 entries reached eight hours, longest 7h01m. So on this week's tape it earns \$0.00 and is pure insurance. Evidence: §1 §3 (by hold time) and §1 §8 row 2 (PROMISING); §3.4(a), where the recovery number §1 deferred is actually computed. · verification: Re-run the 8h sweep against trades; the four cases are enumerable by id. · mode: build it, expect nothing this week	Lifetime it recovers +\$1,720.00 over 28 sessions = +\$61.43/ session, which is 26% of the bucket's −\$6,612.44 headline — not all of it. This week: \$0.00/day.	Arms: On the clock. It closes a position at 8h regardless of P&L. Exit: The stopwatch IS the exit. Kill: Four cases, and 68% of the total is one trade. One case gets \$202 worse under the sweep, and at the neighbouring widths where the sweep had work to do this week (4h, 6h) it LOST — \$81.60 and −\$186.40 a session. If the next 8h hold is a winner, this row is in real trouble and should be said so out loud.	Claude — one timer in the rider's manage path One constant. ★ The honest summary is “cheap insurance against a shape that did not recur”, and it is ranked second, not first, for exactly that reason.
3	Re-derive the 05:00-11:00Z window before any hour-of-day rule ships NEEDS REFRESH Recompute it on the current sample and retire the standing figure. The old number was a property of a sample that no longer describes how he trades. Mechanism: The standing figure is 26 of 28 presses inside 05:00-11:00Z. This week it is 16 of 36, spread across 18 distinct hours. Why: He has widened out. Any hour-of-day filter built on the standing figure would be fitted to a concentration that has already dispersed — and NOT-AN-ACTION #1 shows what happens when a time-of-day result from one week is promoted to a rule. Evidence: §2.5 lead 9 (NEEDS REFRESH); §1 §4 (by hour of day). · verification: Re-derive from data/operator_reads.jsonl and trades at 70 entries. · mode: measure	26/28 → 16/36, across 18 distinct hours.	Arms: Nothing arms. It is a measurement. Exit: n/a. Kill: n/a — it is a refresh, and the outcome is the answer either way.	Claude — one re-derivation n/a.
4	The 78 captured narratives — a ~3,160-character read of the whole session is stored with every press and nothing has been done with it UNSPENT Hold them until passes exist, then pair them. The narrative pairs are the actual experiment: the same kind of text, one ending in a press and one in a pass. Mechanism: Every row in data/operator_reads.jsonl carries a full natural-language read of the session at the moment of the press. 78 of them are stored. None has been used for anything. Why: This is the richest thing on the desk and it is currently worth nothing, for exactly the reason MONDAY #1 exists: 78 narratives that all end in “he pressed” cannot teach a boundary. The moment passes start landing, the paired comparison becomes the single best-powered test available here — same instrument, same author, same format, opposite outcome. Evidence: §2.5 lead 10 (UNSPENT); §2.1 for the sample. · verification: Count rows with a non-empty context field, deduped on (kind, ts) — it is 78 today and every one is a positive. The PASS count is printed in the front box of this report, generated from the file at render time. · mode: hold, then measure	78 narratives, ~3,160 characters each, 0 spent.	Arms: Nothing arms. Exit: n/a. Kill: n/a — it is deliberately deferred, and saying so is the point of the row. △ Do NOT mine them for features now: 78 positives with no negatives will produce a confident model of nothing.	Claude — after ~30 passes land n/a. ★ Blocked on MONDAY #1.
5	The scope's 3×ATR leg definition is a test that everything passes REFUTED Move the written definition to 1×ATR. leg_survival.py already defaults to it; the scope document has not caught up. Mechanism: A 3×ATR leg is alive in 97.6% of ALL minutes. A condition that 97.6% of minutes satisfy carries no information. Why: This matters because every “a leg was running” number in §2 and §3 is computed on the 1× definition, and a reader checking it against the scope's 3× text would conclude the sections had used a loose filter. They used the tighter one. The document is what is wrong.	97.6% of all minutes at 3×; the 1× definition (1×ATR travelled in 8 minutes, dead on a 1×ATR give-back) is alive in 74.6%, which is what §2's P = 0.008 is measured against.	Arms: Nothing arms. It is a documentation fix. Exit: n/a. Kill: n/a.	Claude — one line in the scope n/a.

Evidence: §2.5 lead 7 (REFUTED); §2.3 for the definition actually used.
· verification: `scripts/leg_survival.py` — its default is already `1x`.
· mode: fix the doc

6 operator_reads.jsonl must be de-duplicated before anything counts it FIX — INSTRUMENT De-duplicate at read time, key on (kind, ts), and write the rule down beside the file. Counting lines instead of records nearly doubled the sample — and the key the report has been quoting silently loses a real press. Mechanism: The file holds 135 lines. Under (kind, request.raw) that is 77 records; counted by the presses' own timestamps it is 78. The gap is exactly one thing: two unlabelled buys on 2026-09-14, at 14:51:43 and 14:53:06 — 83 seconds apart, so two real presses — which carry an empty request and so collapse into one record under the named key. Why: Counting lines would have reported 65 buy rows where there are 38 buy presses. Rev 1 of the report printed 77 in Section 1's housekeeping note and 78 in §2.1 under a key that produces 77, and then carried the 78 through eight more figures — in the one quantity this very row warns that “every sample size on this card moves if this is got wrong”. Both counts are defensible; using both without saying so is not. Keying the helper on (kind, ts) makes 78 the answer to every phrasing of the question. Evidence: §1's closing housekeeping note to §2 (both keys); §2.1, the sample ladder. · verification: Count lines against distinct (kind, ts) and against distinct (kind, request.raw); they differ by one and the row above says which one. △ Whole-row JSON equality is NOT the de-dup key — all 135 lines differ on it; the rule has to be written down rather than rediscovered. · mode: fix it	135 lines & 78 presses by timestamp, 77 records by (kind, request.raw); 65 buy rows & 38 buy presses. Entry presses are 36 under either key.	Arms: Nothing arms. Exit: n/a. Kill: n/a — it is a defect.	Claude — one reader helper beside the file n/a.
7 Section 4's gold lead, priced — and then PARKED, with the condition that would revive it PARKED Nothing on Monday. The lead is recorded with a number and a revival condition so it stops being re-floated informally every time an event day is red. Mechanism: Section 4's gold caption floated turning the event-day stand-aside into an event-day OPPORTUNITY on MGC and stopped there — no row on this card carried it, §4 had no shortlist, and MONDAY #3 said the opposite with no cross-reference. It is now measured: the 30-minute window from each of the four event releases, on both instruments, at MNQ \$2.00 and MGC \$10.00 a point, from capture.db 5-second bars. Why: The size is real and the direction is not. Gold's event window is worth 1.97x MNQ's to a 1-lot on the hindsight range — \$813.00 an event against \$413.50 — because a smaller point move carries a five-times multiplier. But a 1-lot long held through the release loses on both tapes: −\$660.00 on MGC and −\$114.00 on MNQ across the same four events, two up and two down. And the same two: MNQ and MGC moved the SAME direction on all four releases, so gold is not a second opinion, it is the same bet with a bigger multiplier. Evidence: §4.1 row 1 (the ranked shortlist Section 4 did not have in Rev 1); §3.4(b). · verification: Re-runnable from data/capture.db 5s bars for both symbols against data/event_calendar.json. △ MGC is NOT in the parquet lake — the lake holds MNQ only — so MGC history is capped at 2026-08-04 and every gold figure here must name capture.db as its bar source. · mode: park it	Four event days, 1 lot, 30 minutes from the release. Hindsight range: MNQ \$1,654.00 total / \$413.50 an event; MGC \$3,252.00 / \$813.00. Held-long move: MNQ −\$114.00, MGC −\$660.00.	Arms: Nothing arms. It is parked. Exit: n/a. Kill: Revive only when BOTH hold: (a) ten or more event days of MGC tape exist — roughly 2026-02 at the current calendar rate — AND (b) something that is not a direction call is written down first, because four events say the direction is a coin flip and the money is in the range. Uncertainty: four events; MGC tape starts 2026-08-04; the desk has never sent an order to gold and holds zero MGC execution record; and the whole range column is a hindsight ceiling nobody captures. △ Do NOT invent an MGC gate for this before then, and never by porting an MNQ one.	Claude — re-measure when the tape reaches ten event days n/a — nothing ships.

HOLD · HOLD — TWO FINDINGS THAT ARE TRUE, INTERESTING, AND MUST NOT BE FITTED · 2 PLAYS

Both rows here are things it would be easy and wrong to act on. The first is the cleanest description in the whole report — your losers are wrong from the first minute, 11.2 points of best-moment against 59.2 for your winners — and turning it into a bail rule means backtesting an exit, which the standing rule forbids. Its kill criterion is unusual and honest: it is on the reader, not on the data. The second is the one place in Section 2 where two legitimate readings

of the same 36 trades point in opposite directions, on 12-trade cells. Re-measure at 70 entries; write no band.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
1	★ Your losers are wrong from the first minute — log it, watch it, do NOT fit it WATCH Carry this as a description and nothing more. It is the cleanest finding in the report and turning it into a bail rule means backtesting an exit, which the standing rule forbids. Mechanism: Measured off the 5-second tape from fill to final lot out, per decision. Why: Your 25 winners had a median best-moment of 59.2 points in your favour and you booked a median 70% of it. Your 14 losers had a median best moment of 11.2 points; 11 of 14 never once showed you 25 points and 7 never showed ten. You are not giving winners back — you are occasionally getting on the wrong side and staying there. Evidence: \$1 \$2 (every entry, with best/worst/kept) and \$1 \$8 row 5 (WATCH). · verification: Re-derivable per decision from the 5s tape against trades. · mode: watch	59.2pt vs 11.2pt median best-moment; 70% of the best moment kept on winners; 11 of 14 losers never showed 25 points.	Arms: Nothing arms. Exit: n/a. Kill: n/a. △ The kill criterion for this row is on the person reading it: the standing rule forbids a 120th exit calibration, and this shape is precisely the bait for one. If it is still true at 100 entries it becomes a candidate; at 39 it is a picture.	Garrath — noticing, not acting n/a.
2	The volatility preference — the description and the dollars point opposite ways OPEN — CONFLICTED Write no ATR band. Re-measure at 70 entries. This is the one place in \$2 where two honest readings of the same 36 trades disagree, and the sample cannot settle it. Mechanism: ATR percentile at press vs at a randomly chosen control minute, and post-hoc terciles of the same 36 trades by ATR. Why: He presses at ATR percentile 71.7 against a baseline of 49.3 (within-session rank 62.9, CI [54.0, 71.4]) — a real preference for a lively tape. But the money went the other way: the quiet tercile made +\$2,762 and the lively one −\$1,262. Those are 12-trade cells, chosen after the fact, and the medians barely move. Evidence: \$2.4 (what the dollars did — post-hoc, three-deep, and not a rule) and \$2.5 lead 6. · verification: Re-run the terciles at 70 entries. · mode: hold	ATR pct 71.7 vs 49.3; quiet tercile +\$2,762, lively tercile −\$1,262, on 12-trade cells.	Arms: Nothing arms. Exit: n/a. Kill: Either reading could die at 70 entries and that is the point of holding it. Acting on the dollar half today would be fitting a 12-trade cell.	Claude — re-measure at 70 entries n/a.

NOT-AN-ACTION · NOT-AN-ACTION — FOUR THINGS NOW RULED OUT, PRINTED WITH THEIR AUTOPSIES · 4 PLAYS

On a week this thin these are the most valuable output in the document. NOT-AN-ACTION #1 is the most tempting idea in the report and the one you are most likely to want after a week like this: only trade the European morning. It was +\$2,445.51 over 20 entries this week and it is **−\$779.48 over the record, where one trade on 10 September flips its sign**. Two more (position-in-range, VWAP stretch) are the numbers everyone expects to matter, measured against 5,255 control minutes on two independent rulers, and both confidence intervals contain 50 — he is neither buying breakouts nor fading them. The fourth withdraws a guard from this book. **Every one of these is scoped:** “refuted as a separator of his presses from random minutes” is not “meaningless”, and the rows say which they mean.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
1	★ “Only trade the European morning” — the most tempting idea in the report, and it does not survive its own sample REFUTED Do NOT build this filter. This week it looks like the whole edge. Over the record it is negative, and a single trade flips its sign. Mechanism: Split the record at 13:00Z and compare. Why: This week, 05:00–12:59Z was +\$2,445.51 over 20 entries and after 13:00Z was −\$553.02 over 16. Over the whole confirmed-manual record the SAME morning window is −\$779.48 — and one trade on 10 September flips its sign. A filter fitted on a week that one trade can reverse is not a filter.	+\$2,445.51 / 20 this week; −\$779.48 lifetime, one trade from changing sign.	Arms: Nothing arms — that is the point. Exit: n/a. Kill: It is already killed. ★ If it ever comes back, it comes back through BUILD #3’s re-derivation at 70 entries, not through another good morning.	n/a n/a.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	Evidence: \$1 \$4 (by hour of day) and \$1 \$8 row 4 (REFUTED as a rule). · verification: Re-derive the split from trades with entry_source='manual' . · mode: do not build			
2	Position-in-range as an entry condition — the number everyone expects to matter, and it does not REFUTED Stop building it. Two independent rulers, 5,255 control minutes, one answer. Mechanism: Rank the press inside its own session's range, against control minutes from the same sessions. Why: Ranked inside its own session the median press sits at the 54.2nd percentile, CI [44.6, 63.5] ; on a clean-session ruler, 56.7 [47.5, 66.2]. Both intervals contain 50. He is not buying breakouts and he is not buying dips — where price sits in the day's range does not separate his presses from picking a minute at random. Evidence: \$2.2 (the one comparison legitimate without negatives) and \$2.5 lead 4. · verification: Re-runnable against the control-minute set for the same four sessions. · mode: do not build	54.2 [44.6, 63.5] and 56.7 [47.5, 66.2]; both contain 50. 5,255 control minutes.	Arms: Nothing arms. Exit: n/a. Kill: It is already killed. Δ Standing caveat: refuted as a separator of his presses from random minutes. It is not a claim that range position is meaningless in general.	n/a n/a. ★ SATURDAY #2 fixes the instrument that computes day range — and this conclusion was re-checked against the fix and does not move.
3	VWAP stretch as an entry condition — he is neither fading nor chasing VWAP REFUTED Stop building it. Same ruler, same answer as NOT-AN-ACTION #2. Mechanism: Signed distance from session VWAP in ATR units at the press, ranked against control minutes. Why: Within-session rank 54.0, CI [45.6, 62.4] — contains 50. The press median stretch is +5.21 ATR against a control baseline of +5.41 . There is no stretch preference to find. Evidence: \$2.2 and \$2.5 lead 5. · verification: Same control-minute set as NOT-AN-ACTION #2. · mode: do not build	rank 54.0 [45.6, 62.4]; press median +5.21 ATR vs baseline +5.41.	Arms: Nothing arms. Exit: n/a. Kill: Already killed.	n/a n/a.
4	★ Do NOT extend the \$250 daily loss limit to this book — it is listed as already automated and it is worse than nothing here REFUTED Leave it pre-registered and frozen on the automated single-lot book where its evidence came from. Do not point it at the manual book. Two separate defects in it are SATURDAY #3 and SATURDAY #4. Mechanism: A \$250 session loss limit that benches the book for the rest of the day. Why: Applied to this book it measures −\$67.09 a session over 29 sessions. It trips on 10 of them, and the trades opened after those trips net +\$1,945.50 . On 6 of the 10 trips he had already stopped anyway, so the entire measured cost sits in four sessions. It is one-directional by construction: the rule can only ever remove trades, and on this book the trades it removes made money. Evidence: \$3.4(c) (the one that is worse than nothing) and \$3.5 row 5. · verification: Re-run the limit against trades per session — and note SATURDAY #4 first: the guard's own P&L reader counts fabricated fills, so its live trips are not the same population as this replay. · mode: do not extend	−\$67.09 a session over 29 sessions; this week −\$31.40/day; the post-trip trades net + \$1,945.50.	Arms: Nothing arms — and the point of this row is that nothing should. Exit: n/a. Kill: Scoped to THIS book. It is not a verdict on the loss limit where it currently runs, and it does not say remove it from there.	n/a n/a.

STANDING CAVEATS ON EVERY DOLLAR ON THIS CARD

★ **The biggest caveat is not a number, it is a missing half.** There are **36 recorded presses and 0 recorded passes**. Every description in Sections 1 and 2 is fitted on one side of a decision boundary, and the only tests in this document that escape that are the ones run against control minutes — the leg, the range position, the VWAP stretch, the page band. Those are labelled where they occur. Everything else is a description of what you did, not a discovery about when to do it.

The fee is \$1.50 per round trip and the multipliers are not the same. MNQ is \$2.00 a point; **MGC is \$10.00.** Every dollar on this card is MNQ at \$2.00, because every trade in the report week was MNQ. The only gold in this document is the FOMC chart in Section 4, which carries no fills and is priced at \$10.00 in its own caption — pricing gold with the Nasdaq multiplier would understate it fivefold.

And every dollar here is a BOOKED dollar, not a repriced one. That is unusual for this report and it is worth saying: the money in Sections 1 and 4 is what the venue actually paid, straight out of trades, net of fees. The modelled numbers are confined to three rows — the eight-hour stopwatch, the loss-limit replay and the page-band sweep — and each of those carries its own uncertainty line in-row.

6 method wounds are open and they bound everything above. There are **36 recorded presses and ZERO recorded passes**, so every behavioural finding in this report is fitted on one side of a decision boundary. Nothing here can distinguish “this is what his entries look like” from “this is what the tape looks like” except the handful of tests that are explicitly run against control minutes, and those are labelled where they occur —

MONDAY #1. the week is 39 entries and five of them ARE the week.

The best five made +\$3,162.48 against a week total of +\$2,916.49; the other 34 net -\$245.99. Any per-entry statistic in Section 1 is a statistic about a distribution whose mass sits in five draws, and the three worst (-\$1,288.00, -\$1,272.00, -\$762.00) are -\$3,322.00 between them — **HOLD #1. one fill in the week is not real — exactly one, and it is an EXIT.** All 124 fill prices of the week (62 legs, entry and exit) were priced against the last MNQ print at the fill second: **123 sit within 0.03% of the tape** (median gap 0.50pt) and all 40 entry-group prices are within 0.023%. The single exception is the already-tagged

BADFILL:paper_engine_0.1pct_fabrication_3of4_lots_20260915 — and the fabrication is on its EXIT (21.69pt, 0.074% adverse), not its entry. It is excluded from every total in this report and it is NOT excluded by the live \$250 loss limit, whose own P&L reader has no data_quality filter and tripped on it. △ Rev 1 of Section 2 said every P&L figure in the week was contaminated; that claim is measured and withdrawn — **SATURDAY #4. three different denominators appear in this document and they are three different instruments, not an error.** Section 1 counts **39 entries**

over 5 sessions from the trade record; Section 2 counts **36 entries out of 78 presses over 4 sessions** from `operator_reads.jsonl`, which began recording mid-week; Section 3 counts **34 entries over 4 days** because `gazbot7-leg-watch` started at 05:55:43Z on 2026-09-15. ★ **And §2's 36 INCLUDES the BADFILL entry that §1's 39 excludes** — deliberately, because the tape at that moment was real even though the fill price was not; every DOLLAR cell in §2.4 is computed on 35 and says so. Each section names its own window; do not carry a rate across them — **BUILD #6. the 15-minute page band was chosen off this week's own curve**, on 34 presses and four days. Its $P = 0.0009$ is a real permutation result against a real null, it is **0.0036 once corrected for winning a four-band sweep** (which it survives), and it is still in-sample. It is now PRE-REGISTERED in writing — band, coverage target, kill line and the date it is re-scored — before it ships rather than after, which is the only thing that can turn next week's count into a test rather than a second fit — **MONDAY #2. the event-day split was chosen after seeing the losses**, on four event days, one of which was positive at +\$224.50 — and the FOMC damage was ONE entry at T+52, not the session. **Its headline “+\$166.60 a day” prices the whole-DAY unit and only the upside of it:** the same stand-aside deletes \$2,267.51 of winners across those four days, and the three units (§3.4b) differ from each other by a single trade. The row survives as a page and would not survive as a bench, which is exactly what it says — **MONDAY #3.**

What did not run this cycle, named exactly. 4 of 4 fragments are finished work and sit in this document in full. Nothing on the running order predates this cycle, so no section carries the red “this is last cycle's text” banner this week — every fragment above was written tonight.

Nothing is missing — every section on the running order rendered from a real fragment, so there are no red placeholder boxes to hunt for. ★ **No play on this card is sourced from any section carrying a stale banner**, and the loop on LAST week's card is closed immediately below this table.



LAST WEEK’S CARD — and why it is not graded row by row this time

The card is only worth writing if somebody checks it afterwards, so this section exists every week. **This week it has to report a gap in itself.** The previous card is reports/friday_v7/plays.json.pre-0918 — **58 plays** about the paper tournament, the shadow board, the router review and the greenfield labs. **This report is about a different subject.** It was refocused on 2026-09-18, by instruction, onto one question: your own trading and whether it can be automated. Not one phase ran this cycle that could check a tournament row, a shadow arm or a router setting.

So those 58 rows are UNGRADED, and that is a different thing from done or not done. Grading them from memory would be fabrication — this build has shipped a hand-copied status line before and it was wrong within a week. The card has not been deleted: it is on disk at the path above, every row still carries its own verification command, and if the tournament comes back so does the ledger. ★ **What must NOT happen is that they quietly disappear.** “We checked and it was done” and “nobody checked” are opposite instructions, and this paragraph is the second of those two.

One row from that card is gradeable, because this week’s own evidence speaks to it directly — and the grade is more interesting than a tick.

PLAY	WHAT IT ASKED FOR	STATUS	THE CHECK THAT DECIDED IT
SATURDAY 5	Move the rider's flatten earlier and put the page UPSTREAM of it — the row that bit on 2026-09-04, when four lots rode an unreachable gateway across the weekend	NO LONGER BITING — but not verified as FIXED	The desk is flat into this weekend: the last leg closed 18:33:23Z and no position was opened after it. desk_reconcile at 23:01:02Z reads venue +0 = tournament +0 + rider +0 → unaccounted +0 with ok true, and day_rider_state.json carries venue_net 0 / venue_ok true at 23:01:06Z. That is the outcome the row wanted, and it agrees with the WEEKEND CARRY CHECK at the top of this report — both are read from the same files at render time rather than copied forward. △ But it is an OUTCOME, not a check: the

PLAY	WHAT IT ASKED FOR	STATUS	THE CHECK THAT DECIDED IT
			final position was closed by hand at 18:33, more than two hours before either eod-flatten shot, so the flatten path was never exercised this week and nothing here says it works. Carried, ungraded.

THE HONEST LIMIT ON THIS TABLE

It contains **one row out of fifty-eight**. That is not a score, it is a sample of one chosen because the evidence happened to exist. Do not read the other fifty-seven as passed, failed or abandoned — read them as **not looked at this week**, which is what they are.

§ The evidence behind those plays

Everything from here down is the working, and it is four sections that build on each other in order. **Section 1** is the record: every entry you made this week, by day, by hold time, by hour, by event day and by how it ended, plus what the book looks like now. **Section 2** asks what the tape looked like at the moment you pressed — the leg, where price sat in the range, the VWAP stretch, the volatility — each one measured against control minutes from the same sessions, which is the only comparison that is legitimate when every recorded example is a “yes”. **Section 3** is the verdict: what can be automated today, what cannot, and exactly how far off the rest is. And **Section 4** draws it — all five sessions with every one of your fills marked on the day’s own price path, plus the FOMC hour on MNQ and on gold. Read Section 4 alongside Section 1: the shape that cost the most money this week is the hardest thing to see in a table and the easiest to see on a chart.

★ **What is deliberately NOT here.** No census section, no gate rehab, no shadow board, no router review, no greenfield lab. Those are the

tournament report and this is not it. Their fragments are on disk and a month old, and the build will not stitch them — see the red box at the top of this page. ★ **One precise exception, because “we did not print it” and “we may not cite it” are different rules:** the run census itself was re-frozen tonight against this week’s tape, and Section 4 cites it twice — once to rank the FOMC hour’s runs against the week’s other 66, and once to withdraw a Rev 1 caption that mis-read census_summary.json’s top lists as a ranking of the week when its own _note says they are the 51 runs the desk SAT OUT. A fresh, on-subject, named citation is fine; a month-old section stitched under this masthead is not. If you are looking for a gate’s P&L this week, the answer is in the front table and it is that the automated desk took **nothing** — and §3.4(d) now shows, with timestamps, that this was a nightly decision rather than a stuck switch.

Fragment freshness is stated once, above — see “what did not run this cycle, named exactly” on the front page. It is not repeated here.

The rule behind all of that: **a gap is never a null result.** “We looked and found nothing” and “nobody looked” are opposite instructions for next week, so nothing is silently omitted, nothing stale is silently stitched, and where a fresh section contradicts itself the build says so in a banner rather than quietly picking a side.

WHAT IS IN HERE

1. **Front** — [the week that was & the week coming](#), [the honest headline](#), [WHAT TO ACTUALLY DO](#), [last week’s card](#)
2. **Section 1** — [THE RECORD](#) — [every entry you made this week, and what the book now looks like](#)
3. **Section 2** — [THE CONDITIONS](#) — [what the tape looked like at the moment you pressed](#)
4. **Section 3** — [THE AUTOMATION GAP](#) — [what can be automated today, what cannot, and how far off](#)
5. **Section 4** — [THE RUN CHARTS](#) — [every fill of the week, marked on the day's own price path](#)

S1 THE RECORD — Monday 14 to Friday 18 September 2026

Garrath — **you made +\$2,916.49 this week** across **39 entries**, every one of them 4 lots, every one of them opened by your hand and stamped manual in the database. That is **+\$583.30 a day at 4 lots** over five sessions, and it is comfortably the best week this book has had — the previous best was +\$1,617.50 and the week before this one was −\$3,049.49. It also turned the whole book around: before Monday the rider was **−\$2,591.43 over 53 entries** for its entire life, and it now stands at **+\$325.06 over 92** — see §6 for what that 92 excludes and what the unfiltered book reads. Six weeks of work were in the red; one week put it in the black.

Two warnings before the tables. First, **five trades are the week** — the best five made +\$3,162.48, which is more than the whole week, and the other 34 entries together came to −\$245.99. Second, the money is grouped by **DECISION, not by database row**: the week produced 62 rows in trades but only 40 entries, because a scale-out books one row per lot leaving. Counting rows would have inflated the trade count by 1.55×

★★ IF YOU READ NOTHING ELSE — FIVE THINGS

1. The week, in one line. +\$2,916.49 on 39 entries. 25 winners worth +\$8,232.99, 14 losers worth −\$5,316.50. Median entry +\$160.00. **Three green days and two red** — Mon +\$839.49, Tue −\$114.00, Wed −\$833.00, Thu +\$1,477.50, Fri +\$1,546.50 — and the bad one was Wednesday, **FOMC day**. (Tuesday is red by \$114.00, which is one scratch trade wide; Wednesday is the day that cost money.)

2. Your losers are wrong from the first minute. This is the cleanest thing in the whole section and it is pure description, no model. Across the week your 25 winners had a median best-moment of **59.2 points** in your favour and you booked a median **70% of it**. Your 14 losers had a median best moment of **11.2 points** — 11 of 14 never once showed you 25 points, and 7

never showed you ten. You are not giving winners back. You are occasionally getting on the wrong side and staying there.

3. The morning carried it, but do not turn that into a rule. 05:00–12:59Z was **+\$2,445.51 over 20 entries** this week; after 13:00Z you were **–\$553.02 over 16**. But over the whole life of the confirmed-manual record the same morning window is **–\$779.48** — one trade on 10 September flips its sign. Point estimate first, uncertainty second: the morning is where you were this week, not a proven edge.

4. The eight-hour hole did not reopen. Over the book's life, **4 entries held past eight hours, 0 winners, –\$6,612.44** — and every other entry you have ever made adds to **+\$6,937.50**. This week **none of the 39 entries reached eight hours**. Longest was 7h 01m. That is the single biggest behavioural change in the record.

5. You are still the exit. 38 of the 39 entries ended on your `MANUAL_CLAIM`; exactly one ended any other way — the **–\$1,288.00** long on Tuesday that the 20:40Z clock flattened. The ladder took a first lot off on 14 of 39, but it did not finish a single one — all 14 still ended on your button.

1 • The week, day by day

Window: entries opened on or after 2026-09-14, gate `day_rider`, MNQ at \$2.00/point and \$1.50 a round trip. EXCLUDE: rows are dropped; the one BADFILL: entry is shown but kept out of every total — it is a real trade whose price came from the paper engine fabricating 3 of its 4 lots.

DAY	ENTRIES	WON	NET	BIGGEST WINNER	BIGGEST LOSER	WHAT HAPPENED
Mon 14 Sep	4	3	+\$839.49	13:18Z LONG + \$607.48	13:49Z LONG – \$162.00	Two morning entries, then the 13:18 runner — and a re-entry 31 min later that lasted 2m50s
Tue 15 Sep	6	4	–\$114.00	06:56Z SHORT + \$623.00	13:39Z LONG – \$1,288.00	Five entries made +\$1,174 by lunchtime; then one long

DAY	ENTRIES	WON	NET	BIGGEST WINNER	BIGGEST LOSER	WHAT HAPPENED
						held 13:39→the 20:40Z clock
Wed 16 Sep	6	3	−\$833.00	13:13Z LONG + \$819.00	18:52Z LONG − \$1,272.00	FOMC. Held long through the decision, bought again at T+52 with the slide still running
Thu 17 Sep	8	5	+\$1,477.50	11:55Z LONG + \$608.50	13:26Z LONG − \$441.00	Four of the eight took a first lot at the 100-pt rung; the overnight long ran 6h35m for + \$272
Fri 18 Sep	15	10	+\$1,546.50	13:55Z SHORT + \$434.50	16:11Z SHORT − \$208.00	Fifteen entries — the busiest day on record, and ten of them winners
WEEK	39	25	+\$2,916.49	+\$819.00	−\$1,288.00	+\$583.30/day at 4 lots

The week is five trades. Ranked by money, the top five are +\$819.00, +\$623.00, +\$608.50, +\$607.48, +\$504.50 — +\$3,162.48 together, against a week total of +\$2,916.49. The remaining 34 entries net −\$245.99: a scratch. The three worst are −\$1,288.00, −\$1,272.00, −\$762.00 — −\$3,322.00 — and two of those three are the two "held it" trades, Tuesday's clock-flat and Wednesday's FOMC re-entry.

2 • Every entry you made this week

One row per **decision**. "Best" is the furthest the trade ever went your way while you were in it, "Worst" the furthest against — both measured off the 5-second tape from the moment you got filled to the moment the last lot left. "Kept" is what fraction of the best moment you actually banked. Times are UTC; Paris is UTC+2.

OPENED (Z)	SIDE	LOTS	HELD	\$ NET	PTS	BEST	WORST	KEPT	HOW IT ENDED
Mon 14 09:09	SHORT	4	1h 23m	+\$145.01	+18.9	66.2	−27.8	29%	your hand, after 1 lot at the ladder

OPENED (Z)	SIDE	LOTS	HELD	\$ NET	PTS	BEST	WORST	KEPT	HOW IT ENDED
Mon 14 10:33	LONG	4	43m	+\$249.00	+31.9	36.4	-2.4	88%	your hand in one go
Mon 14 13:18	LONG	4	30m	+\$607.48	+76.7	113.1	-22.9	68%	your hand, after 1 lot at the ladder
Mon 14 13:49	LONG	4	3m	-\$162.00	-19.5	22.0	-21.0	—	your hand in one go
Tue 15 05:36	LONG	4	17m	+\$444.00	+56.2	61.5	-12.2	91%	your hand, after 1 lot at the ladder
Tue 15 06:56	SHORT	4	1h 14m	+\$623.00	+78.6	108.1	-41.9	73%	your hand, after 1 lot at the ladder
Tue 15 08:20	LONG	4	11m	(-\$304.50)	—	—	—	—	BADFILL — shown, not counted. The paper engine fabricated 3 of the 4 lots at exactly 0.1% adverse
Tue 15 08:38	SHORT	4	54m	+\$172.00	+22.2	50.8	0.0	44%	your hand in one go
Tue 15 10:29	LONG	4	1h 04m	+\$343.00	+43.6	66.5	-7.0	66%	your hand, after 1 lot at the ladder
Tue 15 12:24	SHORT	4	26m	-\$408.00	-50.2	5.2	-64.2	—	your hand in one go
Tue 15 13:39	LONG	4	7h 01m	-\$1,288.00	-160.2	15.8	-228.2	—	the 20:40Z clock in one go
Wed 16 06:12	LONG	4	1h 58m	-\$228.00	-27.8	30.5	-39.0	—	your hand in one go
Wed 16 08:11	SHORT	4	4h 13m	-\$762.00	-94.5	6.5	-107.0	—	your hand in one go
Wed 16 13:13	LONG	4	2h 35m	+\$819.00	+103.1	146.9	-32.6	70%	your hand, after 2 lots at the ladder
Wed 16 16:43	SHORT	4	41m	+\$160.00	+20.8	30.2	-22.5	69%	your hand in one go
Wed 16 18:52	LONG	4	16m	-\$1,272.00	-158.2	0.8	-170.8	—	your hand in one go
Wed 16 19:08	SHORT	4	31m	+\$450.00	+57.0	152.0	-11.8	38%	your hand, after 2 lots at the ladder
Thu 17 00:25	LONG	4	6h 35m	+\$272.00	+34.8	46.5	-64.2	75%	your hand in one go
	LONG	4	1h 37m	+\$193.00	+24.9	57.0	-24.8	44%	your hand, after 1 lot

OPENED (Z)	SIDE	LOTS	HELD	\$ NET	PTS	BEST	WORST	KEPT	HOW IT ENDED
Thu 17 07:07									at the ladder
Thu 17 11:55	LONG	4	21m	+\$608.50	+76.8	98.8	-8.5	78%	your hand, after 1 lot at the ladder
Thu 17 12:19	LONG	4	21m	+\$504.50	+63.8	73.1	-7.2	87%	your hand, after 1 lot at the ladder
Thu 17 13:26	LONG	4	1h 28m	-\$441.00	-54.4	21.5	-121.0	—	your hand in one go
Thu 17 15:27	LONG	4	1h 35m	+\$382.50	+48.6	64.0	-4.2	76%	your hand, after 1 lot at the ladder
Thu 17 19:35	SHORT	4	2m	-\$35.00	-3.6	1.0	-5.5	—	your hand in one go
Thu 17 19:38	LONG	4	12m	-\$7.00	-0.1	4.6	-15.1	—	your hand in one go
Fri 18 02:37	LONG	4	50m	+\$338.00	+43.0	45.2	-11.0	95%	your hand in one go
Fri 18 04:50	LONG	4	56m	+\$414.00	+52.5	61.0	-3.2	86%	your hand, after 1 lot at the ladder
Fri 18 06:06	LONG	4	59m	+\$409.50	+51.9	59.2	-3.0	88%	your hand, after 1 lot at the ladder
Fri 18 07:10	LONG	4	33m	+\$62.00	+8.5	33.5	-14.8	25%	your hand in one go
Fri 18 08:11	SHORT	4	57m	+\$176.00	+22.8	33.0	-18.5	69%	your hand in one go
Fri 18 09:31	SHORT	4	41m	+\$75.00	+10.1	28.0	-10.2	36%	your hand in one go
Fri 18 10:15	LONG	4	44m	+\$65.00	+8.9	41.8	-2.5	21%	your hand in one go
Fri 18 11:01	SHORT	4	1m	+\$2.00	+1.0	1.8	-6.5	57%	your hand in one go
Fri 18 11:21	SHORT	4	16m	-\$198.00	-24.0	6.8	-25.0	—	your hand in one go
Fri 18 12:16	SHORT	4	1h 07m	-\$30.00	-3.0	21.0	-14.2	—	your hand in one go
Fri 18 13:55	SHORT	4	20m	+\$434.50	+55.1	67.8	-10.8	81%	your hand, after 1 lot at the ladder
Fri 18 14:23	LONG	4	50m	-\$117.50	-13.9	28.6	-40.4	—	your hand in one go
Fri 18 16:02	SHORT	4	6m	-\$160.00	-19.2	3.2	-20.0	—	your hand in one go
Fri 18 16:11	SHORT	4	56m	-\$208.00	-25.2	53.0	-30.0	—	your hand in one go
Fri 18 17:07	LONG	4	1h 26m	+\$284.00	+36.2	48.8	-14.2	74%	your hand in one go

OPENED (Z)	SIDE	LOTS	HELD	\$ NET	PTS	BEST	WORST	KEPT	HOW IT ENDED
39 entries	23L / 16S	4 each	med 50m	\$2,916.49	+393.8	1811	—	22%	38 by hand, 1 by the clock

The bottom row is worth a second look. Added up, the tape offered **1811 points** across your 39 trades and you took **394** — 22%. That number is NOT a criticism and it is not a target: it includes the losers, where the "best moment" was a point or two of noise, and nobody exits at the high. On the 25 winners alone you kept a median **70%**, and 17 of 25 kept two-thirds or more. Read it as evidence you are not the problem on the way out.

Sides: **23 longs for +\$2,479.98** and **16 shorts for +\$436.51**. You were more often long and the longs paid better this week, but on 39 trades that gap is one Wednesday afternoon wide — do not build on it.

3 · By hold time — the standing finding, refreshed

HELD FOR	ENTRIES (LIFE)	WON	NET (LIFE)	\$/ENTRY	THIS WEEK	NET THIS WEEK
under 15 min	19	7/19	−\$1,553.50	−\$81.76	5	−\$362.00
15-60 min	34	26/34	+\$6,063.98	+\$178.35	21	+\$2,965.98
1-3 hours	19	14/19	+\$2,998.51	+\$157.82	10	+\$2,090.51
3-8 hours	16	11/16	−\$571.49	−\$35.72	3	−\$1,778.00
over 8 HOURS	4	0/4	−\$6,612.44	−\$1,653.11	—	—
UNDER 3 hours	72	47/72	+\$7,508.99	+\$104.29	36	+\$4,694.49
OVER 3 hours	20	11/20	−\$7,183.93	−\$359.20	3	−\$1,778.00

The standing finding survives the week and grew: **under three hours + \$7,508.99 over 72 entries, over three hours −\$7,183.93 over 20** (it was +\$5,262/44 against −\$5,406/17 when this was last written — the week added 36 short holds worth +\$4,694.49 and 3 long ones worth −\$1,778.00).

△ THE CONFOUND, STATED PLAINLY — AND WHERE IT STOPS

Hold time is partly an OUTCOME, not only a choice. A trade that goes your way hits a rung and closes; a trade that goes against you gets held while you wait for it to come back. So part of that gap is

simply "losers last longer", which is true of almost every discretionary book ever measured, and it does not mean a stopwatch would have made the money.

The over-8h bucket is not that shape. 4 entries, **0 winners**, – \$6,612.44, and their envelopes say why: the 21 Aug long never printed a single tick in profit (best moment –3.4 points) and the 10 Sep long showed +21.8 points before going 444 against. Those are not trades held in hope of a recovery that nearly came. They are positions nobody was watching. That is abandonment, and a stopwatch does address it.

4 · By hour of day — does the edge live in part of the window?

First, a correction to the shape of the question. The "05:00–11:00Z almost exclusively" line comes from the 36 captured presses, and those only start on 14 September. Across the 44 entries the database can prove were yours, the spread is much wider — this week alone you opened trades in **18 different hours of the clock, from 00:25Z to 19:38Z**. $\triangle$ A correction to the correction, from Rev 1: that sentence used to read “from 09:09Z to 17:07Z”, which is not a clock range at all — it is the week's first entry (Monday 09:09) and its last (Friday 17:07), which is what `min()` and `max()` return when you run them over a full timestamp instead of over the time of day. It made the paragraph argue for a nine-hour window while claiming eighteen hours. The figures above are `substr(opened_at, 12, 8)`: earliest 00:25:01Z, latest 19:38:32Z.

BLOCK (UTC)	PARIS	ENTRIES	WON	NET THIS WEEK	\$/ENTRY	WHAT SITS IN THERE
Overnight 00-04Z	02:00-06:59	3	3	+\$1,024.00	+\$341.33	Asia / pre-Europe
Early Europe 05-08Z	07:00-10:59	9	7	+\$1,089.50	+\$121.06	the London open block
Late Europe 09-12Z	11:00-14:59	11	8	+\$1,356.01	+\$123.27	the pre-US grind
US open 13-16Z	15:00-18:59	11	5	+\$26.98	+\$2.45	

BLOCK (UTC)	PARIS	ENTRIES	WON	NET THIS WEEK	\$/ENTRY	WHAT SITS IN THERE
						the cash open and the first three hours
US afternoon 17-23Z	19:00-01:59	5	2	−\$580.00	−\$116.00	into the 20:40Z flatten
05:00-12:59Z — ★ the European morning, and the window the refutation is scored on	07:00-14:59 Paris	20	15	+\$2,445.51	+\$122.28	where the week was made
Everything before 13:00Z — the same window plus the three overnight entries	before 15:00 Paris	23	18	+\$3,469.51	+\$150.85	the 00-04Z row added back in
13:00Z and after	15:00 Paris on	16	7	−\$553.02	−\$34.56	holds both big losers
WEEK — the five blocks, and the two windows that reconcile to it	—	39	25	+\$2,916.49	+\$74.78	23 before 13:00Z + 16 after = 39 ✓

⚠ **WHICH WINDOW IS THE REFUTATION SCORED ON? 05:00-12:59Z, AND THE TWO DIFFER BY \$1,024**

Rev 1 printed a row labelled “Before 13:00Z” that was really the 05:00-12:59Z filter, so the three overnight entries (00:25, 02:37, 04:50 — **all three winners, +\$1,024.00 between them**) were in neither summary row and the table did not add up to the week. Both rows are now printed and labelled, and the WEEK line reconciles. **NOT-AN-ACTION #1 — “only trade the European morning” — is scored on 05:00-12:59Z**, because that is the window the claim names and the only one with an out-of-sample leg. Scored on everything before 13:00Z it looks \$1,024 better this week and it is still refuted by the lifetime figure below, so nothing about the verdict turns on the choice; the labelling does.

Point estimate first. This week the morning block made +\$2,445.51 on 20 entries and the afternoon lost −\$553.02 on 16. **Uncertainty second, and it is large.** Take every entry the database can prove was yours, not just this week's, and the same morning block is **−\$779.48 over 25 entries** while the afternoon is −\$553.02 over 16. The sign of the morning flips entirely,

because one trade — 10 September, 05:40Z, held 13h34m, −\$2,656.00 — is bigger than the whole week's morning profit. **One trade decides this cut.** That is not an hour-of-day edge; that is a hold-time problem wearing an hour-of-day costume.

★ **And the two sides of that split are not equally tested — only one of them is tested at all.** “Every entry the database can prove was yours” adds **five** pre-week entries to the morning and **none** to the afternoon: all five provable pre-week presses (10 Sep 05:40, and 11 Sep at 05:57, 08:59, 09:12, 10:39) fall inside 05:00–12:59Z. So the afternoon's −\$553.02 over 16 is this week's sixteen entries restated, not a wider sample — it cannot flip, because there is nothing outside the week for it to flip against. That makes the refutation stronger, not weaker: **the morning claim is the only side that has ever been looked at outside this week, and looking is what killed it.**

5 · By event day — refreshed with this week's FOMC

DATE	EVENT	ENTRIES	WON	NET	THE TRADE THAT DID IT
Wed 12 Aug	CPI	1	1	+\$224.50	One entry, worked
Fri 04 Sep	NFP	1	0	−\$341.00	One long held 8h44m — the NFP day is also an abandonment day
Fri 11 Sep	CPI	4	1	−\$568.99	Four entries, one winner
Wed 16 Sep	FOMC	6	3	−\$833.00	Long held through the 18:00Z decision, then bought again at T+52 with the slide still running: −\$1,272.00 on that one trade against a day of −\$833.00
EVENT DAYS	4 days	12	5	−\$1,518.49	−\$126.54 an entry
ORDINARY DAYS	25 days	80	53	+\$1,843.55	+\$23.04 an entry

The event side reproduces the standing figure to the cent: **−\$1,518.49 over 12 entries.** The ordinary side has moved because the week added

Thursday and Friday — it now reads +\$1,843.55 over 80 entries rather than the +\$1,076/67 quoted in the build queue.

⚠ **FOUR DAYS. READ IT AS A WARNING LIGHT, NEVER A SWITCH**

This is **4 days of trading** and the split was chosen after seeing the losses. FOMC Wednesday was the only red day of the week — but it was red because of one entry at T+52 into a slide that ran for another 33 minutes, and you had already said yourself: "if i knew the fed was making a decision i probably wouldve held off." The calendar and the T-60/T-15 pages are already built and already running. That is the right shape: it tells you an event is coming and refuses to tell you which way. Nothing here justifies benching anything automatically.

6 · How they ended — your hand, a ladder rung, or the clock

Your ladder is set to **\$100 / \$200 / \$400 / \$600 a lot** and takes exactly one lot at each rung. Here is what actually fires.

ENDING	ENTRIES THIS WEEK	NET	ENTRIES (LIFE)	NET (LIFE)	READ
Your hand — MANUAL_CLAIM	38	+\$4,204.49	81	+\$1,103.56	you decided it was over
A ladder rung took the last lot	—	—	7	+\$341.50	the machine finished it
The 20:40Z clock — CLOCK_FLAT*	1	-\$1,288.00	3	-\$1,024.50	nobody decided; the flatten did
Cross-desk flatten	—	—	1	-\$95.50	the other book's guard
LIFETIME, on the starred basis	39	+\$2,916.49	92	+\$325.06	81 + 7 + 3 + 1 = 92 ✓
CLOCK_FLAT_BUG — the row the asterisk excludes	—	—	1	-\$69.00	2026-08-21 09:11, flat 60 seconds later. A software flatten, not a decision
LIFETIME, unfiltered (gate LIKE 'day_rider%', data_quality IS NULL)	39	+\$2,916.49	93	+\$256.06	the same query with nothing taken out

★ **WHAT THE ASTERISK ON CLOCK_FLAT* MEANS — IT WAS UNDEFINED IN REV 1 AND IT SETS FOUR OTHER NUMBERS**

CLOCK_FLAT* is CLOCK_FLAT plus CLOCK_FLAT_RECON, and it **excludes the 2026-08-21 09:11 entry, −\$69.00, whose exit reason is CLOCK_FLAT_BUG** — a position opened and flattened by software sixty seconds later. It is excluded because it is not an ending anybody chose, and this table is about who decided. **The exclusion is defensible and it was undisclosed, which is worse than the exclusion.** It is what makes the lifetime book read **92 entries / +\$325.06** rather than **93 / +\$256.06**, and that 92 and that +\$325.06 are quoted on the front page, in this section's opening paragraph, in §7 and in §3 of Section 3. Both figures are now printed above so the two can be told apart. Everything else on the row is unchanged: with the bug row put back, *CLOCK_FLAT** reads 4 entries / −\$1,093.50 instead of 3 / −\$1,024.50. △ Note the same 09:11 decision also has a second leg tagged *BADFILL:stale_price_30pt_outside_book_20260821* at −\$81.50, which the *data_quality IS NULL* filter has always dropped — so no reading of that morning is a clean one.

You are the exit, and it is not close. 81 of the 92 entries this book has ever made ended on your claim button. The ladder has taken at least one lot on 31 of them — a \$100 or \$200 target, or an older trail — but the **\$400 and \$600 rungs have never once fired in the entire record** — not a single *TARGET_400* or *TARGET_600* row exists. You always get there first.

△ **ONE TEMPTING NUMBER IN HERE IS A TAUTOLOGY. DO NOT ACT ON IT.**

The 14 entries that touched a rung this week made +\$6,377.99; the 25 you exited in one go lost −\$3,461.50. That looks like "scaling out works". **It is not a finding.** A rung can only fire on a trade that had already gone \$100 a lot your way — so that split is very close to "winners won". Their median best-moment was 67.1 points against 28.0. There is no lever in it.

The one exit this week that was not yours cost $-\$1,288.00$: Tuesday's 13:39Z long, held 7h01m, flattened by the 20:40Z clock. Its envelope shows a best moment of **15.8 points** and a worst of **228.2** — it was never really a trade that worked and then failed. Whether the clock made it better or worse is a question about repricing an exit, and this desk does not reprice exits (119 calibrations, none reproduced your claiming). What can be said without a model: it spent seven hours going one way.

7 · Whose trades are these — and why the week is the honest number

PROVENANCE	ENTRIES	NET	DATES	WHAT THE DATABASE ACTUALLY SAYS
Confirmed yours — entry_source='manual'	44	$-\$308.50$	2026-09-10 → 2026-09-18	You pressed BUY or SELL. 39 of these 44 are this week
UNKNOWN — entry_source IS NULL	48	$+\$633.56$	2026-08-06 → 2026-09-10	Predate the column. Nothing recorded who opened them, and nothing ever will

This matters more than it looks. The 48 unknown-source entries contribute $+\$633.56$ and the 44 that carry your stamp contribute $-\$308.50$ — so **the book's $+\$325.06$ lifetime net is not attributable to anybody in particular** ($44 + 48 =$ the 92-entry basis defined in §6; $-\$308.50 + \$633.56 = +\$325.06$). Everything the database can prove is your hand is those 44 entries, and 39 of them happened this week. Strip the week out and the 5 provable pre-week entries are $-\$3,224.99$.

One shape in the unknown rows is worth naming without drawing a conclusion from it: for **12 straight sessions from 10 to 26 August the book took exactly one counted entry a day, and every one of them opened in the 13:00Z hour** — 20 of the 48 unknown entries sit in that single hour. The day rider's own documented automatic behaviour is one entry per session, no entry after 15:00Z. I am not claiming those rows were the machine's; the record does not say, and the standing rule is that a NULL stays a NULL. But it is why this section leads with the week, where every

single entry carries your stamp, and treats the lifetime tables as context rather than as your track record.

8 · Ranked shortlist — what this section says to do next

#	WHAT	MEASURED ON	THE NUMBER	VERDICT
1	Record the passes. The button shipped Friday; nothing has been written to it yet. Until it has, every table in this section is a description of trades you took and nothing here can become a rule.	36 presses, 40 claims, 0 passes	One side of a boundary. The fitting threshold is 30 labelled examples and it was written for positives AND negatives, so the real count against 30 is zero.	BLOCKER
2	A hard flat-out at eight hours. Needs no model, no judgement and no fitting — it is a stopwatch.	4 entries, 29 sessions	0 winners from 4, – \$6,612.44. Every other entry in the record nets +\$6,937.50. That is –\$228.02 of drag per session traded. It did NOT recur this week (0 of 39) — so on this week's data it would have earned nothing. What it recovers is not computed here; that needs the exit repriced (\$3).	PROMISING
3	Keep standing aside on event days. Calendar and T-60/T-15 pages already built and running; this is about whether to keep trusting them.	4 event days, 12 entries	–\$1,518.49 on event days against + \$1,843.55 over 80 ordinary entries. FOMC Wednesday was the week's only red day. Four days, and the split was chosen after seeing it.	PROMISING
4	"Only trade the European morning." Tempting after this week. Do not do it.	20 entries this week, 25 provable ones over the book's life	+\$2,445.51 this week; –\$779.48 over the confirmed-manual record, where one trade on 10 Sep flips the sign. A filter fitted on a week that a single trade can reverse.	REFUTED as a rule
5	Watch the "it has never gone my way" shape. Your losers do not start well and then fail — they are wrong immediately.	14 losers this week	Median best-moment 11.2 points against 59.2 for winners; 11 of 14 never showed 25 points and 7 never showed ten. ▲ Turning this into a bail rule means backtesting an exit, and the standing rule forbids a 120th	WATCH

#	WHAT	MEASURED ON	THE NUMBER	VERDICT
			calibration. Log it, watch it, do not fit it.	
6	Chase one execution defect. A TARGET_100 rung on 11 Sep booked -\$160.50 — a \$100-a-lot rung (that is +50 points on MNQ) that booked a fill 79.5 points BELOW the entry (trade id 868, long at 29320.25, out at 29240.75).	1 row, outside this week	Every other TARGET_100 in the record booked between +\$97.50 and +\$132.25. This one landed 129.5 points from the level it names. Small money, but a target exit is reporting a price its own trigger cannot have seen.	DEFECT

So what do you actually do on Monday? Nothing changes about how you trade — this section found no rule that survives its own uncertainty, and it would be dishonest to hand you one off 39 entries. Three things it does say: press PASS when you look and walk away, because that is the only thing standing between a description and a model; the eight-hour hole is the one number in the record big enough to be worth a stopwatch; and the morning result, lovely as it is, is one week.

One housekeeping note for §2, and it has to be said twice because the answer depends on the key. data/

operator_reads.jsonl holds **135 lines**. Every press is written twice (PathModified fires twice per write, ~1.4s apart), so the real figure is either **78 presses counted by their own timestamps** or **77 records under the (kind, request.raw) key** — the key cannot separate the two unlabelled 09-14 buys at 14:51:43 and 14:53:06, which are 83 seconds apart and are two real presses. **78 and 77 are both right; they are answers to different questions**, and Rev 1 printed one in this note and the other in §2.1 without saying so. Either way the **entry** count is 36 and counting raw lines would report 65 buy rows where there are 38 buy presses (37 under the key). See BUILD #6: the de-dup helper should key on (kind, ts) so the count stops depending on which sentence you read.

2 WHAT THE TAPE LOOKED LIKE WHEN HE PRESSED

Garrath — **the sample is now 78 recorded presses, and 36 of them are entries.** That is the headline number of this section, because the fitting threshold this desk wrote down for itself is 30 and the positives have crossed it. **The negatives have not: there are ZERO.** The PASS button went in on Wednesday and has never been pressed, so every row in the file is a "yes". That is not a technicality — it is the whole ceiling on this section. With one side of the boundary you can describe what your entries look like; you cannot learn what separates them from the moments you looked and walked away. So this is a description, and I have not tuned a single threshold on it.

One thing separates your presses from ordinary minutes, and it is the leg. At 33 of 36 entries a leg was already running — the desk's own definition, $1 \times \text{ATR}$ travelled in 8 minutes, dead on a $1 \times \text{ATR}$ give-back — against **74.6%** of all minutes on the same four sessions (session-blocked permutation, $P = 0.008$). And where a leg was running you took **its side 30 times out of 33**, against a coin-flip null of 50%: 18 of 19 buys into an UP leg, 12 of 14 sells into a DOWN leg. **You are a continuation trader and the tape says so without being asked.** You also do not wait for the leg to mature: median leg age at your press is **14 minutes**, and the median age of a randomly chosen in-leg minute is 13. You press almost as soon as the thing exists.

And the honest half, which is bigger. The number everyone expects to matter — where price sits in the day's range — **does not separate your presses from noise at all.** Ranked inside its own session, the median press sits at the **54th percentile** of all that session's minutes, 95% CI **[44.6, 63.5]**. That interval contains 50. Re-measured with a clean 22:00Z session ruler instead of the desk's own, it is 56.7 [47.5, 66.2] — the same answer. VWAP stretch is the same story: 54.0 [45.6, 62.4]. **"He buys high in the range" is true and meaningless, because so does every other**

minute of these four days. That comparison is the one thing this dataset can do without negatives, and it kills two filters before anyone builds them.

FIVE THINGS TO READ BEFORE ANY NUMBER BELOW

(1) 135 rows in the file are 78 presses — or 77 records, and the difference is one row of the ladder below. PathModified fires twice per write, so nearly every press is recorded twice ~1.4s apart with an identical request.raw. Deduping on (kind, request.raw) gives 77; deduping on the press's own timestamp gives 78, and 78 is the number used throughout this section. They differ by exactly one thing: the two unlabelled 09-14 buys, 83 seconds apart, which carry no request.raw for the key to separate them by. Counting raw rows would have inflated the sample by 73%.

(2) The ruler is tape_reader.context(), re-run at every minute. I did not invent a second definition of ATR or position-in-range for the baseline — I re-executed the capture's own code path at all 5,255 minutes of the four sessions, so press and baseline are measured identically. Validated press-by-press against the stored facts first: position-in-range agrees to a median 0.003 (p90 0.011), ATR to 0.45pt. The residual is capture write-lag, not a different calculation.

(3) Its "day so far" is not a day. Before 13:30Z the session window is the last 24 hours; after 13:30Z it is every minute with a time-of-day past 13:30 inside the last 24 hours — which splices yesterday afternoon onto today's open. The 2026-09-14 14:51 read says "session bars 203": that is Sunday 22:00-24:00Z plus Monday 13:30-14:51Z, with an eleven-hour hole in the middle, and the "day open" it showed him was Sunday's globex print. 24 of the 36 entry presses were taken before 13:30Z. Every conclusion below is re-checked against a clean 22:00Z-anchored session and none of them moves — but the number on his screen is not the day's range.

(4) The desk's own direction detector was BLIND at two thirds of these presses. drift threw no 5s rows since 13:30Z today at 51 of the 78 presses (24 of the 36 entries), because it is hard-anchored to the RTH open and he trades the European morning. At the 12 entries where it did read, it returned confirmed=false **all 12 times**. Whatever he is reading,

the automatic path could not have seen it — and for most of the week could not even look.

(5) The fills WERE checked against the tape, and the fabrication is one leg, on an EXIT. Rev 1 of this section said “the paper engine fills lots 2–4 at exactly 0.1% adverse” and concluded that every P&L figure here was contaminated. **That is not what the tape says, and the claim is withdrawn.** I priced all 124 fills of the week (62 legs × entry and exit) against the last MNQ print at the fill second in capture.db. The table below is the whole answer.

CHECKED AGAINST THE LAST PRINT AT THE FILL SECOND	RESULT	READ
Fill prices checked (62 legs, entry + exit)	124	every leg of the week, nothing sampled
Gap from the tape — median / p90	0.0017% / 0.0089% 0.50pt / 2.62pt	ordinary multi-lot slippage on a 29,000-point instrument
Within 0.03% of the tape	123 / 124	—
All 40 entry-group prices	max 6.75pt = 0.023%	the worst is 09-16 19:08 SHORT. No entry is anywhere near 0.1%
Reaching the 0.1% class	1 / 124	trade id 895, and it is the EXIT: out at 29,269.3125 against a last print of 29,291.00 — 21.69pt = 0.074% adverse. Its ENTRY is clean (1.88pt = 0.006%)

The mechanism in the tag is right and the side in it is wrong. Three of four lots leaving at 0.1% adverse off a reference of 29,291.3 blends to **29,269.28**; the booked exit is **29,269.3125**, a 0.28-point miss. I swept all 124 fills for that signature and it matches **exactly one** — id 895, already tagged BADFILL:paper_engine_0.1pct_fabrication_3of4_lots_20260915. **So: the week's dollars are venue dollars, with one flagged exception that every total in this report already excludes.** That is the same thing the front page says, and Rev 1 had the two sections contradicting each other on the single question of whether the money is real. Dollars are still secondary here — this section is about the tape, not the P&L — but they are no longer void.

2.1 · THE SAMPLE — what is actually on record

LADDER	COUNT	WHAT IT MEANS
Rows in data/operator_reads.jsonl	135	raw, double-fired
Distinct presses — by their own TIMESTAMP	78	

LADDER	COUNT	WHAT IT MEANS
		the number used everywhere in this section. Duplicates are the double PathModified write, ~1.4s apart
Distinct records — under the (kind, request.raw) key	77	△ the key collapses two real presses into one — see the row below. Both counts are right; they answer different questions
— ENTRY presses with a recorded side	36	21 BUY / 15 SELL. This is the fitting count, and it is over 30. Identical under either key
— CLAIM presses	40	3 of them name a lot count; the rest claim the book. Identical under either key
— buy presses whose label the rider ate	2 by timestamp 1 by the key	this single row is the whole 78-vs-77 gap. 2026-09-14T14:51:43 and 14:53:06 — 83 seconds apart, so two real presses — but both carry an empty request, so (kind, request.raw) cannot tell them apart. Predates the ExecStartPre snapshot fix; side UNKNOWN and it stays unknown
— PASS presses (the negatives)	0	the button exists since 2026-09-18 and has never been pressed. This is the binding constraint on everything below
— of which the BADFILL entry §1 excludes	1	09-15 08:20 BUY, trade id 895. Its TAPE facts are real, so it stays in every tape comparison below; its DOLLARS are not, so §2.4 now excludes it and reads 35. See the box under §2.4
Sessions covered	4	09-15 (7 entry presses — §1's clean day table says 6; the seventh is the BADFILL) · 09-16 (6) · 09-17 (8) · 09-18 (15)
Entry presses matched to a real entry	36 / 36	every one found an entry in trades inside +180s. No press was refused
Entries at 4 lots	36 / 36	there is no single-lot-vs-multi-lot comparison of DECISIONS to make
Baseline minutes for the comparison	5,255	every minute of the same four sessions, same ruler

One correction to carry into §1 and §3. The standing figure is that he trades 05:00–11:00Z almost exclusively — 26 of 28 presses. On this week's 36 that is **16 of 36**, spread across 18 different hours from 00Z to 19Z. The old window was a property of a smaller, older sample, not of his trading. Any rule built on it should be re-derived before it ships.

2.2 · THE ONE COMPARISON THAT IS LEGITIMATE WITHOUT NEGATIVES

Press population against ALL MINUTES of the same four sessions. The right-hand column is the one that matters: each press is ranked inside its own session's minute distribution, so a day on which he pressed fifteen times cannot drag the answer. If he pressed at random moments those ranks average 50.

AT THE MOMENT HE PRESSES	PRESS MEDIAN (P25-P75)	ALL 5,255 MINUTES (P25-P75)	WITHIN-SESSION RANK, 95% CI	READ
Position in the day's range (the desk's own ruler)	0.708 (0.48-0.92)	0.689 (0.45-0.85)	54.2 [44.6, 63.5]	NO SEPARATION. The interval contains 50
Position in the day's range (clean 22:00Z session, running extremes)	0.672 (0.34-0.92)	0.727 (0.35-0.90)	56.7 [47.5, 66.2]	same answer on a better ruler — so it is not a ruler artefact
VWAP stretch, in ATR	+5.21 (-0.73-+21.1)	+5.41 (-1.57-+14.4)	54.0 [45.6, 62.4]	NO SEPARATION. He is not fading stretch and he is not chasing it
ATR, points	8.81 (6.77-11.67)	7.93 (6.45-10.18)	61.5 [52.2, 70.4]	modest, clears 50
ATR percentile vs its own last 24h	71.7 (40.5-85.2)	49.3 (22.1-76.4)	62.9 [54.0, 71.4]	he presses into livelier- than-usual tape. 14 of 36 presses sit above the 75th ATR percentile, where only 25.9% of minutes do

Two filters died here and that is worth more than a finding. Position in range and VWAP stretch are the two conditions every hand-built entry gate on this desk has reached for. On 36 of his own presses, measured against the only honest control available, neither of them knows where he is. Nobody should spend another cycle on them. What survives is a weak volatility preference: he shows up when the tape is moving faster than it has been.

2.3 · THE LEG — the only thing that clearly separates a press from a minute

Leg definition is `scripts/leg_survival.py`'s, unchanged: a leg begins when the last 8 minutes have travelled $1 \times \text{ATR}$ in one direction, and dies when price gives back $1 \times \text{ATR}$ from its extreme. It is evaluated causally — at every minute I know only what was knowable then — and a press only counts as in-leg if the leg had already been detected, 8 minutes after its own start.

LEG STATE AT THE PRESS	PRESSES	ALL MINUTES, SAME SESSIONS	READ
A leg was already running	33 / 36 = 91.7% [78.2, 97.1]	74.6% (press-weighted)	session-blocked permutation, 20,000 draws: P = 0.008
— and he took ITS side	30 / 33 = 90.9% [76.4, 96.9]	50% (coin flip)	18 of 19 buys into an UP leg; 12 of 14 sells into a DOWN leg
Leg age at the press	14 min (11-17)	13 min (10-19)	he does not wait. Detection needs 8 minutes, so he is pressing ~6 minutes after the leg becomes knowable
Travel banked when he presses	28.5pt = 2.9 ATR	22.0pt = 2.5 ATR	

LEG STATE AT THE PRESS	PRESSES	ALL MINUTES, SAME SESSIONS	READ
			slightly further along than a random in-leg minute, but inside the same spread
Distance to the extreme he is trading toward		10.6 ATR	— he is not a breakout buyer. Only 5 of 36 presses are within 1 ATR of the running high (long) or low (short)
Same test with the scope's 3×ATR give-back		91.7%	97.6% REFUTED as an instrument. A 3×ATR leg is alive in 97.6% of all minutes, so "a leg was running" carries no information at all. Use the 1× definition

It survives leave-one-session-out. Drop any one of the four days and the leg-alive rate is 89.7–95.2% and the with-the-leg rate is 88.5–92.6%. It is not one trending Thursday carrying the result. What it is not, however, is an independent sample of 36: the presses cluster (7 of 35 consecutive entry presses came within 30 minutes of the one before, usually a claim followed by a re-entry within seconds), and four sessions is four sessions. Treat the interval as optimistic and the point estimate as the thing worth keeping.

The claims tell the same story more weakly: a leg was running at 35 of 40 claim presses (87.5%) against the same 74.6% baseline. That is a statement about when he claims, not a rule that reproduces what he claims — and per the standing limit, no attempt to reproduce the exits was made or should be.

2.4 • WHAT THE DOLLARS DID — post-hoc, three-deep, and not a rule

★ **The denominator changes here, and only here.** The 36 entry presses above include the 09-15 08:20 BADFILL, whose tape facts are real; its price is not, so **every dollar cell below is computed on 35**. On that basis the entries made +\$2,077.00 across four sessions at 4 lots = +\$519.25/session. (Rev 1 printed +\$1,772.50 / +\$443 a session, which is this figure plus the fabricated −\$304.50 — \$1 excluded that row from every total and \$2 silently put it back.) Two cuts fall out of the tape state above. Both are post-hoc, both have a cell of n=3 or n=12, and I am showing them so they can be tested later, not acted on now.

CUT (POST-HOC) — 35 ENTRIES, BADFILL EXCLUDED	N	TOTAL, 4 LOTS	MEDIAN	WINNERS	WHAT IT IS WORTH
Pressed WITH the running leg	29	+\$3,693.00	+\$176.00	20 / 29	the population he actually trades. The BADFILL sat in this cell; with it, 30 and +\$3,388.50
Pressed AGAINST the running leg	3	-\$1,515.00	-\$208.00	0 / 3	an anecdote, not a finding. Three trades. But all three lost, and the worst (-\$1,272, 09-16 18:52, buying into a DOWN leg at ATR 37.7) is one of the week's two largest losses. Unaffected by the exclusion
No leg running	3	-\$101.00	+\$2.00	2 / 3	nothing to say at n=3. Unaffected
ATR percentile terciles: low / mid / high (re-cut on 35)	12 / 11 / 12	+\$2,762 / -\$233 / -\$453	+\$311 / +\$2 / +\$174	9 / 6 / 7	this still REVERSES the description, and it is weaker than Rev 1 said. He prefers lively tape (\$2.2) and this week the money came from the quiet third. The BADFILL sat in the HIGH tercile: with it the cells read +\$2,762 / +\$272 / -\$1,262 — so a fabricated price was carrying a third of the “lively tape loses money” story. Post-hoc terciles, 11-12 trades a cell, medians that barely move. Do not act on it — re-measure it
ALL 35	35	+\$2,077.00	—	22 / 35	+\$519.25 a session at 4 lots

★ THE THREE DENOMINATORS, RECONCILED IN ONE PLACE — THEY ARE THREE DIFFERENT INSTRUMENTS

§1 counts 39 entries over 5 sessions from the trade record, and drops the BADFILL from every total. **§2 counts 36 entry presses over 4 sessions** from `operator_reads.jsonl`, which only began recording on 15 September — and **§2's 36 INCLUDES the BADFILL entry that §1's 39 excludes**, because what §2 measures at that moment (the leg, the range position, the ATR) is real tape and was really there. **§3 counts 34 entries over 4 days**,

because gazbot7-leg-watch started at 05:55:43Z on 15 September. The one place the inclusion would have mattered is the money, and the money in \$2.4 is now computed on **35**. Do not carry a rate across the three.

2.4b · FIVE TRADES ARE THE WEEK — so do those five look different at the press?

This report carries the “five trades are the week” caveat in six places and then never asks the obvious follow-up: were those five a different kind of moment, or the same moment that happened to keep going? The instrument to ask with is already built, so here is the answer, with **n=5** stamped on every cell. Top five of \$2's 35 clean presses (+\$3,005.00) against the other thirty (−\$928.00). **AUC** is the chance that a randomly chosen one of the five is higher on that feature than a randomly chosen one of the thirty: **0.50 is a coin flip**.

AT THE MOMENT HE PRESSED	TOP 5 +\$3,005.00	OTHER 30 −\$928.00	AUC	READ
Leg age at the press	14 min	14 min	0.46	identical to the minute
ATR percentile	57.3	71.7	0.49	if anything the big winners came in quieter tape — and the interval is a coin flip
Position in the day's range	0.798	0.708	0.53	no separation, same as \$2.2 against the whole tape
VWAP stretch, in ATR	+10.88	+5.21	0.51	no separation
Travel already banked by the leg	2.50 ATR	2.99 ATR	—	the winners were LESS far along, not more
A leg was running, and he took its side	5 / 5	24 / 30	—	the only cell that looks like anything — and on the 83% base rate across all 35, five out of five happens 39% of the time by chance. It is not a signal
Hold time (an OUTCOME, not a condition)	30.5 min	52.2 min	0.49	the winners were SHORTER, because a winner hits a rung and closes. This is the confound \$1 \$3 names, not a finding

The verdict is a null, and printing it is the point. Every feature this desk can measure at the press is **indistinguishable between the five trades that made the week and the thirty that did not** — every AUC between 0.46 and 0.53. **This cannot support a rule at this n and nothing here should be built on.**

Re-measure at 70 entries. What it does buy is a closed question: the big winners were not a different kind of moment that a filter could have found, so the honest reading of “five trades are the week” is size dispersion in one population, not two populations. $\triangle$ And note what is not being claimed: five draws cannot detect anything short of an enormous effect, so this is “we looked and found nothing”, never “there is nothing”.

2.5 • RANKED SHORTLIST

LEAD	VERDICT	MEASURED	NEXT STEP
1 • PRESS THE PASS BUTTON. The negative examples	HIGHEST VALUE ON THE DESK	36 positives, 0 negatives. The 30-press threshold is crossed on one side of the boundary and cannot be crossed on the other by waiting	every time you look and do NOT trade, press it. ~30 passes over two weeks makes the decision boundary learnable for the first time in eleven months. Nothing else in this section is unblocked without it
2 • The continuation description — a live 1×ATR leg, taken on its own side	PROMISING	33/36 in a leg (vs 74.6% of minutes, $P=0.008$); 30/33 with it (vs 50%); holds leave-one-session-out	as an ALERT, not a filter. §3 scores leg_watch's pages against these 36 presses — did the machine see the leg before he did? That is measurable now, without negatives
3 • drift is blind before 13:30Z	FIX — DEFECT	51 of 78 presses, 24 of 36 entries, returned no 5s rows since 13:30Z today. The other 12 all read confirmed=false	anchor the detector's day to the CME session (22:00Z) or fall back to the trailing 24h, as tape_reader already does for its bars. Until then the auto path cannot even look at two thirds of his trading
4 • Position-in-range as an entry condition	REFUTED	within-session rank 54.2 [44.6, 63.5]; 56.7 [47.5, 66.2] on a clean session ruler. Both contain 50	stop building it. Two independent rulers, 5,255 control minutes, one answer
5 • VWAP stretch as an entry condition	REFUTED	rank 54.0 [45.6, 62.4]; press median +5.21 ATR against a baseline of +5.41	same. He is neither fading nor chasing VWAP
6 • The volatility preference	OPEN — CONFLICTED	he presses at ATR percentile 71.7 vs 49.3 (rank 62.9 [54.0, 71.4]) — but on the 35 clean entries the quiet tercile made +\$2,762 and the lively one −\$453 (it read −\$1,262 in Rev 1, with a fabricated fill inside it)	the description and the dollars still point opposite ways, on 11-12-trade cells, and less far apart than Rev 1 said. Re-measure at 70 entries before anyone writes an ATR band
7 • The 3×ATR leg definition	REFUTED	alive in 97.6% of ALL minutes; a test that everything passes	the scope's own text should move to 1×. leg_survival.py already defaults to it
8 • tape_reader's session window	FIX — INSTRUMENT	"day open/high/low" splices yesterday afternoon onto today's RTH; 09-14's 203 "session bars" were Sunday night plus Monday morning across an 11-hour hole	no conclusion here changes when it is fixed (checked), so this is about what he is SHOWN. Anchor to 22:00Z
9 • The 05:00-11:00Z window	NEEDS REFRESH	standing figure 26 of 28 presses; this week 16 of 36, across 18 distinct hours	re-derive before any hour-of-day rule ships. It was a property of the old sample

LEAD	VERDICT	MEASURED	NEXT STEP
10 · The 78 captured narratives	UNSPENT	a ~3,160-character read of the whole session is stored with every press and nothing has been done with it	once passes exist, the narrative pairs are the actual experiment: the same text, one ending in a press and one in a pass

The one-line summary, in your units. 36 entry presses over 4 sessions; on the 35 whose price is venue truth, **+\$519.25 a session at 4 lots**. Where in the range you press and how far from VWAP you press are indistinguishable from picking a minute at random. What you actually do, 90% of the time, is take the side of a leg that is about 14 minutes old — and the desk's own leg detector can see that leg roughly six minutes before you press it.

3 CAN THIS BE AUTOMATED?

Garrath — **the honest answer is that one half of it already is, and it is switched off by a single number.** leg_watch saw a leg running at **33 of your 34** entries this week and took your side at 26 of them. It paged you ahead of **3 of them within half an hour, and 8 within the hour** — the window matters and both figures appear below, so they are named here rather than left to collide. That is not a detection failure and it is not a modelling problem: the watcher's first page fires at **60 minutes** of leg age, the median leg dies at **22**, and the median leg you press into is **28 minutes old**. The machine is watching the right thing and speaking forty minutes after you have already acted.

And that is measurable without a single negative example. Scored against a session-blocked permutation — your own presses shuffled to random minutes of the same days — **the alerting policy you are running today does not separate from chance:** the 25 pages it actually sent preceded **24%** of your entries where shuffling gives **18%**, $P = 0.24$. Move the first page from 60 minutes of leg age to **15** and the same detector precedes **82%** against a chance rate of 56%, **$P = 0.0009$** . The price is your phone going from **7 to 32 pages a day**. That is the whole trade, it is a parameter, and nothing about your judgement has to be learned to make it.

The other half has not moved, and one thing on the "already running" list is worse than nothing. There are still **zero** recorded passes, so there is still no decision boundary to learn. Your exits remain unreproduced after 119 calibrations and I did not run a 120th. And the **\$250 daily loss limit** — listed as a piece that is already automated — turns out to be measured at **−\$67.09 a session** when applied to your book, cannot reach your book anyway, fired twice this week benching nothing, and paged you both times claiming it had benched six gates.

FOUR THINGS TO READ BEFORE ANY NUMBER BELOW

(1) The window is four days, not five. gazbot7-leg-watch started on 2026-09-15 and its log begins **05:55:43Z** that morning. **34 of your 39 countable entries** fall inside it; the four on Monday and the 05:36 Tuesday long do not, and are excluded from every alert score here rather than counted as misses. That is deployment, not truncation.

(2) One "alert" in the log is not an alert, and dropping it changes the headline. data/event_alert.log carries 2026-09-18T07:04:03Z PAGED T-15 FOMC. The timer is OnCalendar=*:0/5:30; journald shows the service ran at 07:05:30 and logged "ok — nothing in a band"; the only FOMC in the calendar is **2026-09-16**; and a T-15 band cannot open two days late. It is a hand-run test that wrote into the live log, and scripts/score_alerts.py credited it with leading your 07:10 Friday entry. It is dropped here, so this section reads **4 led, not 5**. The event-alert channel sent you **nothing real all week**.

(3) "He ignored it" and "he never saw it" are different, and the split is measured. scripts/presence.py reads nginx, and counts only a phone session as looking — a desktop tab polls 134 times a minute and would manufacture presence indefinitely. Of the 20 pages you did not act on, **15 fired while your phone was demonstrably active**. Those 15 are rejections. The other 5 are alerts nobody saw.

(4) No threshold below is a recommendation until it is pre-registered — and the headline P is the WINNING band's, uncorrected. The 15-minute page band is the best point on a curve I drew from 34 of your presses. Its lower neighbour agrees (10 min: $P = 0.015$); its upper neighbour

at 30 min does not ($P = 0.43$), so the effect turns on somewhere between 15 and 30 minutes and 15 is one side of an edge, not the middle of a plateau. **$P = 0.0009$ is the permutation P of the band that won a four-band sweep, and nothing in it pays for having chosen it.** Four bands were tried (60/30/15/10), so the sweep-wide figure is $0.0009 \times 4 = 0.0036$ by Bonferroni (Šidák: $1 - (1 - 0.0009)^4 = 0.0036$, the same to four places). **It still clears comfortably at any conventional level, so saying so strengthens the row rather than killing it** — but the 10-minute band does not survive the same correction ($0.015 \times 4 = 0.060$), which is a second reason not to go past 15. A threshold picked off its own curve still has to be frozen and measured forward before it earns anything; the pre-registration is written out under §3.2 rather than promised.

3.1 · What is already automatable — the gap is DELIVERY, not DETECTION

Two different questions, and the desk has been answering only one of them. Did the machine see it? is about leg OPEN lines, which leg_watch writes silently by design — 49 a day, most of them dead inside half an hour. Did the machine tell him? is about ALERT lines, of which there were 25 all week. Everything worth having lives in the space between those two numbers.

ON YOUR 34 ENTRIES INSIDE THE LOG WINDOW	COUNT	READ
A leg was already running when you pressed	33 / 34 97%	the live watcher's own log, not a replay. \$2 got 33 of 36 by re-running the leg code causally — the same answer from a different direction
— and it was running YOUR way	26 / 34 76%	you are a continuation trader; \$2 measured the same thing at 30 of 33
A page reached your phone first (same side, within 30 min)	3 / 34 9%	15 Sep 10:29 (–1m), 18 Sep 16:11 (–4m), 17 Sep 00:25 (–23m). The best of them landed 58 seconds ahead — that is simultaneity, not a warning
A page reached your phone first (same side, within 60 min)	8 / 34 24%	the other five all landed 41–58 min out. Shuffled to random minutes, chance gives 18% — $P = 0.24$
Legs detected, 2026-09-15 05:55Z → 2026-09-18 21:00Z	176	49/day. Median life 22 min, p75 41, p90 60
Median age of the same-side leg you pressed into	28 min	p25 7, p75 42. The first page band is 60 min

The brief says the desk beat you to a leg on 15 September and it did — but look at what the line says.

06:50:38Z leg OPEN DOWN from 29362.75 (silent – 17 legs/session, most die young), and you sold at 06:56:28. The detector was 5m 50s early and the word silent is in the log line. It knew, and the delivery policy decided you did not need to hear it. That entry made **+\$623.00**.

3.2 · The page band, priced — coverage against noise, with a null

Every row is the same detector. The only thing that changes is **how old a leg has to be before your phone rings**. "Chance" is your own presses shuffled to random minutes of the same sessions, 20,000 draws, sides kept — the only control that is legitimate when there are no recorded passes.

FIRST PAGE AT	PAGES A DAY	PRECEDED YOUR ENTRY (≤60 MIN)	CHANCE RATE	P	READ
60 min — the band that runs today	5	3 / 34 9%	12%		0.86 NOTHING. The 60-minute band alone is behind chance
30 min	19	14 / 34 41%	38%		0.43 still nothing. Halfway is not halfway
15 min	32	28 / 34 82%	56%	sweep-wide	0.0009 0.0036 the best point on the curve. 26 points of separation over chance, and it survives being corrected for having won a four-band sweep
10 min	41	29 / 34 85%	68%	sweep-wide	0.015 0.060 one more entry covered, nine more pages a day, and the null catches up. Does not survive the correction. Past the knee
The correction, stated once	Four bands were swept; the P column's left figure is the winning band's own permutation P and pays nothing for the choosing. Bonferroni over the four: 15 min 0.0036, 10 min 0.060, 30 min 1.0, 60 min 1.0. The recommendation is unchanged				—

★ PRE-REGISTRATION — WRITTEN BEFORE IT SHIPS, BECAUSE OTHERWISE NEXT WEEK IS A SECOND FIT

Rev 1 put “freeze it and measure forward” on the card as a kill criterion after the change ships, which is not a pre-registration — it is a promise to

look. Here is the frozen version, and nothing about it may be re-chosen next Friday:

FIELD	FROZEN VALUE
Band	first page at 15 minutes of leg age, every leg, no travel floor, no other filter. Destination: dashboard. The 60-minute phone page is unchanged
The measured quantity	fraction of his entries preceded by a same-side page within 60 minutes, on entries opened after the band change lands
Coverage target	$\geq 56\%$ — the shuffled chance rate at this band, not the 82% in-sample point. Anything at or below 56% forward is the null
Kill line	forward coverage below 50% over the first 30 forward entries, OR he stops reading the pages. Either puts the band back to 60 and the row dies
Re-scored on	Friday 2026-10-16 — four report cycles, ~120 entries at this week's rate, in next Friday's §3 under this heading
What may NOT be changed	the band, the 60-minute coverage window, the same-side rule, or the target. Changing any of them turns a test back into a fit, and this row's only claim to significance is that it is measured on data it was not drawn from

Every row in that table is a model of a policy, so here is the world. The rows above re-band the detector cleanly — a page at leg age T , every leg, no other filter. What actually shipped this week was 25 pages across the 60/90/120/180 bands with a travel floor on top, **7 a day**, and measured the same way they preceded **8 of your 34 entries (24%) against a chance rate of 18%, $P = 0.24$.** Both readings say the same thing and it does not depend on which one you prefer: **the alerting you are running today is not distinguishable from telling you to look at random moments of the same day.**

The shape is what matters, not the winning row. The separation is not a spike on one threshold — it builds from nothing at 60 min, is still nothing at 30, arrives at 15 and starts being eaten by its own null at 10. That is a knee, and a knee is the shape of a real effect; a single significant cell surrounded by noise is the shape of a fitted one. It is also why **32 pages a day is the honest cost and not a footnote.** You told the desk you look periodically when you have time. Thirty-two pages a day is not periodic, and the first thing that happens to an alarm nobody can answer is that it becomes wallpaper.

So the recommendation is the band, not the bell. Move the detection threshold to 15 minutes, and put the live leg list on the dashboard you already have open — presence says you are on it — rather than on the phone. Keep the phone for the 60-minute band that exists now. That buys the 82% coverage where

you are already looking and leaves your pocket alone. It is also, not coincidentally, the thing that fixes §1's blocker — see 3.4.

3.2b • But HOW FAR AHEAD does a 15-minute page land — and is the leg still alive when he presses?

This is the poke the section's own numbers invite, so here is the answer before anyone else asks it. Coverage is scored on a ≤60-minute window; the median leg dies at 22 minutes; and §3.1 dismisses the current band's pages precisely because “the other five all landed 41–58 min out”. If the covered entries are mostly covered by pages that arrived forty minutes early about legs that were dead by the press, then 82% is bookkeeping rather than a warning.

The instrument is the live journal, not a replay. Every leg below is a leg OPEN line gazbot7-leg-watch actually wrote (176 of them) and its life is the leg OVER after Nmin line that closes it (165 closed). The watcher counts its bands from the leg's start, and writes the OPEN line 8 minutes later at detection — the 09-15 leg that opened at 09:36:07Z logged quiet 30min at 09:58 and ALERT 60min at 10:28, which fixes the offset exactly. A band-T page therefore lands T–8 minutes after the journal line. That 8-minute offset is also why §3.1 says the median leg dies at 22 minutes and the journal's own OVER lines say 30: §3.1 measures life from DETECTION, the log measures it from the leg's START. Same legs, same deaths, two origins — 22 + 8 = 30, and p75 41 + 8 = 50, and p90 60 + 8 = 70, which is the whole distribution reconciling term by term.

FIRST PAGE AT	PAGES A DAY	COVERED (≤60 MIN, SAME SIDE)	MEDIAN LEAD	PAGED LEG STILL ALIVE AT THE PRESS	READ
60 min — the band that runs today	11.0	10 / 34 29%	31.4 min	5 / 10	⚠ three numbers exist for this band and they are three different things. 3 of 34 is §3.2's clean re-band scored at ≤30 min; 8 of 34 is what the pages that actually went out achieved (60/90/120/180 bands with a travel floor, 25 pages, ≤60 min); 10 of 34 is this journal re-band — every leg that reached 60 minutes, no travel floor, so it pages more often

FIRST PAGE AT	PAGES A DAY	COVERED (≤60 MIN, SAME SIDE)	MEDIAN LEAD	PAGED LEG STILL ALIVE AT THE PRESS	READ
					than the live policy does
30 min	26.2	20 / 34 59%	22.6 min	14 / 20	—
15 min	45.5	29 / 34 85%	26.5 min	19 / 29 66%	the live journal says 85% where §3.2's replay says 82% — two instruments, one answer
10 min	48.0	29 / 34 85%	26.4 min	21 / 29	no more entries covered; only more pages

AT THE 15-MINUTE BAND, THE THREE QUESTIONS THAT DECIDE WHETHER 85% IS REAL	ANSWER	READ
Was the paged leg still alive when he pressed? (the watcher's own 3×ATR rule, from its own OVER line)	19 / 29 66%	The main worry does not hold. Two thirds of the pages were about a leg that was still running when he acted
Was a same-side 1×ATR leg alive at the press? (§2.3's stricter ruler)	23 / 29 79%	the stricter definition agrees, and matches §2's 33-of-36 base rate
How many pages landed inside 20 minutes?	10 / 29	This is the honest cost. Median lead is 26.5 min (p25 17.5, p75 45.8) — longer than the median leg has left to live
Coverage if a page only counts when it lands inside 20 minutes	10 / 34 29%	it collapses. The 85% is bought with a wide window, not with punctual pages
... inside 30 minutes	16 / 34 47%	half-way between, as you would expect from the lead distribution

So the answer is: it is a real warning, and it is early rather than late. The failure mode the question was hunting — pages about dead legs — is **not** what is happening: 19 of 29 paged legs were still running at the press and 23 of 29 had a live same-side 1×ATR leg. What is true is that a 15-minute band delivers its median page **26 minutes** before he acts, and only a third of them inside twenty. **That qualifies MONDAY #2 without killing it**, and it is exactly why the row's second half — put it on the dashboard, not the phone — matters more than the band: a warning that arrives half an hour early is useful on a screen you are already looking at, and is noise in your pocket. △ Two entries out of 34 are the whole difference between this table and §3.2's; the replay and the journal are two views of one four-day week, not independent confirmations.

3.3 · What is NOT automatable, and exactly how far off it is

§1 and §2 both stop at the same wall, so here is the wall with numbers on it instead of an assertion.

WHAT A RULE WOULD NEED	HELD TODAY	NEEDED	DISTANCE, IN WEEKS AT THE CURRENT RATE
ENTRY rule — presses he made	36 §2's count	30	cleared
ENTRY rule — passes he made	0	30	INFINITE. The button shipped 2026-09-18 and has never been pressed. Nothing else on this box records a look that ended in no trade
ALERT-selection rule — pages he took	5	30	~3.6 weeks at 7/week. This is the binding side here, not the negatives
ALERT-selection rule — pages he refused while looking	15	30	~0.7 weeks at 21/week. The first true negatives this desk has ever held
EXIT rule	119 attempts 86% capture	—	NO ROUTE. 81 of your 92 entries ended on your claim button, the \$400 and \$600 rungs have never fired once, and the standing rule forbids a 120th calibration. I did not run one

Read the two middle rows together, because they invert the usual complaint. For the question "would he have bought here", the missing side is the negatives and it is missing absolutely. For the question "is this leg worth telling him about", the negatives are already arriving faster than the positives — 21 a week against 7 — because every page he declines while his phone is awake is a label. The alert-selection problem is **about four weeks from being a properly posed learning problem**, and the press problem is not on a clock at all until the PASS button gets used.

And there is a second wall behind the first, worth knowing before anyone celebrates 30 passes. §2 measured what the box can actually see at the moment you press: position-in-range and VWAP stretch are indistinguishable from a random minute, and drift, the desk's own direction detector, could not even read at 24 of 36 entries and returned confirmed=false at all 12 where it could. Strip those out and the features that survive are **a running same-side leg and ATR percentile**. Two. Thirty passes buys you a boundary drawn in a two-dimensional space — which is worth having, and is not a model of your judgement.

3.4 • The pieces that need no model at all, each scored on this week

Three rules that require nothing to be learned about you. A stopwatch, a calendar, and a loss cap. Each is priced in your units: dollars a day at four lots.

(a) The hold-time stopwatch — and the number \$1 handed to this section

\$1 measured the over-8-hour bucket at **0 winners from 4** and **−\$6,612.44** and said explicitly that what a stopwatch recovers was not computed there. It is computed here: flatten each of those four at exactly T+8h at the tape price. **The delta depends only on exit prices, so the paper engine's fabricated entry fills cancel out of it.**

ENTRY	HELD	ACTUAL	AT +8H	DELTA	WOULD HAVE BEEN
2026-08-21 13:13 LONG	70.8h	−\$2,149.44	29370.00	+\$1,168.00	−\$981.44
2026-09-01 08:24 LONG	9.0h	−\$1,466.00	29192.25	+\$522.00	−\$944.00
2026-09-04 05:29 LONG	8.7h	−\$341.00	29580.50	−\$202.00	−\$543.00 — the stopwatch made this one worse
2026-09-10 05:40 LONG	13.6h	−\$2,656.00	29184.25	+\$232.00	−\$2,424.00
TOTAL over 28 sessions	—	−\$6,612.44	—	+\$1,720.00	+\$61.43 a session traded

So the eight-hour hole is a \$1,720 opportunity, not a \$6,612 one. All four were already deep underwater by hour eight — the stopwatch recovers **26%** of the damage, one of the four gets worse by \$202, and **68% of the total is a single trade** (21 August). That is a much smaller prize than the bucket's headline, and it is the right number to carry forward.

Swept across every width, split at Monday. **Each column divides by the sessions in its own window**, so the 8-hour rung reads +\$74.78 here against the +\$61.43 above — same \$1,720, 23 sessions rather than 28.

CAP	BEFORE THIS WEEK 52 ENTRIES / 23 SESSIONS	THIS WEEK 39 ENTRIES / 5 SESSIONS	READ
2 hours	+\$69.07/sess	+\$2.60/sess	4 entries touched this week
3 hours	+\$152.93/sess	+\$24.80/sess	3 touched
4 hours	+\$117.24/sess	−\$81.60/sess	sign flips on the new data
6 hours	+\$156.15/sess	−\$186.40/sess	sign flips harder
8 hours — the pre-registered rung	+\$74.78/sess	\$0.00/sess	nothing to do. Your longest hold this week was 7h01m

The honest read, both halves. On the 23 sessions the finding came from the sweep is positive at every width from 2h to 8h, and stays positive at 10h and 12h on the whole book — a smooth band, not a spike, which is the shape of something real. On the first five sessions of new data it earned nothing at 8h because you

never held that long, and where it did have something to do — 4h and 6h — **it lost money**. Two or three entries either way; that is noise-sized, not a refutation. But it is the first forward week and **it did not confirm**.

One trade settles the question of whether a shorter cap would have rescued the week. Tuesday's 13:39Z long, $-\$1,288.00$, held 7h01m into the 20:40Z clock: capped at 1h it is $-\$1,180$, at 2h $-\$1,274$, at 3h $-\$1,294$, at 4h $-\$1,486$, at 6h $-\$1,764$. **No stopwatch of any width saves it**. It was wrong from the first hour and the tape happened to come back into the clock. Whatever that trade needs, it is not a timer.

(b) The event-day stand-aside

WINDOW	DAYS	ENTRIES	NET	PER SESSION AVOIDED
This week — FOMC Wednesday		1	6	$-\$833.00$ +\$166.60/day across the 5-day week. The only red day of the week
Whole book — 12 Aug CPI, 4 Sep NFP, 11 Sep CPI, 16 Sep FOMC		4	12	$-\$1,518.49$ +\$54.23 a session across 28; +\$379.62 per event day avoided
Ordinary days		25	80	$+\$1,843.55$ \$1's figure, reproduced

Point estimate first: +\$166.60 a day this week, +\$54.23 a session over the book. Uncertainty second, and it is the size of the finding: **four days**, one of them (12 Aug CPI) positive at $+\$224.50$, and the split was chosen after the losses were visible. Three of four negative is not a distribution. It also does not need to be a switch — the T-60 and T-15 pages already exist, they name the event and refuse to name a direction, and the Wednesday damage was **one entry at T+52 into a slide that ran another 33 minutes**, not the day. Keep the page. Do not build the bench.

★ “**+\$166.60 A DAY**” IS ONE OF THREE ANSWERS, AND THE ROW WAS PRICING ONLY THE UPSIDE OF THE AVOIDANCE

The $+\$166.60$ is FOMC Wednesday's $-\$833.00$ spread across five sessions — in a week that made **+\$583.30 a day**. It prices what a whole-day stand-aside saves and never prices what it costs, and Wednesday also held the week's single biggest winner (**+\$819.00 at 13:13**) and 3 of its 6 entries won. Since the row's own argument is that the damage was one entry at T+52 rather than the session, the whole day is the wrong unit — so here are

all three units side by side, every one at 4 lots and every one as \$/day against the week's +\$583.30.

WHAT YOU STAND ASIDE FROM	ENTRIES REMOVED THIS WEEK	THEIR NET	\$/DAY AFTER	DELTA	WHAT IT ALSO REMOVES
The whole event DAY — the unit the row is priced in today	6	−\$833.00	+\$749.90	+\$166.60	the +\$819.00 at 13:13 and the + \$450.00 at 19:08 — \$1,429 of winners deleted to avoid \$2,262 of losers
T−60 to T+60 around the release (17:00-19:00Z)	1	−\$1,272.00	+\$837.70	+\$254.40	nothing. It removes the 18:52 entry and keeps both the +\$819.00 and the +\$450.00
Only entries opened AFTER the release	2	−\$822.00	+\$747.70	+\$164.40	the +\$450.00 at 19:08 — the reversal that recovered a third of the damage

Over the whole book (29 sessions traded, 4 event days) the ranking reverses: whole-day **+\$52.36/session**, T±60 **+\$43.86**, after-the-release **+\$20.60**. The reason is the cost side this week hides: across four event days a whole-day stand-aside removes **\$2,267.51 of winners** (+\$819.00 on 16 Sep, +\$614.01 on 11 Sep, +\$224.50 on 12 Aug) to avoid \$3,786.00 of losers, and on **12 August it deletes a winning day outright** — one entry, +\$224.50, and nothing else traded that day.

The difference between the three IS the size of the finding, and the difference is one trade. T±60 wins this week by exactly the 18:52 long and nothing else; over the book it is the same single entry. **So this is not evidence for switching the unit** — n = 1 is not a window study, and picking the window that isolates the one big loser after seeing it is the definition of fitting. What it is evidence for is that **“+\$166.60 a day” should never have been quoted as the value of the policy**: it is the value of one unit, on one event, with the winners it also deletes left out of the arithmetic. The row stays a **page**, the page already fires at T−60 and T−15, and no bench of any width is justified on four event days.

(c) The \$250 daily loss limit — the one that is worse than nothing

This is listed in the brief as a piece that is already automated and running. It is running. Here is what it did.

WHAT WAS CHECKED	WHAT WAS FOUND
Can it reach your book?	No. <code>daily_loss_limit.py</code> writes off to six names in <code>data/gate_switches.env</code> — <code>grind_long</code> , <code>capitulation_long</code> , <code>abs_veto_long</code> , <code>exhaustion_short</code> , <code>abs_veto_short</code> , <code>rgv_short</code> . <code>day_rider</code> is not among them and its switch lives in a different file, <code>data/day_rider.env</code> , which this script never opens
Did it fire this week?	Twice — 2026-09-15T20:40:40Z and 2026-09-16T12:24:40Z. Both journal records end "benched": []. All six gates were already =off, so it changed nothing and wrote no line to the trial log
What did it tell you?	Two critical pages reading "All six gates benched; they re-arm at the 22:00Z reopen." That sentence is a hard-coded string; the code never consults its own changed list. It benched none and said it benched six
What number fired it on Tuesday?	<code>pnl_usd</code> : -418.50. Your real session was -\$114.00; the other - \$304.50 is the BADFILL: entry whose price the paper engine fabricated. Its query sums every closed row with no <code>data_quality</code> filter. Without the fabricated fill it would not have fired
And if it COULD stop you — would you want it to?	No. Walked across 29 sessions of your book it trips on 10, and what you opened afterwards on those days nets +\$1,945.50. Stopping you there costs -\$67.09 a session. This week alone: - \$157.00, or - \$31.40 a day

Be precise about what is refuted. The \$250 level was derived on the single-lot automated book — 43 sessions, +\$35/day to +\$87/day — and it is pre-registered and frozen for 60 forward sessions on that book. Nothing here touches that. What is refuted is **extending it to your hand**: on the four sessions where it trips and you keep going, you traded back +\$789.00, +\$1,084.50, +\$157.00 and -\$85.00. Your losing sessions are not the runaway shape the limit was built for — the 20:40Z clock and your own claim button are already doing that job. On this book it is a rule that only ever removes trades, and the trades it removes made money.

(d) The automated desk took NOTHING all week — was that a decision or a fault?

The front table prints “opened by any automated gate | 0 | \$0.00” and moves on, and the premise of this whole report (“it is about you and not about it”) depends on the answer. **It is a decision, it is re-taken every single tick, and the gates are automatically re-armed every night.** Here is the machinery, with timestamps.

GATE	SWITCH NOW	WHO SET IT, AND WHEN	LAST TIME IT TRADED
grind_long	off		2026-09-03

GATE	SWITCH NOW	WHO SET IT, AND WHEN	LAST TIME IT TRADED
		the ROUTER, 2026-09-14T15:45:17Z — 3m 05s after the open-hour watcher armed it at 15:42Z on a confirmed break	
abs_veto_long	off	the ROUTER, 2026-09-14T15:45:17Z — the same tick, the same arm	2026-09-11
capitulation_long	off	the ROUTER, 2026-09-17T22:05:23Z — and at 22:00Z on 09-14, 09-15 and 09-16 before that	2026-09-11
exhaustion_short	off	the ROUTER, 2026-09-17T22:05:23Z — same nightly pattern	2026-09-11
abs_veto_short	off	the ROUTER, 2026-09-17T22:05:23Z — same nightly pattern	2026-09-10
rgv_short	off	reactivate_gates.py's HOLD set — never auto-armed, for the 2026-08-09 churn reason recorded in the file's header	2026-07-30

AND THIS IS WHAT HAPPENS EVERY NIGHT	THE RECORD
Three of the six are AUTOMATICALLY RE-ARMED at 22:00Z	gazbot7-gate-reactivate.timer, nightly. Its own log on 09-15, 09-16 and 09-17: “gates auto-reactivated at Paris midnight (1-day off expired): abs_veto_short, capitulation_long, exhaustion_short”. abs_veto_long and grind_long are start-benched by policy (“router arms on evidence”); rgv_short is HELD
... and the router benches them again within minutes	data/router_trial_log.txt: changed: {'capitulation_long': 'off', 'exhaustion_short': 'off', 'abs_veto_short': 'off'} at 22:00:14Z on 09-14, 22:00:13Z on 09-15, 22:00:13Z on 09-16 and 22:05:23Z on 09-17. Four nights, four re-arms, four benches
The only ARM anywhere in the report week	the open-hour watcher, 2026-09-14T15:42Z, grind_long + abs_veto_long on a confirmed break (ER 0.37, ATR 19, bias +96pt UP, new extreme). Benched 3 minutes later
The router's stated reason, unchanged to the last tick	22:35:11Z tonight: “no change, all six gates stay off ... UNTRADEABLE METER is a CONFIRMED STAY-OUT for a forty-fifth consecutive cycle”, and it says in the same sentence that it holds full arm/bench authority and is declining on the tape, not because it is forbidden
Tonight's 22:00Z re-arm	did not happen and correctly said so: “VENUE SHUT (Fri 22:00Z) — already all off, nothing to do”

So: benched deliberately, by the router, every session, and re-armed for it to re-decide every night. Nobody turned the tournament off and forgot. **The last fill any gate booked was 2026-09-11** — the Friday before this report week — so the automated desk's silence starts one trading day before the window this report grades, which is why the refocus is not an artefact of a stuck switch. **Re-arming is therefore NOT a play:** the mechanism already re-arms three gates a night and the router benches them on a stated reading of the tape. $\triangle$ The one thing worth putting a name on is what that reading rests on: 45 consecutive

cycles of an **untradeable-day meter**, which is a hindsight measure of the day just traded and not a forecast of the next one. Whether a gate should be armed is a tournament question and the tournament is not in this report — but “the router keeps saying stay out” is a decision this desk should re-examine the next time the tournament is reviewed, not a fact about the tape.

3.5 · Ranked shortlist

#	WHAT	MEASURED — THE NUMBER, THEN THE UNCERTAINTY	VERDICT
1	Drop leg_watch's first page band from 60 minutes to 15, and route it to the dashboard rather than the phone. A parameter change. No model, no fitting on your judgement, no negatives required	The 25 pages actually sent preceded 24% of your entries against a chance rate of 18% — P = 0.24; the 60-minute band on its own reads 9% against 12%, P = 0.86. Either way, no separation. At 15 minutes the same detector precedes 82% against 56%, P = 0.0009 — and 0.0036 corrected for the four-band sweep it won, which still clears; 10 min reads 0.015 raw and 0.060 corrected, 30 min does not clear either way (0.43) — an edge between 15 and 30, not a plateau. The live journal reaches the same place independently: 29 of 34, 85%. Uncertainty, and there are now three: 34 presses over four days; the 15 was chosen off this curve; and the median page lands 26.5 minutes ahead, with only 10 of 29 inside twenty minutes — restrict coverage to pages inside 20 min and it falls to 29% (§3.2b). The paged leg was still alive at 19 of 29 presses, so this is an early warning, not a stale one. Cost is real and is the reason for the dashboard half: 7 → 32 pages a day. Pre-registered under §3.2: band 15, target ≥56%, kill below 50%, re-scored Friday 2026-10-16. Nothing in it may be re-chosen next week	PROMISING the best thing in this section
2	Keep standing aside on event days — as a page, never as a bench. The calendar and the T-60/T-15 pages are already built and running	+\$166.60/day this week if you sit out the whole DAY (FOMC Wednesday, 6 entries, — \$833.00, the worst day of the week) — but that is one of three units and it prices only the upside. Sitting out T±60 around the release is +\$254.40/day this week and keeps the +\$819.00; sitting out only after the release is +\$164.40. Over the book the order reverses: + \$52.36 / +\$43.86 / +\$20.60 a	PROMISING as a page only

#	WHAT	MEASURED — THE NUMBER, THEN THE UNCERTAINTY	VERDICT
		session. Uncertainty: four event days, one of them positive at + \$224.50, split chosen after the losses — and the entire difference between the three units is ONE trade, the 18:52 long. The whole-day version also deletes \$2,267.51 of winners across four event days, and deletes 12 August, a winning day, outright	
3	Turn the paged legs into labels. Every page you decline while your phone is awake is a negative example. This is the only item that unblocks \$1's BLOCKER without waiting on the PASS button	Already held: 5 positives / 15 negatives from 25 pages, with 5 more that nobody saw. Negatives accrue at 21/week, positives at 7/week — so positives are the binding side and the pool reaches 30/30 in about 3.6 weeks. At a 15-minute band it reaches it in roughly one. Uncertainty: \$0/day directly — this buys a dataset, not a trade	ENABLER
4	The eight-hour stopwatch. \$1 ranked it second and deferred the recovery number to here	\$0.00/day this week — 0 of 39 entries reached 8h, longest 7h01m. Lifetime it recovers + \$1,720.00 over 28 sessions = + \$61.43/session, which is 26% of the bucket's -\$6,612.44 headline, not all of it. Uncertainty: four cases, 68% of the total is one trade, one case gets \$202 worse, and where the sweep had work to do this week (4h, 6h) it lost -\$81.60 and -\$186.40 a session. Cheap insurance against a shape that did not recur	PROMISING forward week neutral
5	Do NOT extend the \$250 daily loss limit to your book. It stays pre-registered and frozen on the automated single-lot book where its evidence came from	-\$67.09 a session over 29 sessions — it trips on 10 and the trades you opened afterwards net +\$1,945.50. This week: -\$31.40/day. Uncertainty: one-directional, because the rule can only ever remove trades — and on 6 of the 10 trips you had already stopped, so the entire cost sits in 4 sessions	REFUTED for this book only
6	Defect — the loss limit pages a claim its own code contradicts. The notify string says "All six gates benched"; the journal says "benched": []. Drive the sentence from changed	Fired twice this week, benched zero both times, paged you both times. Every gate it can reach was already =off, so the guard is structurally incapable of doing the thing it announces	FIX — DEFECT
7	Defect — the loss limit counts fabricated fills. session_pnl() sums every closed row with no data_quality filter	Tuesday's trip read -\$418.50: -\$114.00 real, -\$304.50 fabricated by the paper engine. Below the limit on the real number. A guard whose trigger can be moved by a price that never printed	FIX — DEFECT

#	WHAT	MEASURED — THE NUMBER, THEN THE UNCERTAINTY	VERDICT
8	Defect — a hand-run wrote a live PAGED line. Give event_alert.py a --dry-run that logs nowhere, or stamp manual runs	2026-09-18T07:04:03Z PAGED T-15 FOMC, off the *:0/5:30 cadence, two days after the only FOMC in the calendar, with the timer's own run at 07:05:30 logging "nothing in a band". It cost this section one false LED and would have inflated the alerting score by 25%	FIX — DEFECT
9	Put the PASS count where he already looks. Everything in this report is blocked on a button, and nothing in it measures whether the button gets pressed — the only verification offered for MONDAY #1 was a one-line python snippet a human has to remember to run	Passes recorded today: 0. At an observed rate of zero a week, 30 never lands — there is no date to quote, and that is the honest form of the answer. From this cycle the count is generated into the front box of every Friday report from operator_reads.jsonl at render time, so a zero next Friday is a visible fact rather than a thing nobody checked. What is still missing is the dashboard tile — one number, beside the metric strip he already reads. Uncertainty: none; it is a counter. This is the same distinction this report draws about last week's 58 ungraded plays: "we checked and it was zero" and "nobody checked" are opposite instructions	FIX — INSTRUMENT unblocks nothing, measures everything
10	The automated desk is benched by DECISION, nightly — not by a stuck switch. Recorded here because the premise of this whole report depends on it and Rev 1 asserted it without evidence	All six gates =off. Three of them (abs_veto_short, capitulation_long, exhaustion_short) are auto-re-armed at 22:00Z every night by gazbot7-gate-reactivate.timer and re-benched by the router within minutes — 22:00:14Z 09-14, 22:00:13Z 09-15, 22:00:13Z 09-16, 22:05:23Z 09-17. The only arm in the week (open-hour watcher, 09-14 15:42Z, grind_long + abs_veto_long) was benched 3m05s later. Last fill by any gate: 2026-09-11. Uncertainty: the router's stated reason is 45 consecutive cycles of an untradeable-day meter, which is hindsight about the day just traded, not a forecast — a tournament question, and the tournament is not in this report	ANSWERED no action

So — can it be automated? Telling you when to look: yes, and the machinery is already built and already running with one number set **wrong**. That is worth doing on Monday and it costs a config change. **Deciding what to take: no, and not because the answer is no — because the question**

has never been asked in a form that has two sides. You have 36 yeses and zero noes, and the fastest route to the second kind is not the PASS button (which needs you to remember) but the 21 pages a week you are already declining with your phone in your hand. **Deciding when to get out: no route at all**, and 119 calibrations is enough evidence to stop looking for one. The stopwatch and the calendar are worth keeping for what they are — small, model-free, and honest about their four data points. And the one guard everybody assumed was protecting you is watching a book it cannot touch, announcing actions it did not take, on a number a fabricated fill can move.

4 THE RUN CHARTS — every fill of the week, marked on the day's own price path

Five sessions, one chart each, **00:00-21:00 UTC** — wide enough that not one of your 61 legs falls off the edge, including the 00:25 and 02:37 overnight entries. **A triangle is an entry** (apex up = long, apex down = short), **a circle is an exit**, the line between them is the hold, and green/red is the money. The colour pair is validated for deuteranopia and colour is never the only carrier: shape gives direction and every labelled exit carries a signed dollar figure. Hover any marker for the full fill.

⚠ **These are LEGS, not entries.** A scale-out books one marker per lot leaving, so Monday's four decisions appear as nine markers. Count triangles that share an entry price as one decision — or read Section 1's tables, which count decisions throughout. **Two charts follow the five sessions:** the FOMC hour up close on MNQ, and the same two hours on **gold**, which carries no fills because neither you nor the desk trades it — it is there because FOMC moved MGC further than it moved the Nasdaq.



Four decisions, nine booked legs - the difference is the scale-out, which books one row per lot leaving. A 09:09 short scaled out twice into the European morning and then gave \$107.25 back on the remaining two lots; the 10:33 long was opened twelve seconds after that claim, straight out of it and on the opposite side. The day is the 13:18 long: four lots, four separate exits between 13:45 and 13:48 for +\$607.48 together. Then the 13:49 re-entry - twenty seconds after the last lot left, 123 points above the entry it had just closed out of - lasted 2m50s and cost -\$162.00.



The clearest picture in this report of the one shape that hurts. Four entries before noon made +\$1,582.00 - the 06:56 short alone scaled out four times for +\$623.00 - and the 12:24 short gave -\$408.00 of it back. Then the 13:39 long is the single flat line running from 13:39 all the way to the right-hand edge: seven hours and one minute, closed not by a decision but by the

20:40Z clock, -\$1,288.00. One held trade turned the best morning of the week into a red day. △ A seventh row that day - 08:20, four lots, -\$304.50 - is NOT on this chart and not in the -\$114.00: it is the BADFILL entry whose price the paper engine fabricated on 3 of its 4 lots.



The only FOMC in the report week and the worst day in it. The morning was already -\$990.00 before the decision: a 06:12 long held two hours for -\$228.00 and an 08:11 short held four hours into a rising tape for -\$762.00. The 13:13 long then scaled out three times for +\$819.00, the biggest single decision of the week. The decision lands at 18:00Z, in the gap between the 17:25 claim and the 18:52 entry - and the 18:52 long, placed 52 minutes AFTER the release with the slide still running, was out 16 minutes later for -\$1,272.00. The short that followed it at 19:08 recovered +\$450.00 of it in half an hour. Read the two together: he was on the wrong side of the move, then on the right side of the same move, thirteen minutes apart.




The cleanest day of the week and the one that shows the ladder working. The chart opens with a position already on - the 00:25 long, held 6h35m through the Asian session for +\$272.00 - and then four of the eight entries take their first lot off at the TARGET_100 rung while the tape grinds 330 points higher. 11:55 and 12:19 are back-to-back continuation entries into the same leg, +\$507.00 and +\$375.38 on the runner lots. The only real loser is the 13:26 long, held 1h28m across the US open for -\$441.00, and the two 19:35/19:38 entries at the right edge are a short and a long thirty-two seconds apart, together -\$42.00: a flip, priced.



Fifteen entries, the busiest day on record for this book, and ten of them winners. The whole of the morning is long and the whole of the afternoon is short, which is the continuation description in Section 2 drawn as a picture: he takes the side of the leg that is running. The first four entries ride a 240-

point overnight and European rally for +\$1,223.50; the 11:21 and 12:16 shorts fade the turn early and lose -\$228.00 between them; the 13:55 short catches it properly for +\$434.50. The last three entries are the afternoon chopping him - -\$117.50, -\$160.00, -\$208.00 - and then the 17:07 long closes the week at +\$284.00. The desk went flat at 18:33Z and stayed flat.



This is the window the event calendar prices. The 30 minutes from the release ran 2.7-3.5x a normal 30-minute window on both FOMC days in our MNQ tape - 98th and 100th percentile. RANKED OVER ALL 67 RUNS of this week's census, this one chart holds the week's two biggest: -229 points from 18:42, which he FOUGHT (long into it at 18:52, out 16 minutes later for -\$1,272.00), and -194 points from 19:04, which he CAUGHT (short at 19:08 for +\$450.00). It also holds ranks 4 and 5 - -144 points from 18:22 and +139 points from 19:26 - so FOUR of the week's top five runs happened inside these ninety minutes. [△](#) Rev 1 said census_summary.json 'picks two of the week's top three runs' and named the 18:22 and 19:26 moves. That file cannot rank the week: its own _note reads 'sat-out runs by |move| desc' and its top15/top25 lists cover only the 51 runs the desk SAT OUT, of 67. Those two moves are ranks 1 and 2 of the sat-out list, not of the week. The ranking above is re-derived over the full 67 from the census run table. He was flat through the release itself - the previous position was claimed at 17:25 and the next was opened at 18:52.



The same decision on the other instrument we capture, at the gold multiplier. NAME THE BAR SOURCE, because the number is meaningless without it: capture.db 5-second MGC bars, which start 2026-08-04. The 18:00-18:30Z window ran a 66.1-point high-low range. Against the 31 PRIOR 18:00-18:30Z windows in that tape that is 6.5x the median and 2.8x the maximum - which is where the front page's 6.5x and event_calendar.json's standing '6.6x its median / 2.8x its previous MAXIMUM' note come from - and that baseline is the same half-hour of the day over 31 days, not all windows, which neither of them says. $\triangle$ Against ALL 1,423 prior 30-minute windows it is 6.1x the median and the 99.6th percentile - large, but NOT 'the largest event response in either tape': NFP on 4 September ran MGC 110.1 points, 10.1x its median and the 99.9th percentile, and that claim is withdrawn. There is not one marker on this chart because the desk does not trade gold and neither does he. It is here for one reason, and that reason now has a number and a verdict attached - see the shortlist below.

4.1 • Ranked shortlist — what Section 4 says to do next

Rev 1 of this section floated a genuinely new lead in the gold caption — “if the event-day stand-aside on the card is ever turned into an event-day OPPORTUNITY, this is the instrument the measurement says to look at first” — and then stopped. No row on the action card carried it, this was the one section with no shortlist, and MONDAY #3 tells him to keep standing aside

on event days with no cross-reference either way. **A lead with no verdict and no kill criterion is not a finding, so here is the measurement.**

What the 30 minutes from each release were worth to **one lot**, on both instruments, at **MNQ \$2.00** and **MGC \$10.00** a point. Bar source: capture.db 5-second bars for both symbols. Range is the hindsight ceiling of the window — the perfect entry and the perfect exit, which nobody gets. Move is what a 1-lot long bought at the release and held for thirty minutes actually paid.

EVENT WINDOW (30 MIN FROM THE RELEASE)	MNQ RANGE PT → \$1 LOT	MGC RANGE PT → \$1 LOT	MNQ MOVE \$1 LOT	MGC MOVE \$1 LOT
Wed 12 Aug — CPI, 12:30Z	205.00 → \$410.00	51.10 → \$511.00	+\$31.00	+\$102.00
Fri 04 Sep — NFP, 12:30Z	218.00 → \$436.00	110.10 → \$1,101.00	−\$260.00	−\$792.00
Fri 11 Sep — CPI, 12:30Z	271.75 → \$543.50	97.90 → \$979.00	+\$163.50	+\$452.00
Wed 16 Sep — FOMC, 18:00Z	132.25 → \$264.50	66.10 → \$661.00	−\$48.50	−\$422.00
FOUR EVENT DAYS	\$1,654.00 \$413.50 an event	\$3,252.00 \$813.00 an event	−\$114.00	−\$660.00

LEAD	VERDICT	THE NUMBER, THEN THE UNCERTAINTY	REVIVAL CONDITION
1 • Trade the event window on GOLD. The idea the Rev 1 caption floated and did not price	PARKED	The size is real and the direction is not. Gold's event window is worth 1.97× MNQ's to a 1-lot — \$813.00 against \$413.50 an event, because a smaller point move carries a five-times multiplier. But a 1-lot long held through the release loses on both tapes: −\$660.00 on MGC and −\$114.00 across the same four events. Two up, two down — and the same two: MNQ and MGC moved the SAME direction on all four releases, so gold is not a second opinion, it is the same bet with a bigger multiplier. Uncertainty: four events; MGC tape only starts 2026-08-04; the desk has never traded gold and holds no MGC execution record at all; and the whole range column is a hindsight ceiling nobody can capture	Re-open when BOTH: (a) ten or more event days of MGC tape exist — ~2026-02 at the current calendar rate — AND (b) something that is not a direction call has been written down first, because four events say the direction is a coin flip and the money is in the range. Do NOT invent an MGC gate for it before then, and never by porting an MNQ one
2 • The cross-reference MONDAY #3 was missing. Section 4 floats an event-day OPPORTUNITY; the card says	NOT A CONTRADICTION	They are about different books. MONDAY #3 is about his MNQ hand, where four event days ran −\$1,518.49	n/a — this row exists so the two are read together. §3.4(b) prices the stand-aside three

LEAD	VERDICT	THE NUMBER, THEN THE UNCERTAINTY	REVIVAL CONDITION
keep standing aside on event days		over 12 entries and the damage was one entry at T+52. This row is about an instrument nobody trades, with zero entries and zero P&L. Standing aside on MNQ event days and measuring MGC event windows are compatible; doing the second is not a reason to stop doing the first	ways; this prices the instrument

So Section 4's honest output is one parked lead and a correction. The charts do the job they were built for — the 13:39 Tuesday flat line and the FOMC ninety minutes are the two shapes this week's money turns on, and neither is visible in a table. What the gold chart is not is a trade: it is 2.0× the dollar range and a coin-flip direction, on four events, in a market this desk has never sent an order to.