

GAZBOT V7 — the Friday report

Week ending **Fri 2026-09-25** · Mon 21 – Fri 25 September 2026 · published 2026-09-25
· MNQ, with MGC captured · 14 plays · front matter + 4 body sections

★★★ WHY YOU ALMOST READ LAST WEEK'S REPORT TONIGHT — AND THE GUARD THAT NOW STOPS IT

This report was one build away from being the 2026-09-18 edition, re-published under tonight's date. You are owed the mechanism, because you have been handed a silently stale build before and a promise is not a guard.

What happened. `friday_phases._report_week()` resolved the report week as “the most recent Friday whose 22:00Z session cut has already passed”.

The cron starts at **21:05Z on Friday — 55 minutes before that cut** — so on Friday night it always rolled back seven days. Measured at 21:05Z tonight it returned **WEEK = 2026-09-18** while the clock said 2026-09-25. Every phase was handed the wrong week, assemble's output resolved to `weekly_2026-09-18.html` which already existed, and the runner's freshness check passed it because **a git operation on 09-24 09:43 had touched every fragment**, making all four look newer than the 09-18T22:00Z cutoff. The log read sections: 0 built, 4 already had, 0 missing then HAVE assemble

– skipping. That is the exact failure the checkpoint's own docstring says it exists to prevent: “a way to publish last week twice”.

The two fixes, both on disk. One — on a Friday the report week is **today**, because a Friday-night run is about the week ending tonight. That is the only run instant that changes; a Saturday or Sunday resume already named this Friday, so the fix makes the cron AGREE with its own resume rather than contradict it. Two — an mtime cannot say what week a file is ABOUT, so every artifact now carries a **week stamp** written by the runner when the phase succeeds, and a stamp that disagrees with the report week is **STALE no matter how new the file is**. Verified tonight: with mtimes untouched, all four body phases flipped from HAVE to todo on the stamp alone.

△ **What that means for this document.** The front matter, the money, the action card, last week's graded ledger and the run charts are all rebuilt on **2026-09-21 → 2026-09-25**. The four long analytical body sections were written about 09-14→09-18 and **could not be re-run tonight** — each is a 45-60 minute job. They are stitched in **under a red banner naming the week they grade**, because the alternative was to publish nothing, and their tape-reading work is still the best on the desk. **Do not read a dollar figure**

out of Sections 1-4 as this week's. Every number in the front matter is this week's.

★ **THE WEEKEND CARRY CHECK — THE DESK IS FLAT, AND HERE IS WHAT SAYS SO**

Nothing is riding over this weekend. The last leg on the book closed at **20:37:52Z on Friday 25 September** (\$22.00 at 30,912.75) and no position was opened after it. `data/day_rider_state.json` reads `closed true`; its venue read at 22:07:06Z returned **net 0** with `venue_ok true` — taken after that final claim, which is what makes it evidence rather than an assumption. `core_health` at 22:07:55Z reads `flat true`, `halted false`, and `desk_reconcile` at 22:07:32Z reads `venue +0 = tournament +0 + rider +0 → unaccounted +0`.

⚠ **One field in that file will mislead you and it is named here on purpose:** `lots_open` reads **4**, not 0. It is the size of the LAST position, not the open size, and it is never cleared. `closed` is the field that decides, and the venue line is what confirms it. (`weekend_carry_box()` in this build checks `lots_open` first and is only saved from shouting by the `closed` test beside it.)

Related: LAST WEEK'S CARD below grades the flatten row this outcome speaks to — and grades it **ungraded**, because the final position was closed by hand more than two hours before either eod-flatten shot, so the flatten path was never exercised this week and nothing here says it works.

◆ **THE WEEK THAT WAS, AND THE WEEK COMING**

Everything in this section is read from `data/event_calendar.json` (regenerated 2026-09-25 21:27 UTC) and the times are computed into **Paris** from the New York wall-clock the calendar stores. **The impact**

column is OUR measurement on OUR tape, not a general claim about these releases — and where we have not measured one, it says so.

The week that was — 21 to 25 September 2026

No scheduled event in the report week.

The week coming — 28 to 2 October 2026

1 scheduled event(s) in the week of 28 September.

DATE	TIME	EVENT	INSTRUMENTS	WHAT IT DID TO OUR TAPE
Fri 2 Oct 2026	14:30 Paris 12:30Z · 08:30 ET	NFP Employment Situation (Sep)	MNQ, MGC	MNQ 218-220pt in the 30-minute window (100th percentile both times); MGC 3-6x

And the seven weeks after that — so nothing arrives as a surprise

Every confirmed date the calendar holds out to 13 November. The **28 October FOMC prints at 19:00 Paris, not 20:00** — the US and EU clocks change three days apart and the table below computes each row rather than offsetting it.

DATE	TIME	EVENT	INSTRUMENTS	WHAT IT DID TO OUR TAPE
Wed 14 Oct 2026	14:30 Paris 12:30Z · 08:30 ET	CPI CPI (Sep)	MNQ, MGC	MNQ 205-272pt in the 30-minute window (95th-100th percentile against a 69pt median); MGC 3-6x
Tue 20 Oct 2026	22:05 Paris 20:05Z · 16:05 ET	EARNINGS NFLX earnings	MNQ	NOT MEASURED on our tape. Carried because it moves the index, not because we have tested it — treat as unquantified
Wed 28 Oct 2026	19:00 Paris 18:00Z · 14:00 ET	FOMC FOMC decision	MNQ, MGC	MNQ ran 2.7-3.5x a normal 30-minute window on both FOMC days in our tape (98th/100th percentile); MGC 6.6x its median and 2.8x its previous MAXIMUM
Wed 28 Oct 2026	21:05 Paris 20:05Z · 16:05 ET	EARNINGS TSLA earnings (time UNCONFIRMED —	MNQ	NOT MEASURED on our tape. Carried because it moves the index, not

DATE	TIME	EVENT	INSTRUMENTS	WHAT IT DID TO OUR TAPE
		assumed after close) TIME UNCONFIRMED		because we have tested it — treat as unquantified
Thu 29 Oct 2026	21:05 Paris 20:05Z · 16:05 ET	EARNINGS AAPL earnings (time UNCONFIRMED — assumed after close) TIME UNCONFIRMED	MNQ	NOT MEASURED on our tape. Carried because it moves the index, not because we have tested it — treat as unquantified
Thu 29 Oct 2026	21:05 Paris 20:05Z · 16:05 ET	EARNINGS AMZN earnings (time UNCONFIRMED — assumed after close) TIME UNCONFIRMED	MNQ	NOT MEASURED on our tape. Carried because it moves the index, not because we have tested it — treat as unquantified
Wed 4 Nov 2026	22:05 Paris 21:05Z · 16:05 ET	EARNINGS GOOG earnings (time UNCONFIRMED — assumed after close) TIME UNCONFIRMED	MNQ	NOT MEASURED on our tape. Carried because it moves the index, not because we have tested it — treat as unquantified
Wed 4 Nov 2026	22:05 Paris 21:05Z · 16:05 ET	EARNINGS GOOGL earnings (time UNCONFIRMED — assumed after close) TIME UNCONFIRMED	MNQ	NOT MEASURED on our tape. Carried because it moves the index, not because we have tested it — treat as unquantified
Wed 4 Nov 2026	22:05 Paris 21:05Z · 16:05 ET	EARNINGS MSFT earnings (time UNCONFIRMED — assumed after close) TIME UNCONFIRMED	MNQ	NOT MEASURED on our tape. Carried because it moves the index, not because we have tested it — treat as unquantified
Wed 4 Nov 2026	22:05 Paris 21:05Z · 16:05 ET	EARNINGS META earnings (time UNCONFIRMED — assumed after close) TIME UNCONFIRMED	MNQ	NOT MEASURED on our tape. Carried because it moves the index, not because we have tested it — treat as unquantified
Fri 6 Nov 2026	14:30 Paris 13:30Z · 08:30 ET	NFP Employment Situation (Oct)	MNQ, MGC	MNQ 218-220pt in the 30-minute window (100th percentile both times); MGC 3-6×
Tue 10 Nov 2026	14:30 Paris 13:30Z · 08:30 ET	CPI CPI (Oct)	MNQ, MGC	MNQ 205-272pt in the 30-minute window (95th-100th percentile against a 69pt median); MGC 3-6×

The macro backdrop — and what it is NOT for

★ **READ THIS BEFORE THE TABLE.** These four were measured against our own instruments on **2026-09-18** and the result is a warning label, not a signal: **all four CO-MOVE with the tape and NONE of them predicts the next day.** Same-day, VIX against MNQ is $r = -0.669$ and DXY against MGC is -0.408 ; oil is unrelated to either. Next-day: nothing. **So this table EXPLAINS the week you just had. It does not forecast the one coming,** and no play on the card above is sourced from it. `scripts/macro_watch.py` says the same thing in its own docstring — “IT MEASURES NO EDGE”.

⚠ **And the window is one day, not the week.** `data/macro_watch.jsonl` was started on the morning of 2026-09-18 and holds **175 hourly reads across 8 day(s)** — 07:22Z to 21:52Z. Below is Friday. It is not a week of context and it must not be read as one; next week’s report gets the first full five days. The two same-day correlations quoted above were measured separately on 2026-09-18 and are carried here as a stated result — **this build does not re-derive them, and thirteen hourly rows could not.**

SERIES	LAST	VS PREV CLOSE	RANGE ON THE DAY
Oil CLX6 — front WTI future	94.61	+0.00%	88.79 – 97.88 (2 polls returned no quote)
The dollar DXZ6 — dollar index future	100.77	-0.25%	99.92 – 101.09
The 10-year ZNZ6 — note future — ⚠ PRICE, so DOWN = yields UP	104.86	+0.00%	104.59 – 106.30 (1 poll returned no quote)
VIX VIX — ⚠ DELAYED ~15 min on this account	14.87	-5.11%	14.15 – 16.48

★★ **IF YOU READ ONLY THIS — FOUR THINGS, IN ORDER**

1 — ★ THE FINDING THAT SURVIVED THE WEEK: this book is five draws, in either direction. Last week that was written as a caveat on a win. It is not a caveat, it is the shape of the book, and this week proves it by inverting the sign with the concentration intact. Same ruler, both weeks, entries not rows:

WEEK	ENTRIES	TOP 5	THE OTHER N	WEEK NET
14-18 Sep	39	\$3,162.48	-\$245.99 over 34	\$2,916.49
21-25 Sep	99	\$3,623.00	-\$3,983.00 over 94	-\$360.00

In both weeks **five entries out of 39 and 99 are larger than the entire week**. That is a far stronger statement than “last week’s upside was concentrated”, and it is the only headline claim from the 09-18 edition that survives being handed this week’s data. **Two weeks is still two weeks** — but it tells you where to look: the tails are the book, and the middle is noise around zero.

2 — You had a losing week, and it was NOT you getting worse. — \$360.00 across 99 entries and 396 contracts, against \$2,916.49 over 39 last week. But the per-entry economics barely moved — what moved was the mix:

ENTRIES THAT...	N LAST WK	\$/ENTRY LAST WK	N THIS WK	\$/ENTRY THIS WK
reached a profit rung (TARGET_*)	14	\$455.57	20	\$487.82
reached no rung at all	25	-\$138.46	79	-\$128.06

Both buckets pay almost exactly what they paid last week. What collapsed is how often you reached a rung at all: **14/39 = 36%** last week against **20/99 = 20%** this week. You traded **2.5× as much** (7.8 entries a day → 19.8) and the extra entries were the ones that never got anywhere. Δ That drop is **P = 0.054** on a two-proportion test — right on the line, on two weeks, so treat it as the best available reading and not as established. **The honest baseline is \$/entry at 4 lots, never \$/day**, because the days are not comparable once the entry count doubles.

3 — The automated desk did NOT take nothing this week — and it was the only part of the book that made money. The gates took **15 signals** (30 legs, one lot each) for **\$257.00**, while your hand lost **-\$360.00** on 396 contracts. **Do not read that as the machine beating you** — normalised to a single stated unit it is **\$8.57 per contract** for the gates against **-\$0.91 per contract** for you, and then two disqualifiers land on it: the gates are **one lot and paper**, and **8 of 15 signals are abs_veto_long** — strip that one family and the rest of the automated book is **-\$47.00 over 7 signals**. It fails strip-the-best-family, so the verdict is **PROMISING** on a 15-signal week and **nothing here justifies moving size from one book to the other**.

★ **PASSES RECORDED: 3** — and the answer to MONDAY #1 is **yes, the habit has started**. First 2026-09-24, latest 2026-09-25, across **2 of the week's 5 sessions. 3 of 30. A 4th pass row is NOT counted**: the 09-24 09:56Z record carries BUTTON AUDIT - not a real pass in its own request line — a wiring check, not a decision. It is excluded here and named so the count cannot drift. At **3.0 a week** over the 7 days the button has existed, **30 lands around 2026-11-27** — about 9 more weeks. △ **Read that date as one observation, not a trend**: all 3 landed in 2 sessions at the very end of the week, so the rate rests on 2 days. Generated from data/operator_reads.jsonl at render time.

4 — The 15-minute page band is PROVEN — in one direction only. MONDAY #2 of the last card shipped nothing, so its forward week was available and unspent. It is spent now. Replaying the unchanged detector out of data/leg_watch.log over all five sessions: at a 15-minute band a same-side page precedes **82.8% of your 99 entries** within the hour, against a session-blocked shuffled rate of **63.4%** — a lift of **+19.4 points**, with **all five leave-one-day-out folds positive** (+13.1 to +23.4) and 15/10 forming a plateau rather than a step. The in-sample claim was 82% on 34 presses; the forward number on **99 entries** is 82.8%. **That replicates.** △ **And in the other direction it is at chance**: the same band fires **42.6 pages a session** and only 28.2% of them are followed by a trade against a 29.5% shuffled rate — **a lift of -1.3**. So it is a good description of where you trade and **wallpaper as a phone alert**. Dashboard only. Pre-registered tonight in data/prereg_leg_band.json, with the forward leg (next 30 entries) unrun and named.

★ THE HONEST HEADLINE — the biggest week this book has ever traded, and it gave back everything it had made

Your own hand booked **−\$360.00** across **99 entries** and **396 contracts** — **−\$3.64 an entry**, on a **53% win rate**. **3 sessions green, 2 red.** △ **Win rate is not the story and is not a kill criterion on this desk** — you won more than half your entries and still lost money, which is the mirror image of the 39% week that made **+\$625**.

△ **The week's two red sessions, together, wiped out the whole of last week.** **Thu 24 Sep −\$1,533.00; Tue 22 Sep −\$1,396.50;** that is **−\$2,929.50** across two days against last week's **\$2,916.49** across all five — they gave back last week's entire profit and \$13.01 more. △ Neither session on its own is larger than last week's gain; it takes both. And **one entry cost −\$1,838.50** (SHORT 4 lots at 30,765.94, opened 12:46Z on Tuesday, closed by hand) — worse than the whole week's net, and worse than the **−\$1,288.00 CLOCK_FLAT** that the 09-18 report called the single worst decision of its week.

The book, life to date. Your hand is now **191 entries and −\$8.44** since 2026-08-10. It stood at **+\$351.56 over 92 entries** into this week: **this week handed back the book's entire lifetime profit** and left it flat. That is the cleanest one-line summary of the five days, and it is the number to hold against any per-entry statistic in this report.

Three things to hold while reading the rest. One — rows are not entries. The clean week is **153 rows** in trades across both books and only **99 entries by hand**, because a scale-out books one row per lot leaving. Two — two fills this week were not real, and they are not the same KIND of unreal. See the note below the table; this matters more than it did last week. Three — the two books are different sizes. Your hand trades 4 lots, the gates 1. Never net them into one \$/day.

★★★ DATA_QUALITY IS NOW A MIXED COLUMN — TWO UNRELATED MEANINGS, AND THE NEW ONE IS THE DANGEROUS ONE

The 09-18 report treated BADFILL as one known mechanism: the paper engine inventing a price on 3 of 4 lots. **A second and completely different class appeared on 09-24** and nothing in the previous edition anticipates it:

TAG	ID	LOTS	\$	WHAT IS ACTUALLY WRONG
BADFILL:paper_engine_0 . 1pct_fabrication_3of4_ lots_20260915		895	4	−\$304.50
BADFILL:assistant_test _claim_NOT_operator_de cision_20260924		1019	4	−\$842.50

Why the second is worse. A bad price is visible — it disagrees with the tape. A trade that was never your decision looks perfectly normal, and it sits in the same column the first one uses. **Anything that reports your \$/day or your win rate on 09-24 without the filter attributes an assistant’s test claim to your judgement** — which is the exact defect the tag was written to prevent. It is −\$842.50 over 4 lots and it is excluded from every total in this front matter: **−\$360.00 is the clean figure, −\$1,202.50 is what an unfiltered reader sees.**

△ **And the filter is data_quality IS NULL, which is correct** — checked tonight, the table holds **1,086 NULL, 9 tagged and ZERO empty-string** rows, so there is no empty-string trap here. The trap is elsewhere: **the live \$250 loss limit has no data_quality filter at all** and counts this −\$842.50 straight into its own trigger, exactly as it counted the −\$304.50 on 09-15. That is SATURDAY #4 from the last card and it is **still not done** — graded below.

There is no single honest total, so here are all of them, with what separates each from the next.

WHAT YOU ARE LOOKING AT	ROWS	NET \$	WHY IT DIFFERS FROM THE ROW ABOVE
Raw trades rows, 2026-09-21 → 2026-09-25	154	−\$945.50	everything the book recorded, both books, including the fabricated one

WHAT YOU ARE LOOKING AT	ROWS	NET \$	WHY IT DIFFERS FROM THE ROW ABOVE
... of which: the assistant_test_claim row of 24 September	1	−\$842.50	a REAL fill that was never your decision. Dropped from every total below
Clean ledger (data_quality IS NULL) — the number this report means	153	−\$103.00	the week's money across BOTH books, 426 contracts
... opened by hand (entry_source = 'manual') — 99 entries, 4 lots each	123	−\$360.00	your book. −\$0.91 per contract
... opened by an automated gate — 15 signals, 1 lot each	30	\$257.00	the gates DID trade this week. \$8.57 per contract — but paper, one lot, and 8/15 from one family

⚠ **The classifier trap, for whoever writes the next query.** Those automated rows carry entry_source IS NULL — so IS NULL happens to find them this week. **Do not reuse it.** entry_source only began being written on 2026-09-10, so **88 of his OWN day_rider legs from 08-10 to 09-10 also carry NULL.** Any window reaching back past 09-10 must split on gate NOT LIKE 'day_rider%', which is what the two functions behind this table now do.

Your week, day by day — entries, not legs

SESSION	ENTRIES	CONTRACTS	NET \$	WON	\$/ENTRY
Mon 21 Sep	8	32	\$1,190.00	5/8	\$148.75
Tue 22 Sep	8	32	−\$1,396.50	4/8	−\$174.56
Wed 23 Sep	14	56	\$404.50	6/14	\$28.89
Thu 24 Sep	34	136	−\$1,533.00	17/34	−\$45.09
Fri 25 Sep	35	140	\$975.00	20/35	\$27.86

⚠ **Entries, by the day the position was OPENED.** The clean-ledger table above counts **legs** by the day they CLOSED and reads 153 across both books — both are right about different things and the gap is the scale-out plus the gates. Note the shape of the week: the two biggest sessions by entry count (35 and 34 entries) are the last two, and they are also where the PASS button was finally used.

The automated book, collapsed to independent signals

GATE FAMILY	SIGNALS	NET \$	\$/SIGNAL
abs_veto_long	8	\$304.00	\$38.00
abs_veto_short	1	\$123.00	\$123.00
capitulation_long	3	\$70.50	\$23.50
grind_long	3	−\$240.50	−\$80.17

△ **30 legs are NOT 30 observations.** Every gate runs as _A and _B and those are the two halves of one scale-out, so the week's 30 legs collapse to **15 signals** — and they pair exactly: 15 A+B pairs and 0 left over, which is the check that the key is right. **grind_long is the only loser** and it lost on every signal it took. Strip abs_veto_long and the remainder is −\$47.00.

How the money left the book — both books, legs

EXIT	LEGS	NET \$	PER LEG
TARGET_100	20	\$2,194.12	\$109.71
TARGET	13	\$638.00	\$49.08
TARGET_200	1	\$200.75	\$200.75
STOP	14	−\$527.00	−\$37.64
MANUAL_CLAIM	105	−\$2,608.87	−\$24.85

MANUAL_CLAIM is your own hand on the button and it is where the money went: **105 legs for −\$2,608.87.** TARGET_100/TARGET_200 are the per-lot profit ladder and every one of them is green. TARGET and STOP belong to the gates, not to you — the gates are the only book on the desk with a stop, and 14 stops cost them −\$527.00 while their targets made +\$838.75. **CLOCK_FLAT does not appear at all this week:** the 20:40Z timer never had to close a position, because you closed them all yourself.

★ WHAT TO ACTUALLY DO

Garrath — this is the whole answer to “so what do I actually DO?”, and it is at the front. Every row below mirrors reports/friday_v7/plays.json, which is the reconciled source of truth for this week; the same file drives the [Monday playbook](#), so the two cannot drift apart. ★ **Almost nothing on this card changes how you trade, and that is the honest answer to this week.** You had the best week this book has ever had — and a week is 39 entries, five of which ARE the week. The automated desk took **nothing**, so there is no gate to re-tune and no threshold that this week has earned the right to move. What the card is instead: **one habit, one config change, six instrument faults, six measurements and four**

refutations. 2 of the rows are NOT-AN-ACTION: refutations printed with their autopsies, because on a week like this the most valuable output is a list of things that have now been ruled out. ★ The first of them is the idea you are most likely to want after a week like this — only trade the European morning — and it dies on its own sample: +\$2,445.51 over 20 entries this week, -\$779.48 over the record, one trade from changing sign. 14 plays, five windows, in the order they need doing. Every play carries its **evidence tier** and its **kill criterion** — if a play has no way to be proven wrong it should not be on the card.

THE THREE THAT MATTER MOST

1 · ★★ drift is still blind before 13:30Z — and this week it was blind on 42 of your 104 presses (SATURDAY #1). One constant. `src/gazbot7/drift.py` line 50: `OPEN_UTC_MIN = 13*60+30` anchors the detector to the US cash open, so every call before 13:30Z returns before the 13:30 UTC open and no read at all. Anchor it to the **CME session (22:00Z reopen)**, which is the session your book actually trades. **42 of 104 presses (40.4%) blind.** 62 returned a real read. Count it yourself: every blind press carries the same *drift_error* string.

2 · ★★ The \$250 loss limit will bench your whole gate book over a trade that was never your decision (SATURDAY #2). Add and `data_quality is null` to `session_pnl()` in `scripts/daily_loss_limit.py` (line 49). One clause. **-\$842.50 on one row, against a -\$250 trigger.** That single excluded row is 3.4× the trigger on its own.

3 · ★★ KEEP PRESSING PASS — it started this week, and 3 of 30 is still blocking everything (MONDAY #1). Every time you look at the tape and do not trade, press **PASS**. You started on Thursday. Keep going — 27 more. **3 of 30**, first 2026-09-24, in **2 of 5 sessions**. At the observed 3.0/week over the 7 days the button has existed, 30 lands around **late November** — △ **read that as one observation, not a trend:** two sessions produced all three.

Generated from the top of the card below, so the two cannot drift apart. Everything else, including every refuted idea and why it died, is in the five windows underneath.

MONDAY · MONDAY 2026-09-28, AT THE DESK — ★ ONE HABIT TO KEEP, ONE CONFIG TO SHIP, ONE THING TO LEAVE ALONE · 3 PLAYS

★★ THIS CARD IS MOSTLY LAST WEEK’S CARD, AND THAT IS THE FINDING. Six of the nine rows on the 09-18 card were **not done**, including all three it marked ★ as its most serious defects — and they have now survived a week in which you took **99 entries and lost \$360**. A defect that lives through a week like that outranks any new measurement, so the SATURDAY window below is ranked by what each one demonstrably cost rather than by how interesting it is. ★ **MONDAY #1 is still ranked first and still has no dollar figure**. You pressed PASS for the first time on Thursday: **3 real declines against 104 presses**. It is working and it is still blocking — 27 more and this desk can, for the first time, learn what separates a moment you take from a moment you leave. ★ **MONDAY #2 is the one row on this card that PASSED a robustness test**, out-of-sample, on 2.5× the sample it was fitted on.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
1	★★ KEEP PRESSING PASS — it started this week, and 3 of 30 is still blocking everything BLOCKER Every time you look at the tape and do not trade, press PASS. You started on Thursday. Keep going — 27 more. Mechanism: data/operator_reads.jsonl now holds 104 buys, 109 claims and 4 pass rows for 09-21&rarr;09-25. △ One of those four is not a pass: the 09-24 09:56:03Z row says BUTTON AUDIT - not a real pass in its own request line. Three real declines. Why: MONDAY #1 of the last card had a kill criterion — “any number above zero next Friday is this row working” — and tonight IS next Friday and the answer is three. The row is working but not yet unblocking. The fitting threshold is 30 labelled examples and it was written for positives AND negatives; at 3 negatives against 104 presses nothing can separate, and this report again hands you no entry rule. Evidence: ★ PASSES RECORDED line on the front page; ↗ LAST WEEK’S CARD, MONDAY 1. · verification: <code>python3 -c "import json;print(sum(1 for l in open('data/operator_reads.jsonl') if json.loads(l)['kind']=='pass'))"</code> — it reads 4 tonight, 3 of which are real. The front page generates this line every week. · mode: do it	3 of 30, first 2026-09-24, in 2 of 5 sessions. At the observed 3.0/ week over the 7 days the button has existed, 30 lands around late November — △ read that as one observation, not a trend: two sessions produced all three.	Arms: Nothing arms. It is a habit and it is still the cheapest item on this card. Exit: n/a. Kill: If 30 passes land and the press/pass boundary STILL does not separate on any feature, that is a real answer and this desk stops trying to model entry selection.	Garrath — at the desk, every session n/a. ★ Still nothing else in this report is unblocked without it.
2	Ship the 15-minute leg band — to the DASHBOARD, and never to the phone PROVEN Set LEG_LOUD to include 15 and route the 15-minute line to the dashboard only. Do not page it. The pre-registration is already written: data/prereg_leg_band.json. Mechanism: scripts/leg_watch.py still reads LEG_LOUD="60,120,180". The band never shipped, so its out-of-sample week was available and unspent — it is spent now. Why: ★ This is the one candidate in the report that PASSED a real robustness test. Replaying the unchanged detector out of data/leg_watch.log over all five sessions: a same-side 15-minute page precedes 82.8% of your 99 entries within the hour against a session-blocked shuffled rate of 63.4% — +19.4 points. The in-sample claim was 82% on 34 presses; the forward number on 99 entries, 2.5× the fitting sample, is 82.8%. All five leave-one-day-out folds stay positive (+13.1 to +23.4), and 15/10 form a plateau (82.8/83.8) rather than a step, while 30 falls to 50.5%. △ And it is dead as a page: at 15 minutes it fires 42.6 times a session and only 28.2% of those pages are followed by a trade against a 29.5% shuffled rate — lift -1.3, i.e. exactly chance. Coverage and paging	82.8% vs 63.4% shuffled (+19.4) on 99 entries / 222 legs / 5 sessions. Page direction: 28.2% vs 29.5% (-1.3) on 213 pages.	Arms: On the next restart of gazbot7-leg-watch. Exit: n/a — it is an alert, not a position. Kill: Pre-registered: coverage below 70%, or lift at or under +5 points, over 30 forward entries. △ The page-side kill has ALREADY triggered, which is why this ships to a dashboard and not to your phone.	Garrath — one env var and one destination Restore LEG_LOUD="60,120,180".

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	are opposite directions of the same join and only one of them replicates. Evidence: ★ IF YOU READ ONLY THIS #4; data/prereg_leg_band.json; ↩ LAST WEEK'S CARD, MONDAY 2. · verification: Re-run the replay in data/prereg_leg_band.json against data/Leg_watch.log and the trade record. The forward leg — the next 30 entries after 2026-09-25 — is pre-registered and unrun. · mode: do it			
3	Keep standing aside on event days — and the next one is NFP, Friday 2026-10-02 MEASURED Unchanged: stand aside as a PAGE, never as a bench. Nothing is switched off. Mechanism: data/event_calendar.json, regenerated tonight at 21:27Z, 55 events, human-verified. Why: MONDAY #3 of the last card was done — the calendar is live and regenerating. It was also untested this week, because the calendar reports no scheduled event at all in 09-21&rrarr;09-25. So the row carries forward unchanged and un-revalidated, and it matters again in seven days. Evidence: ♦ THE WEEK THAT WAS, AND THE WEEK COMING. · verification: The calendar is regenerated and human-verified; the event-day split is re-derivable from trades against it. · mode: do it	0 events this week. 1 event on 2026-10-02 (NFP). The evidence behind the stand-aside is still the same four event days it was in September — it did not grow.	Arms: Nothing arms. A page, not a bench. Exit: n/a. Kill: If the event-day split stops separating once the sample reaches ~10 event days, drop it. It is at four and it has not moved.	Garrath n/a.

SATURDAY · SAT 2026-09-26 / SUN 2026-09-27 — BEFORE THE SUNDAY 22:00Z REOPEN. ★ SIX ROWS, AND FIVE OF THEM ARE LAST WEEK'S ROWS AGAIN · 6 PLAYS

★ **Nothing in this window has a clock on it** — check the carry box at the top of the report for the desk's state, which is read from the state files at render time rather than asserted. **What IS here is the same six instrument faults as last week, re-ranked by evidence.** Two of them stopped being theoretical this week. drift was **blind on 42 of your 104 presses**, printing no 5s rows since 13:30Z today into your own tape reader on 40% of your decisions. And the \$250 loss limit — which has no data_quality filter — will now count a **−\$842.50 trade that was never your decision** towards a −\$250 trigger that switches off the whole gate book. △ **SATURDAY 6 is graded UNGRADED for an unusual reason:** last week's verification command named a file that does not exist, so the check was never runnable. The path is corrected on this card.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
1	★★ drift is still blind before 13:30Z — and this week it was blind on 42 of your 104 presses FIX — DEFECT One constant. src/gazbot7/drift.py line 50: OPEN.UTC.MIN = 13*60+30 anchors the detector to the US cash open, so every call before 13:30Z returns before the 13:30 UTC open and no read at all. Anchor it to the CME session (22:00Z reopen), which is the session your book actually trades. Mechanism: This was SATURDAY #1 on the 09-18 card, marked ★ and called “the biggest single instrument fault in the report”. It was not done. Last commit to drift.py is 2026-08-05. Why: It stopped being latent this week. Of the 104 presses recorded in operator_reads.jsonl between 09-21 and 09-25, 42 carry drift_error: no 5s rows since 13:30Z today (venue shut?) — the direction meter was dark on 40% of your own decisions, and the tape reader printed that error to you in the moment. A defect that survived a 99-entry losing week outranks any new measurement, which is why it is rank 1 rather than rank 4. Evidence: ↩ LAST WEEK'S CARD, SATURDAY 1 (graded NOT DONE tonight). · verification: grep -n OPEN.UTC.MIN src/gazbot7/drift.py — it reads 13*60+30 today. After the fix, call drift.read() at 08:00Z Monday and it returns a real read. · mode: do it	42 of 104 presses (40.4%) blind. 62 returned a real read. Count it yourself: every blind press carries the same drift_error string.	Arms: Nothing arms. It is a constant. Exit: n/a. Kill: If the CME-anchored read measures no better than the 13:30-anchored one on the 62 presses where BOTH are available, the anchor is not the problem and this row is wrong. That is a real test and it has not been run — the fix is justified by the 42 blind presses, not by a claim that the read is good.	Garrath — one constant, one commit Restore the constant. No state, no migration.
2			Arms: It is already armed — the limit runs live.	Garrath — one clause

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	★★ The \$250 loss limit will bench your whole gate book over a trade that was never your decision FIX — DEFECT Add and <code>data_quality</code> is null to <code>session_pnl()</code> in <code>scripts/daily_loss_limit.py</code> (line 49). One clause. Mechanism: <code>session_pnl()</code> sums where <code>closed_at</code> is not null and <code>closed_at >= ?</code> with no quality filter, so every tagged row counts towards a <code>−\$250</code> trigger that writes off to every gate. Why: This was SATURDAY #4 and it was not done. It has since acquired a worse input. On 09-15 it counted a fabricated PRICE (<code>−\$304.50</code>). On 09-24 the book gained id 1019, <code>−\$842.50</code> over 4 lots, tagged BADFILL:assistant_test_claim_NOT_operator_decision_20260924 — the fill was real and only the authorship was not. An unfiltered reader attributes it to your judgement AND counts it into a guard that switches the desk off. Evidence: ★ THE HONEST HEADLINE — the mixed-column box; ↗ LAST WEEK'S CARD, SATURDAY 4. · verification: <code>sqlite3 data/gazbot7.db "select id,pnl_usd,data_quality from trades where data_quality is not null and closed_at>='2026-09-21'", then re-run session_pnl() for 2026-09-24 with and without the clause.</code> · mode: do it	−\$842.50 on one row, against a <code>−\$250</code> trigger. That single excluded row is 3.4× the trigger on its own.	Exit: n/a. Kill: None. This is a filter that every other reader on the desk already applies; there is no reading under which the guard should price a trade nobody chose to make.	Remove the clause.
3	★ The \$250 loss limit still pages you a sentence its own journal contradicts FIX — DEFECT <code>scripts/daily_loss_limit.py</code> line ~98 computes <code>changed = bench_all(...)</code> and then pages the literal string “All six gates benched”. Page changed instead — it is already in scope on the line above. Mechanism: <code>bench_all()</code> returns the gates it actually switched. The notify call ignores the return value and asserts six. Why: SATURDAY #3, not done, last commit 2026-09-13. A guard that tells you something different from what it wrote to its own journal is worse than one that says nothing, because you will act on the page and not on the journal. Evidence: ↗ LAST WEEK'S CARD, SATURDAY 3. · verification: <code>grep -n 'All six gates' scripts/daily_loss_limit.py, then compare against the benched key in data/daily_loss_limit_state.json.</code> · mode: do it	One line. changed is computed and discarded.	Arms: Already live. Exit: n/a. Kill: None — it is a string that disagrees with a variable three lines above it.	Garrath — one line Trivial.
4	★ A TARGET_100 rung booked a fill 79.5 points BELOW its entry — and it is still counted as clean DEFECT Decide what trade id 868 is, then tag it or fix the label. Do not leave it as a clean row. Mechanism: <code>LONG 1 @ 29,320.25 &rarr; 29,240.75, <code>−79.5pt, <code>−\$160.50, exit_reason = TARGET_100</code>. A \$100/lot profit rung cannot book a loss of that size — either the rung fired at the wrong price or the reason is mislabelled.</code></code> Why: SATURDAY #6, not done. <code>data_quality</code> IS NULL on it today, so it is inside every total this desk prints, including the ones in this report. This is cheap to close and it is corroding the one column the desk uses to say what is real. Evidence: ↗ LAST WEEK'S CARD, SATURDAY 6. · verification: <code>sqlite3 data/gazbot7.db "select * from trades where id=868", then the rider journal at 2026-09-11T12:30:06Z.</code> · mode: do it	−\$160.50 on one leg. Small money, and that is not the point: it is a clean row that is provably not clean.	Arms: n/a. Exit: n/a. Kill: If the rider journal shows a legitimate TARGET_100 at that price, the label is right and the PRICE path is the bug — a different row, but still a row.	Garrath — one query and one decision n/a.
5	event_alert.py still has no dry-run, and NFP is next Friday FIX — DEFECT Add <code>--dry-run</code> to <code>scripts/event_alert.py</code> so a hand-run cannot write a live PAGED line into the record. Mechanism: The script has no <code>argparse</code> at all and calls <code>log(f"PAGED ...")</code> unconditionally. Why: SATURDAY #5, not done, last commit 2026-09-18 — the day the card was written. It moves up the card this week for a calendar reason rather than a new measurement: NFP falls on Friday 2026-10-02, the first scheduled event since FOMC on 09-16, so the event path gets exercised for real in seven days. Evidence: ♦ THE WEEK THAT WAS, AND THE WEEK COMING; ↗ LAST WEEK'S CARD, SATURDAY 5. · verification: <code>grep -n 'dry.run\' argparse' scripts/event_alert.py — no hits today.</code> · mode: do it	0 events in 09-21&rarr;09-25 (calendar regenerated 21:27Z tonight, 55 events). 1 event in the week after next: NFP, Fri 2026-10-02.	Arms: n/a. Exit: n/a. Kill: None.	Garrath Trivial.
6	tape_reader's “day open/high/low” splices yesterday afternoon onto this morning — and last week's check named the wrong file	Unmeasured. No figure is offered, deliberately.	Arms: n/a. Exit: n/a. Kill: If the first bar at 07:00Z Monday already reads 22:00Z	Garrath n/a.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	FIX — INSTRUMENT Anchor scripts/tape_reader.py's session window to the 22:00Z CME reopen. △ The path is scripts/tape_reader.py, not src/gazbot7/tape_reader.py as the 09-18 card said. Mechanism: Same species as SATURDAY #1: a session boundary that is not the session your book trades. Why: This row is UNGRADED, and the reason is itself the finding: the 09-18 card's verification command names a file that does not exist, so the check could not be run as written — which means it was not run when the row was written. The underlying fault is still plausible and still unmeasured. It ranks last in this window because, unlike the five above it, nothing this week shows it costing anything. Evidence: ↩ LAST WEEK'S CARD, SATURDAY 2. · verification: grep -n 'day_open\' session' scripts/tape_reader.py, then print the session bar count and first bar timestamp at 07:00Z Monday — the first bar should be 22:00Z Sunday, not Friday afternoon. · mode: check it first		Sunday, there is no fault and the row closes.	

BUILD · THE MEASUREMENT QUEUE — TWO ROWS, AND ONE OF THEM WAS BUILT TONIGHT · 2 PLAYS

This window is short this week, deliberately. The 09-18 card carried seven BUILD rows, all of them derived from the four analytical sections — and those sections describe last week. Rather than restate seven measurements against a week they were not measured on, they are carried in one honest row (HOLD #1) and this window holds only what this cycle can actually stand behind. **BUILD #1 is the press/pass comparison**, wired end to end and published as **UNGRADEABLE** with its n printed — 3 against 104 — so that the week the count reaches 30 the answer is one query and not one project. **BUILD #2 is the report's own machinery**, and it is on the card because it nearly cost you this entire document.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
1	The press/pass comparison, wired end to end tonight so that the week 30 lands the answer is one query MEASURED Nothing to build — it is built. When the pass count reaches 30, re-run the comparison below. Mechanism: The three real passes were pulled from operator_reads.jsonl and put beside the 104 presses of the same week on the features Section 2 measured. Why: △ Published as UNGRADEABLE with the n printed, rather than not run. 3 against 104 cannot separate anything and must not be shown as a filter or a threshold. What it establishes is that the pipeline works: the passes carry the same tape_reader context the presses do, so the query is written and waiting. Evidence: ★ PASSES RECORDED; the run-chart 'first two PASS presses' chart. · verification: The three pass rows are at 2026-09-24T15:10:57Z, 2026-09-24T16:53:34Z and 2026-09-25T18:49:41Z. The audit row at 09-24T09:56:03Z is excluded by its own tag. · mode: park it	The three passes sit INSIDE the press distribution on every feature measured. Median position-in-range 77% (pass) vs 73% (press); VWAP stretch +4.4 ATR vs +4.6; ATR 15.7pt vs 12.8pt; drift confirmed on 0/3 passes and 13/104 presses. The boundary needs 30 and has 3.	Arms: n/a. Exit: n/a. Kill: n/a — it is a measurement, and its result this cycle is 'not enough data', stated as such.	the desk n/a.
2	★★ DONE TONIGHT — the report nearly published last week twice, and both halves of the guard are now on disk FIX — DEFECT Already shipped. Recorded here so the next reader can see what changed and why. Mechanism: friday_phases._report_week() resolved the report week as “the most recent Friday whose 22:00Z cut has passed”; the cron starts at 21:05Z on Friday, 55 minutes before that cut, so on Friday night it always rolled back seven days. Measured tonight it returned 2026-09-18. The runner then skipped every phase because a git operation on 09-24 09:43 had made all four fragments newer than the 09-18T22:00Z staleness cutoff. Why: Three fixes, all verified: (1) on a Friday the report week is today — the only run instant that changes, and it makes the cron agree with its own Saturday resume instead of contradicting it. (2) every artifact now carries a week stamp written by the runner on success, and a stamp that disagrees is STALE regardless of mtime. (3) the build now	Verified with mtimes untouched: all four body phases flipped HAVE &rrarr; todo on the stamp alone. The publish guard refuses --slug 2026-09-18 with a named error.	Arms: Next Friday's cron run. Exit: n/a. Kill: If a future run legitimately needs to rebuild an older week, pass --slug explicitly and the week stamps will force a full rebuild rather than reuse — which is the correct behaviour, not a bug.	done 2026-09-25 Both fixes are small and local.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	refuses to publish at all if the newest clean trade is newer than the masthead week. Evidence: The red box at the top of the front page. · verification: <code>python3 -c "import sys;sys.path.insert(0,'scripts/friday');import friday_phases as f;print(f.WEEK)" &rrar;</code> <code>2026-09-25. python3 scripts/friday_v7_build.py --slug 2026-09-18 &rrar;</code> refuses. · mode: already done			

HOLD · HOLD — THE ELEVEN ROWS THAT ARE CARRIED BUT NOT RE-DERIVED, NAMED RATHER THAN QUIETLY DROPPED · 1 PLAY

One row, and it is an admission. Eleven rows from the 09-18 card — the European-morning filter, the position-in-range and VWAP nulls, the best-moment split, the stopwatch, the 78 captured narratives — are all conclusions about **39 entries over five different days**. A 99-entry week that could have confirmed or destroyed several of them arrived, and none of them was put in front of it, because the four sections that own them could not be re-run tonight. **They are carried and labelled, never restated as if they were this week's**. The rule going in: a row that survives two cycles without being re-derived is not a finding, it is a habit.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
1	★ Eleven rows from the 09-18 card are CARRIED but NOT re-derived on this week's data — read them as last week's CARRIED Do not act on them this week. They are on disk at <code>reports/friday_v7/plays.json.pre-0925</code>. Mechanism: The 09-18 card carried 7 BUILD, 2 HOLD and 4 NOT-AN-ACTION rows, all derived from the 09-14&rrar;09-18 analysis in Sections 1–4. Why: Those four analytical sections could not be re-run tonight — each is a 45–60 minute job and the week-resolution fault was not found until the tail. Their conclusions are about 39 entries over five different days. Several of them (the European-morning filter, the position-in-range and VWAP nulls, the best-moment split) are exactly the kind of finding a 99-entry week could confirm or destroy, and none of them has been put in front of it. Evidence: ↶ LAST WEEK'S CARD; the red box at the top of the front page. · verification: <code>python3 -c "import json;print(len(json.load(open('reports/friday_v7/plays.json.pre-0925'))))"</code> · mode: park it	11 rows carried, 0 re-derived. The 09-18 card's own evidence base grew by 2.5× this week and was not spent.	Arms: n/a. Exit: n/a. Kill: They are re-derived on 09-21&rrar;09-25 next cycle, or they come off the card. A row that survives two cycles without being re-derived is not a finding, it is a habit.	next cycle n/a.

NOT-AN-ACTION · NOT-AN-ACTION — TWO READINGS OF THIS WEEK THAT ARE TEMPTING AND WRONG · 2 PLAYS

Both of these are the conclusions a losing week invites, and both fail on the data. The first is “the machine beat me” — the gates made +\$257 on one lot while your hand lost \$360 on four. Normalised it is +\$8.57 against −\$0.91 a contract, and then it **fails strip-the-best-family**: take out `abs_veto_long` and the rest of the automated book is −\$47.00 over 7 signals. The second is “I got worse” — strip the five best and five worst entries from each week and the middle of this week is still **+\$12.79 an entry and +\$1,138.00 in total**. The whole move is in the tails. **Neither row asks you to do anything, and that is the point.**

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
1	The gates made +\$257 while you lost \$360 — that is NOT the machine beating you, and nothing here moves size PROMISING Do nothing. Do not move size from your book to the gate book on this evidence. Mechanism: The automated gates took 15 independent signals (30 legs — every gate runs <code>_A/_B</code> as one scale-out, and all 30 pair	15 signals, +\$257.00, +\$17.13/signal. Strip the best family: −\$47.00 over 7. grind_long lost on 3 of 3 signals (6 of 6 legs).	Arms: n/a — they are already running on paper. Exit: n/a. Kill: Cause of death if it dies: strip-the-best-trades. It has already failed that on a 15-signal week. Revival condition: the same gates, live-sized, over a month,	the desk n/a.

#	WHAT TO DO	THE NUMBER BEHIND IT	ARMS WHEN · EXIT · KILL CRITERION	OWNER & REVERT
	exactly) for +\$257.00, one lot each. Your hand took 99 entries at 4 lots for -\$360.00. Why: Normalise before comparing, or it is the report's own "a \$ figure without its baseline" trap. Per contract it is +\$8.57 (gates) against -\$0.91 (you). Then two disqualifiers land: the gates are one lot and paper, and the book fails strip-the-best-family — <code>abs_veto_long</code> is 8 of the 15 signals and +\$304.00 of the +\$257.00, so removing it leaves -\$47.00 over 7 signals. Evidence: ★ IF YOU READ ONLY THIS #3; the automated-book table on the front page. · verification: <code>sqlite3 data/gazbot7.db "select gate,count(*),sum(pnl_usd) from trades where closed_at>='2026-09-21' and gate not like 'day_rider%' group by gate"</code> · mode: park it		still positive after stripping their best family.	
2	Your losing week was NOT a deterioration in the middle of the distribution — do not go looking for one NULL Do nothing. Resist the read that you got worse at trading this week. Mechanism: Strip the top 5 and bottom 5 entries from each week and compare what is left. Why: Last week the middle 29 entries made +\$135.35 each; this week the middle 89 made +\$12.79 each — lower, but still positive, and the middle of this week is +\$1,138.00 in total. The week is — \$360.00 because the tails are -\$1,498.00 net (top 5 +\$3,623.00 against bottom 5 -\$5,121.00). The whole move is in the tails. Evidence: ★ IF YOU READ ONLY THIS #1 and #2. · verification: <code>Group by (opened_at, entry_price, side), clean, manual; sort by pnl; strip 5 from each end.</code> · mode: park it	Middle-of-distribution \$/entry: + \$135.35 & rarr; + \$12.79. Tails: + \$3,162.48/-\$4,171.00 & rarr; + \$3,623.00/-\$5,121.00.	Arms: n/a. Exit: n/a. Kill: n/a — it is a null result, and the honest form of it is that two weeks cannot tell you whether the mix shift (36% of entries reaching a rung & rarr; 20%) is skill or tape. P = 0.054, right on the line.	the desk n/a.

STANDING CAVEATS ON EVERY DOLLAR ON THIS CARD

★ **The biggest caveat is not a number, it is a missing half.** There are **36 recorded presses and 0 recorded passes**. Every description in Sections 1 and 2 is fitted on one side of a decision boundary, and the only tests in this document that escape that are the ones run against control minutes — the leg, the range position, the VWAP stretch, the page band. Those are labelled where they occur. Everything else is a description of what you did, not a discovery about when to do it.

The fee is \$1.50 per round trip and the multipliers are not the same. MNQ is \$2.00 a point; **MGC is \$10.00**. Every dollar on this card is MNQ at \$2.00, because every trade in the report week was MNQ. The only gold in this document is the FOMC chart in Section 4, which carries no fills and is priced at \$10.00 in its own caption — pricing gold with the Nasdaq multiplier would understate it fivefold.

And every dollar here is a BOOKED dollar, not a repriced one. That is unusual for this report and it is worth saying: the money in Sections 1 and 4 is what the venue actually paid, straight out of trades, net of fees. The modelled numbers are confined to three rows — the eight-hour stopwatch,

the loss-limit replay and the page-band sweep — and each of those carries its own uncertainty line in-row.

7 method wounds are open and they bound everything above. this report was one build away from being last week's report, re-published tonight. `_report_week()` rolled the report week back seven days for any run starting before 22:00Z on a Friday — and the cron starts at 21:05Z — while the runner's mtime-based freshness check could not tell a rebuilt fragment from one a git operation had touched. Both halves are fixed and verified tonight, and the build now refuses to publish if the newest trade is newer than the masthead week. ★ **But the four long analytical sections could NOT be re-run in the time left**, so Sections 1–4 describe 2026-09-14→09-18 and carry a red banner saying so. Every number in the front matter, the card, the graded ledger and the run charts is this week's. **Do not read a dollar figure out of the body sections as this week's — BUILD #2. there are 3 recorded passes against 104 recorded presses**, so every behavioural finding on this desk is still fitted on one side of a decision boundary. Three is progress from zero and it is not a sample: measured against the same week's presses the three passes sit inside the press distribution on every feature (median 77% of day range against 73%, +4.4 ATR from VWAP against +4.6). **The boundary needs 30 and has 3**, and this report again hands you no entry rule — **MONDAY #1. this book is five draws in either direction, and two weeks now show it.** The best five entries made +\$3,623.00 against a week total of −\$360.00; the other 94 net −\$3,983.00. Last week the identical shape ran the other way: +\$3,162.48 top five, −\$245.99 for the other 34, on a +\$2,916.49 week. **Any per-entry statistic in this document describes a distribution whose mass sits in five draws**, and the middle of the distribution is positive in both weeks — **NOT-AN-ACTION #2. the entry count went 39 → 99 and the two weeks are not comparable on a \$/day basis.** 7.8 entries a day became 19.8. The per-entry economics of each bucket barely moved — entries that reached a profit rung paid +\$455.57 last week and +\$487.82 this week; entries that reached no rung cost −\$138.46 and −\$128.06 — but the share reaching a rung fell from **36% to 20%**, at **P = 0.054**, which is on the line and not established. **The baseline in this report is \$/entry at 4 lots, never \$/day — NOT-AN-ACTION #2. data_quality is now a MIXED column with two unrelated meanings**, and the new one is the more dangerous. A fabricated PRICE (the paper engine, 09-15, −\$304.50) disagrees with the

tape and is findable. A fabricated DECISION (09-24, id 1019, -\$842.50, assistant_test_claim_NOT_operator_decision) had a REAL fill and only the authorship was invented — it looks perfectly normal. Both live in the same column. Everything in this report filters on data_quality IS NULL, which was CHECKED tonight and is correct (1,086 NULL, 9 tagged, **zero empty-string** rows). **The live \$250 loss limit does not filter at all — SATURDAY #2. the automated book is 15 signals, not 30 legs.** Every gate runs as _A and _B and those are the two halves of one scale-out, so counting 30 legs as 30 observations double-counts the week. Collapsed properly it is 15 signals for +\$257.00, one lot each, on paper — and 8 of the 15 are abs_veto_long. **Strip that one family and the remainder is — \$47.00 — NOT-AN-ACTION #1. the 15-minute page band replicated in ONE direction and is at chance in the other,** and both numbers have to be quoted together. It precedes 82.8% of his entries against a session-blocked 63.4% (+19.4, 5/5 leave-one-day-out folds positive) — that is the direction the claim was made in and it held on 99 entries. But only 28.2% of its 213 pages are followed by a trade against a 29.5% shuffled rate, on 42.6 pages a session. **Coverage is not a page rate; co-fire is asymmetric and must be quoted in the direction it is used — MONDAY #2.**

What did not run this cycle, named exactly. 1 of 4 fragments are finished work and sit in this document in full. 3 fragment(s) predate this cycle entirely (Section 1, Section 2 and Section 3) and are stitched IN FULL under a red dated banner that says so at the top of the section. No play on this card is sourced from one. Nothing is missing — every section on the running order rendered from a real fragment, so there are no red placeholder boxes to hunt for. ★ No play on this card is sourced from any section carrying a stale banner, and the loop on LAST week's card is closed immediately below this table.



LAST WEEK’S CARD — graded, nine rows, every one from disk

The card is only worth writing if somebody checks it afterwards. **Nine rows were checkable and nine were checked: 2 done, 7 not done.** Every row below names the command that decided it, and every one of those commands came off the card itself — no memory, no inference.

★★★ THE HEADLINE FROM THIS TABLE, AND IT SETS THIS WEEK’S RANKING

All three rows the card marked with a ★ as its most serious defects are still open. drift is still blind before 13:30Z and was blind on **42 of his 104 presses this week**. The \$250 loss limit still pages a sentence its own journal contradicts, and still has no data_quality filter — so it now counts a **−\$842.50 trade that was never his decision** into a −\$250 bench trigger. **A defect that survived a 99-entry losing week outranks any new measurement**, so those rows are carried to the TOP of this week’s card rather than restated here.

△ **And the one row that was graded last week is the one row that shipped.** The habit rows moved; the code rows did not. That is a pattern worth naming: nothing on this card gets done unless it is either one tap at the desk or somebody edits a file, and six weeks of “SATURDAY” rows have now accumulated without a Saturday.

PLAY	WHAT IT ASKED FOR	STATUS	THE CHECK THAT DECIDED IT
MONDAY 1	★★ PRESS THE PASS BUTTON — the row the whole report was blocked on	WORKING — 3 of 30	operator_reads.jsonl holds 4 pass rows, first 2026-09-24. △ One is not a pass: the 09-24 09:56:03Z row says BUTTON AUDIT - not a real pass in its own request line. So 3 real declines, in 2 of 5 sessions. The card's kill criterion was “any number above zero next Friday”, and tonight IS next Friday. The habit has started; the boundary is still blocked — 3 negatives cannot separate anything against 104 presses.
MONDAY 2	Drop leg_watch's first page from 60 minutes to 15, and route it to the dashboard	NOT DONE — but now SCORED, and it was right	scripts/leg_watch.py still reads LEG_LOUD="60,120,180"; last commit 2026-09-15, and data/prereg_leg_band.json did not exist. The band never shipped

PLAY	WHAT IT ASKED FOR	STATUS	THE CHECK THAT DECIDED IT
			and its forward week was spent doing nothing. Scored tonight instead, on the unchanged detector over all 5 sessions: 82.8% of his 99 entries are preceded by a same-side 15-min page within the hour vs a 63.4% shuffled rate (+19.4, 5/5 LODO folds positive). The in-sample 82% on 34 presses replicated on 99. Δ As a PAGE it is dead: 42.6 pages/session, page→trade lift −1.3. Dashboard only.
MONDAY 3	Keep standing aside on event days — as a PAGE, never a bench	DONE — and untested this week	data/event_calendar.json was regenerated tonight (generated_at 2026-09-25T21:27Z, 55 events) and it is human-verified. It reports NO scheduled event in 09-21→09-25, so the stand-aside had nothing to stand aside from and the row is untested, not vindicated. Δ It bites next week: NFP, Friday 2026-10-02 is the next event on the calendar.
SATURDAY 1	★ drift is blind before 13:30Z — and that is two thirds of his trading	NOT DONE — and it bit 42 times this week	src/gazbot7/drift.py still carries OPEN_UTC_MIN = 13*60+30 and still returns before the 13:30 UTC open; last commit 2026-08-05. This is no longer a latent fault: of his 104 recorded presses this week, 42 carry drift_error: no 5s rows since 13:30Z today (venue shut?) — the direction meter was blind on 40% of his own decisions, including the 09:56Z button audit. The card called this “the biggest single instrument fault in the report” and it went through a 99-entry week untouched.
SATURDAY 2	tape_reader's day open/high/low splices yesterday afternoon onto this morning	UNGRADED — the row names a file that is not there	The card's verification says to print tape_reader's session bar count, but there is no src/gazbot7/tape_reader.py — the module is scripts/tape_reader.py. A verification command that names the wrong path cannot be run, which is its own finding: this row was written without executing its own check. Carried forward with the path corrected.
SATURDAY 3	★ The \$250 daily loss limit pages you a claim its own journal contradicts	NOT DONE	scripts/daily_loss_limit.py computes changed = bench_all(...) and then pages the literal string “All six gates benched” instead of reading changed. Last commit 2026-09-13, before the card. One line, still there.

PLAY	WHAT IT ASKED FOR	STATUS	THE CHECK THAT DECIDED IT
SATURDAY 4	★ The loss limit's trigger can be moved by a price that never printed	NOT DONE — and this week it got WORSE	session_pnl() still sums where closed_at is not null and closed_at >= ? with no data_quality filter. It counted the -\$304.50 fabricated price on 09-15; this week it will count id 1019, -\$842.50, assistant_test_claim_NOT_operator_decision — a trade that was never his decision — straight into a -\$250 trigger that benches the whole gate book. The defect has acquired a second and worse input.
SATURDAY 5	A hand-run of event_alert.py wrote a live PAGED line into the record	NOT DONE — no dry-run flag exists	scripts/event_alert.py has no --dry-run and no argparse at all; it still calls log(f"PAGED ...") unconditionally. Last commit 2026-09-18, the day the card was written, and nothing since.
SATURDAY 6	A TARGET_100 rung booked a fill 79.5 points BELOW the entry it was supposed to take \$100 a lot above	NOT DONE — still untagged	Trade id 868 is unchanged: LONG 1 @ 29,320.25 → 29,240.75, -79.5pt, -\$160.50, exit_reason = TARGET_100, and data_quality IS NULL — so it is still counted as clean in every total on this desk, including the ones in this report. Either the rung fired at the wrong price or the reason is mislabelled; neither has been decided.

THE HONEST LIMIT ON THIS TABLE

These are the **nine rows of the 2026-09-18 card**. The card BEFORE it — reports/friday_v7/plays.json.pre-0918, **58 plays** about the paper tournament, the shadow board, the router and the greenfield labs — is still **UNGRADED**, and that is a different thing from done or not done. No phase has run since 09-18 that could check a tournament row, a shadow arm or a router setting. The file is on disk, every row still carries its verification command, and if the tournament comes back so does its ledger. **Read those 58 as not looked at** — not as passed.

⚠ **SATURDAY 2 is graded UNGRADED for an unusual reason** and it is worth one sentence: its verification command names src/gazbot7/tape_reader.py, which does not exist — the module is scripts/tape_reader.py. The check could not be run as written, which means it was

never run when the row was written. That is a small fault in the card itself, and the path is corrected on this week's.

§ The evidence behind those plays

Everything from here down is the working, and it is four sections that build on each other in order. **Section 1** is the record: every entry you made this week, by day, by hold time, by hour, by event day and by how it ended, plus what the book looks like now. **Section 2** asks what the tape looked like at the moment you pressed — the leg, where price sat in the range, the VWAP stretch, the volatility — each one measured against control minutes from the same sessions, which is the only comparison that is legitimate when every recorded example is a “yes”. **Section 3** is the verdict: what can be automated today, what cannot, and exactly how far off the rest is. And **Section 4** draws it — all five sessions with every one of your fills marked on the day's own price path, plus the FOMC hour on MNQ and on gold. Read Section 4 alongside Section 1: the shape that cost the most money this week is the hardest thing to see in a table and the easiest to see on a chart.

★ **What is deliberately NOT here.** No census section, no gate rehab, no shadow board, no router review, no greenfield lab. Those are the tournament report and this is not it. Their fragments are on disk and a month old, and the build will not stitch them — see the red box at the top of this page. ★ **One precise exception, because “we did not print it” and “we may not cite it” are different rules:** the run census itself was re-frozen tonight against this week's tape, and Section 4 cites it twice — once to rank the FOMC hour's runs against the week's other 66, and once to withdraw a Rev 1 caption that mis-read `census_summary.json`'s top lists as a ranking of the week when its own `_note` says they are the 51 runs the desk SAT OUT. A fresh, on-subject, named citation is fine; a month-old

section stitched under this masthead is not. If you are looking for a gate's P&L this week, the answer is in the front table and it is that the automated desk took **nothing** — and §3.4(d) now shows, with timestamps, that this was a nightly decision rather than a stuck switch.

Fragment freshness is stated once, above — see “what did not run this cycle, named exactly” on the front page. It is not repeated here.

The rule behind all of that: **a gap is never a null result**. “We looked and found nothing” and “nobody looked” are opposite instructions for next week, so nothing is silently omitted, nothing stale is silently stitched, and where a fresh section contradicts itself the build says so in a banner rather than quietly picking a side.

WHAT IS IN HERE

1. **Front** — [the week that was & the week coming, the honest headline, WHAT TO ACTUALLY DO, last week's card](#)
 2. **Section 1** — [THE RECORD — every entry you made this week, and what the book now looks like](#)
 3. **Section 2** — [THE CONDITIONS — what the tape looked like at the moment you pressed](#)
 4. **Section 3** — [THE AUTOMATION GAP — what can be automated today, what cannot, and how far off](#)
 5. **Section 4** — [THE RUN CHARTS — every fill of the week, marked on the day's own price path](#)
-

★ THIS SECTION IS AN EARLIER CYCLE'S — IT DID NOT REGENERATE THIS WEEK

Section 1 — THE RECORD — every entry you made this week, and what the book now looks like was not rebuilt by the 2026-09-25 run: it did not regenerate this cycle. The text below is the `op_record.html` on disk, last touched **2026-09-24 09:43 UTC**, and **it describes the week to Friday 2026-09-18, not the week to 2026-09-25.**

△ **Its mtime is NOT evidence of its vintage** — a git operation on 2026-09-24 touched every fragment in this directory, which is exactly how an earlier week nearly shipped under tonight's masthead. The week stamp above is what this banner trusts.

It is reproduced in full rather than deleted, because the working in it is real and nothing else on the desk records it. **But every number, date and verdict below is a PRIOR-WEEK number.** Do not line one up against a number in the front matter, the action card, the graded ledger or the run charts, all of which are this week's; do not carry a conclusion out of it onto the action card; and where it disagrees with the front matter, the front matter wins.

And read it as a GAP, not as a null result. The question this section asks was not asked again this week. "We looked and found nothing" and "nobody looked" are opposite instructions for next week.

S1 THE RECORD — Monday 14 to Friday 18 September 2026

Garrath — **you made +\$2,916.49 this week** across **39 entries**, every one of them 4 lots, every one of them opened by your hand and stamped manual in the database. That is **+\$583.30 a day at 4 lots** over five sessions, and it is comfortably the best week this book has had — the previous best was +\$1,617.50 and the week before this one was −\$3,049.49. It also turned the whole book around: before Monday the rider was **−\$2,591.43 over 53 entries** for its entire life, and it now stands at **+\$325.06 over 92** — see §6

for what that 92 excludes and what the unfiltered book reads. Six weeks of work were in the red; one week put it in the black.

Two warnings before the tables. First, **five trades are the week** — the best five made +\$3,162.48, which is more than the whole week, and the other 34 entries together came to −\$245.99. Second, the money is grouped by **DECISION, not by database row**: the week produced 62 rows in trades but only 40 entries, because a scale-out books one row per lot leaving. Counting rows would have inflated the trade count by 1.55×

★★ IF YOU READ NOTHING ELSE — FIVE THINGS

1. The week, in one line. +\$2,916.49 on 39 entries. 25 winners worth +\$8,232.99, 14 losers worth −\$5,316.50. Median entry +\$160.00. **Three green days and two red** — Mon +\$839.49, Tue −\$114.00, Wed −\$833.00, Thu +\$1,477.50, Fri +\$1,546.50 — and the bad one was Wednesday, **FOMC day**. (Tuesday is red by \$114.00, which is one scratch trade wide; Wednesday is the day that cost money.)

2. Your losers are wrong from the first minute. This is the cleanest thing in the whole section and it is pure description, no model. Across the week your 25 winners had a median best-moment of **59.2 points** in your favour and you booked a median **70% of it**. Your 14 losers had a median best moment of **11.2 points** — 11 of 14 never once showed you 25 points, and 7 never showed you ten. You are not giving winners back. You are occasionally getting on the wrong side and staying there.

3. The morning carried it, but do not turn that into a rule. 05:00–12:59Z was **+\$2,445.51 over 20 entries** this week; after 13:00Z you were −\$553.02 over 16. But over the whole life of the confirmed-manual record the same morning window is −\$779.48 — one trade on 10 September flips its sign. Point estimate first, uncertainty second: the morning is where you were this week, not a proven edge.

4. The eight-hour hole did not reopen. Over the book's life, **4 entries held past eight hours, 0 winners, −\$6,612.44** — and every other entry you have ever made adds to +\$6,937.50. This week **none of the 39 entries**

reached eight hours. Longest was 7h 01m. That is the single biggest behavioural change in the record.

5. You are still the exit. 38 of the 39 entries ended on your MANUAL_CLAIM; exactly one ended any other way — the −\$1,288.00 long on Tuesday that the 20:40Z clock flattened. The ladder took a first lot off on 14 of 39, but it did not finish a single one — all 14 still ended on your button.

1 • The week, day by day

Window: entries opened on or after 2026-09-14, gate day_rider, MNQ at \$2.00/point and \$1.50 a round trip. EXCLUDE: rows are dropped; the one BADFILL: entry is shown but kept out of every total — it is a real trade whose price came from the paper engine fabricating 3 of its 4 lots.

DAY	ENTRIES	WON	NET	BIGGEST WINNER	BIGGEST LOSER	WHAT HAPPENED
Mon 14 Sep	4	3	+\$839.49	13:18Z LONG + \$607.48	13:49Z LONG − \$162.00	Two morning entries, then the 13:18 runner — and a re-entry 31 min later that lasted 2m50s
Tue 15 Sep	6	4	−\$114.00	06:56Z SHORT + \$623.00	13:39Z LONG − \$1,288.00	Five entries made +\$1,174 by lunchtime; then one long held 13:39→the 20:40Z clock
Wed 16 Sep	6	3	−\$833.00	13:13Z LONG + \$819.00	18:52Z LONG − \$1,272.00	FOMC. Held long through the decision, bought again at T+52 with the slide still running
Thu 17 Sep	8	5	+\$1,477.50	11:55Z LONG + \$608.50	13:26Z LONG − \$441.00	Four of the eight took a first lot at the 100-pt rung; the overnight long ran 6h35m for + \$272
Fri 18 Sep	15	10	+\$1,546.50	13:55Z SHORT + \$434.50	16:11Z SHORT − \$208.00	Fifteen entries — the busiest day on record, and ten of them winners

DAY	ENTRIES	WON	NET	BIGGEST WINNER	BIGGEST LOSER	WHAT HAPPENED
WEEK	39	25	+\$2,916.49	+\$819.00	−\$1,288.00	+\$583.30/day at 4 lots

The week is five trades. Ranked by money, the top five are +\$819.00, +\$623.00, +\$608.50, +\$607.48, +\$504.50 — +\$3,162.48 together, against a week total of +\$2,916.49. The remaining 34 entries net −\$245.99: a scratch. The three worst are −\$1,288.00, −\$1,272.00, −\$762.00 — −\$3,322.00 — and two of those three are the two "held it" trades, Tuesday's clock-flat and Wednesday's FOMC re-entry.

2 · Every entry you made this week

One row per **decision**. "Best" is the furthest the trade ever went your way while you were in it, "Worst" the furthest against — both measured off the 5-second tape from the moment you got filled to the moment the last lot left. "Kept" is what fraction of the best moment you actually banked. Times are UTC; Paris is UTC+2.

OPENED (Z)	SIDE	LOTS	HELD	\$ NET	PTS	BEST	WORST	KEPT	HOW IT ENDED
Mon 14 09:09	SHORT	4	1h 23m	+\$145.01	+18.9	66.2	−27.8	29%	your hand, after 1 lot at the ladder
Mon 14 10:33	LONG	4	43m	+\$249.00	+31.9	36.4	−2.4	88%	your hand in one go
Mon 14 13:18	LONG	4	30m	+\$607.48	+76.7	113.1	−22.9	68%	your hand, after 1 lot at the ladder
Mon 14 13:49	LONG	4	3m	−\$162.00	−19.5	22.0	−21.0	—	your hand in one go
Tue 15 05:36	LONG	4	17m	+\$444.00	+56.2	61.5	−12.2	91%	your hand, after 1 lot at the ladder
Tue 15 06:56	SHORT	4	1h 14m	+\$623.00	+78.6	108.1	−41.9	73%	your hand, after 1 lot at the ladder
Tue 15 08:20	LONG	4	11m	(−\$304.50)	—	—	—	—	BADFILL — shown, not counted. The paper

OPENED (Z)	SIDE	LOTS	HELD	\$ NET	PTS	BEST	WORST	KEPT	HOW IT ENDED
									engine fabricated 3 of the 4 lots at exactly 0.1% adverse
Tue 15 08:38	SHORT	4	54m	+\$172.00	+22.2	50.8	0.0	44%	your hand in one go
Tue 15 10:29	LONG	4	1h 04m	+\$343.00	+43.6	66.5	-7.0	66%	your hand, after 1 lot at the ladder
Tue 15 12:24	SHORT	4	26m	-\$408.00	-50.2	5.2	-64.2	—	your hand in one go
Tue 15 13:39	LONG	4	7h 01m	-\$1,288.00	-160.2	15.8	-228.2	—	the 20:40Z clock in one go
Wed 16 06:12	LONG	4	1h 58m	-\$228.00	-27.8	30.5	-39.0	—	your hand in one go
Wed 16 08:11	SHORT	4	4h 13m	-\$762.00	-94.5	6.5	-107.0	—	your hand in one go
Wed 16 13:13	LONG	4	2h 35m	+\$819.00	+103.1	146.9	-32.6	70%	your hand, after 2 lots at the ladder
Wed 16 16:43	SHORT	4	41m	+\$160.00	+20.8	30.2	-22.5	69%	your hand in one go
Wed 16 18:52	LONG	4	16m	-\$1,272.00	-158.2	0.8	-170.8	—	your hand in one go
Wed 16 19:08	SHORT	4	31m	+\$450.00	+57.0	152.0	-11.8	38%	your hand, after 2 lots at the ladder
Thu 17 00:25	LONG	4	6h 35m	+\$272.00	+34.8	46.5	-64.2	75%	your hand in one go
Thu 17 07:07	LONG	4	1h 37m	+\$193.00	+24.9	57.0	-24.8	44%	your hand, after 1 lot at the ladder
Thu 17 11:55	LONG	4	21m	+\$608.50	+76.8	98.8	-8.5	78%	your hand, after 1 lot at the ladder
Thu 17 12:19	LONG	4	21m	+\$504.50	+63.8	73.1	-7.2	87%	your hand, after 1 lot at the ladder
Thu 17 13:26	LONG	4	1h 28m	-\$441.00	-54.4	21.5	-121.0	—	your hand in one go
Thu 17 15:27	LONG	4	1h 35m	+\$382.50	+48.6	64.0	-4.2	76%	your hand, after 1 lot at the ladder
Thu 17 19:35	SHORT	4	2m	-\$35.00	-3.6	1.0	-5.5	—	your hand in one go
	LONG	4	12m	-\$7.00	-0.1	4.6	-15.1	—	

OPENED (Z)	SIDE	LOTS	HELD	\$ NET	PTS	BEST	WORST	KEPT	HOW IT ENDED
Thu 17 19:38									your hand in one go
Fri 18 02:37	LONG	4	50m	+\$338.00	+43.0	45.2	-11.0	95%	your hand in one go
Fri 18 04:50	LONG	4	56m	+\$414.00	+52.5	61.0	-3.2	86%	your hand, after 1 lot at the ladder
Fri 18 06:06	LONG	4	59m	+\$409.50	+51.9	59.2	-3.0	88%	your hand, after 1 lot at the ladder
Fri 18 07:10	LONG	4	33m	+\$62.00	+8.5	33.5	-14.8	25%	your hand in one go
Fri 18 08:11	SHORT	4	57m	+\$176.00	+22.8	33.0	-18.5	69%	your hand in one go
Fri 18 09:31	SHORT	4	41m	+\$75.00	+10.1	28.0	-10.2	36%	your hand in one go
Fri 18 10:15	LONG	4	44m	+\$65.00	+8.9	41.8	-2.5	21%	your hand in one go
Fri 18 11:01	SHORT	4	1m	+\$2.00	+1.0	1.8	-6.5	57%	your hand in one go
Fri 18 11:21	SHORT	4	16m	-\$198.00	-24.0	6.8	-25.0	—	your hand in one go
Fri 18 12:16	SHORT	4	1h 07m	-\$30.00	-3.0	21.0	-14.2	—	your hand in one go
Fri 18 13:55	SHORT	4	20m	+\$434.50	+55.1	67.8	-10.8	81%	your hand, after 1 lot at the ladder
Fri 18 14:23	LONG	4	50m	-\$117.50	-13.9	28.6	-40.4	—	your hand in one go
Fri 18 16:02	SHORT	4	6m	-\$160.00	-19.2	3.2	-20.0	—	your hand in one go
Fri 18 16:11	SHORT	4	56m	-\$208.00	-25.2	53.0	-30.0	—	your hand in one go
Fri 18 17:07	LONG	4	1h 26m	+\$284.00	+36.2	48.8	-14.2	74%	your hand in one go
39 entries	23L / 16S	4 each	med 50m	+\$2,916.49	+393.8	1811	—	22%	38 by hand, 1 by the clock

The bottom row is worth a second look. Added up, the tape offered **1811 points** across your 39 trades and you took **394** — 22%. That number is NOT a criticism and it is not a target: it includes the losers, where the "best moment" was a point or two of noise, and nobody exits at the high. On the 25 winners alone you kept a median **70%**, and 17 of 25 kept two-thirds or more. Read it as evidence you are not the problem on the way out.

Sides: **23 longs for +\$2,479.98** and **16 shorts for +\$436.51**. You were more often long and the longs paid better this week, but on 39 trades that gap is one Wednesday afternoon wide — do not build on it.

3 · By hold time — the standing finding, refreshed

HELD FOR	ENTRIES (LIFE)	WON	NET (LIFE)	\$/ENTRY	THIS WEEK	NET THIS WEEK
under 15 min	19	7/19	−\$1,553.50	−\$81.76	5	−\$362.00
15-60 min	34	26/34	+\$6,063.98	+\$178.35	21	+\$2,965.98
1-3 hours	19	14/19	+\$2,998.51	+\$157.82	10	+\$2,090.51
3-8 hours	16	11/16	−\$571.49	−\$35.72	3	−\$1,778.00
over 8 HOURS	4	0/4	−\$6,612.44	−\$1,653.11	—	—
UNDER 3 hours	72	47/72	+\$7,508.99	+\$104.29	36	+\$4,694.49
OVER 3 hours	20	11/20	−\$7,183.93	−\$359.20	3	−\$1,778.00

The standing finding survives the week and grew: **under three hours + \$7,508.99 over 72 entries, over three hours −\$7,183.93 over 20** (it was +\$5,262/44 against −\$5,406/17 when this was last written — the week added 36 short holds worth +\$4,694.49 and 3 long ones worth −\$1,778.00).

△ THE CONFOUND, STATED PLAINLY — AND WHERE IT STOPS

Hold time is partly an OUTCOME, not only a choice. A trade that goes your way hits a rung and closes; a trade that goes against you gets held while you wait for it to come back. So part of that gap is simply "losers last longer", which is true of almost every discretionary book ever measured, and it does not mean a stopwatch would have made the money.

The over-8h bucket is not that shape. 4 entries, **0 winners**, −\$6,612.44, and their envelopes say why: the 21 Aug long never printed a single tick in profit (best moment −3.4 points) and the 10 Sep long showed +21.8 points before going 444 against. Those are not trades

held in hope of a recovery that nearly came. They are positions nobody was watching. That is abandonment, and a stopwatch does address it.

4 · By hour of day — does the edge live in part of the window?

First, a correction to the shape of the question. The "05:00–11:00Z almost exclusively" line comes from the 36 captured presses, and those only start on 14 September. Across the 44 entries the database can prove were yours, the spread is much wider — this week alone you opened trades in **18 different hours of the clock, from 00:25Z to 19:38Z**. $\triangle$ A correction to the correction, from Rev 1: that sentence used to read “from 09:09Z to 17:07Z”, which is not a clock range at all — it is the week's first entry (Monday 09:09) and its last (Friday 17:07), which is what `min()` and `max()` return when you run them over a full timestamp instead of over the time of day. It made the paragraph argue for a nine-hour window while claiming eighteen hours. The figures above are `substr(opened_at, 12, 8)`: earliest 00:25:01Z, latest 19:38:32Z.

BLOCK (UTC)	PARIS	ENTRIES	WON	NET THIS WEEK	\$/ENTRY	WHAT SITS IN THERE
Overnight 00-04Z	02:00-06:59	3	3	+\$1,024.00	+\$341.33	Asia / pre-Europe
Early Europe 05-08Z	07:00-10:59	9	7	+\$1,089.50	+\$121.06	the London open block
Late Europe 09-12Z	11:00-14:59	11	8	+\$1,356.01	+\$123.27	the pre-US grind
US open 13-16Z	15:00-18:59	11	5	+\$26.98	+\$2.45	the cash open and the first three hours
US afternoon 17-23Z	19:00-01:59	5	2	-\$580.00	-\$116.00	into the 20:40Z flatten
05:00-12:59Z — ★ the European morning, and the window the refutation is scored on	07:00-14:59 Paris	20	15	+\$2,445.51	+\$122.28	where the week was made
Everything before 13:00Z — the same window plus the	before 15:00 Paris	23	18	+\$3,469.51	+\$150.85	the 00-04Z row added back in

BLOCK (UTC)	PARIS	ENTRIES	WON	NET THIS WEEK	\$/ENTRY	WHAT SITS IN THERE
three overnight entries						
13:00Z and after	15:00 Paris on	16	7	−\$553.02	−\$34.56	holds both big losers
WEEK — the five blocks, and the two windows that reconcile to it	—	39	25	+\$2,916.49	+\$74.78	23 before 13:00Z + 16 after = 39 ✓

△ **WHICH WINDOW IS THE REFUTATION SCORED ON? 05:00-12:59Z, AND THE TWO DIFFER BY \$1,024**

Rev 1 printed a row labelled “Before 13:00Z” that was really the 05:00-12:59Z filter, so the three overnight entries (00:25, 02:37, 04:50 — **all three winners, +\$1,024.00 between them**) were in neither summary row and the table did not add up to the week. Both rows are now printed and labelled, and the WEEK line reconciles. **NOT-AN-ACTION no. 1 (last week’s card) — “only trade the European morning” — is scored on 05:00-12:59Z**, because that is the window the claim names and the only one with an out-of-sample leg. Scored on everything before 13:00Z it looks \$1,024 better this week and it is still refuted by the lifetime figure below, so nothing about the verdict turns on the choice; the labelling does.

Point estimate first. This week the morning block made +\$2,445.51 on 20 entries and the afternoon lost −\$553.02 on 16. **Uncertainty second, and it is large.** Take every entry the database can prove was yours, not just this week's, and the same morning block is **−\$779.48 over 25 entries** while the afternoon is −\$553.02 over 16. The sign of the morning flips entirely, because one trade — 10 September, 05:40Z, held 13h34m, −\$2,656.00 — is bigger than the whole week's morning profit. **One trade decides this cut.** That is not an hour-of-day edge; that is a hold-time problem wearing an hour-of-day costume.

★ **And the two sides of that split are not equally tested — only one of them is tested at all.** “Every entry the database can prove was yours” adds **five** pre-week entries to the morning and **none** to the afternoon: all five provable pre-week presses (10 Sep 05:40, and 11 Sep at 05:57, 08:59,

09:12, 10:39) fall inside 05:00–12:59Z. So the afternoon's −\$553.02 over 16 is this week's sixteen entries restated, not a wider sample — it cannot flip, because there is nothing outside the week for it to flip against. That makes the refutation stronger, not weaker: **the morning claim is the only side that has ever been looked at outside this week, and looking is what killed it.**

5 · By event day — refreshed with this week's FOMC

DATE	EVENT	ENTRIES	WON	NET	THE TRADE THAT DID IT
Wed 12 Aug	CPI	1	1	+\$224.50	One entry, worked
Fri 04 Sep	NFP	1	0	−\$341.00	One long held 8h44m — the NFP day is also an abandonment day
Fri 11 Sep	CPI	4	1	−\$568.99	Four entries, one winner
Wed 16 Sep	FOMC	6	3	−\$833.00	Long held through the 18:00Z decision, then bought again at T+52 with the slide still running: −\$1,272.00 on that one trade against a day of −\$833.00
EVENT DAYS	4 days	12	5	−\$1,518.49	−\$126.54 an entry
ORDINARY DAYS	25 days	80	53	+\$1,843.55	+\$23.04 an entry

The event side reproduces the standing figure to the cent: **−\$1,518.49 over 12 entries**. The ordinary side has moved because the week added Thursday and Friday — it now reads +\$1,843.55 over 80 entries rather than the +\$1,076/67 quoted in the build queue.

△ **FOUR DAYS. READ IT AS A WARNING LIGHT, NEVER A SWITCH**

This is **4 days of trading** and the split was chosen after seeing the losses. FOMC Wednesday was the only red day of the week — but it was red because of one entry at T+52 into a slide that ran for another 33 minutes, and you had already said yourself: "if i knew the fed was making a decision i probably wouldve held off." The calendar and the

T-60/T-15 pages are already built and already running. That is the right shape: it tells you an event is coming and refuses to tell you which way. Nothing here justifies benching anything automatically.

6 · How they ended — your hand, a ladder rung, or the clock

Your ladder is set to **\$100 / \$200 / \$400 / \$600 a lot** and takes exactly one lot at each rung. Here is what actually fires.

ENDING	ENTRIES THIS WEEK	NET	ENTRIES (LIFE)	NET (LIFE)	READ
Your hand — MANUAL_CLAIM	38	+\$4,204.49	81	+\$1,103.56	you decided it was over
A ladder rung took the last lot	—	—	7	+\$341.50	the machine finished it
The 20:40Z clock — CLOCK_FLAT*	1	-\$1,288.00	3	-\$1,024.50	nobody decided; the flatten did
Cross-desk flatten	—	—	1	-\$95.50	the other book's guard
LIFETIME, on the starred basis	39	+\$2,916.49	92	+\$325.06	81 + 7 + 3 + 1 = 92 ✓
CLOCK_FLAT_BUG — the row the asterisk excludes	—	—	1	-\$69.00	2026-08-21 09:11, flat 60 seconds later. A software flatten, not a decision
LIFETIME, unfiltered (gate LIKE 'day_rider%', data_quality IS NULL)	39	+\$2,916.49	93	+\$256.06	the same query with nothing taken out

★ WHAT THE ASTERISK ON CLOCK_FLAT* MEANS — IT WAS UNDEFINED IN REV 1 AND IT SETS FOUR OTHER NUMBERS

CLOCK_FLAT* is CLOCK_FLAT plus CLOCK_FLAT_RECON, and it **excludes the 2026-08-21 09:11 entry, -\$69.00, whose exit reason is CLOCK_FLAT_BUG** — a position opened and flattened by software sixty seconds later. It is excluded because it is not an ending anybody chose, and this table is about who decided. **The exclusion is defensible and it was undisclosed, which is worse than the exclusion.** It is what

makes the lifetime book read **92 entries / +\$325.06** rather than **93 / +\$256.06**, and that 92 and that +\$325.06 are quoted on the front page, in this section's opening paragraph, in §7 and in §3 of Section 3. Both figures are now printed above so the two can be told apart. Everything else on the row is unchanged: with the bug row put back, *CLOCK_FLAT** reads 4 entries / −\$1,093.50 instead of 3 / −\$1,024.50. △ Note the same 09:11 decision also has a second leg tagged *BADFILL:stale_price_30pt_outside_book_20260821* at −\$81.50, which the *data_quality IS NULL* filter has always dropped — so no reading of that morning is a clean one.

You are the exit, and it is not close. 81 of the 92 entries this book has ever made ended on your claim button. The ladder has taken at least one lot on 31 of them — a \$100 or \$200 target, or an older trail — but the **\$400 and \$600 rungs have never once fired in the entire record** — not a single *TARGET_400* or *TARGET_600* row exists. You always get there first.

△ **ONE TEMPTING NUMBER IN HERE IS A TAUTOLOGY. DO NOT ACT ON IT.**

The 14 entries that touched a rung this week made +\$6,377.99; the 25 you exited in one go lost −\$3,461.50. That looks like "scaling out works". **It is not a finding.** A rung can only fire on a trade that had already gone \$100 a lot your way — so that split is very close to "winners won". Their median best-moment was 67.1 points against 28.0. There is no lever in it.

The one exit this week that was not yours cost −\$1,288.00: Tuesday's 13:39Z long, held 7h01m, flattened by the 20:40Z clock. Its envelope shows a best moment of **15.8 points** and a worst of **228.2** — it was never really a trade that worked and then failed. Whether the clock made it better or worse is a question about repricing an exit, and this desk does not reprice exits (119 calibrations, none reproduced your claiming). What can be said without a model: it spent seven hours going one way.

7 · Whose trades are these — and why the week is the honest number

PROVENANCE	ENTRIES	NET	DATES	WHAT THE DATABASE ACTUALLY SAYS
Confirmed yours — entry_source='manual'	44	-\$308.50	2026-09-10 → 2026-09-18	You pressed BUY or SELL. 39 of these 44 are this week
UNKNOWN — entry_source IS NULL	48	+\$633.56	2026-08-06 → 2026-09-10	Predate the column. Nothing recorded who opened them, and nothing ever will

This matters more than it looks. The 48 unknown-source entries contribute +\$633.56 and the 44 that carry your stamp contribute -\$308.50 — so **the book's +\$325.06 lifetime net is not attributable to anybody in particular** (44 + 48 = the 92-entry basis defined in §6; -\$308.50 + \$633.56 = +\$325.06). Everything the database can prove is your hand is those 44 entries, and 39 of them happened this week. Strip the week out and the 5 provable pre-week entries are -\$3,224.99.

One shape in the unknown rows is worth naming without drawing a conclusion from it: for **12 straight sessions from 10 to 26 August the book took exactly one counted entry a day, and every one of them opened in the 13:00Z hour** — 20 of the 48 unknown entries sit in that single hour. The day rider's own documented automatic behaviour is one entry per session, no entry after 15:00Z. I am not claiming those rows were the machine's; the record does not say, and the standing rule is that a NULL stays a NULL. But it is why this section leads with the week, where every single entry carries your stamp, and treats the lifetime tables as context rather than as your track record.

8 · Ranked shortlist — what this section says to do next

#	WHAT	MEASURED ON	THE NUMBER	VERDICT
1	Record the passes. The button shipped Friday;		One side of a boundary. The fitting threshold is	BLOCKER

#	WHAT	MEASURED ON	THE NUMBER	VERDICT
	nothing has been written to it yet. Until it has, every table in this section is a description of trades you took and nothing here can become a rule.	36 presses, 40 claims, 0 passes	30 labelled examples and it was written for positives AND negatives, so the real count against 30 is zero.	
2	A hard flat-out at eight hours. Needs no model, no judgement and no fitting — it is a stopwatch.	4 entries, 29 sessions	0 winners from 4, – \$6,612.44. Every other entry in the record nets +\$6,937.50. That is –\$228.02 of drag per session traded. It did NOT recur this week (0 of 39) — so on this week's data it would have earned nothing. What it recovers is not computed here; that needs the exit repriced (\$3).	PROMISING
3	Keep standing aside on event days. Calendar and T-60/T-15 pages already built and running; this is about whether to keep trusting them.	4 event days, 12 entries	–\$1,518.49 on event days against + \$1,843.55 over 80 ordinary entries. FOMC Wednesday was the week's only red day. Four days, and the split was chosen after seeing it.	PROMISING
4	"Only trade the European morning." Tempting after this week. Do not do it.	20 entries this week, 25 provable ones over the book's life	+\$2,445.51 this week; –\$779.48 over the confirmed-manual record, where one trade on 10 Sep flips the sign. A filter fitted on a week that a single trade can reverse.	REFUTED as a rule
5	Watch the "it has never gone my way" shape. Your losers do not start well and then fail — they are wrong immediately.	14 losers this week	Median best-moment 11.2 points against 59.2 for winners; 11 of 14 never showed 25 points and 7 never showed ten. ⚠ Turning this into a bail rule means backtesting an exit, and the standing rule forbids a 120th calibration. Log it, watch it, do not fit it.	WATCH
6	Chase one execution defect. A TARGET_100 rung on 11 Sep booked –\$160.50 — a \$100-a-lot rung (that is +50 points on MNQ) that booked a fill 79.5 points BELOW the entry (trade id 868, long at 29320.25, out at 29240.75).	1 row, outside this week	Every other TARGET_100 in the record booked between +\$97.50 and +\$132.25. This one landed 129.5 points from the level it names. Small money, but a target exit is reporting a price its own trigger cannot have seen.	DEFECT

So what do you actually do on Monday? Nothing changes about how you trade — this section found no rule that survives its own uncertainty, and it would be dishonest to hand you one off 39 entries. Three things it does say: press PASS when you look and walk away, because that is the only thing standing between a description and a model; the eight-hour hole is the one number in the record big enough to be worth a stopwatch; and the morning result, lovely as it is, is one week.

One housekeeping note for §2, and it has to be said twice because the answer depends on the key. data/

operator_reads.jsonl holds **135 lines**. Every press is written twice (PathModified fires twice per write, ~1.4s apart), so the real figure is either **78 presses counted by their own timestamps** or **77 records under the (kind, request.raw) key** — the key cannot separate the two unlabelled 09-14 buys at 14:51:43 and 14:53:06, which are 83 seconds apart and are two real presses. **78 and 77 are both right; they are answers to different questions**, and Rev 1 printed one in this note and the other in §2.1 without saying so. Either way the **entry** count is 36 and counting raw lines would report 65 buy rows where there are 38 buy presses (37 under the key). See BUILD no. 6 (last week's card): the de-dup helper should key on (kind, ts) so the count stops depending on which sentence you read.

★ THIS SECTION IS AN EARLIER CYCLE'S — IT DID NOT REGENERATE THIS WEEK

Section 2 — THE CONDITIONS — what the tape looked like at the moment you pressed was not rebuilt by the 2026-09-25 run: it did not regenerate this cycle. The text below is the `op_conditions.html` on disk, last touched **2026-09-24 09:43 UTC**, and **it describes the week to Friday 2026-09-18, not the week to 2026-09-25.**

△ **Its mtime is NOT evidence of its vintage** — a git operation on 2026-09-24 touched every fragment in this directory, which is exactly how an earlier week nearly shipped under tonight's masthead. The week stamp above is what this banner trusts.

It is reproduced in full rather than deleted, because the working in it is real and nothing else on the desk records it. **But every number, date and verdict below is a PRIOR-WEEK number.** Do not line one up against a number in the front matter, the action card, the graded ledger or the run charts, all of which are this week's; do not carry a conclusion out of it onto the action card; and where it disagrees with the front matter, the front matter wins.

And read it as a GAP, not as a null result. The question this section asks was not asked again this week. "We looked and found nothing" and "nobody looked" are opposite instructions for next week.

2 WHAT THE TAPE LOOKED LIKE WHEN HE PRESSED

Garrath — **the sample is now 78 recorded presses, and 36 of them are entries.** That is the headline number of this section, because the fitting threshold this desk wrote down for itself is 30 and the positives have crossed it. **The negatives have not: there are ZERO.** The PASS button went in on Wednesday and has never been pressed, so every row in the file is a "yes". That is not a technicality — it is the whole ceiling on this section. With one side of the boundary you can describe what your entries look like; you cannot learn what separates them from the moments you looked and

walked away. So this is a description, and I have not tuned a single threshold on it.

One thing separates your presses from ordinary minutes, and it is the leg. At 33 of 36 entries a leg was already running — the desk's own definition, $1 \times \text{ATR}$ travelled in 8 minutes, dead on a $1 \times \text{ATR}$ give-back — against **74.6%** of all minutes on the same four sessions (session-blocked permutation, $P = 0.008$). And where a leg was running you took **its side 30 times out of 33**, against a coin-flip null of 50%: 18 of 19 buys into an UP leg, 12 of 14 sells into a DOWN leg. **You are a continuation trader and the tape says so without being asked.** You also do not wait for the leg to mature: median leg age at your press is **14 minutes**, and the median age of a randomly chosen in-leg minute is 13. You press almost as soon as the thing exists.

And the honest half, which is bigger. The number everyone expects to matter — where price sits in the day's range — **does not separate your presses from noise at all.** Ranked inside its own session, the median press sits at the **54th percentile** of all that session's minutes, 95% CI **[44.6, 63.5]**. That interval contains 50. Re-measured with a clean 22:00Z session ruler instead of the desk's own, it is 56.7 [47.5, 66.2] — the same answer. VWAP stretch is the same story: 54.0 [45.6, 62.4]. **"He buys high in the range" is true and meaningless, because so does every other minute of these four days.** That comparison is the one thing this dataset can do without negatives, and it kills two filters before anyone builds them.

FIVE THINGS TO READ BEFORE ANY NUMBER BELOW

(1) 135 rows in the file are 78 presses — or 77 records, and the difference is one row of the ladder below. PathModified fires twice per write, so nearly every press is recorded twice $\sim 1.4\text{s}$ apart with an identical request.raw. Deduping on (kind, request.raw) gives **77**; deduping on the press's own timestamp gives **78**, and 78 is the number used throughout this section. They differ by exactly one thing: the two unlabelled 09-14 buys, 83

seconds apart, which carry no request . raw for the key to separate them by. Counting raw rows would have inflated the sample by 73%.

(2) The ruler is `tape_reader.context()`, re-run at every minute. I did not invent a second definition of ATR or position-in-range for the baseline — I re-executed the capture's own code path at all 5,255 minutes of the four sessions, so press and baseline are measured identically. Validated press-by-press against the stored facts first: position-in-range agrees to a median 0.003 (p90 0.011), ATR to 0.45pt. The residual is capture write-lag, not a different calculation.

(3) Its "day so far" is not a day. Before 13:30Z the session window is the last 24 hours; after 13:30Z it is every minute with a time-of-day past 13:30 inside the last 24 hours — which splices yesterday afternoon onto today's open. The 2026-09-14 14:51 read says "session bars 203": that is Sunday 22:00–24:00Z plus Monday 13:30–14:51Z, with an eleven-hour hole in the middle, and the "day open" it showed him was Sunday's globex print. 24 of the 36 entry presses were taken before 13:30Z. Every conclusion below is re-checked against a clean 22:00Z-anchored session and none of them moves — but the number on his screen is not the day's range.

(4) The desk's own direction detector was BLIND at two thirds of these presses. drift threw no 5s rows since 13:30Z today at 51 of the 78 presses (24 of the 36 entries), because it is hard-anchored to the RTH open and he trades the European morning. At the 12 entries where it did read, it returned `confirmed=false` **all 12 times**. Whatever he is reading, the automatic path could not have seen it — and for most of the week could not even look.

(5) The fills WERE checked against the tape, and the fabrication is one leg, on an EXIT. Rev 1 of this section said "the paper engine fills lots 2–4 at exactly 0.1% adverse" and concluded that every P&L figure here was contaminated. **That is not what the tape says, and the claim is withdrawn.** I priced all 124 fills of the week (62 legs × entry and exit) against the last MNQ print at the fill second in `capture.db`. The table below is the whole answer.

CHECKED AGAINST THE LAST PRINT AT THE FILL SECOND	RESULT
--	--------

READ

124 every leg of the week, nothing sampled
--

CHECKED AGAINST THE LAST PRINT AT THE FILL SECOND	RESULT	READ
Fill prices checked (62 legs, entry + exit)		
Gap from the tape — median / p90	0.0017% / 0.0089% 0.50pt / 2.62pt	ordinary multi-lot slippage on a 29,000-point instrument
Within 0.03% of the tape	123 / 124	—
All 40 entry-group prices	max 6.75pt = 0.023%	the worst is 09-16 19:08 SHORT. No entry is anywhere near 0.1%
Reaching the 0.1% class	1 / 124	trade id 895, and it is the EXIT: out at 29,269.3125 against a last print of 29,291.00 — 21.69pt = 0.074% adverse. Its ENTRY is clean (1.88pt = 0.006%)

The mechanism in the tag is right and the side in it is wrong. Three of four lots leaving at 0.1% adverse off a reference of 29,291.3 blends to **29,269.28**; the booked exit is **29,269.3125**, a 0.28-point miss. I swept all 124 fills for that signature and it matches **exactly one** — id 895, already tagged BADFILL:paper_engine_0.1pct_fabrication_3of4_lots_20260915. **So: the week's dollars are venue dollars, with one flagged exception that every total in this report already excludes.** That is the same thing the front page says, and Rev 1 had the two sections contradicting each other on the single question of whether the money is real. Dollars are still secondary here — this section is about the tape, not the P&L — but they are no longer void.

2.1 · THE SAMPLE — what is actually on record

LADDER	COUNT	WHAT IT MEANS
Rows in data/operator_reads.jsonl	135	raw, double-fired
Distinct presses — by their own TIMESTAMP	78	the number used everywhere in this section. Duplicates are the double PathModified write, ~1.4s apart
Distinct records — under the (kind, request.raw) key	77	△ the key collapses two real presses into one — see the row below. Both counts are right; they answer different questions
— ENTRY presses with a recorded side	36	21 BUY / 15 SELL. This is the fitting count, and it is over 30. Identical under either key
— CLAIM presses	40	3 of them name a lot count; the rest claim the book. Identical under either key
— buy presses whose label the rider ate	2 by timestamp 1 by the key	this single row is the whole 78-vs-77 gap. 2026-09-14T14:51:43 and 14:53:06 — 83 seconds apart, so two real presses — but both carry an empty request, so (kind, request.raw) cannot tell them apart. Pre-dates the ExecStartPre snapshot fix; side UNKNOWN and it stays unknown
— PASS presses (the negatives)	0	

LADDER	COUNT	WHAT IT MEANS
		the button exists since 2026-09-18 and has never been pressed. This is the binding constraint on everything below
— of which the BADFILL entry \$1 excludes		09-15 08:20 BUY, trade id 895. Its TAPE facts are real, so it stays in every tape comparison below; its DOLLARS are not, so \$2.4 now excludes it and reads 35. See the box under \$2.4
Sessions covered		09-15 (7 entry presses — \$1's clean day table says 6; the seventh is the BADFILL) · 09-16 (6) · 09-17 (8) · 09-18 (15)
Entry presses matched to a real entry	36 / 36	every one found an entry in trades inside +180s. No press was refused
Entries at 4 lots	36 / 36	there is no single-lot-vs-multi-lot comparison of DECISIONS to make
Baseline minutes for the comparison	5,255	every minute of the same four sessions, same ruler

One correction to carry into \$1 and \$3. The standing figure is that he trades 05:00–11:00Z almost exclusively — 26 of 28 presses. On this week's 36 that is **16 of 36**, spread across 18 different hours from 00Z to 19Z. The old window was a property of a smaller, older sample, not of his trading. Any rule built on it should be re-derived before it ships.

2.2 • THE ONE COMPARISON THAT IS LEGITIMATE WITHOUT NEGATIVES

Press population against ALL MINUTES of the same four sessions. The right-hand column is the one that matters: each press is ranked inside its own session's minute distribution, so a day on which he pressed fifteen times cannot drag the answer. If he pressed at random moments those ranks average 50.

AT THE MOMENT HE PRESSES	PRESS MEDIAN (P25-P75)	ALL 5,255 MINUTES (P25-P75)	WITHIN-SESSION RANK, 95% CI	READ
Position in the day's range (the desk's own ruler)	0.708 (0.48–0.92)	0.689 (0.45–0.85)	54.2 [44.6, 63.5]	NO SEPARATION. The interval contains 50
Position in the day's range (clean 22:00Z session, running extremes)	0.672 (0.34–0.92)	0.727 (0.35–0.90)	56.7 [47.5, 66.2]	same answer on a better ruler — so it is not a ruler artefact
VWAP stretch, in ATR	+5.21 (–0.73–+21.1)	+5.41 (–1.57–+14.4)	54.0 [45.6, 62.4]	NO SEPARATION. He is not fading stretch and he is not chasing it
ATR, points	8.81 (6.77–11.67)	7.93 (6.45–10.18)	61.5 [52.2, 70.4]	modest, clears 50
ATR percentile vs its own last 24h	71.7 (40.5–85.2)	49.3 (22.1–76.4)	62.9 [54.0, 71.4]	he presses into livelier-than-usual tape. 14 of 36 presses sit above the

AT THE MOMENT HE PRESSES	PRESS MEDIAN (P25-P75)	ALL 5,255 MINUTES (P25-P75)	WITHIN-SESSION RANK, 95% CI	READ
				75th ATR percentile, where only 25.9% of minutes do

Two filters died here and that is worth more than a finding. Position in range and VWAP stretch are the two conditions every hand-built entry gate on this desk has reached for. On 36 of his own presses, measured against the only honest control available, neither of them knows where he is. Nobody should spend another cycle on them. What survives is a weak volatility preference: he shows up when the tape is moving faster than it has been.

2.3 • THE LEG — the only thing that clearly separates a press from a minute

Leg definition is `scripts/leg_survival.py`'s, unchanged: a leg begins when the last 8 minutes have travelled $1 \times \text{ATR}$ in one direction, and dies when price gives back $1 \times \text{ATR}$ from its extreme. It is evaluated causally — at every minute I know only what was knowable then — and a press only counts as in-leg if the leg had already been detected, 8 minutes after its own start.

LEG STATE AT THE PRESS	PRESSES	ALL MINUTES, SAME SESSIONS	READ
A leg was already running	33 / 36 = 91.7% [78.2, 97.1]	74.6% (press-weighted)	session-blocked permutation, 20,000 draws: P = 0.008
— and he took ITS side	30 / 33 = 90.9% [76.4, 96.9]	50% (coin flip)	18 of 19 buys into an UP leg; 12 of 14 sells into a DOWN leg
Leg age at the press	14 min (11-17)	13 min (10-19)	he does not wait. Detection needs 8 minutes, so he is pressing ~6 minutes after the leg becomes knowable
Travel banked when he presses	28.5pt = 2.9 ATR	22.0pt = 2.5 ATR	slightly further along than a random in-leg minute, but inside the same spread
Distance to the extreme he is trading toward	10.6 ATR	—	he is not a breakout buyer. Only 5 of 36 presses are within 1 ATR of the running high (long) or low (short)
Same test with the scope's 3×ATR give-back	91.7%	97.6%	REFUTED as an instrument. A 3×ATR leg is alive in 97.6% of all minutes, so "a leg was running" carries no information at all. Use the 1× definition

It survives leave-one-session-out. Drop any one of the four days and the leg-alive rate is 89.7–95.2% and the with-the-leg rate is 88.5–92.6%. It is not one

trending Thursday carrying the result. What it is not, however, is an independent sample of 36: the presses cluster (7 of 35 consecutive entry presses came within 30 minutes of the one before, usually a claim followed by a re-entry within seconds), and four sessions is four sessions. Treat the interval as optimistic and the point estimate as the thing worth keeping.

The claims tell the same story more weakly: a leg was running at 35 of 40 claim presses (87.5%) against the same 74.6% baseline. That is a statement about when he claims, not a rule that reproduces what he claims — and per the standing limit, no attempt to reproduce the exits was made or should be.

2.4 • WHAT THE DOLLARS DID — post-hoc, three-deep, and not a rule

★ **The denominator changes here, and only here.** The 36 entry presses above include the 09-15 08:20 BADFILL, whose tape facts are real; its price is not, so **every dollar cell below is computed on 35**. On that basis the entries made + \$2,077.00 across four sessions at 4 lots = +\$519.25/session. (Rev 1 printed +\$1,772.50 / +\$443 a session, which is this figure plus the fabricated −\$304.50 — \$1 excluded that row from every total and \$2 silently put it back.) Two cuts fall out of the tape state above. Both are post-hoc, both have a cell of n=3 or n=12, and I am showing them so they can be tested later, not acted on now.

CUT (POST-HOC) — 35 ENTRIES, BADFILL EXCLUDED	N	TOTAL, 4 LOTS	MEDIAN	WINNERS	WHAT IT IS WORTH
Pressed WITH the running leg	29	+\$3,693.00	+\$176.00	20 / 29	the population he actually trades. The BADFILL sat in this cell; with it, 30 and +\$3,388.50
Pressed AGAINST the running leg	3	−\$1,515.00	−\$208.00	0 / 3	an anecdote, not a finding. Three trades. But all three lost, and the worst (−\$1,272, 09-16 18:52, buying into a DOWN leg at ATR 37.7) is one of the week's two largest losses. Unaffected by the exclusion
No leg running	3	−\$101.00	+\$2.00	2 / 3	nothing to say at n=3. Unaffected
ATR percentile terciles: low / mid /	12 / 11 / 12		+\$311 / +\$2 / +\$174	9 / 6 / 7	this still REVERSES the description, and

CUT (POST-HOC) — 35 ENTRIES, BADFILL EXCLUDED	N	TOTAL, 4 LOTS	MEDIAN	WINNERS	WHAT IT IS WORTH
high (re-cut on 35)		+\$2,762 / -\$233 / -\$453			it is weaker than Rev 1 said. He prefers lively tape (\$2.2) and this week the money came from the quiet third. The BADFILL sat in the HIGH tercile: with it the cells read +\$2,762 / +\$272 / -\$1,262 — so a fabricated price was carrying a third of the “lively tape loses money” story. Post-hoc terciles, 11-12 trades a cell, medians that barely move. Do not act on it — re-measure it
ALL 35	35	+\$2,077.00	—	22 / 35	+\$519.25 a session at 4 lots

★ THE THREE DENOMINATORS, RECONCILED IN ONE PLACE — THEY ARE THREE DIFFERENT INSTRUMENTS

\$1 counts 39 entries over 5 sessions from the trade record, and drops the BADFILL from every total. **\$2 counts 36 entry presses over 4 sessions** from `operator_reads.jsonl`, which only began recording on 15 September — and **\$2's 36 INCLUDES the BADFILL entry that \$1's 39 excludes**, because what \$2 measures at that moment (the leg, the range position, the ATR) is real tape and was really there. **\$3 counts 34 entries over 4 days**, because `gazbot7-leg-watch` started at 05:55:43Z on 15 September. The one place the inclusion would have mattered is the money, and the money in \$2.4 is now computed on **35**. Do not carry a rate across the three.

2.4b · FIVE TRADES ARE THE WEEK — so do those five look different at the press?

This report carries the “five trades are the week” caveat in six places and then never asks the obvious follow-up: were those five a different kind of moment, or the same moment that happened to keep going? The instrument to ask with is

already built, so here is the answer, with **n=5** stamped on every cell. Top five of \$2's 35 clean presses (+\$3,005.00) against the other thirty (−\$928.00). **AUC** is the chance that a randomly chosen one of the five is higher on that feature than a randomly chosen one of the thirty: **0.50 is a coin flip**.

AT THE MOMENT HE PRESSED	TOP 5 +\$3,005.00	OTHER 30 −\$928.00	AUC	READ
Leg age at the press	14 min	14 min	0.46	identical to the minute
ATR percentile	57.3	71.7	0.49	if anything the big winners came in quieter tape — and the interval is a coin flip
Position in the day's range	0.798	0.708	0.53	no separation, same as \$2.2 against the whole tape
VWAP stretch, in ATR	+10.88	+5.21	0.51	no separation
Travel already banked by the leg	2.50 ATR	2.99 ATR	—	the winners were LESS far along, not more
A leg was running, and he took its side	5 / 5	24 / 30	—	the only cell that looks like anything — and on the 83% base rate across all 35, five out of five happens 39% of the time by chance. It is not a signal
Hold time (an OUTCOME, not a condition)	30.5 min	52.2 min	0.49	the winners were SHORTER, because a winner hits a rung and closes. This is the confound \$1 \$3 names, not a finding

The verdict is a null, and printing it is the point. Every feature this desk can measure at the press is **indistinguishable between the five trades that made the week and the thirty that did not** — every AUC between 0.46 and 0.53. **This cannot support a rule at this n and nothing here should be built on. Re-measure at 70 entries.** What it does buy is a closed question: the big winners were not a different kind of moment that a filter could have found, so the honest reading of “five trades are the week” is size dispersion in one population, not two populations. △ And note what is not being claimed: five draws cannot detect anything short of an enormous effect, so this is “we looked and found nothing”, never “there is nothing”.

2.5 · RANKED SHORTLIST

LEAD	VERDICT	MEASURED	NEXT STEP
1 · PRESS THE PASS BUTTON. The negative examples	HIGHEST VALUE ON THE DESK	36 positives, 0 negatives. The 30-press threshold is crossed on one side of the boundary and	every time you look and do NOT trade, press it. ~30 passes over two weeks makes the decision boundary learnable for the first

LEAD	VERDICT	MEASURED	NEXT STEP
		cannot be crossed on the other by waiting	time in eleven months. Nothing else in this section is unblocked without it
2 · The continuation description — a live 1×ATR leg, taken on its own side	PROMISING	33/36 in a leg (vs 74.6% of minutes, $P=0.008$); 30/33 with it (vs 50%); holds leave-one-session-out	as an ALERT, not a filter. §3 scores leg_watch's pages against these 36 presses — did the machine see the leg before he did? That is measurable now, without negatives
3 · drift is blind before 13:30Z	FIX — DEFECT	51 of 78 presses, 24 of 36 entries, returned no 5s rows since 13:30Z today. The other 12 all read confirmed=false	anchor the detector's day to the CME session (22:00Z) or fall back to the trailing 24h, as tape_reader already does for its bars. Until then the auto path cannot even look at two thirds of his trading
4 · Position-in-range as an entry condition	REFUTED	within-session rank 54.2 [44.6, 63.5]; 56.7 [47.5, 66.2] on a clean session ruler. Both contain 50	stop building it. Two independent rulers, 5,255 control minutes, one answer
5 · VWAP stretch as an entry condition	REFUTED	rank 54.0 [45.6, 62.4]; press median +5.21 ATR against a baseline of +5.41	same. He is neither fading nor chasing VWAP
6 · The volatility preference	OPEN — CONFLICTED	he presses at ATR percentile 71.7 vs 49.3 (rank 62.9 [54.0, 71.4]) — but on the 35 clean entries the quiet tercile made +\$2,762 and the lively one −\$453 (it read −\$1,262 in Rev 1, with a fabricated fill inside it)	the description and the dollars still point opposite ways, on 11-12-trade cells, and less far apart than Rev 1 said. Re-measure at 70 entries before anyone writes an ATR band
7 · The 3×ATR leg definition	REFUTED	alive in 97.6% of ALL minutes; a test that everything passes	the scope's own text should move to 1×. leg_survival.py already defaults to it
8 · tape_reader's session window	FIX — INSTRUMENT	"day open/high/low" splices yesterday afternoon onto today's RTH; 09-14's 203 "session bars" were Sunday night plus Monday morning across an 11-hour hole	no conclusion here changes when it is fixed (checked), so this is about what he is SHOWN. Anchor to 22:00Z
9 · The 05:00-11:00Z window	NEEDS REFRESH	standing figure 26 of 28 presses; this week 16 of 36, across 18 distinct hours	re-derive before any hour-of-day rule ships. It was a property of the old sample
10 · The 78 captured narratives	UNSPENT	a ~3,160-character read of the whole session is stored with every press and nothing has been done with it	once passes exist, the narrative pairs are the actual experiment: the same text, one ending in a press and one in a pass

The one-line summary, in your units. 36 entry presses over 4 sessions; on the 35 whose price is venue truth, **+\$519.25 a session at 4 lots**. Where in the range you press and how far from VWAP you press are indistinguishable from picking a minute at random. What you actually do, 90% of the time, is take the side of a leg that is about 14 minutes old — and the desk's own leg detector can see that leg roughly six minutes before you press it.

★ THIS SECTION IS AN EARLIER CYCLE'S — IT DID NOT REGENERATE THIS WEEK

Section 3 — THE AUTOMATION GAP — what can be automated today, what cannot, and how far off was not rebuilt by the 2026-09-25 run: it did not regenerate this cycle. The text below is the `automation_gap.html` on disk, last touched **2026-09-24 09:43 UTC**, and **it describes the week to Friday 2026-09-18, not the week to 2026-09-25.**

△ **Its mtime is NOT evidence of its vintage** — a git operation on 2026-09-24 touched every fragment in this directory, which is exactly how an earlier week nearly shipped under tonight's masthead. The week stamp above is what this banner trusts.

It is reproduced in full rather than deleted, because the working in it is real and nothing else on the desk records it. **But every number, date and verdict below is a PRIOR-WEEK number.** Do not line one up against a number in the front matter, the action card, the graded ledger or the run charts, all of which are this week's; do not carry a conclusion out of it onto the action card; and where it disagrees with the front matter, the front matter wins.

And read it as a GAP, not as a null result. The question this section asks was not asked again this week. "We looked and found nothing" and "nobody looked" are opposite instructions for next week.

3 CAN THIS BE AUTOMATED?

Garrath — **the honest answer is that one half of it already is, and it is switched off by a single number.** `leg_watch` saw a leg running at **33 of your 34** entries this week and took your side at 26 of them. It paged you ahead of **3 of them within half an hour, and 8 within the hour** — the window matters and both figures appear below, so they are named here rather than left to collide. That is not a detection failure and it is not a modelling problem: the watcher's first page fires at **60 minutes** of leg age, the median leg dies at **22**, and the median leg you press into is **28 minutes old**. The machine is watching the right thing and speaking forty minutes after you have already acted.

And that is measurable without a single negative example. Scored against a session-blocked permutation — your own presses shuffled to random minutes of the same days — **the alerting policy you are running today does not separate from chance:** the 25 pages it actually sent preceded **24%** of your entries where shuffling gives **18%**, $P = 0.24$. Move the first page from 60 minutes of leg age to **15** and the same detector precedes **82%** against a chance rate of 56%, **$P = 0.0009$** . The price is your phone going from **7 to 32 pages a day**. That is the whole trade, it is a parameter, and nothing about your judgement has to be learned to make it.

The other half has not moved, and one thing on the "already running" list is worse than nothing. There are still **zero** recorded passes, so there is still no decision boundary to learn. Your exits remain unreproduced after 119 calibrations and I did not run a 120th. And the **\$250 daily loss limit** — listed as a piece that is already automated — turns out to be measured at **-\$67.09 a session** when applied to your book, cannot reach your book anyway, fired twice this week benching nothing, and paged you both times claiming it had benched six gates.

FOUR THINGS TO READ BEFORE ANY NUMBER BELOW

(1) The window is four days, not five. gazbot7-leg-watch started on 2026-09-15 and its log begins **05:55:43Z** that morning. **34 of your 39 countable entries** fall inside it; the four on Monday and the 05:36 Tuesday long do not, and are excluded from every alert score here rather than counted as misses. That is deployment, not truncation.

(2) One "alert" in the log is not an alert, and dropping it changes the headline. data/event_alert.log carries 2026-09-18T07:04:03Z PAGED T-15 FOMC. The timer is OnCalendar=*:0/5:30; journald shows the service ran at 07:05:30 and logged "ok — nothing in a band"; the only FOMC in the calendar is **2026-09-16**; and a T-15 band cannot open two days late. It is a hand-run test that wrote into the live log, and scripts/score_alerts.py credited it with leading your 07:10 Friday entry. It is dropped here, so this section reads **4 led, not 5**. The event-alert channel sent you **nothing real all week**.

(3) "He ignored it" and "he never saw it" are different, and the split is measured. scripts/presence.py reads nginx, and counts only a phone session as looking — a desktop tab polls 134 times a minute and would manufacture presence indefinitely. Of the 20 pages you did not act on, **15 fired while your phone was demonstrably active.** Those 15 are rejections. The other 5 are alerts nobody saw.

(4) No threshold below is a recommendation until it is pre-registered — and the headline P is the WINNING band's, uncorrected. The 15-minute page band is the best point on a curve I drew from 34 of your presses. Its lower neighbour agrees (10 min: $P = 0.015$); its upper neighbour at 30 min does not ($P = 0.43$), so the effect turns on somewhere between 15 and 30 minutes and 15 is one side of an edge, not the middle of a plateau. **$P = 0.0009$ is the permutation P of the band that won a four-band sweep, and nothing in it pays for having chosen it.** Four bands were tried (60/30/15/10), so the sweep-wide figure is $0.0009 \times 4 = 0.0036$ by Bonferroni (Šidák: $1 - (1 - 0.0009)^4 = 0.0036$, the same to four places). **It still clears comfortably at any conventional level, so saying so strengthens the row rather than killing it** — but the 10-minute band does not survive the same correction ($0.015 \times 4 = 0.060$), which is a second reason not to go past 15. A threshold picked off its own curve still has to be frozen and measured forward before it earns anything; the pre-registration is written out under §3.2 rather than promised.

3.1 · What is already automatable — the gap is DELIVERY, not DETECTION

Two different questions, and the desk has been answering only one of them. Did the machine see it? is about leg OPEN lines, which leg_watch writes silently by design — 49 a day, most of them dead inside half an hour. Did the machine tell him? is about ALERT lines, of which there were 25 all week. Everything worth having lives in the space between those two numbers.

ON YOUR 34 ENTRIES INSIDE THE LOG WINDOW

COUNT

READ

A leg was already running when you pressed

33 / 34
97%

the live watcher's own log, not a replay. §2
got 33 of 36 by re-running the leg code
causally — the same answer from a
different direction

ON YOUR 34 ENTRIES INSIDE THE LOG WINDOW	COUNT	READ
— and it was running YOUR way	26 / 34 76%	you are a continuation trader; \$2 measured the same thing at 30 of 33
A page reached your phone first (same side, within 30 min)	3 / 34 9%	15 Sep 10:29 (–1m), 18 Sep 16:11 (–4m), 17 Sep 00:25 (–23m). The best of them landed 58 seconds ahead — that is simultaneity, not a warning
A page reached your phone first (same side, within 60 min)	8 / 34 24%	the other five all landed 41-58 min out. Shuffled to random minutes, chance gives 18% — P = 0.24
Legs detected, 2026-09-15 05:55Z → 2026-09-18 21:00Z	176	49/day. Median life 22 min, p75 41, p90 60
Median age of the same-side leg you pressed into	28 min	p25 7, p75 42. The first page band is 60 min

The brief says the desk beat you to a leg on 15 September and it did — but look at what the line says.

06:50:38Z leg OPEN DOWN from 29362.75 (silent – 17 legs/session, most die young), and you sold at 06:56:28. The detector was 5m 50s early and the word silent is in the log line. It knew, and the delivery policy decided you did not need to hear it. That entry made **+\$623.00**.

3.2 • The page band, priced — coverage against noise, with a null

Every row is the same detector. The only thing that changes is **how old a leg has to be before your phone rings**. "Chance" is your own presses shuffled to random minutes of the same sessions, 20,000 draws, sides kept — the only control that is legitimate when there are no recorded passes.

FIRST PAGE AT	PAGES A DAY	PRECEDED YOUR ENTRY (≤60 MIN)	CHANCE RATE	P	READ
60 min — the band that runs today	5	3 / 34 9%	12%	0.86	NOTHING. The 60-minute band alone is behind chance
30 min	19	14 / 34 41%	38%	0.43	still nothing. Halfway is not halfway
15 min	32	28 / 34 82%	56%	0.0009 sweep-wide 0.0036	the best point on the curve. 26 points of separation over chance, and it survives being corrected for having won a four-band sweep
10 min	41	29 / 34 85%	68%	0.015 sweep-wide 0.060	one more entry covered, nine more pages a day, and the

FIRST PAGE AT	PAGES A DAY	PRECEDED YOUR ENTRY (≤ 60 MIN)	CHANCE RATE	P	READ
					null catches up. Does not survive the correction. Past the knee
The correction, stated once	Four bands were swept; the P column's left figure is the winning band's own permutation P and pays nothing for the choosing. Bonferroni over the four: 15 min 0.0036, 10 min 0.060, 30 min 1.0, 60 min 1.0. The recommendation is unchanged				—

★ **PRE-REGISTRATION — WRITTEN BEFORE IT SHIPS, BECAUSE OTHERWISE NEXT WEEK IS A SECOND FIT**

Rev 1 put “freeze it and measure forward” on the card as a kill criterion after the change ships, which is not a pre-registration — it is a promise to look. Here is the frozen version, and nothing about it may be re-chosen next Friday:

FIELD	FROZEN VALUE
Band	first page at 15 minutes of leg age, every leg, no travel floor, no other filter. Destination: dashboard. The 60-minute phone page is unchanged
The measured quantity	fraction of his entries preceded by a same-side page within 60 minutes, on entries opened after the band change lands
Coverage target	$\geq 56\%$ — the shuffled chance rate at this band, not the 82% in-sample point. Anything at or below 56% forward is the null
Kill line	forward coverage below 50% over the first 30 forward entries, OR he stops reading the pages. Either puts the band back to 60 and the row dies
Re-scored on	Friday 2026-10-16 — four report cycles, ~120 entries at this week's rate, in next Friday's §3 under this heading
What may NOT be changed	the band, the 60-minute coverage window, the same-side rule, or the target. Changing any of them turns a test back into a fit, and this row's only claim to significance is that it is measured on data it was not drawn from

Every row in that table is a model of a policy, so here is the world. The rows above re-band the detector cleanly — a page at leg age T, every leg, no other filter. What actually shipped this week was 25 pages across the 60/90/120/180 bands with a travel floor on top, **7 a day**, and measured the same way they preceded **8 of your 34 entries (24%) against a chance rate of 18%, $P = 0.24$.** Both readings say the same thing and it does not depend on which one you prefer: **the alerting you are running today is not distinguishable from telling you to look at random moments of the same day.**

The shape is what matters, not the winning row. The separation is not a spike on one threshold — it builds from nothing at 60 min, is still nothing at 30, arrives at 15 and starts being eaten by its own null at 10. That is a knee, and a knee is the shape of a real effect; a single significant cell surrounded by noise is the shape of a fitted one. It is also why **32 pages a day is the honest cost and not a footnote**. You told the desk you look periodically when you have time. Thirty-two pages a day is not periodic, and the first thing that happens to an alarm nobody can answer is that it becomes wallpaper.

So the recommendation is the band, not the bell. Move the detection threshold to 15 minutes, and put the live leg list on the dashboard you already have open — presence says you are on it — rather than on the phone. Keep the phone for the 60-minute band that exists now. That buys the 82% coverage where you are already looking and leaves your pocket alone. It is also, not coincidentally, the thing that fixes §1's blocker — see 3.4.

3.2b · But HOW FAR AHEAD does a 15-minute page land — and is the leg still alive when he presses?

This is the poke the section's own numbers invite, so here is the answer before anyone else asks it. Coverage is scored on a ≤ 60 -minute window; the median leg dies at 22 minutes; and §3.1 dismisses the current band's pages precisely because “the other five all landed 41–58 min out”. If the covered entries are mostly covered by pages that arrived forty minutes early about legs that were dead by the press, then 82% is bookkeeping rather than a warning.

The instrument is the live journal, not a replay. Every leg below is a leg OPEN line gazbot7-leg-watch actually wrote (176 of them) and its life is the leg OVER after Nmin line that closes it (165 closed). The watcher counts its bands from the leg's start, and writes the OPEN line 8 minutes later at detection — the 09-15 leg that opened at 09:36:07Z logged quiet 30min at 09:58 and ALERT 60min at 10:28, which fixes the offset exactly. A band-T page therefore lands T–8 minutes after the journal line. That 8-minute offset is also why §3.1 says the median leg dies at **22** minutes and the journal's own *OVER* lines say **30**: §3.1 measures life from DETECTION, the log measures it from the leg's START. Same legs, same deaths, two origins — $22 + 8 = 30$, and p75 $41 + 8 = 50$, and p90 $60 + 8 = 70$, which is the whole distribution reconciling term by term.

FIRST PAGE AT	PAGES A DAY	COVERED (≤60 MIN, SAME SIDE)	MEDIAN LEAD	PAGED LEG STILL ALIVE AT THE PRESS	READ
60 min — the band that runs today	11.0	10 / 34 29%	31.4 min	5 / 10	△ three numbers exist for this band and they are three different things. 3 of 34 is §3.2's clean re- band scored at ≤30 min; 8 of 34 is what the pages that actually went out achieved (60/90/120/180 bands with a travel floor, 25 pages, ≤60 min); 10 of 34 is this journal re-band — every leg that reached 60 minutes, no travel floor, so it pages more often than the live policy does
30 min	26.2	20 / 34 59%	22.6 min	14 / 20	—
15 min	45.5	29 / 34 85%	26.5 min	19 / 29 66%	the live journal says 85% where §3.2's replay says 82% — two instruments, one answer
10 min	48.0	29 / 34 85%	26.4 min	21 / 29	no more entries covered; only more pages

AT THE 15-MINUTE BAND, THE THREE QUESTIONS THAT DECIDE WHETHER 85% IS REAL			ANSWER	READ
Was the paged leg still alive when he pressed? (the watcher's own 3×ATR rule, from its own OVER line)			19 / 29 66%	The main worry does not hold. Two thirds of the pages were about a leg that was still running when he acted
Was a same-side 1×ATR leg alive at the press? (§2.3's stricter ruler)			23 / 29 79%	the stricter definition agrees, and matches §2's 33-of-36 base rate
How many pages landed inside 20 minutes?			10 / 29	This is the honest cost. Median lead is 26.5 min (p25 17.5, p75 45.8) — longer than the median leg has left to live
Coverage if a page only counts when it lands inside 20 minutes			10 / 34 29%	it collapses. The 85% is bought with a wide window, not with punctual pages
... inside 30 minutes			16 / 34 47%	half-way between, as you would expect from the lead distribution

So the answer is: it is a real warning, and it is early rather than late. The failure mode the question was hunting — pages about dead legs — is **not** what is happening: 19 of 29 paged legs were still running at the press and 23 of 29 had a live same-side 1×ATR leg. What is true is that a 15-minute band delivers its median page **26 minutes** before he acts, and only a third of them inside twenty.

That qualifies MONDAY no. 2 (last week’s card) without killing it, and it is exactly why the row's second half — put it on the dashboard, not the phone — matters more than the band: a warning that arrives half an hour early is useful on a screen you are already looking at, and is noise in your pocket. △ Two entries out of 34 are the whole difference between this table and §3.2's; the replay and the journal are two views of one four-day week, not independent confirmations.

3.3 · What is NOT automatable, and exactly how far off it is

§1 and §2 both stop at the same wall, so here is the wall with numbers on it instead of an assertion.

WHAT A RULE WOULD NEED	HELD TODAY	NEEDED	DISTANCE, IN WEEKS AT THE CURRENT RATE
ENTRY rule — presses he made	36 §2's count	30	cleared
ENTRY rule — passes he made	0	30	INFINITE. The button shipped 2026-09-18 and has never been pressed. Nothing else on this box records a look that ended in no trade
ALERT-selection rule — pages he took	5	30	~3.6 weeks at 7/week. This is the binding side here, not the negatives
ALERT-selection rule — pages he refused while looking	15	30	~0.7 weeks at 21/week. The first true negatives this desk has ever held
EXIT rule	119 attempts 86% capture	—	NO ROUTE. 81 of your 92 entries ended on your claim button, the \$400 and \$600 rungs have never fired once, and the standing rule forbids a 120th calibration. I did not run one

Read the two middle rows together, because they invert the usual complaint. For the question "would he have bought here", the missing side is the negatives and it is missing absolutely. For the question "is this leg worth telling him about", the negatives are already arriving faster than the positives — 21 a week against 7 — because every page he declines while his phone is awake is a label. The alert-selection problem is **about four weeks from being a properly posed learning problem**, and the press problem is not on a clock at all until the PASS button gets used.

And there is a second wall behind the first, worth knowing before anyone celebrates 30 passes. §2 measured what the box can actually see at the moment you press: position-in-range and VWAP stretch are indistinguishable from a

random minute, and drift, the desk's own direction detector, could not even read at 24 of 36 entries and returned confirmed=false at all 12 where it could. Strip those out and the features that survive are **a running same-side leg** and **ATR percentile**. Two. Thirty passes buys you a boundary drawn in a two-dimensional space — which is worth having, and is not a model of your judgement.

3.4 · The pieces that need no model at all, each scored on this week

Three rules that require nothing to be learned about you. A stopwatch, a calendar, and a loss cap. Each is priced in your units: dollars a day at four lots.

(a) The hold-time stopwatch — and the number \$1 handed to this section

\$1 measured the over-8-hour bucket at **0 winners from 4** and **−\$6,612.44** and said explicitly that what a stopwatch recovers was not computed there. It is computed here: flatten each of those four at exactly T+8h at the tape price. **The delta depends only on exit prices, so the paper engine's fabricated entry fills cancel out of it.**

ENTRY	HELD	ACTUAL	AT +8H	DELTA	WOULD HAVE BEEN
2026-08-21 13:13 LONG	70.8h	−\$2,149.44	29370.00	+\$1,168.00	−\$981.44
2026-09-01 08:24 LONG	9.0h	−\$1,466.00	29192.25	+\$522.00	−\$944.00
2026-09-04 05:29 LONG	8.7h	−\$341.00	29580.50	−\$202.00	−\$543.00 — the stopwatch made this one worse
2026-09-10 05:40 LONG	13.6h	−\$2,656.00	29184.25	+\$232.00	−\$2,424.00
TOTAL over 28 sessions	—	−\$6,612.44	—	+\$1,720.00	+\$61.43 a session traded

So the eight-hour hole is a \$1,720 opportunity, not a \$6,612 one. All four were already deep underwater by hour eight — the stopwatch recovers **26%** of the damage, one of the four gets worse by \$202, and **68% of the total is a single trade** (21 August). That is a much smaller prize than the bucket's headline, and it is the right number to carry forward.

Swept across every width, split at Monday. **Each column divides by the sessions in its own window**, so the 8-hour rung reads +\$74.78 here against the +\$61.43 above — same \$1,720, 23 sessions rather than 28.

CAP	BEFORE THIS WEEK 52 ENTRIES / 23 SESSIONS	THIS WEEK 39 ENTRIES / 5 SESSIONS	READ
2 hours	+\$69.07/sess	+\$2.60/sess	4 entries touched this week
3 hours	+\$152.93/sess	+\$24.80/sess	3 touched
4 hours	+\$117.24/sess	-\$81.60/sess	sign flips on the new data
6 hours	+\$156.15/sess	-\$186.40/sess	sign flips harder
8 hours — the pre-registered rung	+\$74.78/sess	\$0.00/sess	nothing to do. Your longest hold this week was 7h01m

The honest read, both halves. On the 23 sessions the finding came from the sweep is positive at every width from 2h to 8h, and stays positive at 10h and 12h on the whole book — a smooth band, not a spike, which is the shape of something real. On the first five sessions of new data it earned nothing at 8h because you never held that long, and where it did have something to do — 4h and 6h — **it lost money**. Two or three entries either way; that is noise-sized, not a refutation. But it is the first forward week and **it did not confirm**.

One trade settles the question of whether a shorter cap would have rescued the week. Tuesday's 13:39Z long, −\$1,288.00, held 7h01m into the 20:40Z clock: capped at 1h it is −\$1,180, at 2h −\$1,274, at 3h −\$1,294, at 4h −\$1,486, at 6h −**\$1,764. No stopwatch of any width saves it.** It was wrong from the first hour and the tape happened to come back into the clock. Whatever that trade needs, it is not a timer.

(b) The event-day stand-aside

WINDOW	DAYS	ENTRIES	NET	PER SESSION AVOIDED
This week — FOMC Wednesday		1	6	+\$166.60/day across the 5-day week. The only red day of the week
Whole book — 12 Aug CPI, 4 Sep NFP, 11 Sep CPI, 16 Sep FOMC		4	12	+\$54.23 a session across 28; +\$379.62 per event day avoided
Ordinary days		25	80	+\$1,843.55 §1's figure, reproduced

Point estimate first: +\$166.60 a day this week, +\$54.23 a session over the book. Uncertainty second, and it is the size of the finding: **four days**, one of them (12 Aug CPI) positive at +\$224.50, and the split was chosen after the losses were visible. Three of four negative is not a distribution. It also does not need to be a switch — the T-60 and T-15 pages already exist, they name the event and refuse to name a direction, and the Wednesday damage was **one entry at T+52 into a**

slide that ran another 33 minutes, not the day. Keep the page. Do not build the bench.

★ “+\$166.60 A DAY” IS ONE OF THREE ANSWERS, AND THE ROW WAS PRICING ONLY THE UPSIDE OF THE AVOIDANCE

The +\$166.60 is FOMC Wednesday's −\$833.00 spread across five sessions — in a week that made **+\$583.30 a day**. It prices what a whole-day stand-aside saves and never prices what it costs, and Wednesday also held the week's single biggest winner (**+\$819.00 at 13:13**) and 3 of its 6 entries won. Since the row's own argument is that the damage was one entry at T+52 rather than the session, the whole day is the wrong unit — so here are all three units side by side, every one at 4 lots and every one as \$/day against the week's +\$583.30.

WHAT YOU STAND ASIDE FROM	ENTRIES REMOVED THIS WEEK	THEIR NET	\$/DAY AFTER	DELTA	WHAT IT ALSO REMOVES
The whole event DAY — the unit the row is priced in today	6	−\$833.00	+\$749.90	+\$166.60	the +\$819.00 at 13:13 and the +\$450.00 at 19:08 — \$1,429 of winners deleted to avoid \$2,262 of losers
T−60 to T+60 around the release (17:00-19:00Z)	1	−\$1,272.00	+\$837.70	+\$254.40	nothing. It removes the 18:52 entry and keeps both the +\$819.00 and the +\$450.00
Only entries opened AFTER the release	2	−\$822.00	+\$747.70	+\$164.40	the +\$450.00 at 19:08 — the reversal that recovered a third of the damage

Over the whole book (29 sessions traded, 4 event days) the ranking reverses: whole-day **+\$52.36/session**, T±60 **+\$43.86**, after-the-release **+\$20.60**. The reason is the cost side this week hides: across four event days a whole-day stand-aside removes **\$2,267.51 of winners** (+\$819.00 on 16 Sep, +\$614.01 on 11 Sep, +\$224.50 on 12 Aug) to avoid \$3,786.00 of losers, and on **12 August it deletes a winning day outright** — one entry, +\$224.50, and nothing else traded that day.

The difference between the three IS the size of the finding, and the difference is one trade. T±60 wins this week by exactly the 18:52 long and nothing else; over the book it is the same single entry. **So this is not**

evidence for switching the unit — $n = 1$ is not a window study, and picking the window that isolates the one big loser after seeing it is the definition of fitting. What it is evidence for is that “**+\$166.60 a day**” **should never have been quoted as the value of the policy**: it is the value of one unit, on one event, with the winners it also deletes left out of the arithmetic. The row stays a **page**, the page already fires at T-60 and T-15, and no bench of any width is justified on four event days.

(c) The \$250 daily loss limit — the one that is worse than nothing

This is listed in the brief as a piece that is already automated and running. It is running. Here is what it did.

WHAT WAS CHECKED	WHAT WAS FOUND
Can it reach your book?	No. <code>daily_loss_limit.py</code> writes off to six names in <code>data/gate_switches.env</code> — <code>grind_long</code> , <code>capitulation_long</code> , <code>abs_veto_long</code> , <code>exhaustion_short</code> , <code>abs_veto_short</code> , <code>rgv_short</code> . <code>day_rider</code> is not among them and its switch lives in a different file, <code>data/day_rider.env</code> , which this script never opens
Did it fire this week?	Twice — 2026-09-15T20:40:40Z and 2026-09-16T12:24:40Z. Both journal records end "benched": []. All six gates were already =off, so it changed nothing and wrote no line to the trial log
What did it tell you?	Two critical pages reading "All six gates benched; they re-arm at the 22:00Z reopen." That sentence is a hard-coded string; the code never consults its own changed list. It benched none and said it benched six
What number fired it on Tuesday?	<code>pn1_usd</code> : -418.50. Your real session was -\$114.00; the other -\$304.50 is the BADFILL: entry whose price the paper engine fabricated. Its query sums every closed row with no <code>data_quality</code> filter. Without the fabricated fill it would not have fired
And if it COULD stop you — would you want it to?	No. Walked across 29 sessions of your book it trips on 10, and what you opened afterwards on those days nets +\$1,945.50. Stopping you there costs -\$67.09 a session. This week alone: -\$157.00, or -\$31.40 a day

Be precise about what is refuted. The \$250 level was derived on the single-lot automated book — 43 sessions, +\$35/day to +\$87/day — and it is pre-registered and frozen for 60 forward sessions on that book. Nothing here touches that. What is refuted is **extending it to your hand**: on the four sessions where it trips and you keep going, you traded back +\$789.00, +\$1,084.50, +\$157.00 and -\$85.00. Your losing sessions are not the runaway shape the limit was built for — the 20:40Z clock and your own claim button are already doing that job. On this book it is a rule that only ever removes trades, and the trades it removes made money.

(d) The automated desk took NOTHING all week — was that a decision or a fault?

The front table prints “opened by any automated gate | 0 | \$0.00” and moves on, and the premise of this whole report (“it is about you and not about it”) depends on the answer. **It is a decision, it is re-taken every single tick, and the gates are automatically re-armed every night.** Here is the machinery, with timestamps.

GATE	SWITCH NOW	WHO SET IT, AND WHEN	LAST TIME IT TRADED
grind_long	off	the ROUTER, 2026-09-14T15:45:17Z — 3m 05s after the open-hour watcher armed it at 15:42Z on a confirmed break	2026-09-03
abs_veto_long	off	the ROUTER, 2026-09-14T15:45:17Z — the same tick, the same arm	2026-09-11
capitulation_long	off	the ROUTER, 2026-09-17T22:05:23Z — and at 22:00Z on 09-14, 09-15 and 09-16 before that	2026-09-11
exhaustion_short	off	the ROUTER, 2026-09-17T22:05:23Z — same nightly pattern	2026-09-11
abs_veto_short	off	the ROUTER, 2026-09-17T22:05:23Z — same nightly pattern	2026-09-10
rgv_short	off	reactivate_gates.py's HOLD set — never auto-armed, for the 2026-08-09 churn reason recorded in the file's header	2026-07-30

AND THIS IS WHAT HAPPENS EVERY NIGHT	THE RECORD
Three of the six are AUTOMATICALLY RE-ARMED at 22:00Z	gazbot7-gate-reactivate.timer, nightly. Its own log on 09-15, 09-16 and 09-17: “gates auto-reactivated at Paris midnight (1-day off expired): abs_veto_short, capitulation_long, exhaustion_short”. abs_veto_long and grind_long are start-benched by policy (“router arms on evidence”); rgv_short is HELD
... and the router benches them again within minutes	data/router_trial_log.txt: changed: {'capitulation_long': 'off', 'exhaustion_short': 'off', 'abs_veto_short': 'off'} at 22:00:14Z on 09-14, 22:00:13Z on 09-15, 22:00:13Z on 09-16 and 22:05:23Z on 09-17. Four nights, four re-arms, four benches
The only ARM anywhere in the report week	the open-hour watcher, 2026-09-14T15:42Z, grind_long + abs_veto_long on a confirmed break (ER 0.37, ATR 19, bias +96pt UP, new extreme). Benched 3 minutes later
The router's stated reason, unchanged to the last tick	22:35:11Z tonight: “no change, all six gates stay off ... UNTRADEABLE METER is a CONFIRMED STAY-OUT for a forty-fifth consecutive cycle”, and it says in the same sentence that it holds full arm/bench authority and is declining on the tape, not because it is forbidden
Tonight's 22:00Z re-arm	did not happen and correctly said so: “VENUE SHUT (Fri 22:00Z) — already all off, nothing to do”

So: benched deliberately, by the router, every session, and re-armed for it to re-decide every night. Nobody turned the tournament off and forgot. The last fill any gate booked was 2026-09-11 — the Friday before this report week — so the automated desk's silence starts one trading day before the window this report grades, which is why the refocus is not an artefact of a stuck switch. Re-arming is therefore NOT a play: the mechanism already re-arms three gates a night and the router benches them on a stated reading of the tape. ⚠ The one thing worth putting a name on is what that reading rests on: 45 consecutive cycles of an **untradeable-day meter**, which is a hindsight measure of the day just traded and not a forecast of the next one. Whether a gate should be armed is a tournament question and the tournament is not in this report — but “the router keeps saying stay out” is a decision this desk should re-examine the next time the tournament is reviewed, not a fact about the tape.

3.5 · Ranked shortlist

#	WHAT	MEASURED — THE NUMBER, THEN THE UNCERTAINTY	VERDICT
1	Drop leg_watch's first page band from 60 minutes to 15, and route it to the dashboard rather than the phone. A parameter change. No model, no fitting on your judgement, no negatives required	The 25 pages actually sent preceded 24% of your entries against a chance rate of 18% — P = 0.24; the 60-minute band on its own reads 9% against 12%, P = 0.86. Either way, no separation. At 15 minutes the same detector precedes 82% against 56%, P = 0.0009 — and 0.0036 corrected for the four-band sweep it won, which still clears; 10 min reads 0.015 raw and 0.060 corrected, 30 min does not clear either way (0.43) — an edge between 15 and 30, not a plateau. The live journal reaches the same place independently: 29 of 34, 85%. Uncertainty, and there are now three: 34 presses over four days; the 15 was chosen off this curve; and the median page lands 26.5 minutes ahead, with only 10 of 29 inside twenty minutes — restrict coverage to pages inside 20 min and it falls to 29% (§3.2b). The paged leg was still alive at 19 of 29 presses, so this is an early warning, not a stale one. Cost is real and is the reason for the dashboard half: 7 → 32 pages a day. Pre-registered under §3.2: band 15, target ≥56%, kill below 50%, re-scored Friday	PROMISING the best thing in this section

#	WHAT	MEASURED — THE NUMBER, THEN THE UNCERTAINTY	VERDICT
		2026-10-16. Nothing in it may be re-chosen next week	
2	Keep standing aside on event days — as a page, never as a bench. The calendar and the T-60/T-15 pages are already built and running	+\$166.60/day this week if you sit out the whole DAY (FOMC Wednesday, 6 entries, – \$833.00, the worst day of the week) — but that is one of three units and it prices only the upside. Sitting out T±60 around the release is +\$254.40/day this week and keeps the +\$819.00; sitting out only after the release is +\$164.40. Over the book the order reverses: + \$52.36 / +\$43.86 / +\$20.60 a session. Uncertainty: four event days, one of them positive at + \$224.50, split chosen after the losses — and the entire difference between the three units is ONE trade, the 18:52 long. The whole-day version also deletes \$2,267.51 of winners across four event days, and deletes 12 August, a winning day, outright	PROMISING as a page only
3	Turn the paged legs into labels. Every page you decline while your phone is awake is a negative example. This is the only item that unblocks \$1's BLOCKER without waiting on the PASS button	Already held: 5 positives / 15 negatives from 25 pages, with 5 more that nobody saw. Negatives accrue at 21/week, positives at 7/week — so positives are the binding side and the pool reaches 30/30 in about 3.6 weeks. At a 15-minute band it reaches it in roughly one. Uncertainty: \$0/day directly — this buys a dataset, not a trade	ENABLER
4	The eight-hour stopwatch. \$1 ranked it second and deferred the recovery number to here	\$0.00/day this week — 0 of 39 entries reached 8h, longest 7h01m. Lifetime it recovers + \$1,720.00 over 28 sessions = + \$61.43/session, which is 26% of the bucket's –\$6,612.44 headline, not all of it. Uncertainty: four cases, 68% of the total is one trade, one case gets \$202 worse, and where the sweep had work to do this week (4h, 6h) it lost –\$81.60 and – \$186.40 a session. Cheap insurance against a shape that did not recur	PROMISING forward week neutral
5	Do NOT extend the \$250 daily loss limit to your book. It stays pre-registered and frozen on the automated single-lot book where its evidence came from	–\$67.09 a session over 29 sessions — it trips on 10 and the trades you opened afterwards net +\$1,945.50. This week: –\$31.40/day. Uncertainty: one-directional, because the rule can only ever remove trades — and on 6 of the 10 trips you had already	REFUTED for this book only

#	WHAT	MEASURED — THE NUMBER, THEN THE UNCERTAINTY	VERDICT
		stopped, so the entire cost sits in 4 sessions	
6	Defect — the loss limit pages a claim its own code contradicts. The notify string says "All six gates benched"; the journal says "benched": []. Drive the sentence from changed	Fired twice this week, benched zero both times, paged you both times. Every gate it can reach was already =off, so the guard is structurally incapable of doing the thing it announces	FIX — DEFECT
7	Defect — the loss limit counts fabricated fills. session_pnl() sums every closed row with no data_quality filter	Tuesday's trip read −\$418.50: −\$114.00 real, −\$304.50 fabricated by the paper engine. Below the limit on the real number. A guard whose trigger can be moved by a price that never printed	FIX — DEFECT
8	Defect — a hand-run wrote a live PAGED line. Give event_alert.py a --dry-run that logs nowhere, or stamp manual runs	2026-09-18T07:04:03Z PAGED T-15 FOMC, off the *:0/5:30 cadence, two days after the only FOMC in the calendar, with the timer's own run at 07:05:30 logging "nothing in a band". It cost this section one false LED and would have inflated the alerting score by 25%	FIX — DEFECT
9	Put the PASS count where he already looks. Everything in this report is blocked on a button, and nothing in it measures whether the button gets pressed — the only verification offered for MONDAY no. 1 (last week's card) was a one-line python snippet a human has to remember to run	Passes recorded today: 0. At an observed rate of zero a week, 30 never lands — there is no date to quote, and that is the honest form of the answer. From this cycle the count is generated into the front box of every Friday report from operator_reads.jsonl at render time, so a zero next Friday is a visible fact rather than a thing nobody checked. What is still missing is the dashboard tile — one number, beside the metric strip he already reads. Uncertainty: none; it is a counter. This is the same distinction this report draws about last week's 58 ungraded plays: "we checked and it was zero" and "nobody checked" are opposite instructions	FIX — INSTRUMENT unblocks nothing, measures everything
10	The automated desk is benched by DECISION, nightly — not by a stuck switch. Recorded here because the premise of this whole report depends on it and Rev 1 asserted it without evidence	All six gates =off. Three of them (abs_veto_short, capitulation_long, exhaustion_short) are auto-re-armed at 22:00Z every night by gazbot7-gate-reactivate.timer and re-benched by the router within minutes — 22:00:14Z 09-14, 22:00:13Z 09-15, 22:00:13Z 09-16, 22:05:23Z 09-17. The only arm in the week (open-hour watcher, 09-14 15:42Z, grind_long + abs_veto_long) was benched 3m05s later. Last fill by any gate: 2026-09-11. Uncertainty:	ANSWERED no action


#	WHAT	MEASURED — THE NUMBER, THEN THE UNCERTAINTY	VERDICT
		the router's stated reason is 45 consecutive cycles of an untradeable-day meter, which is hindsight about the day just traded, not a forecast — a tournament question, and the tournament is not in this report	

So — can it be automated? Telling you when to look: yes, and the machinery is already built and already running with one number set **wrong**. That is worth doing on Monday and it costs a config change. **Deciding what to take: no, and not because the answer is no — because the question has never been asked in a form that has two sides**. You have 36 yeses and zero noes, and the fastest route to the second kind is not the PASS button (which needs you to remember) but the 21 pages a week you are already declining with your phone in your hand. **Deciding when to get out: no route at all**, and 119 calibrations is enough evidence to stop looking for one. The stopwatch and the calendar are worth keeping for what they are — small, model-free, and honest about their four data points. And the one guard everybody assumed was protecting you is watching a book it cannot touch, announcing actions it did not take, on a number a fabricated fill can move.


4 THE RUN CHARTS — every fill of the week, marked on the day's own price path

Five sessions, one chart each, **00:00-21:00 UTC** — wide enough that not one of your 61 legs falls off the edge, including the 00:25 and 02:37 overnight entries. **A triangle is an entry** (apex up = long, apex down = short), **a circle is an exit**, the line between them is the hold, and green/red is the money. The colour pair is validated for deuteranopia and colour is never the only carrier: shape gives direction and every labelled exit carries a signed dollar figure. Hover any marker for the full fill.


△ **These are LEGS, not entries.** A scale-out books one marker per lot leaving, so Monday's four decisions appear as nine markers. Count triangles that share an entry price as one decision — or read Section 1's tables, which count decisions throughout. **Two charts follow the five sessions:** the FOMC hour up close on MNQ, and the same two hours on **gold**, which carries no fills because neither you nor the desk trades it — it is there because FOMC moved MGC further than it moved the Nasdaq.




The worst single decision of the week is on this chart. The 12:46Z SHORT at 30,765.94 — four lots, closed by hand — cost -\$1,838.50, more than the entire week's net and more than the -\$1,288.00 CLOCK_FLAT that last week's report called its own worst decision. Look at what the path does after the triangle: it goes straight up and keeps going. Nothing closed it but you, and there was no stop to close it.



34 entries — the busiest session this book has ever had — for -\$1,533.00, and the first two PASS presses in the desk's history are on it (violet diamonds). Two things to look for. One: the presses cluster, and the losses cluster with them. Two: the 09:43Z entry that cost -\$842.50 is NOT drawn — it is tagged assistant_test_claim_NOT_operator_decision, a real fill that was never your decision, and it is excluded from every total in this report.




The Sunday-night reopen through Monday. 8 entries, +\$1,190.00 — the best session of the week and the quietest. Every green triangle here reached a profit rung; that is the whole difference between this session and Tuesday's.



14 entries for +\$404.50, and the two biggest winners of the week are both on it (+\$927.50 and +\$644.00, both SHORT, both into the 13:30-14:00Z

open). It also holds a -\$902.00 short at 17:04Z. Same side, same day, opposite outcome — this is the chart that shows the book is five draws rather than a strategy.



35 entries for +\$975.00 and the third PASS press (violet diamond, 18:49Z). The session closed with four lots still open at 20:33Z — the carry check at the top of this report is the authority on what happened to them, not this chart.



One entry, four lots, one leg — it never scaled out and it never reached a rung. The triangle is the 12:46:17Z fill at 30,765.94; the circle is your own hand closing it. This is what -\$1,838.50 looks like as a shape: a short taken into a move that was already going the other way, held while it kept going.

The report's ranking of the week is built on entries like this one, not on the win rate.



These two diamonds are worth more to this desk than any winning trade on the other charts. Every threshold, band and description in this report was fitted on presses only — moments you took. A pass is a moment you looked and declined, and until 09-24 the record contained none. Δ 3 is not 30: measured against the 104 presses of the same week these three sit squarely INSIDE the press distribution on every feature (median 77% of day range vs 73%, +4.4 ATR from VWAP vs +4.6). They separate nothing yet. What they prove is that the pipeline works end to end, so the week the count reaches 30 the answer is one query away.